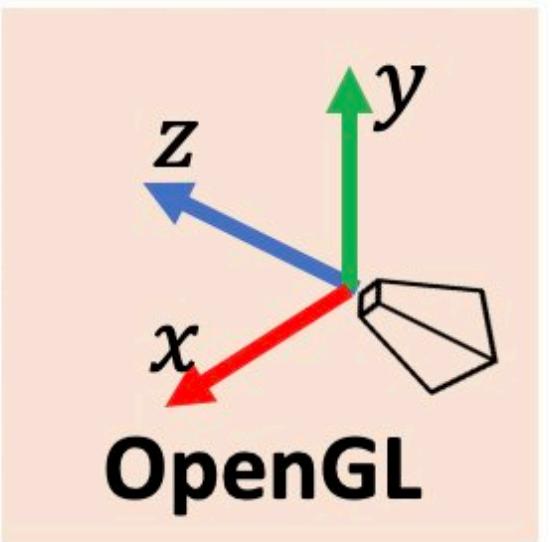


# SCENE REPRESENTATIONS

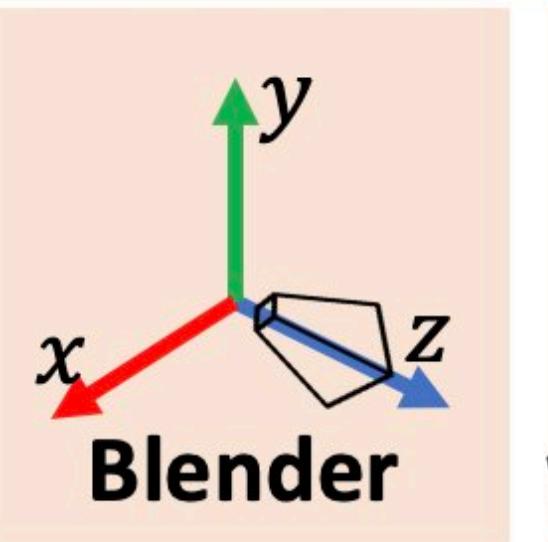
Select all squares with the correct  
**3D coordinate systems**  
If there are none, click skip.

...  
 International Conference on 3D Vision  
@3DVconf

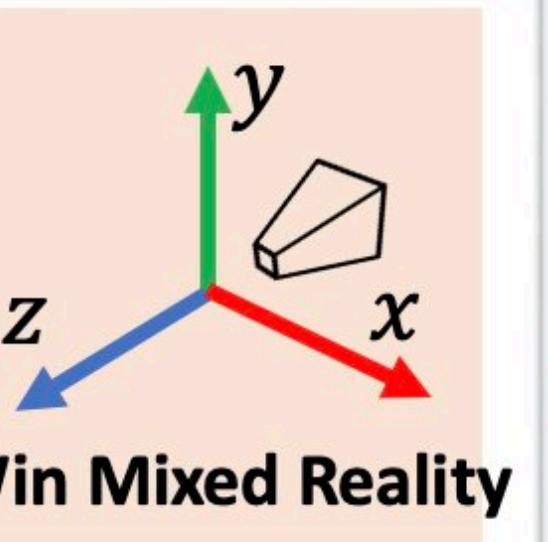
3DV Trivia Time!  
Can you solve the 3DV CAPTCHAs? 🤔



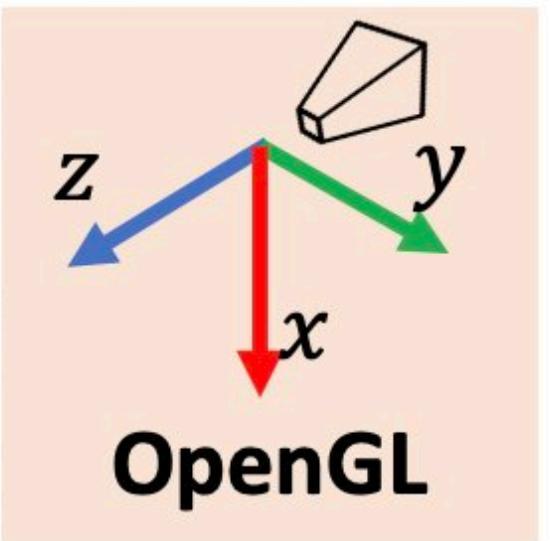
**OpenGL**



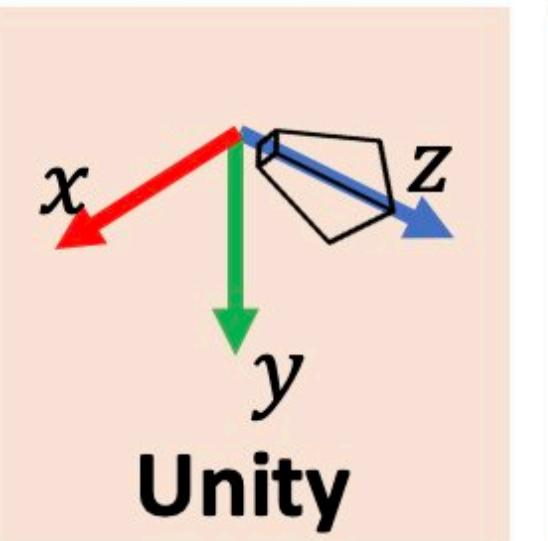
**Blender**



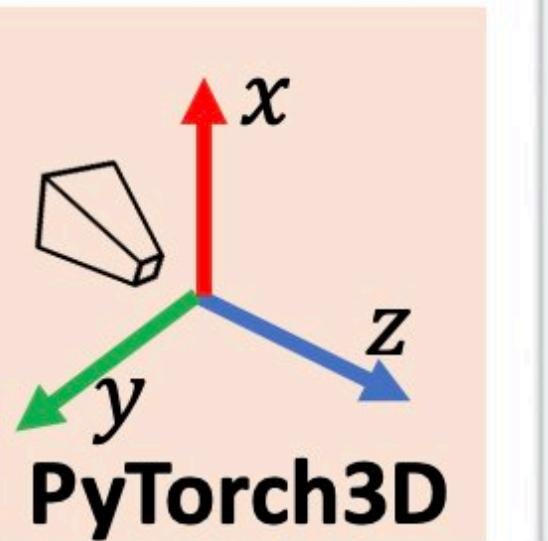
**Win Mixed Reality**



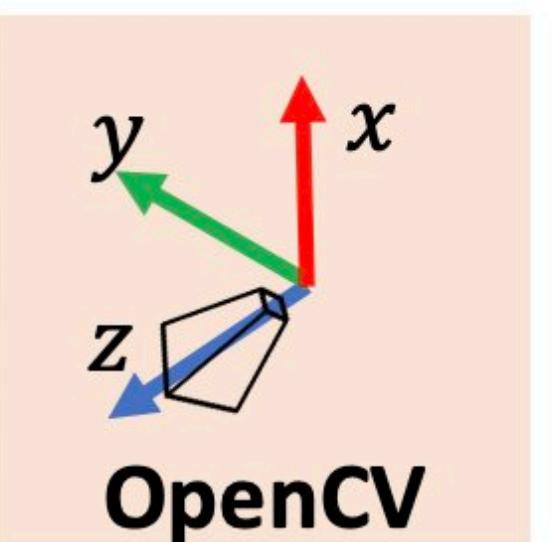
**OpenGL**



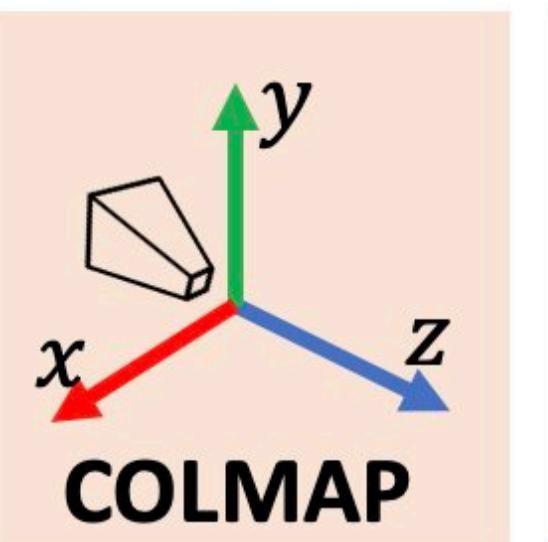
**Unity**



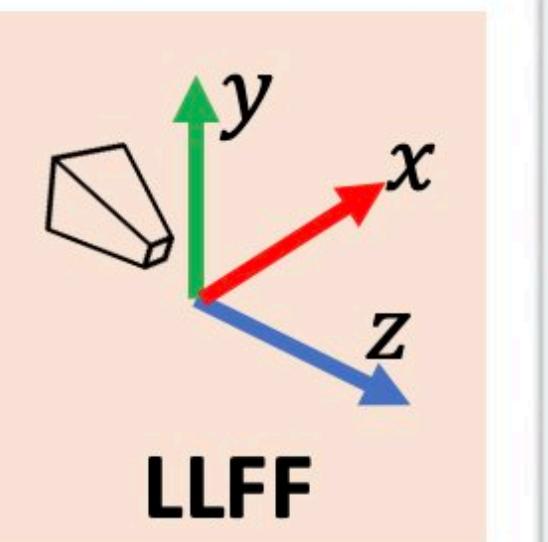
**PyTorch3D**



**OpenCV**



**COLMAP**



**LLFF**

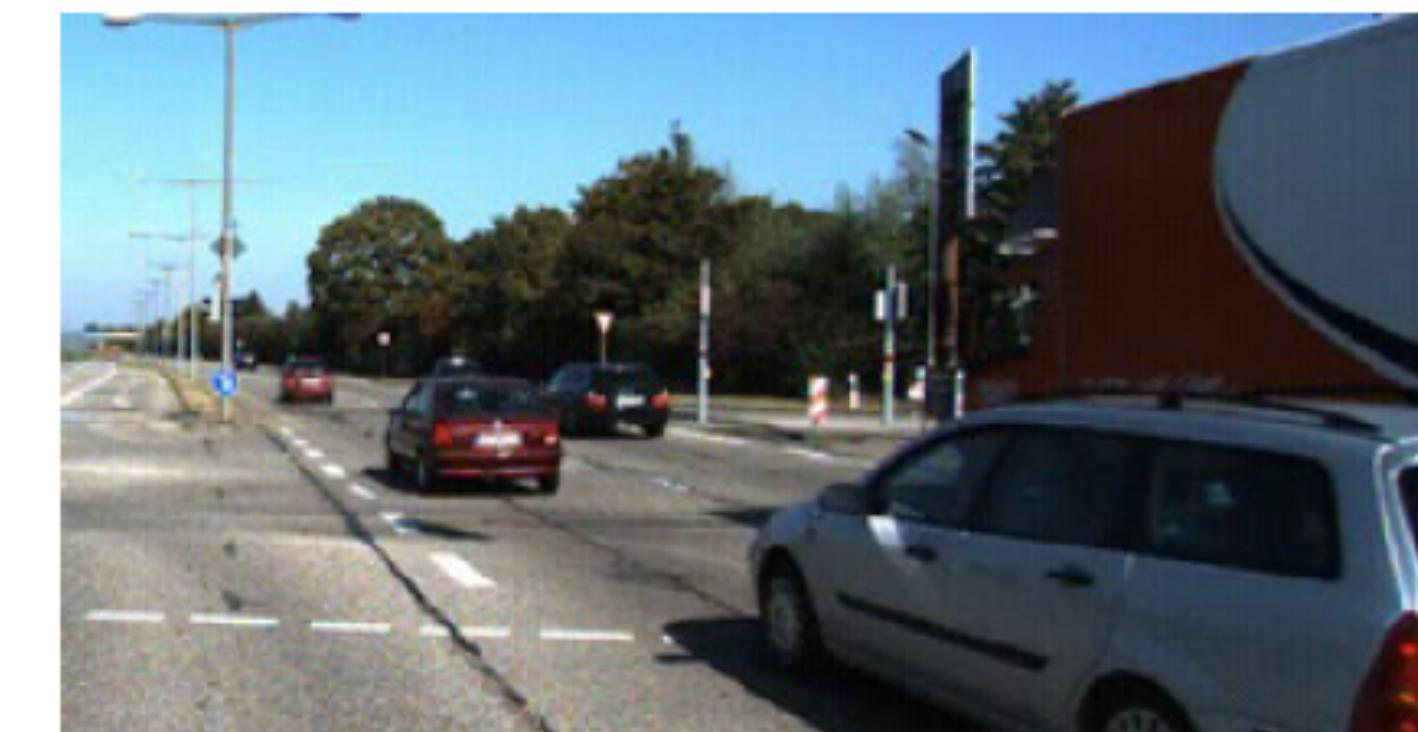


**SKIP**

# Unsupervised Depth Prediction!



Frame at time  $t_1$



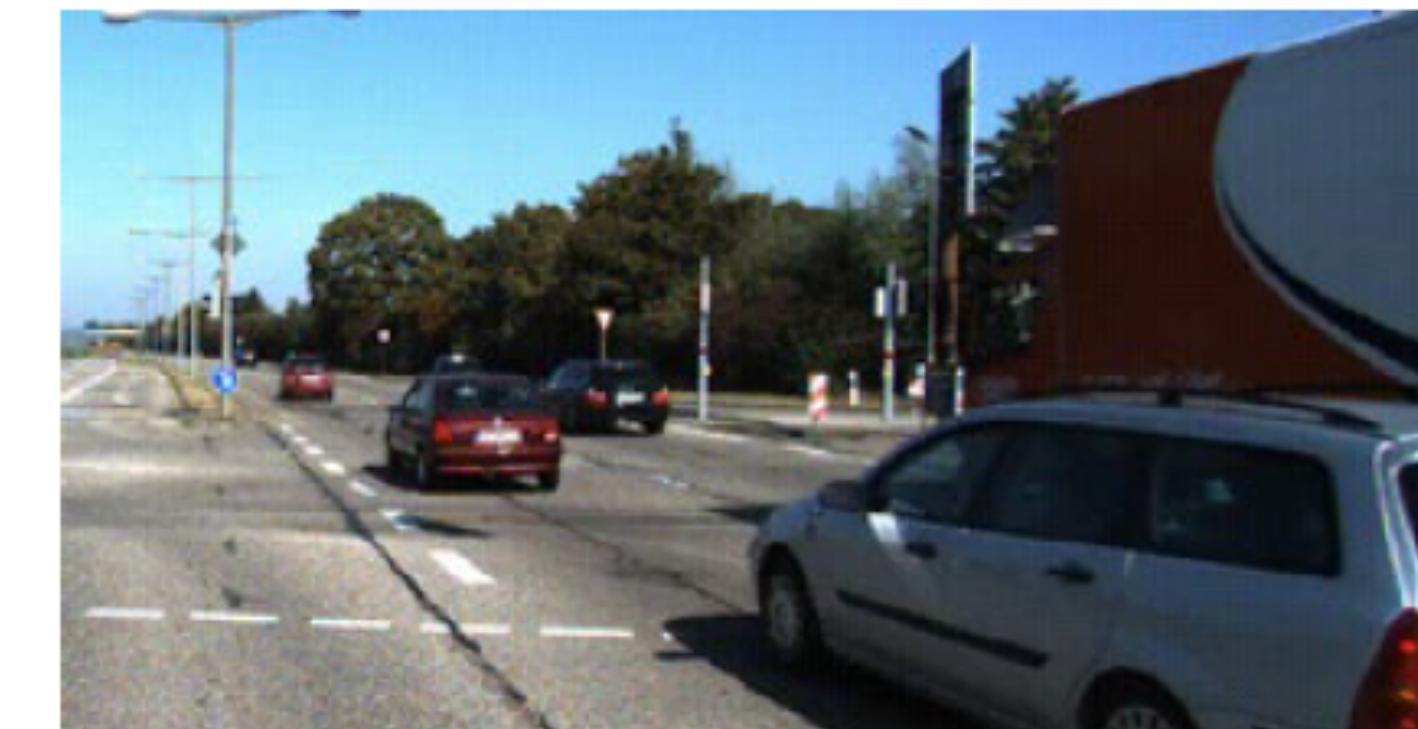
Frame at time  $t_2$

*Goal: Learn Depth and Ego-Motion (relative camera pose) just from video!*

# Unsupervised Depth Prediction!



Frame at time  $t_1$

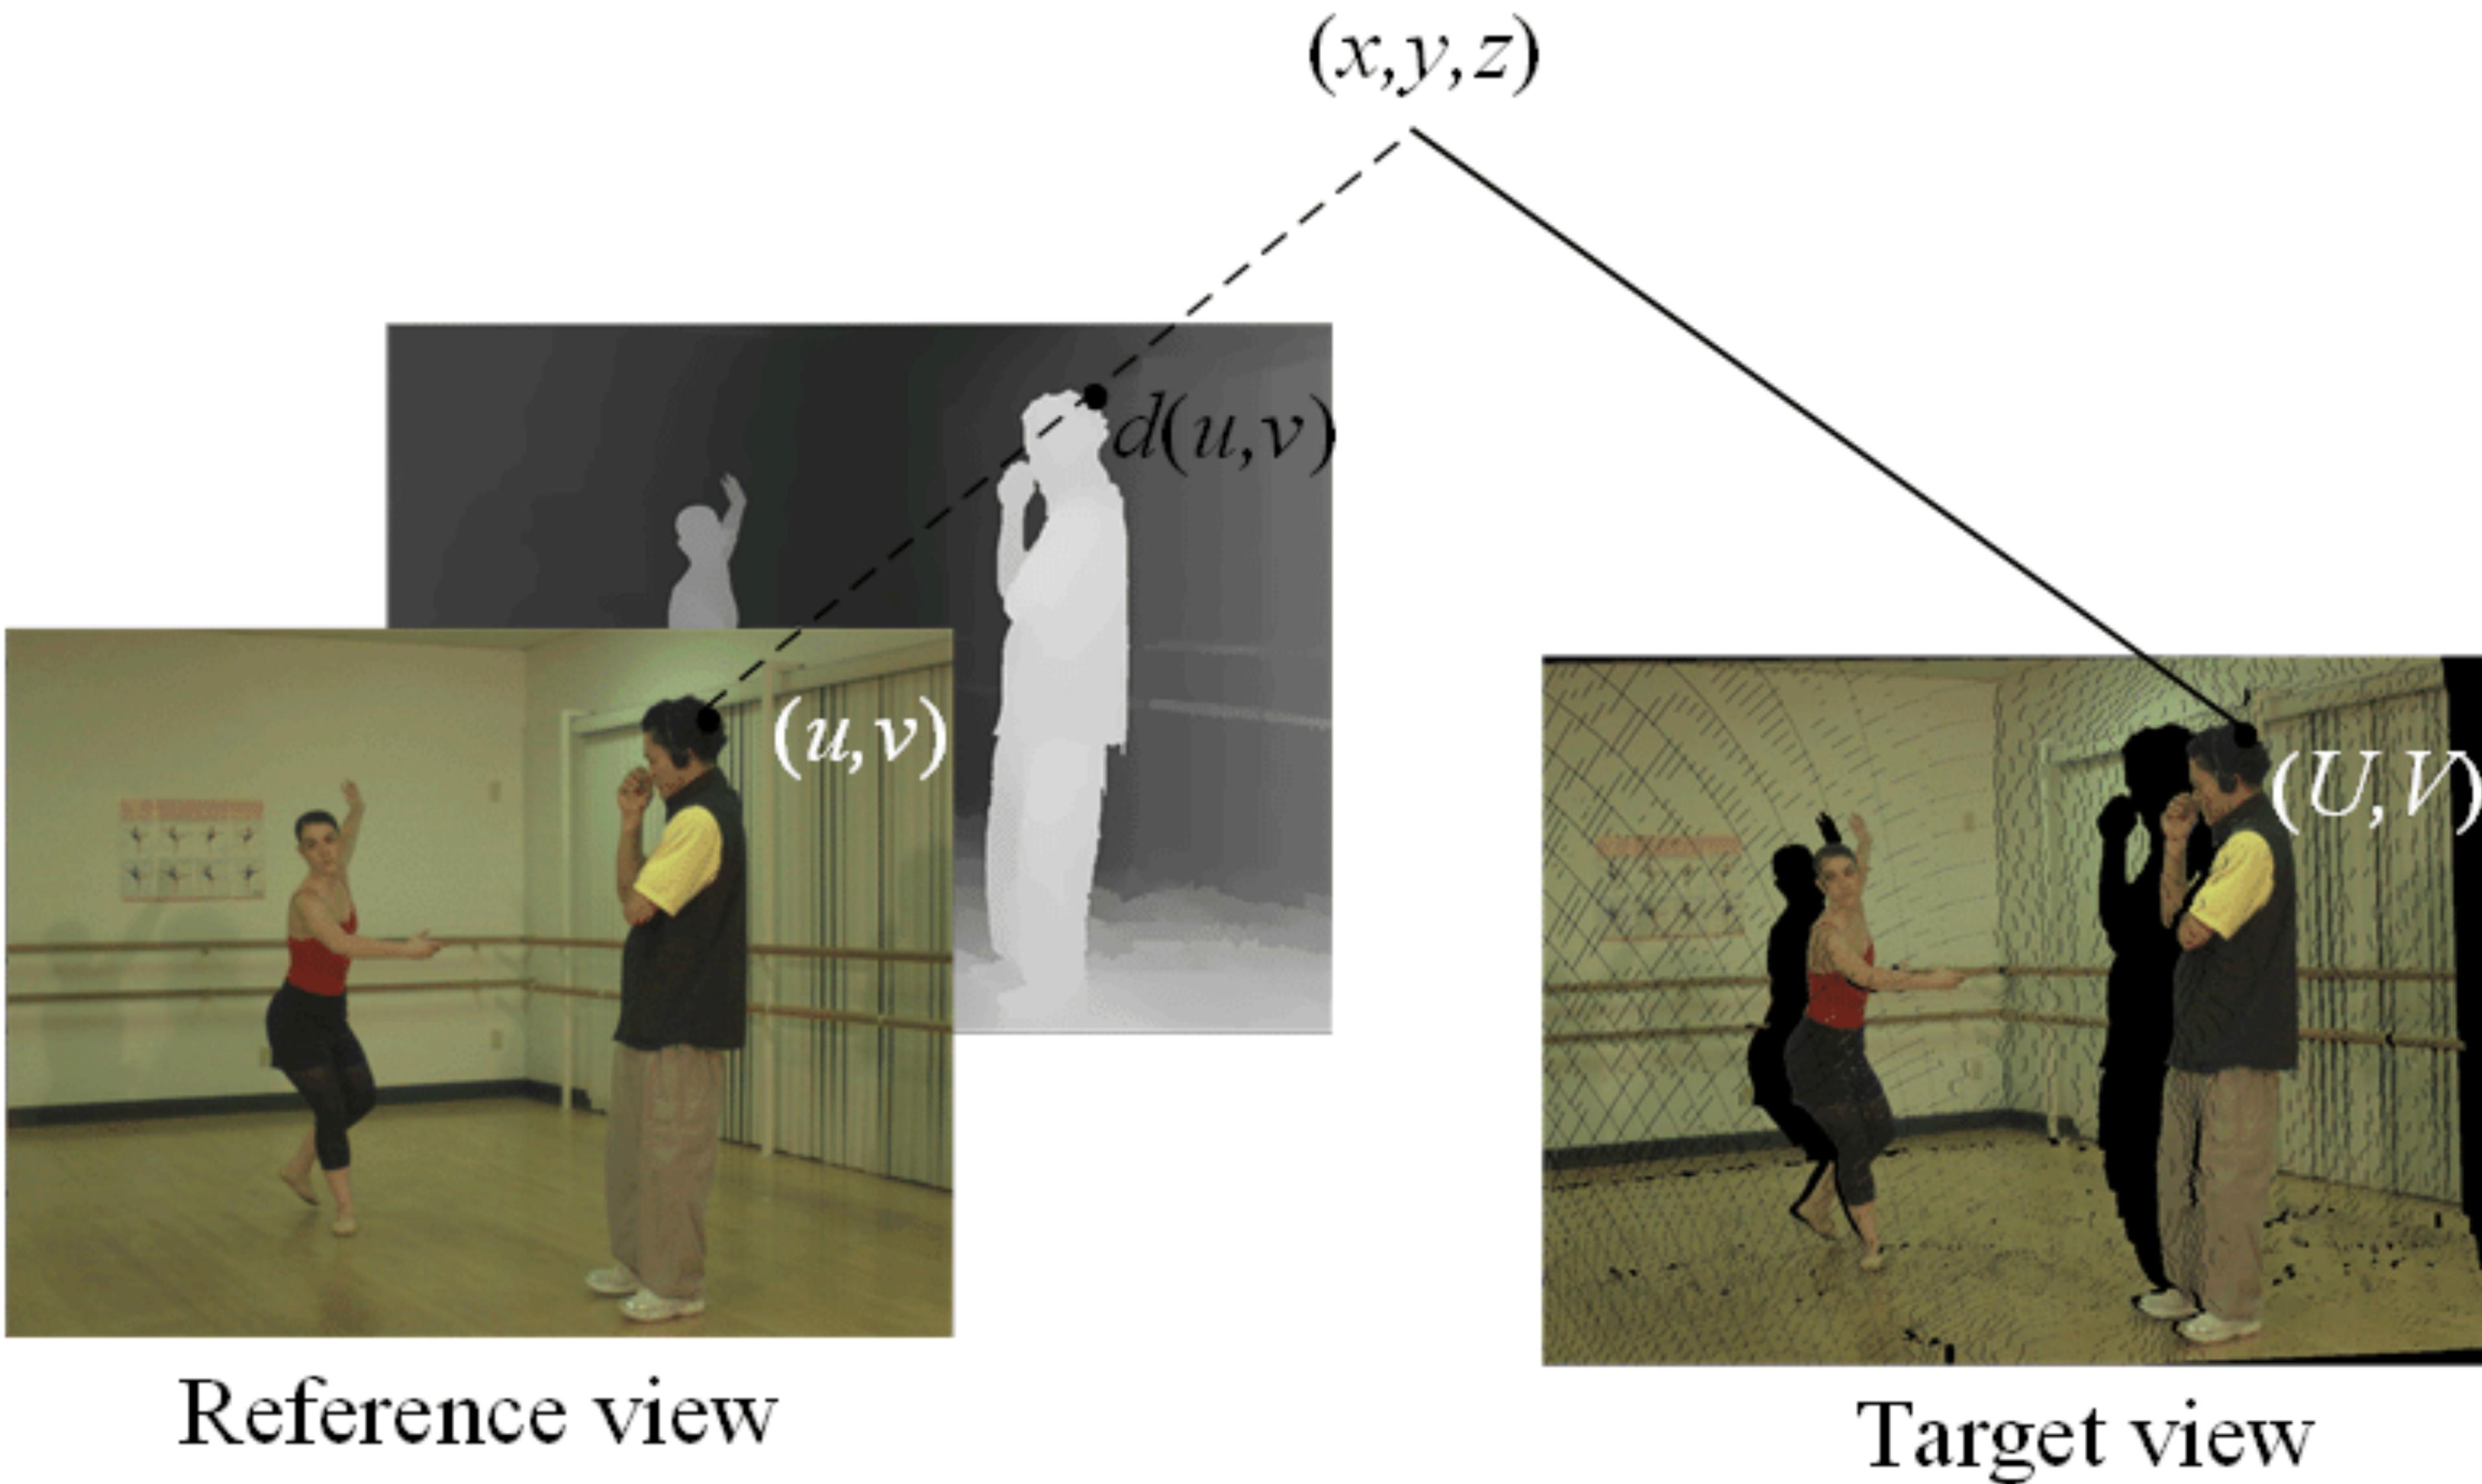


Frame at time  $t_2$

*Goal: Learn Depth and Ego-Motion (relative camera pose) just from video!*

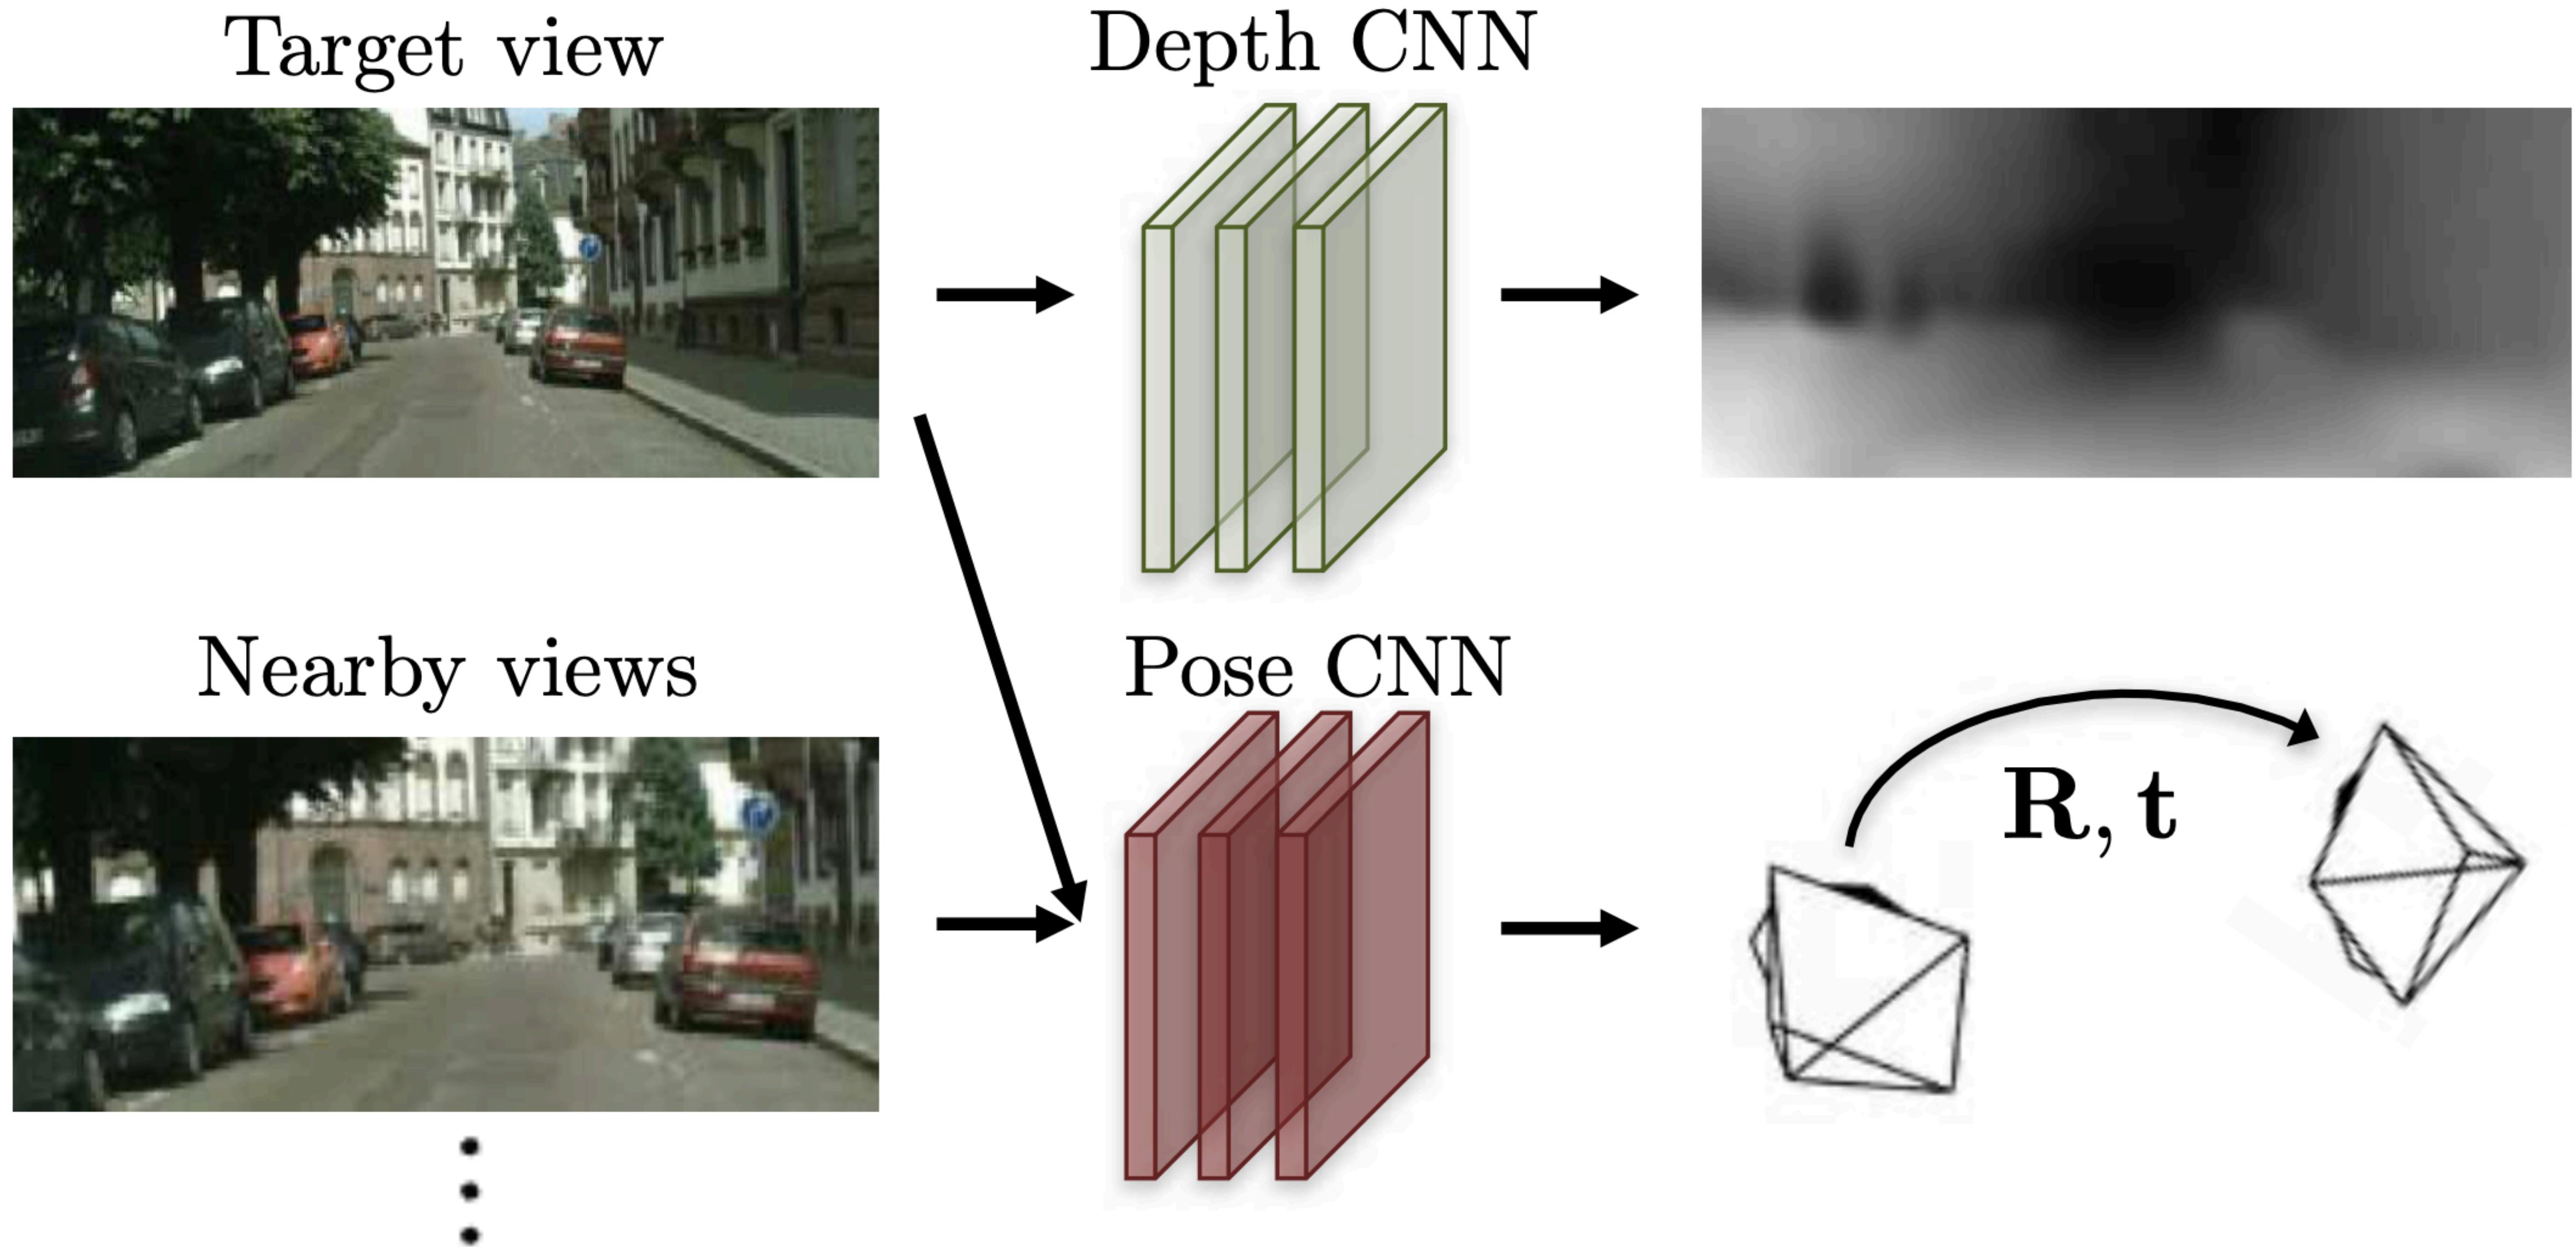
*How?*

# View Synthesis with Depth Maps (Depth Warping)



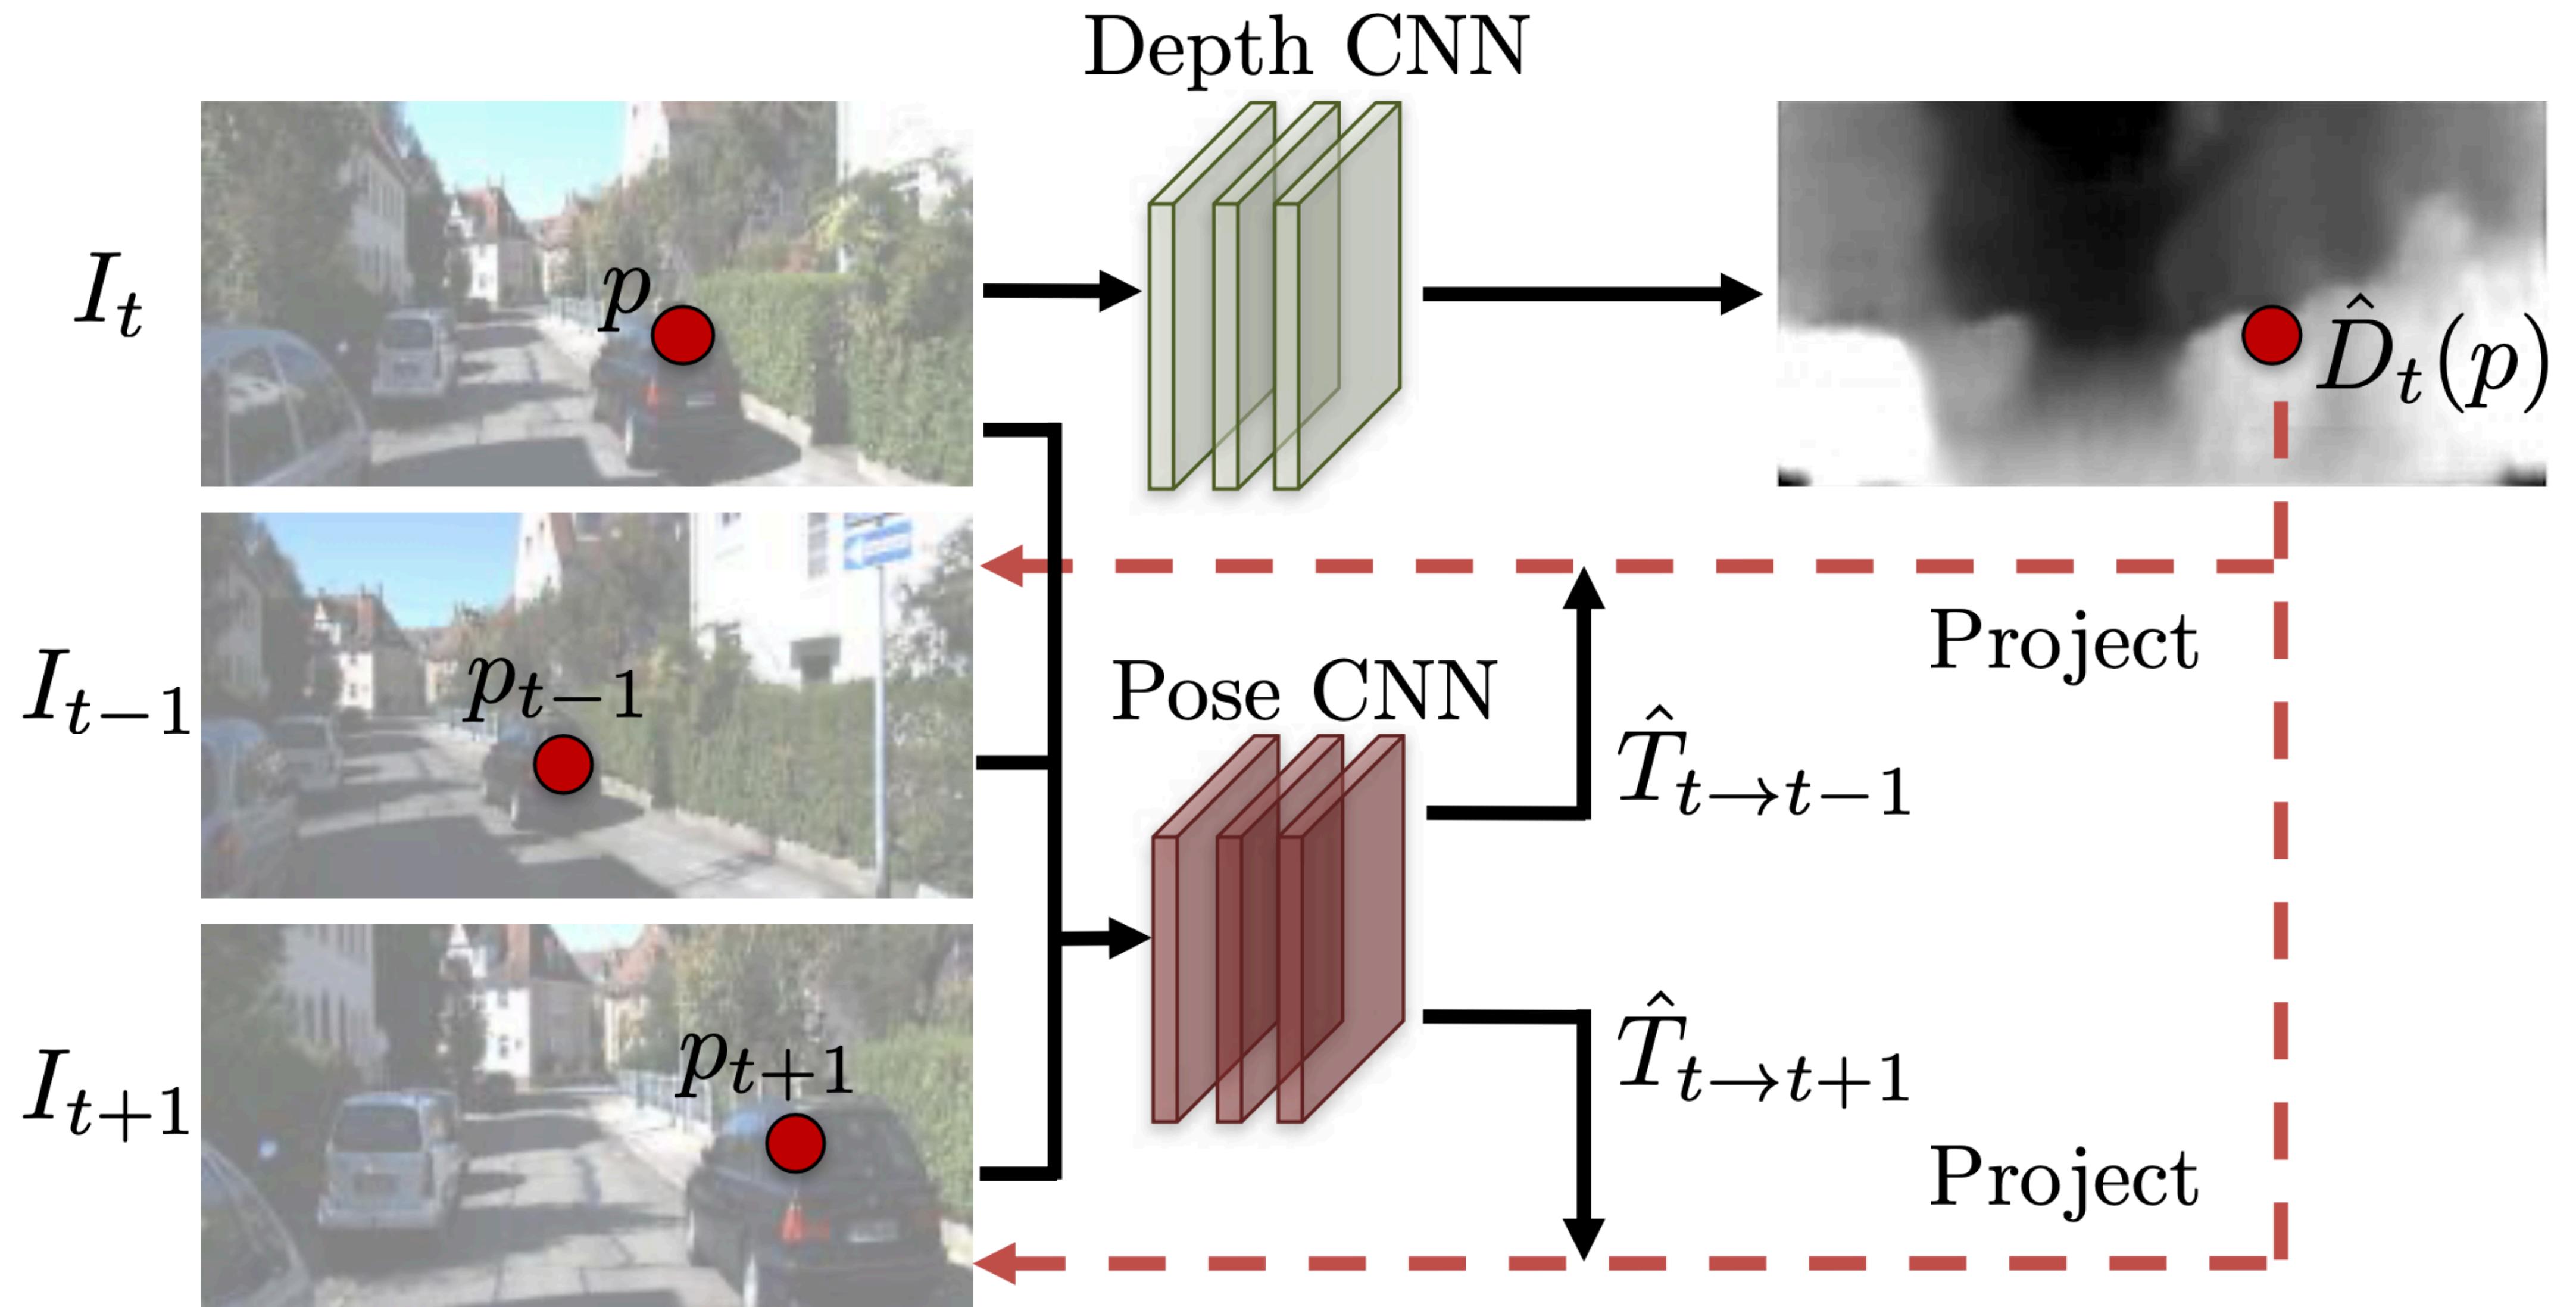
# Unsupervised Depth and Ego-Motion from Video

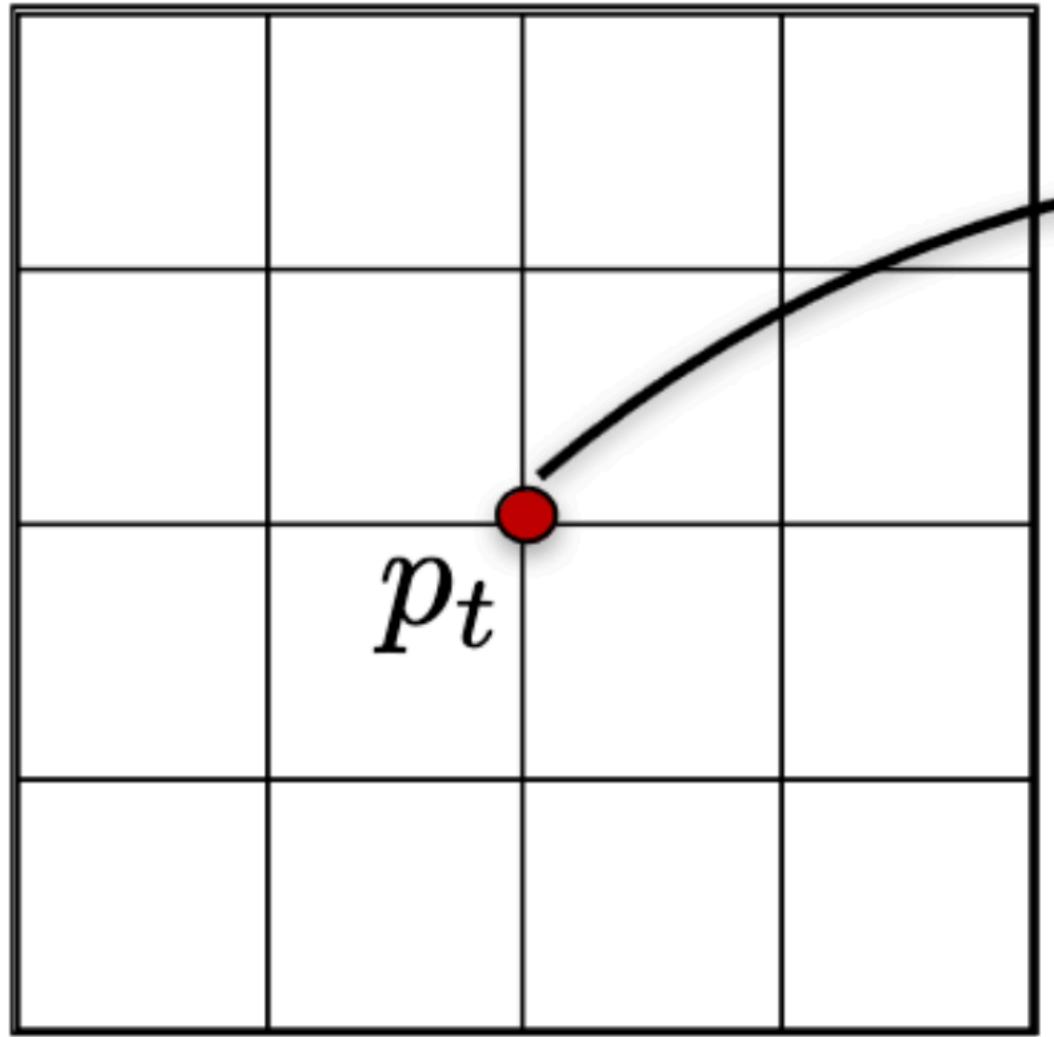
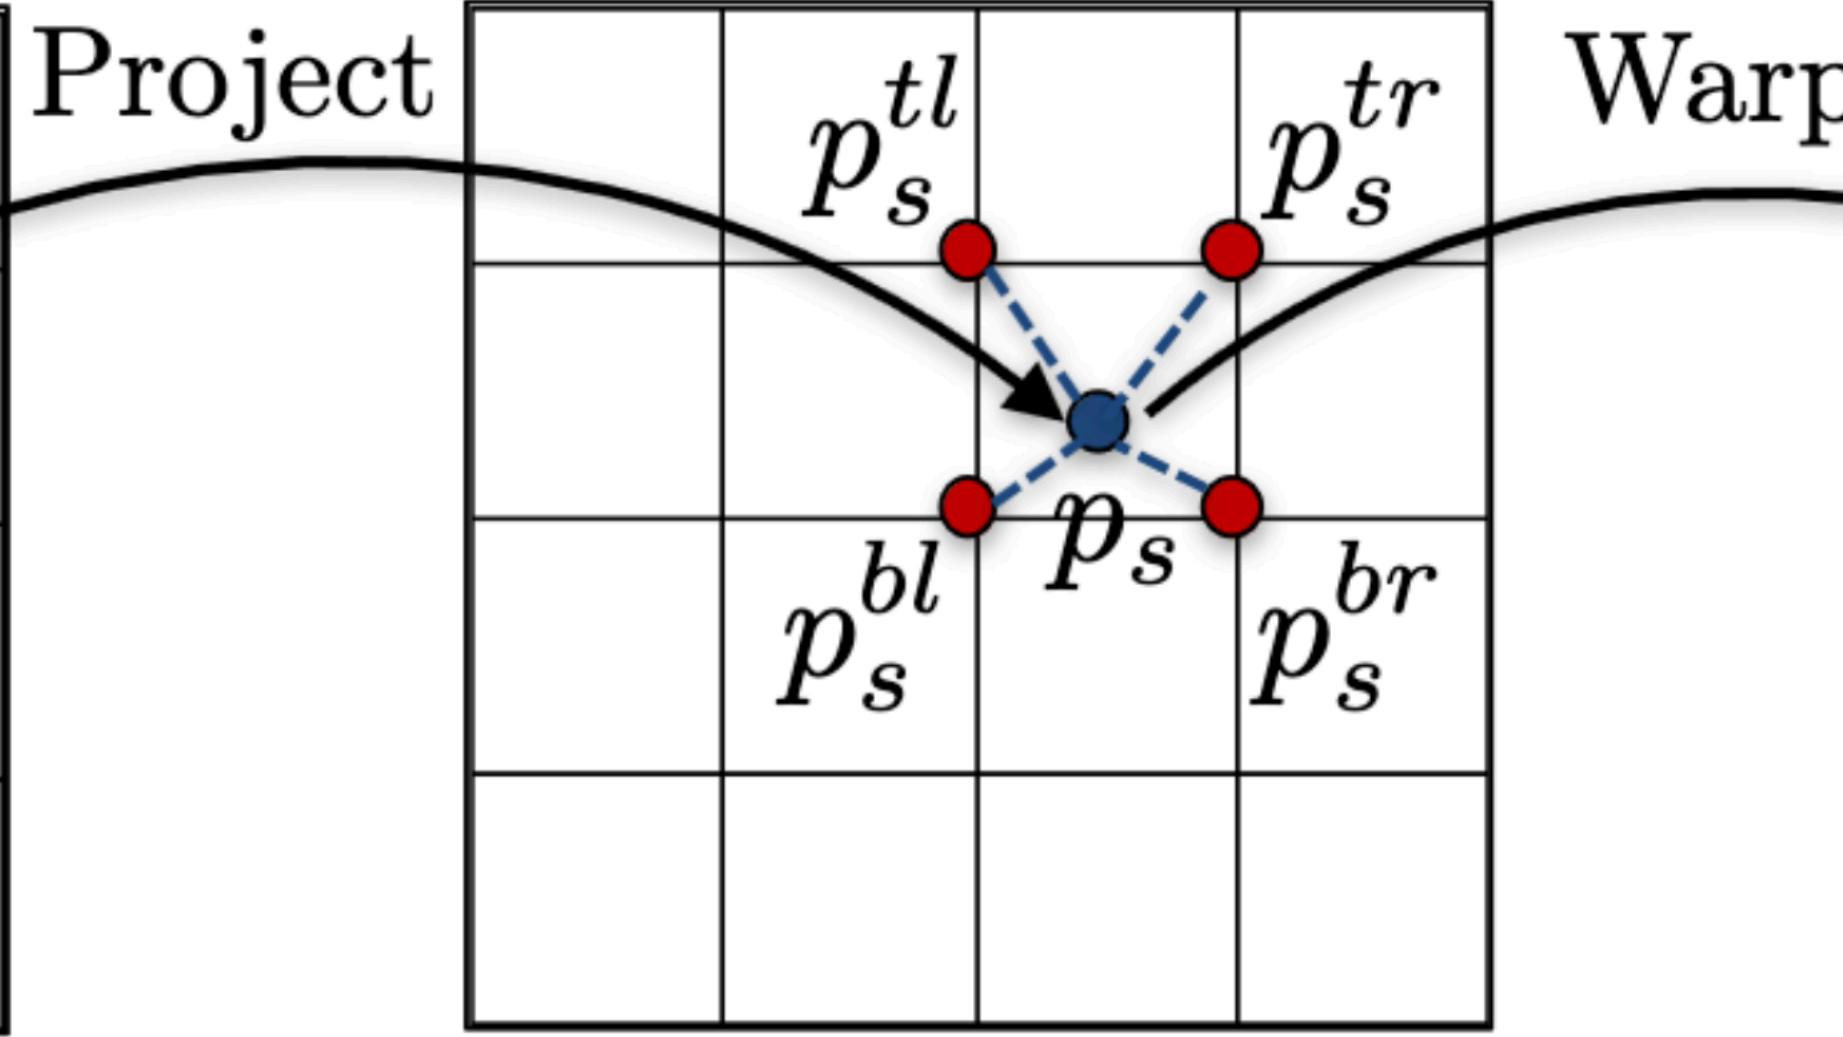
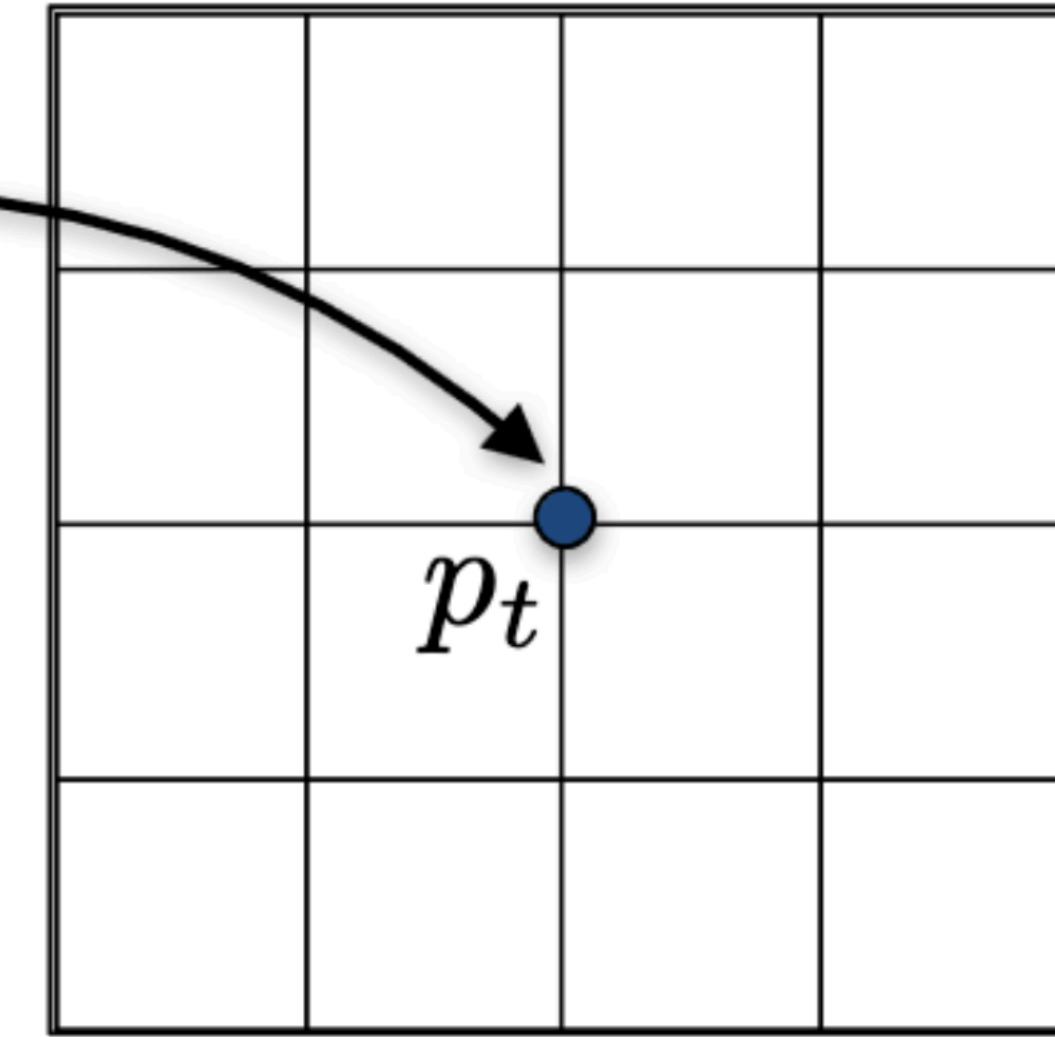
(Zhou et al. 2017)



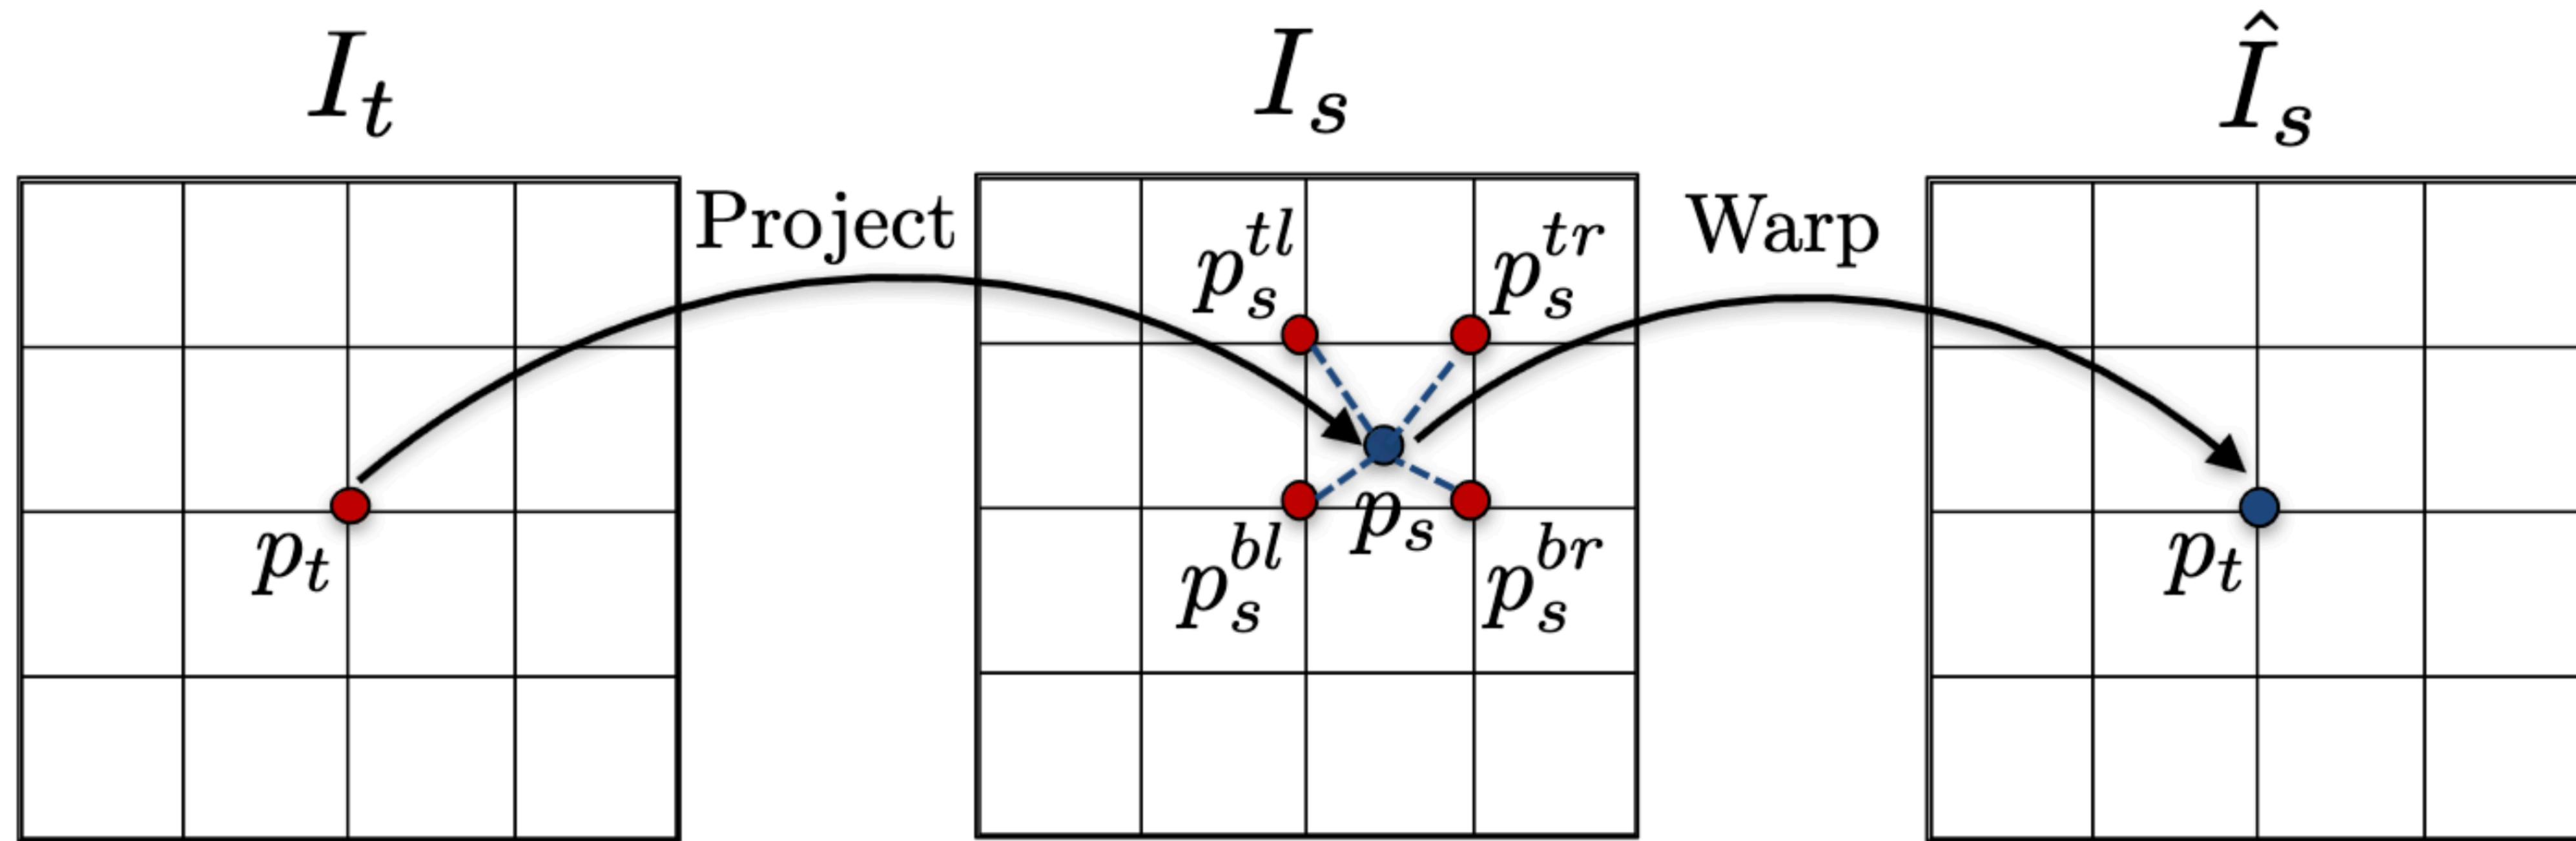
# Unsupervised Depth and Ego-Motion from Video

(Zhou et al. 2017)



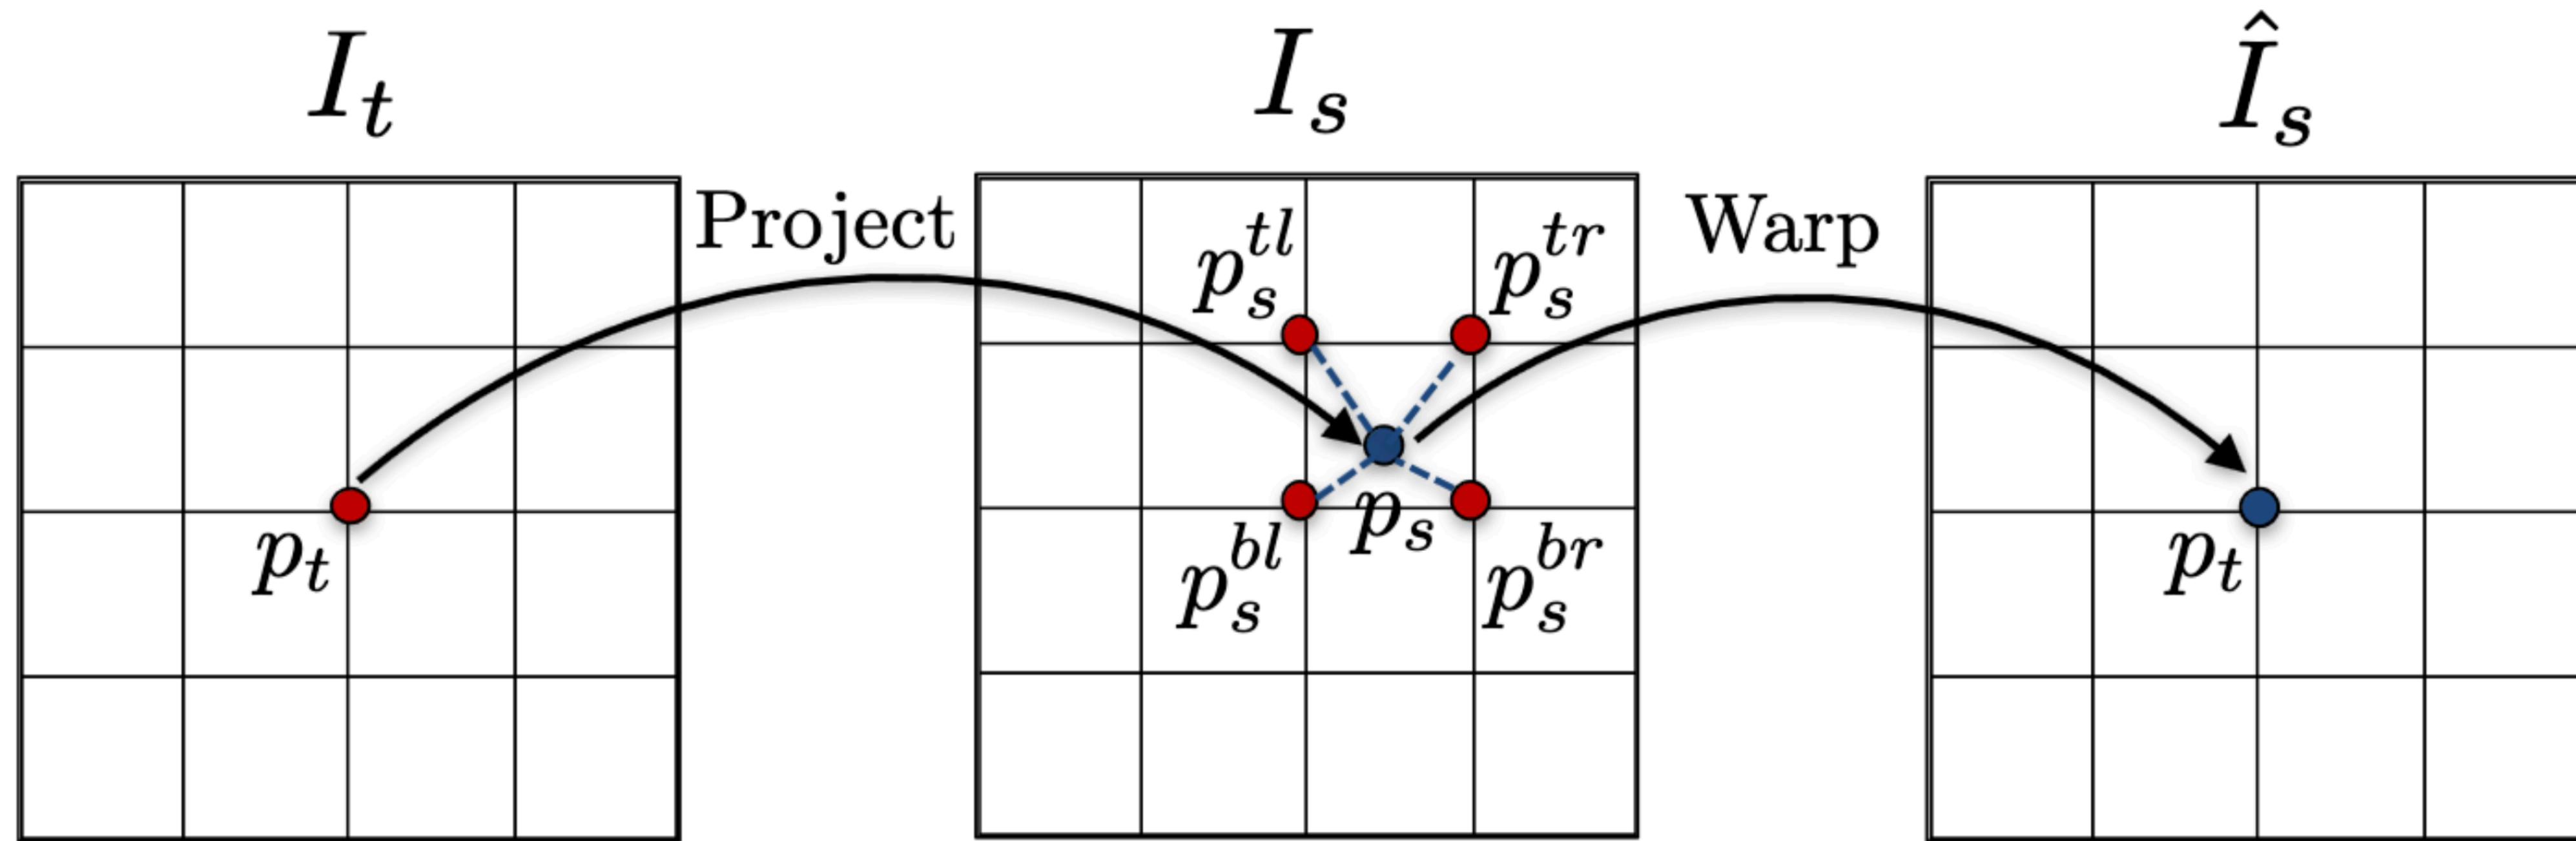
$I_t$  $I_s$  $\hat{I}_s$ 

$$p_s \sim K \hat{T}_{t \rightarrow s} \hat{D}_t(p_t) K^{-1} p_t$$



$$p_s \sim K \hat{T}_{t \rightarrow s} \hat{D}_t(p_t) K^{-1} p_t$$

$$\hat{I}_s(p_t) = I_s(p_s) = \sum_{i \in \{t,b\}, j \in \{l,r\}} w^{ij} I_s(p_s^{ij})$$

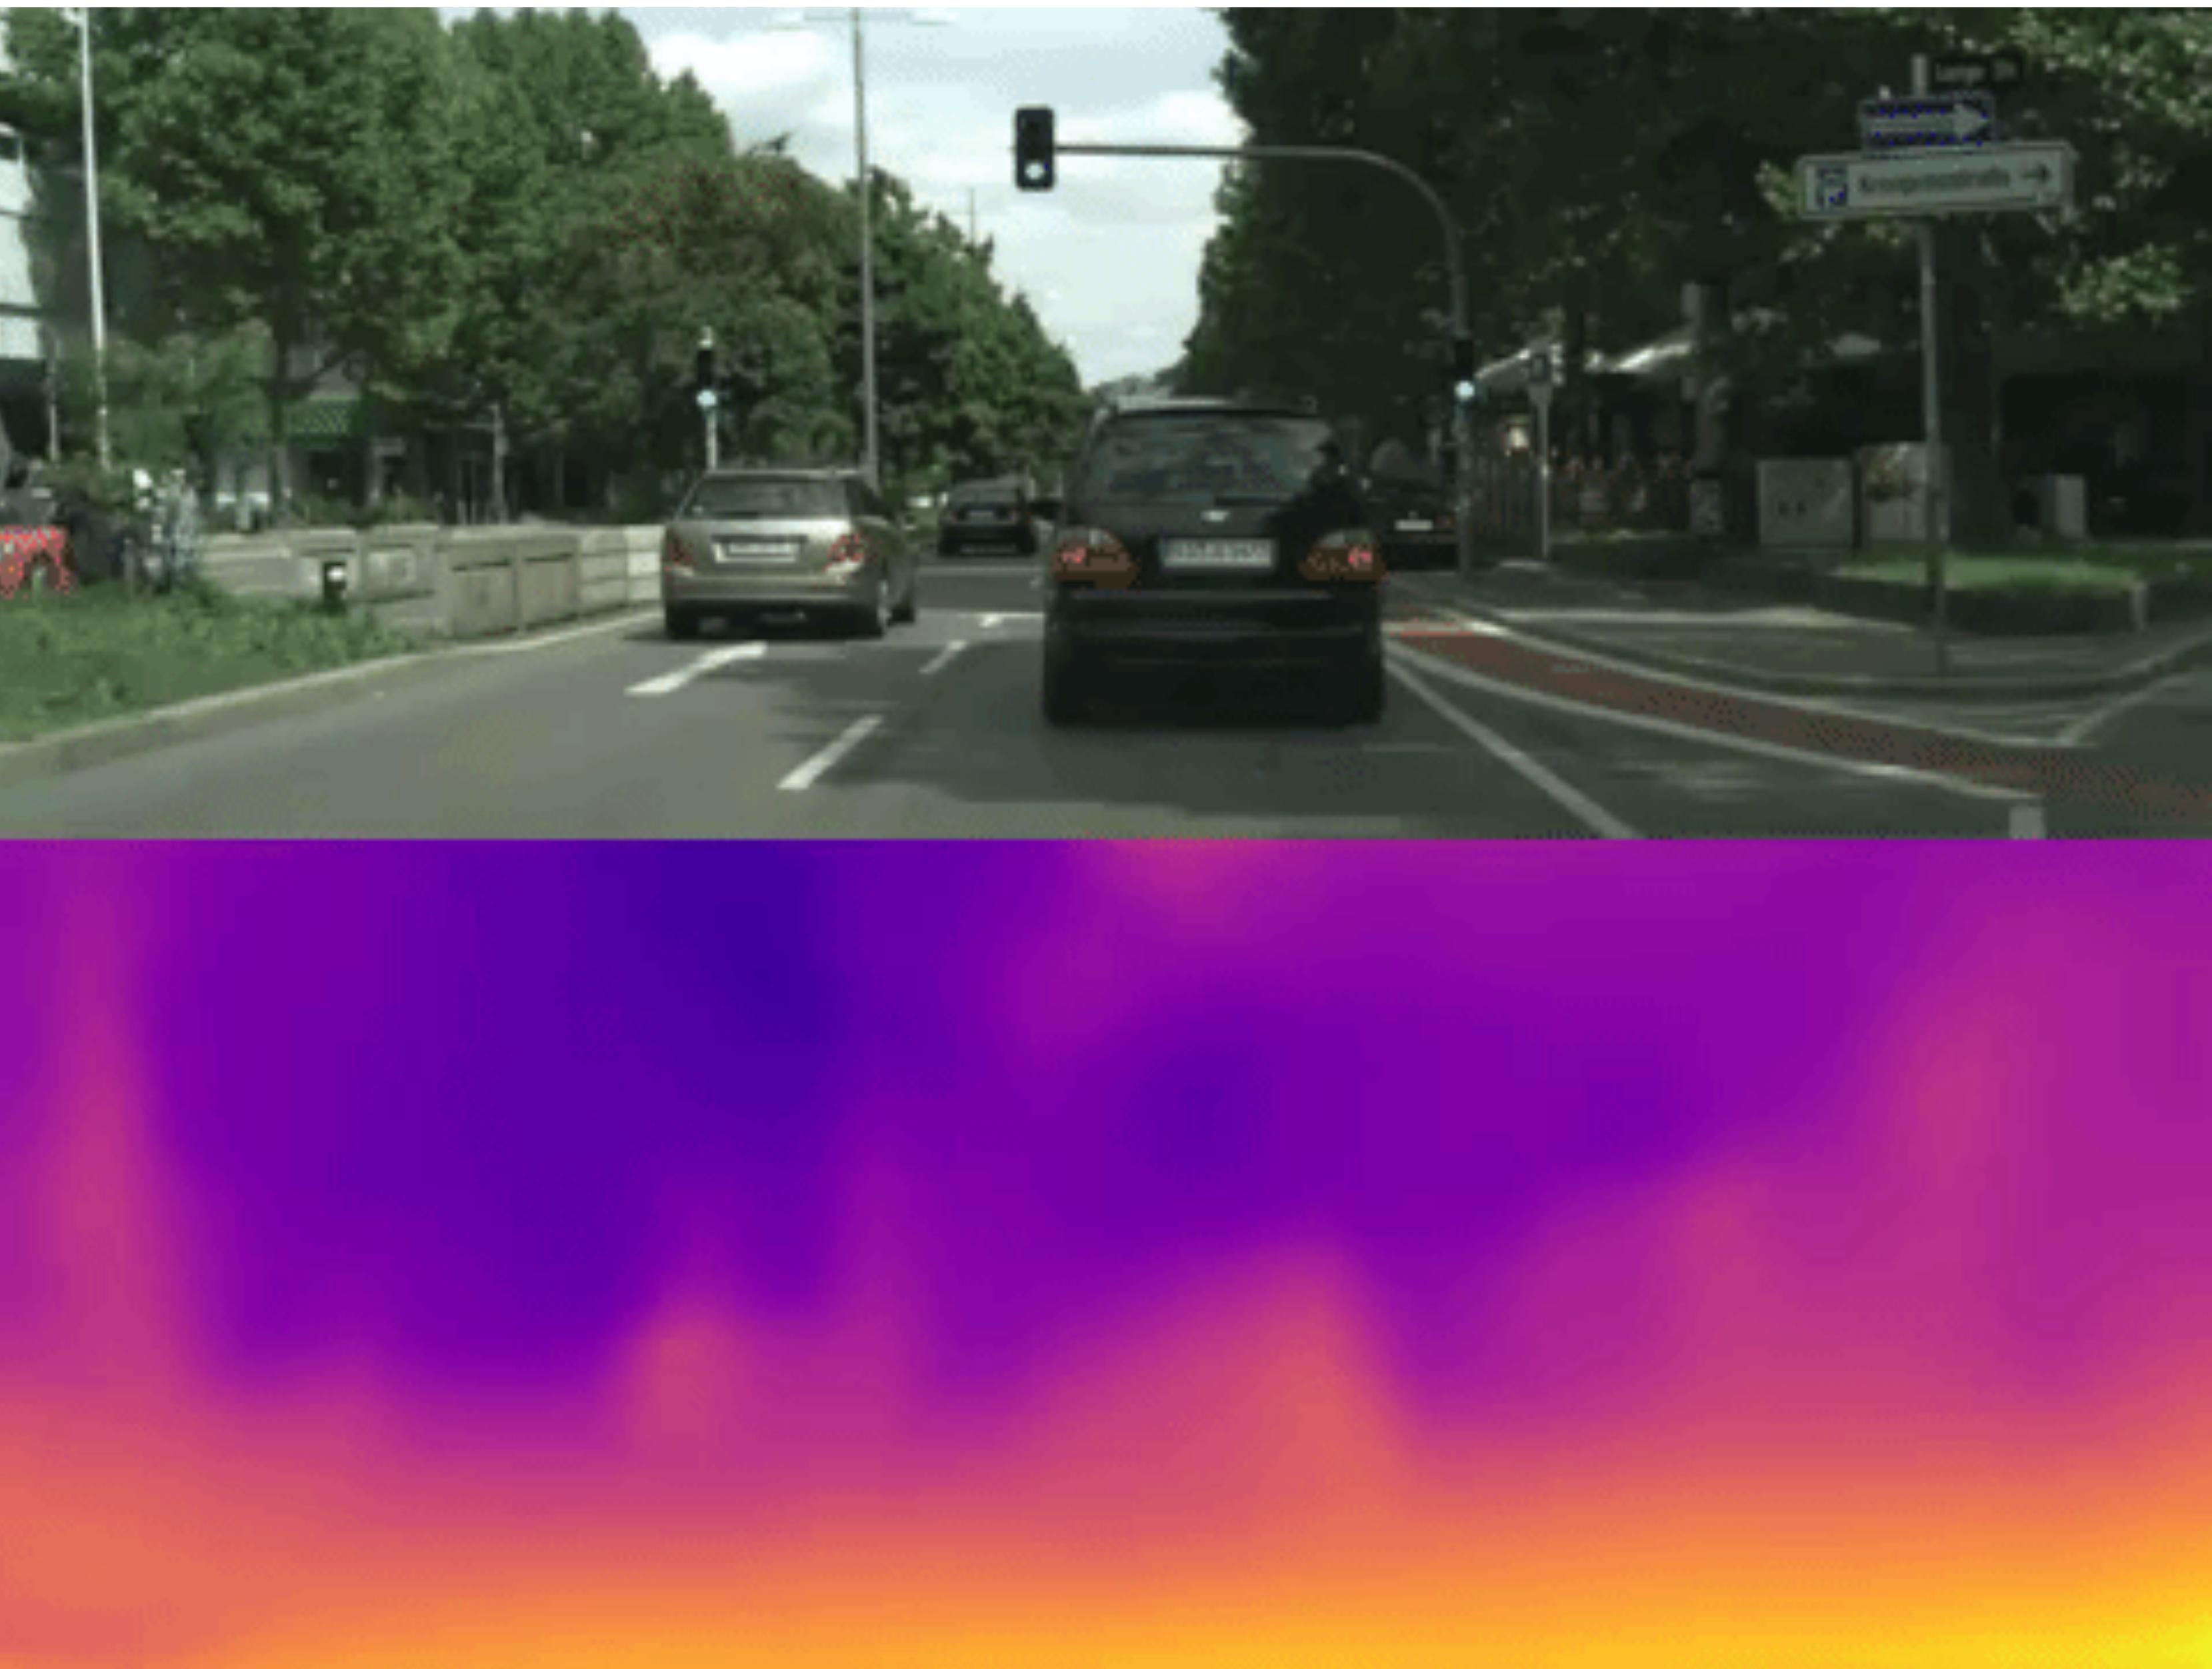


$$p_s \sim K \hat{T}_{t \rightarrow s} \hat{D}_t(p_t) K^{-1} p_t$$

$$\hat{I}_s(p_t) = I_s(p_s) = \sum_{i \in \{t,b\}, j \in \{l,r\}} w^{ij} I_s(p_s^{ij})$$

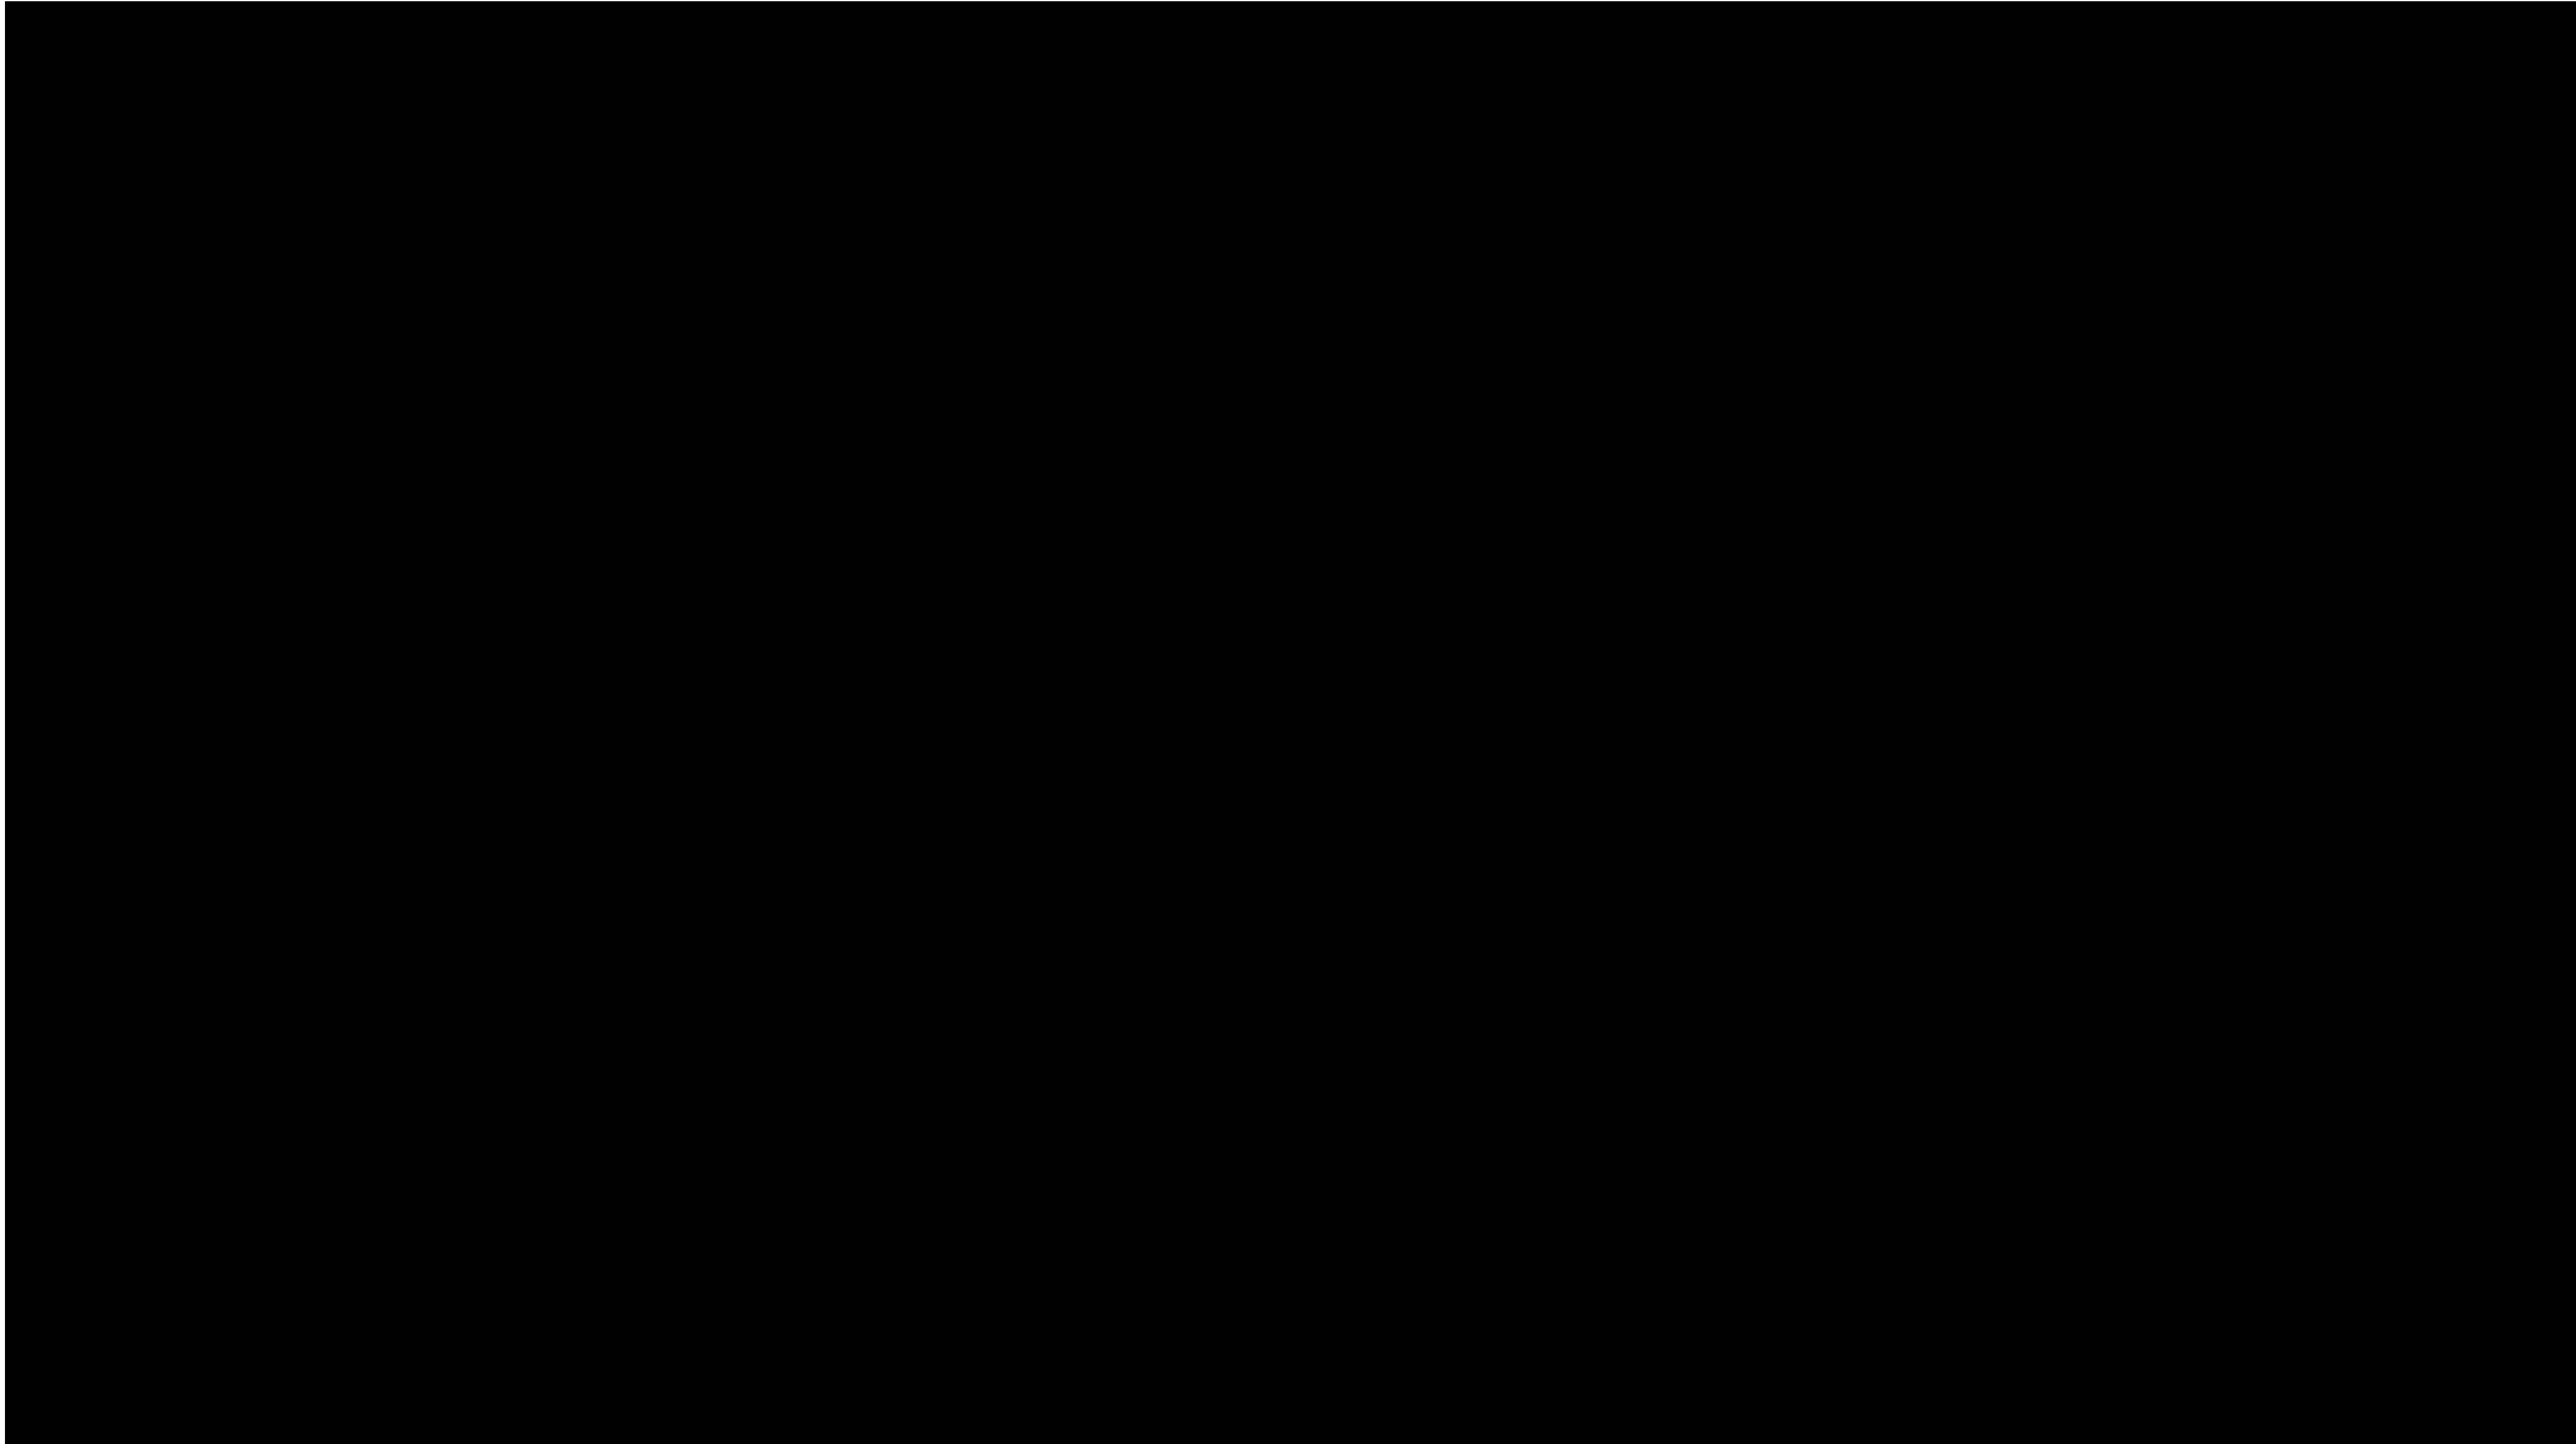
$$\mathcal{L}_{vs} = \sum_s \sum_p |I_t(p) - \hat{I}_s(p)|$$





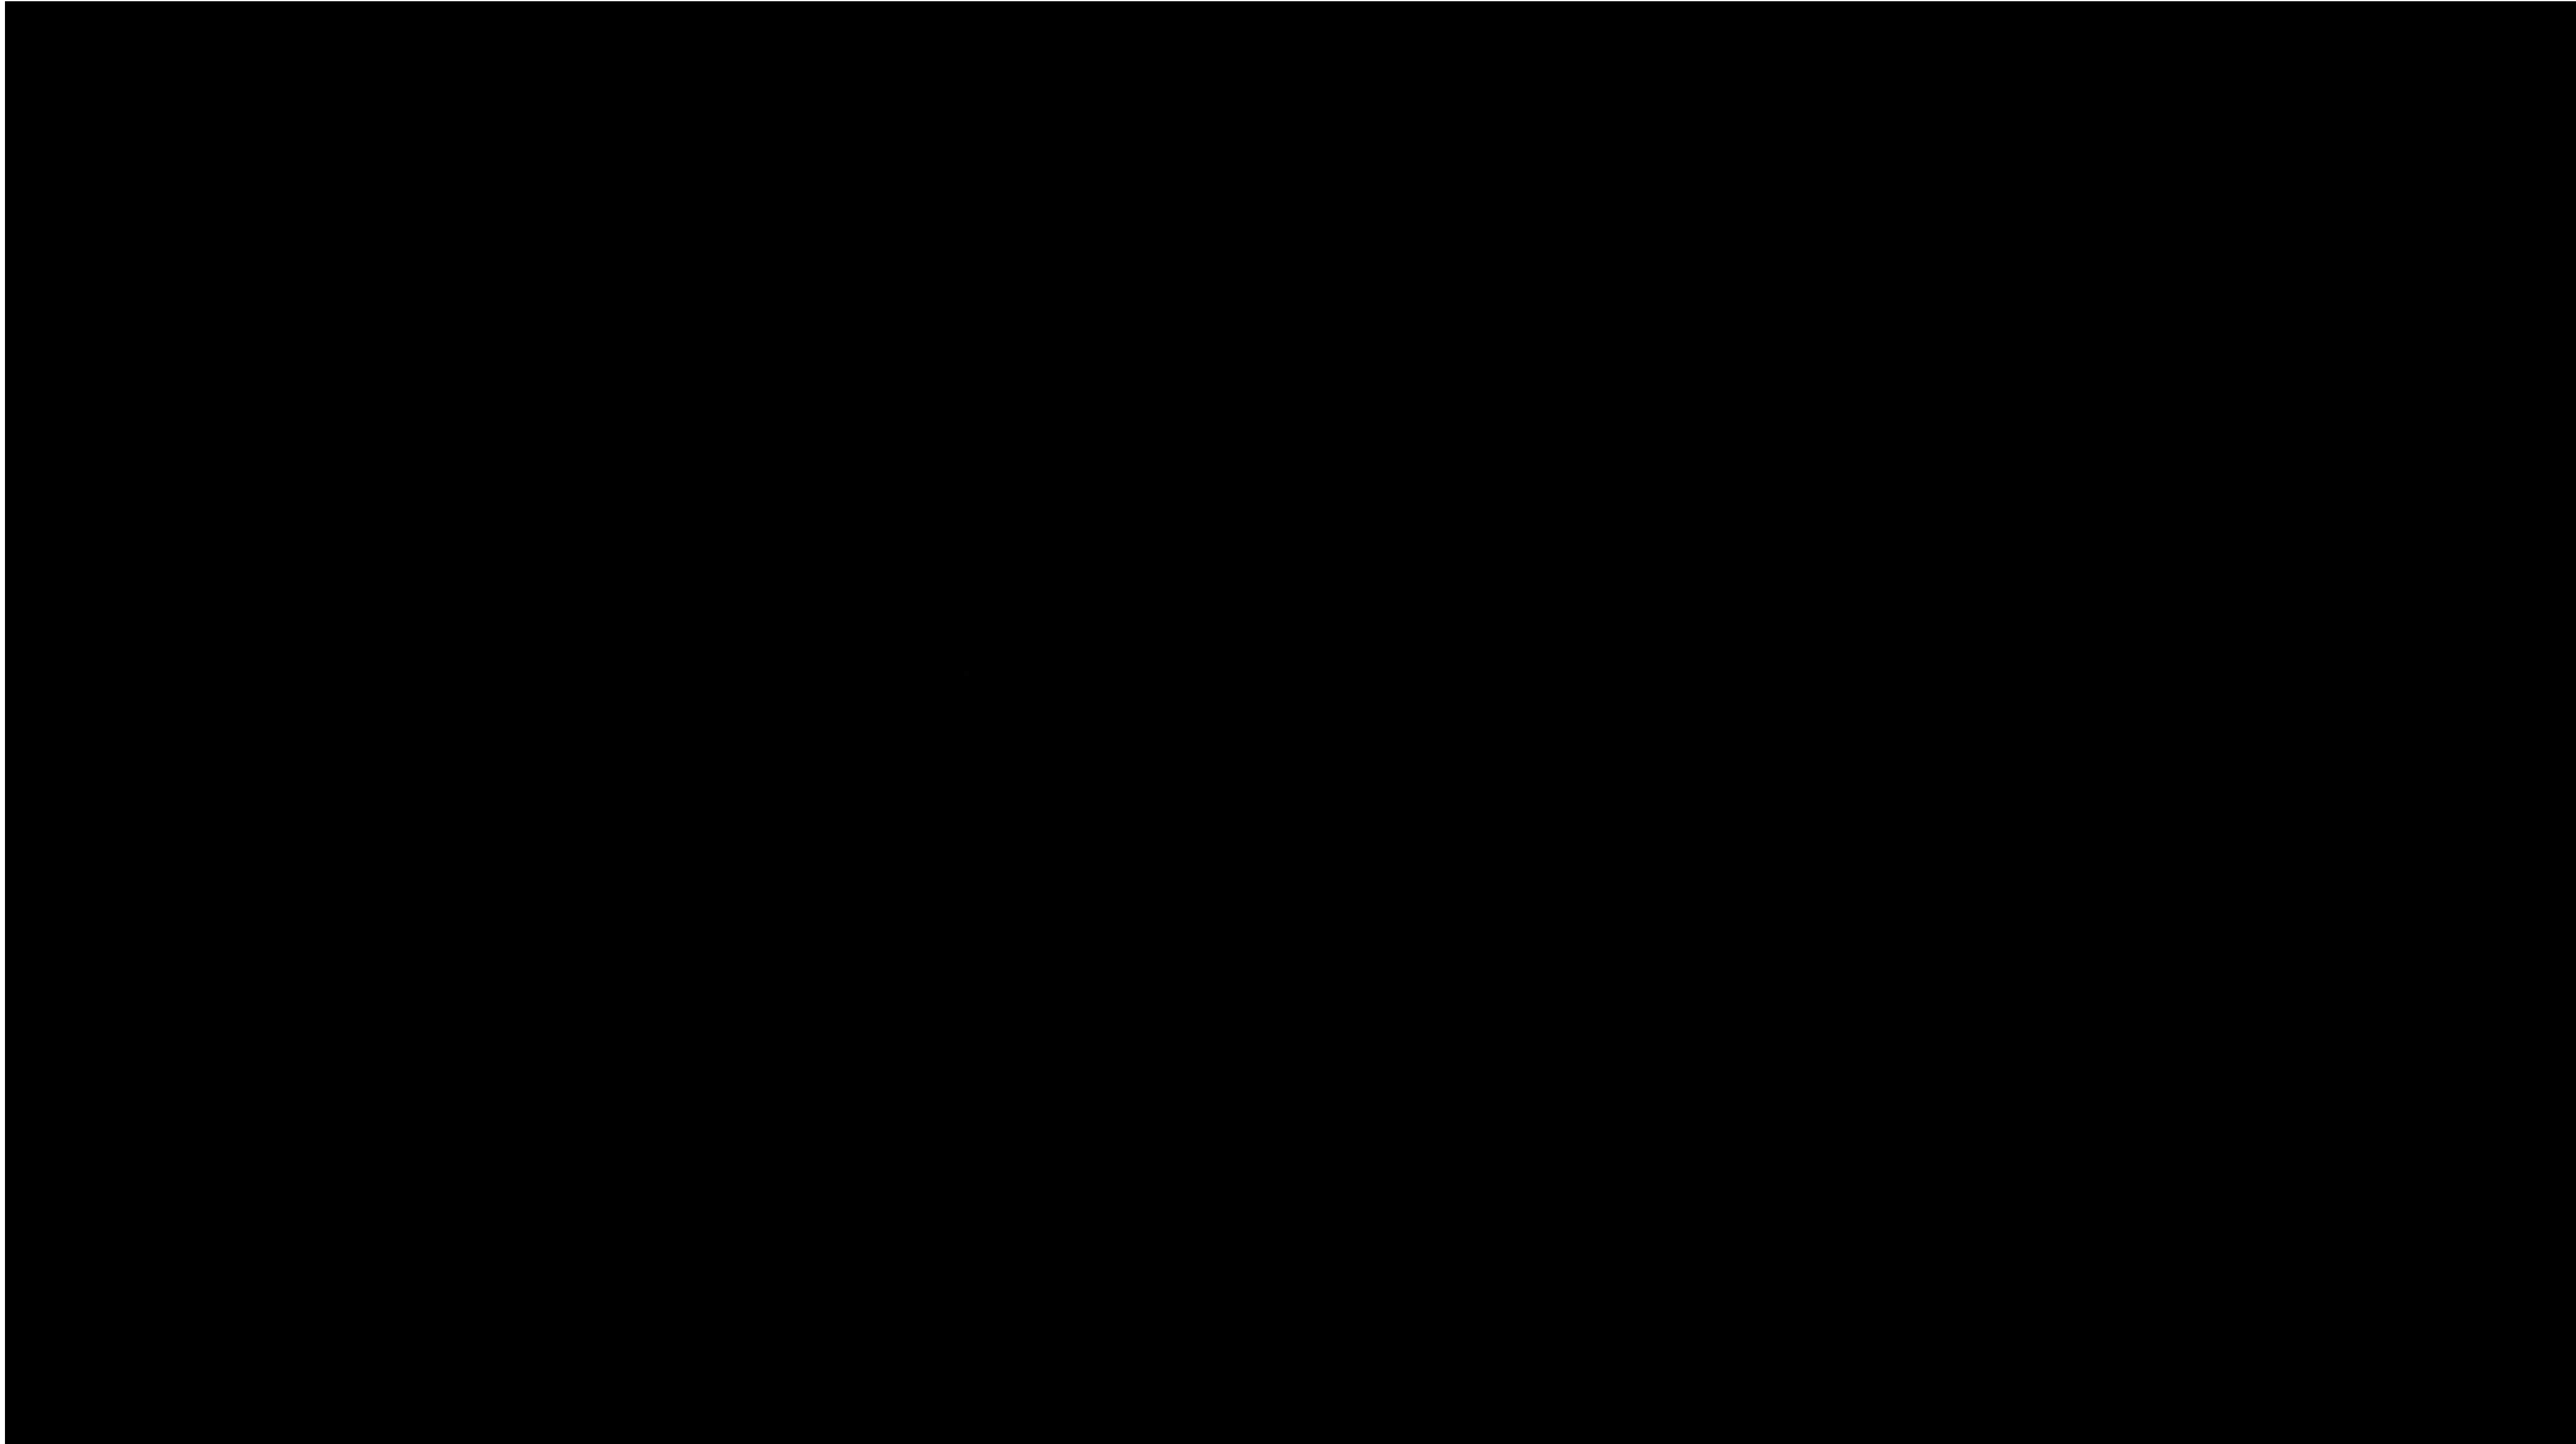
# Self-supervised Learning of Depth and Pose from Video

## Guizilini et al. 2021

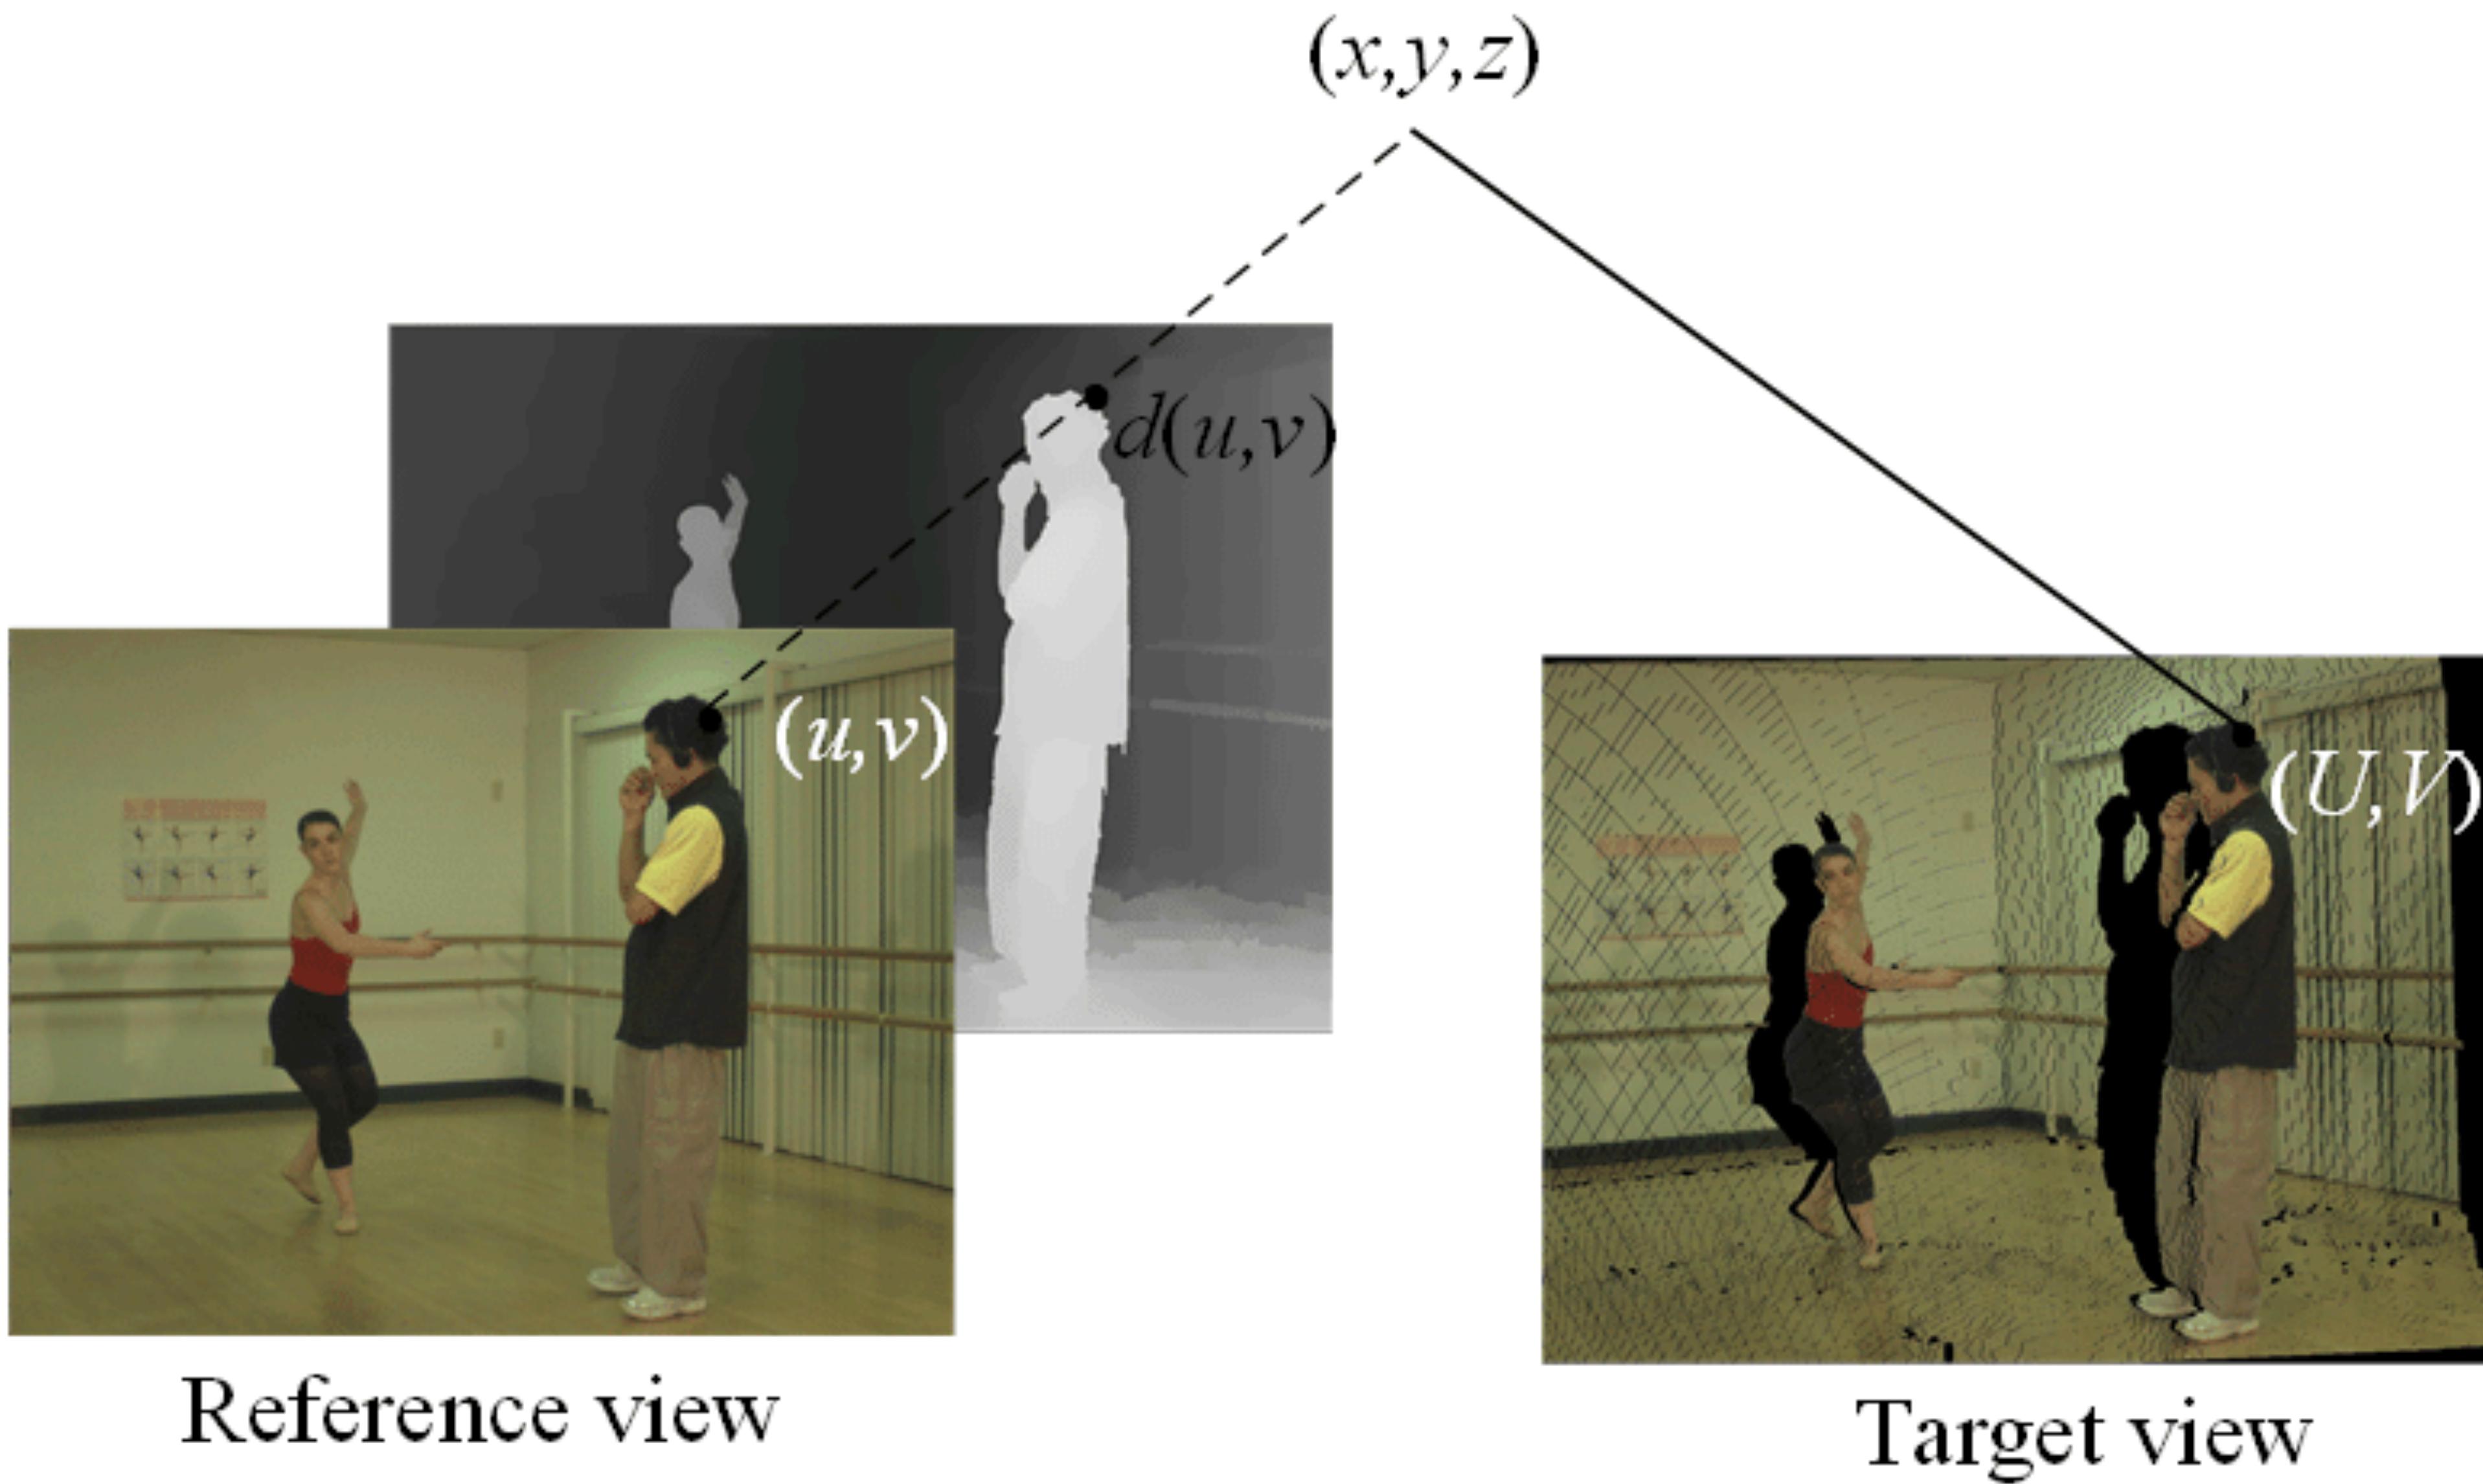


# Self-supervised Learning of Depth and Pose from Video

## Guizilini et al. 2021



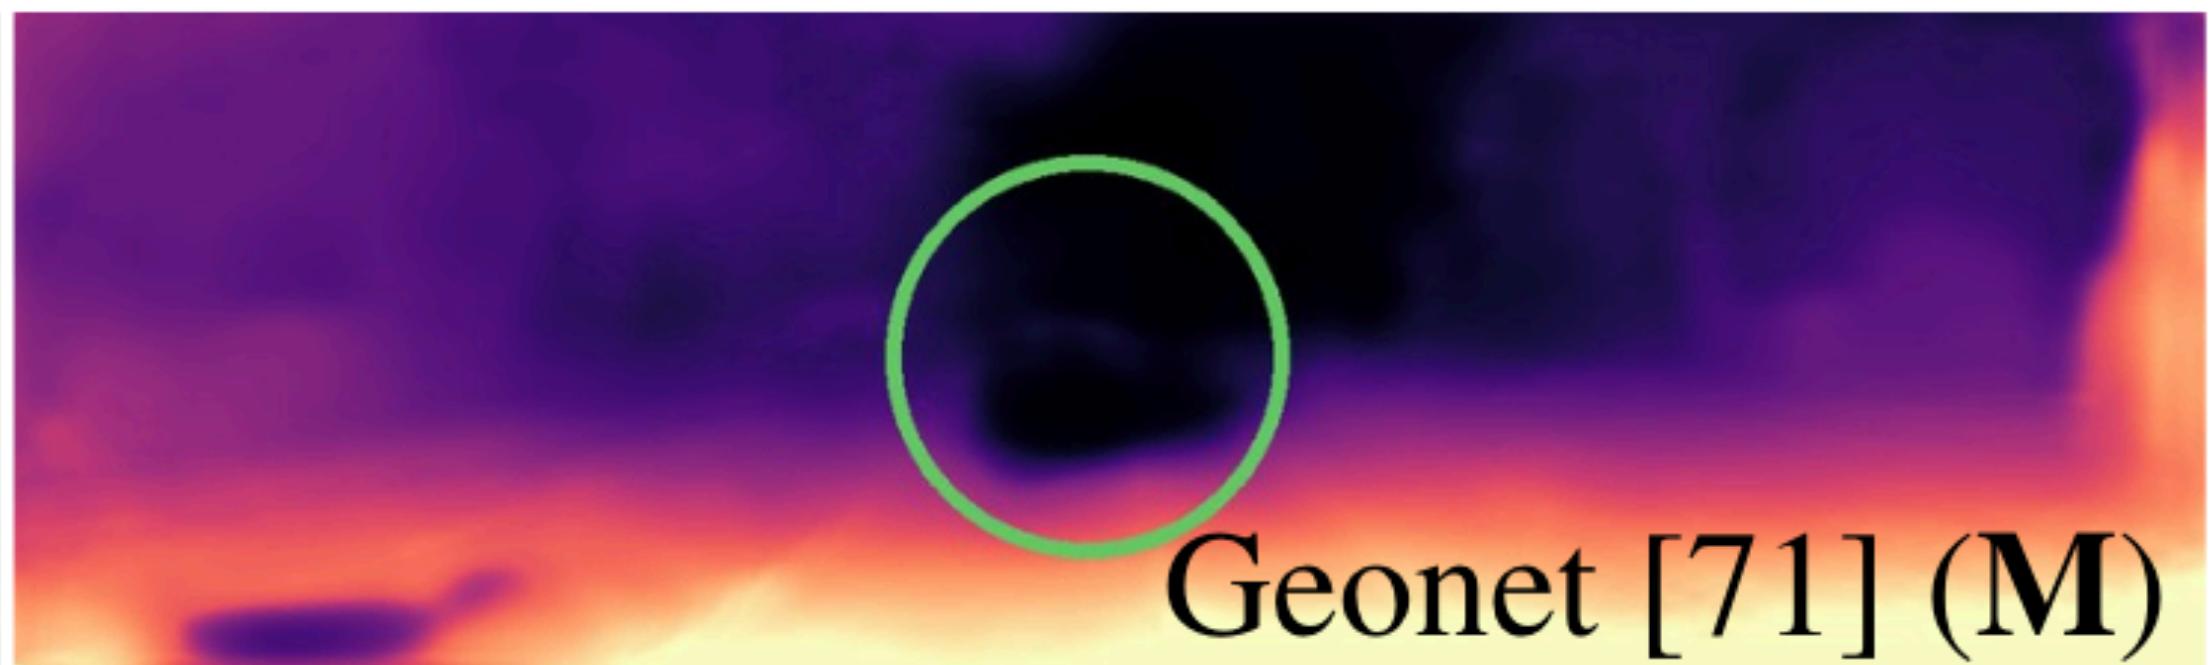
# Can't deal with occlusions / disocclusions



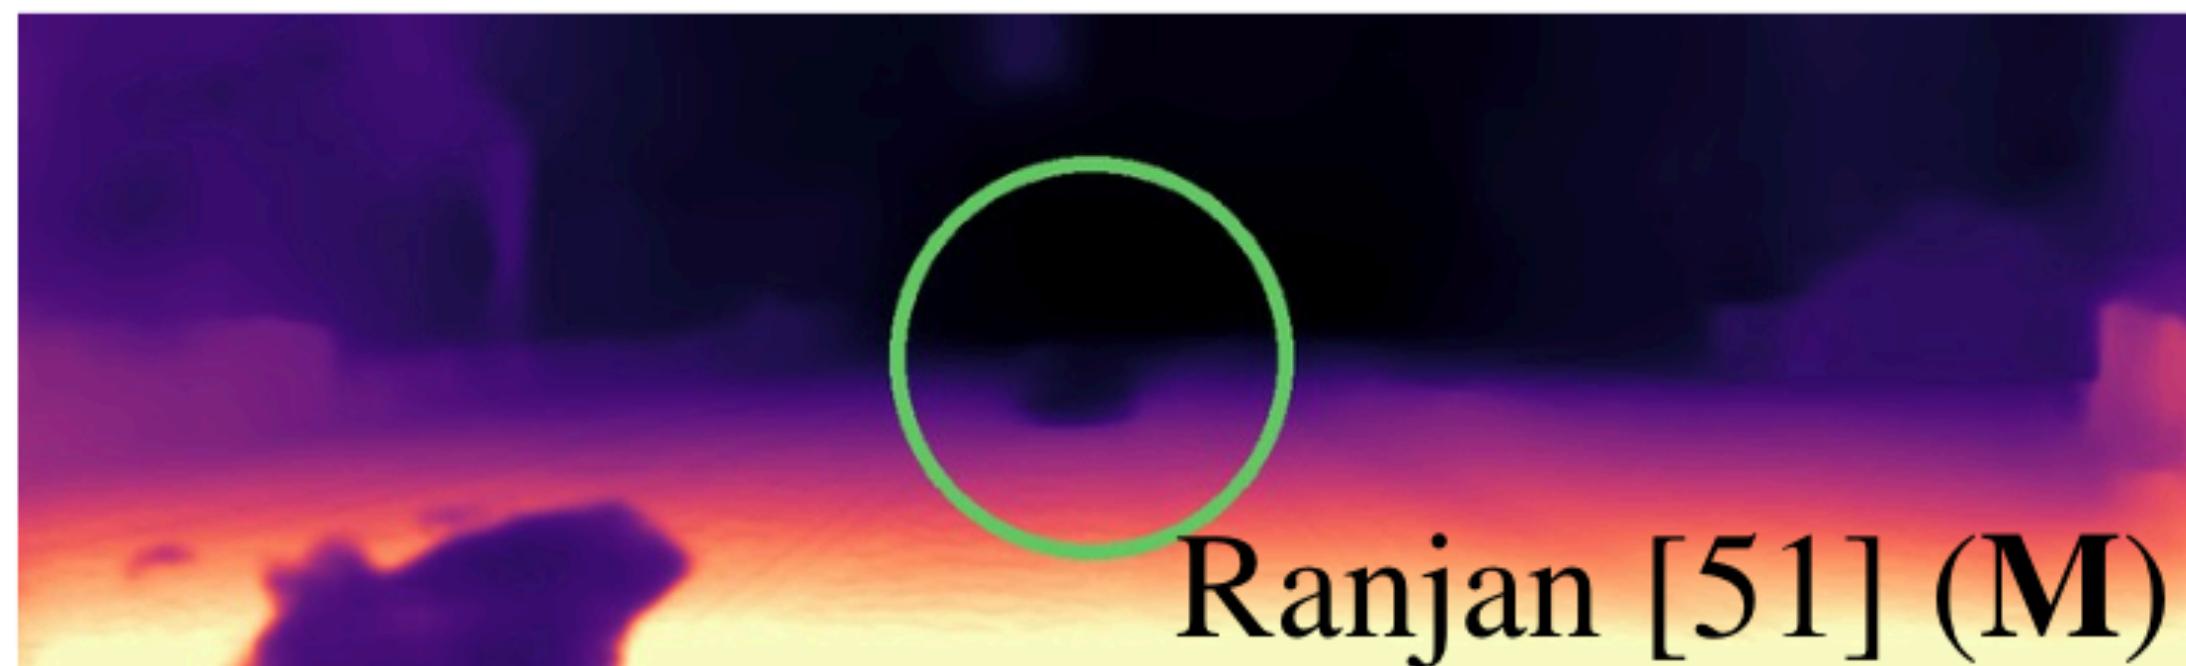
# Assumes static scenes



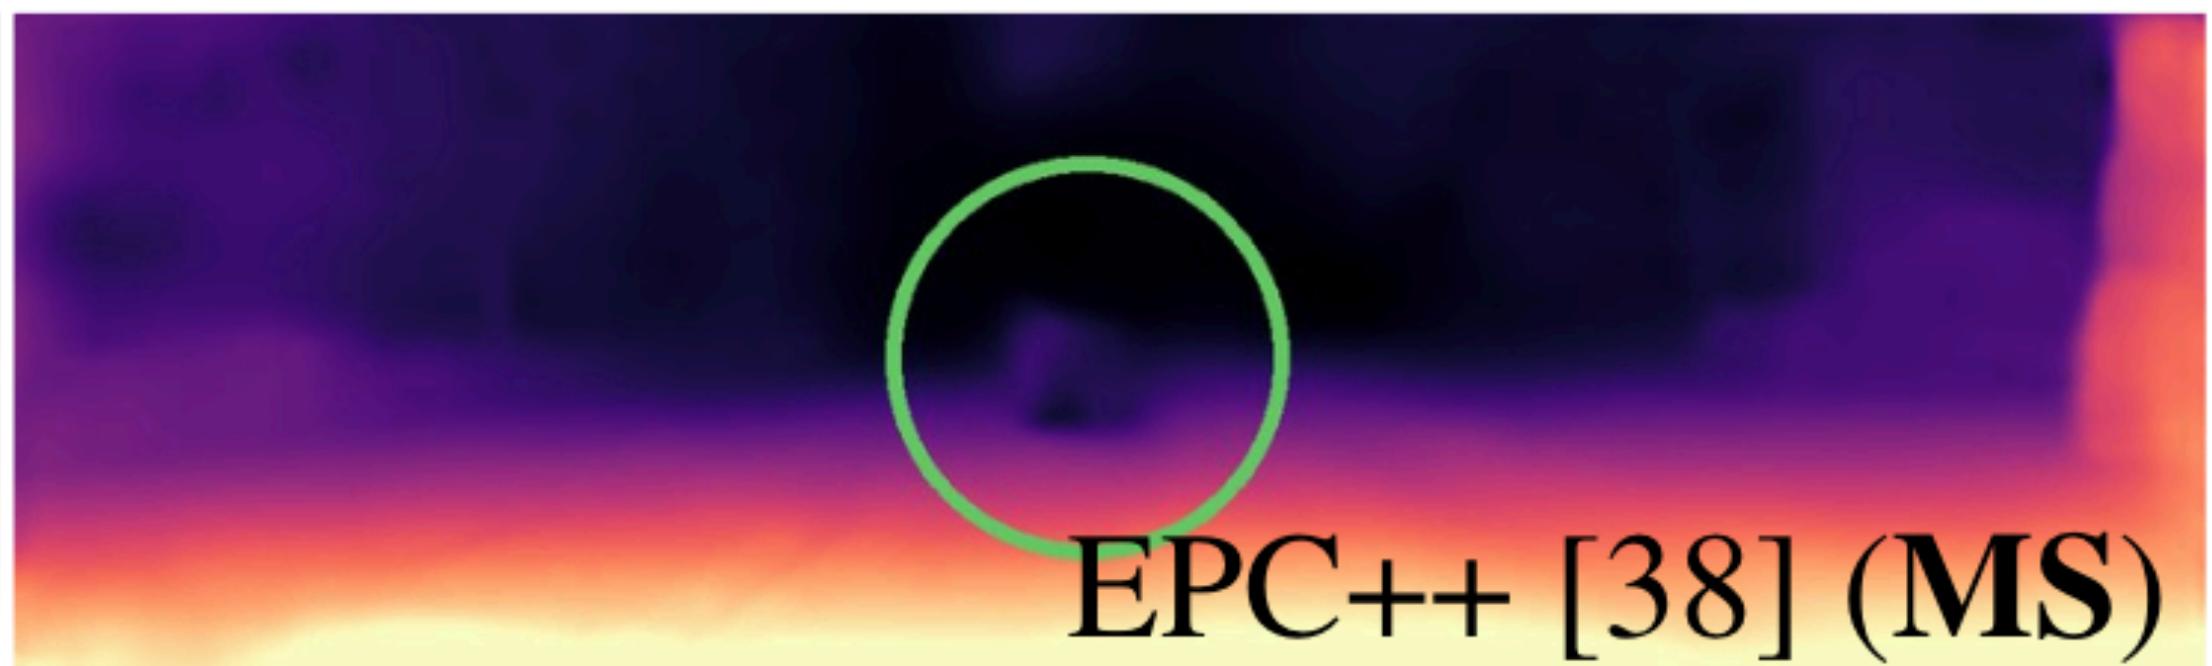
Input



Geonet [71] (M)



Ranjan [51] (M)



EPC++ [38] (MS)

# Digging Into Self-Supervised Monocular Depth Estimation

Clément Godard<sup>1</sup>

Oisin Mac Aodha<sup>2</sup>

Michael Firman<sup>3</sup>

Gabriel Brostow<sup>3,1</sup>

<sup>1</sup>UCL

<sup>2</sup>Caltech

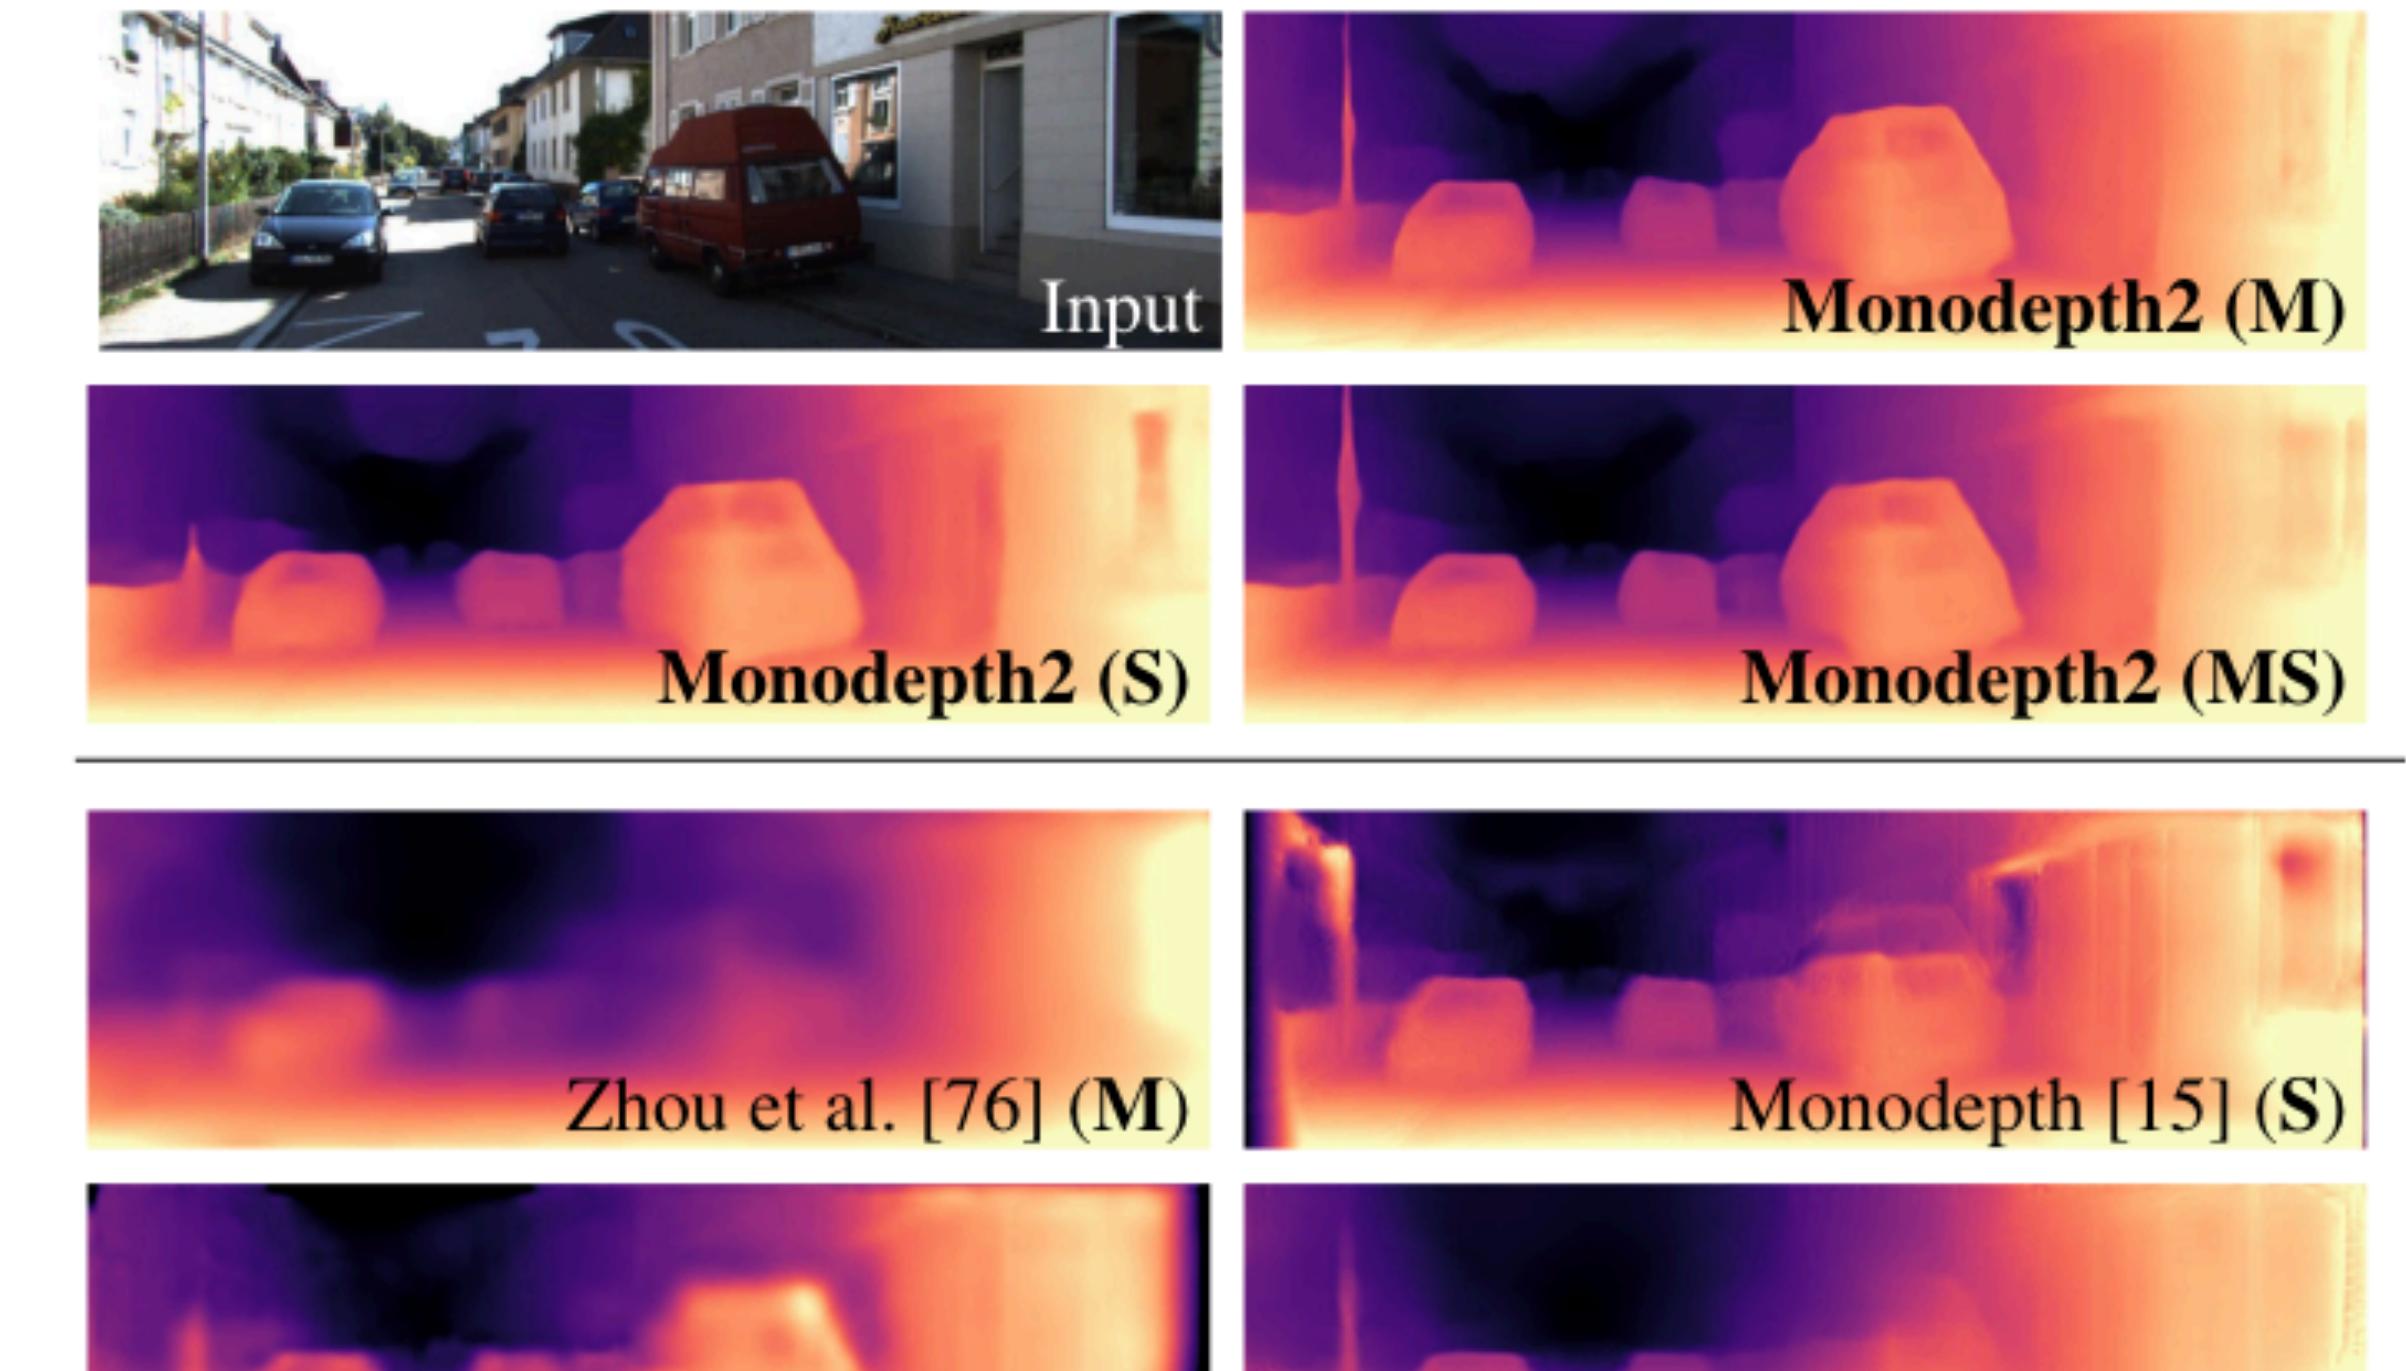
<sup>3</sup>Niantic

[www.github.com/nianticlabs/monodepth2](http://www.github.com/nianticlabs/monodepth2)

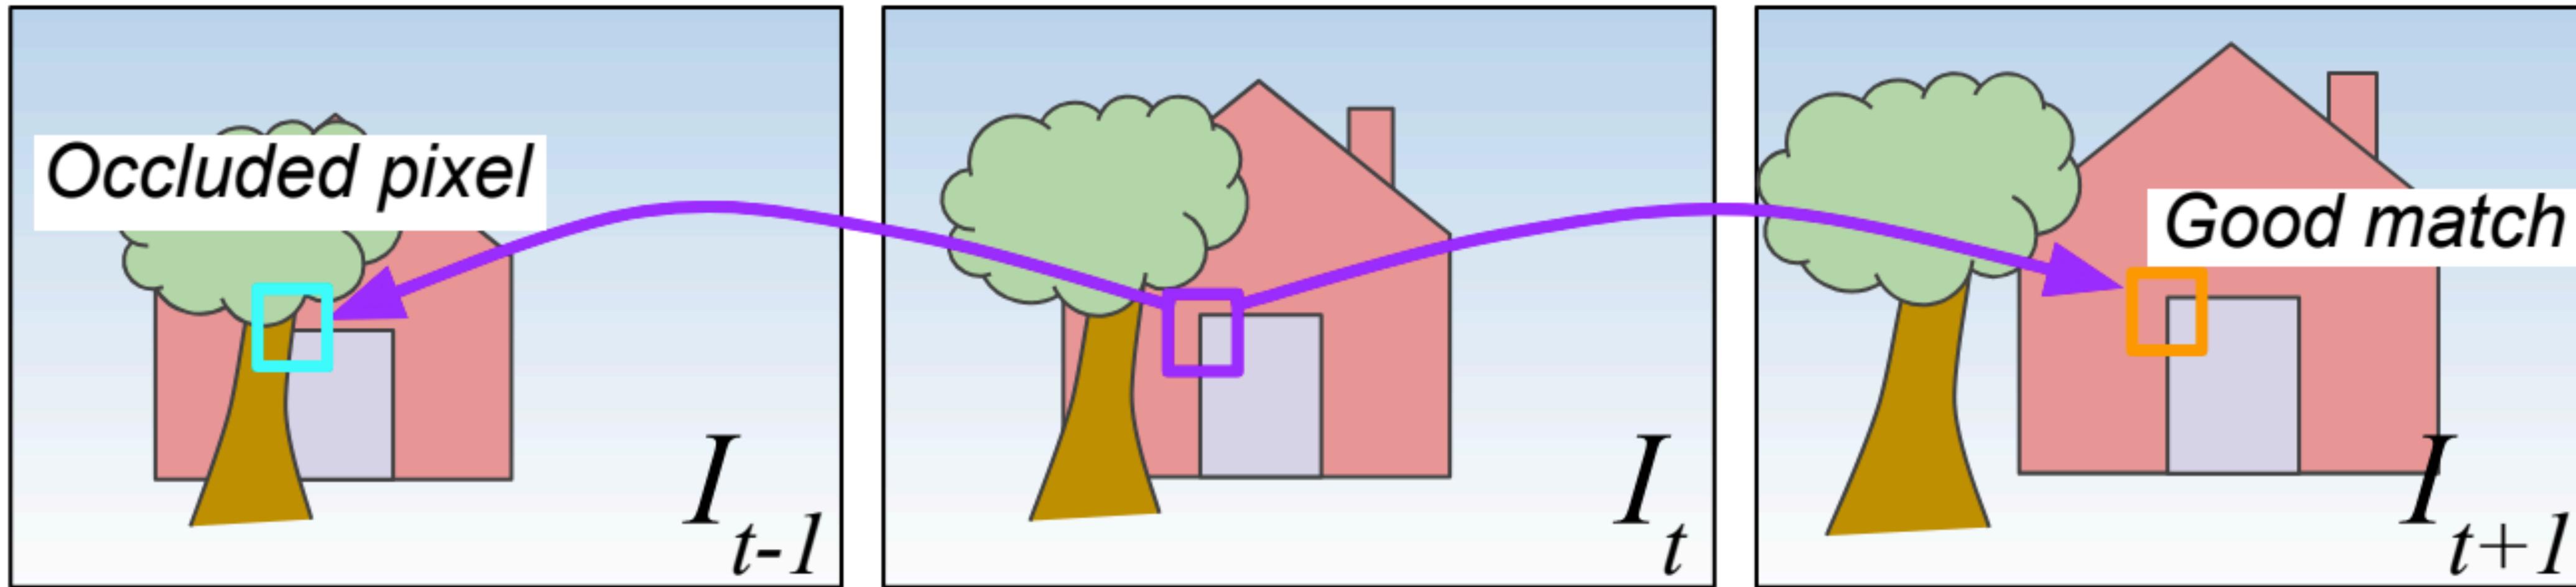
## Abstract

*Per-pixel ground-truth depth data is challenging to acquire at scale. To overcome this limitation, self-supervised learning has emerged as a promising alternative for training models to perform monocular depth estimation. In this paper, we propose a set of improvements, which together result in both quantitatively and qualitatively improved depth maps compared to competing self-supervised methods.*

*Research on self-supervised monocular training usually explores increasingly complex architectures, loss functions, and image formation models, all of which have recently*



# Dealing with Occlusions / Disocclusions



$$pe(\boxed{\text{red}}, \boxed{\text{green}}) = \begin{array}{|c|} \hline \text{error} \\ \hline \end{array}$$
$$pe(\boxed{\text{red}}, \boxed{\text{red}}) = \begin{array}{|c|} \hline \text{error} \\ \hline \end{array}$$

Baseline:  $\text{avg}(\begin{array}{|c|} \hline \text{blue} \\ \hline \end{array}, \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array}) = \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} \quad \times$

Ours:  $\min(\begin{array}{|c|} \hline \text{blue} \\ \hline \end{array}, \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array}) = \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} \quad \checkmark$

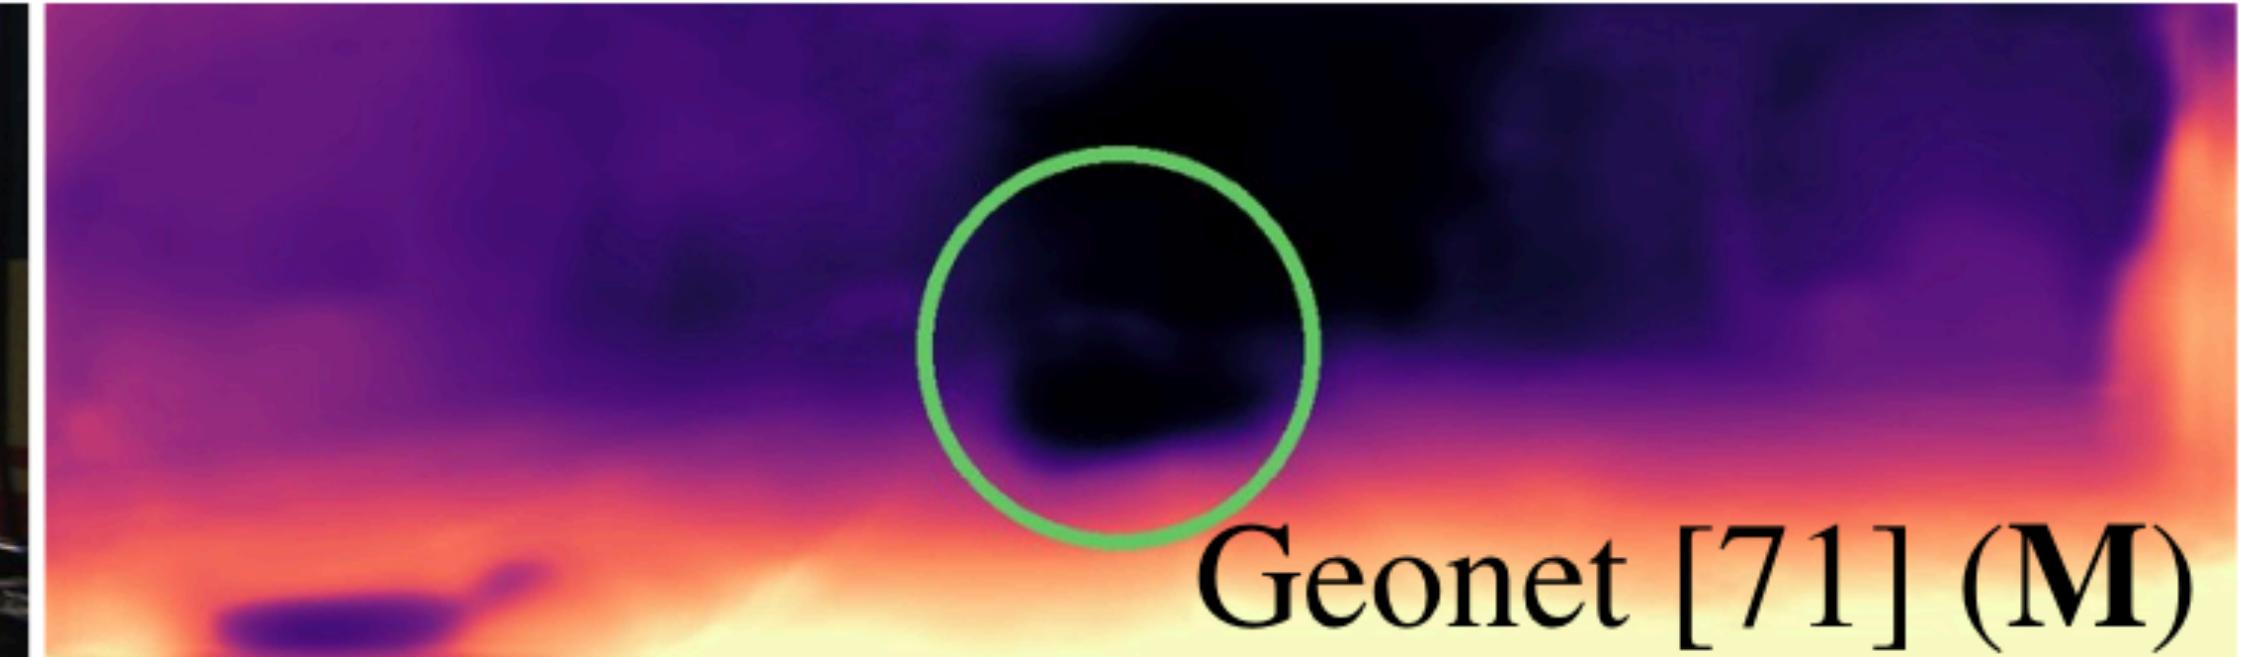
# Dealing with Occlusions / Disocclusions



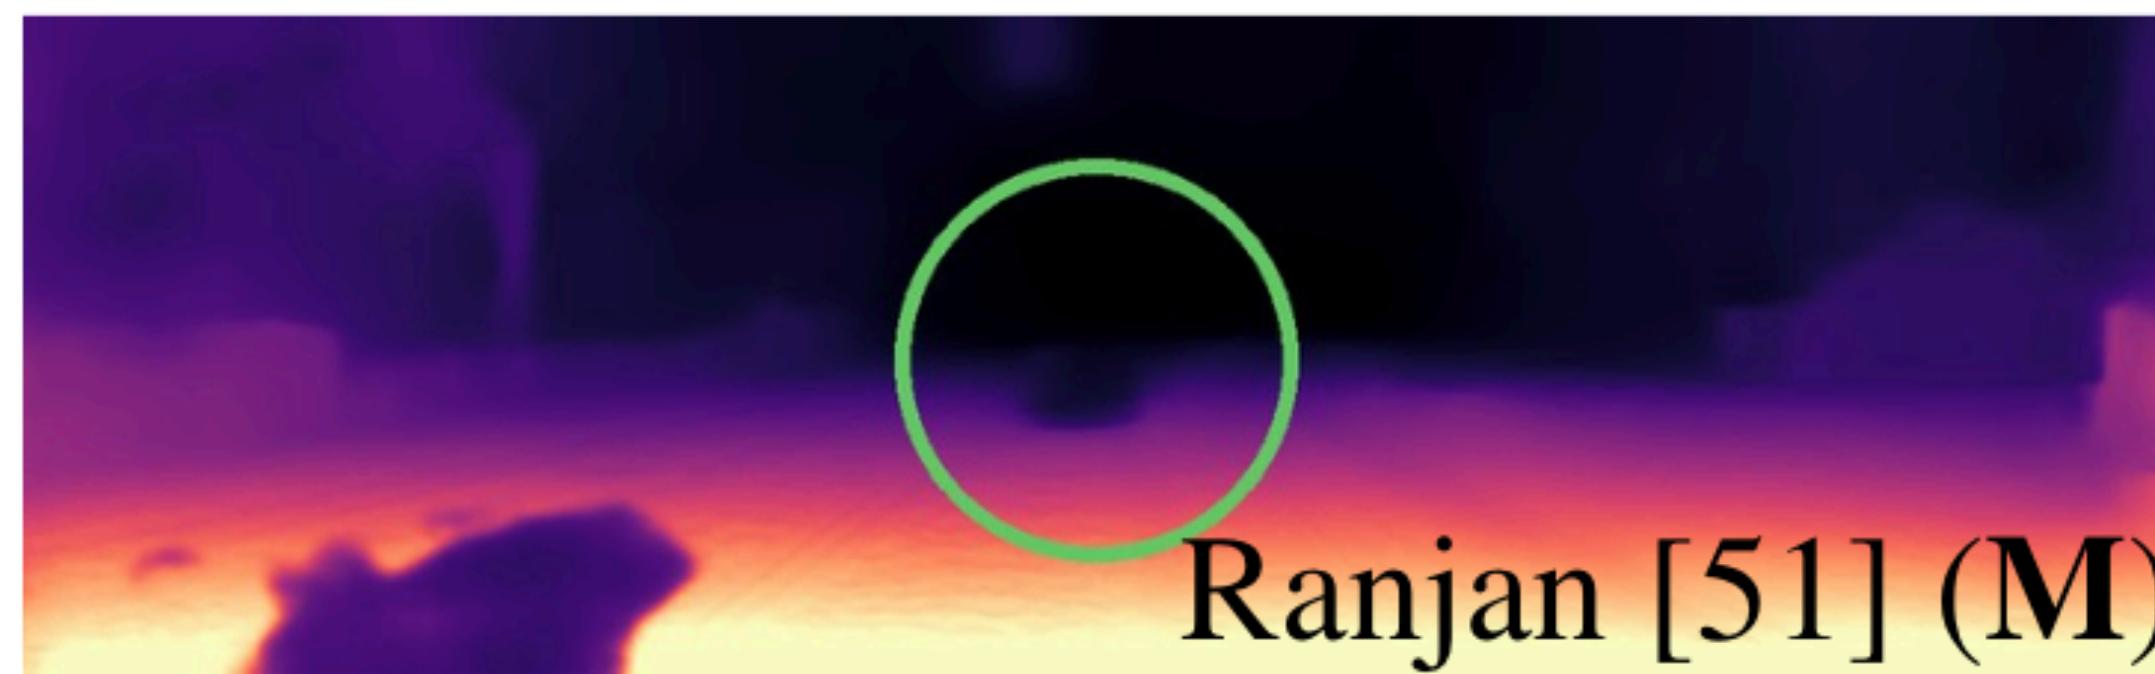
# Dealing with moving objects



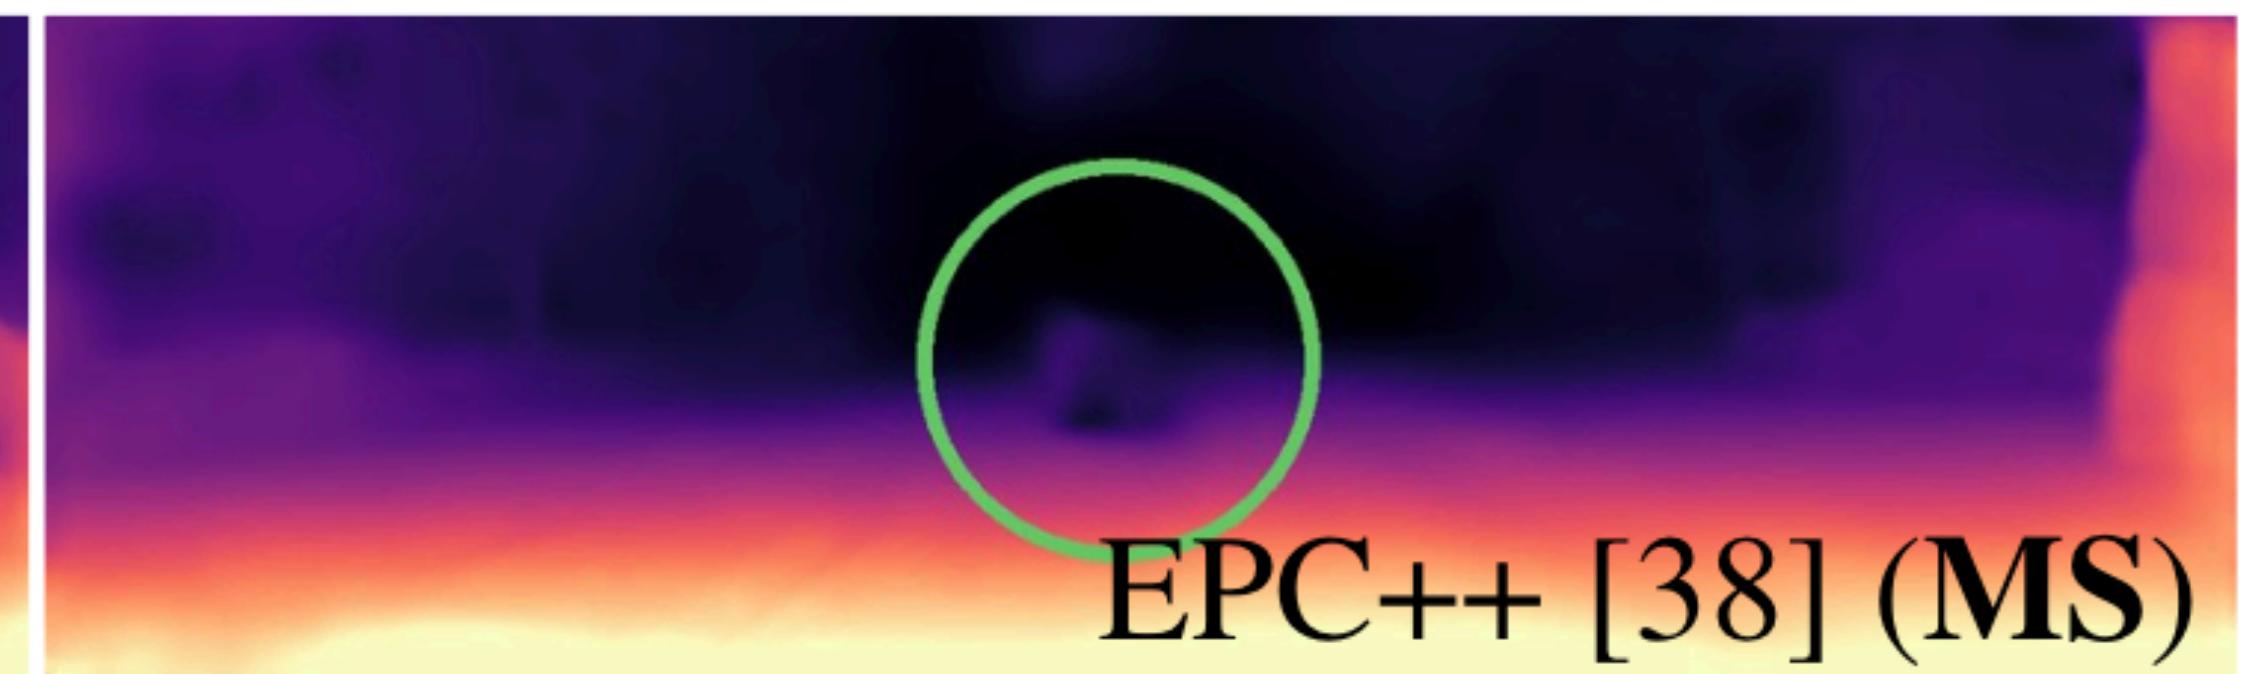
Input



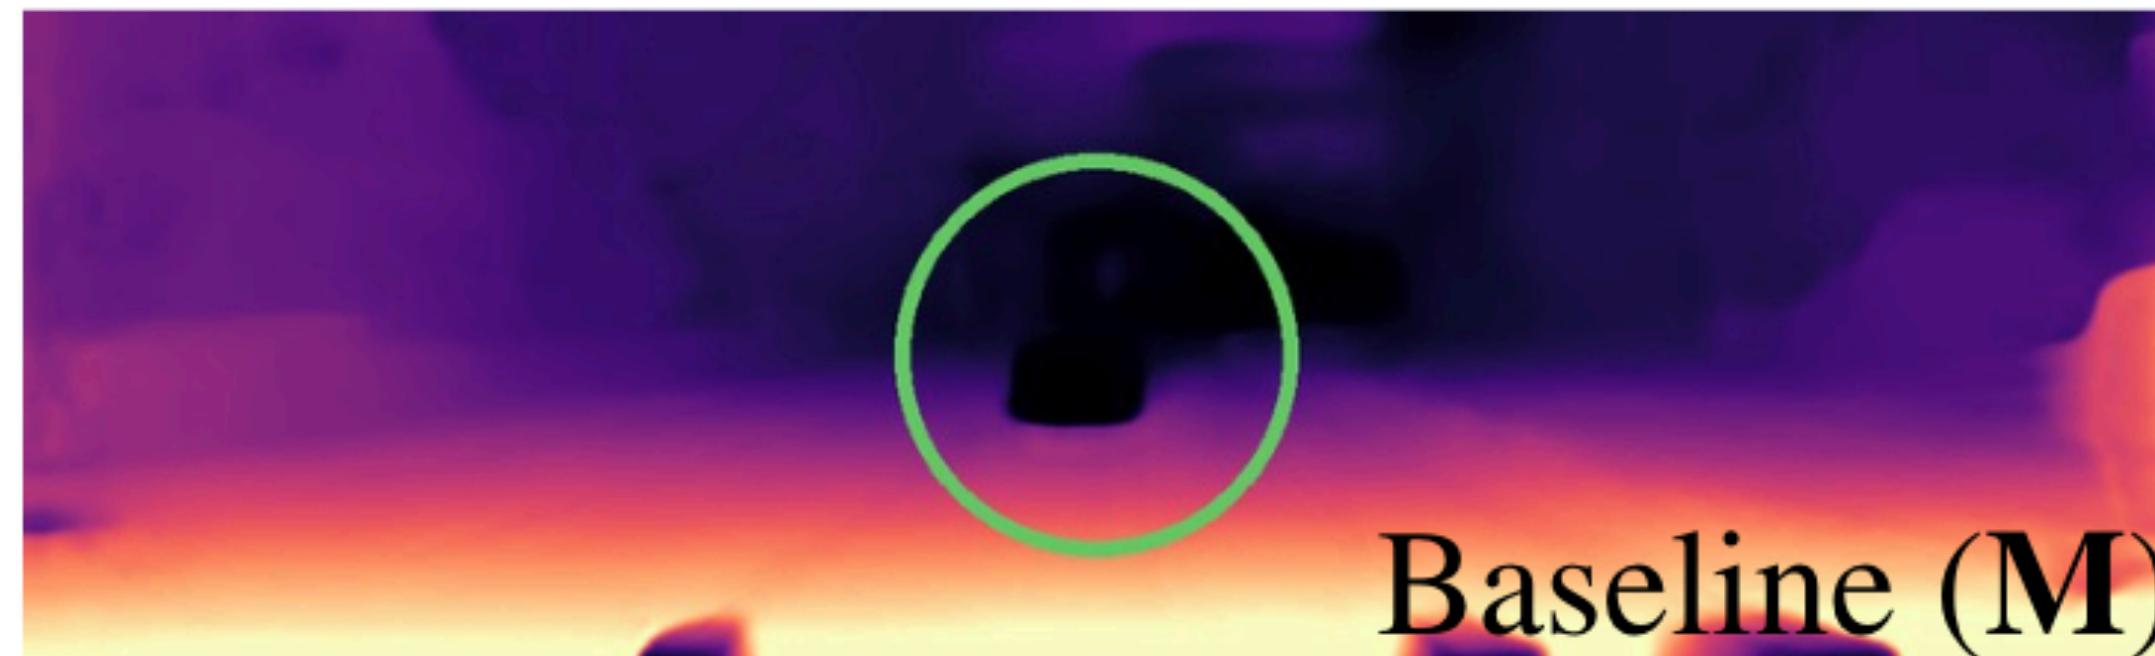
Geonet [71] (M)



Ranjan [51] (M)



EPC++ [38] (MS)



Baseline (M)

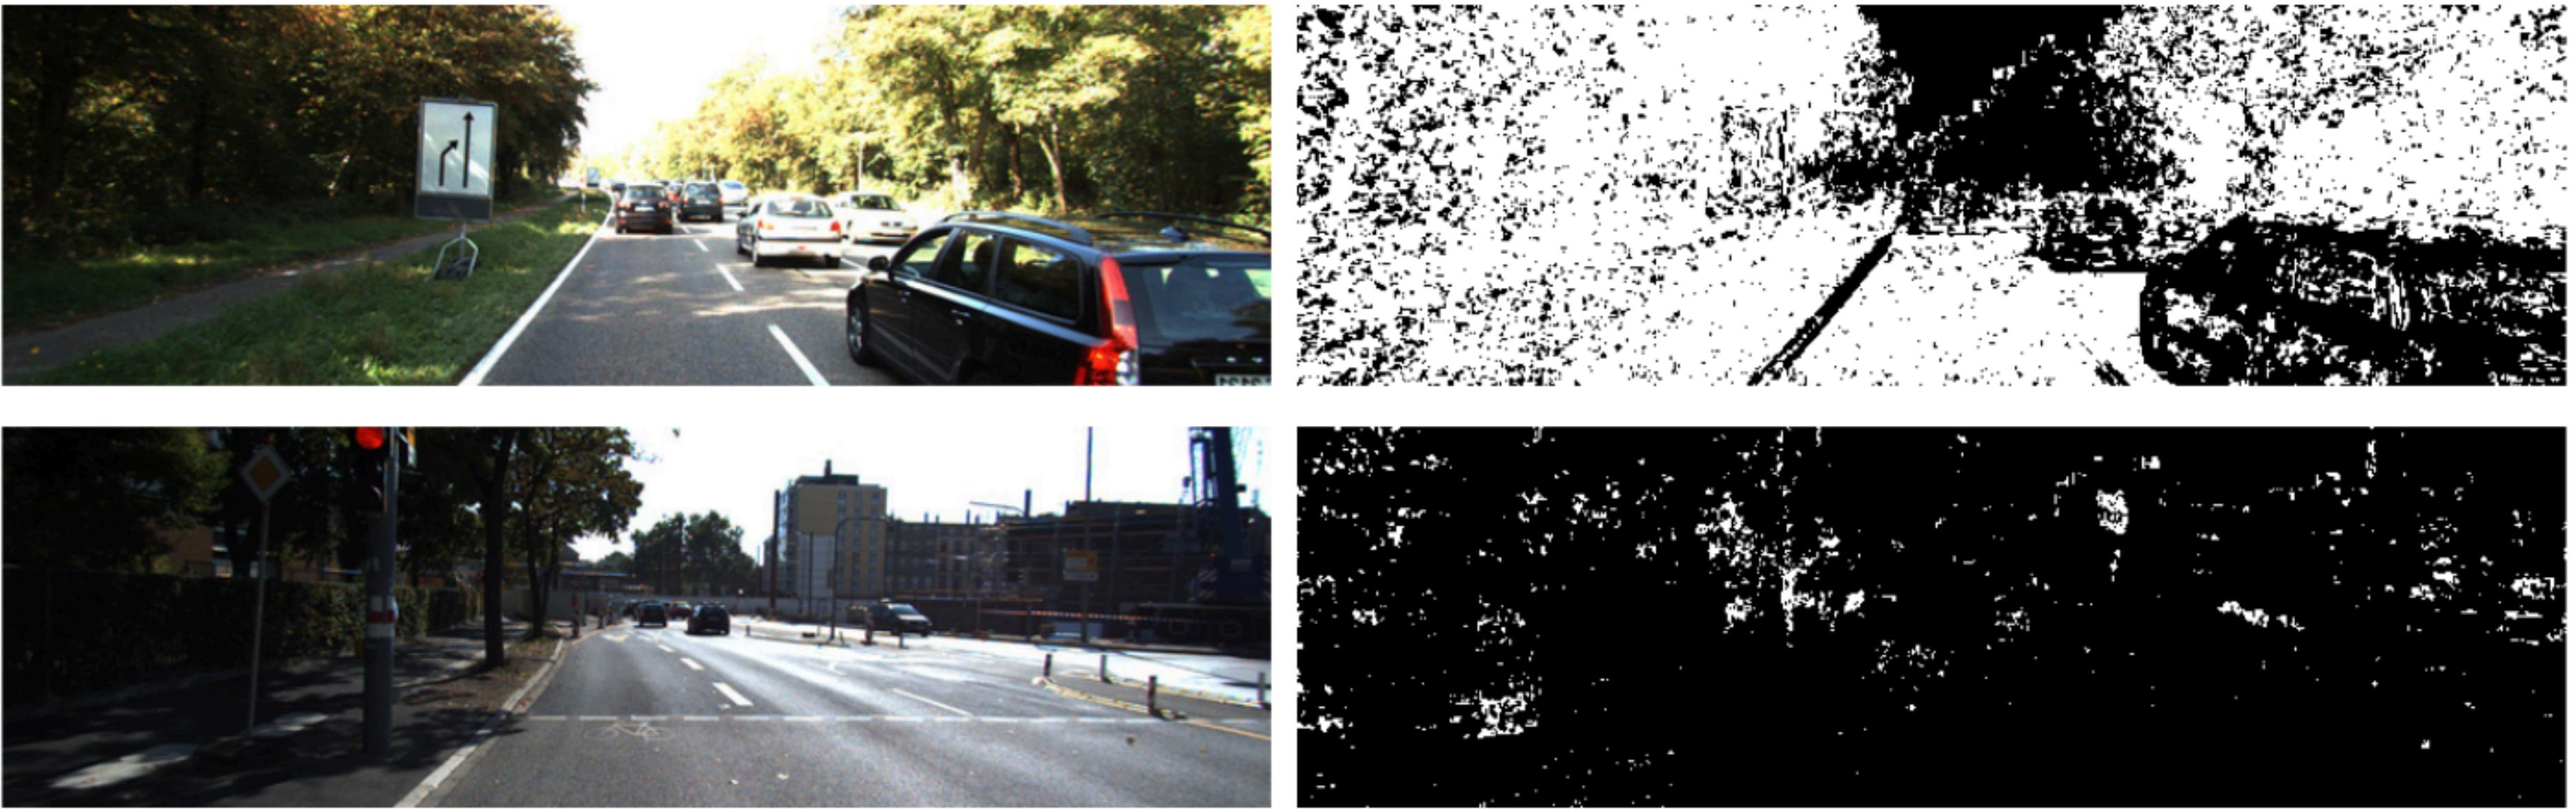


Monodepth2 (M)

# Dealing with moving objects: Predict per-pixel mask $\mu$

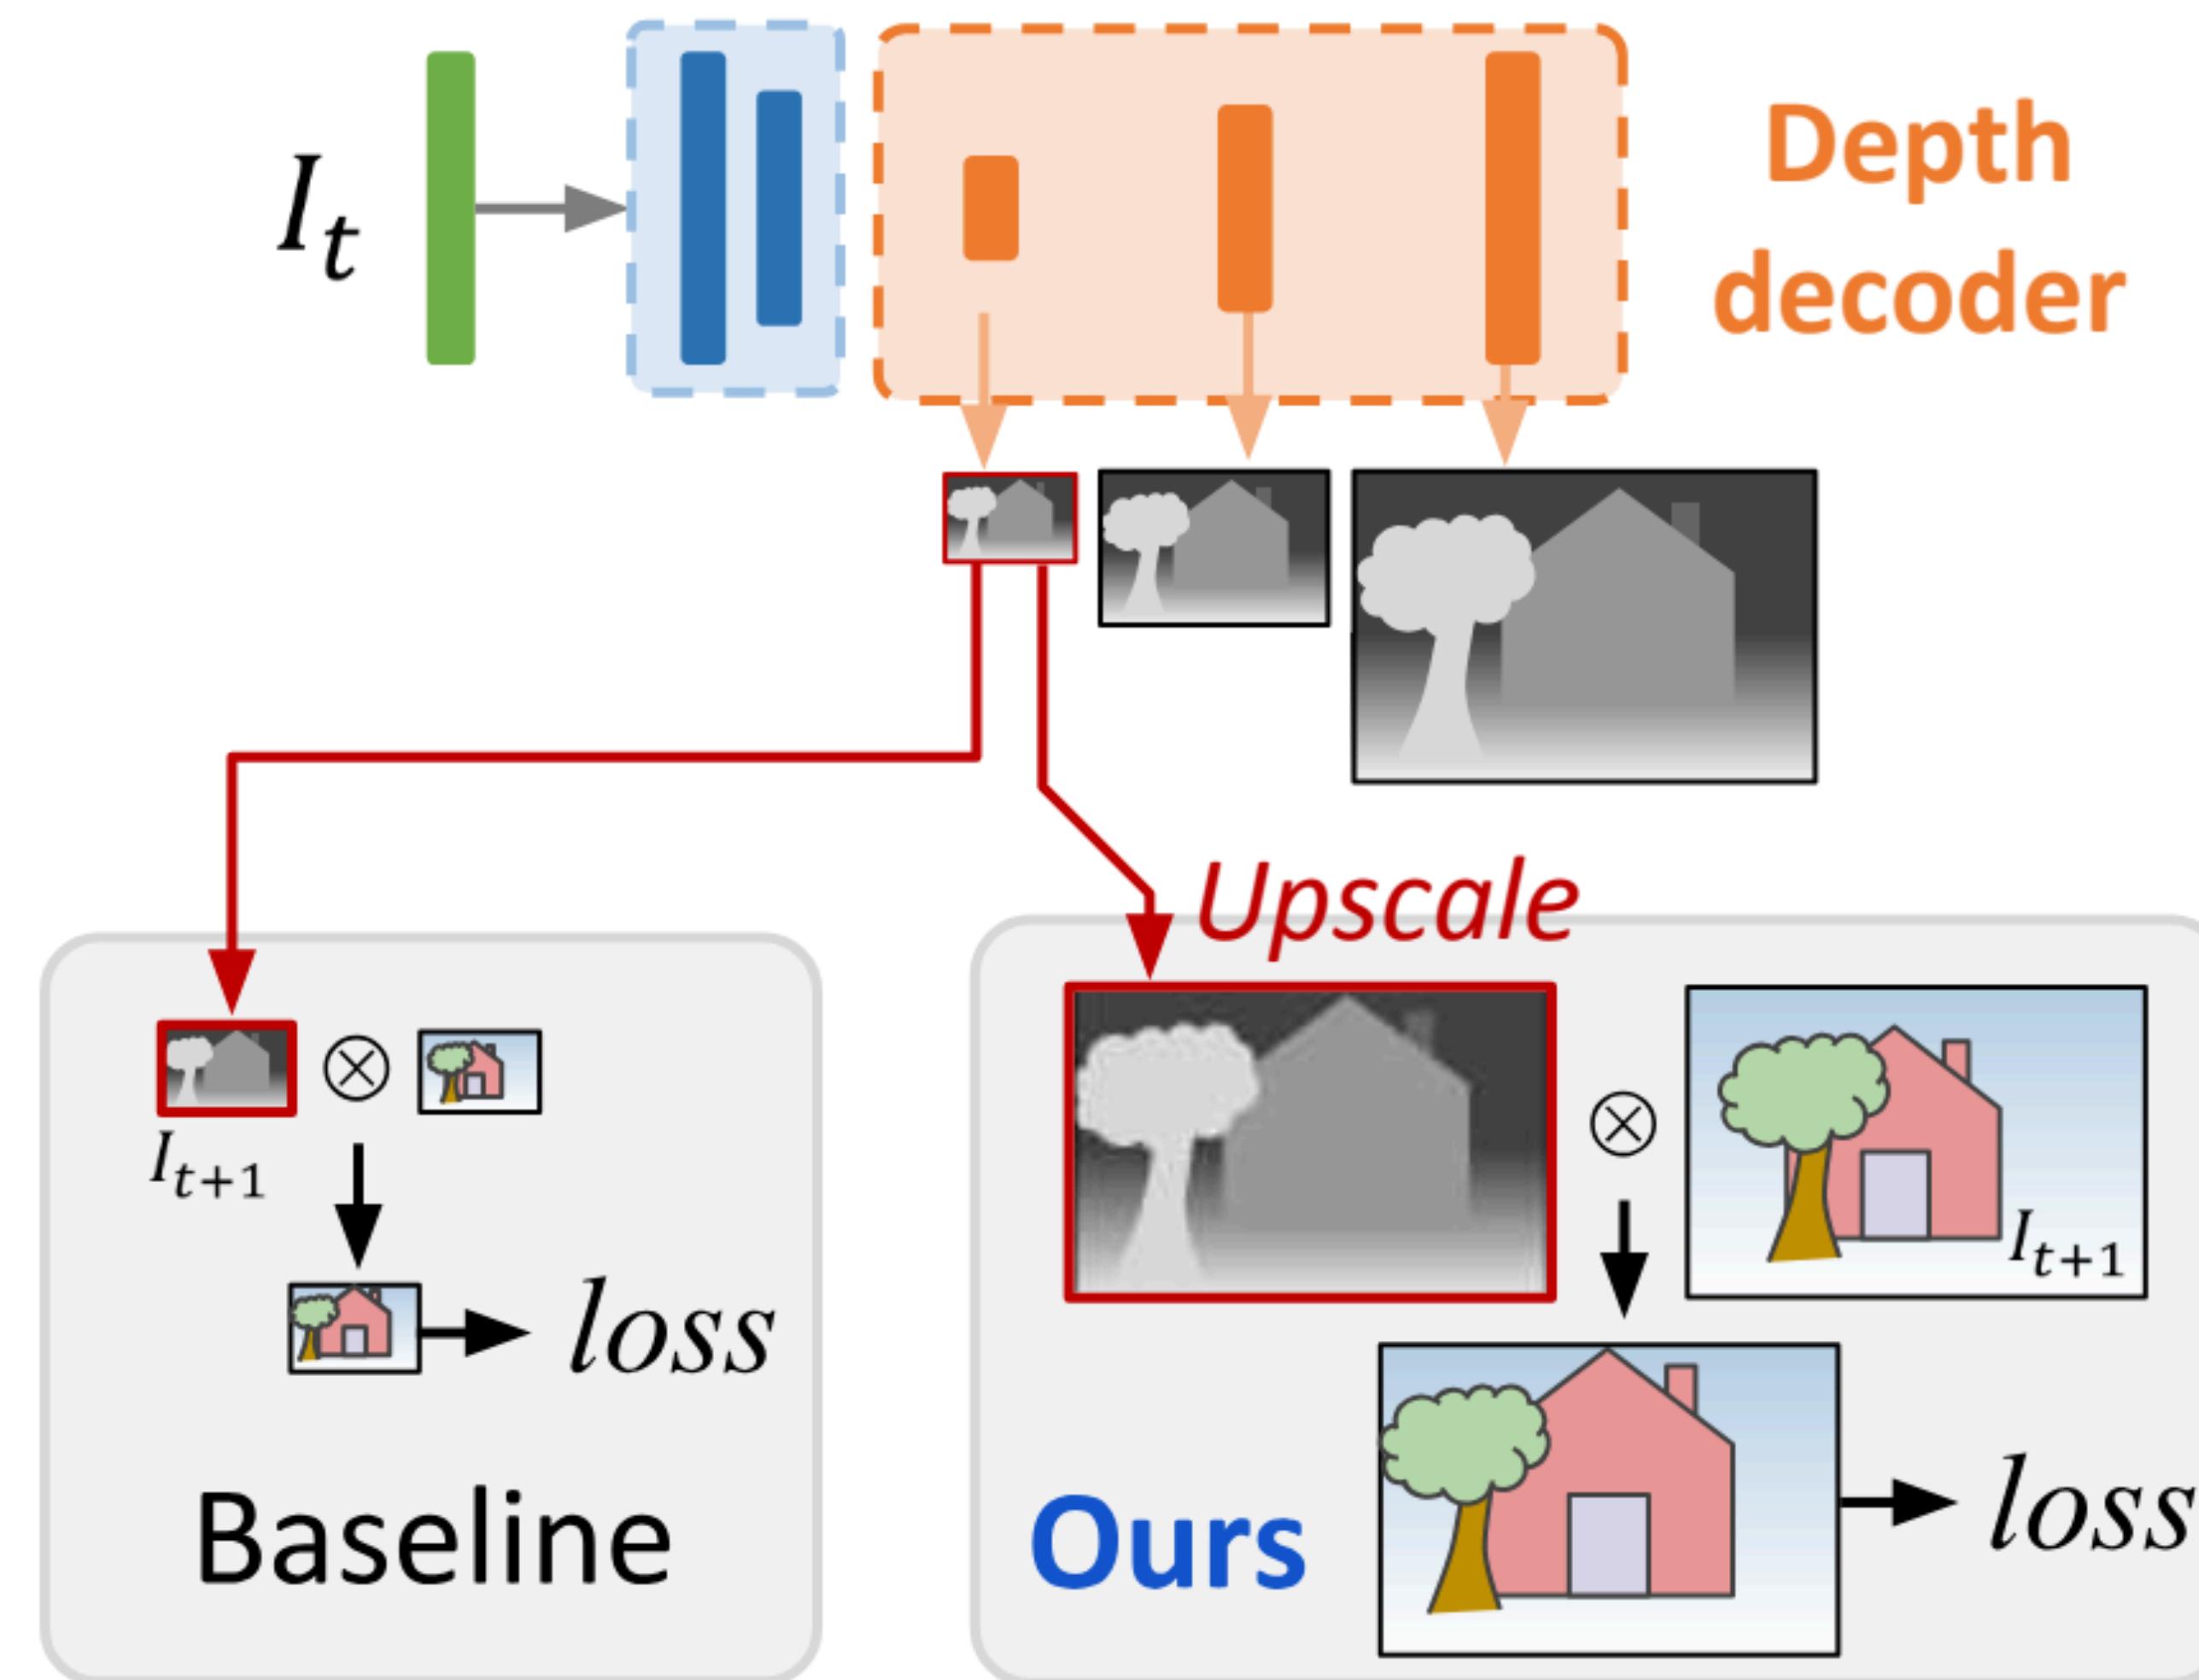
We observe that pixels which remain the same between adjacent frames in the sequence often indicate a static camera, an object moving at equivalent relative translation to the camera, or a low texture region. We therefore set  $\mu$  to only include the loss of pixels where the reprojection error of the warped image  $I_{t' \rightarrow t}$  is lower than that of the original, unwarped source image  $I'_t$ , *i.e.*

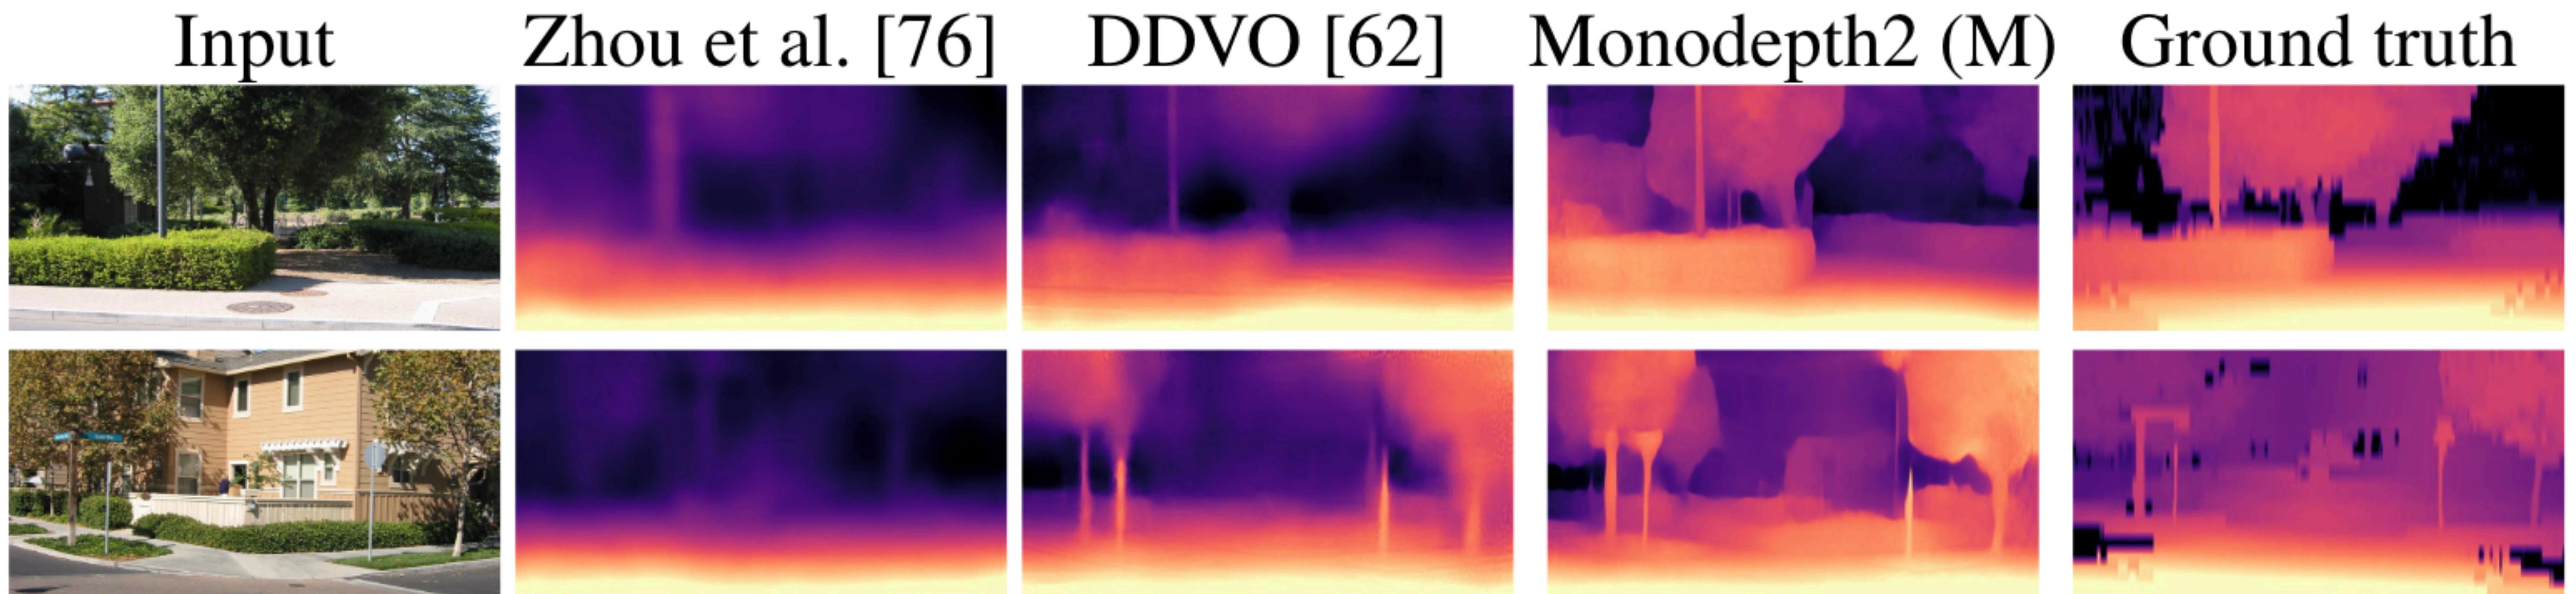
$$\mu = \left[ \min_{t'} pe(I_t, I_{t' \rightarrow t}) < \min_{t'} pe(I_t, I'_{t'}) \right], \quad (5)$$



**Figure 5. Auto-masking.** We show auto-masks computed after one epoch, where black pixels are removed from the loss (*i.e.*  $\mu = 0$ ). The mask prevents objects moving at similar speeds to the camera (top) and whole frames where the camera is static (bottom) from contaminating the loss. The mask is computed from the input frames and network predictions using Eqn. 5.

# Better multi-scale loss



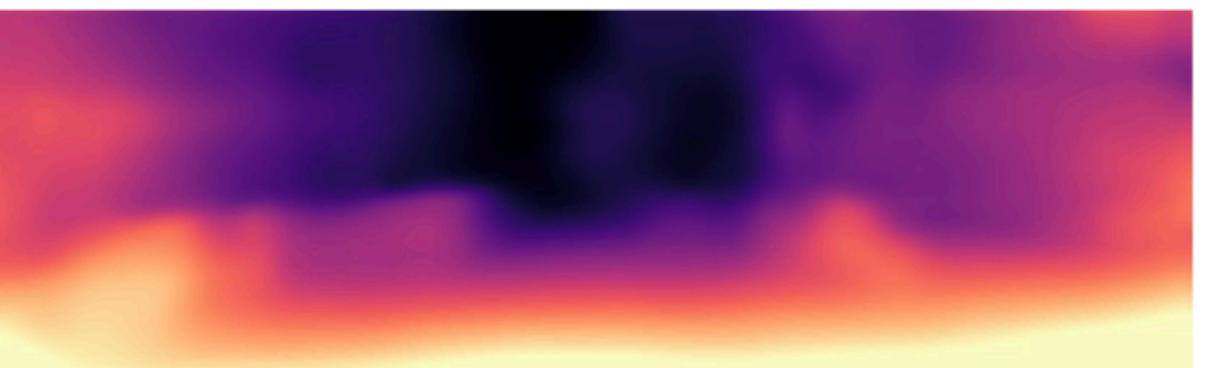
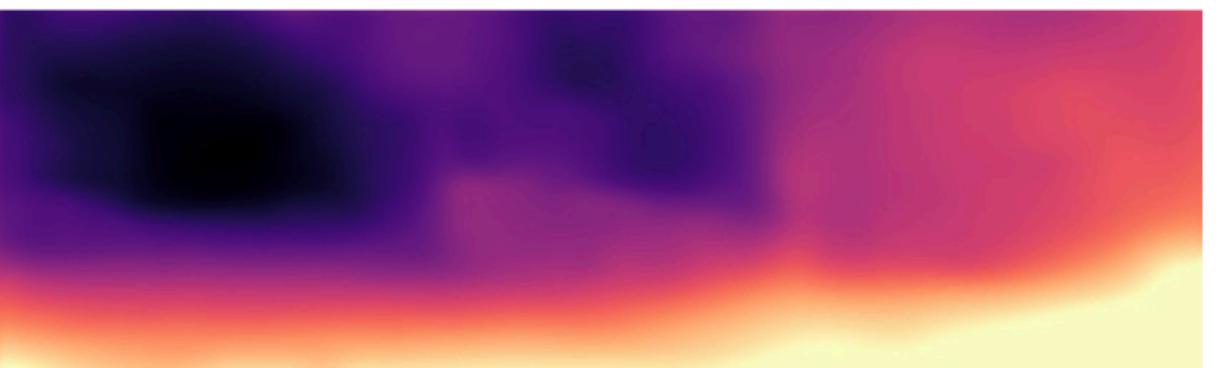
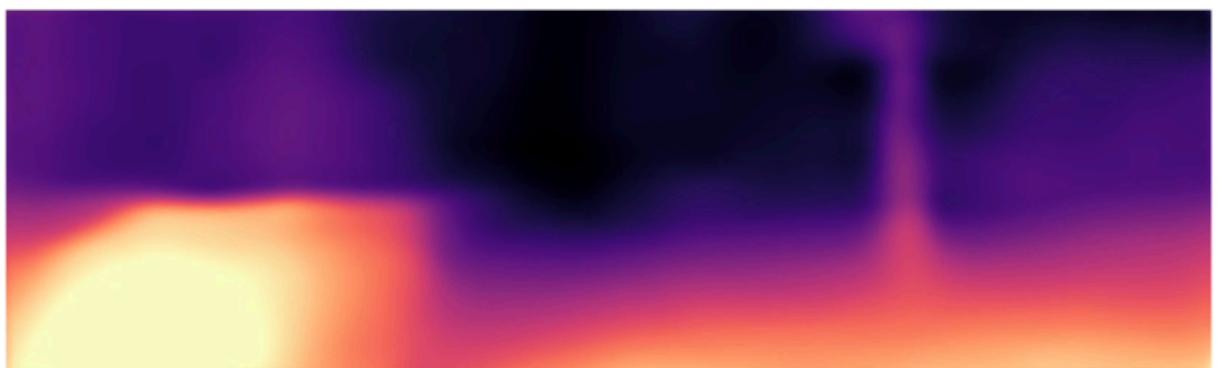
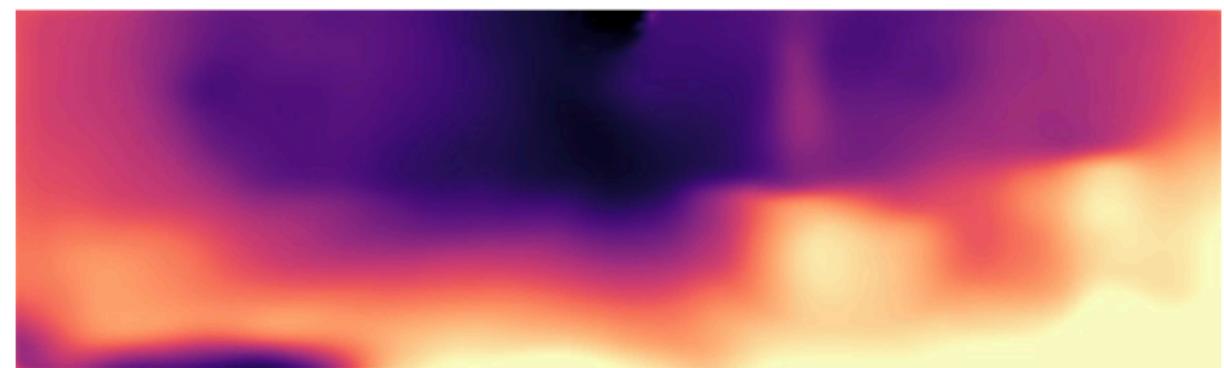


**Figure 6. Qualitative Make3D results.** All methods were trained on KITTI using monocular supervision.

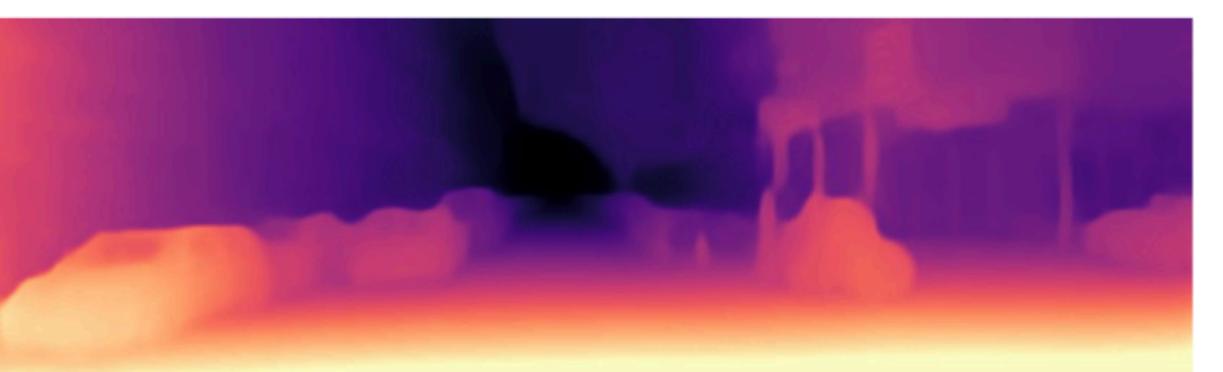
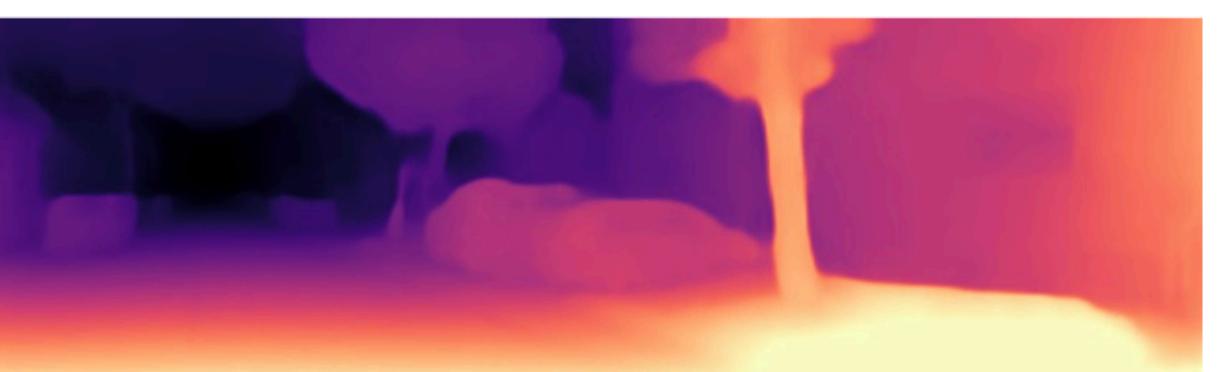
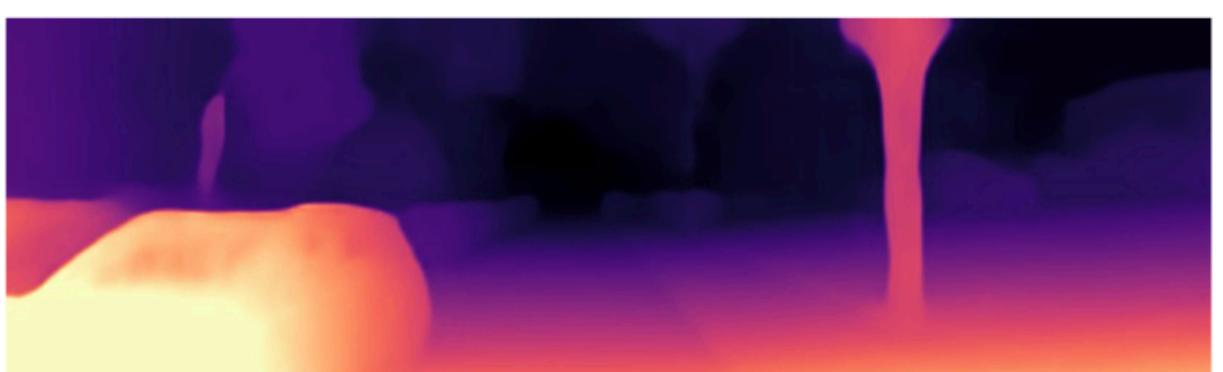
Input

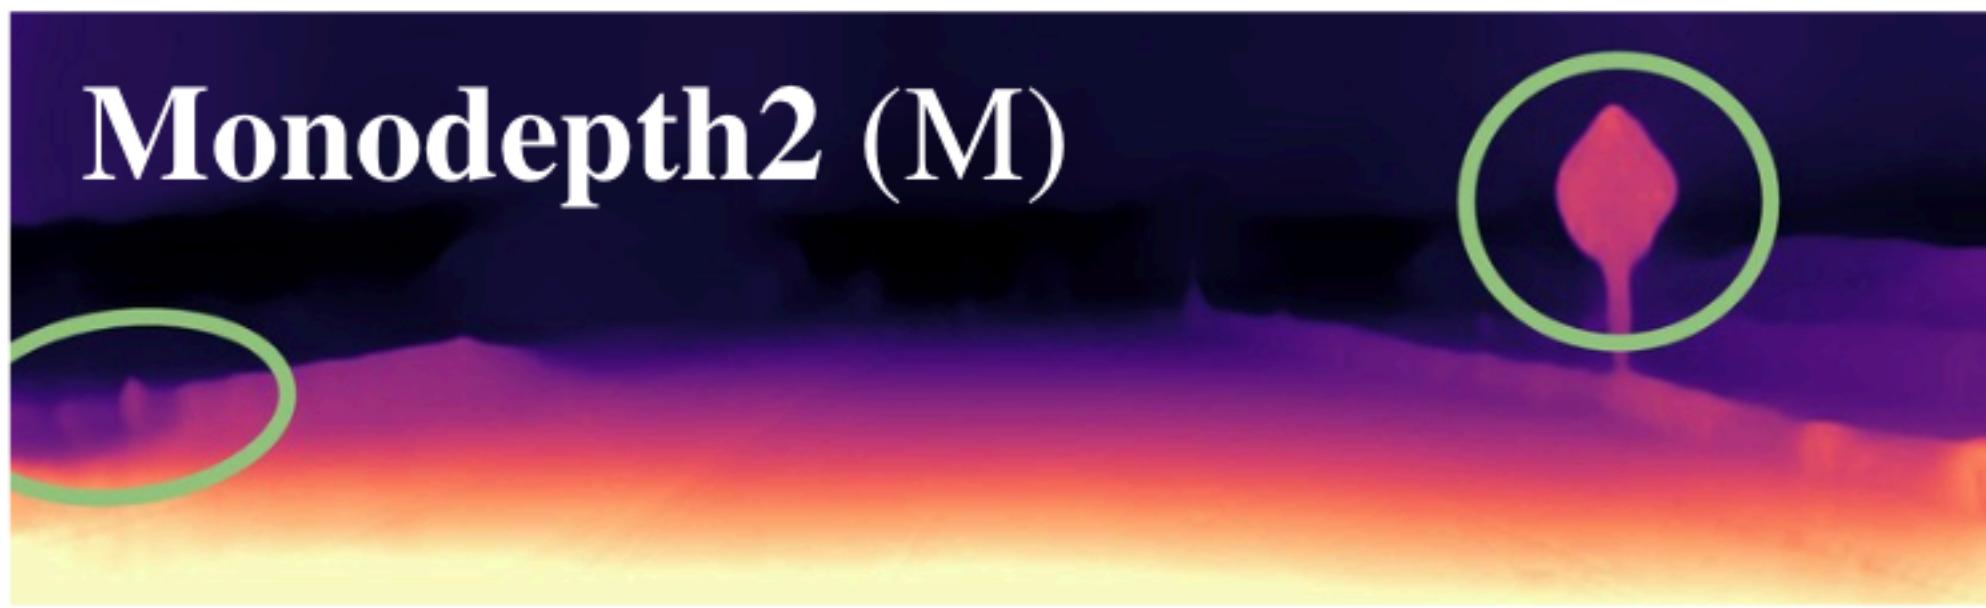


Zhou et al. [76]



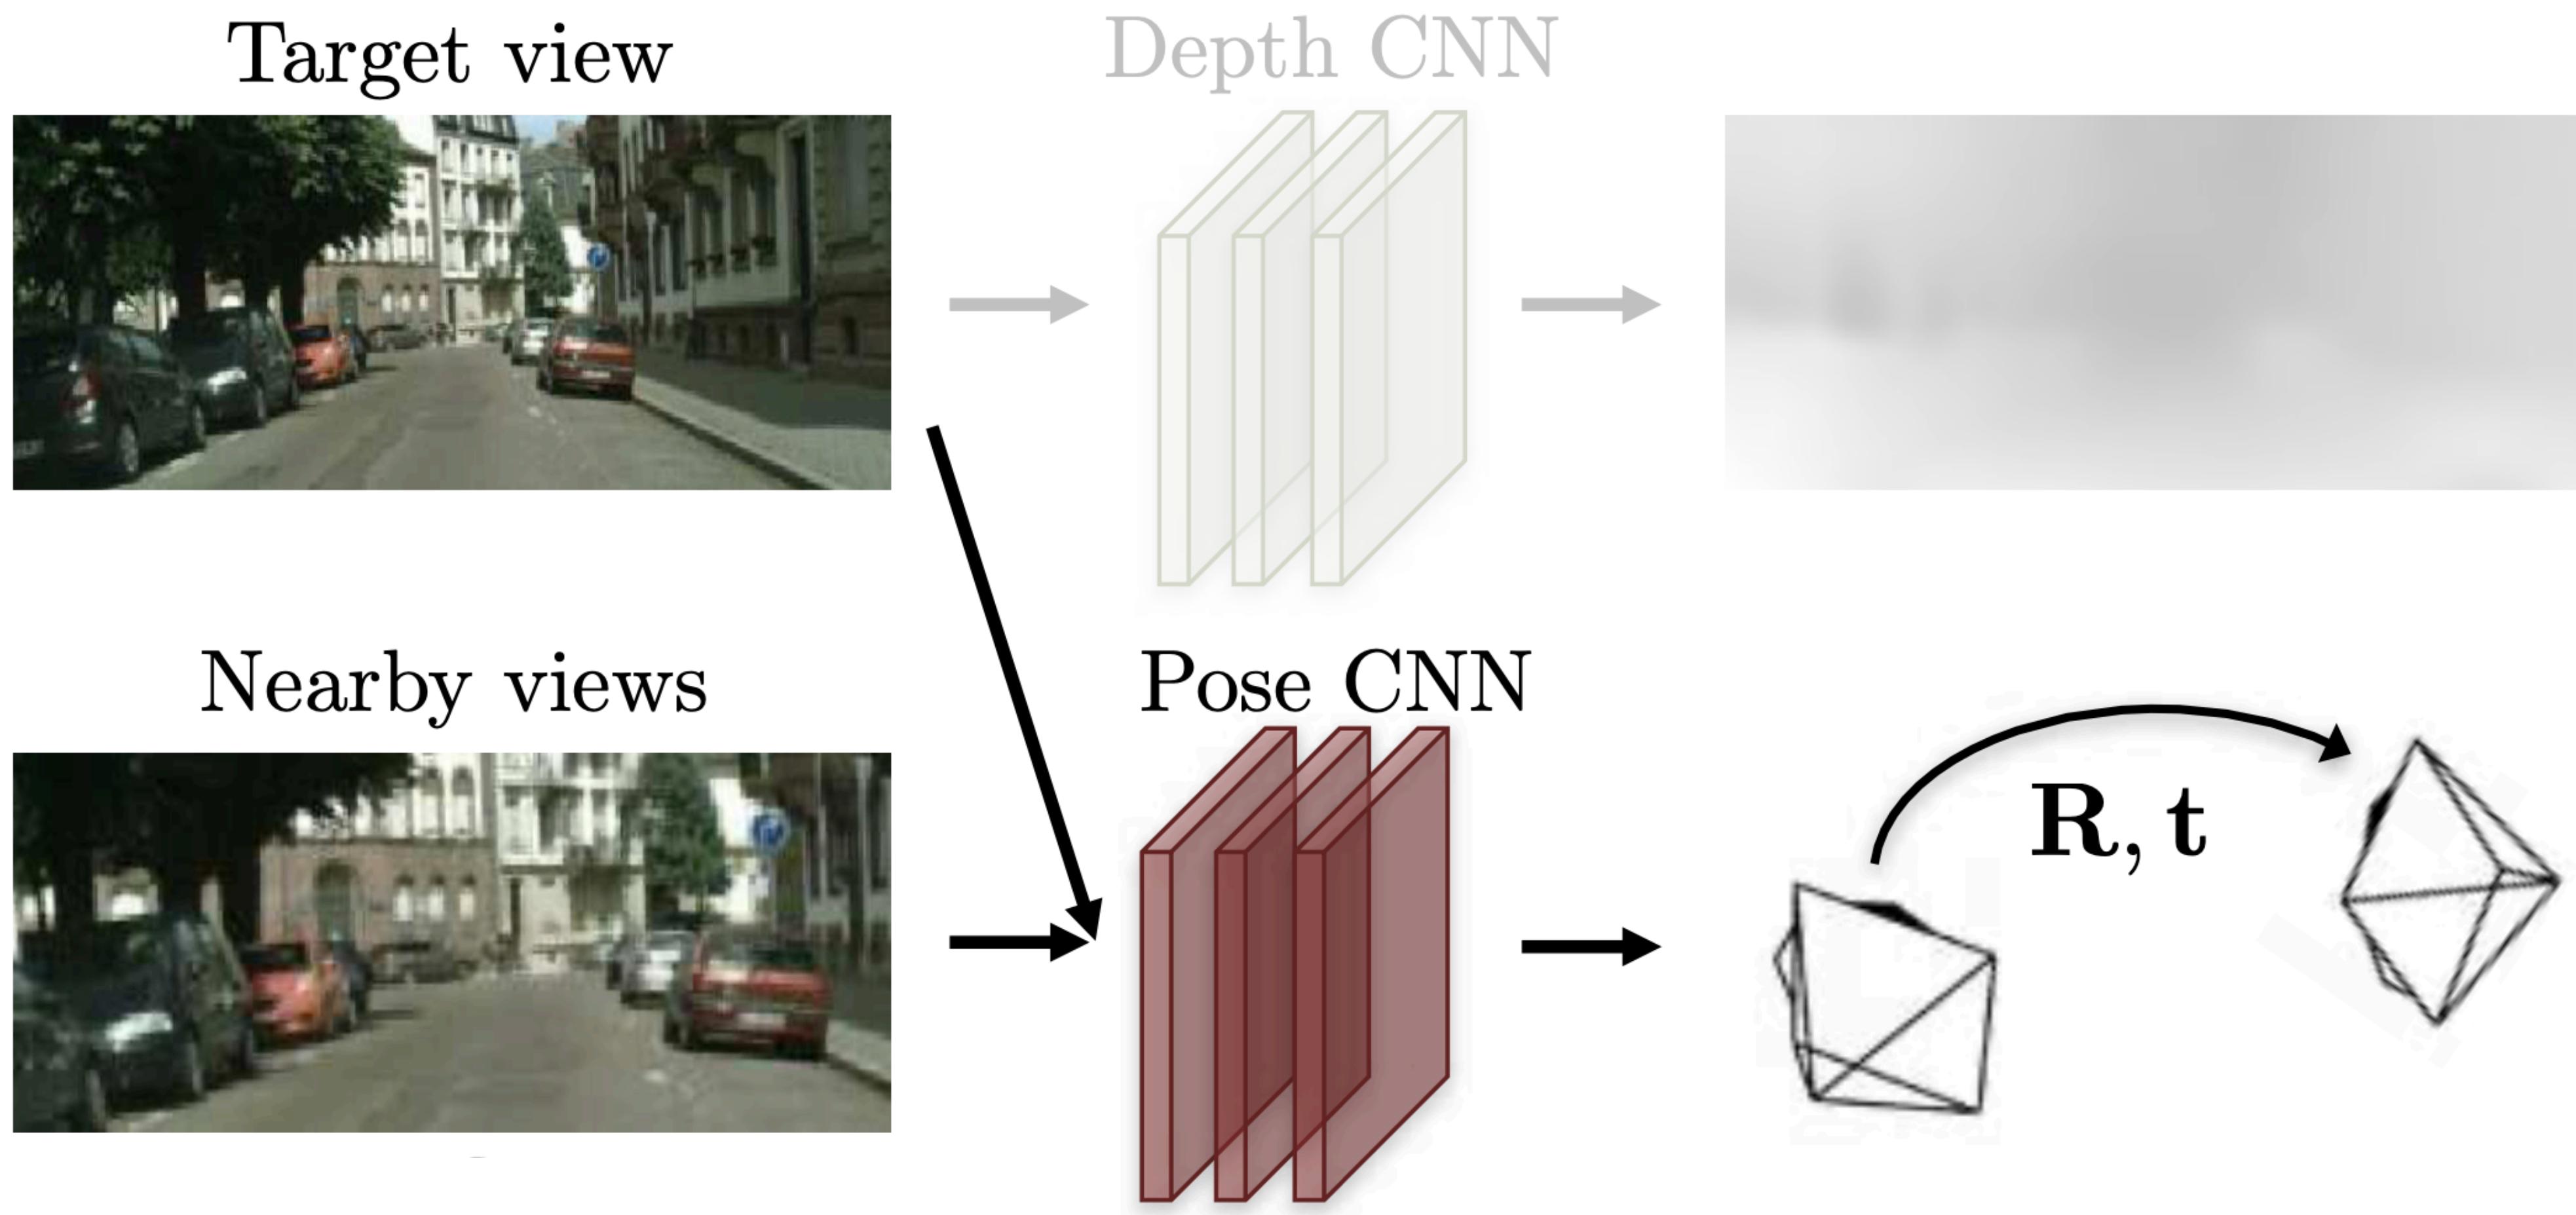
MD2 MS





**Figure 8. Failure cases.** **Top:** Our self-supervised loss fails to learn good depths for distorted, reflective and color-saturated regions. **Bottom:** We can fail to accurately delineate objects where boundaries are ambiguous (left) or shapes are intricate (right).

# Problem: Only works for “simple” camera trajectories :/



CNNs are bad at predicting poses: no single pixel neighborhood informs completely about pose...

3D Scene

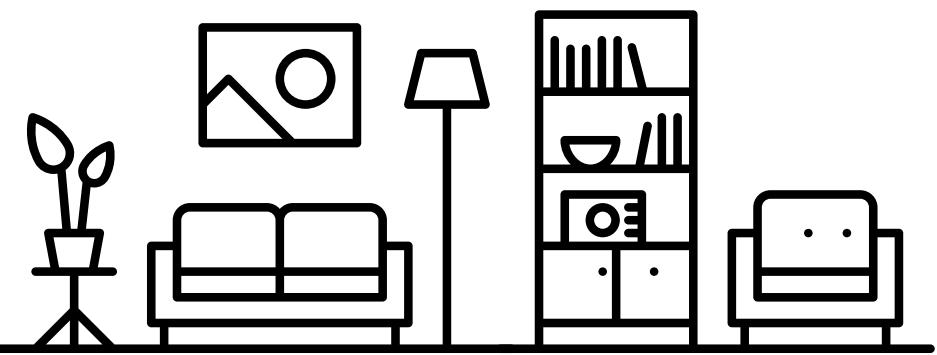


Image  
Formation

Graphics

3D Scene

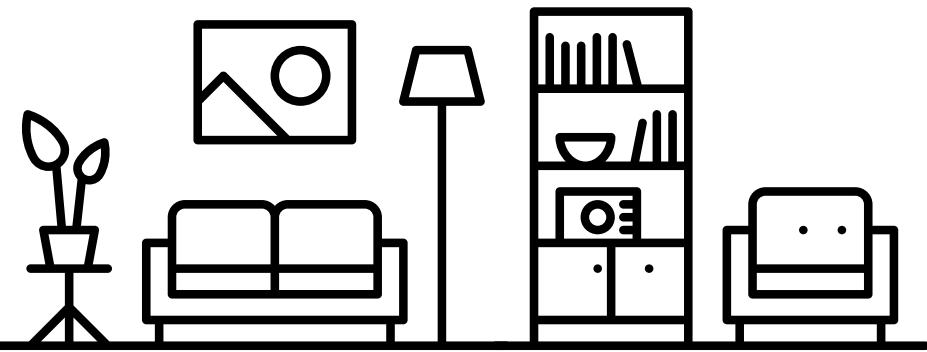


Image  
Formation

Neural Scene  
Representation

**Graphics**

3D Scene

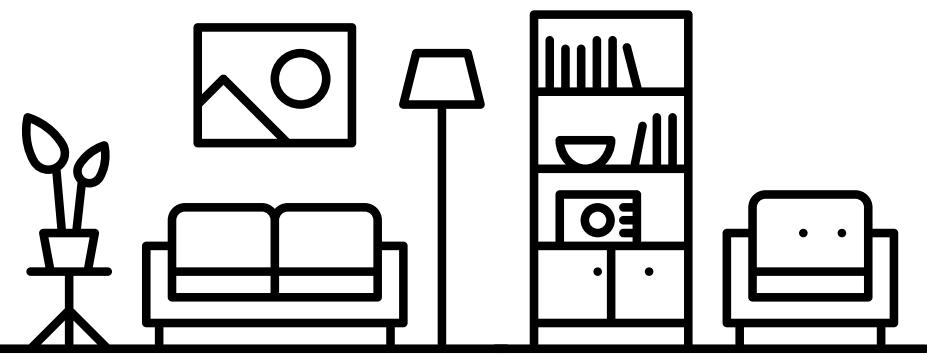


Image  
Formation

Inference

Neural Scene  
Representation

**Graphics**

3D Scene

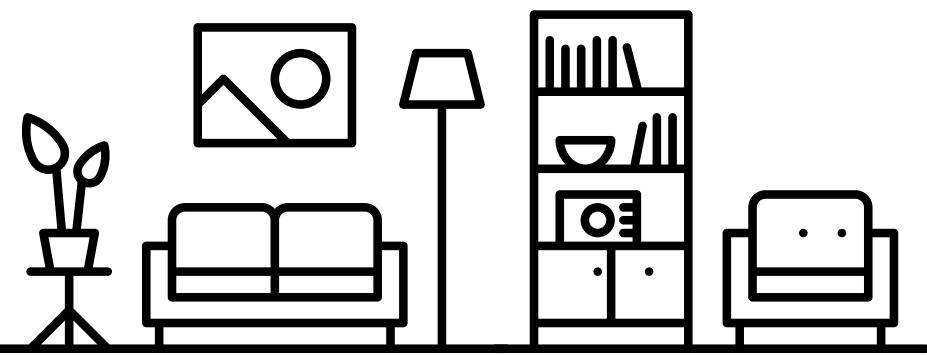


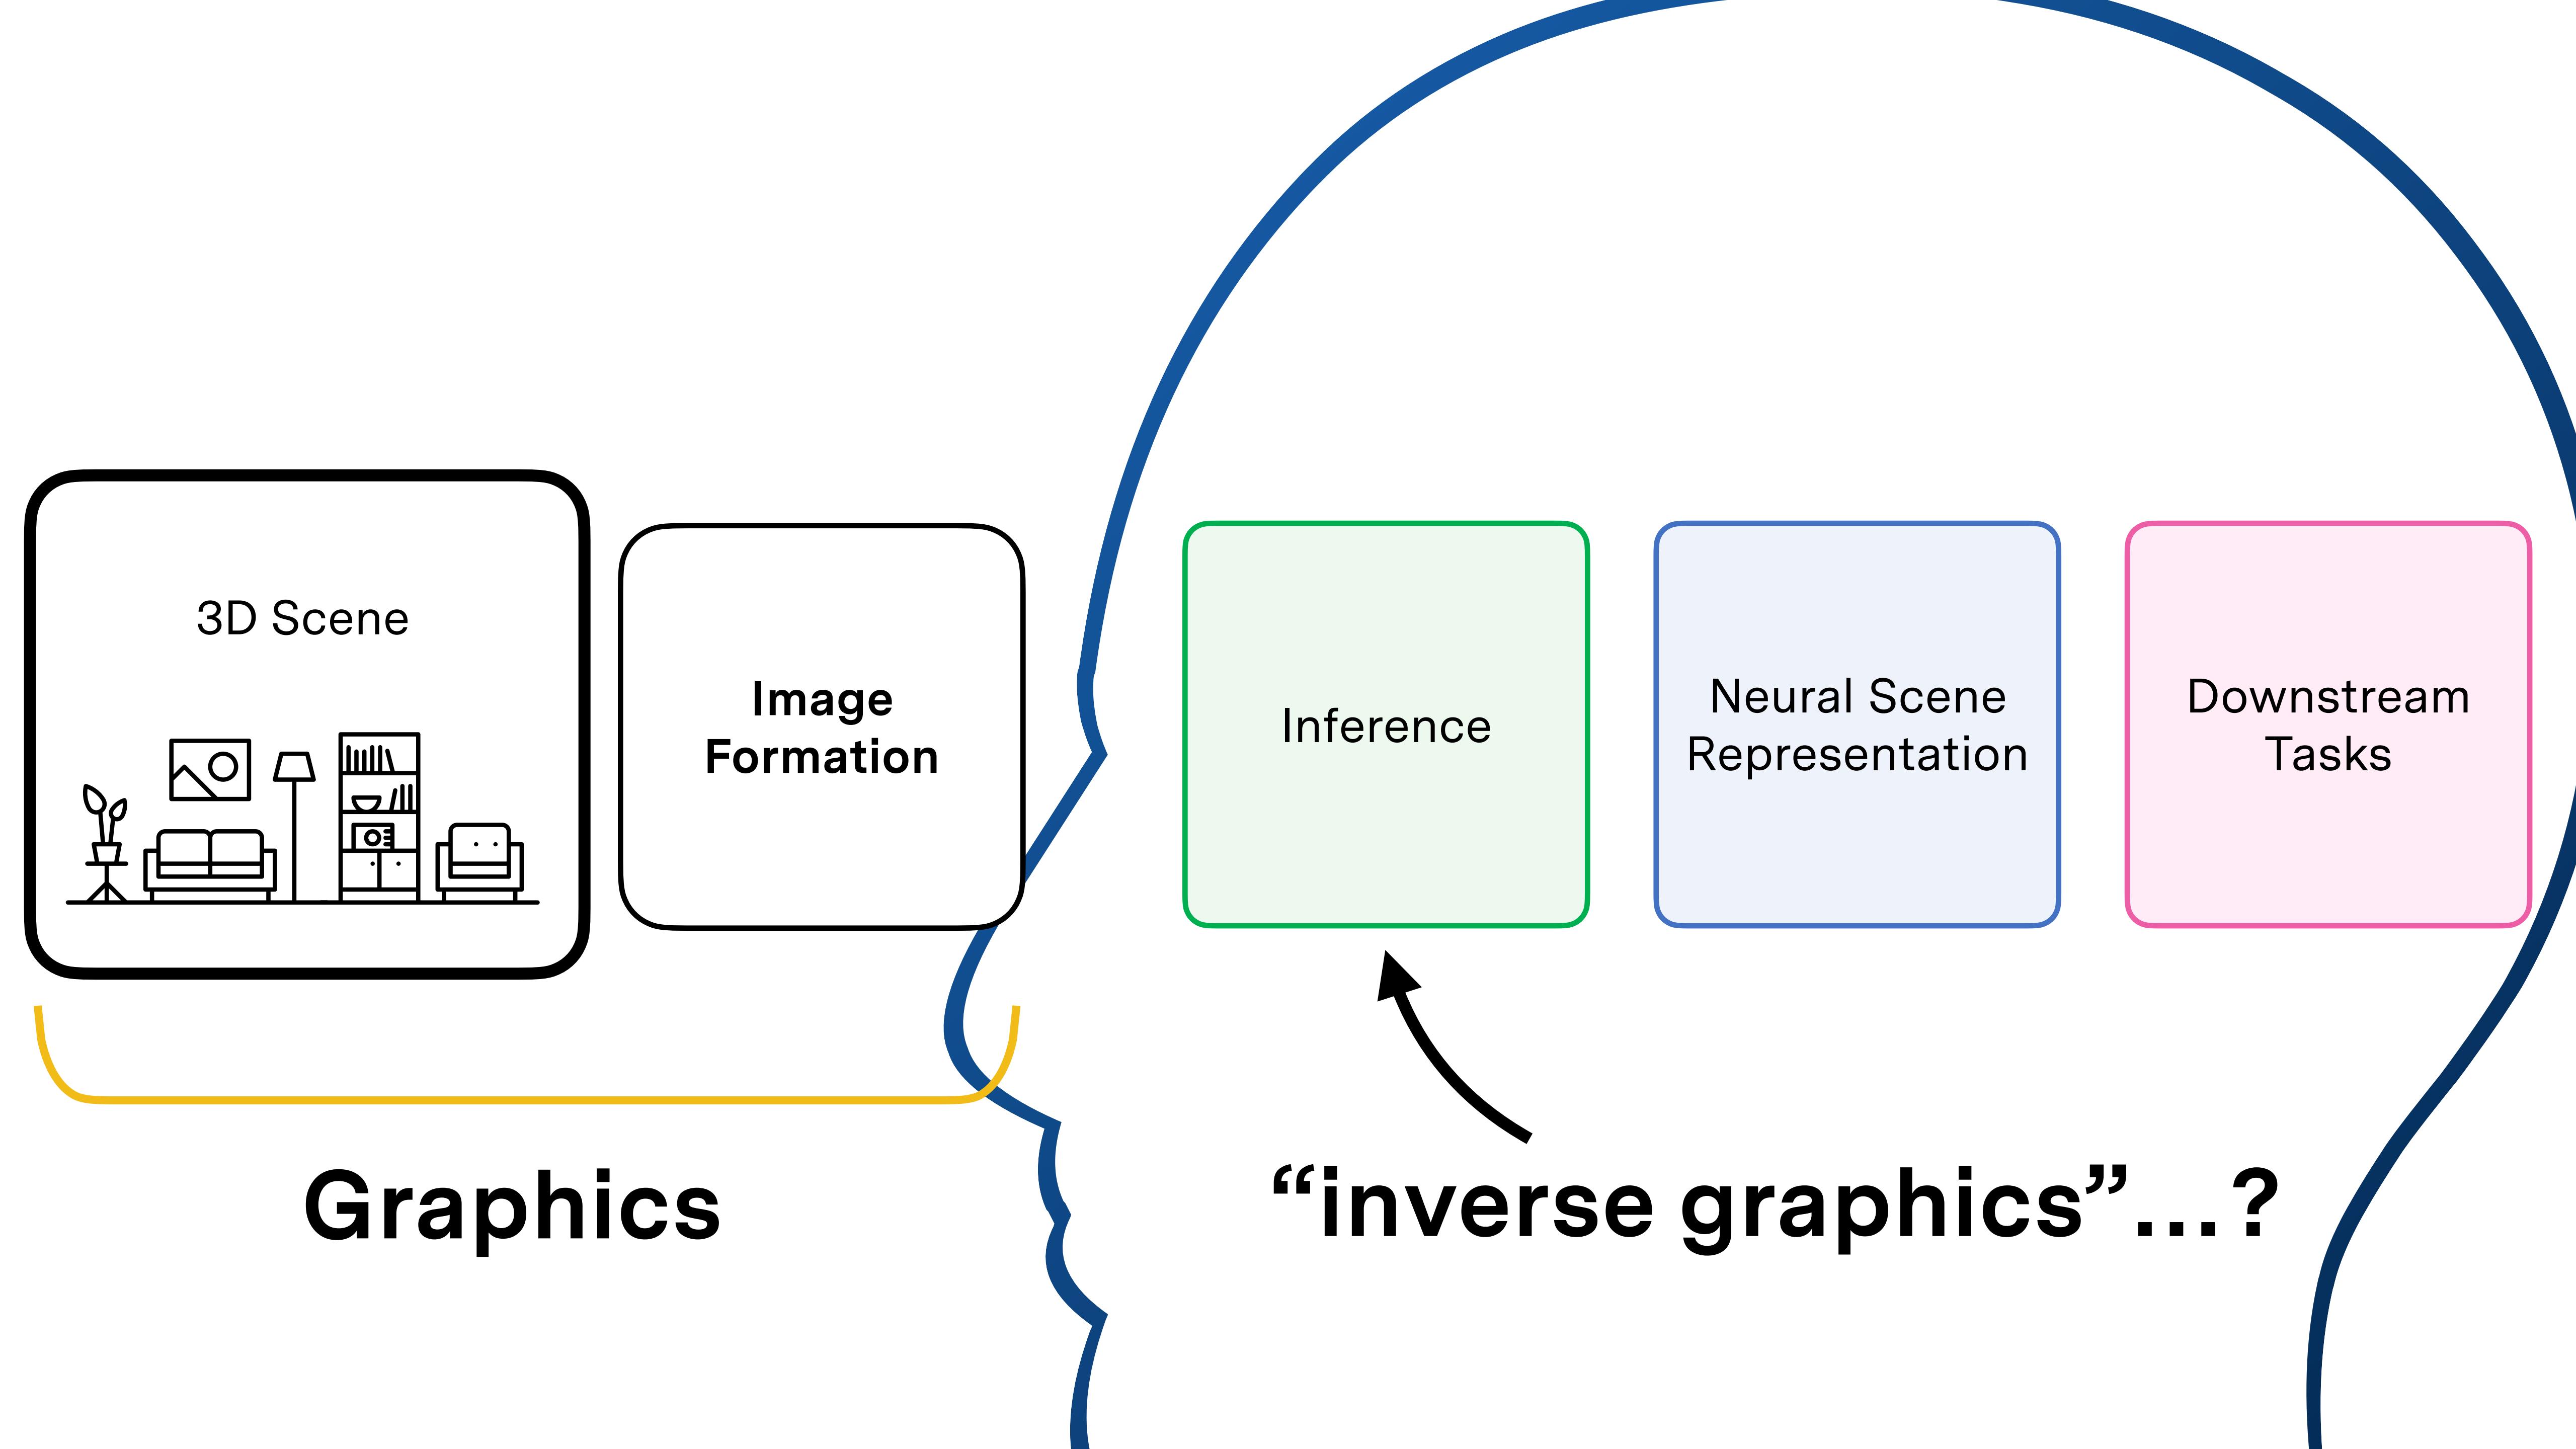
Image Formation

**Graphics**

Inference

Neural Scene Representation

“inverse graphics”...?



# Today: How to computationally represent 3D scenes.

3D Scene

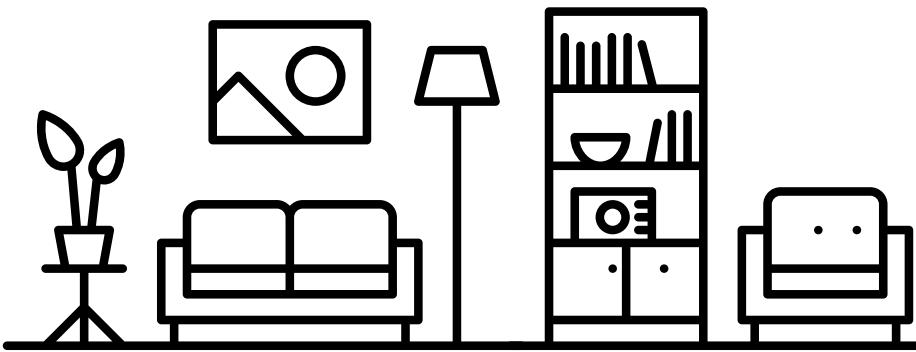


Image  
Formation

Inference

Neural Scene  
Representation

Downstream  
Tasks

Why?

At the end of the day, we want to make predictions about 3D scenes.  
For that, we need to know how we can represent 3D scenes computationally.

What you'll  
learn.

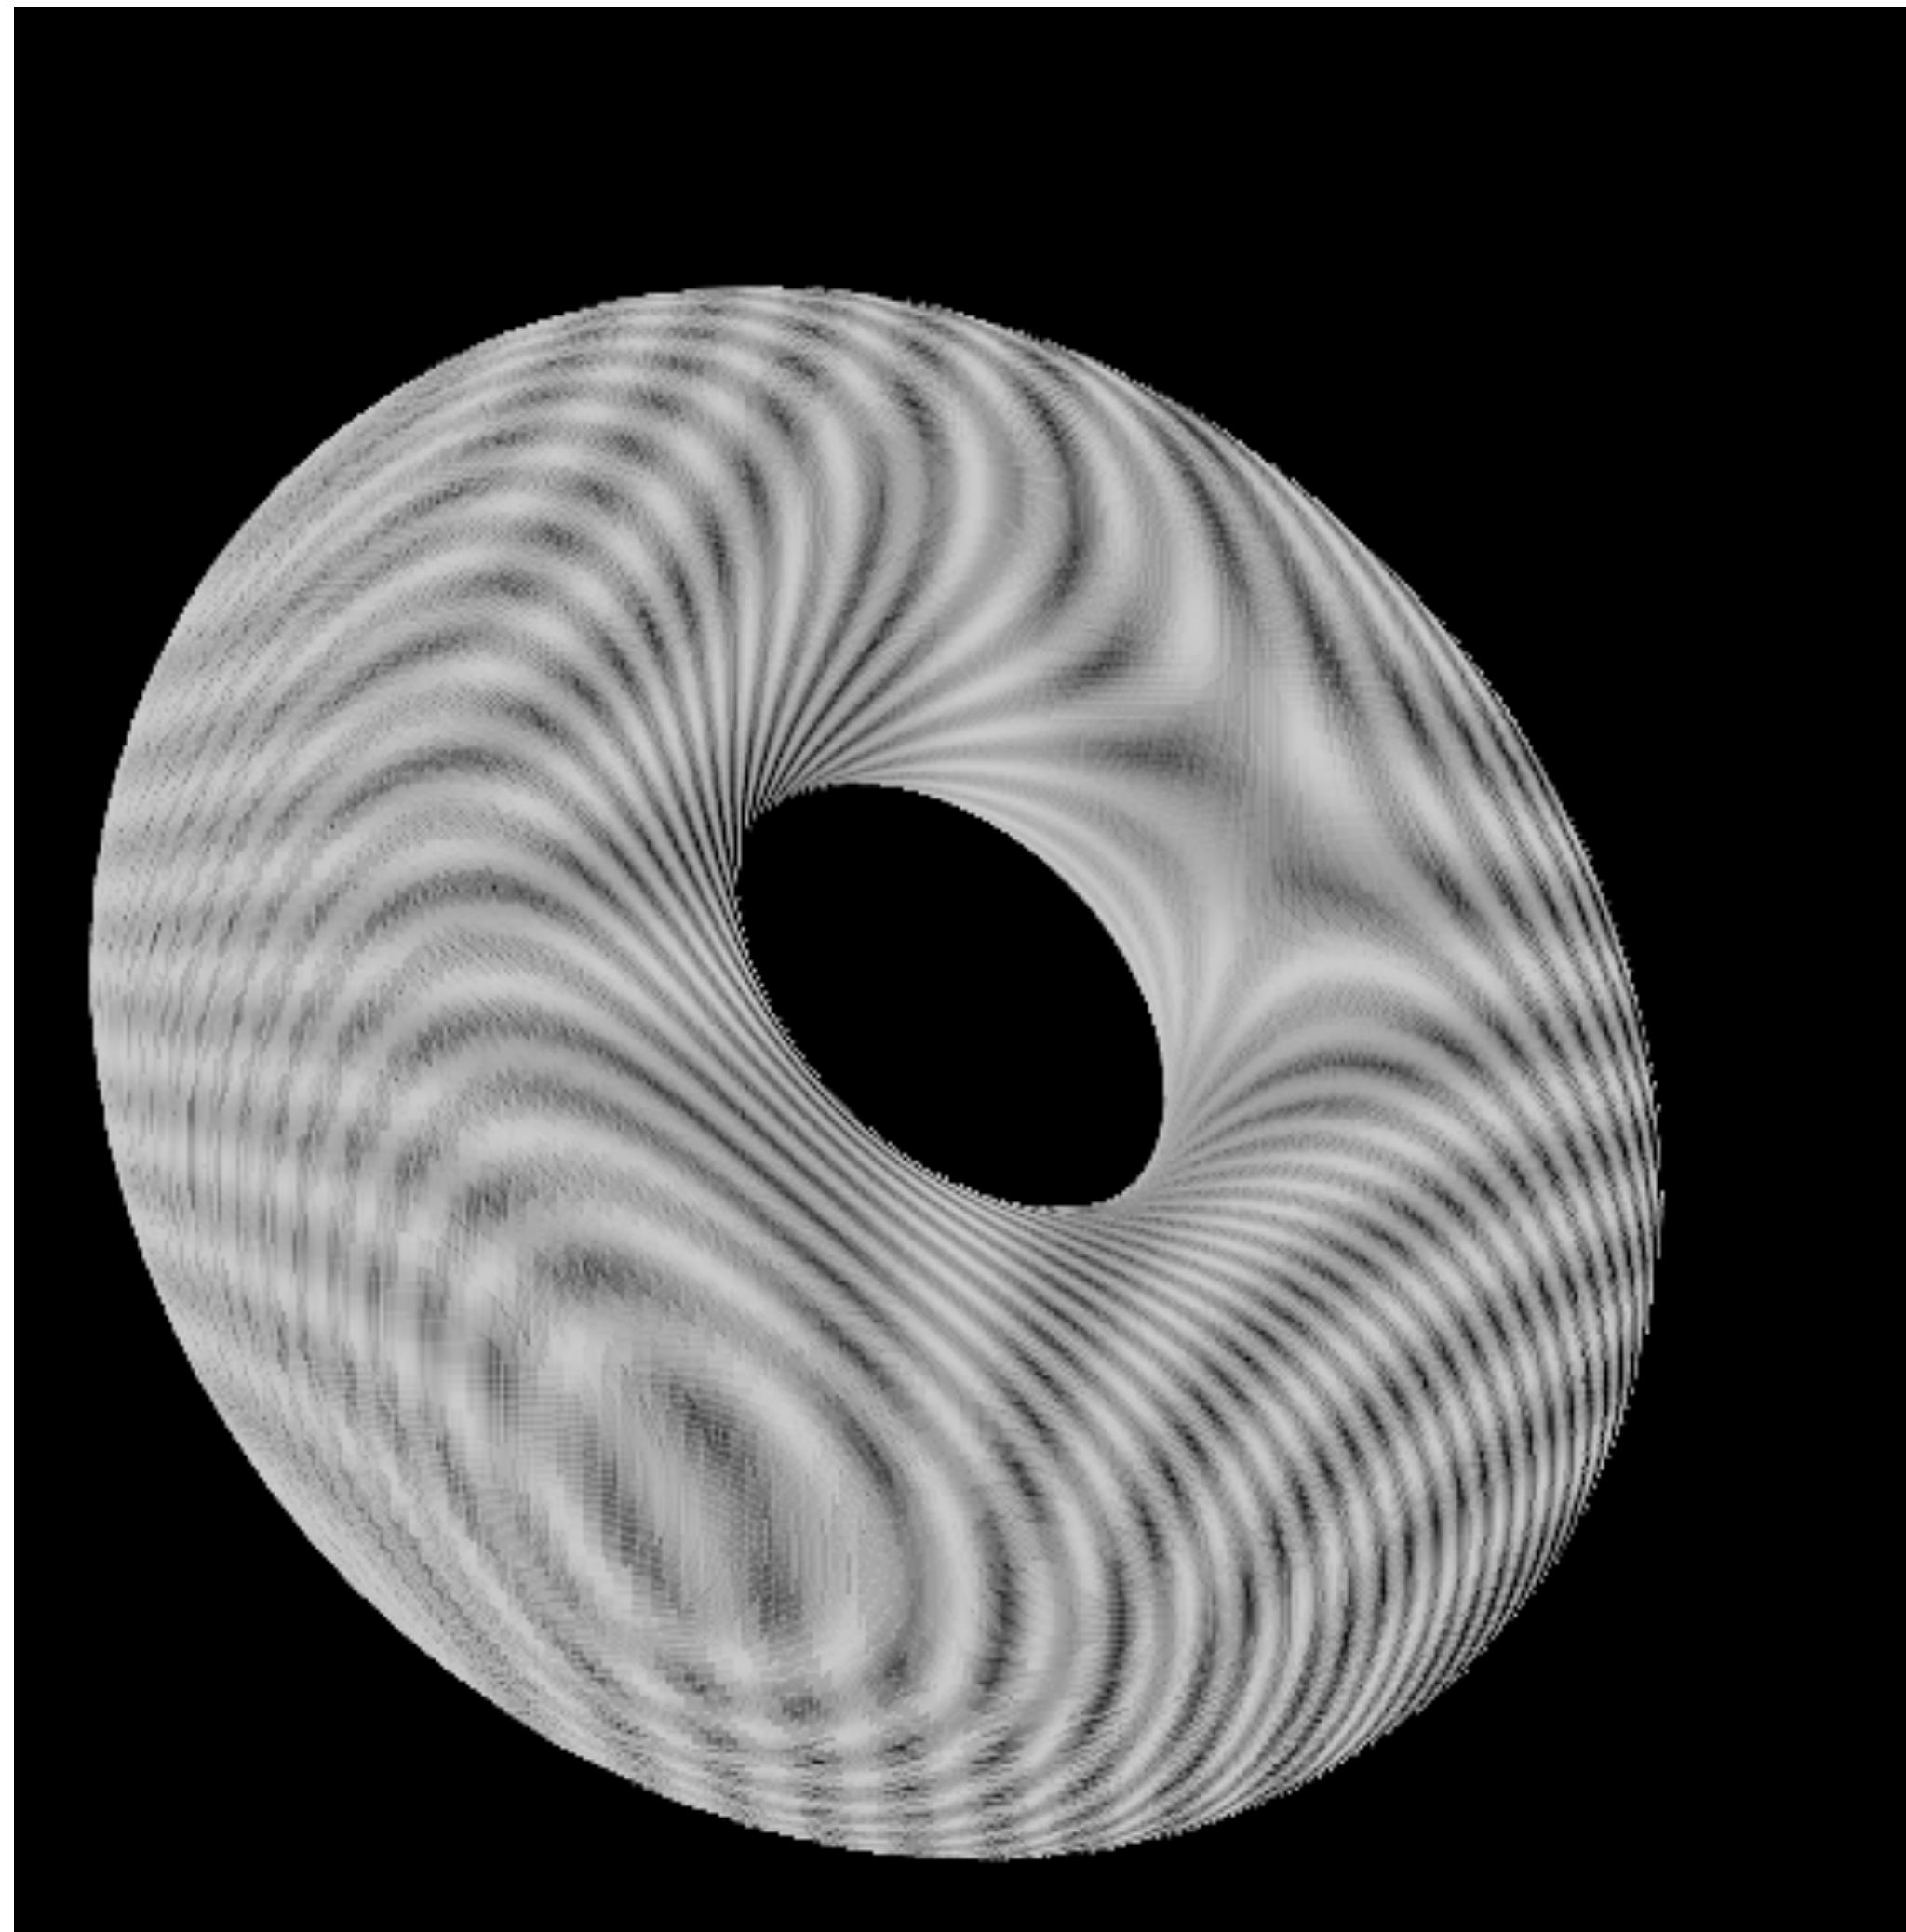
Surface-based representations, volumetric representations, discrete representations, continuous representations.

# Some Slides adapted from...

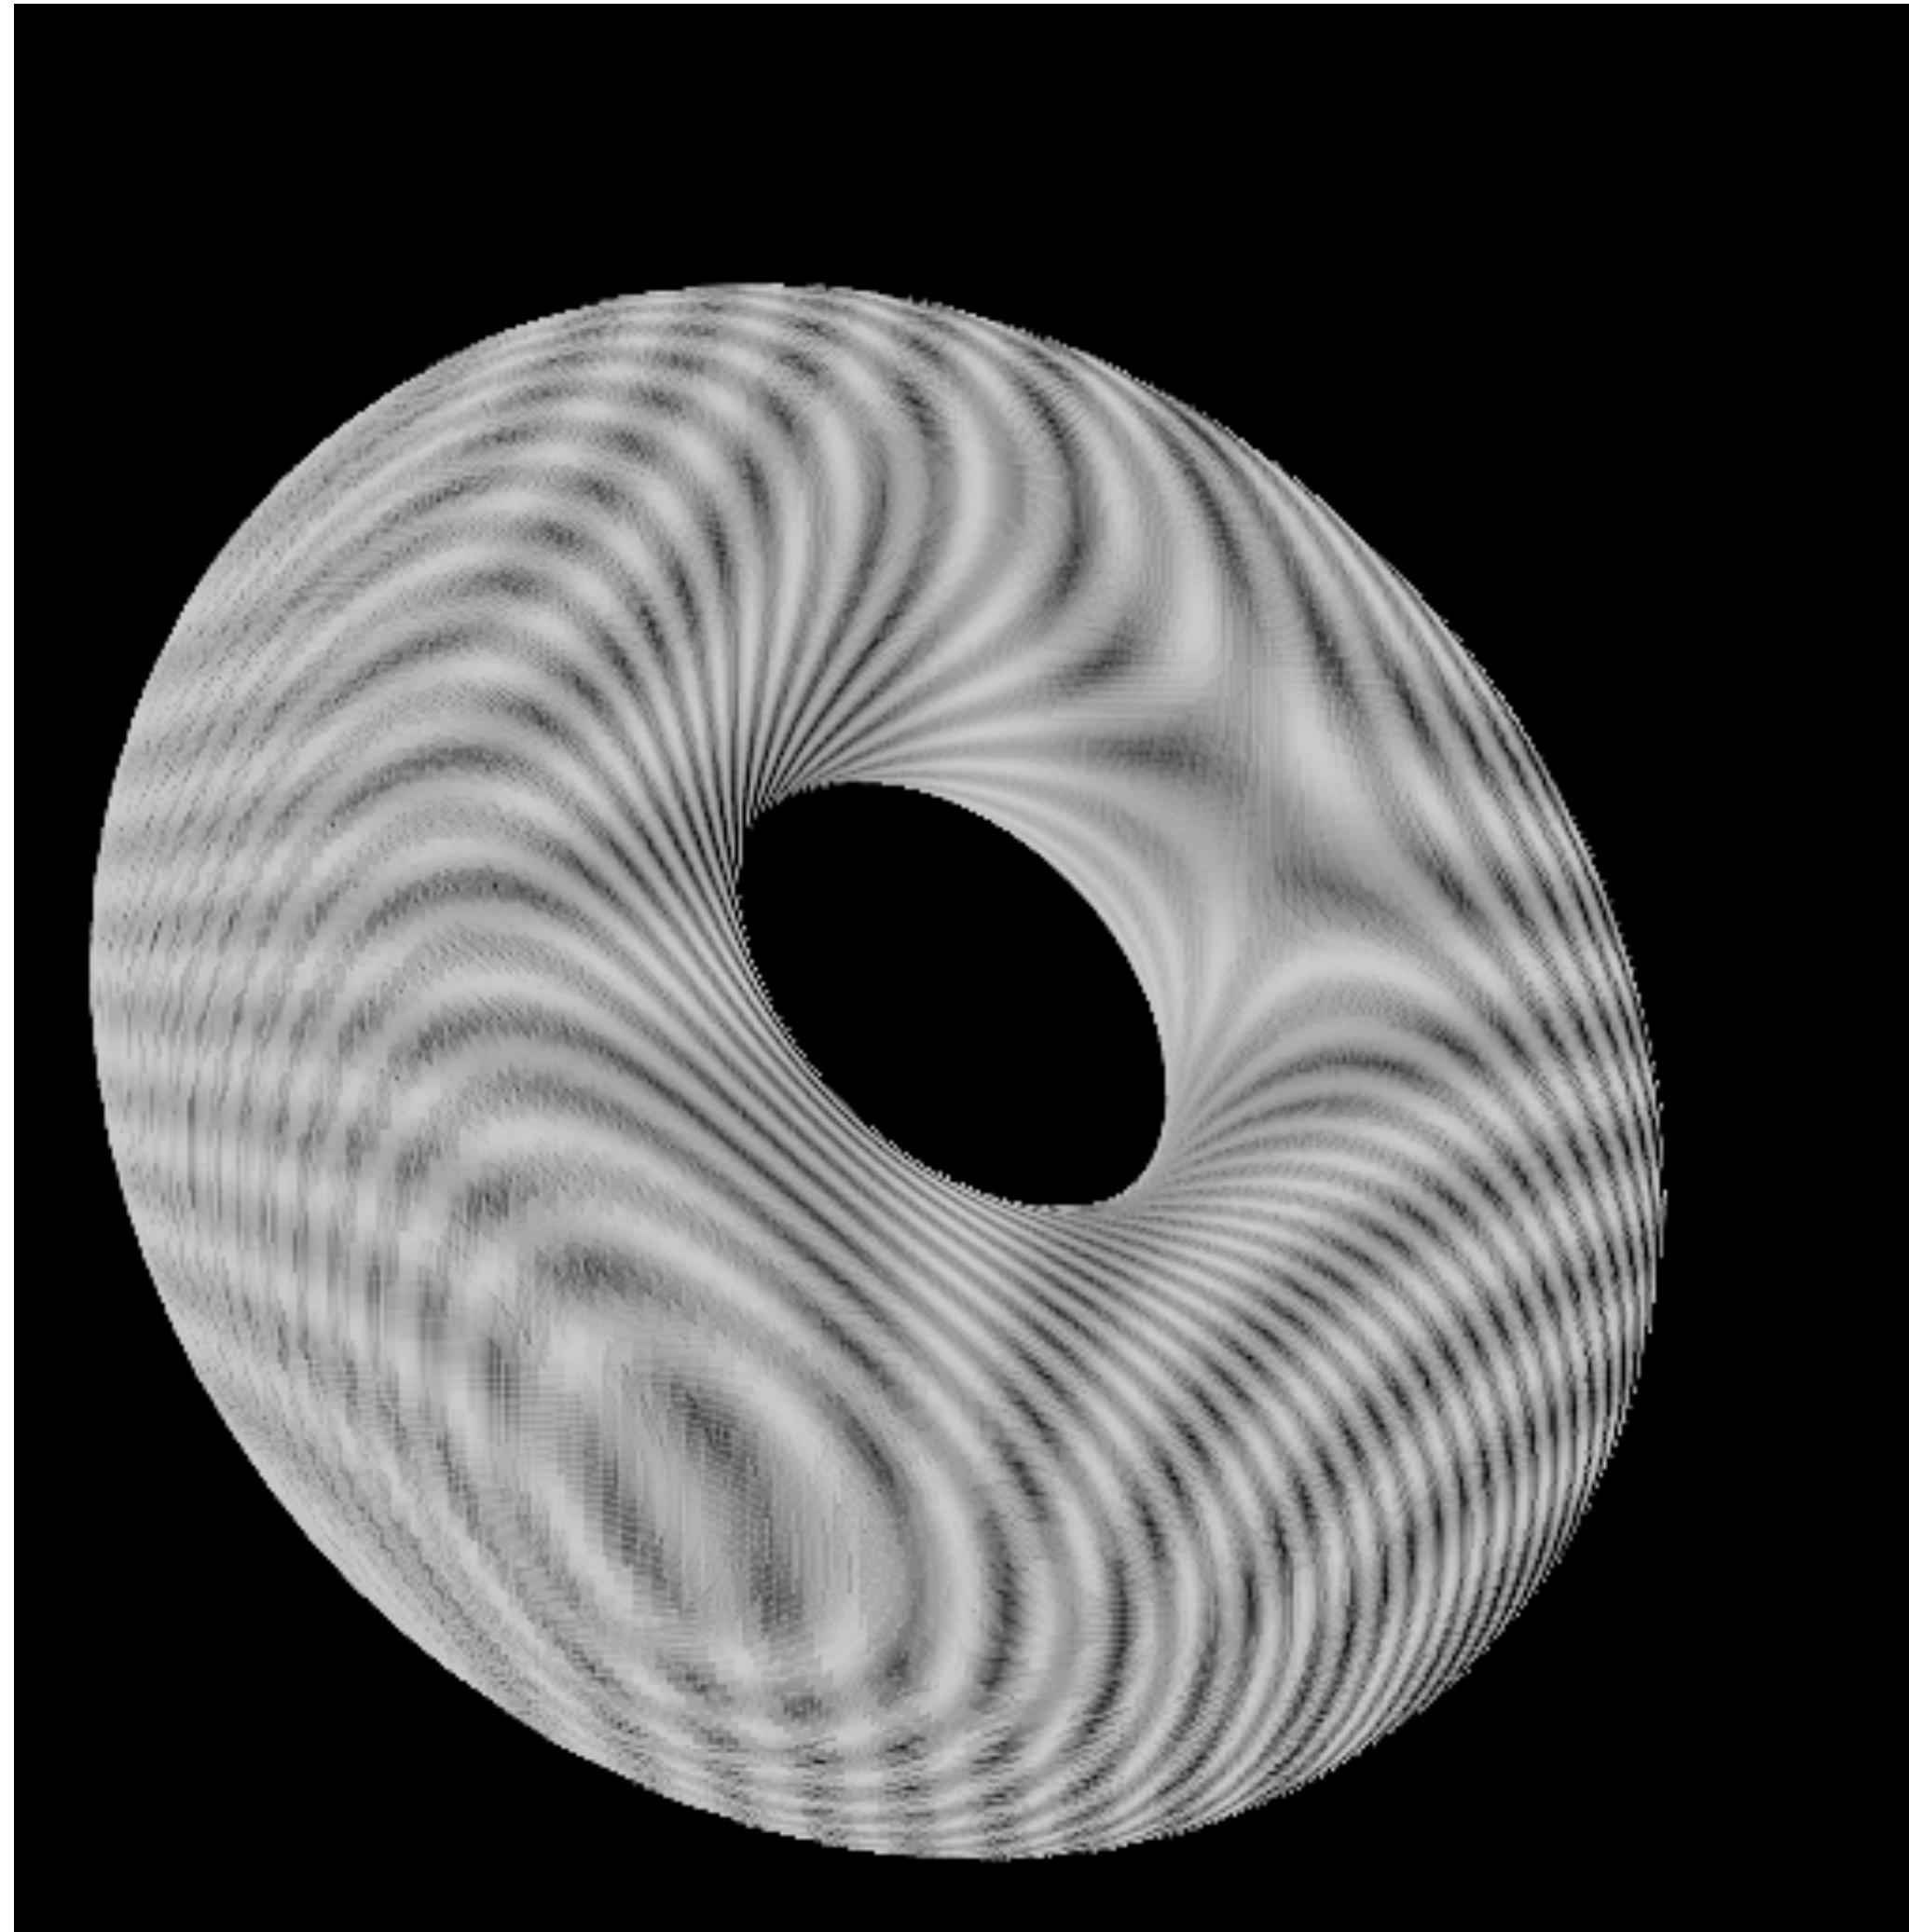
CMU 16-889: Learning for 3D Vision

Prof. Shubham Tulsiani

# Depth Maps – Representing the Visible

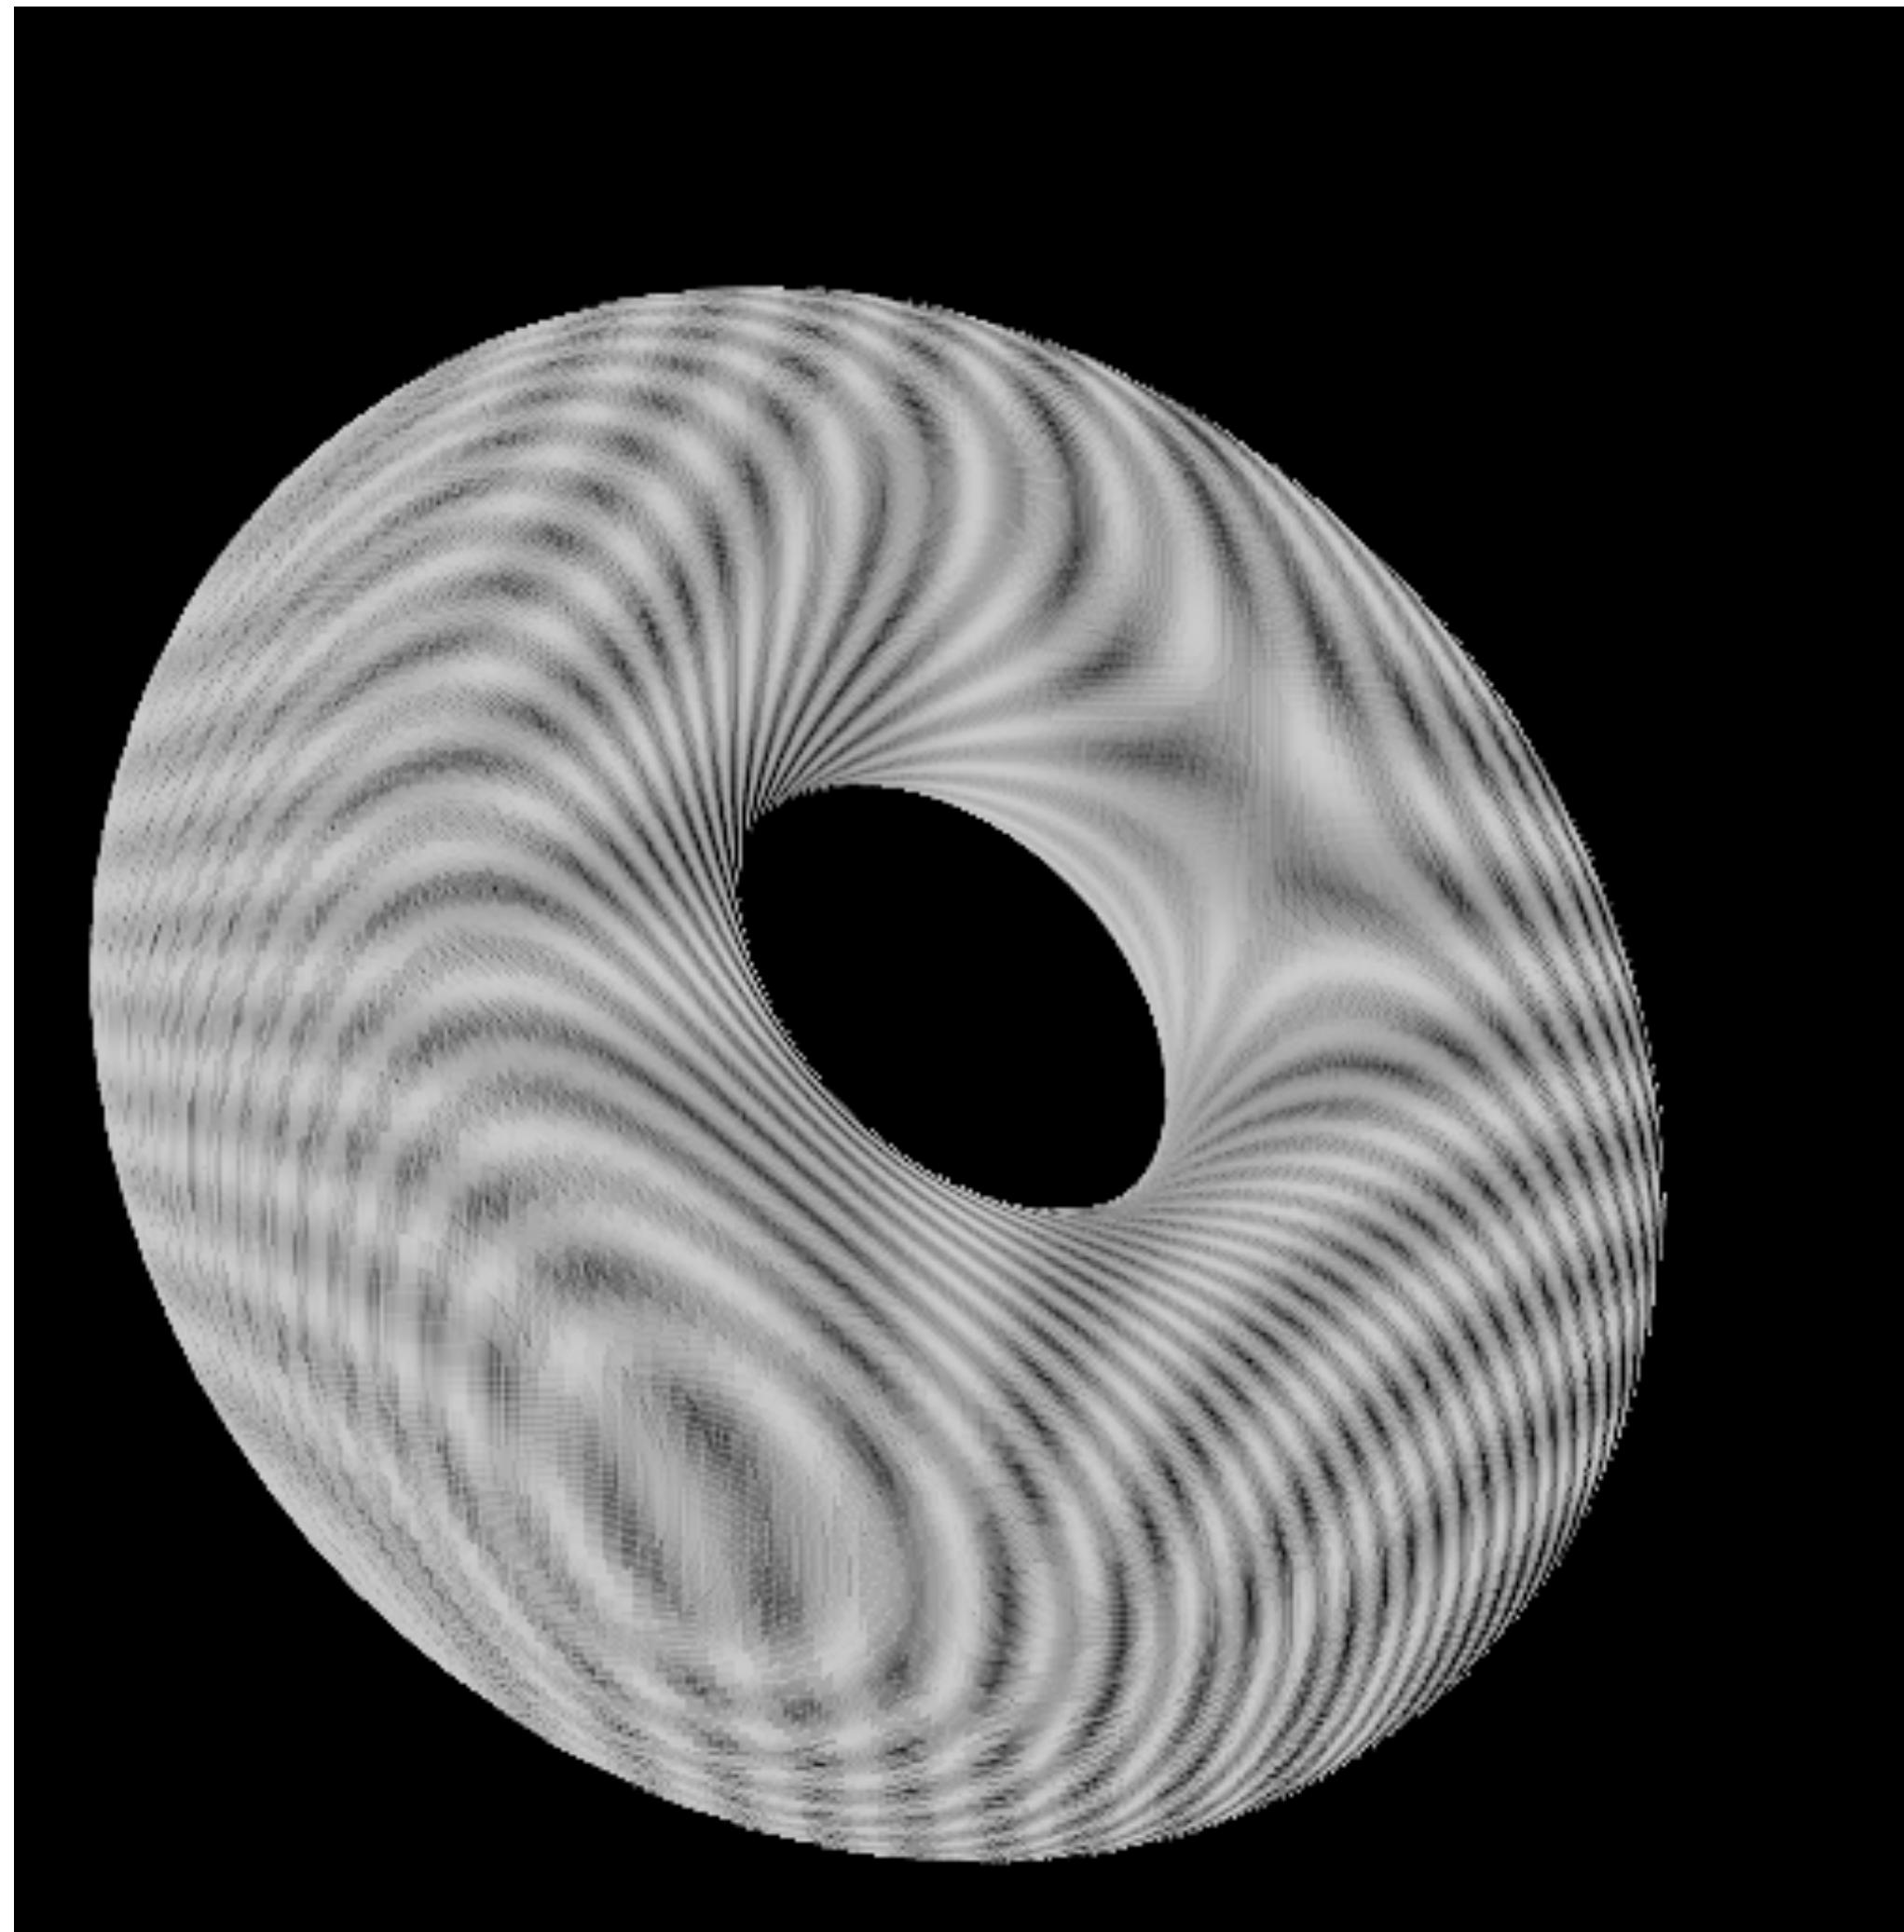


# Depth Maps – Representing the Visible



Does not capture the ‘full’  
3D structure

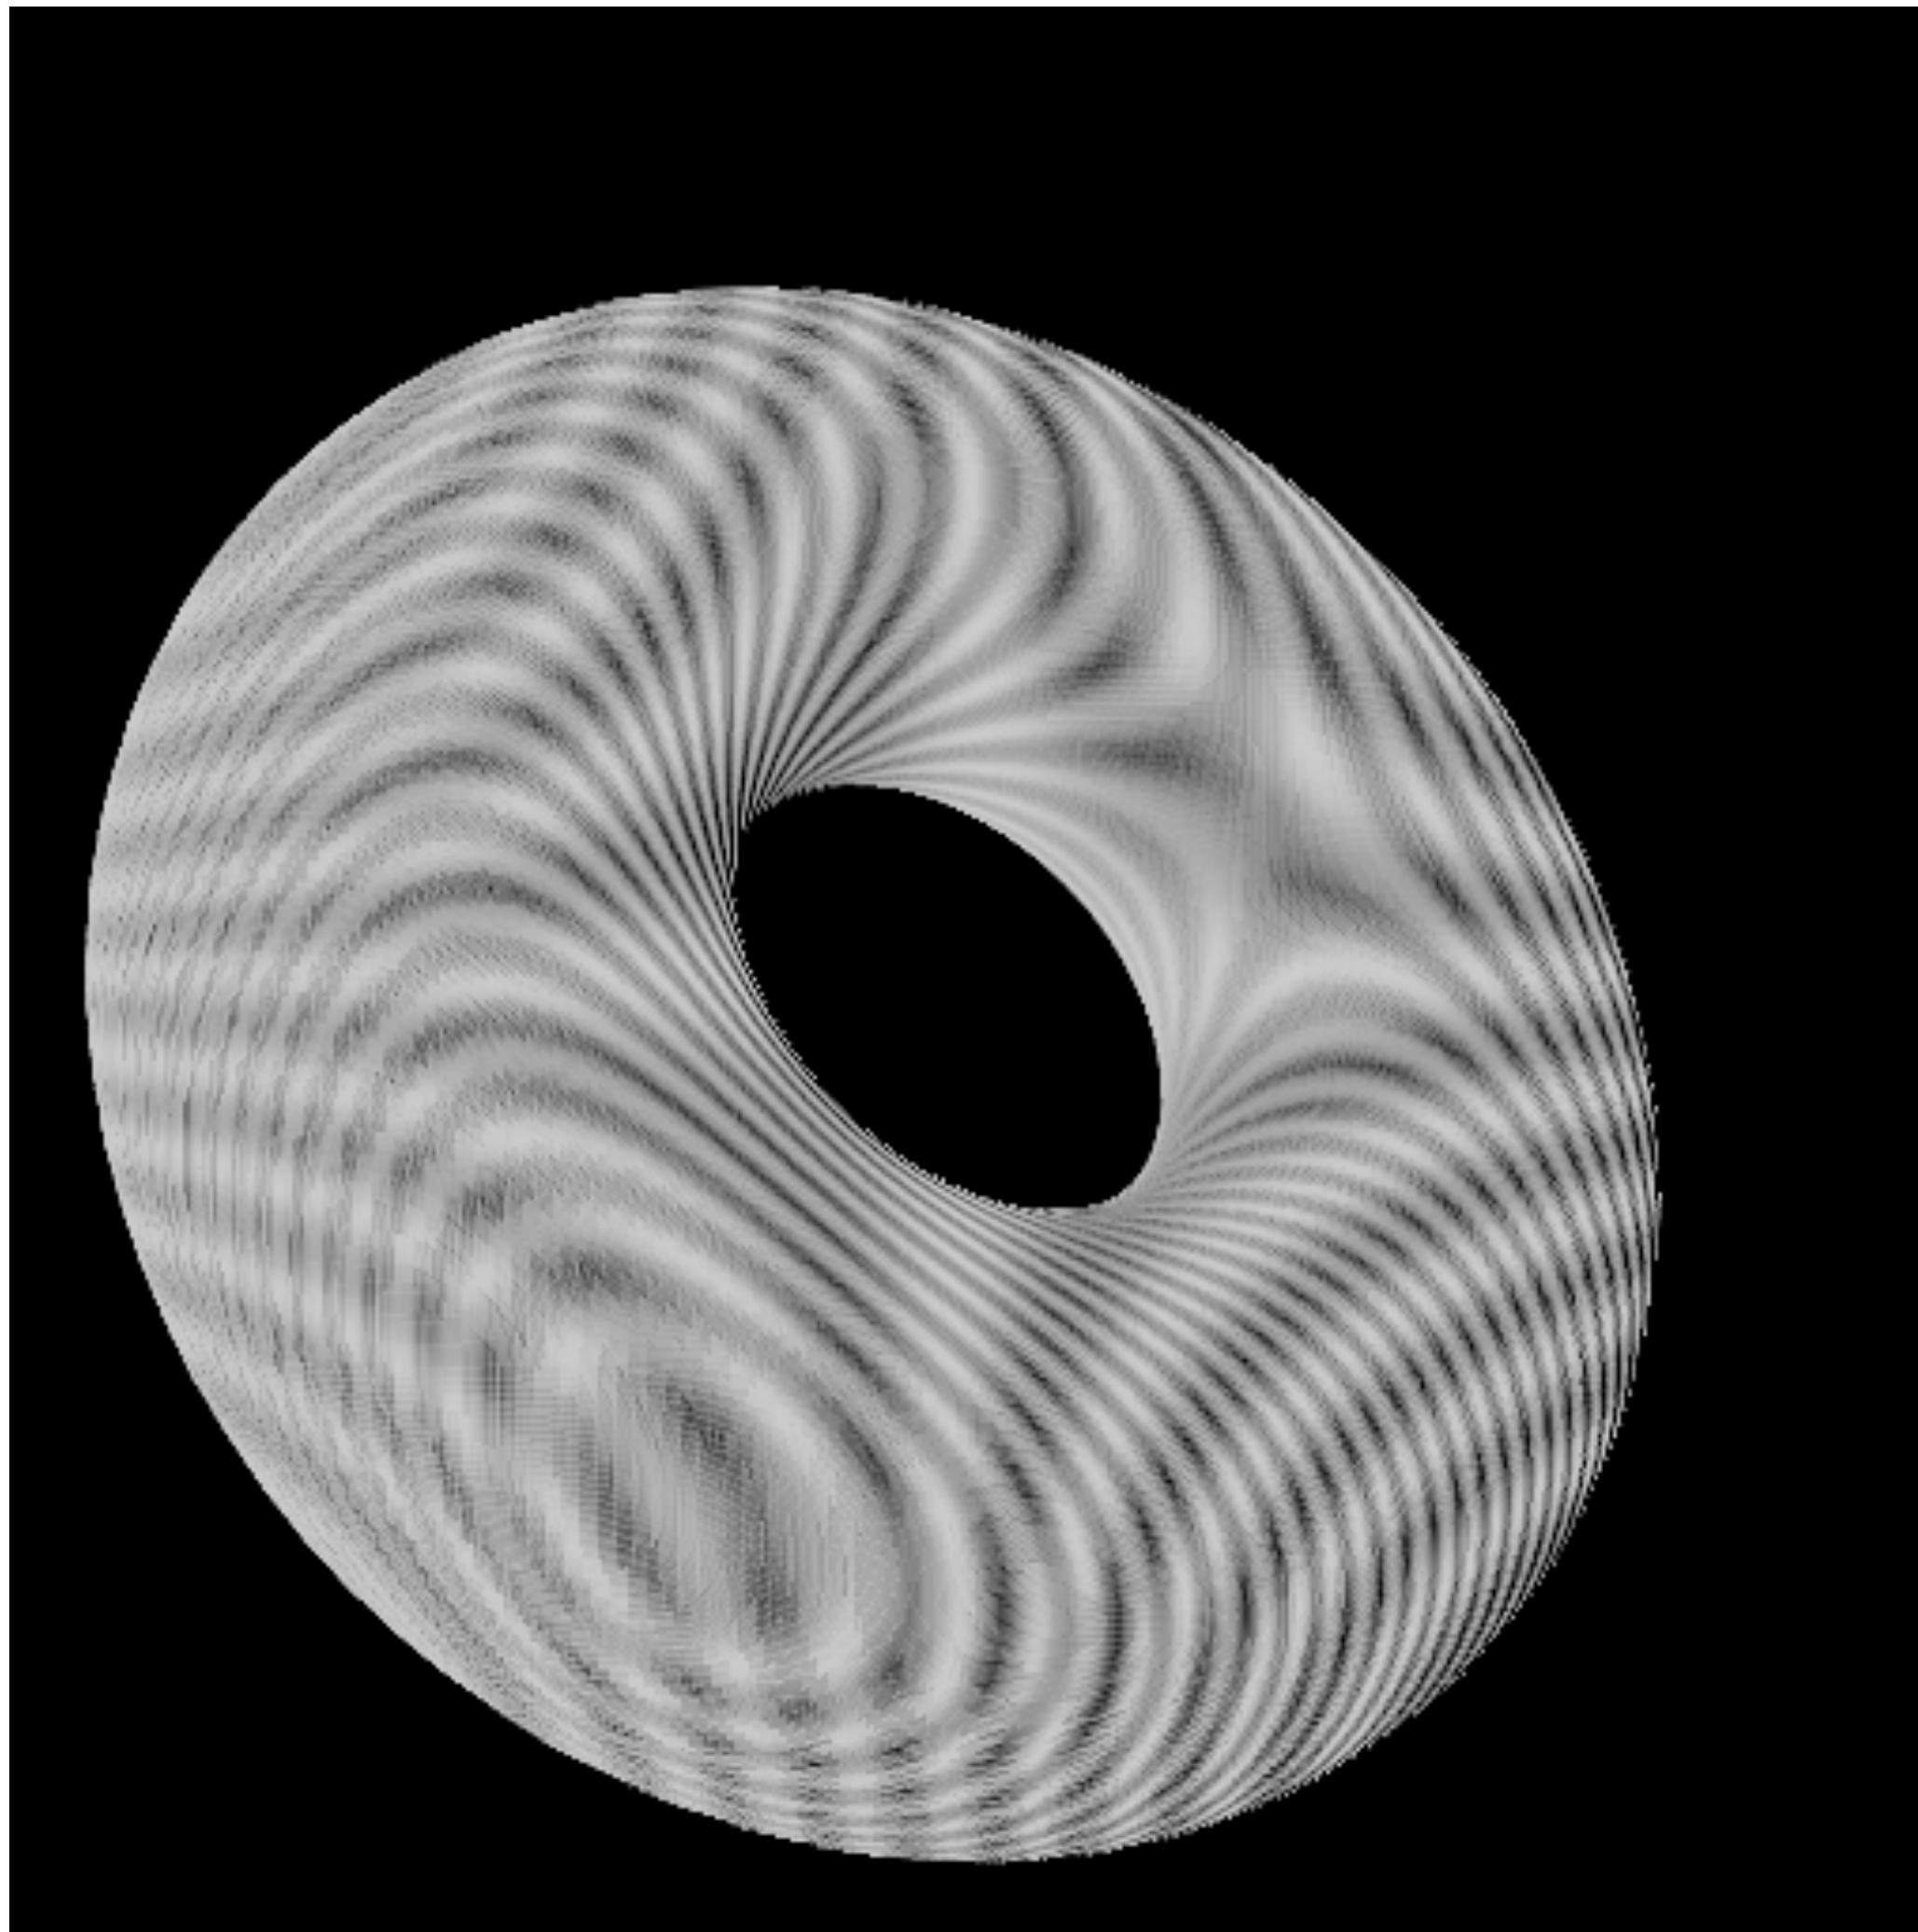
# Depth Maps – Representing the Visible



‘2.5D’ representation –  
properties associated  
to image pixels

Does not capture the ‘full’  
3D structure

# Depth Maps – Representing the Visible

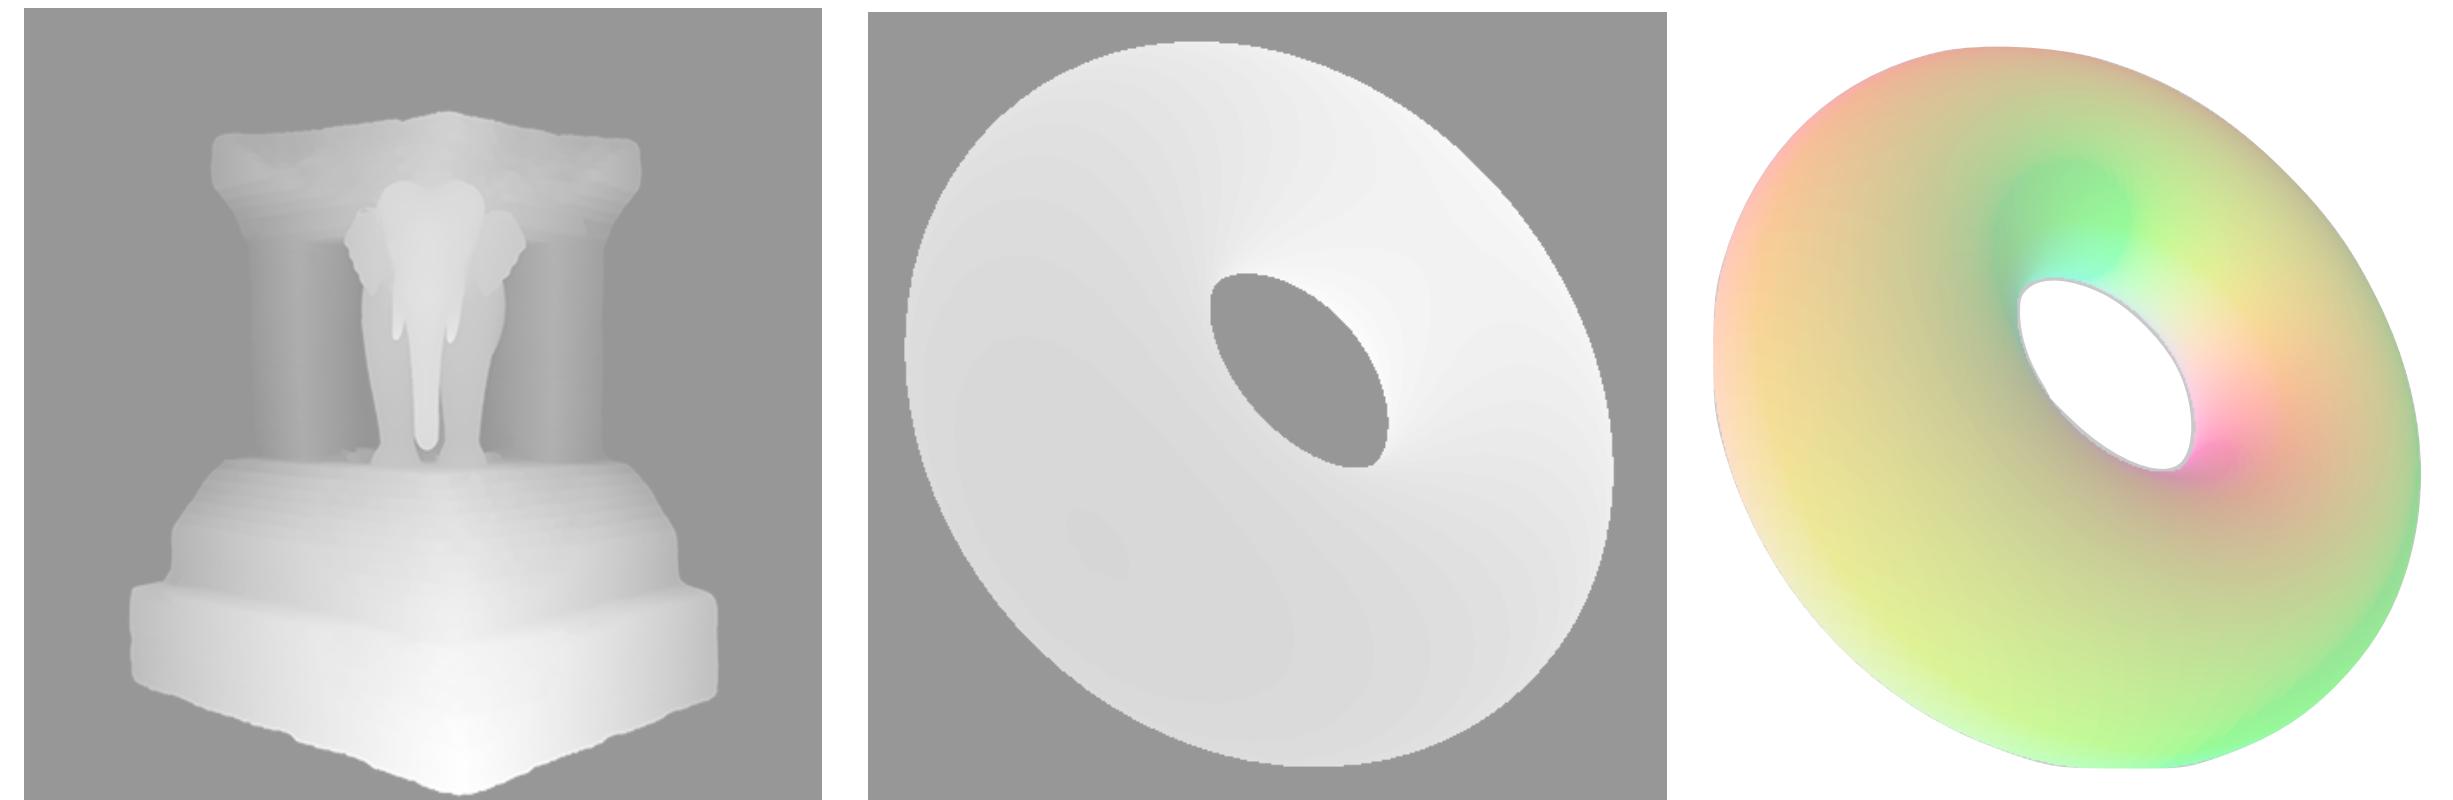


‘2.5D’ representation –  
properties associated  
to image pixels

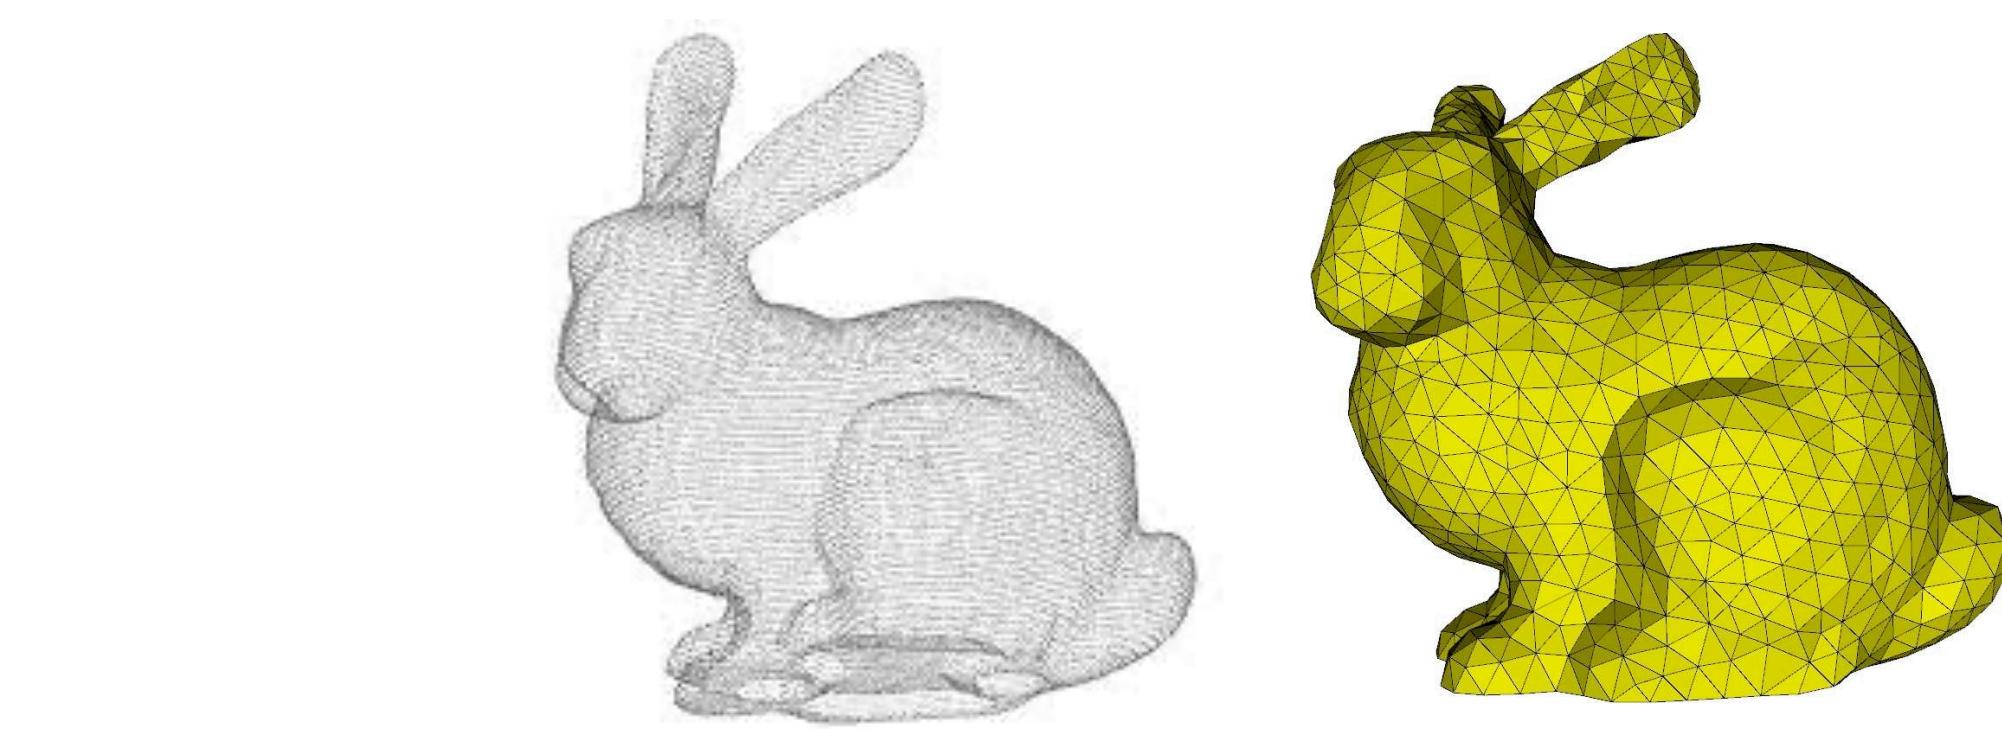
Does not capture the ‘full’  
3D structure

can we relax this ‘image-  
dependence’ constraint?

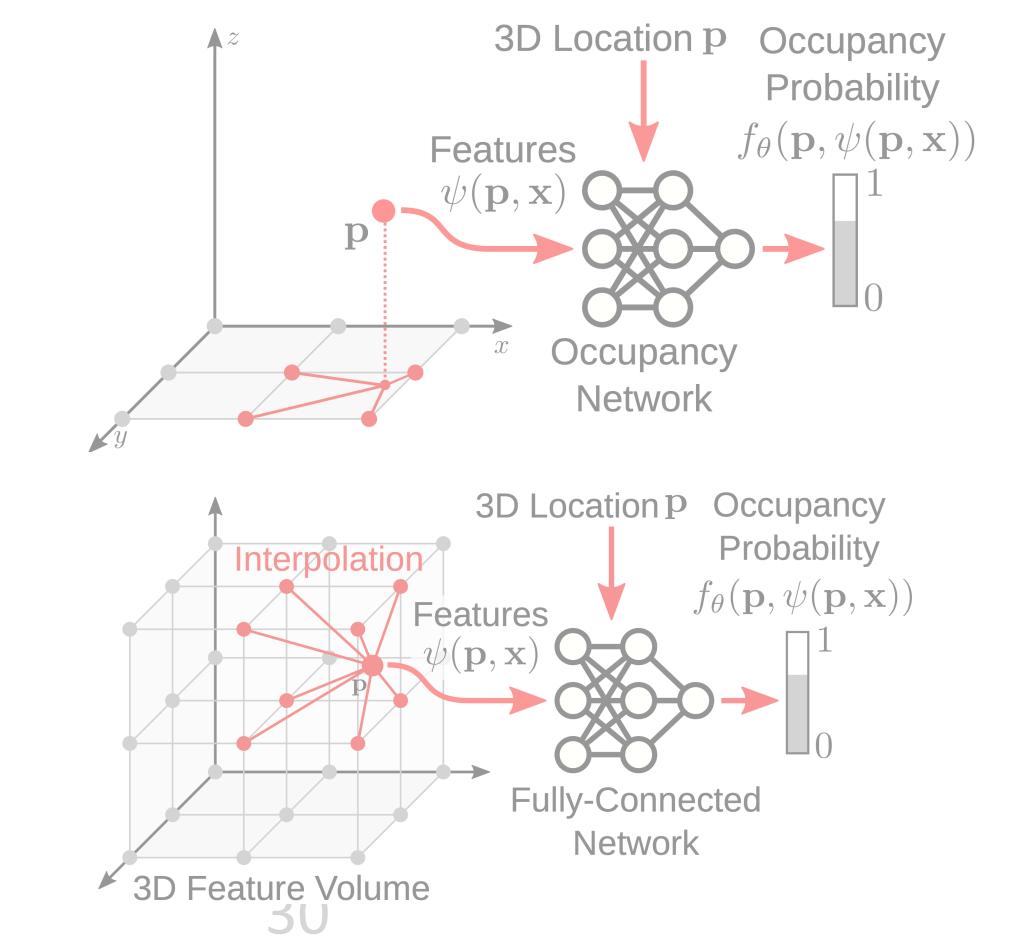
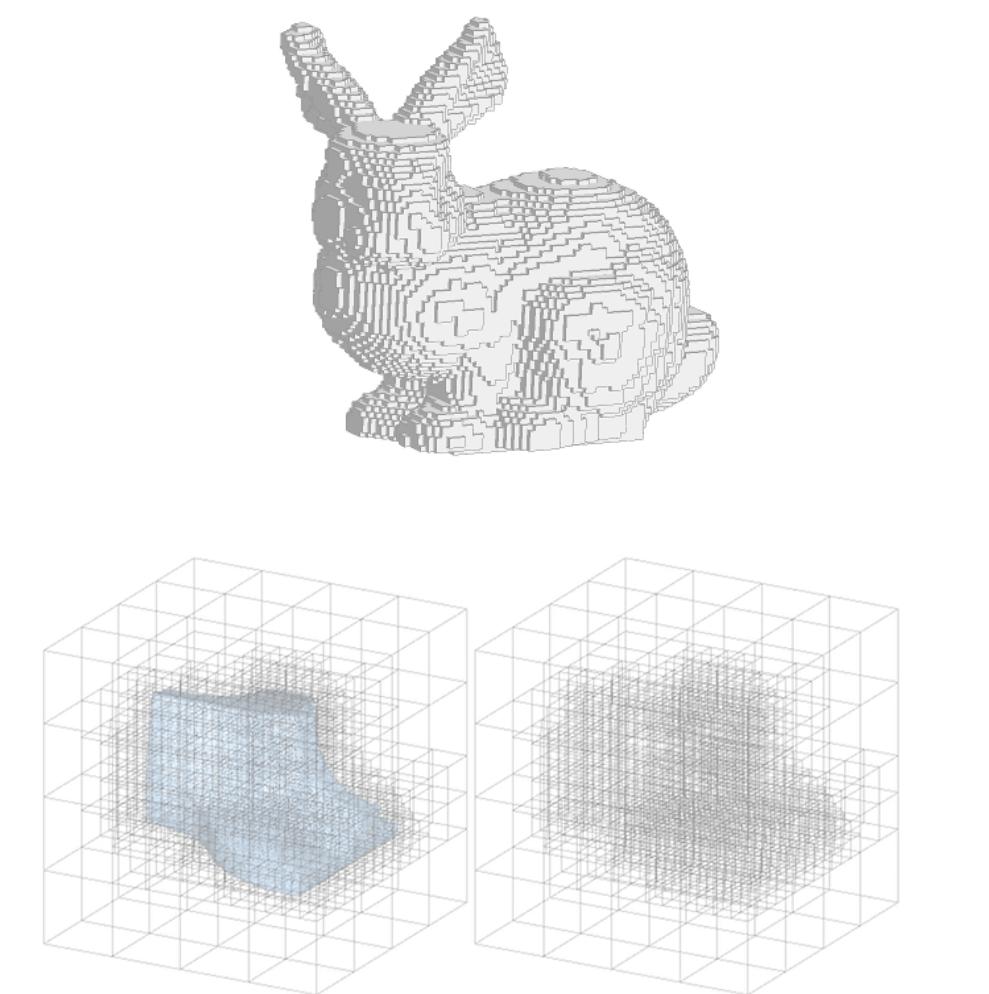
## 2.5D Representations



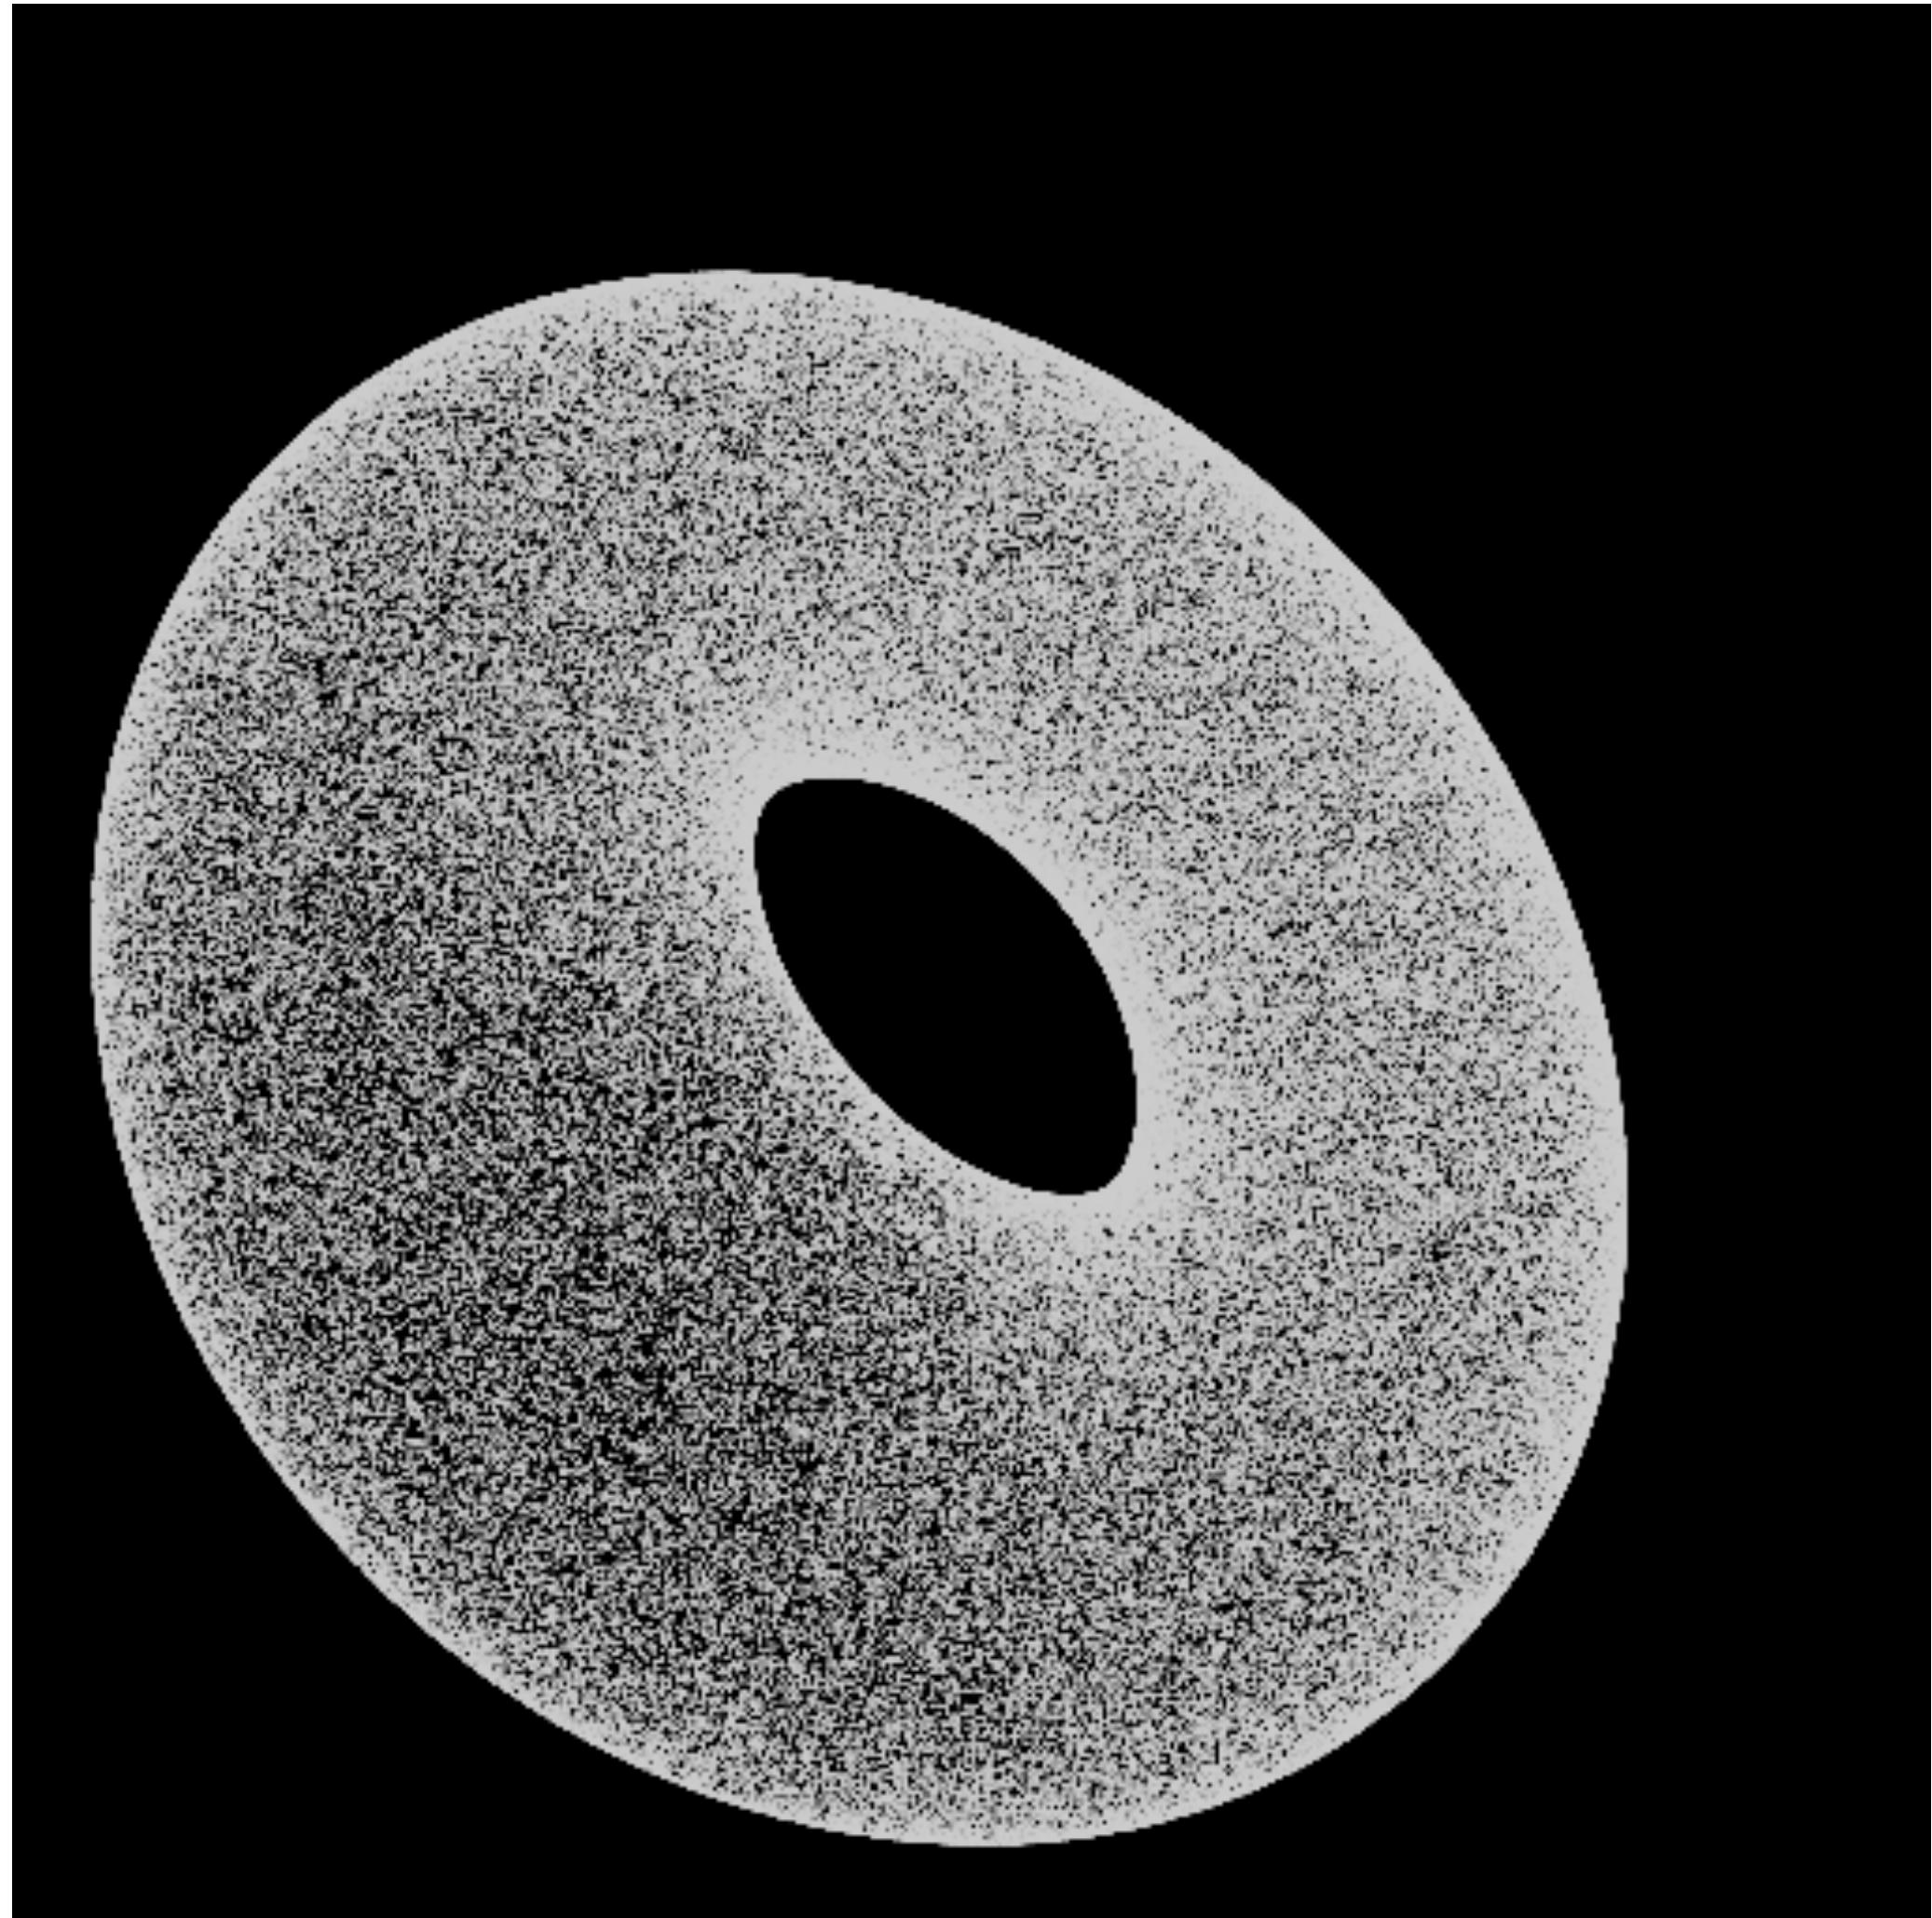
## Surface Representations



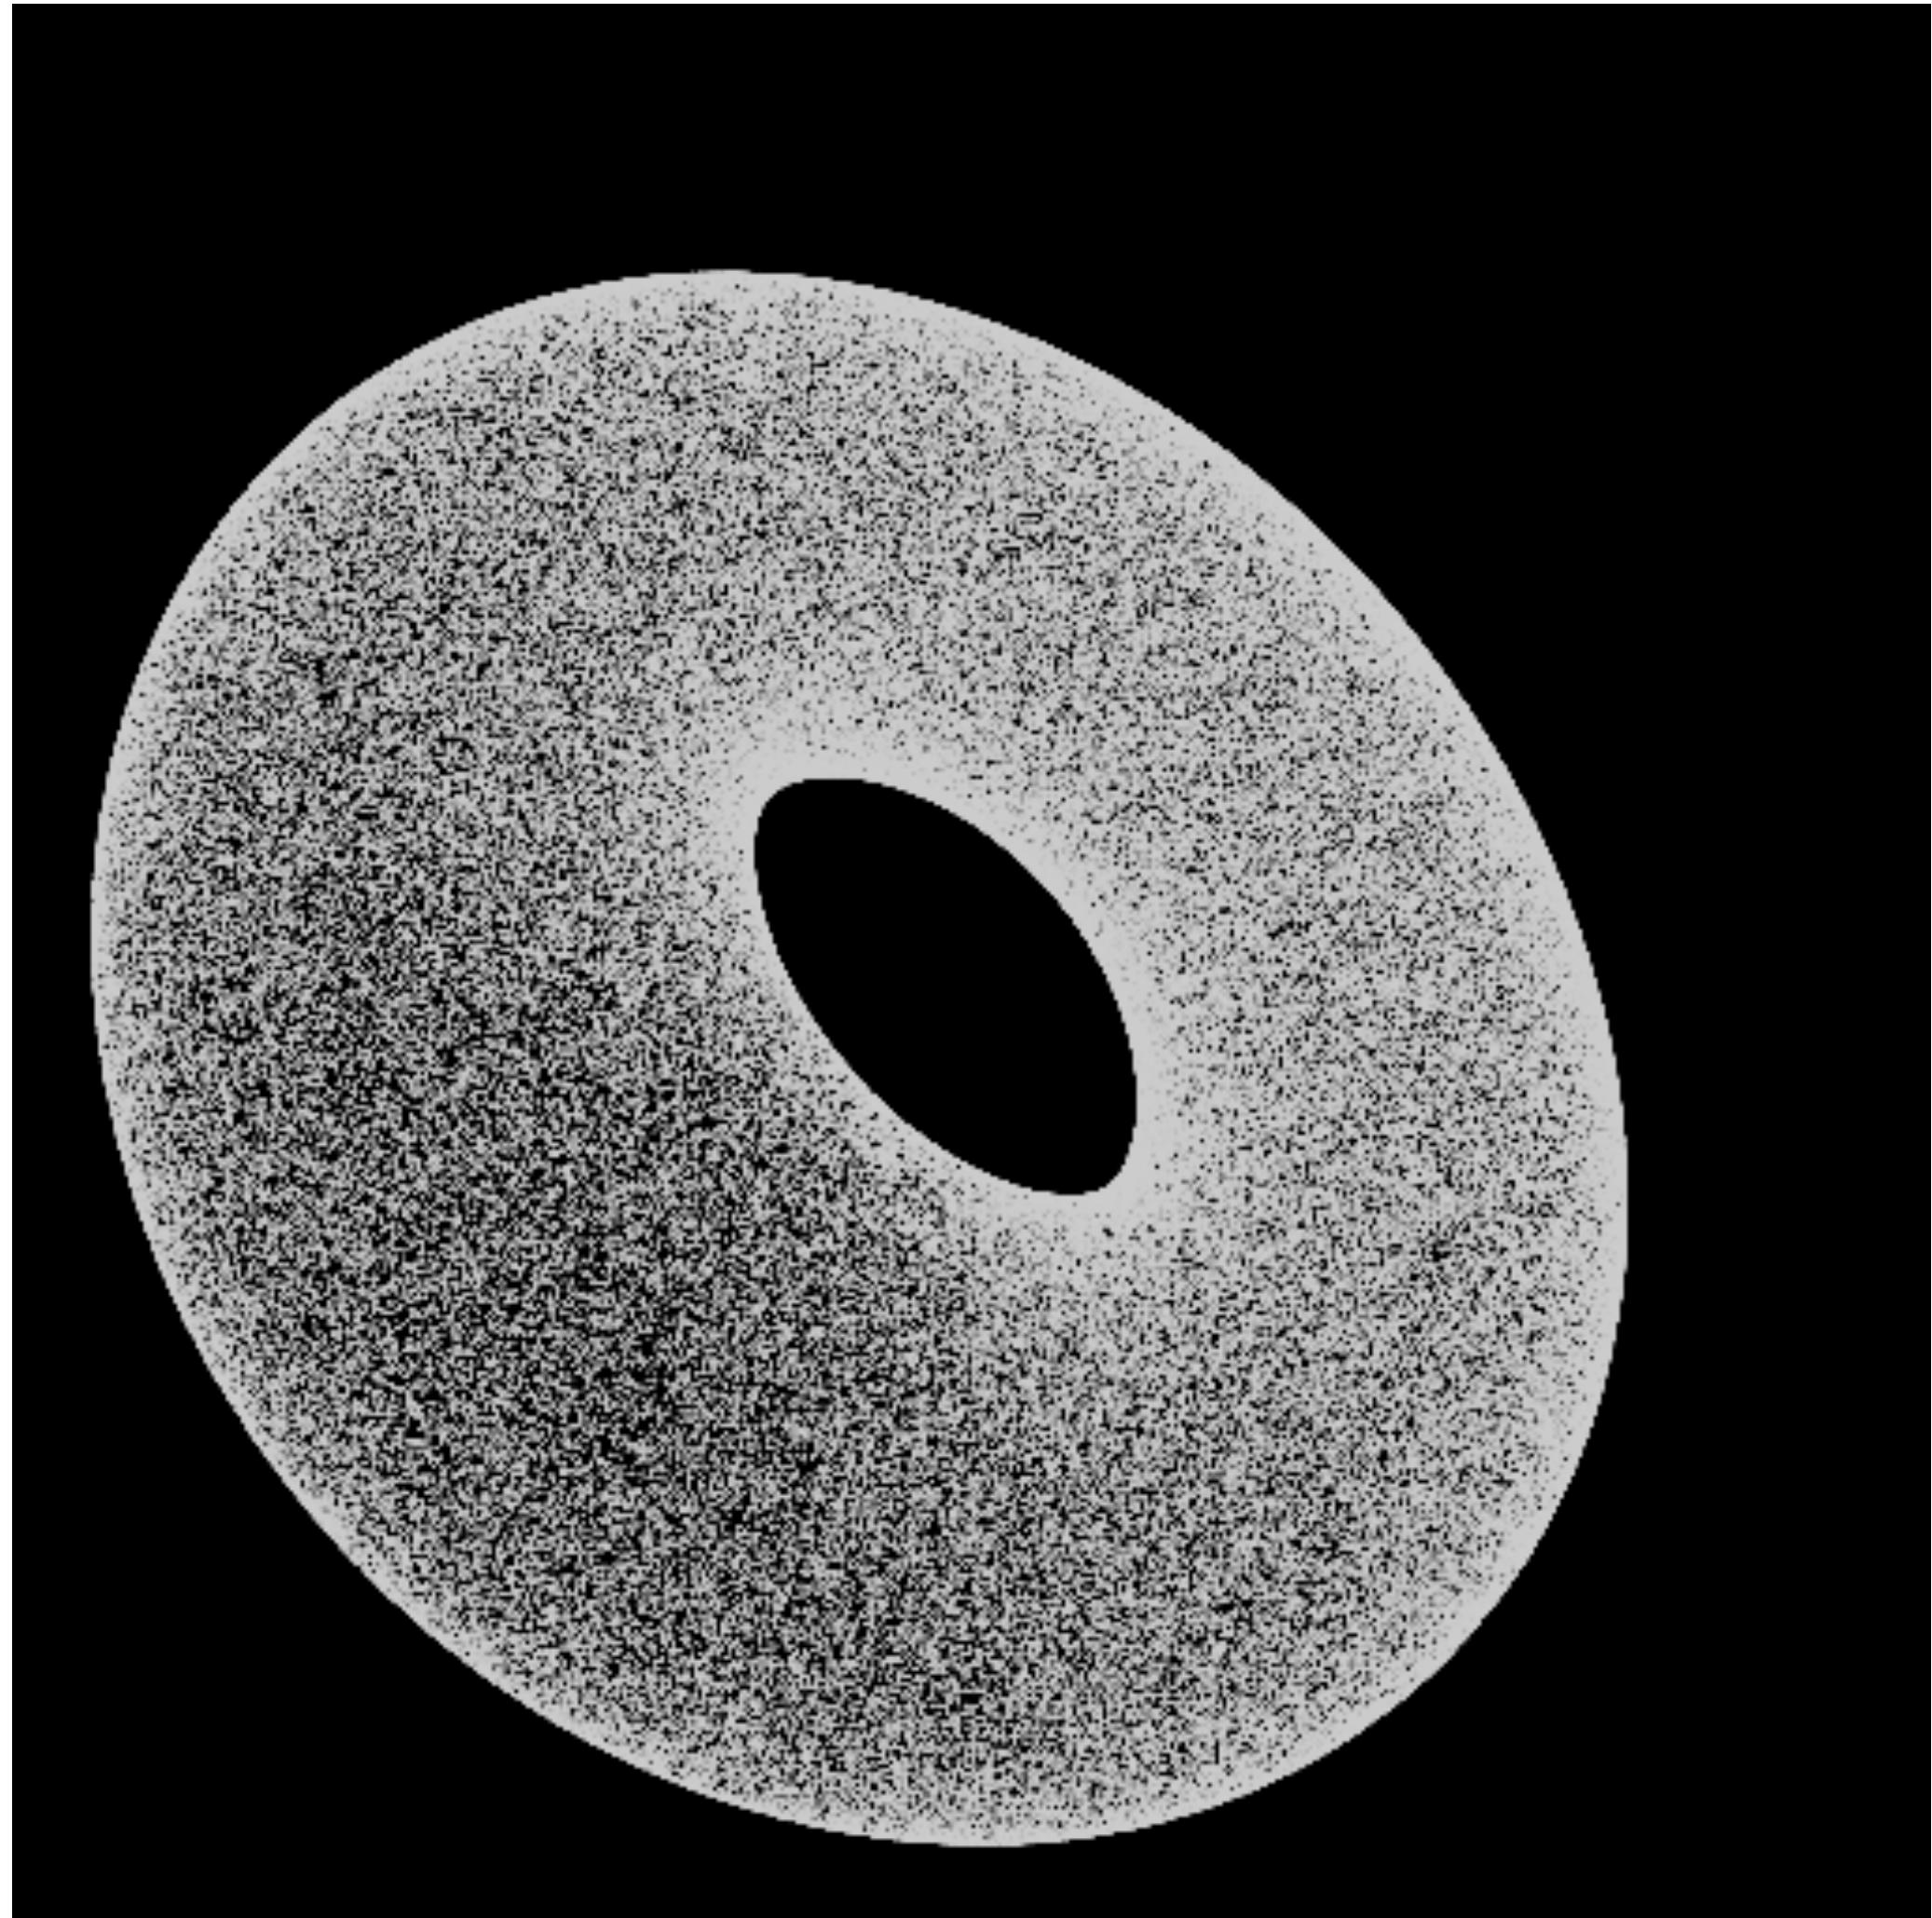
## Field Parameterizations



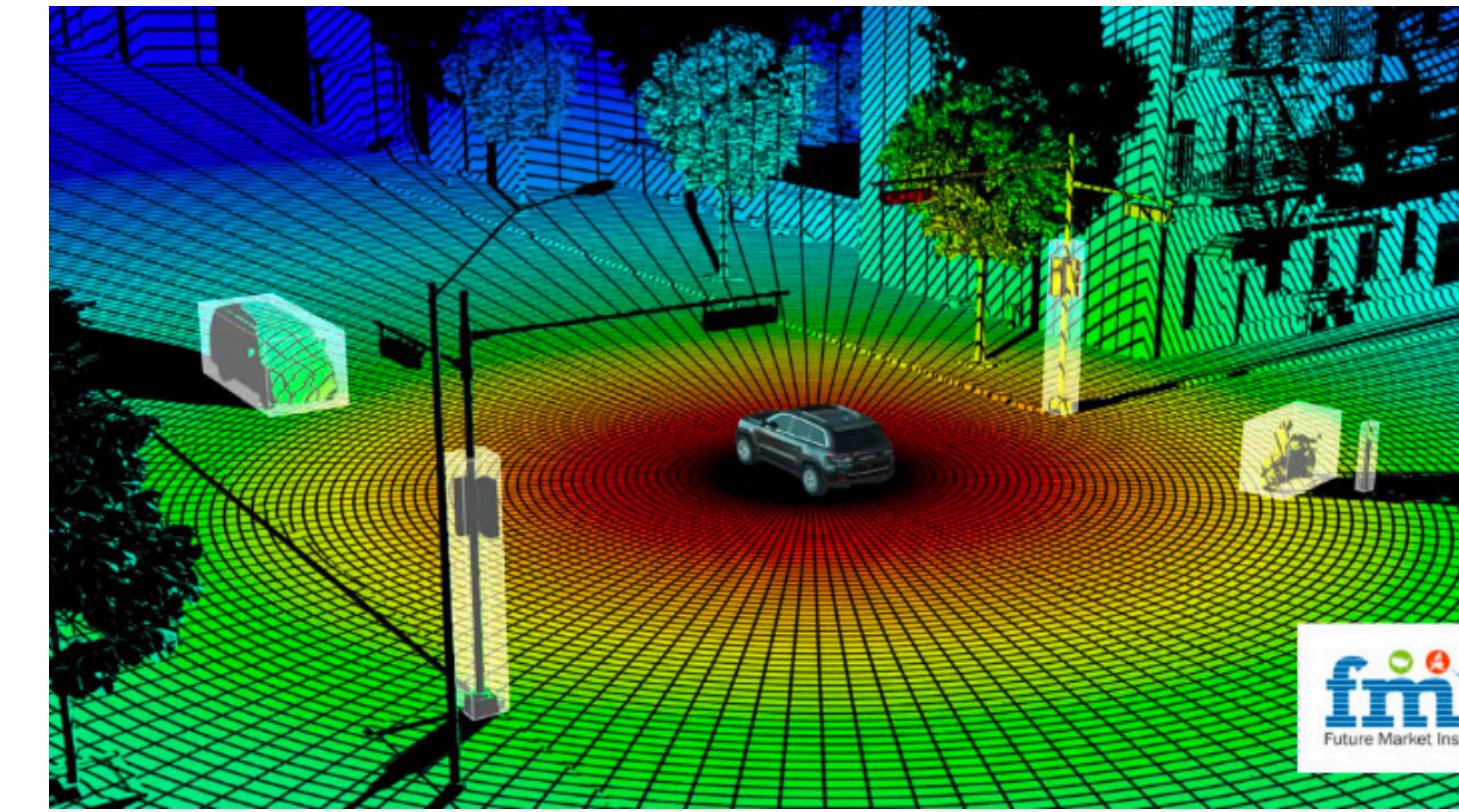
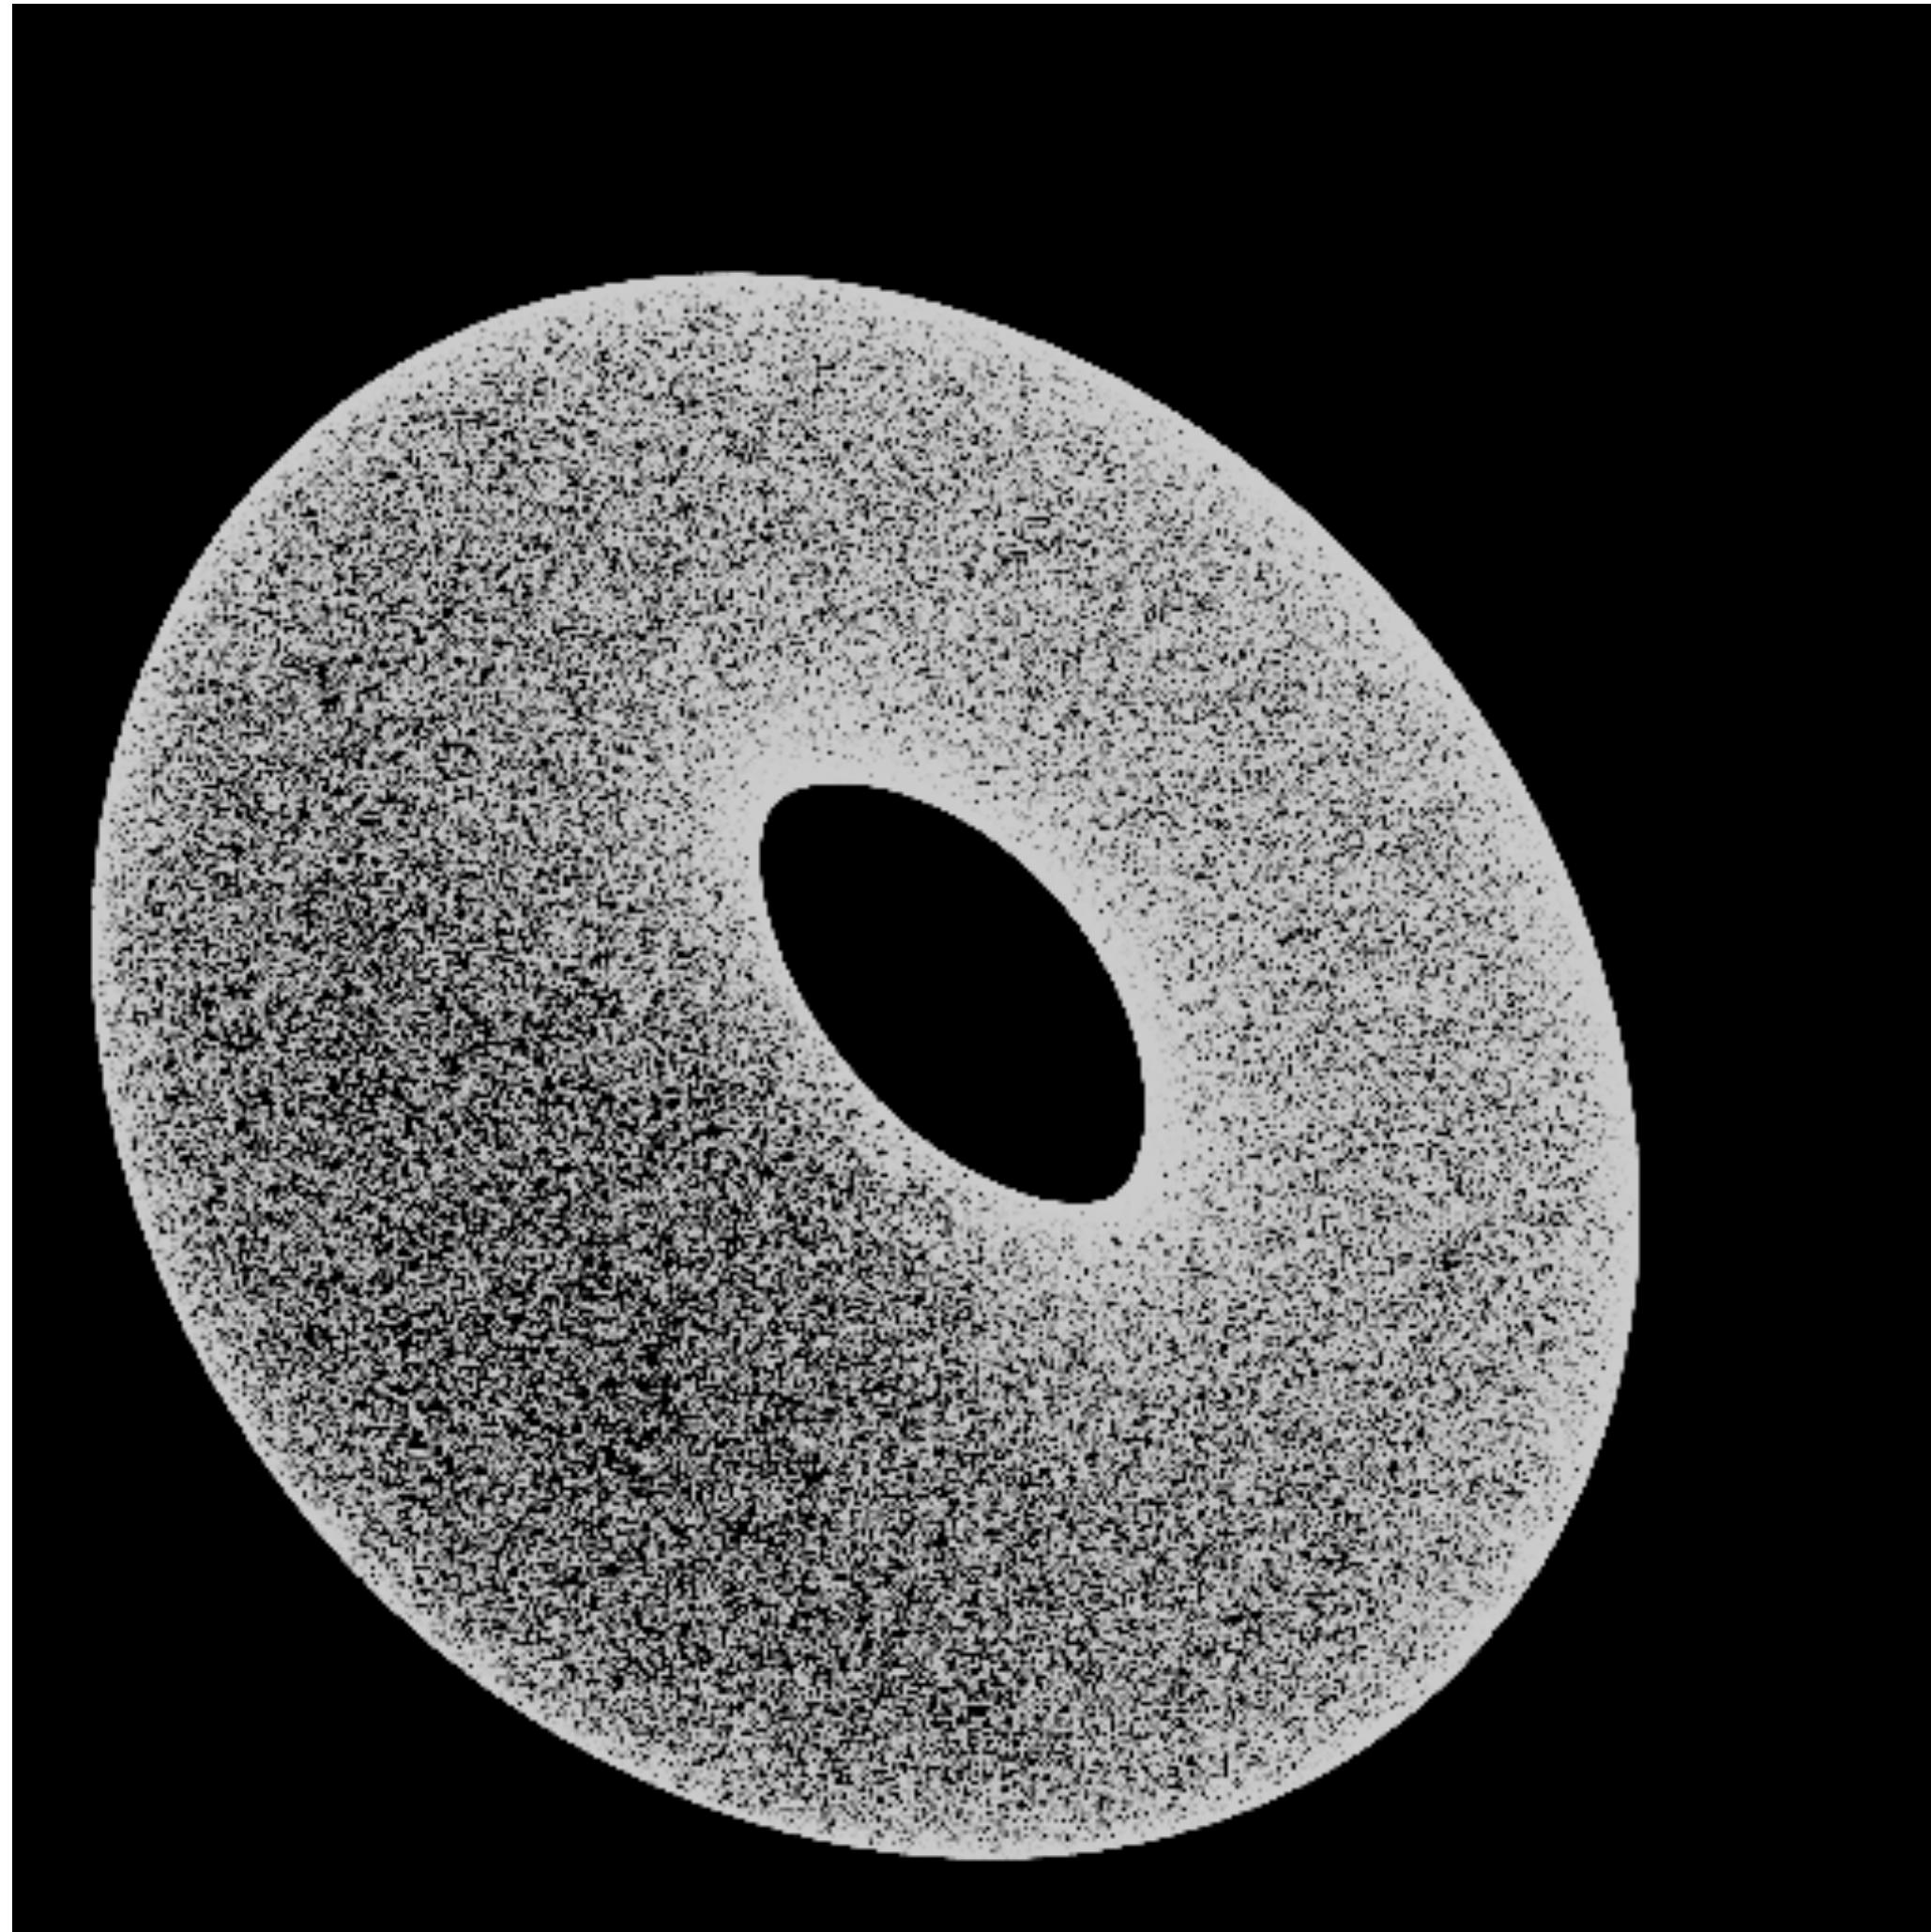
# Point Clouds



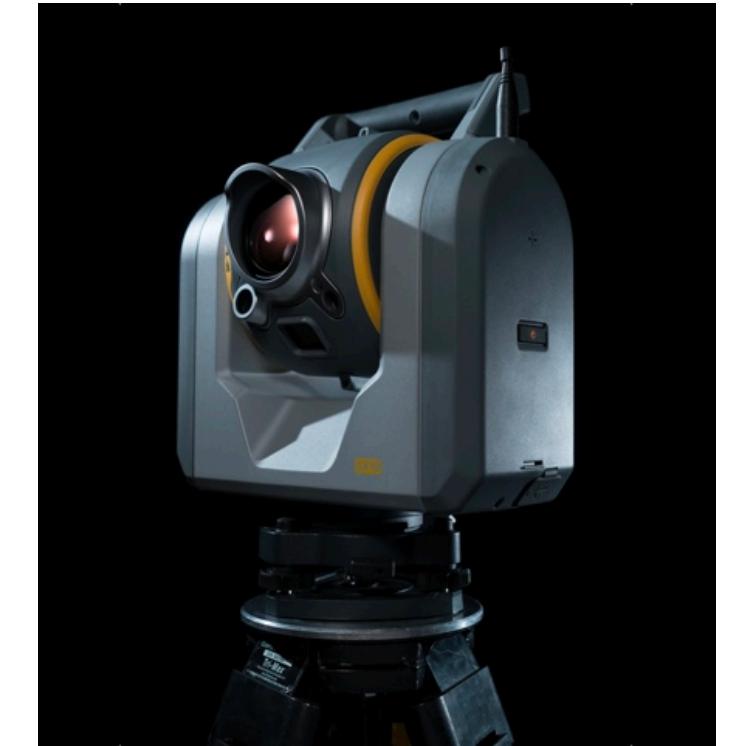
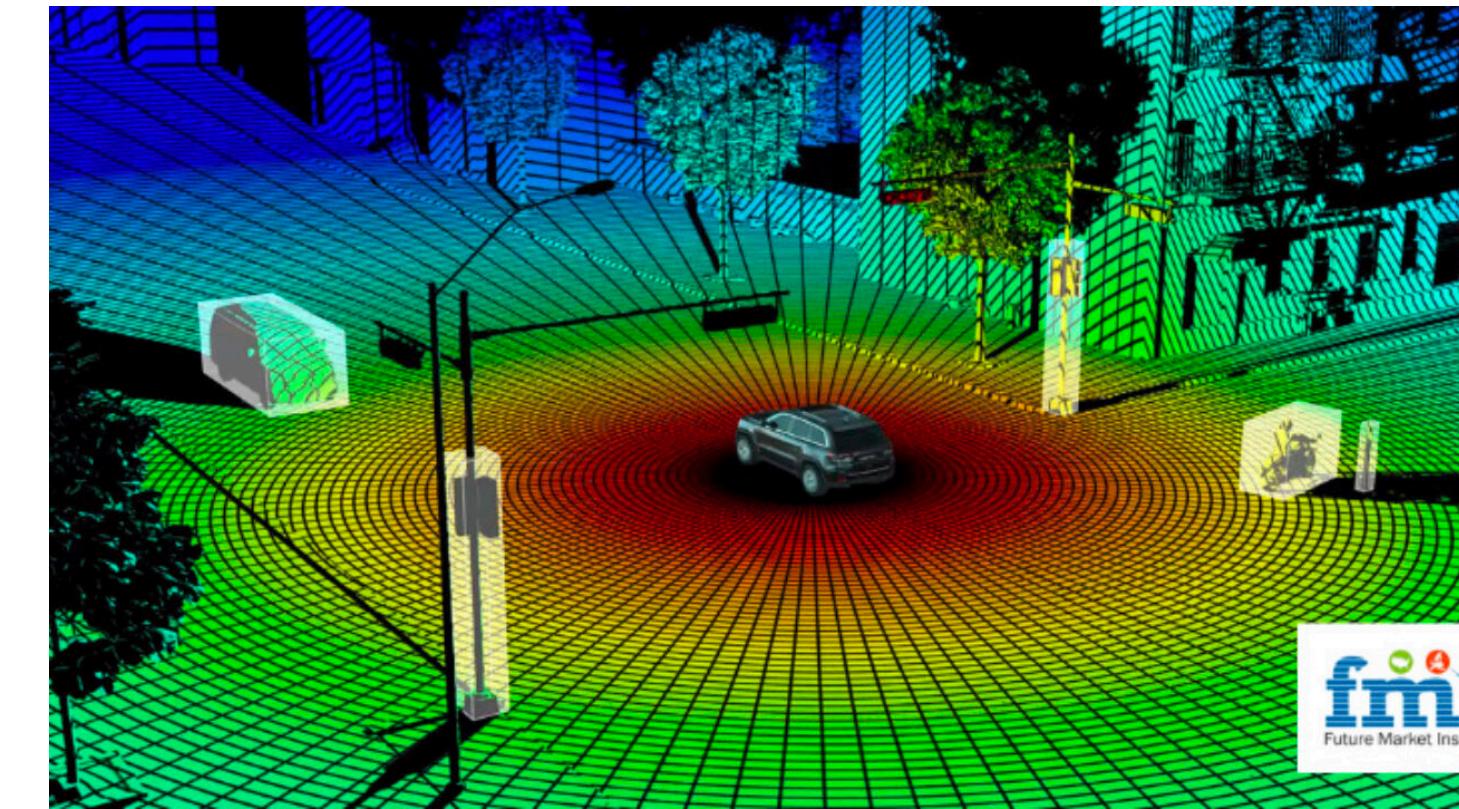
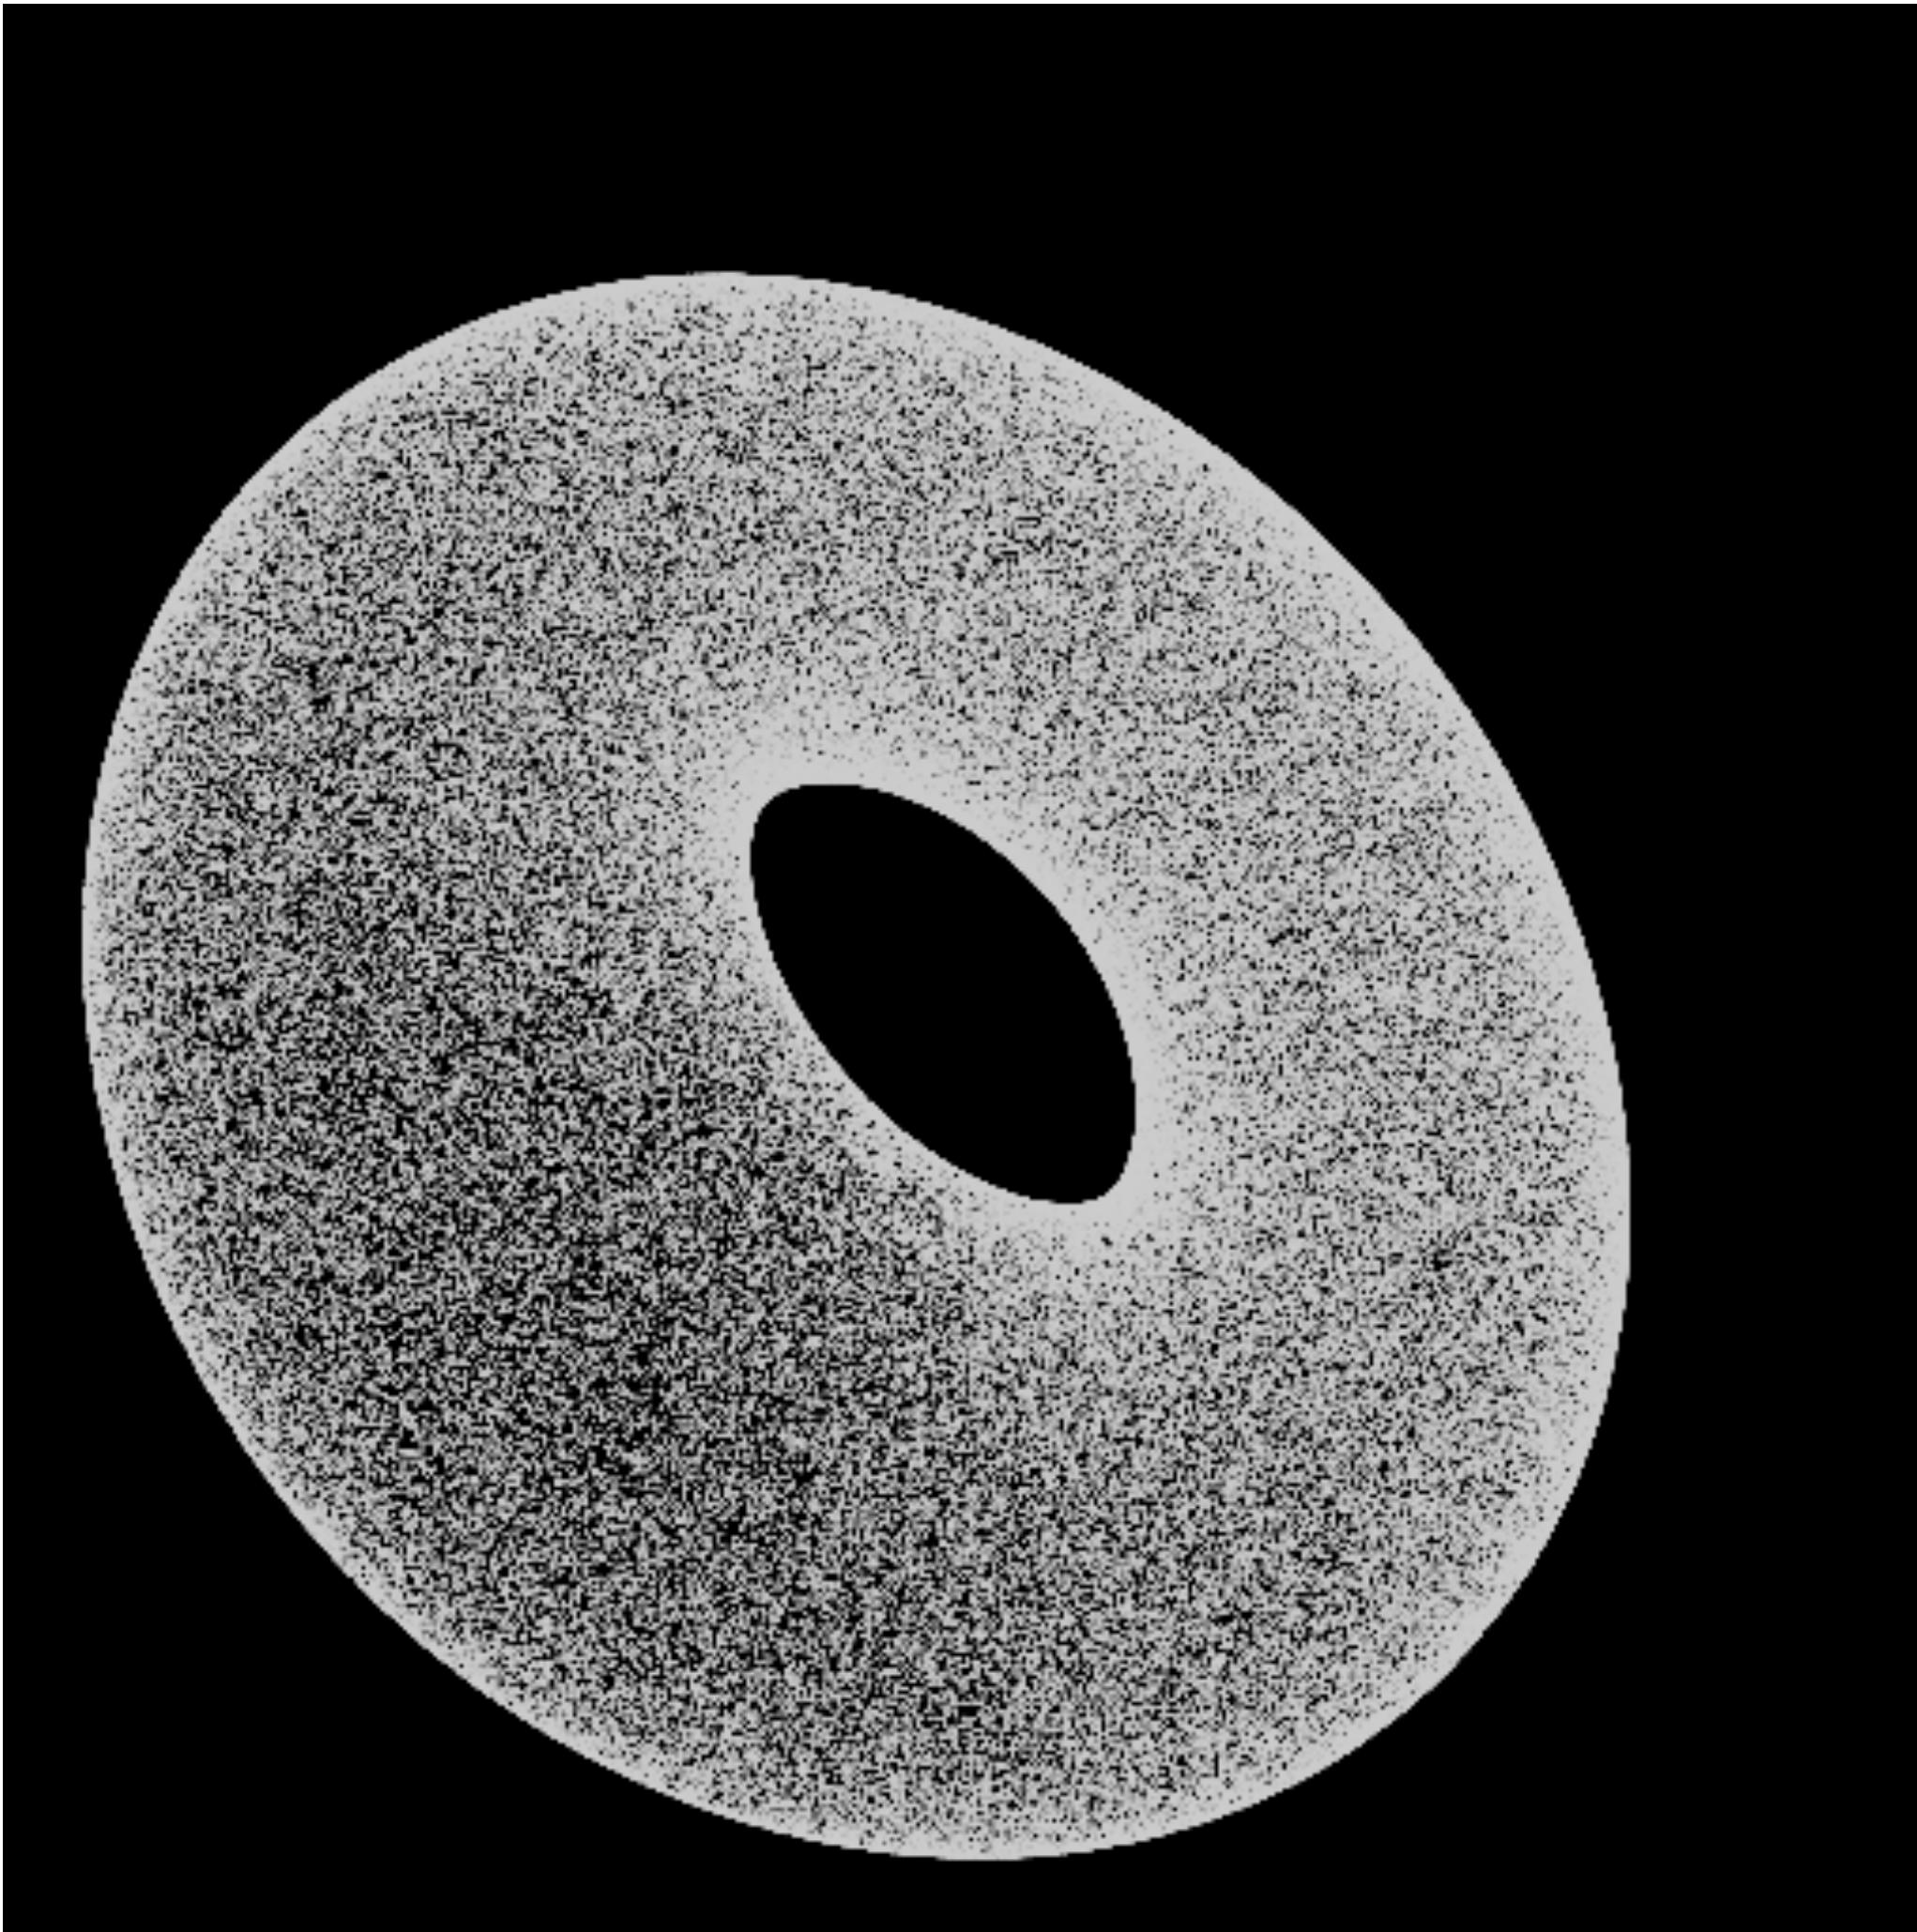
# Point Clouds



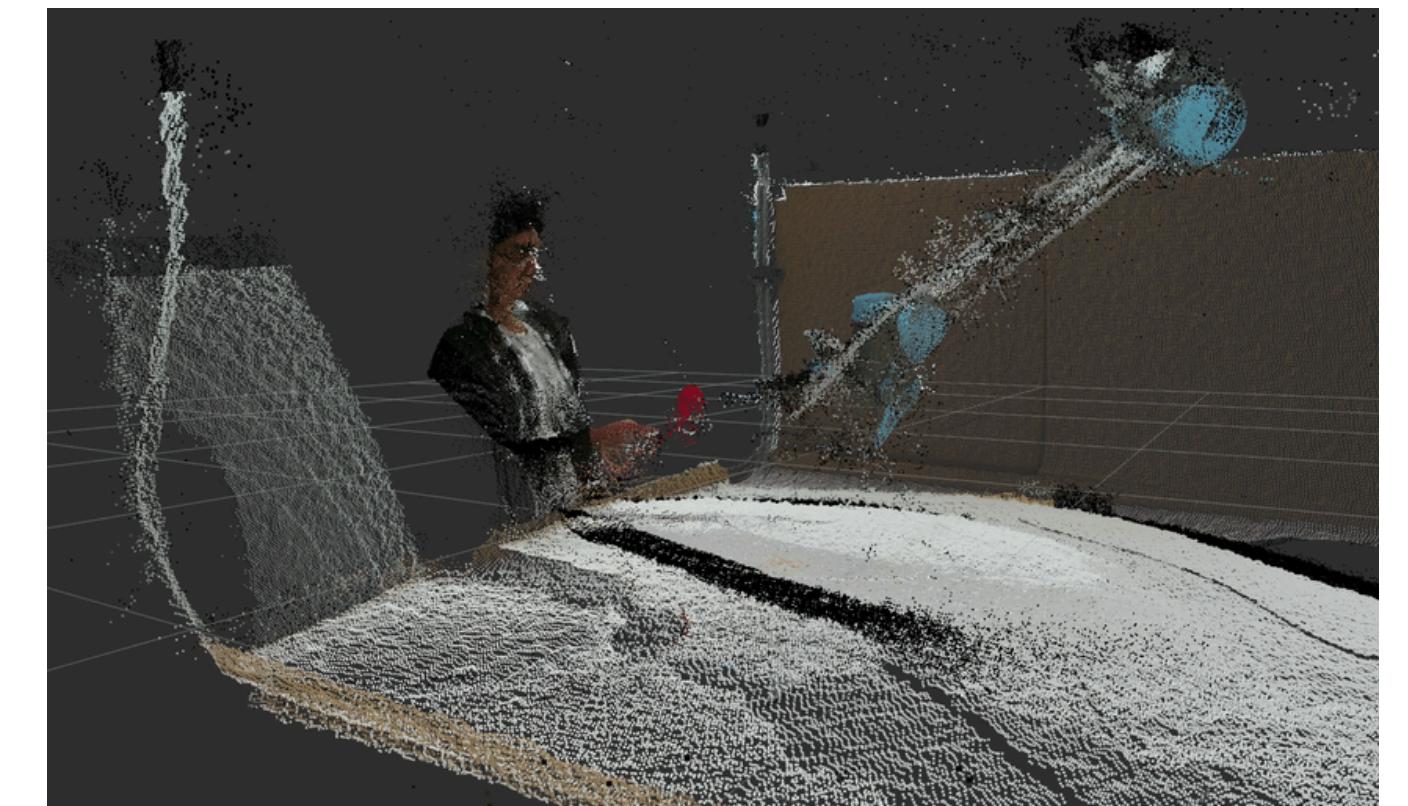
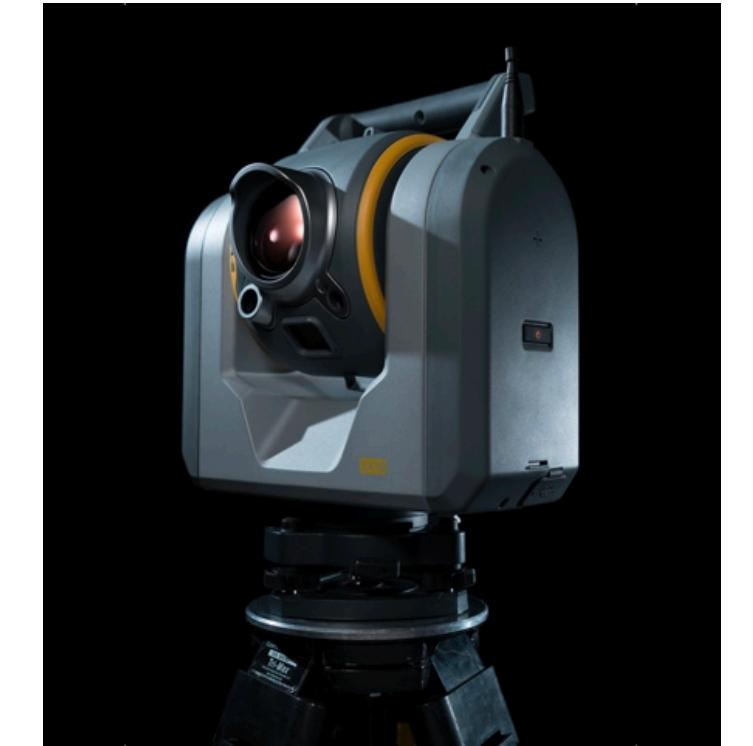
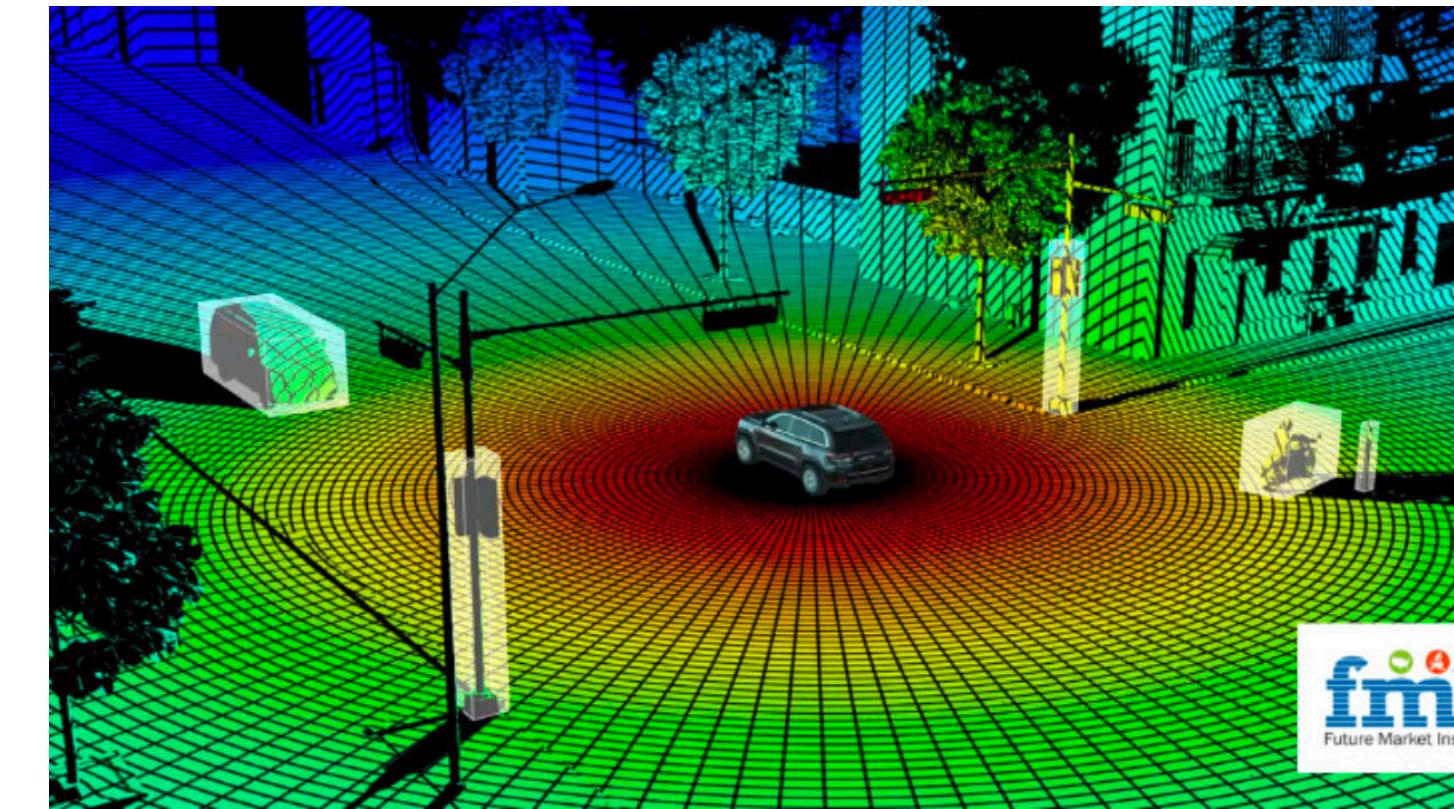
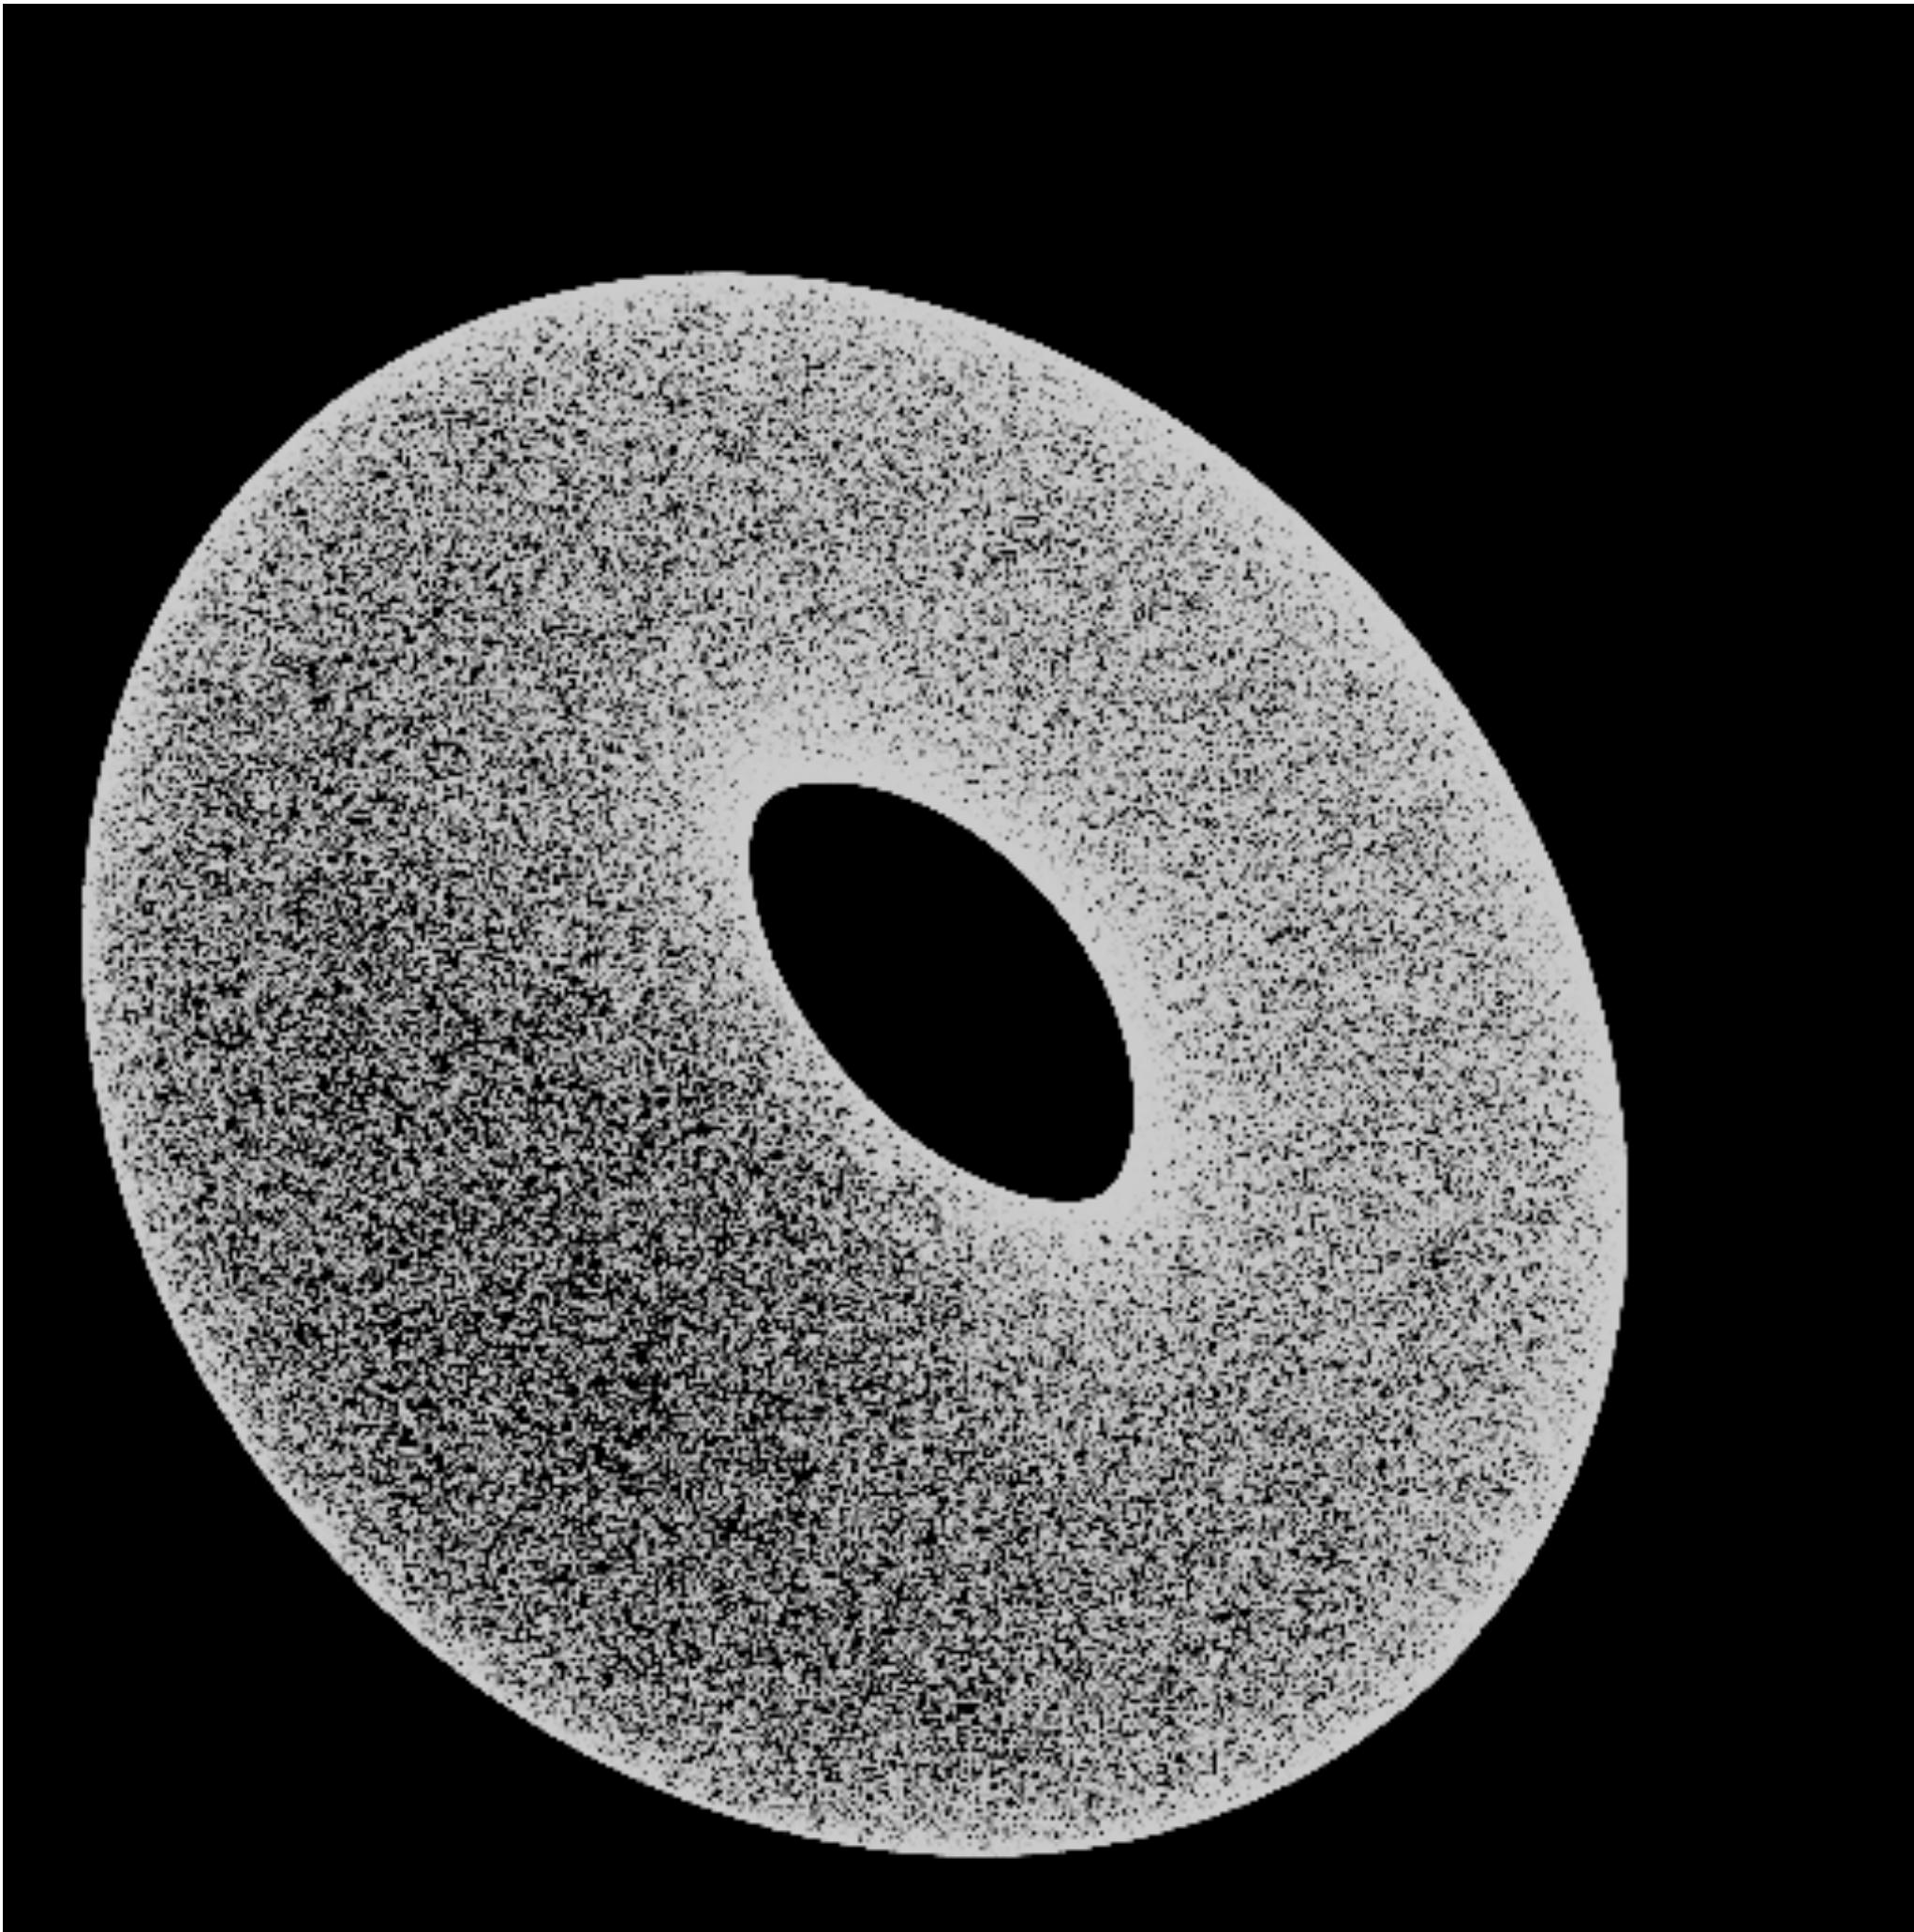
# Point Clouds



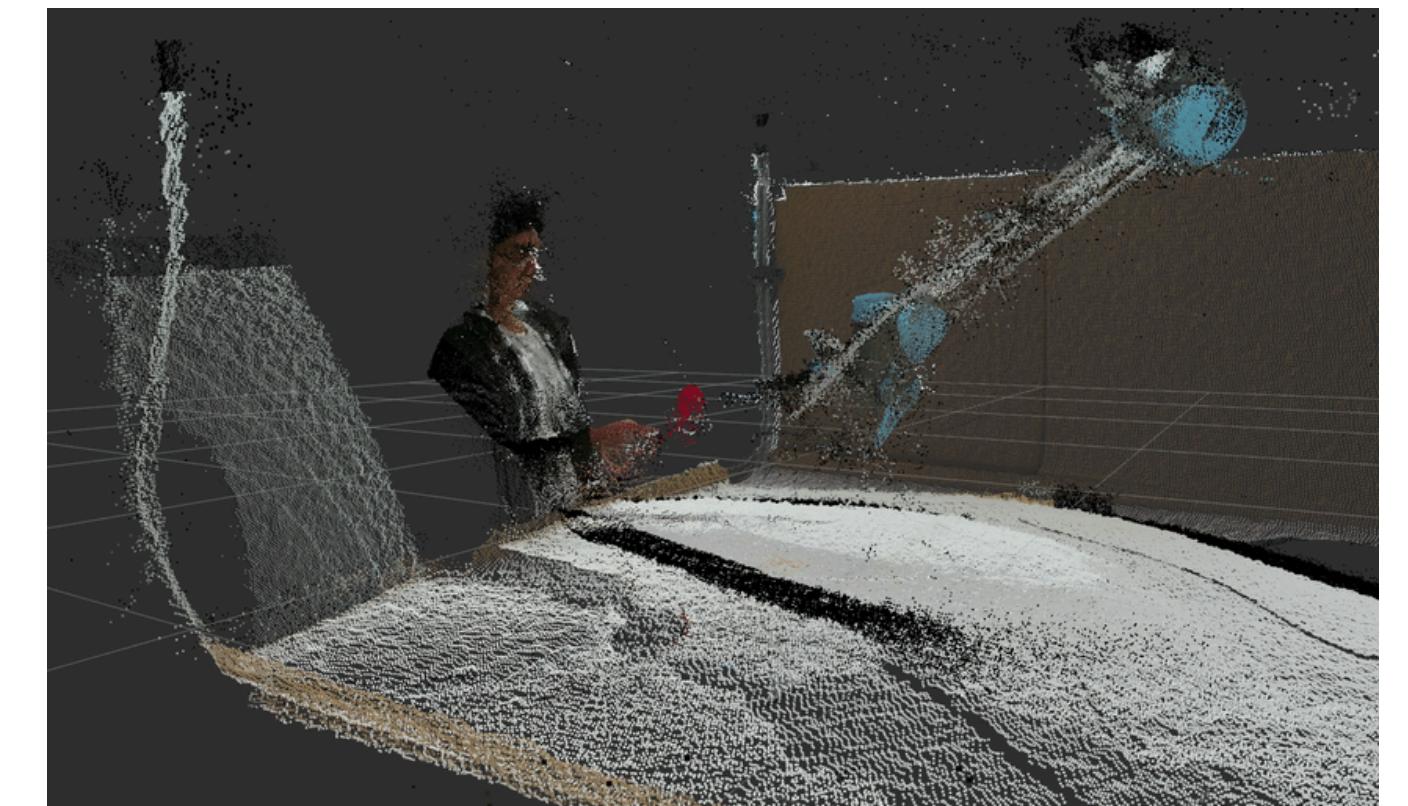
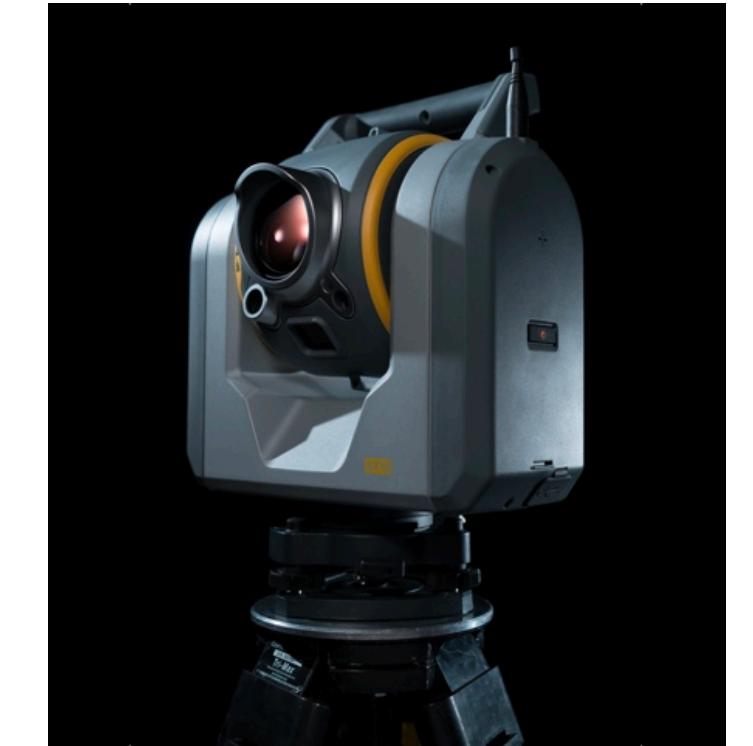
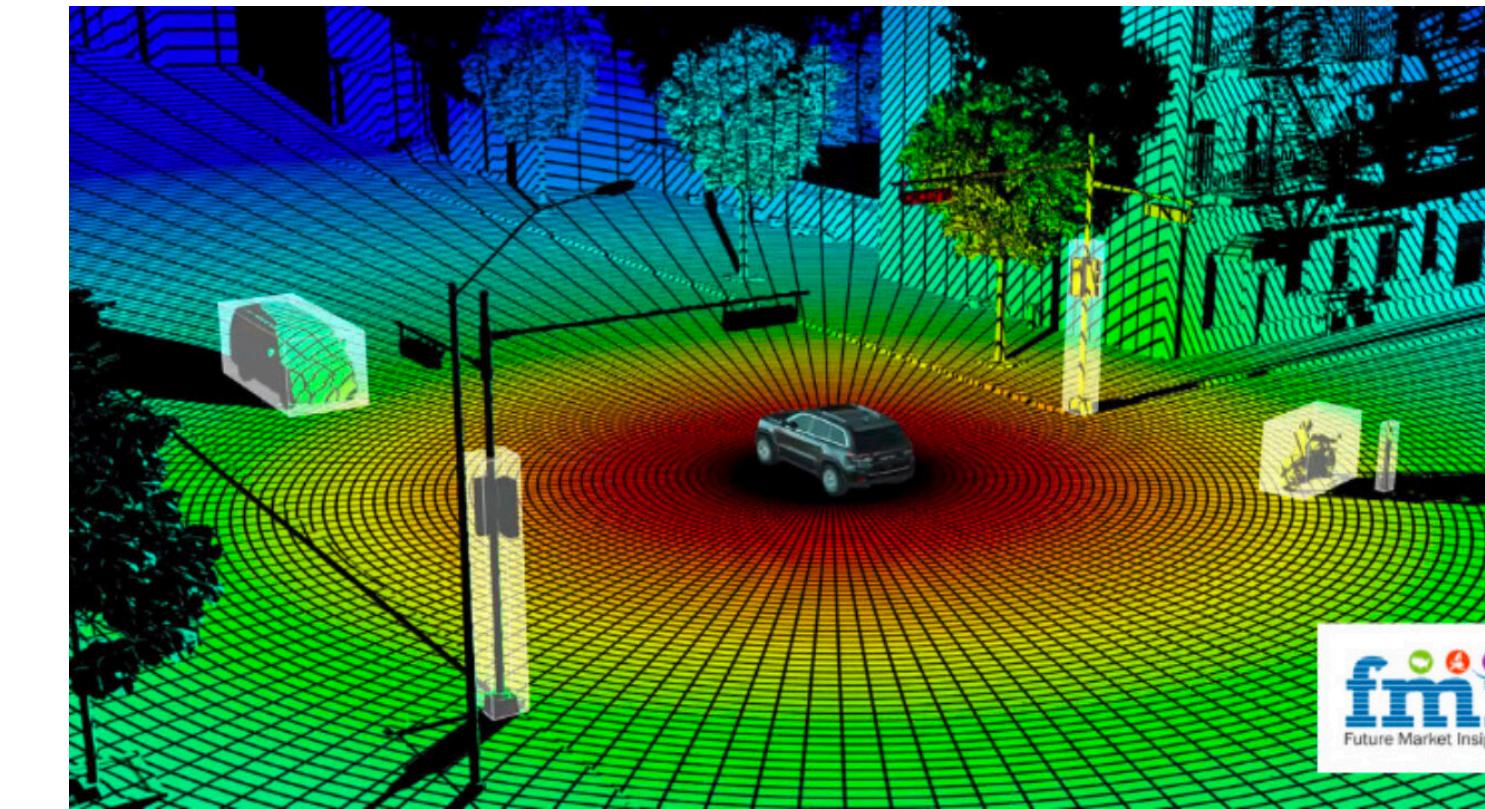
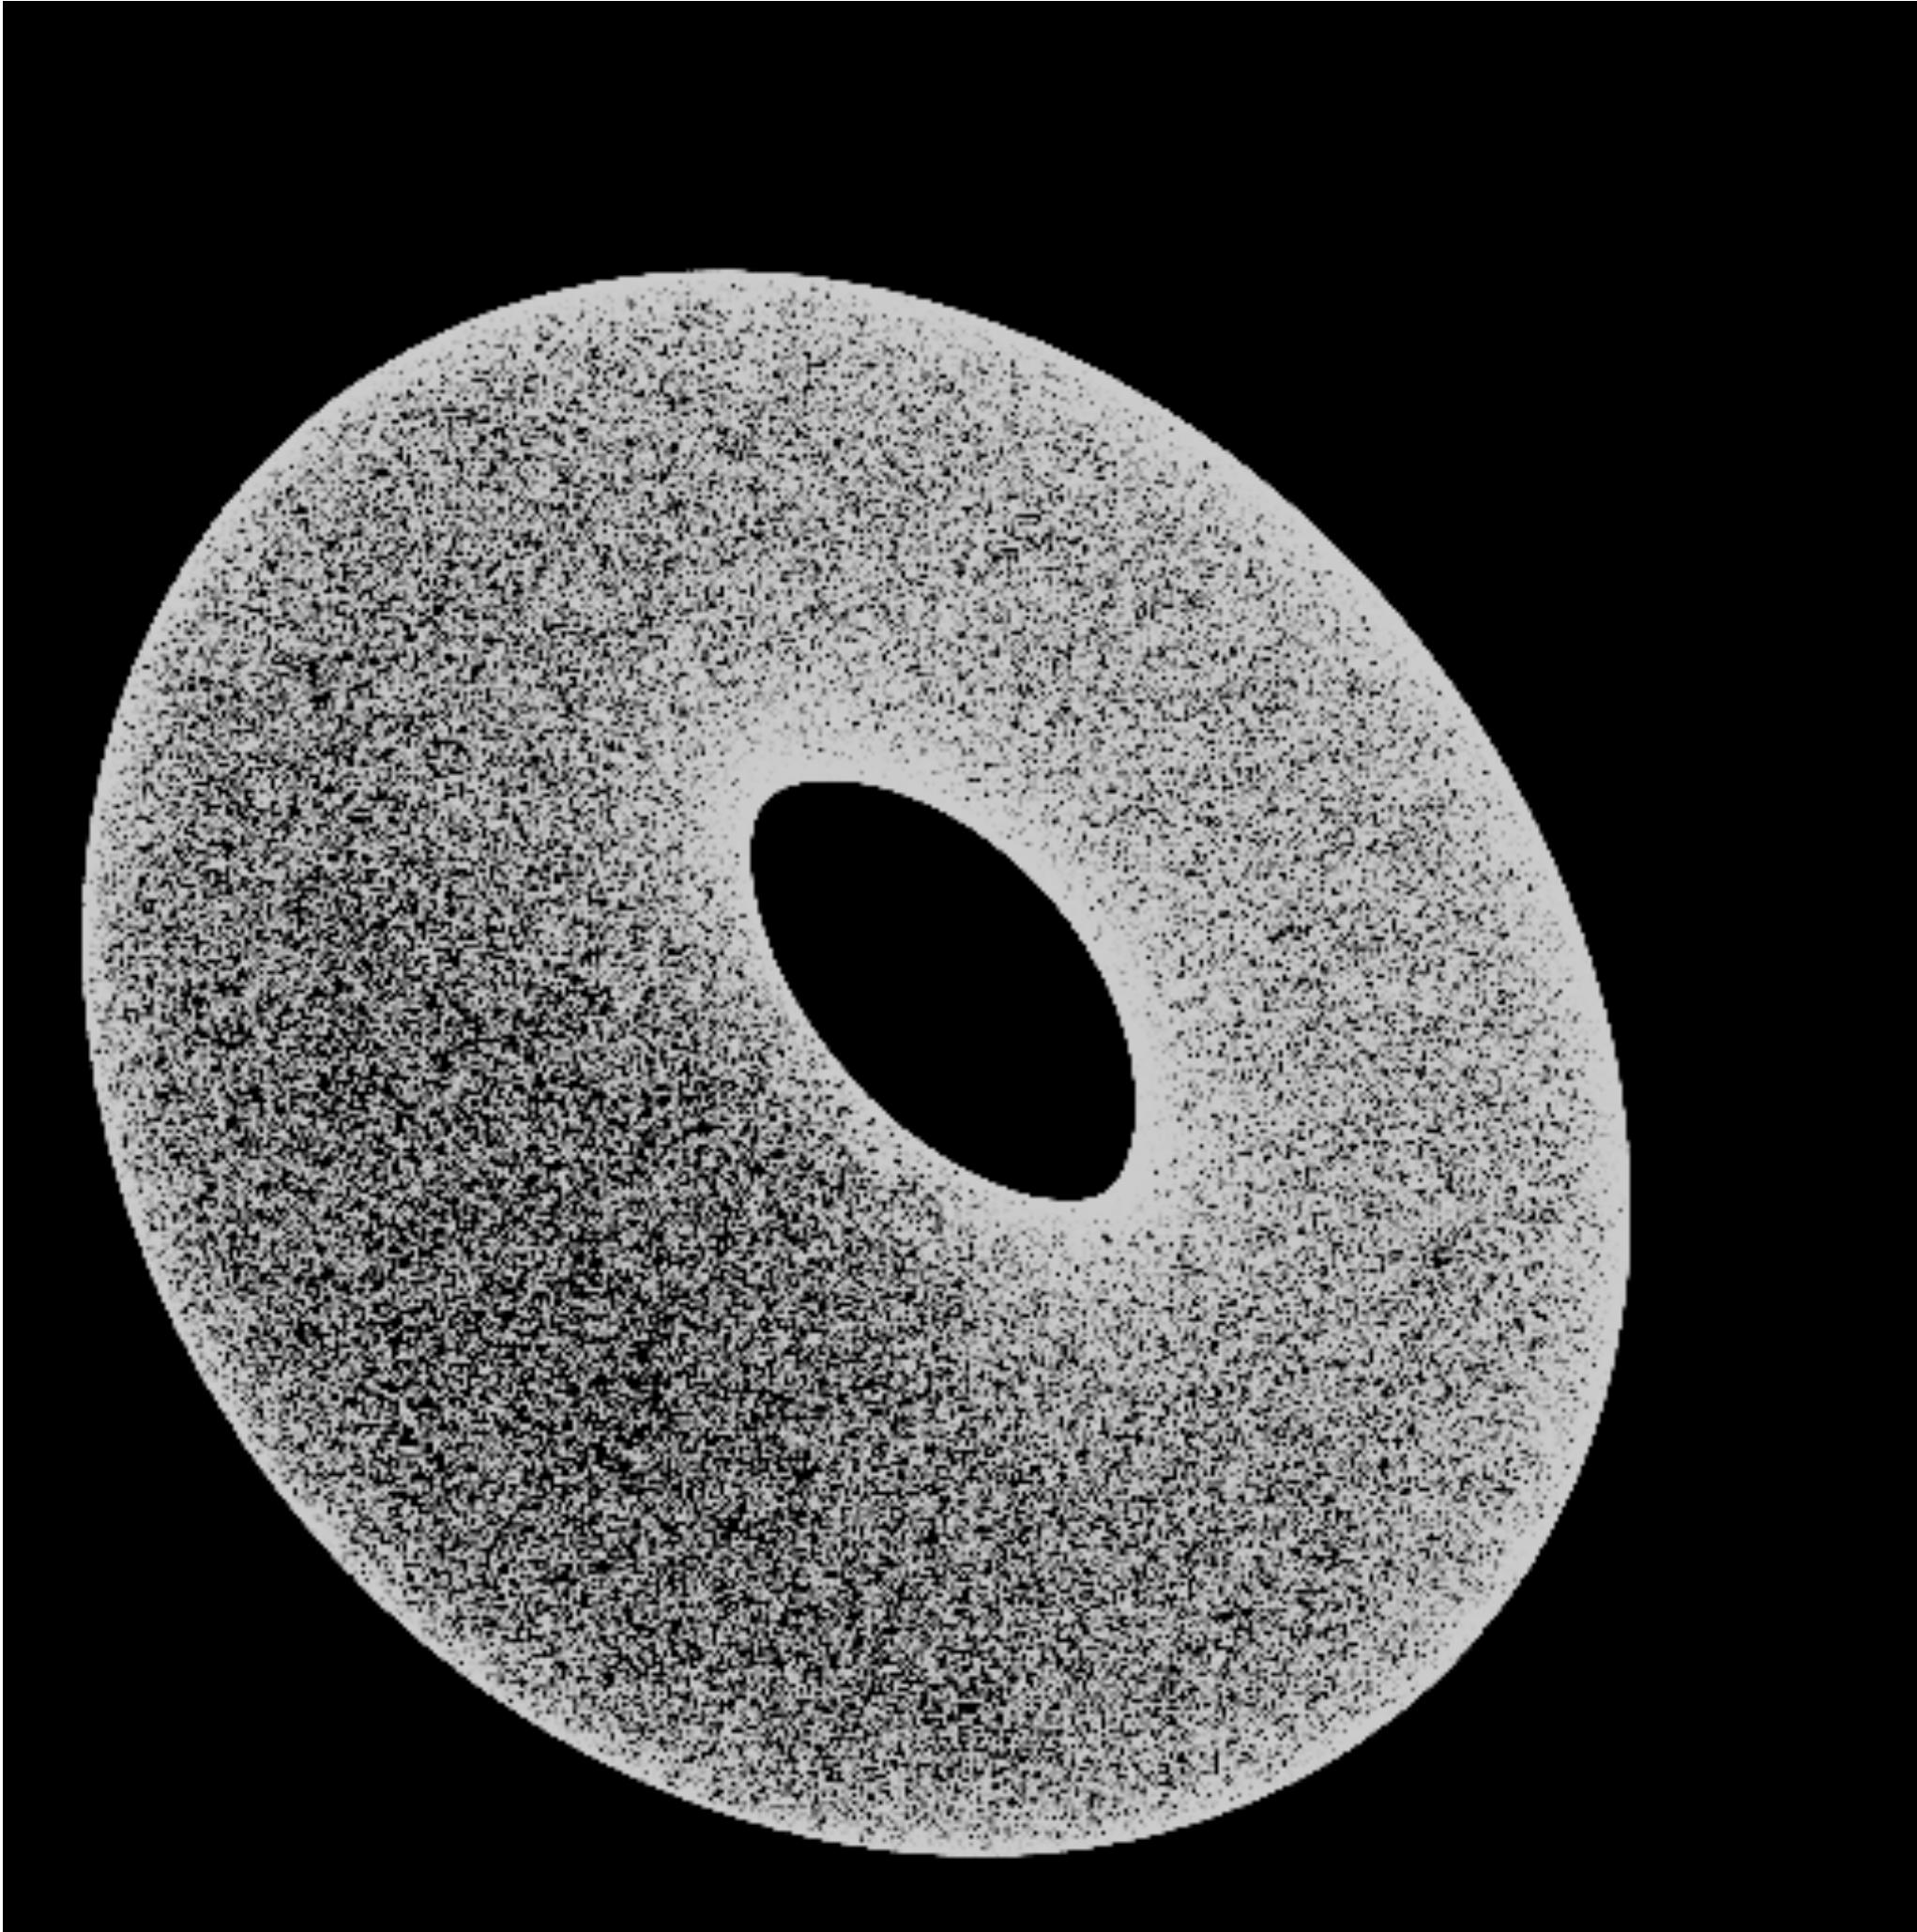
# Point Clouds



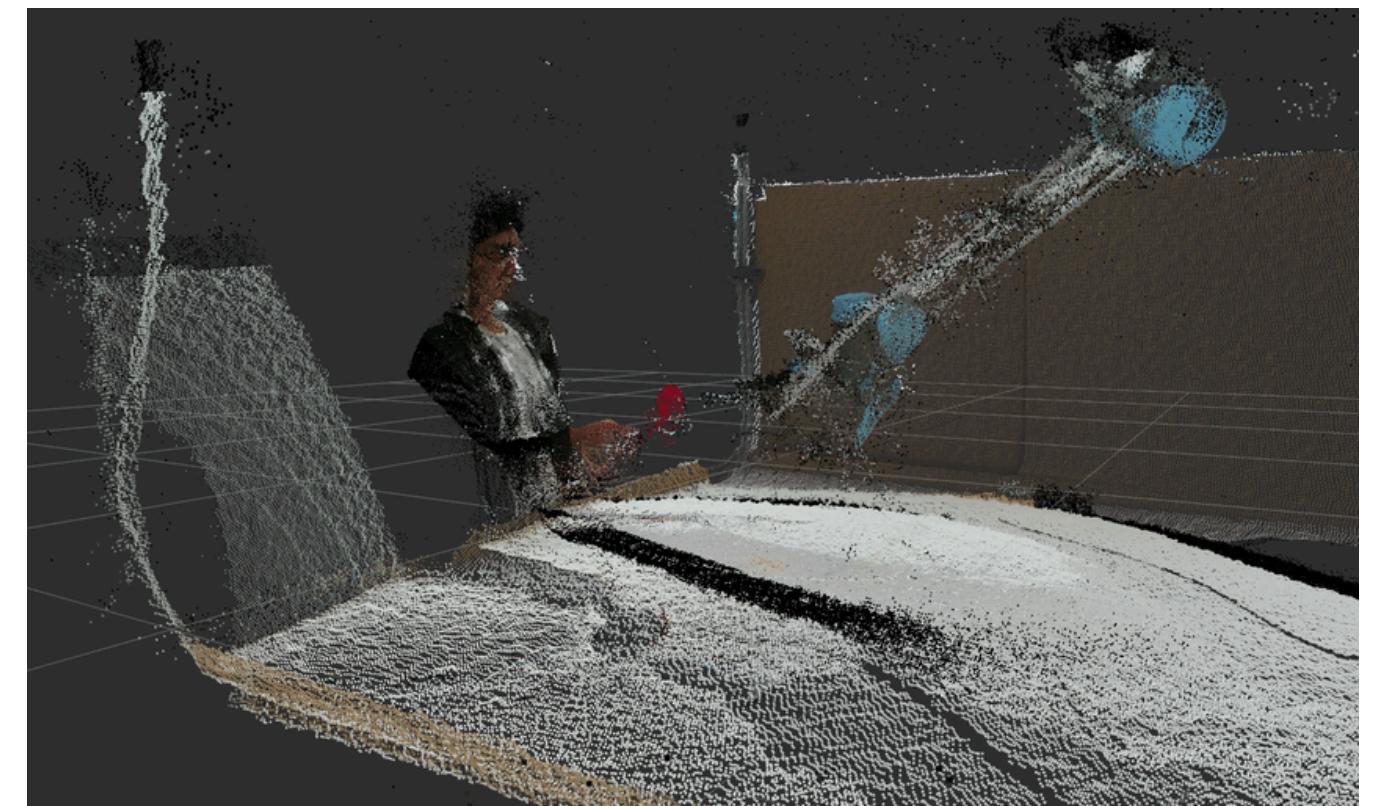
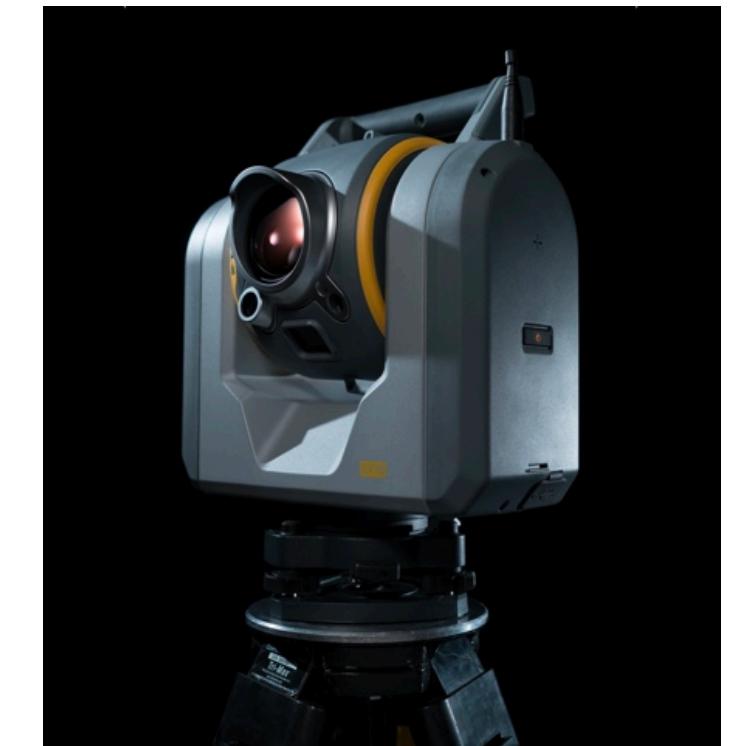
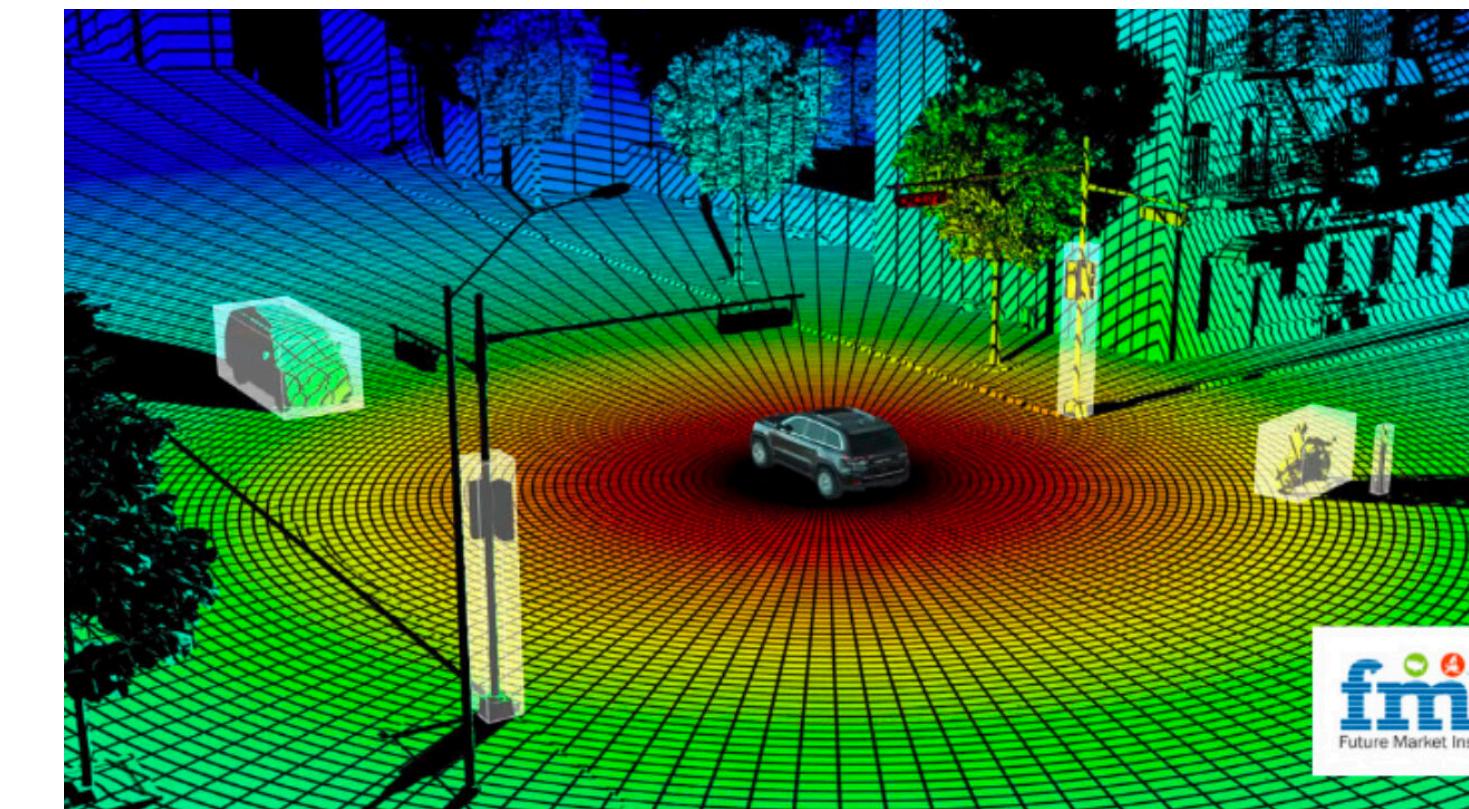
# Point Clouds



# Point Clouds

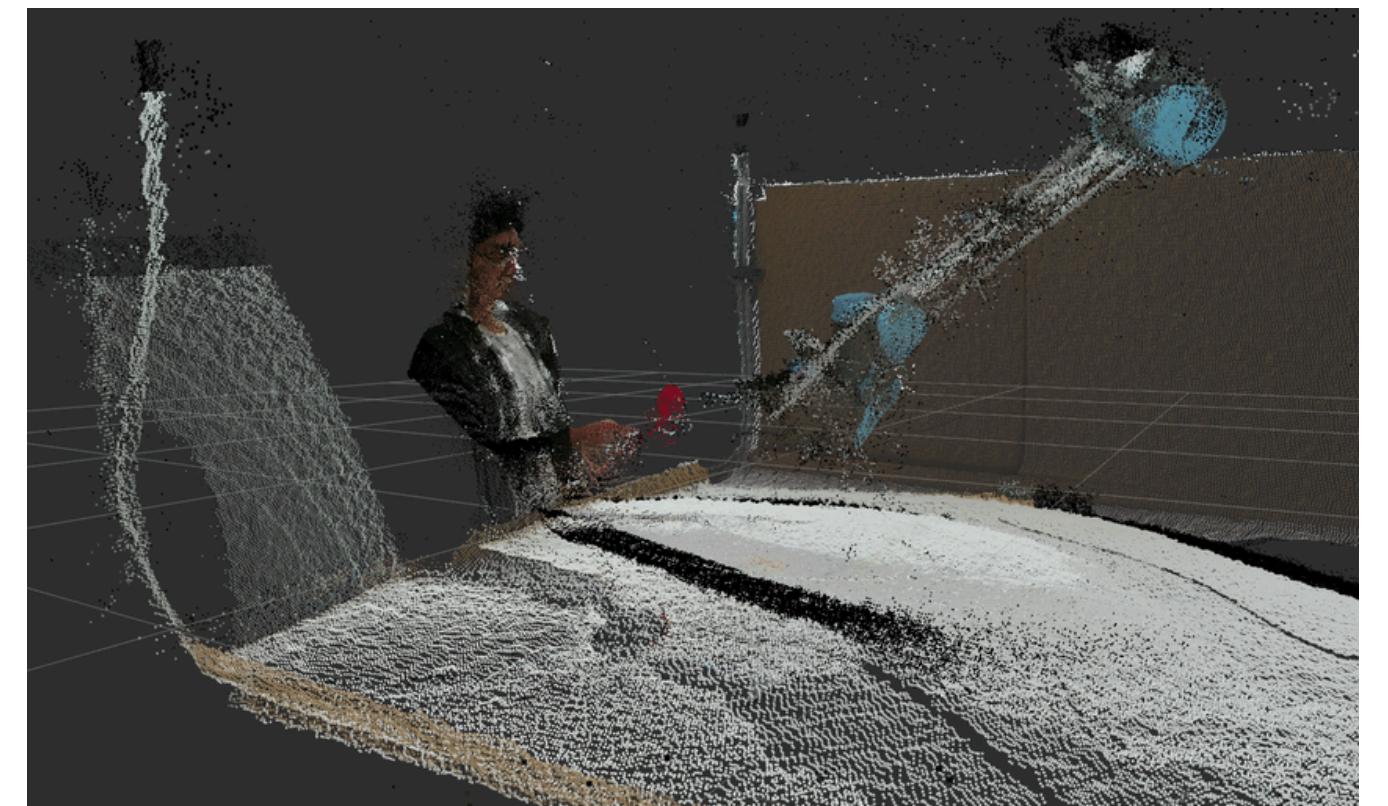
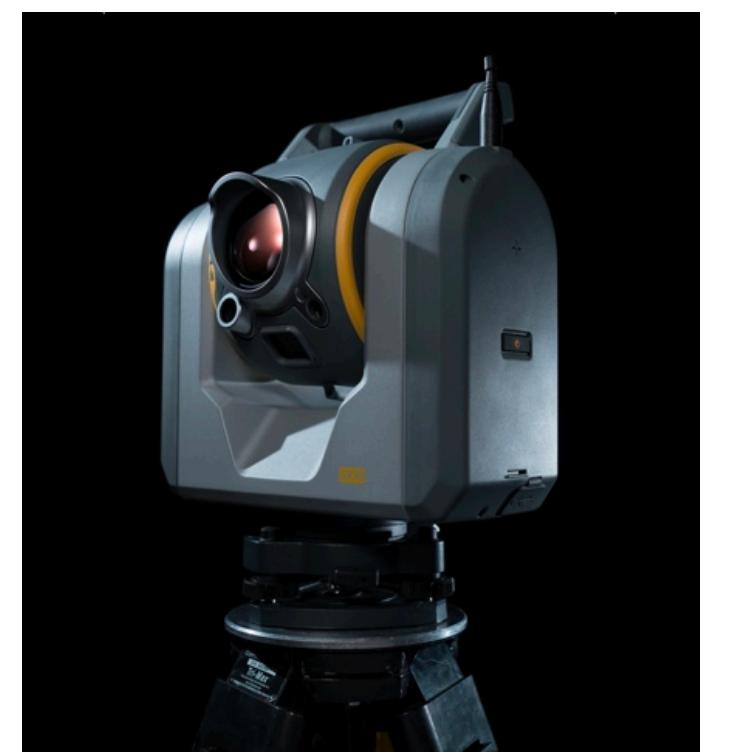
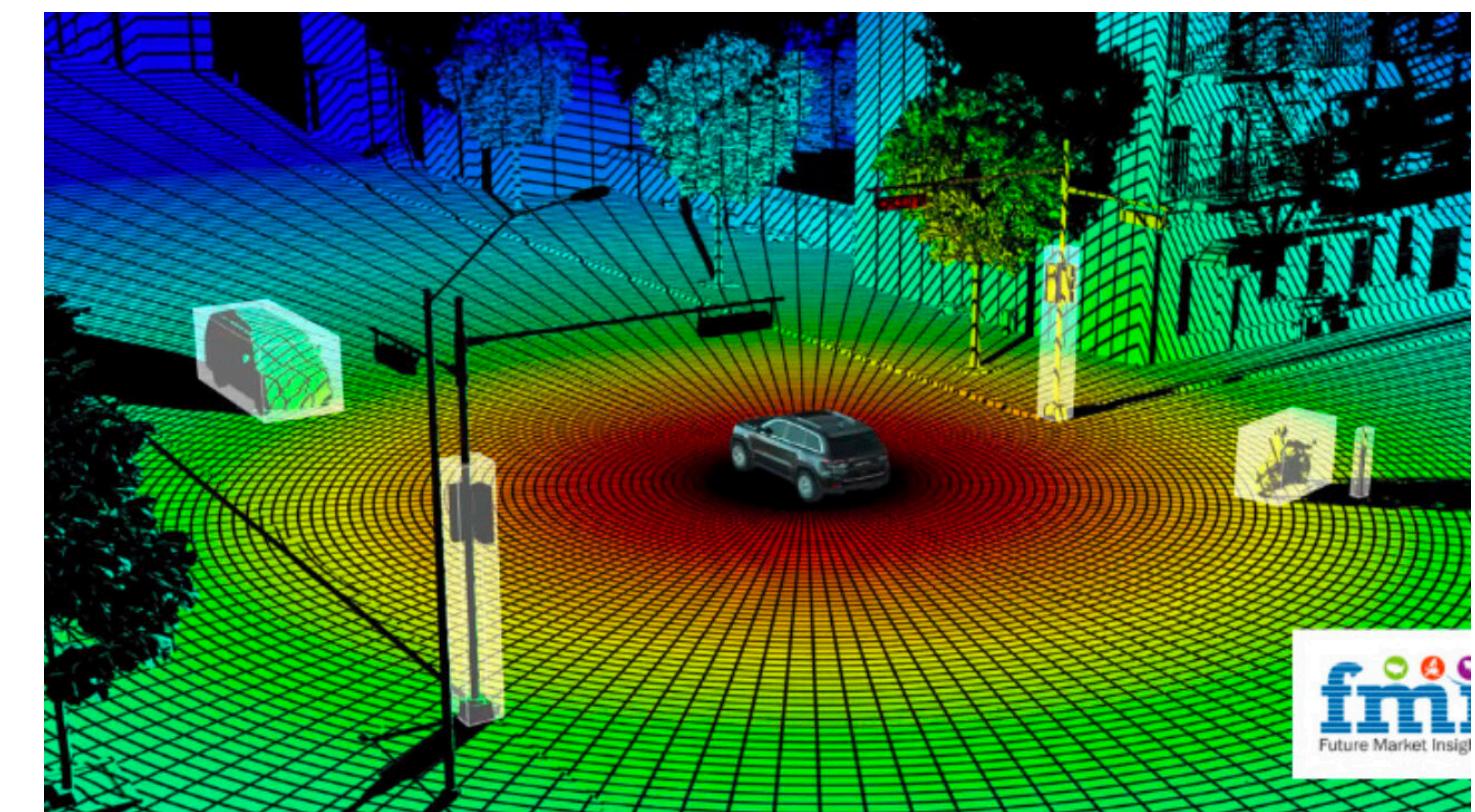


# Point Clouds



# Point Clouds

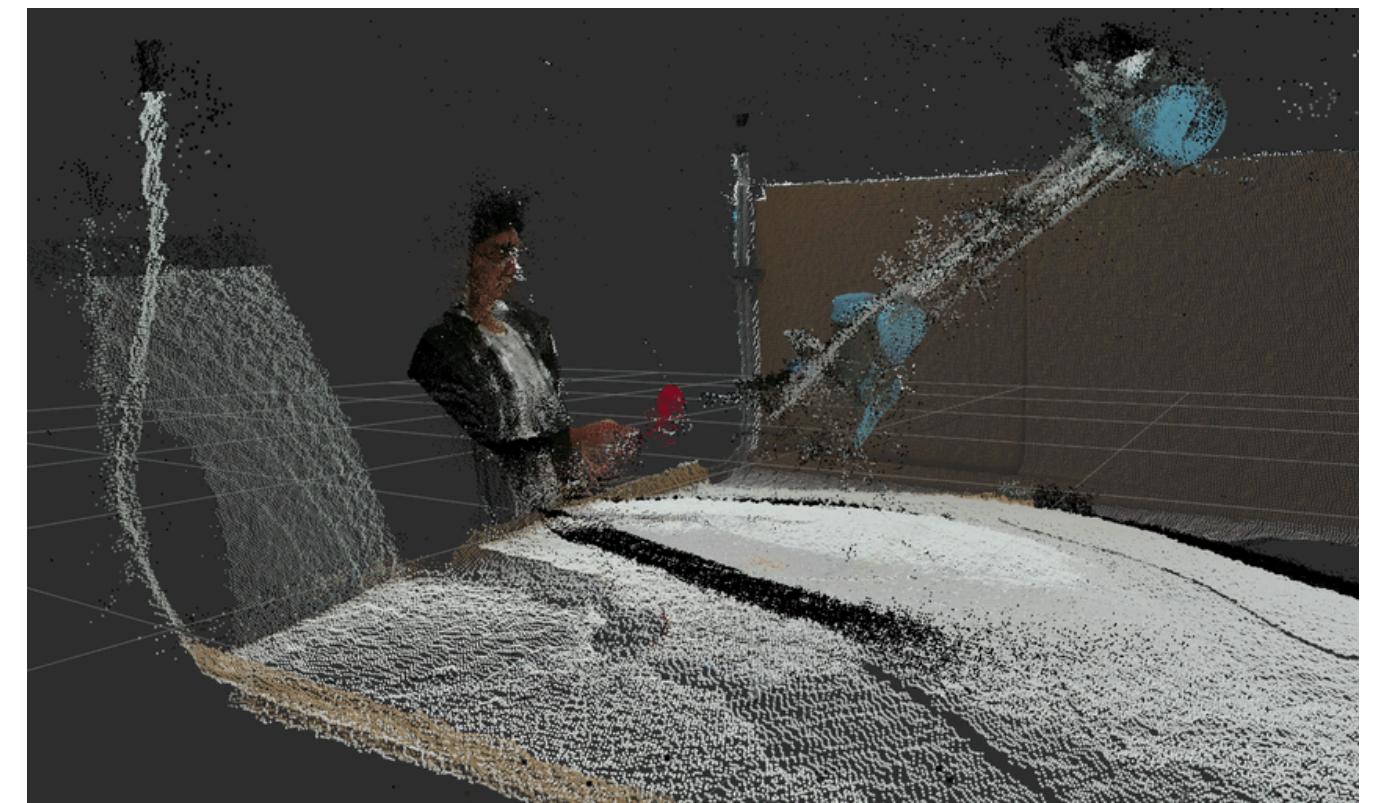
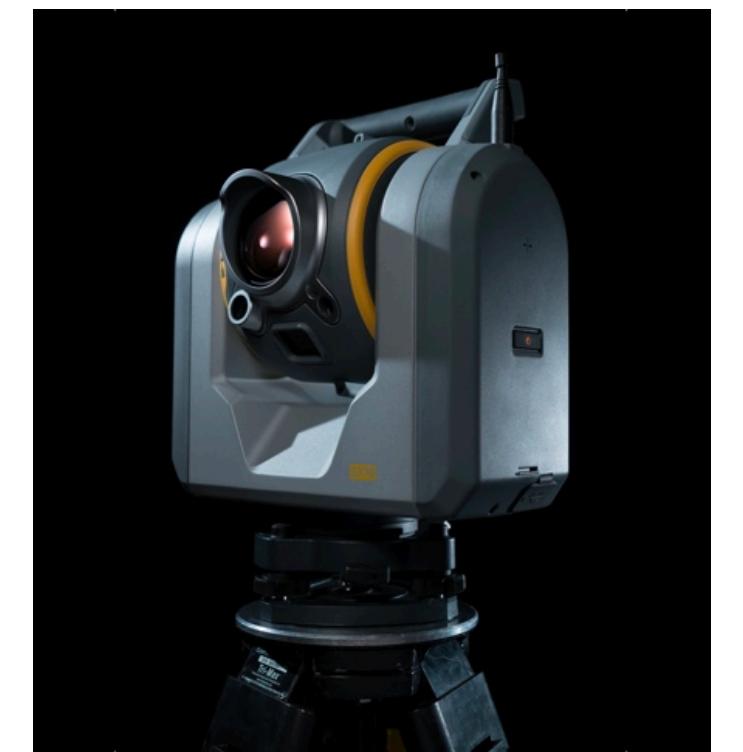
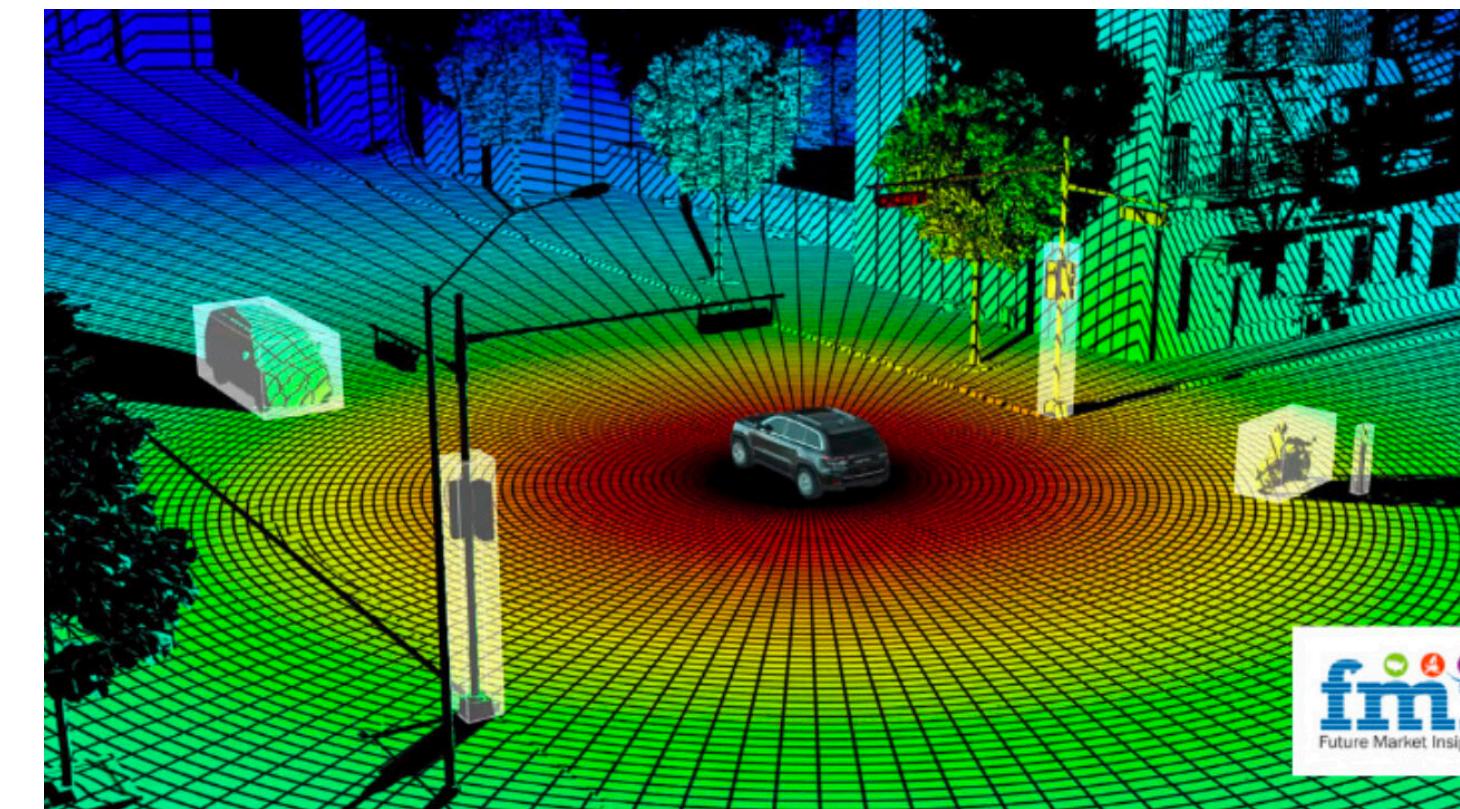
Often acquired in an image-like manner from sensors but can be useful to ‘forget’ this



# Point Clouds

Often acquired in an image-like manner from sensors but can be useful to ‘forget’ this

e.g. point clouds from multiple views can make a single point cloud

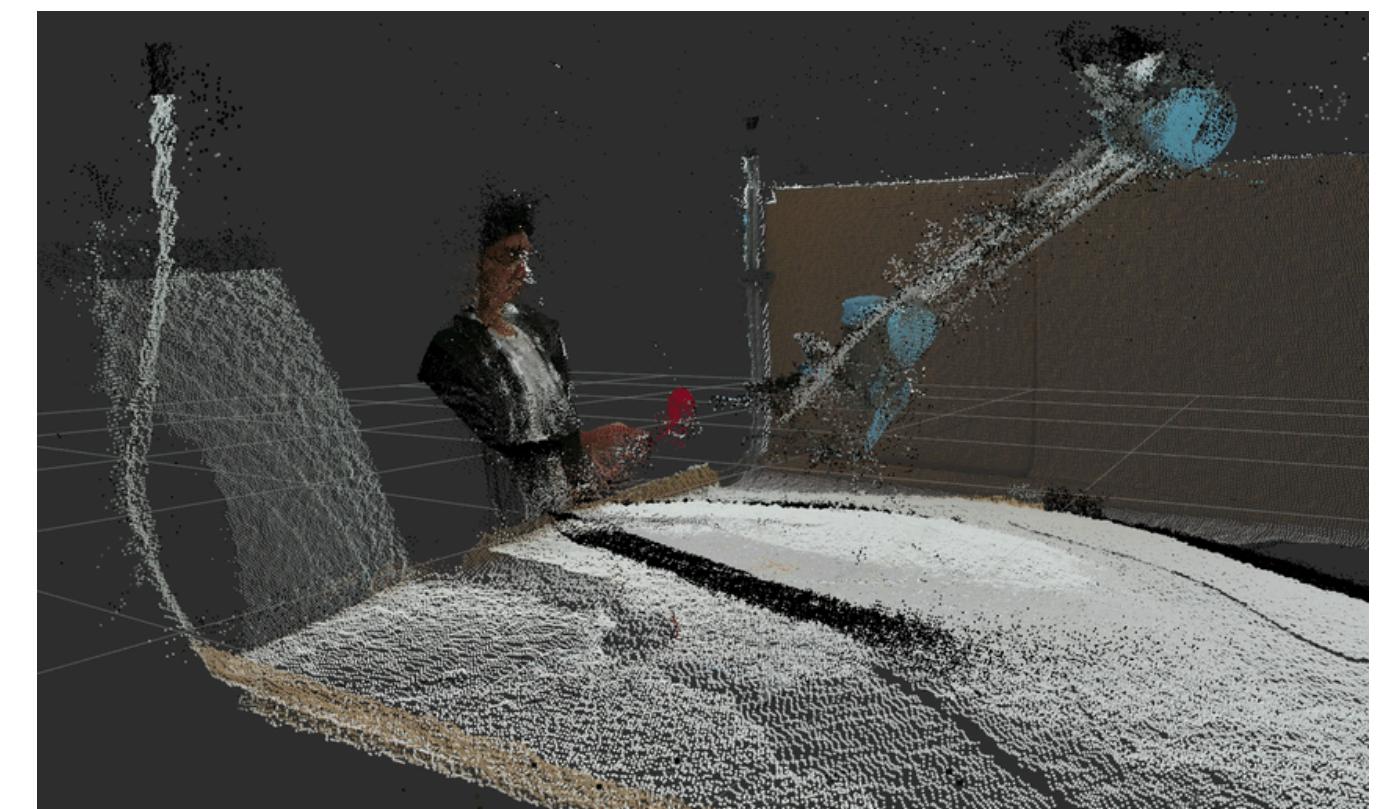
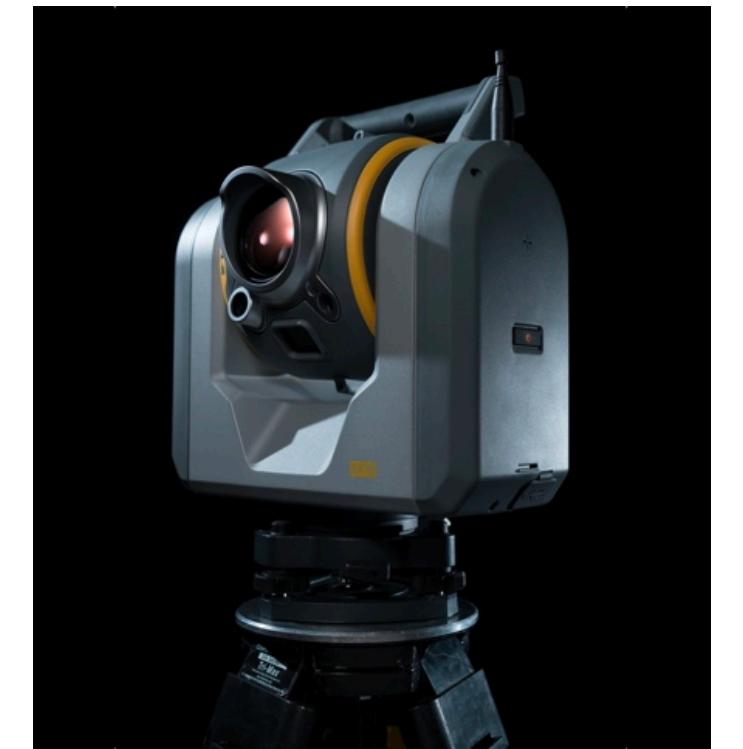
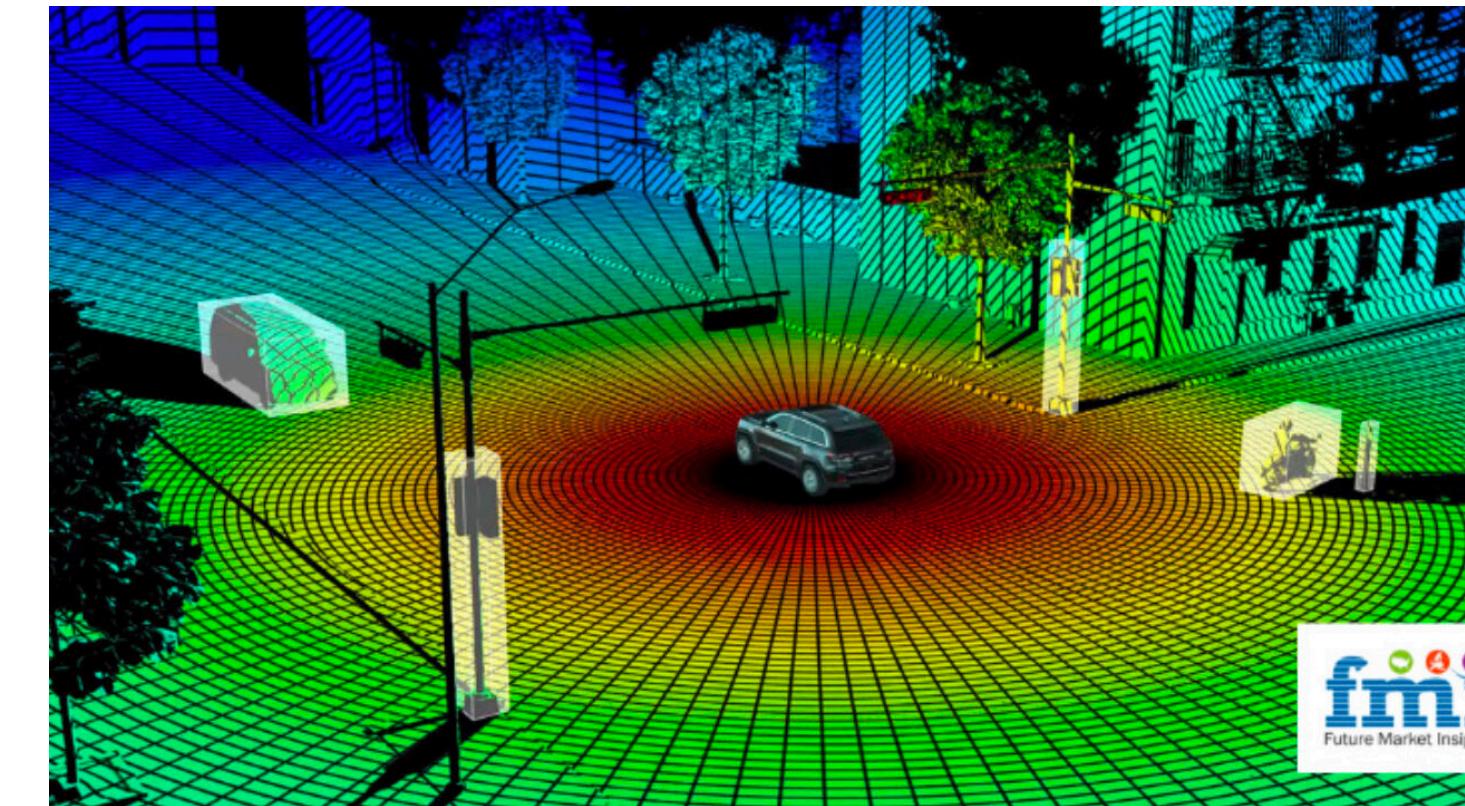


# Point Clouds

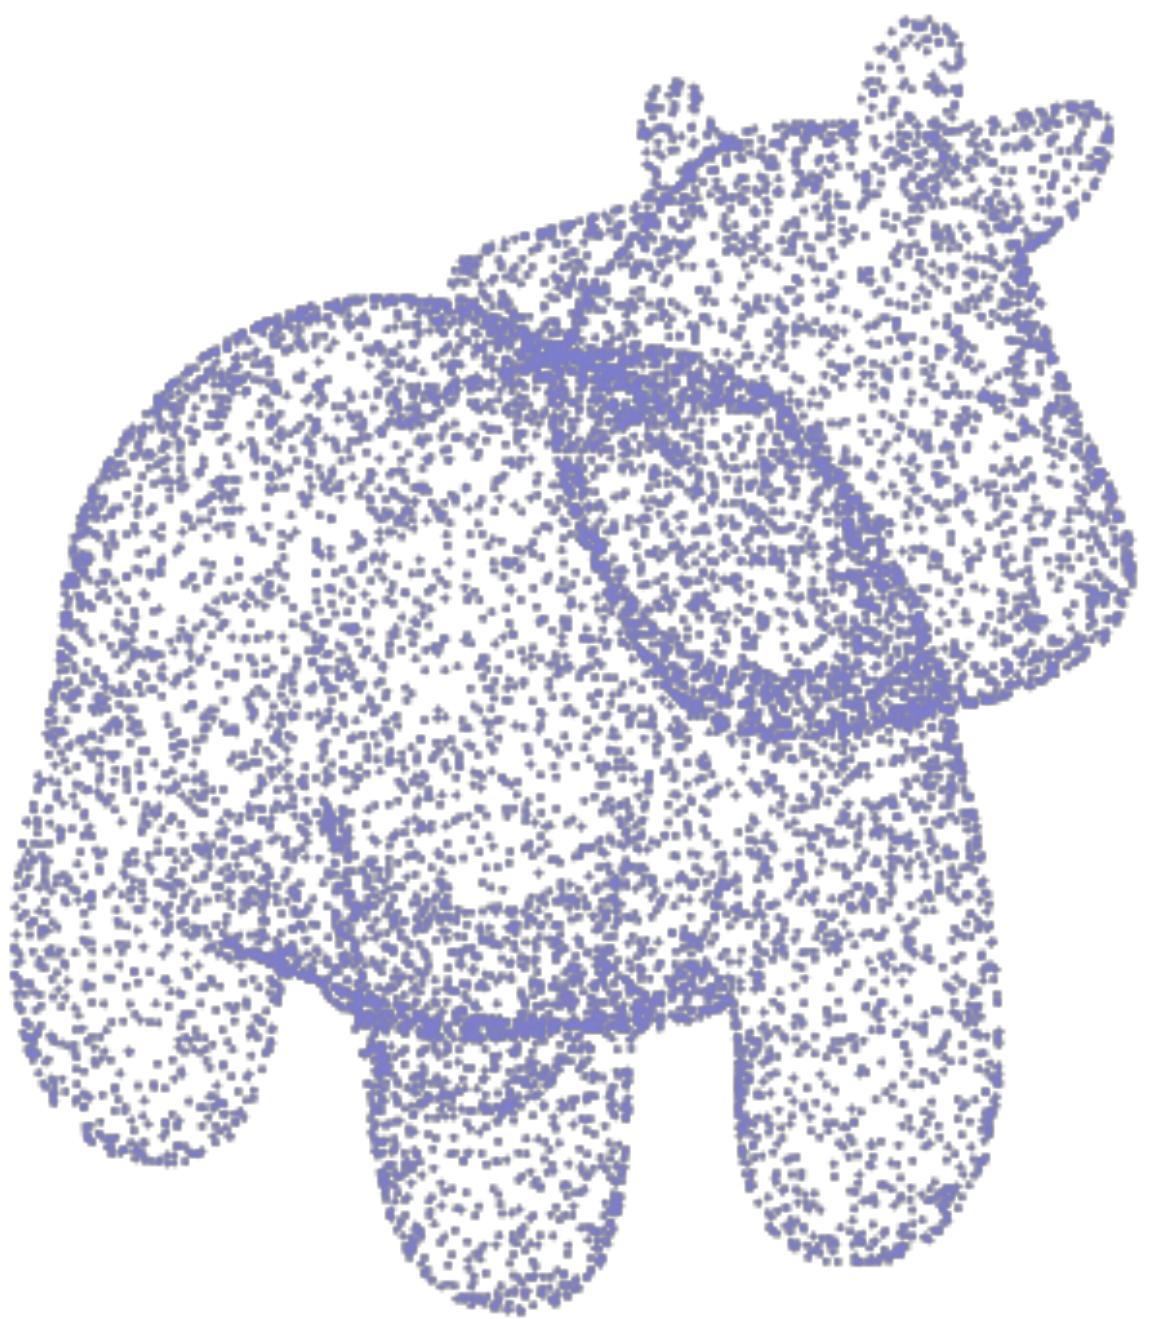
Often acquired in an image-like manner from sensors but can be useful to ‘forget’ this

e.g. point clouds from multiple views can make a single point cloud

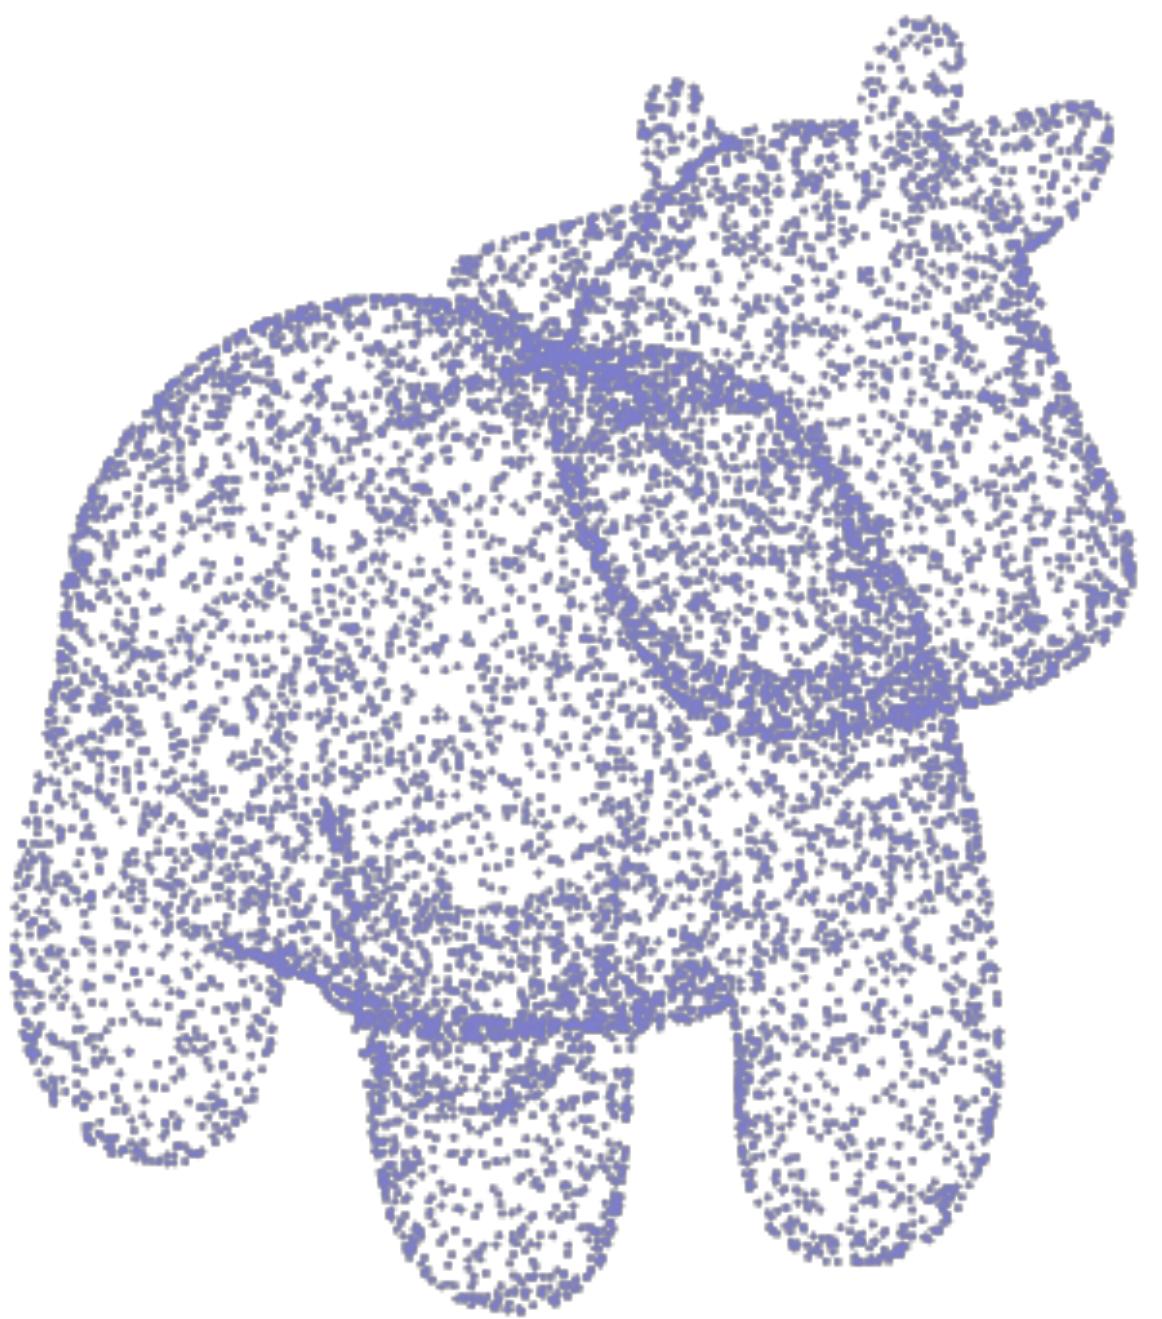
can help unify approaches independent of sensing modality



# Point Clouds

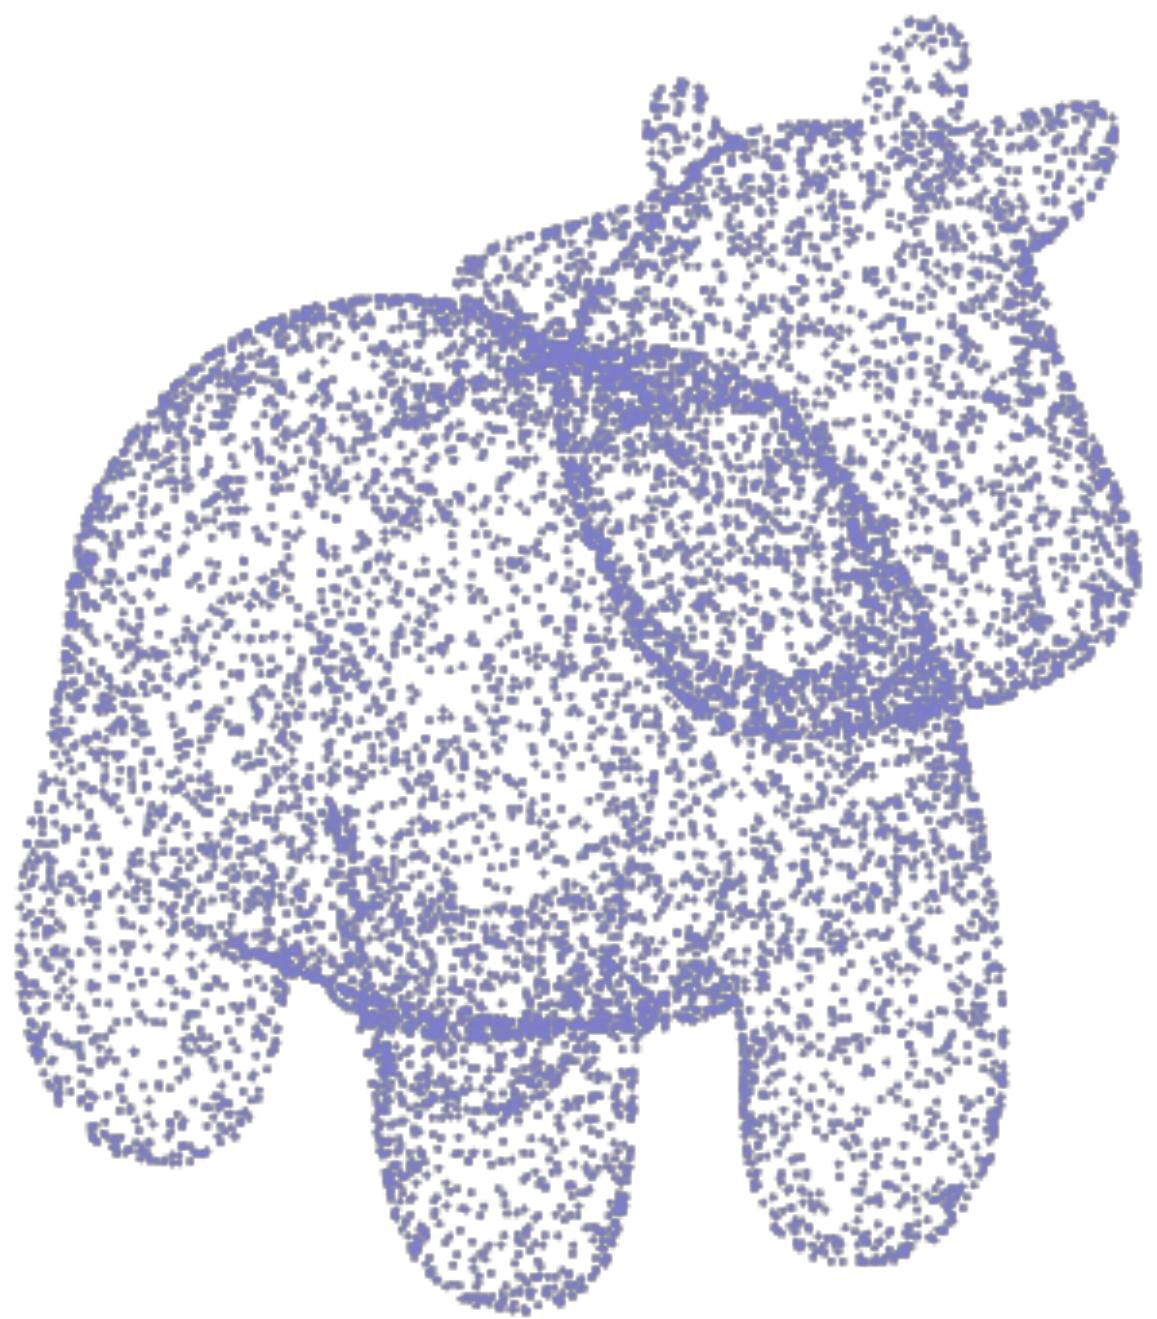


# Point Clouds



$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

# Point Clouds

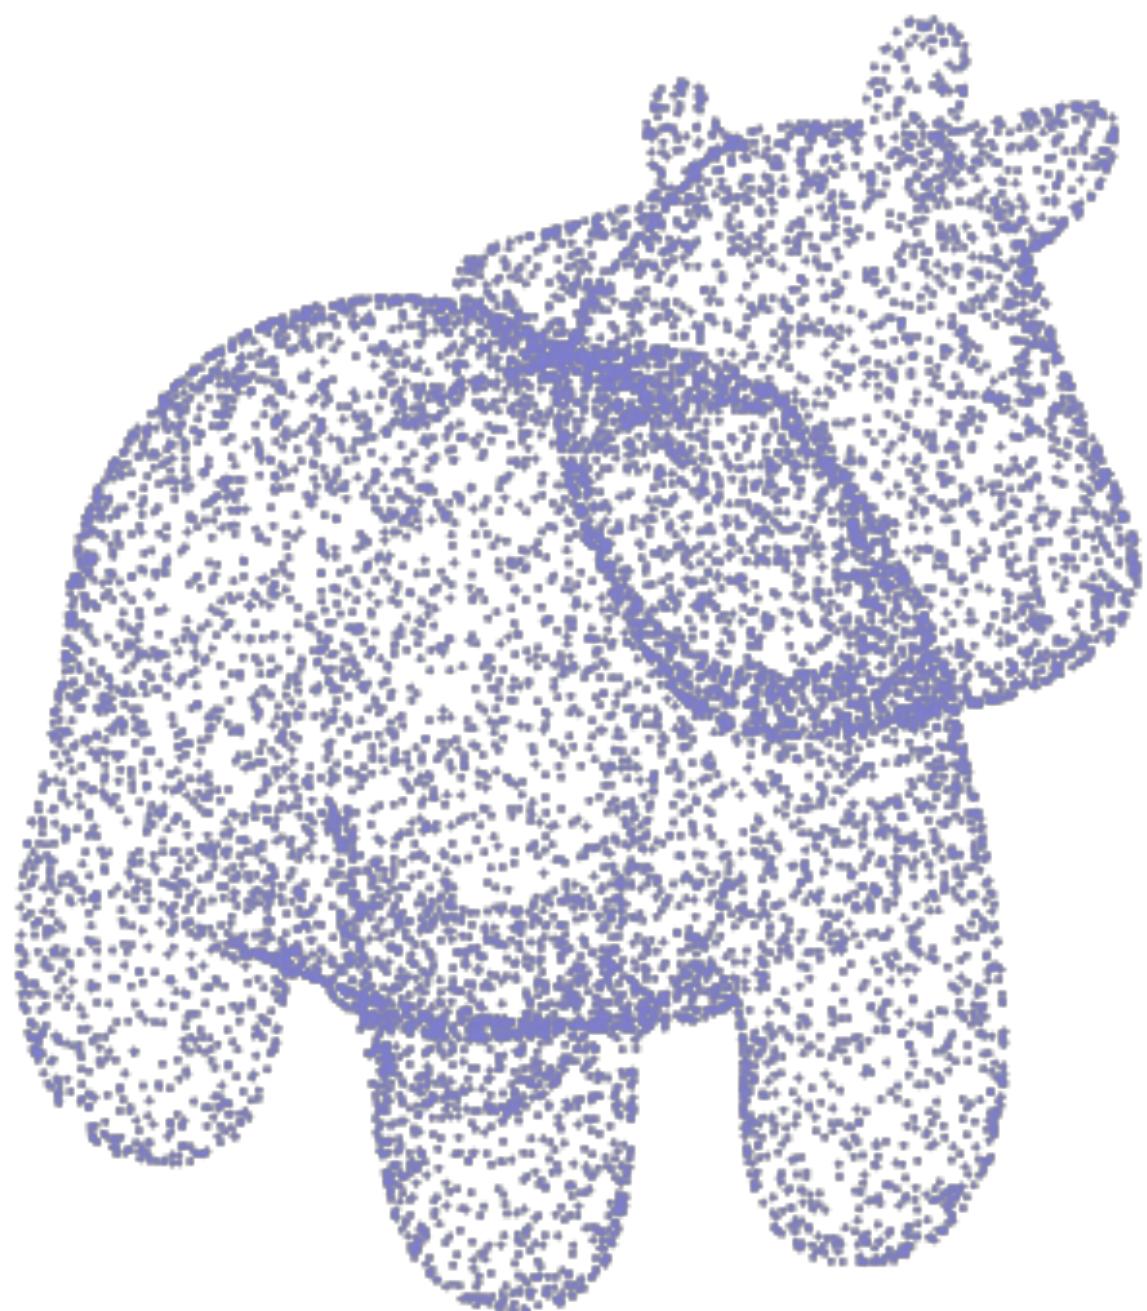


Permutation  $\sigma$



$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

# Point Clouds



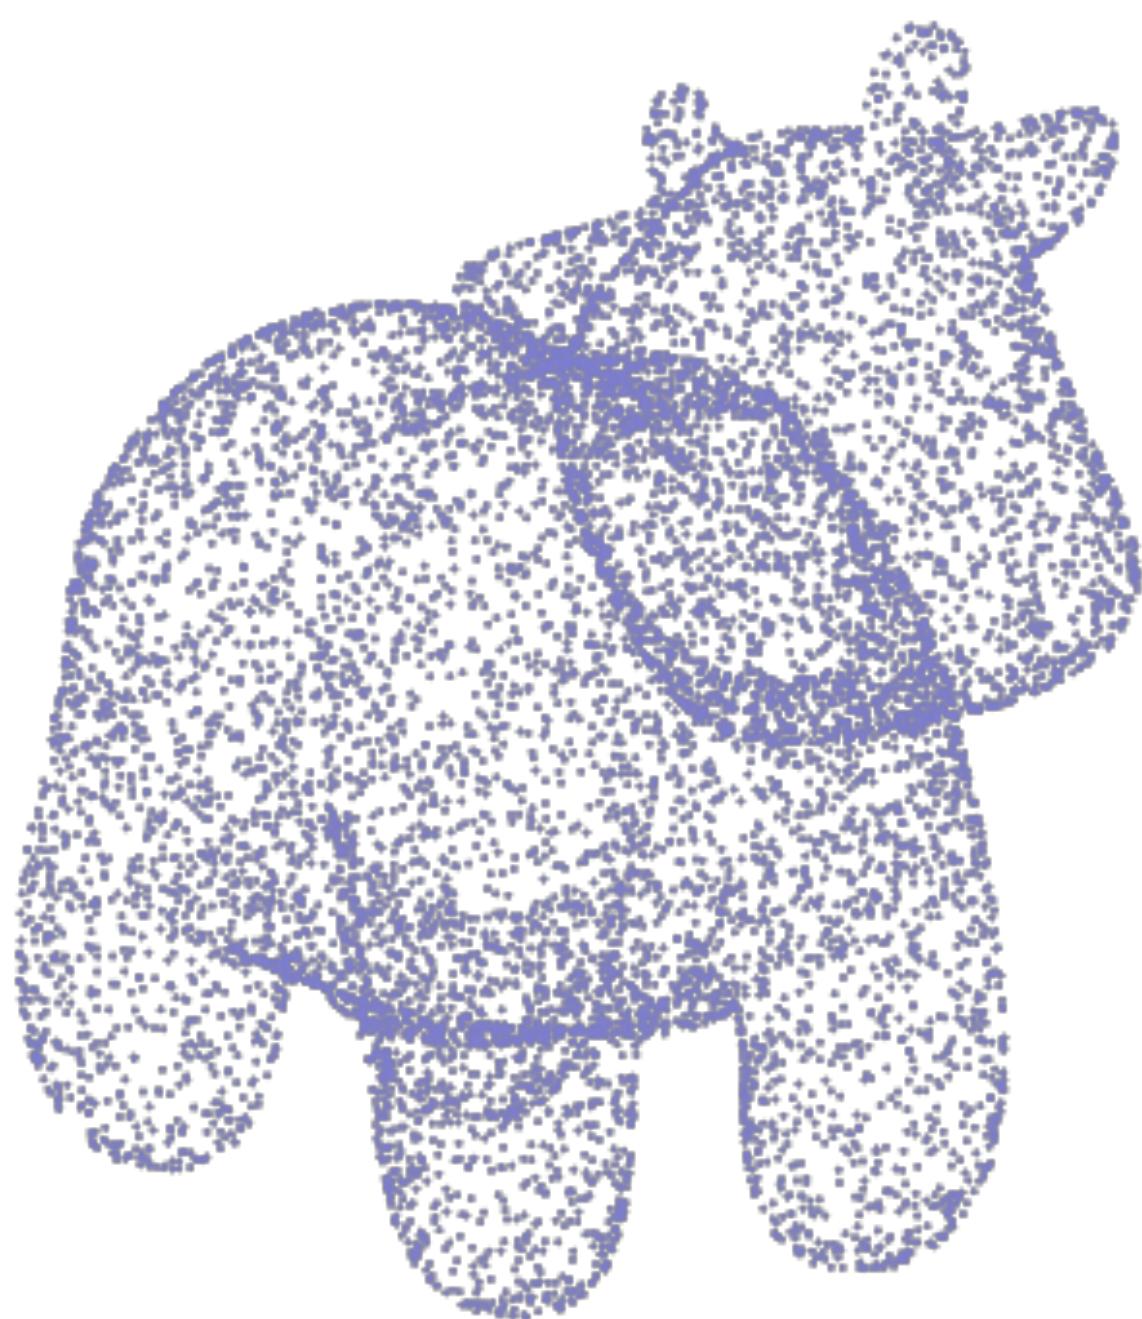
Permutation  $\sigma$



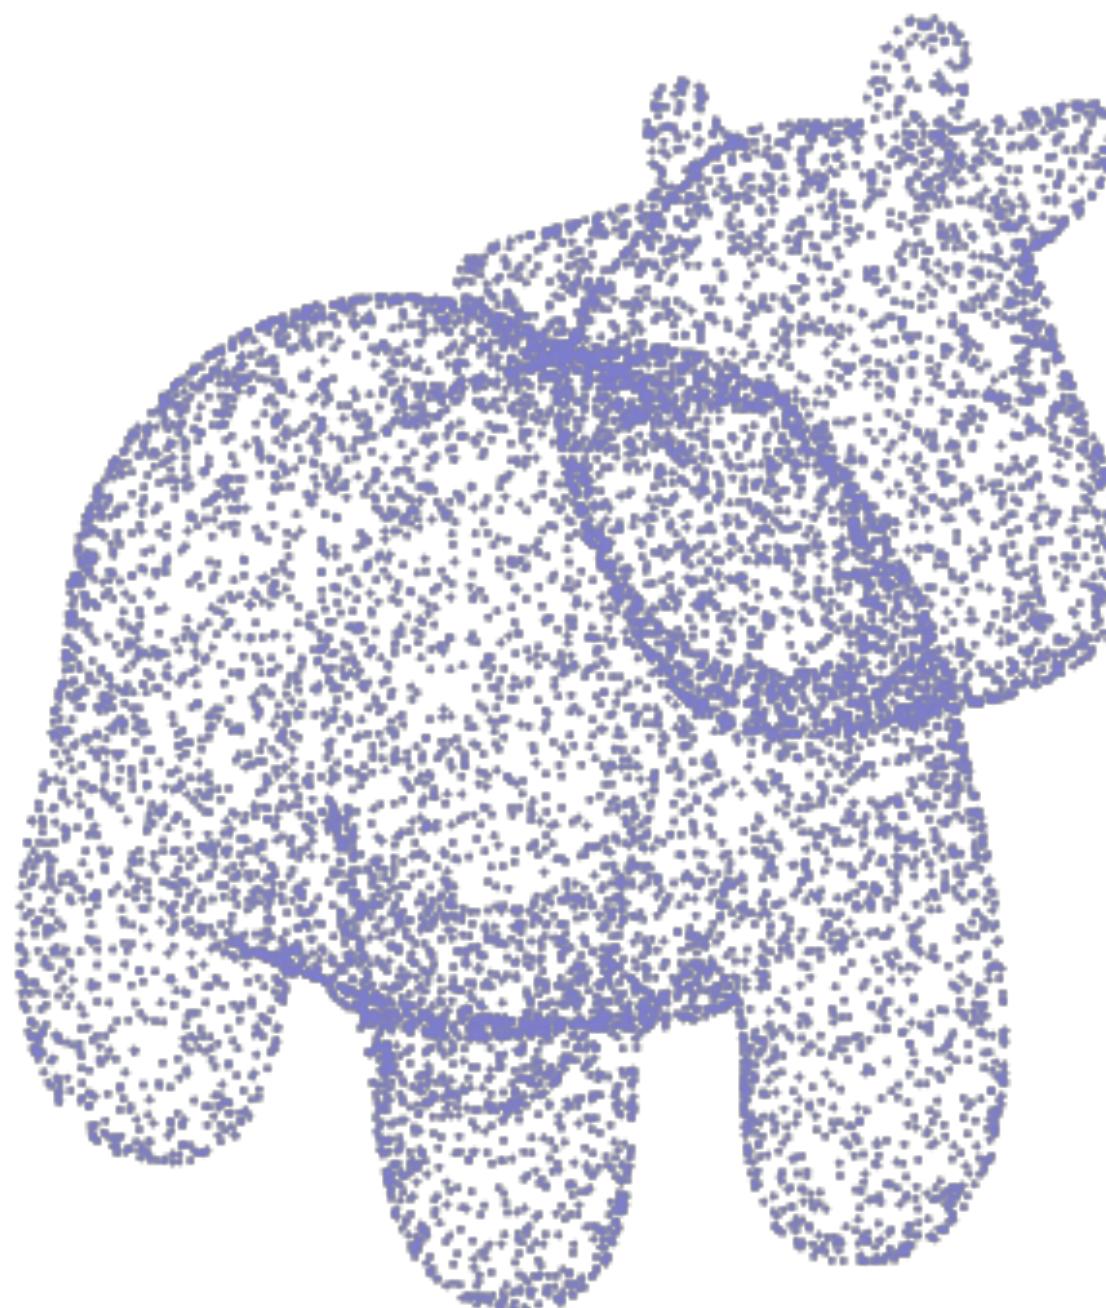
$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

$$\{\mathbf{p}_{\sigma(1)}, \mathbf{p}_{\sigma(2)}, \dots, \mathbf{p}_{\sigma(N)}\}$$

# Point Clouds



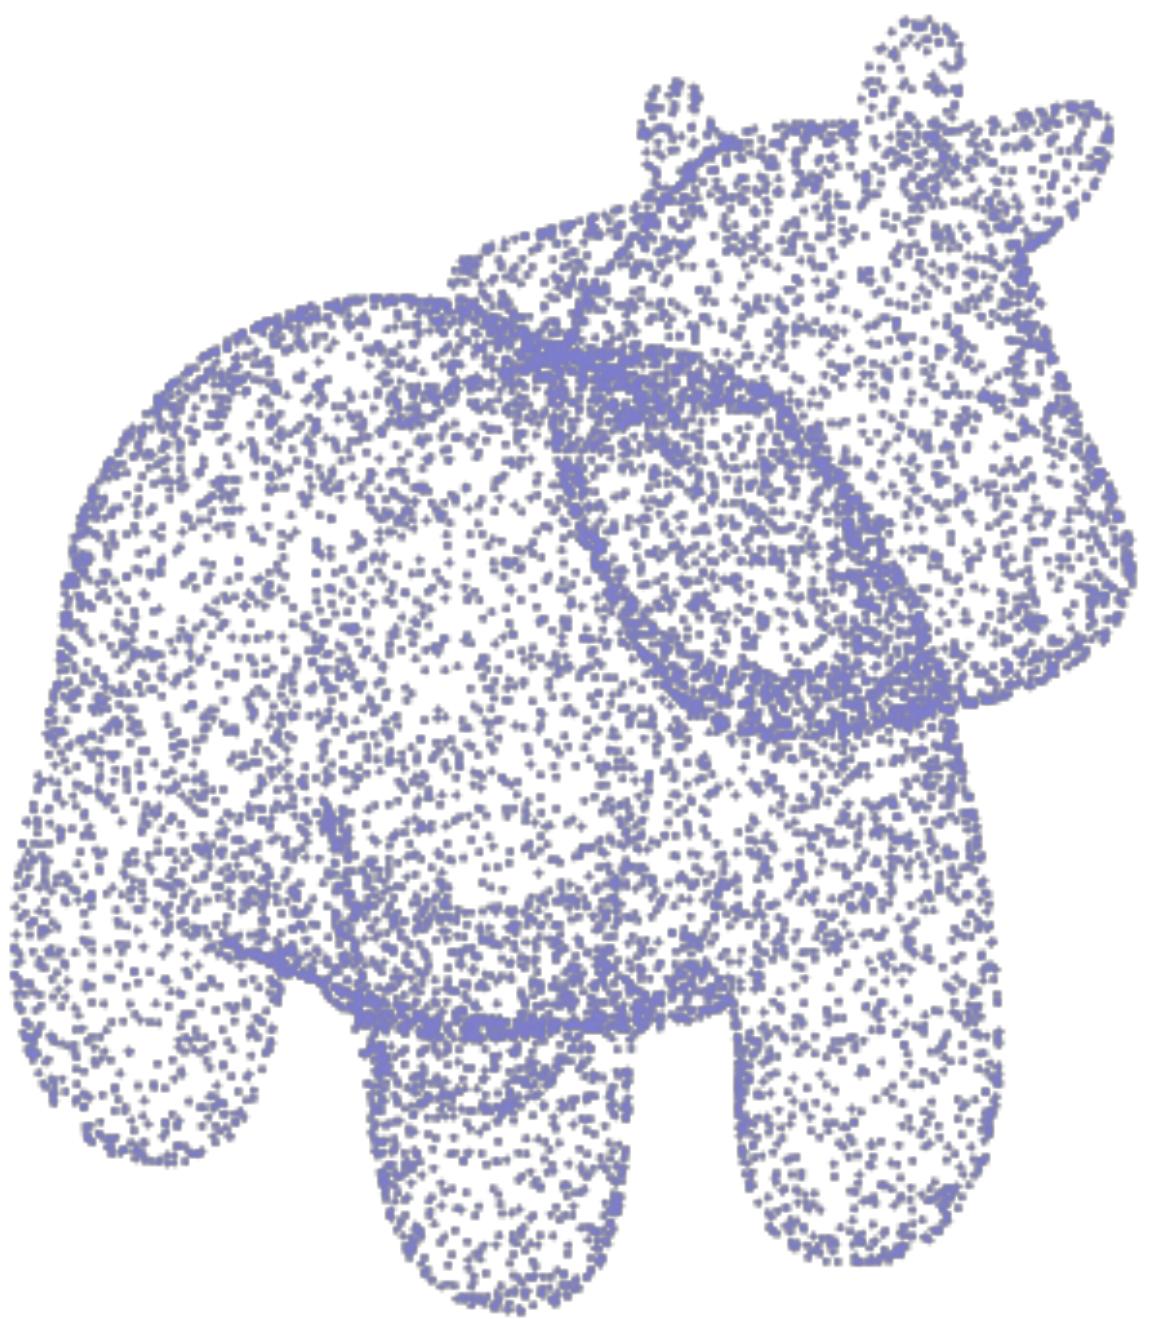
Permutation  $\sigma$   
→



$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

$$\{\mathbf{p}_{\sigma(1)}, \mathbf{p}_{\sigma(2)}, \dots, \mathbf{p}_{\sigma(N)}\}$$

# Point Clouds



$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

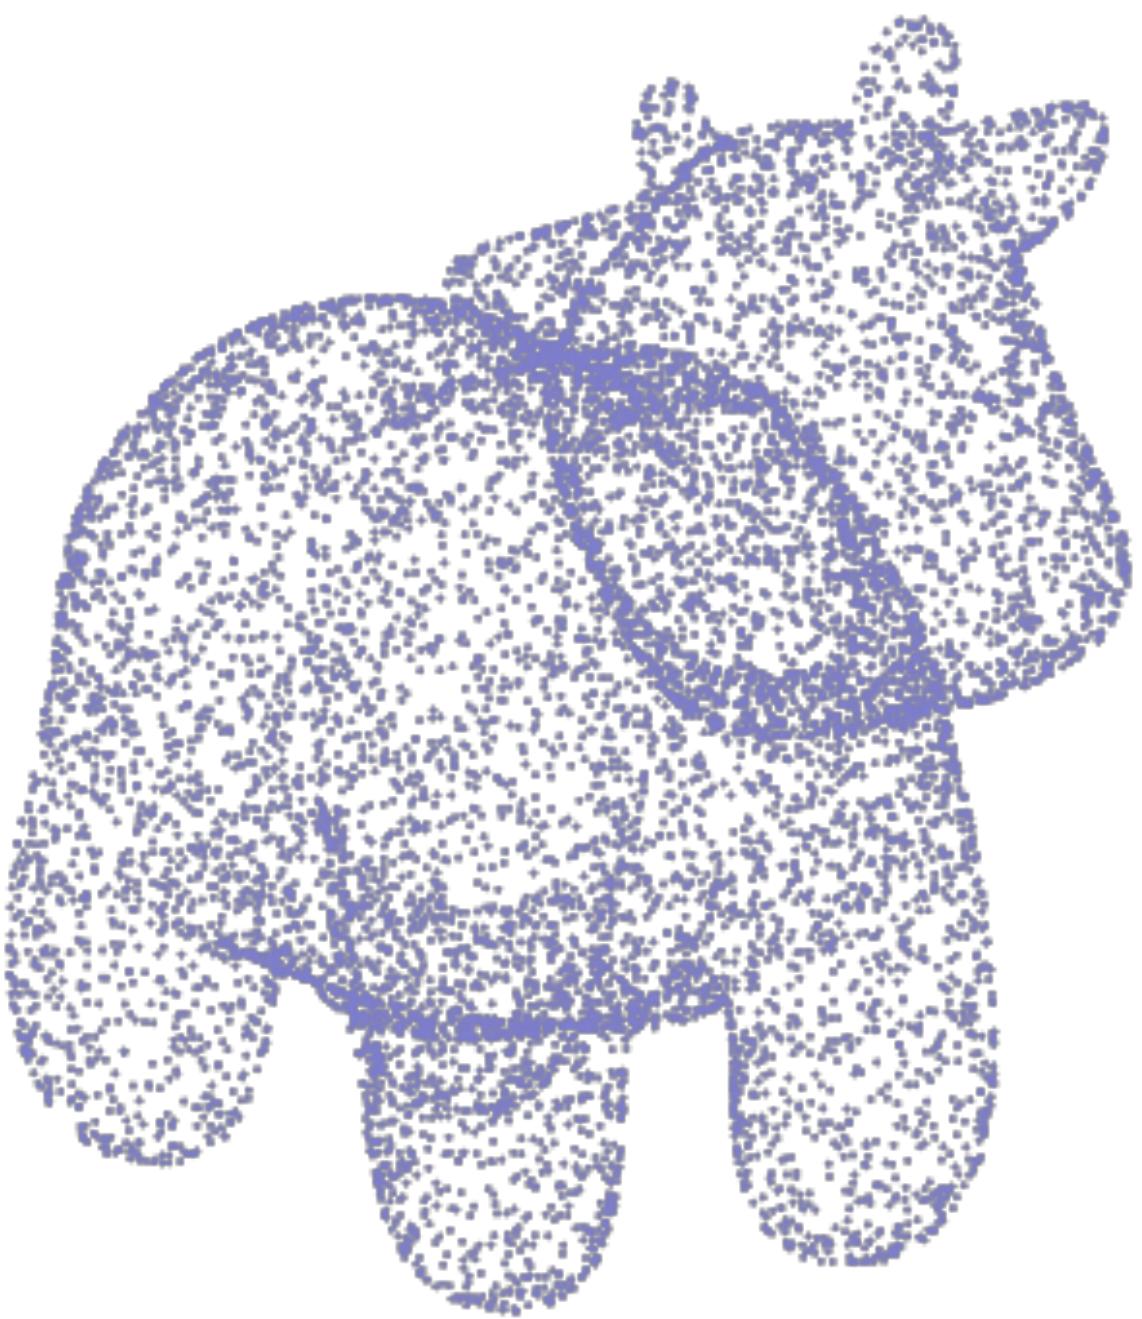
# Point Clouds



$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

**Unordered set of points**

# Point Clouds

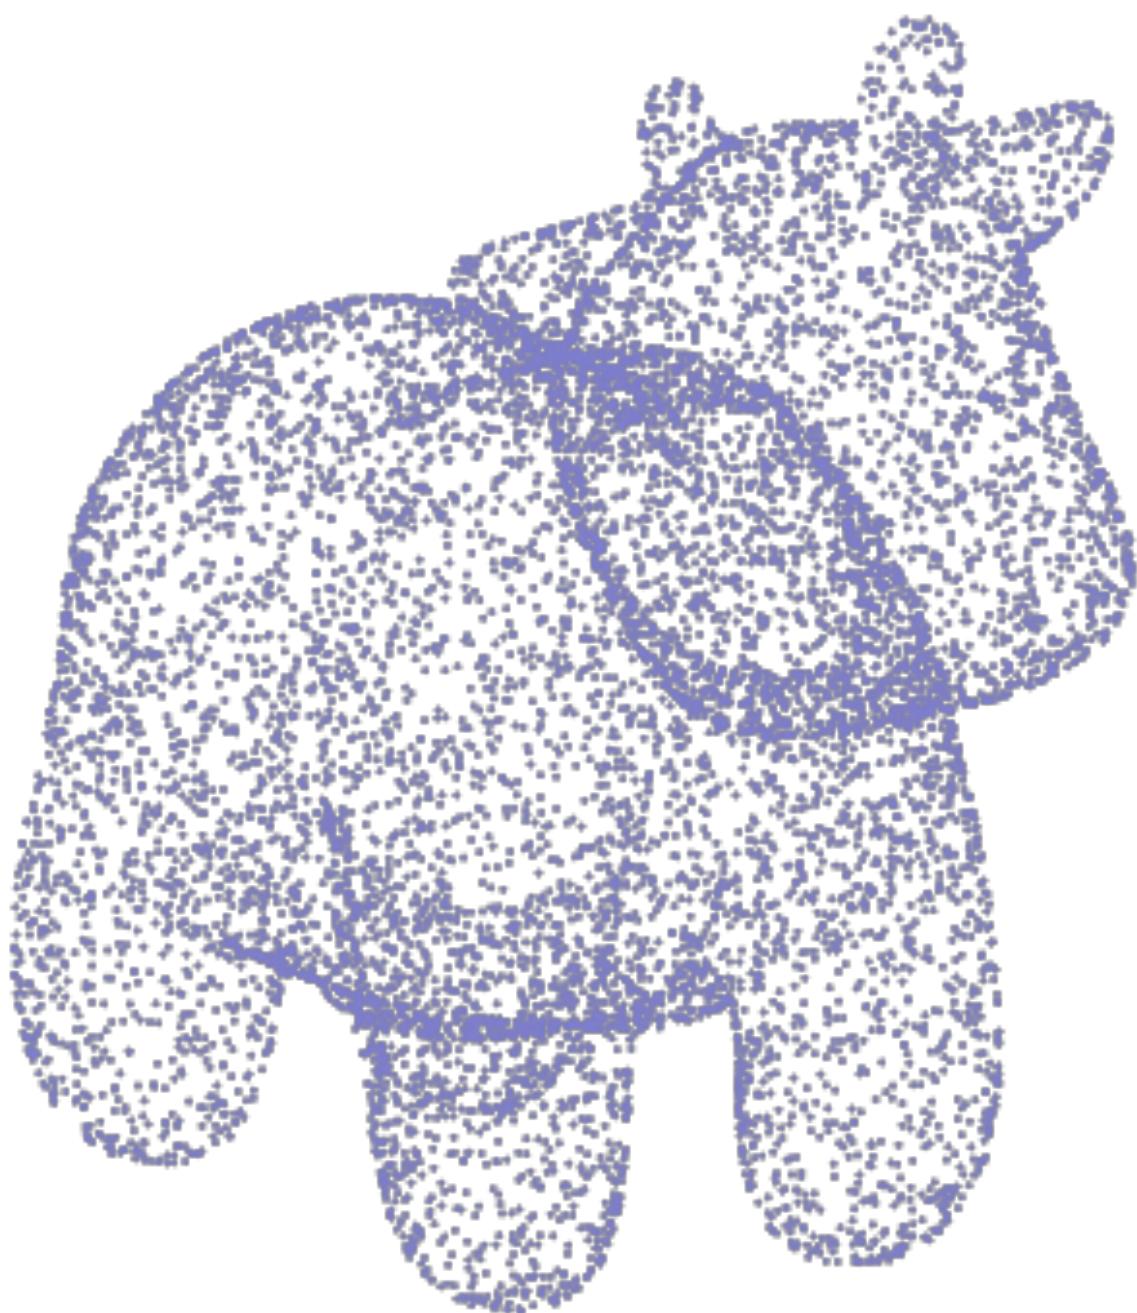


x1,y1,z1
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.

$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

**Unordered set of points**

# Point Clouds



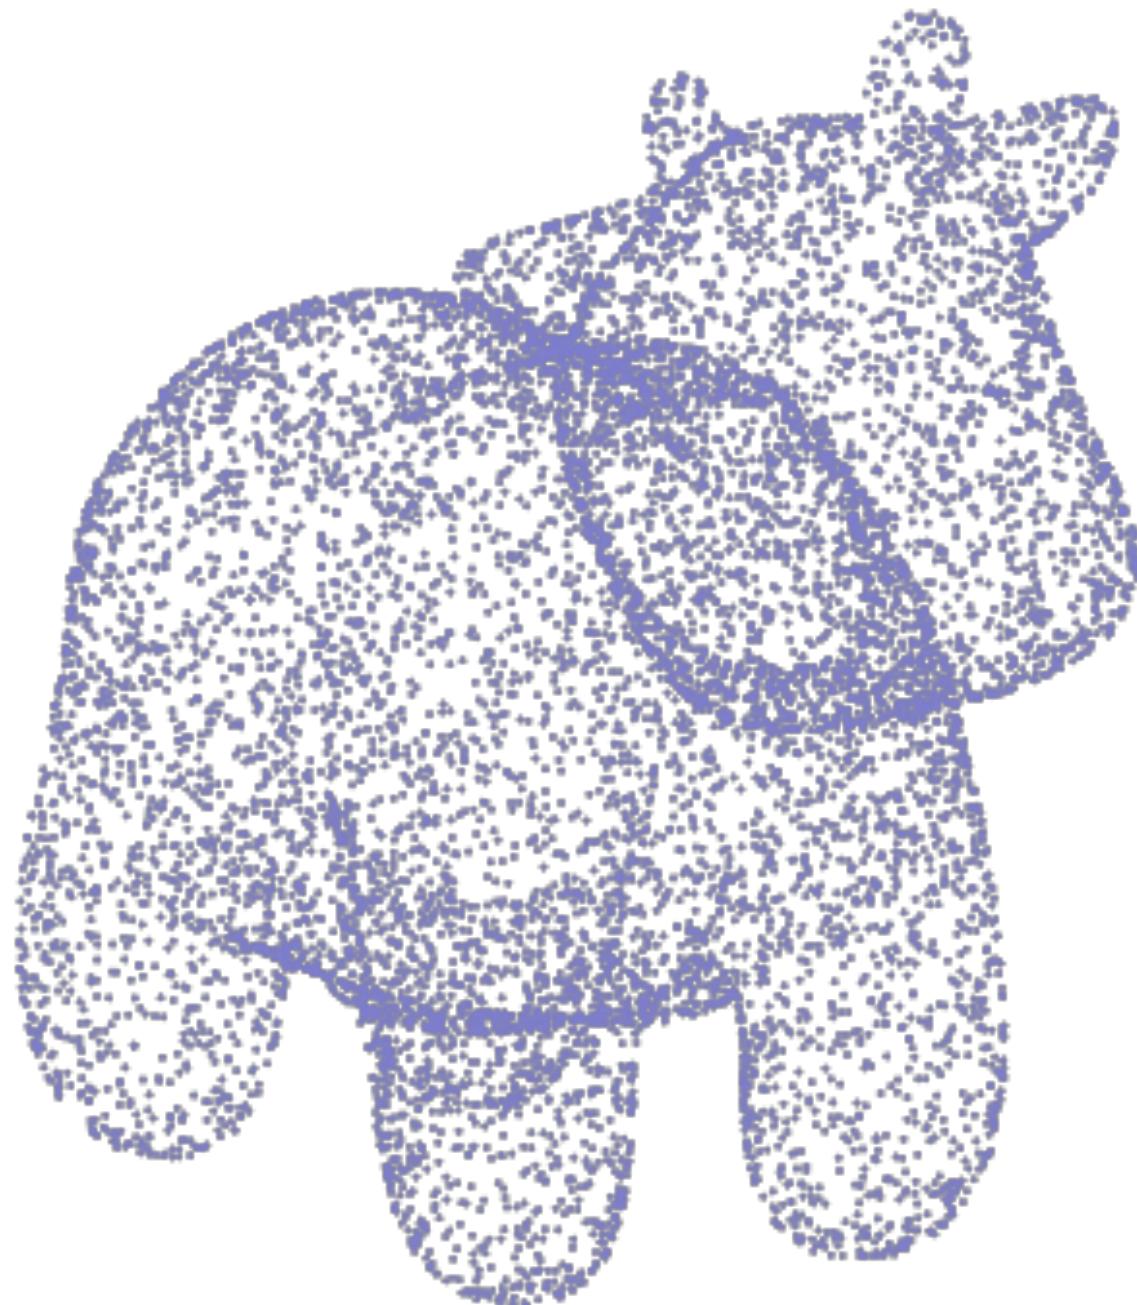
x1,y1,z1
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.

Often represented as a NX3 array, but ordering does **not** matter (unlike images)

$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

**Unordered set of points**

# Point Clouds



$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

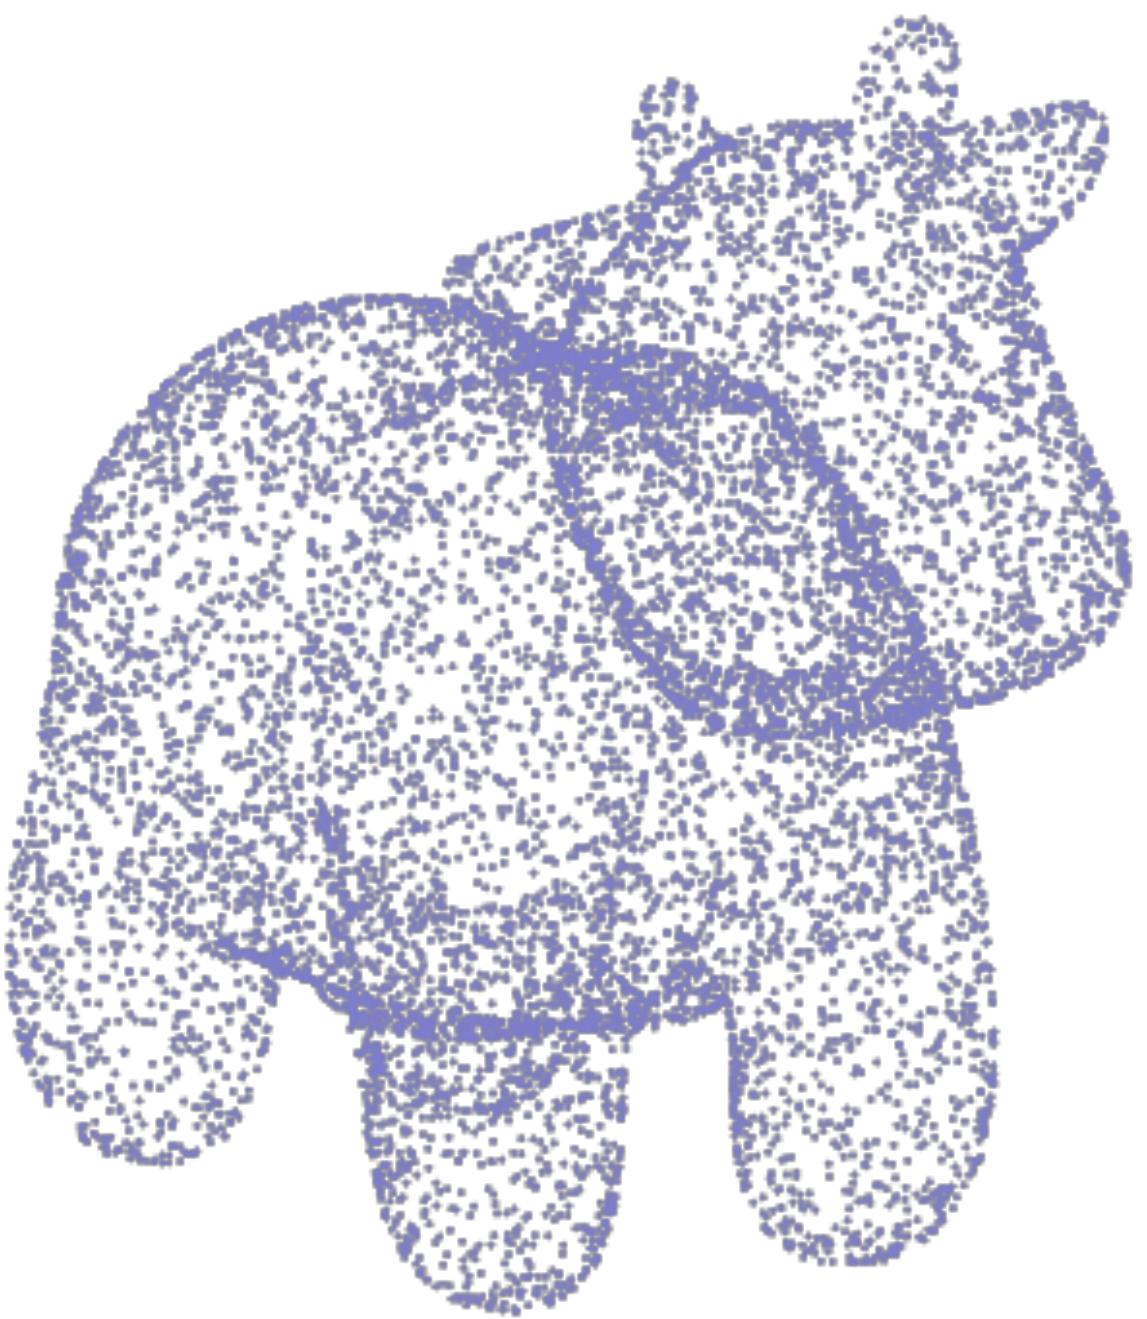
# Unordered set of points

x1,y1,z1

Often represented as a  $N \times 3$  array, but ordering does not matter (unlike images)

Need processing/generation methods that are permutation invariant (e.g. fully connected layer will not work)

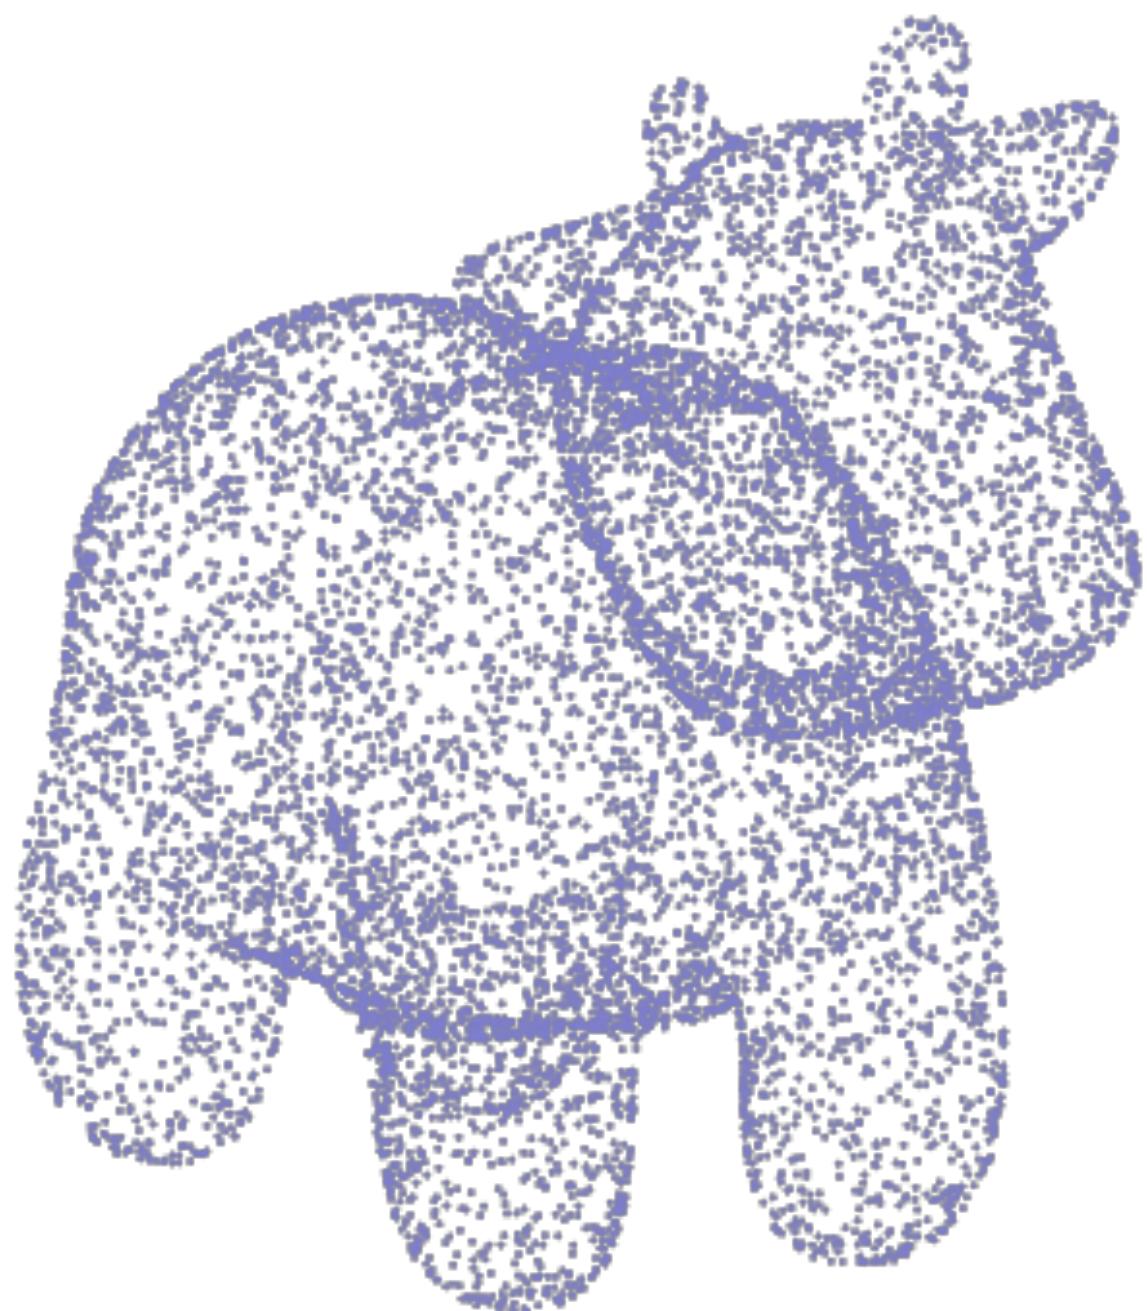
# Point Clouds



$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

**Unordered set of points**

# Point Clouds

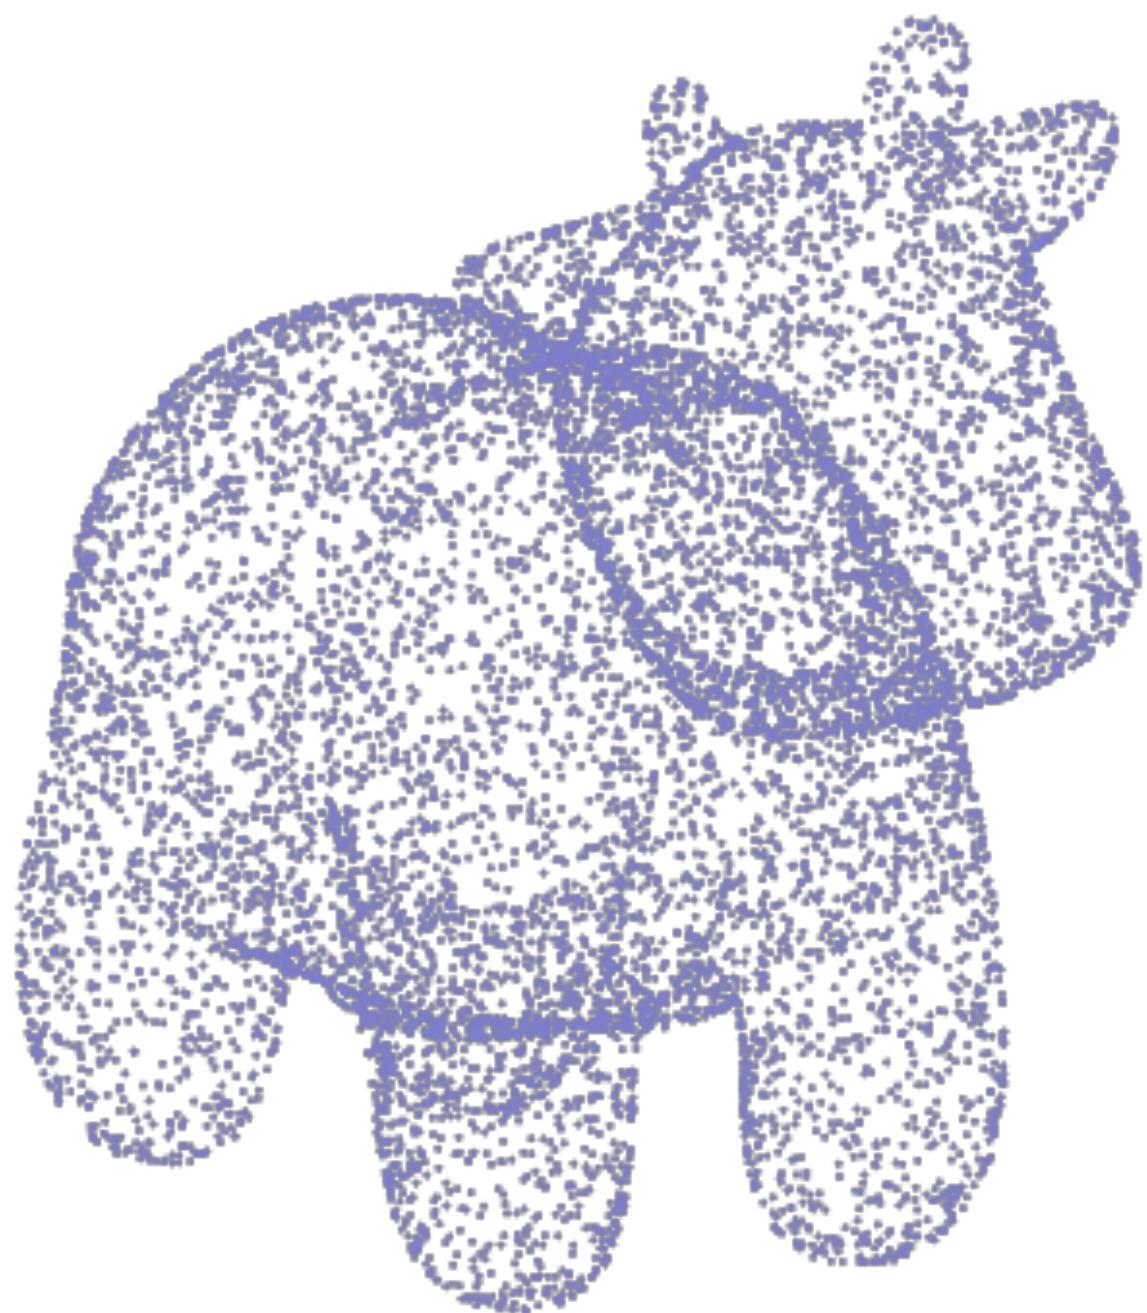


No explicit ‘connectivity’  
information

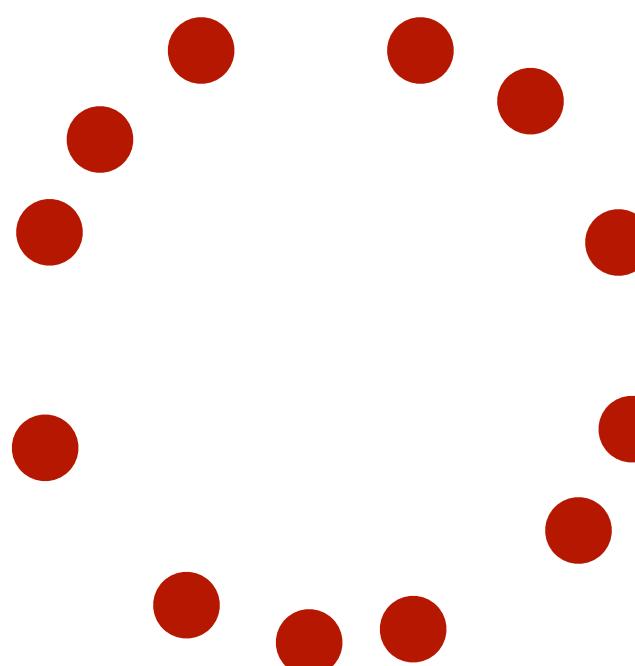
$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

**Unordered set of points**

# Point Clouds



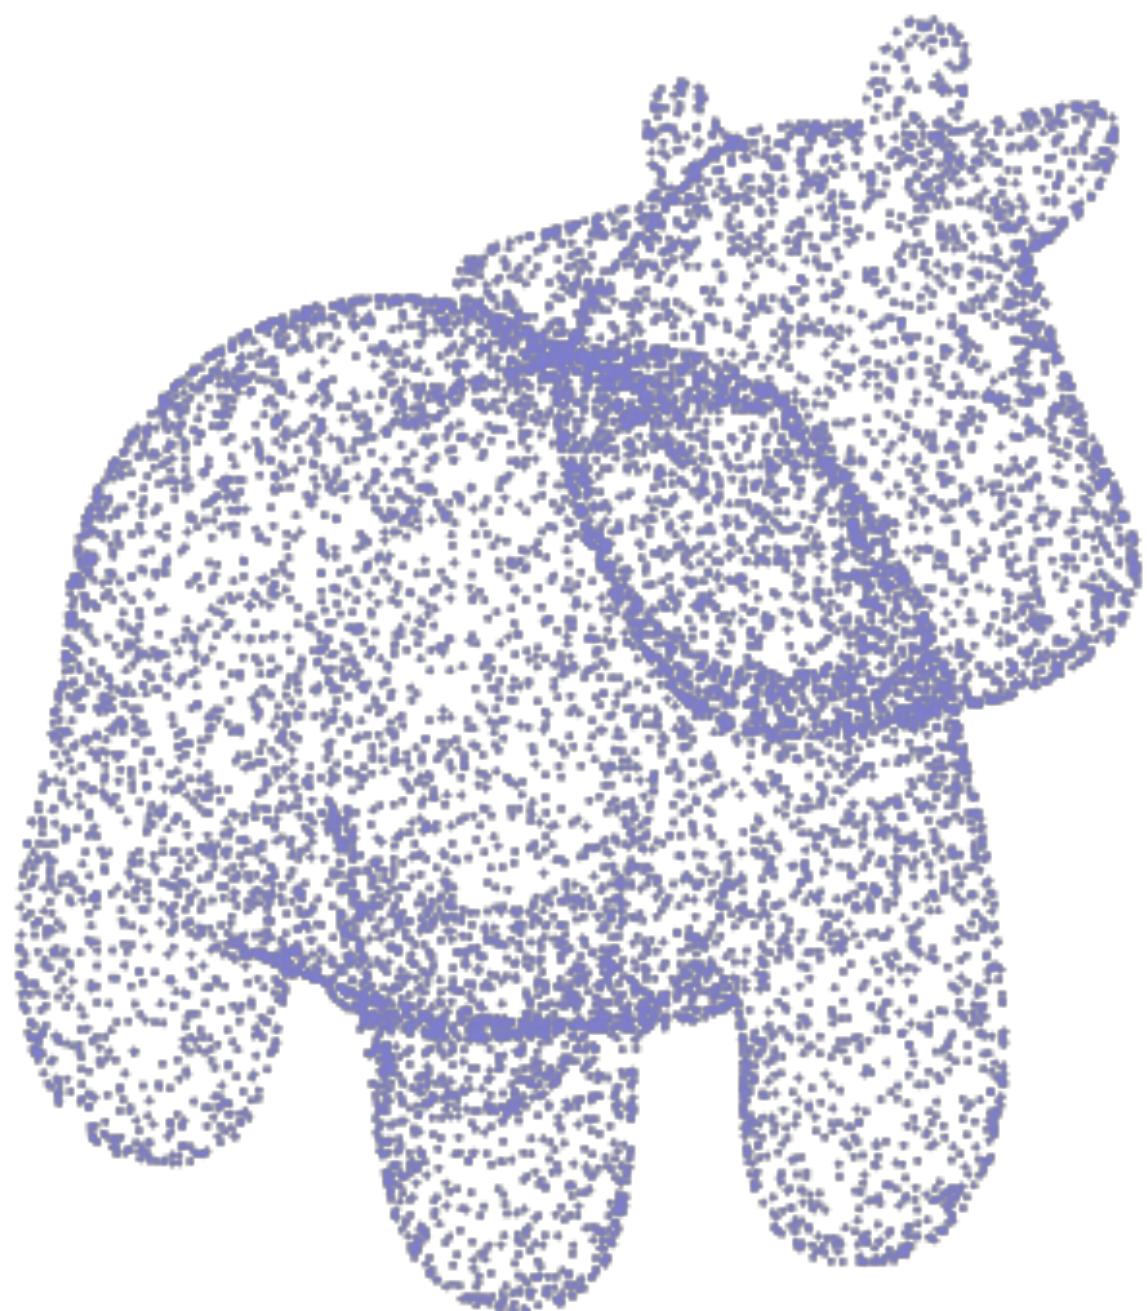
No explicit ‘connectivity’ information



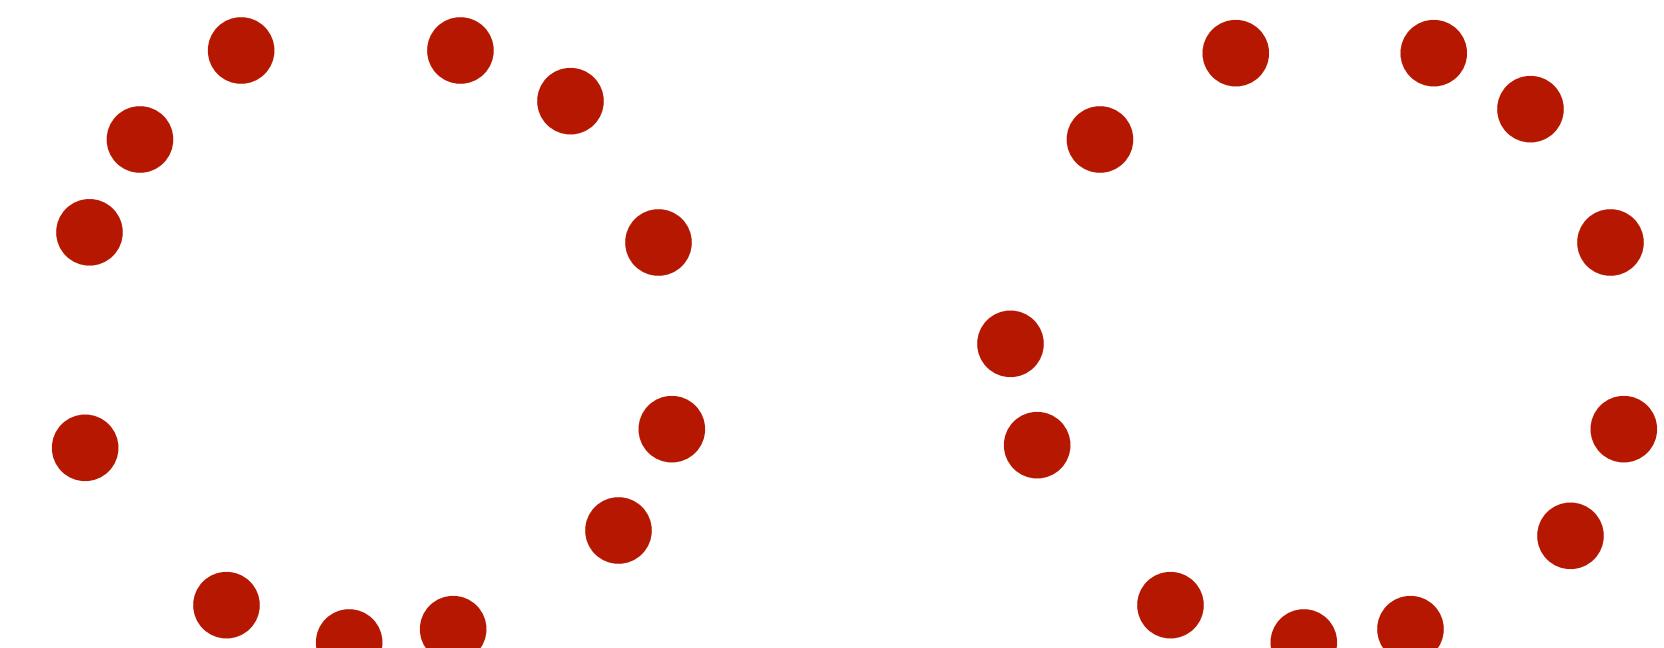
$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

**Unordered set of points**

# Point Clouds



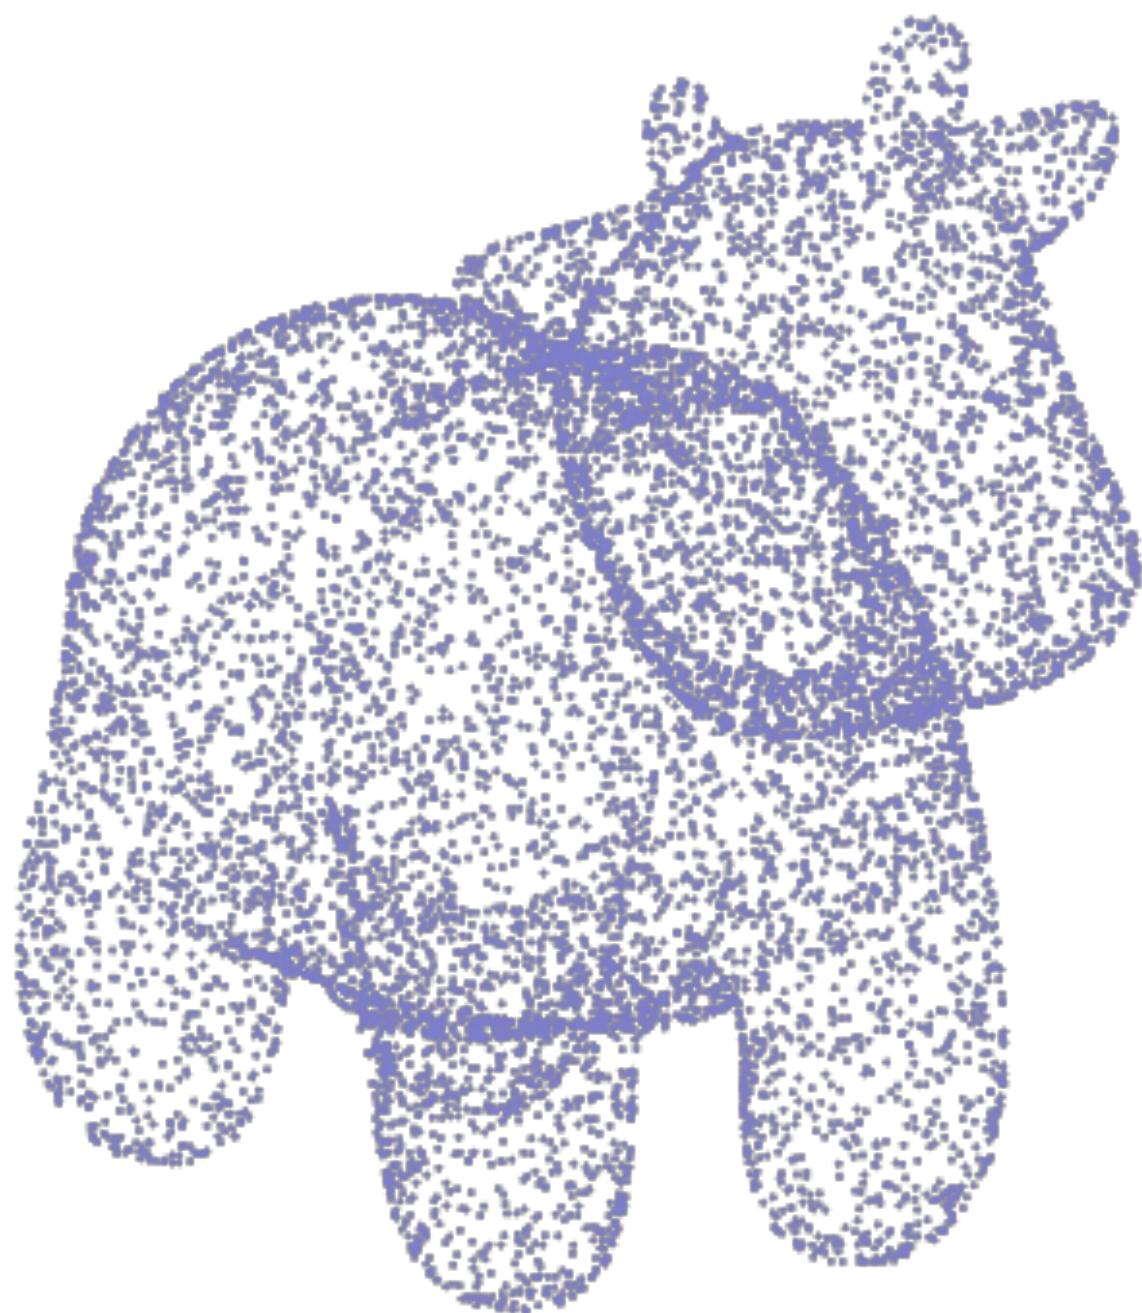
No explicit ‘connectivity’ information



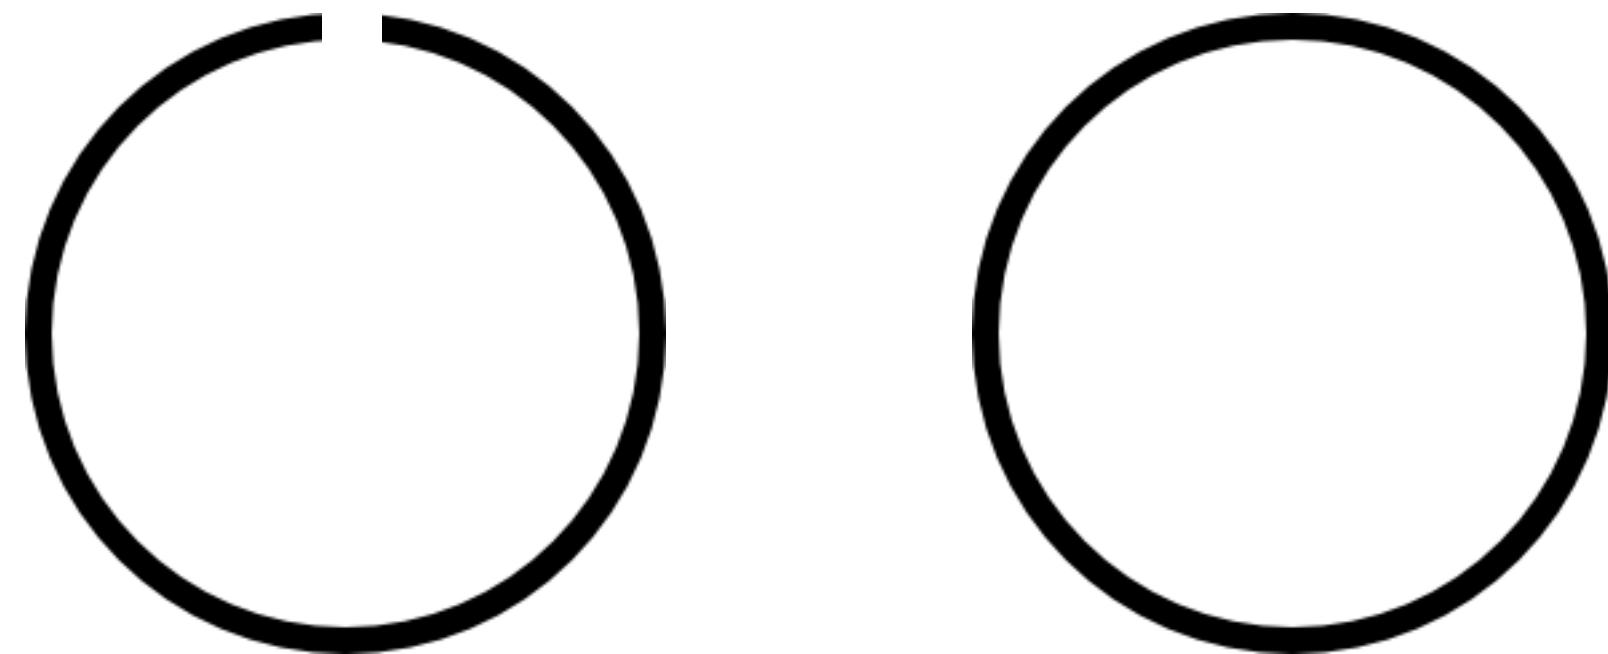
$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

**Unordered set of points**

# Point Clouds



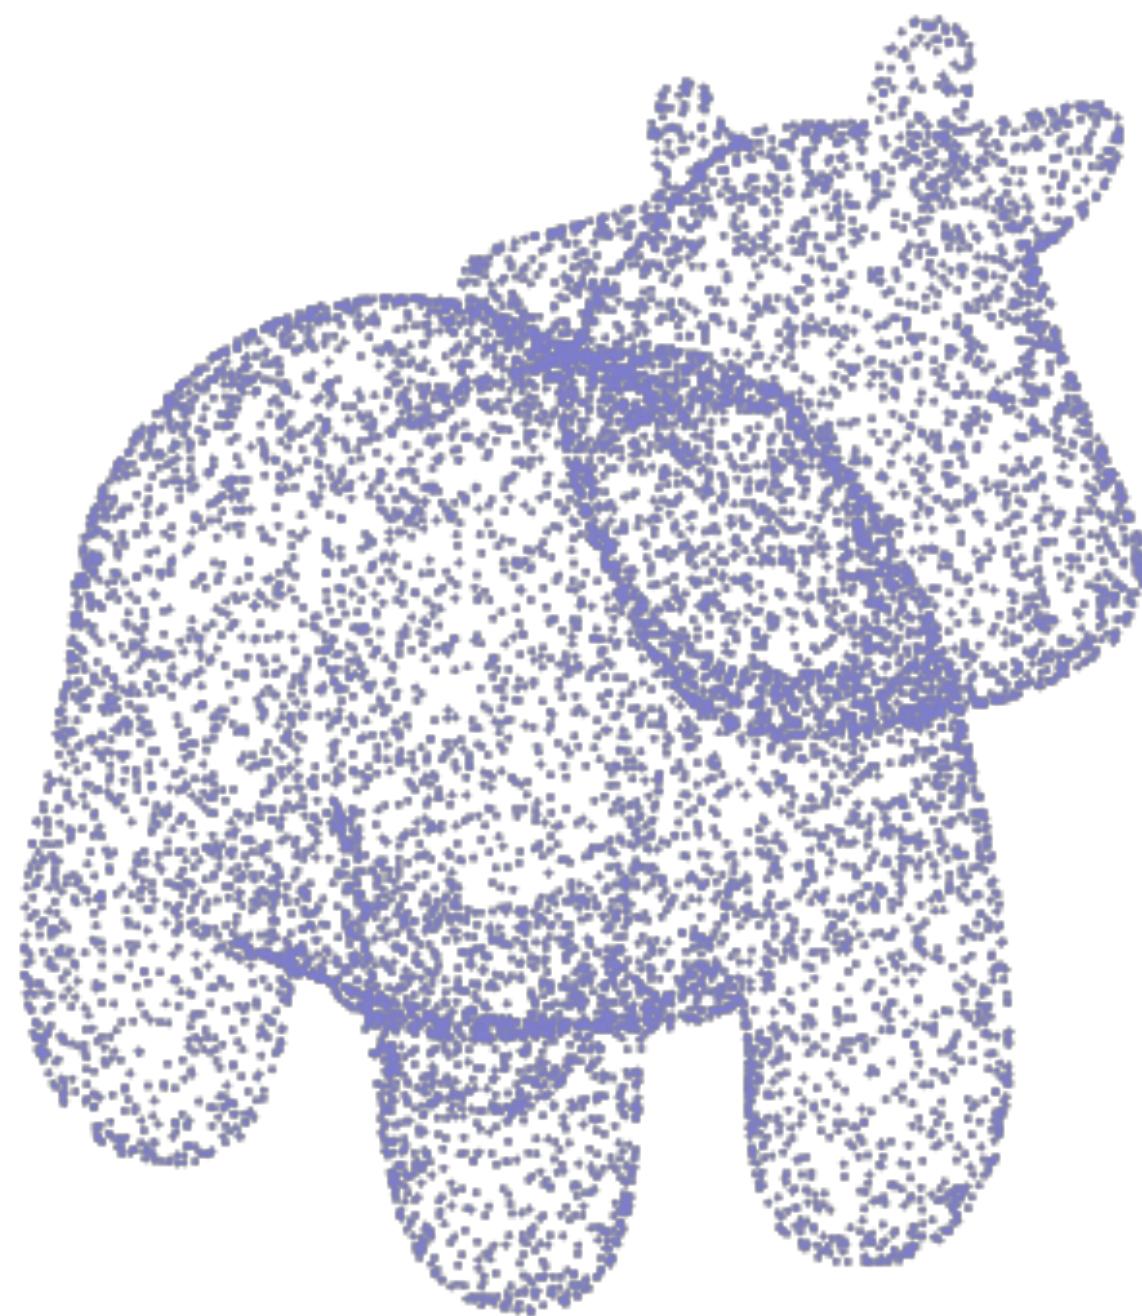
No explicit ‘connectivity’  
information



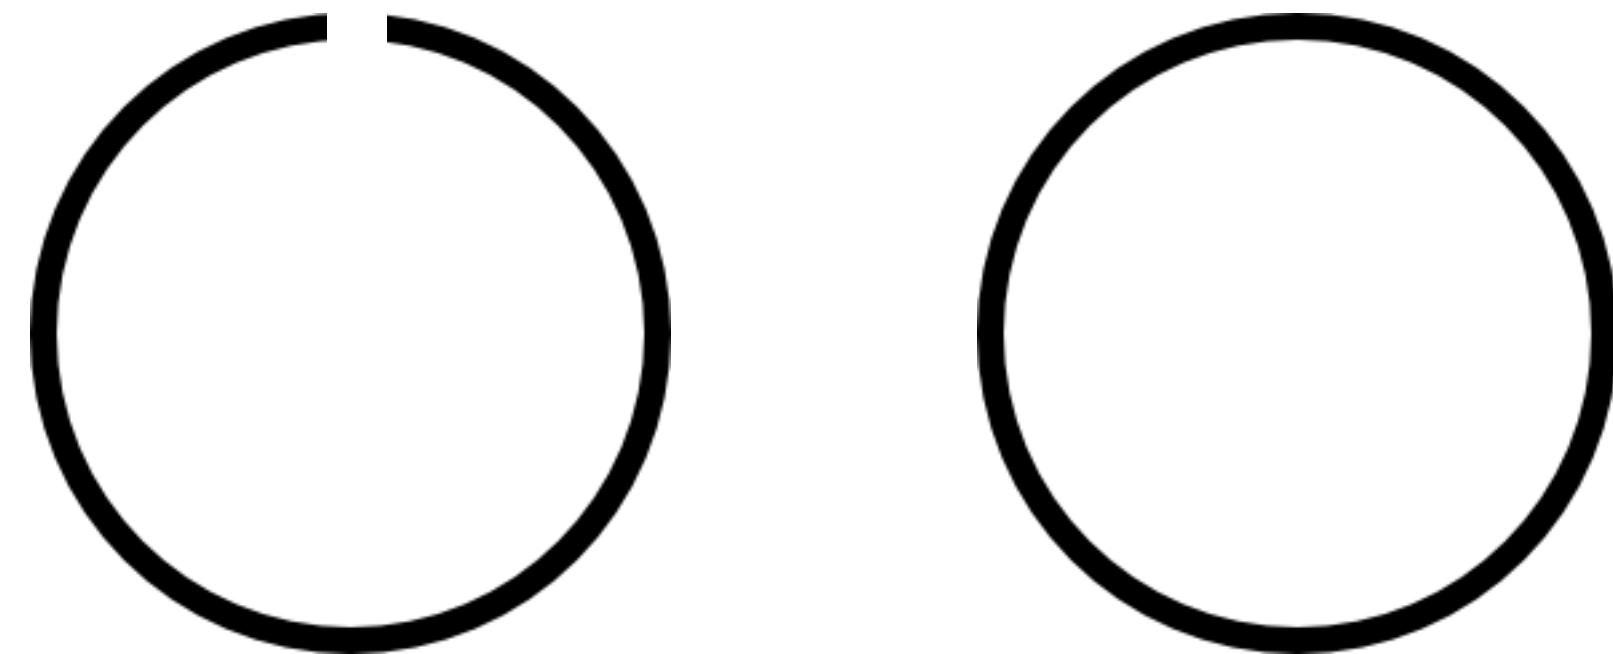
$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

**Unordered set of points**

# Point Clouds



No explicit ‘connectivity’ information



$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

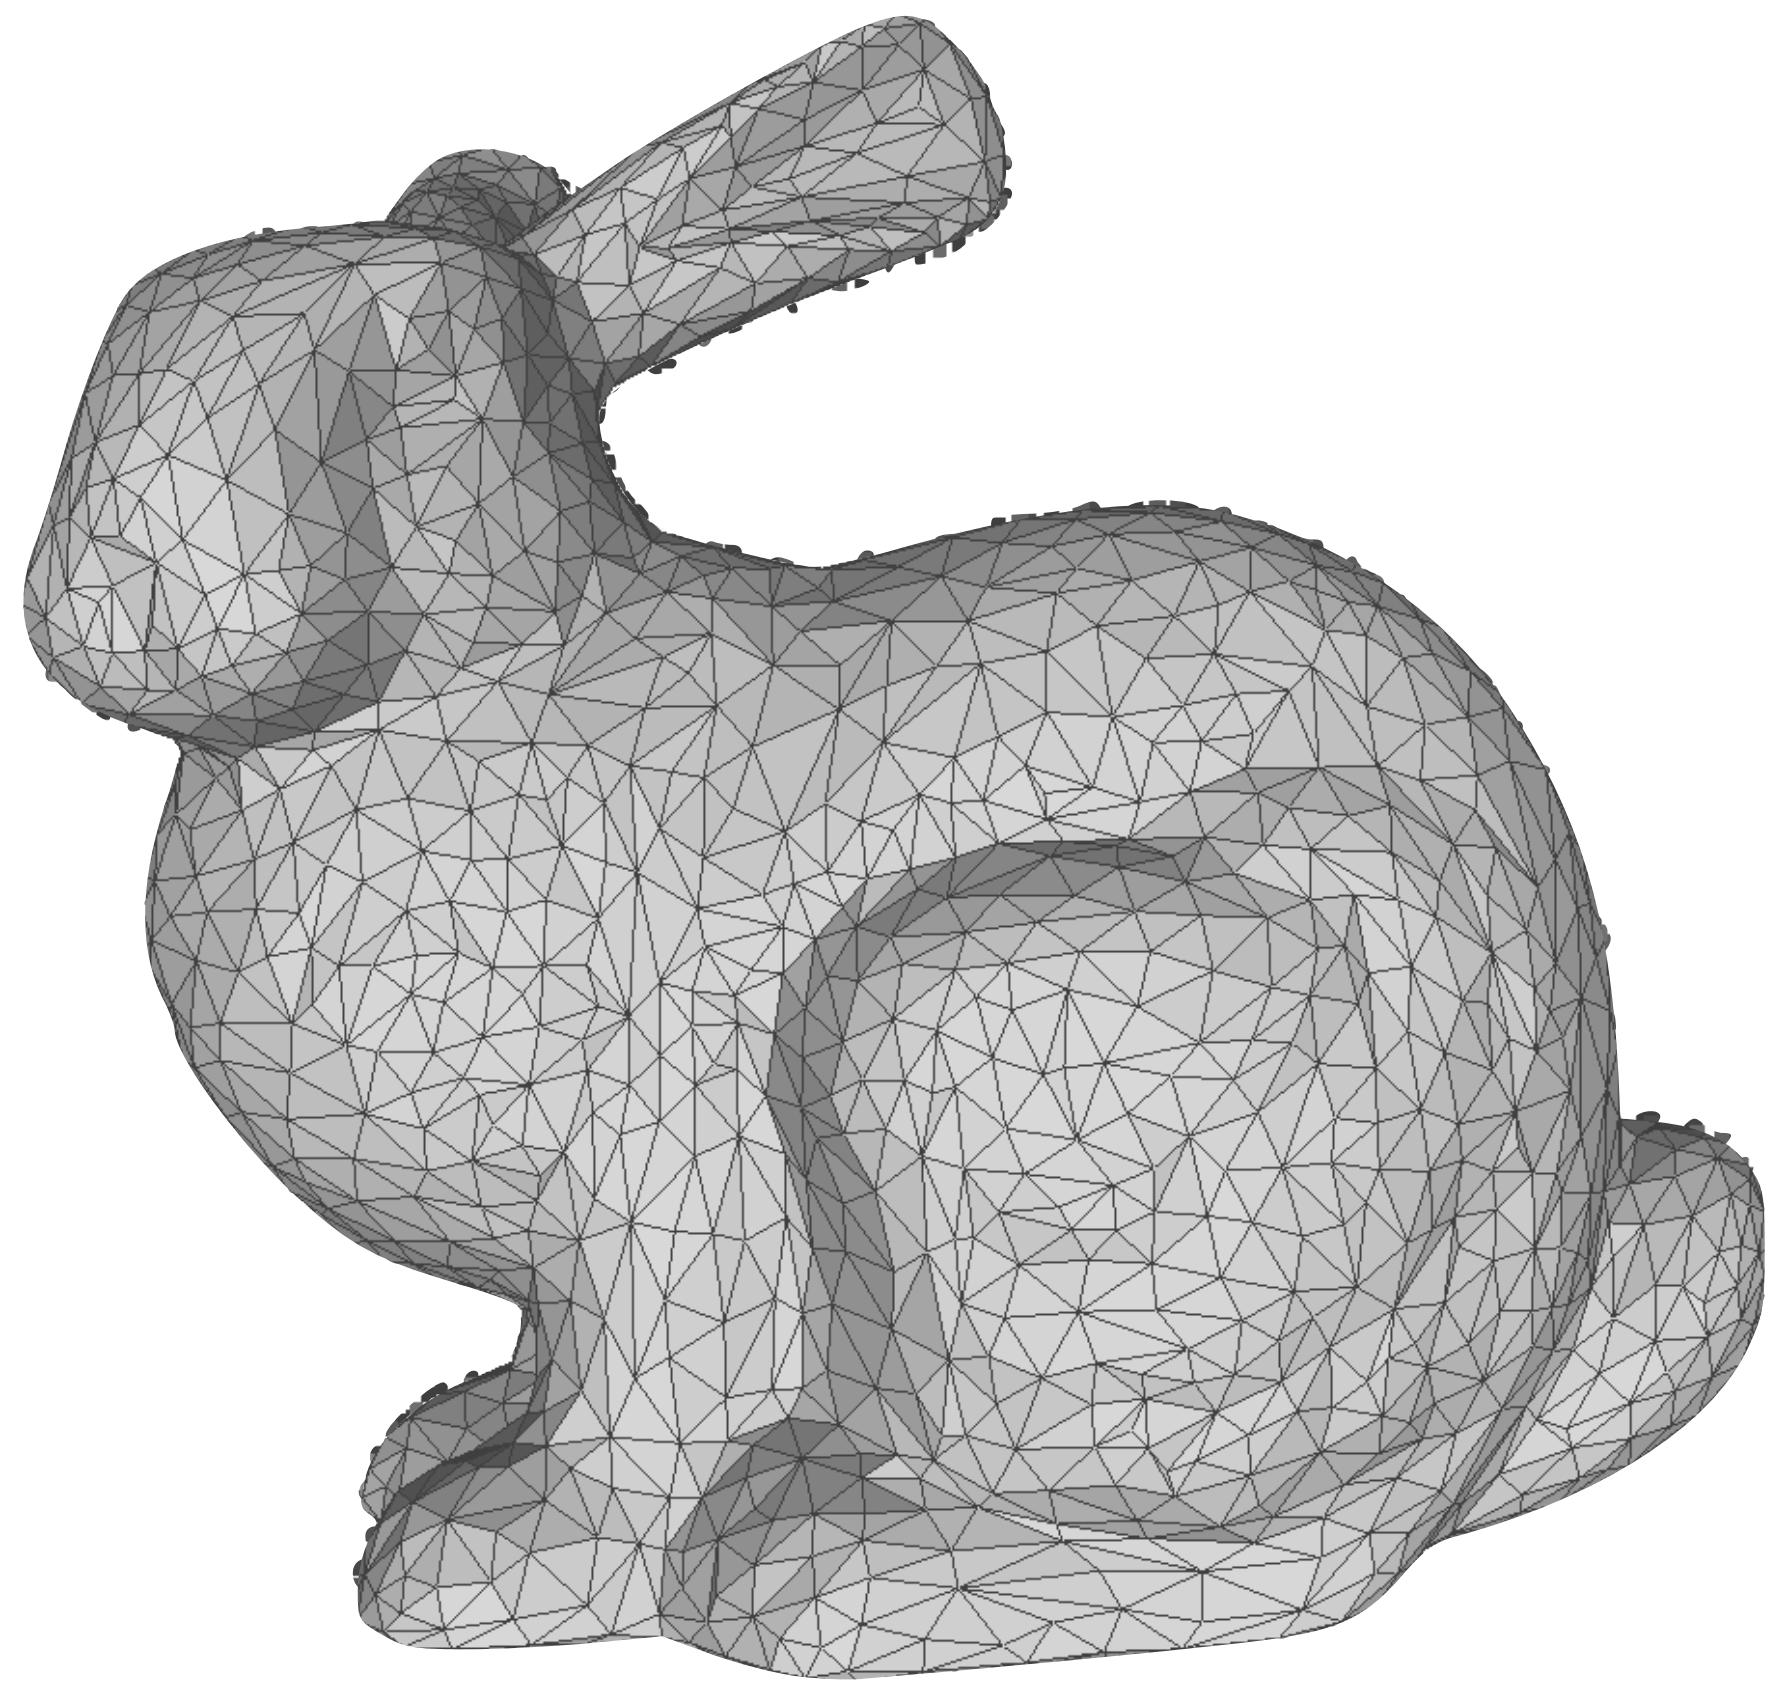
**Unordered set of points**

(yes, we can sample more - but  
more efficient to add edges)

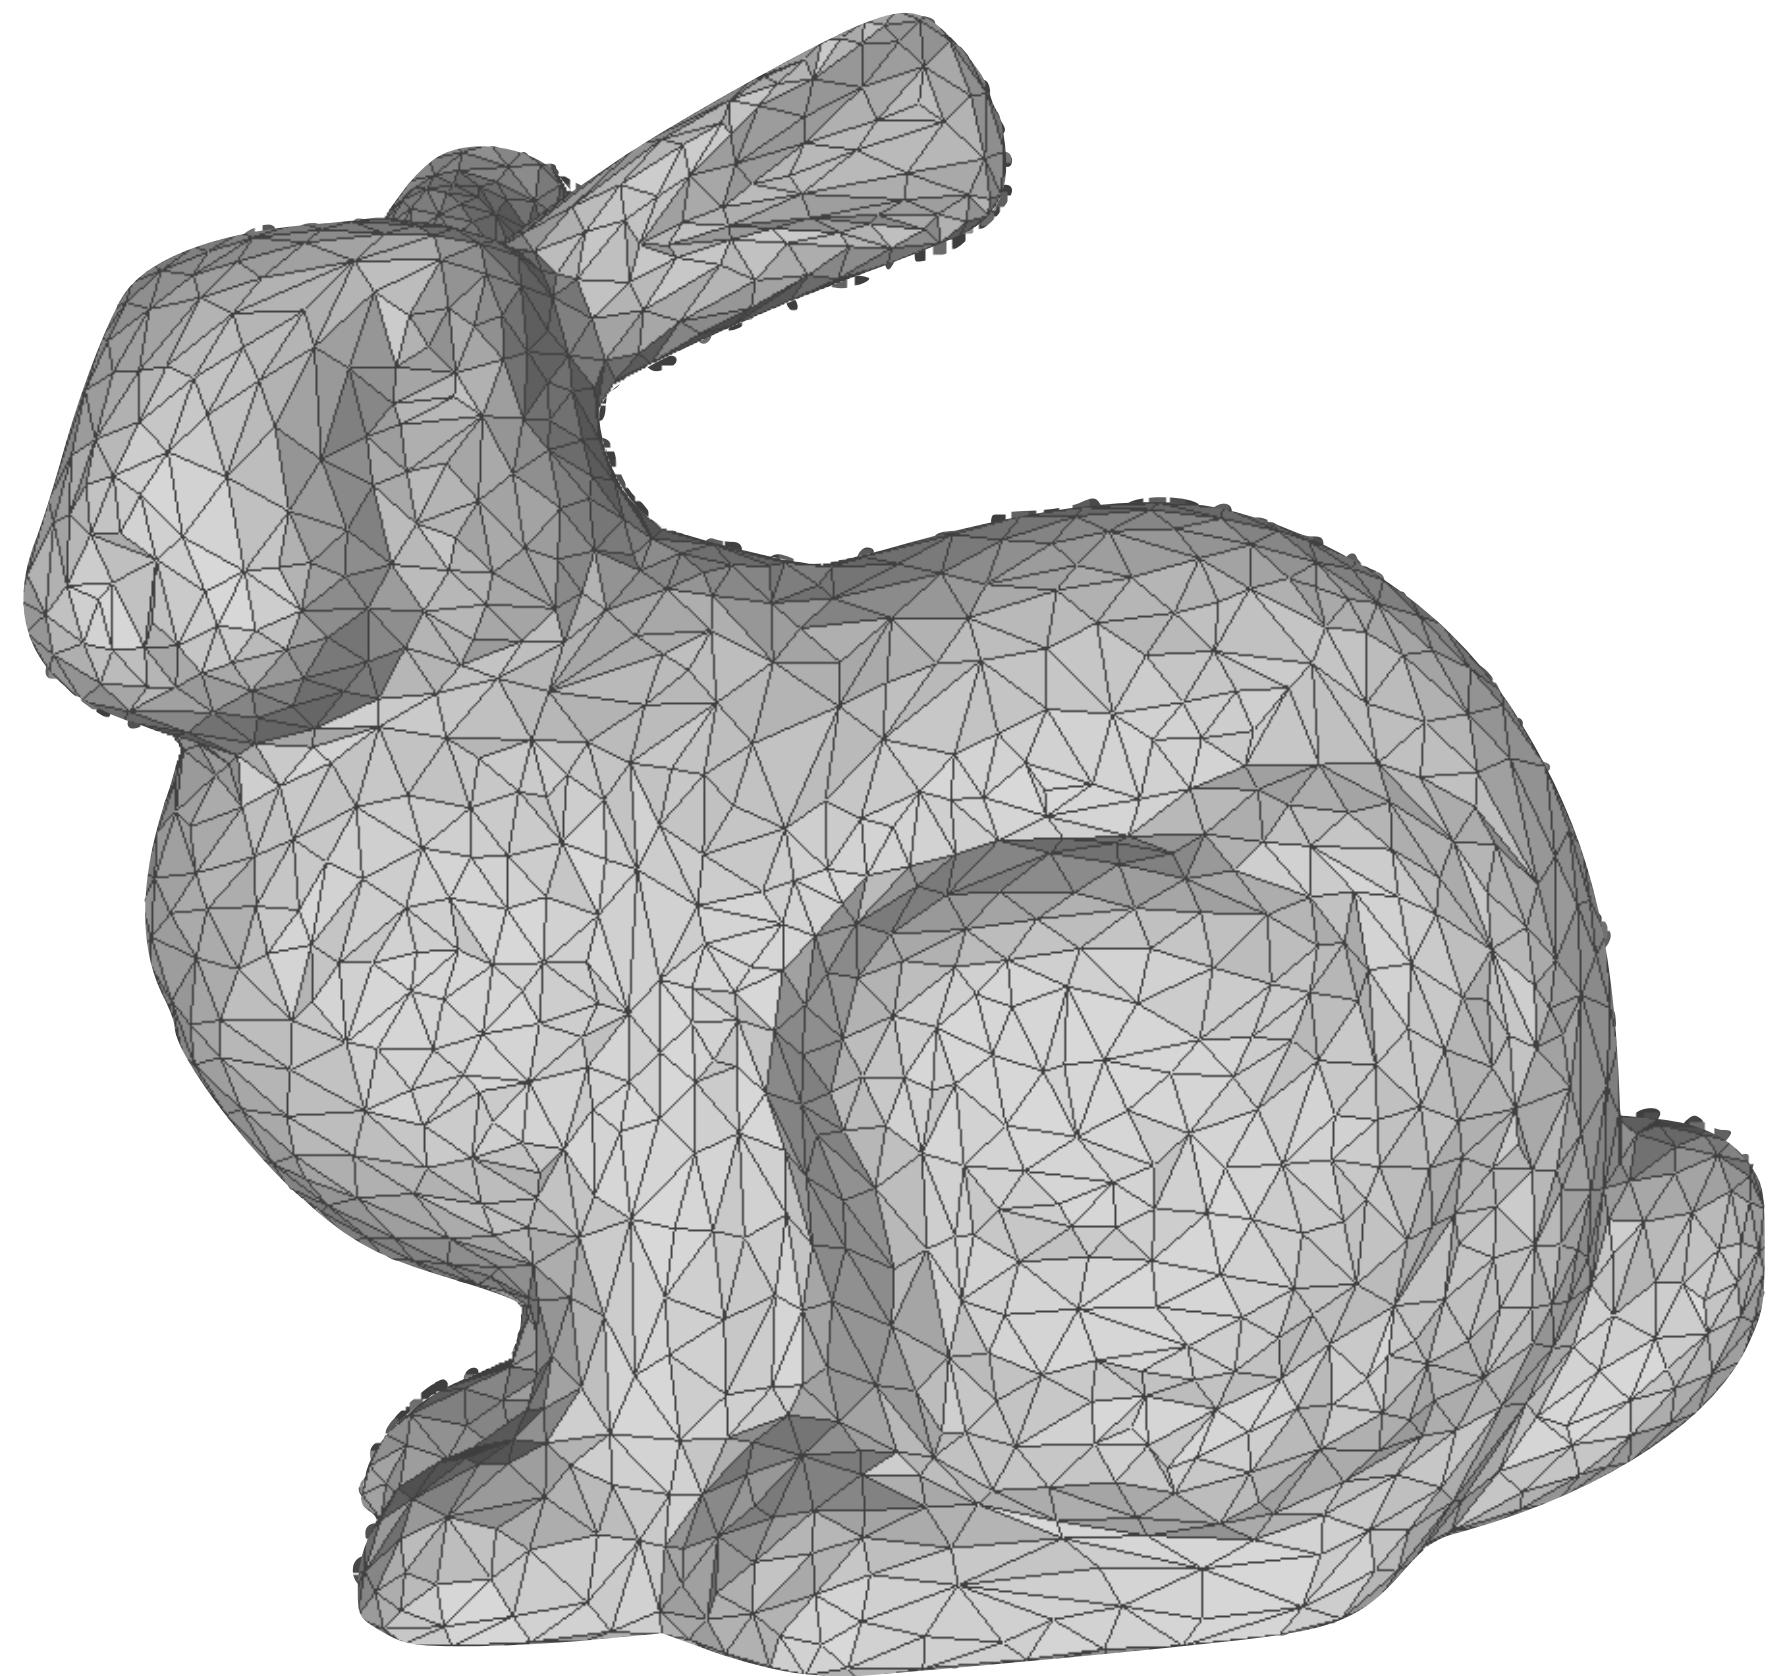
# Meshes



# Meshes

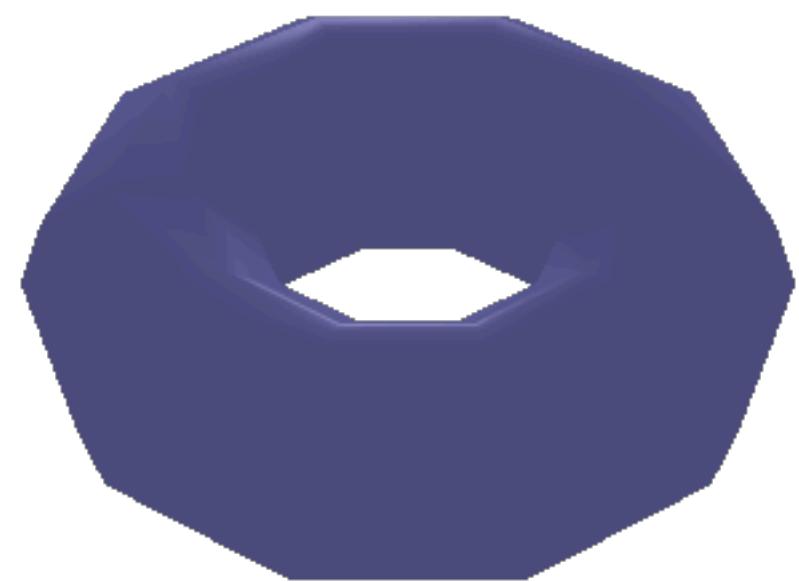
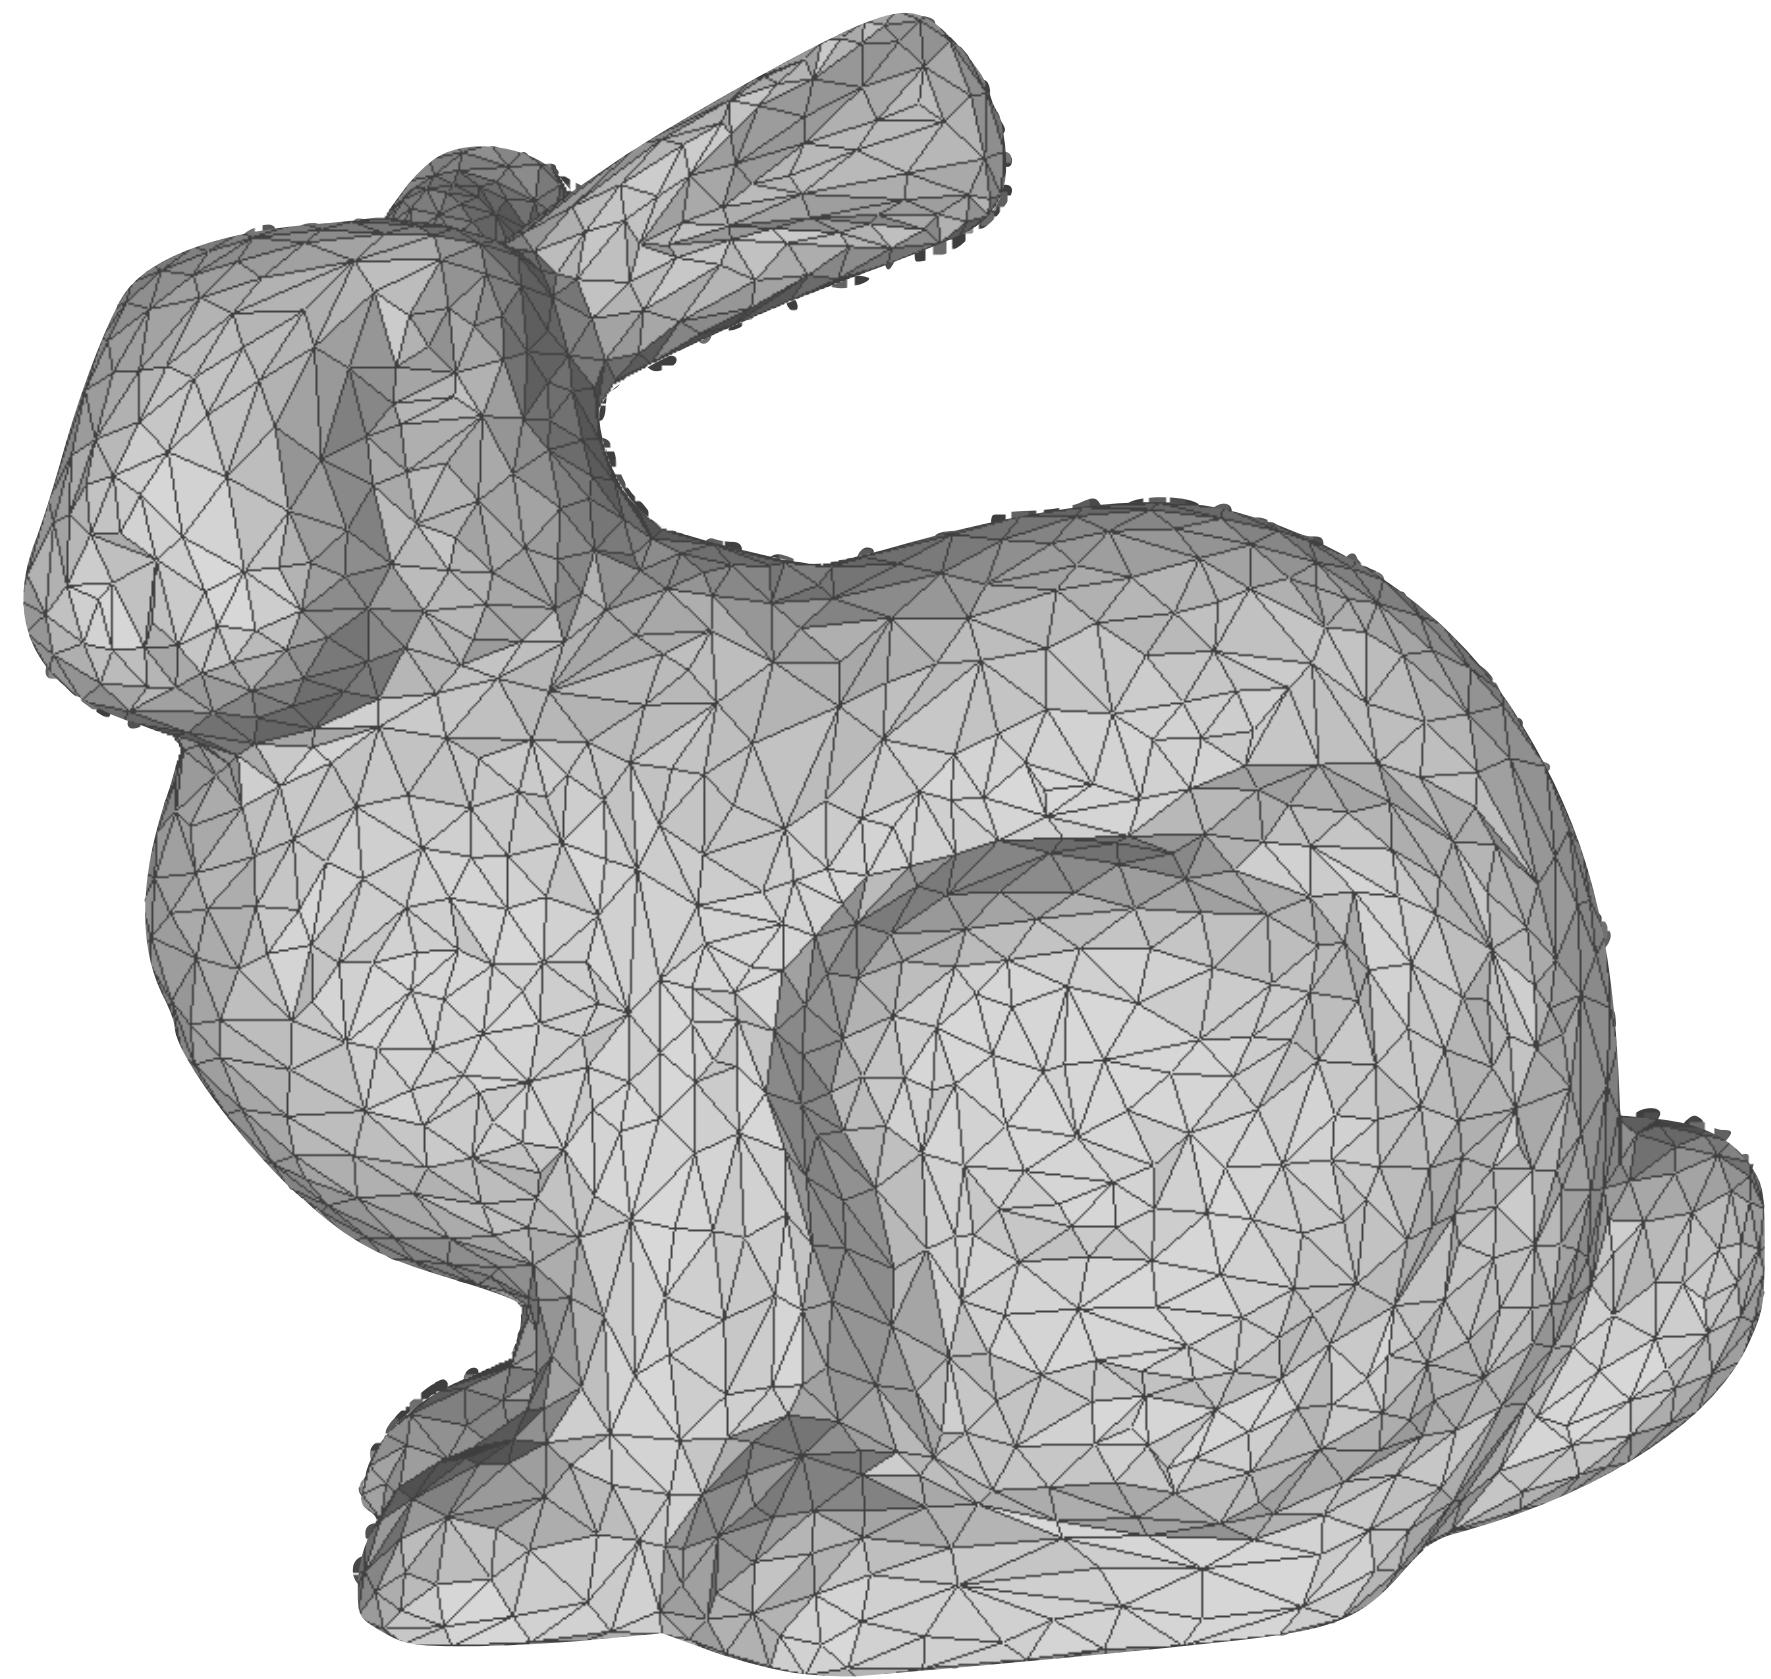


# Meshes



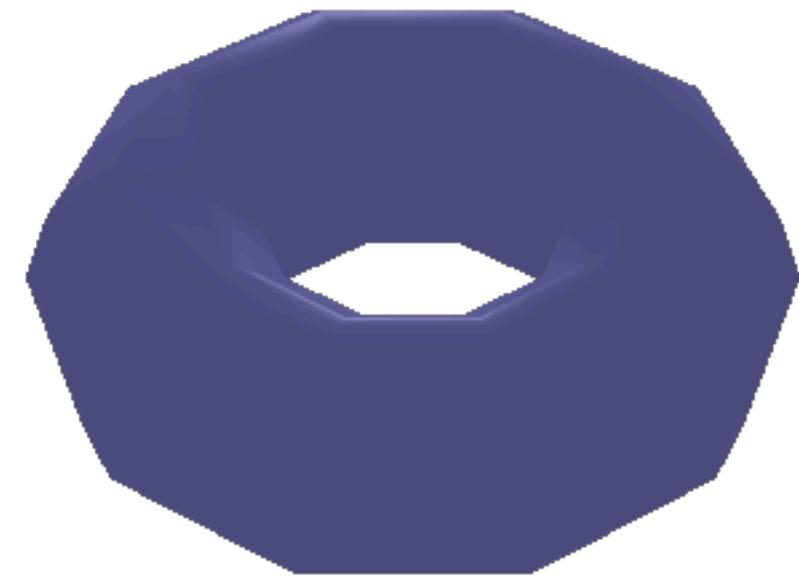
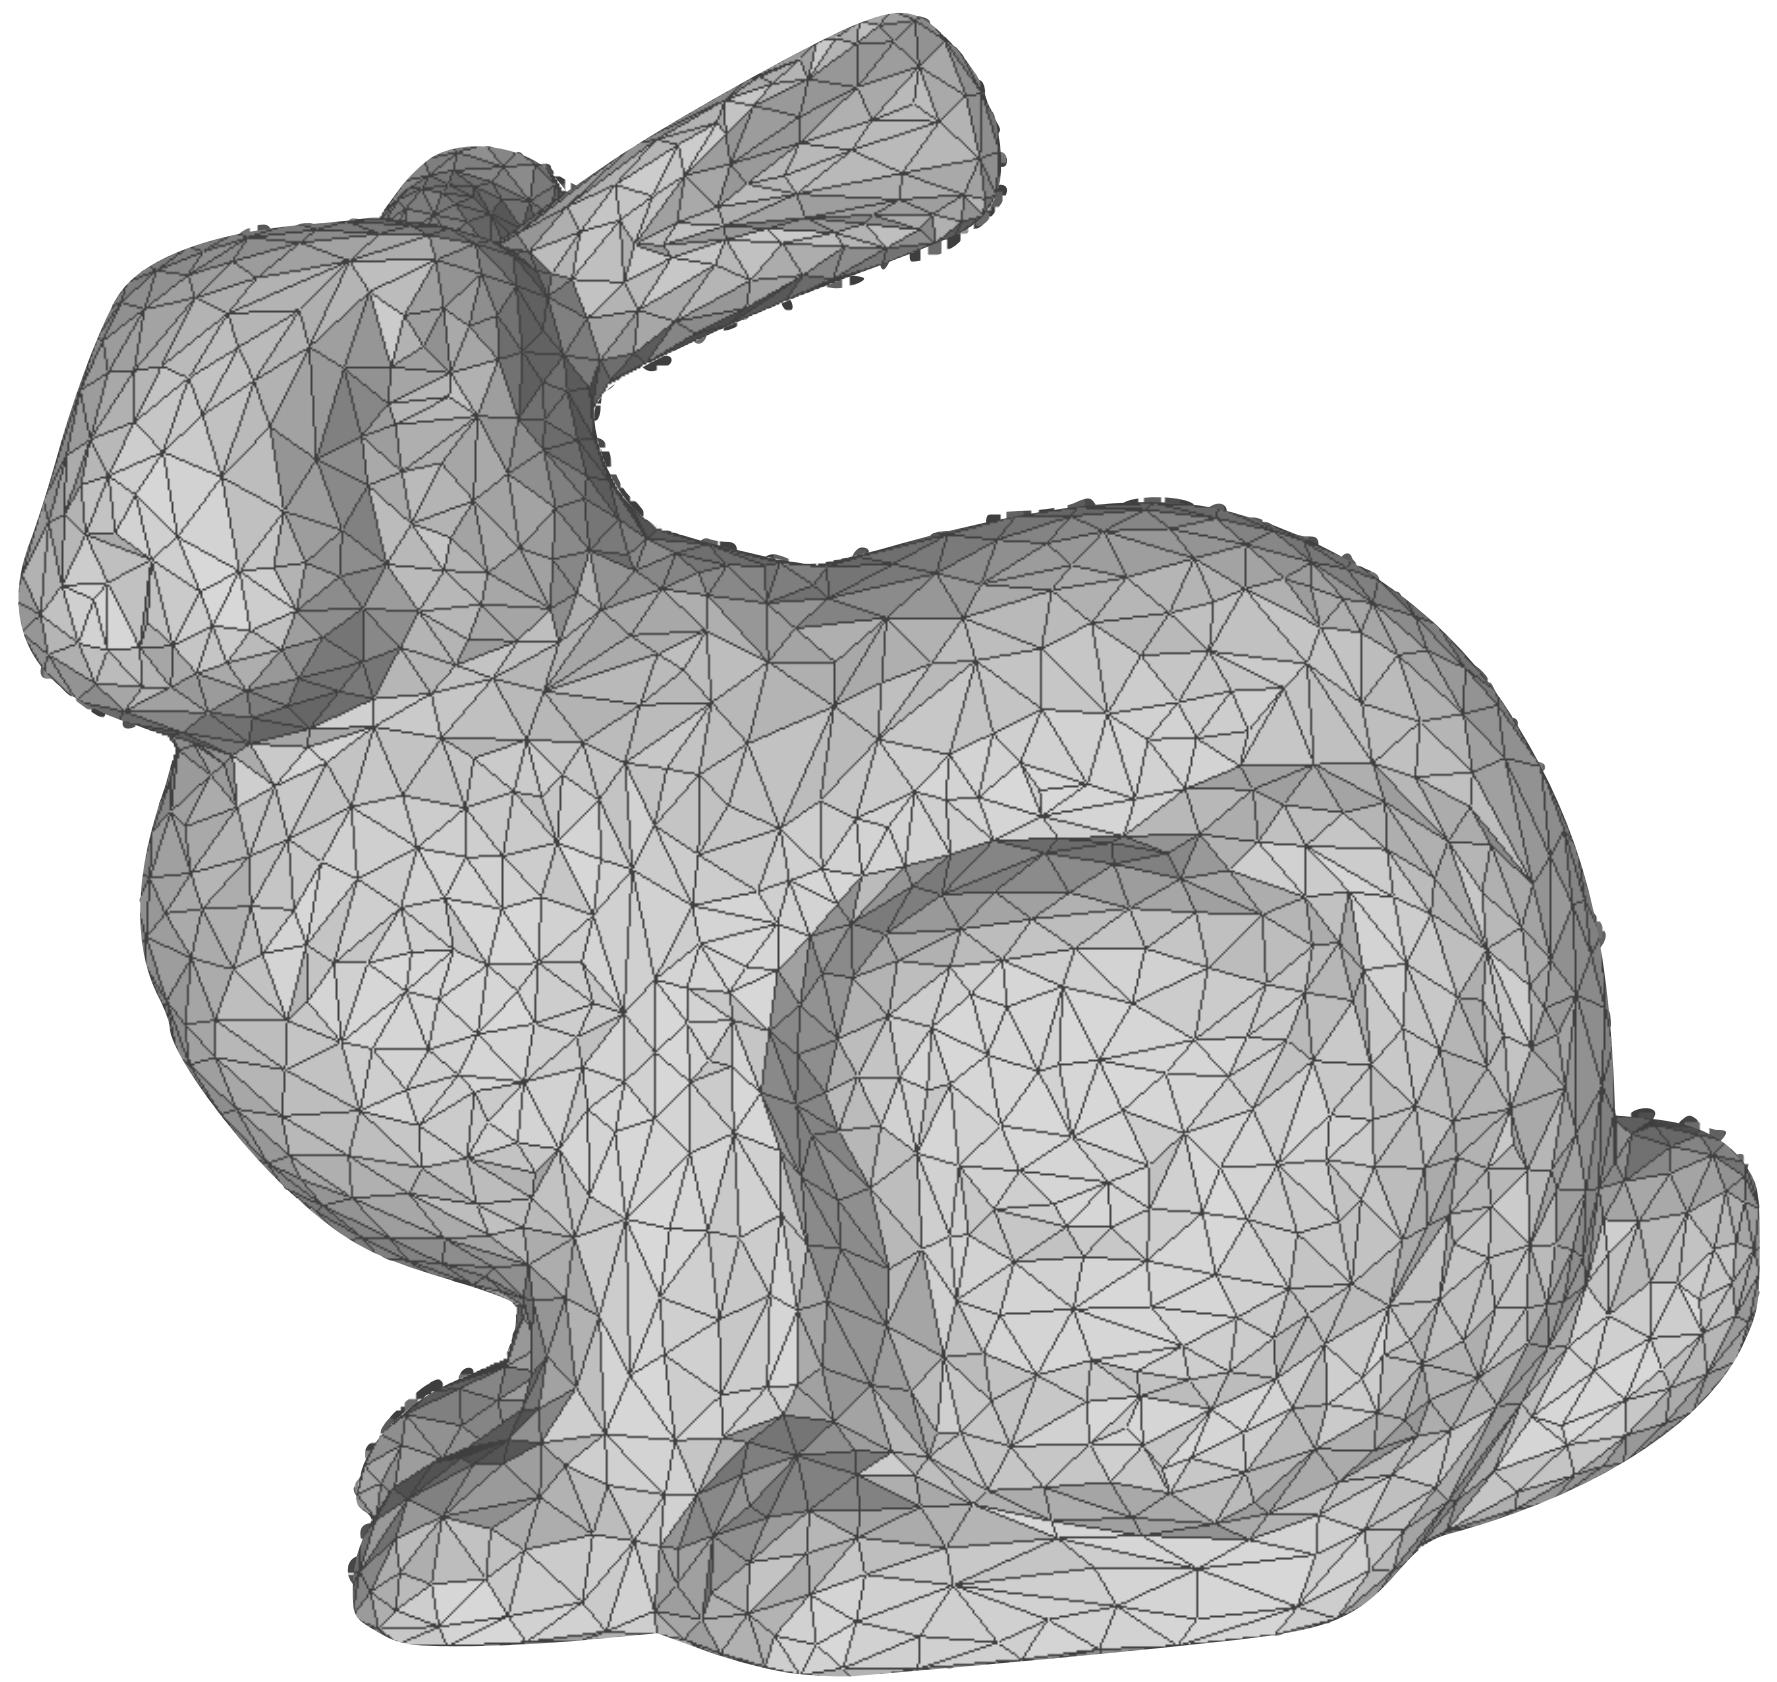
Vertices + Faces

# Meshes



Vertices + Faces

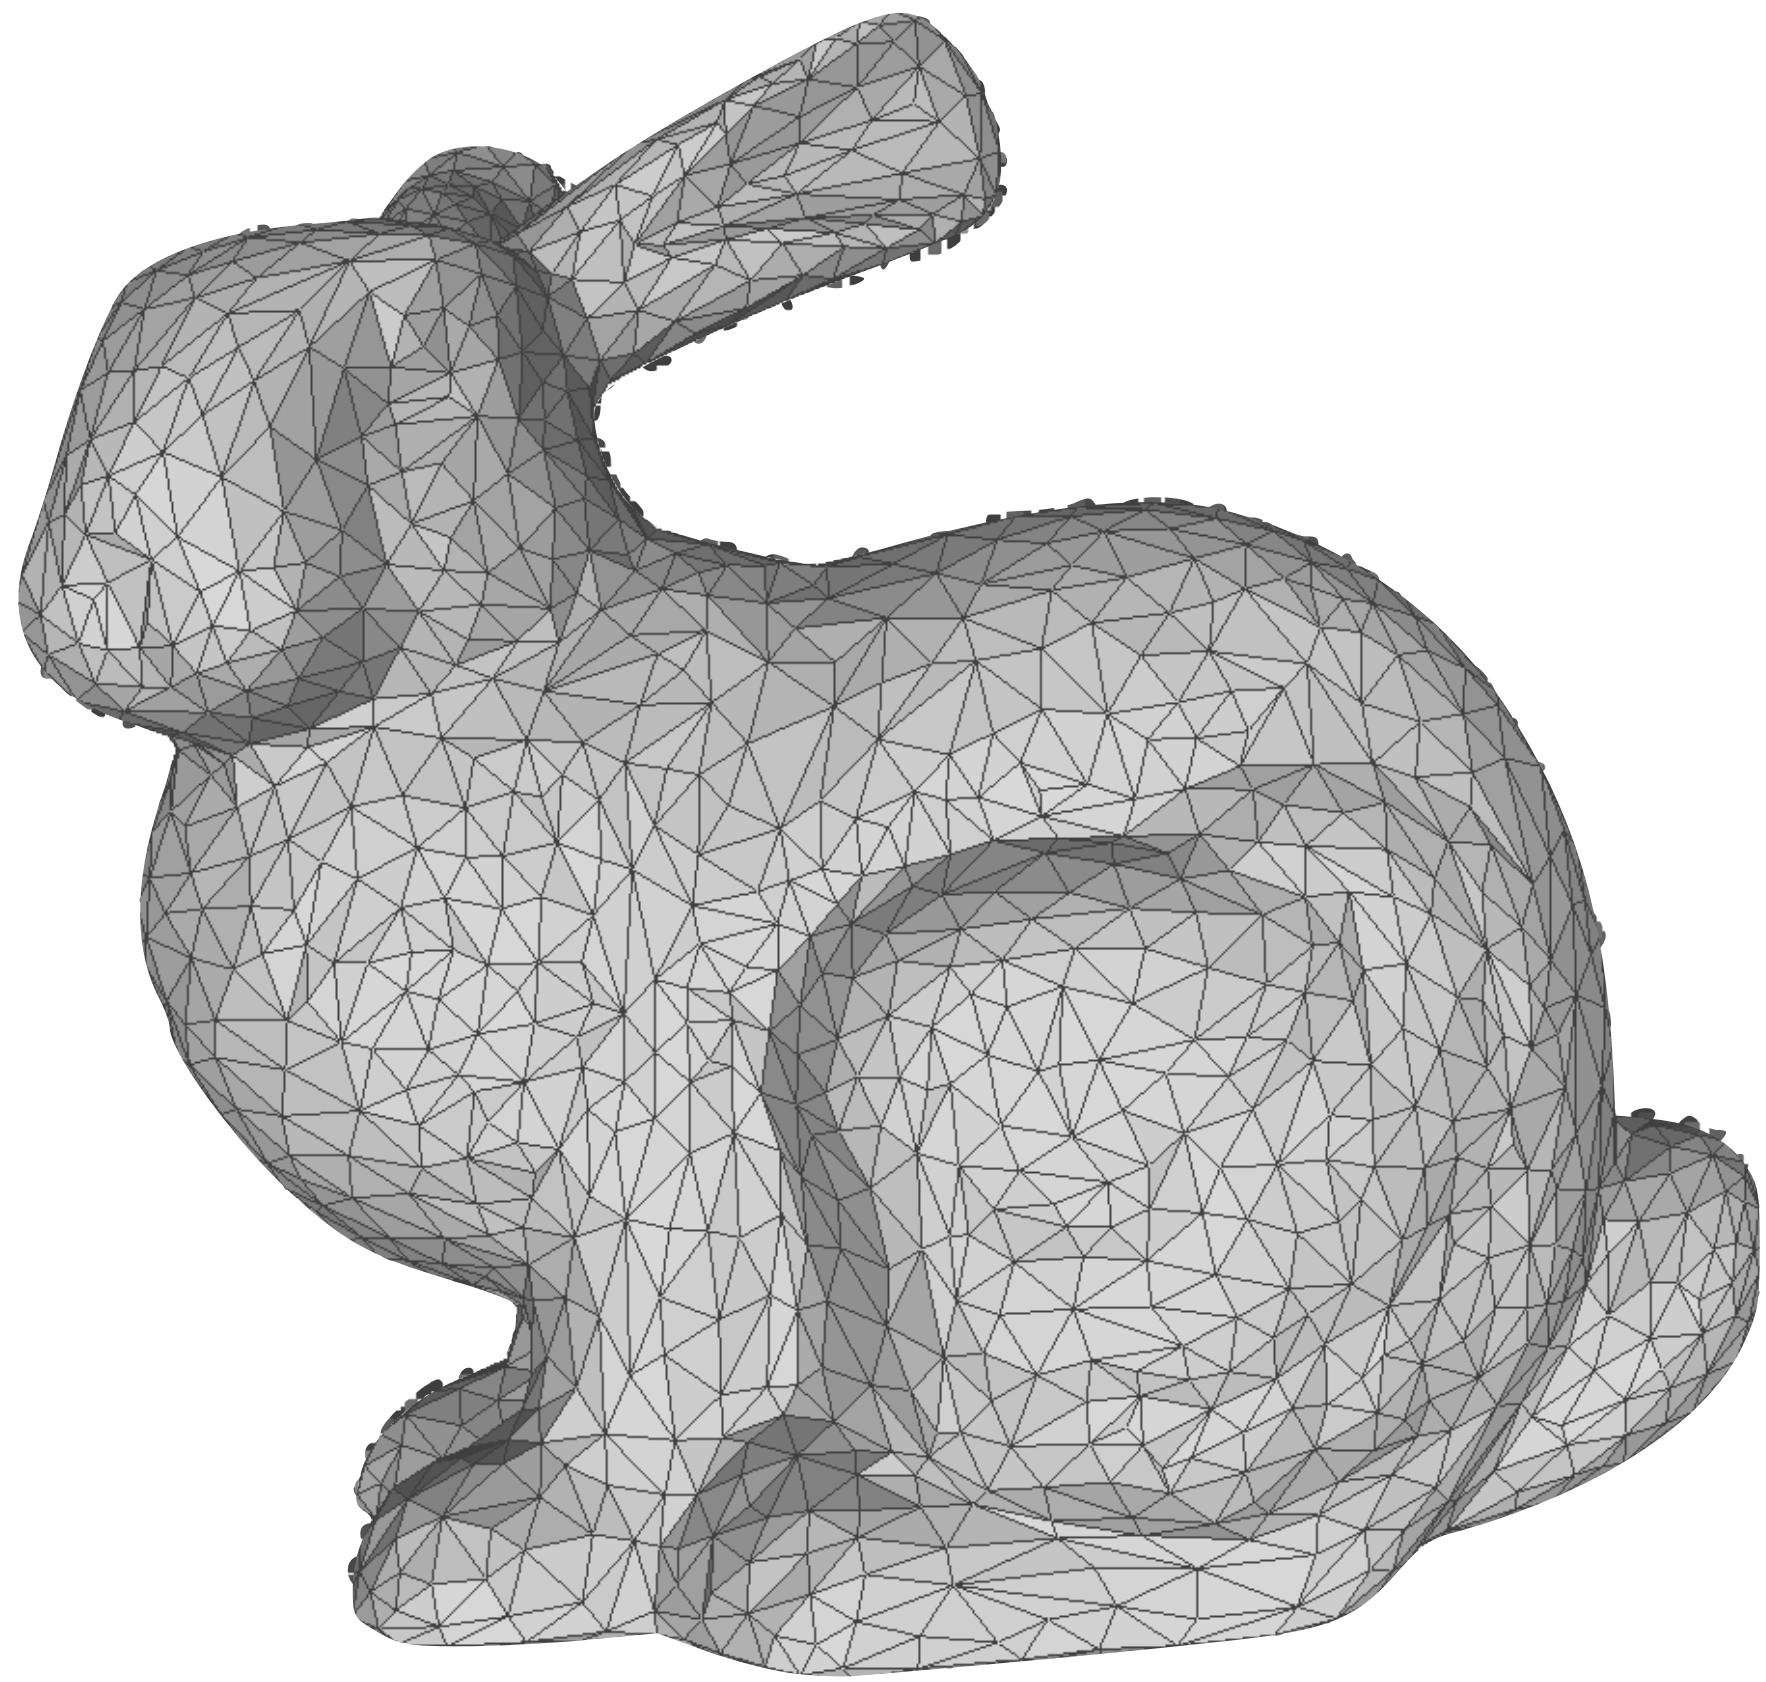
# Meshes



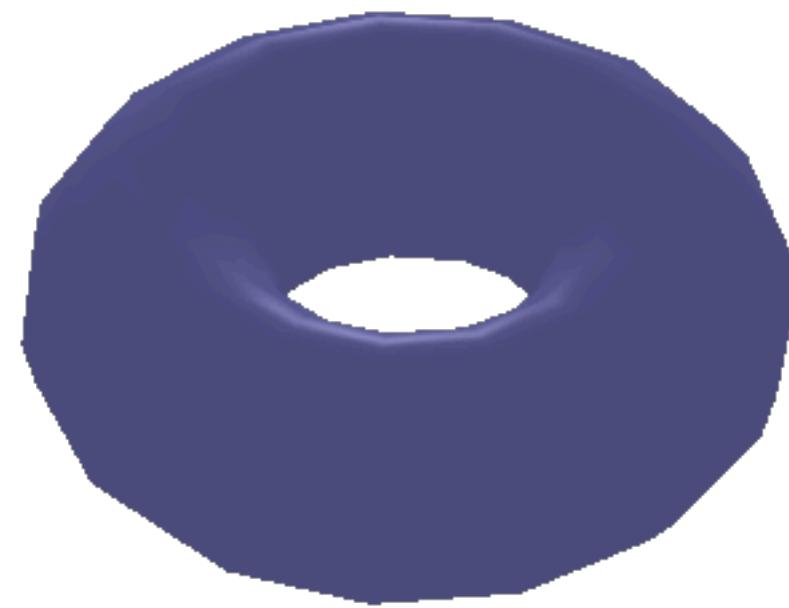
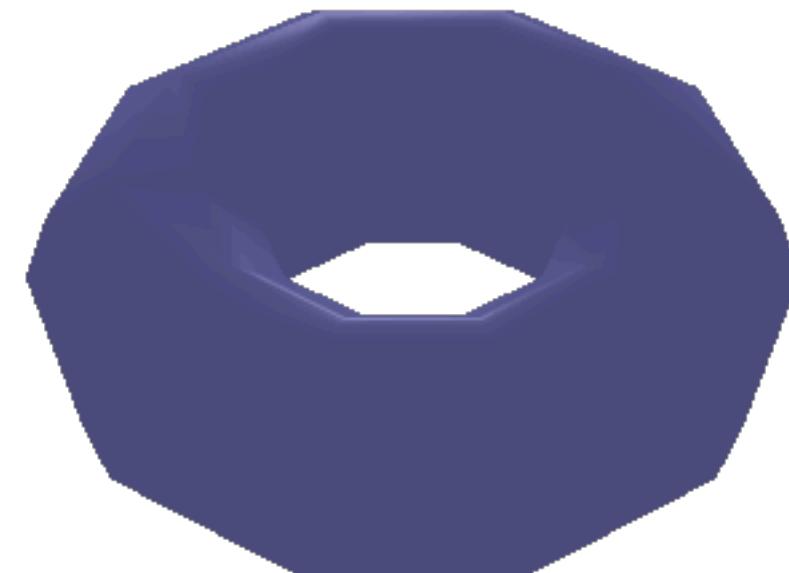
Piecewise linear approximations of underlying surface

Vertices + Faces

# Meshes

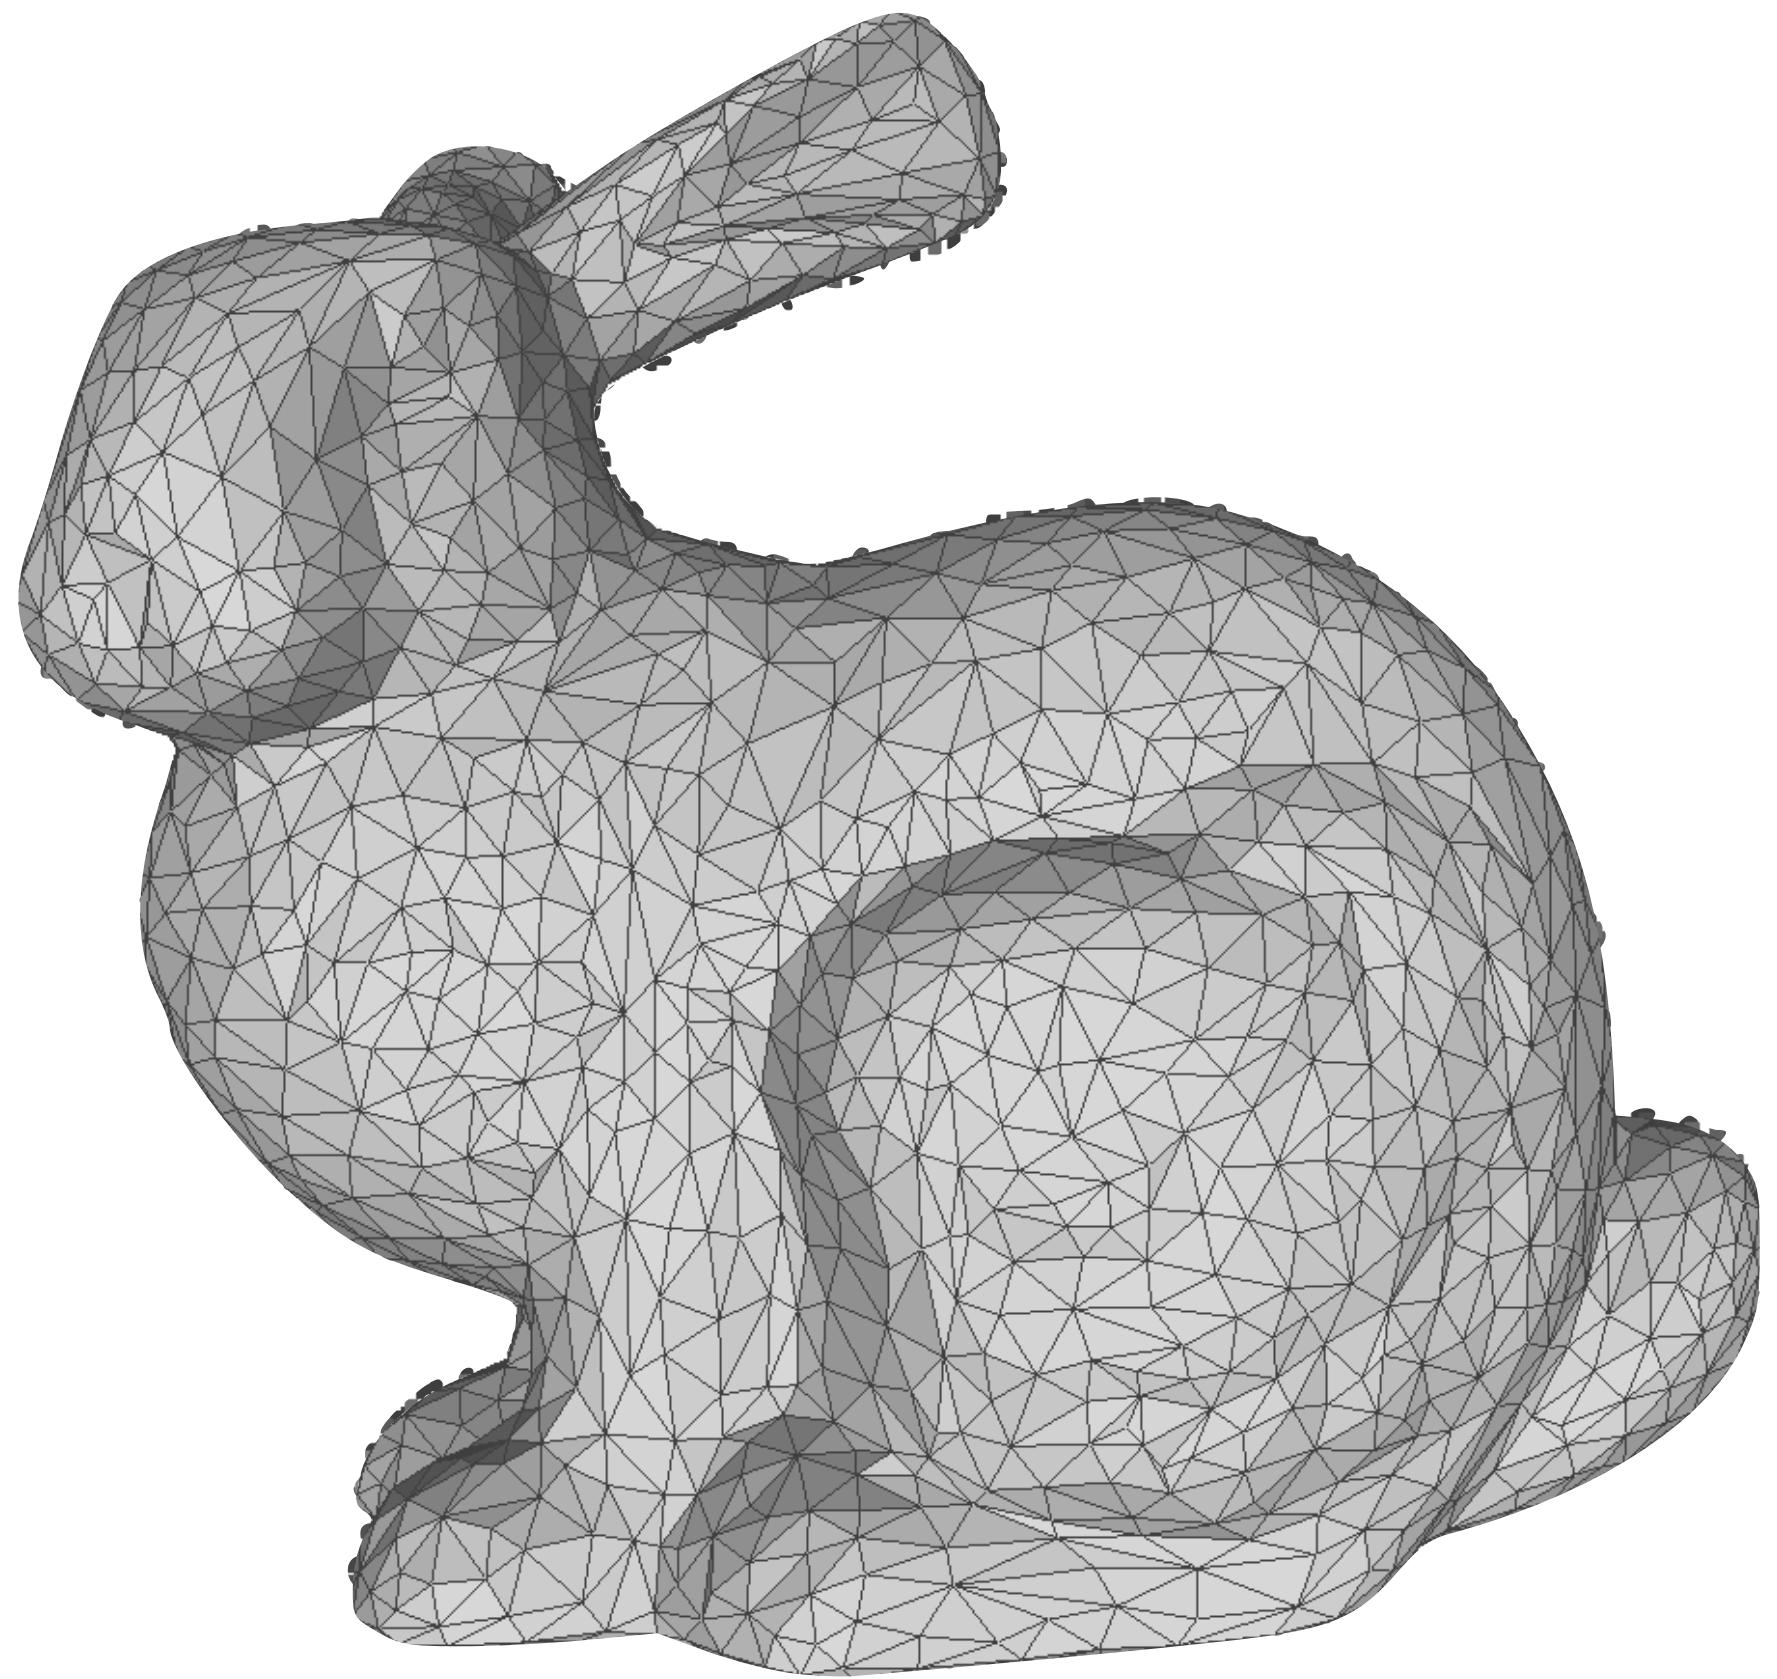


Vertices + Faces

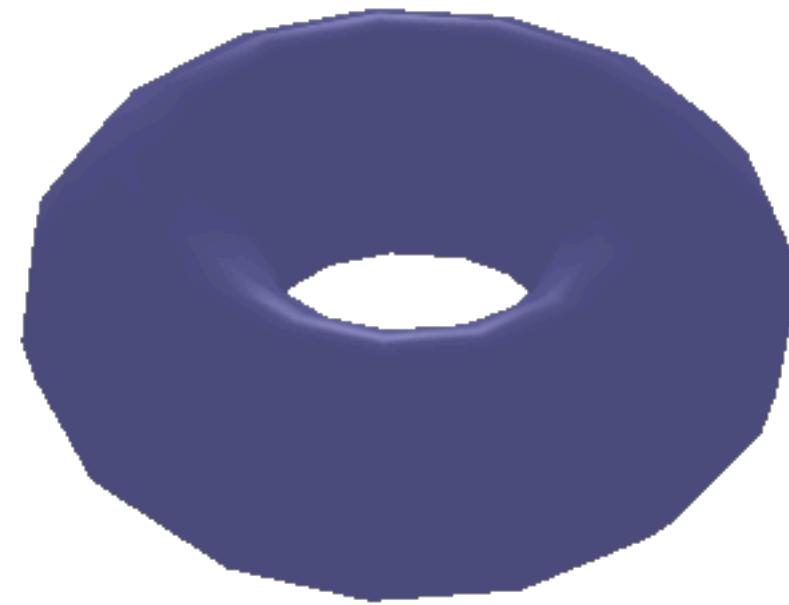
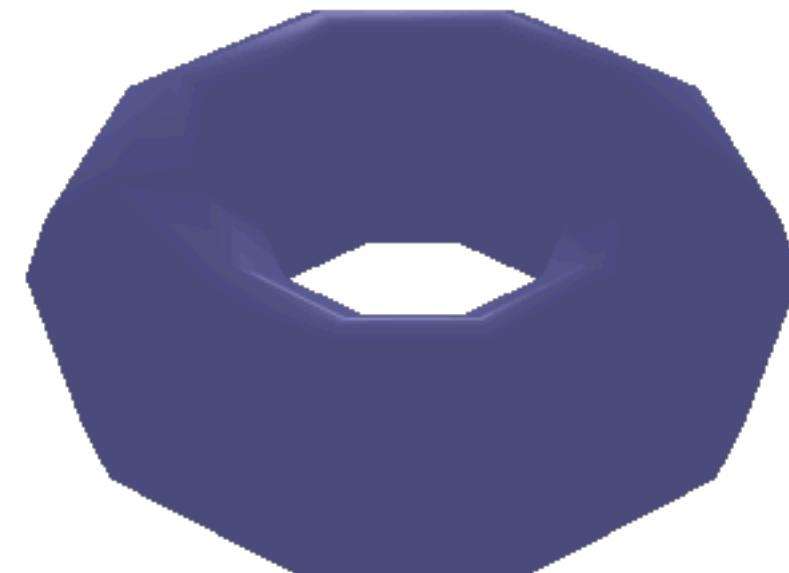


Piecewise linear  
approximations of  
underlying surface

# Meshes



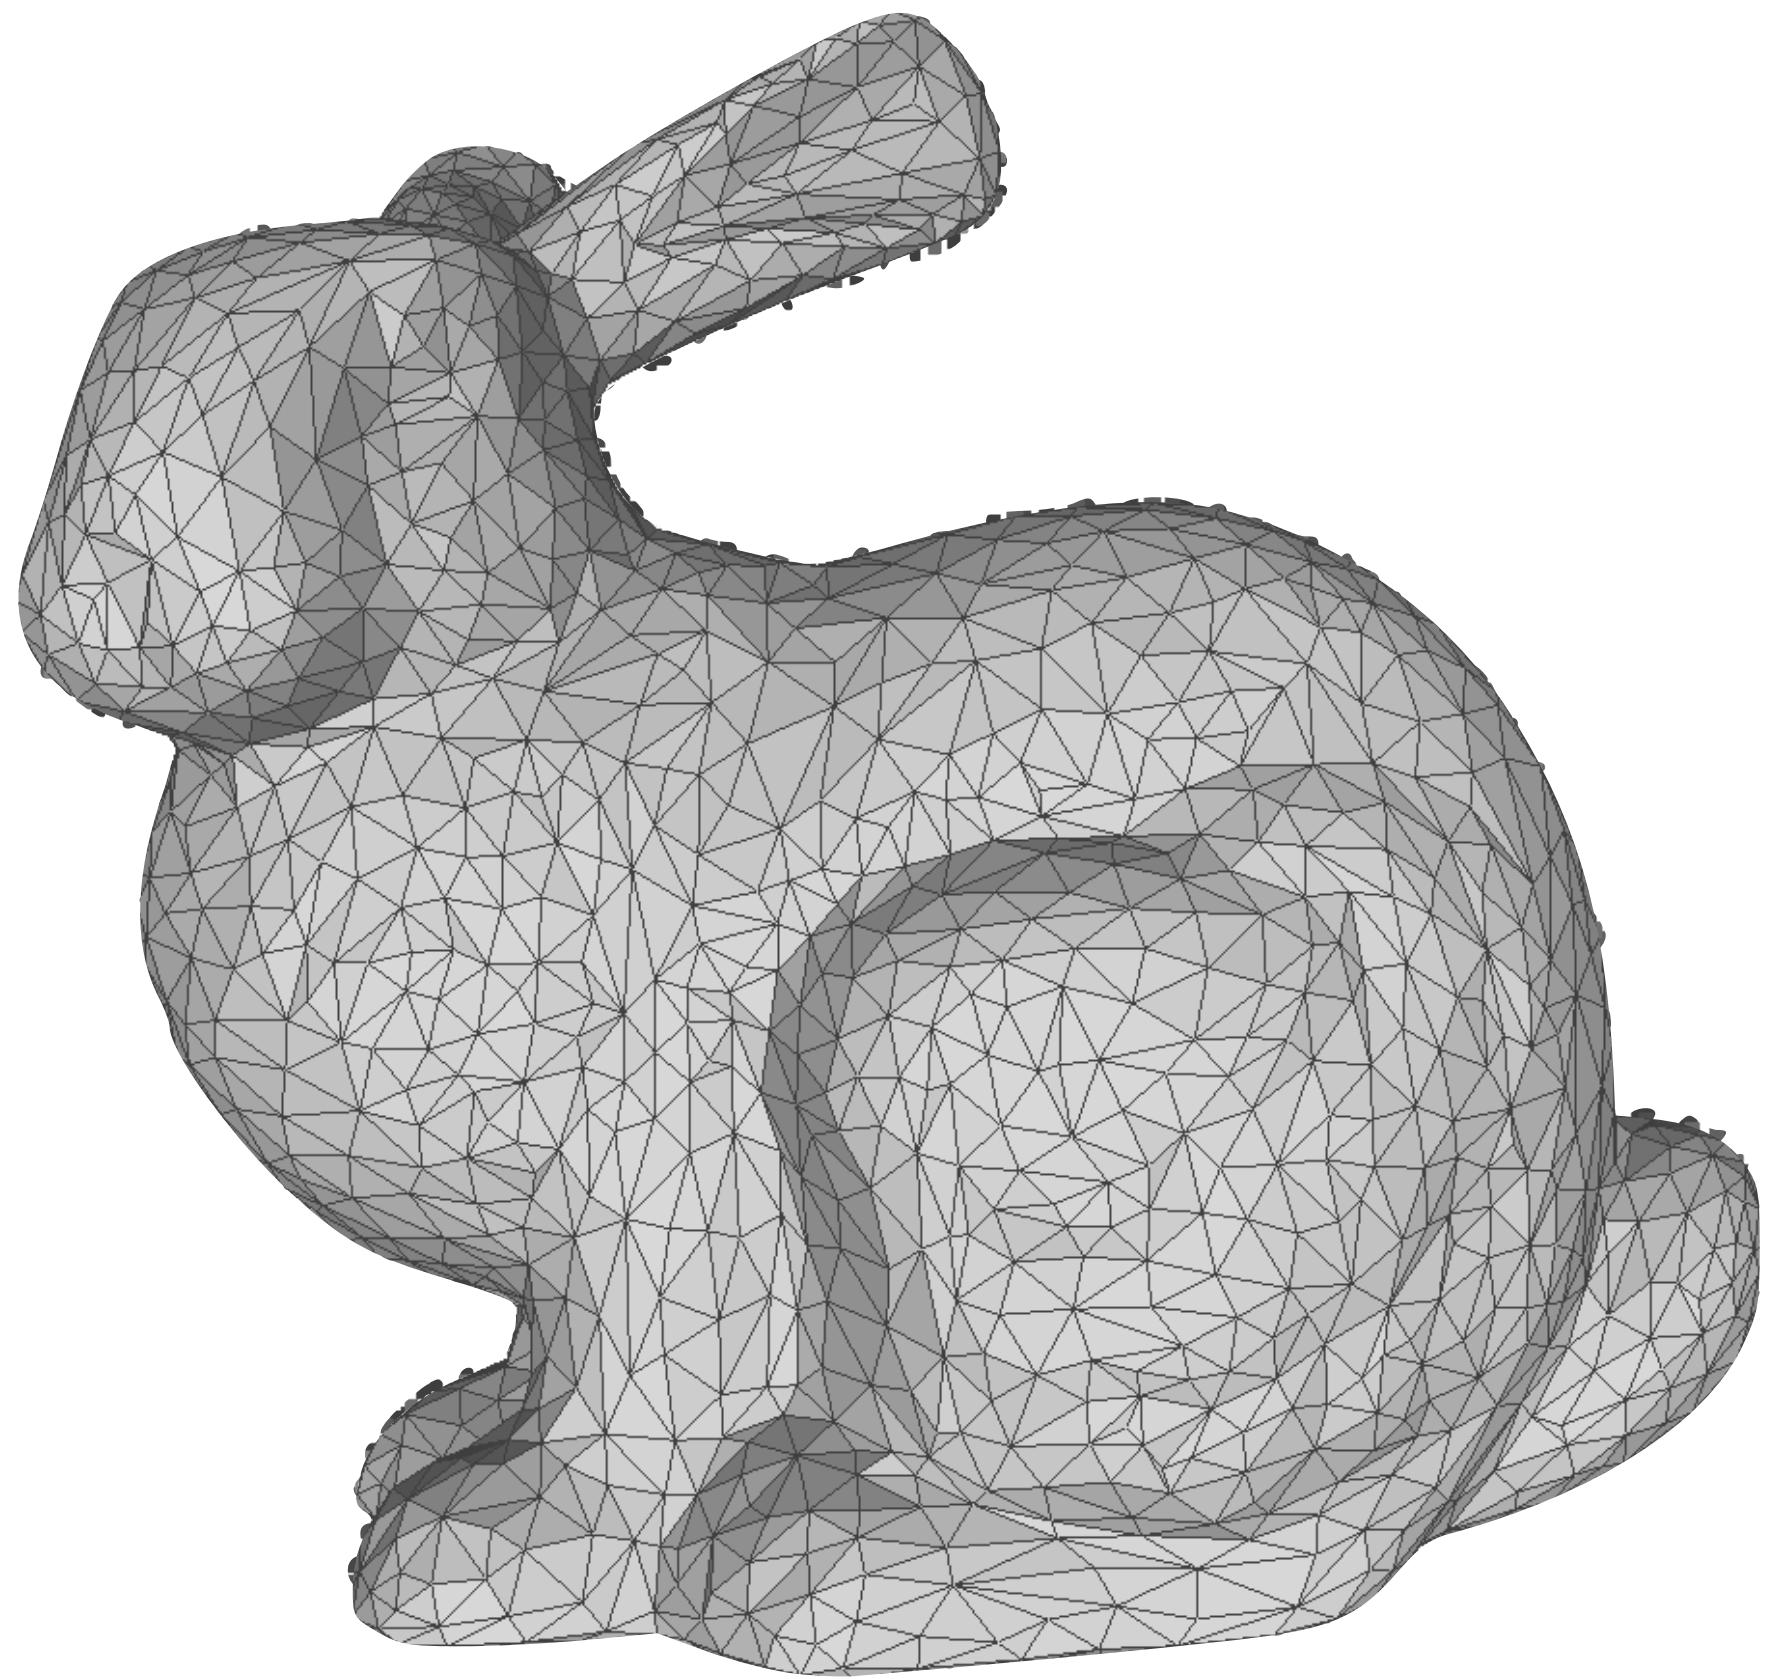
Vertices + Faces



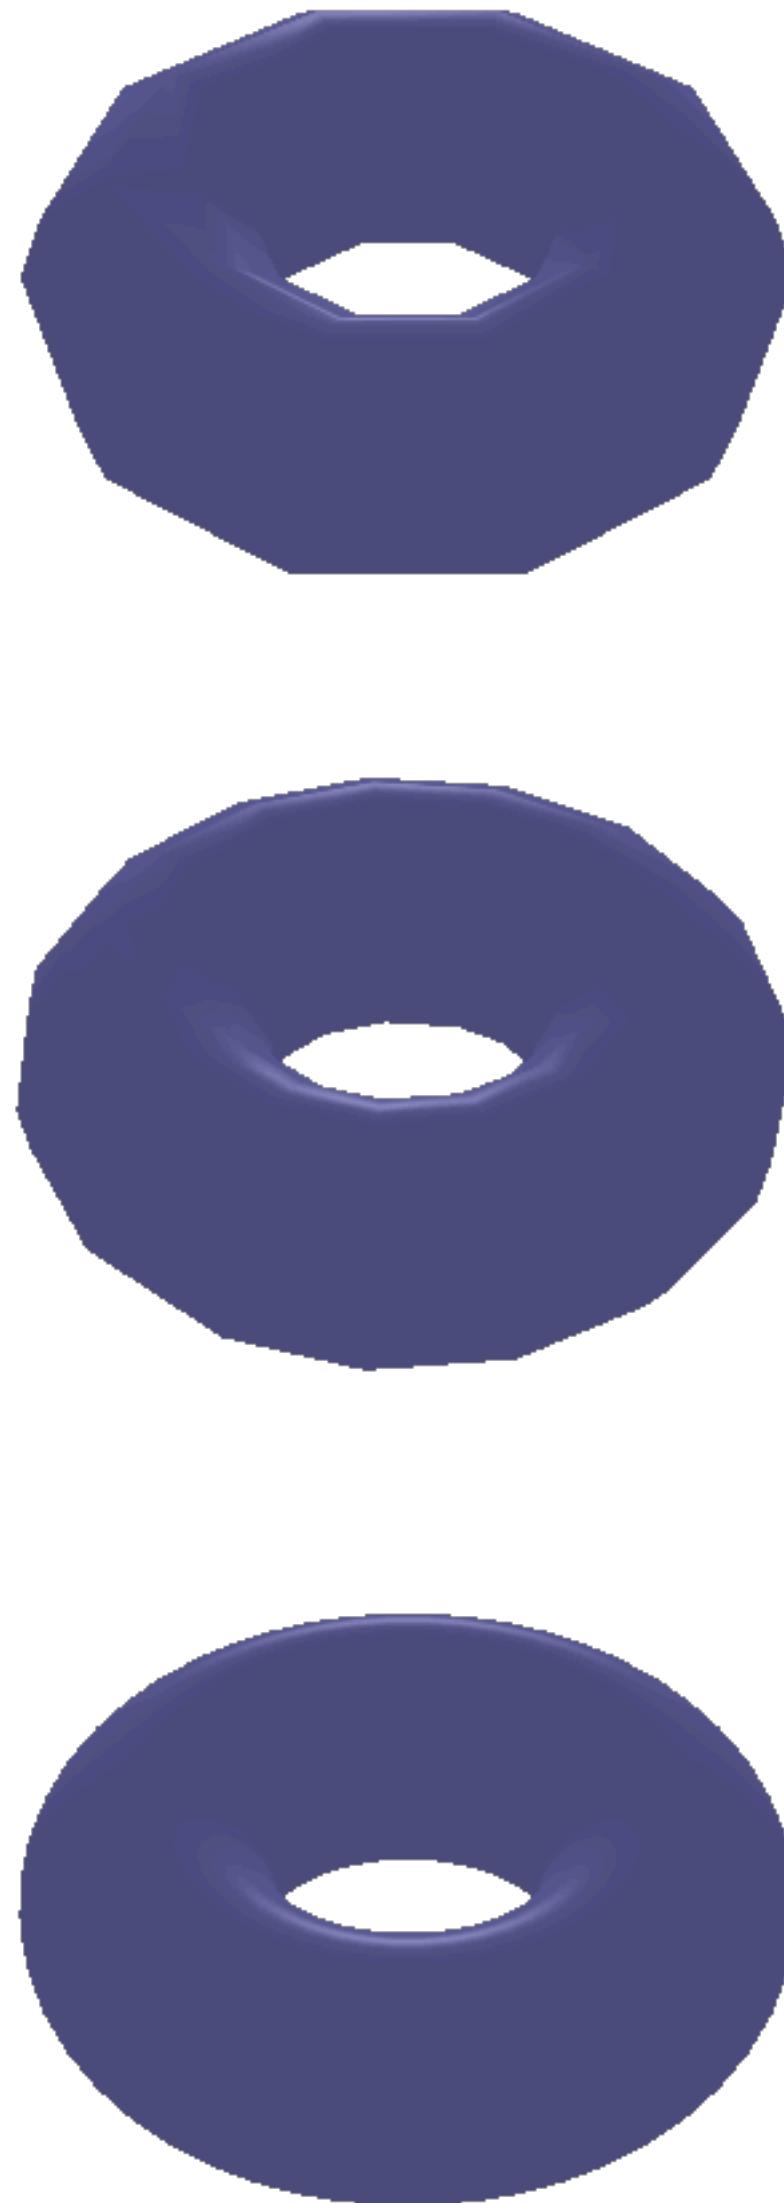
Piecewise linear  
approximations of  
underlying surface

More elements allow  
better approximation

# Meshes



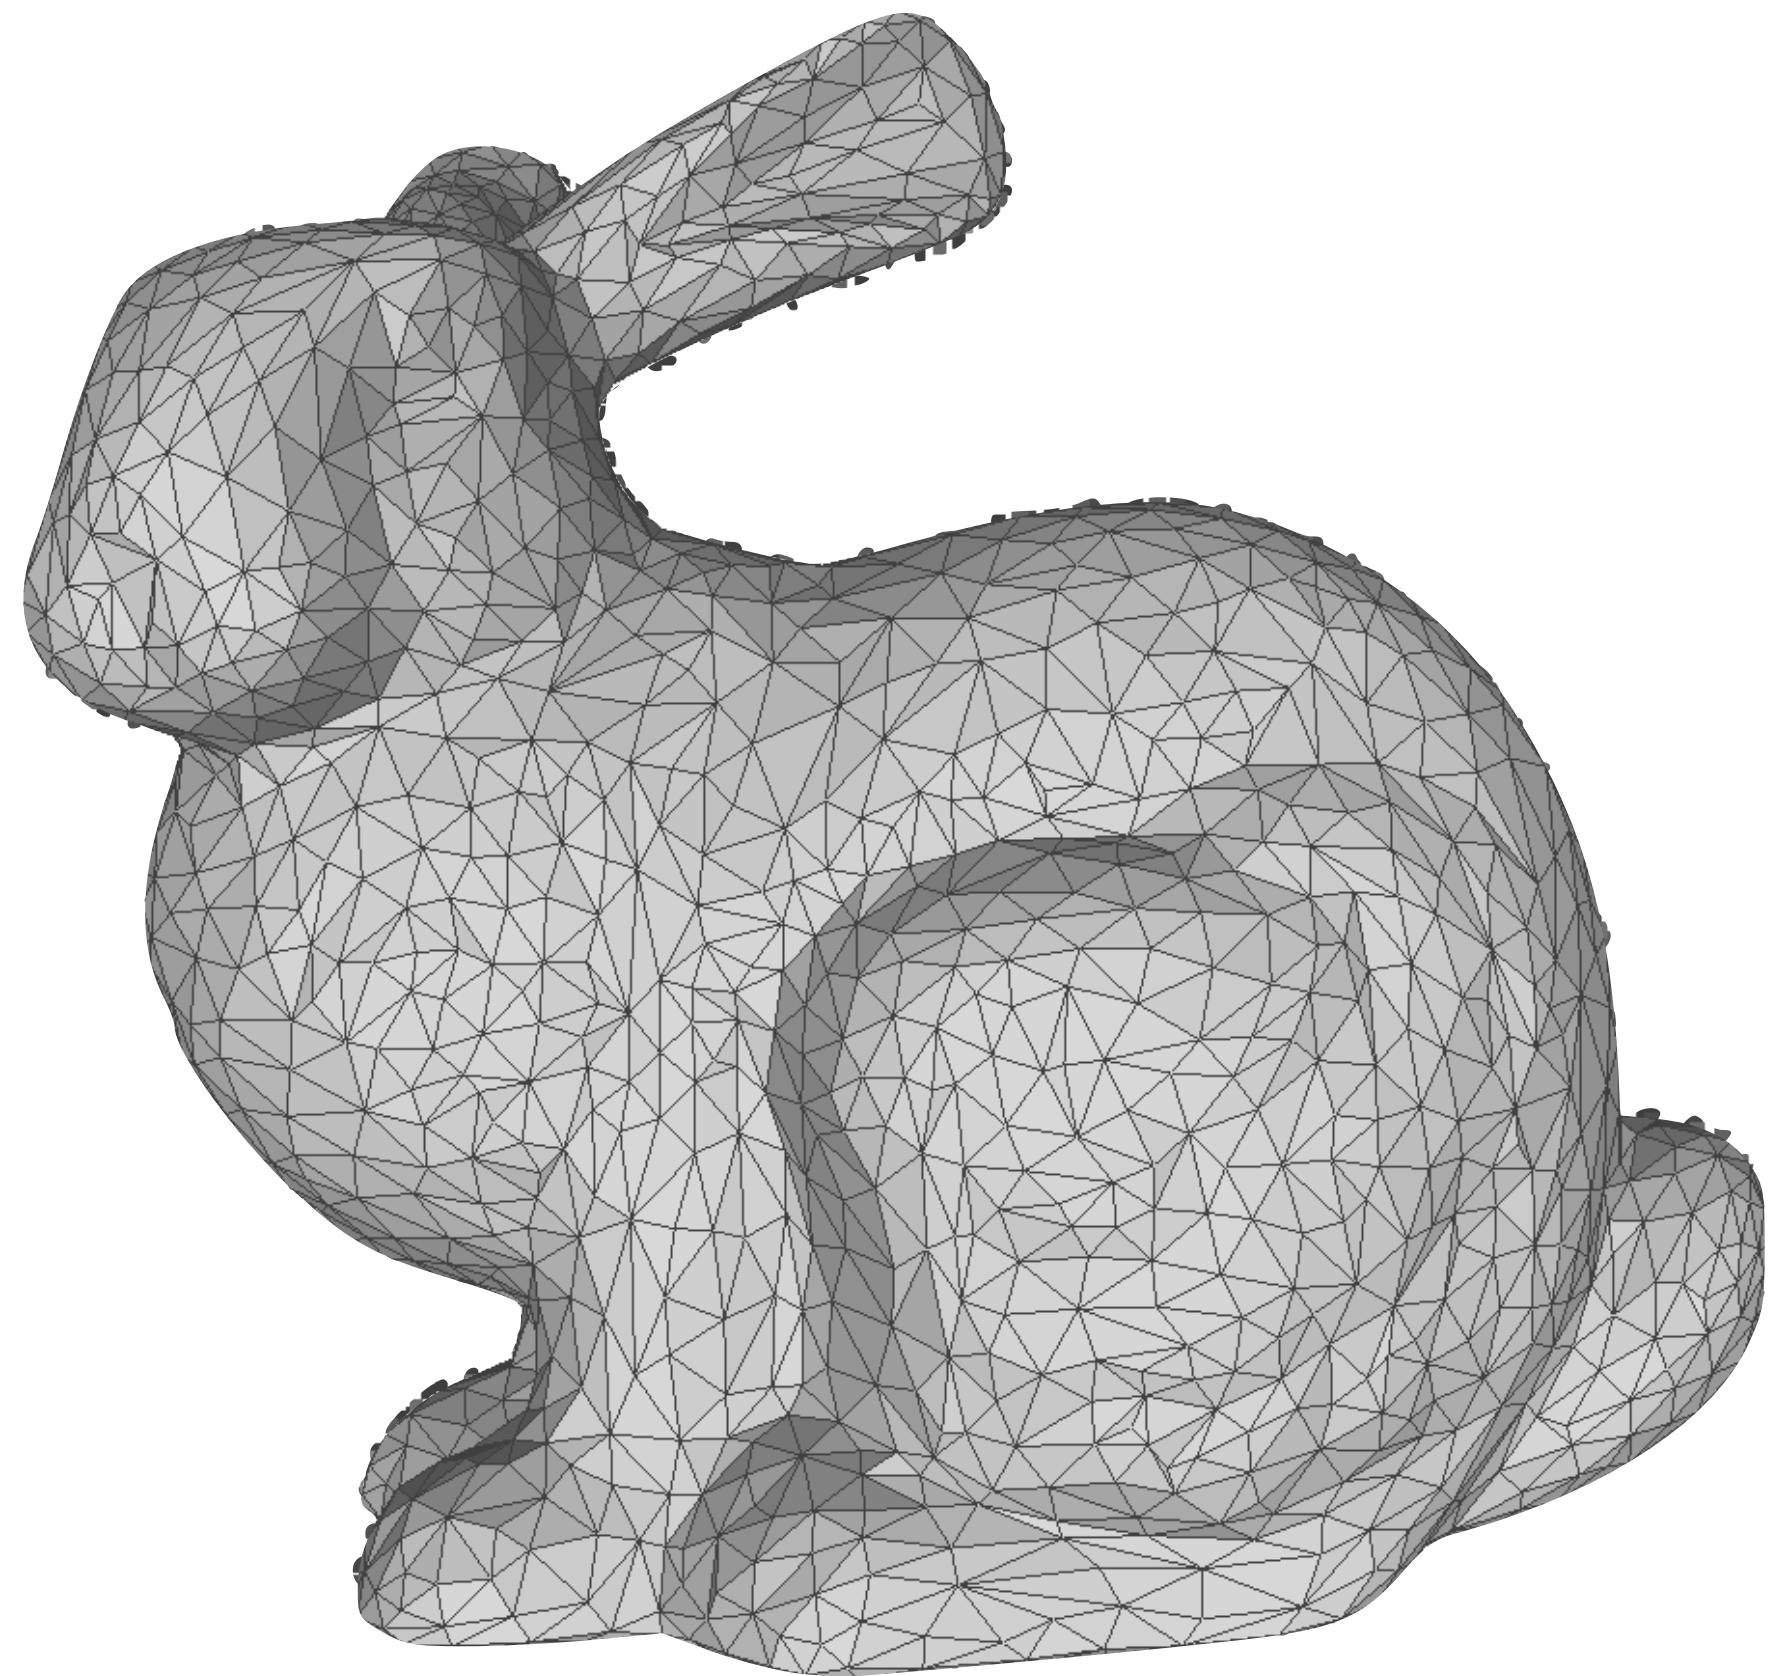
Vertices + Faces



Piecewise linear  
approximations of  
underlying surface

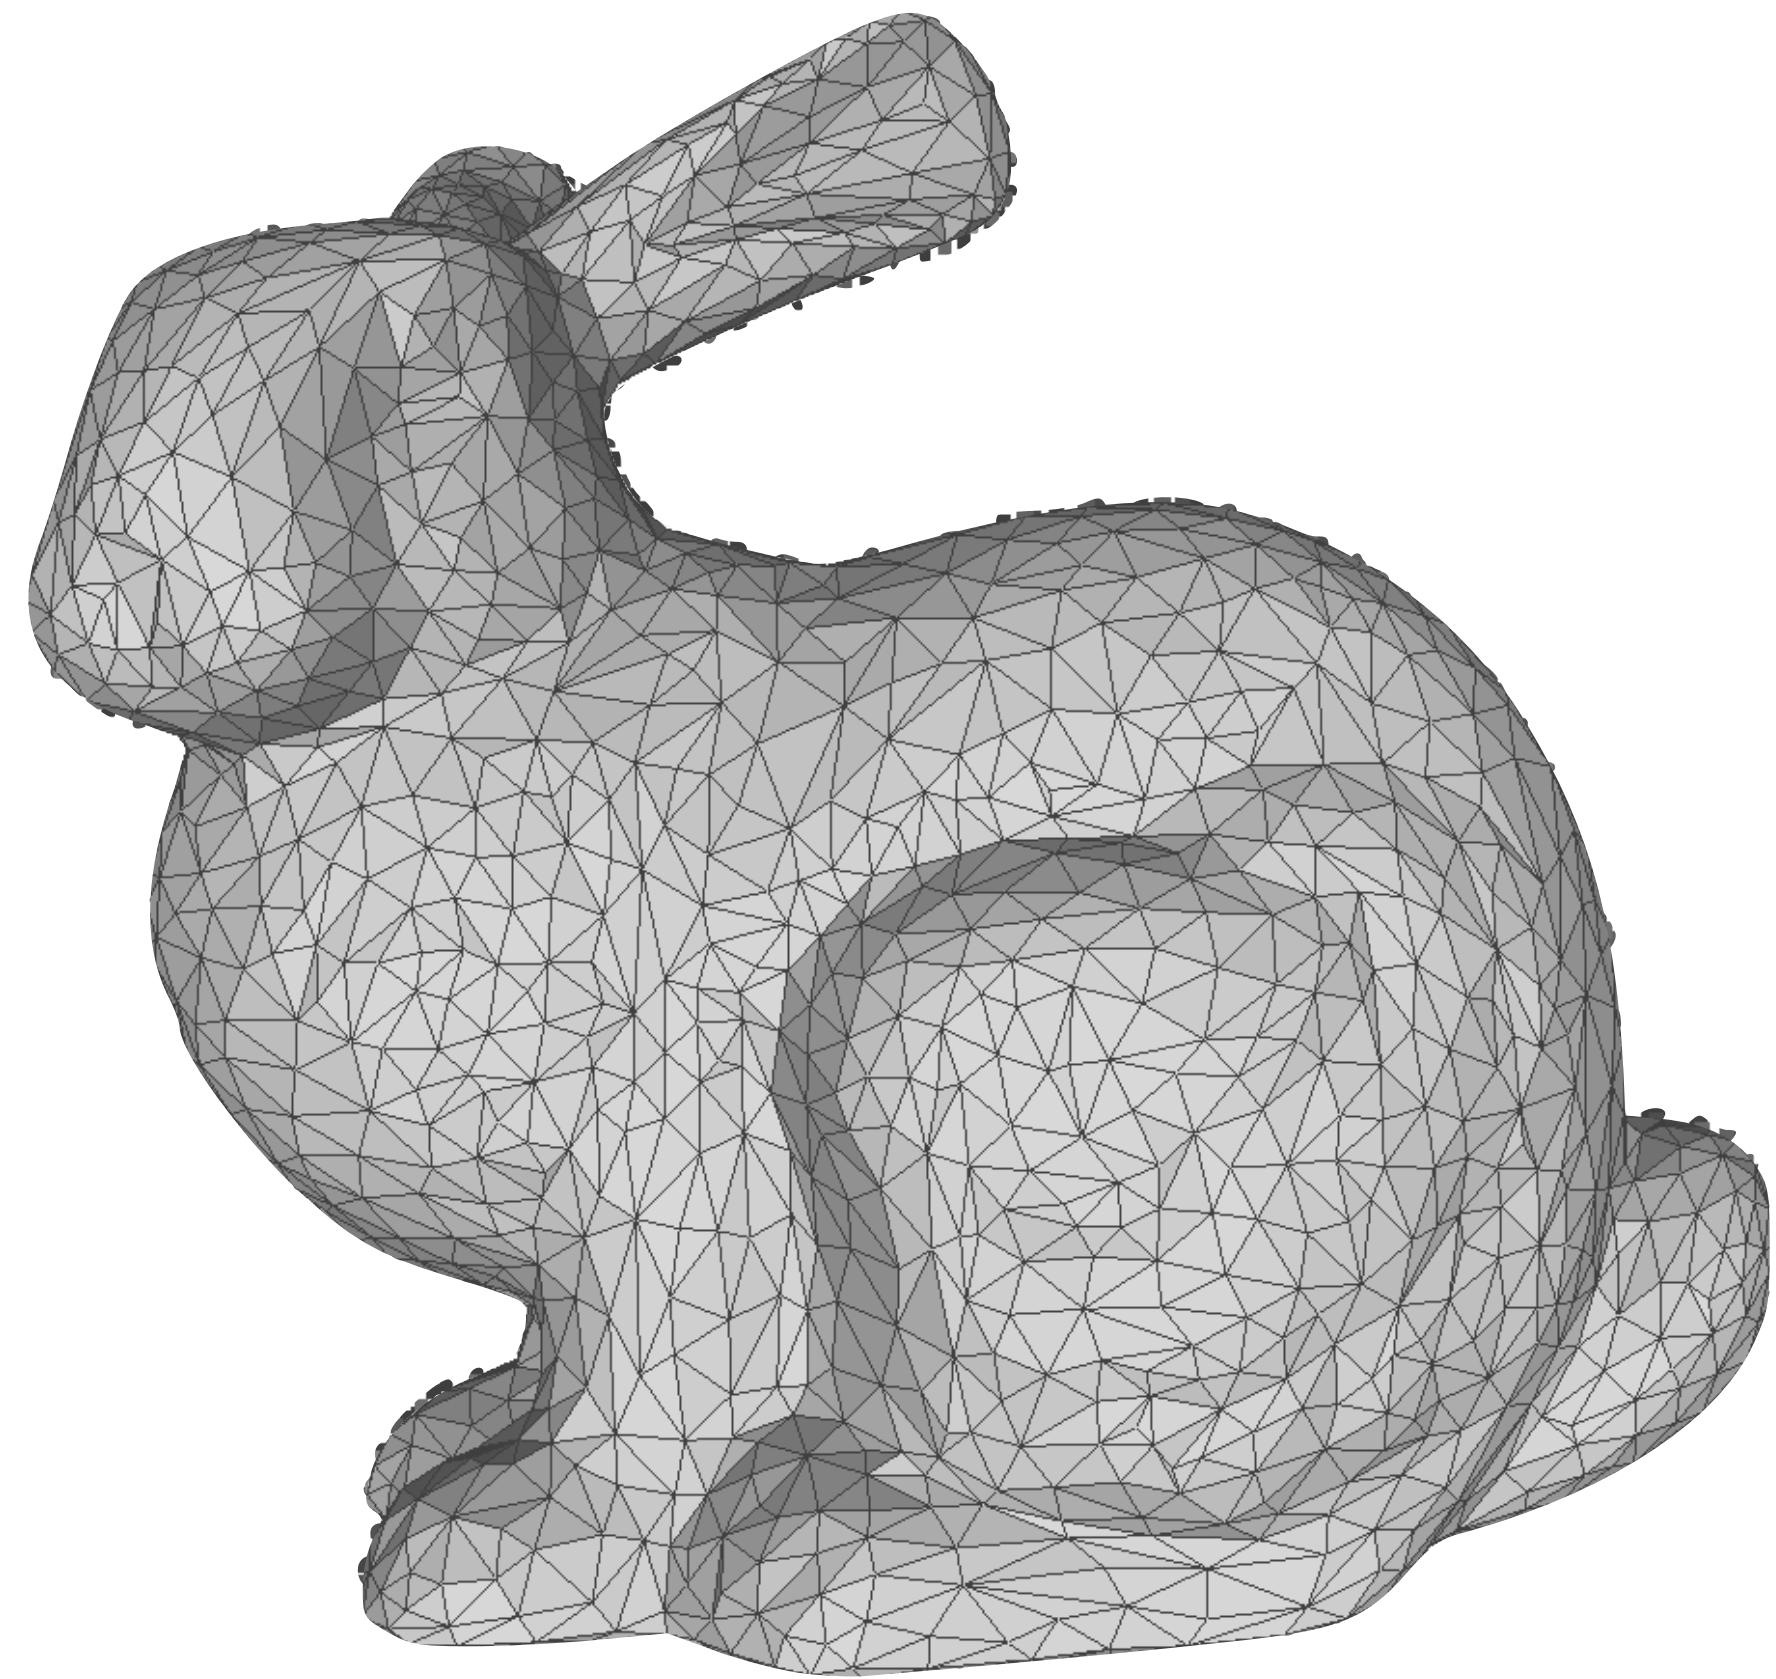
More elements allow  
better approximation

# Meshes

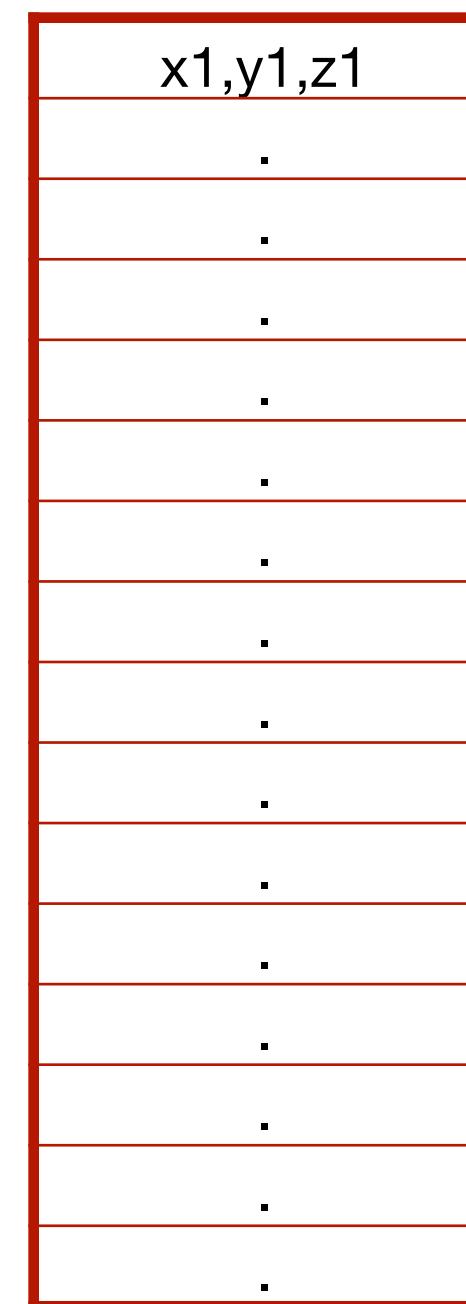


Vertices + Faces

# Meshes

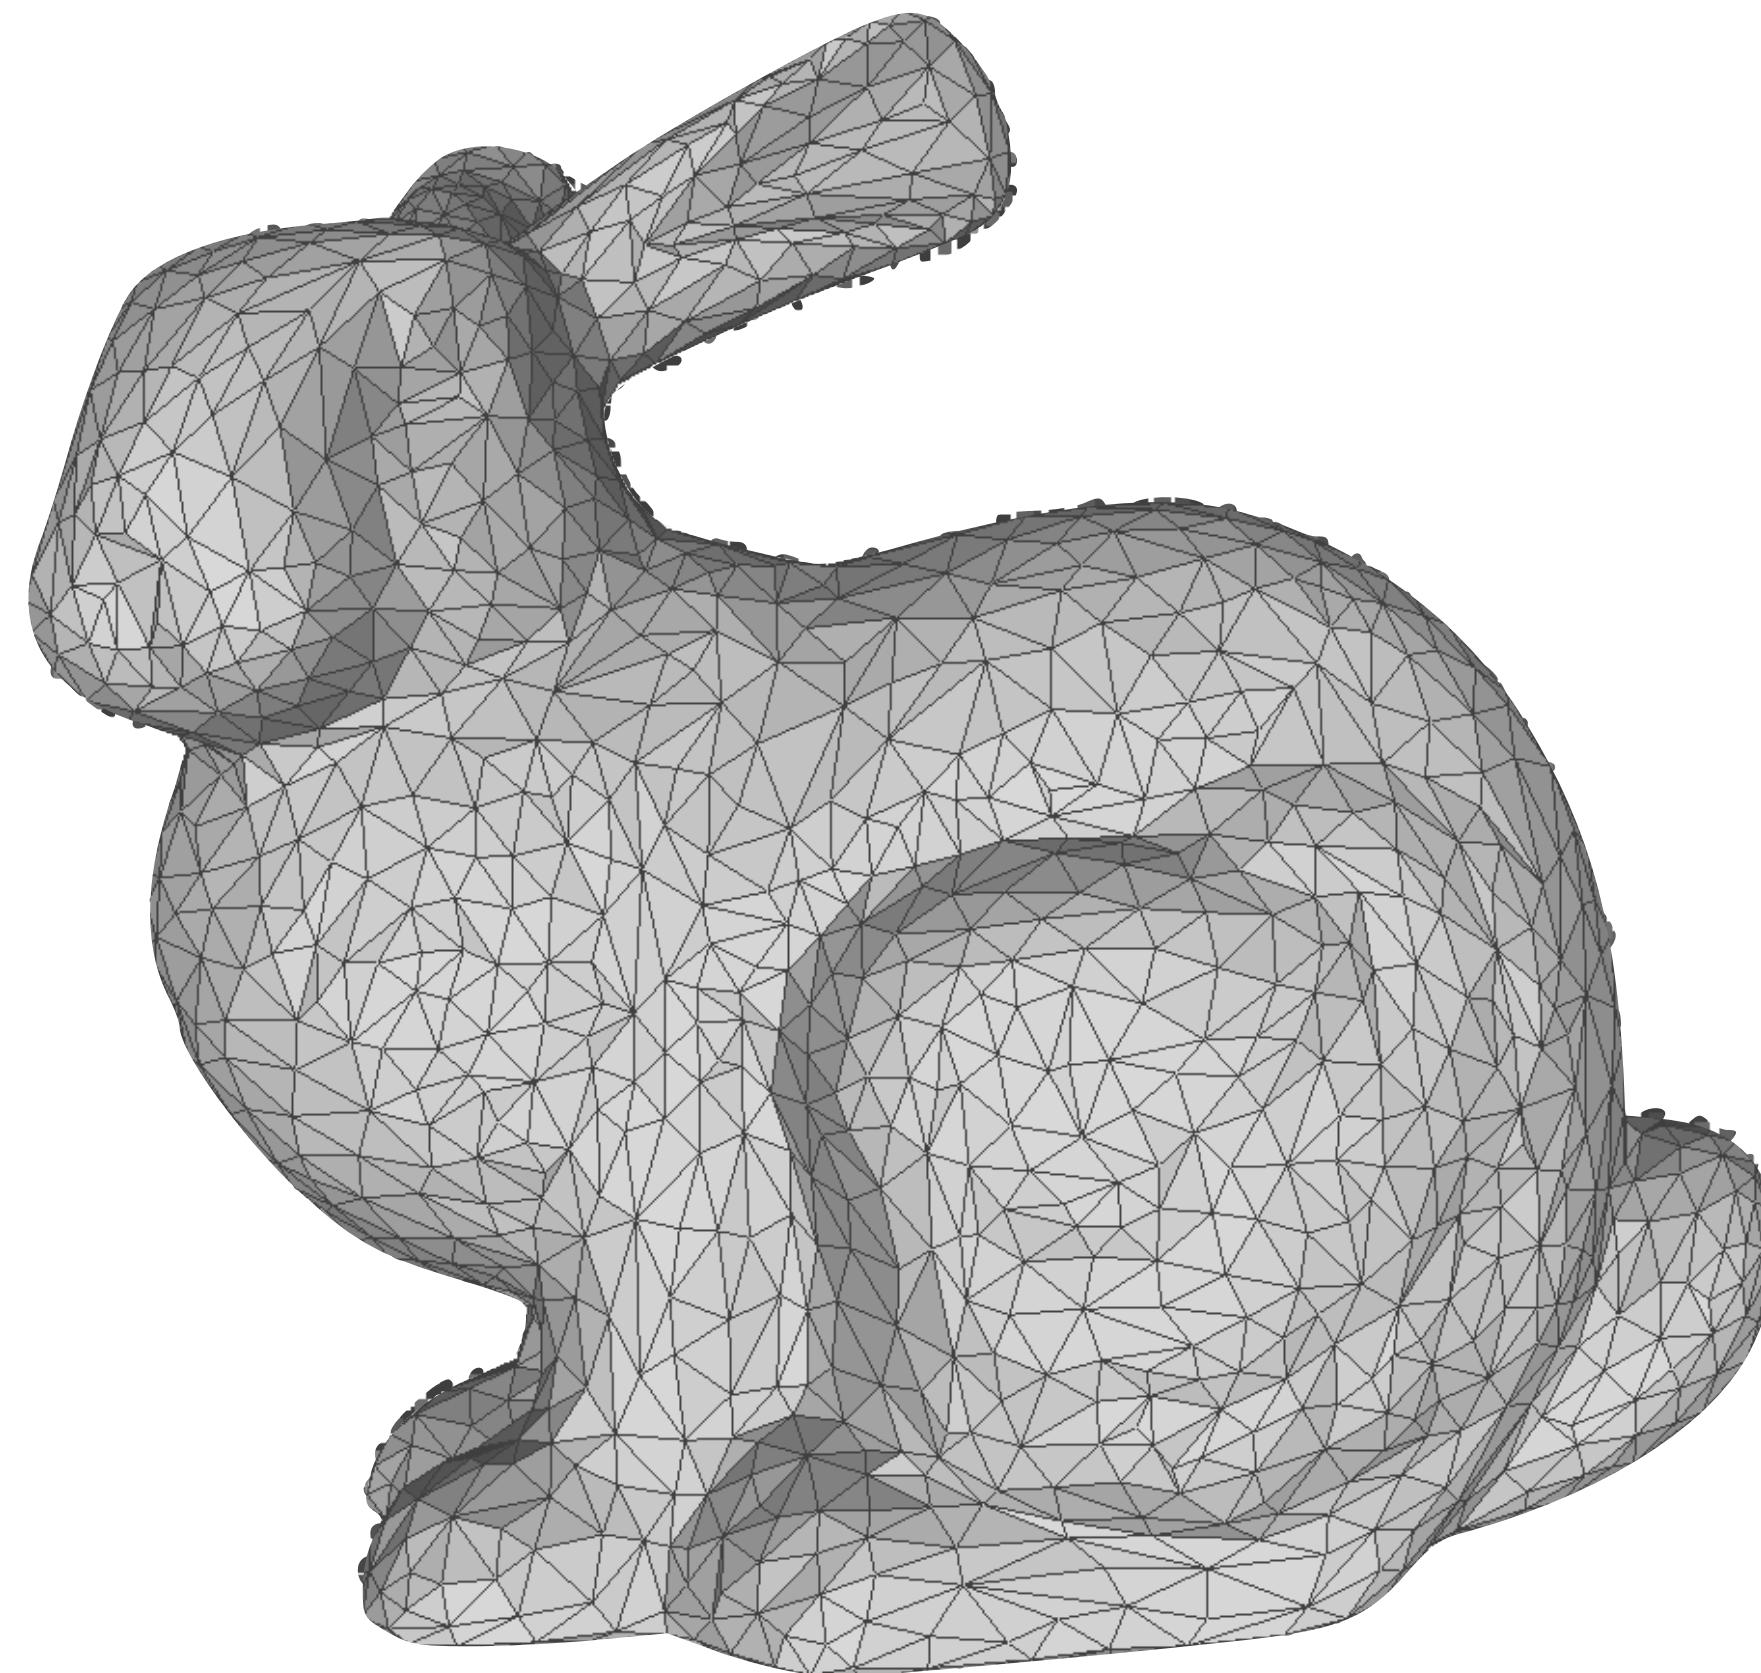


# Vertices



# Vertices + Faces

# Meshes



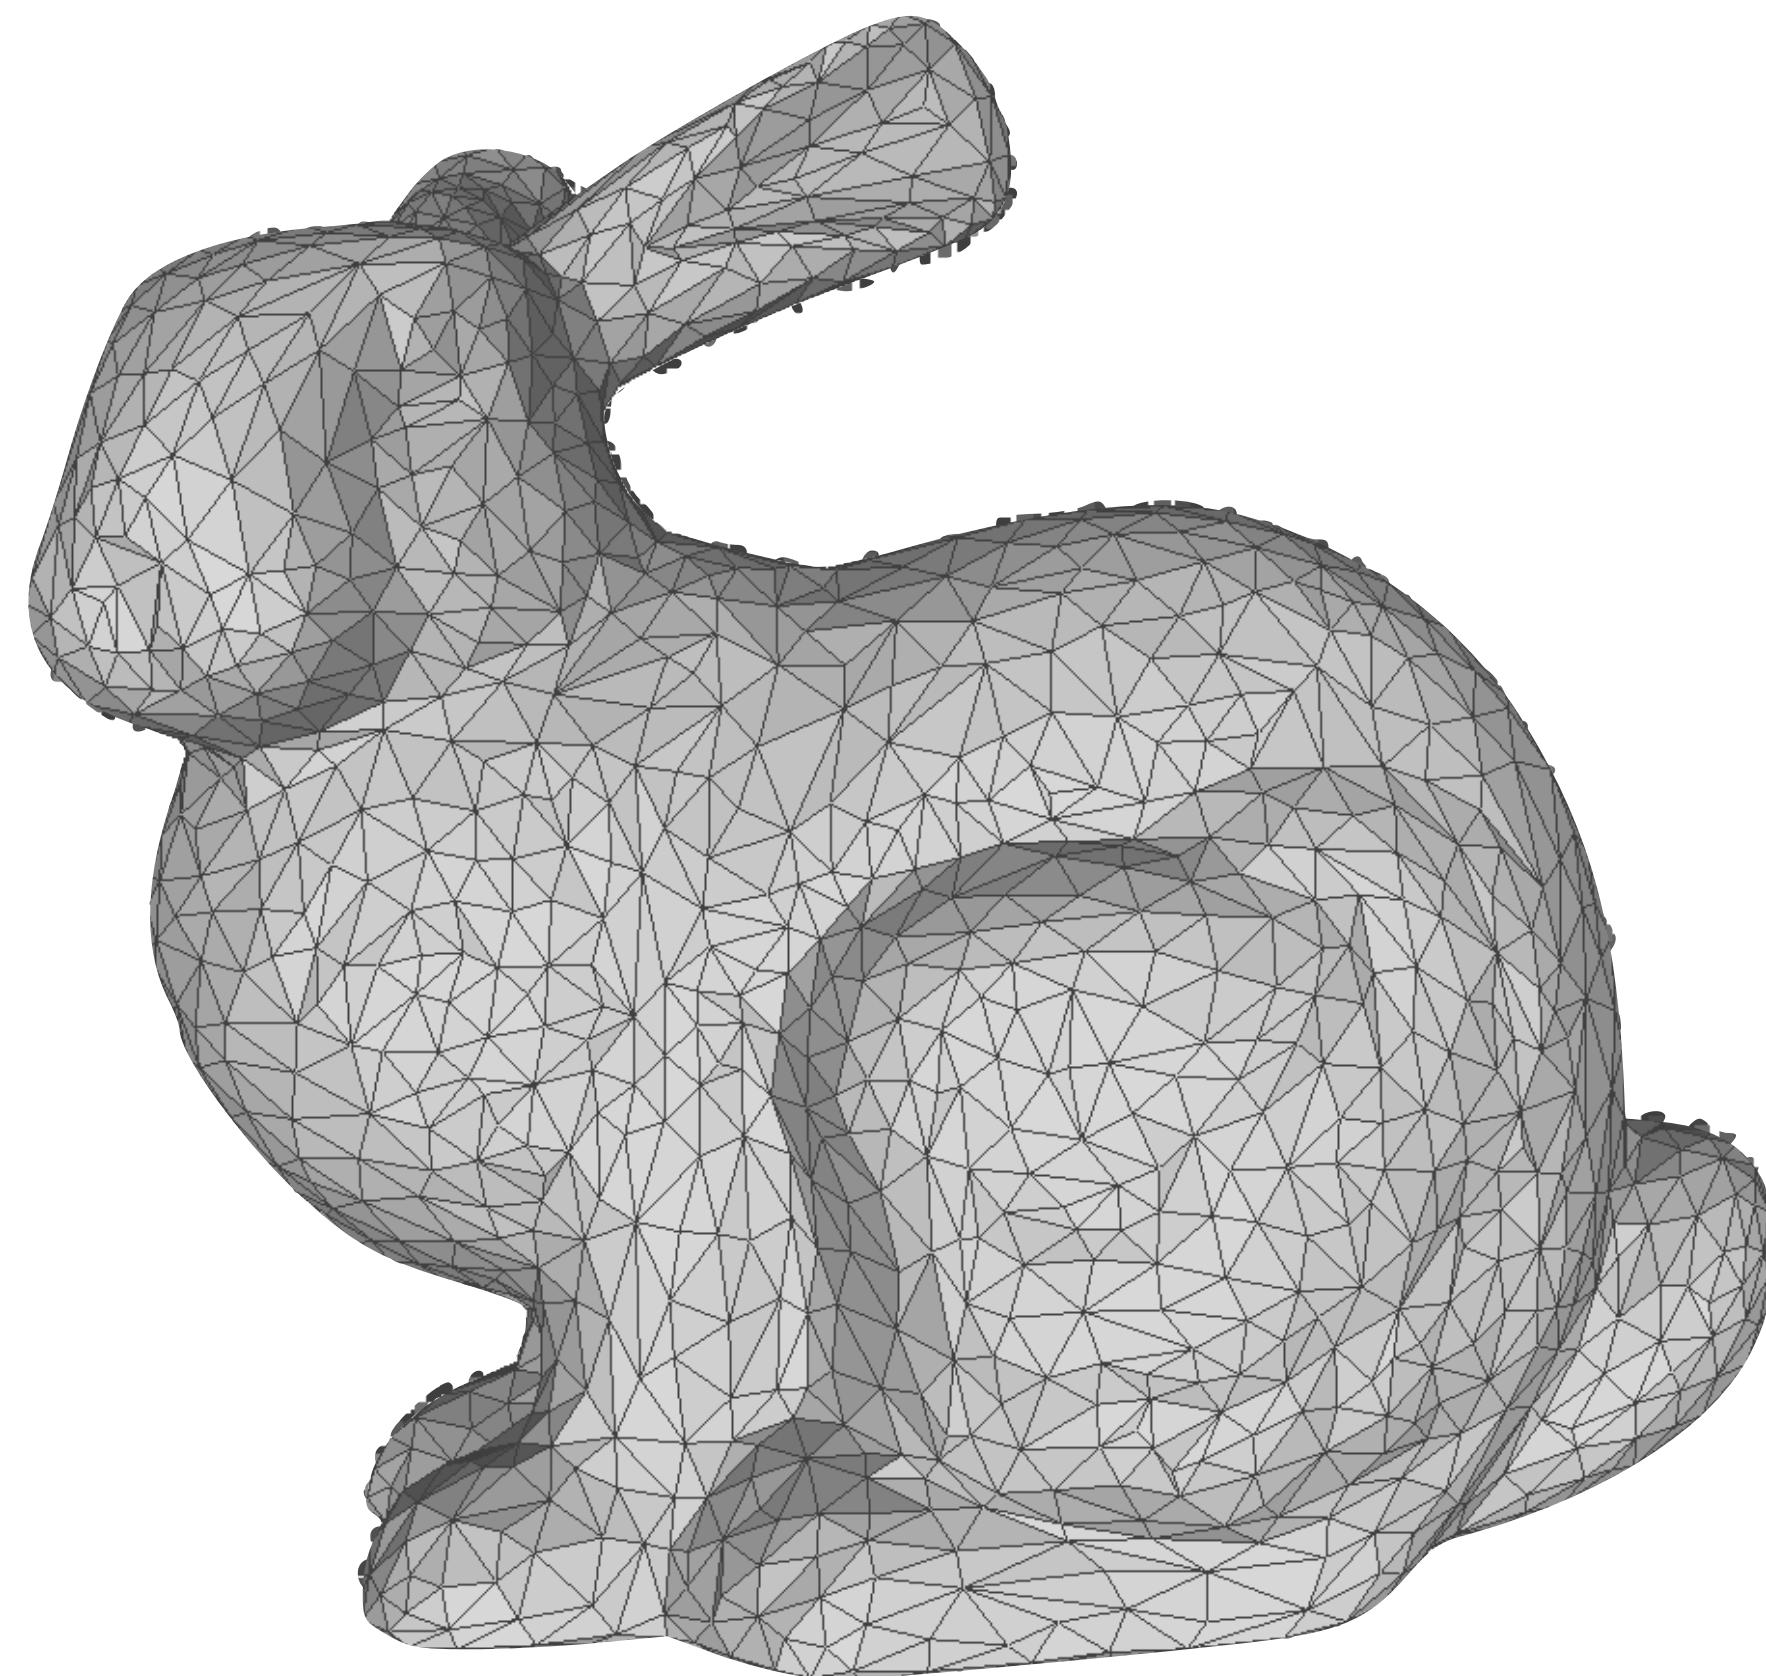
# Vertices + Faces

# Vertices

x1,y1,z1

# Positions of Vertices

# Meshes



# Vertices + Faces

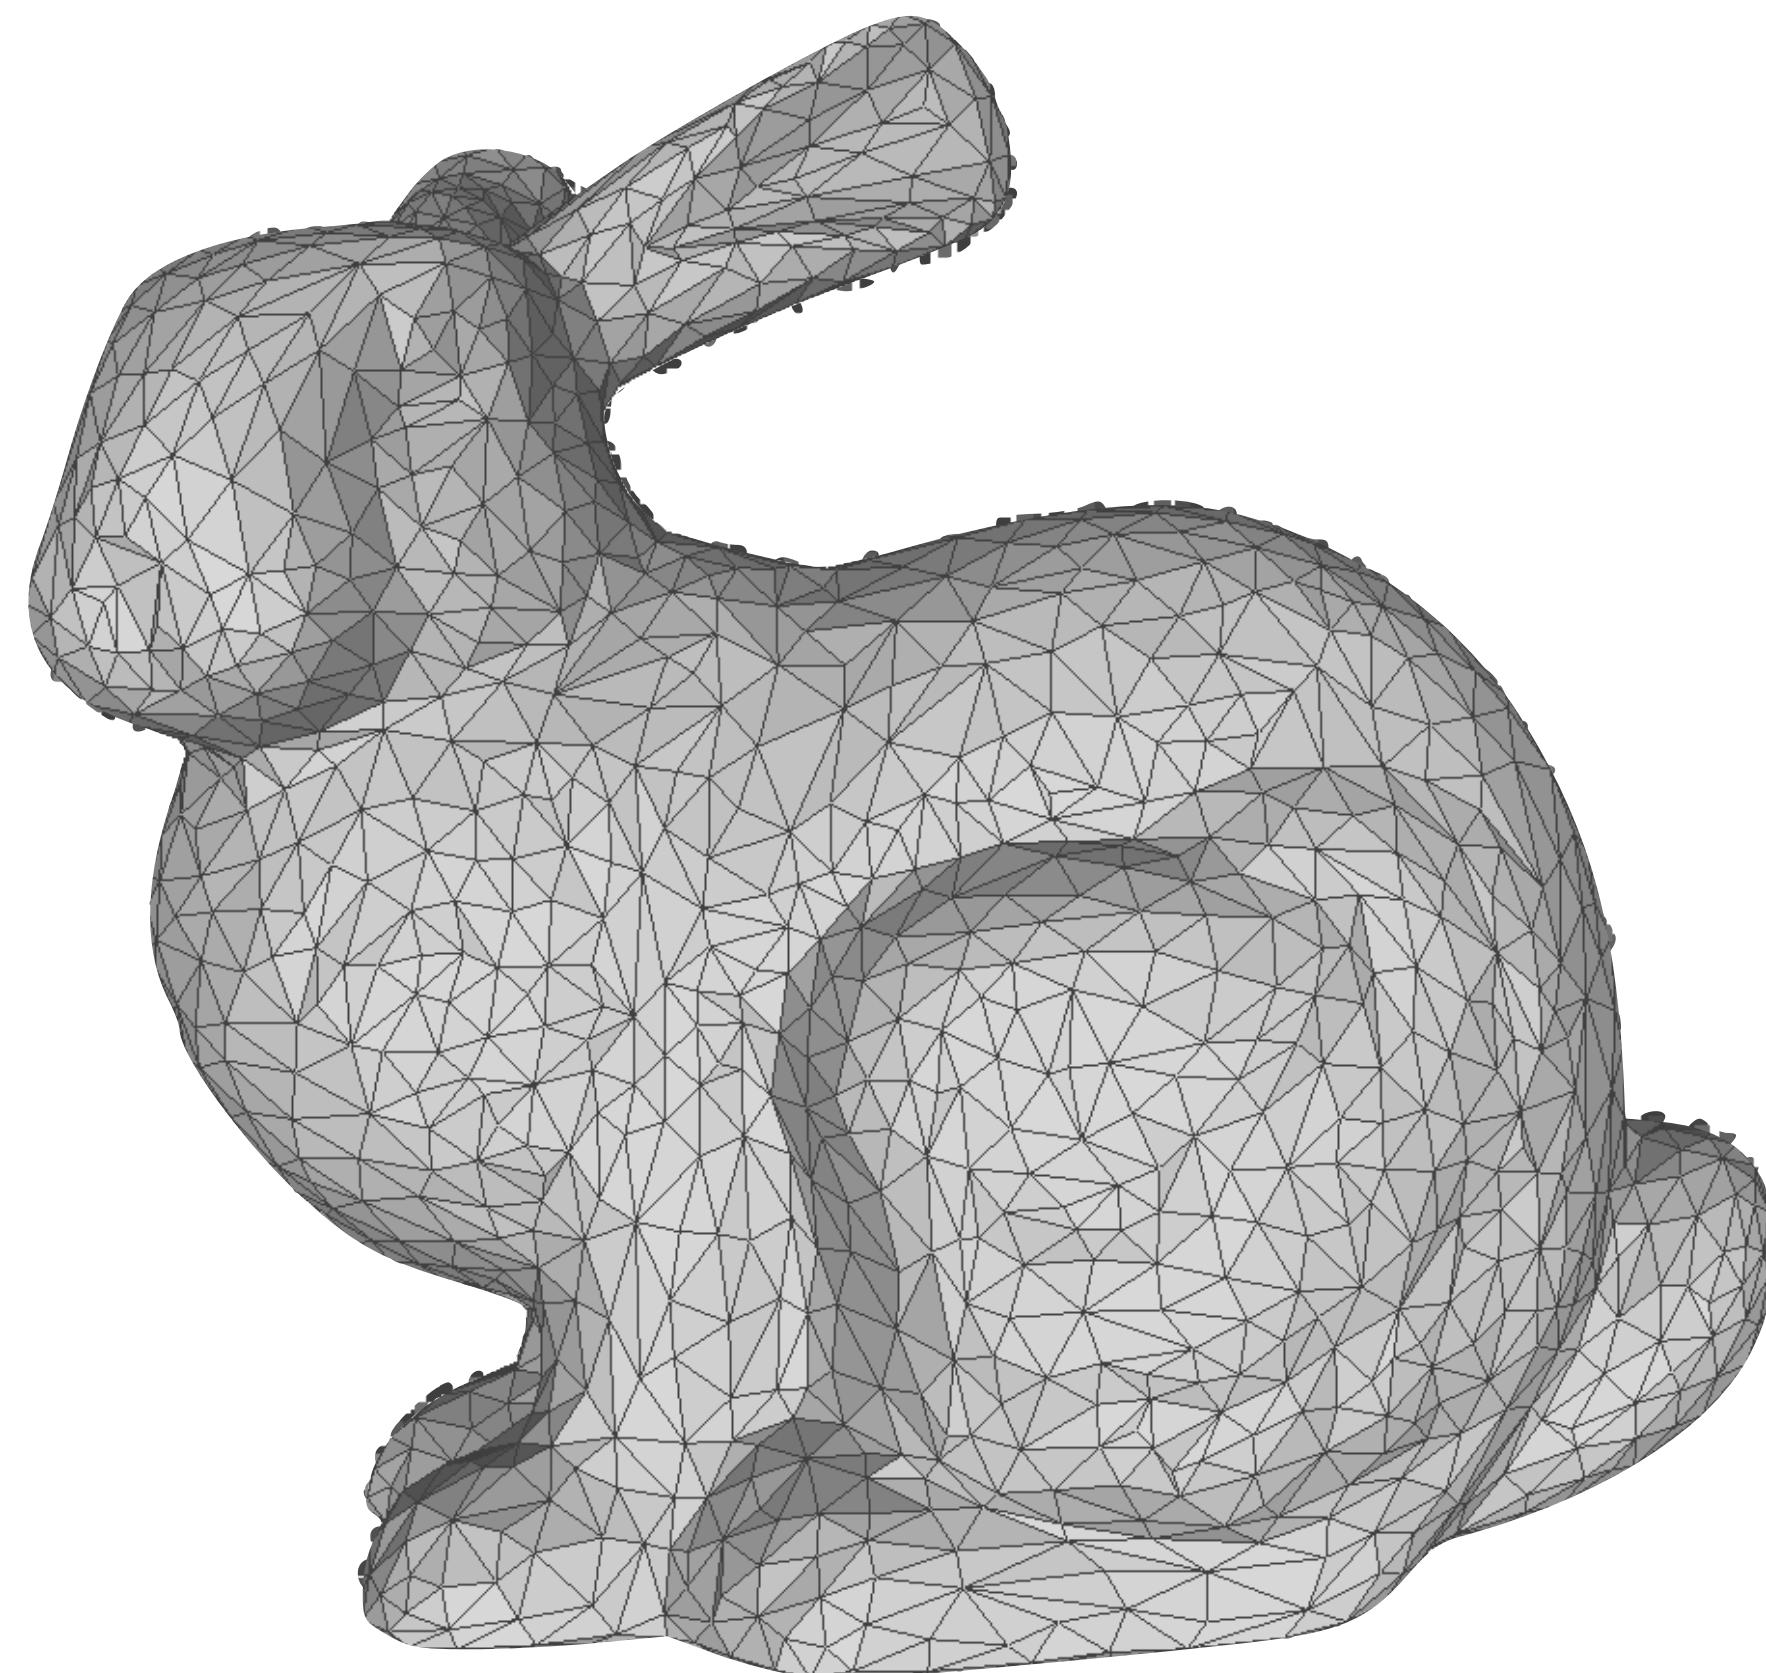
# Vertices

x1,y1,z1

# Faces

# Positions of Vertices

# Meshes



# Vertices + Faces

# Vertices

x1,y1,z1

# Positions of Vertices

# Faces

# Connectivity

(indices of vertices  
that make a ‘face’)

# Vertices + Faces

# Vertices

# Positions of Vertices

Connectivity  
(indices of vertices  
that make a ‘face’)

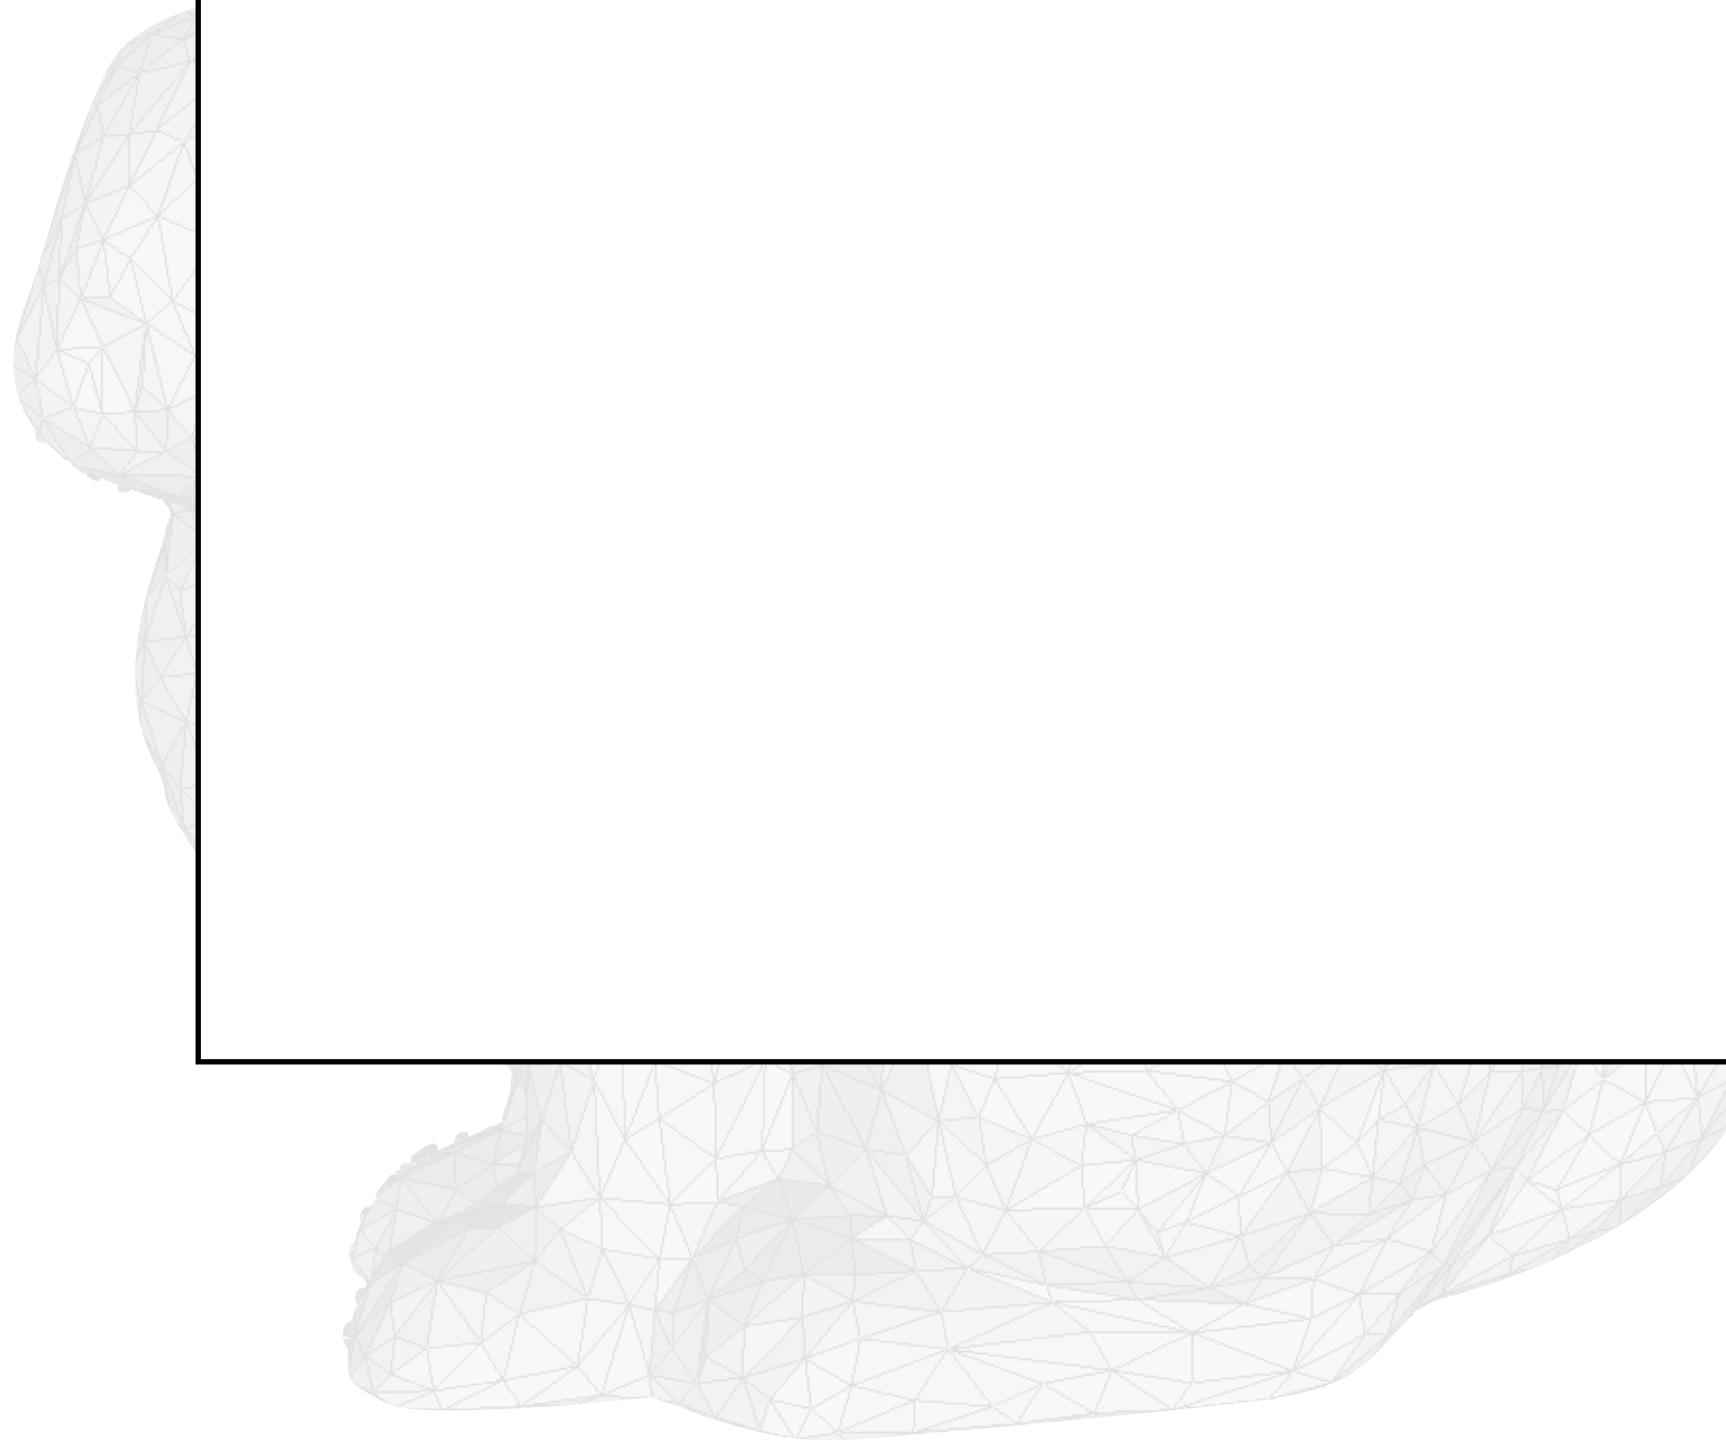
# Meshes

# Faces

## Vertices

i,j,k
.

## Ordering Matters in Face Indices



Vertices + Faces

Positions of Vertices

.
.

Connectivity  
(indices of vertices  
that make a ‘face’)

.
.
.
.

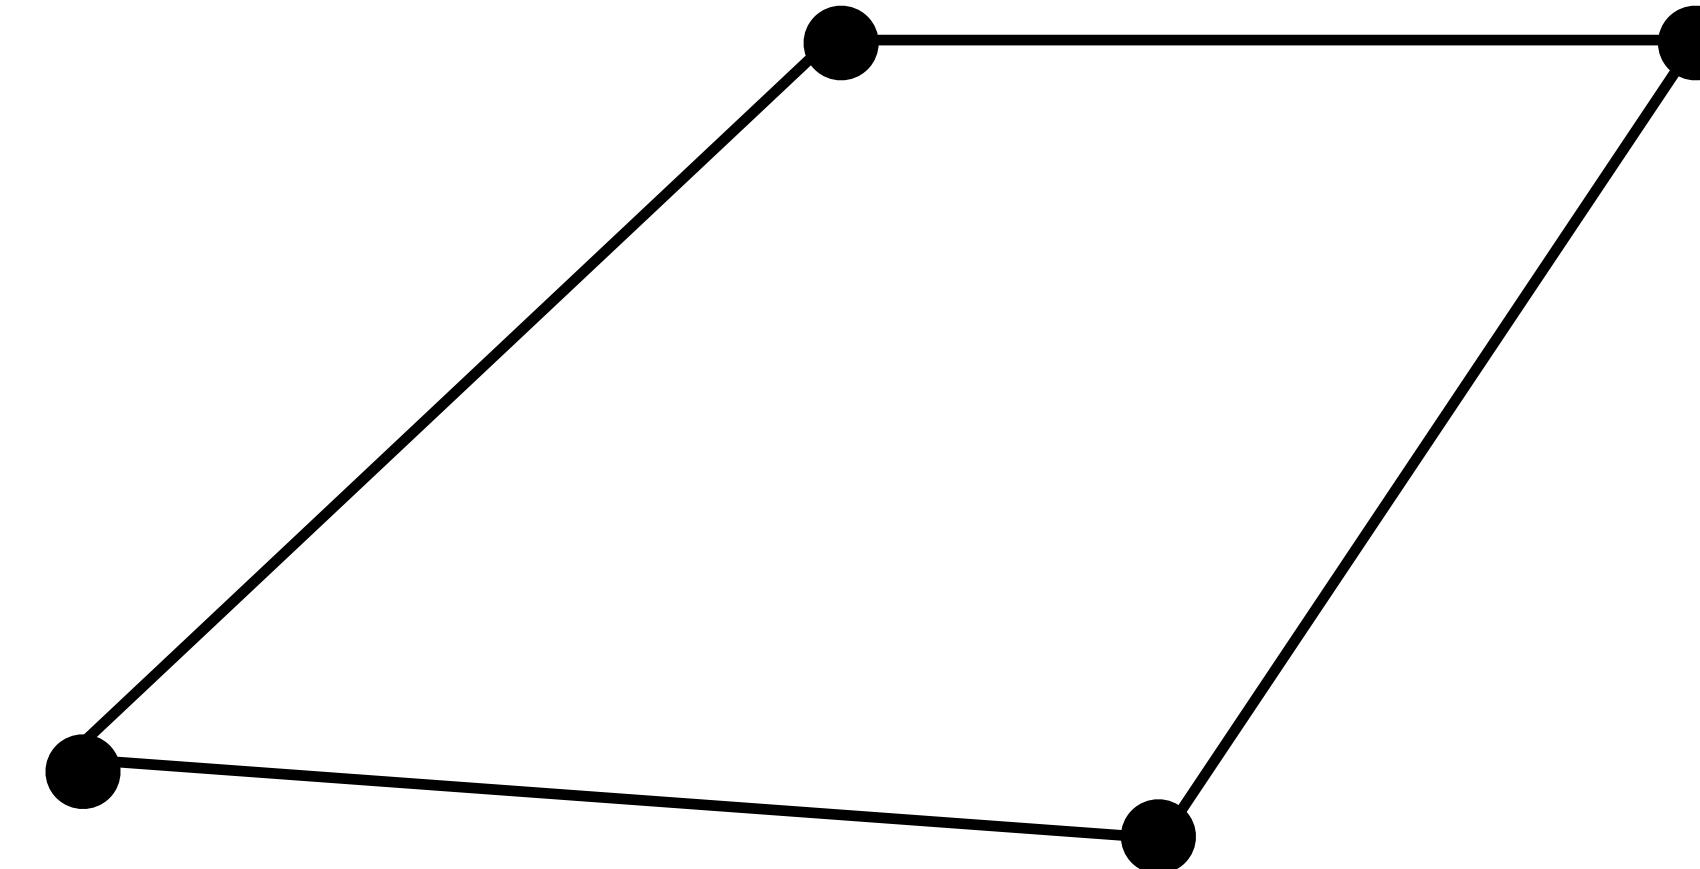
# Meshes

# Faces

## Vertices

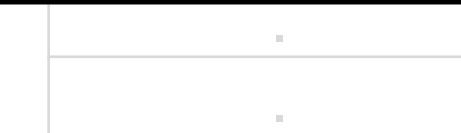
i,j,k

### Ordering Matters in Face Indices



Vertices + Faces

Positions of Vertices



Connectivity  
(indices of vertices  
that make a 'face')



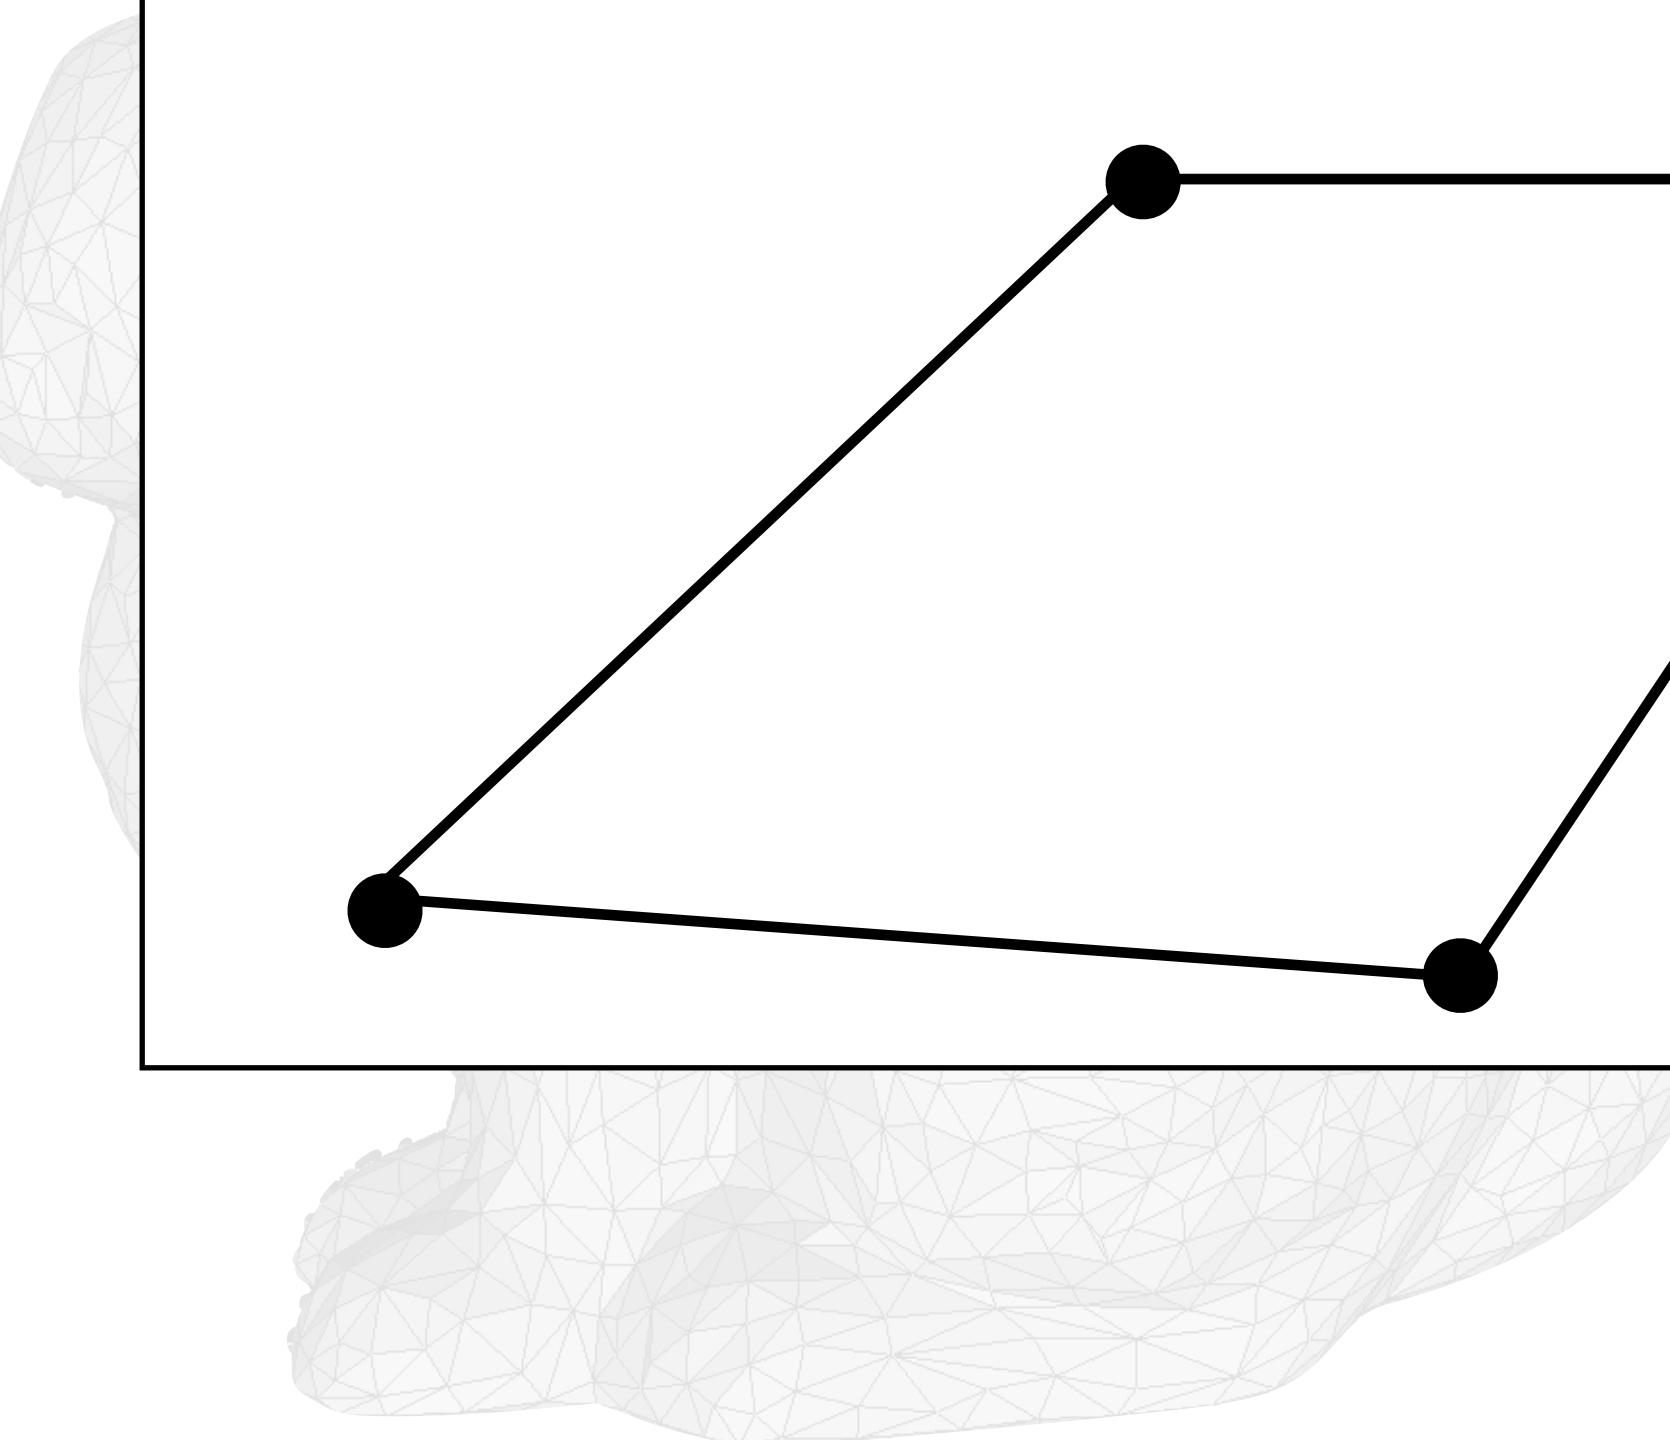
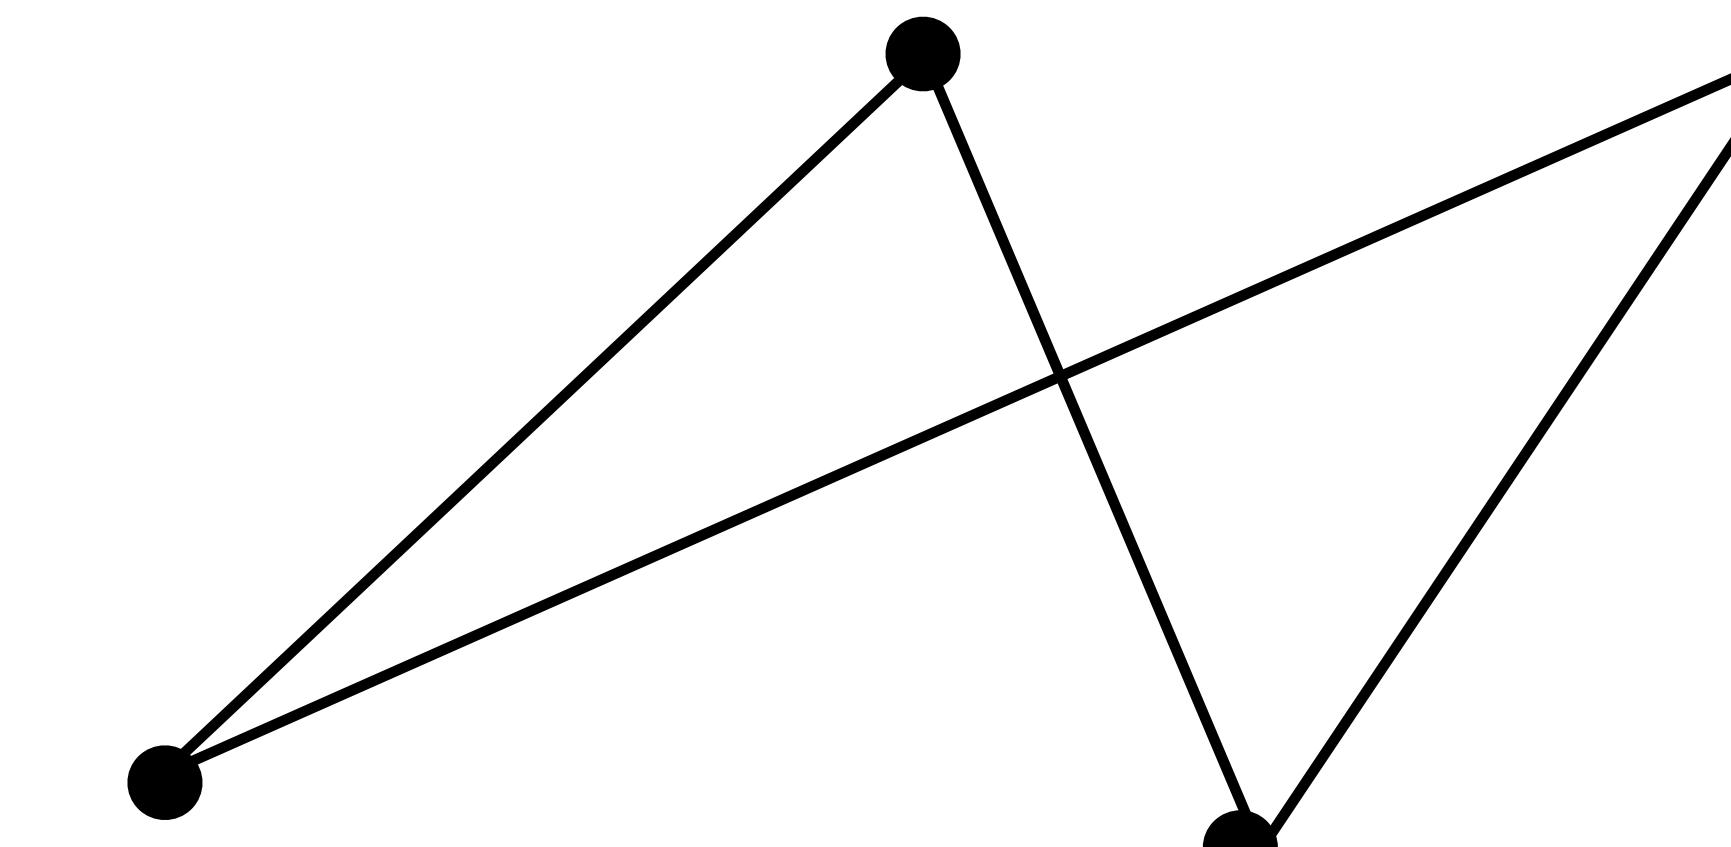
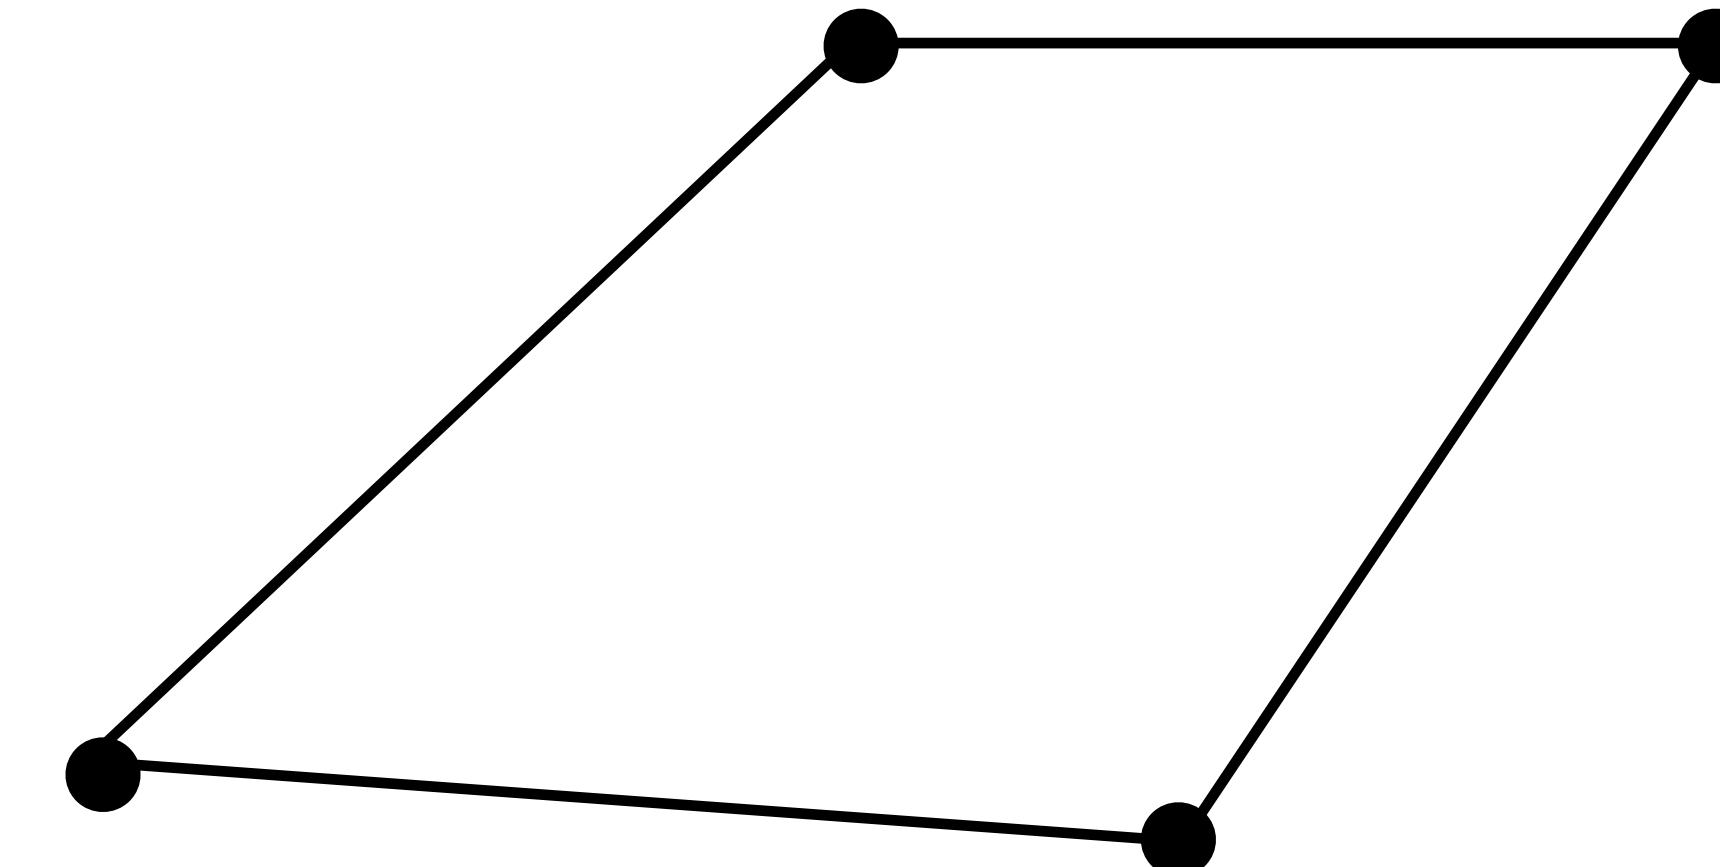
# Meshes

# Faces

## Vertices

i,j,k

### Ordering Matters in Face Indices



Vertices + Faces

Positions of Vertices

Connectivity  
(indices of vertices  
that make a ‘face’)

# Meshes

Faces

Vertices

i,j,k

Arbitrary Polygons can be non-planar



Vertices + Faces

Positions of Vertices

Connectivity  
(indices of vertices  
that make a ‘face’)

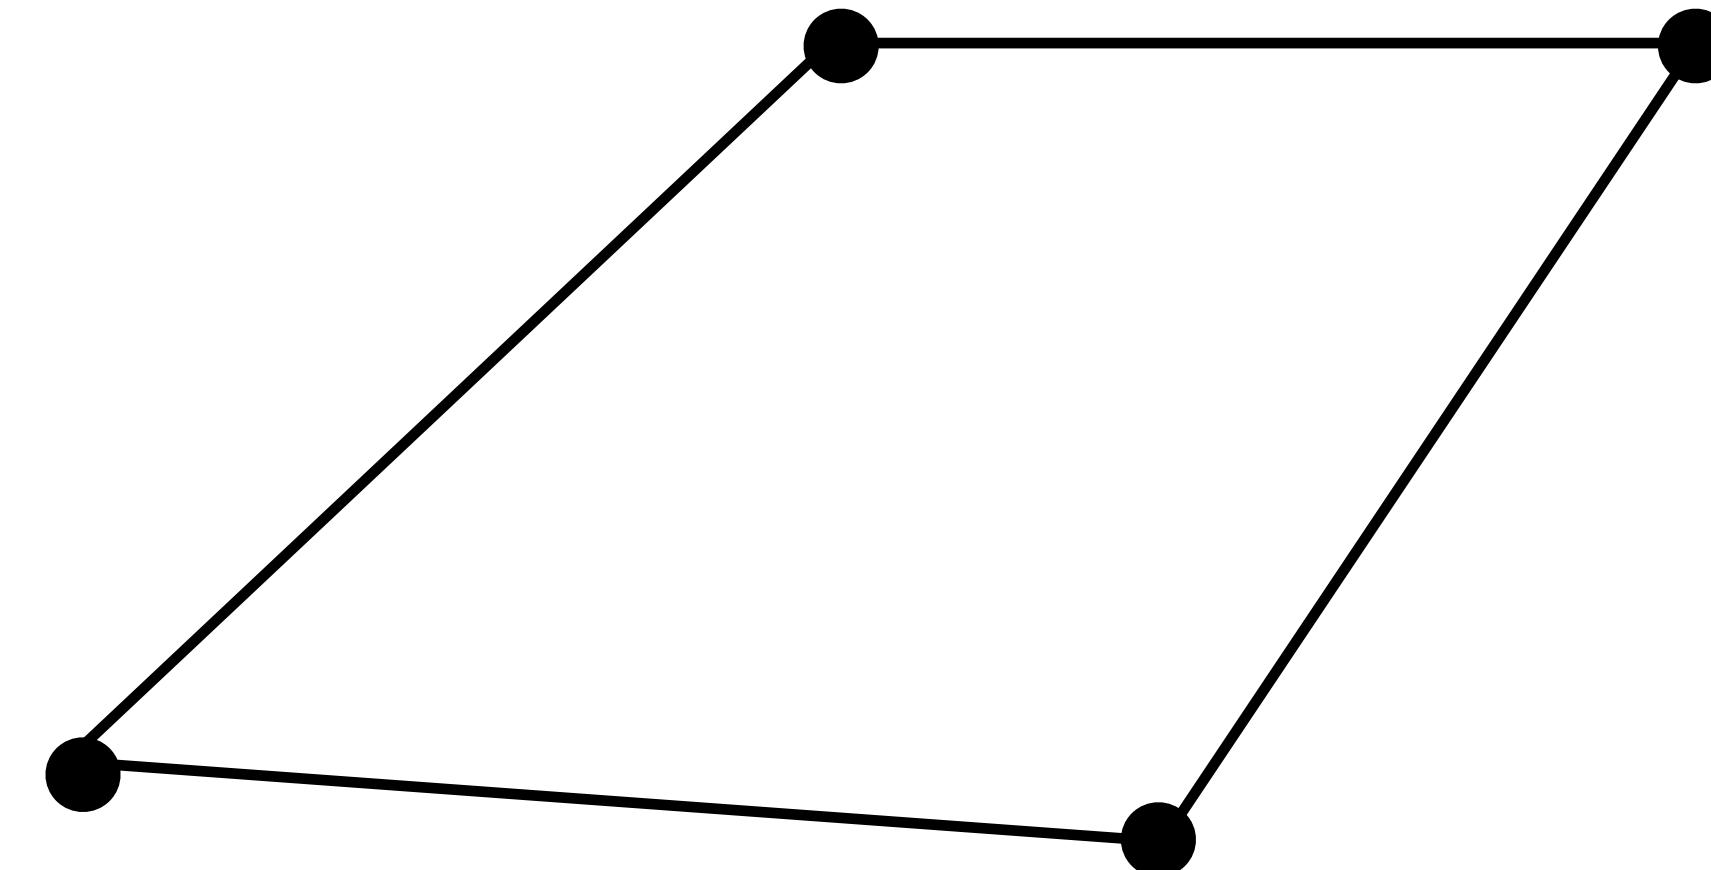
# Meshes

Faces

Vertices

i,j,k

Arbitrary Polygons can be non-planar



Vertices + Faces

Positions of Vertices

Connectivity  
(indices of vertices  
that make a ‘face’)

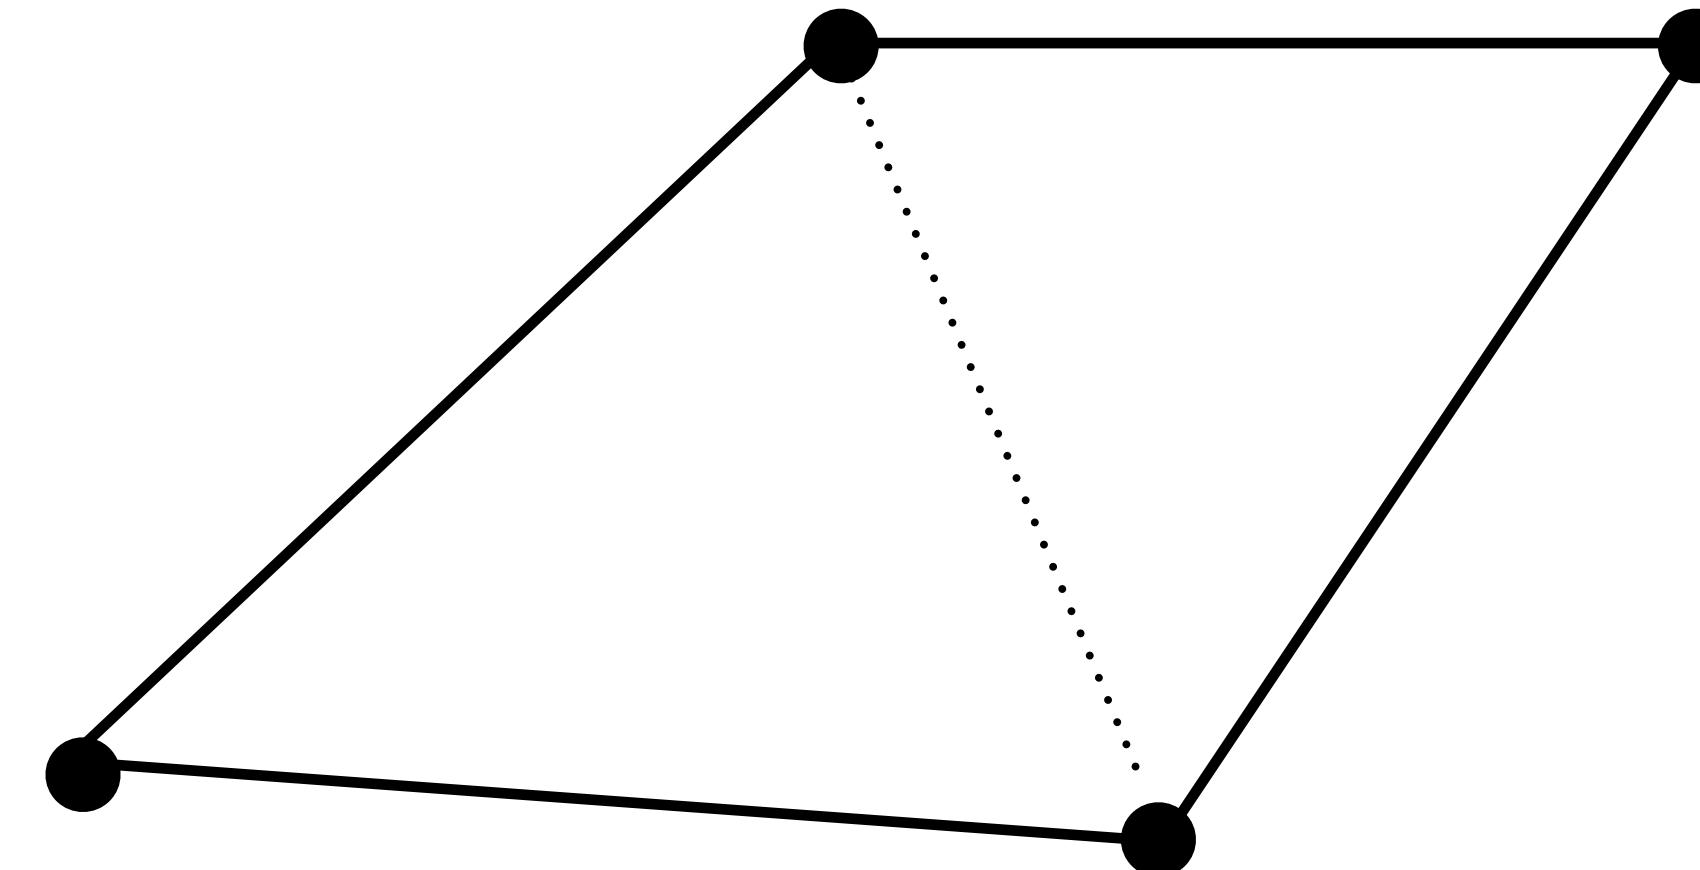
# Meshes

Faces

Vertices

i,j,k

Arbitrary Polygons can be non-planar



Vertices + Faces

Positions of Vertices

Connectivity  
(indices of vertices  
that make a ‘face’)

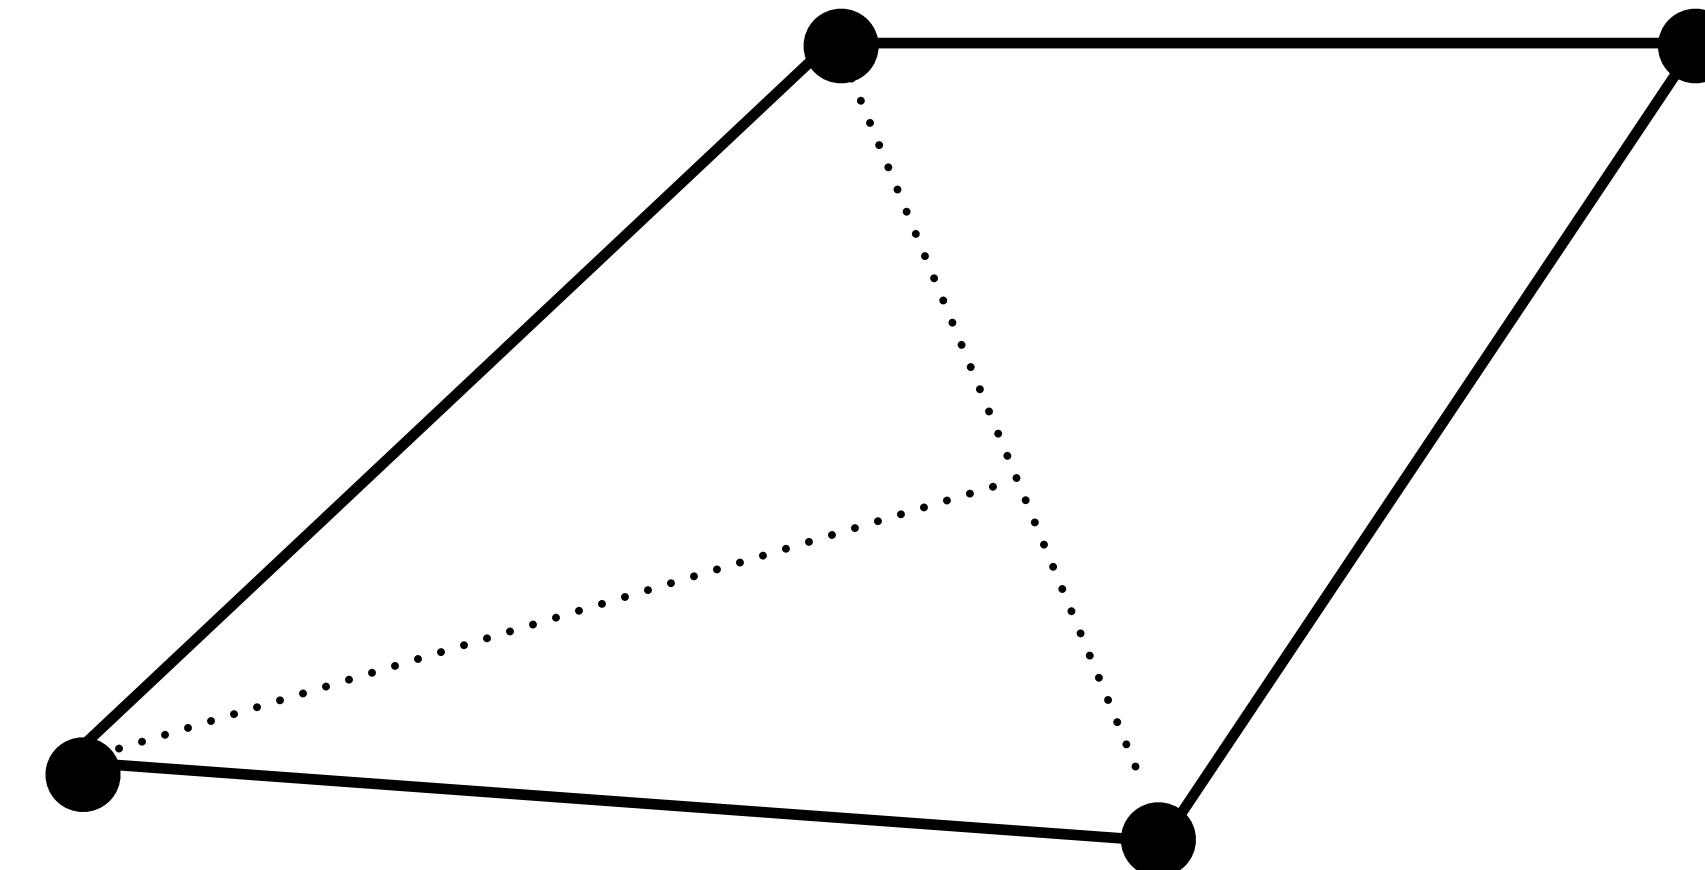
# Meshes

Faces

Vertices

i,j,k

Arbitrary Polygons can be non-planar



Vertices + Faces

Positions of Vertices

Connectivity  
(indices of vertices  
that make a ‘face’)

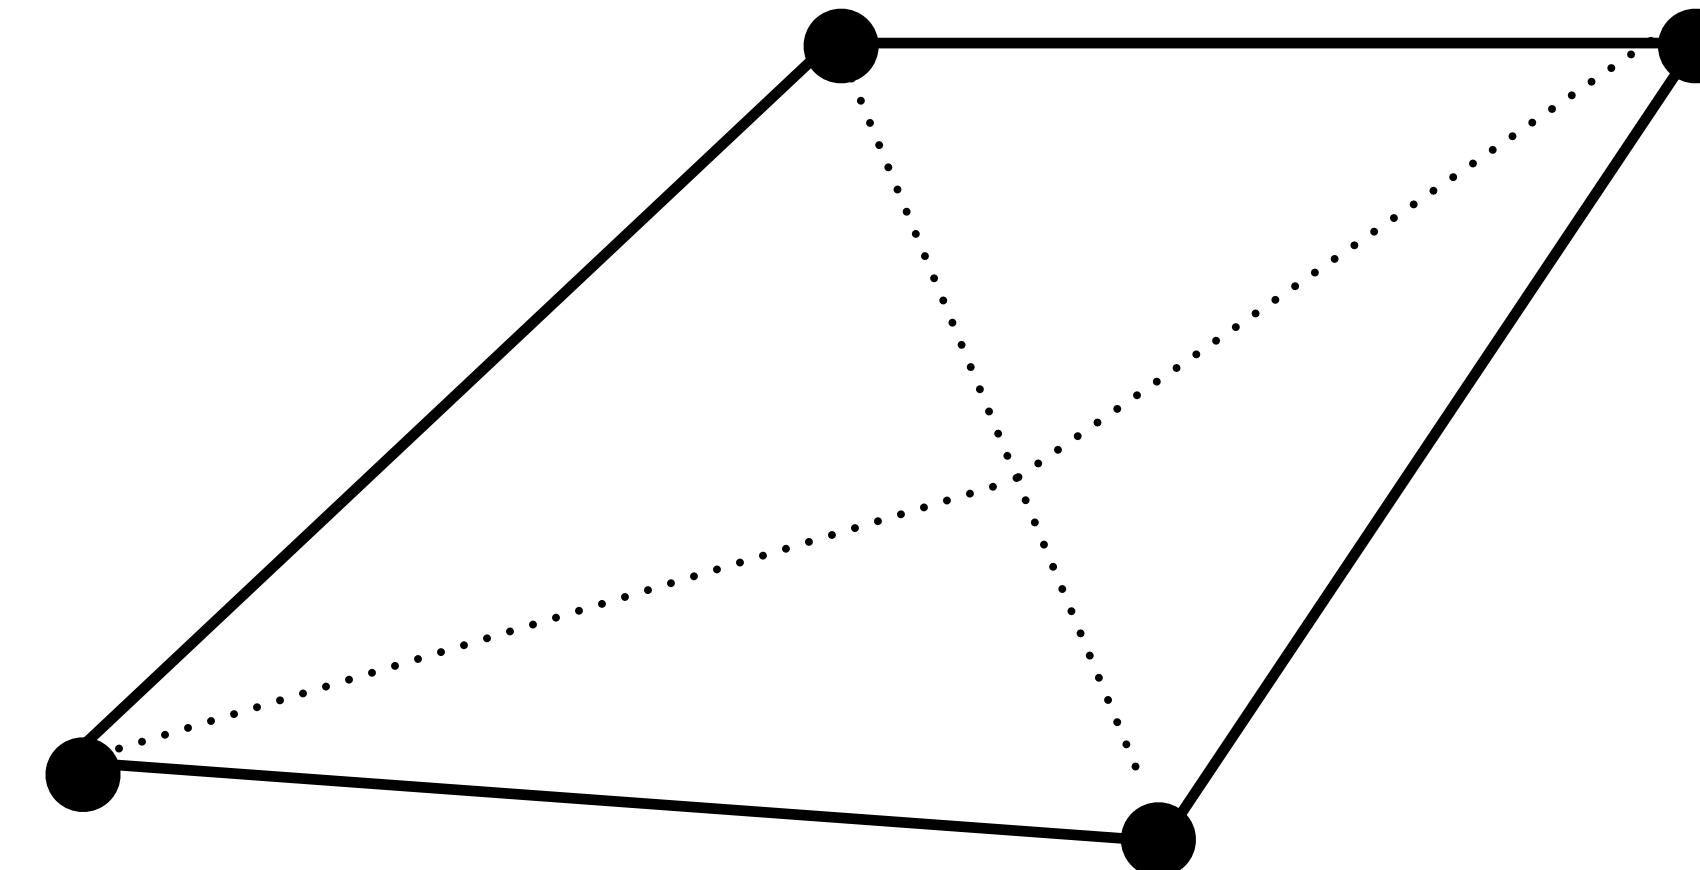
# Meshes

Faces

Vertices

i,j,k

Arbitrary Polygons can be non-planar



Vertices + Faces

Positions of Vertices

Connectivity  
(indices of vertices  
that make a ‘face’)

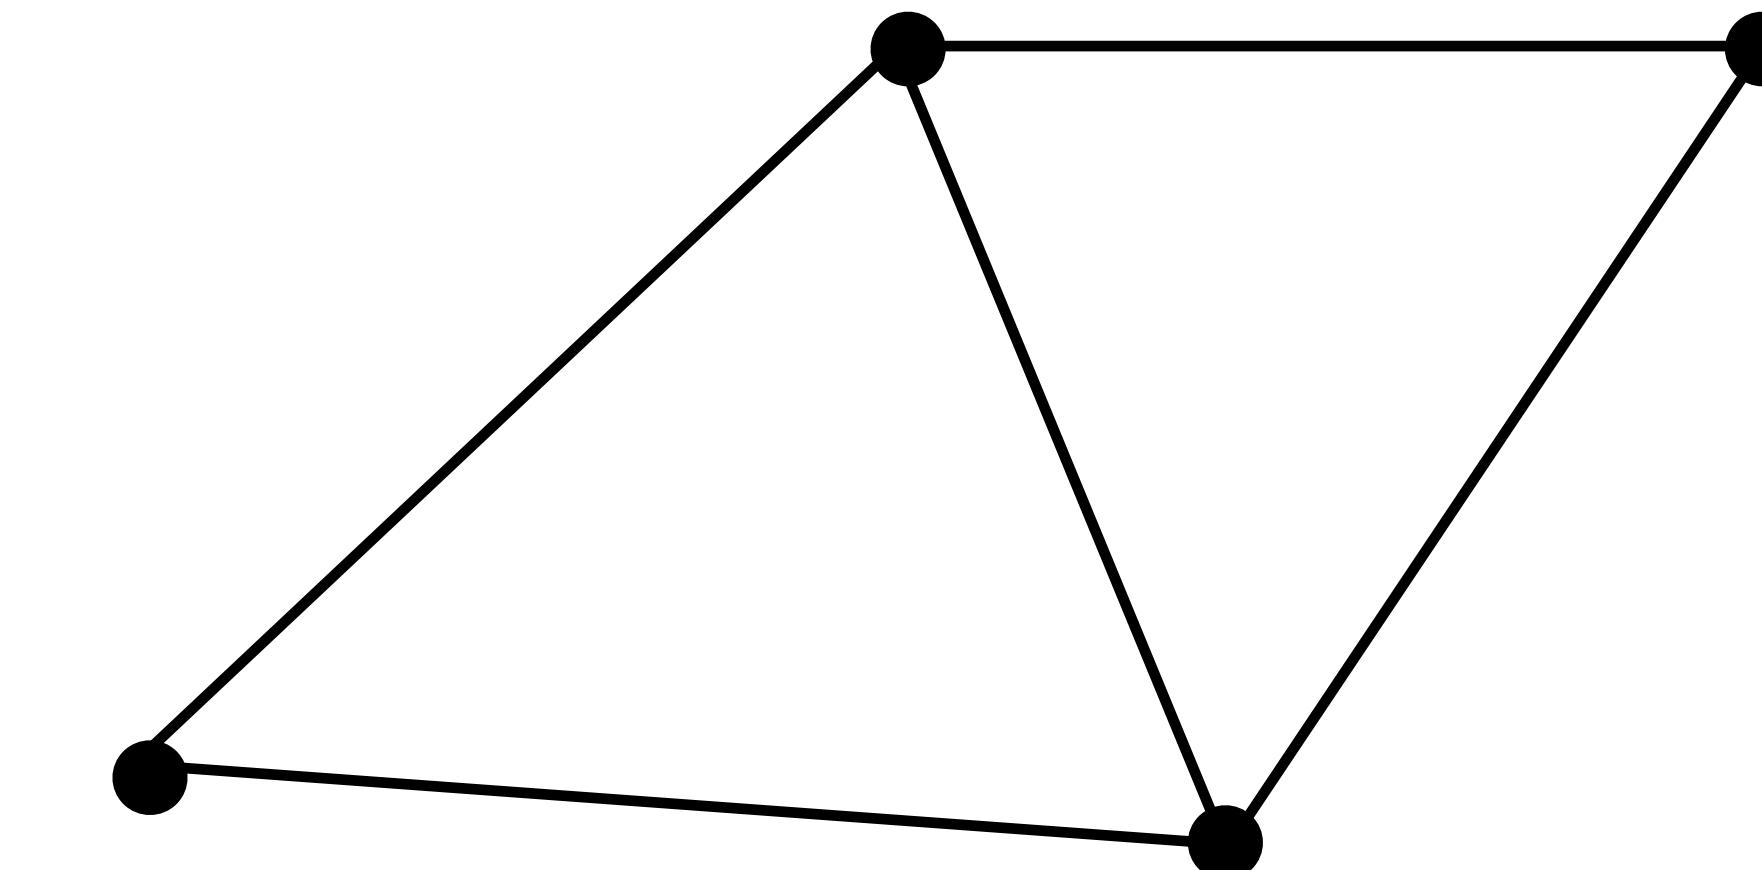
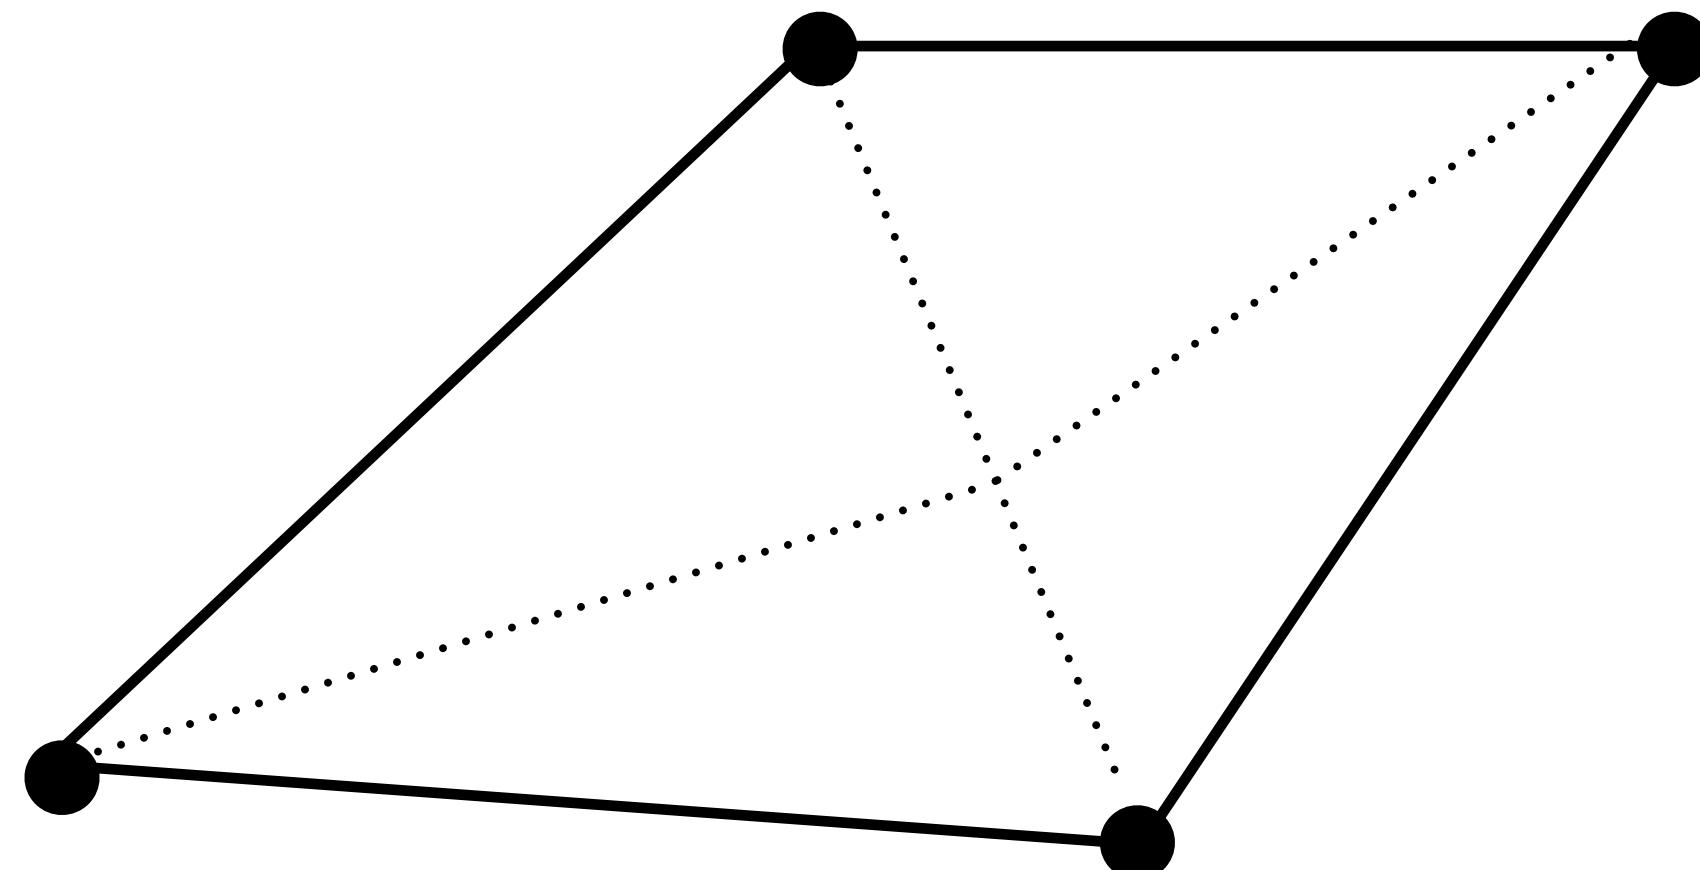
# Meshes

## Faces

### Vertices

i,j,k

Arbitrary Polygons can be non-planar



restrict to *Triangular*  
Meshes

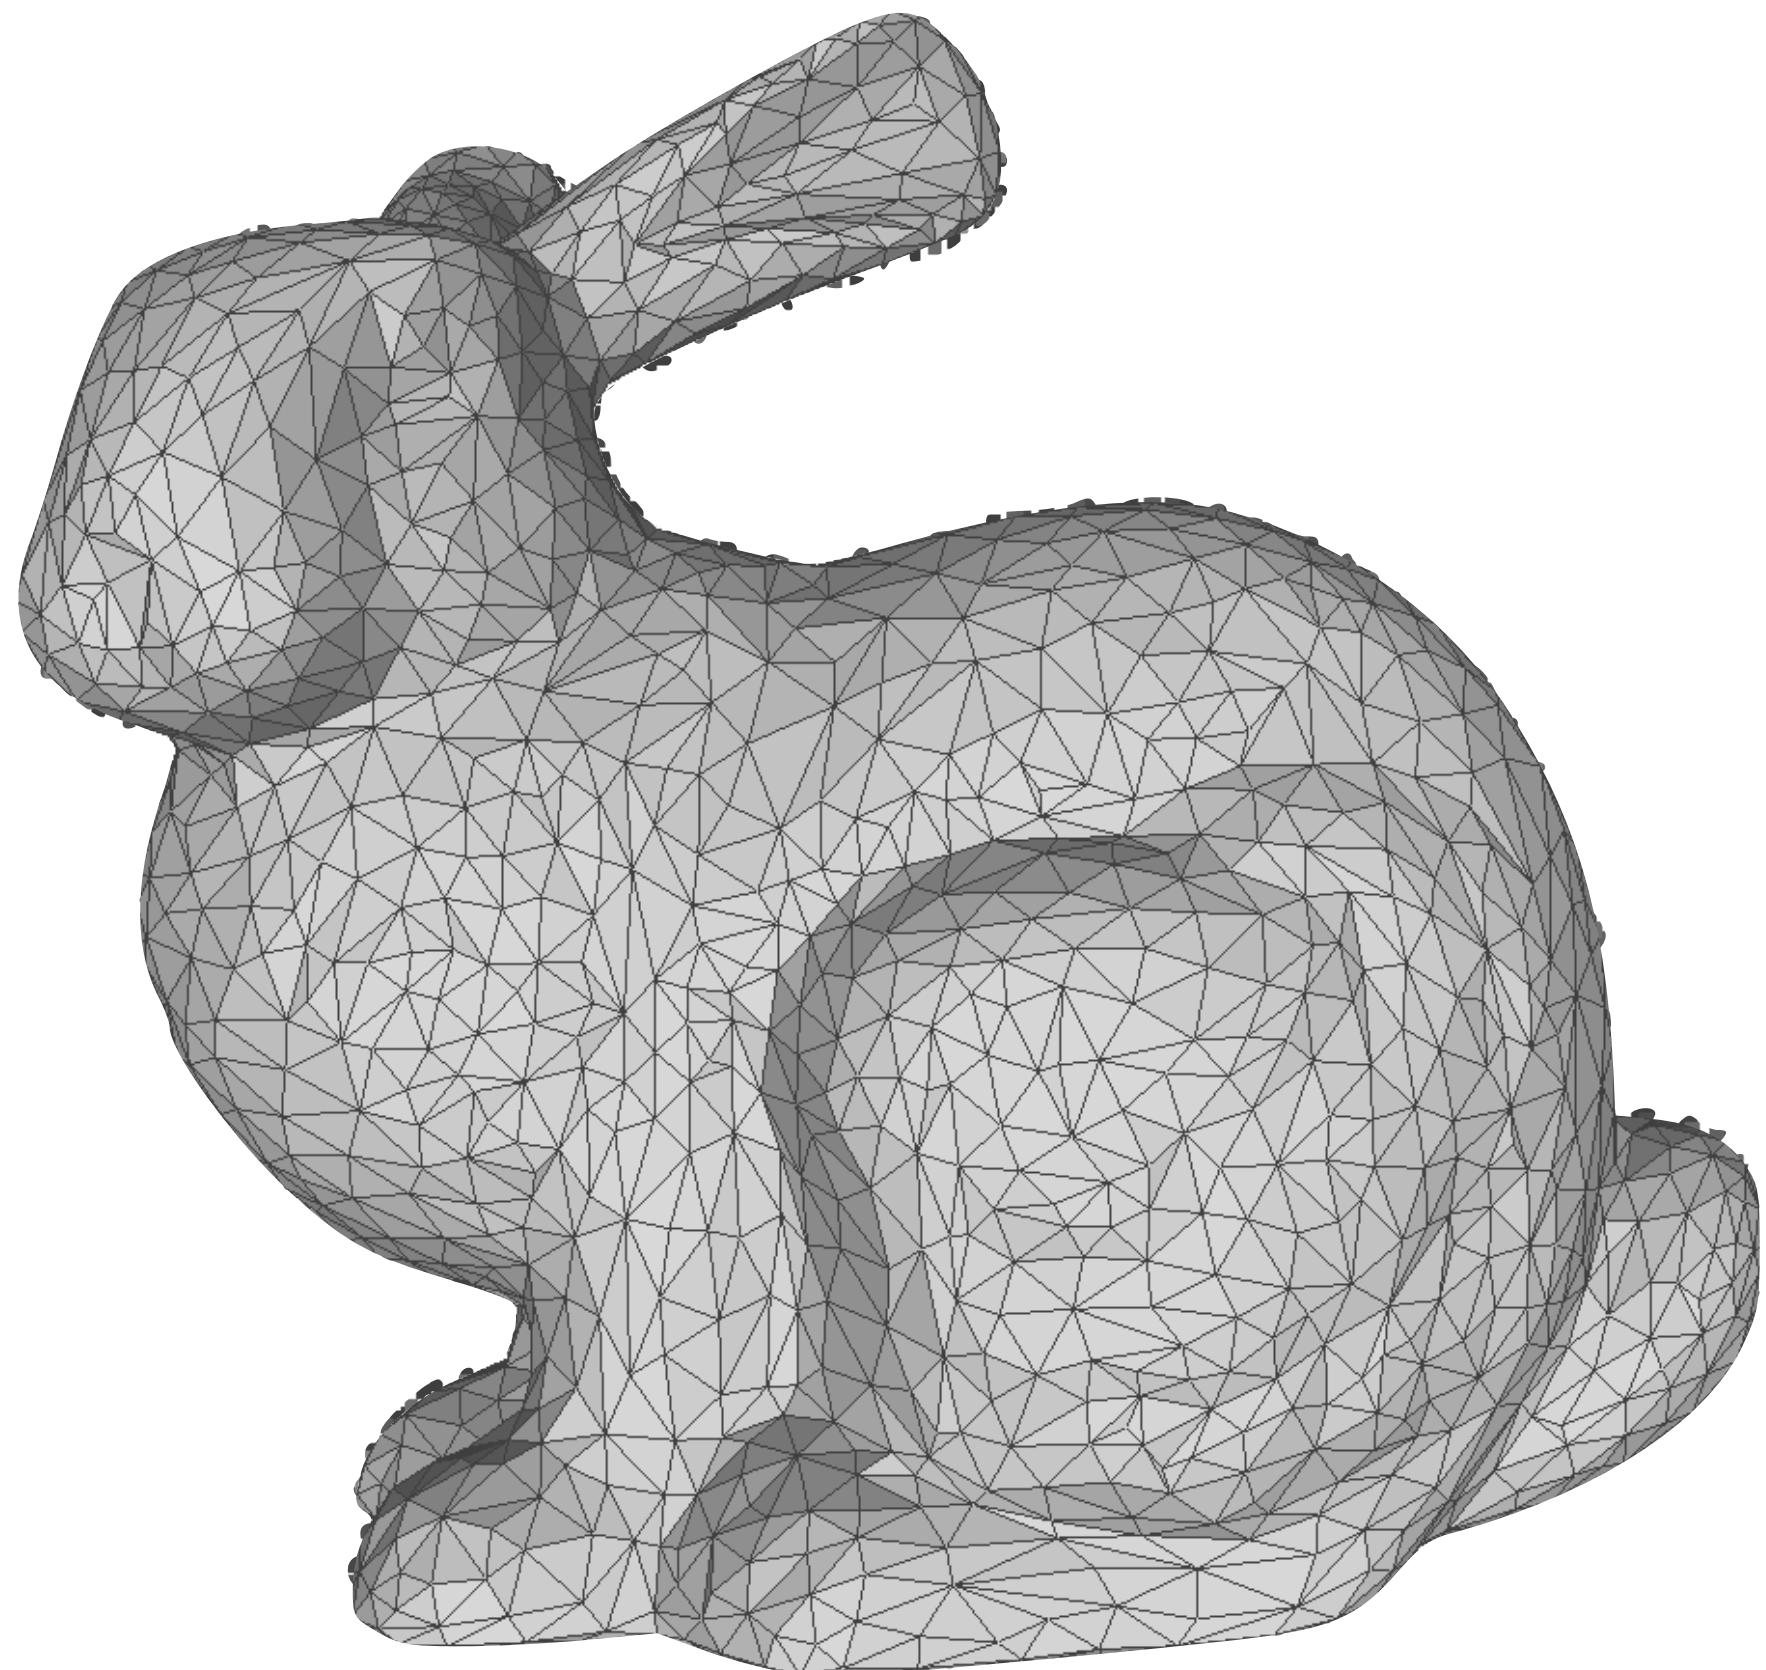
Vertices + Faces

Positions of Vertices

Connectivity  
(indices of vertices  
that make a ‘face’)

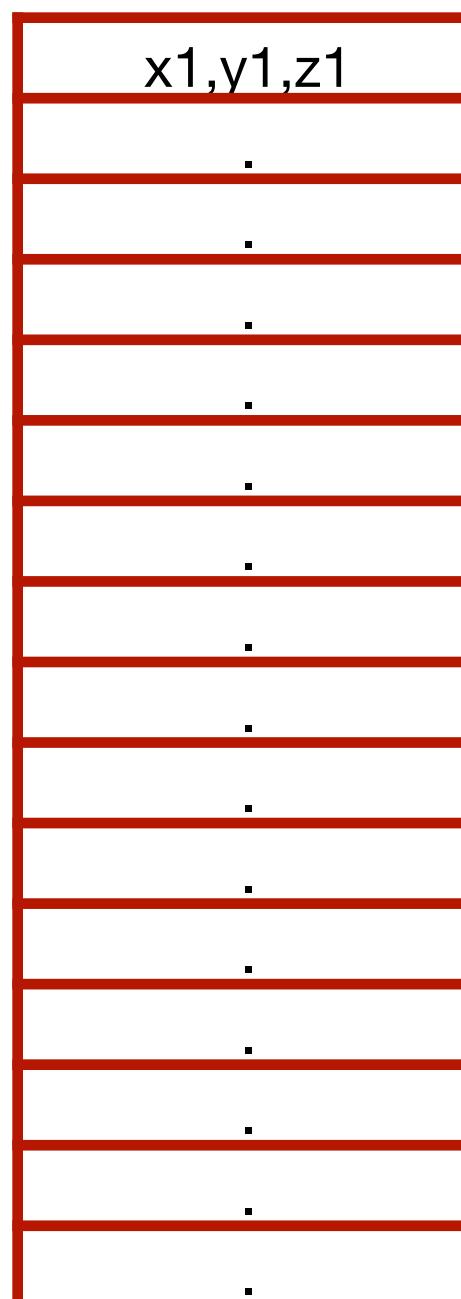
# Triangular Meshes

# Faces



# Vertices + Faces

# Vertices

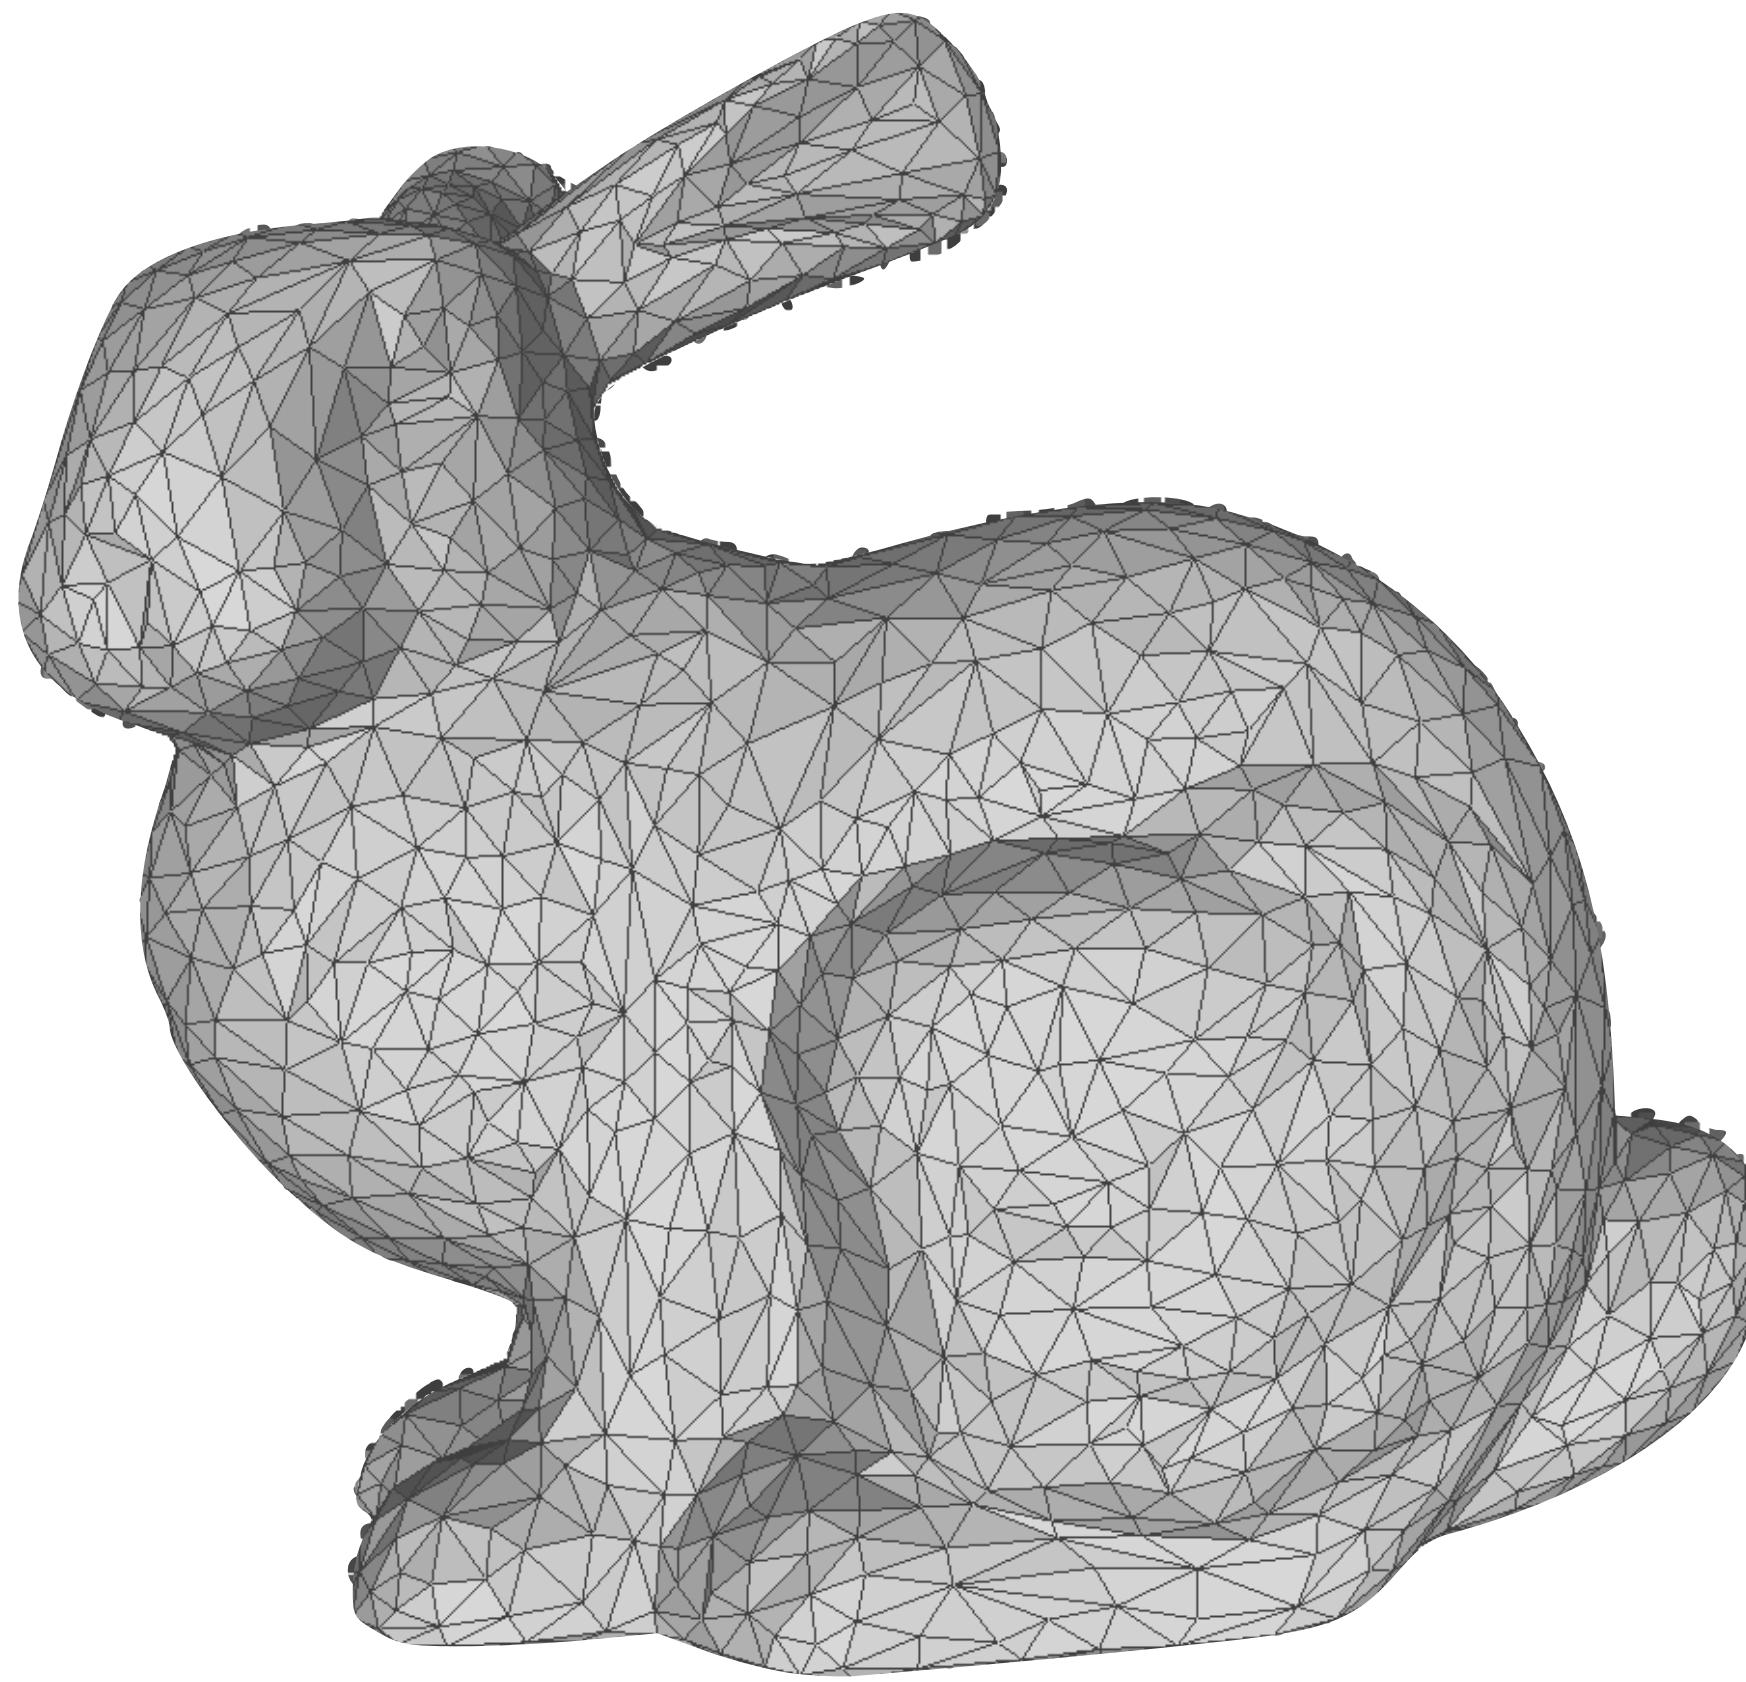


# Positions of Vertices



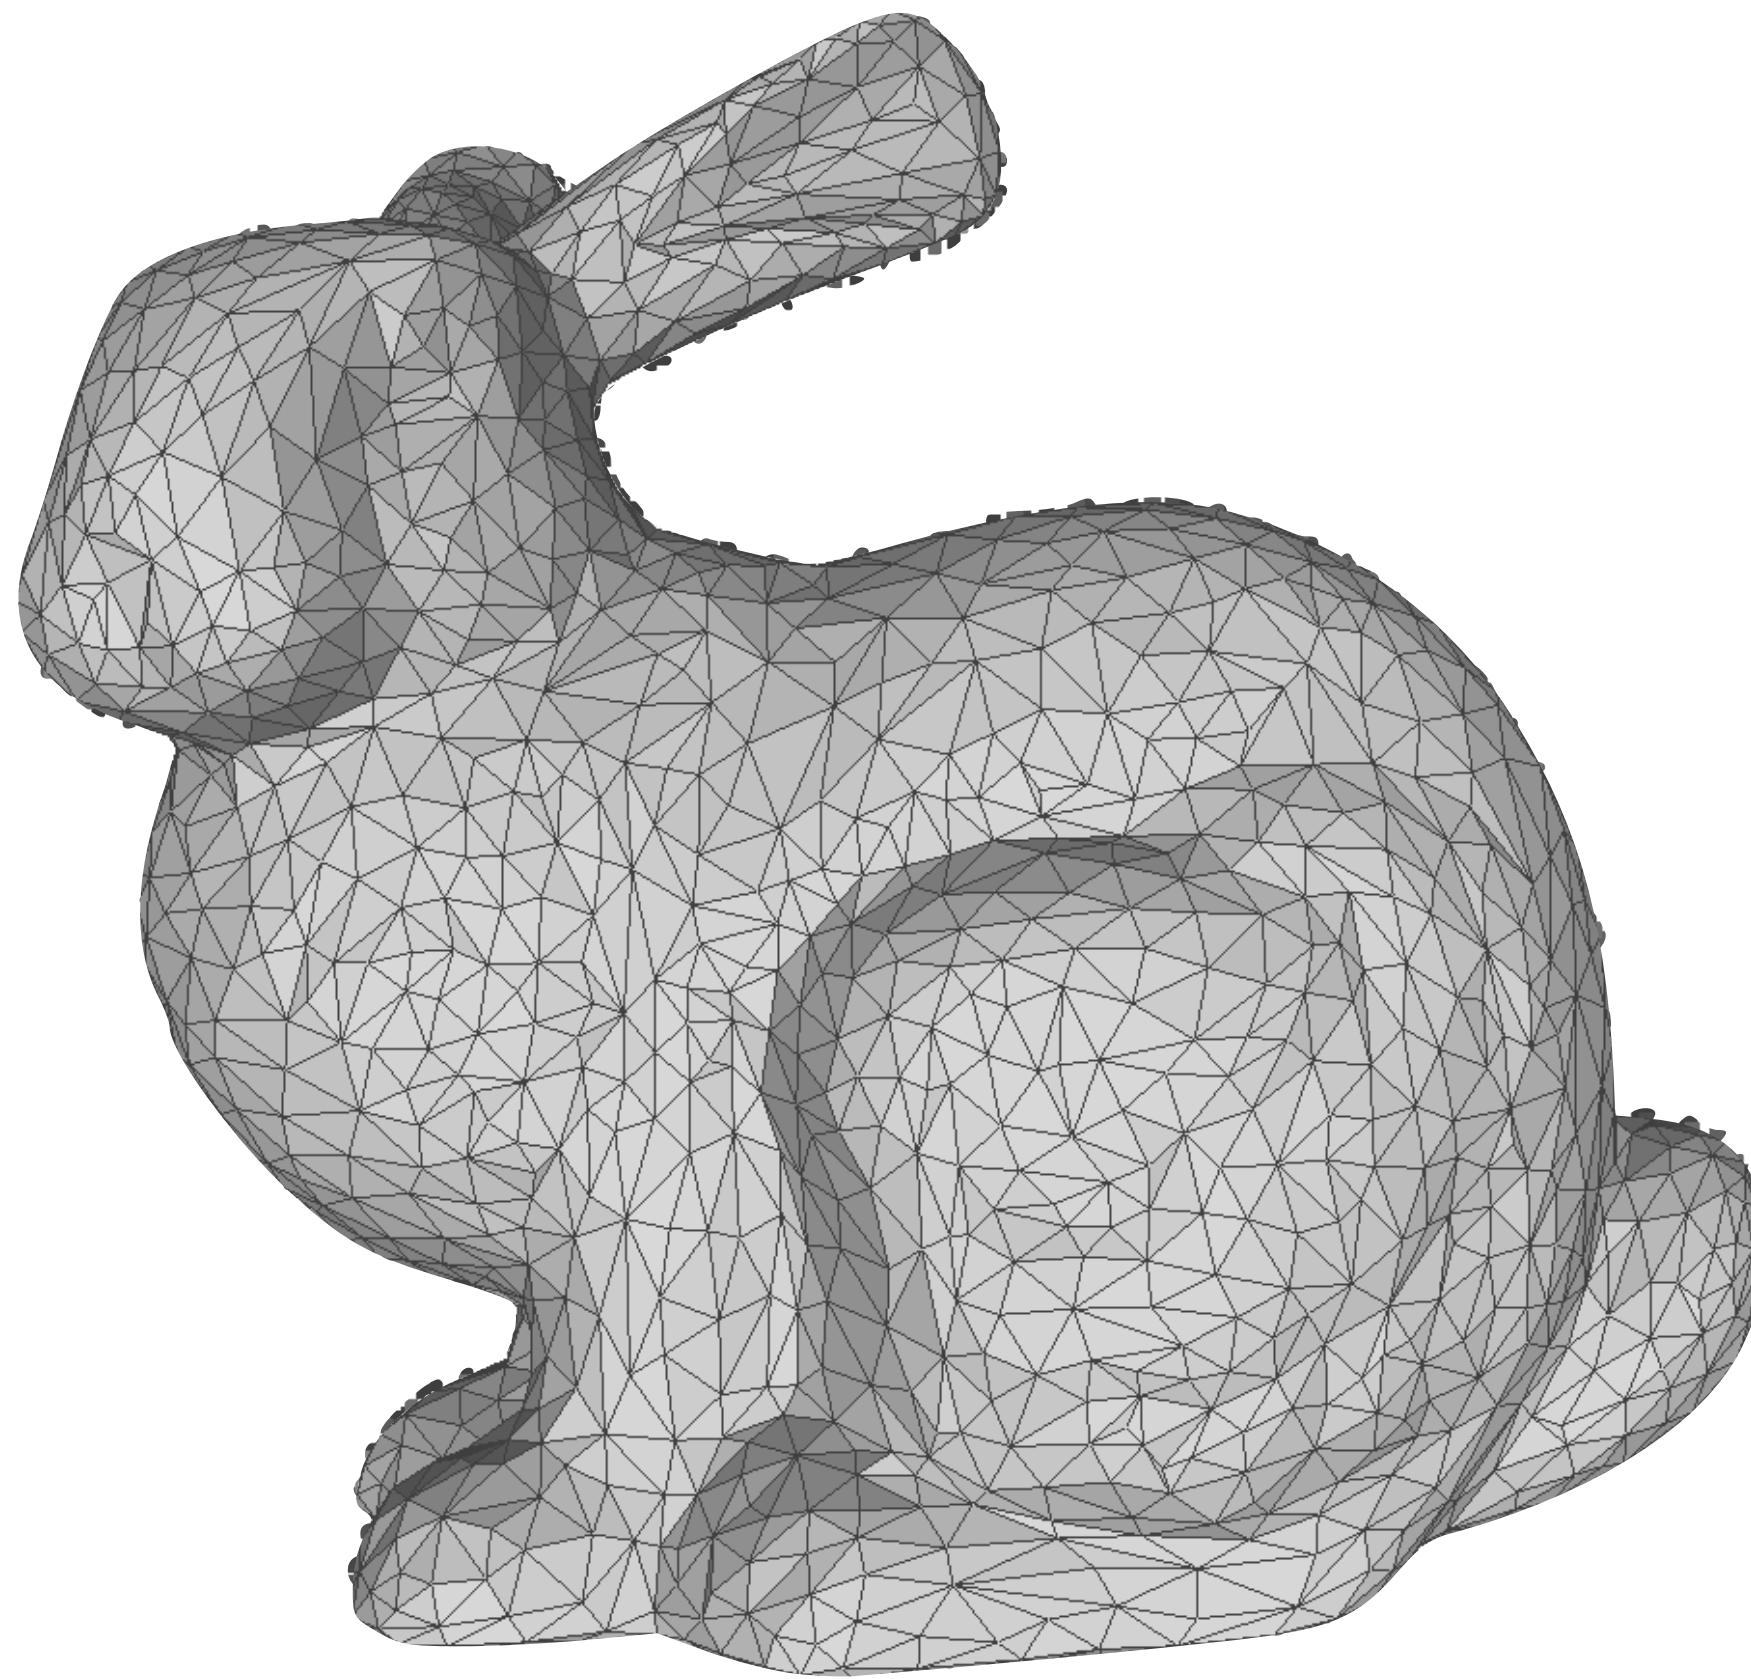
Connectivity  
(indices of **three**  
vertices that make a  
'face')

# Is a point inside or outside the shape?



Vertices + Faces

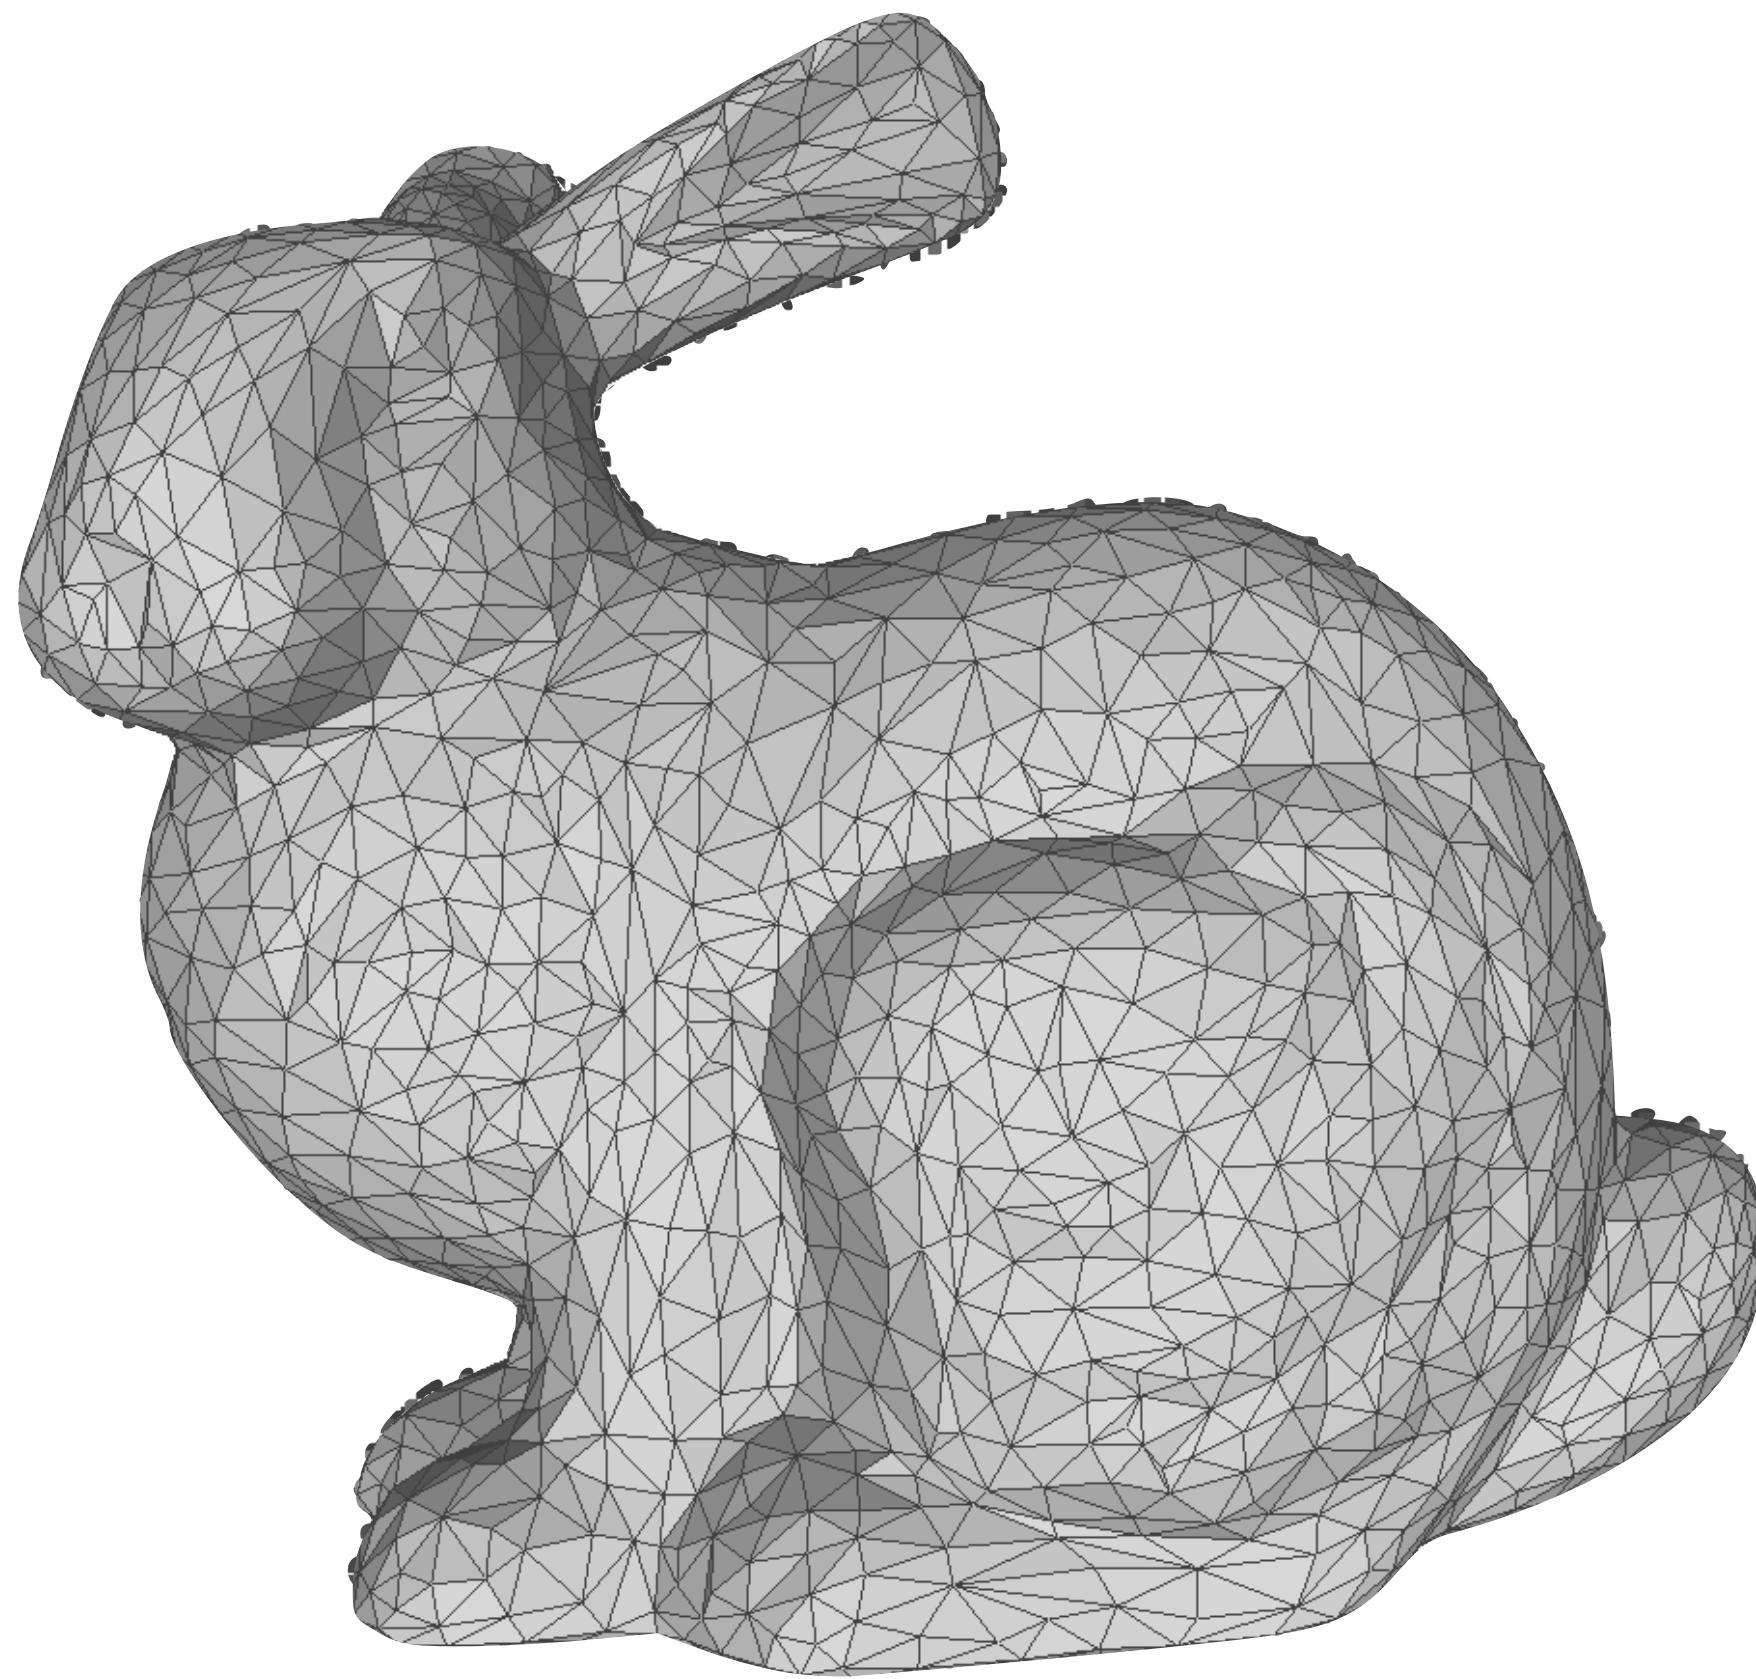
# Is a point inside or outside the shape?



Meshes can be “watertight”, i.e., if you filled the mesh with water, nothing would leak out

Vertices + Faces

# Is a point inside or outside the shape?

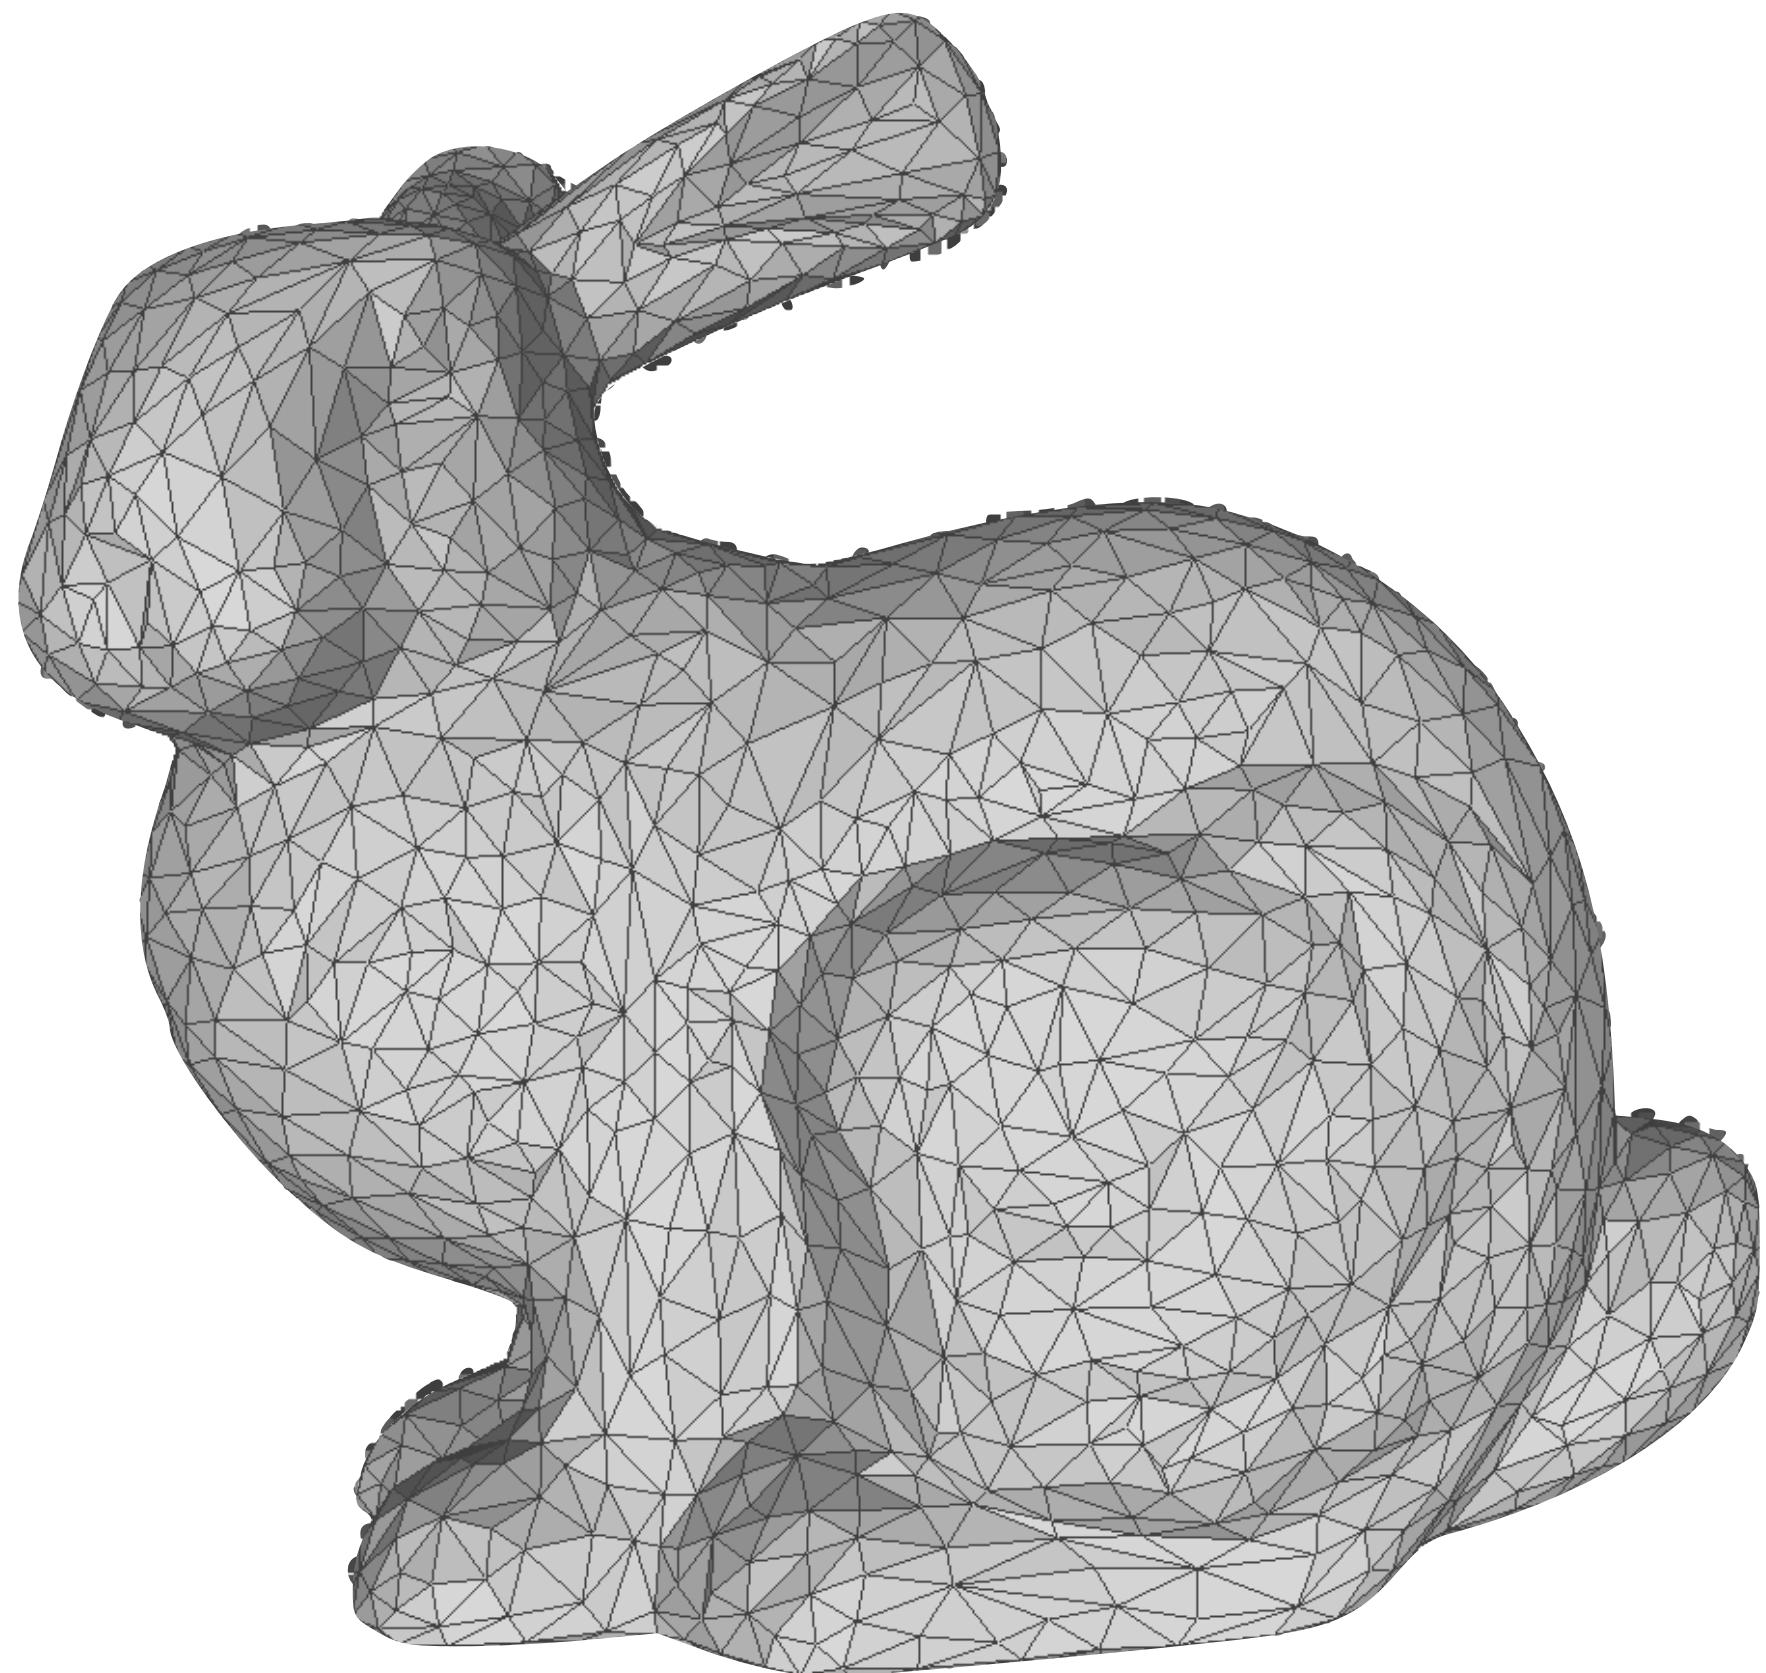


Vertices + Faces

Meshes can be “watertight”, i.e., if you filled the mesh with water, nothing would leak out

If mesh is watertight, can define “inside” and “outside”

# Is a point inside or outside the shape?



Vertices + Faces

Meshes can be “watertight”, i.e., if you filled the mesh with water, nothing would leak out

If mesh is watertight, can define “inside” and “outside”

If not, needs further processing, see “Generalized Winding Numbers”

# Triangular Meshes

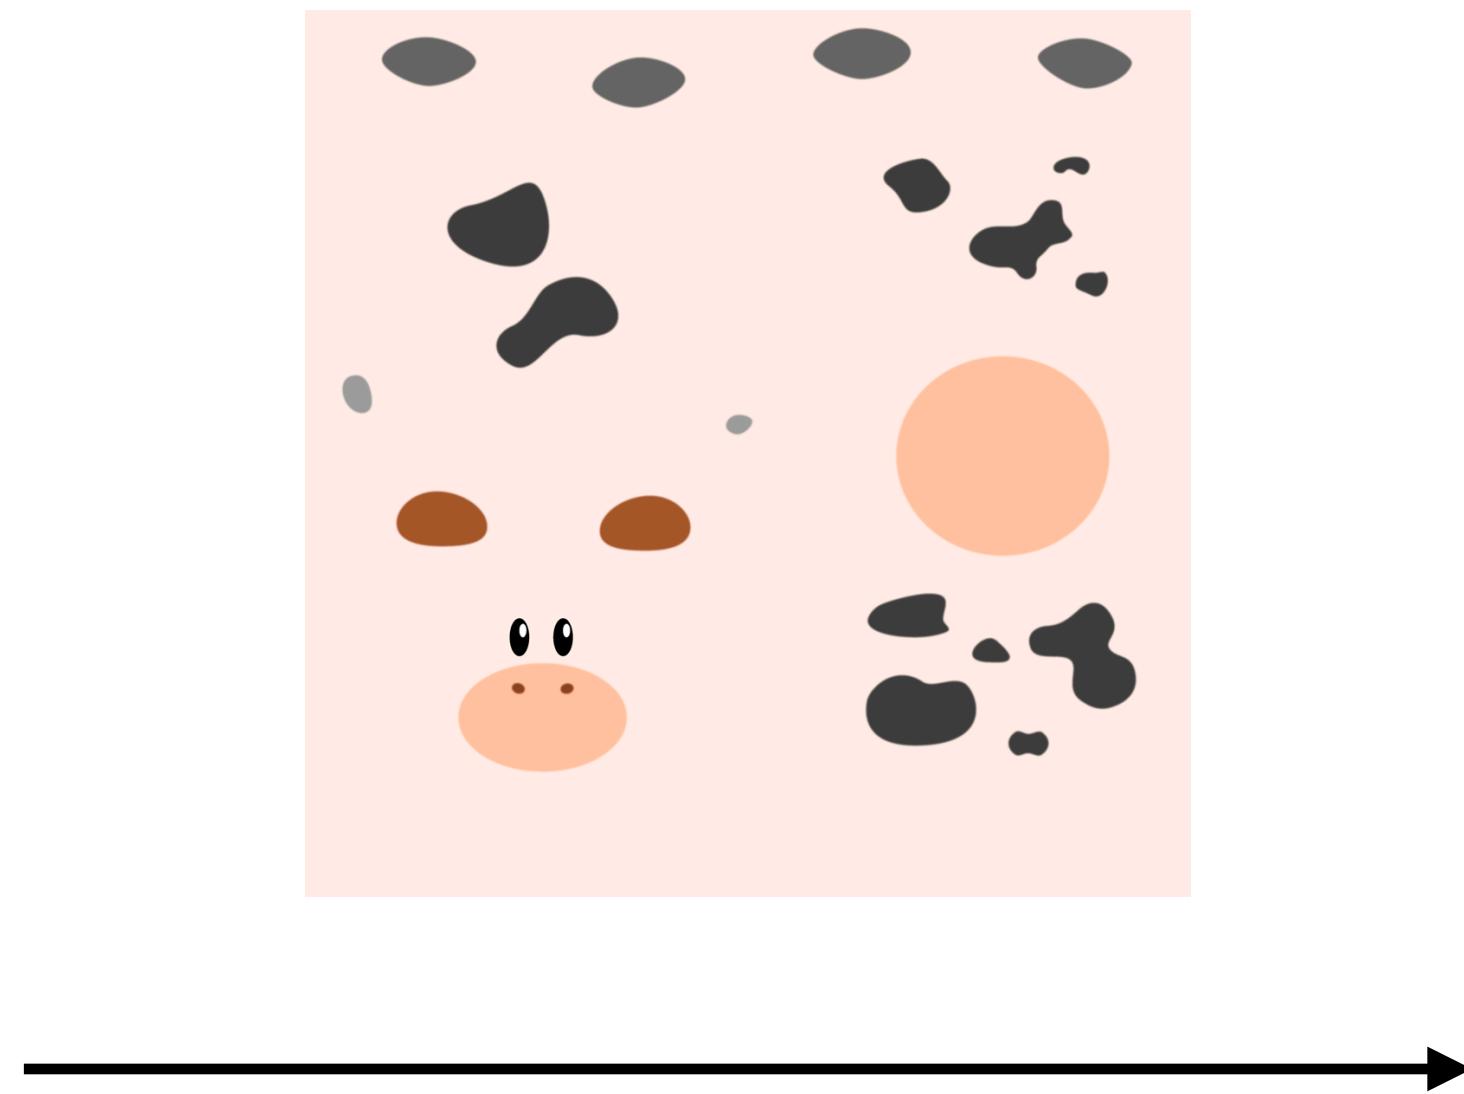


# Triangular Meshes



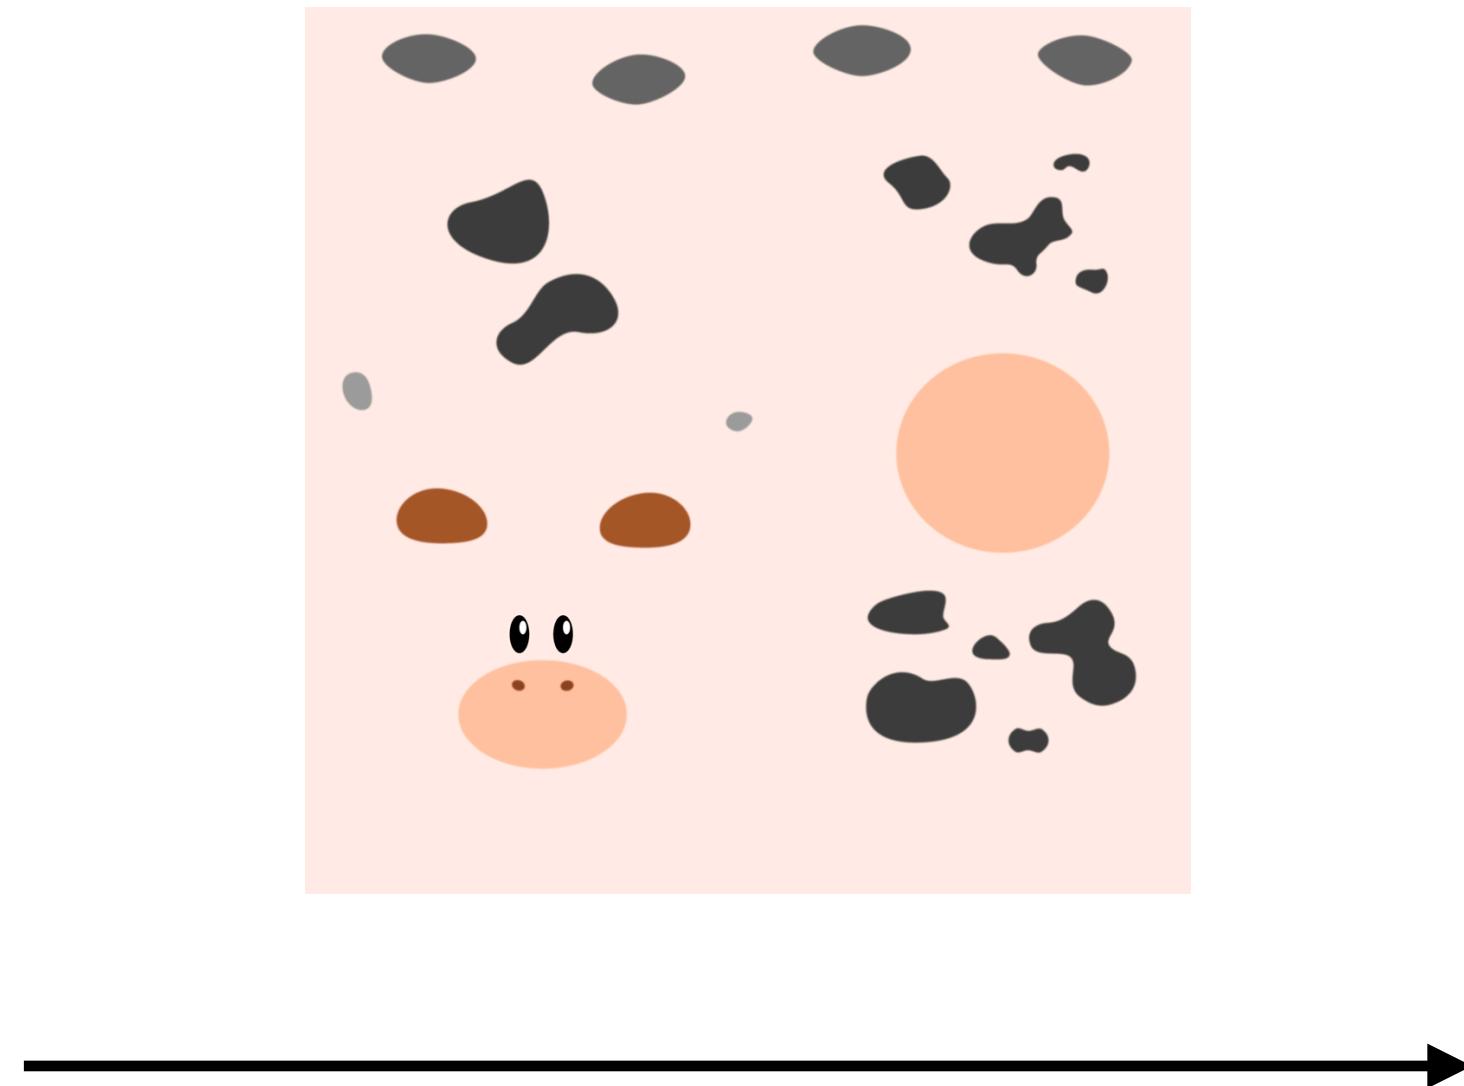
Efficient to compute  
ray-triangle  
intersections

# Triangular Meshes



Efficient to compute  
ray-triangle  
intersections

# Triangular Meshes

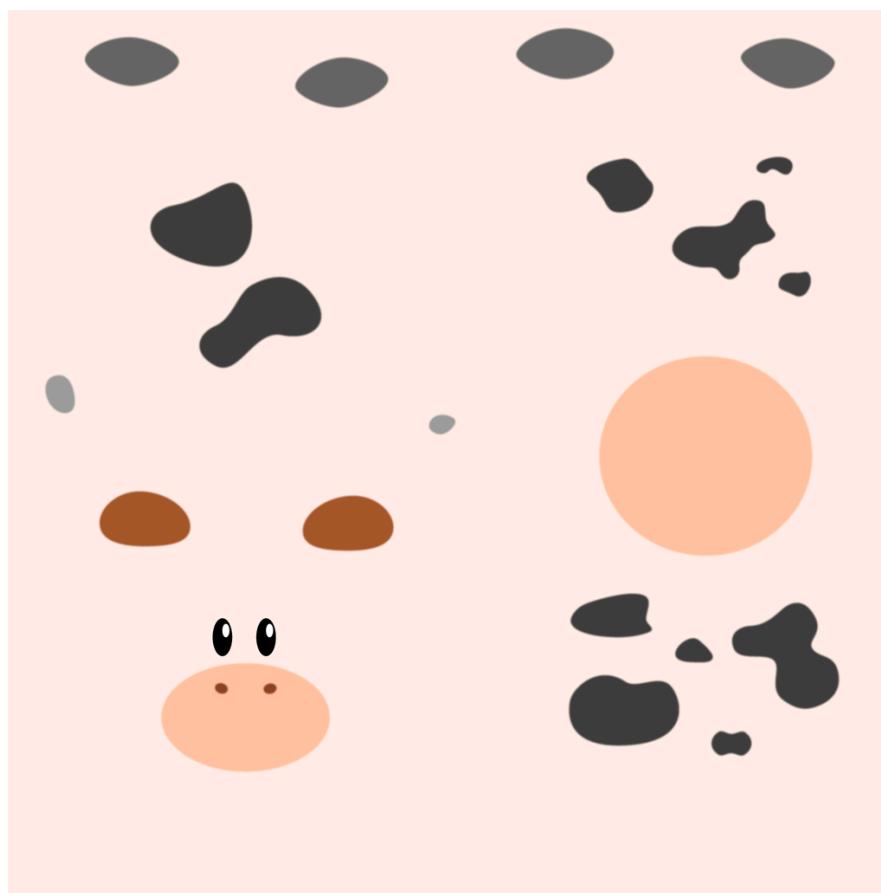


Efficient to compute  
ray-triangle  
intersections

# Triangular Meshes

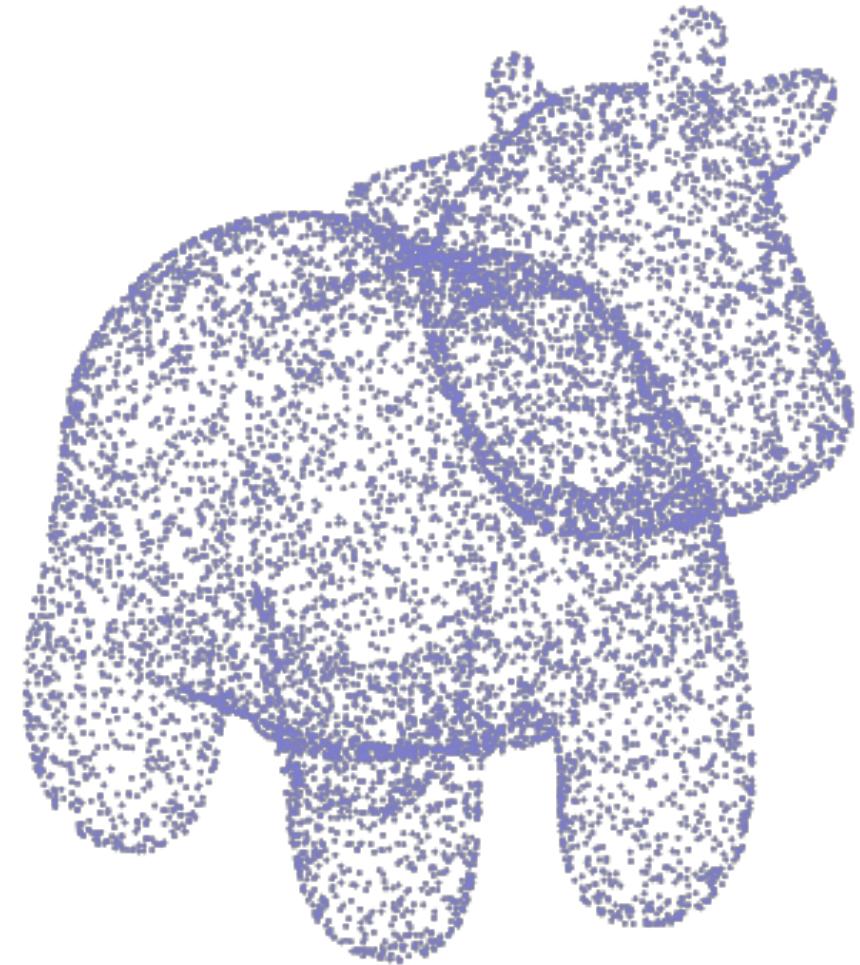


Efficient to compute  
ray-triangle  
intersections

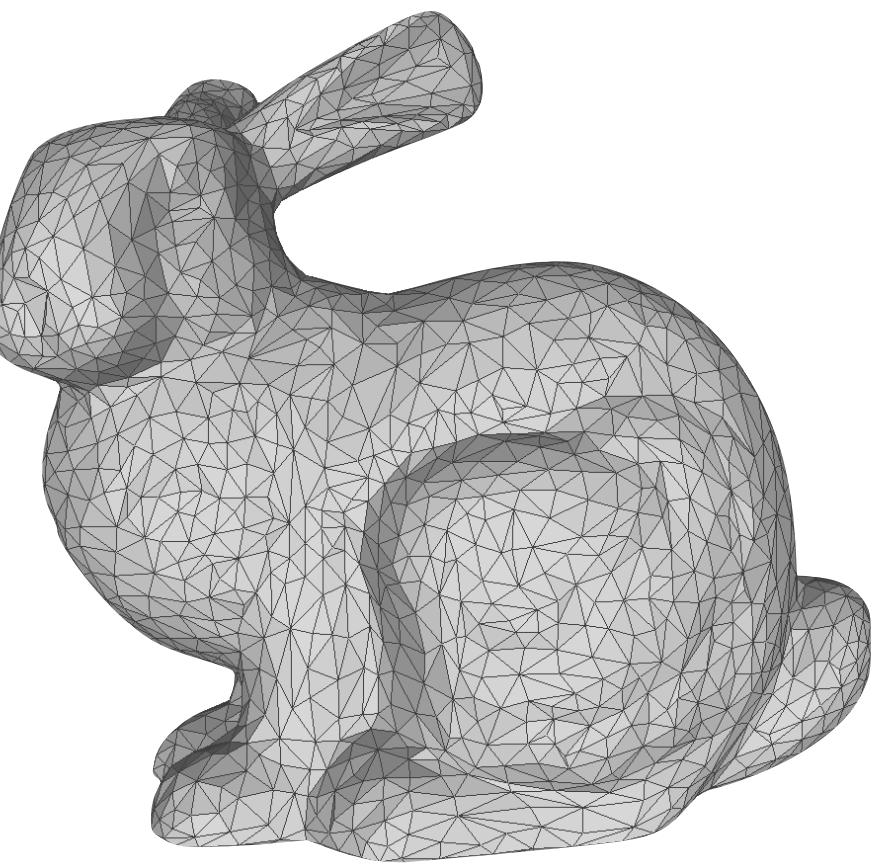
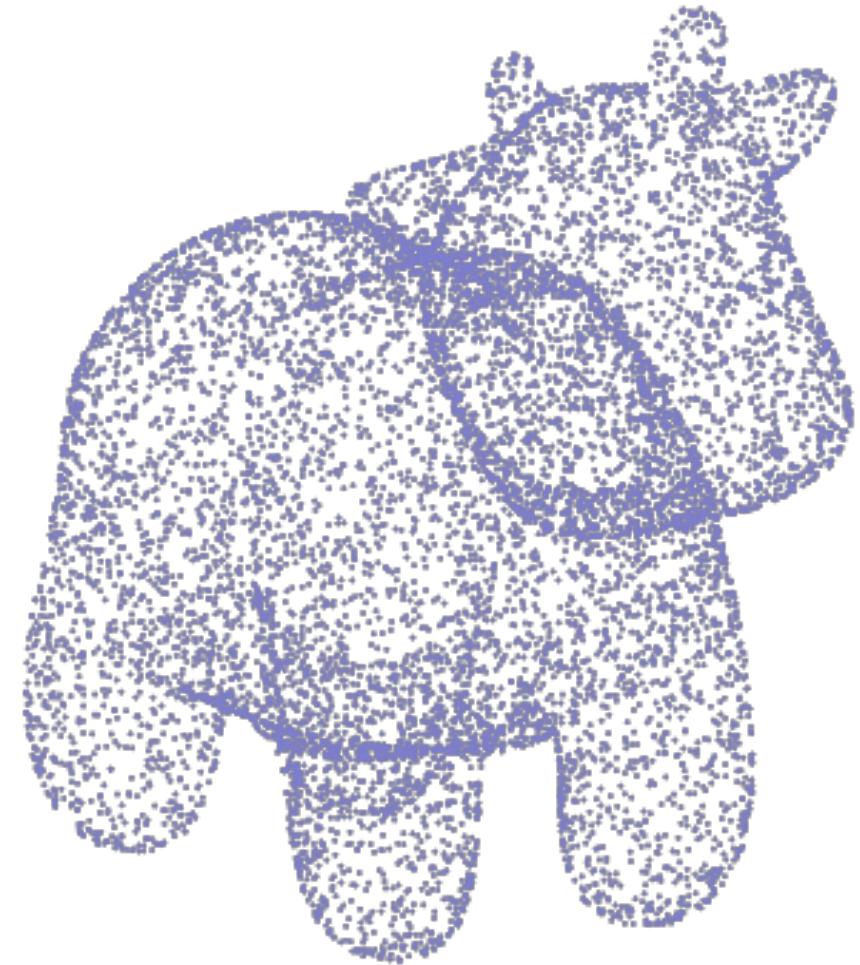


Easy to texture and render  
(common representation  
across graphics)

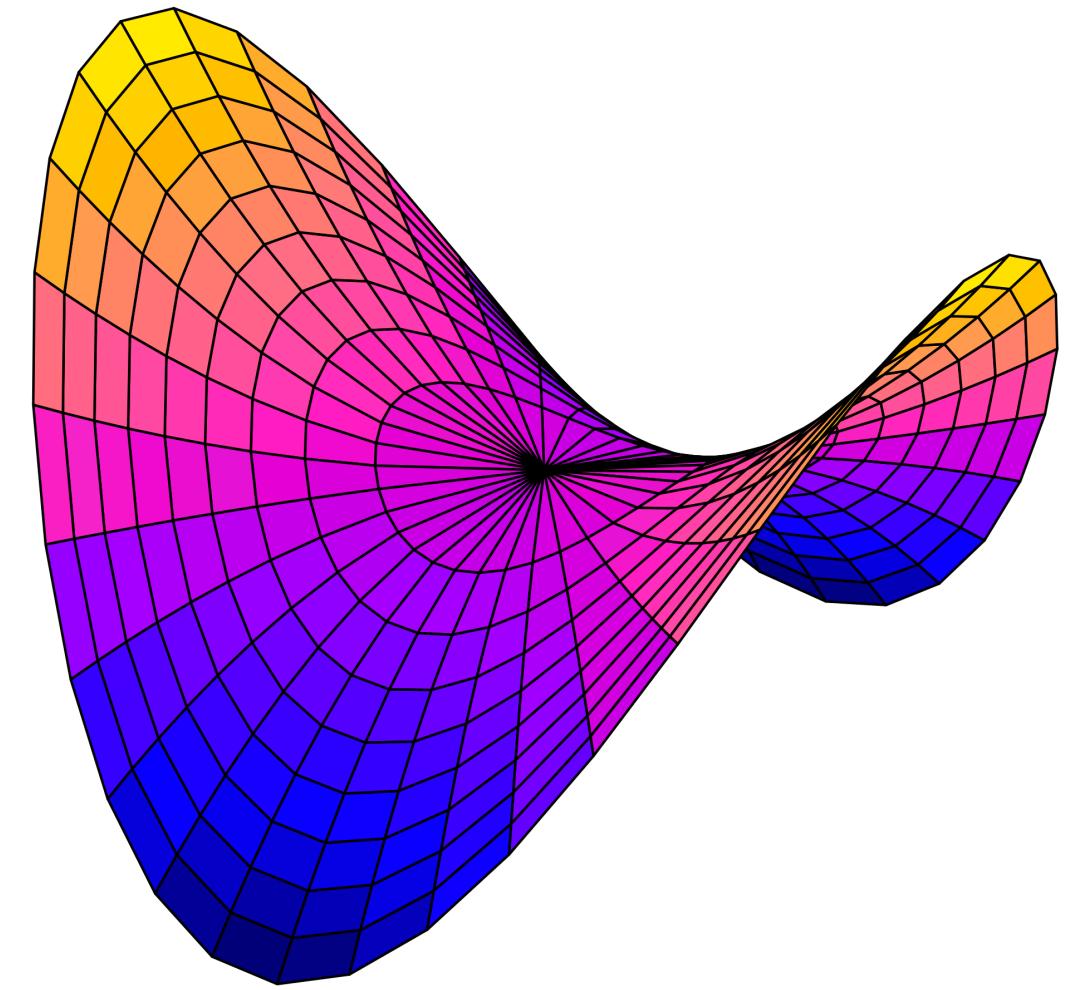
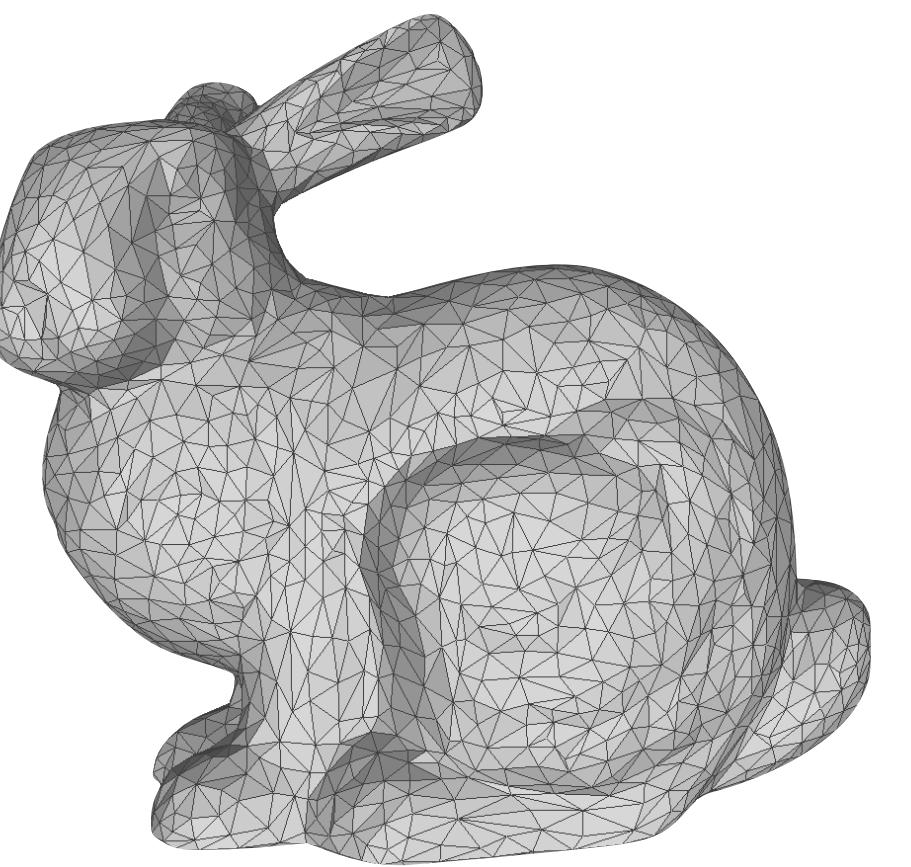
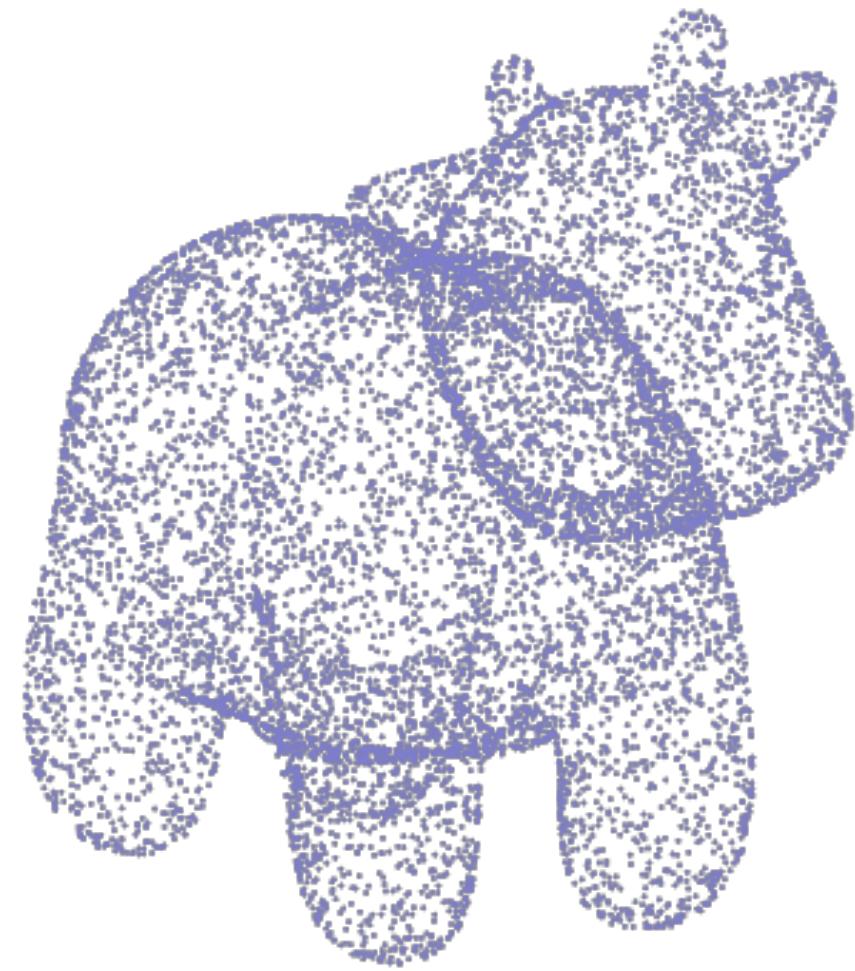
# **Surface Representations**



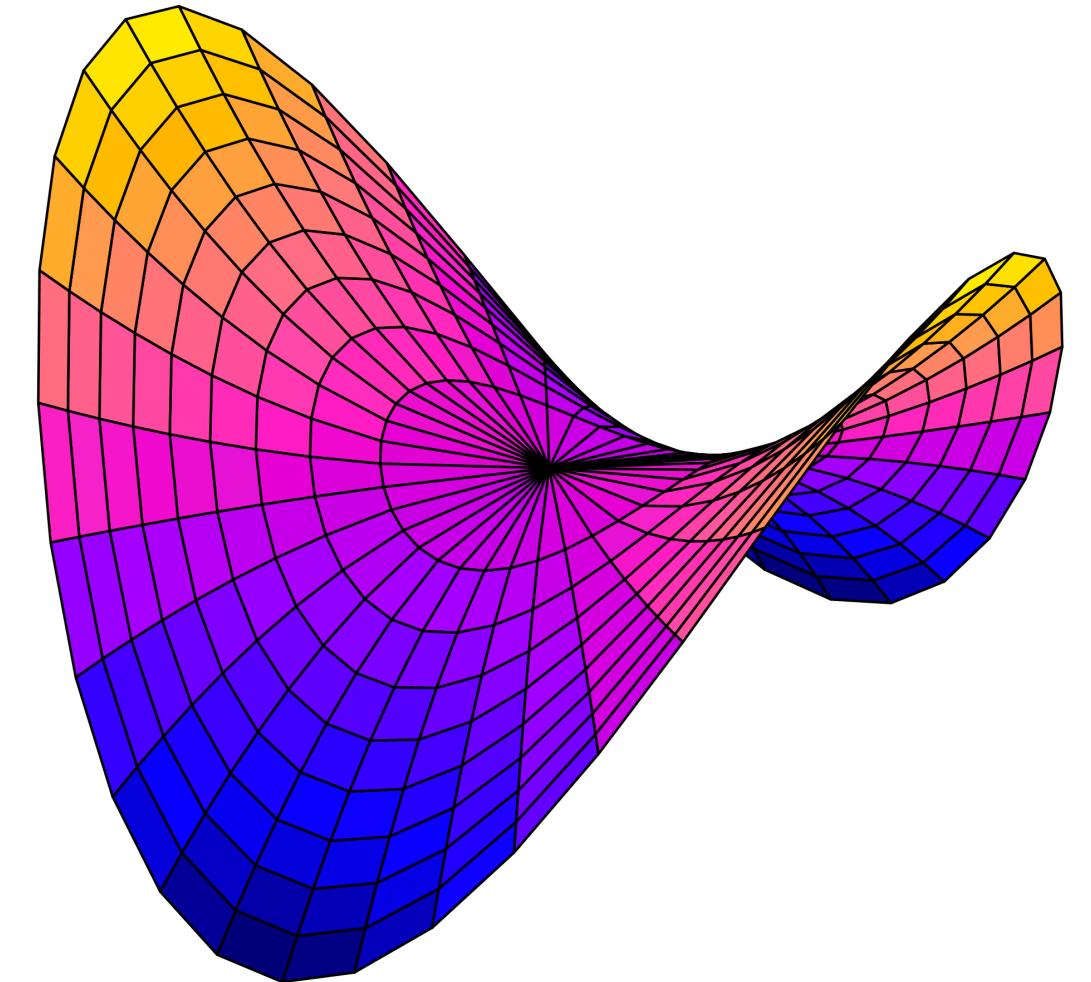
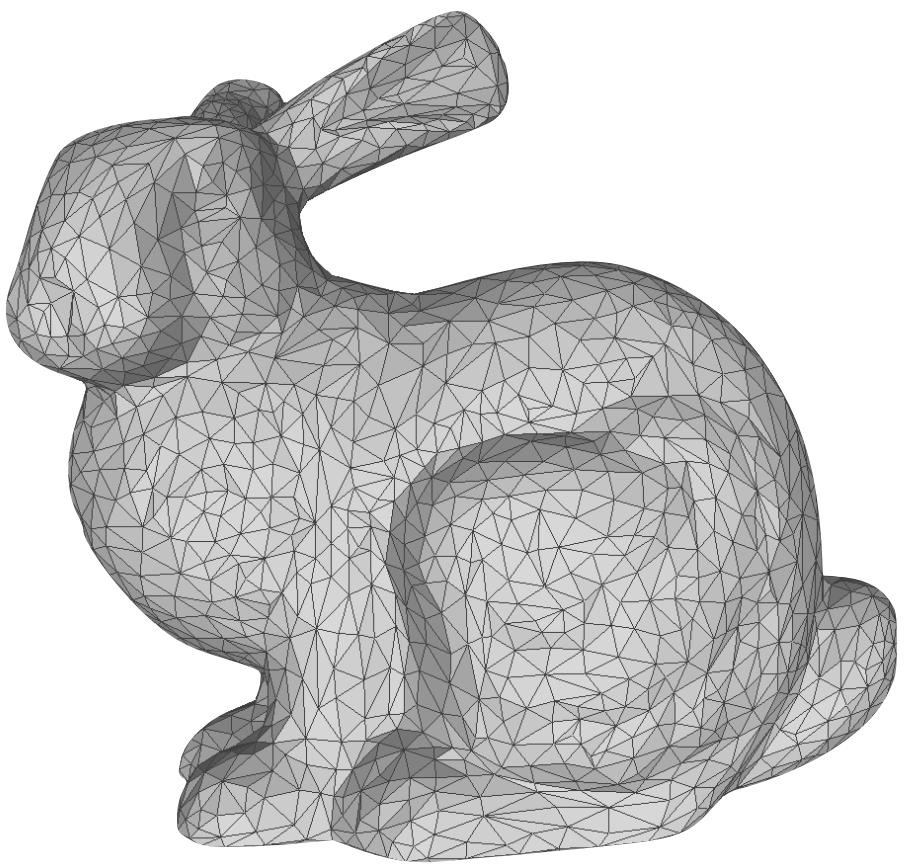
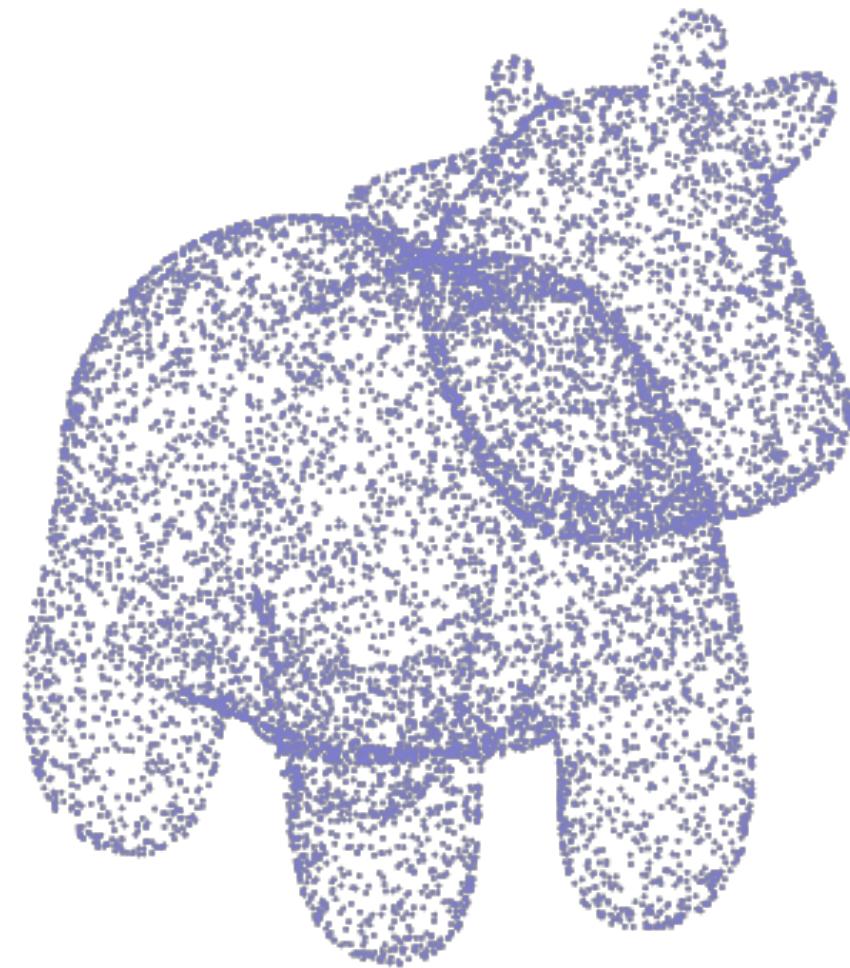
## Surface Representations



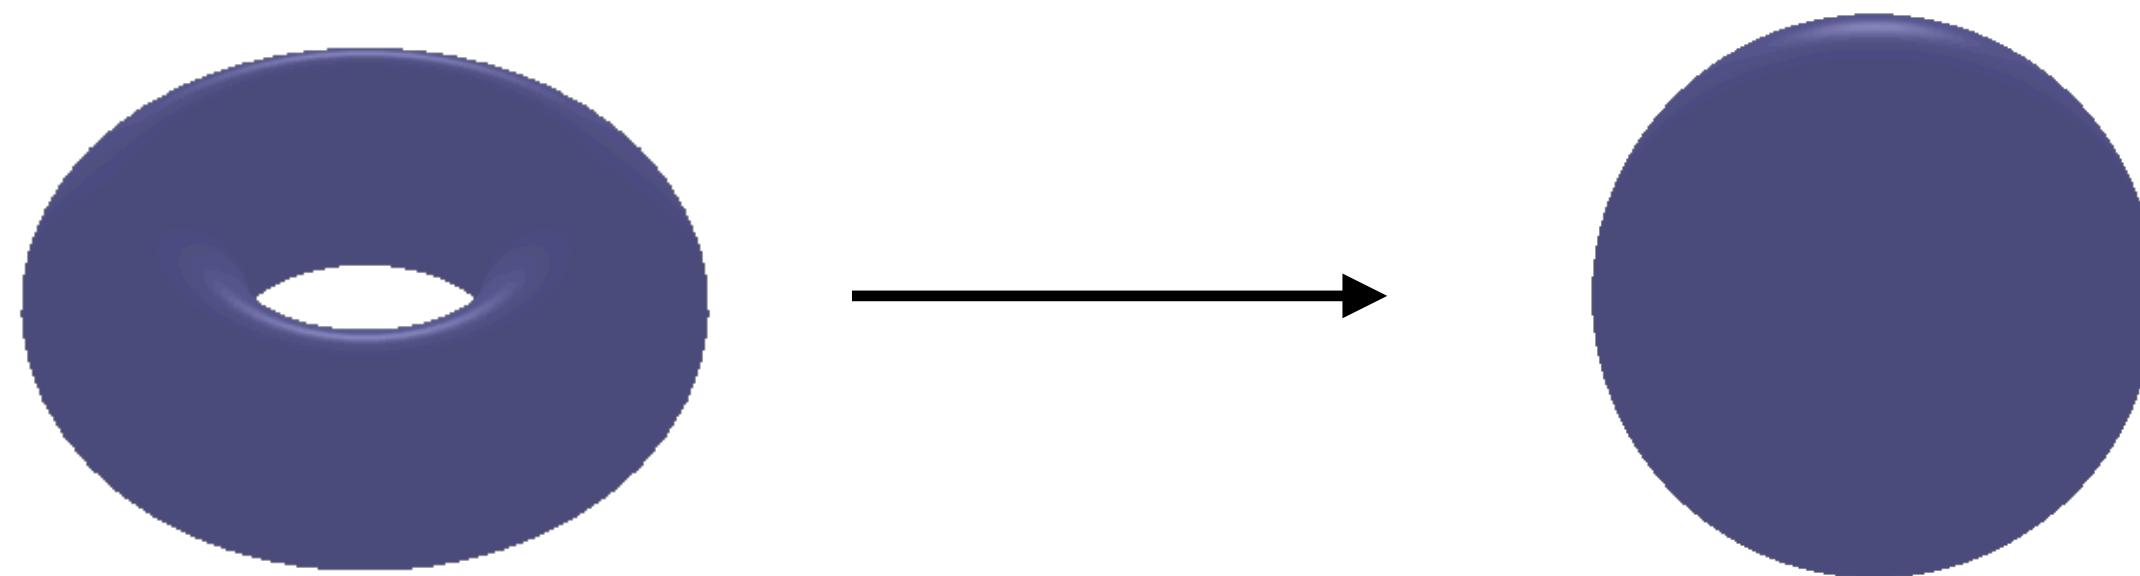
## Surface Representations

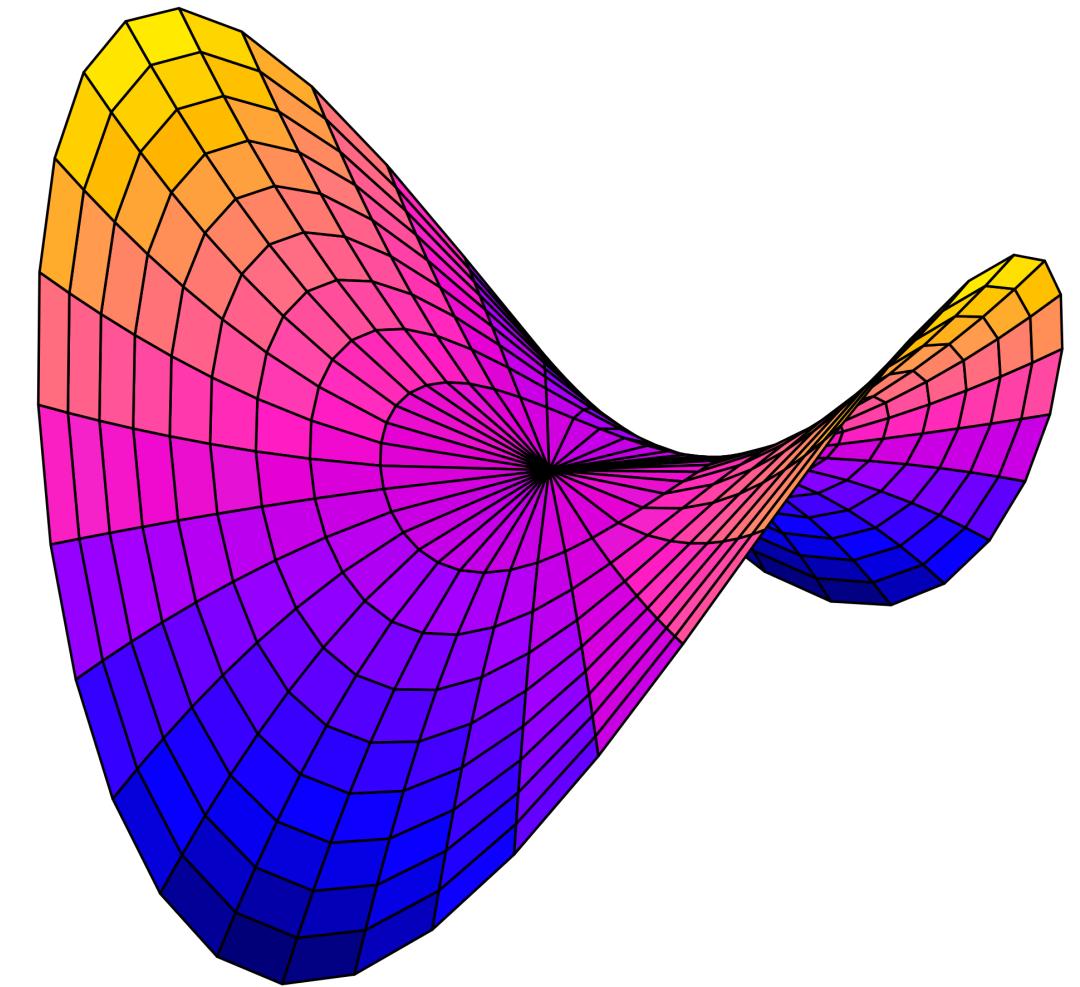
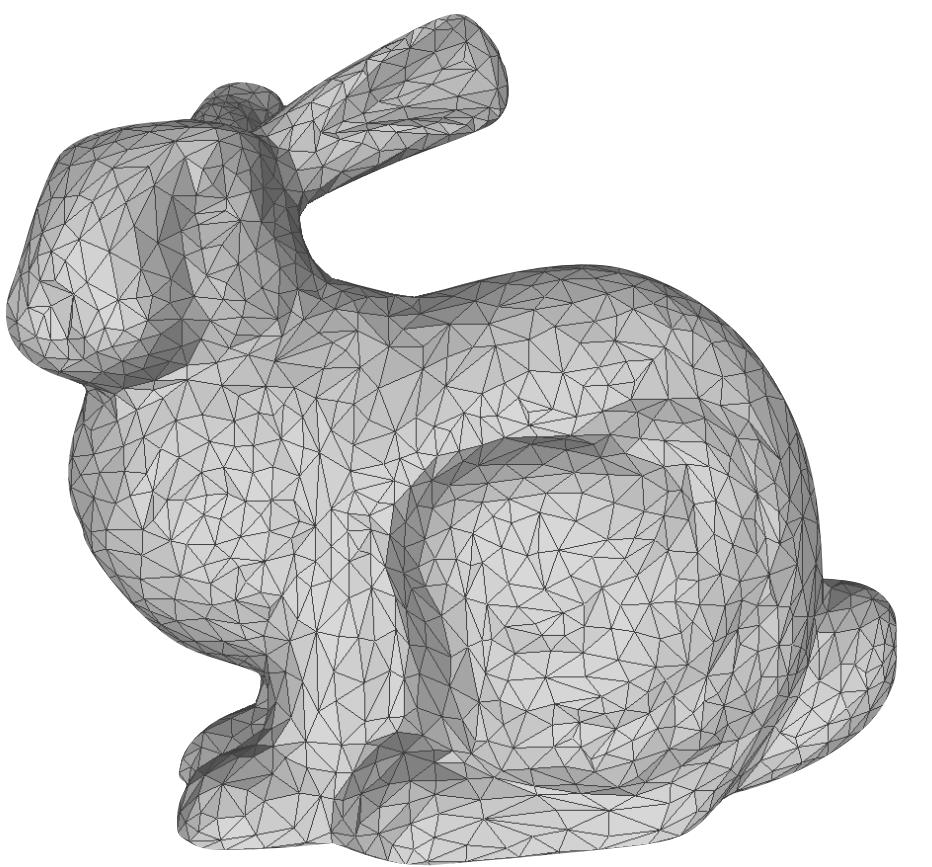
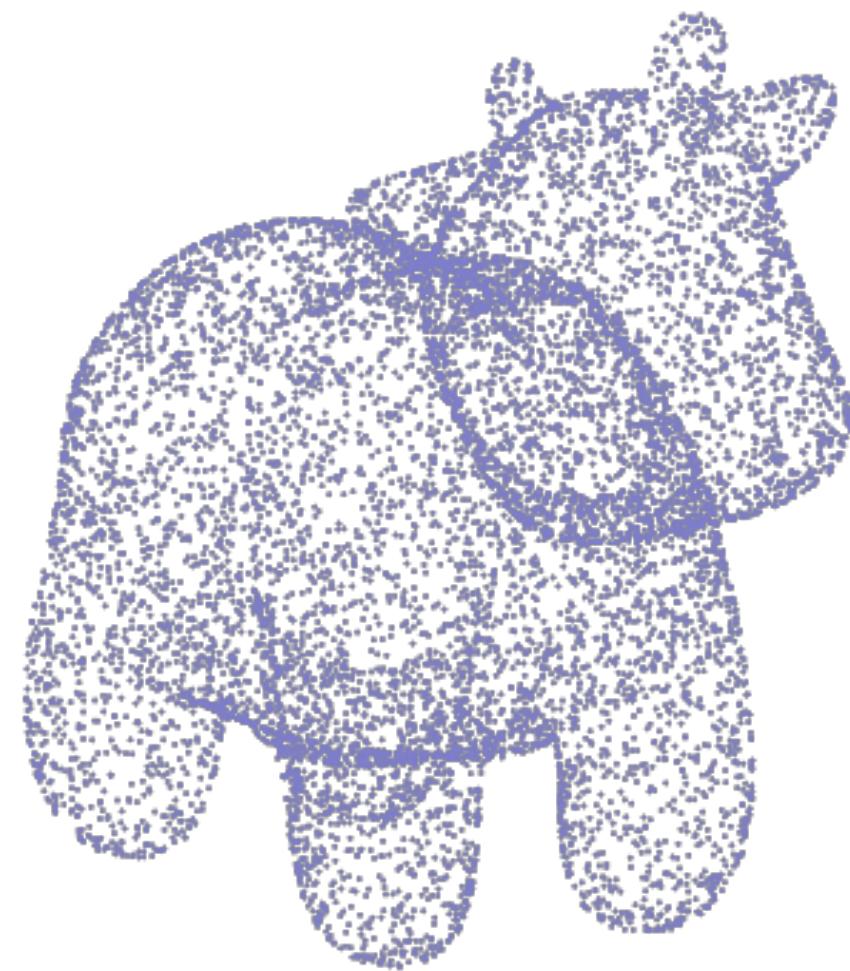


## Surface Representations

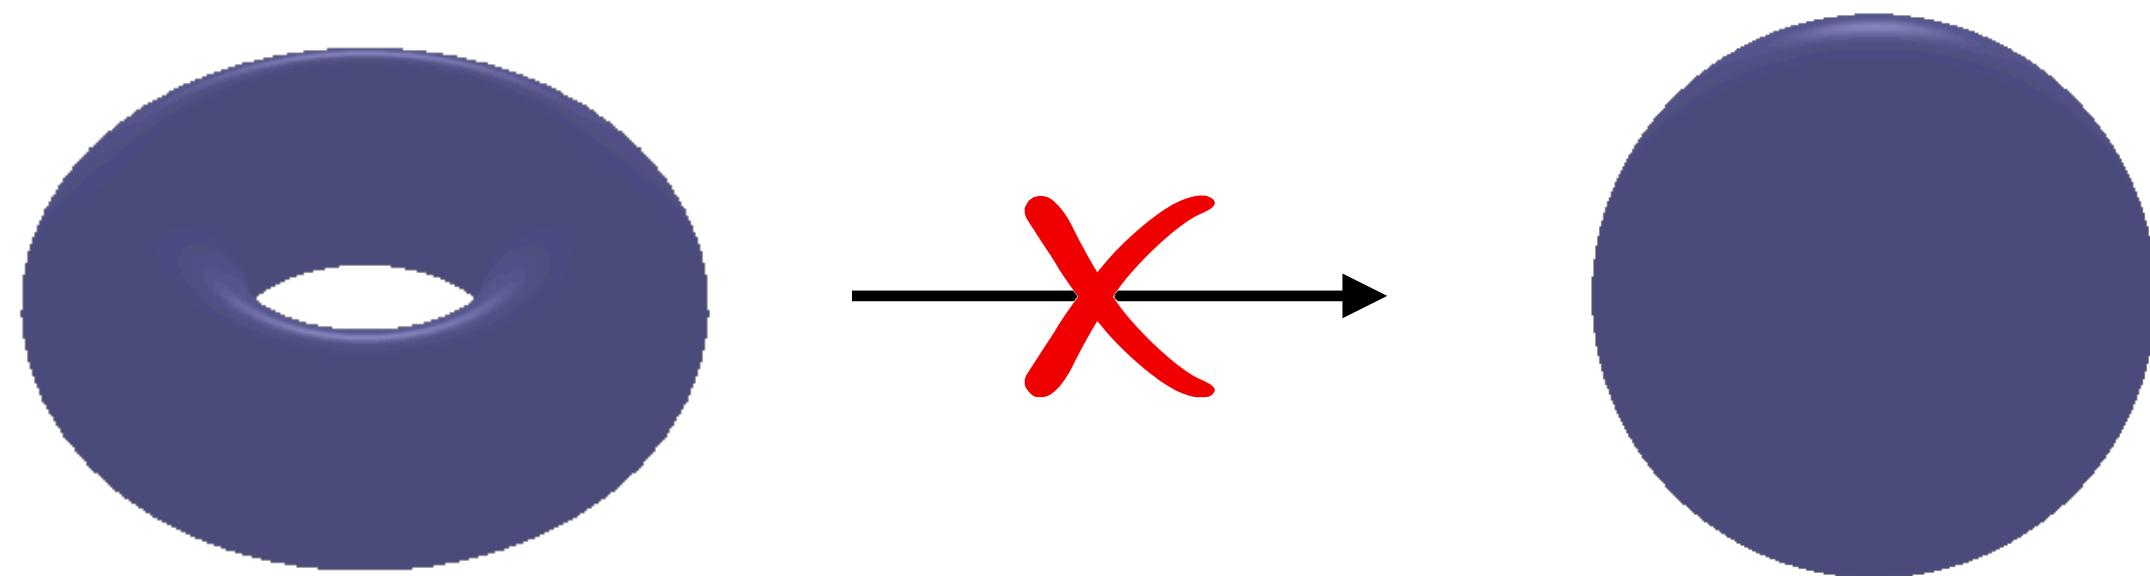


## Surface Representations

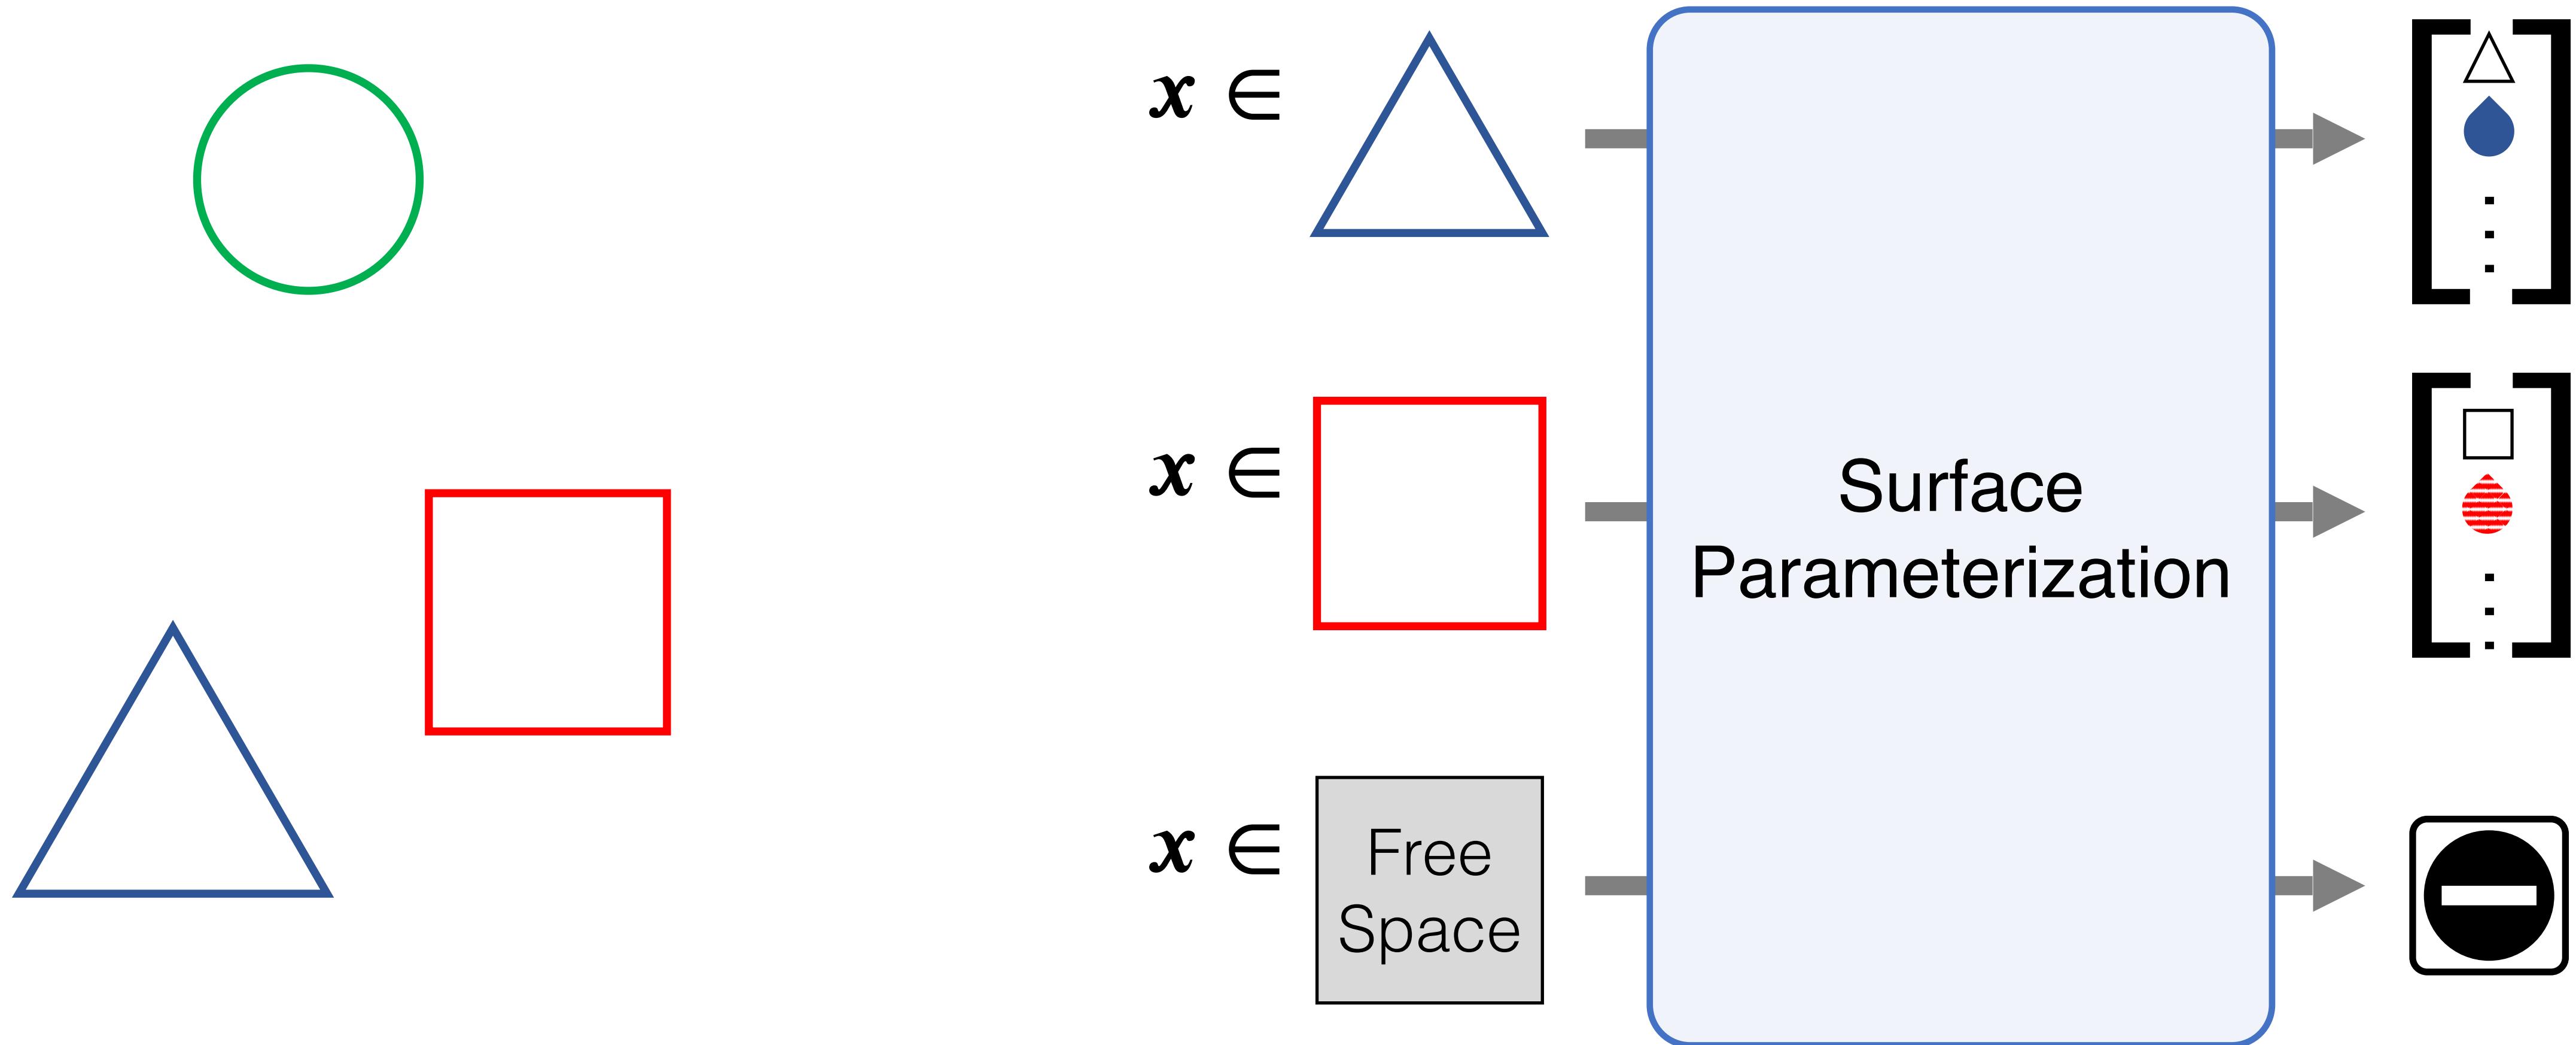




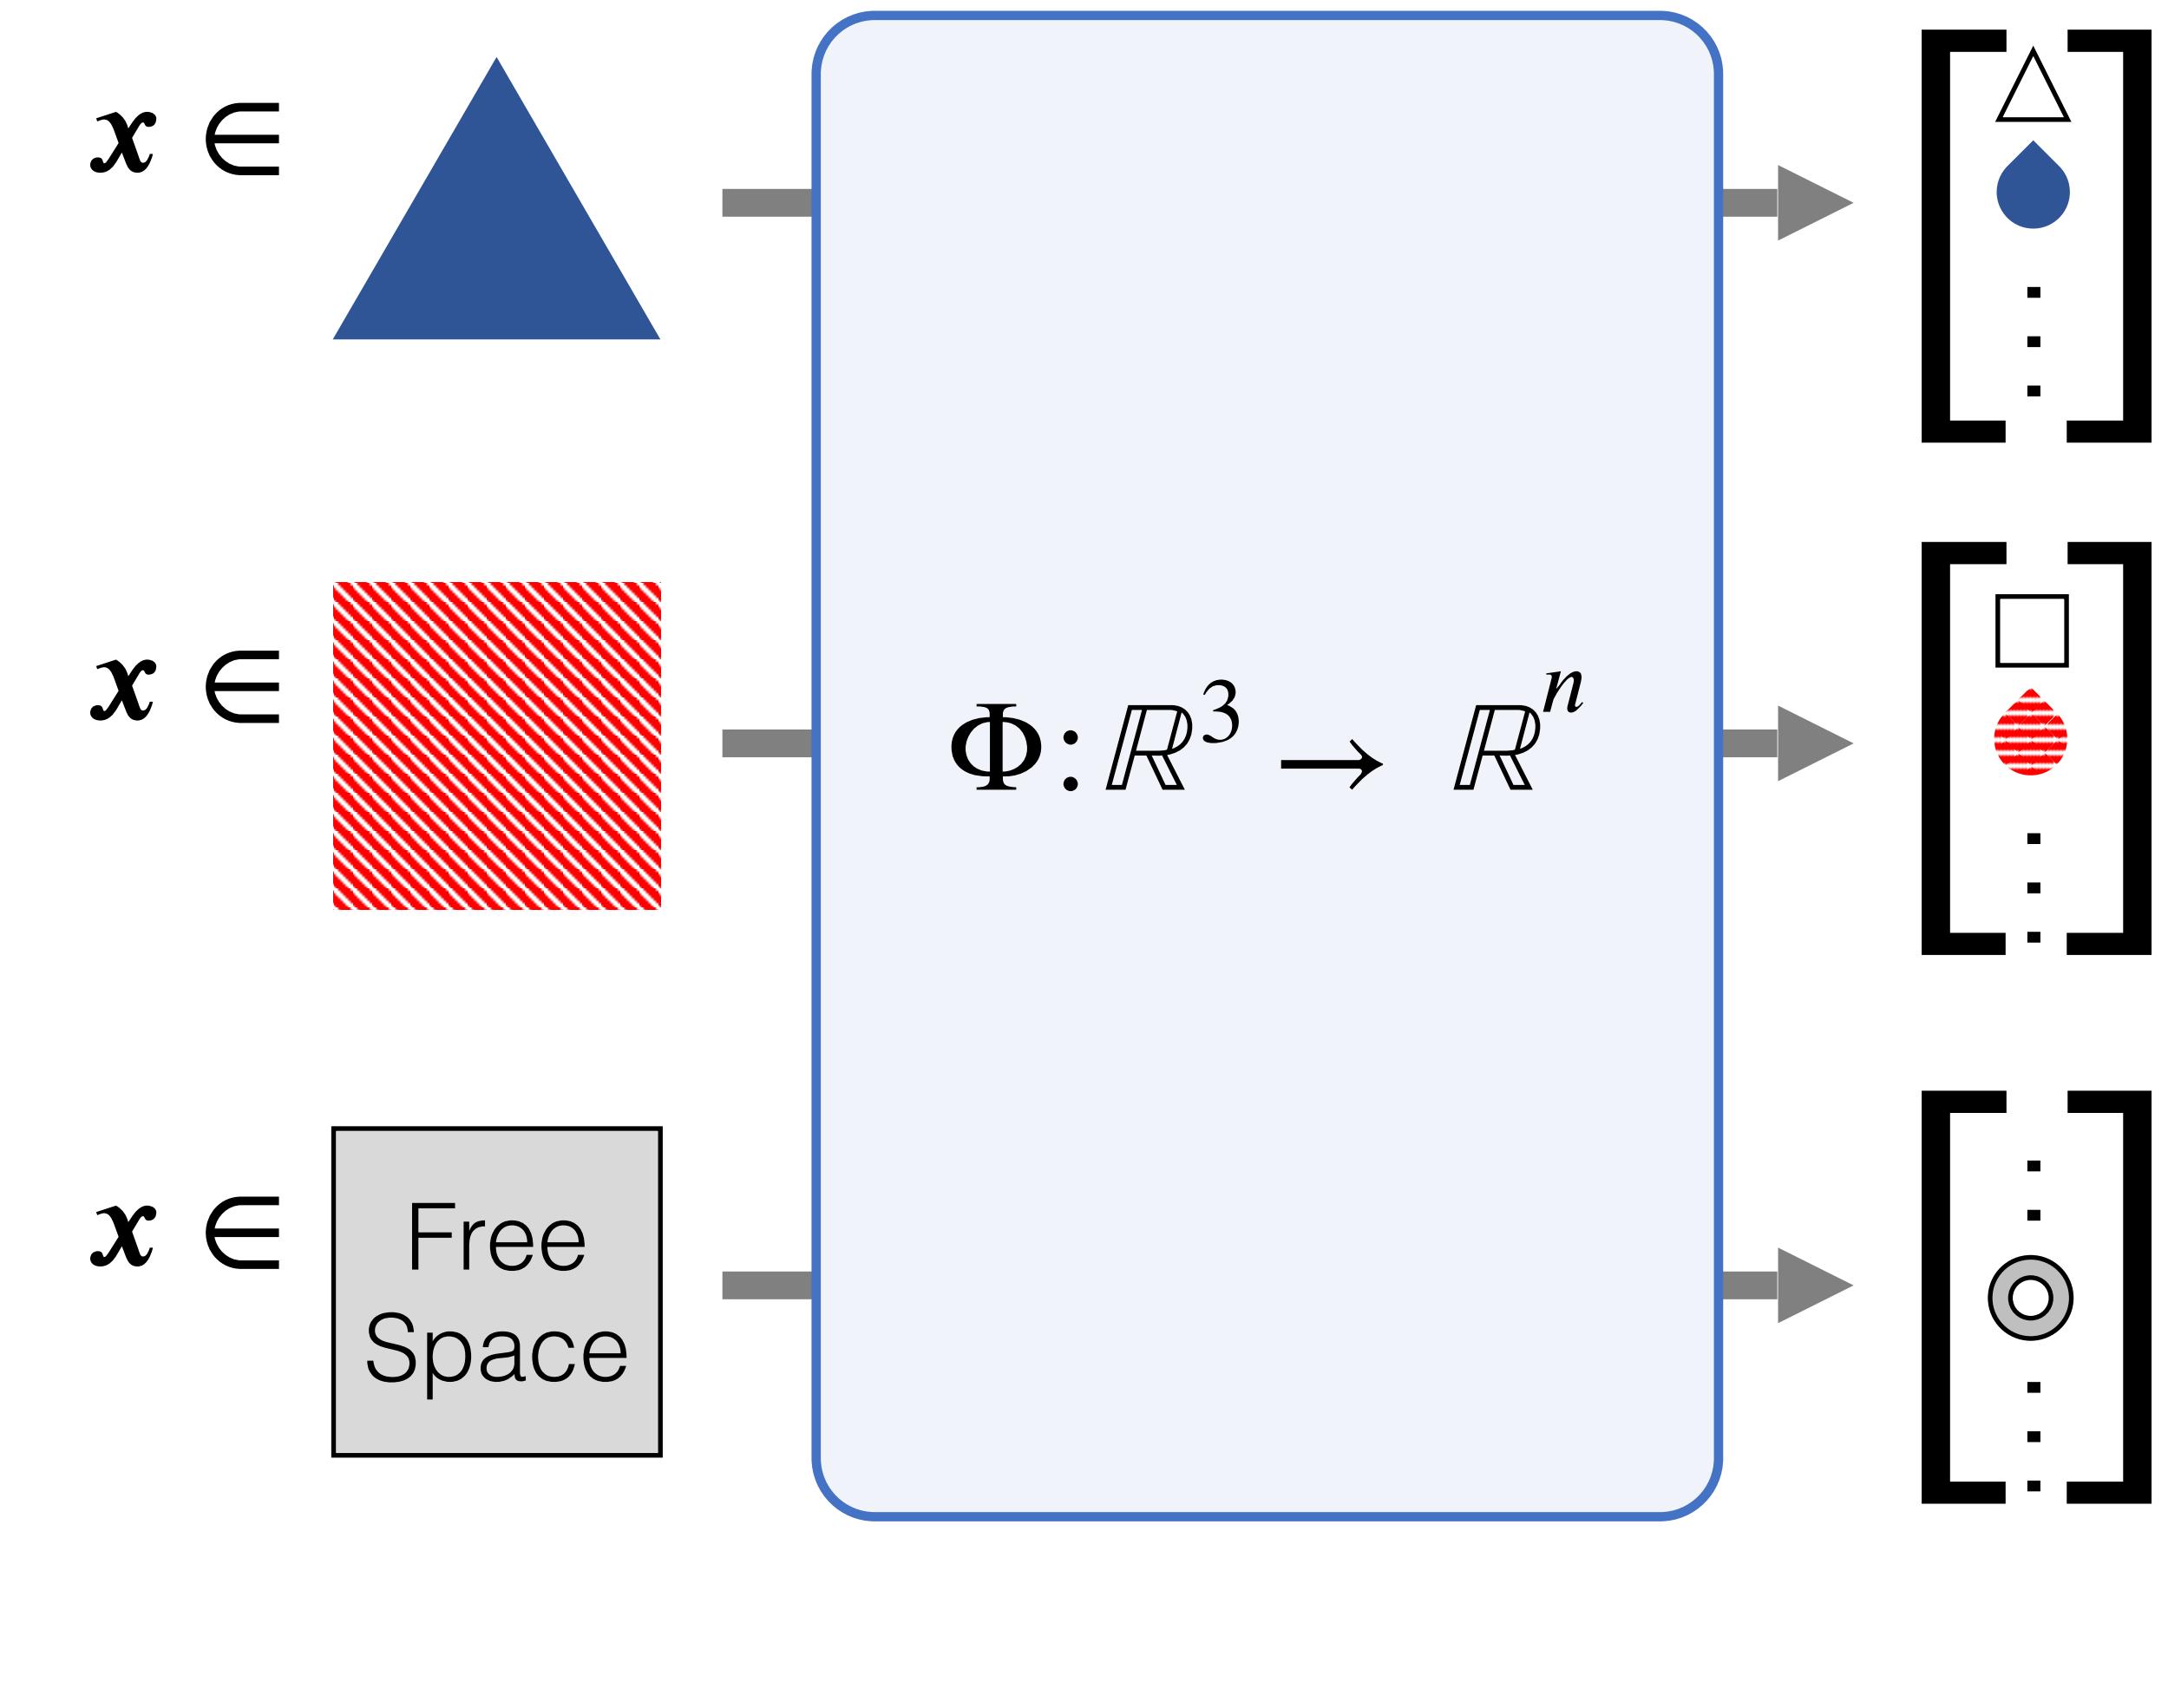
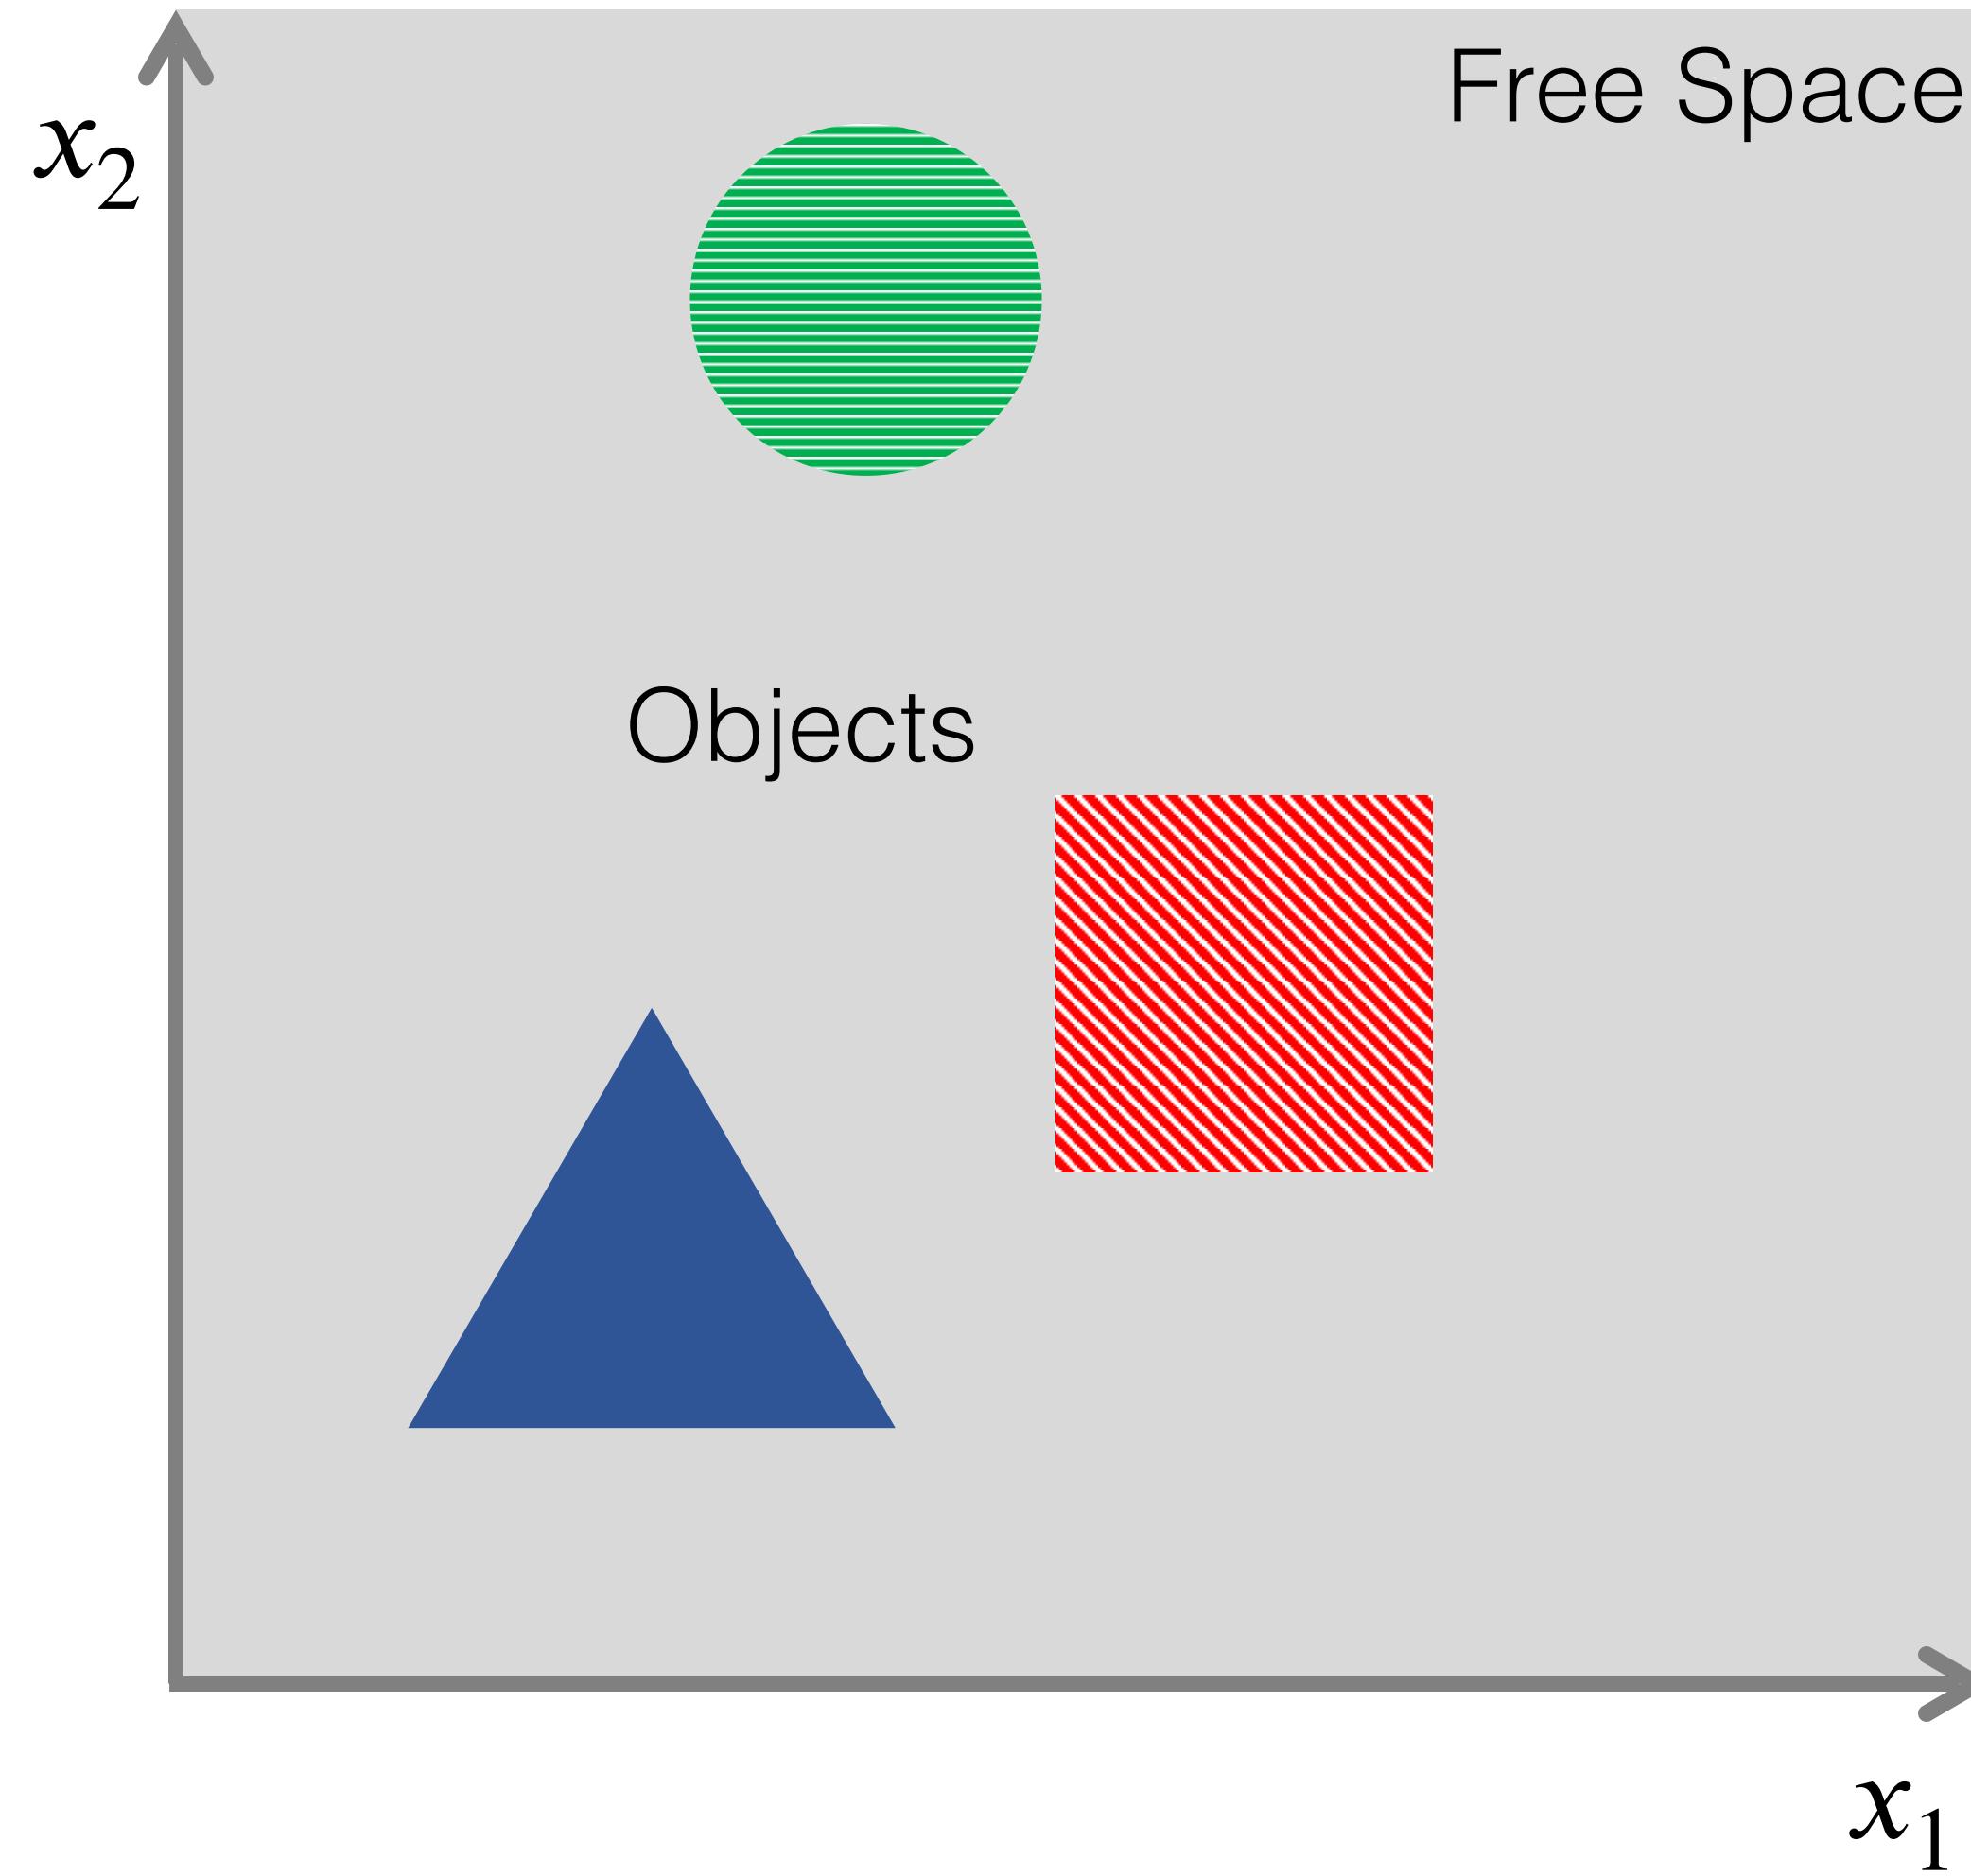
## Surface Representations



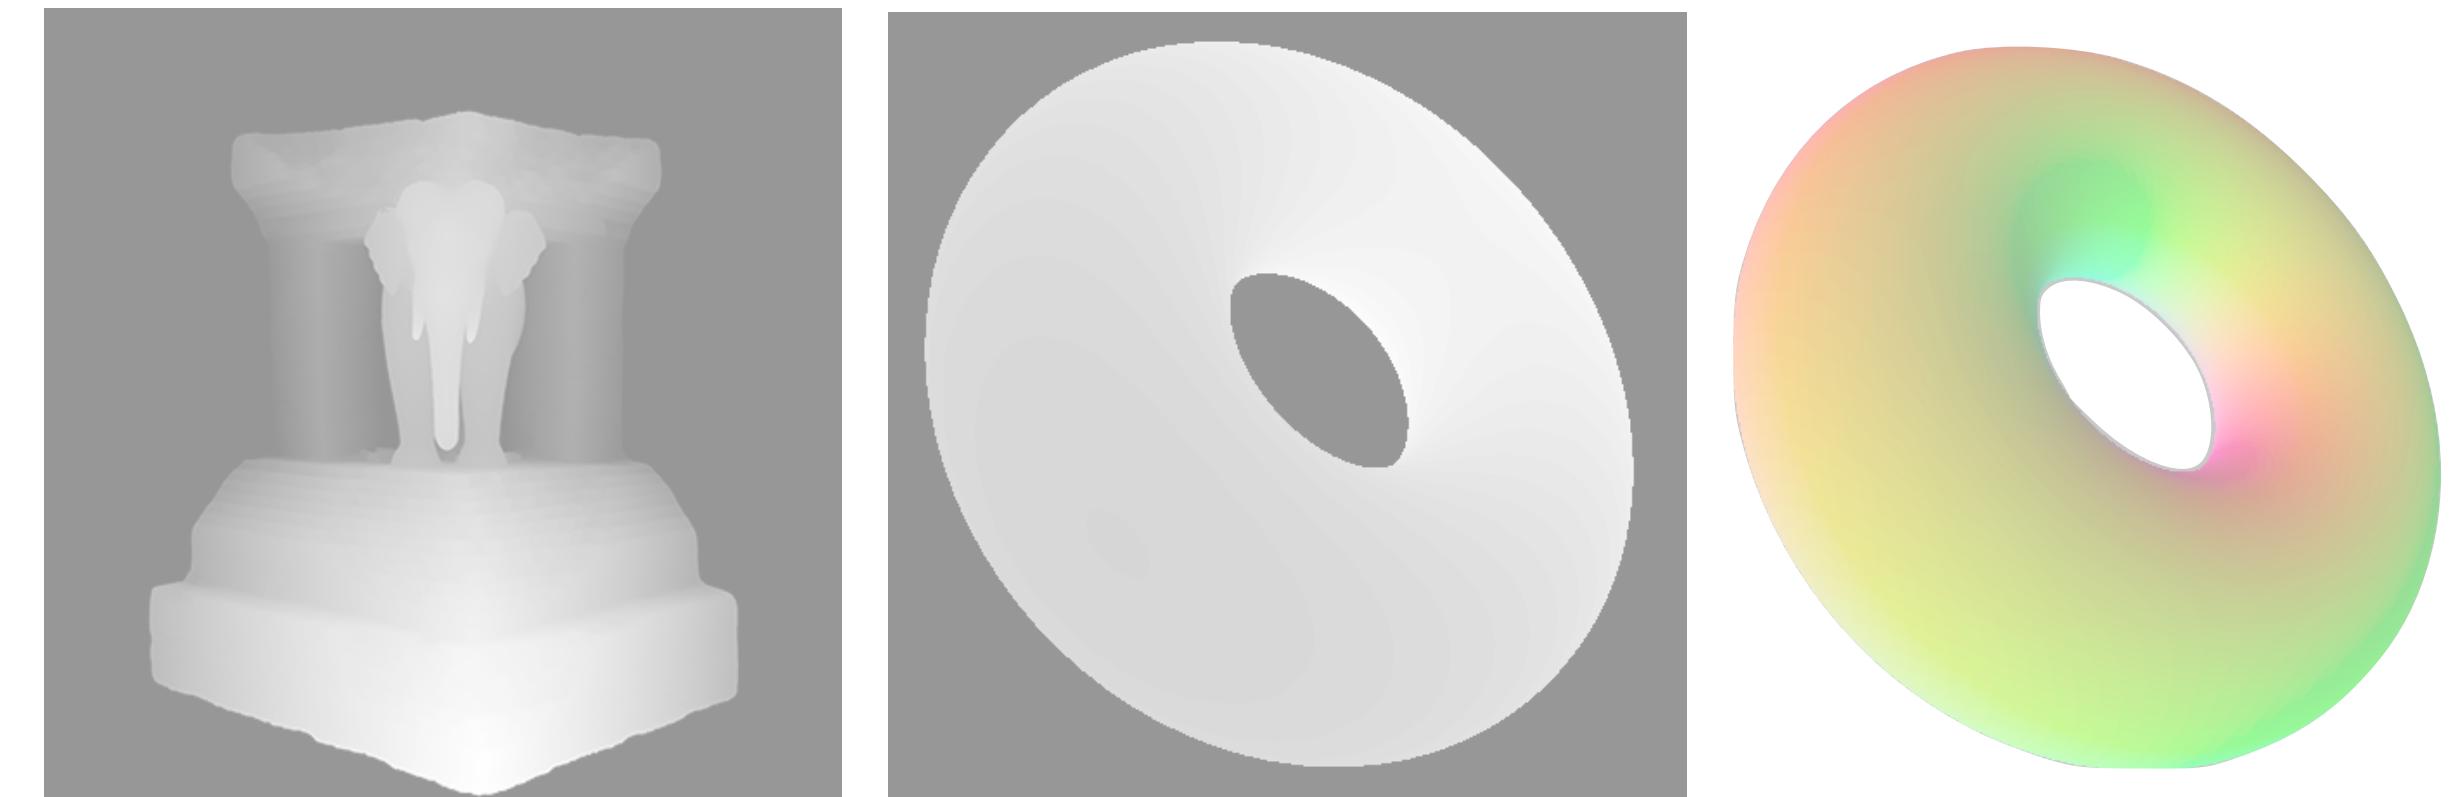
Surface representations only parameterize stuff on surfaces of things.



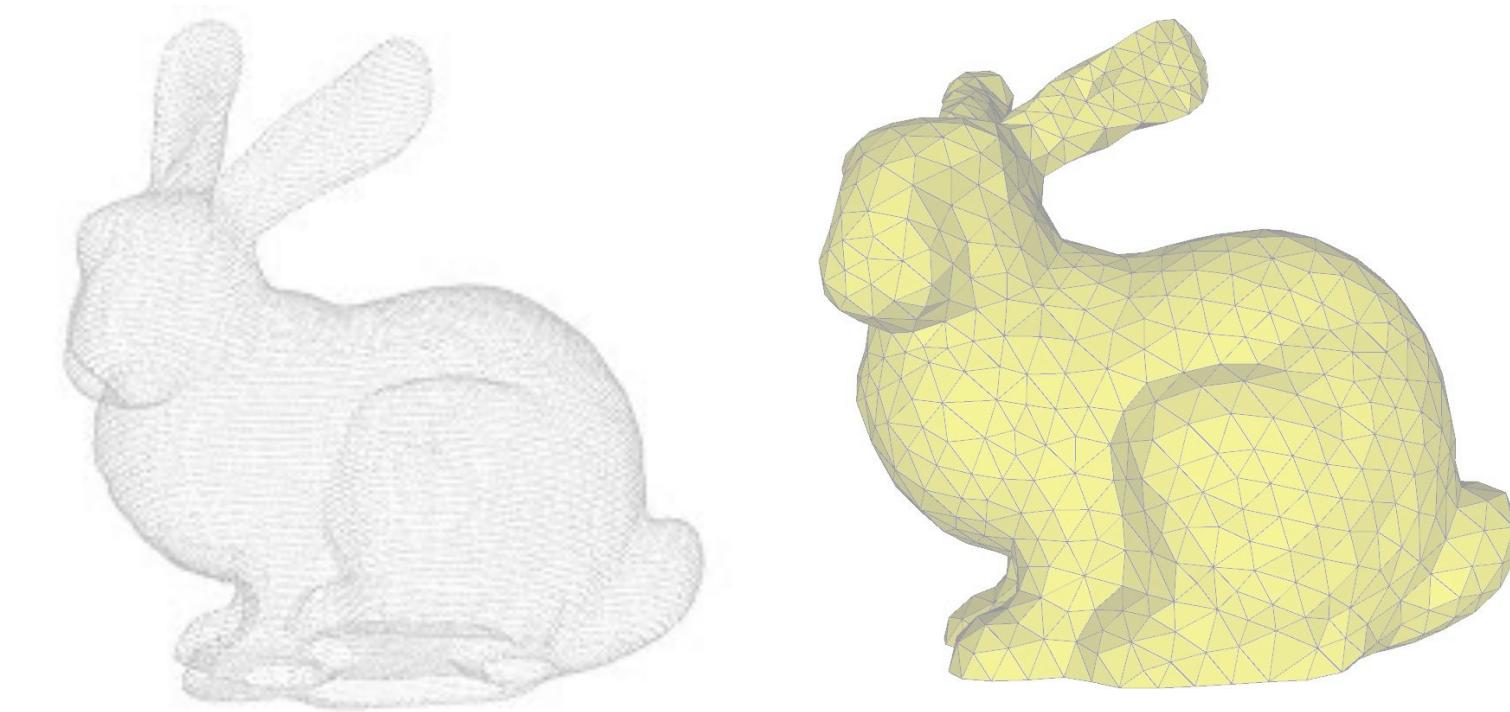
# What if we wanted to parameterize stuff **everywhere**?



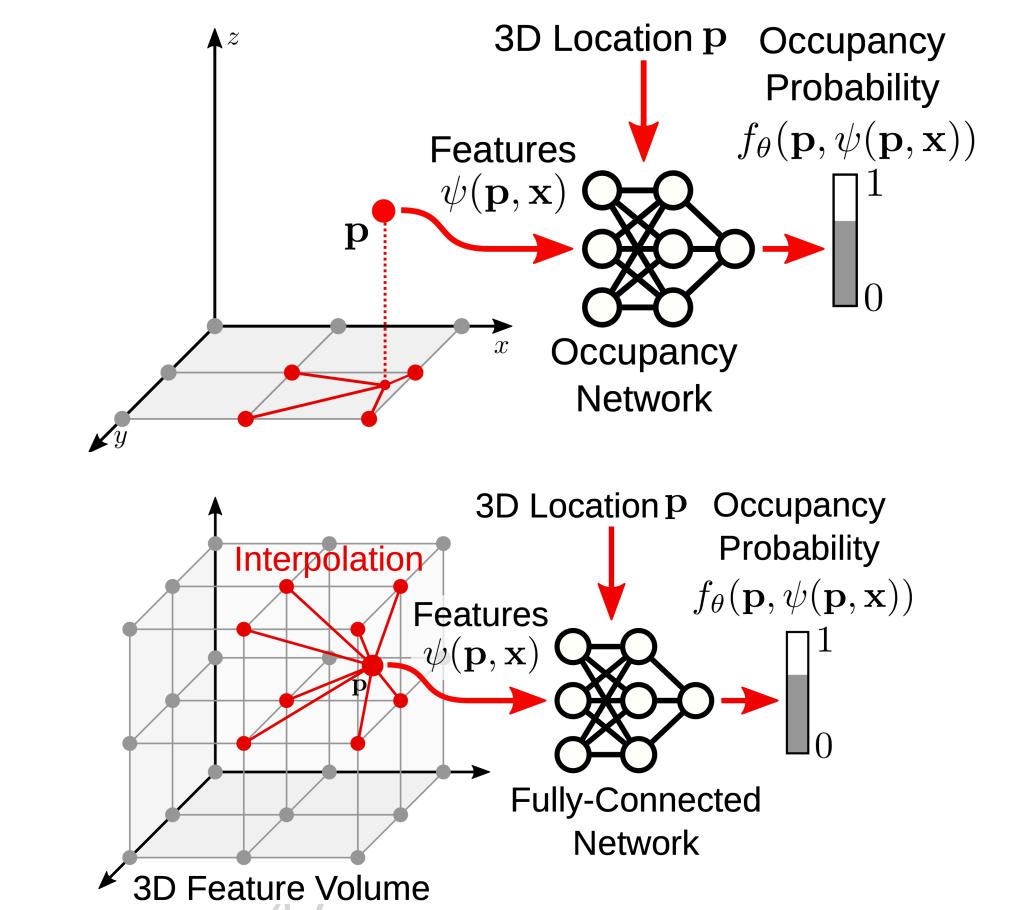
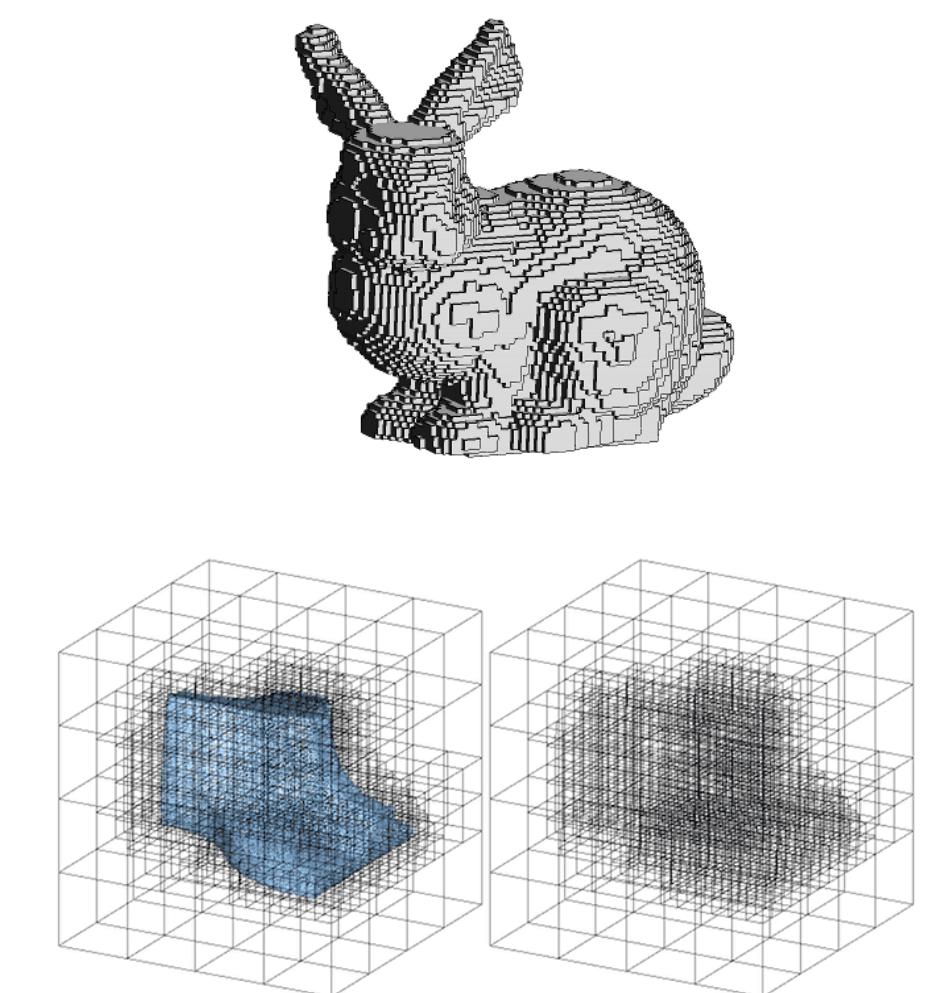
## 2.5D Representations



## Surface Representations



## Field Parameterizations



# Definition: Field (physics)

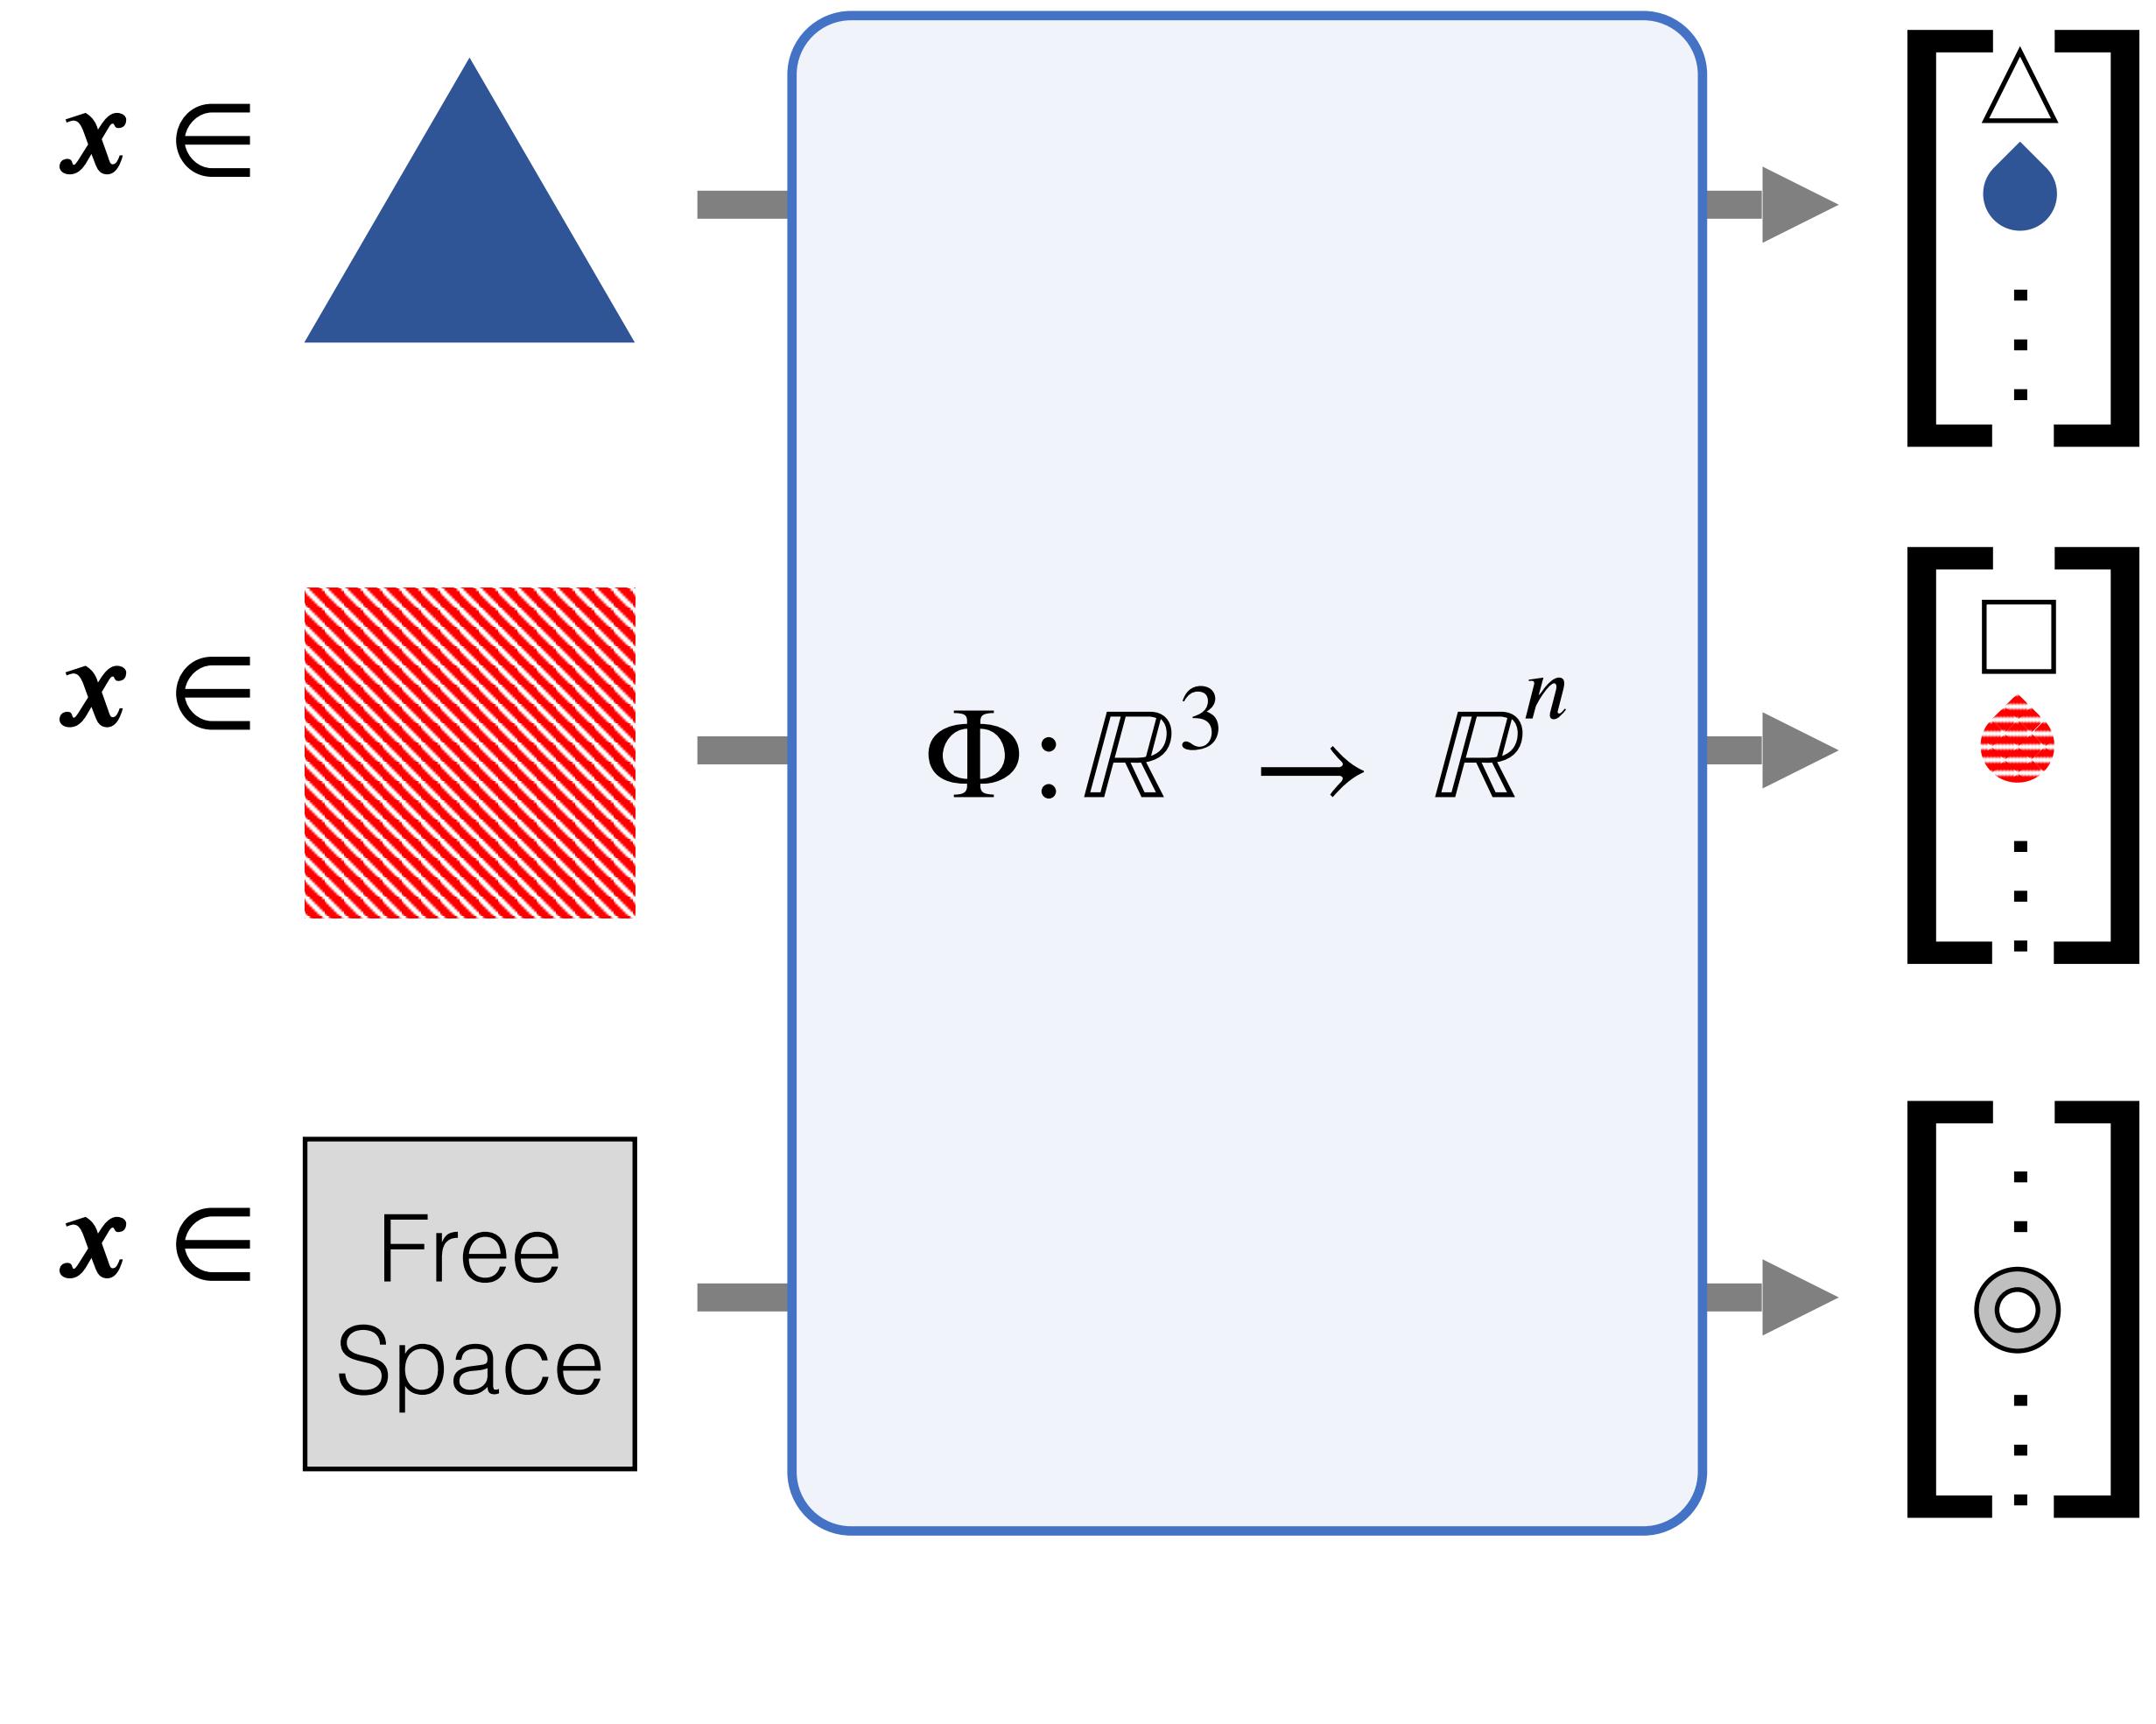
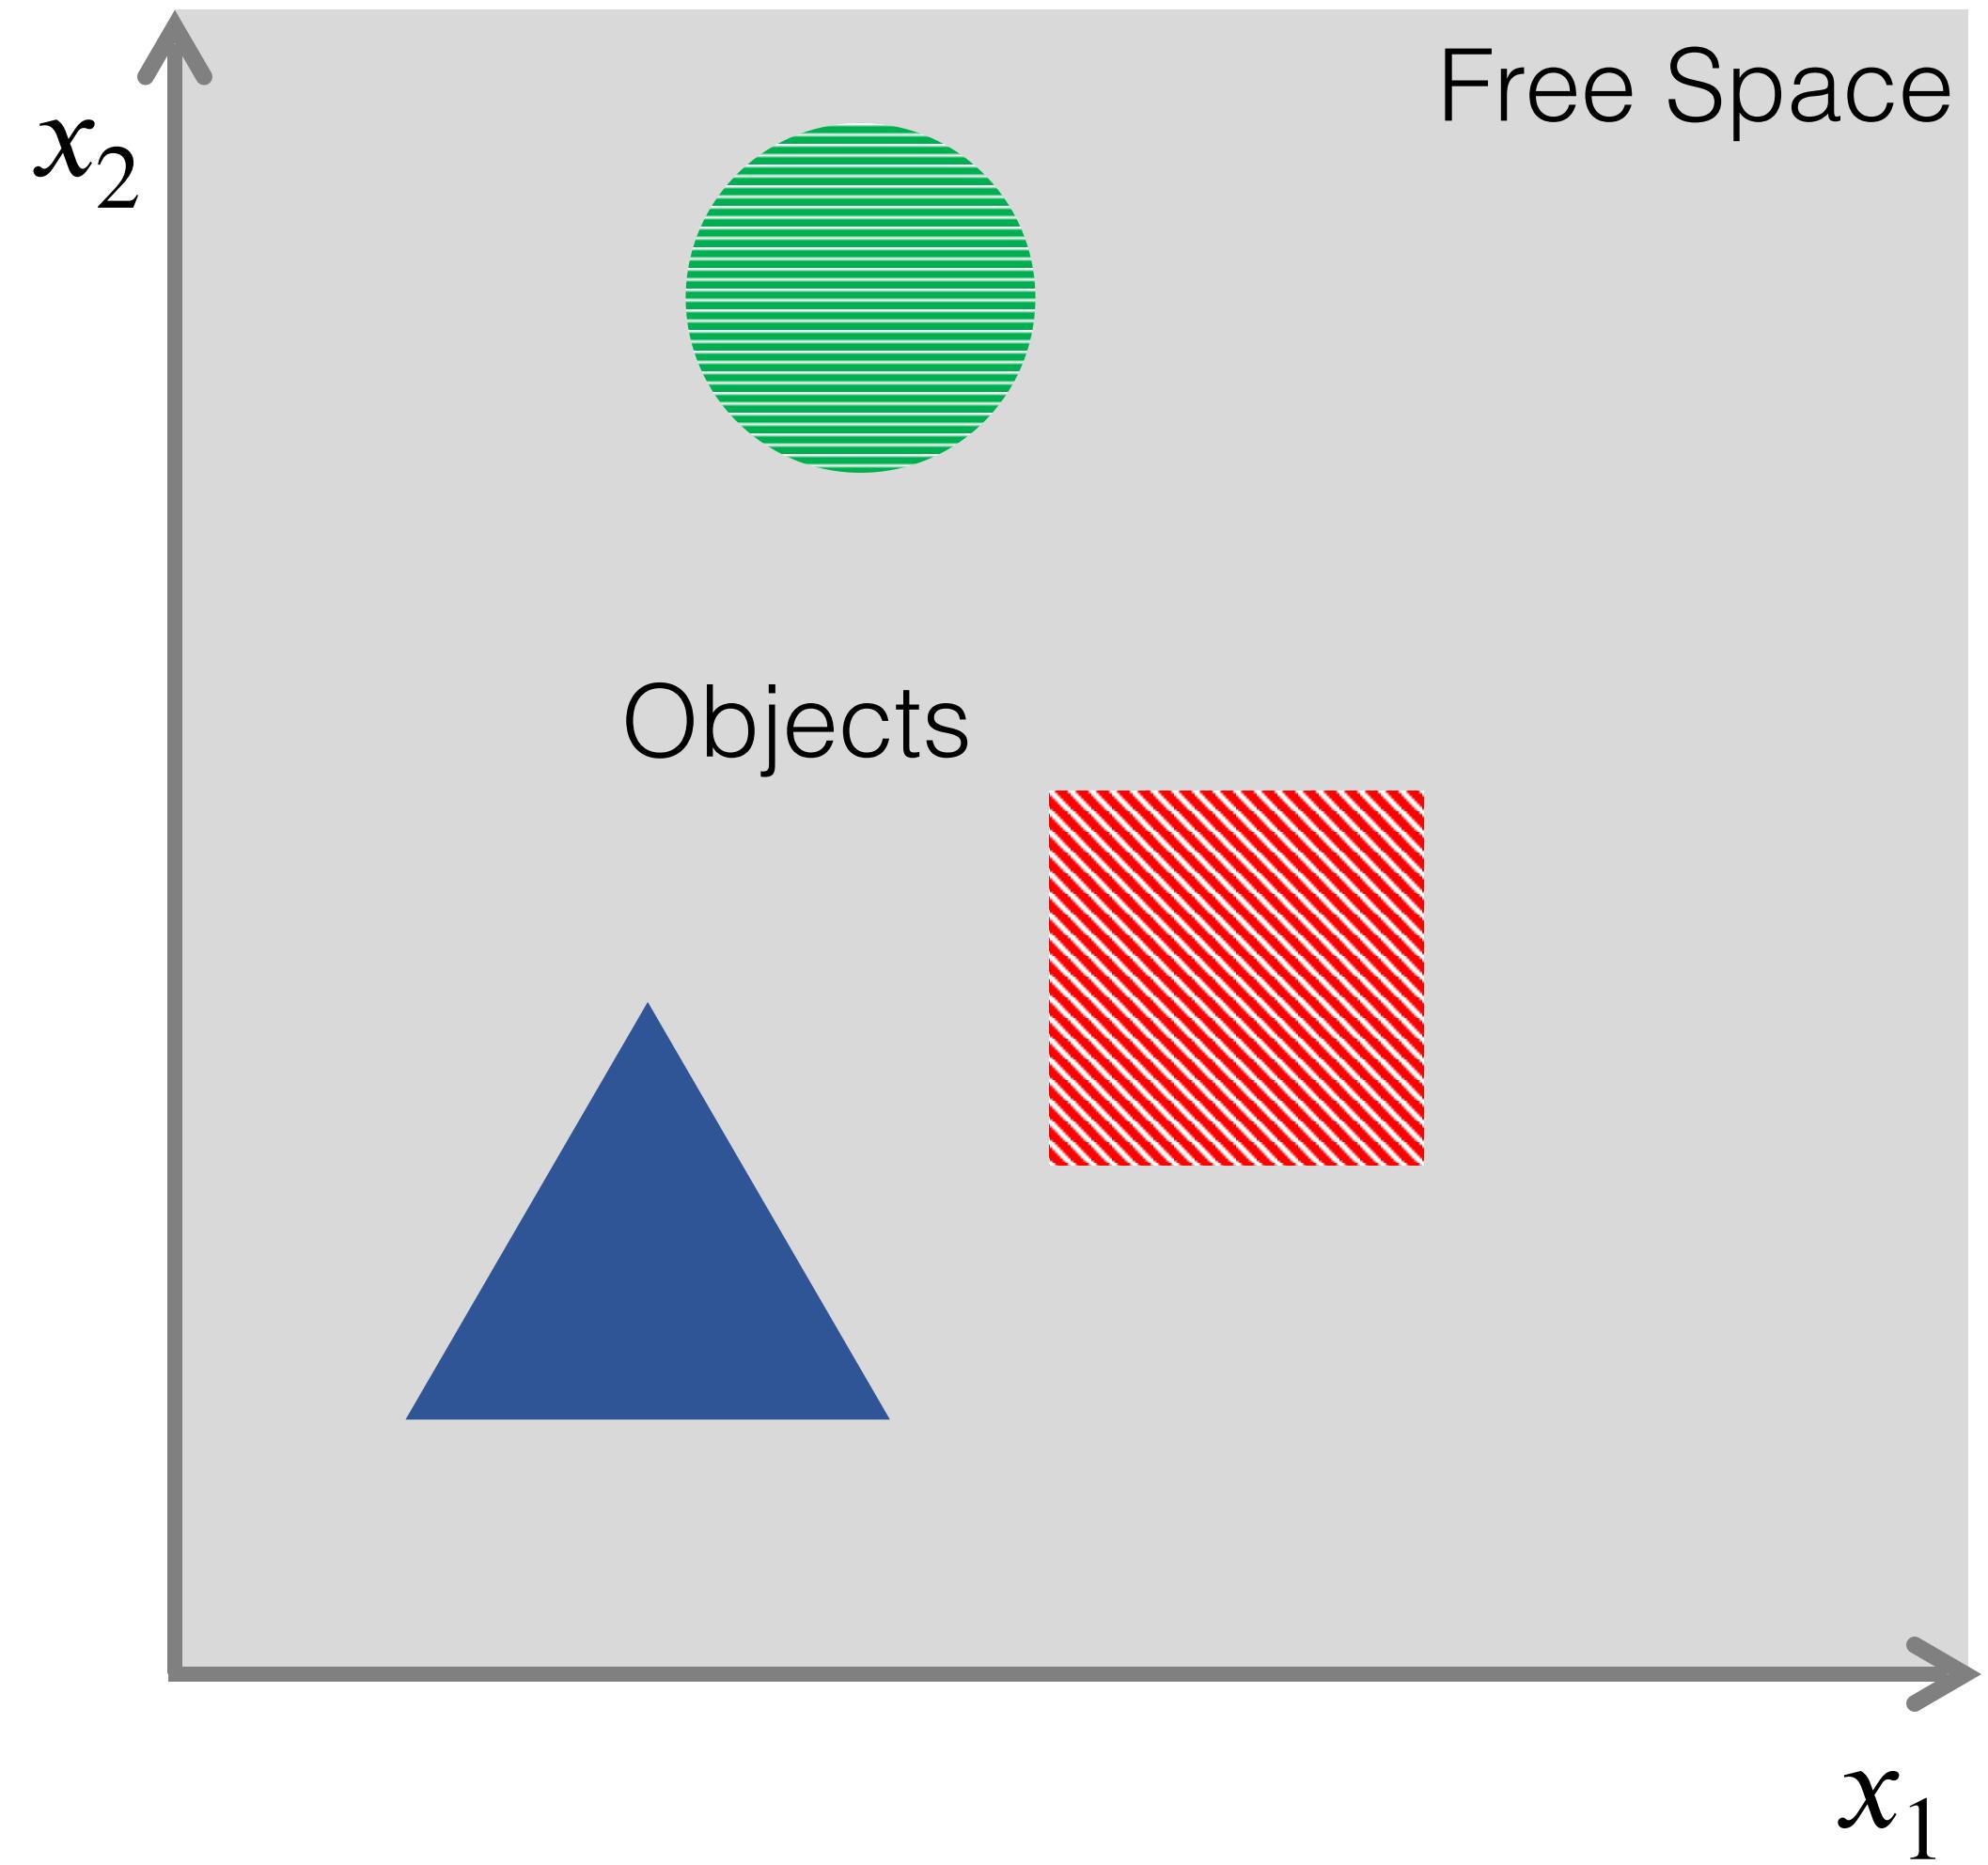
“

A physical quantity represented by a scalar, vector, or tensor, that has a value for each point in space and time.

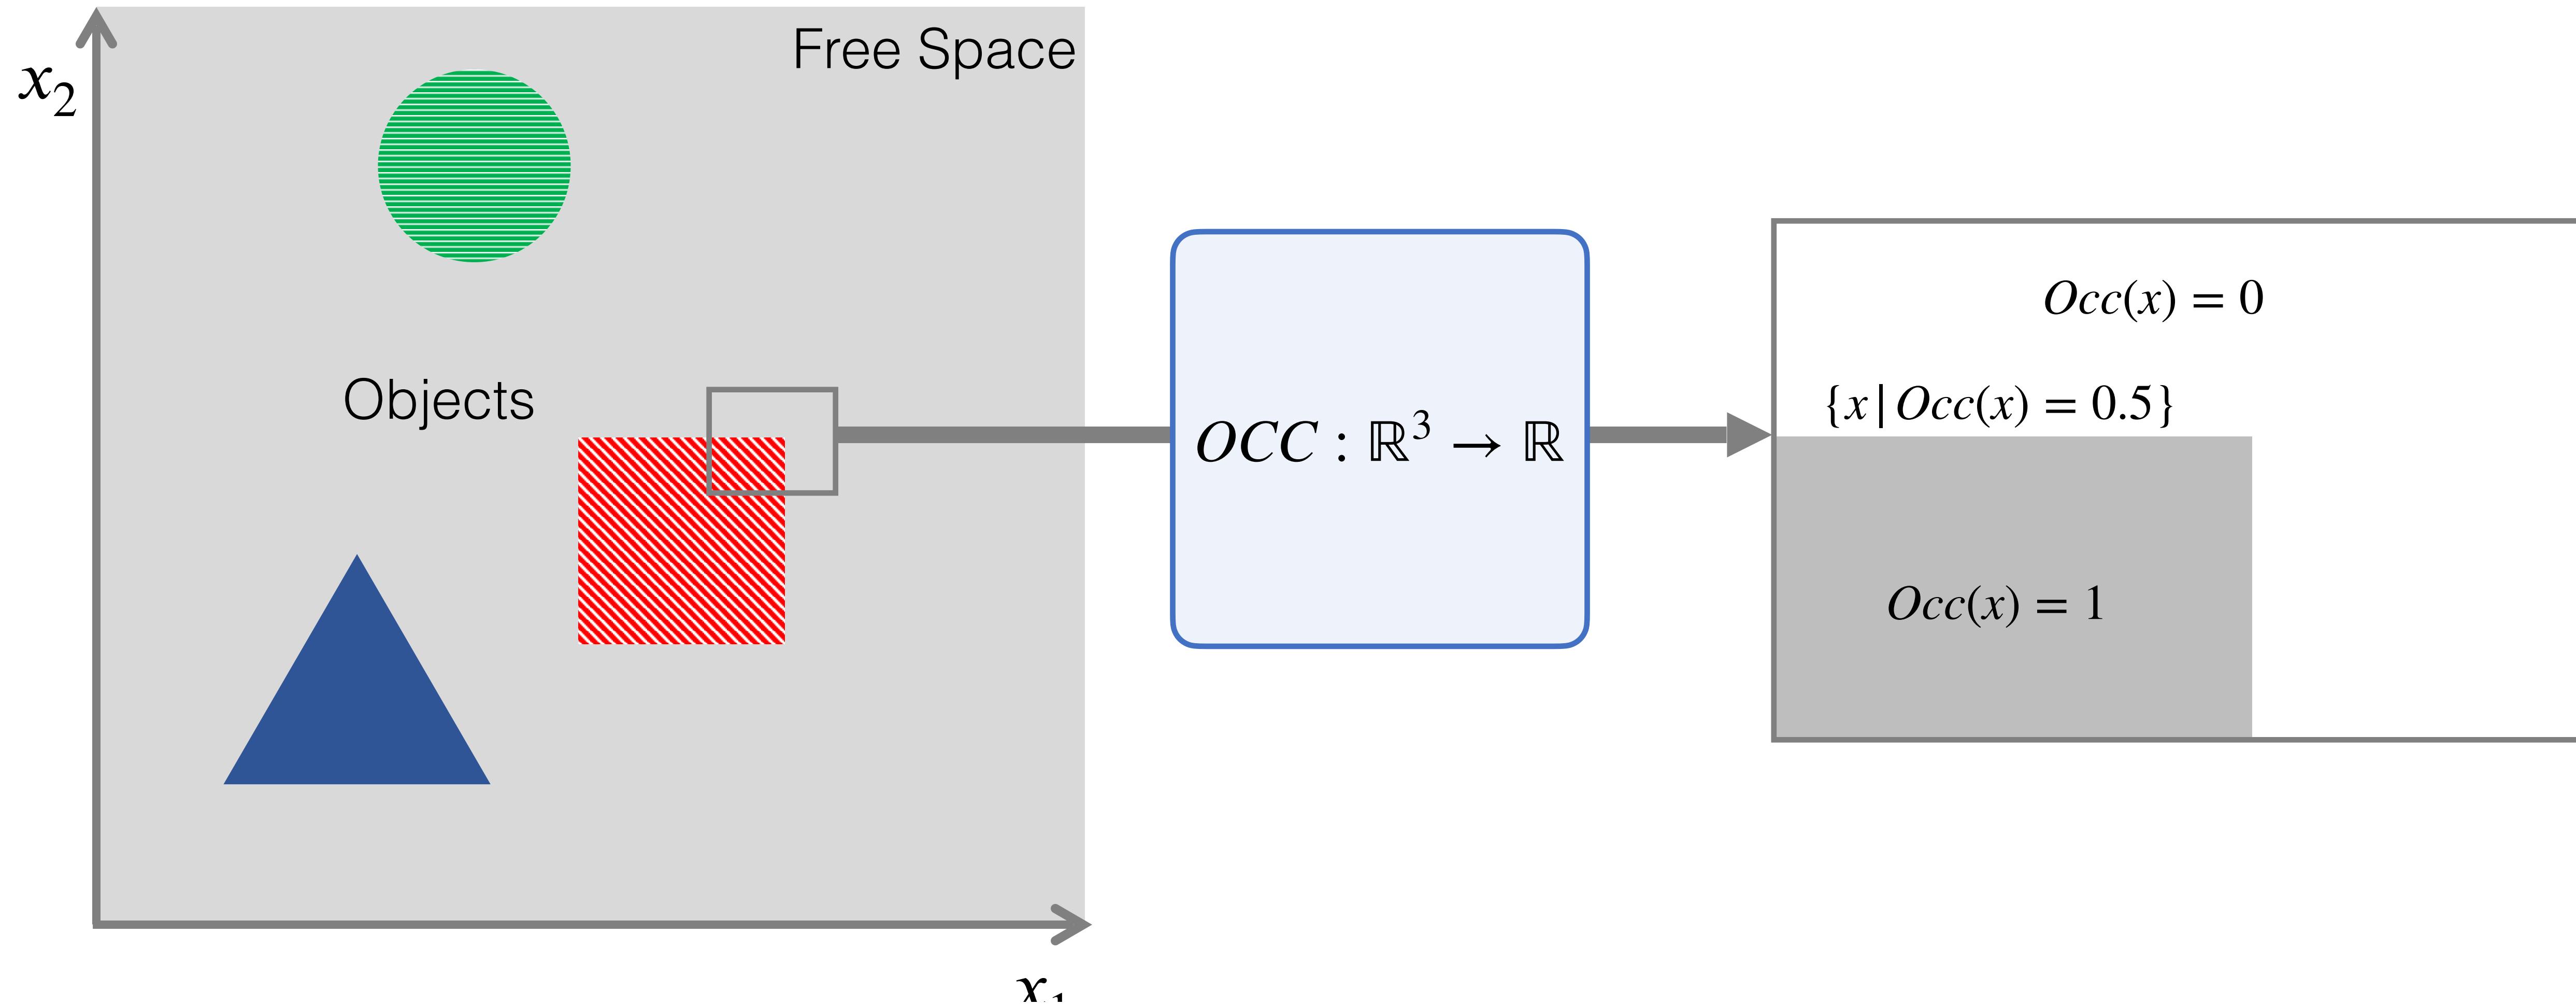
Will use term “field” for any function that takes a space, time, or space-time coordinate as input to map it to some quantity at that coordinate.

$$\Phi: \mathbb{R}^3 \rightarrow \mathbb{R}^n$$

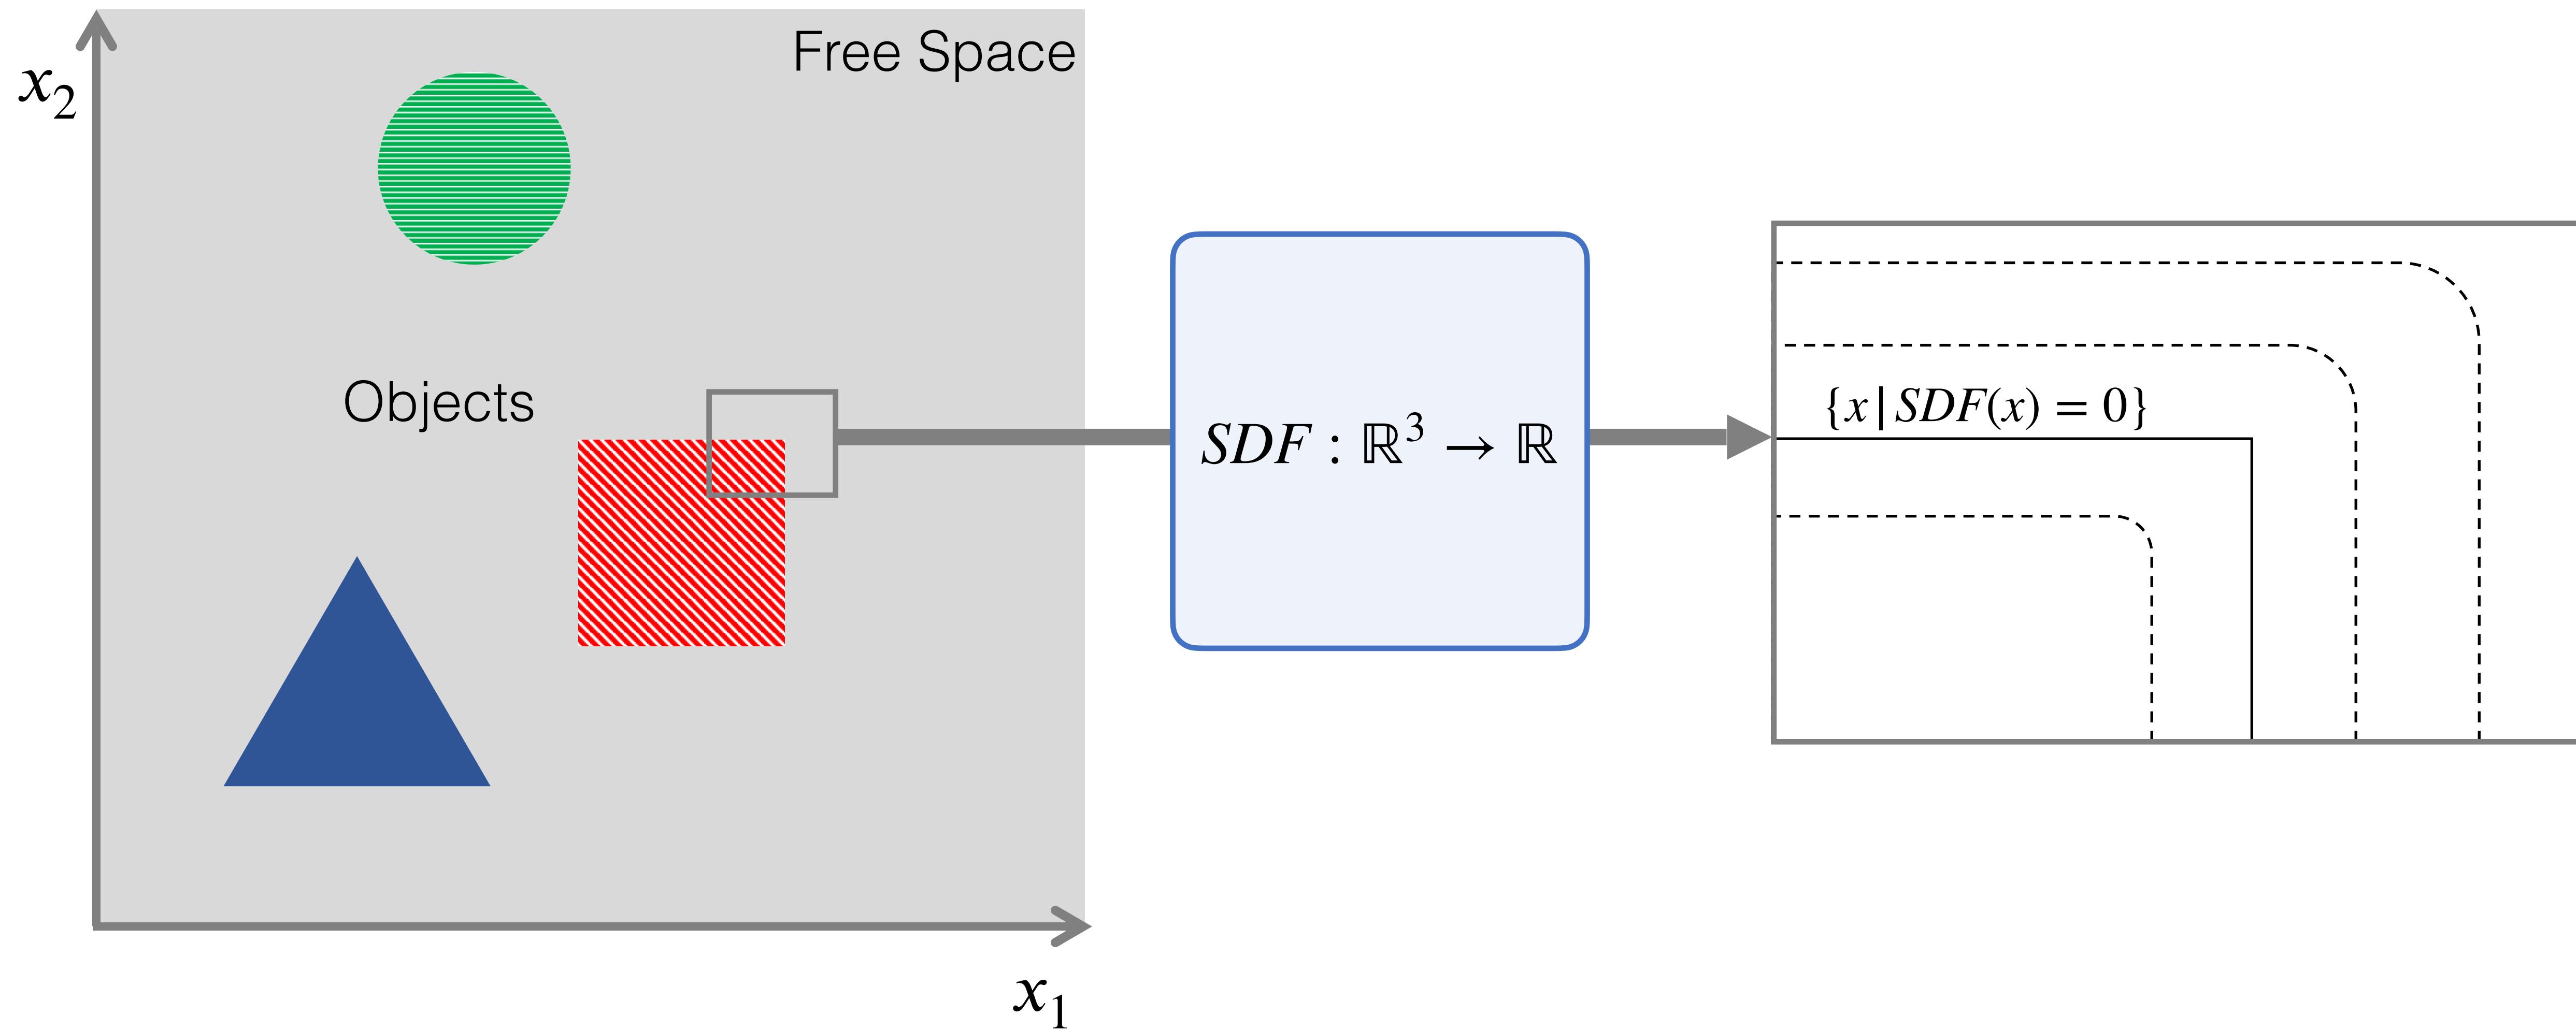
# How to parameterize surfaces?



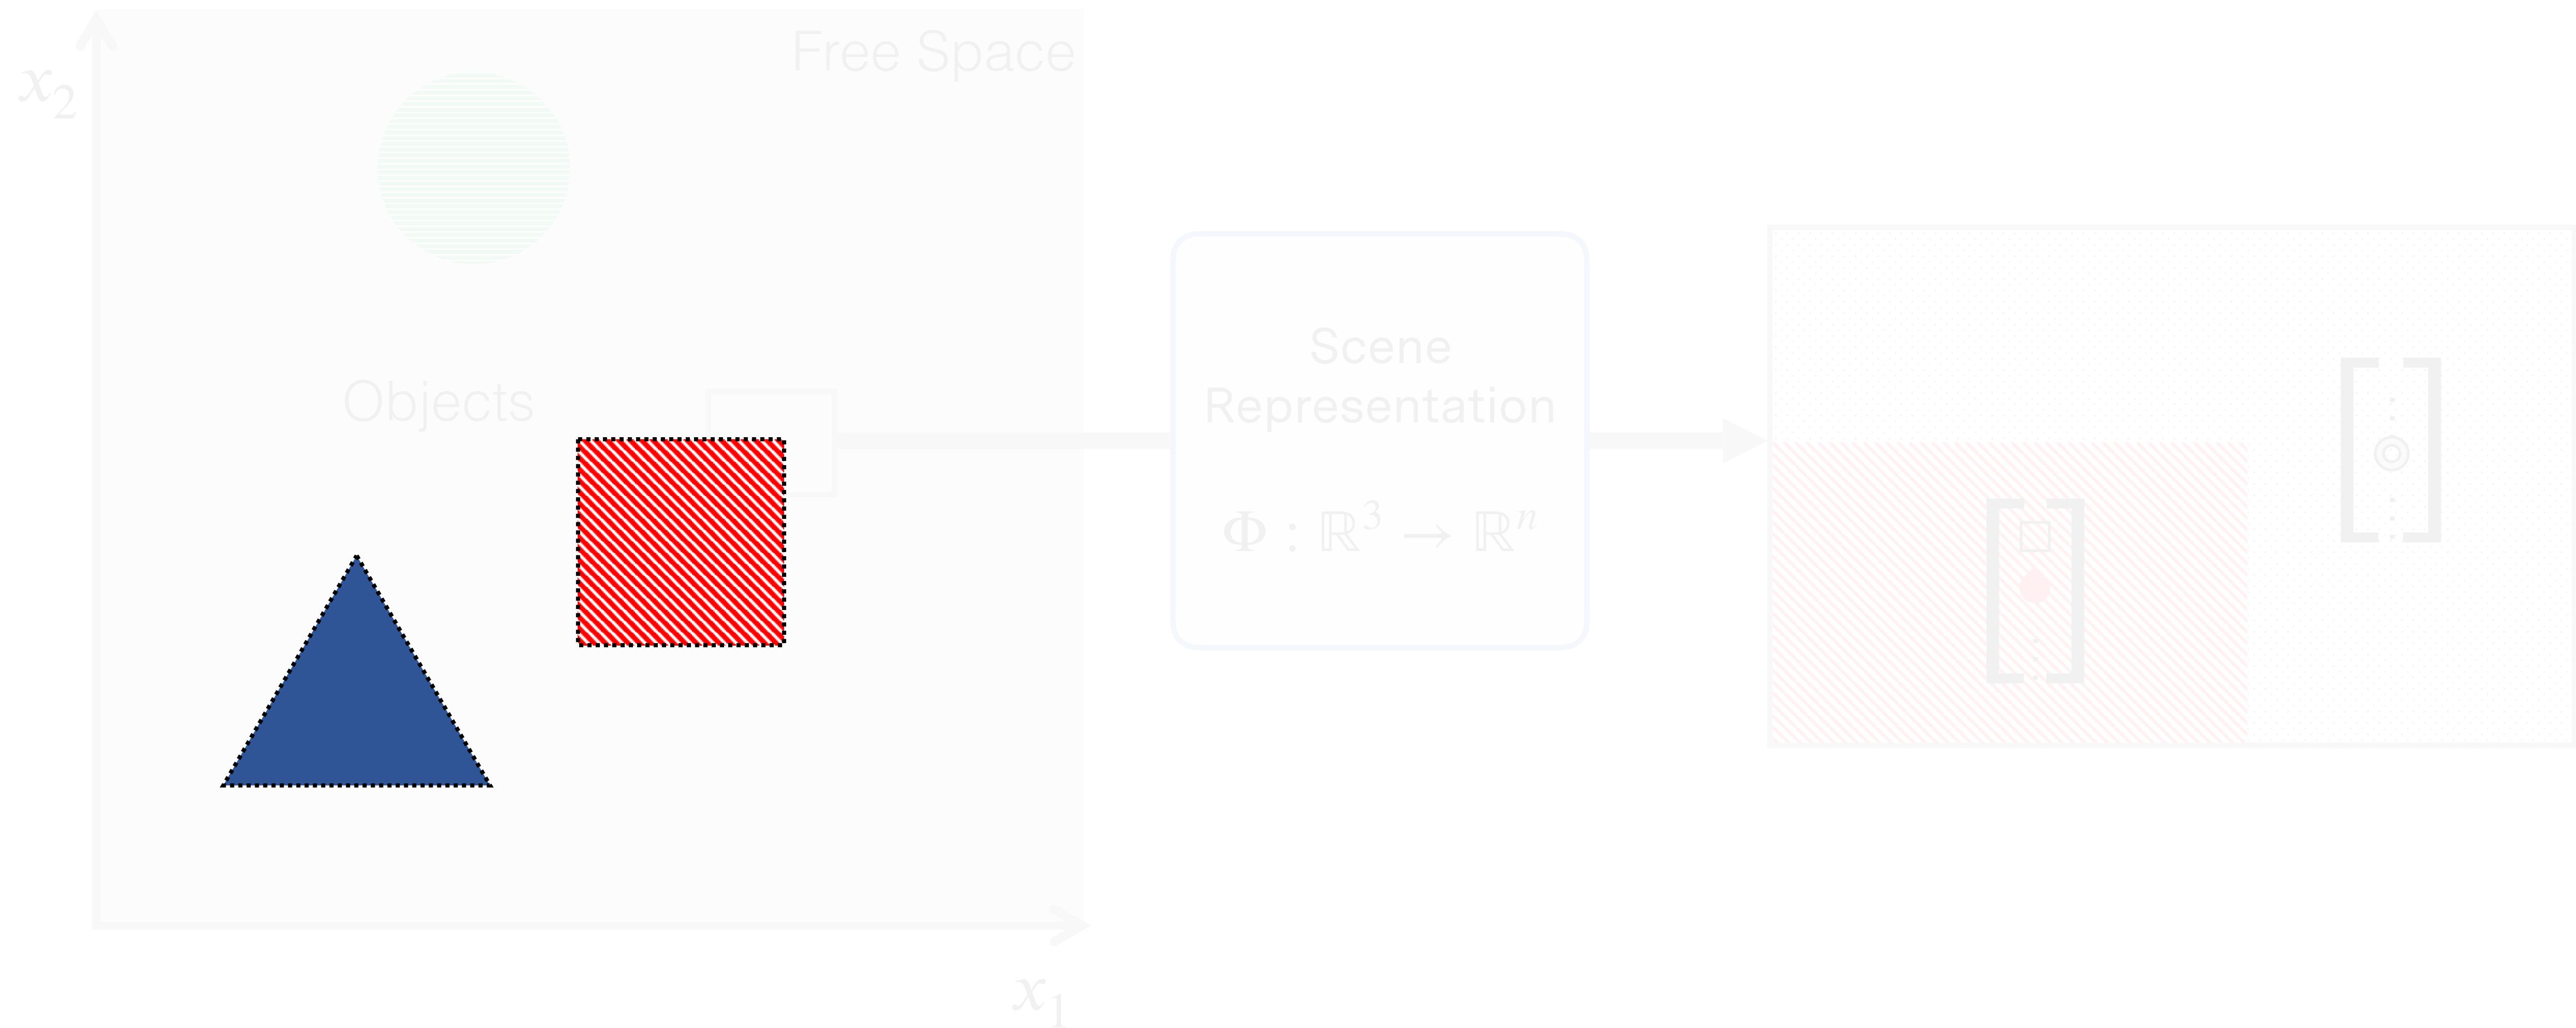
# The Occupancy Field



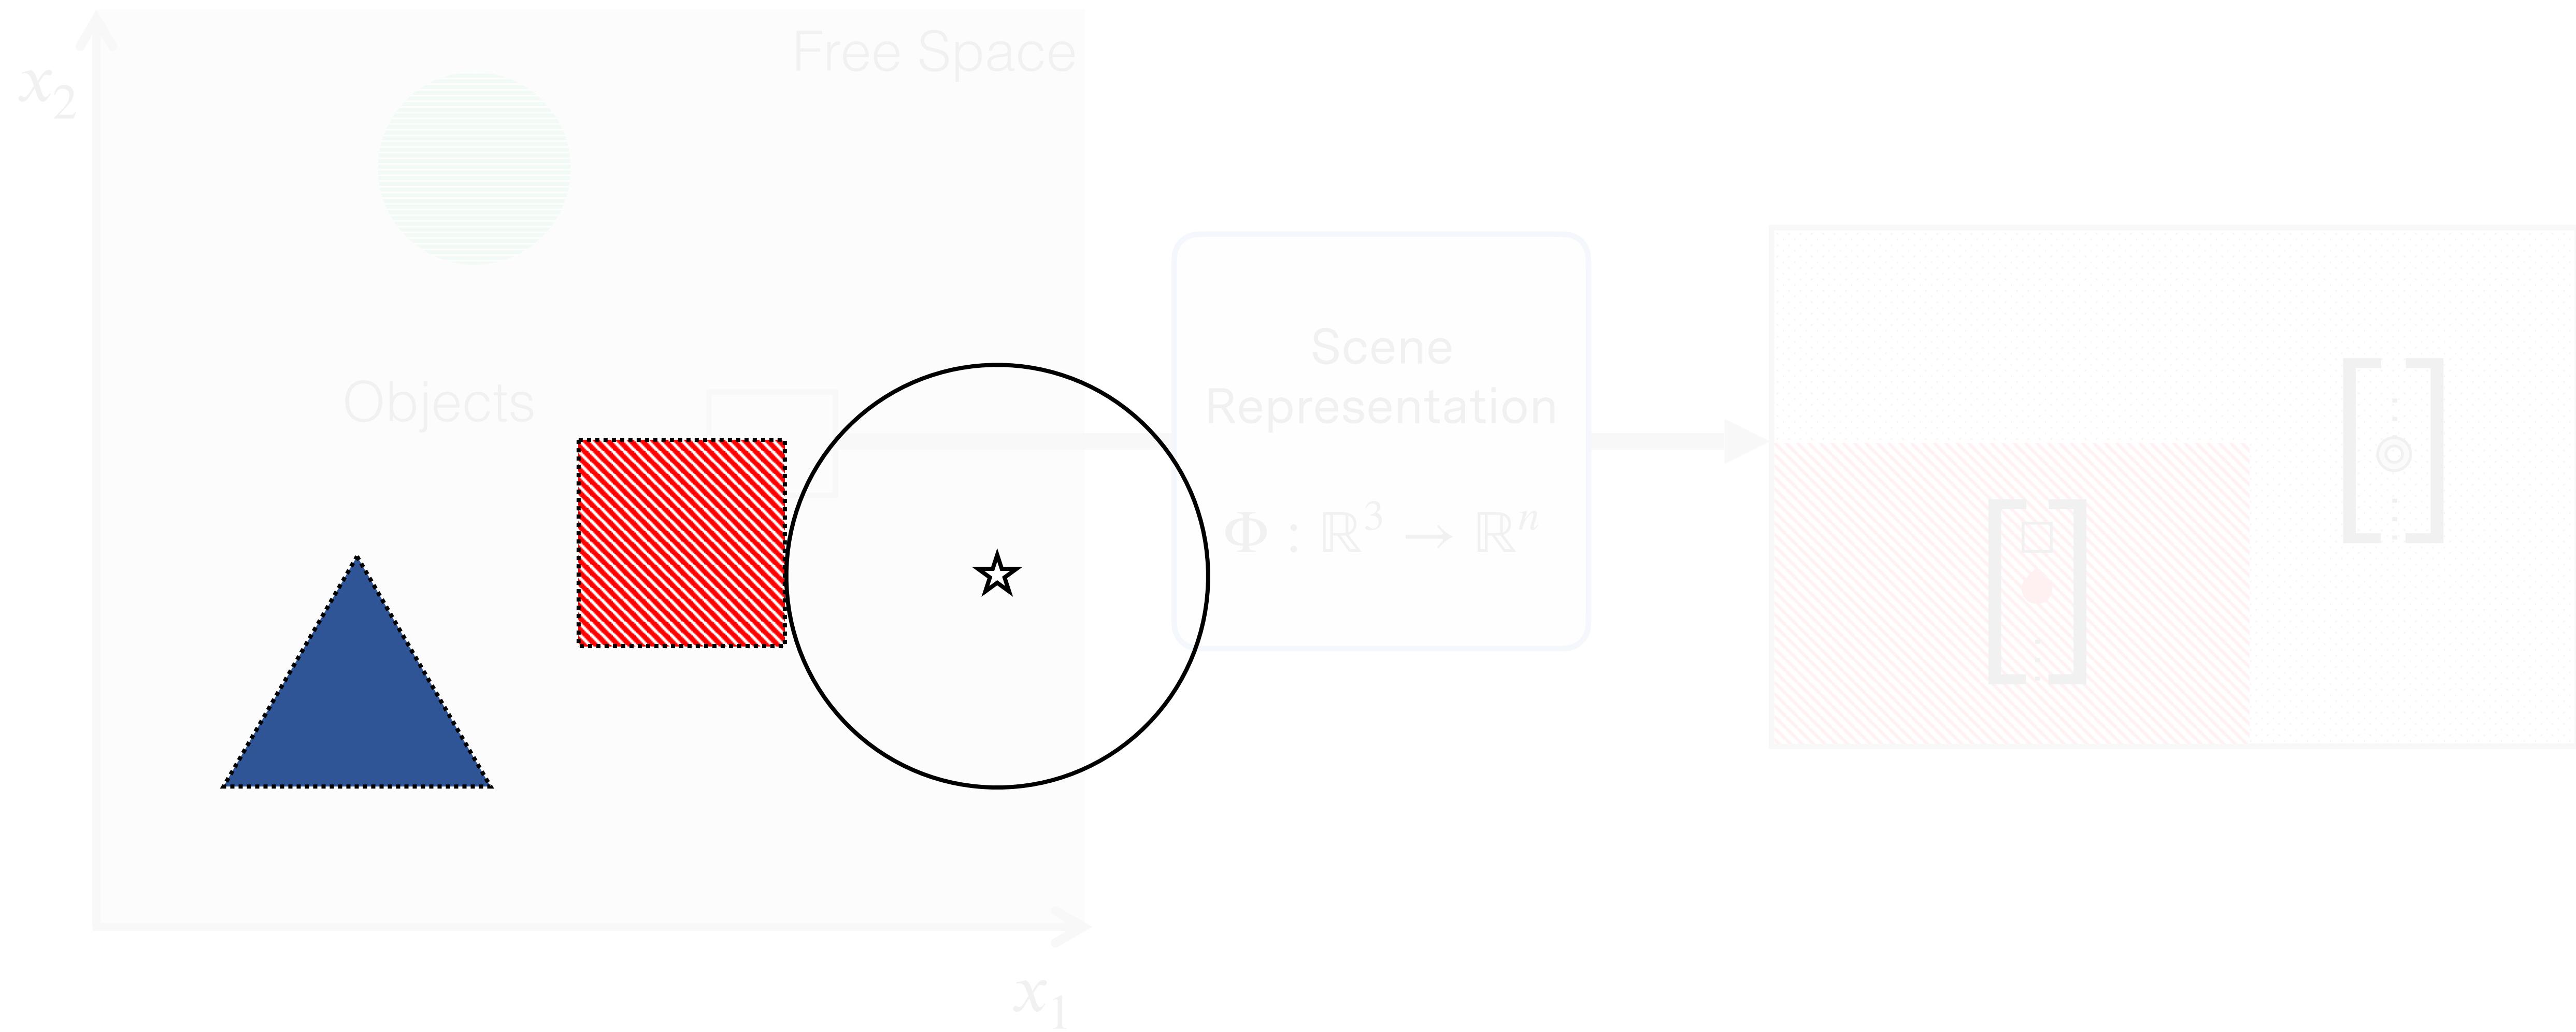
# The Signed Distance Field



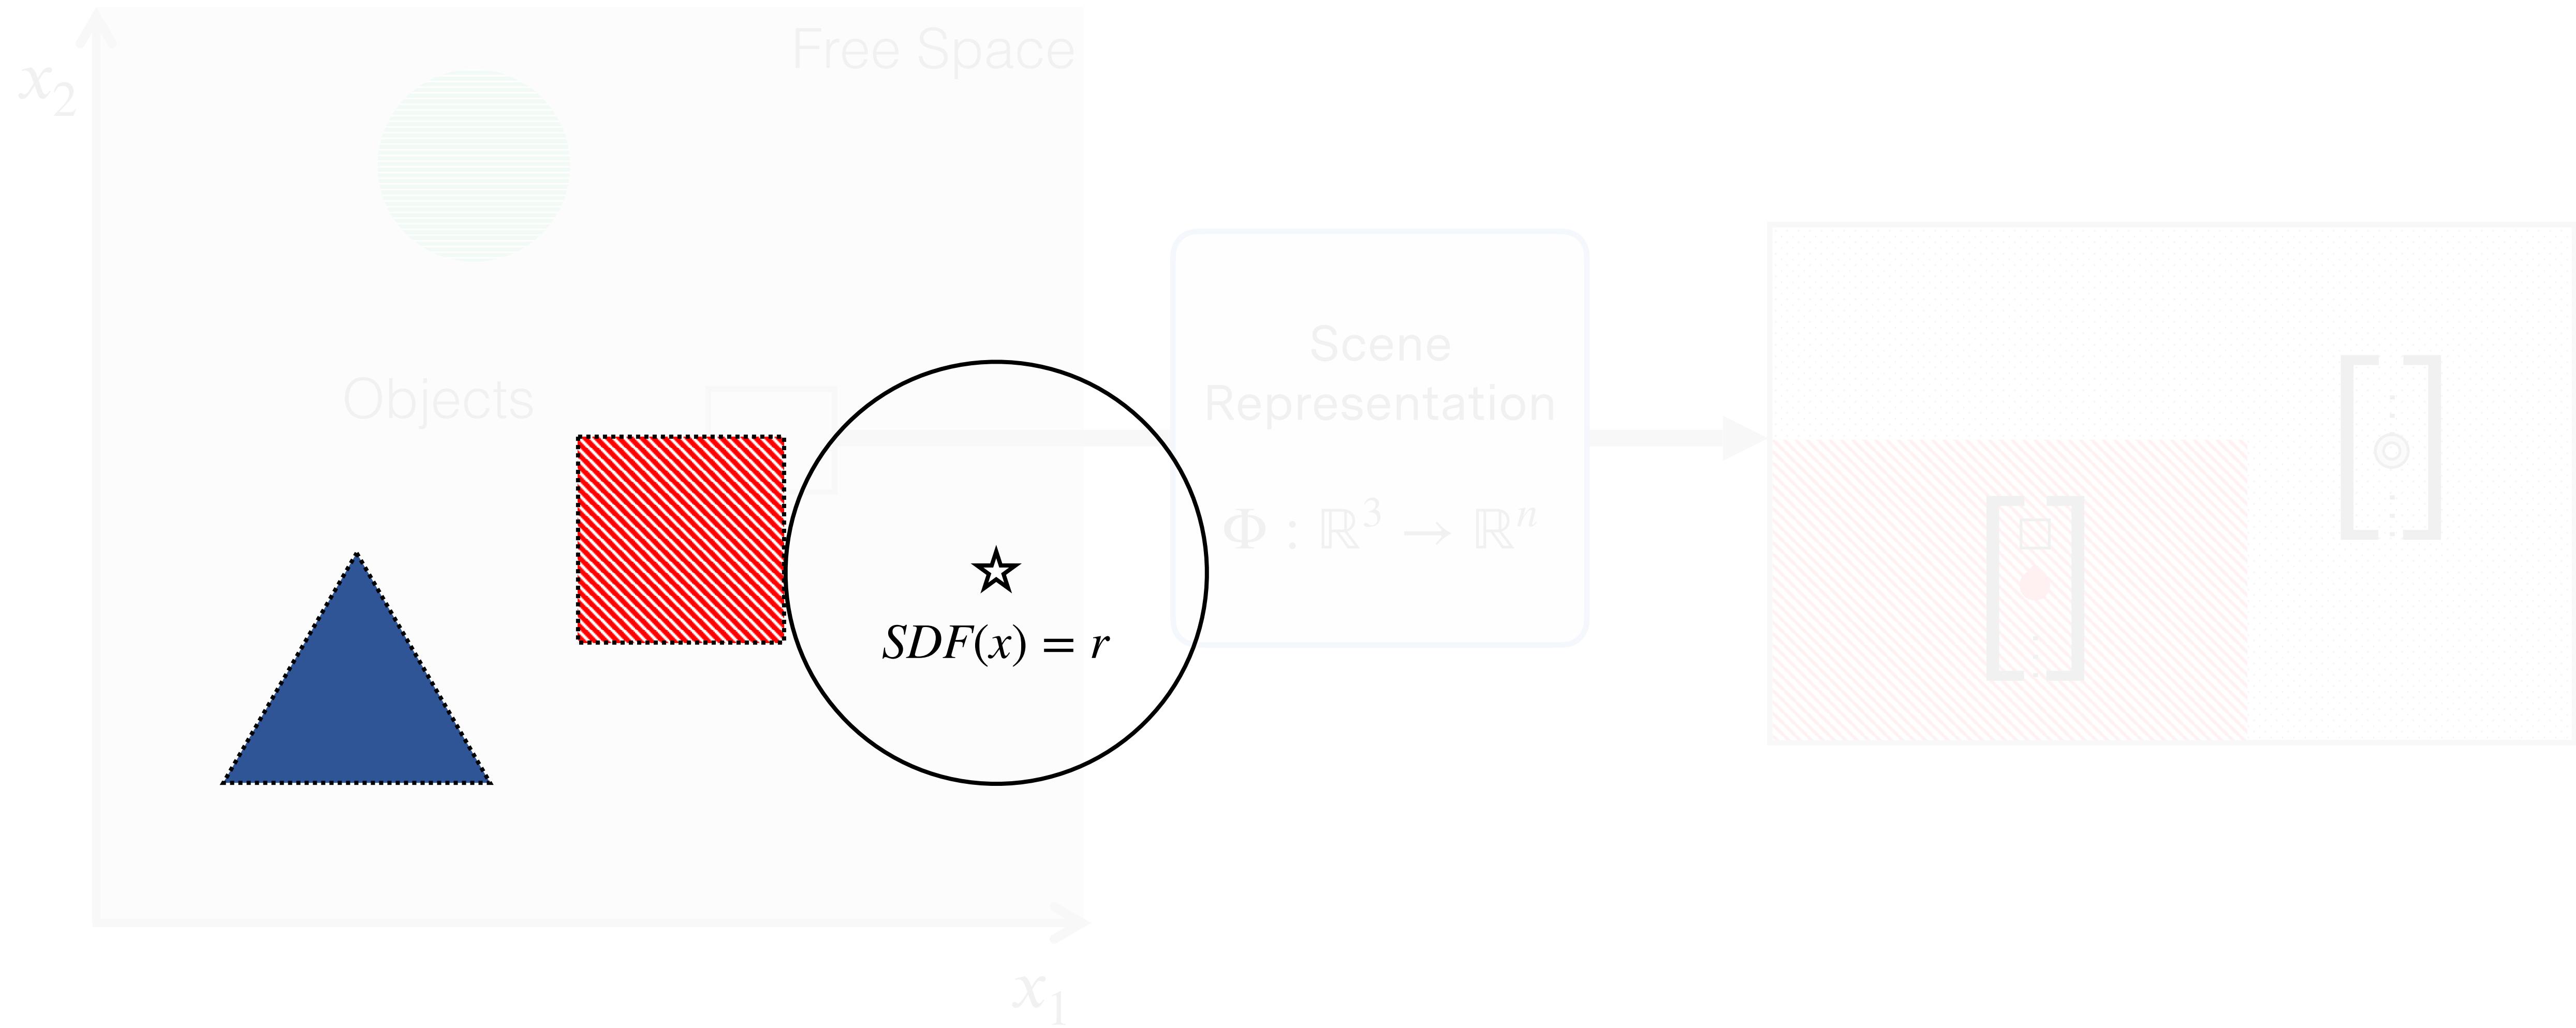
# The Signed Distance Field



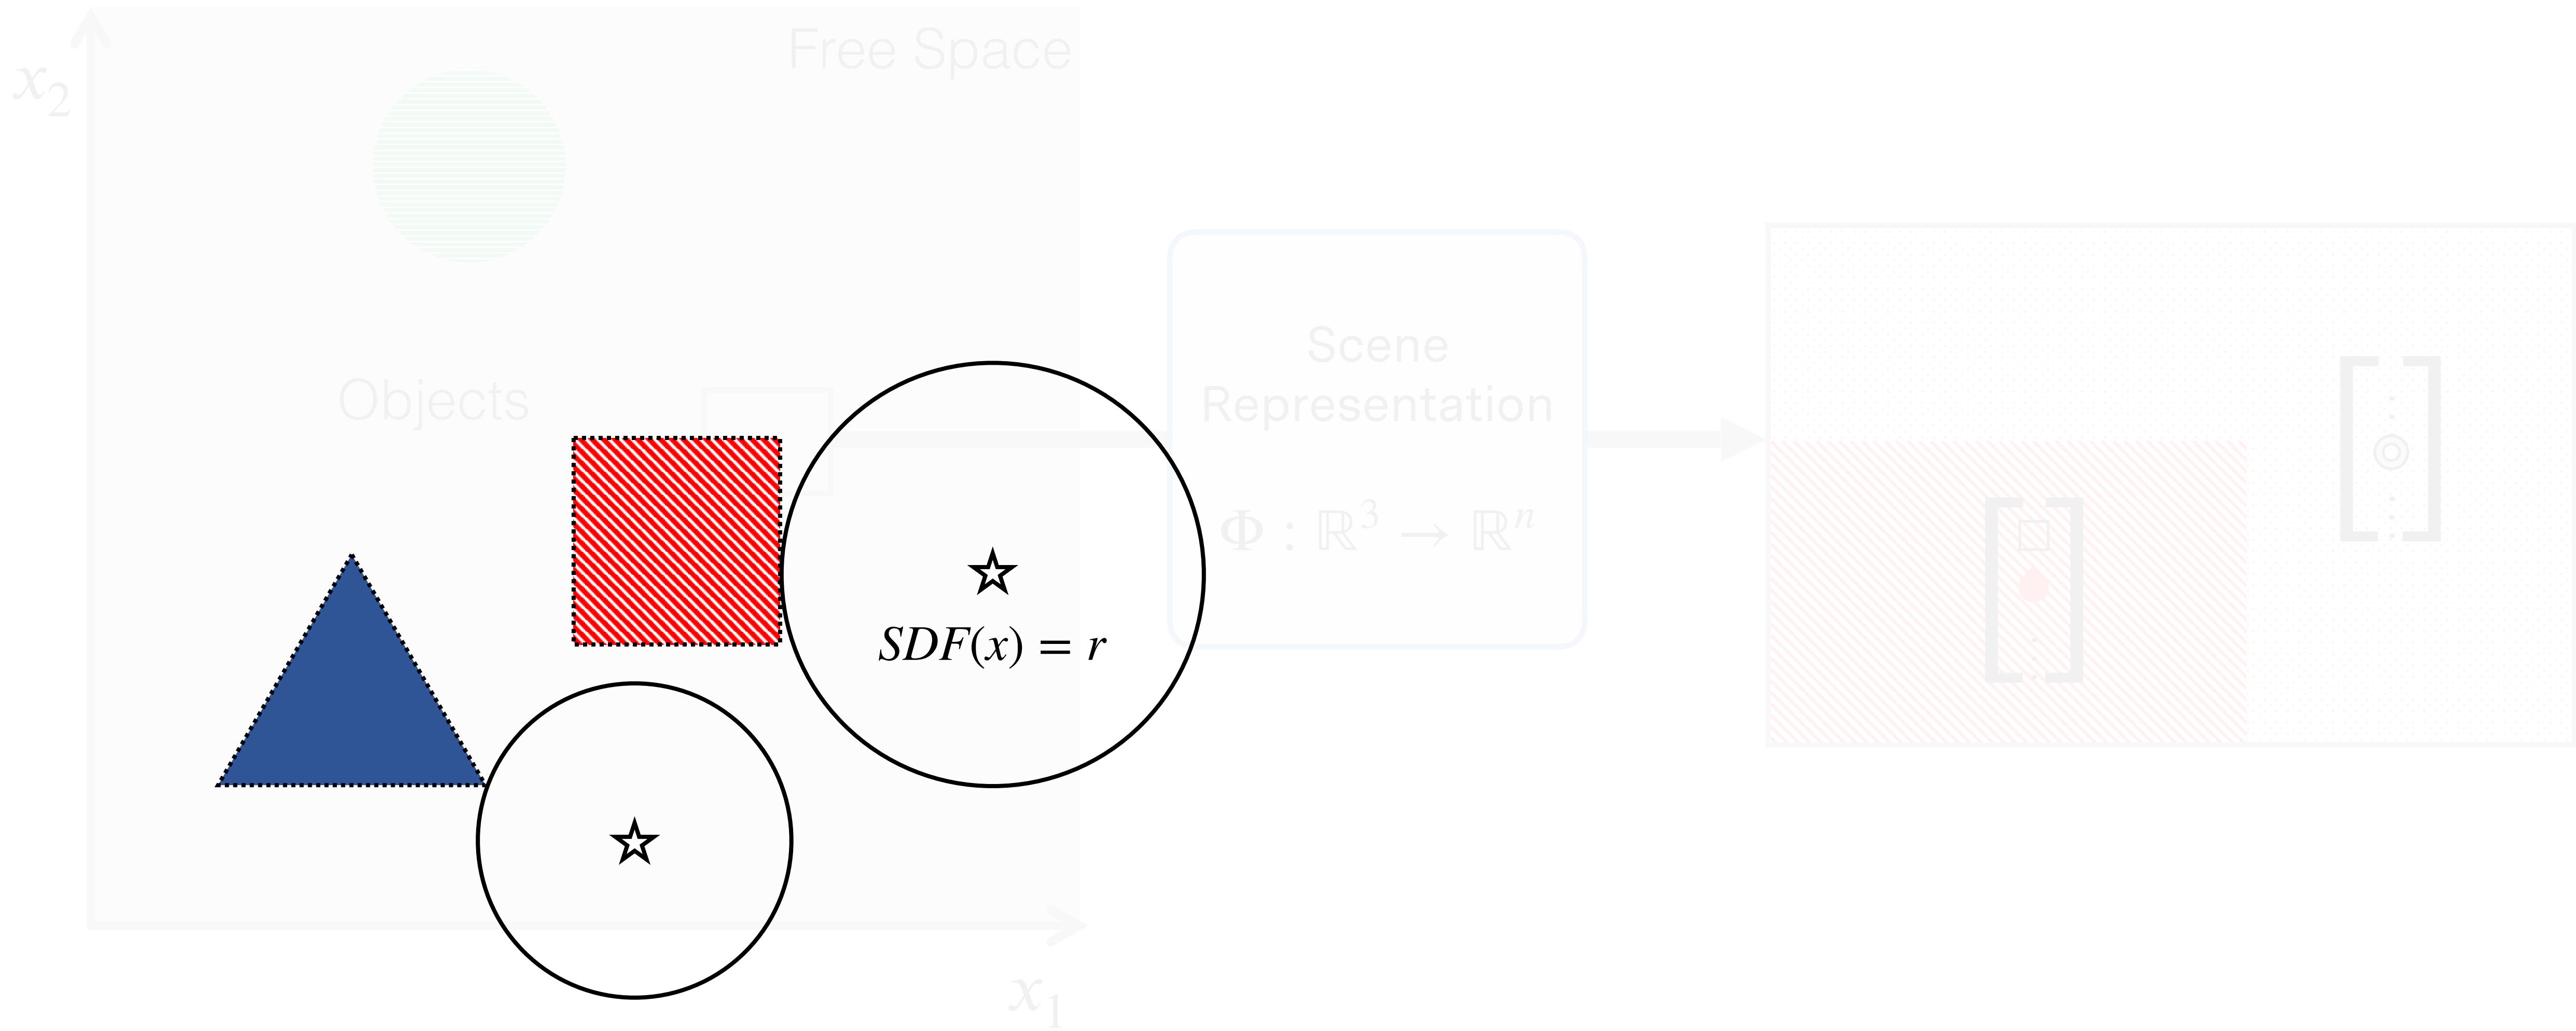
# The Signed Distance Field



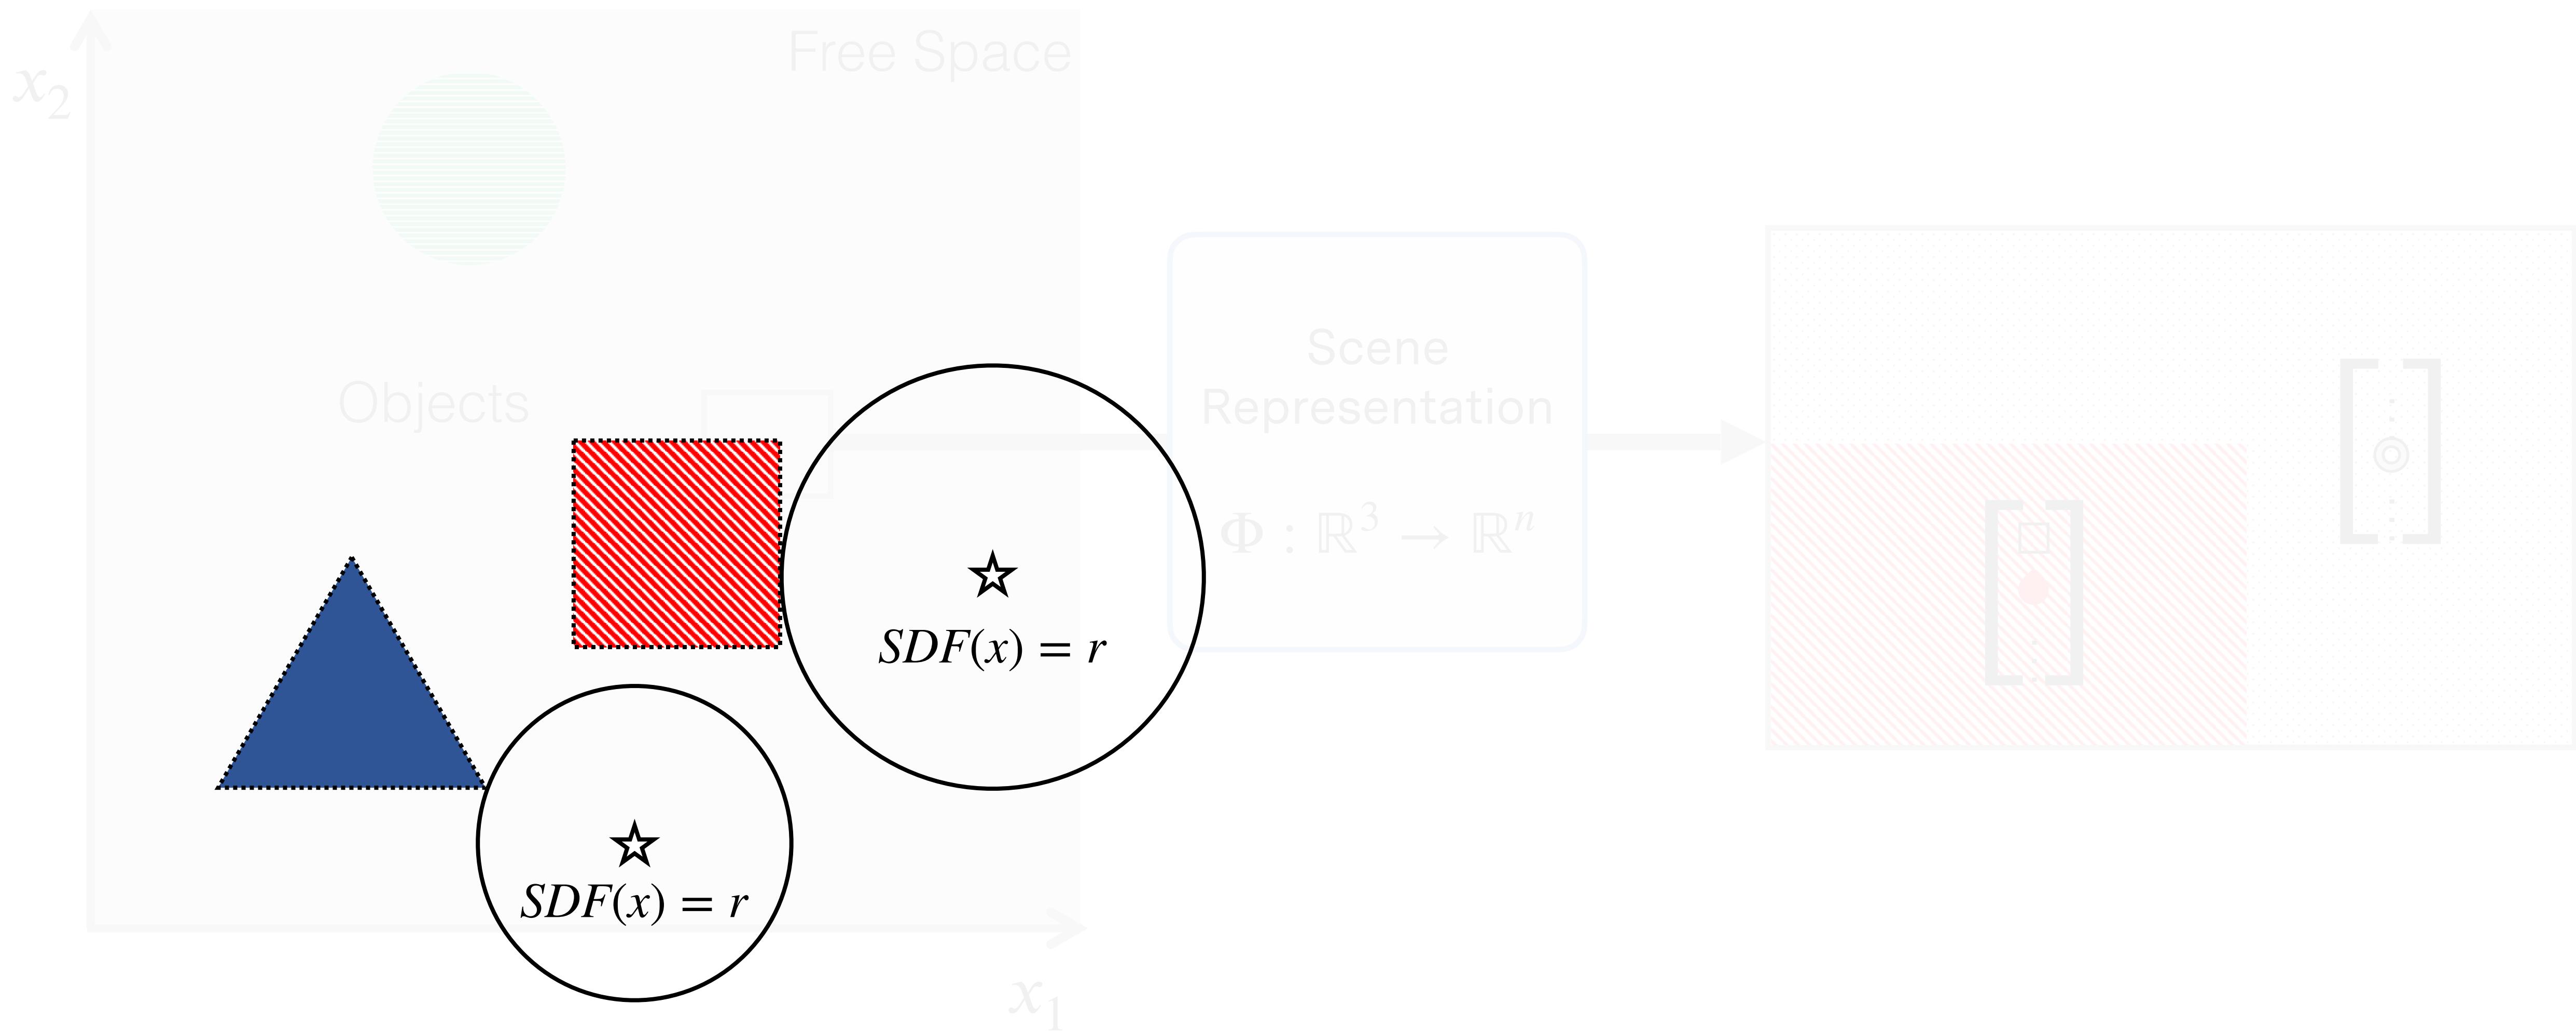
# The Signed Distance Field



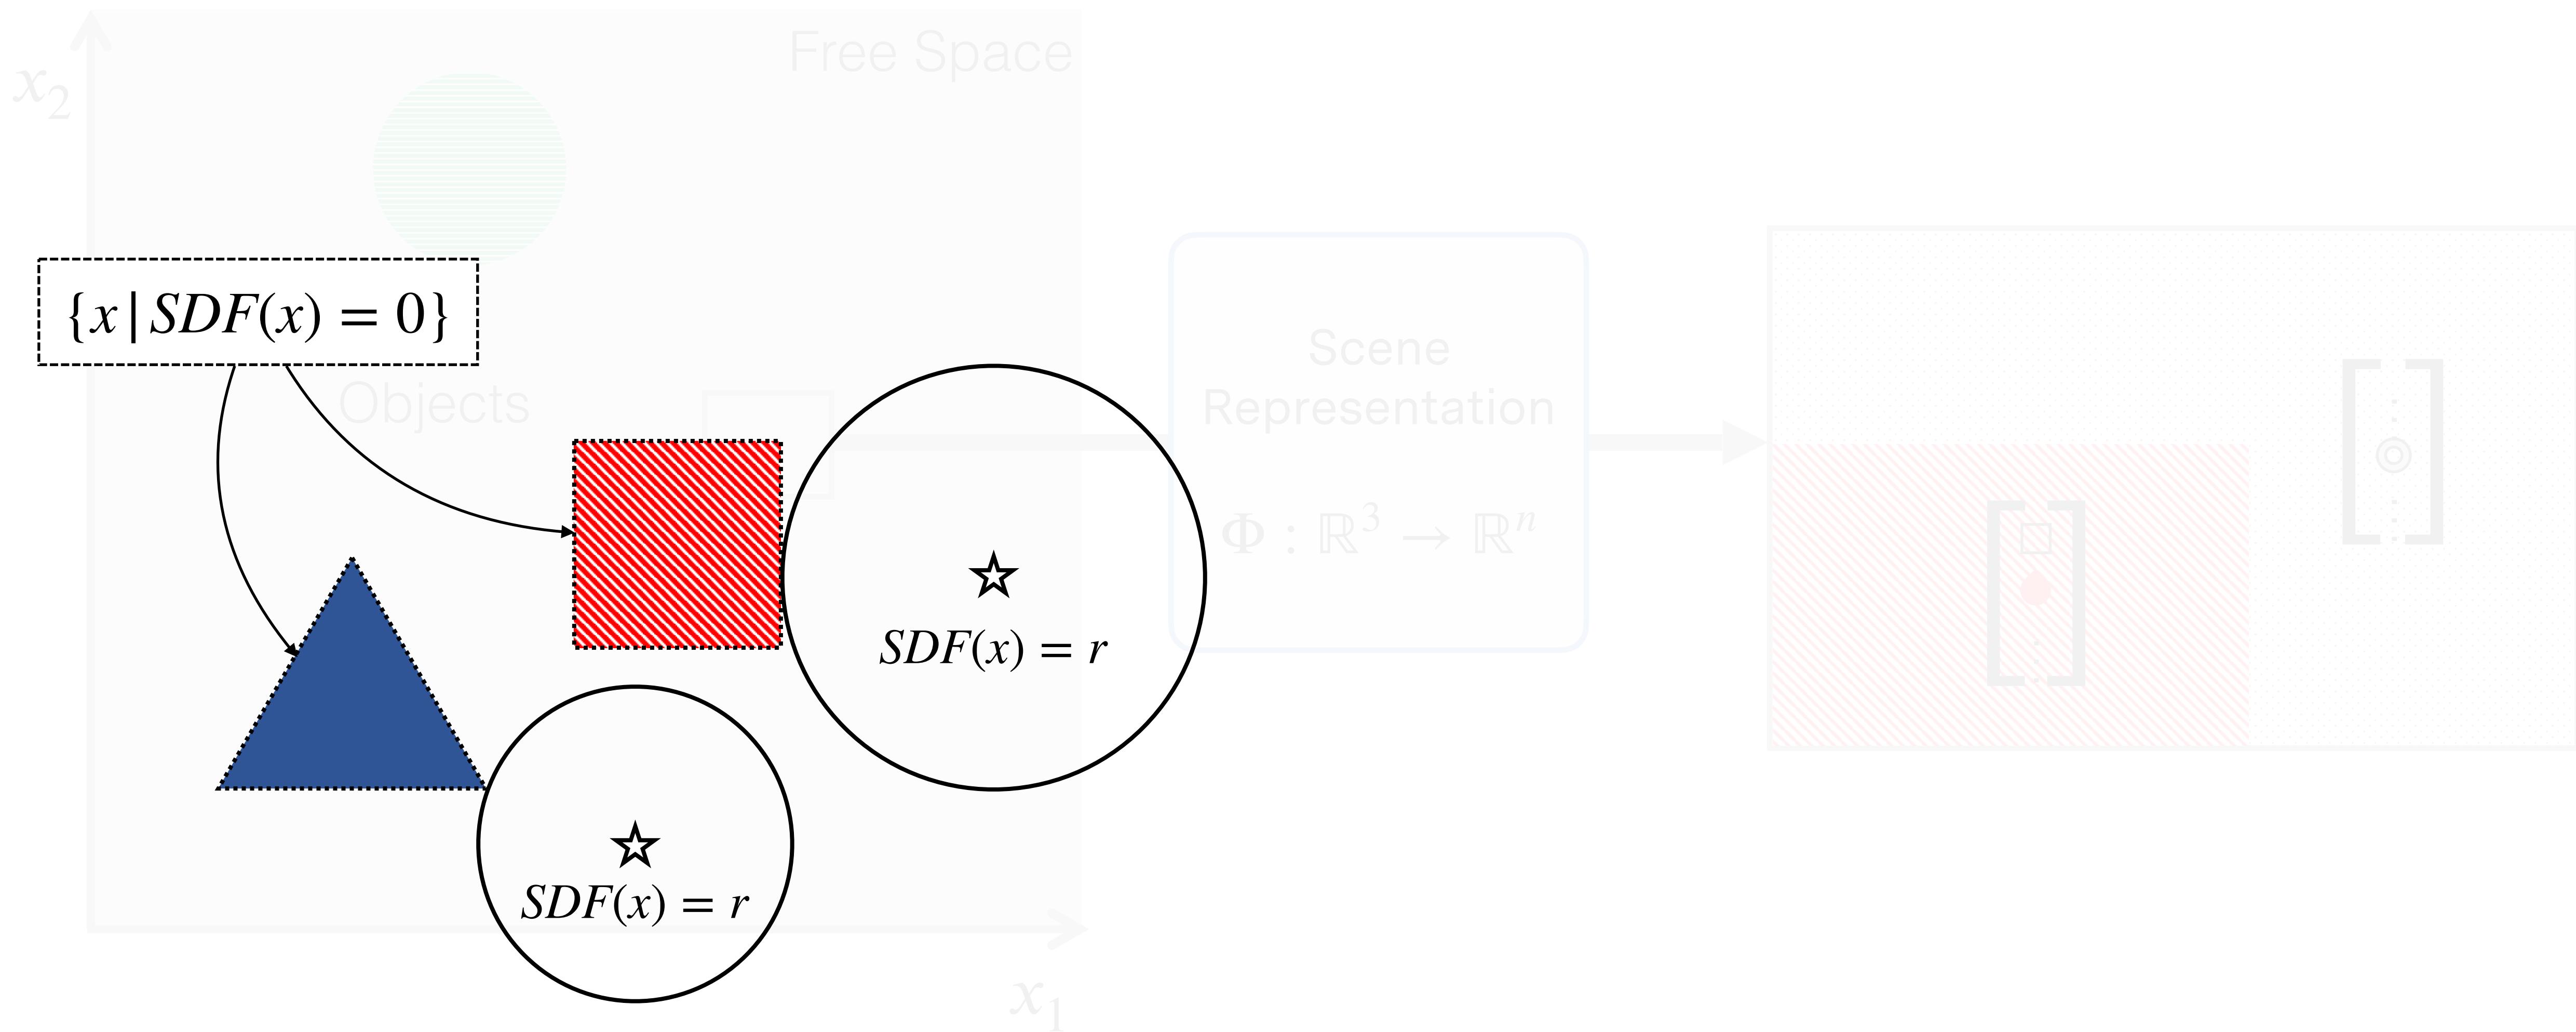
# The Signed Distance Field



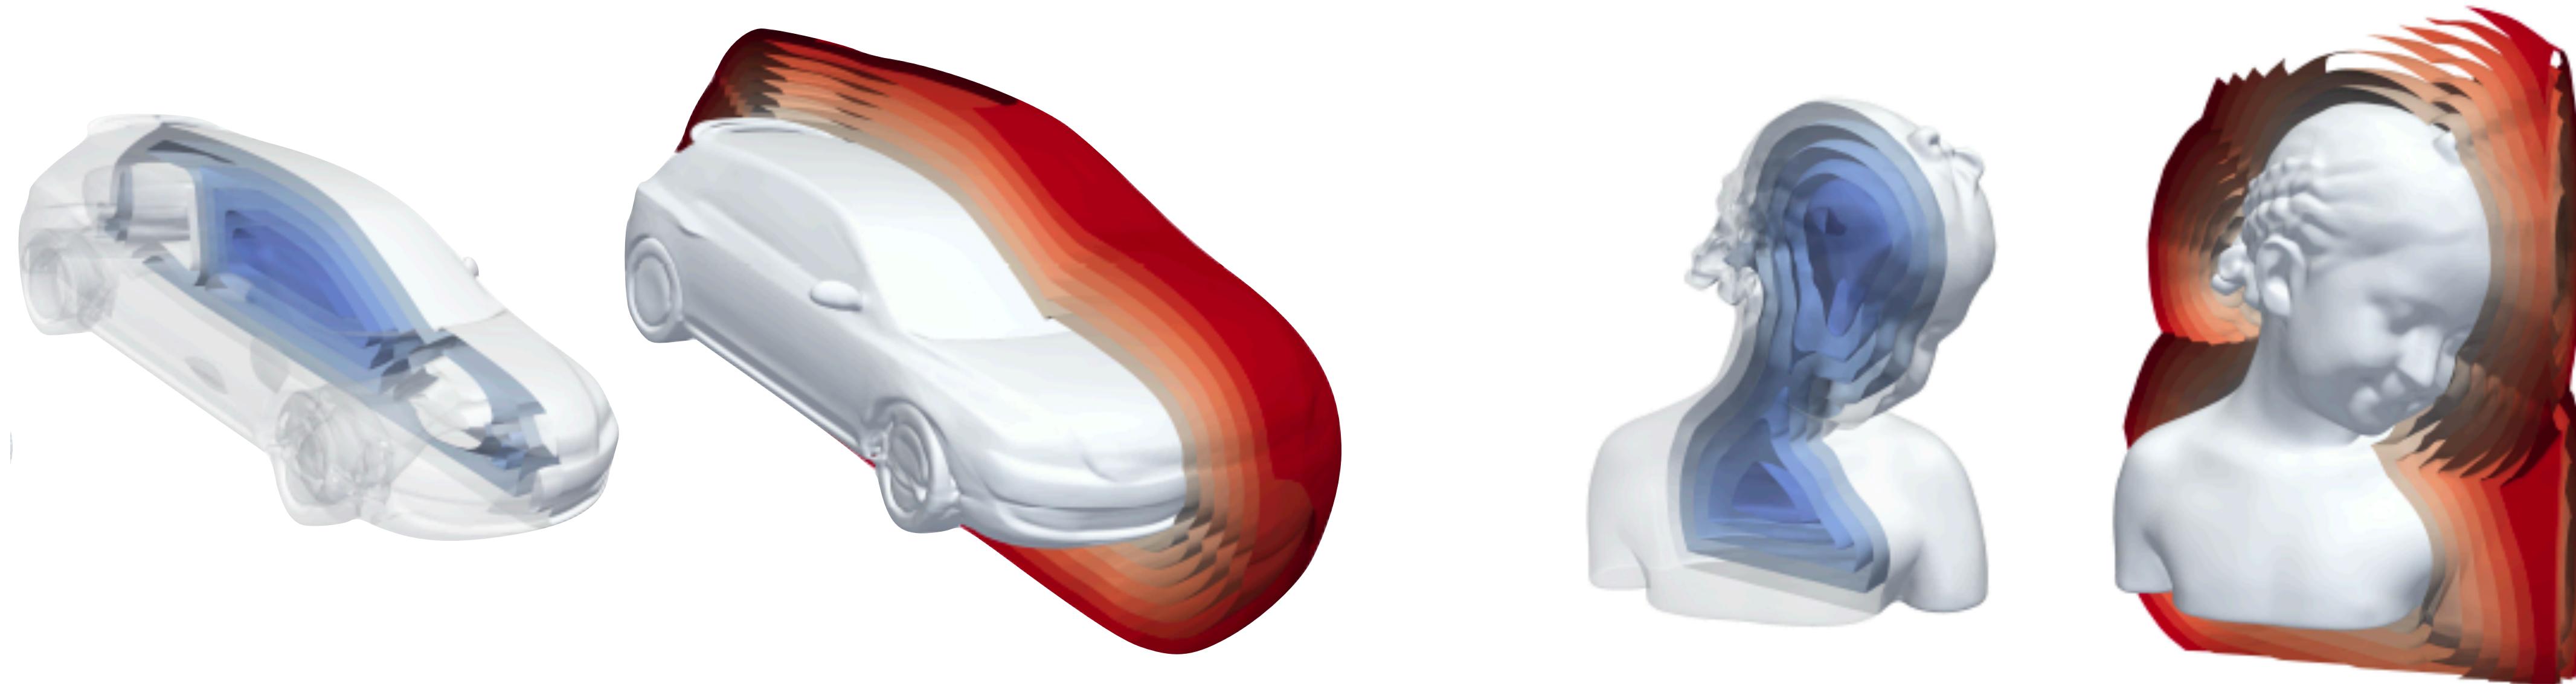
# The Signed Distance Field



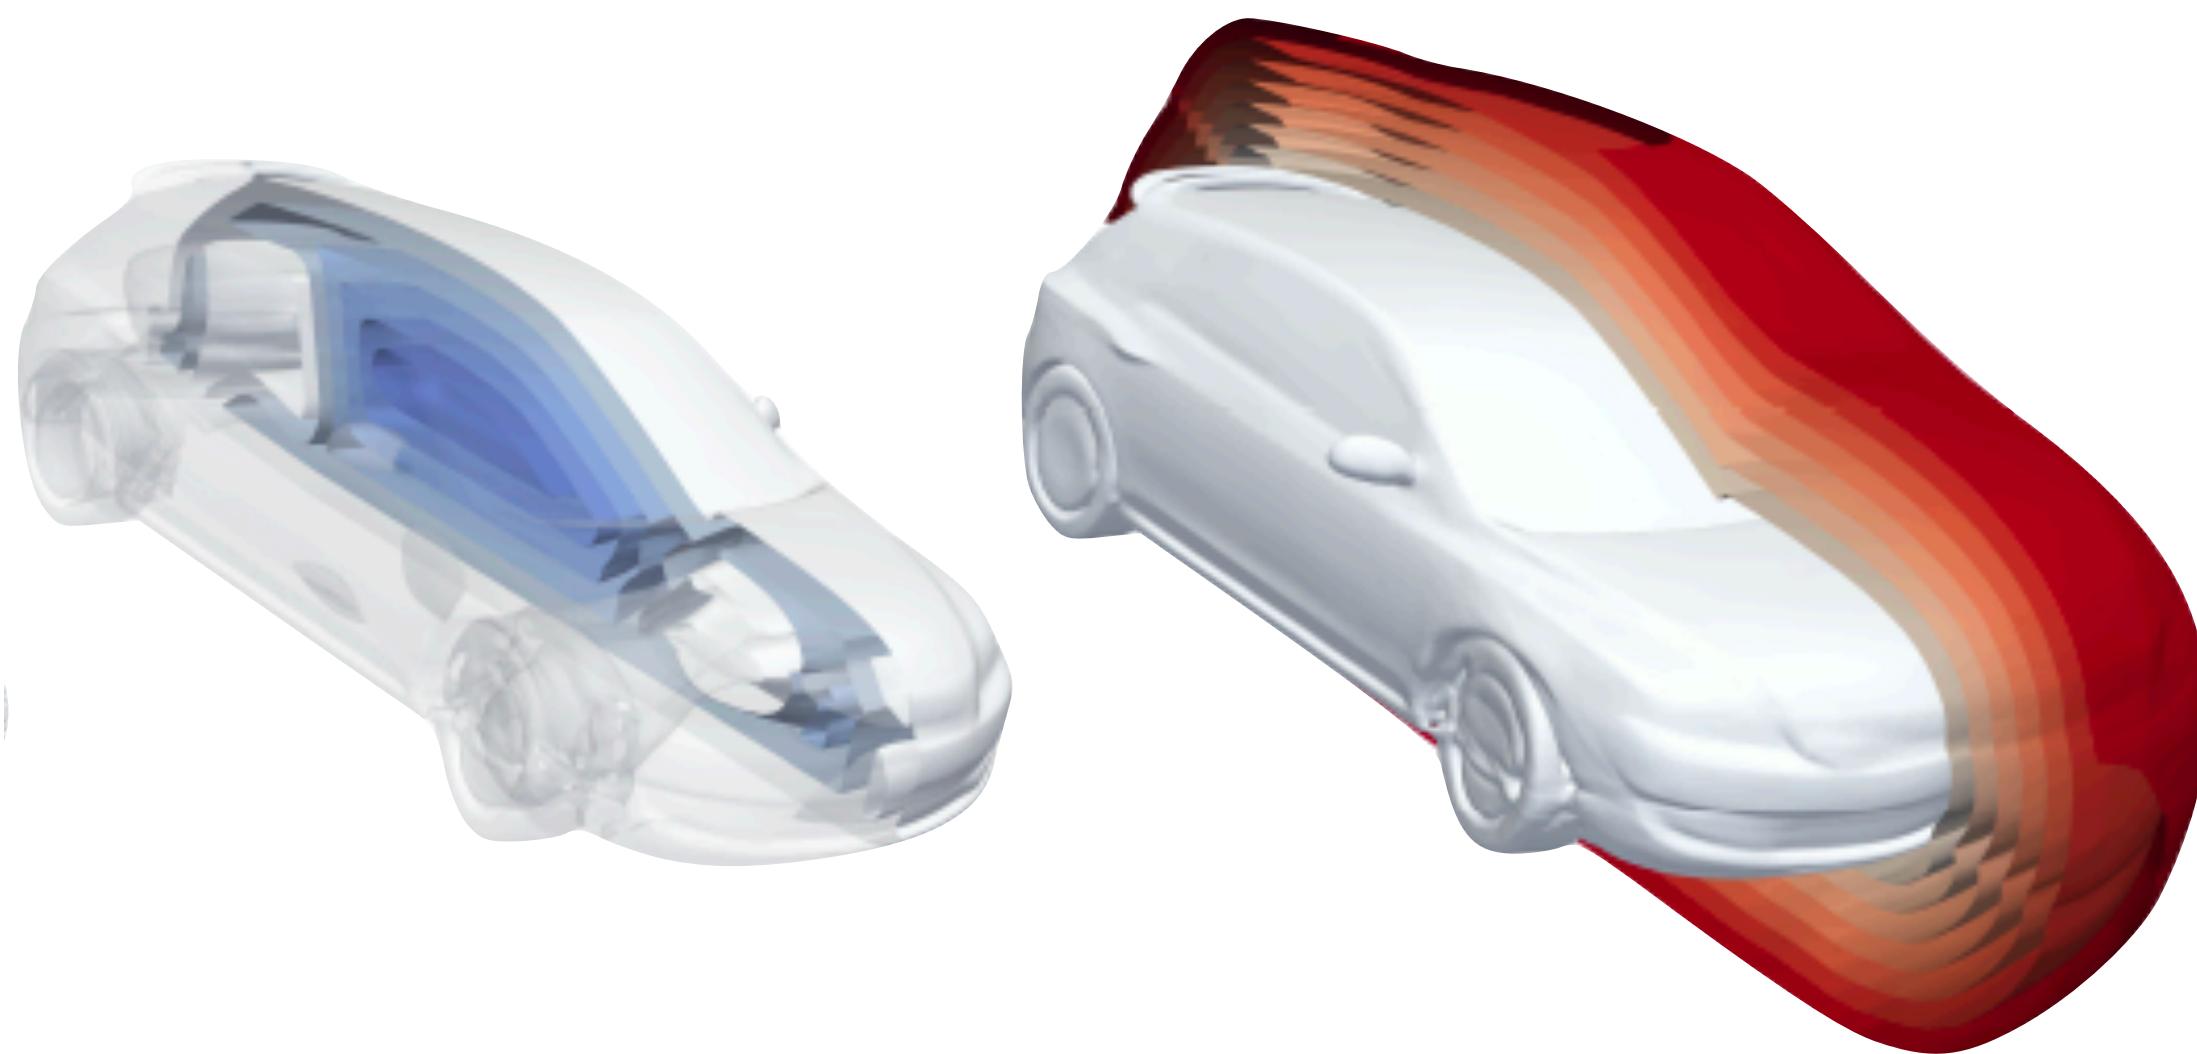
# The Signed Distance Field



# Implicit Surfaces

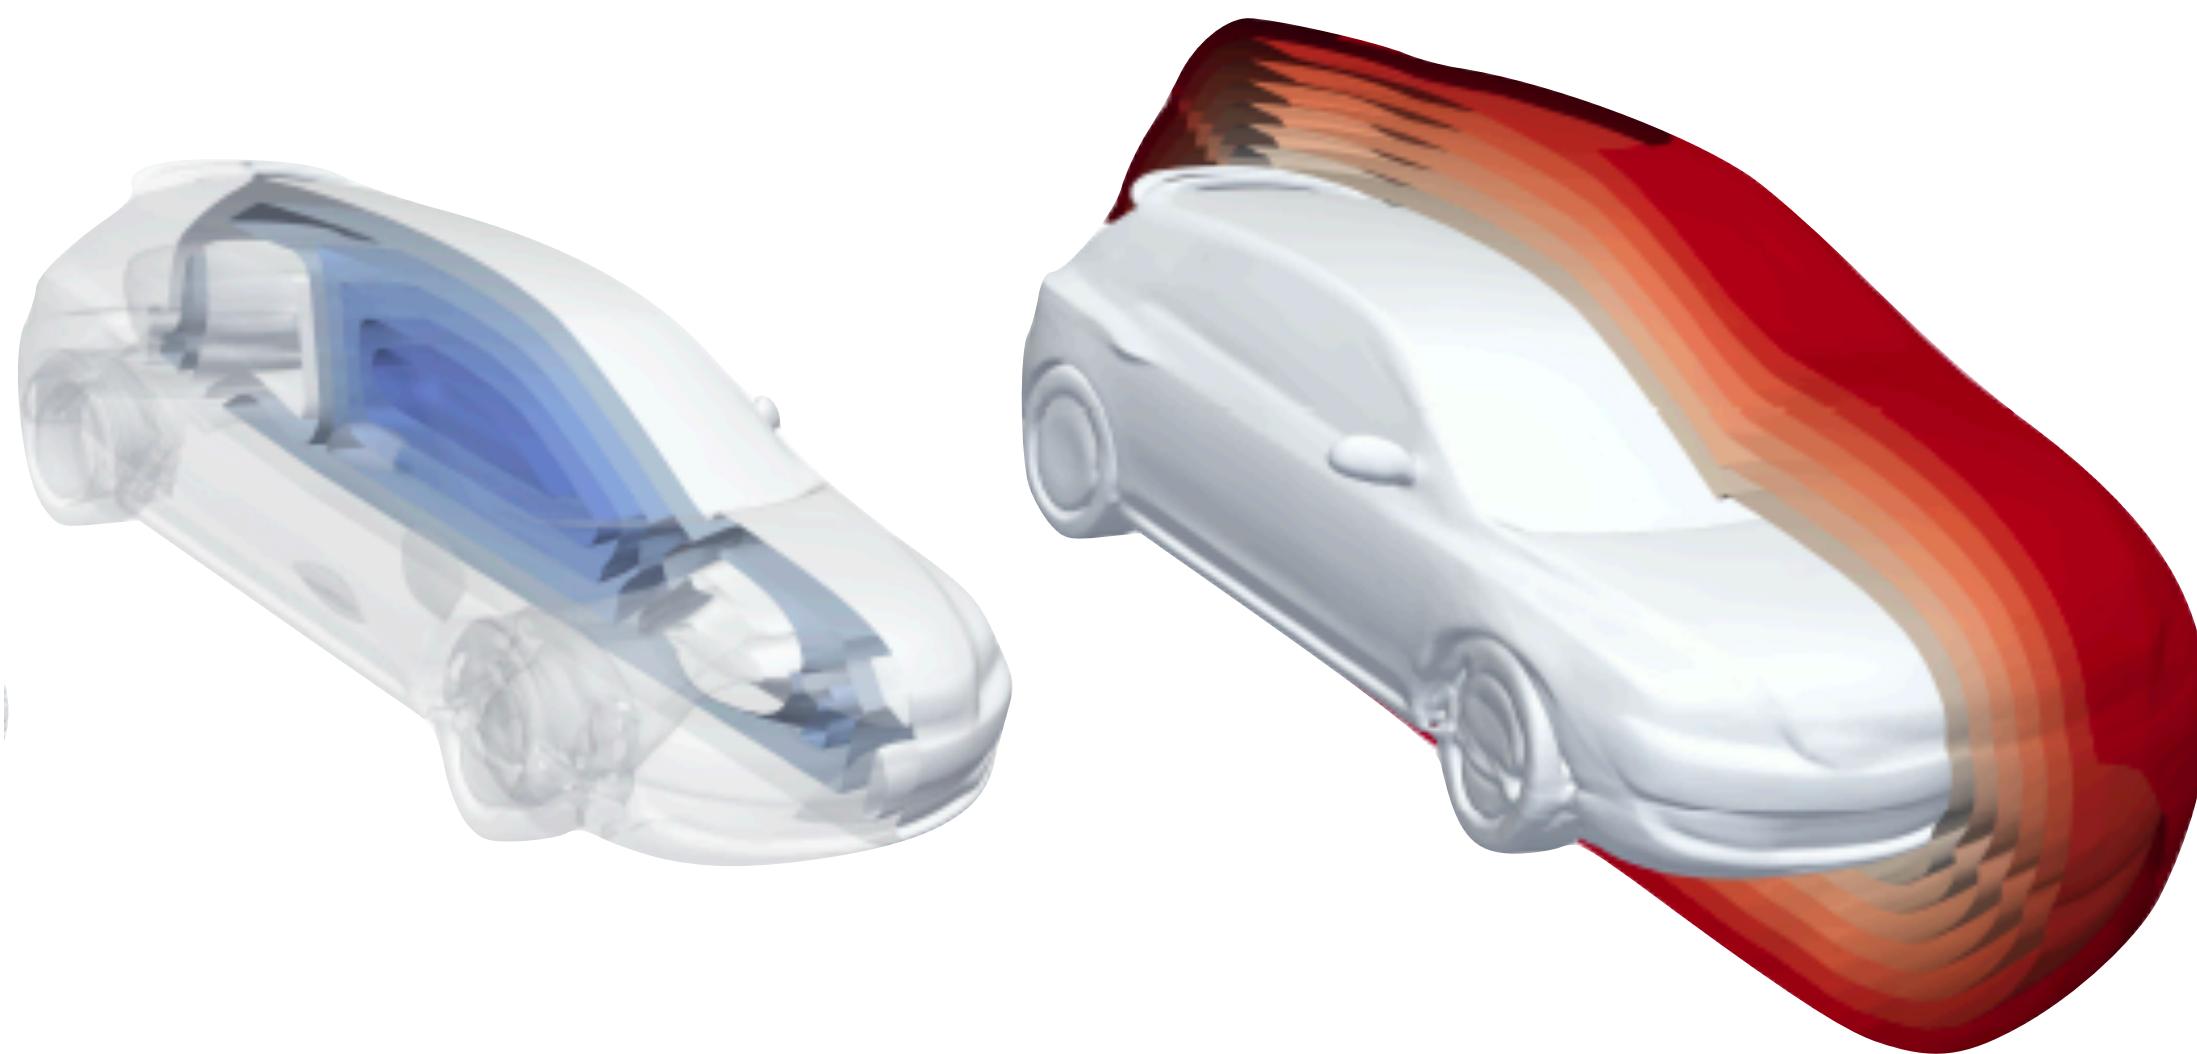


# Implicit Surfaces



Level sets of Occupancy, SDF  
parameterize surface

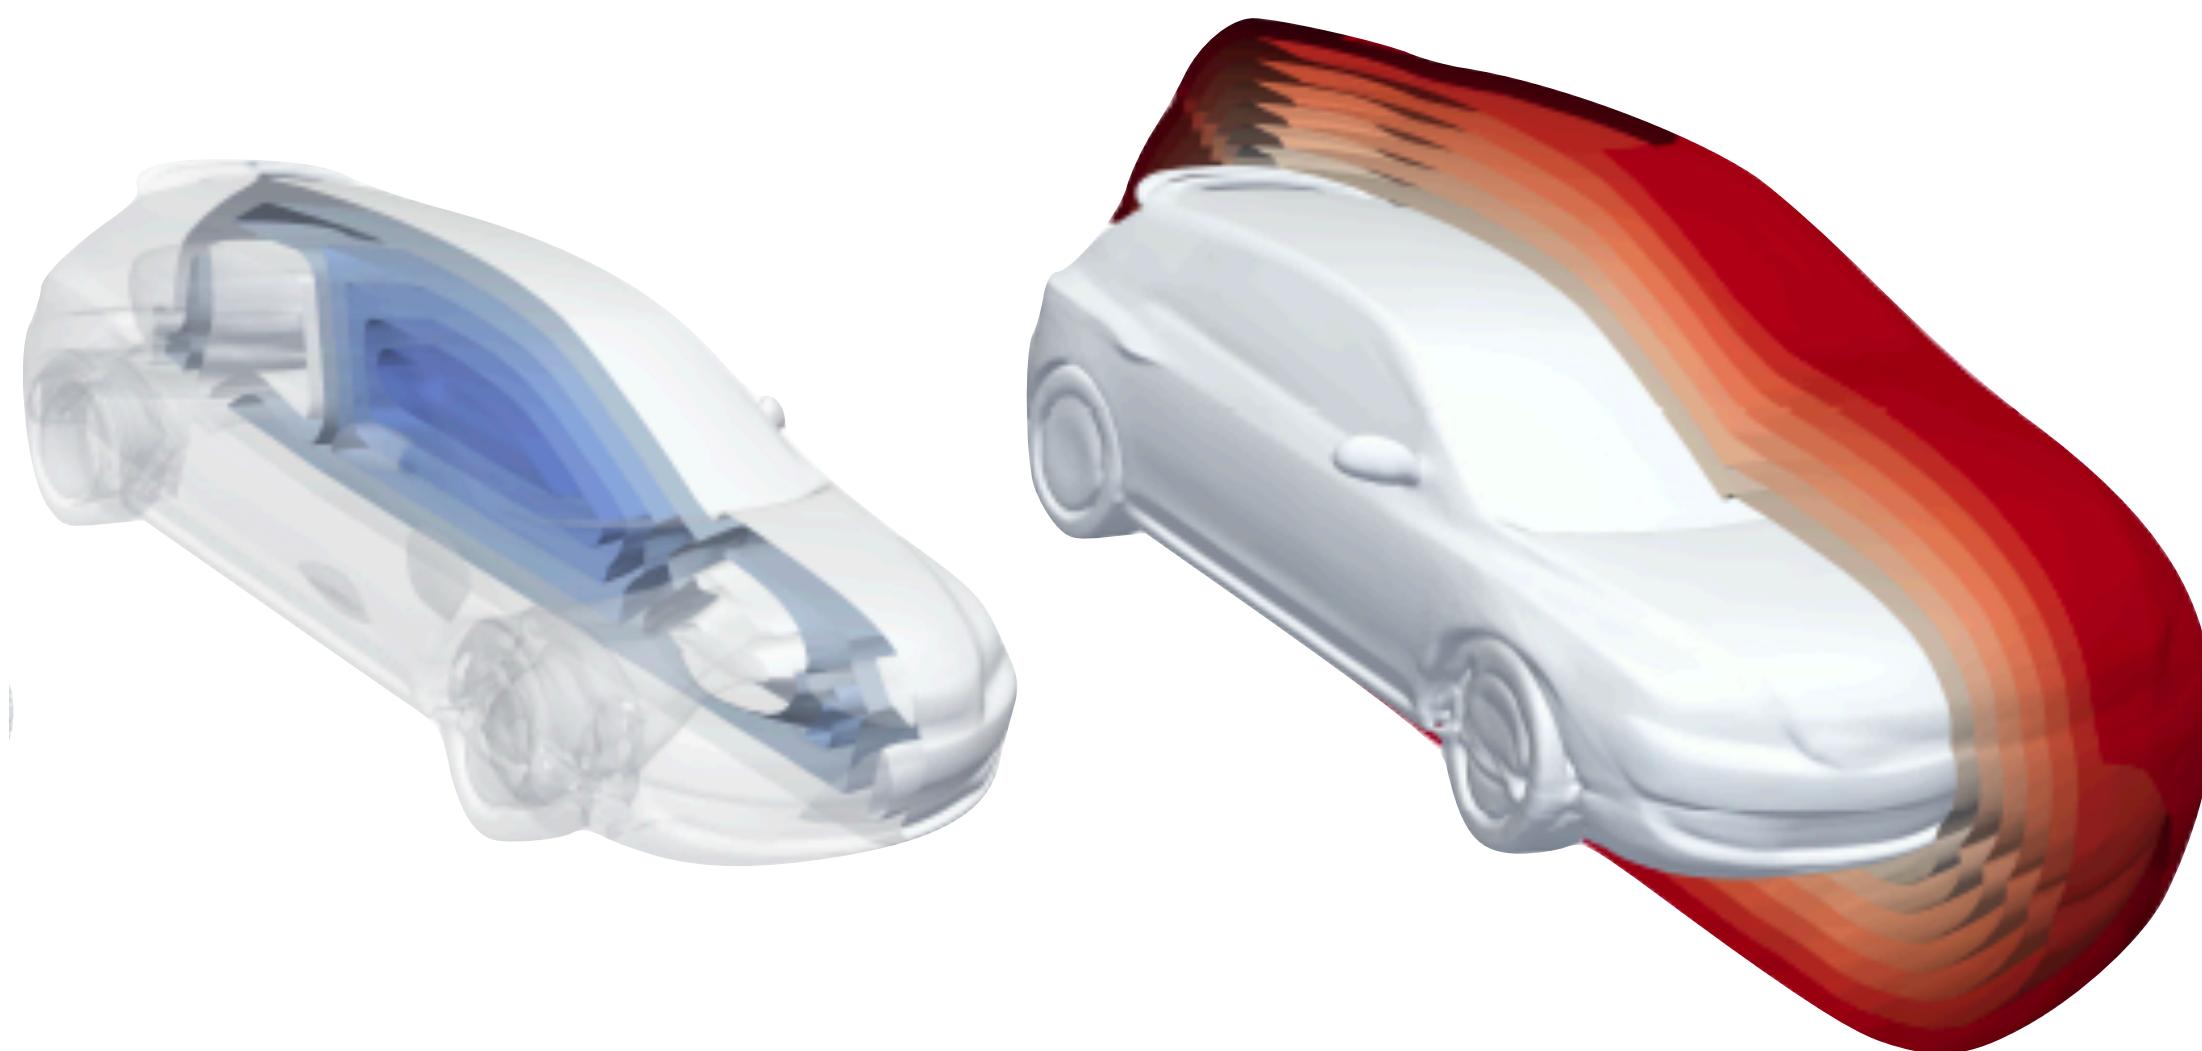
# Implicit Surfaces



Level sets of Occupancy, SDF  
parameterize surface

Easy to answer questions about  
occupancy

# Implicit Surfaces

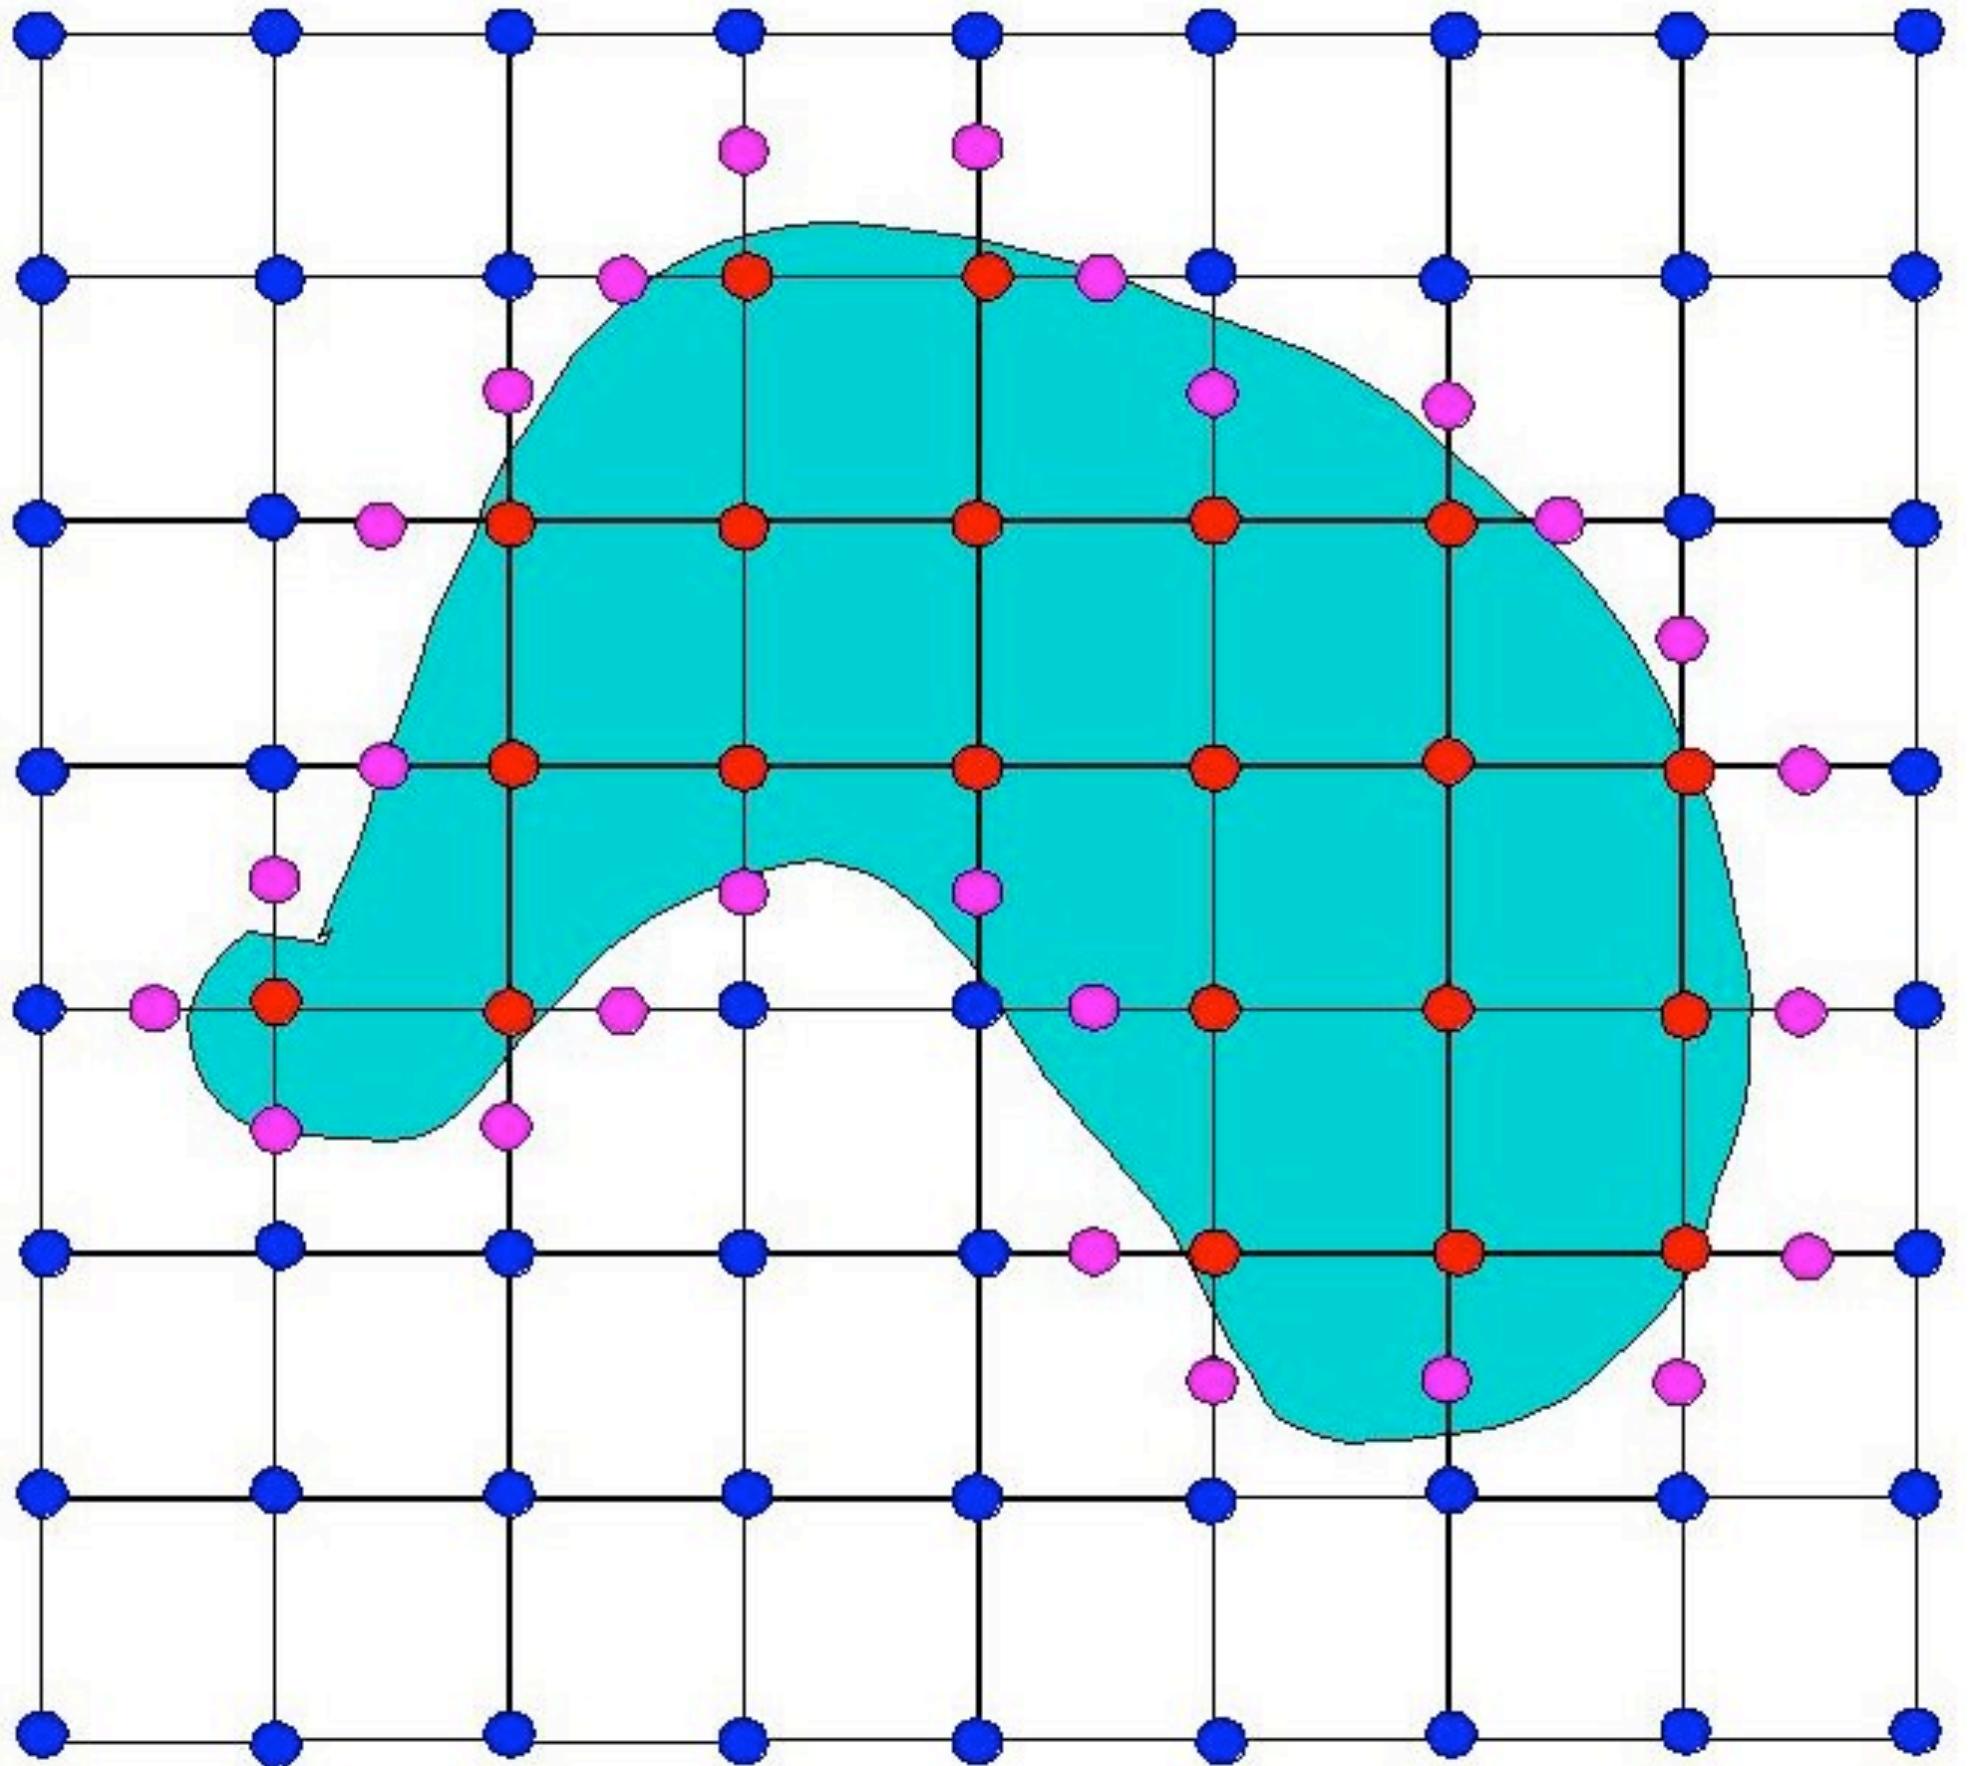


Level sets of Occupancy, SDF  
parameterize surface

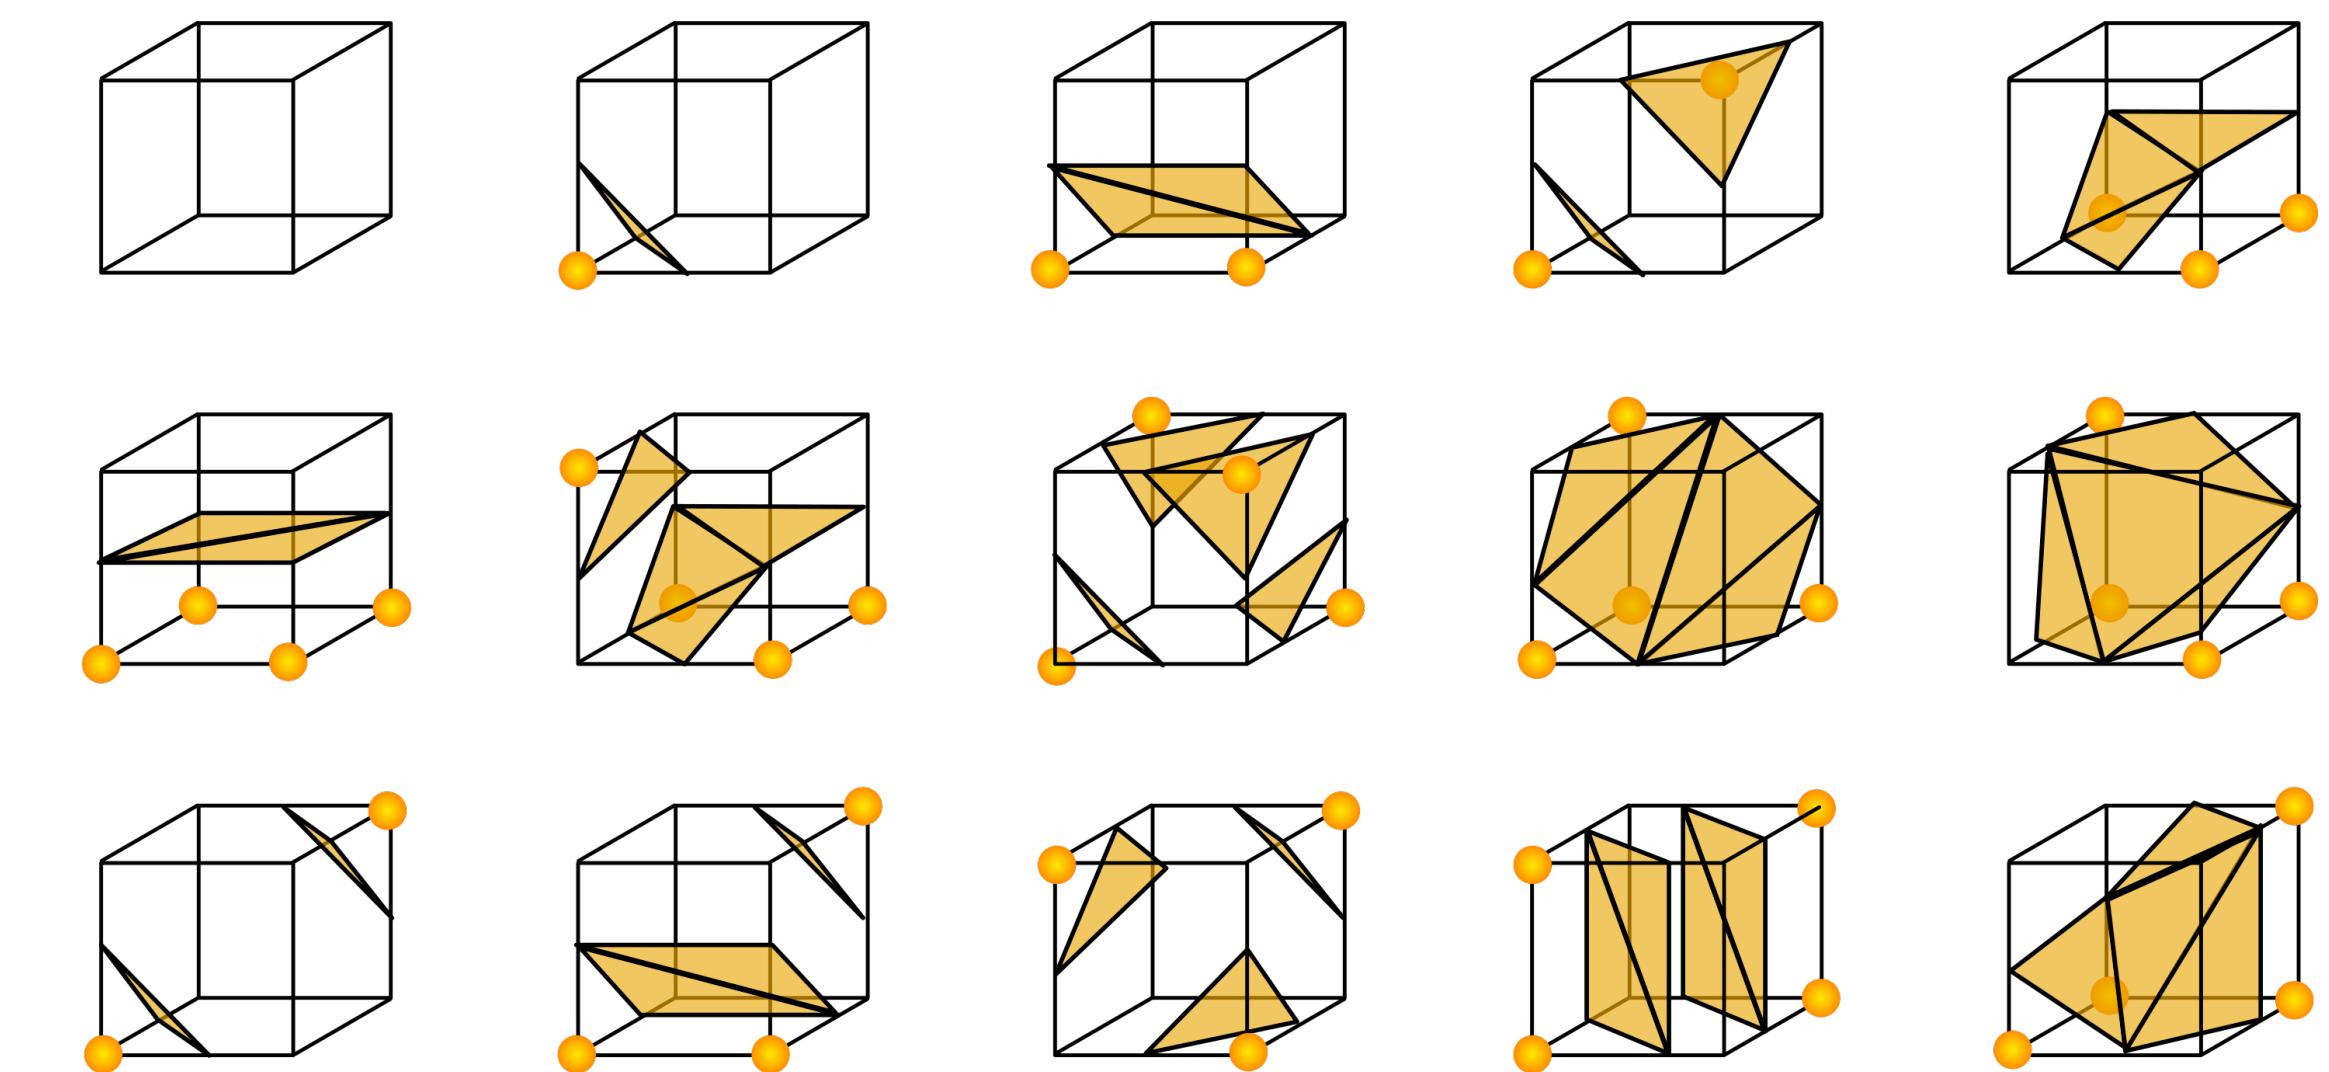
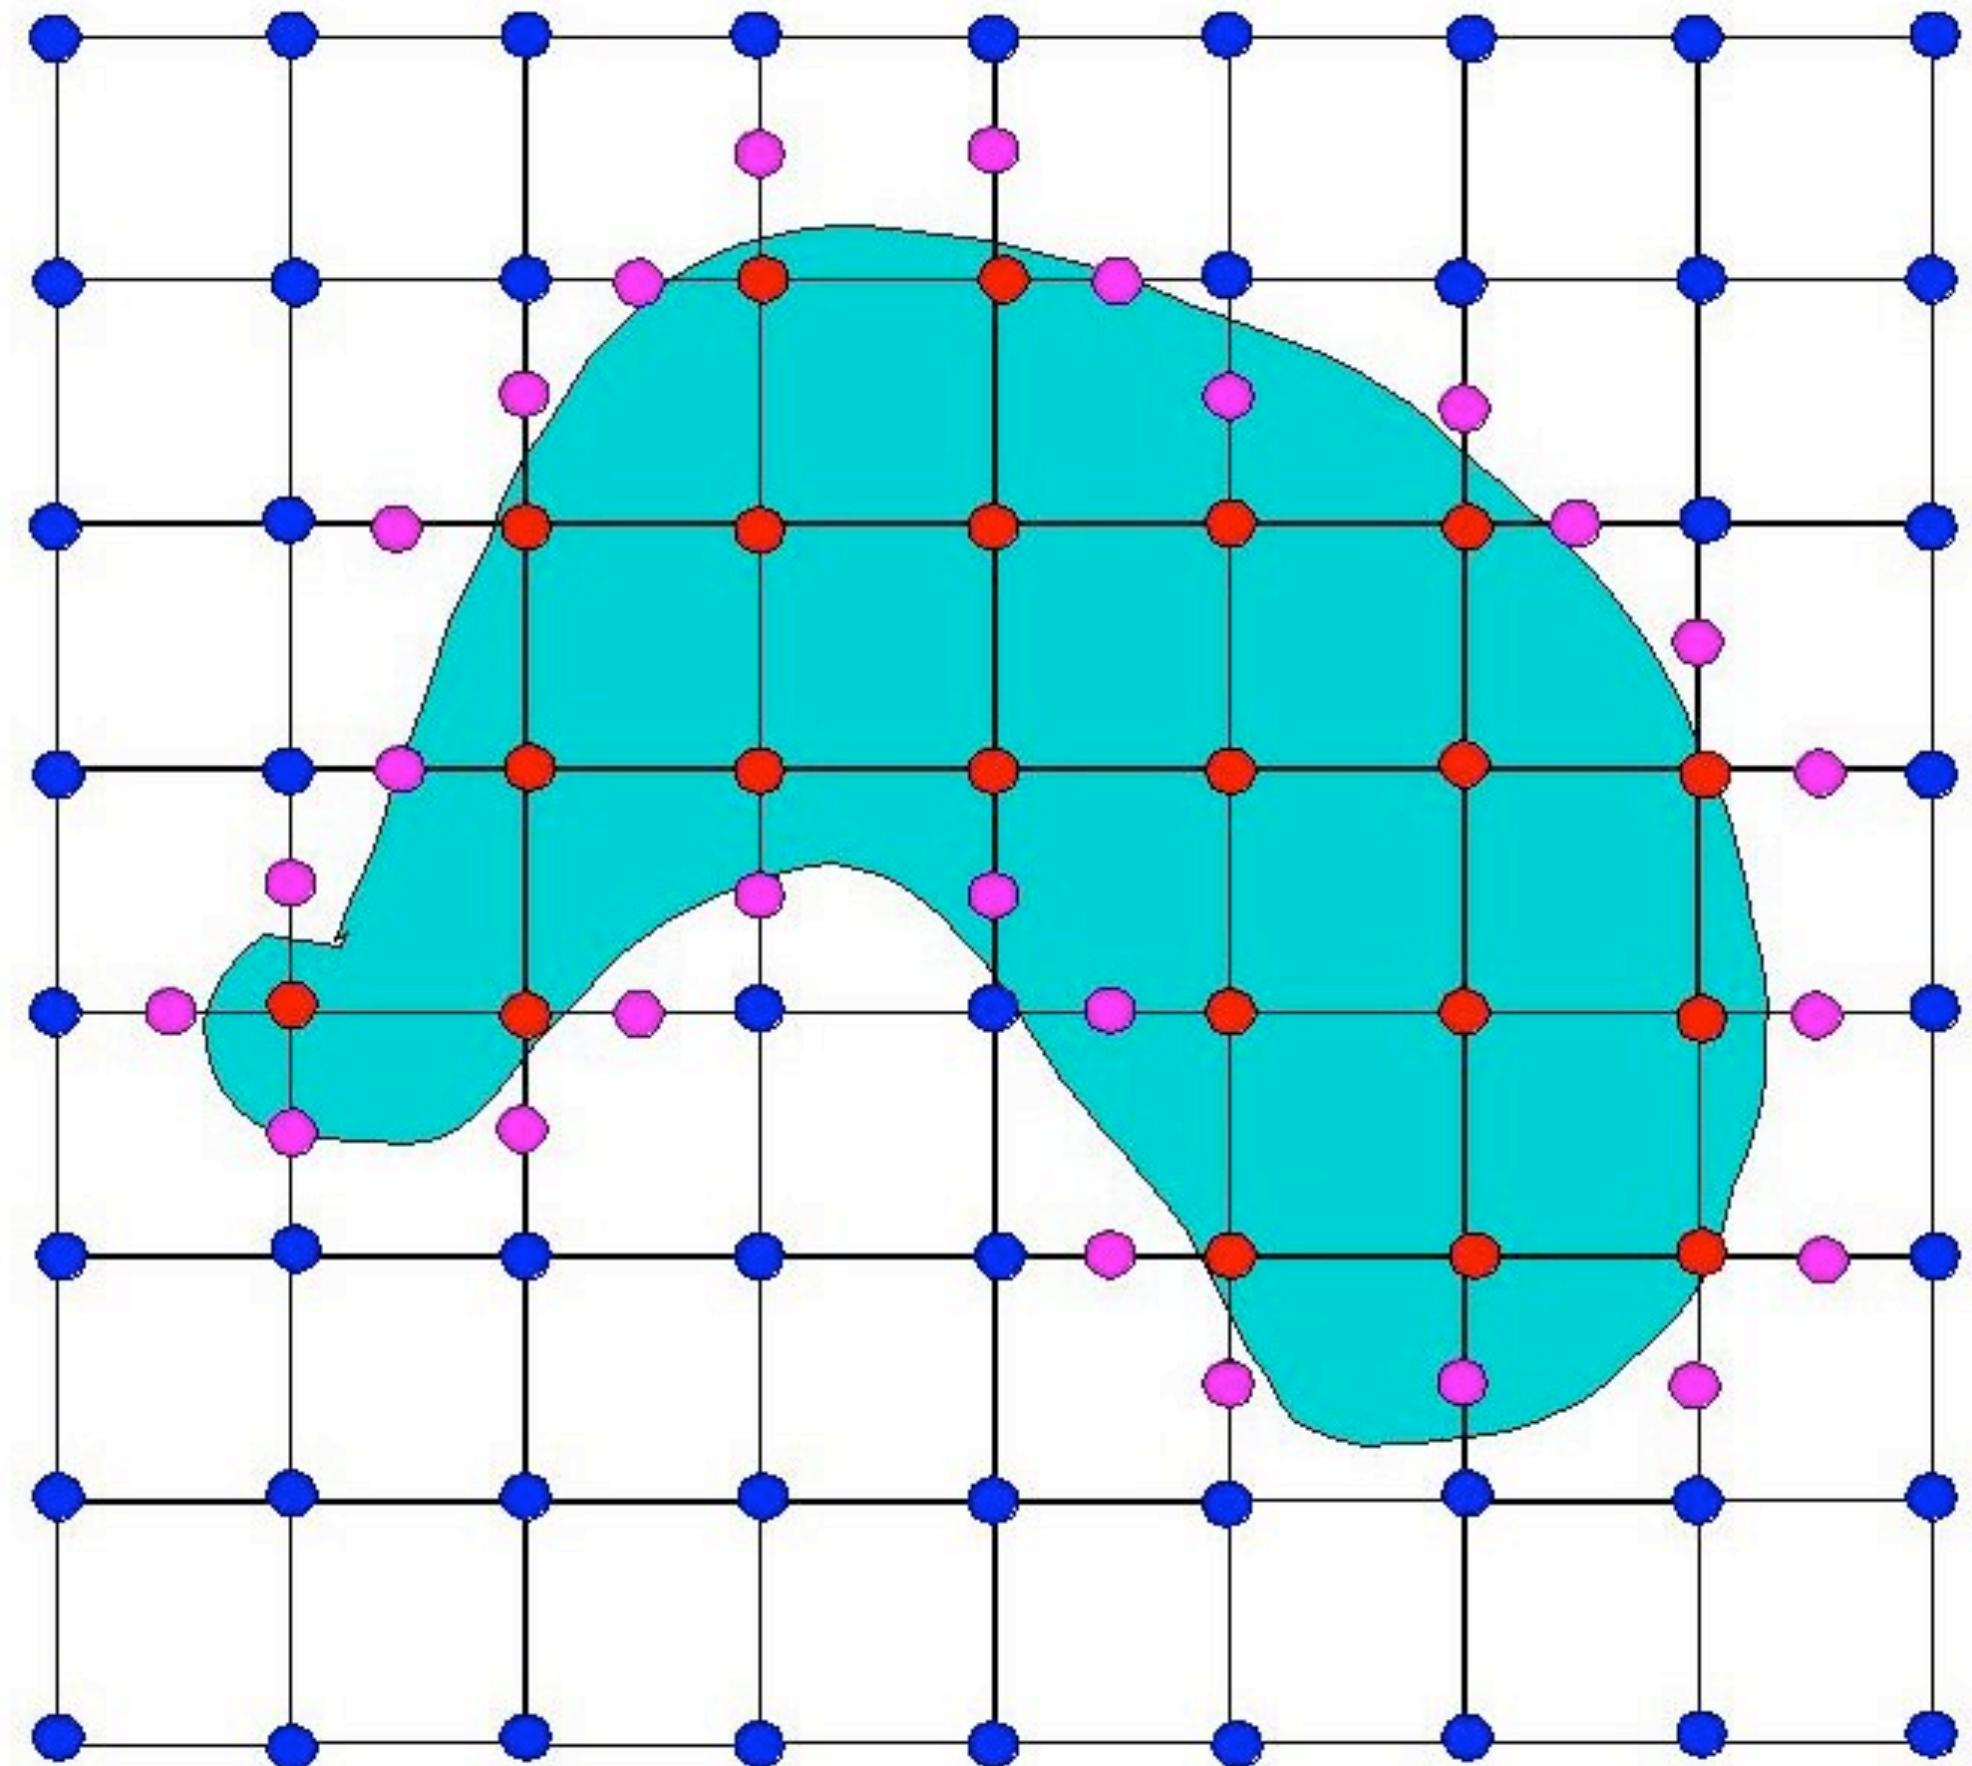
Easy to answer questions about  
occupancy

Difficult to reason about the surface:  
have to first **find** zero-level-set!

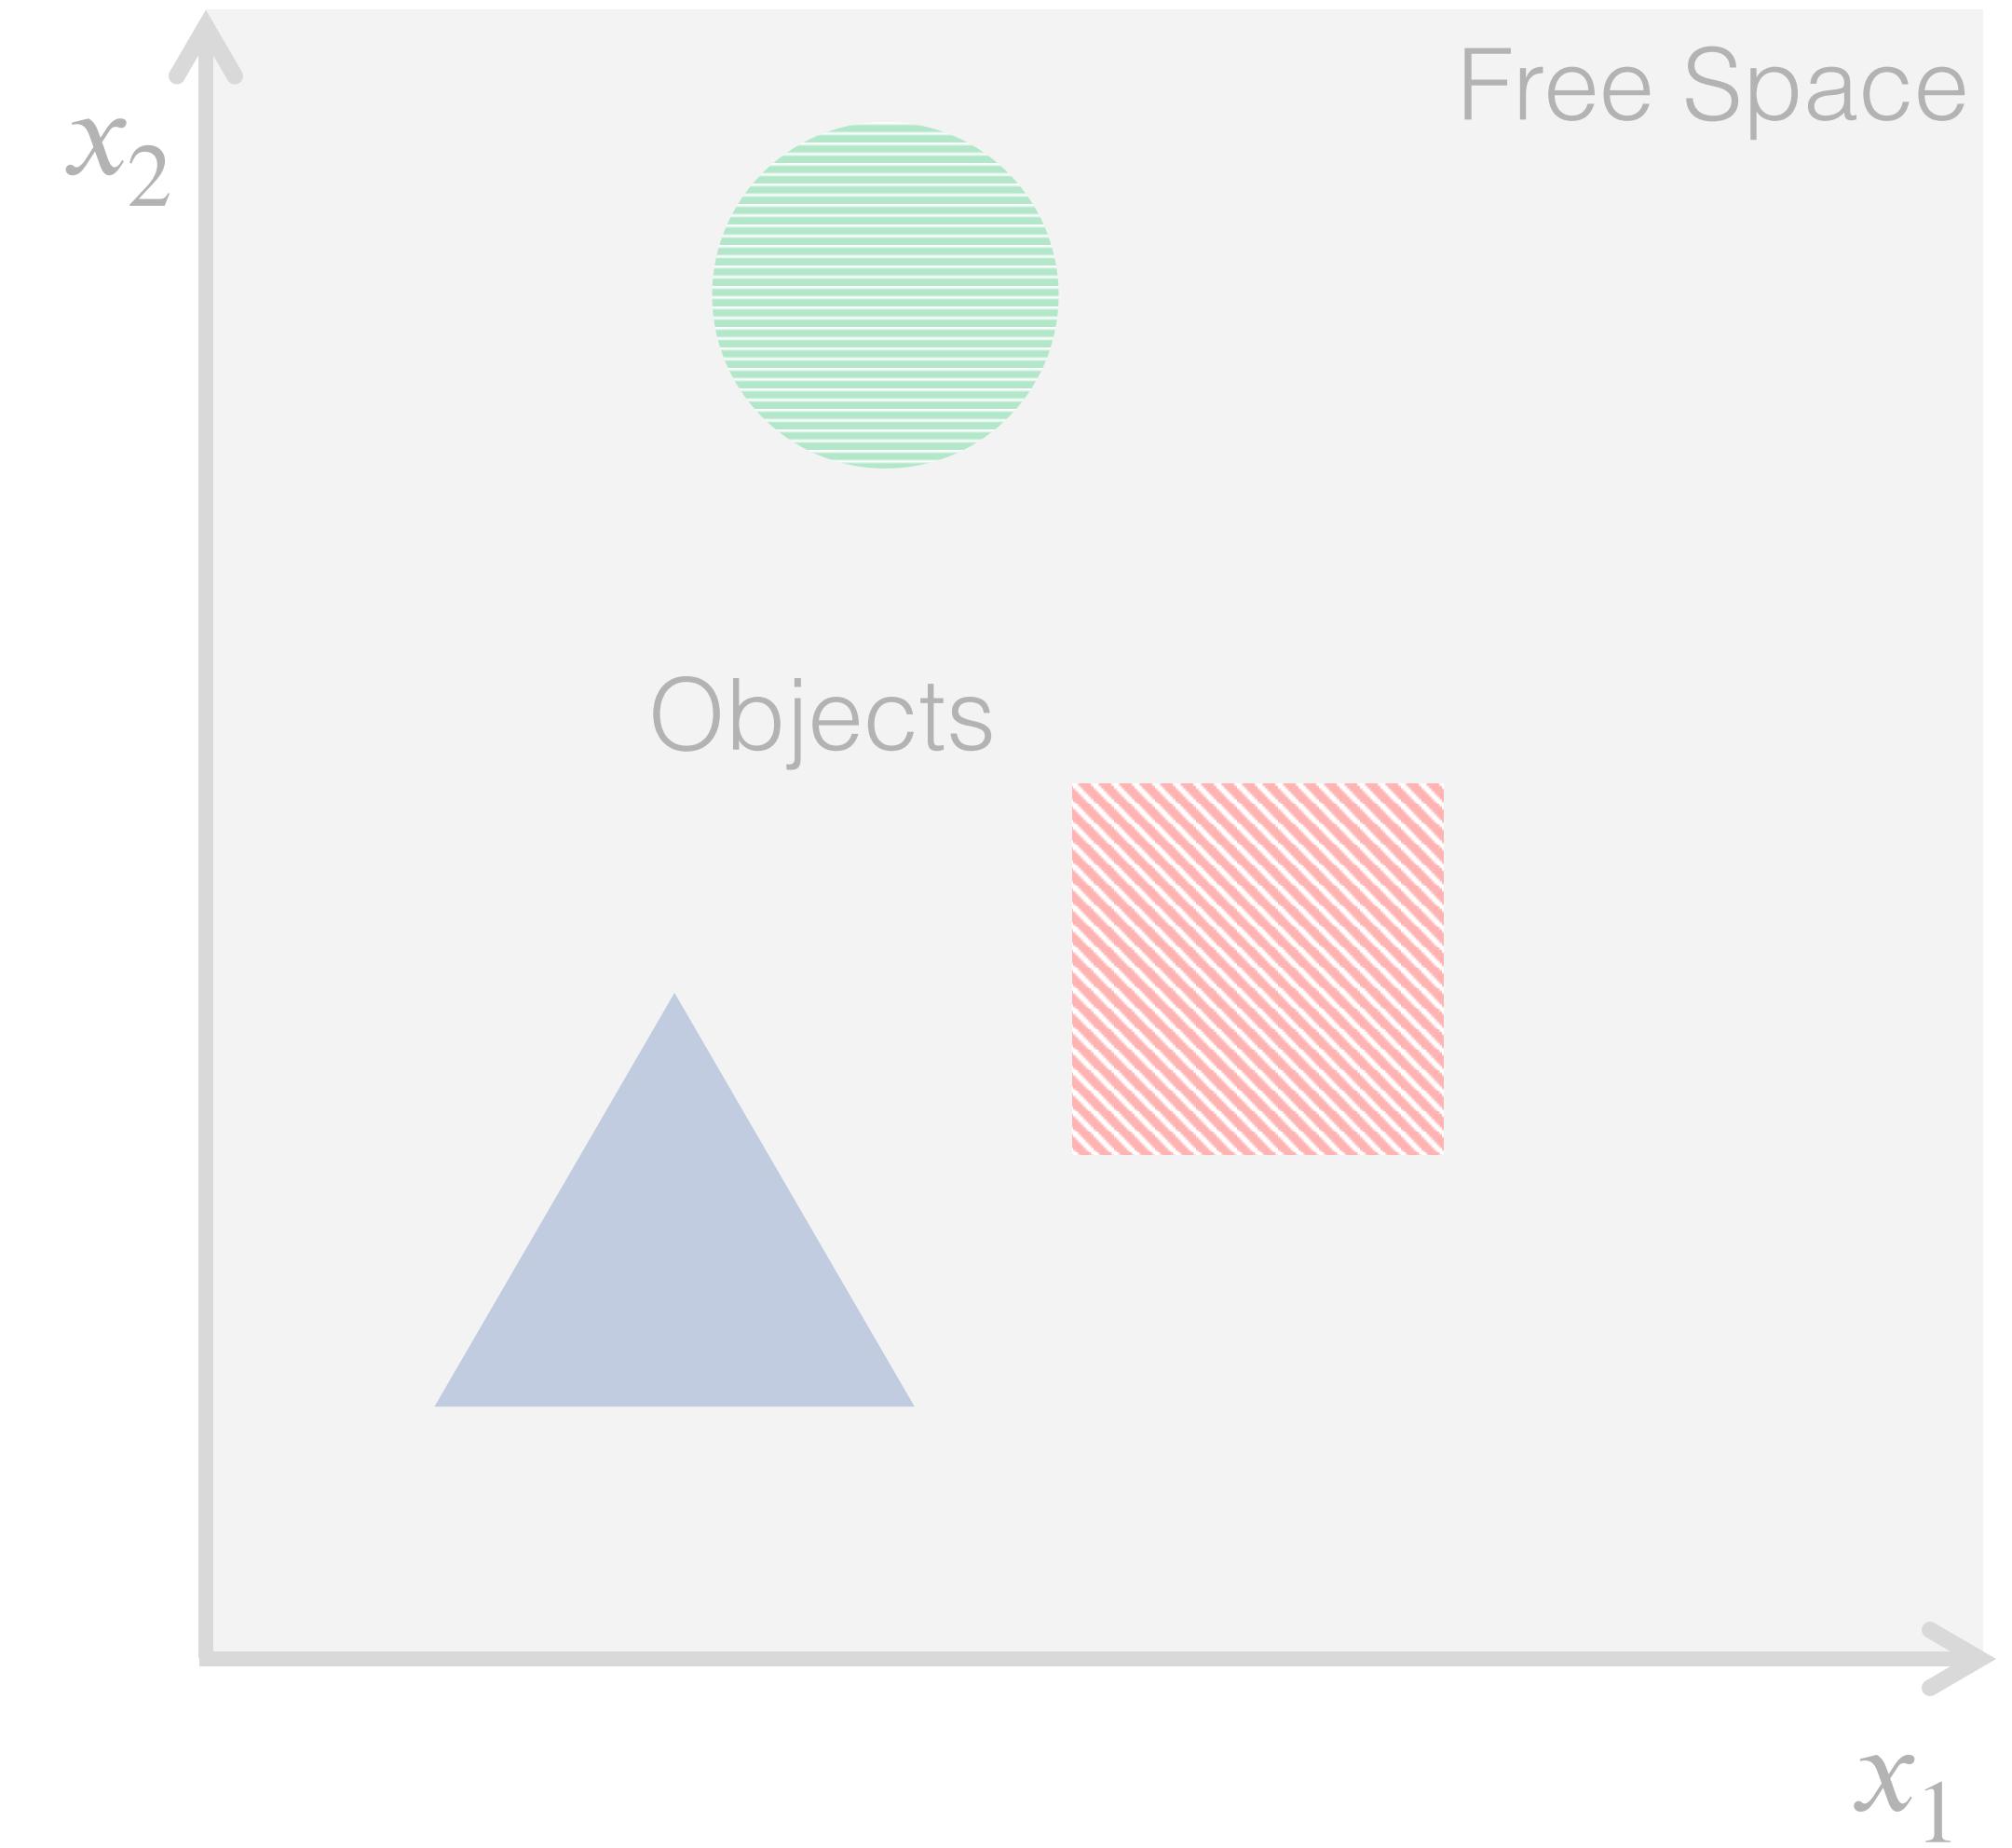
# Marching Cubes: Extracting 2D Mesh from 3D Implicit Representation



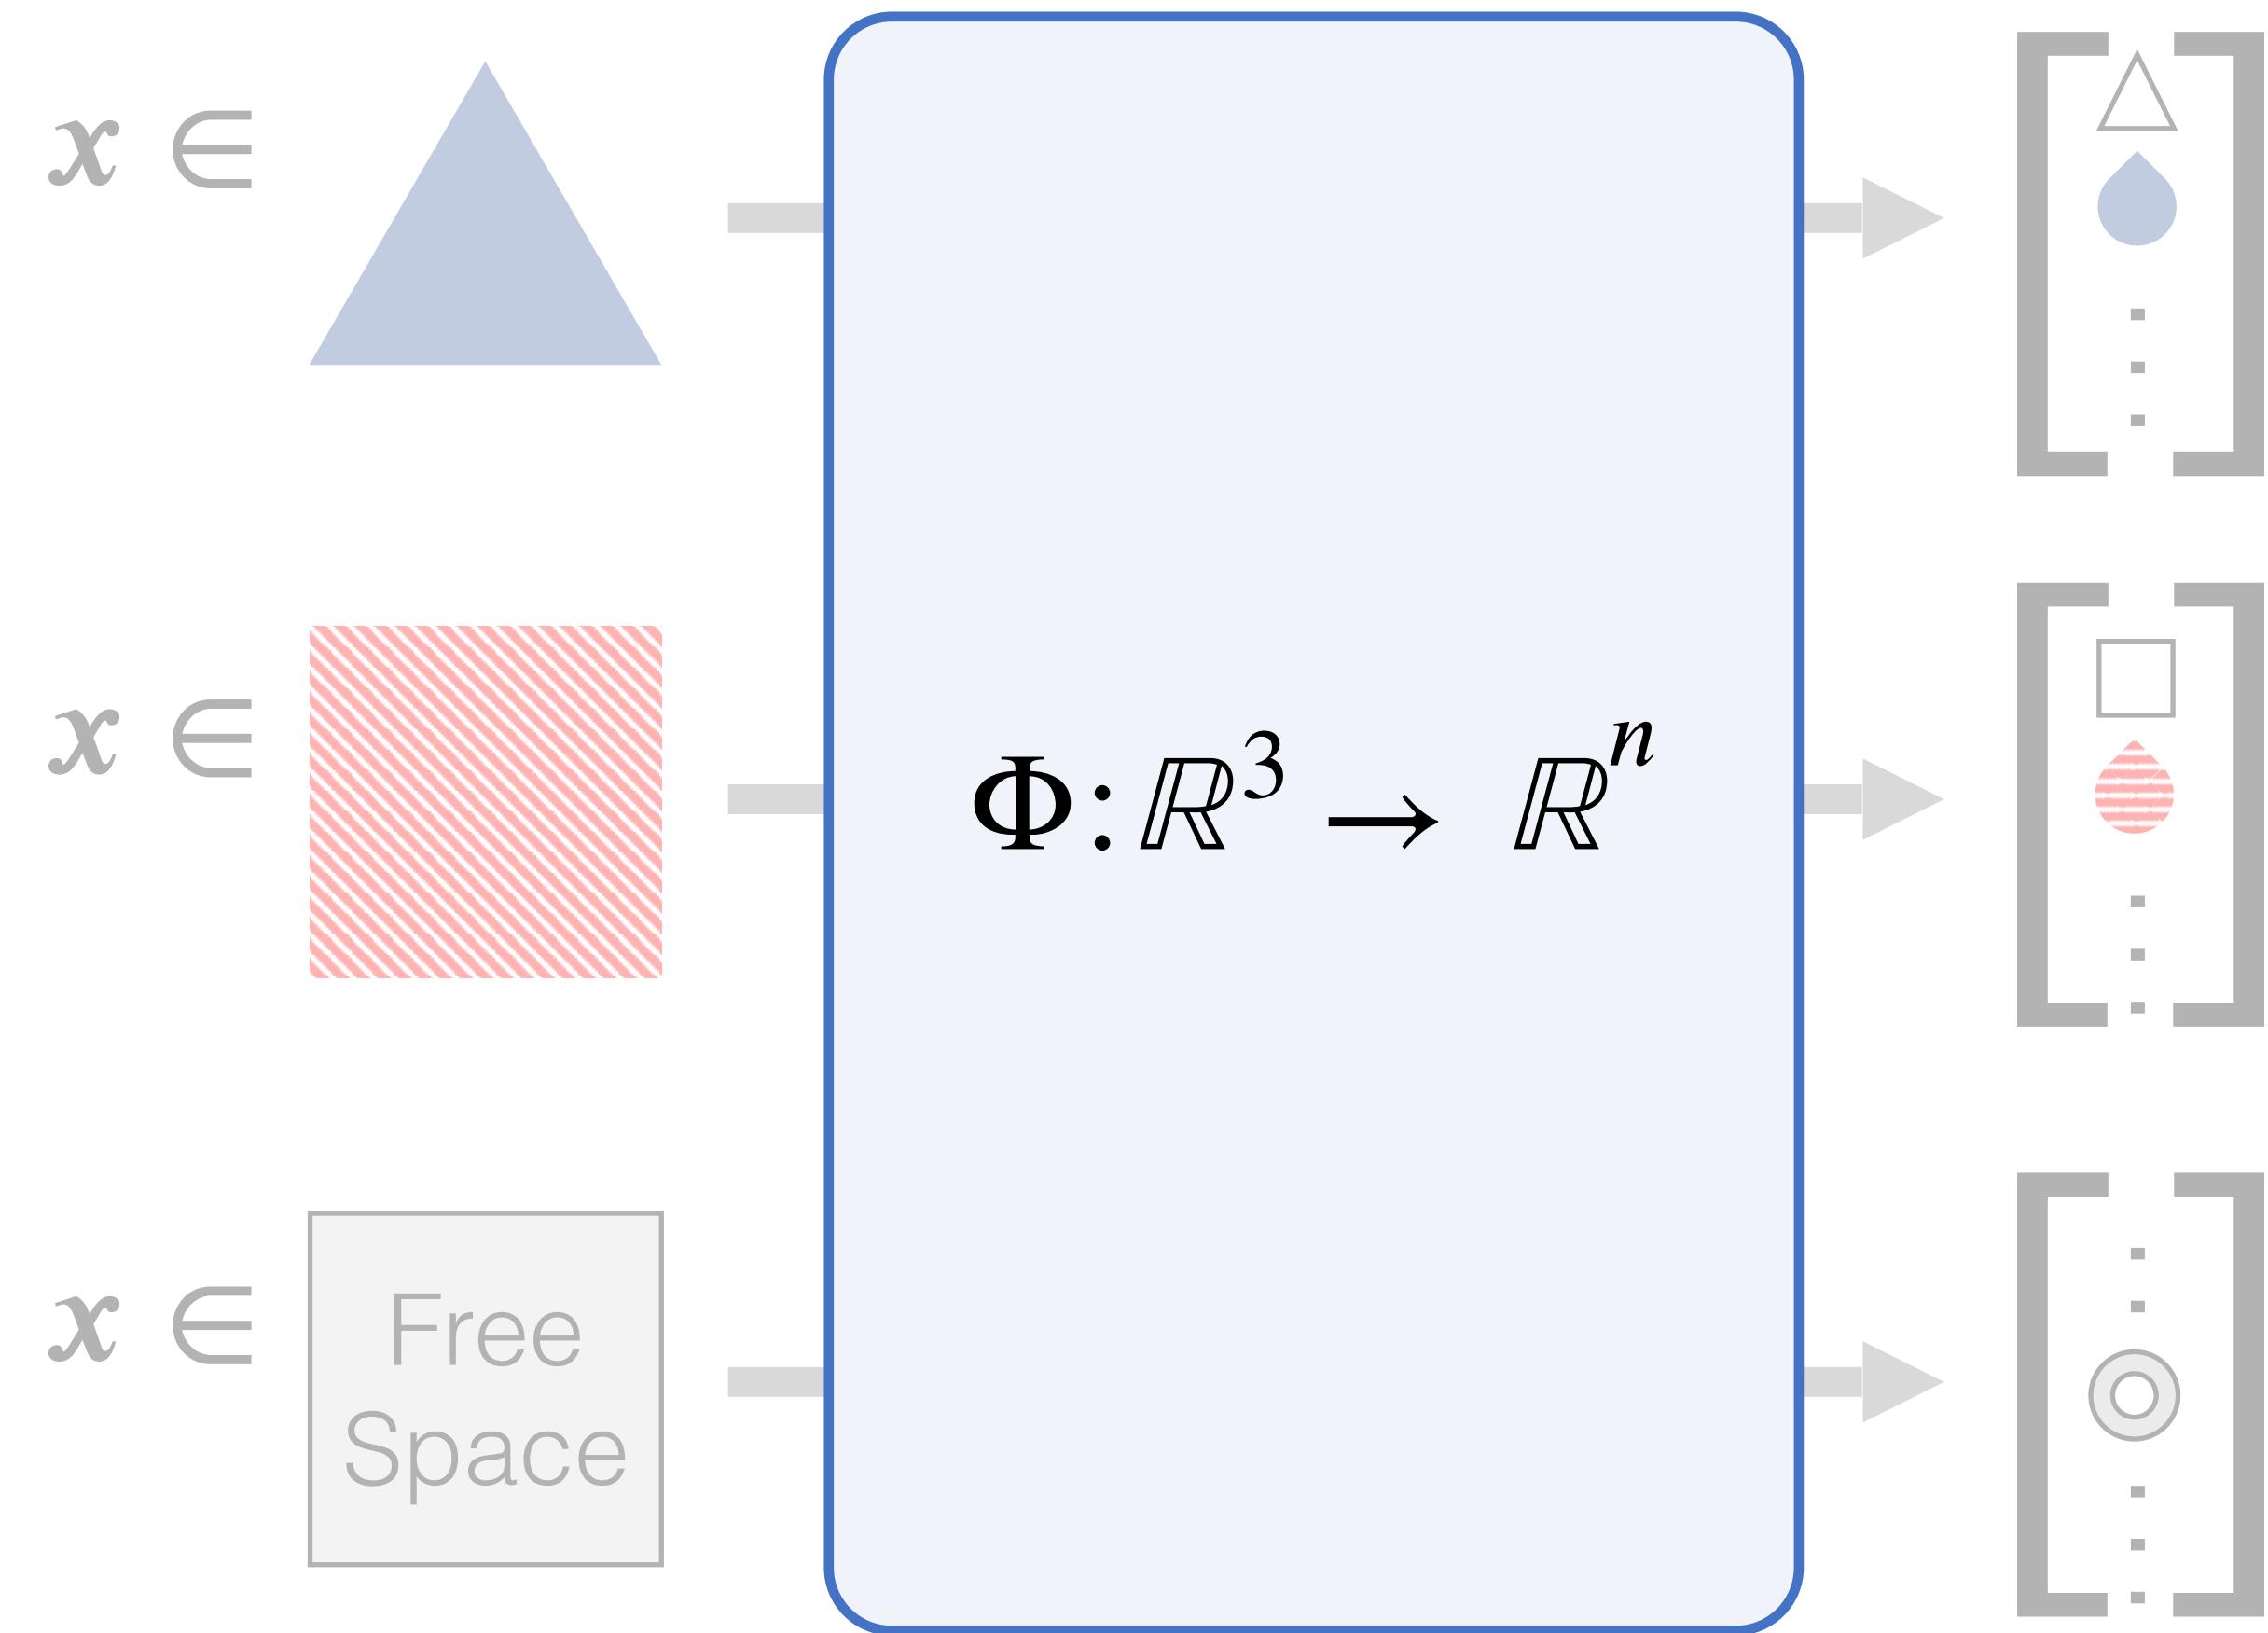
# Marching Cubes: Extracting 2D Mesh from 3D Implicit Representation



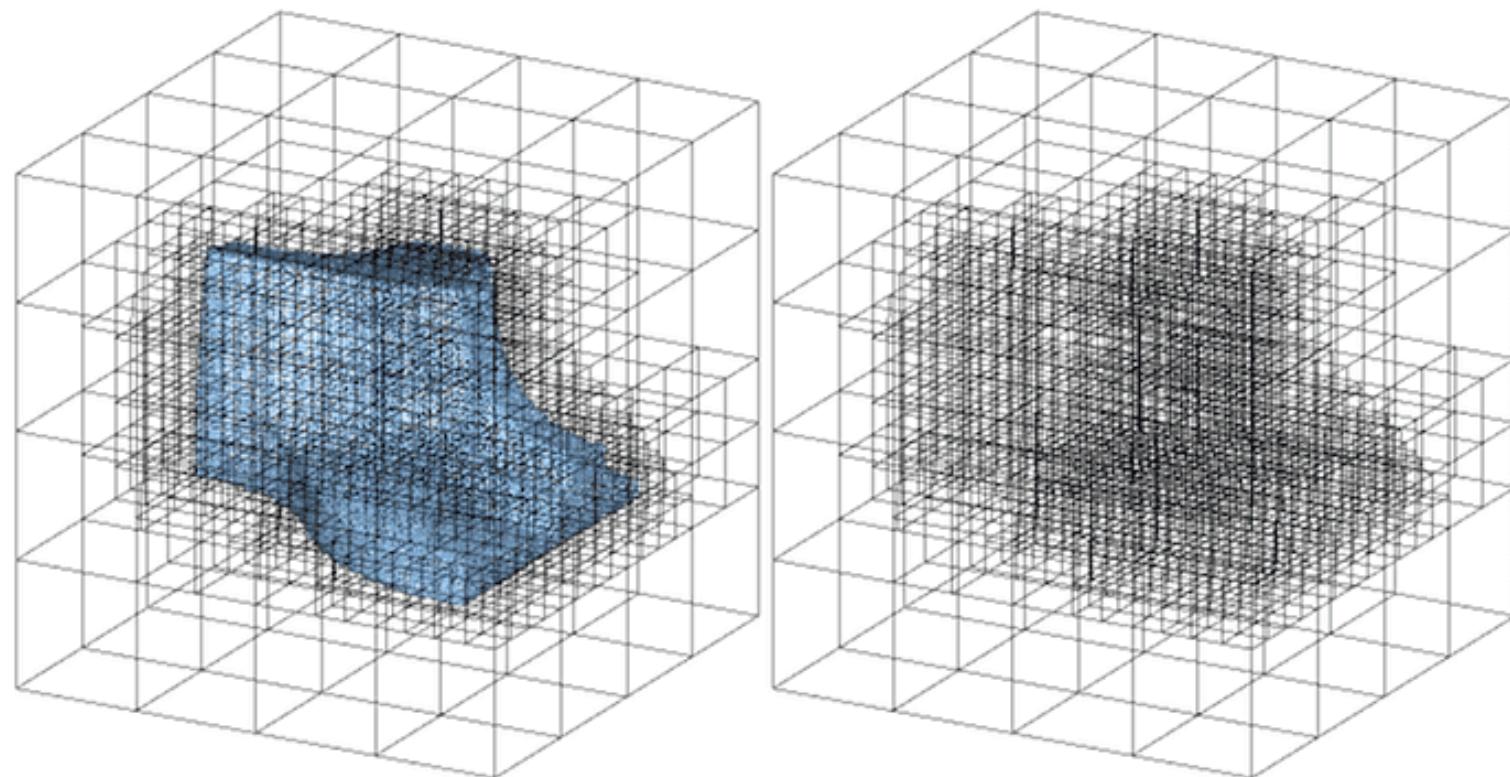
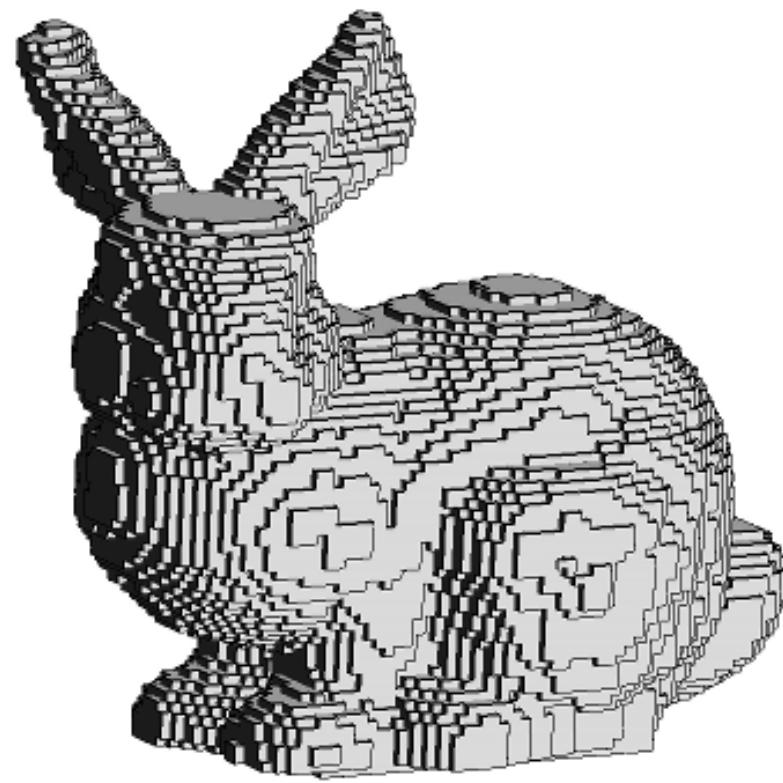
# How to parameterize surfaces?



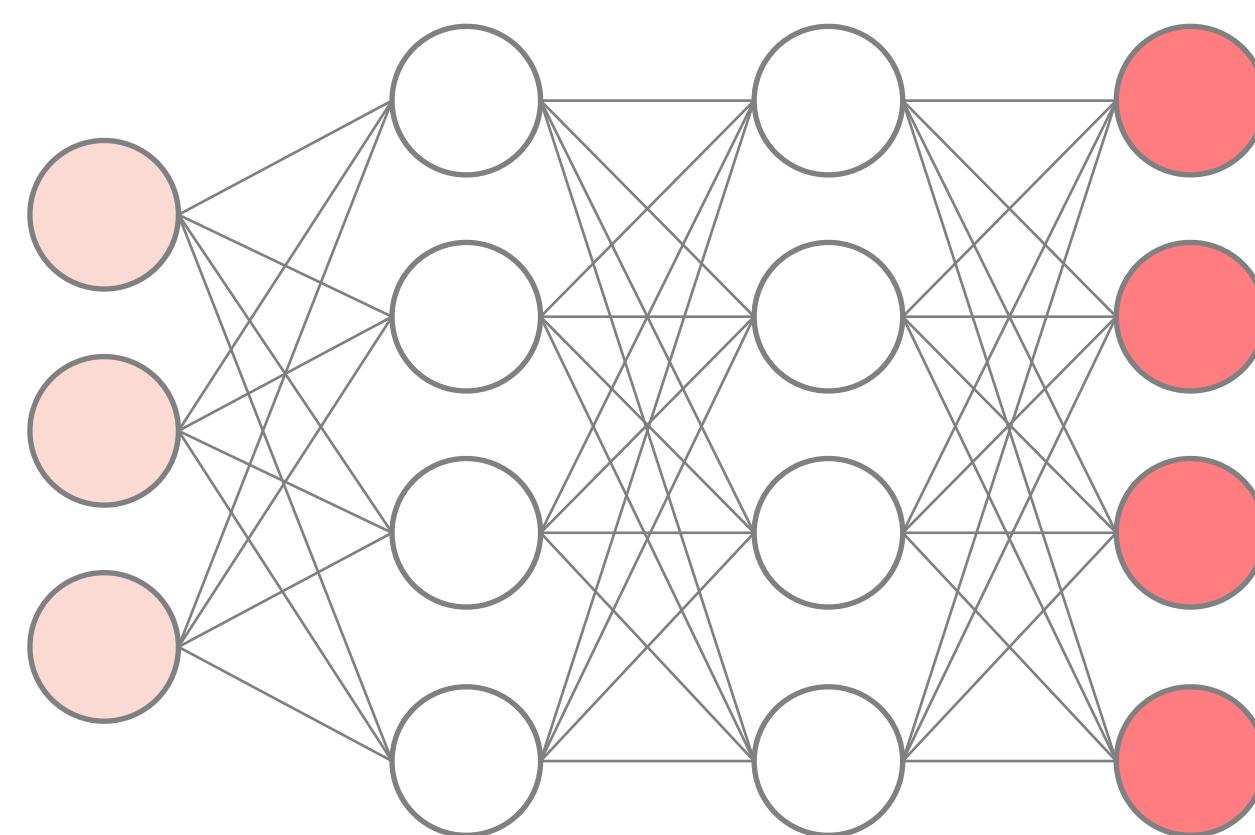
How to parameterize 3D function?



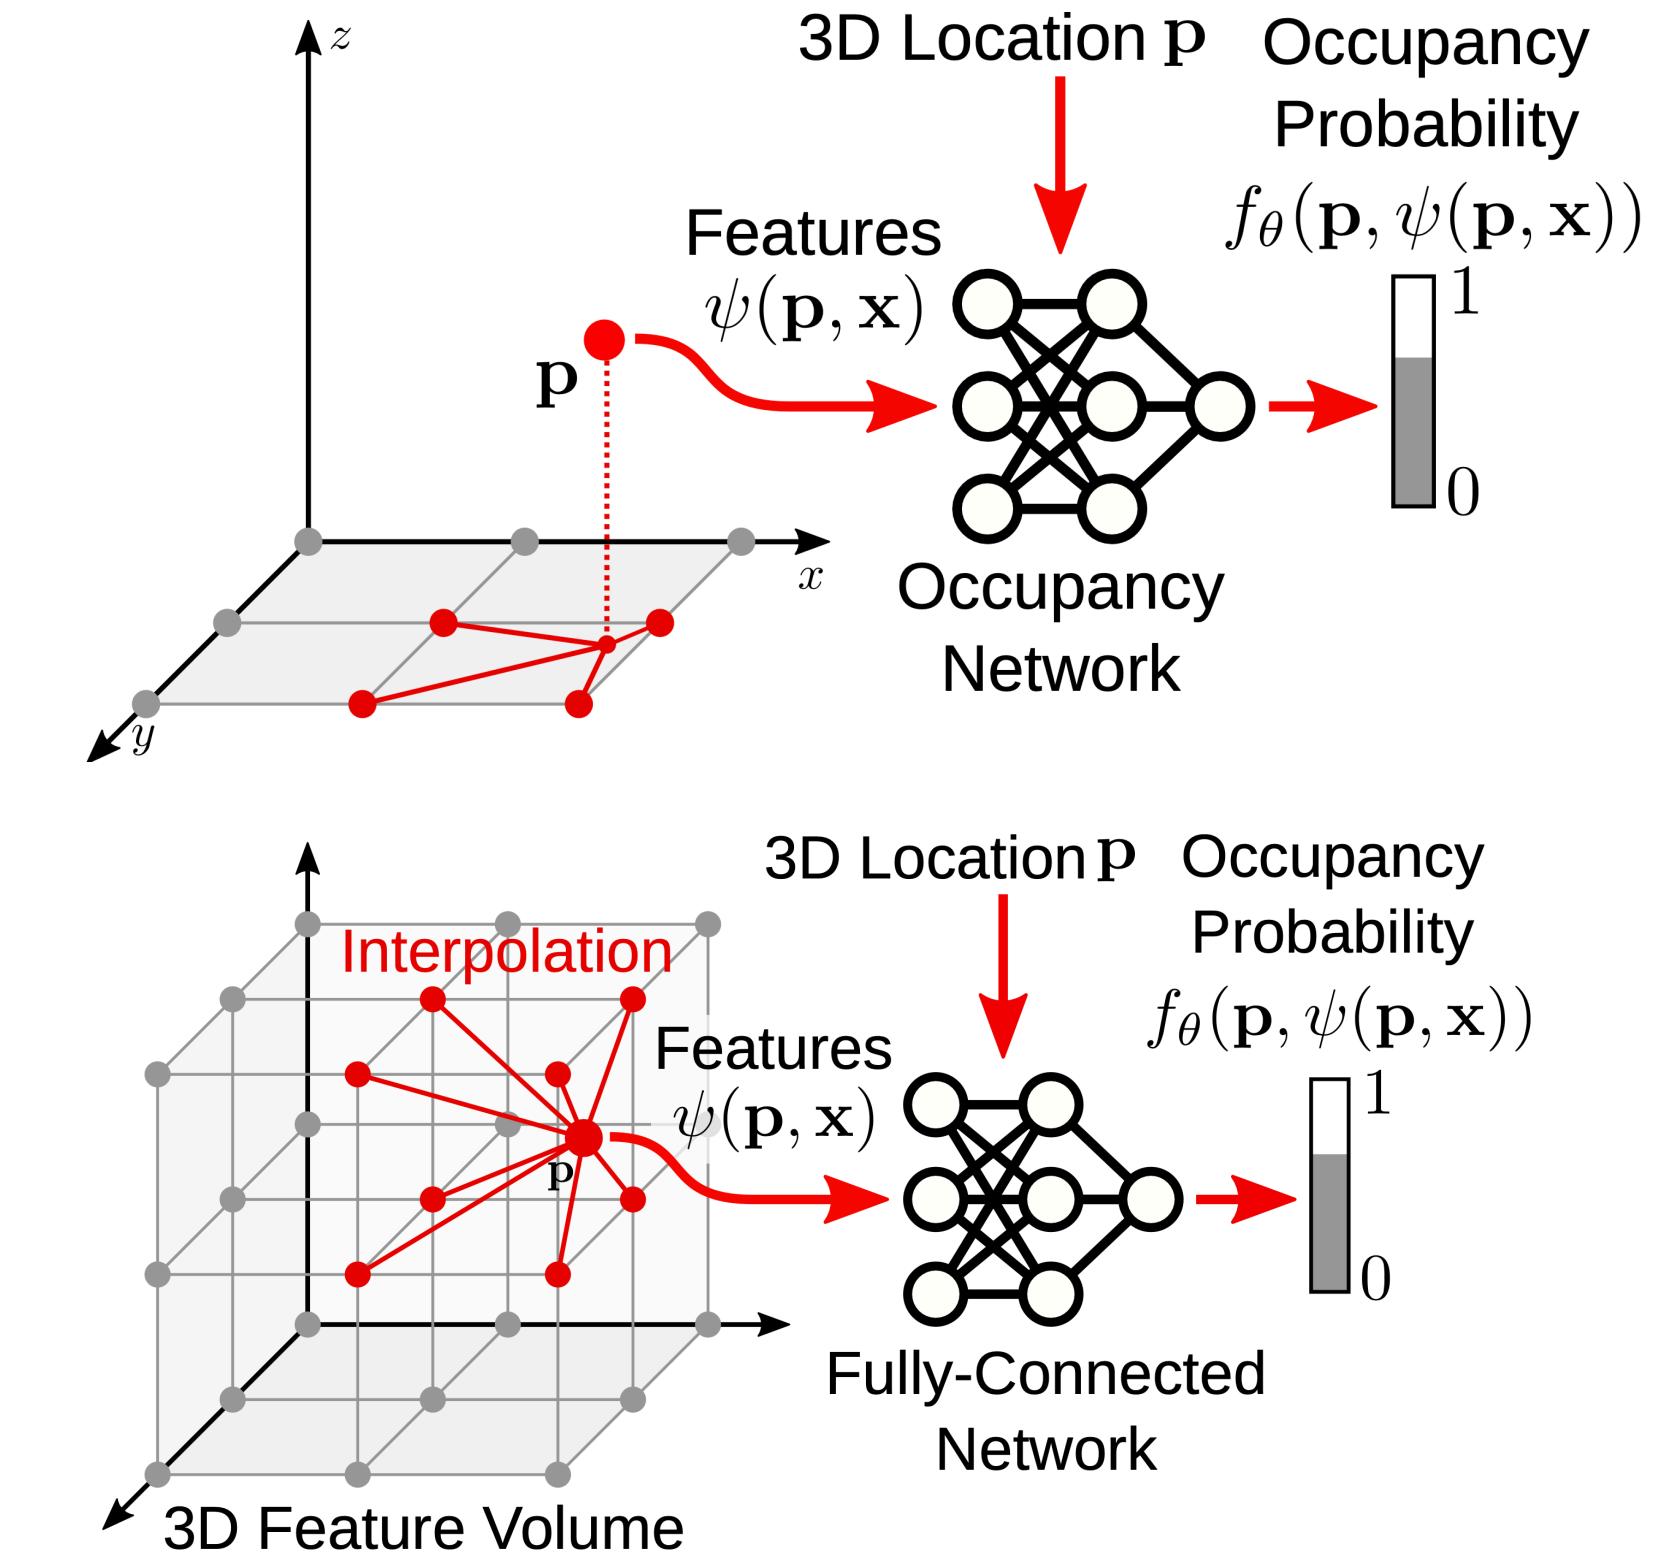
# Parameterizing Fields



Discrete Parameterizations

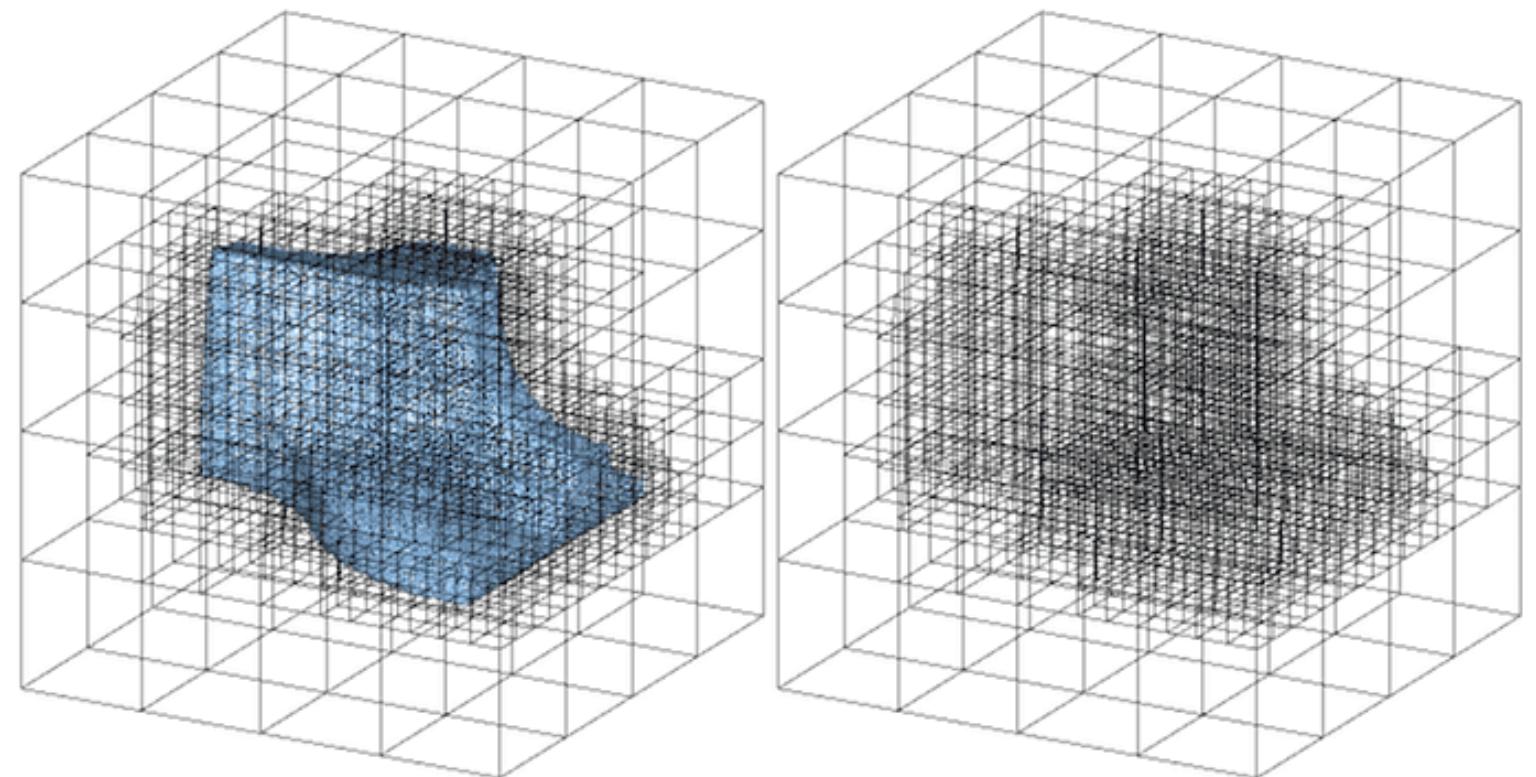
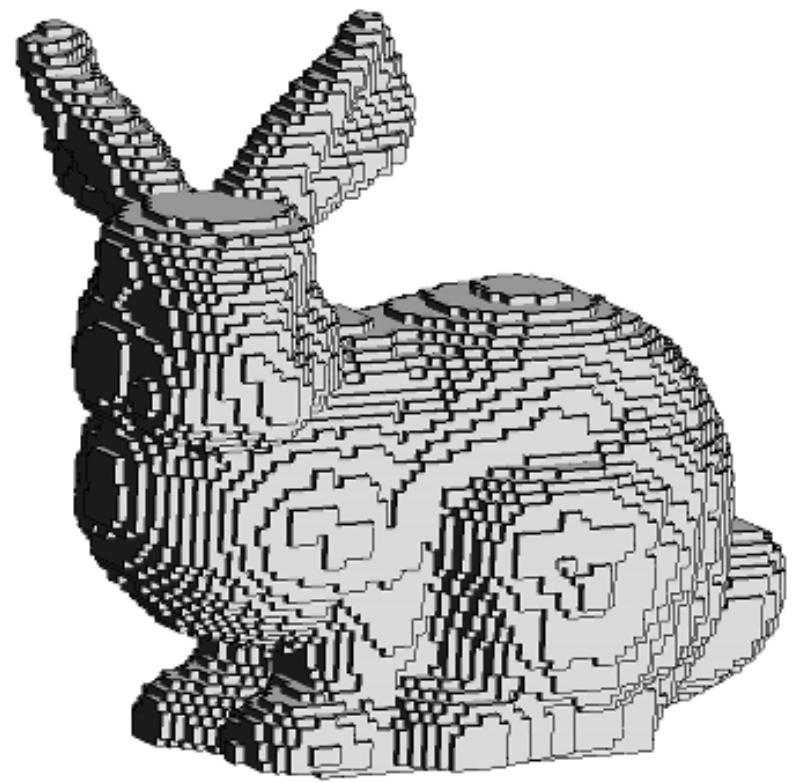


Continuous Parameterizations

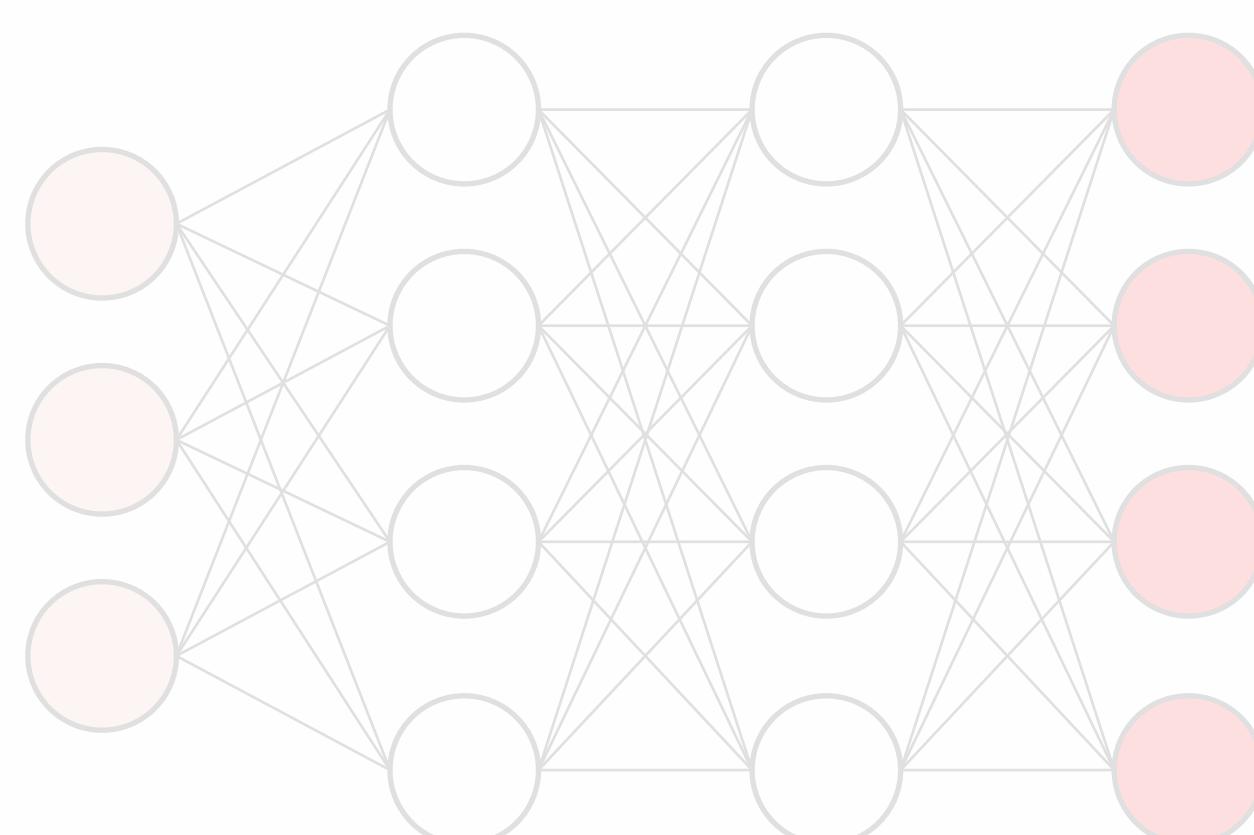


Hybrid Parameterizations

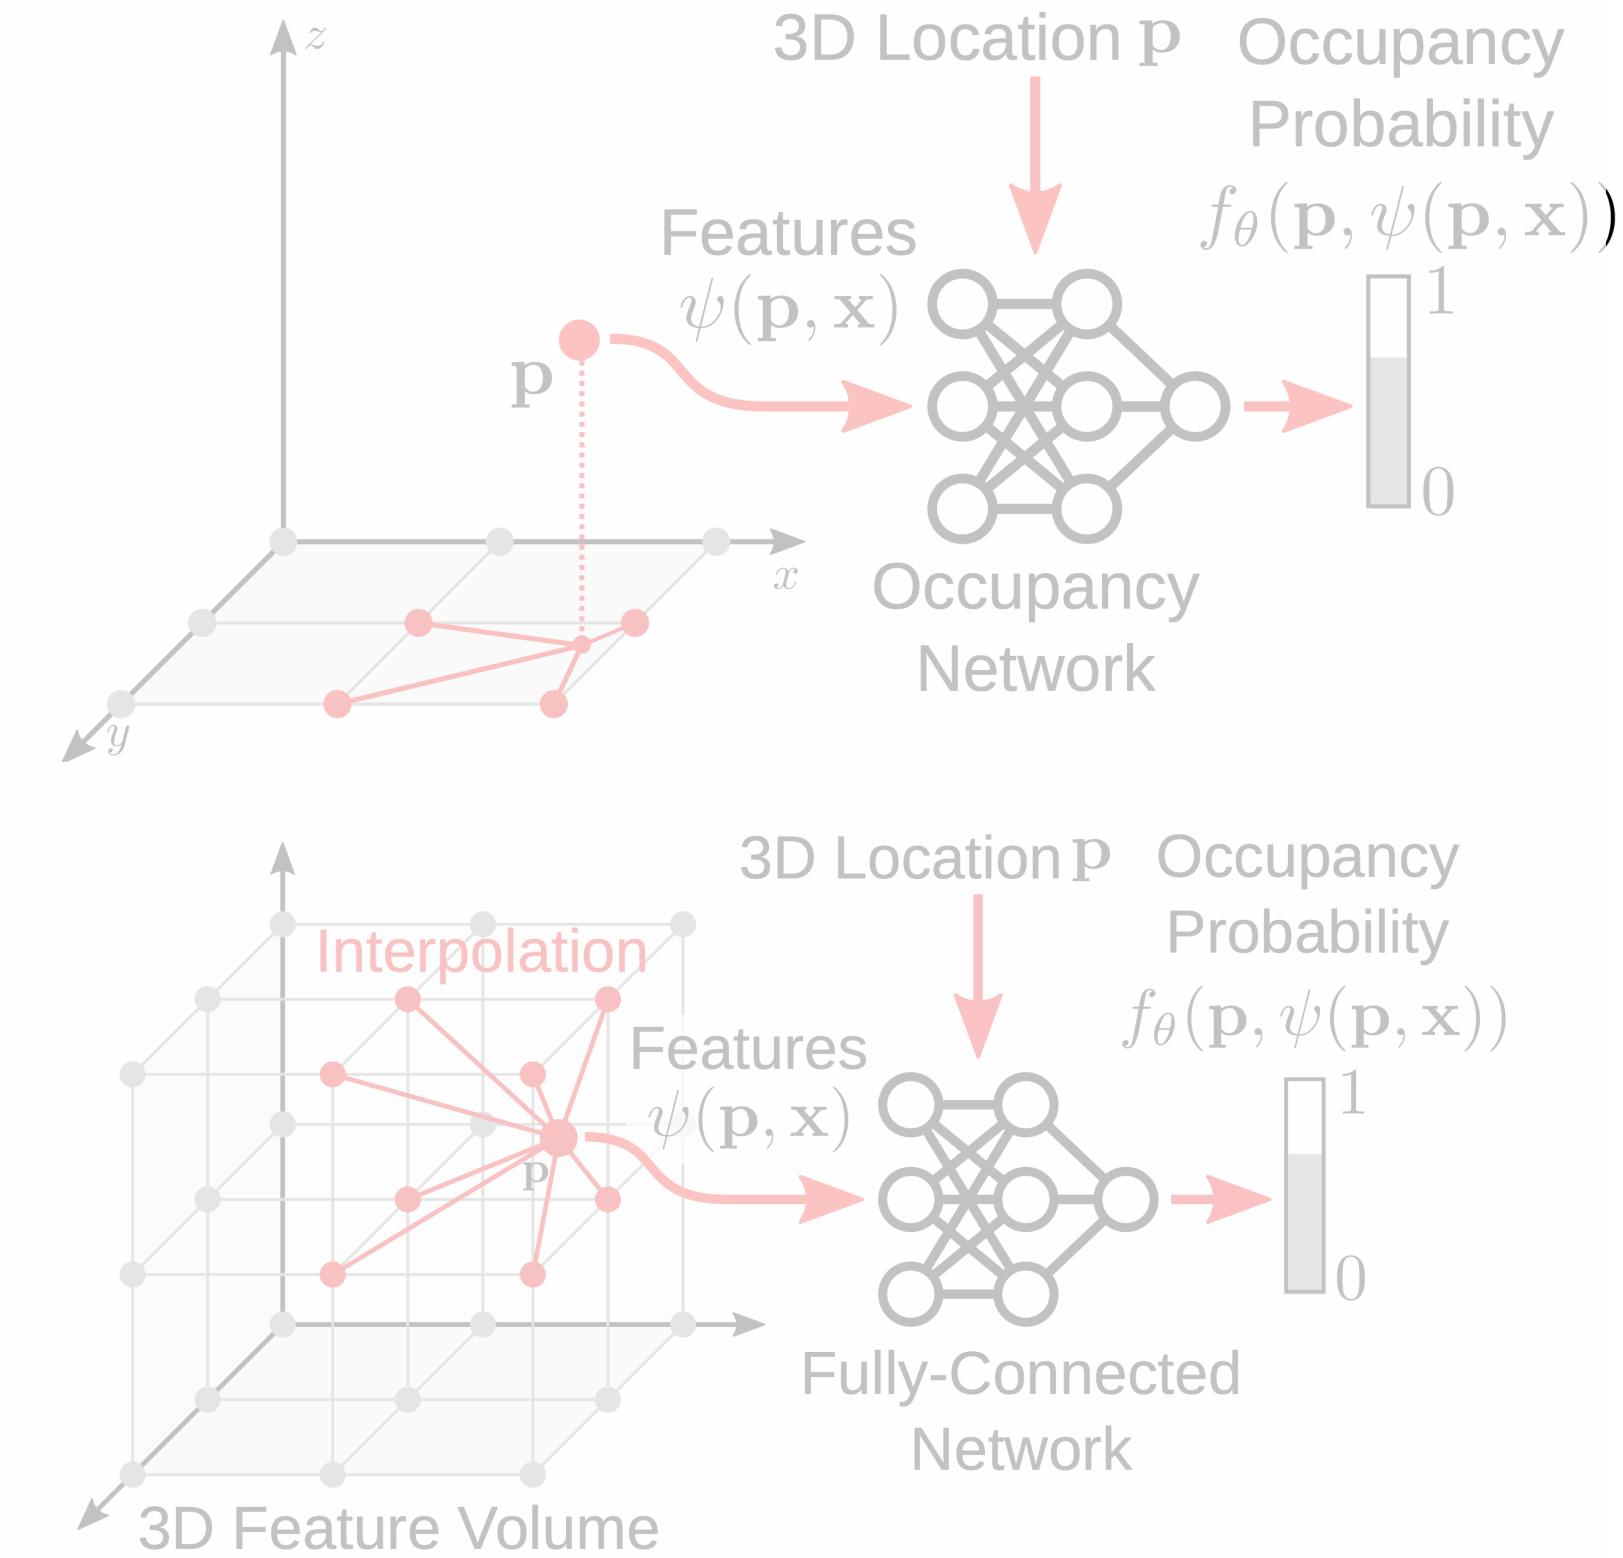
# Voxel Grids



Discrete Parameterizations

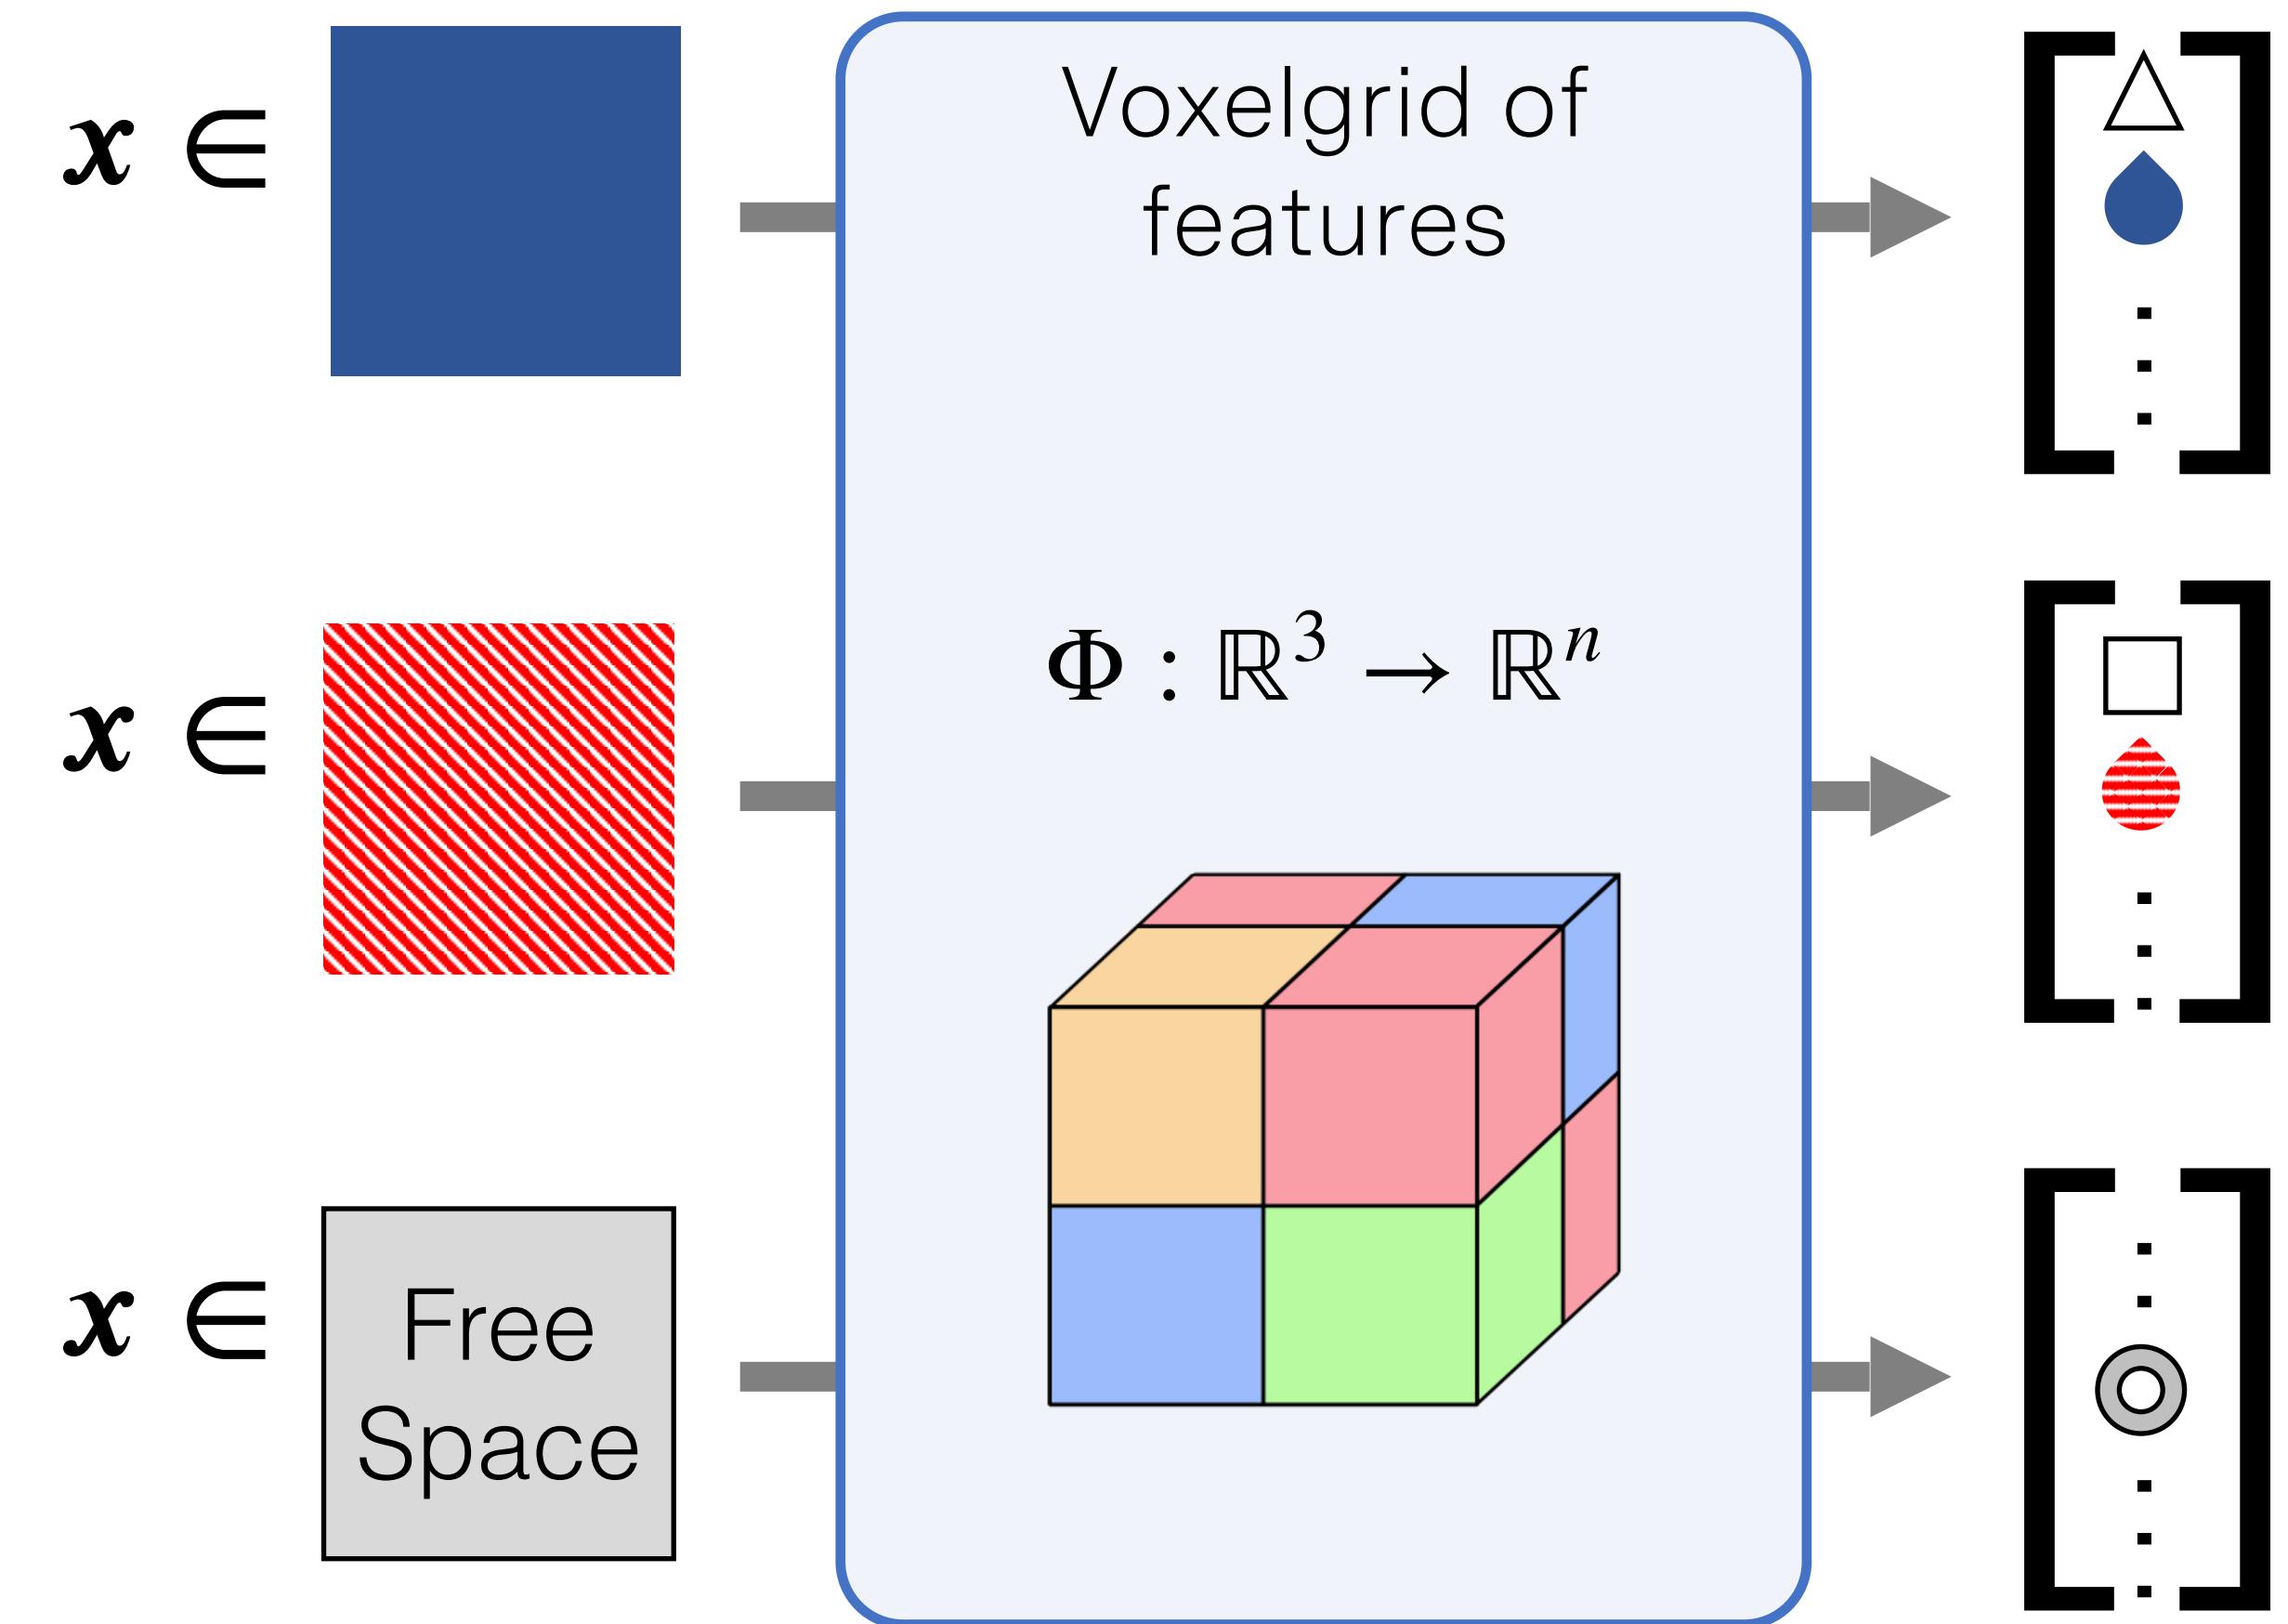
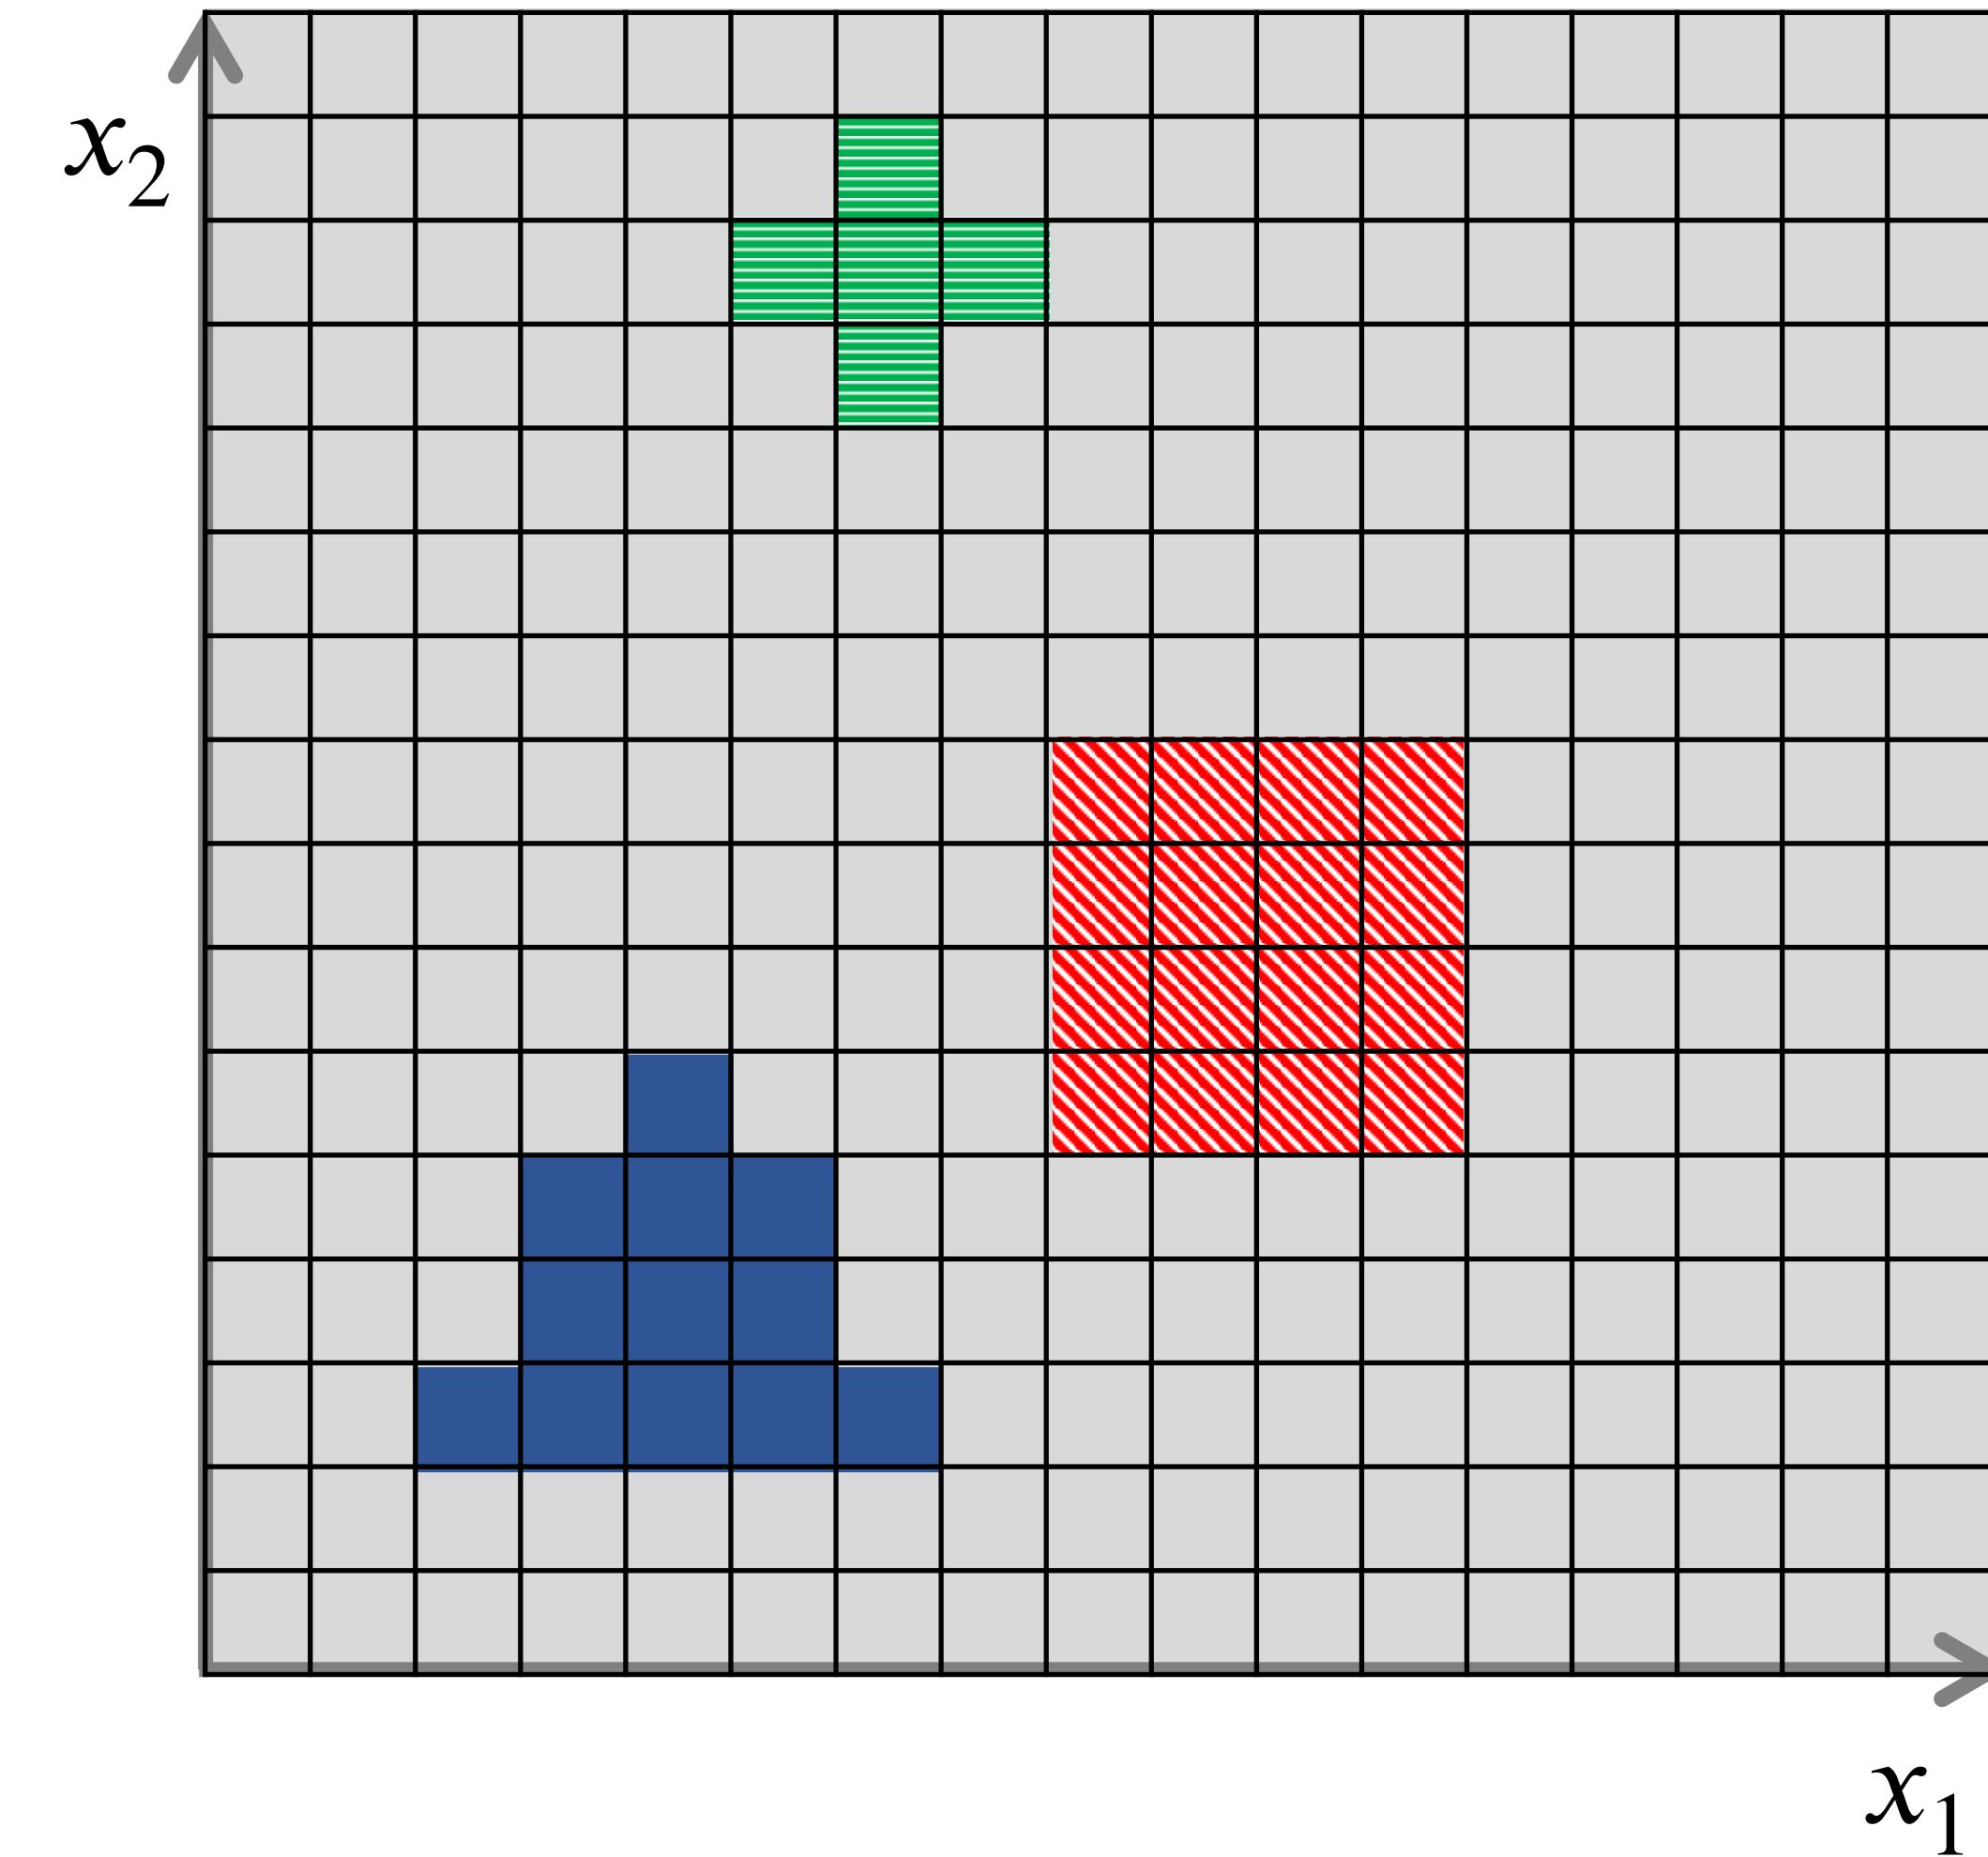


Continuous Parameterizations

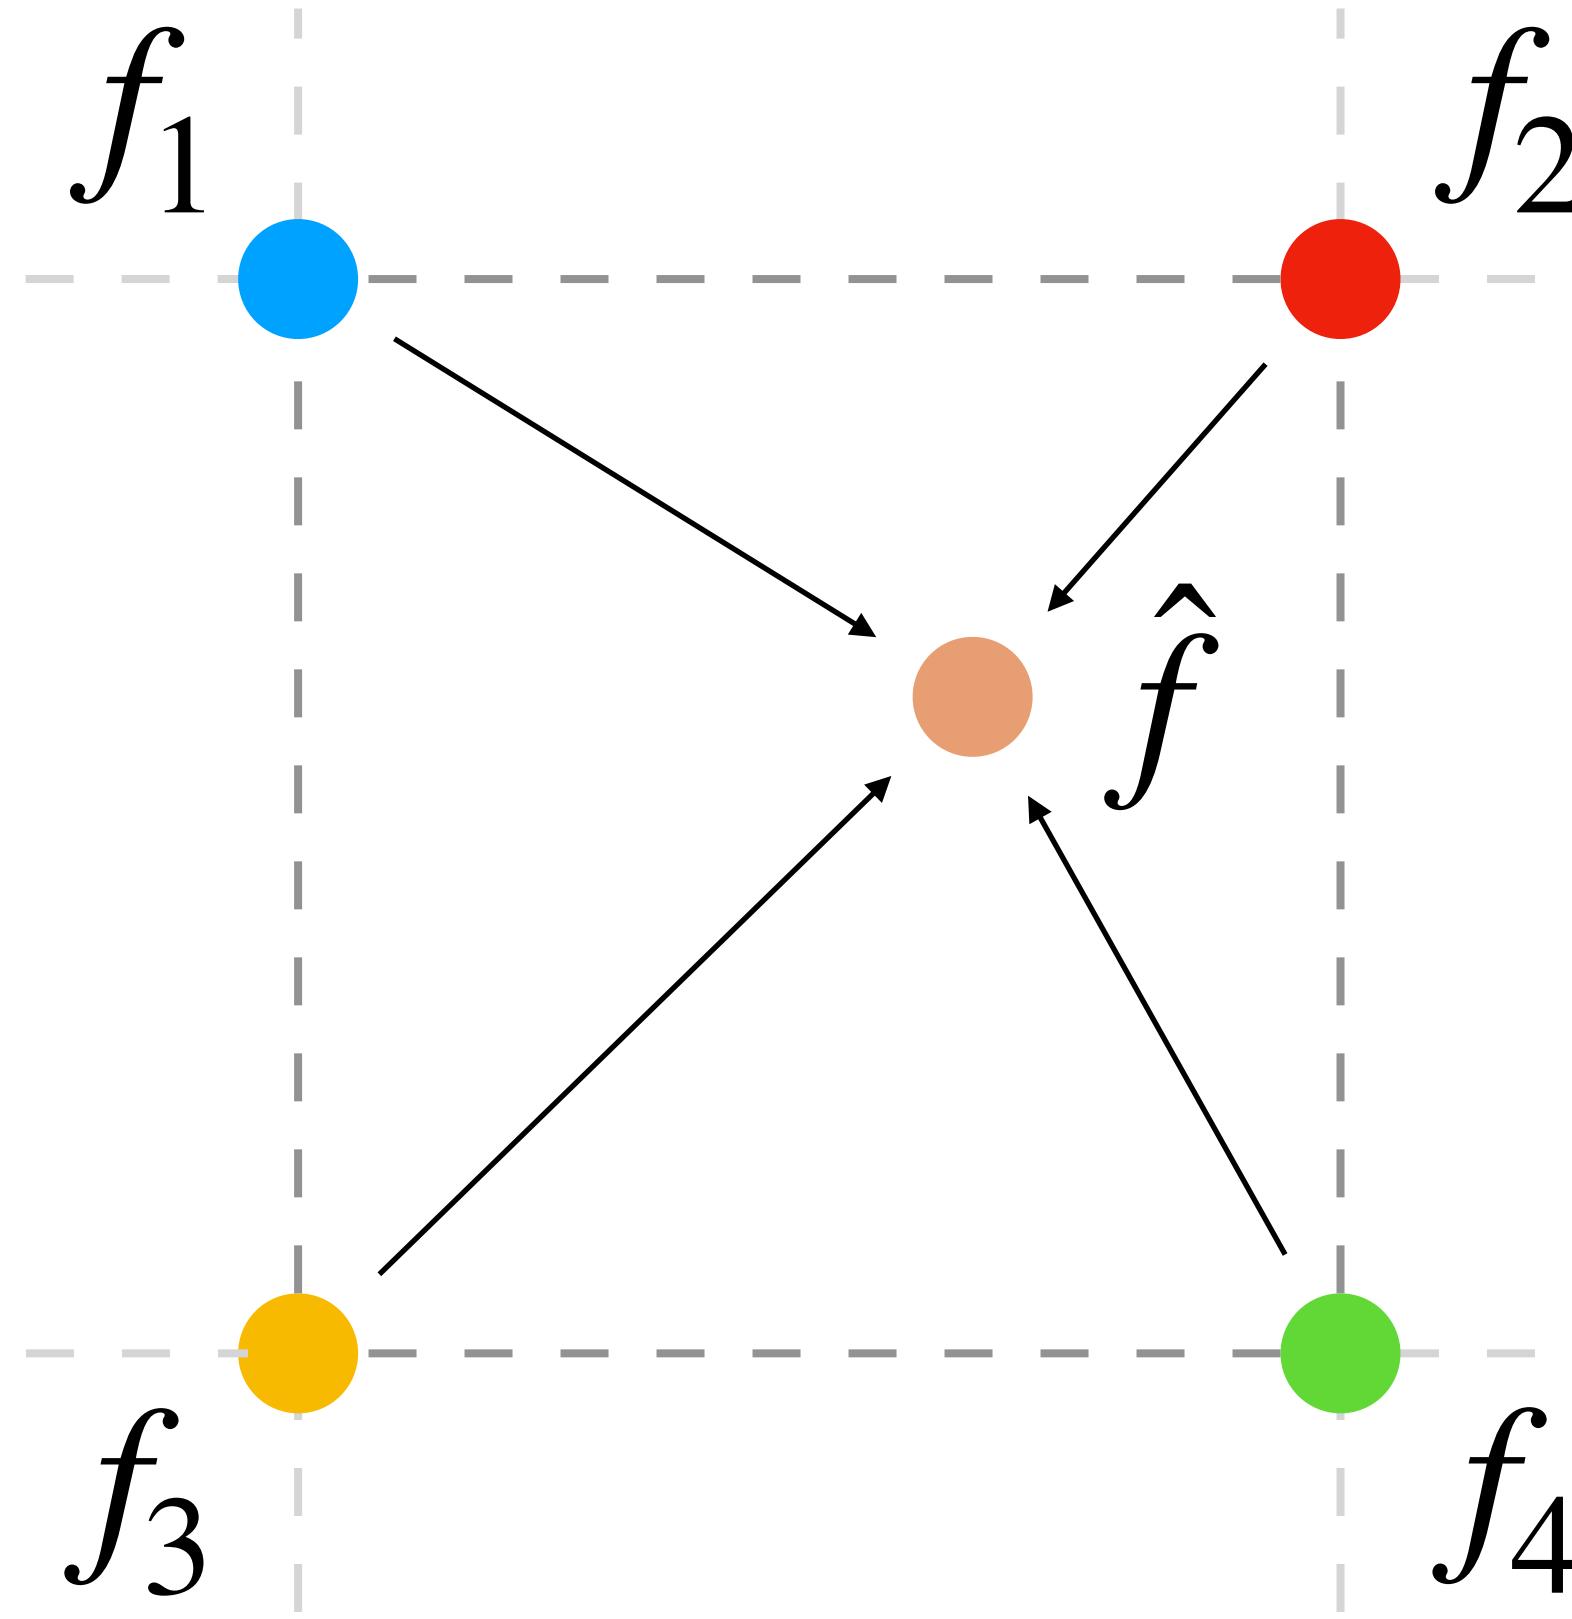


Hybrid Parameterizations

# Voxel Grids



# Interpolation: Querying *continuous* values



- Only stores values at vertex locations, i.e. corners
- Values at intermediate coordinates are defined via interpolation w/ some kernel  $k$

$$\hat{f}(\mathbf{x}) = \sum_{i=1}^N f_i k(\mathbf{x}, \mathbf{x}_i)$$

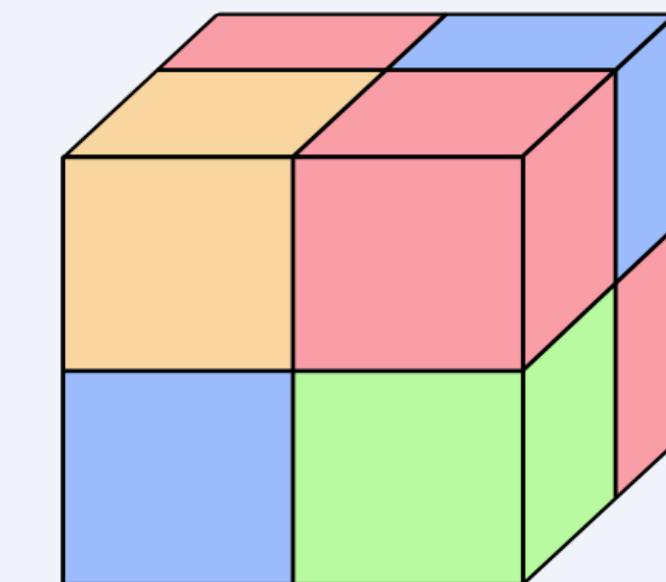
Interpolation in a 2D voxel grid.

# Voxel grids

- For  $d$  spatial dimensions and resolution  $n$ ,  
memory grows  $O(n^d)$
- Fast sampling:  $d$ -linear interpolation  
= index into array  $2^d$  times & weighted sum.
- Convenient processing as it exposes *locality*. Can  
easily run convolutions, nearest-neighbor lookups,  
etc.
- Intractable in higher dimensions :/

Voxelgrid of  
features

$$\Phi : \mathbb{R}^3 \rightarrow \mathbb{R}^n$$



# Multi-Resolution Representations

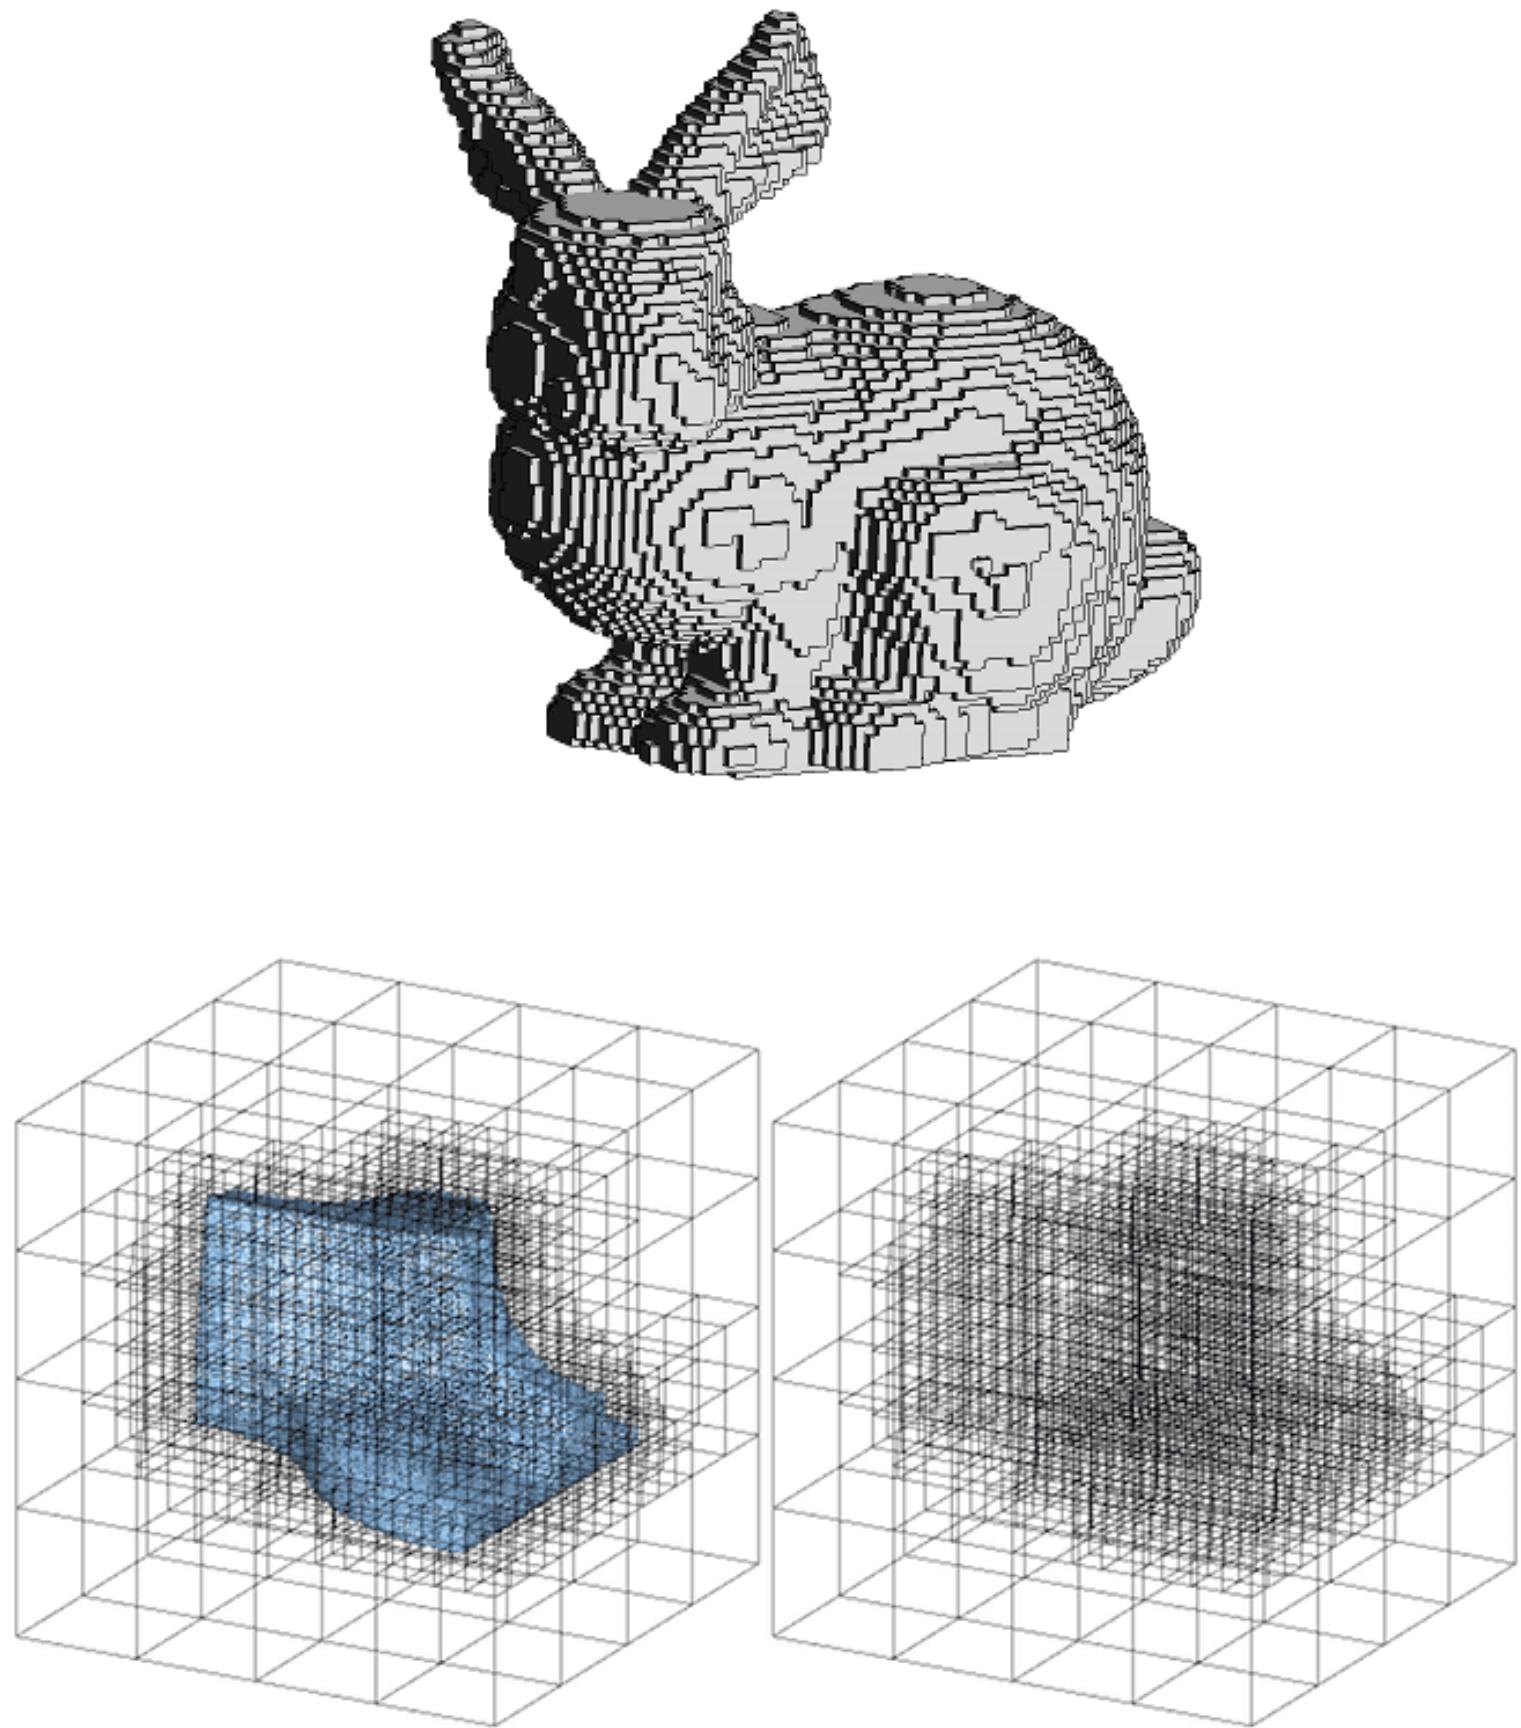
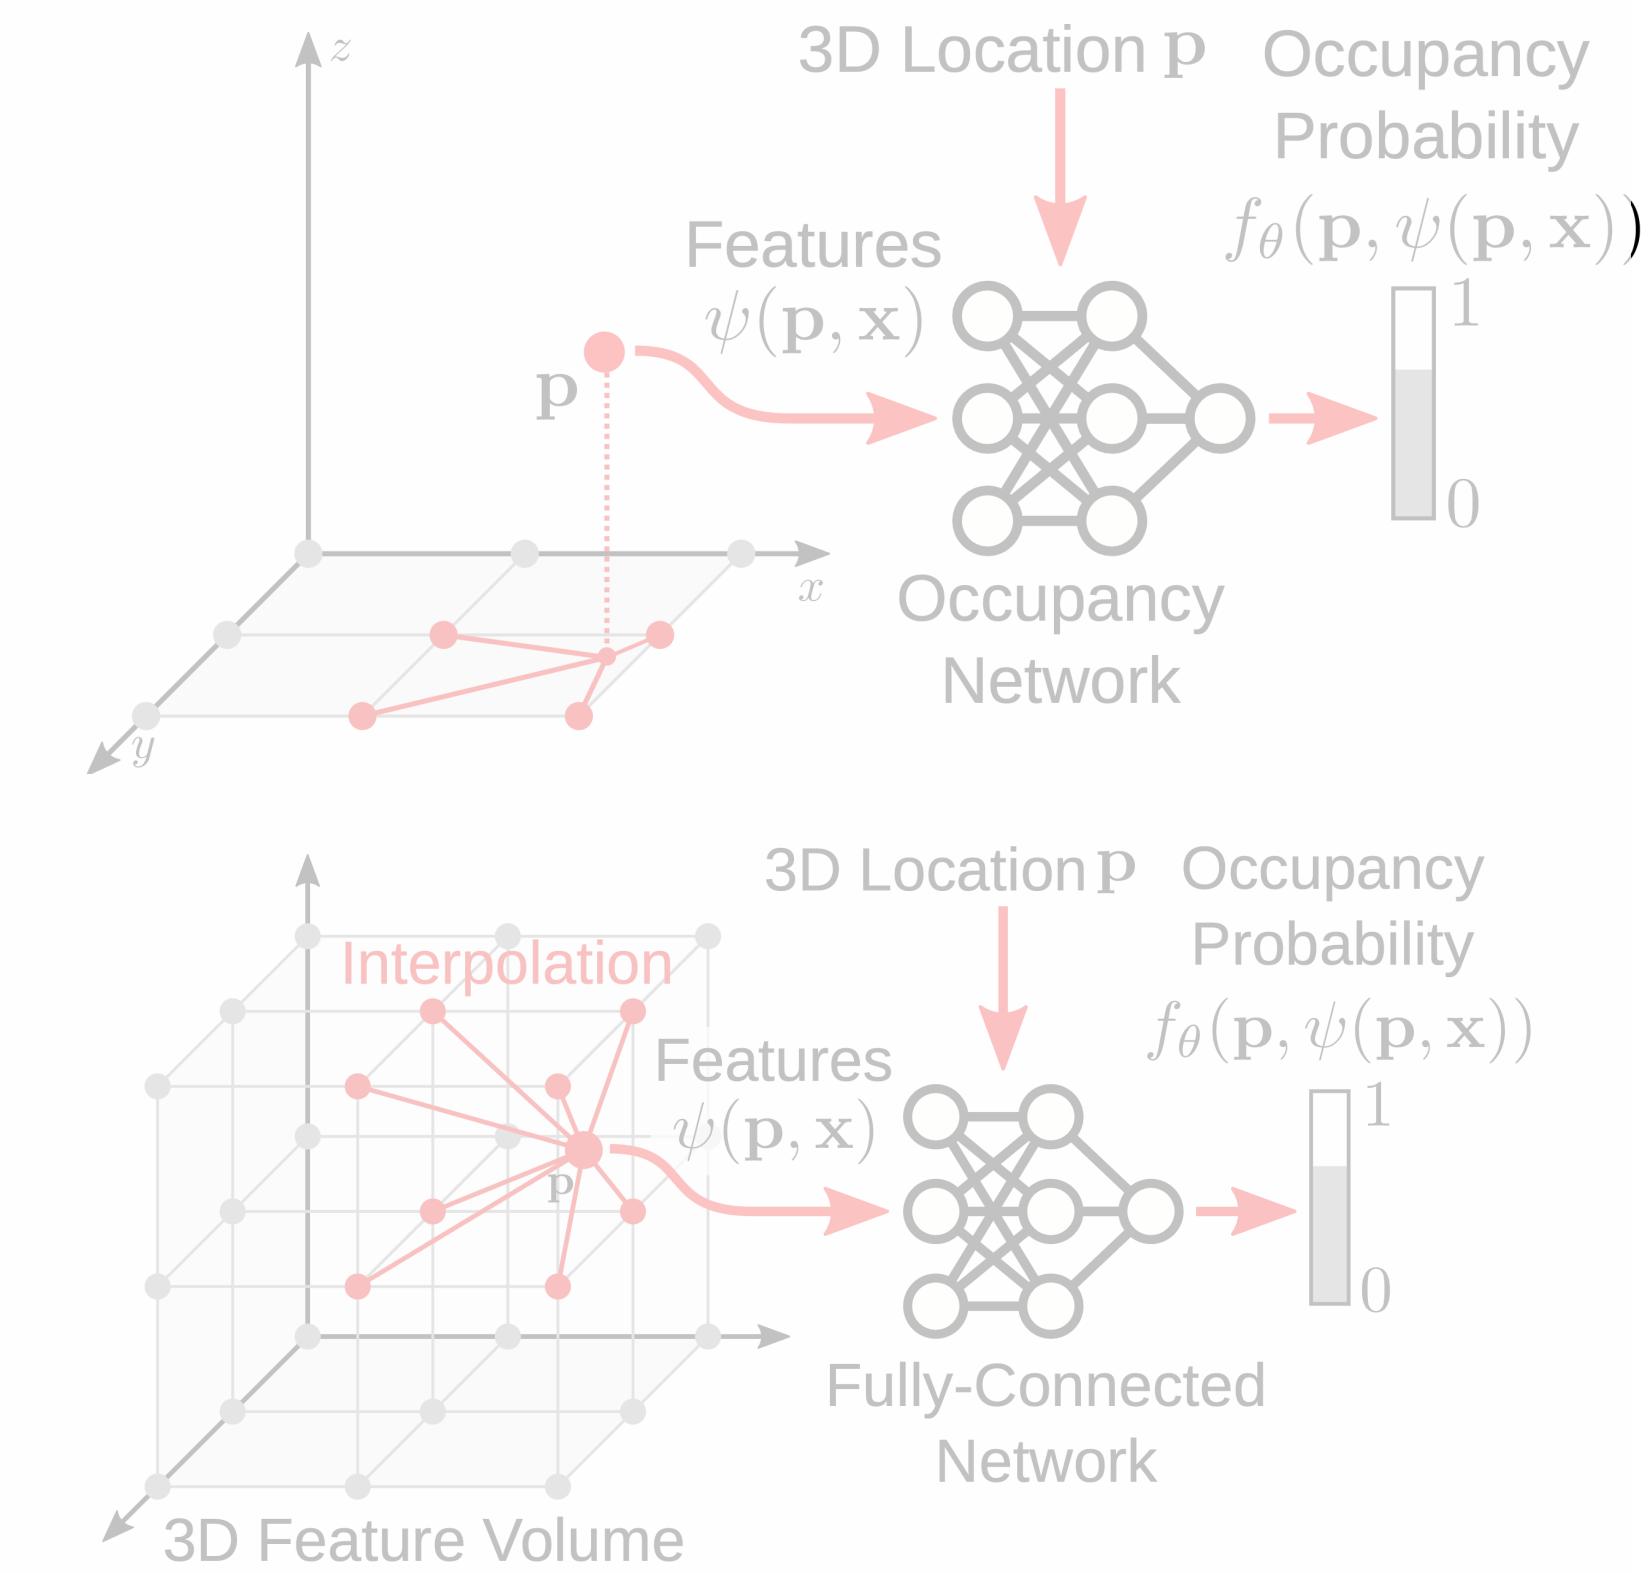
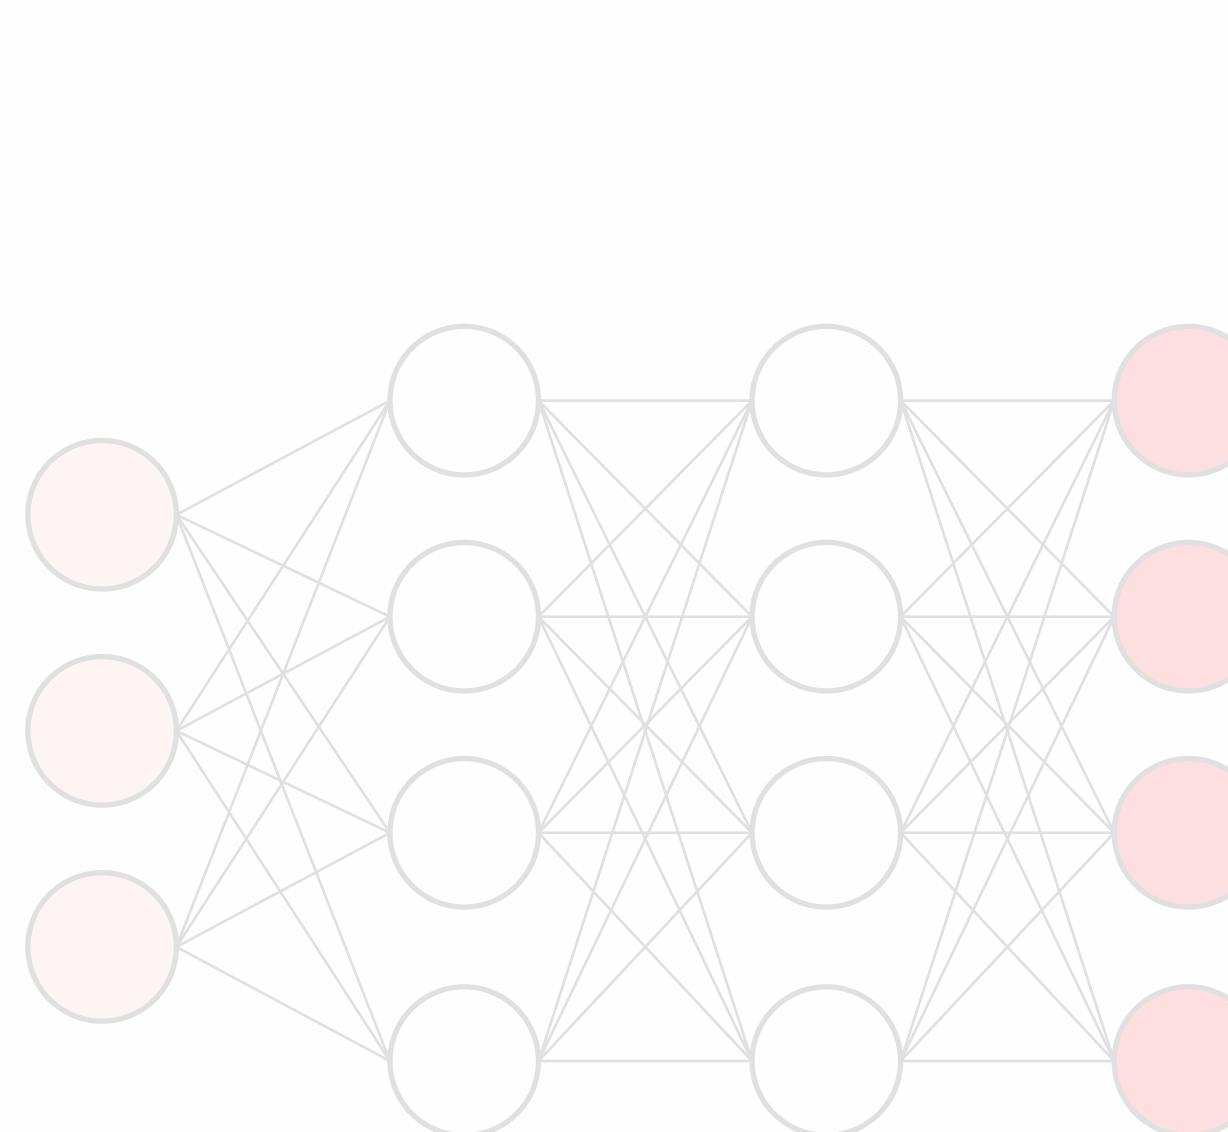


Image Source: <https://doc.cgal.org/latest/Orthree/index.html>

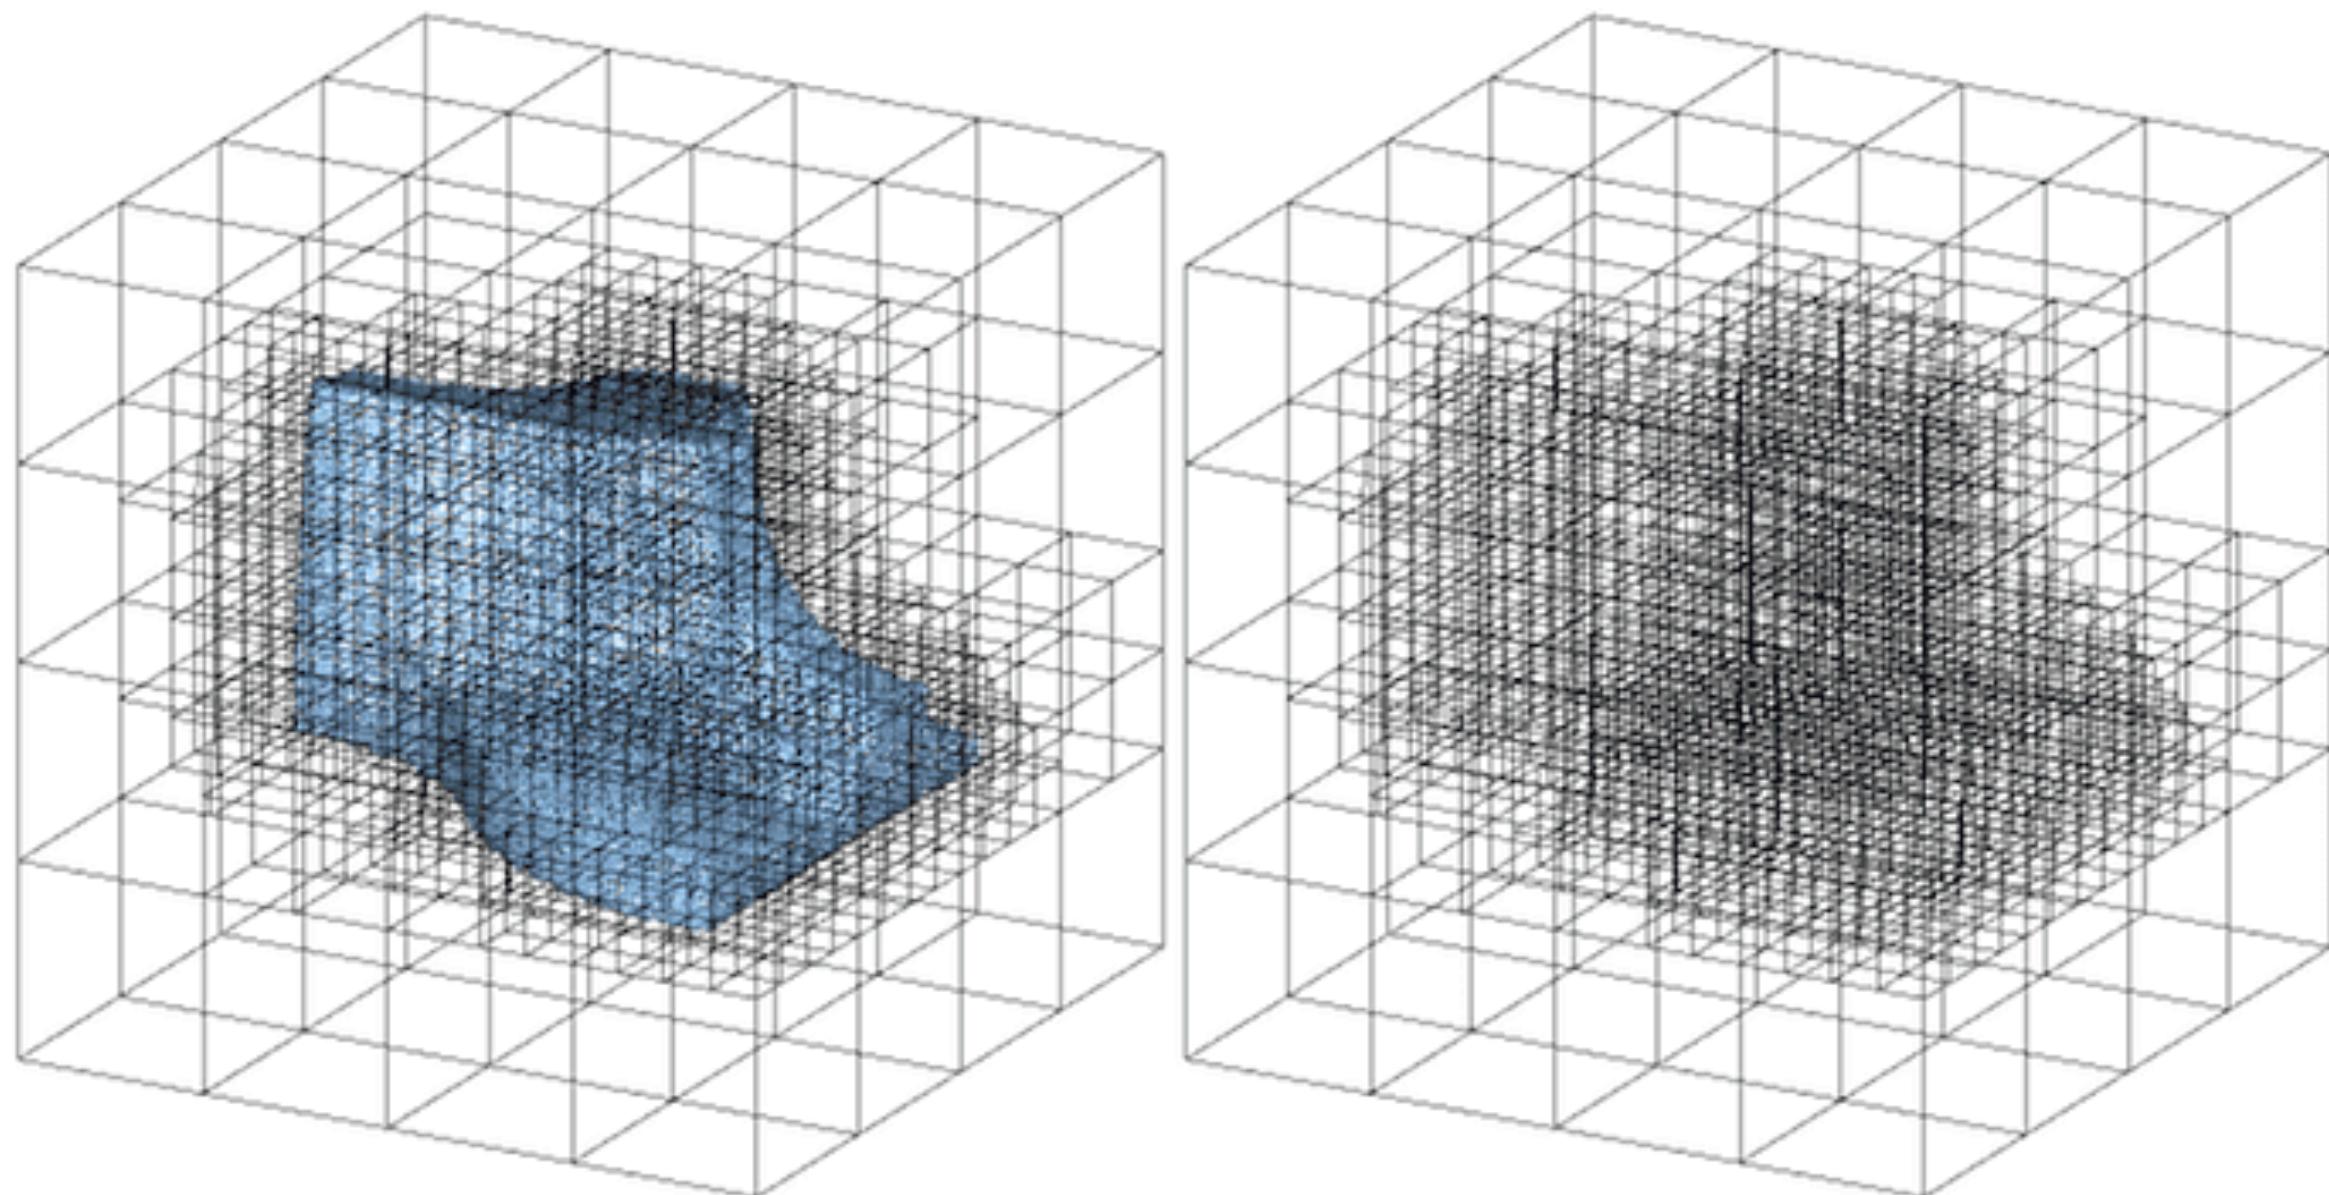
**Discrete Parameterizations**

**Continuous Parameterizations**

**Hybrid Parameterizations**



# Octrees



Divide volume into 8 voxels

If any box has more than  $N$  points in it, divide it into 8 voxels

Do not divide voxels with zero points.

Image Source: <https://doc.cgal.org/latest/Orytree/index.html>

# Octrees

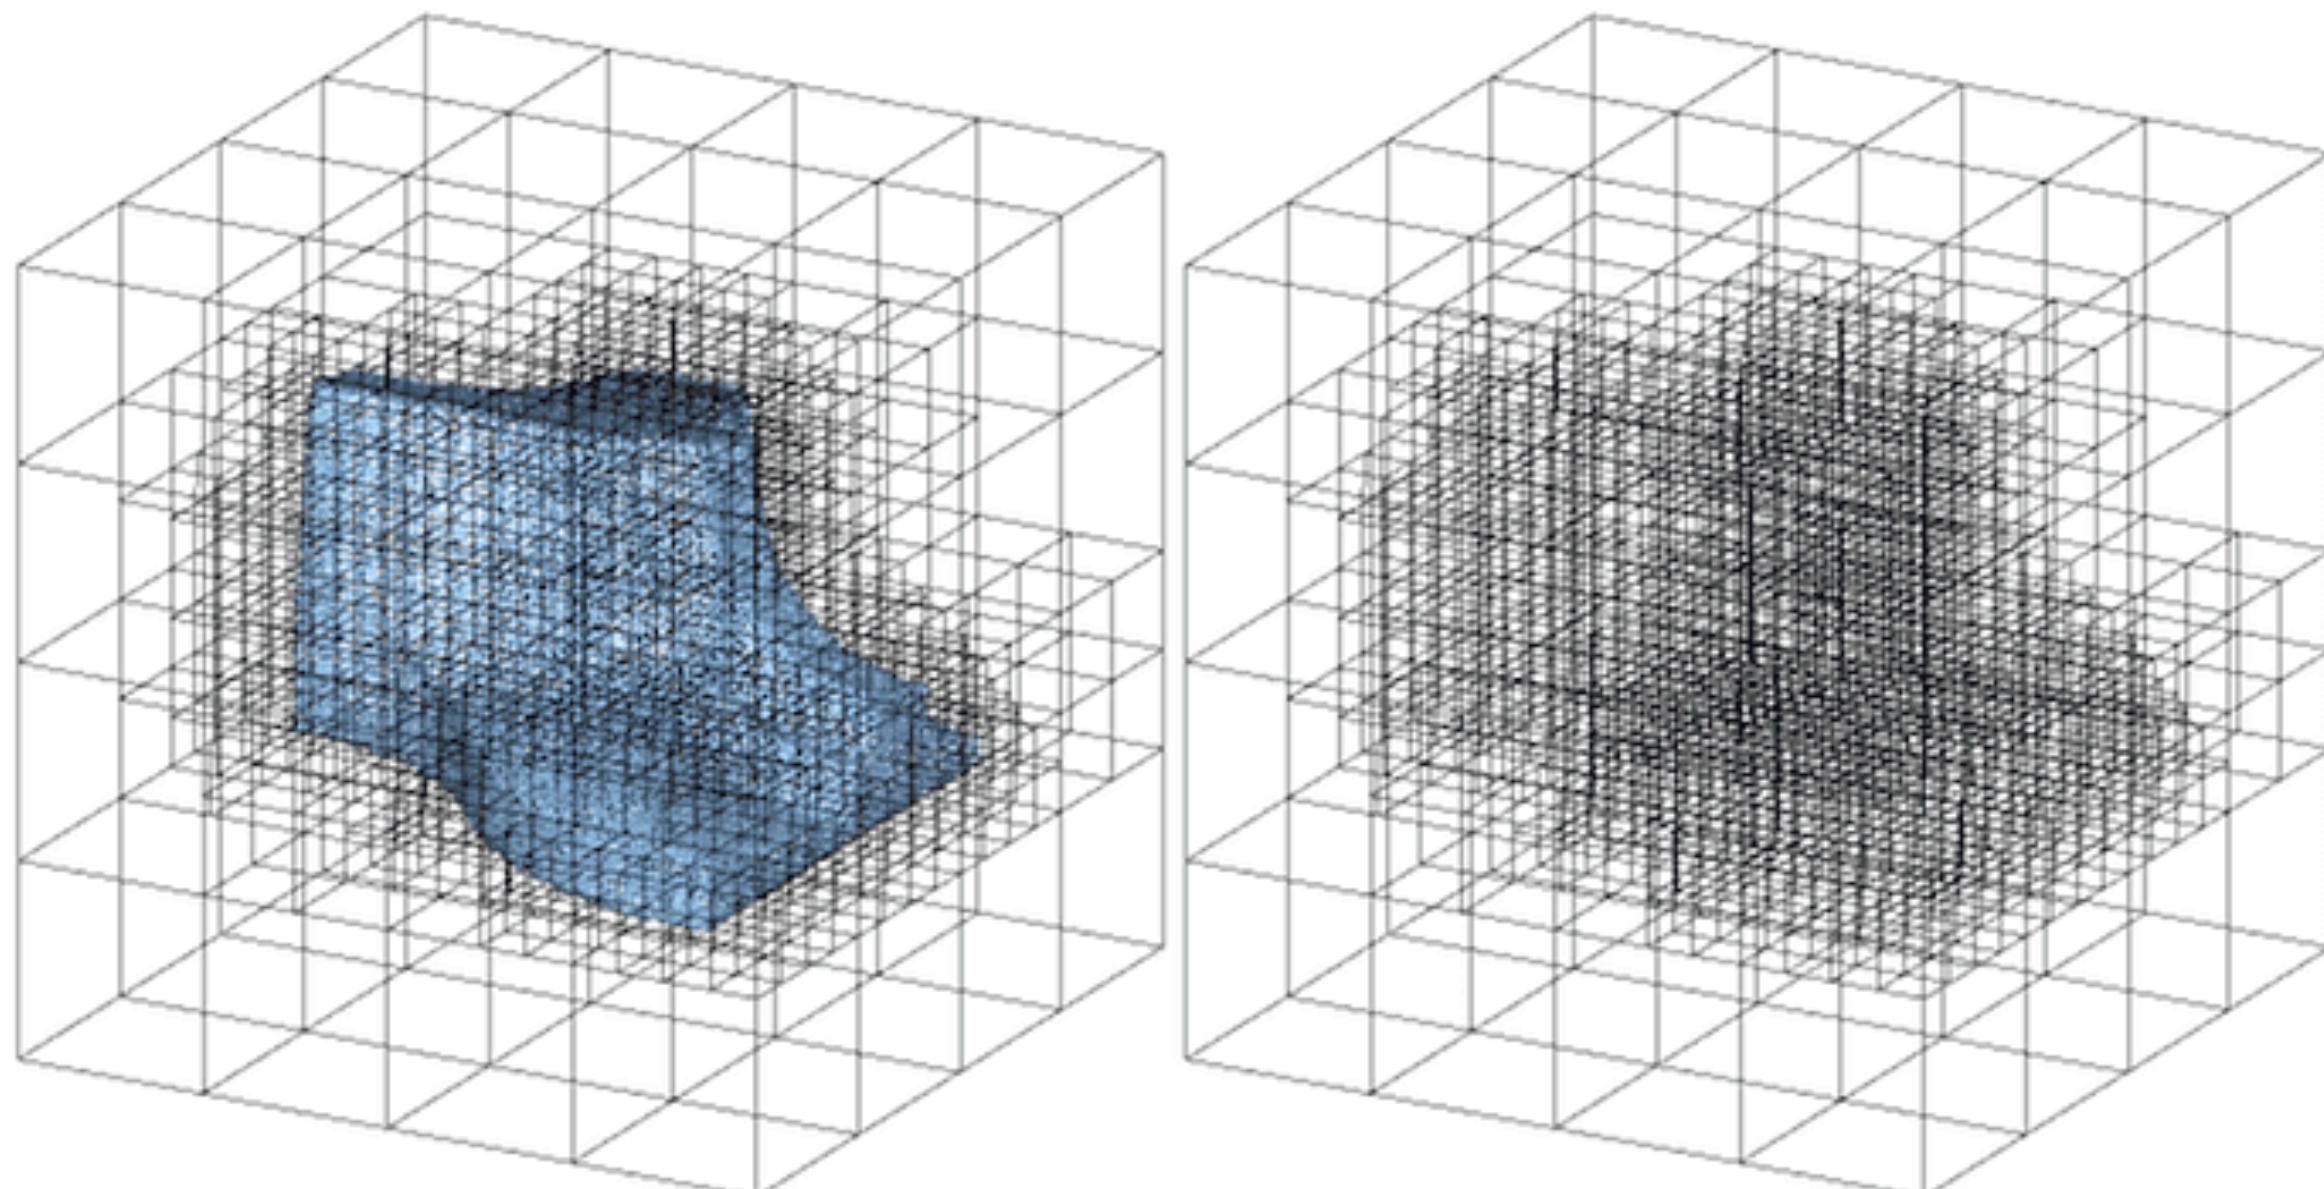


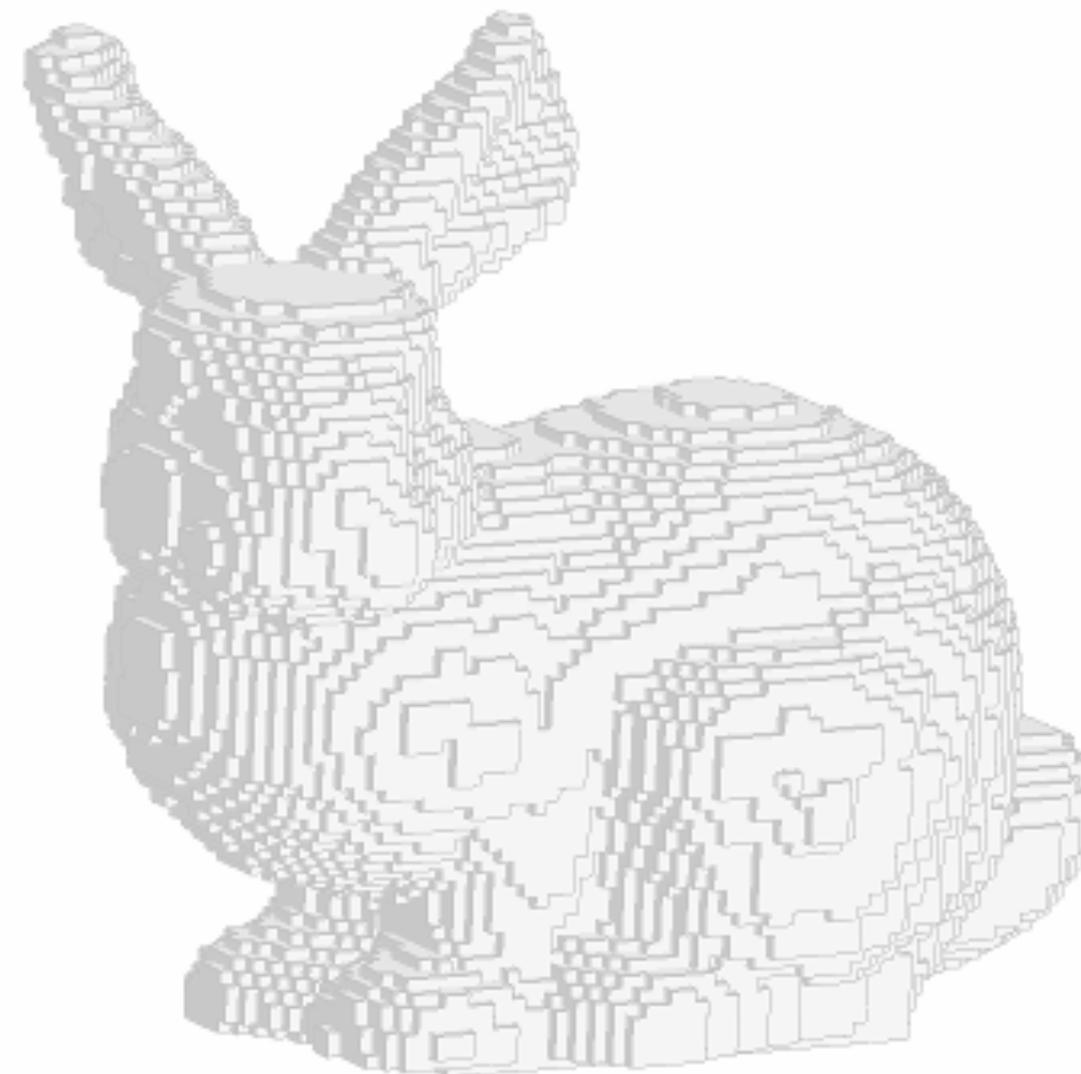
Image Source: <https://doc.cgal.org/latest/Orytree/index.html>

Saves memory (less resolution in empty parts)

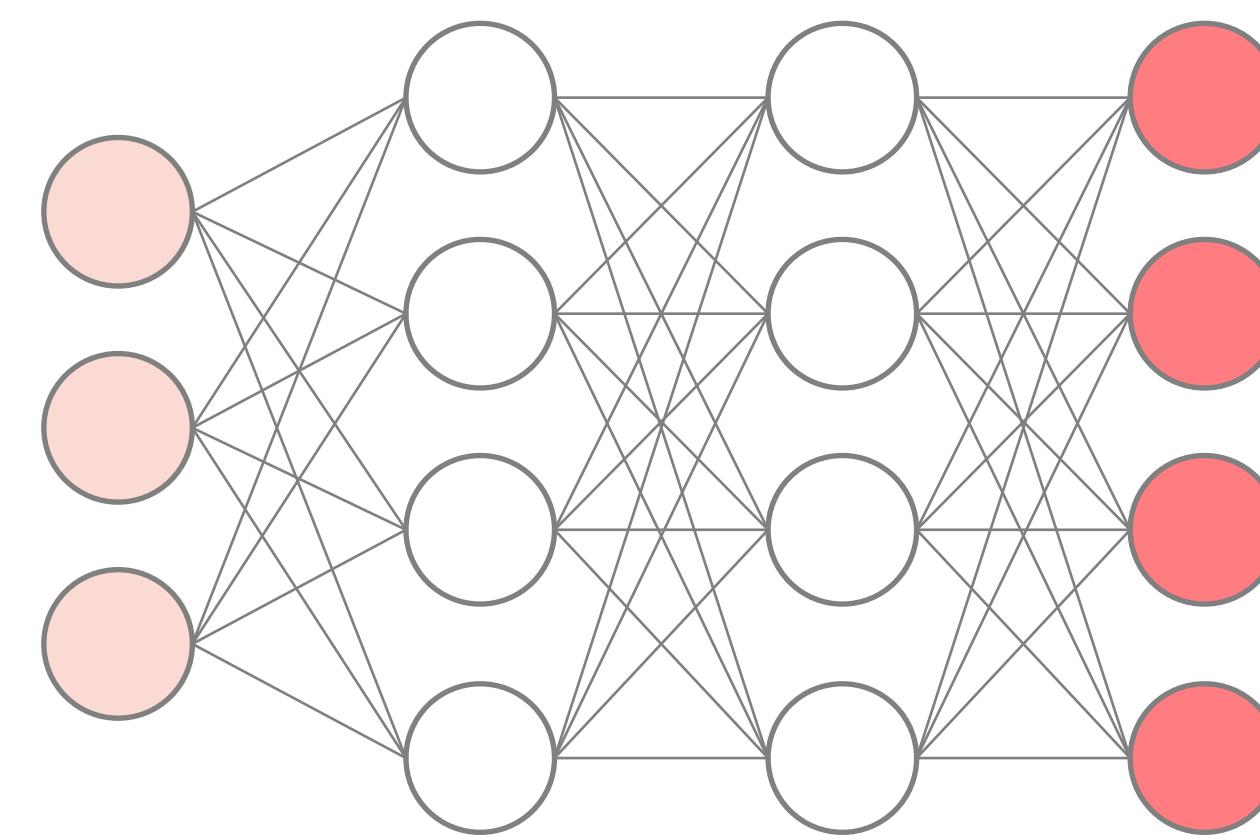
Sub-division is not differentiable:  
can't easily build neural network  
that outputs Octree

To sample point, have to traverse  
multi-resolution hierarchy. Still  
fast, though.

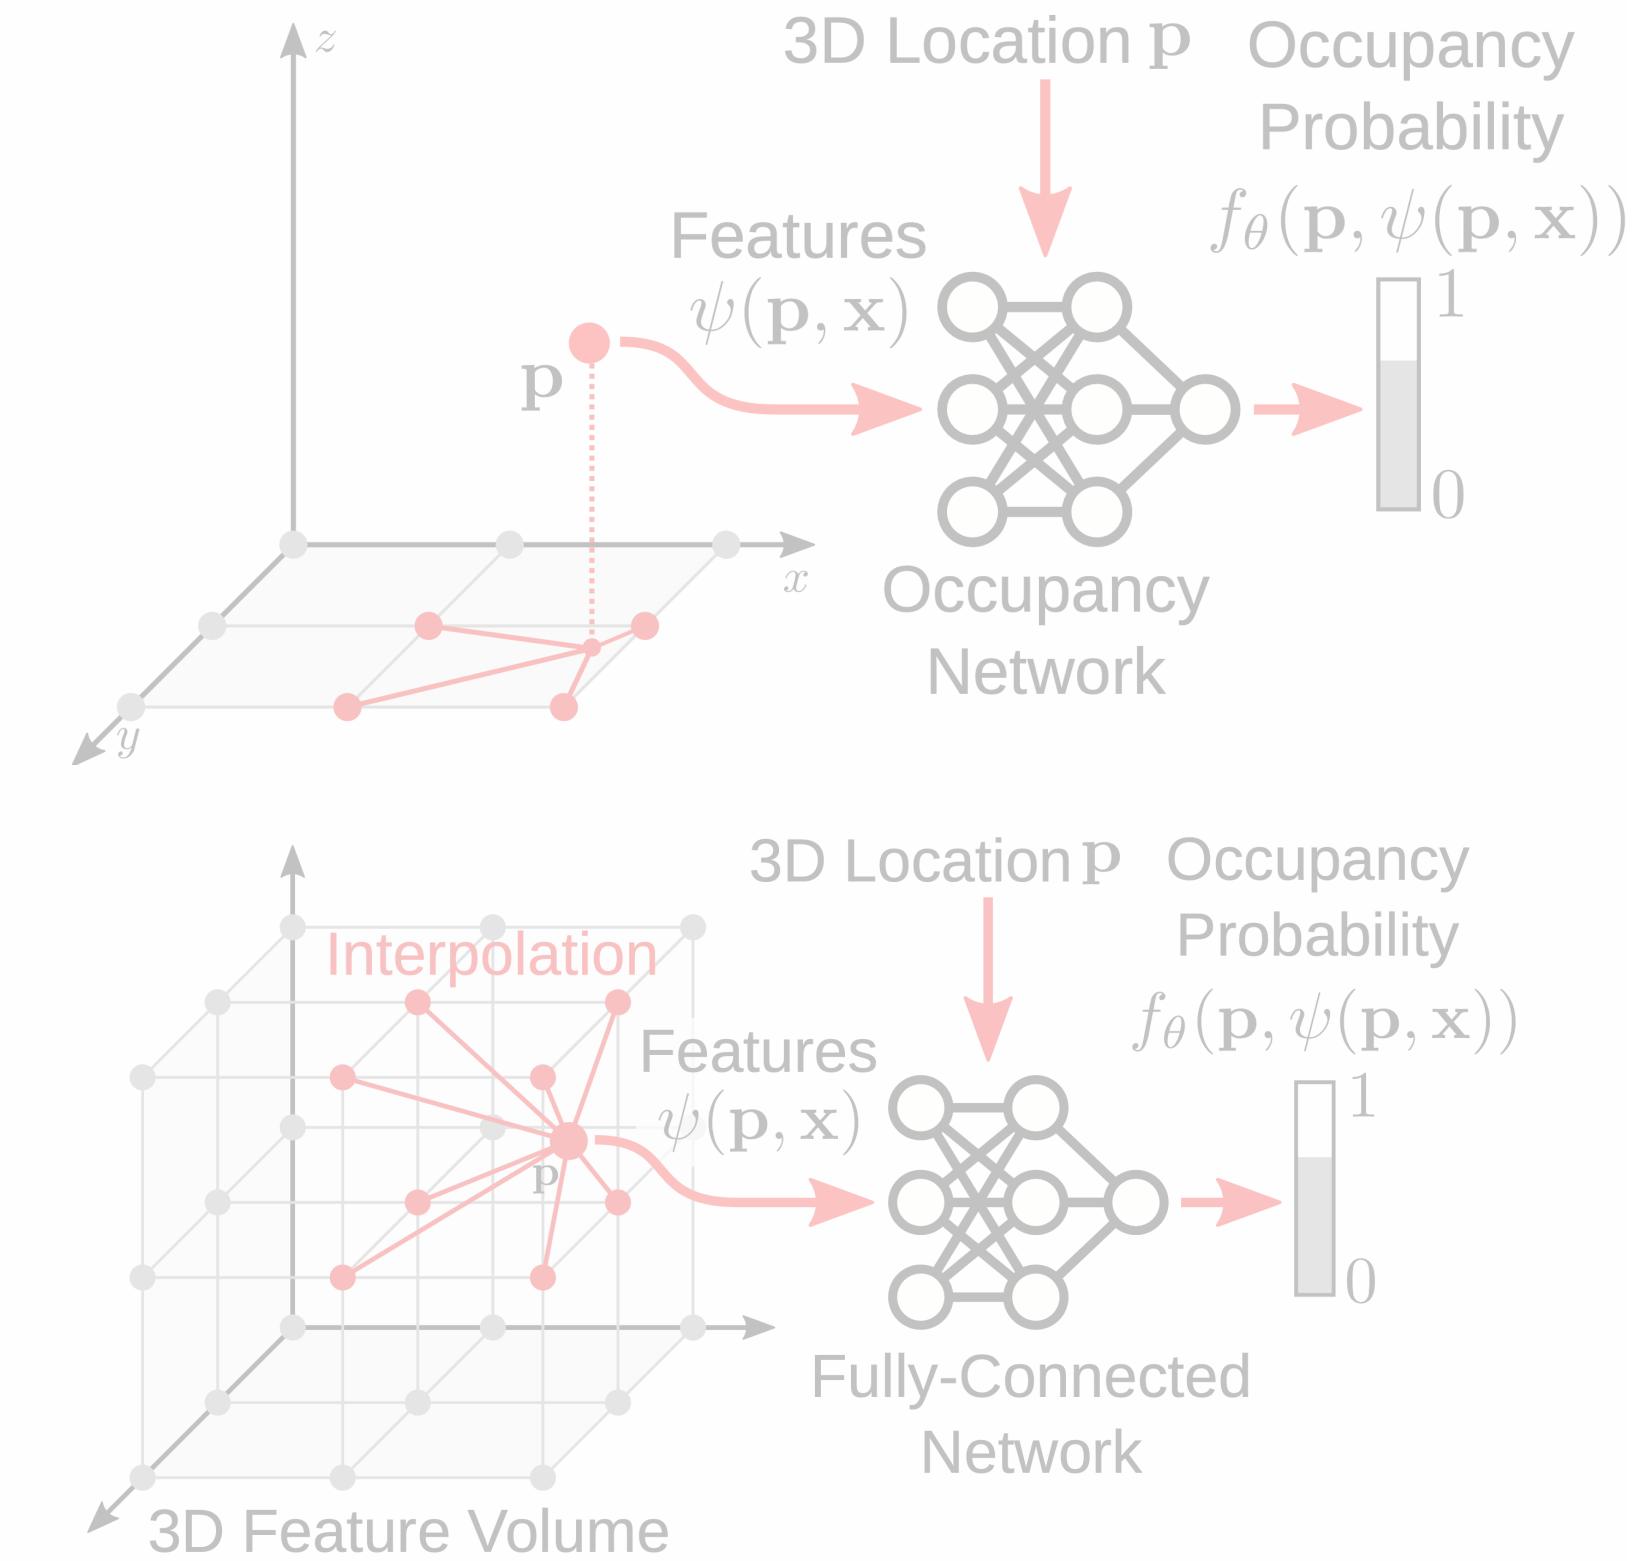
# Neural Fields



Discrete Parameterizations

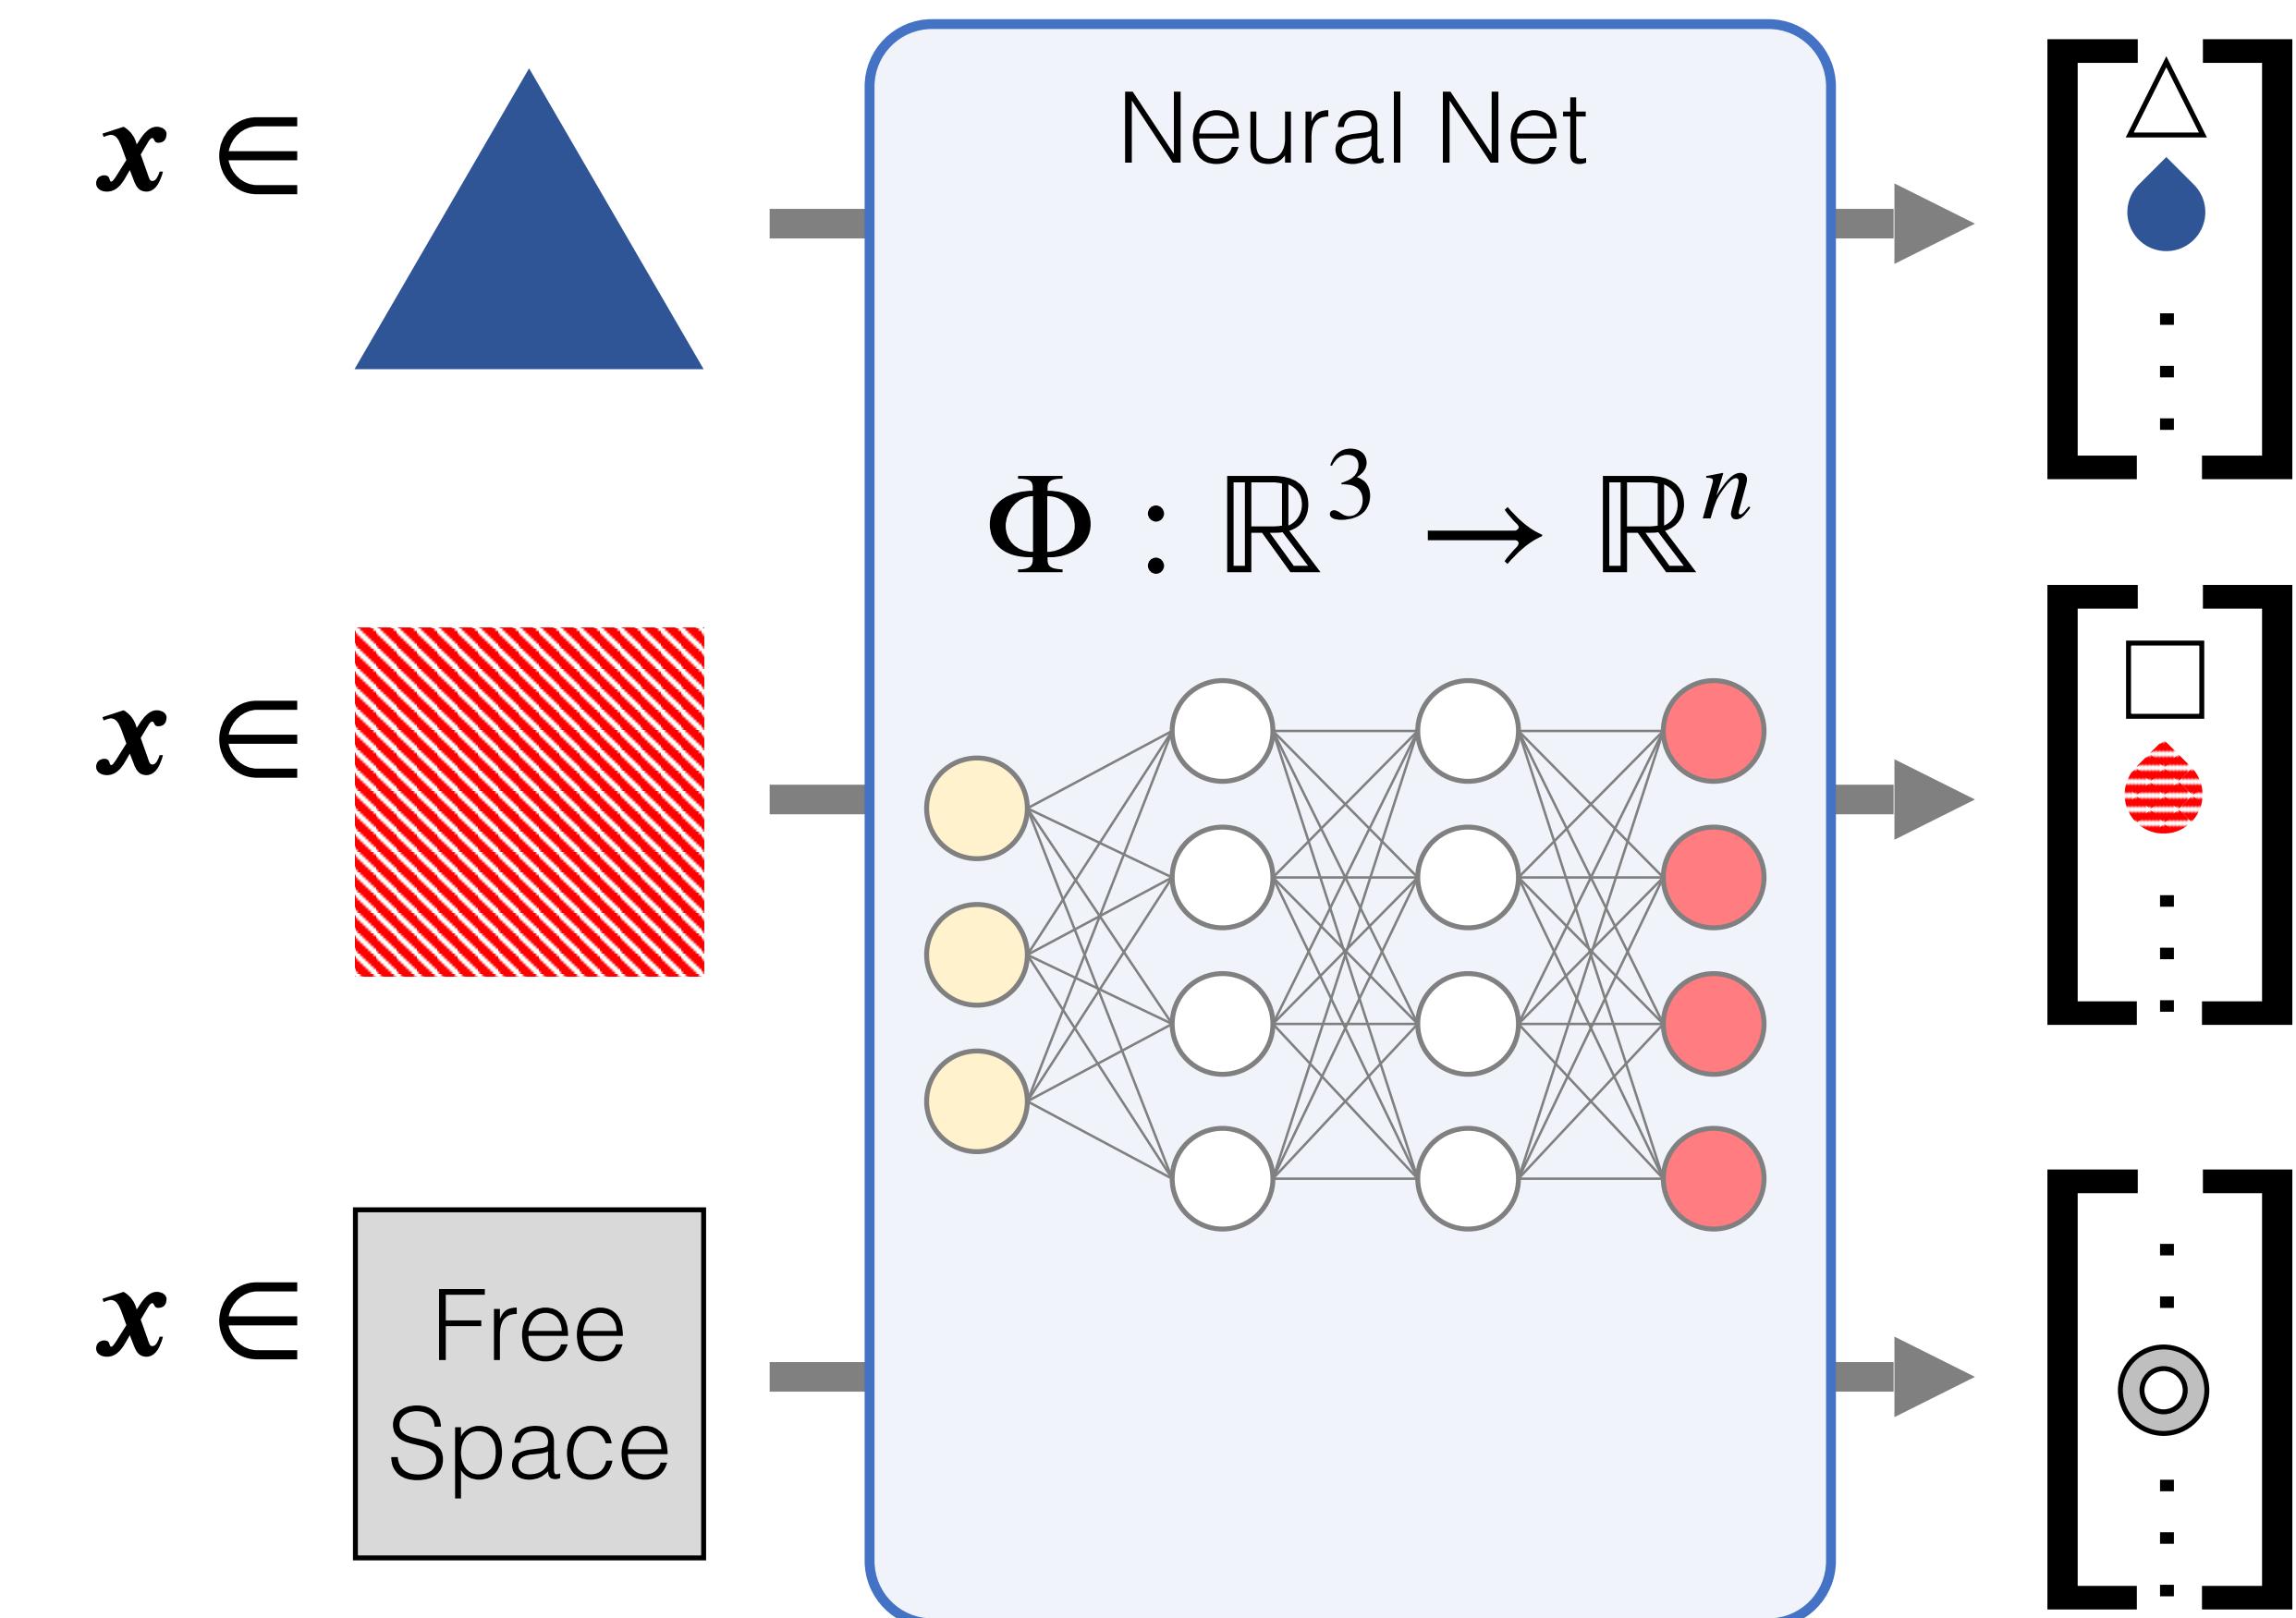
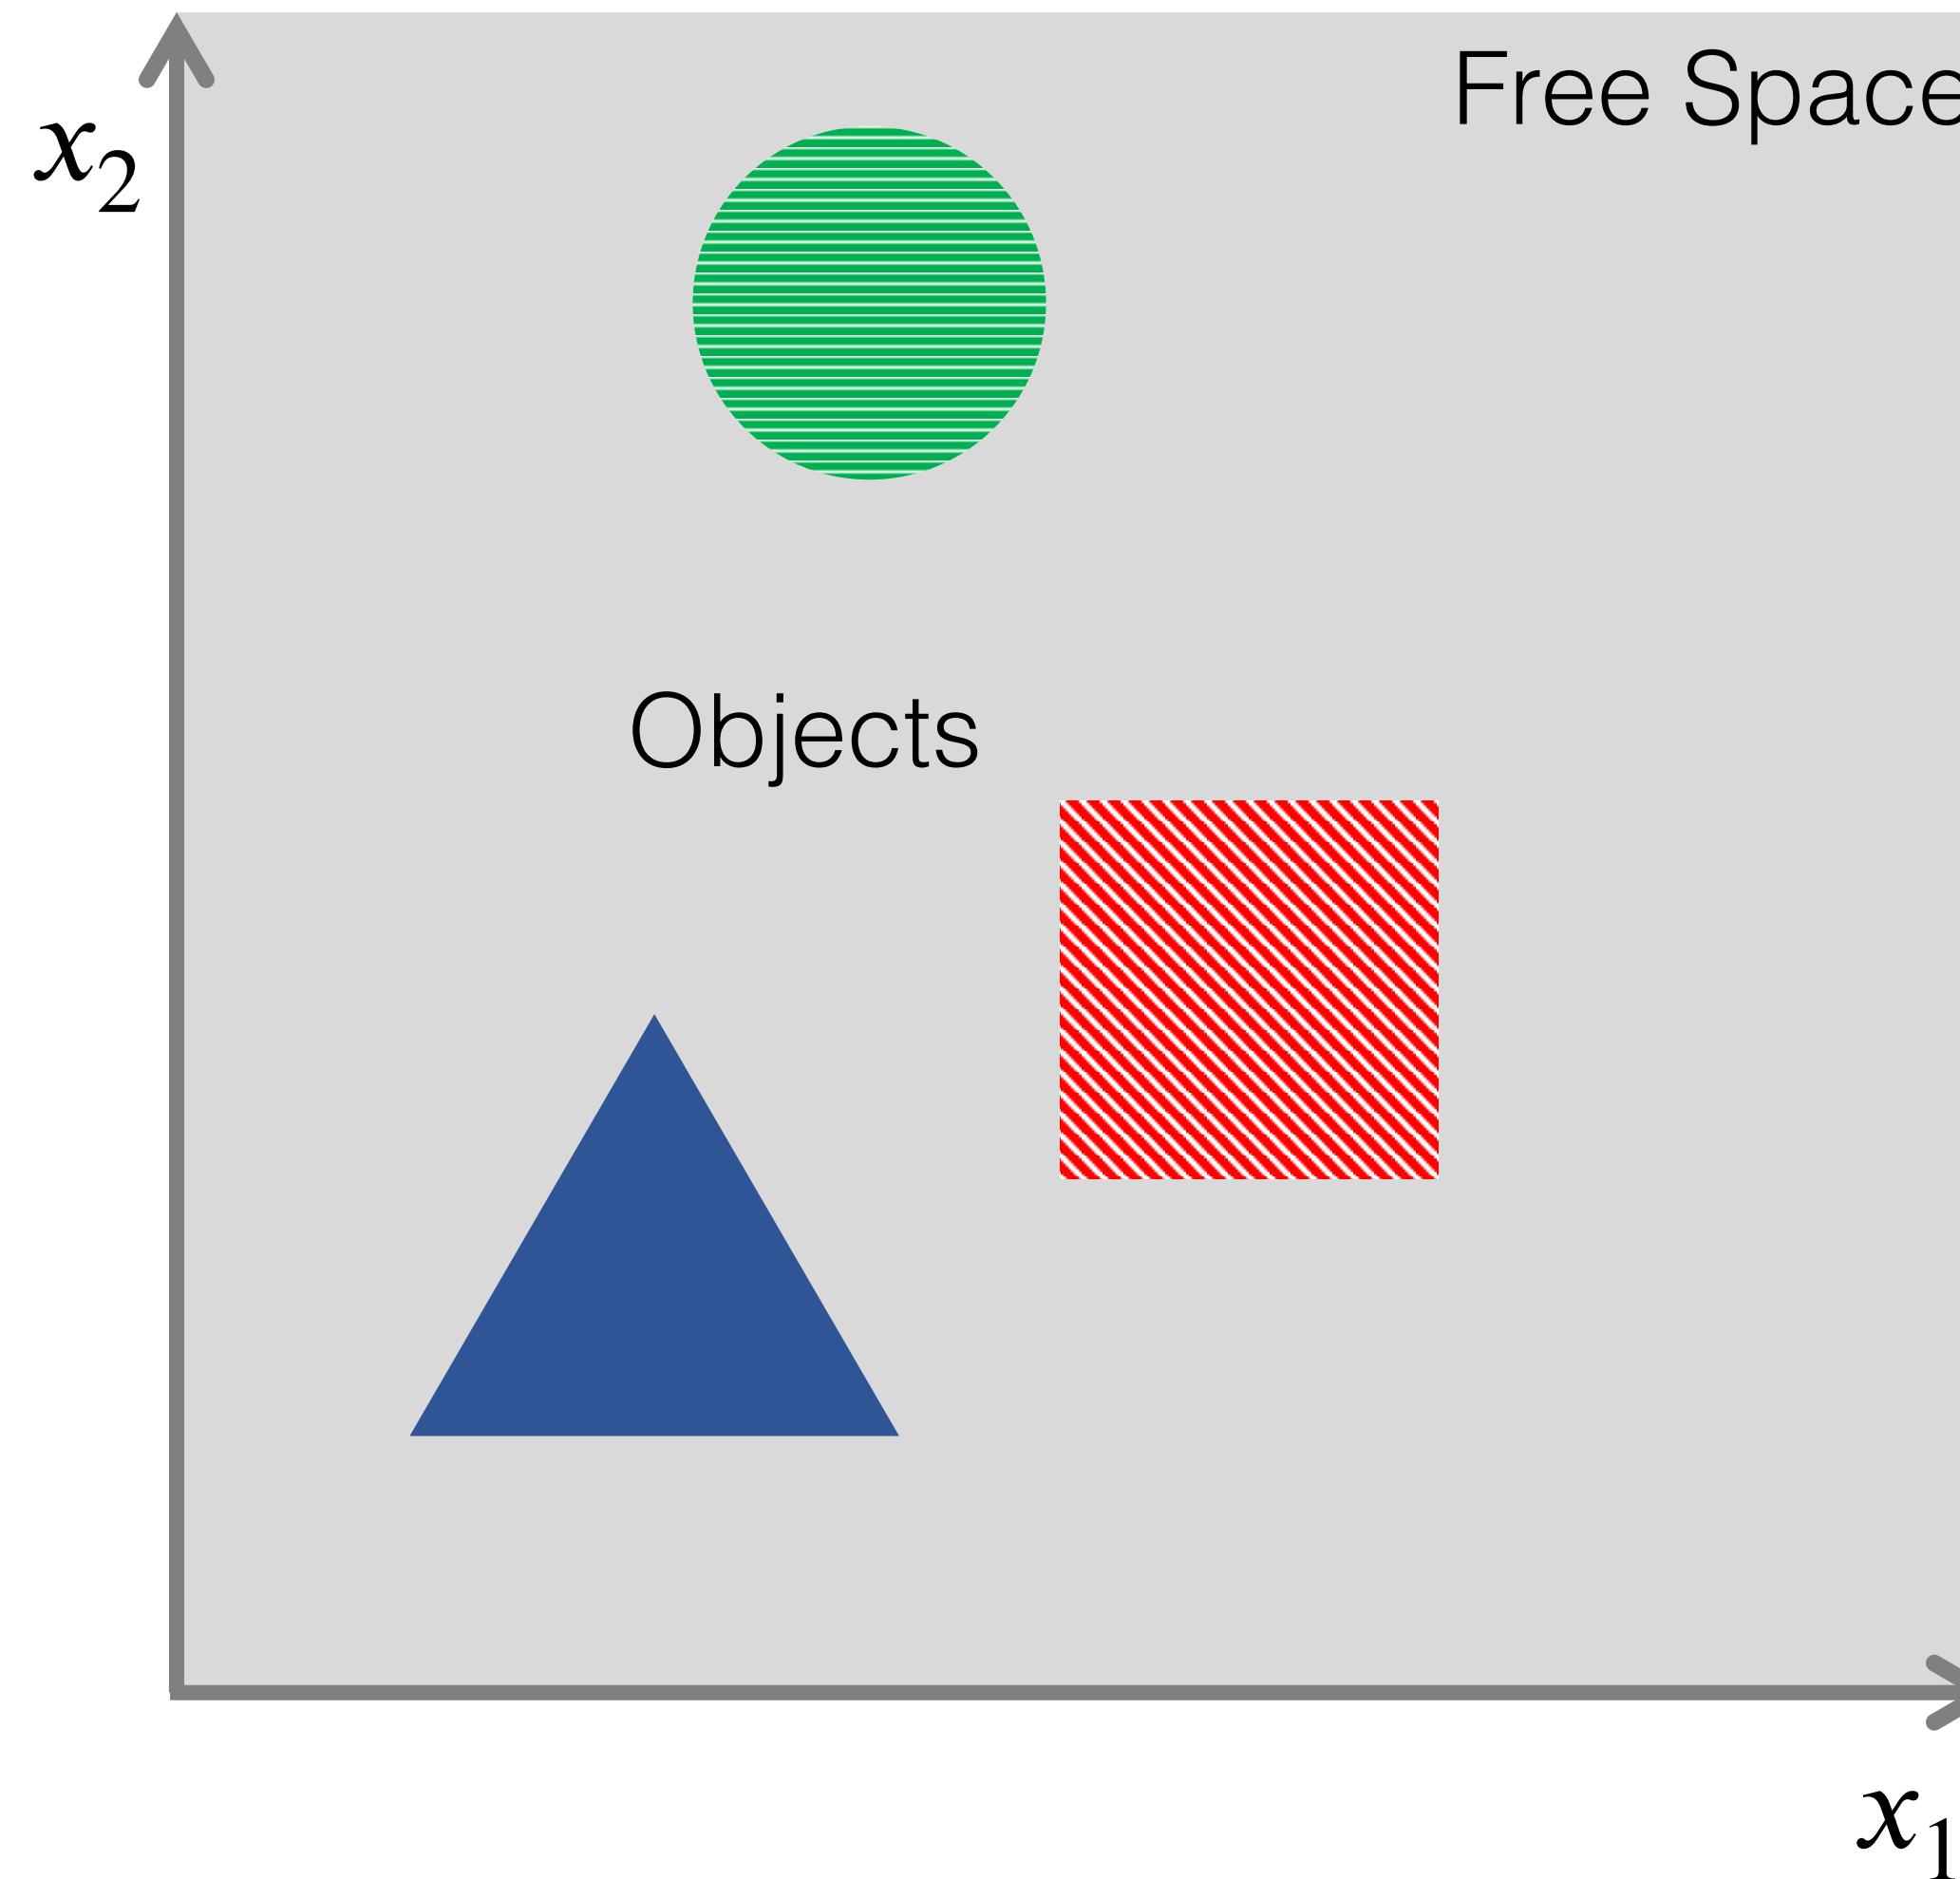


Neural Fields



Hybrid Parameterizations

# Neural fields: Parameterize Field as Neural Network!



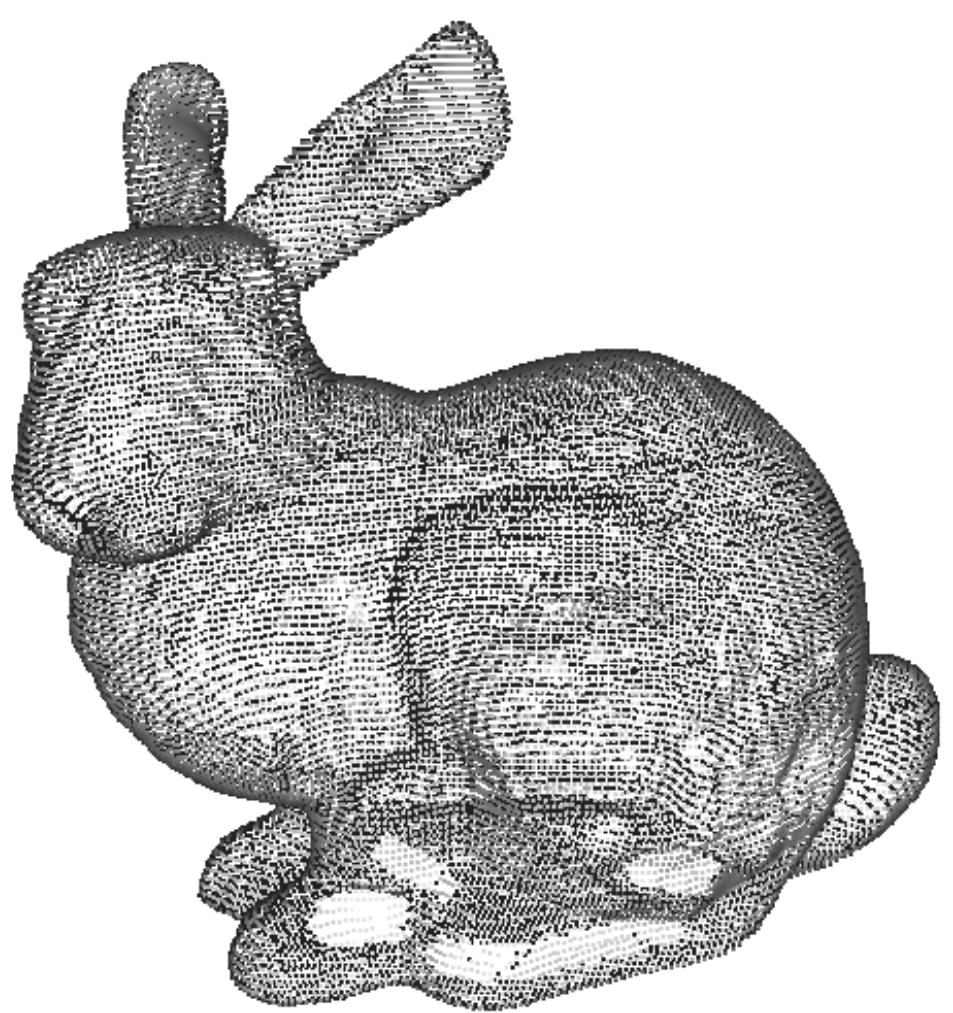
Occupancy Networks: Learning 3D Reconstruction in Function Space, Mescheder et al.

IM-Net: Learning Implicit Fields for Generative Shape Modeling, Chen et al.

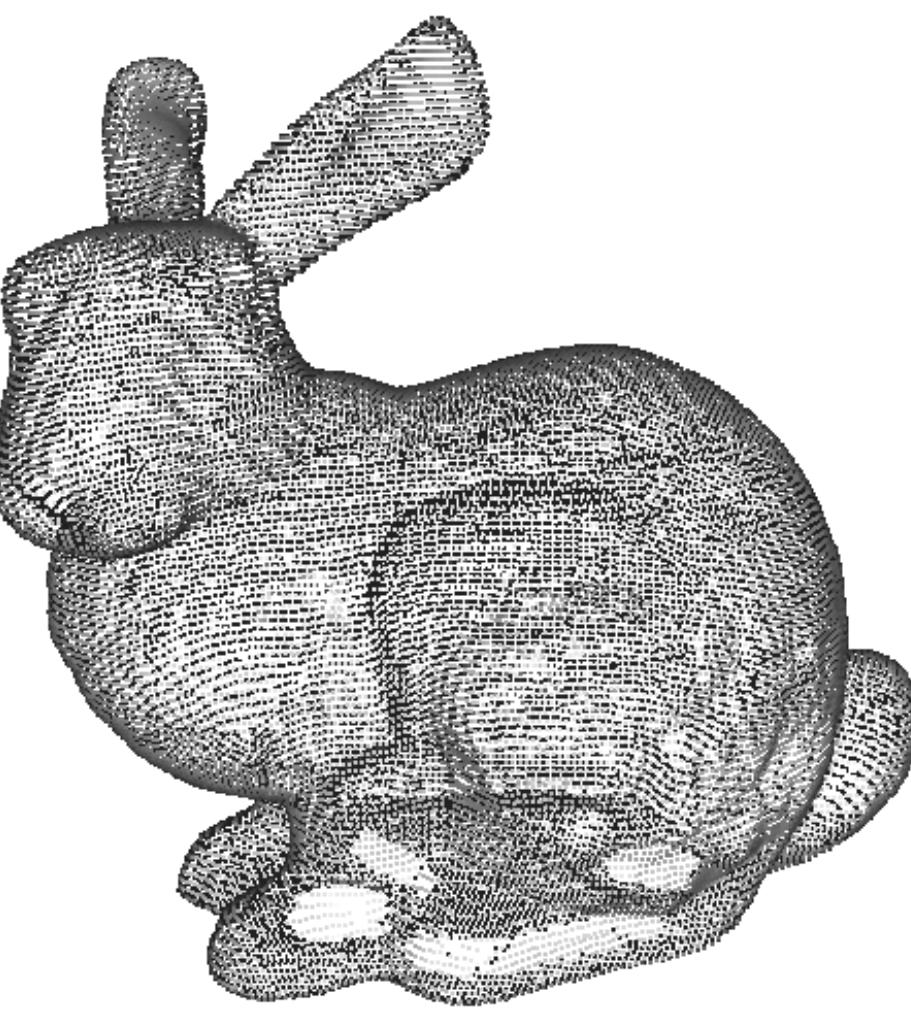
DeepSDF: Learning Continuous Signed Distance Functions for Shape Representation, Park et al. 2019

Sitzmann et al: Scene Representation Networks: Continuous 3D-Structure-Aware Neural Scene Representations, NeurIPS 2019.

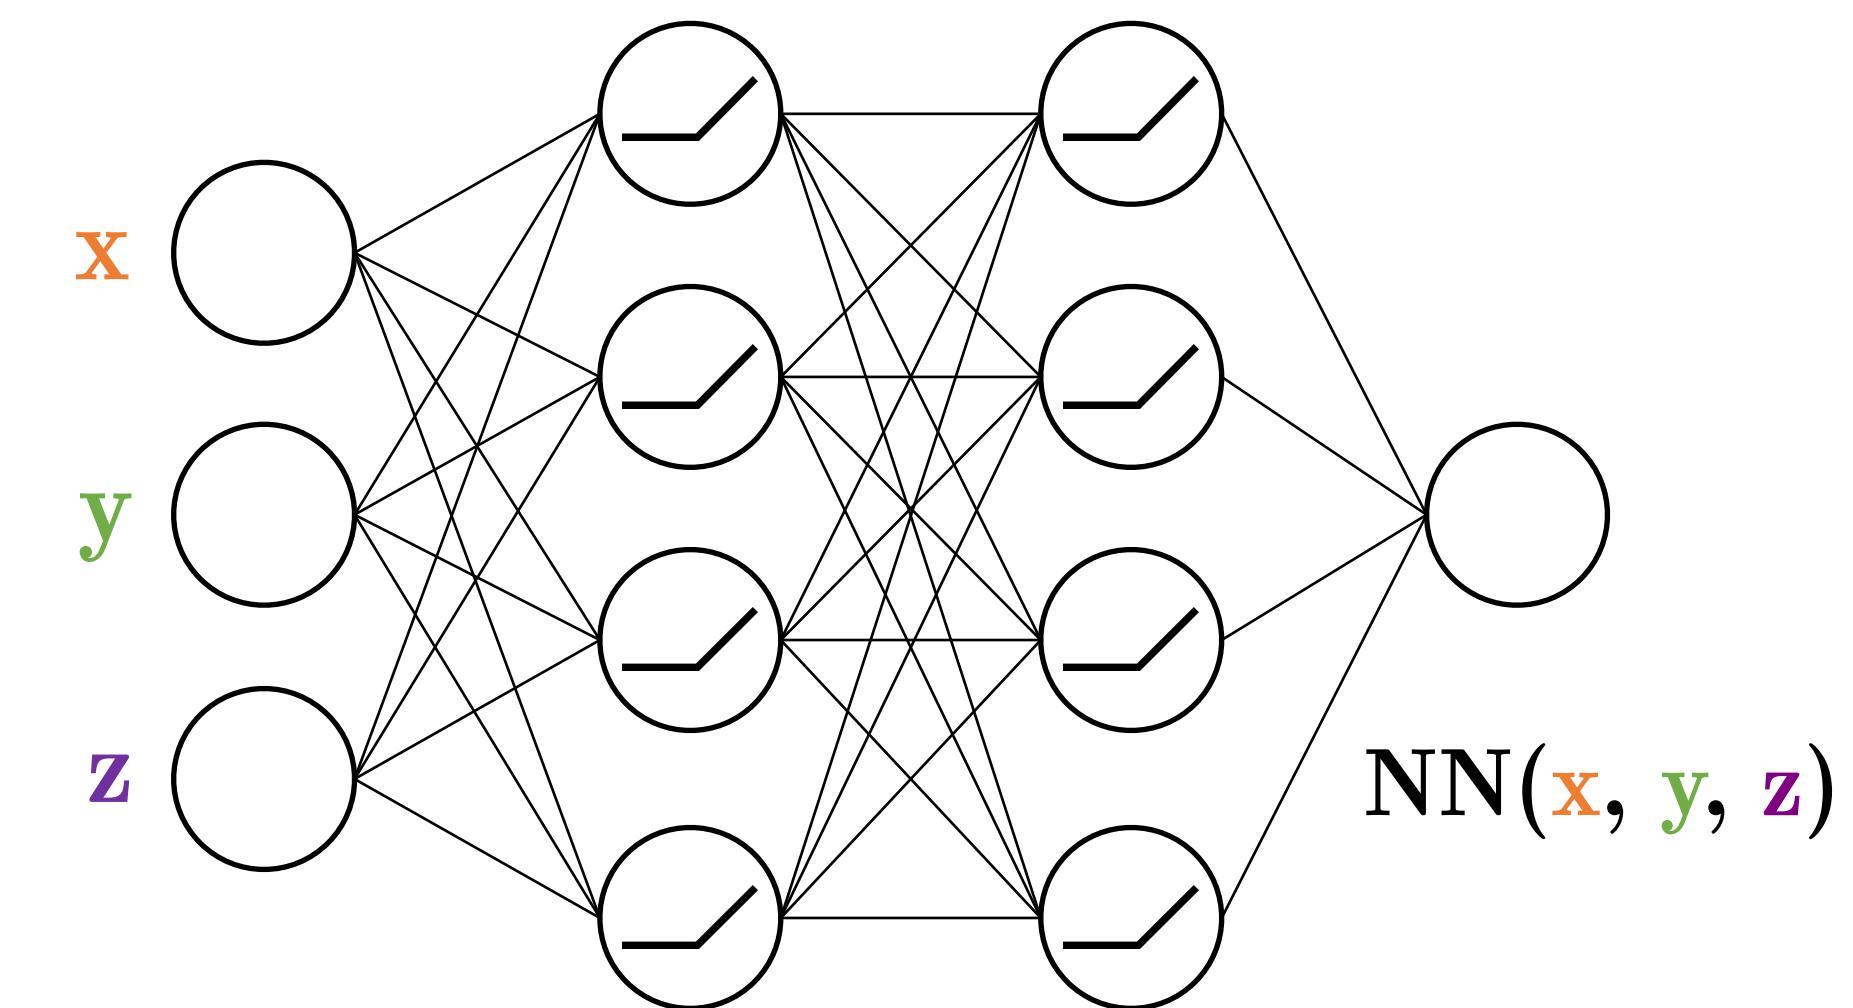
# Shapes



Shapes



ReLU MLP

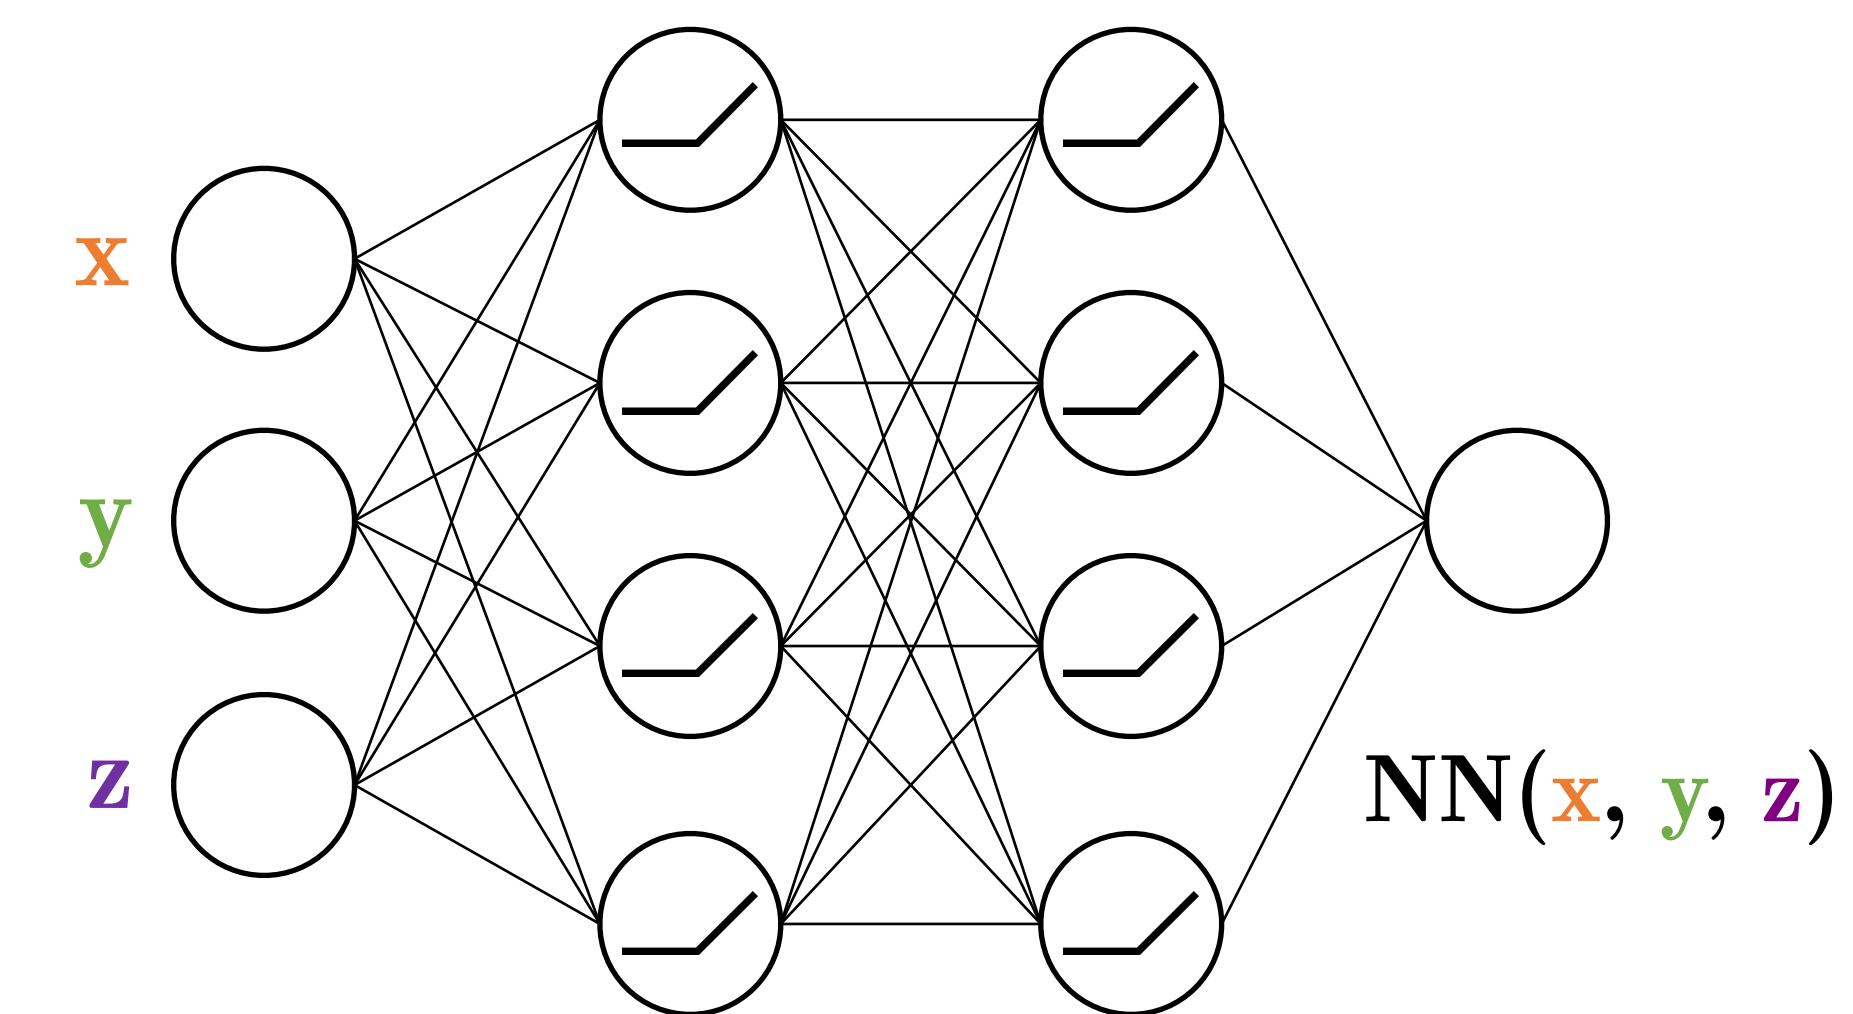


*Mescheder et al. 2018*  
*Park et al. 2018*  
*Chen et al. 2018*

# Shapes



ReLU MLP

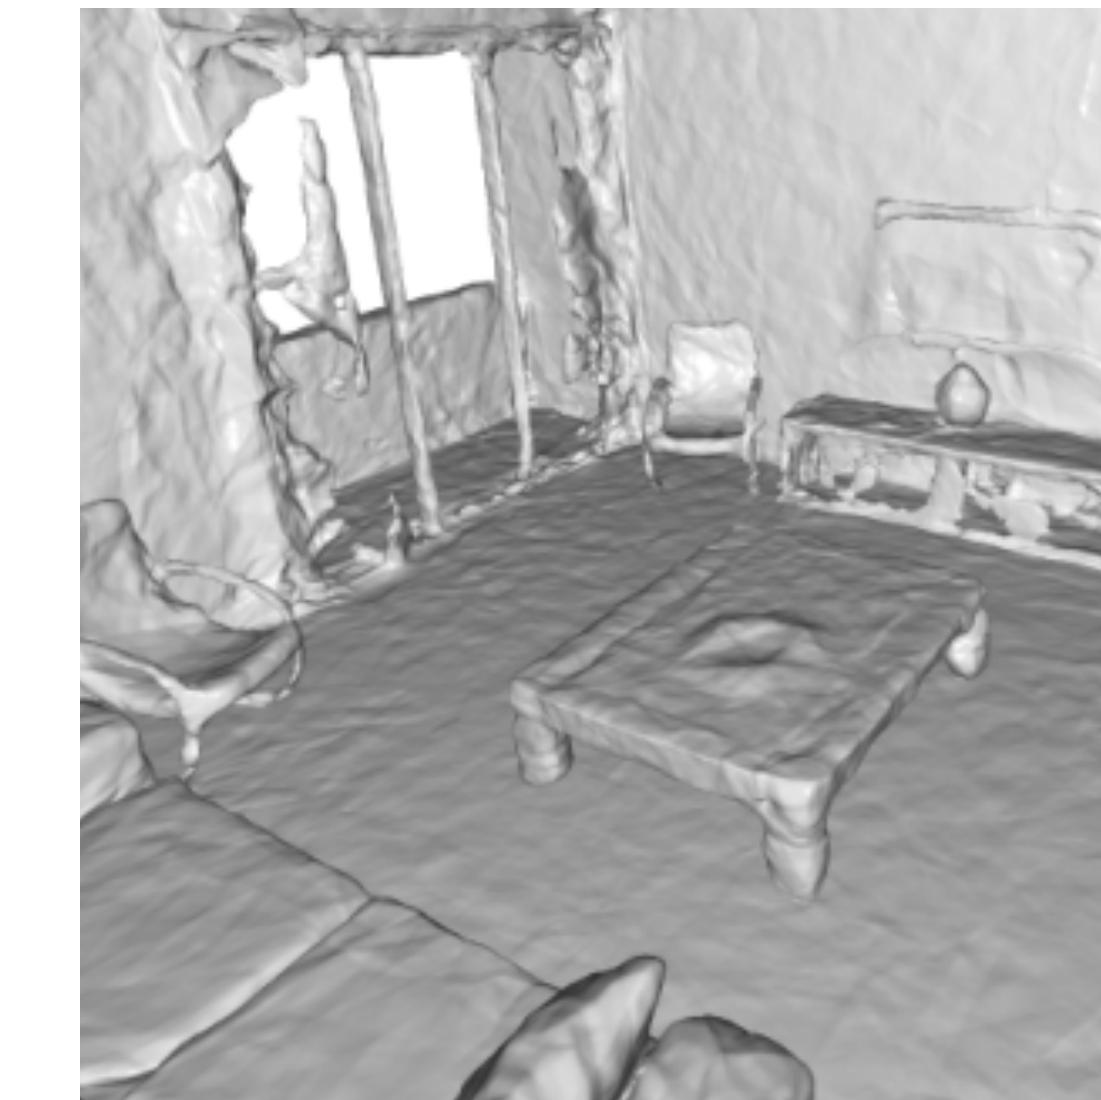
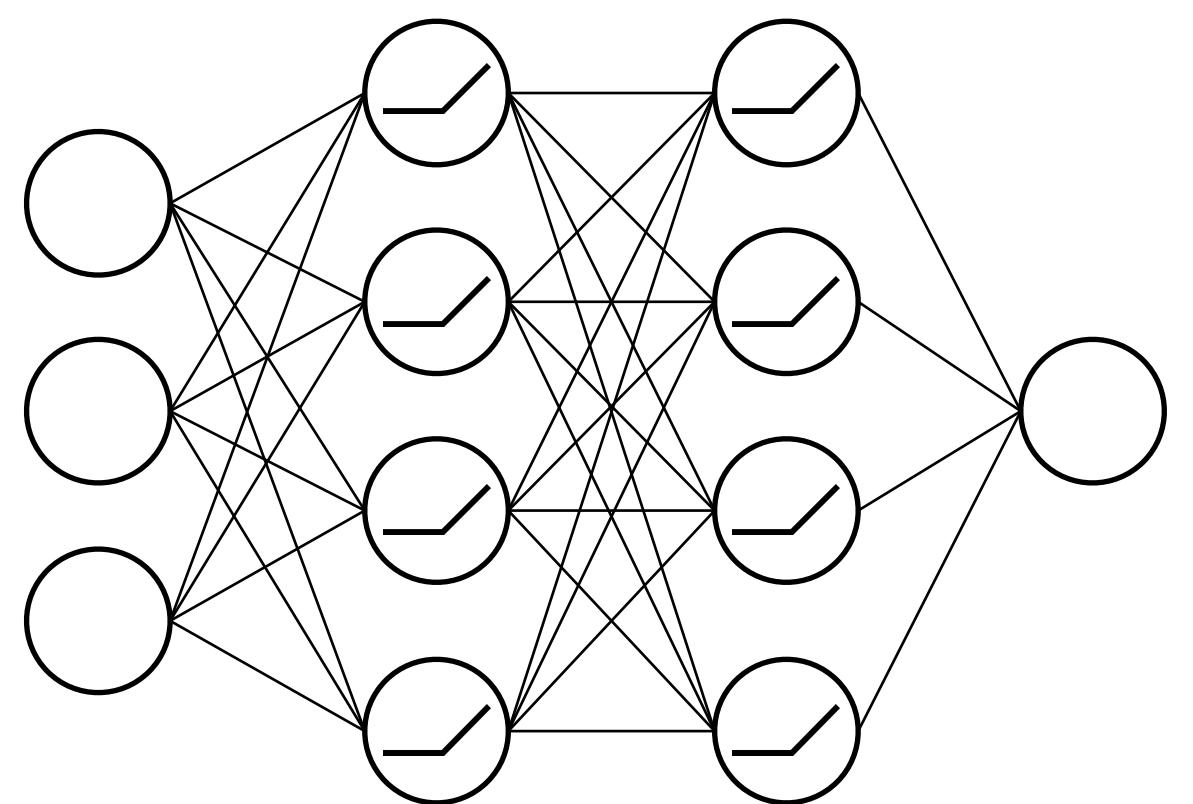


*Mescheder et al. 2018*  
*Park et al. 2018*  
*Chen et al. 2018*

# Shapes



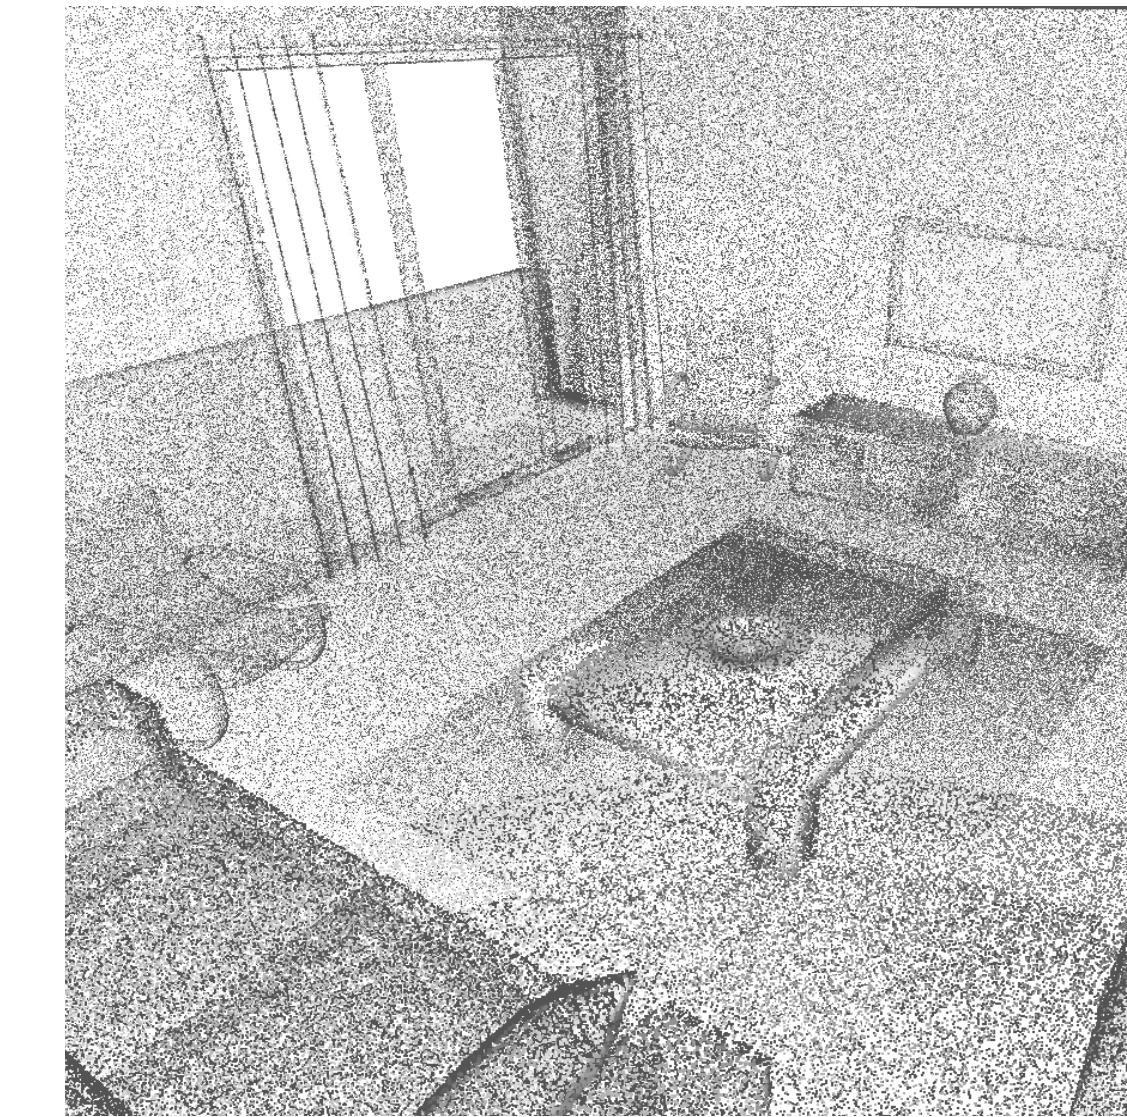
ReLU MLP



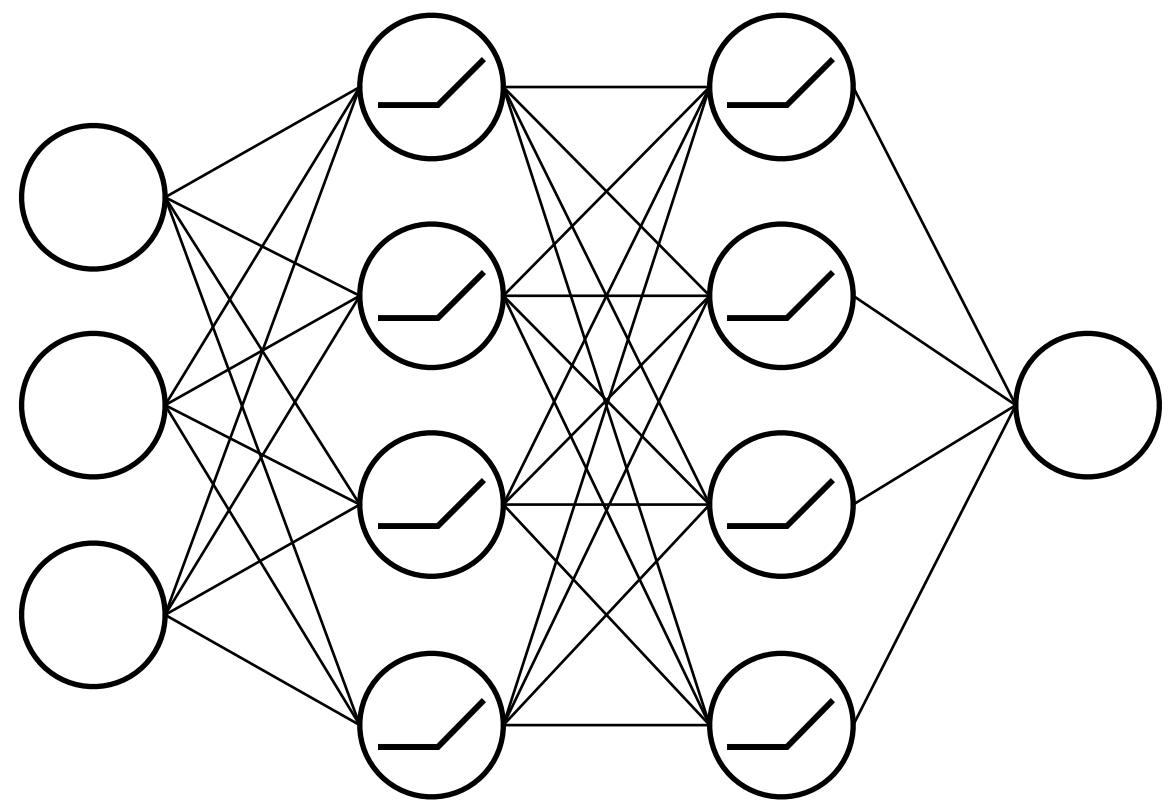
Images



Shapes



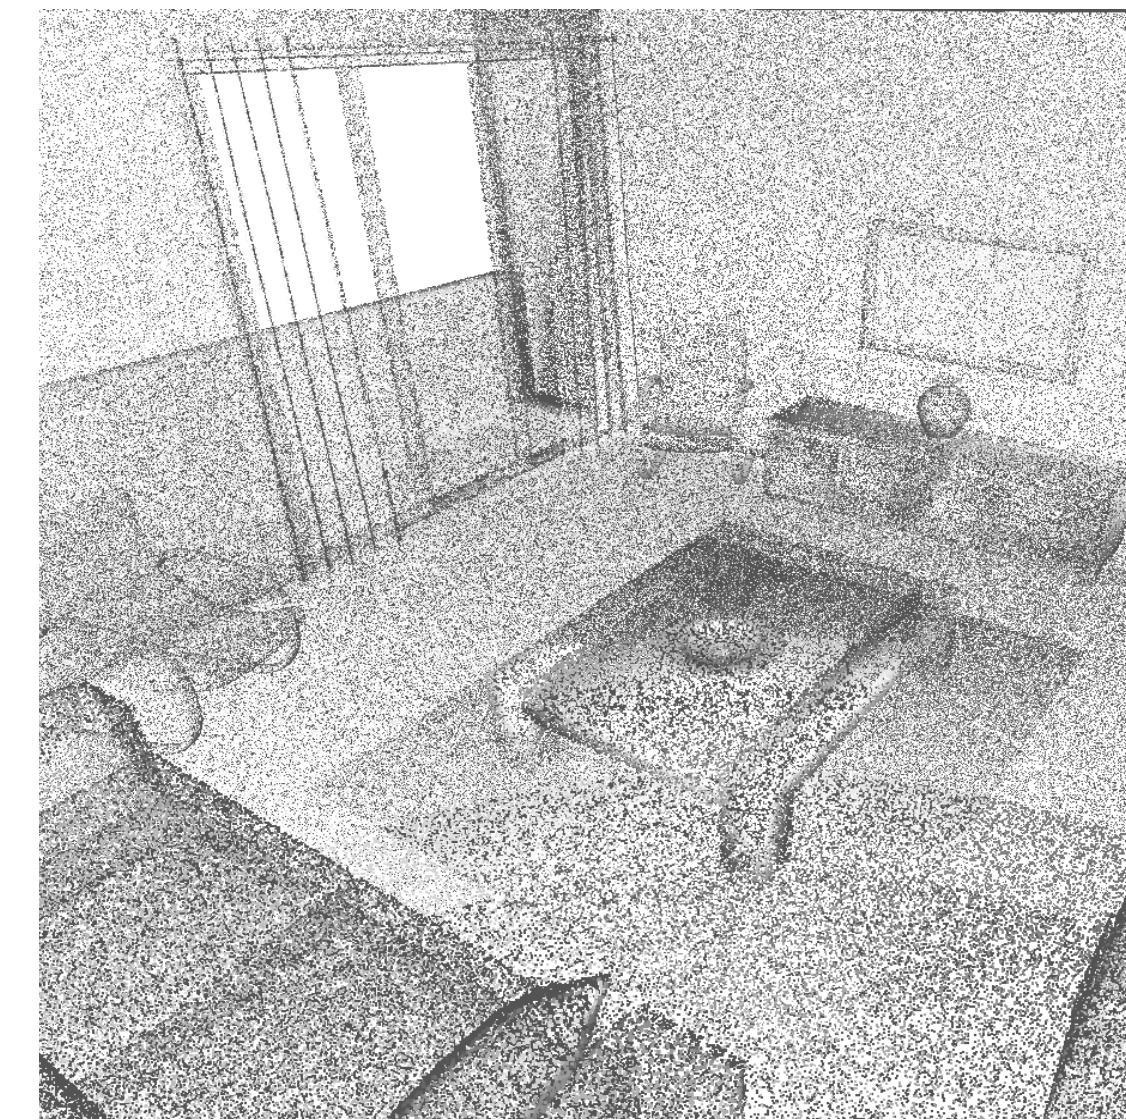
ReLU MLP



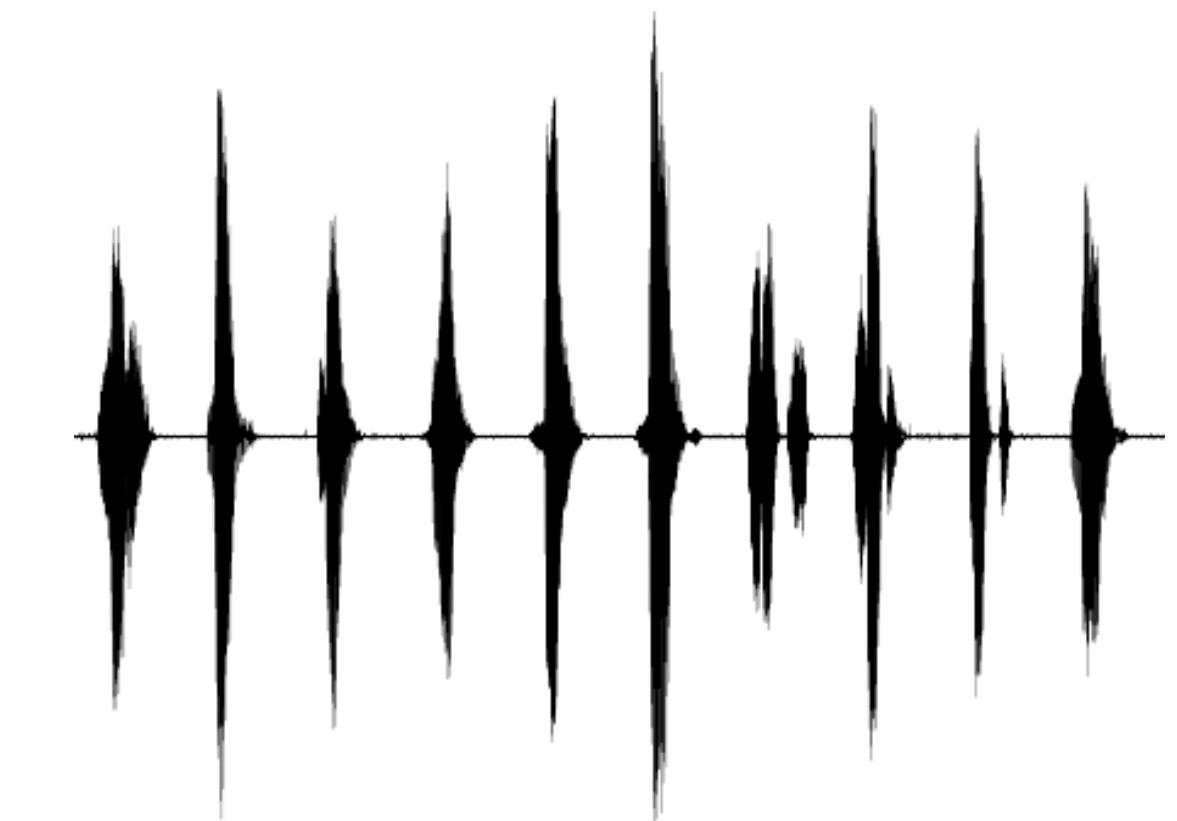
Images



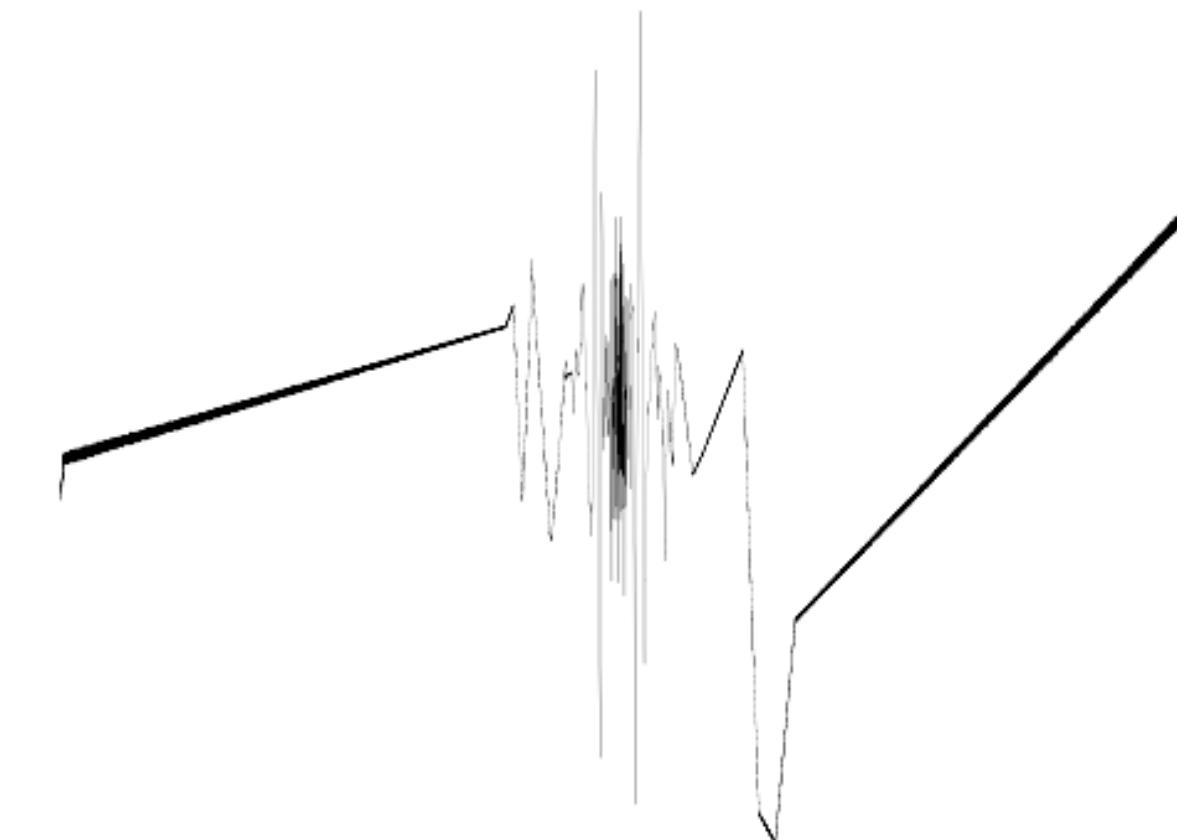
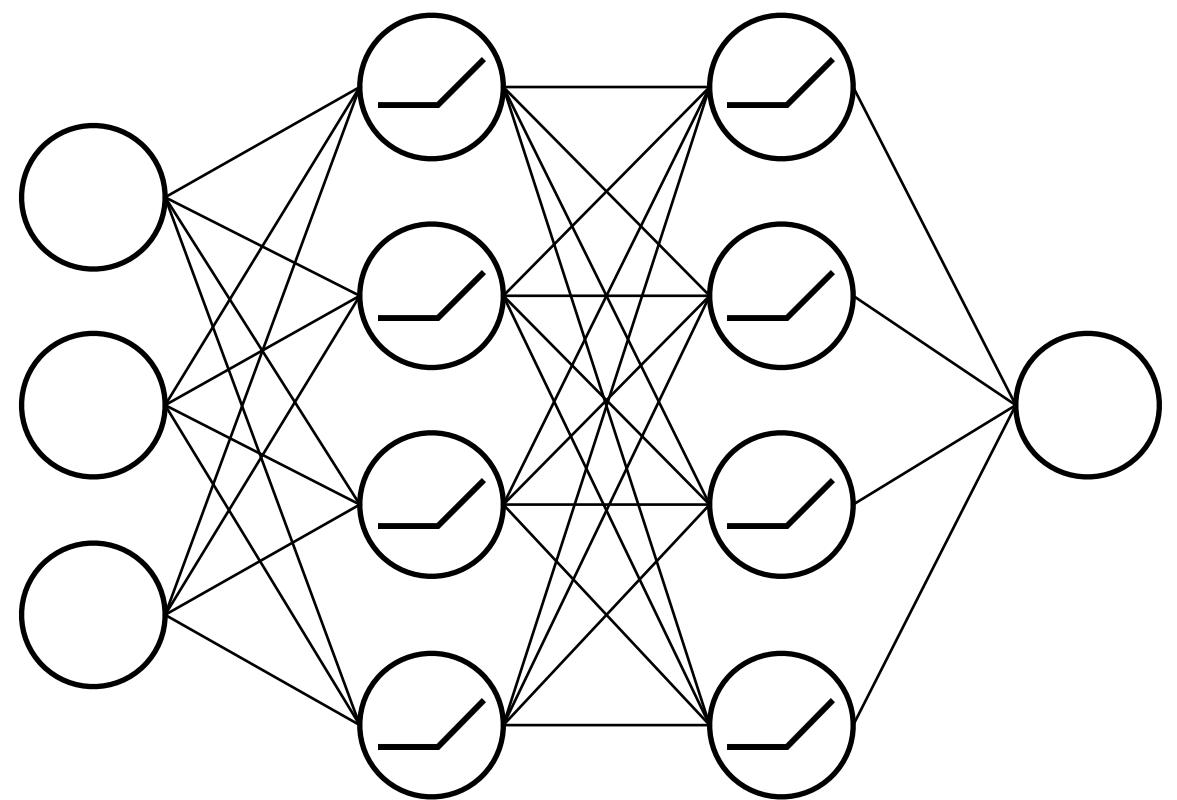
Shapes



Audio



ReLU MLP



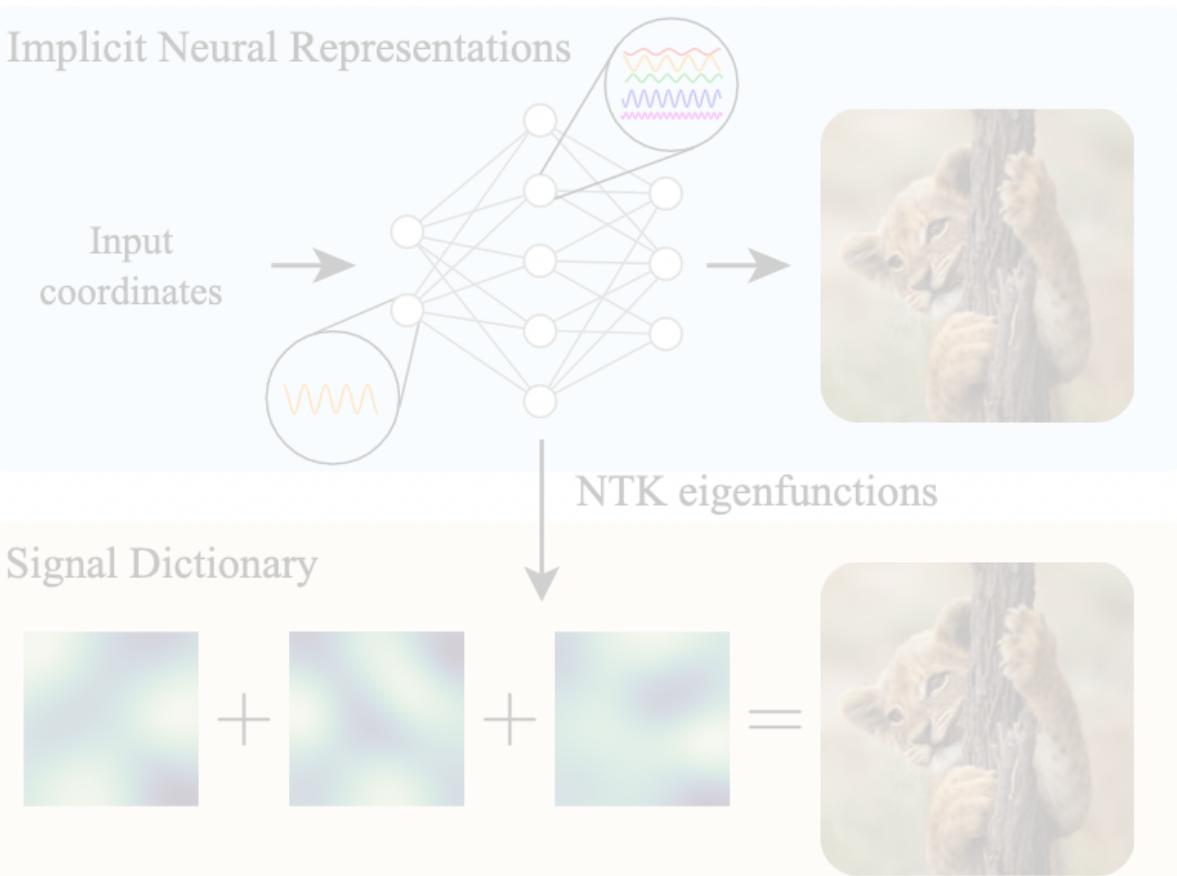
# SIREN: Neural Implicit Representations With Periodic Activation Functions

Sitzmann & Martel et al. NeurIPS 2020



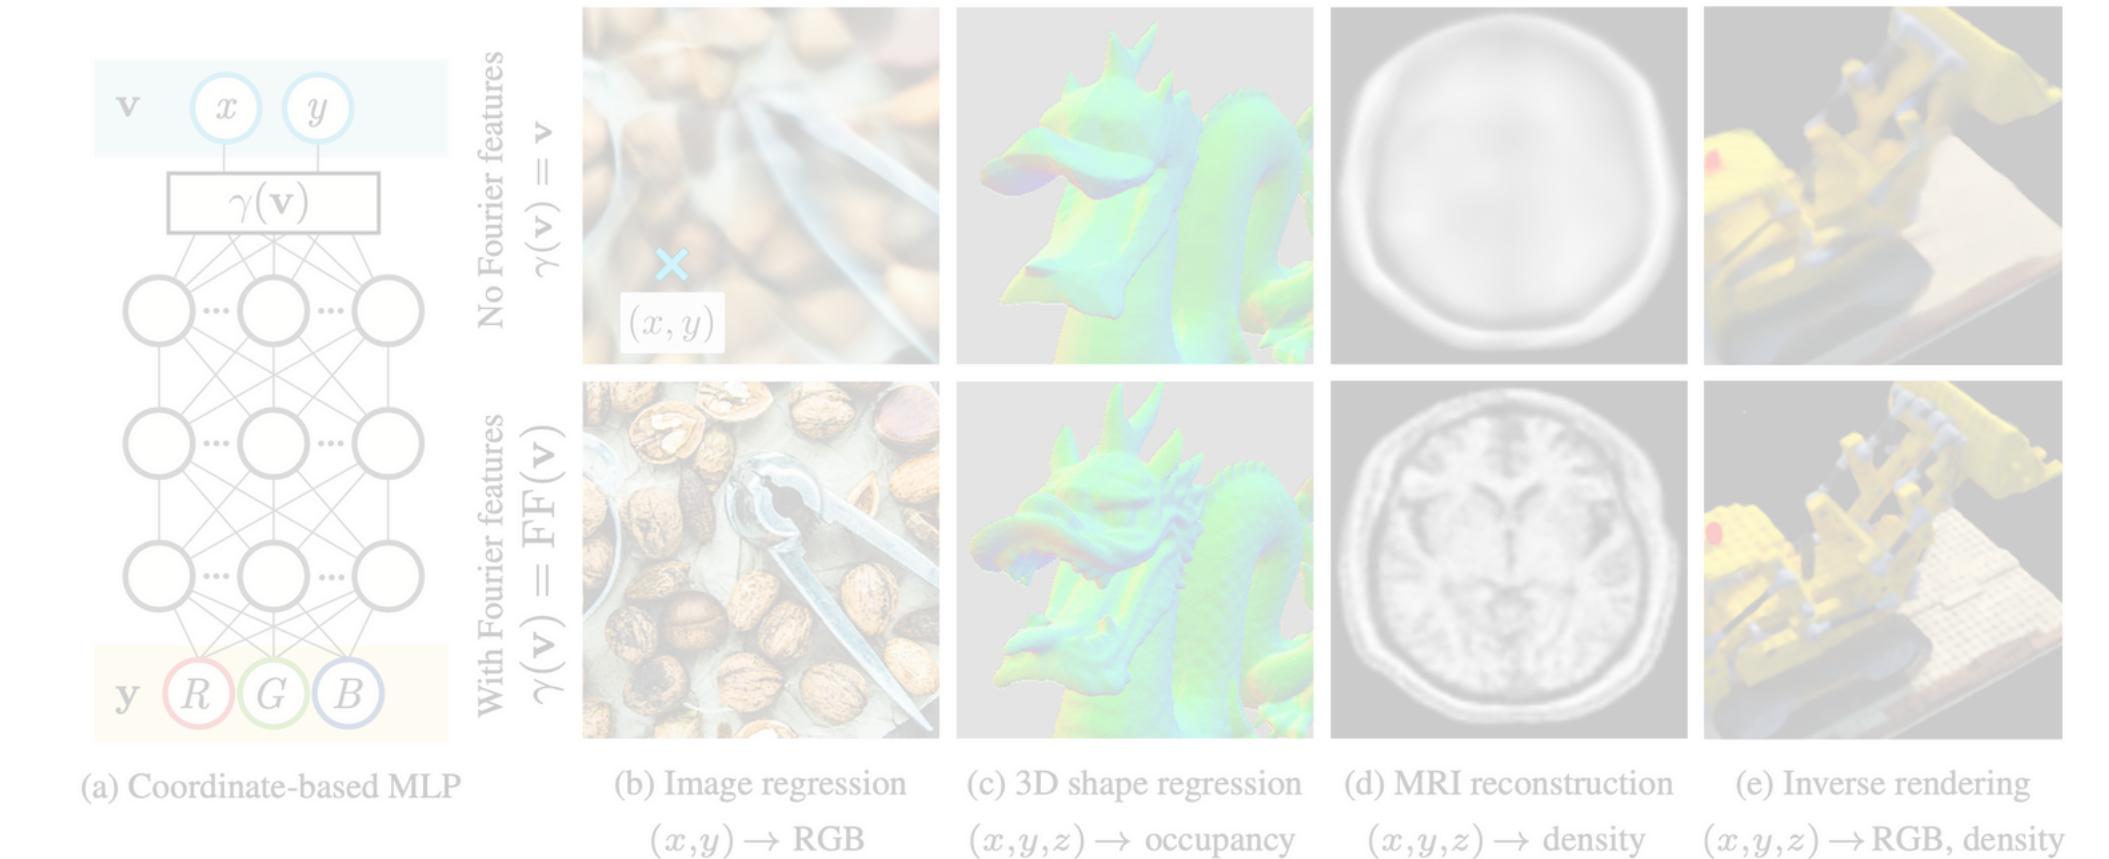
## A Structured Dictionary Perspective on Implicit Neural Representations

Yüce & Ortiz-Jiménez et al. CVPR 2022



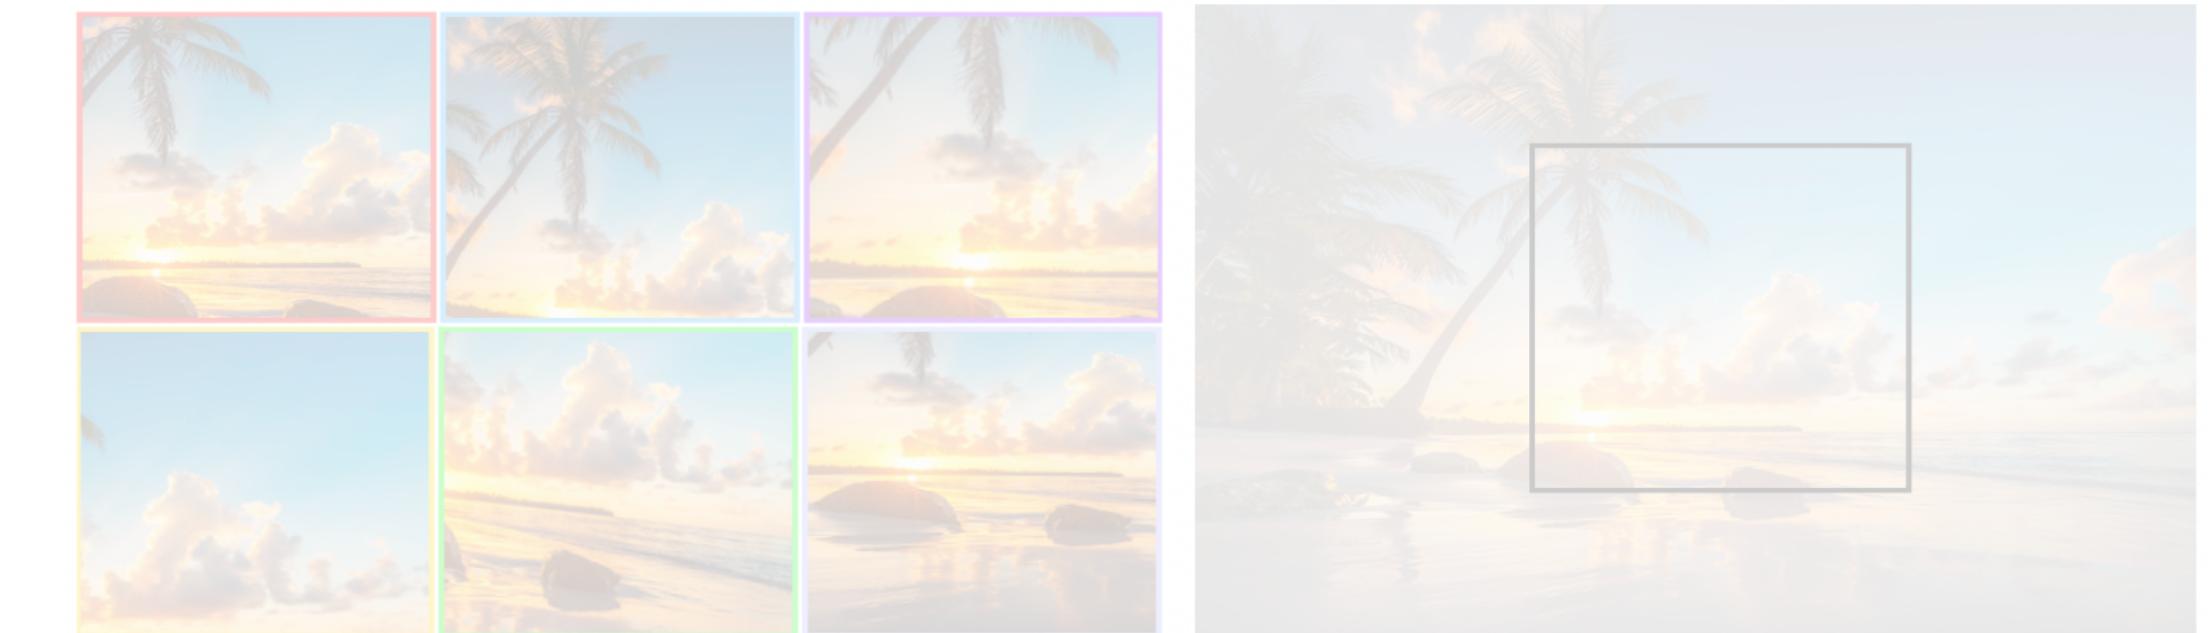
# Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains

Tancik, Srinivasan, Mildenhall NeurIPS 2020



## Gaussian Activated Neural Radiance Fields for High Fidelity Reconstruction & Pose Estimation

Chng et al. ECCV 2022



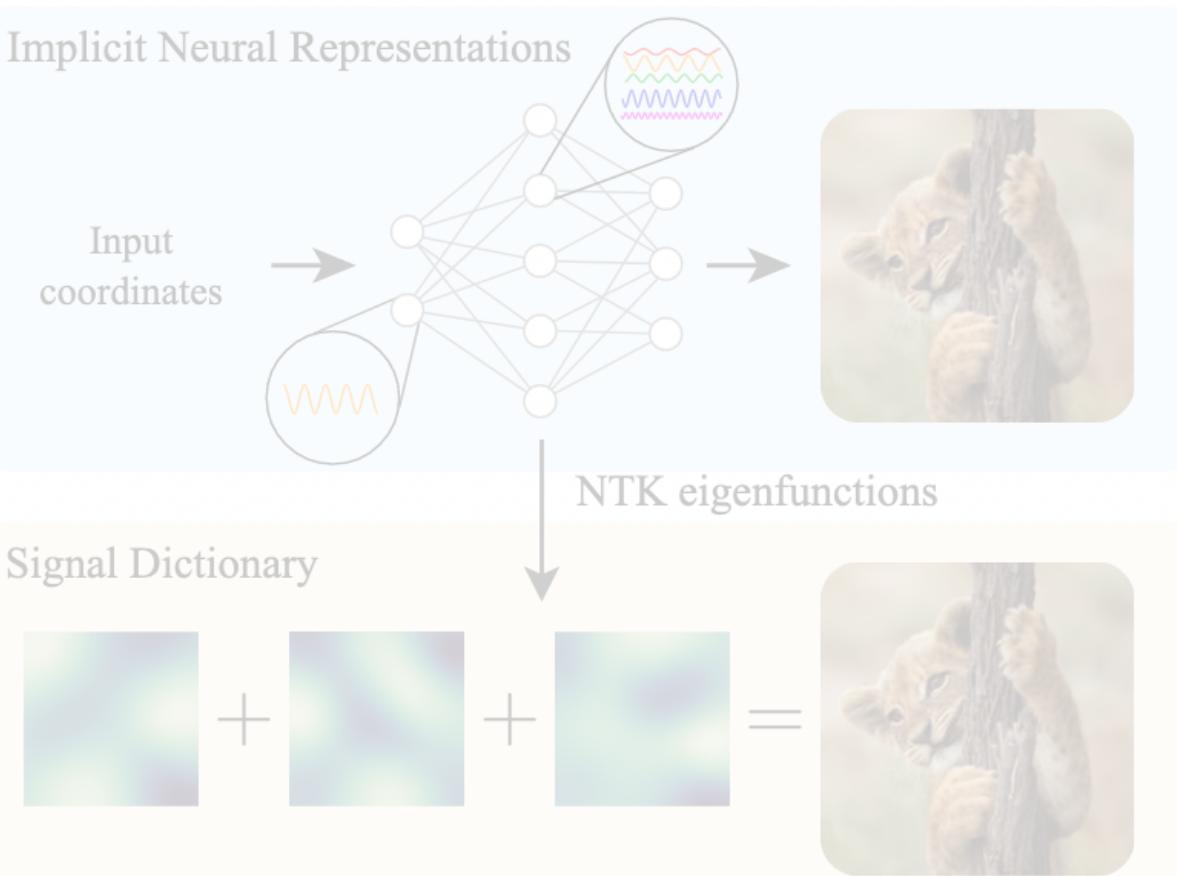
# SIREN: Neural Implicit Representations With Periodic Activation Functions

Sitzmann & Martel et al. NeurIPS 2020



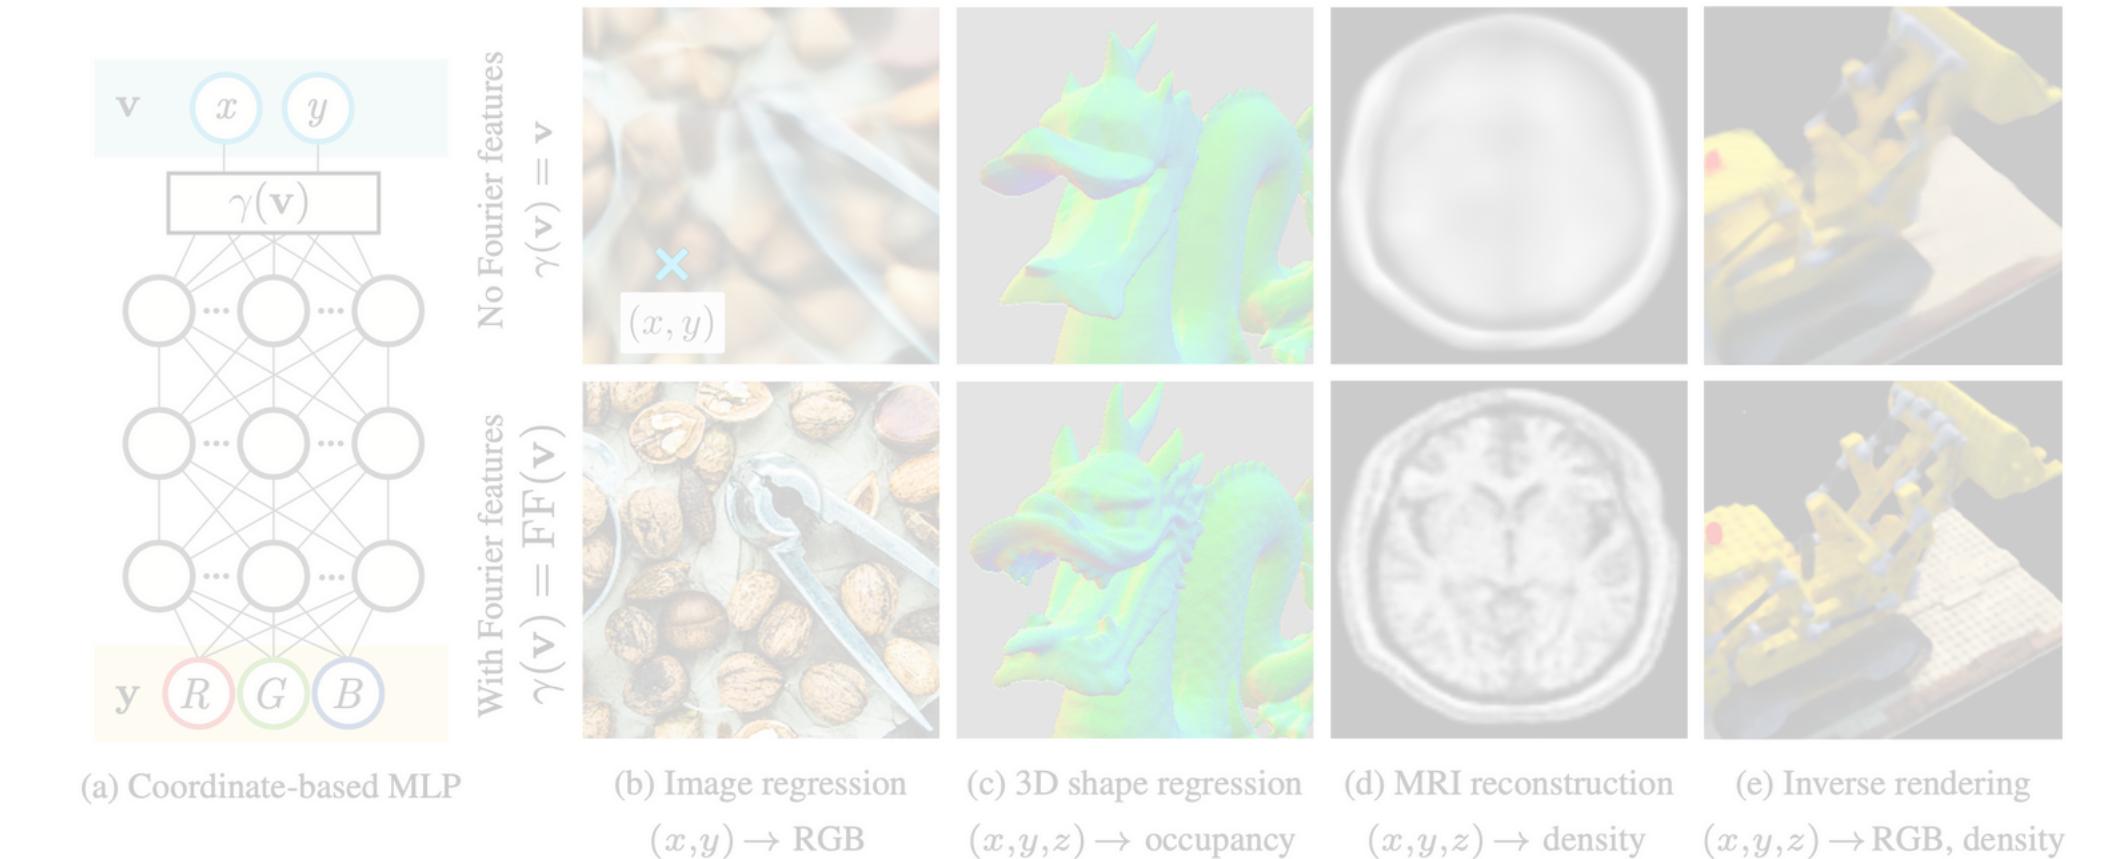
## A Structured Dictionary Perspective on Implicit Neural Representations

Yüce & Ortiz-Jiménez et al. CVPR 2022



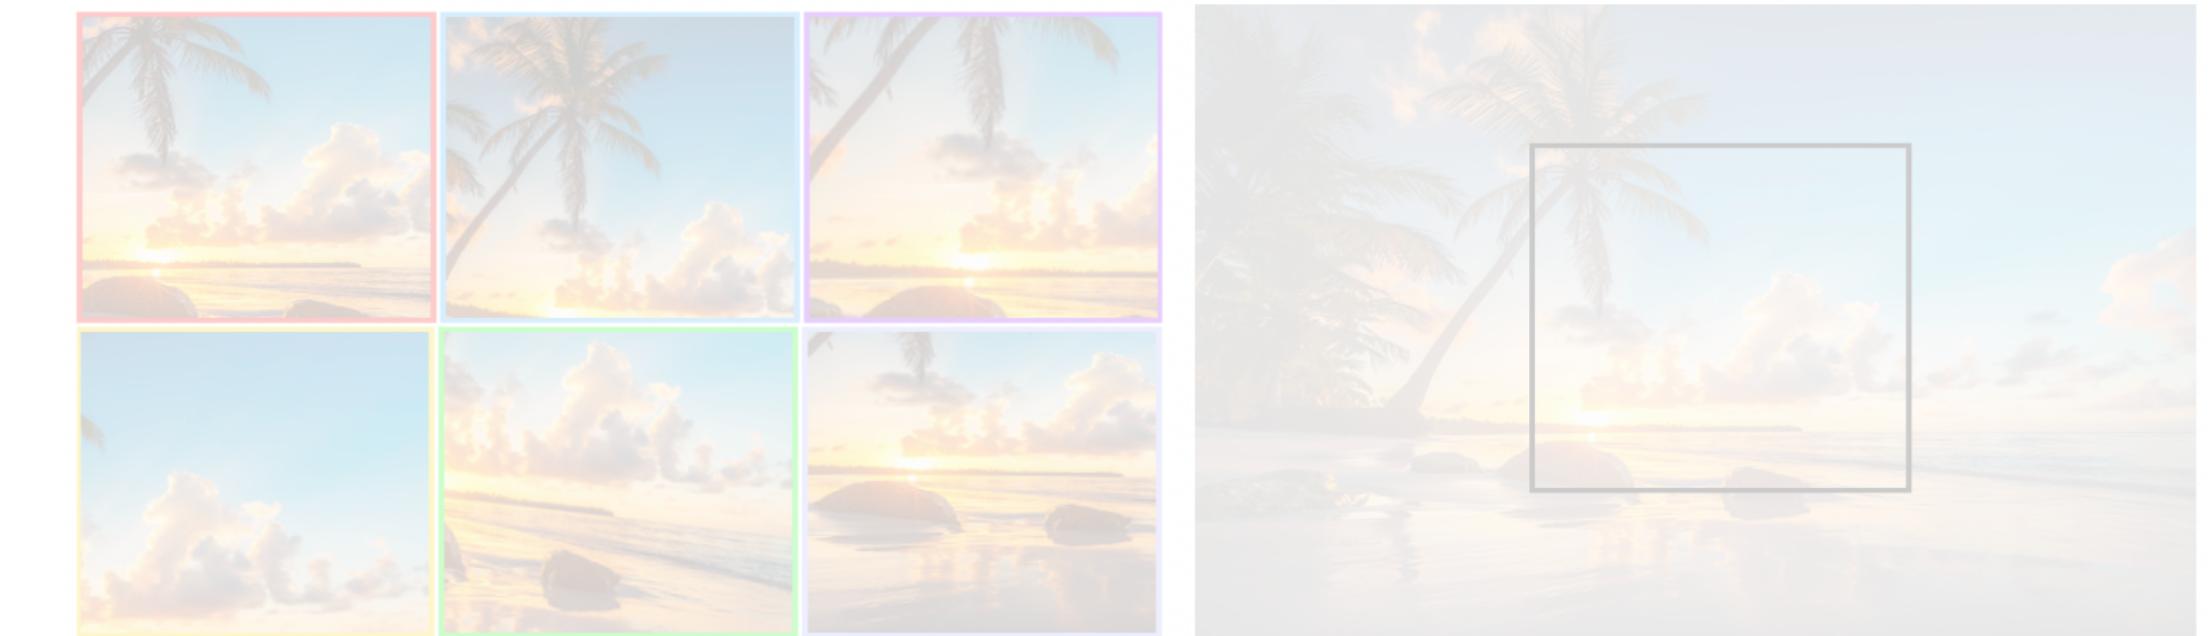
# Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains

Tancik, Srinivasan, Mildenhall NeurIPS 2020



## Gaussian Activated Neural Radiance Fields for High Fidelity Reconstruction & Pose Estimation

Chng et al. ECCV 2022

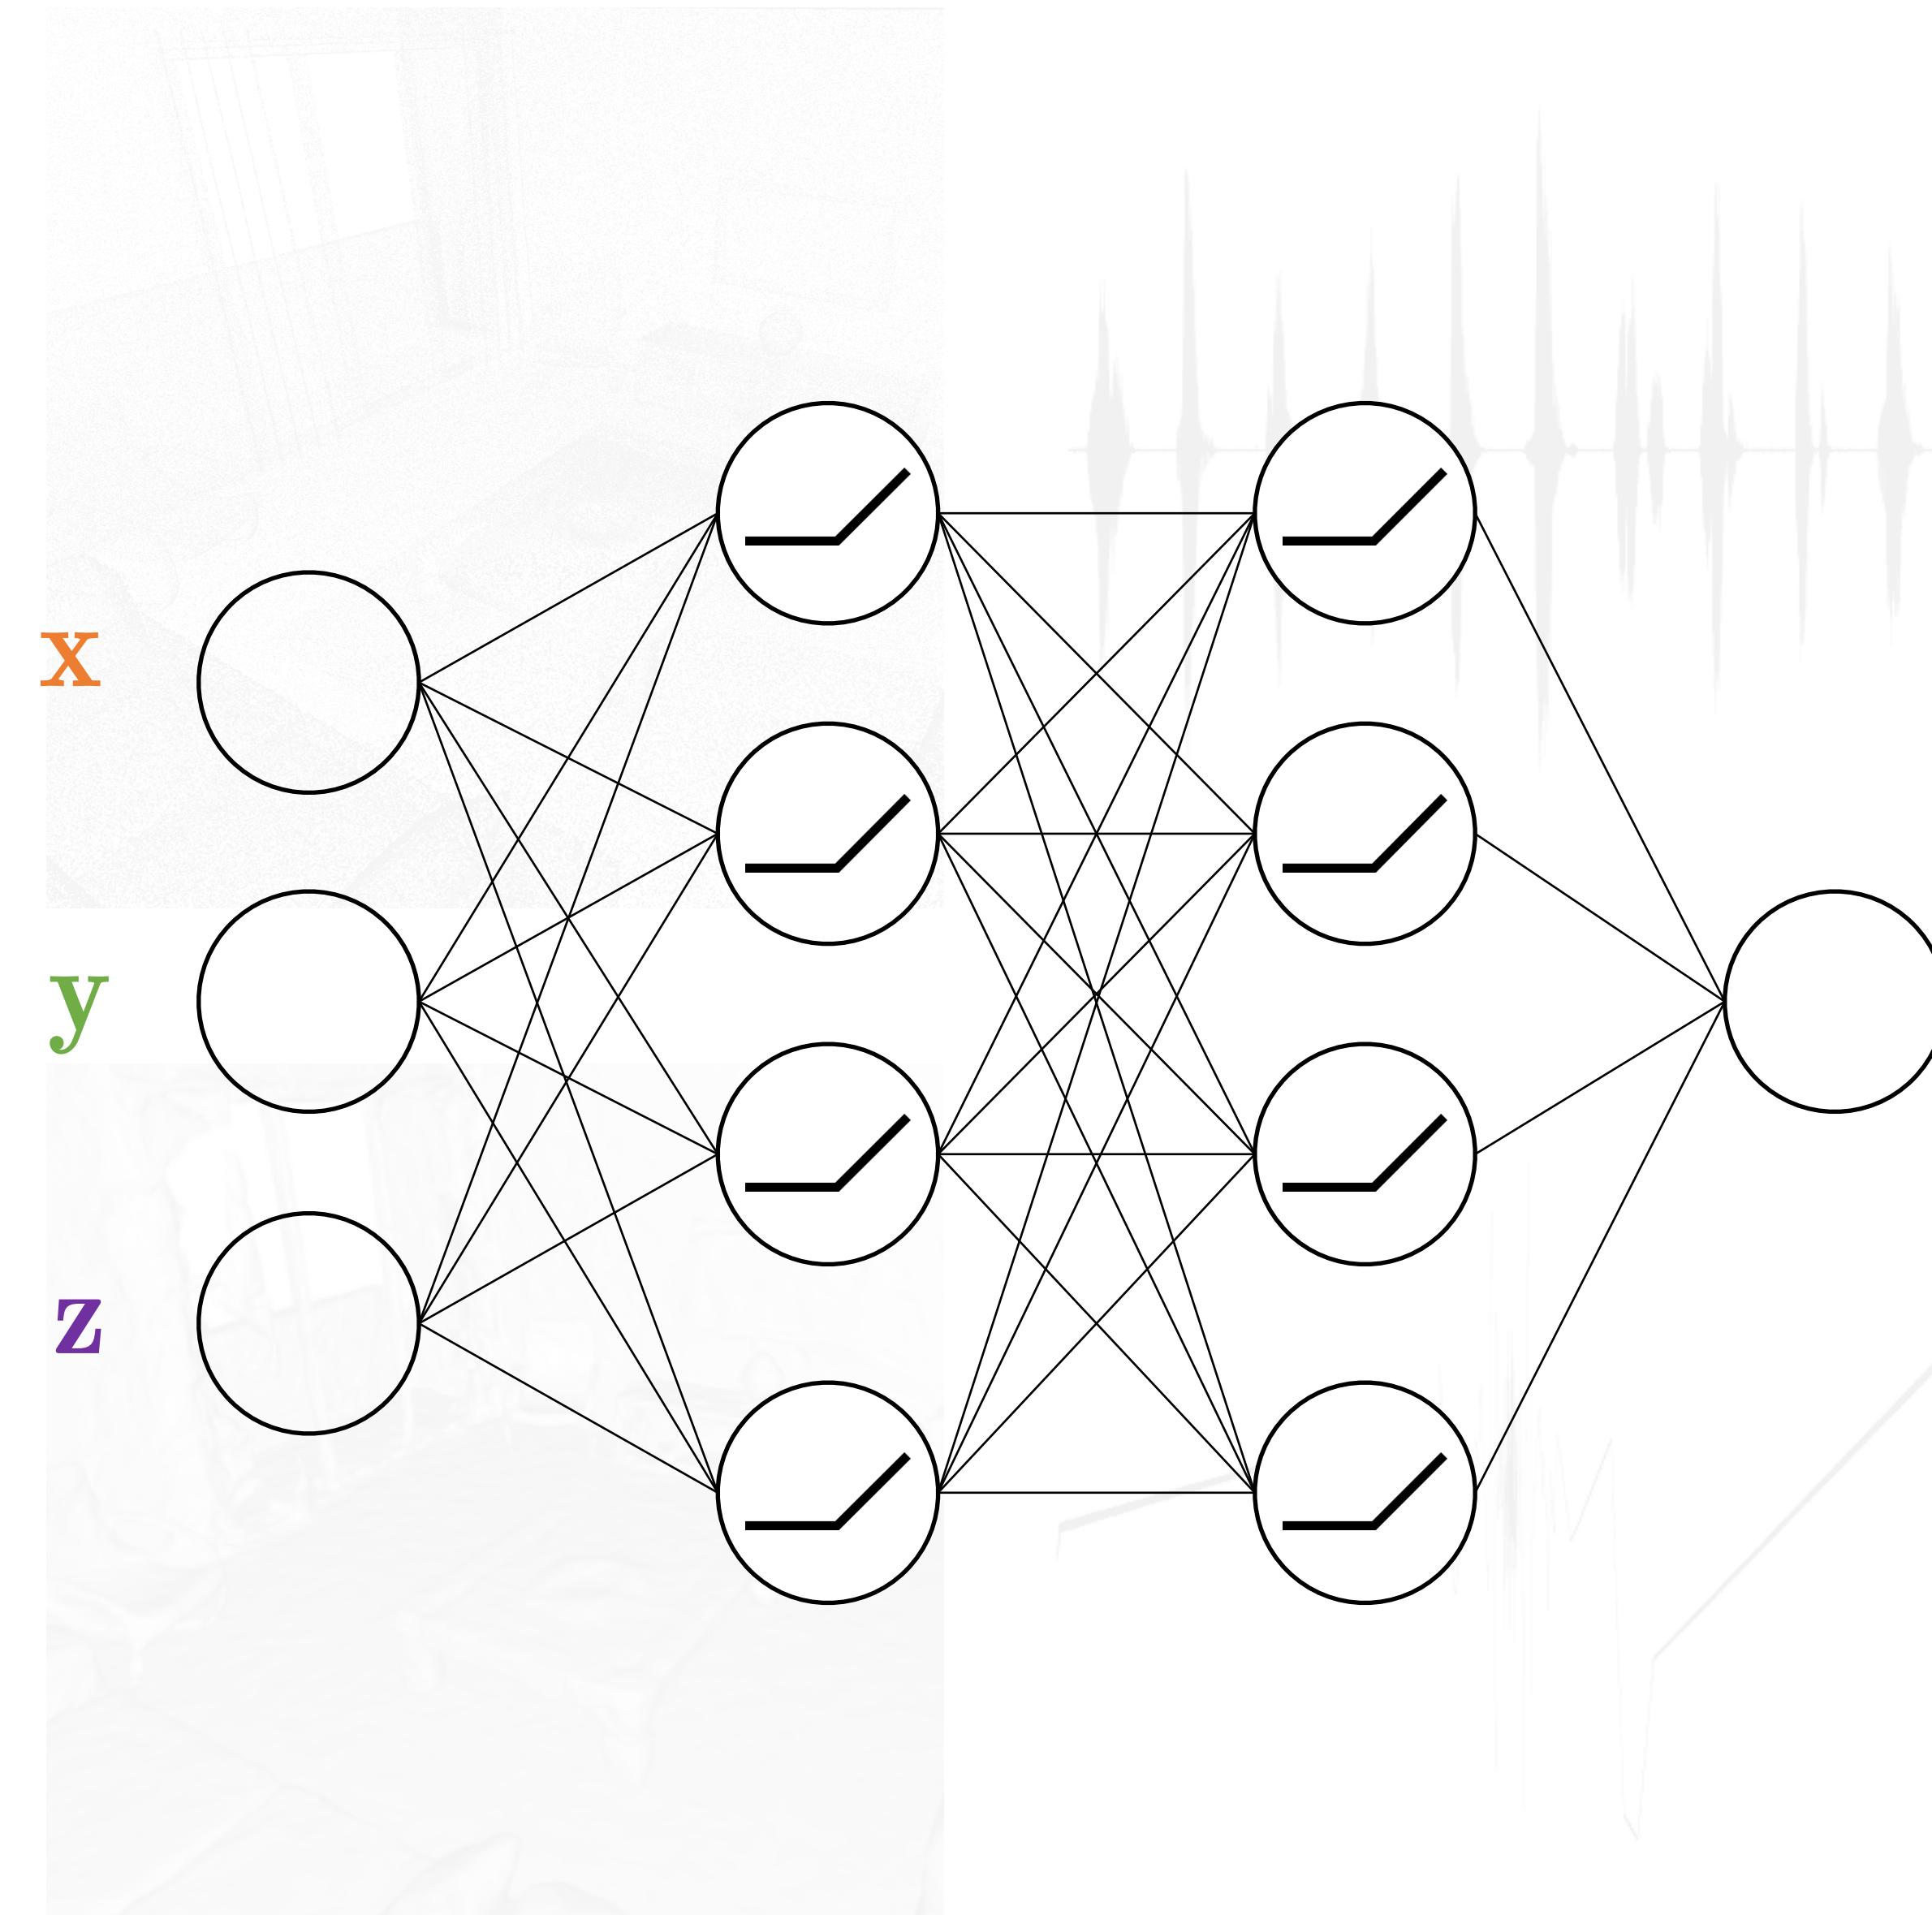


Images

Shapes

Audio

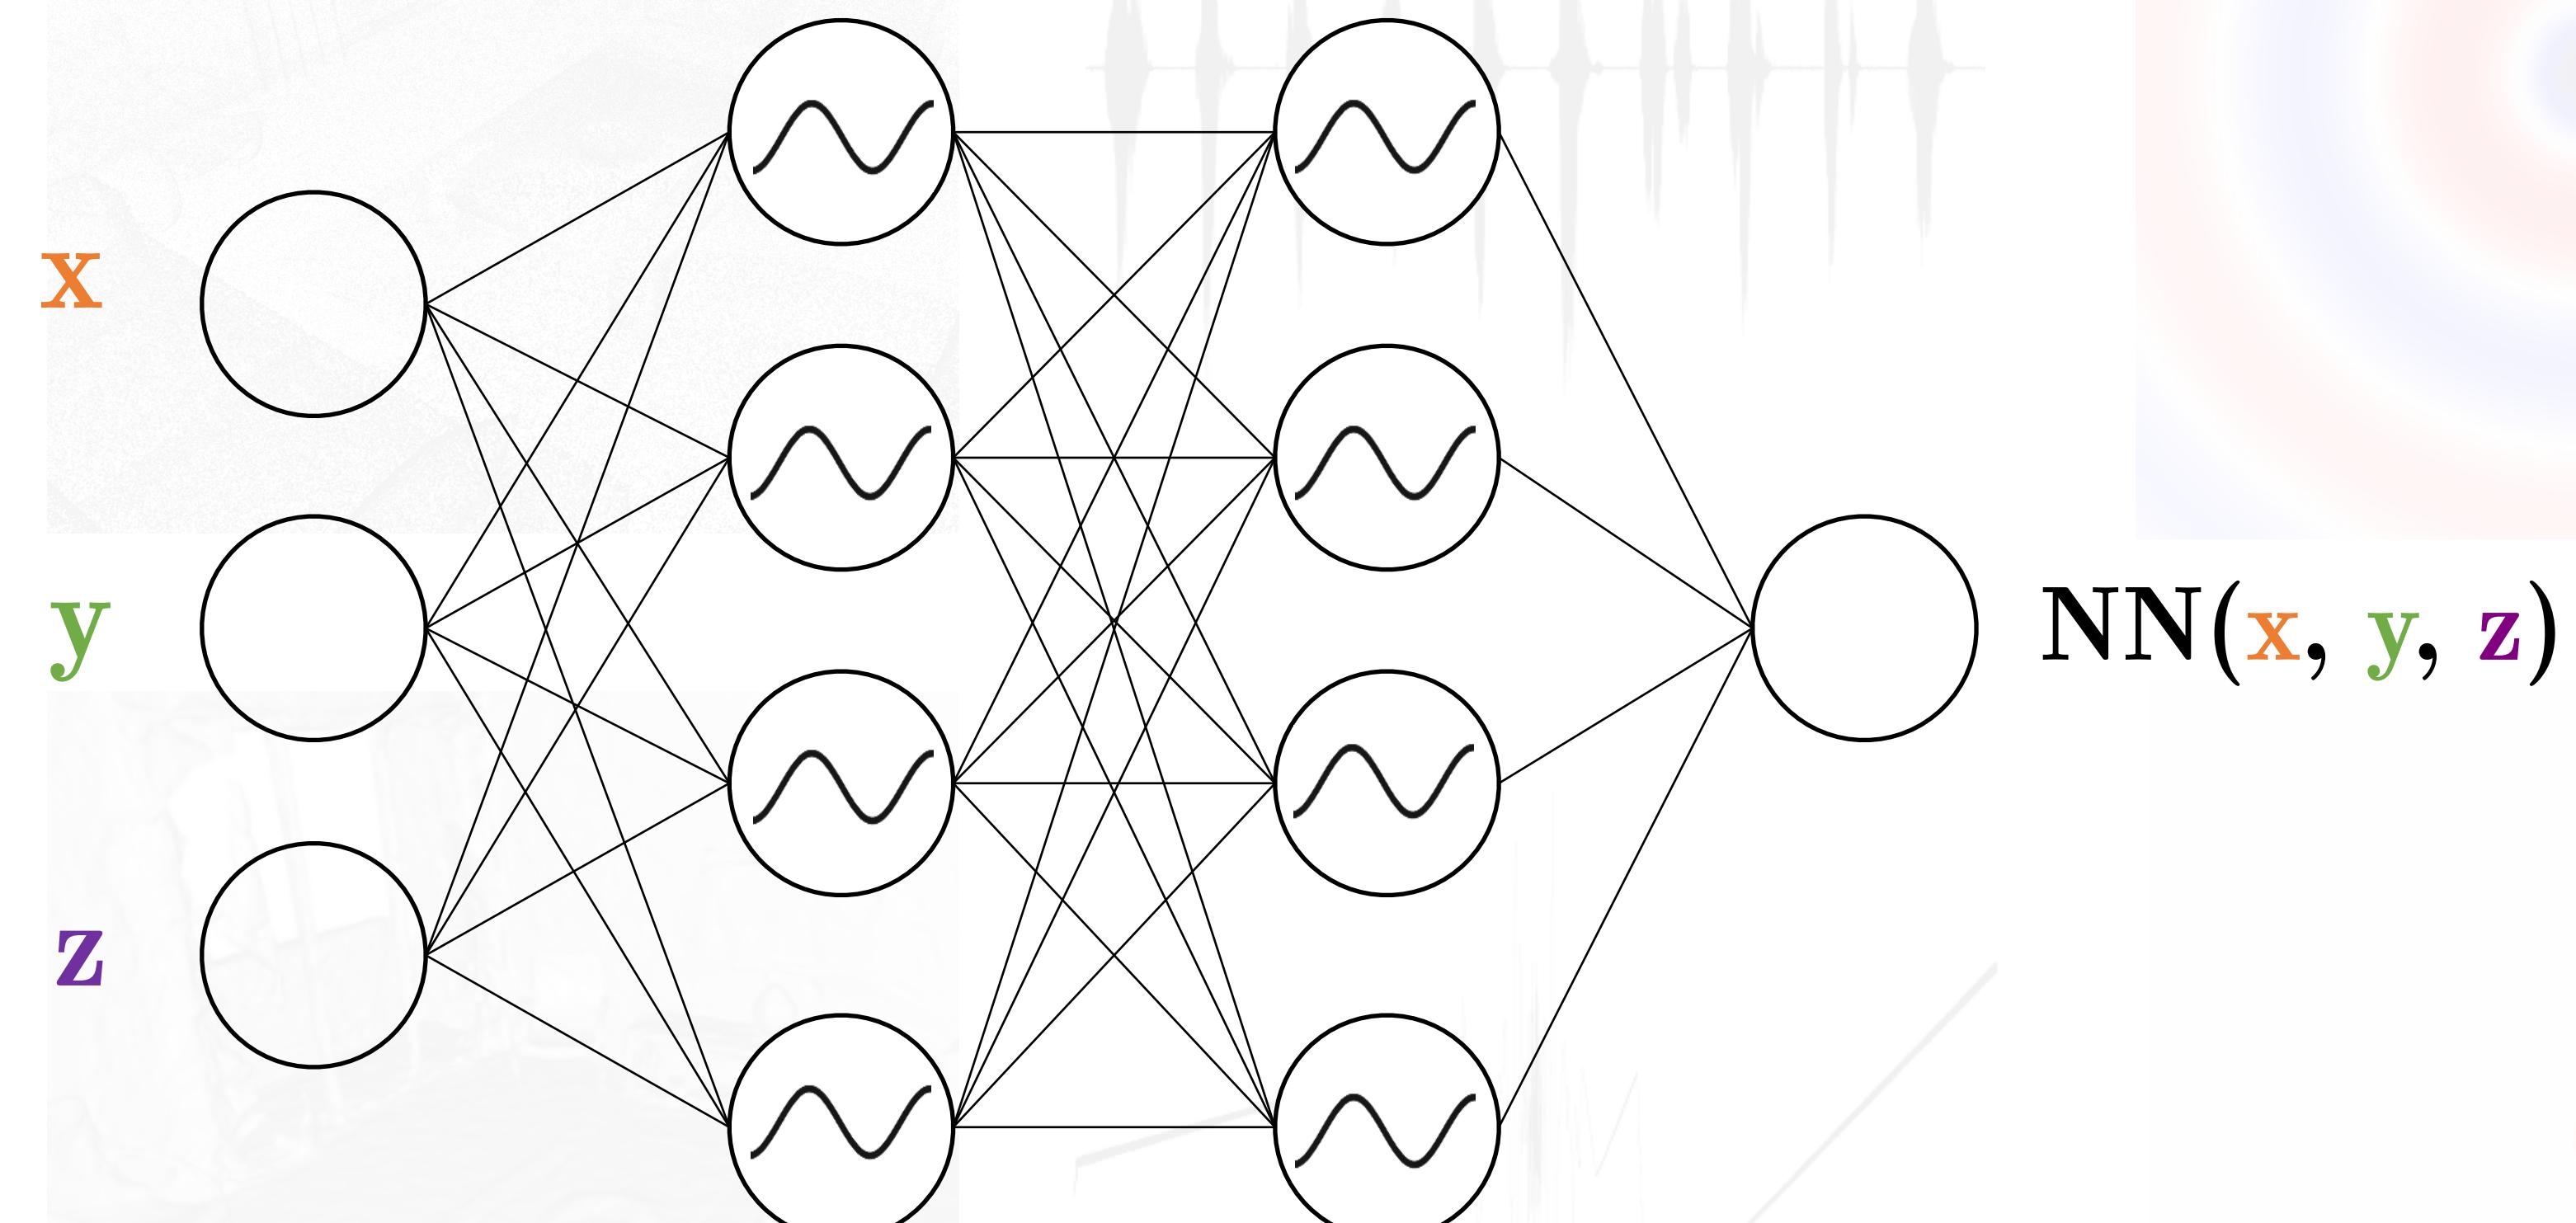
Quantities defined by a  
differential equation



$$\mathbf{NN}(\mathbf{x}, \mathbf{y}, \mathbf{z})$$



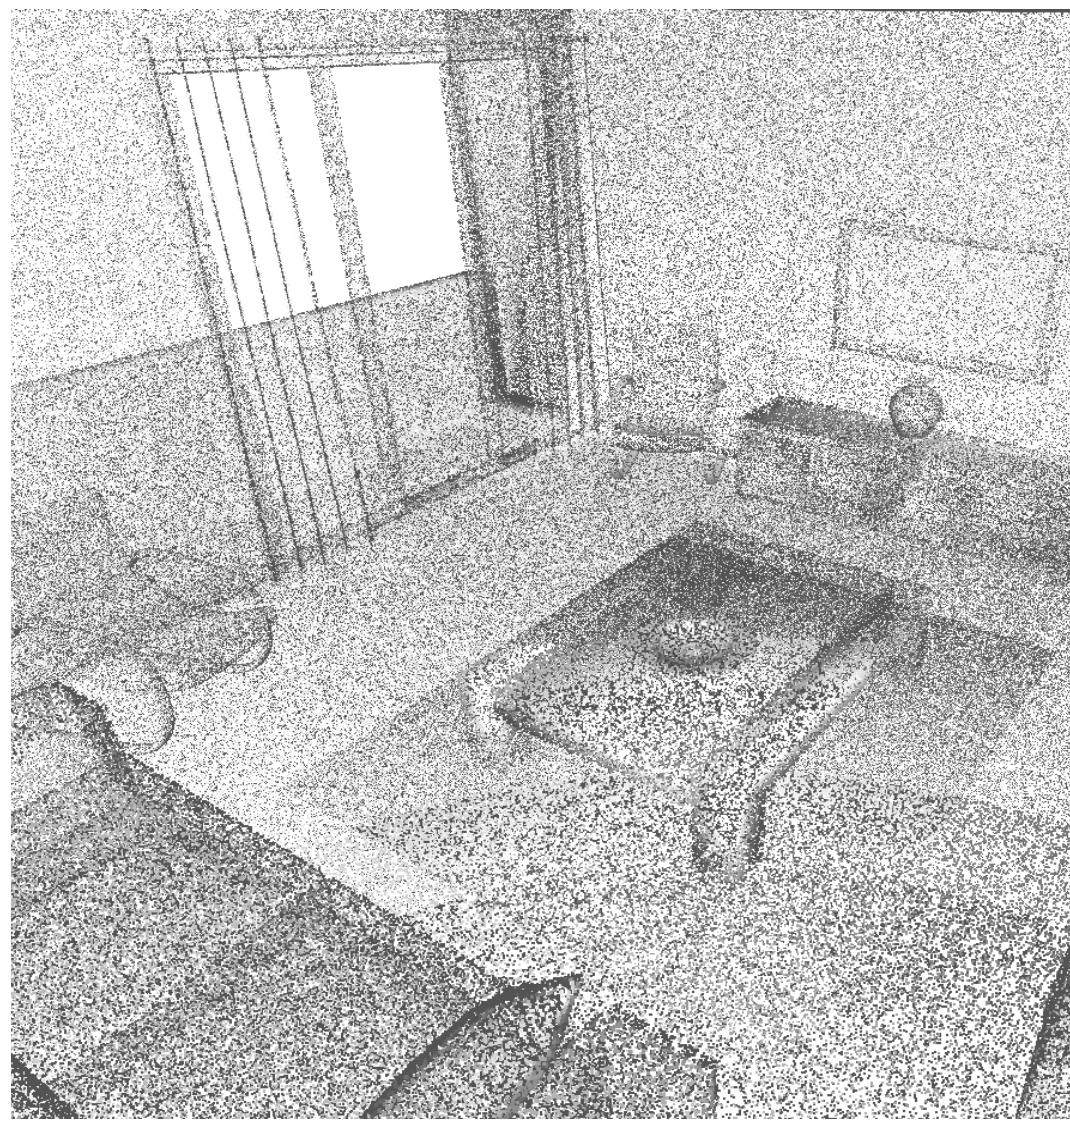
# SIREN: Sinusoidal Representation Networks



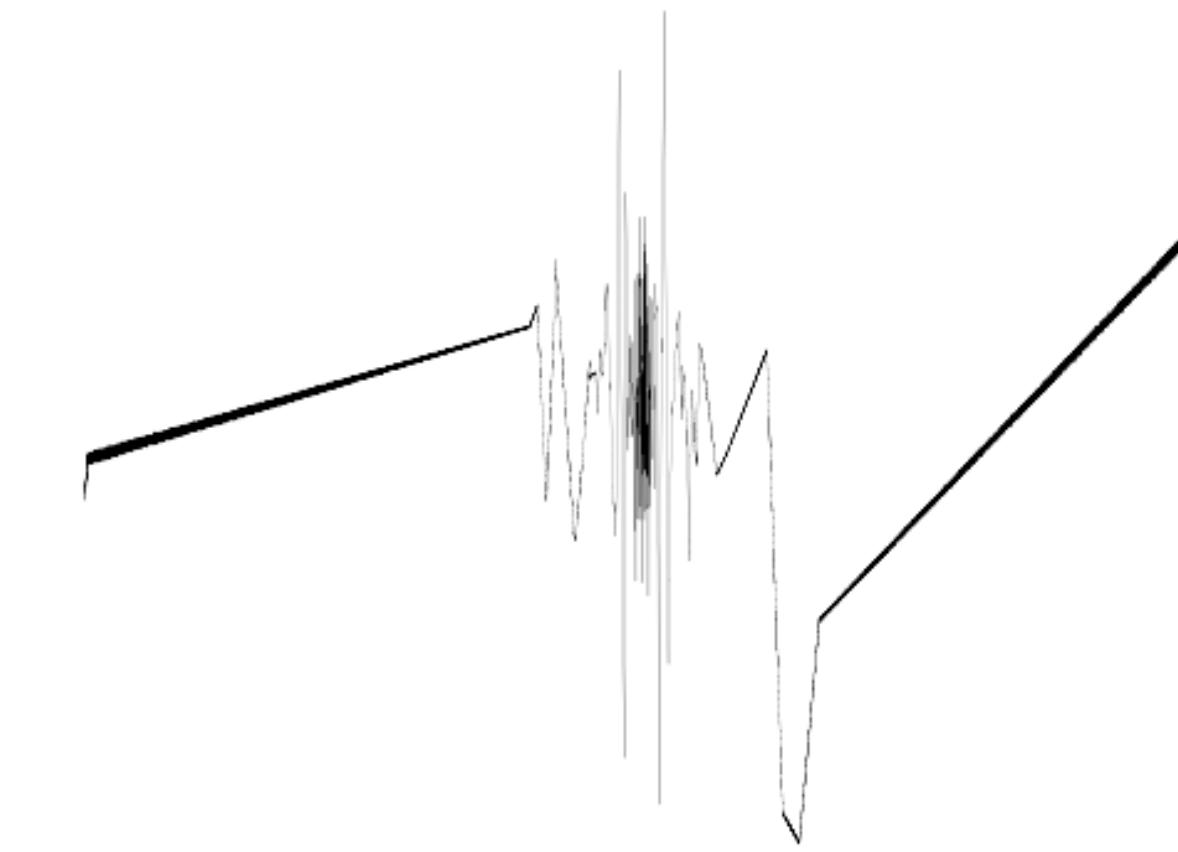
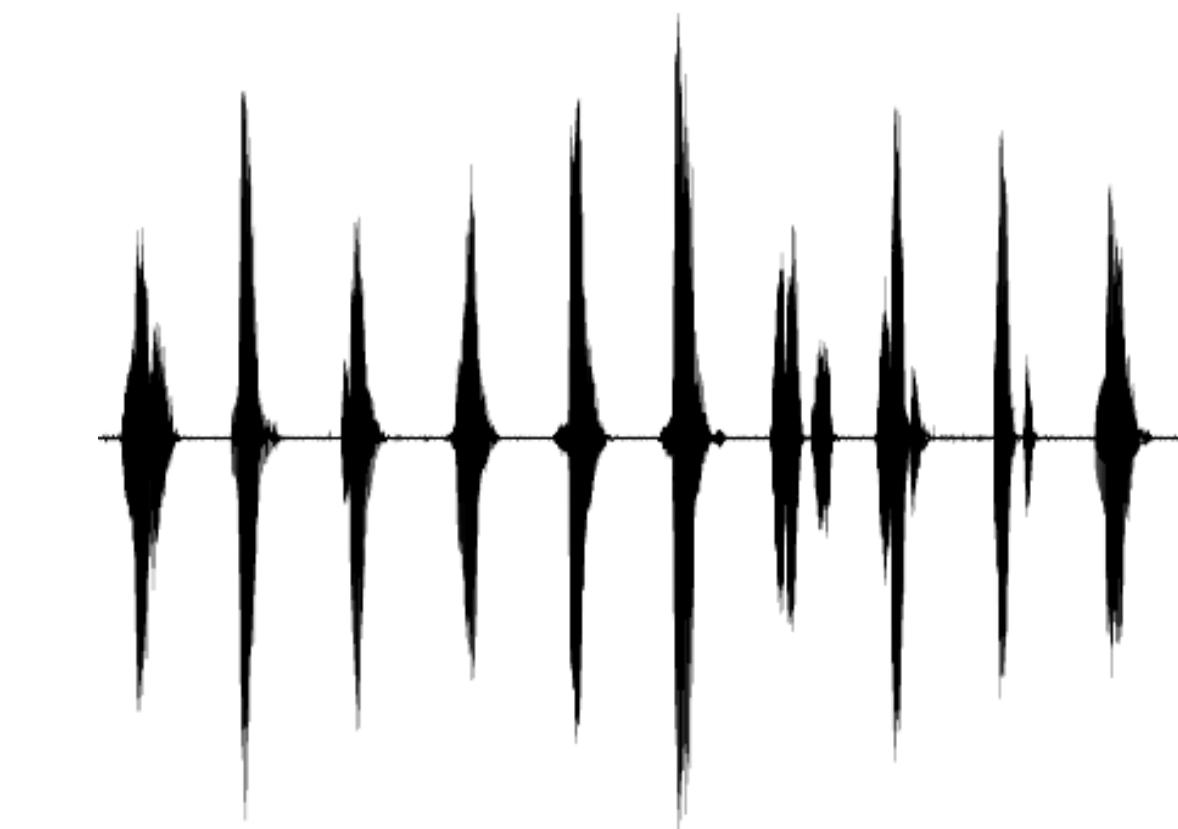
Images



Shapes



Audio



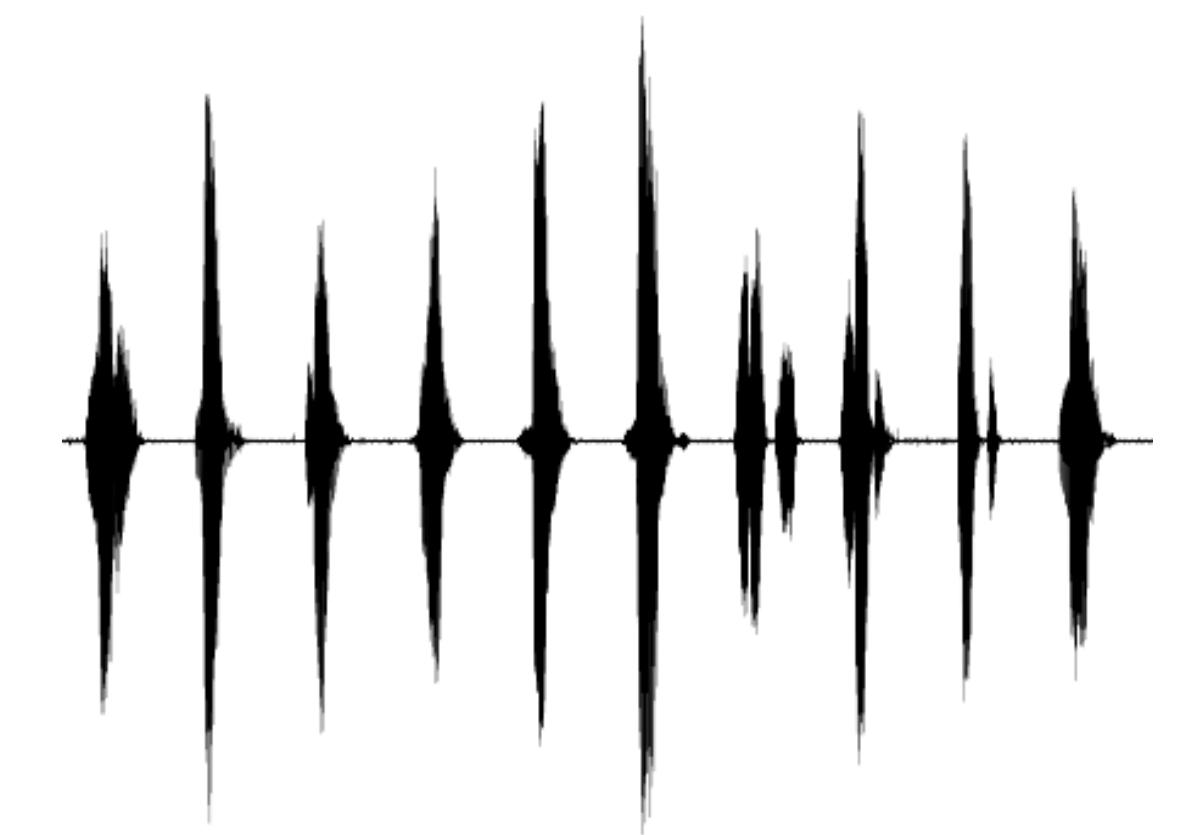
Images



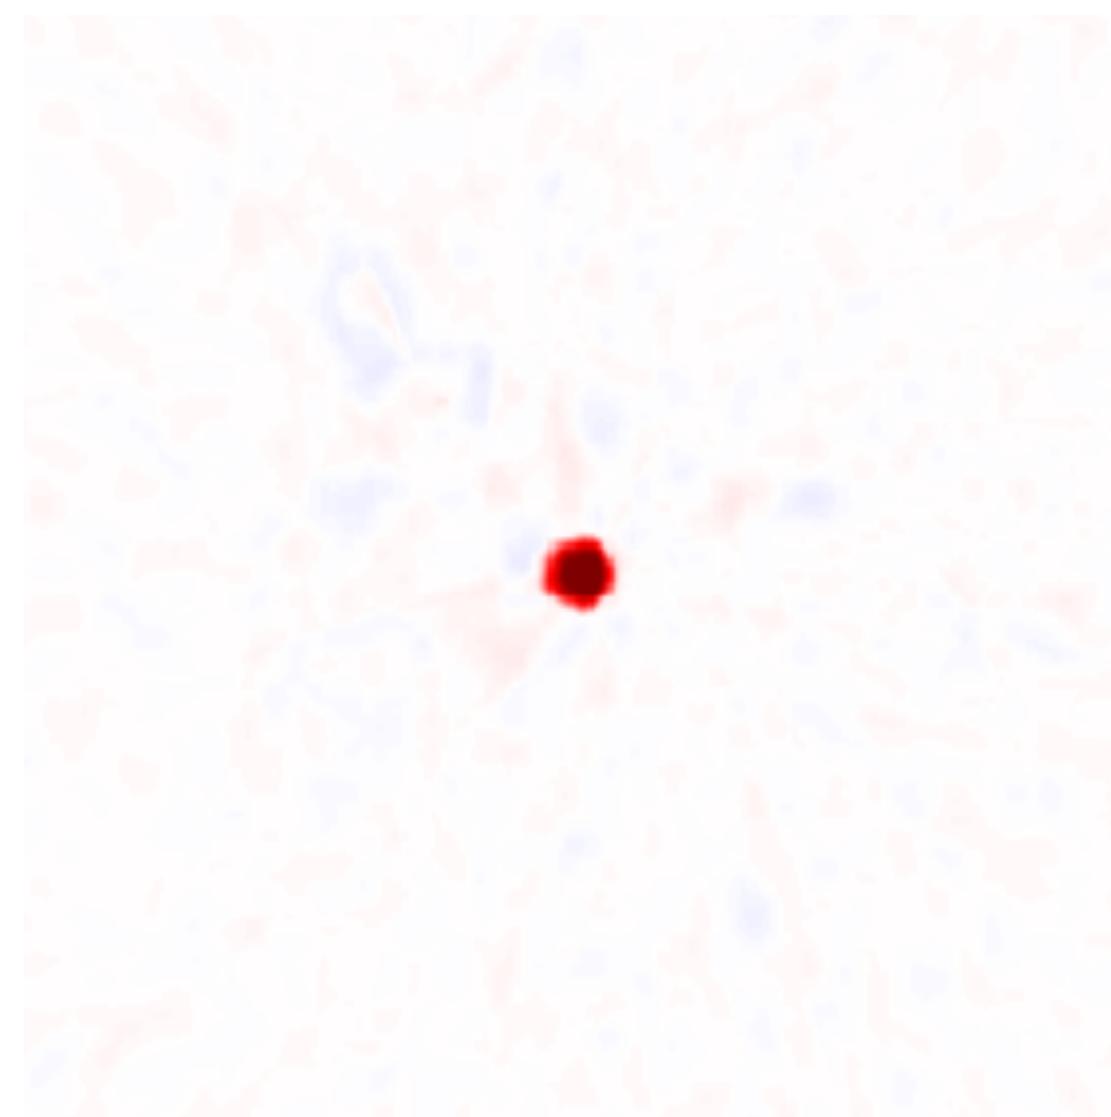
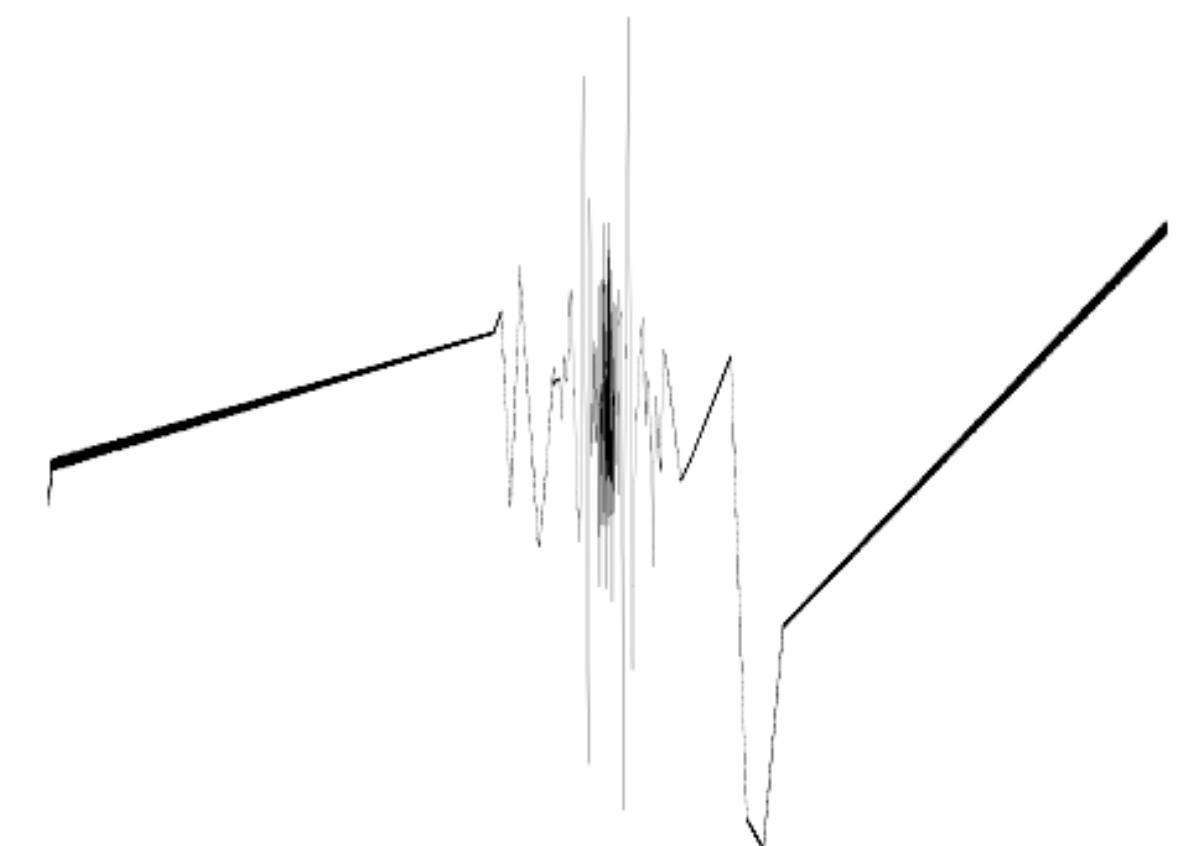
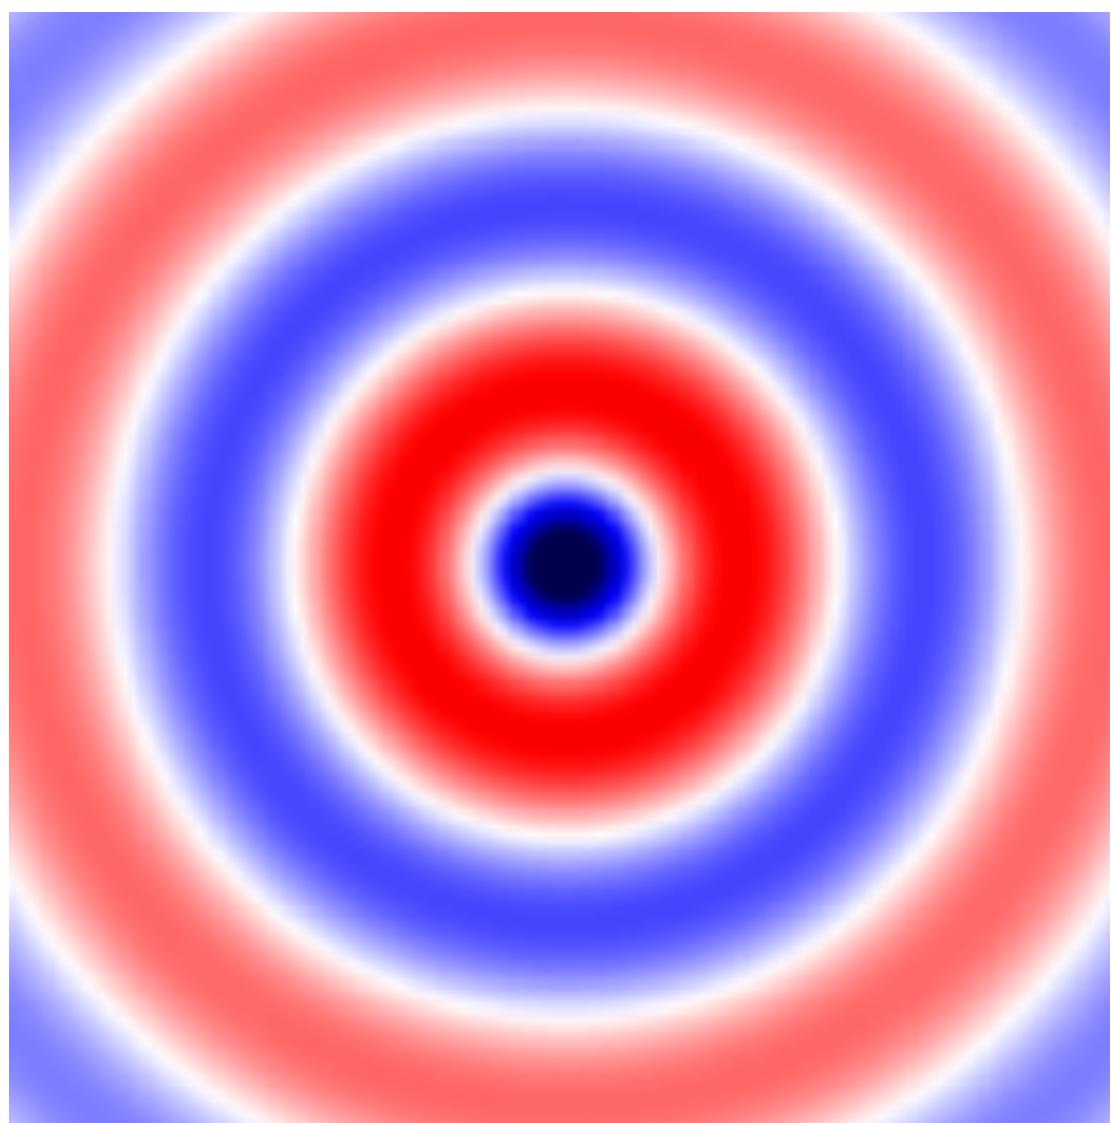
Shapes



Audio



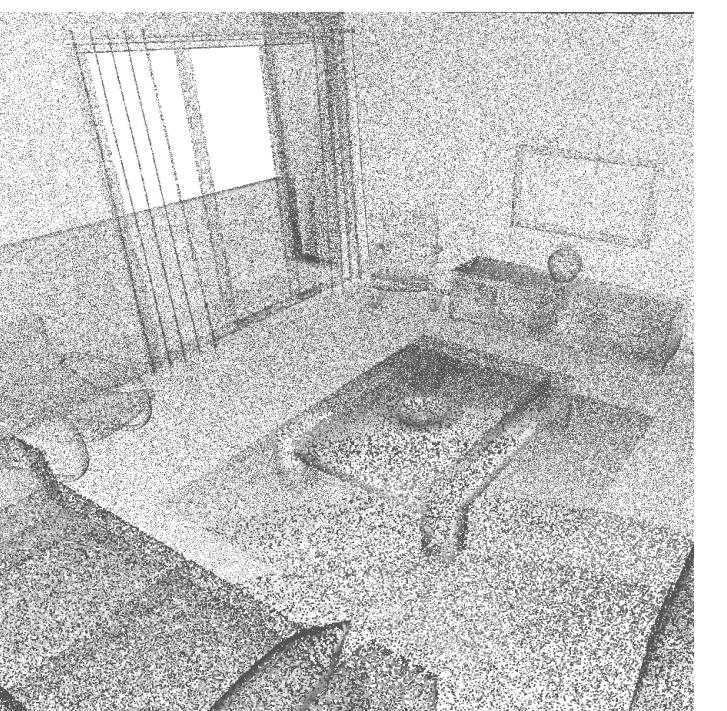
Quantities defined by a differential equation



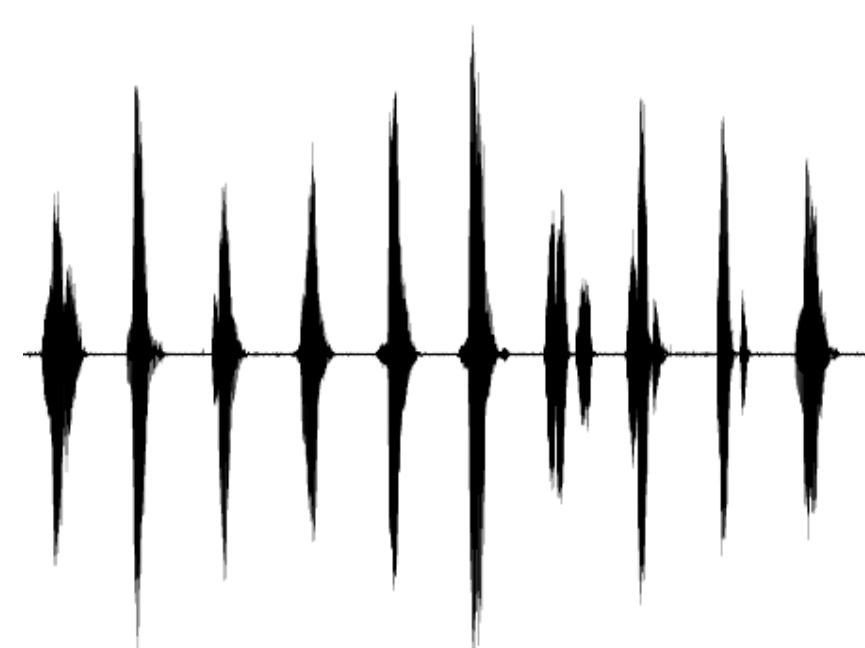
Images



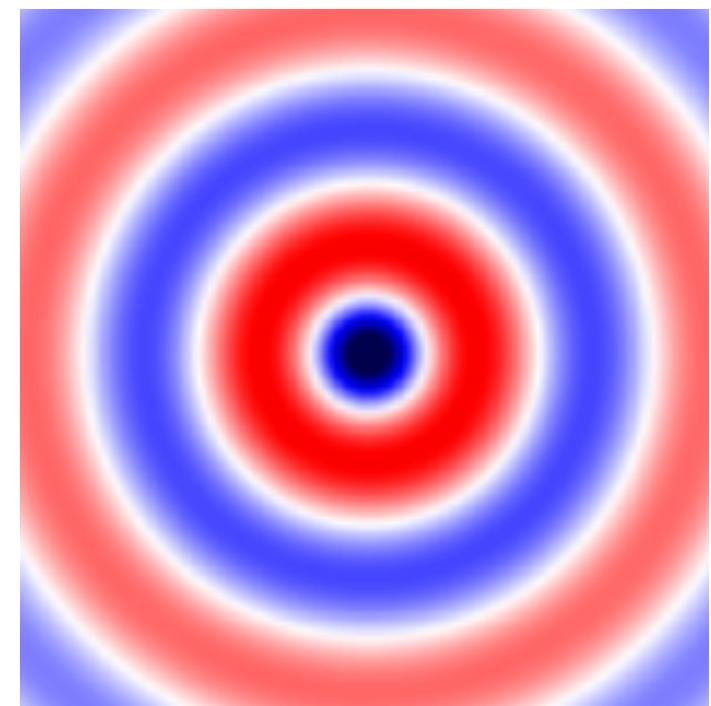
Shapes



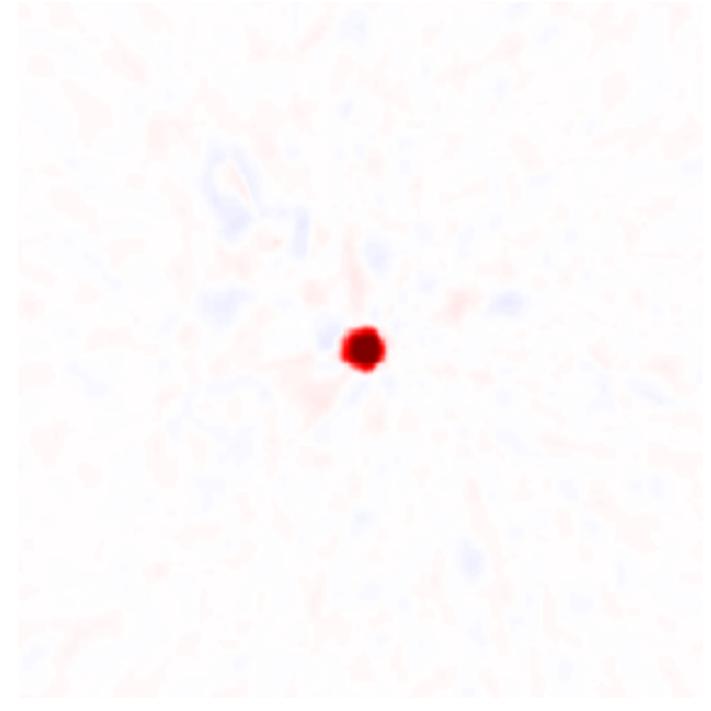
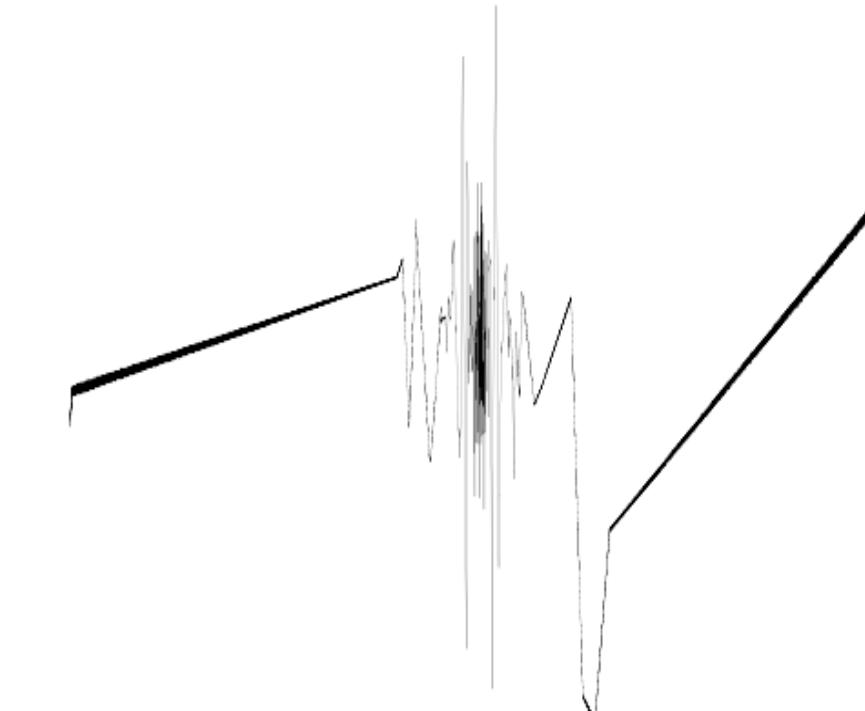
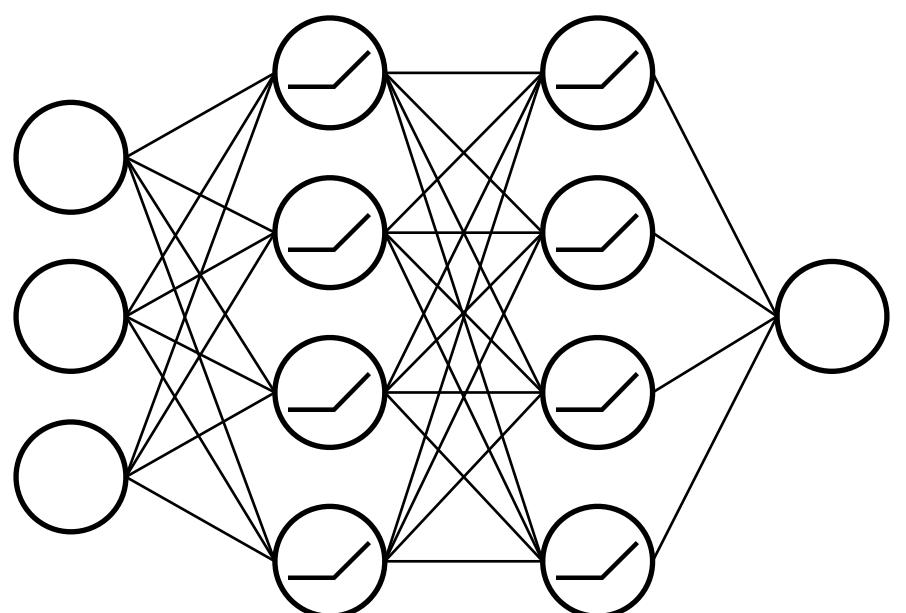
Audio



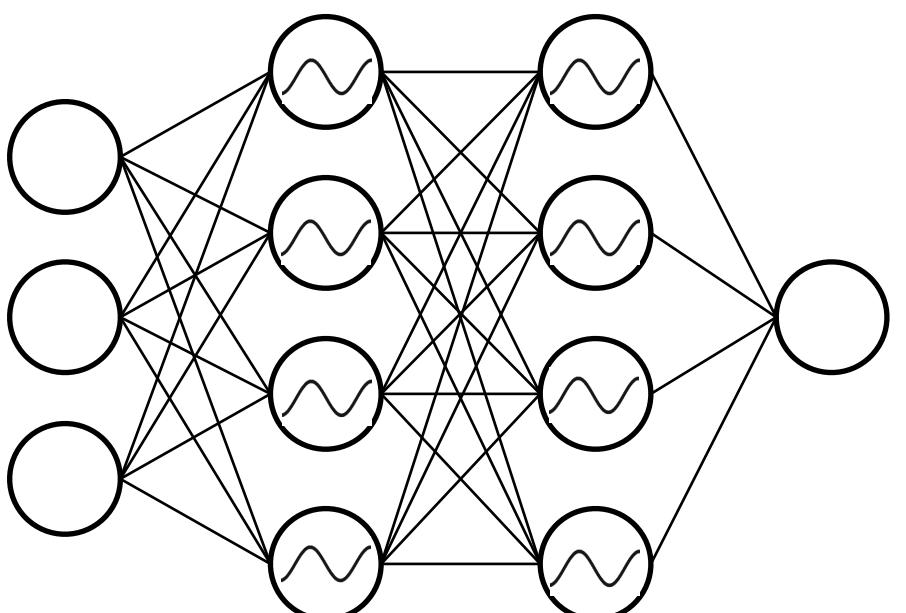
Quantities defined by a differential equation



ReLU MLP



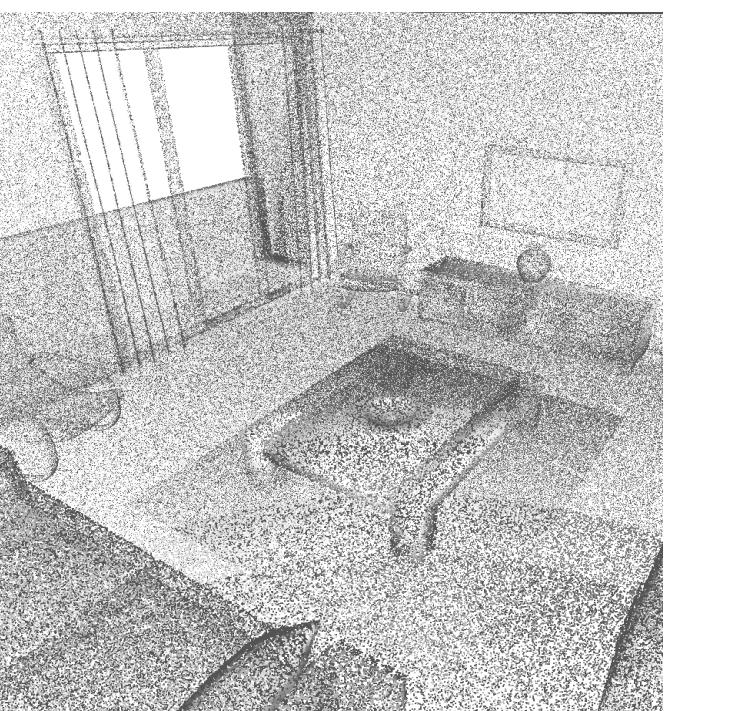
SIREN



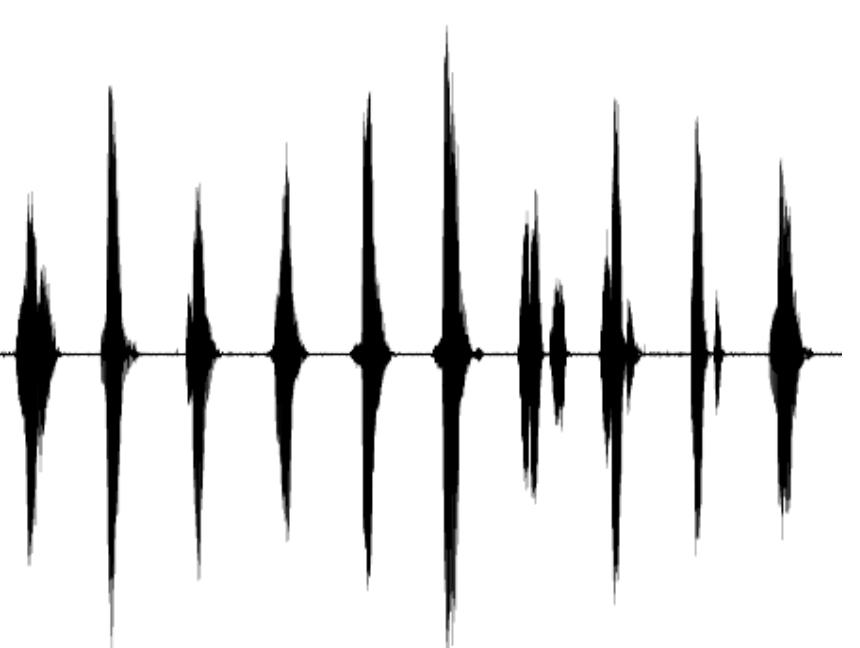
Images



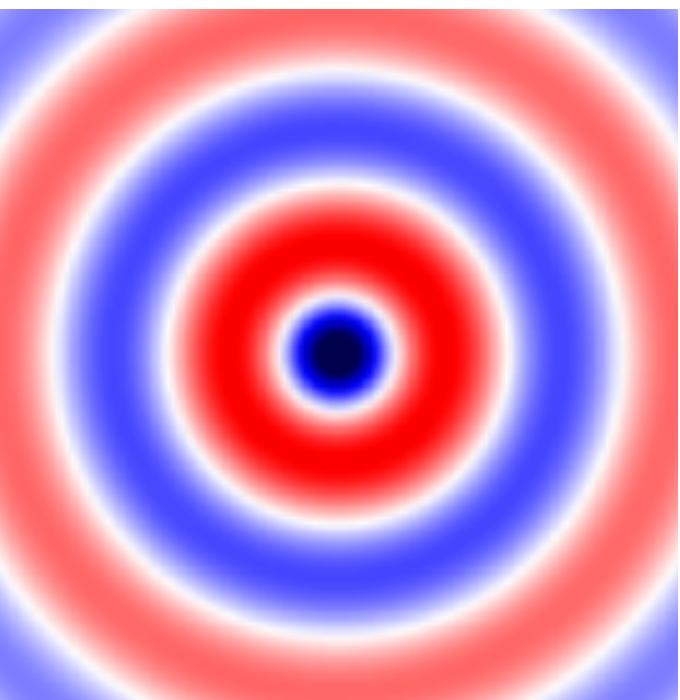
Shapes



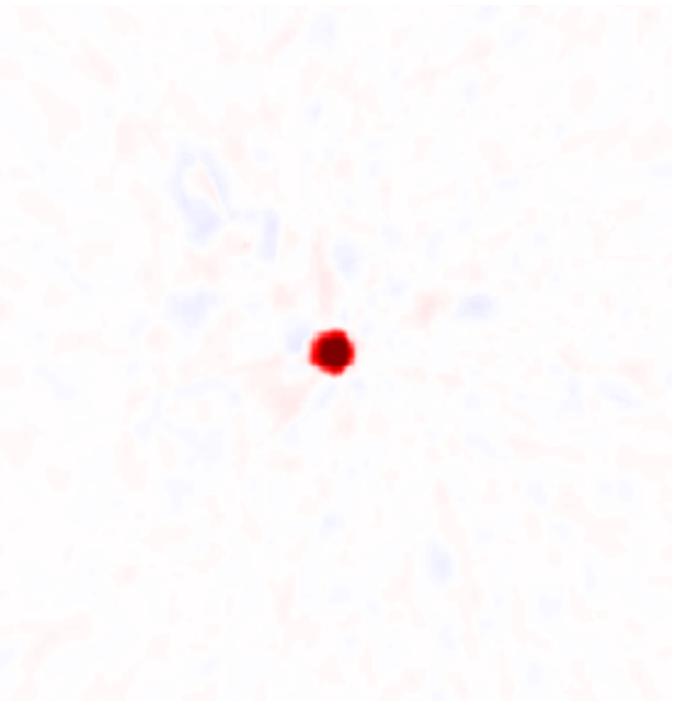
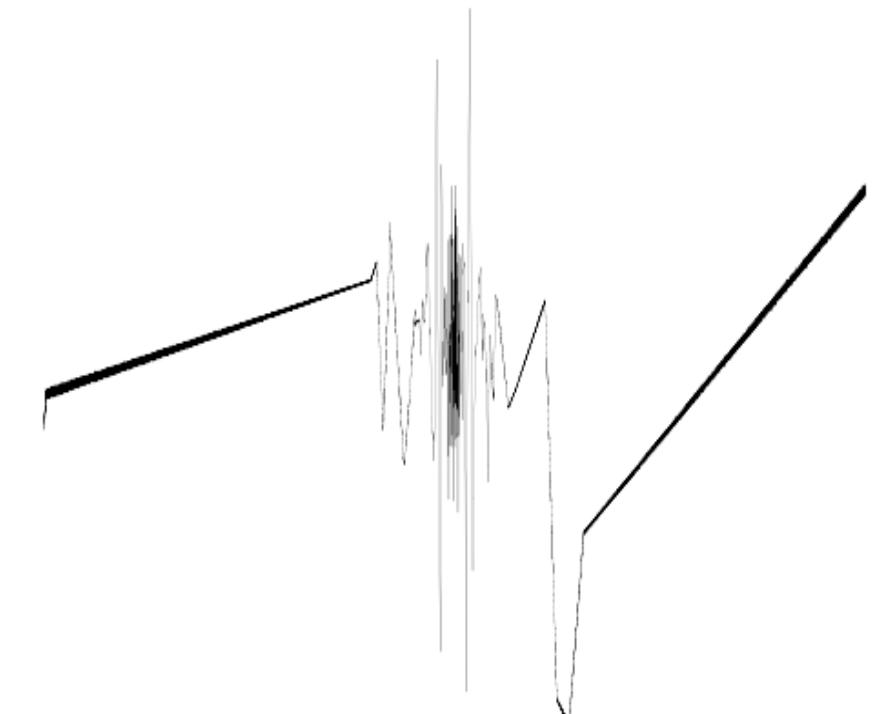
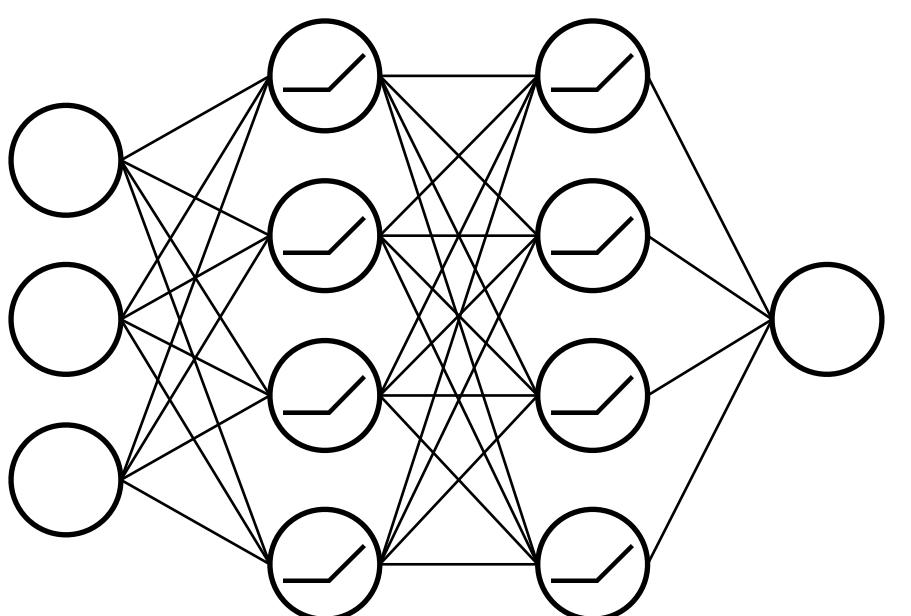
Audio



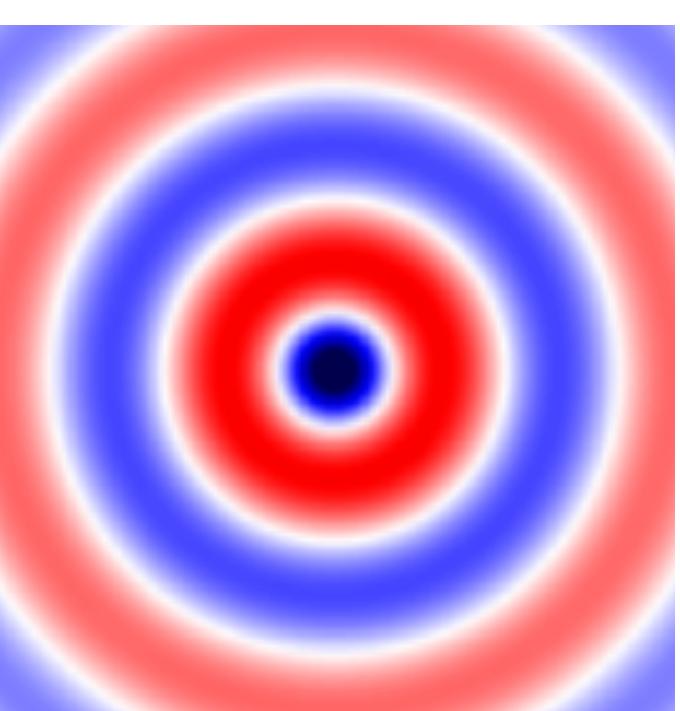
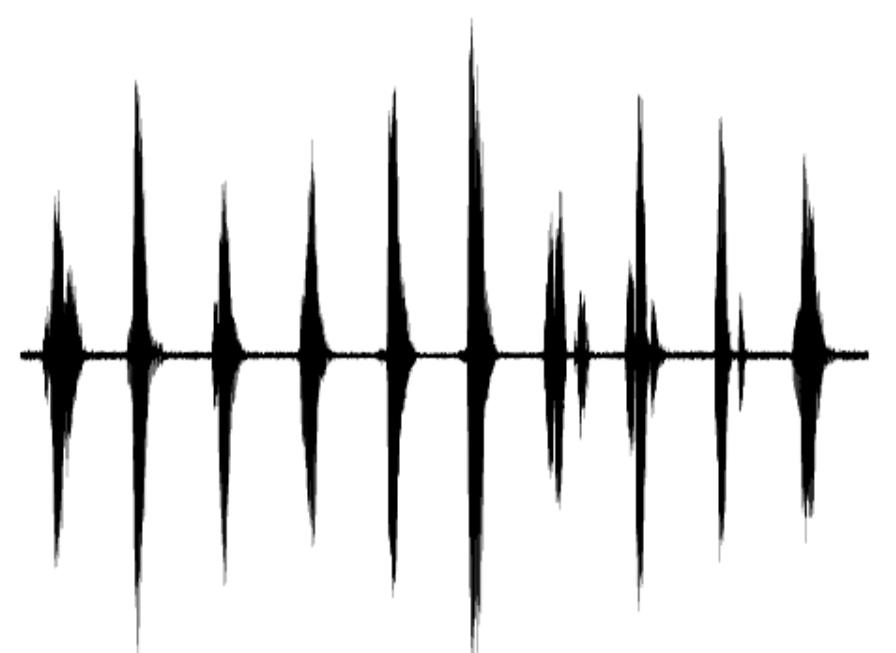
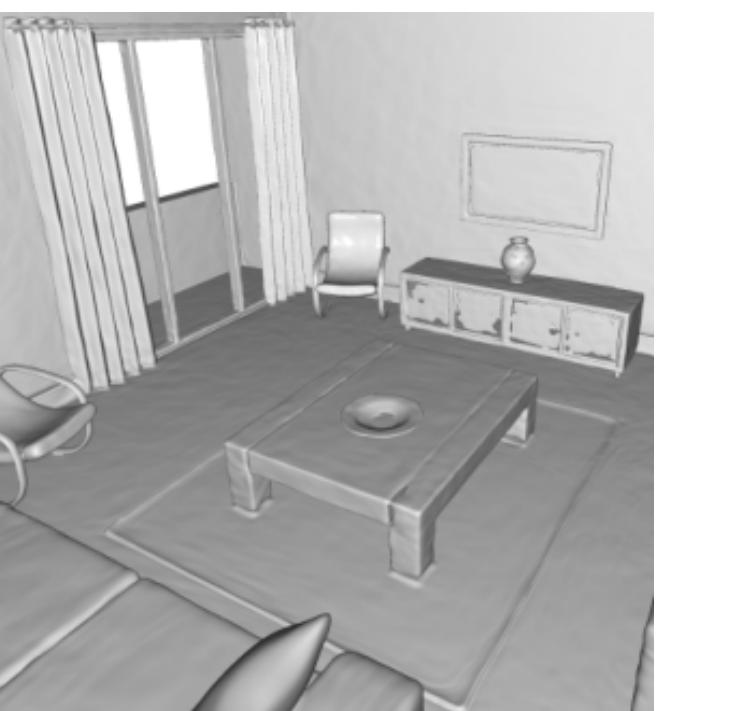
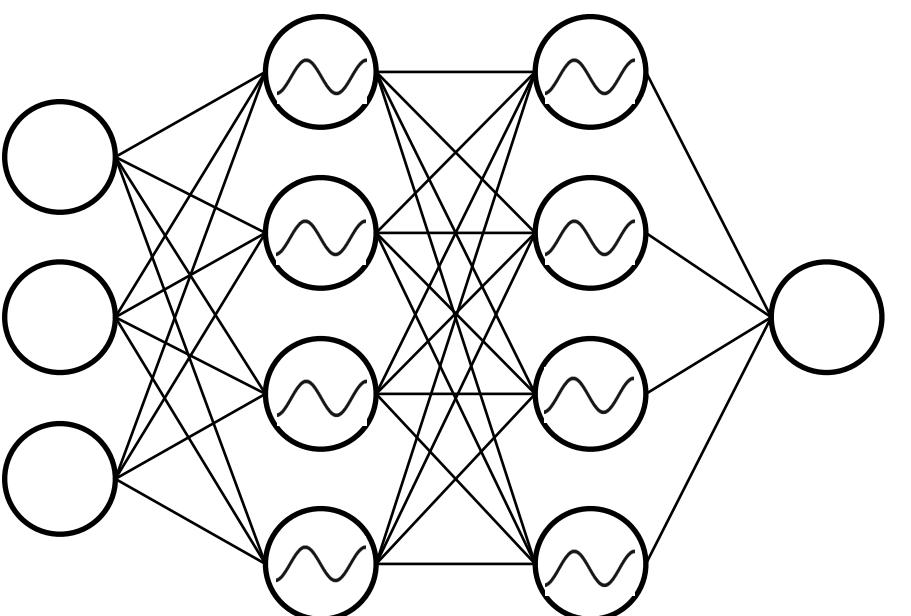
Quantities defined by a differential equation



ReLU MLP



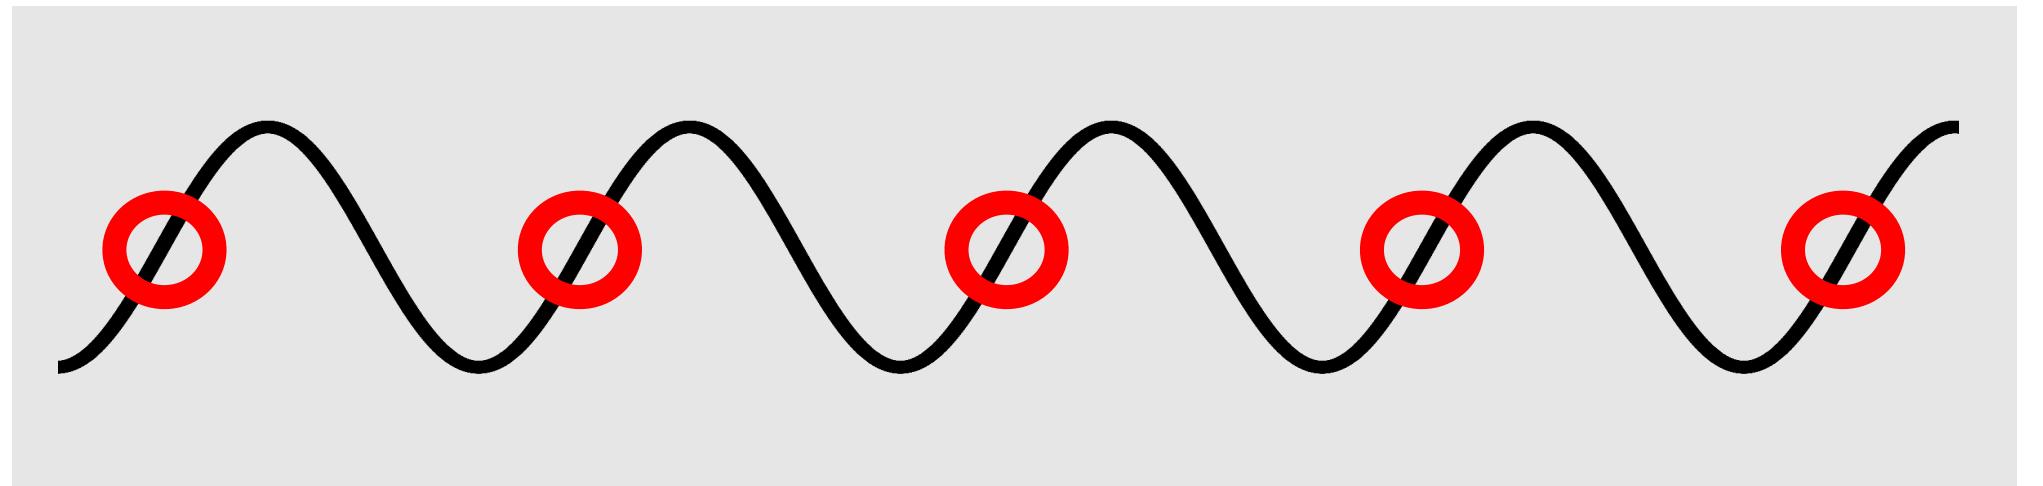
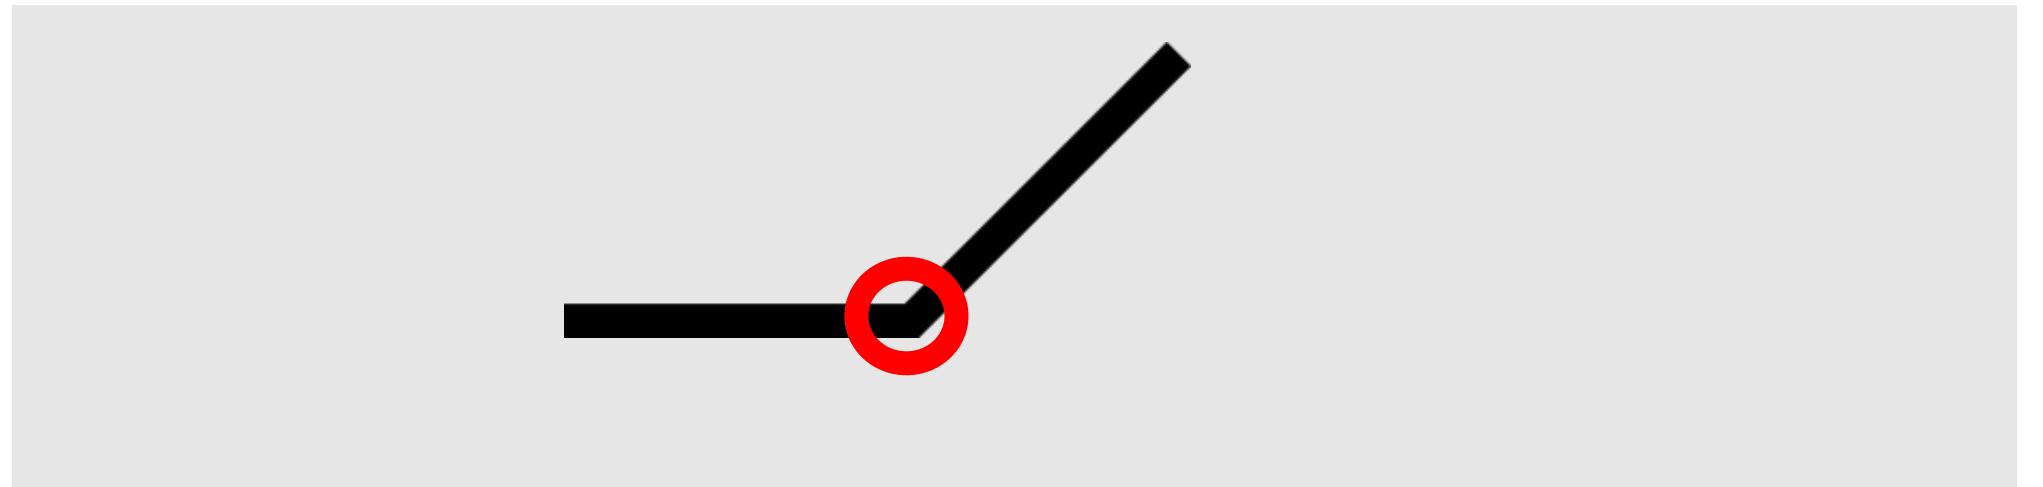
SIREN



# Locality

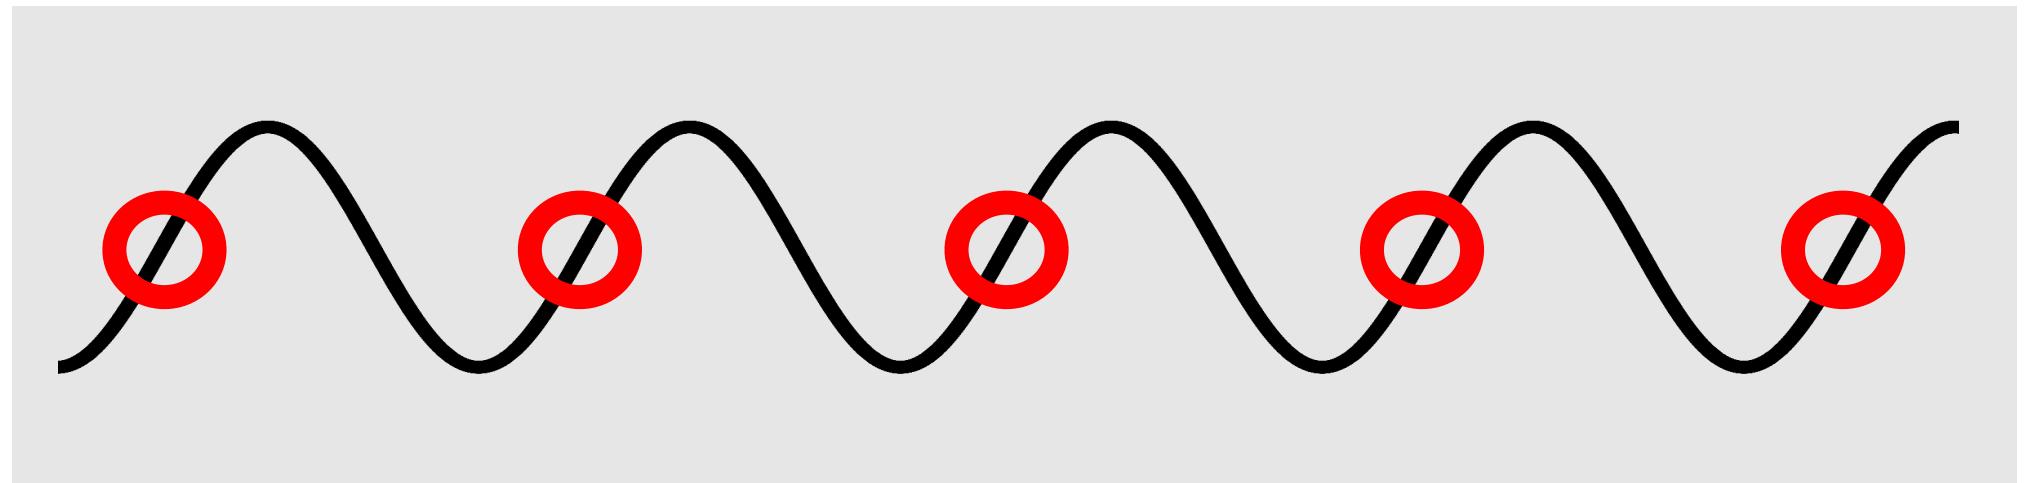
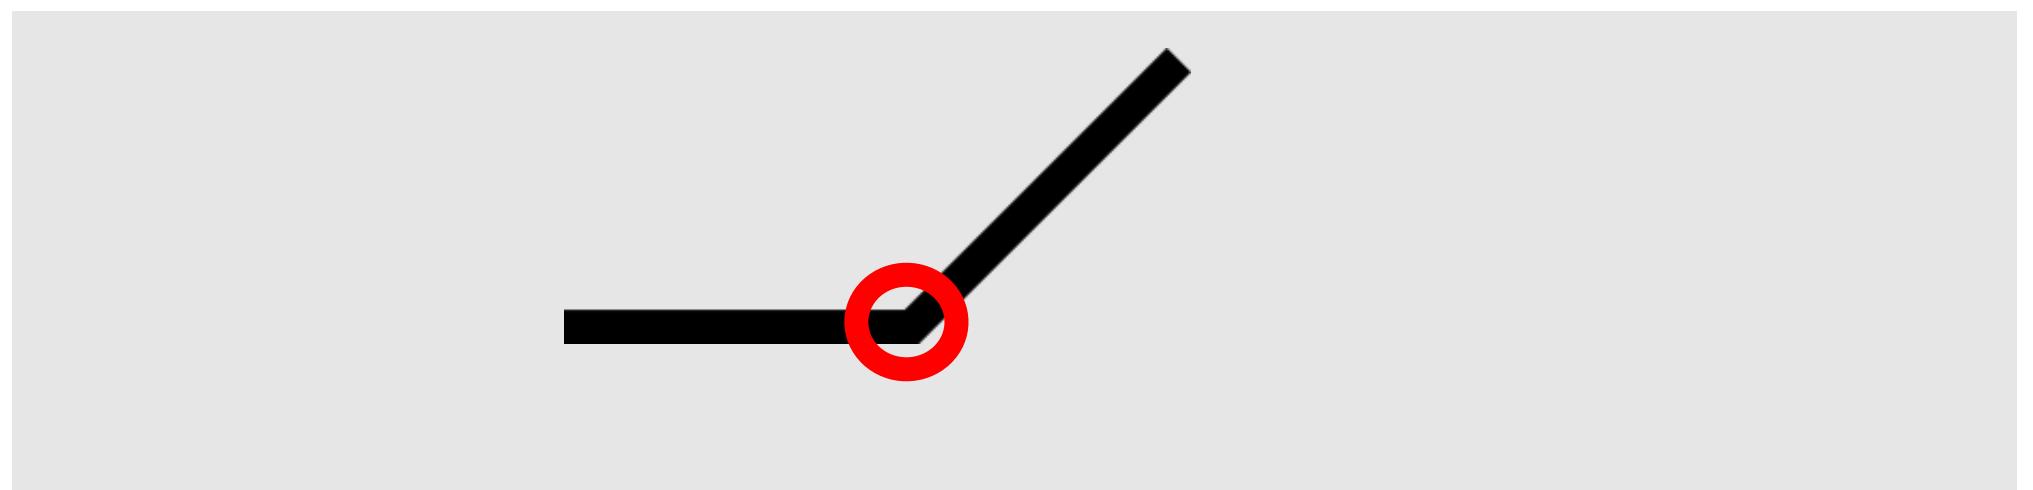
$$\delta/\delta x$$

$$\delta^2/\delta x^2$$



Periodicity allows SIREN to replicate activations across input domain

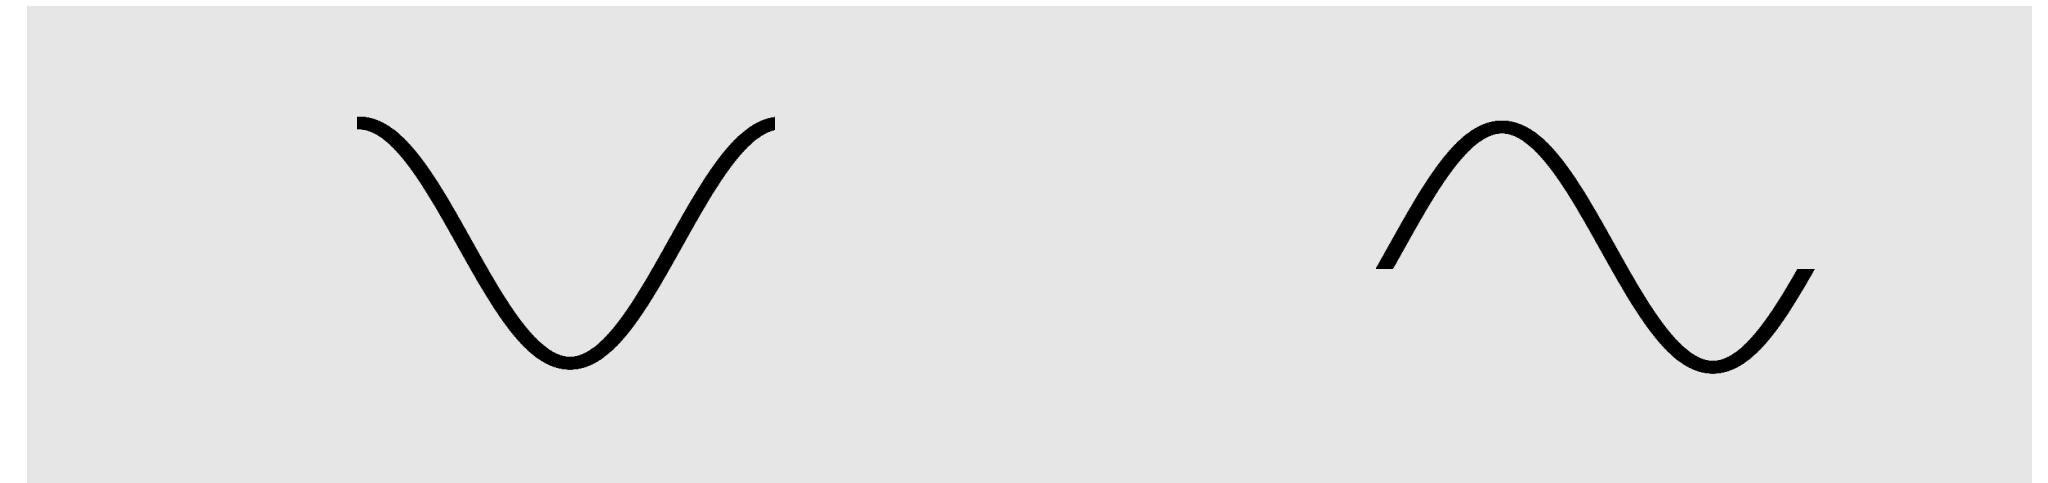
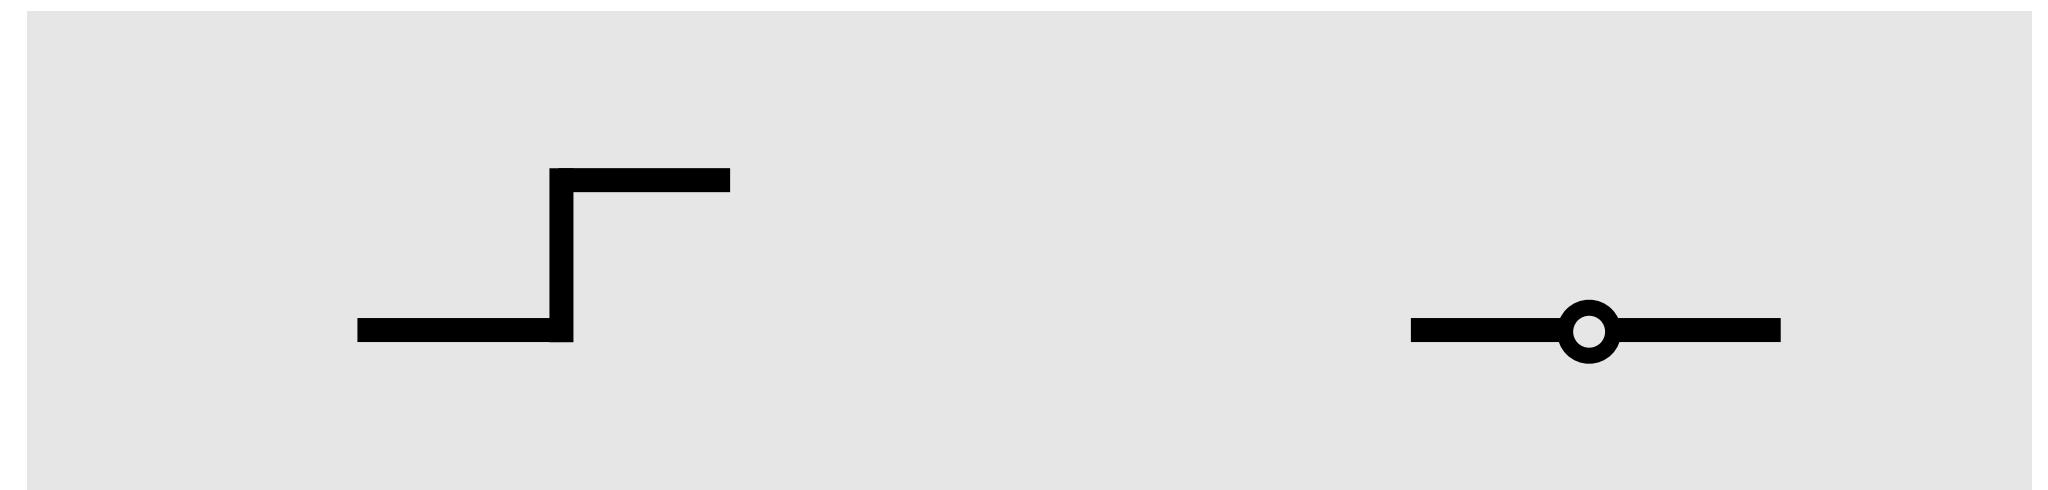
# Locality



Periodicity allows SIREN to replicate activations across input domain

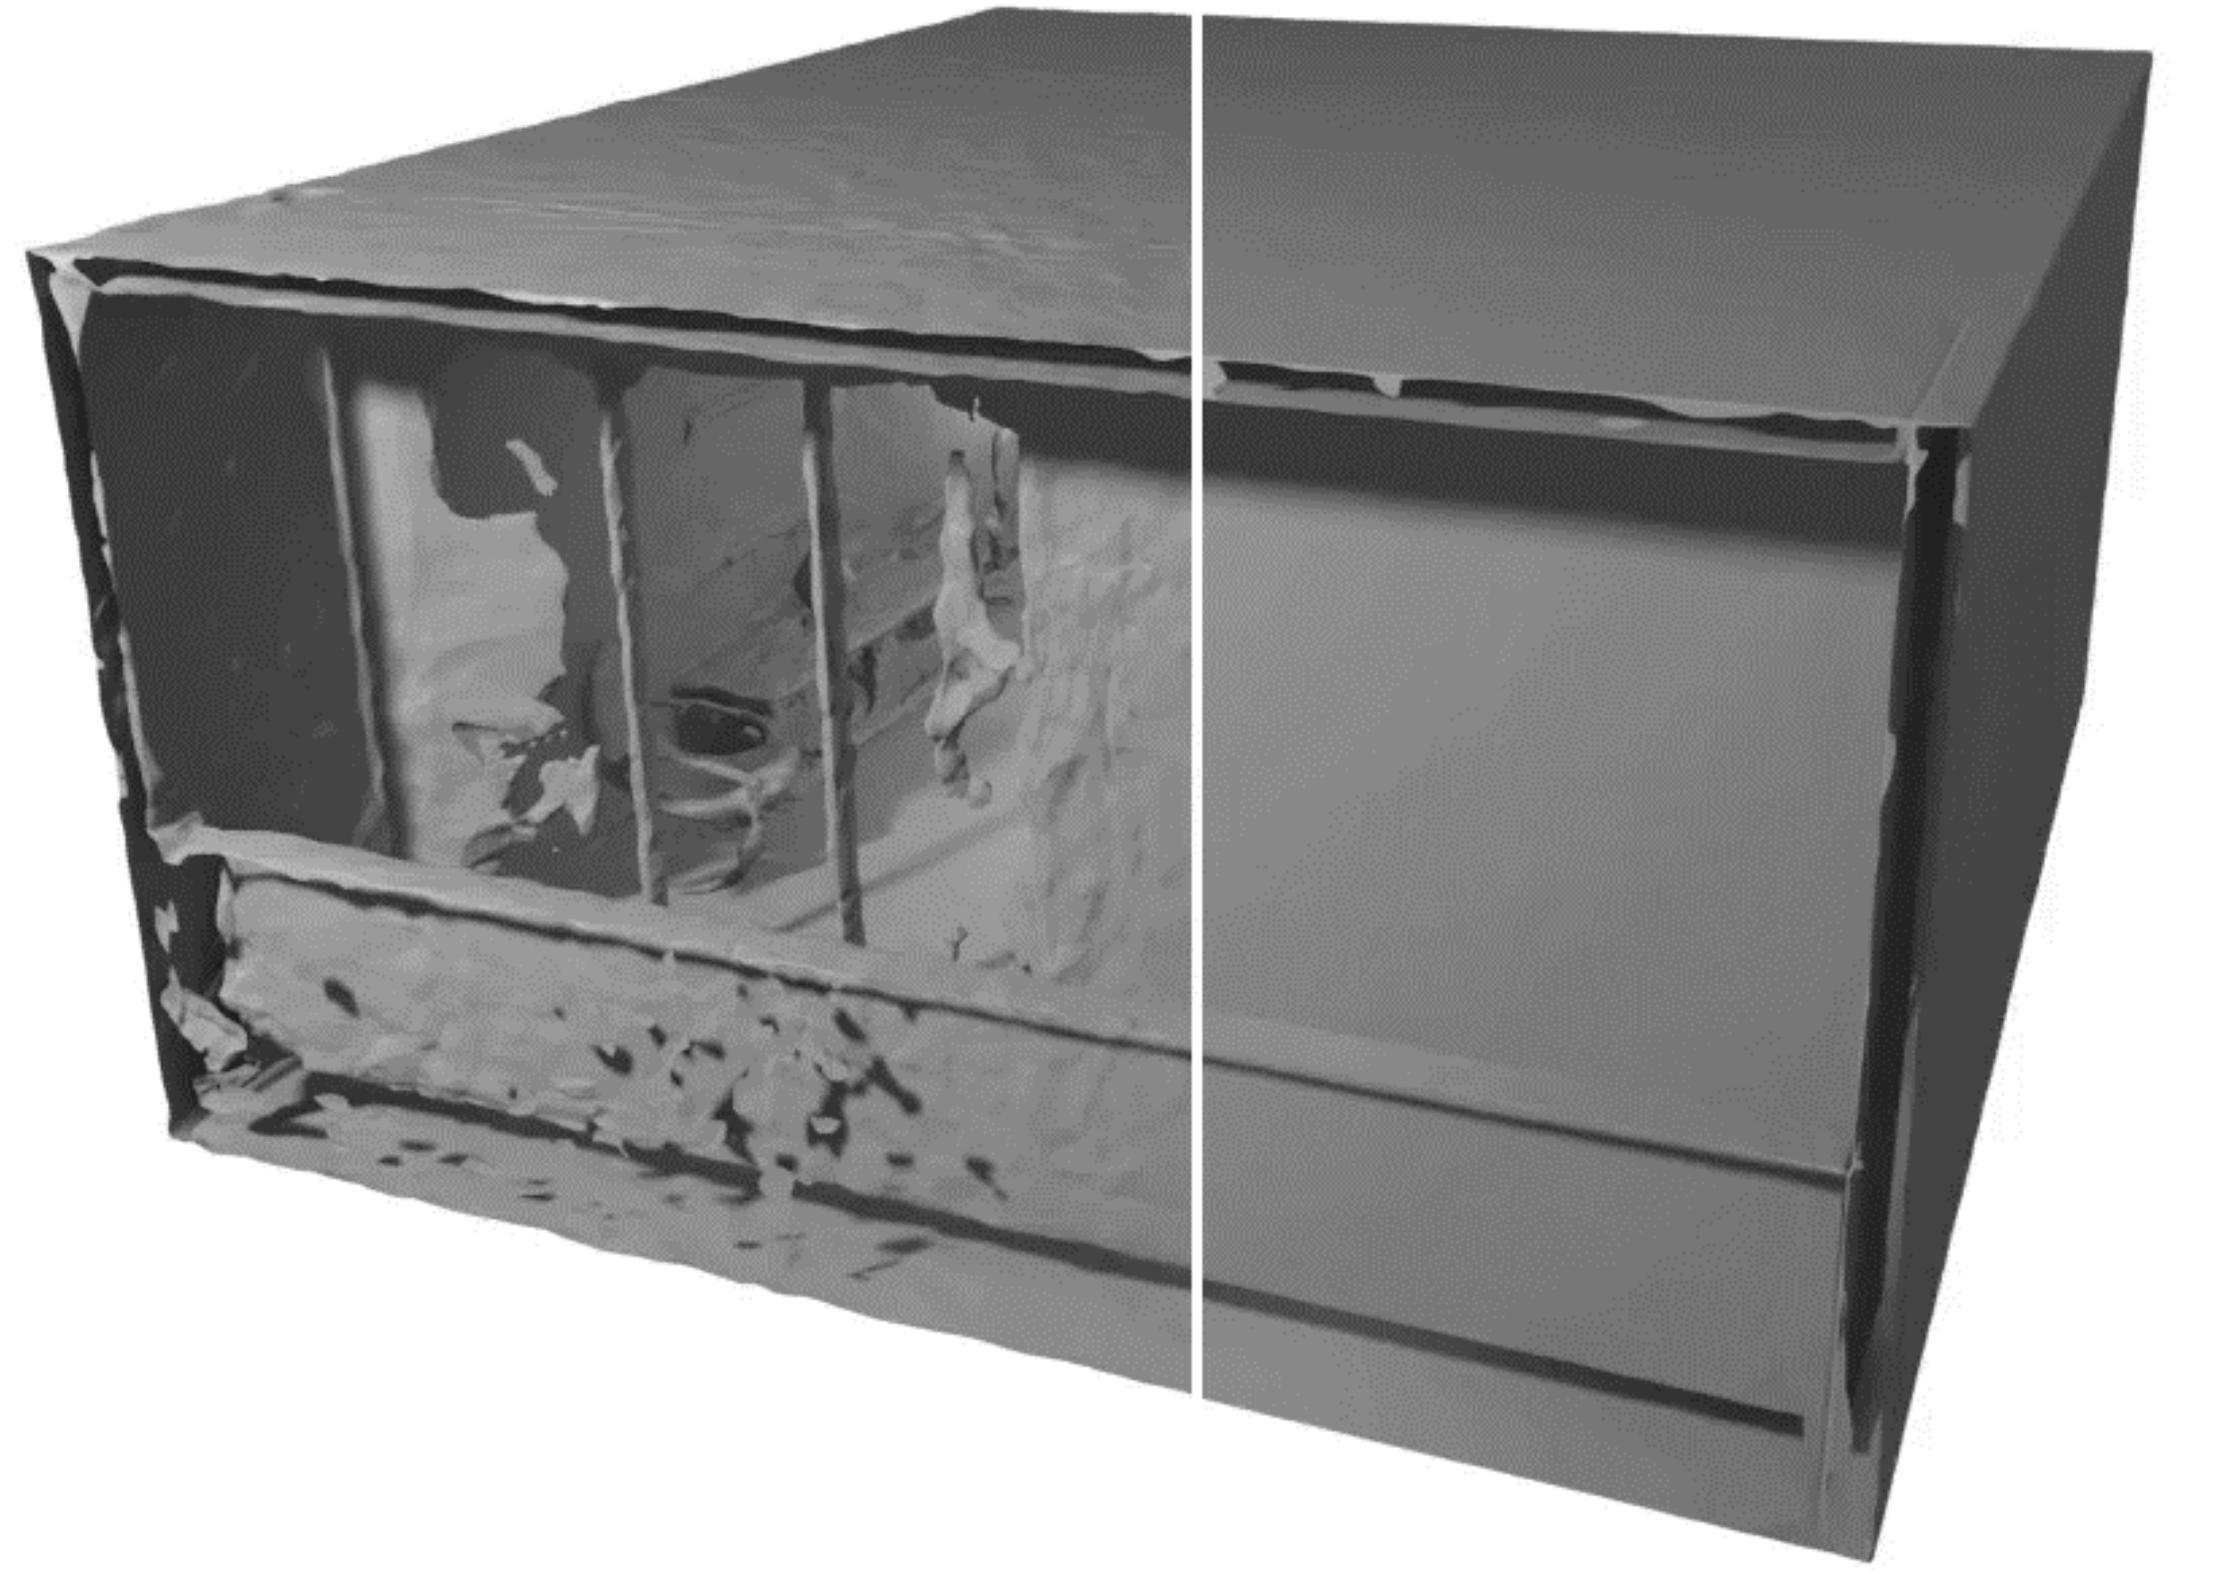
# Derivatives

$$\delta/\delta x$$



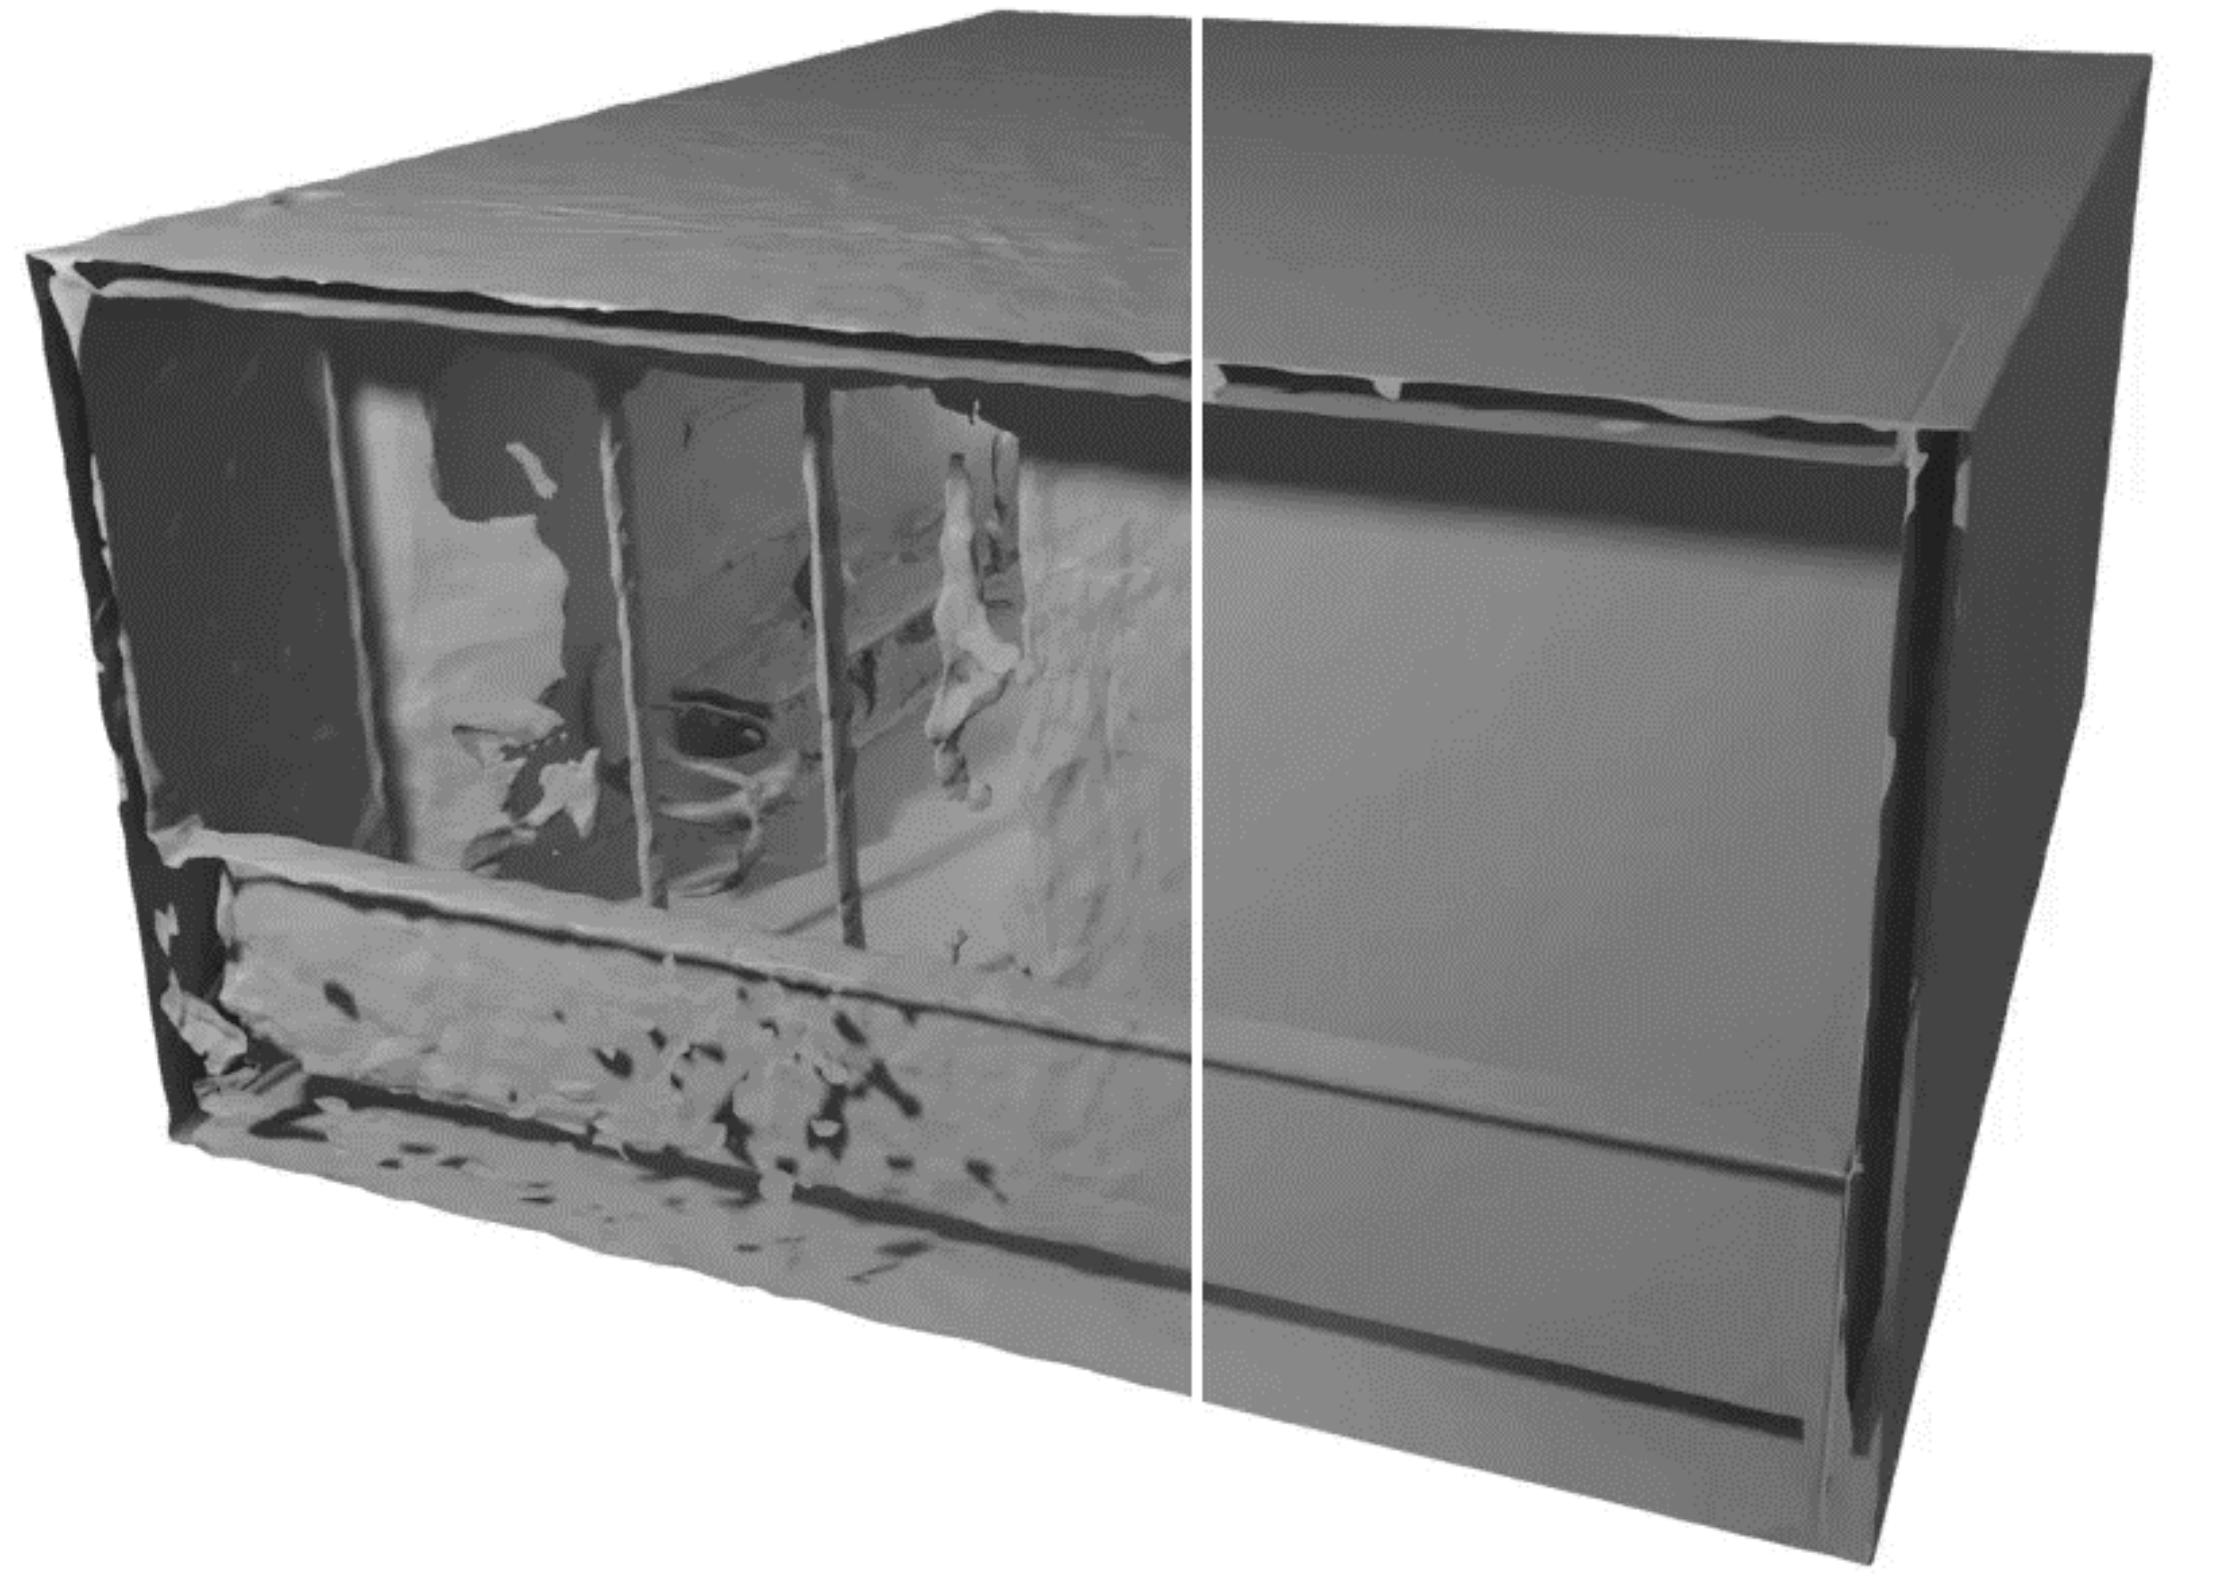
All derivatives exist, are nonzero and bounded by 1

ReLU

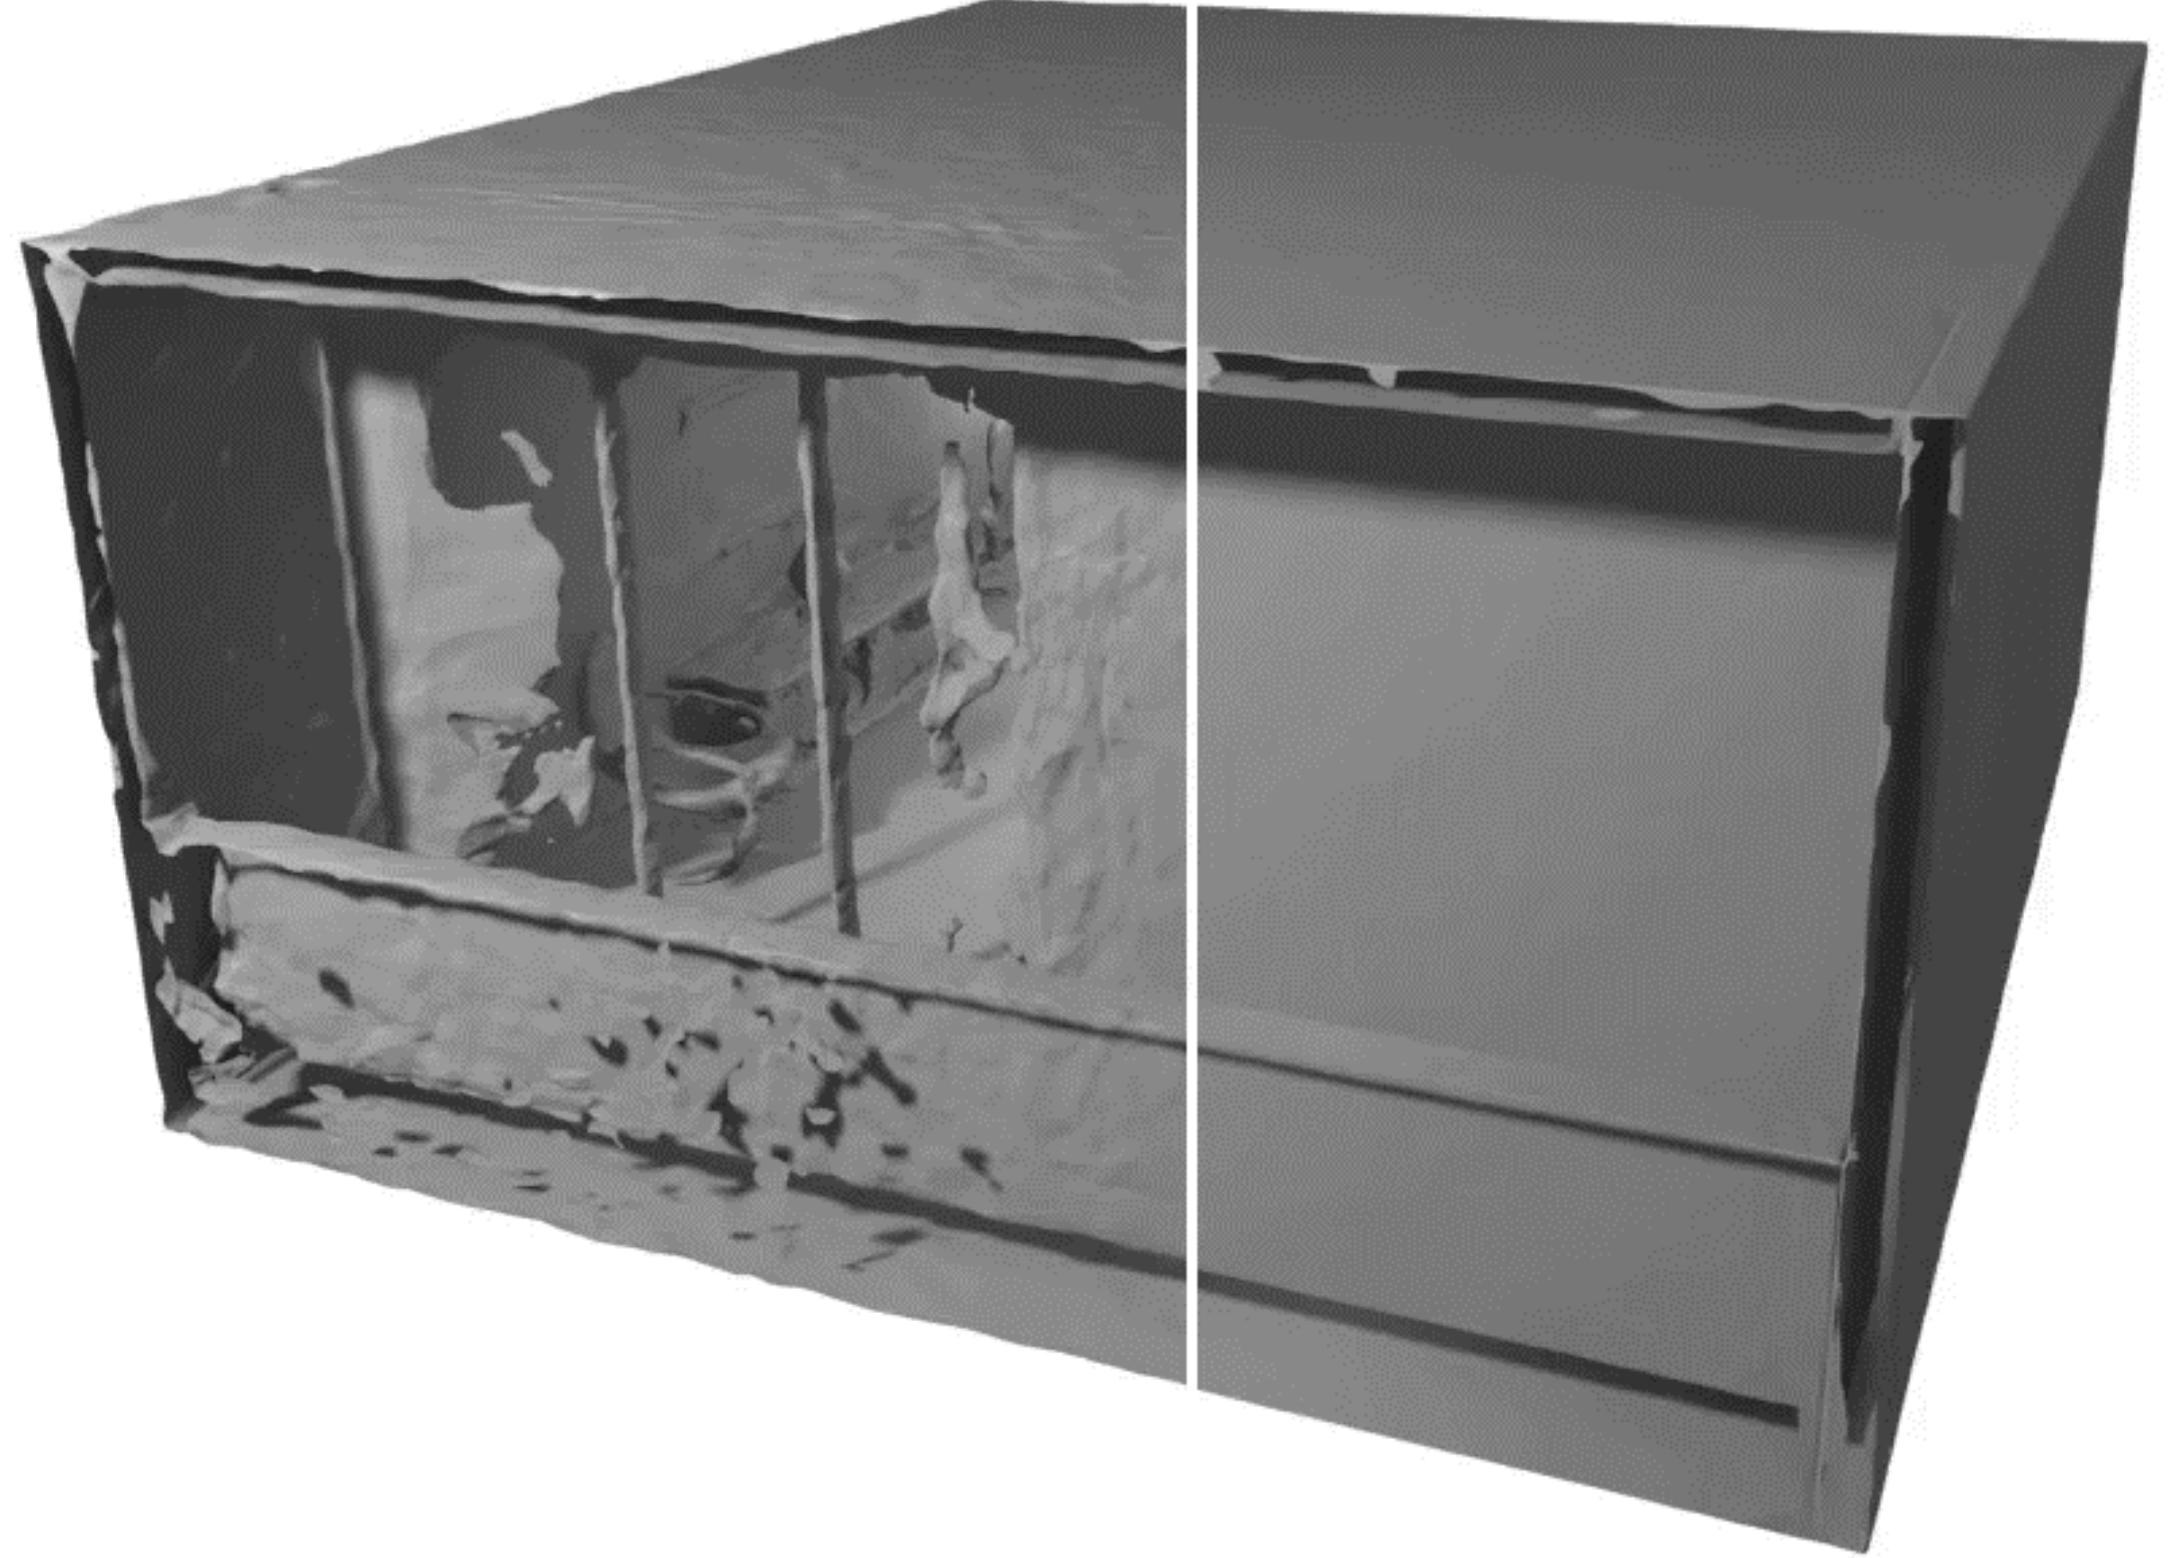


SIREN<sup>74</sup>

ReLU



SIREN<sup>74</sup>



~1MB compared to  
110MB of full mesh!

# Background: Kernel Regression

Given “training” set  $\{(x_i, y_i)\}_i$ , a kernel function makes predictions on a point  $\mathbf{x}$  by interpolating labels  $y_i$  in the training set according to pairwise weights between  $x_i$  to  $\mathbf{x}$  as measured by a kernel function:

$$f(\mathbf{x}) = \sum_{i=1}^n (\mathbf{K}^{-1} \mathbf{y})_i k(\mathbf{x}_i, \mathbf{x})$$

Where  $\mathbf{K}_{ij} = k(\mathbf{x}_i, \mathbf{x}_j)$ ,  $k$  is kernel function with  $k \geq 0$ .

# The Neural Tangent Kernel

$$f(\mathbf{x}) = \sum_{i=1}^n (\mathbf{K}^{-1}\mathbf{y})_i k(\mathbf{x}_i, \mathbf{x})$$

It turns out (Jacot et al. 2018) that in the limit of an infinitely wide MLP  $f(x)$  with initialization  $\theta_0$ , test-time predictions are made according to a kernel:

$$k_{NTK}(\mathbf{x}_1, \mathbf{x}_2) = \langle \nabla_{\theta} f_{\theta_0}(\mathbf{x}_1), \nabla_{\theta} f_{\theta_0}(\mathbf{x}_2) \rangle$$

# The Neural Tangent Kernel

$$f(\mathbf{x}) = \sum_{i=1}^n (\mathbf{K}^{-1}\mathbf{y})_i k(\mathbf{x}_i, \mathbf{x})$$

It turns out (Jacot et al. 2018) that in the limit of an infinitely wide MLP  $f(x)$  with initialization  $\theta_0$ , test-time predictions are made according to a kernel:

$$k_{NTK}(\mathbf{x}_1, \mathbf{x}_2) = \langle \nabla_{\theta} f_{\theta_0}(\mathbf{x}_1), \nabla_{\theta} f_{\theta_0}(\mathbf{x}_2) \rangle$$

In other words: “Neural” implicit representations interpolate the training set according to a relatively simple rule. No “intelligence”, no AI, no magic.

# The Neural Tangent Kernel

$$f(\mathbf{x}) = \sum_{i=1}^n (\mathbf{K}^{-1}\mathbf{y})_i k(\mathbf{x}_i, \mathbf{x})$$

It turns out (Jacot et al. 2018) that in the limit of an infinitely wide MLP  $f(x)$  with initialization  $\theta_0$ , test-time predictions are made according to a kernel:

$$k_{NTK}(\mathbf{x}_1, \mathbf{x}_2) = \langle \nabla_{\theta} f_{\theta_0}(\mathbf{x}_1), \nabla_{\theta} f_{\theta_0}(\mathbf{x}_2) \rangle$$

What we can do is design the kernel and pick the space of  $(\mathbf{x}, \mathbf{y})$ , the space in which our neural network will interpolate.

# The Neural Tangent Kernel

$$f(\mathbf{x}) = \sum_{i=1}^n (\mathbf{K}^{-1}\mathbf{y})_i k(\mathbf{x}_i, \mathbf{x})$$

It turns out (Jacot et al. 2018) that in the limit of an infinitely wide MLP  $f(x)$  with initialization  $\theta_0$ , test-time predictions are made according to a kernel:

$$k_{NTK}(\mathbf{x}_1, \mathbf{x}_2) = \langle \nabla_{\theta} f_{\theta_0}(\mathbf{x}_1), \nabla_{\theta} f_{\theta_0}(\mathbf{x}_2) \rangle$$

We do not know of an equivalent rule for transformers to date - work in progress!

# The Neural Tangent Kernel

$$f(\mathbf{x}) = \sum_{i=1}^n (\mathbf{K}^{-1}\mathbf{y})_i k(\mathbf{x}_i, \mathbf{x})$$

It turns out (Jacot et al. 2018) that in the limit of an infinitely wide MLP  $f(x)$  with initialization  $\theta_0$ , test-time predictions are made according to a kernel:

$$k_{NTK}(\mathbf{x}_1, \mathbf{x}_2) = \langle \nabla_{\theta} f_{\theta_0}(\mathbf{x}_1), \nabla_{\theta} f_{\theta_0}(\mathbf{x}_2) \rangle$$

What does this mean for neural fields?

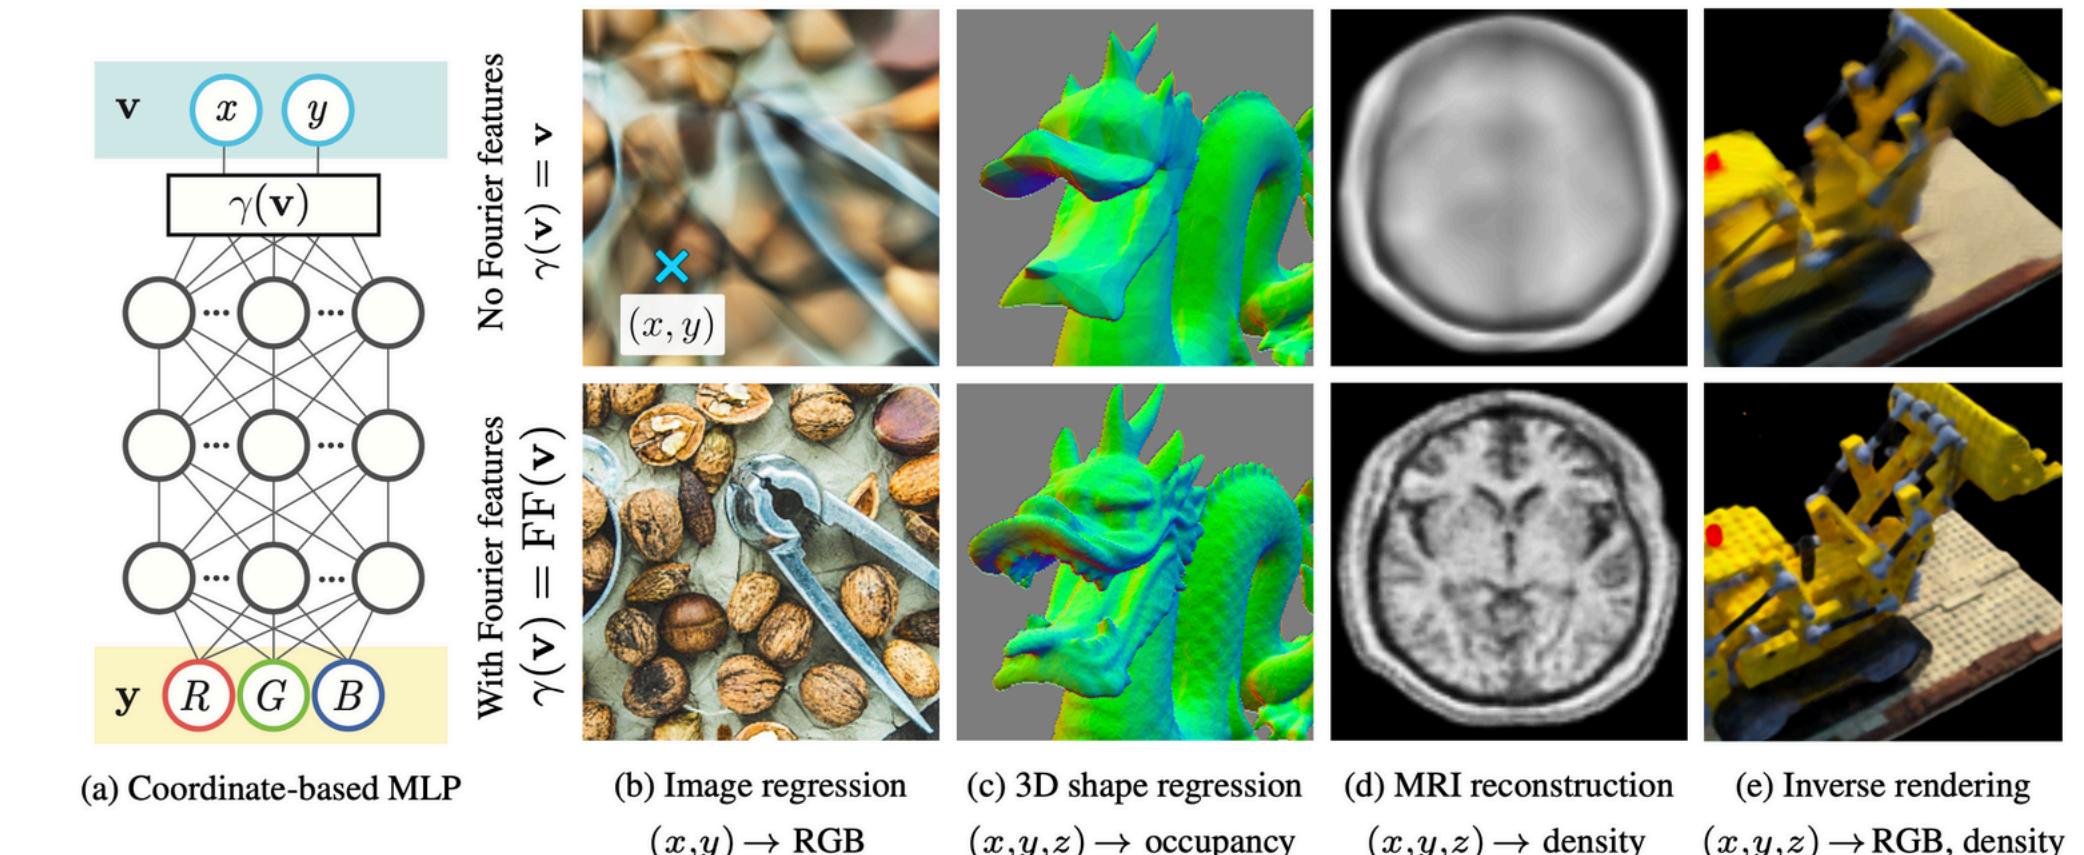
# SIREN: Neural Implicit Representations With Periodic Activation Functions

Sitzmann & Martel et al. NeurIPS 2020



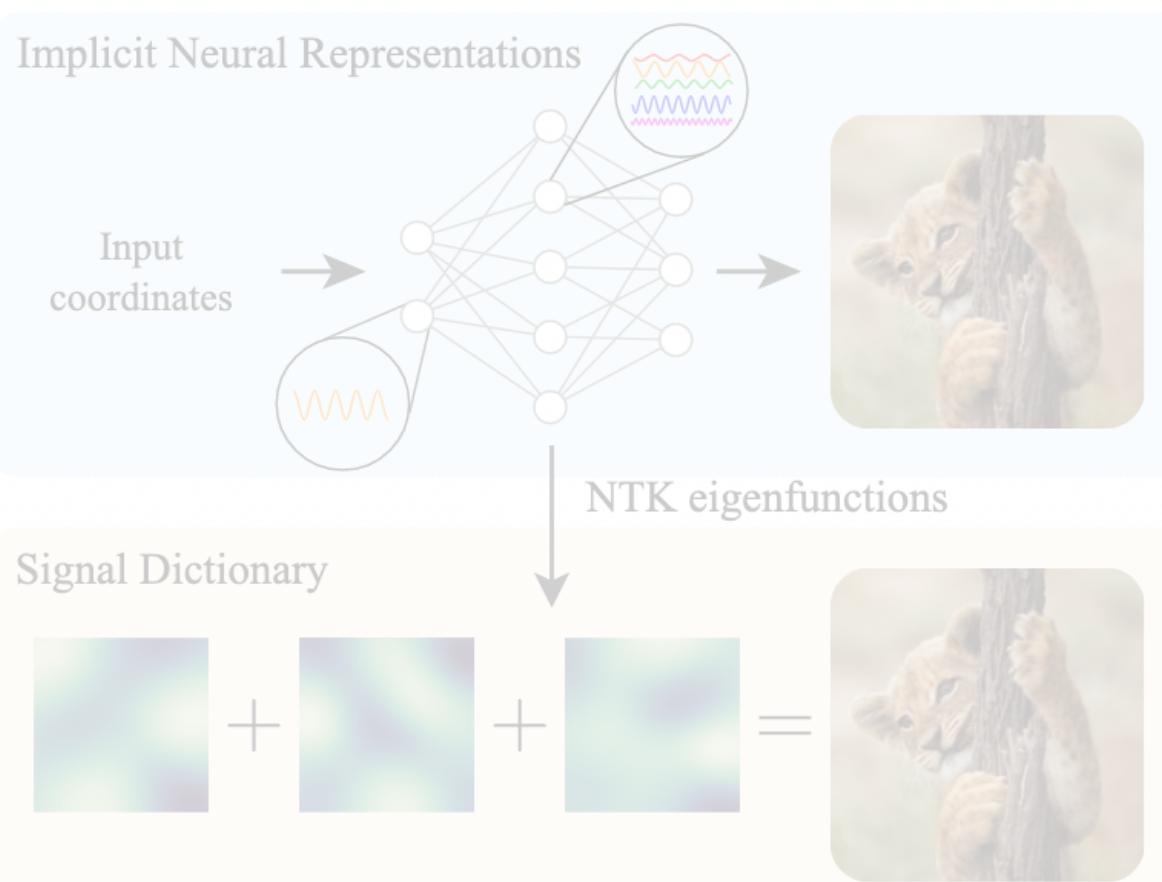
# Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains

Tancik, Srinivasan, Mildenhall et al, NeurIPS 2020



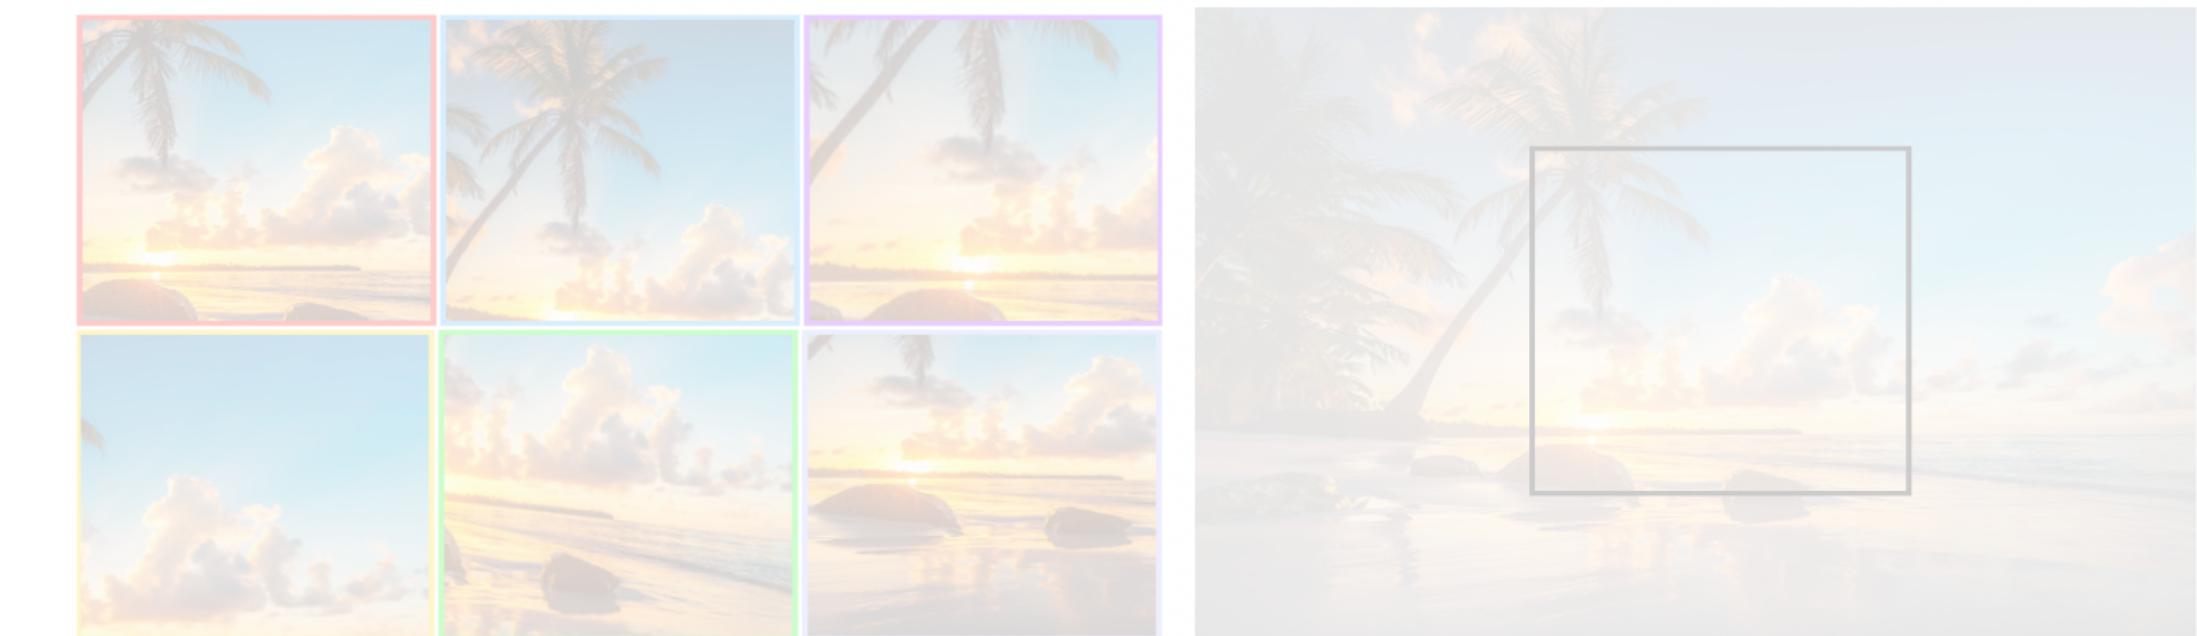
# A Structured Dictionary Perspective on Implicit Neural Representations

Yüce & Ortiz-Jiménez et al. CVPR 2022



# Gaussian Activated Neural Radiance Fields for High Fidelity Reconstruction & Pose Estimation

Chng et al. ECCV 2022



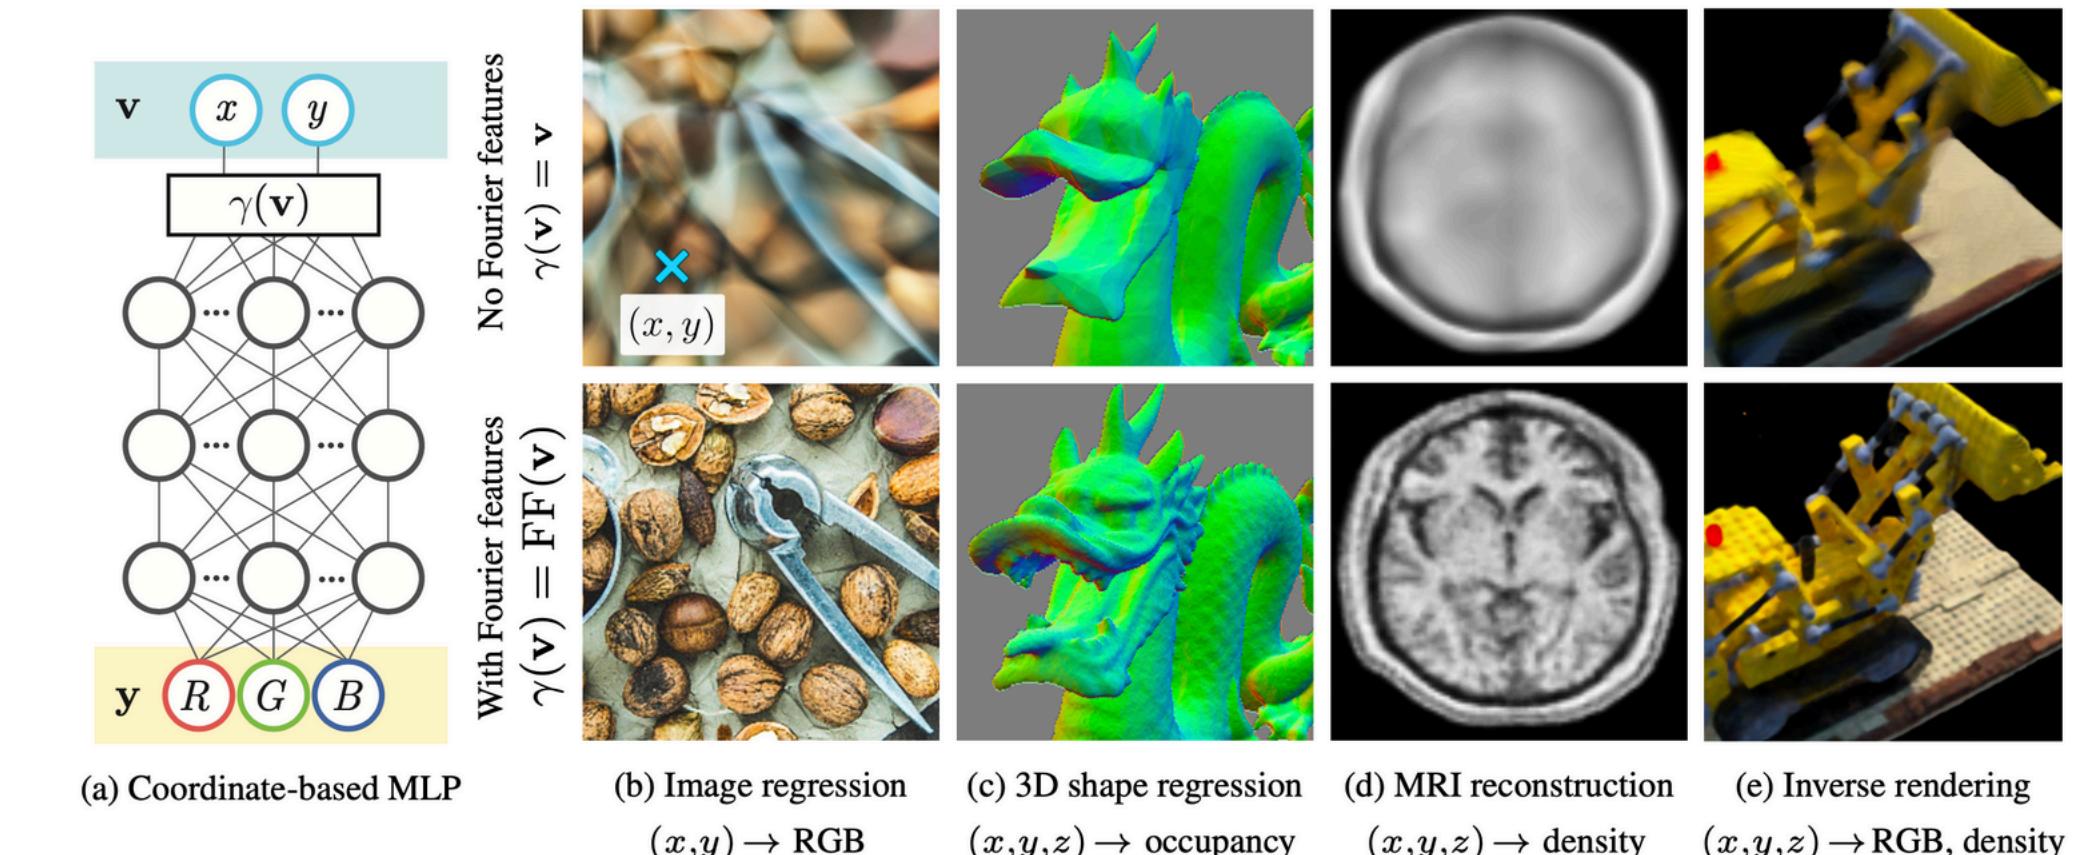
# SIREN: Neural Implicit Representations With Periodic Activation Functions

Sitzmann & Martel et al. NeurIPS 2020



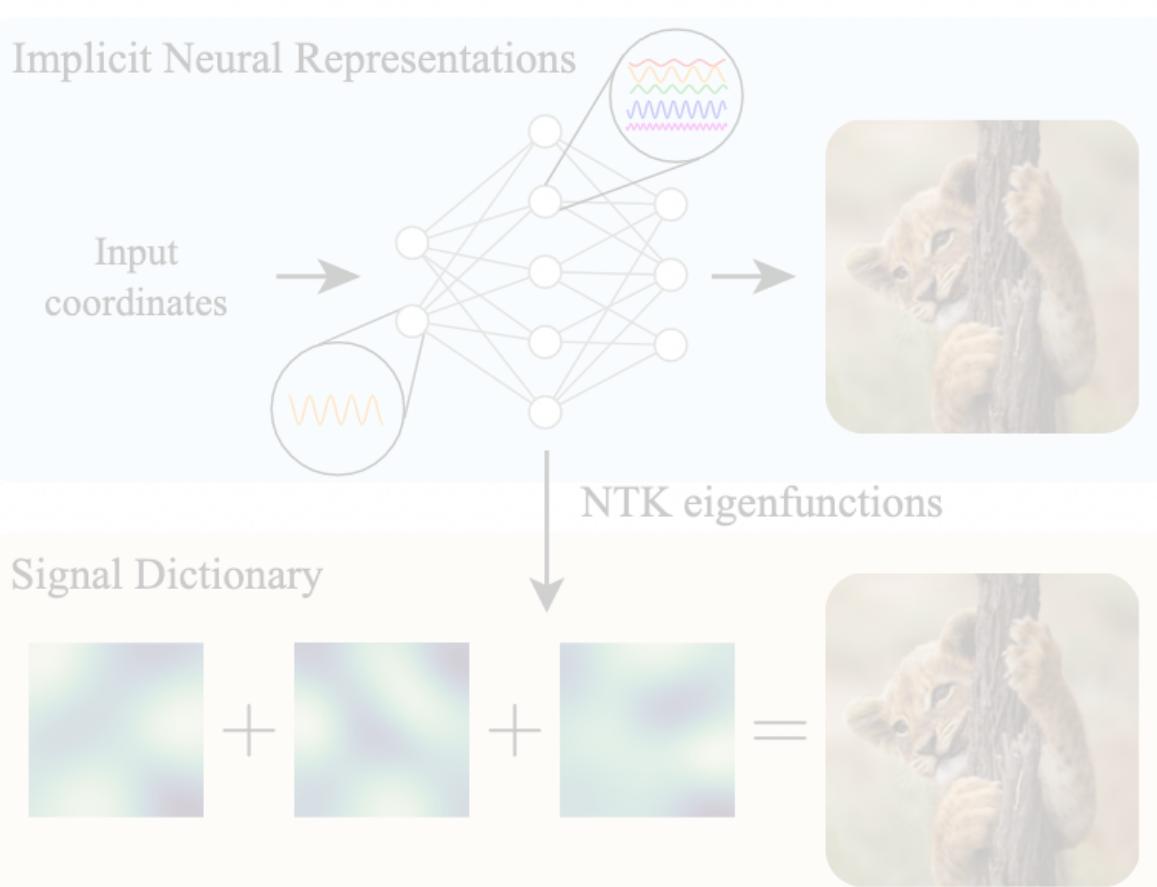
# Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains

Tancik, Srinivasan, Mildenhall et al, NeurIPS 2020



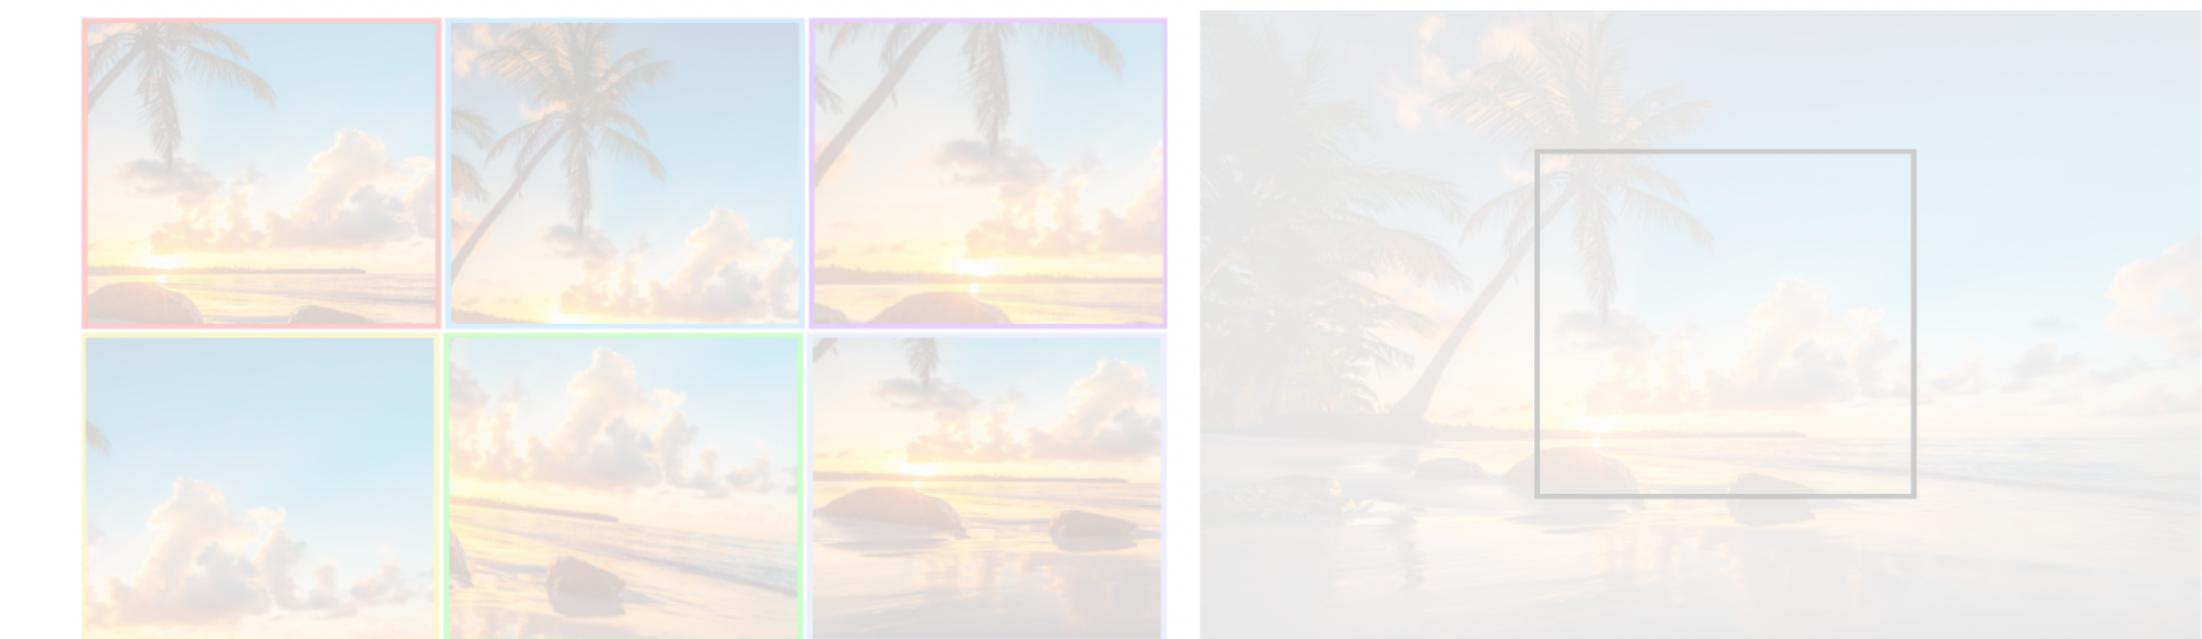
# A Structured Dictionary Perspective on Implicit Neural Representations

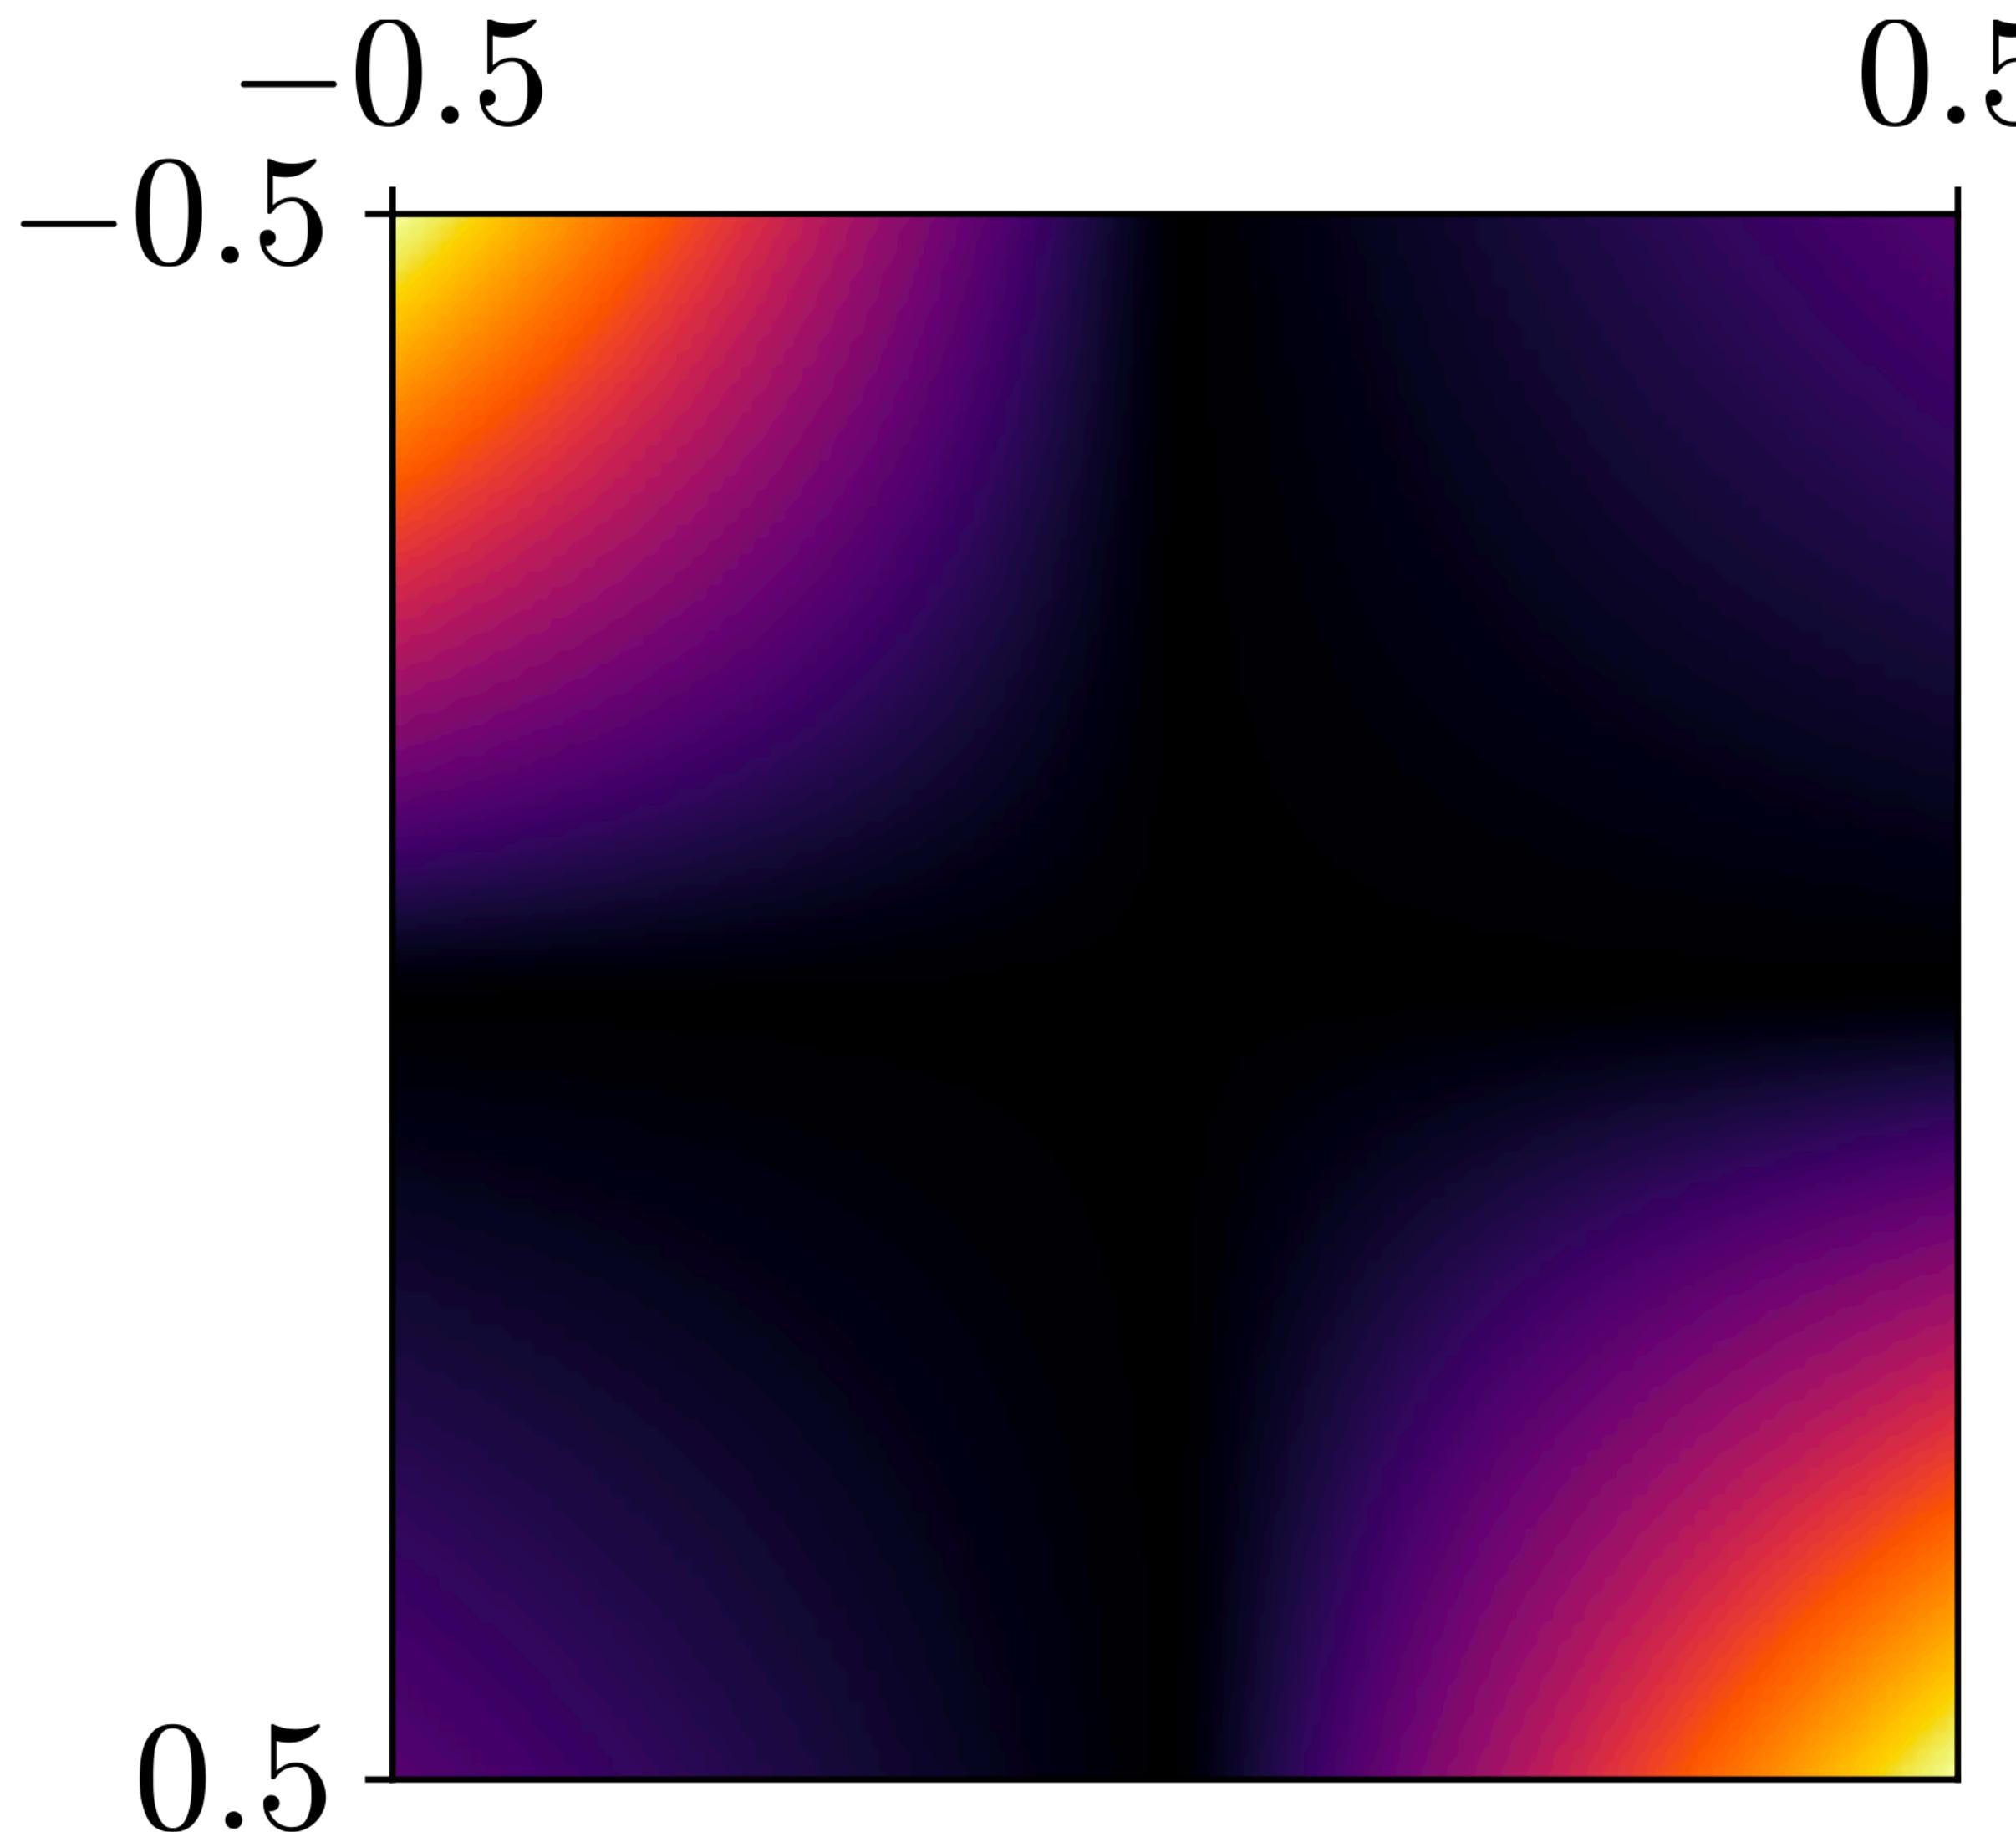
Yüce & Ortiz-Jiménez et al. CVPR 2022



# Gaussian Activated Neural Radiance Fields for High Fidelity Reconstruction & Pose Estimation

Chng et al. ECCV 2022



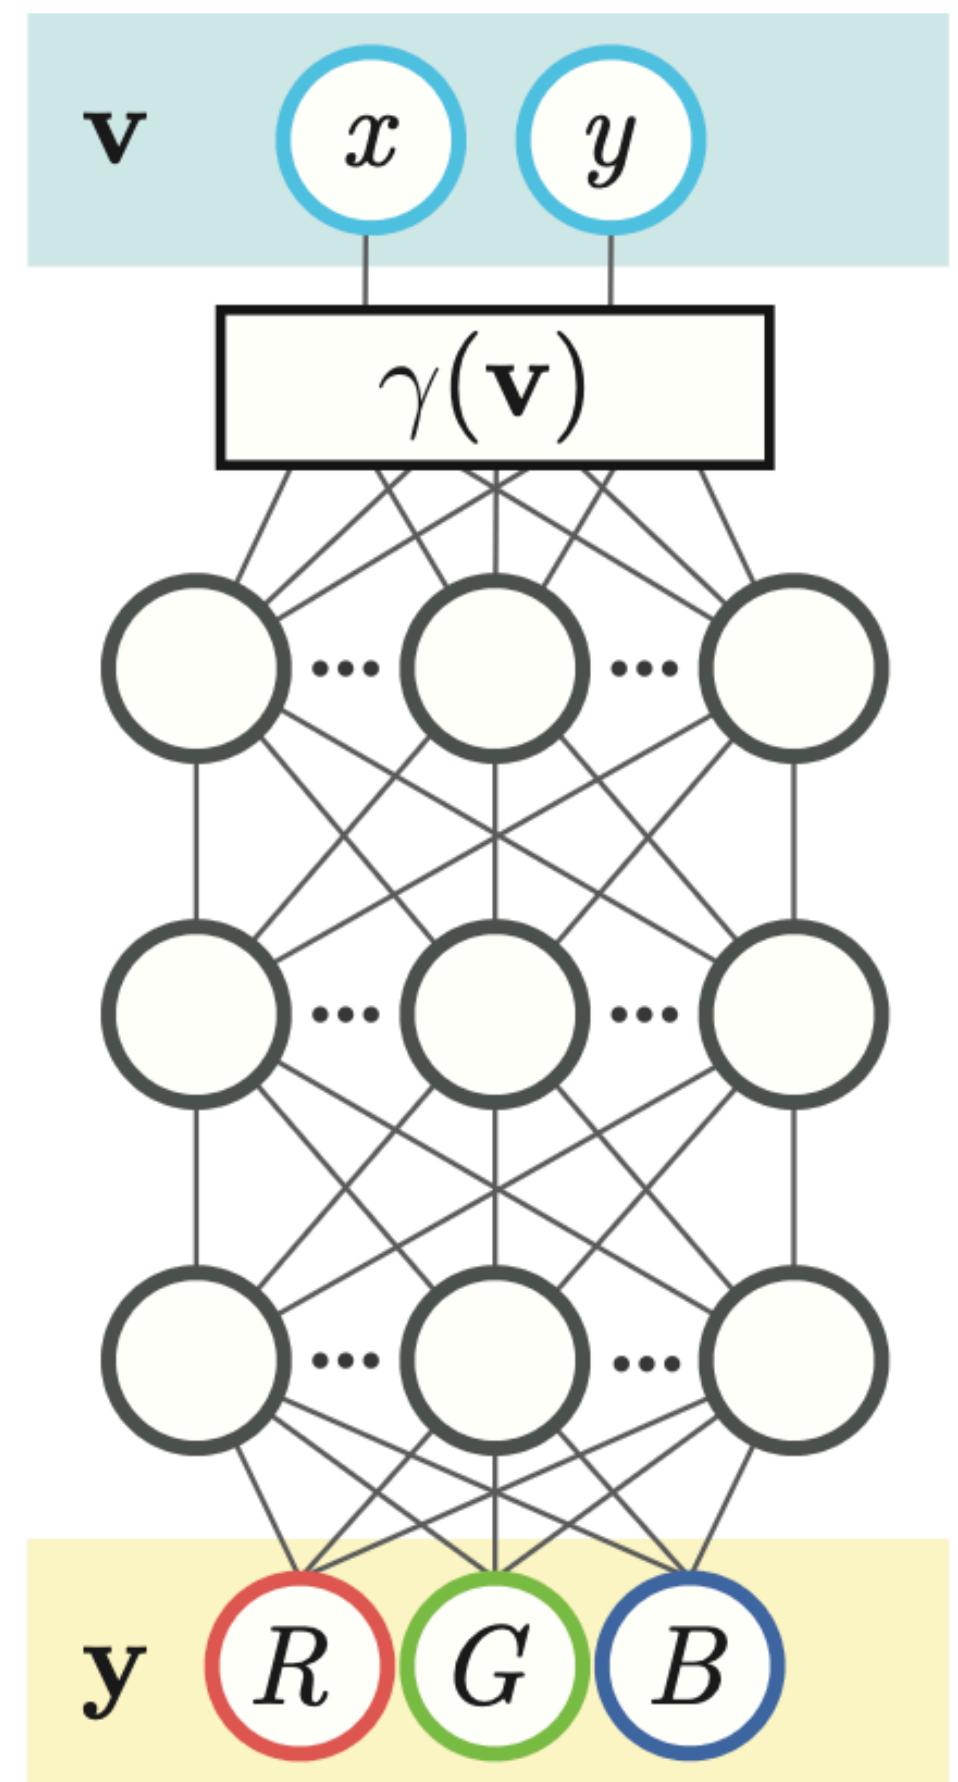


NTK kernel  $k_{NTK}(x_i, x_j)$  for a 4-layer  
ReLU MLP with one scalar input.

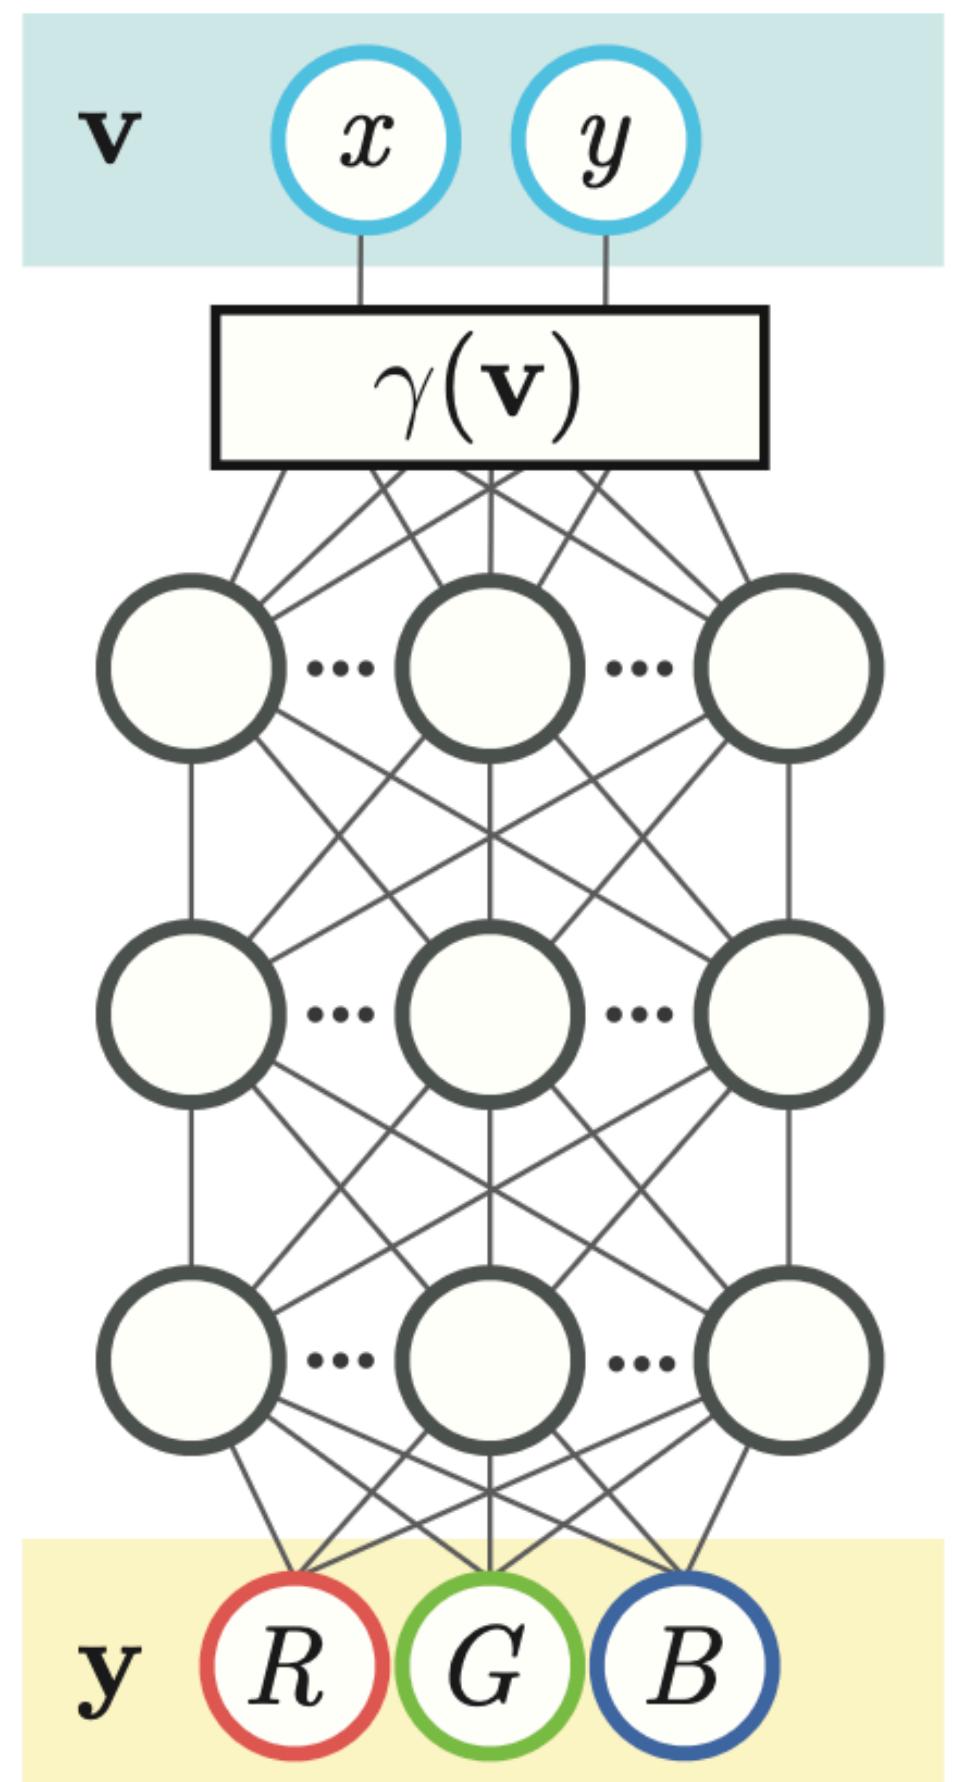
From “Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains”, Tancik, Srinivasan, Mildenhall et al.

NTK of ReLU neural net is quite non-local.

# Feature Embeddings

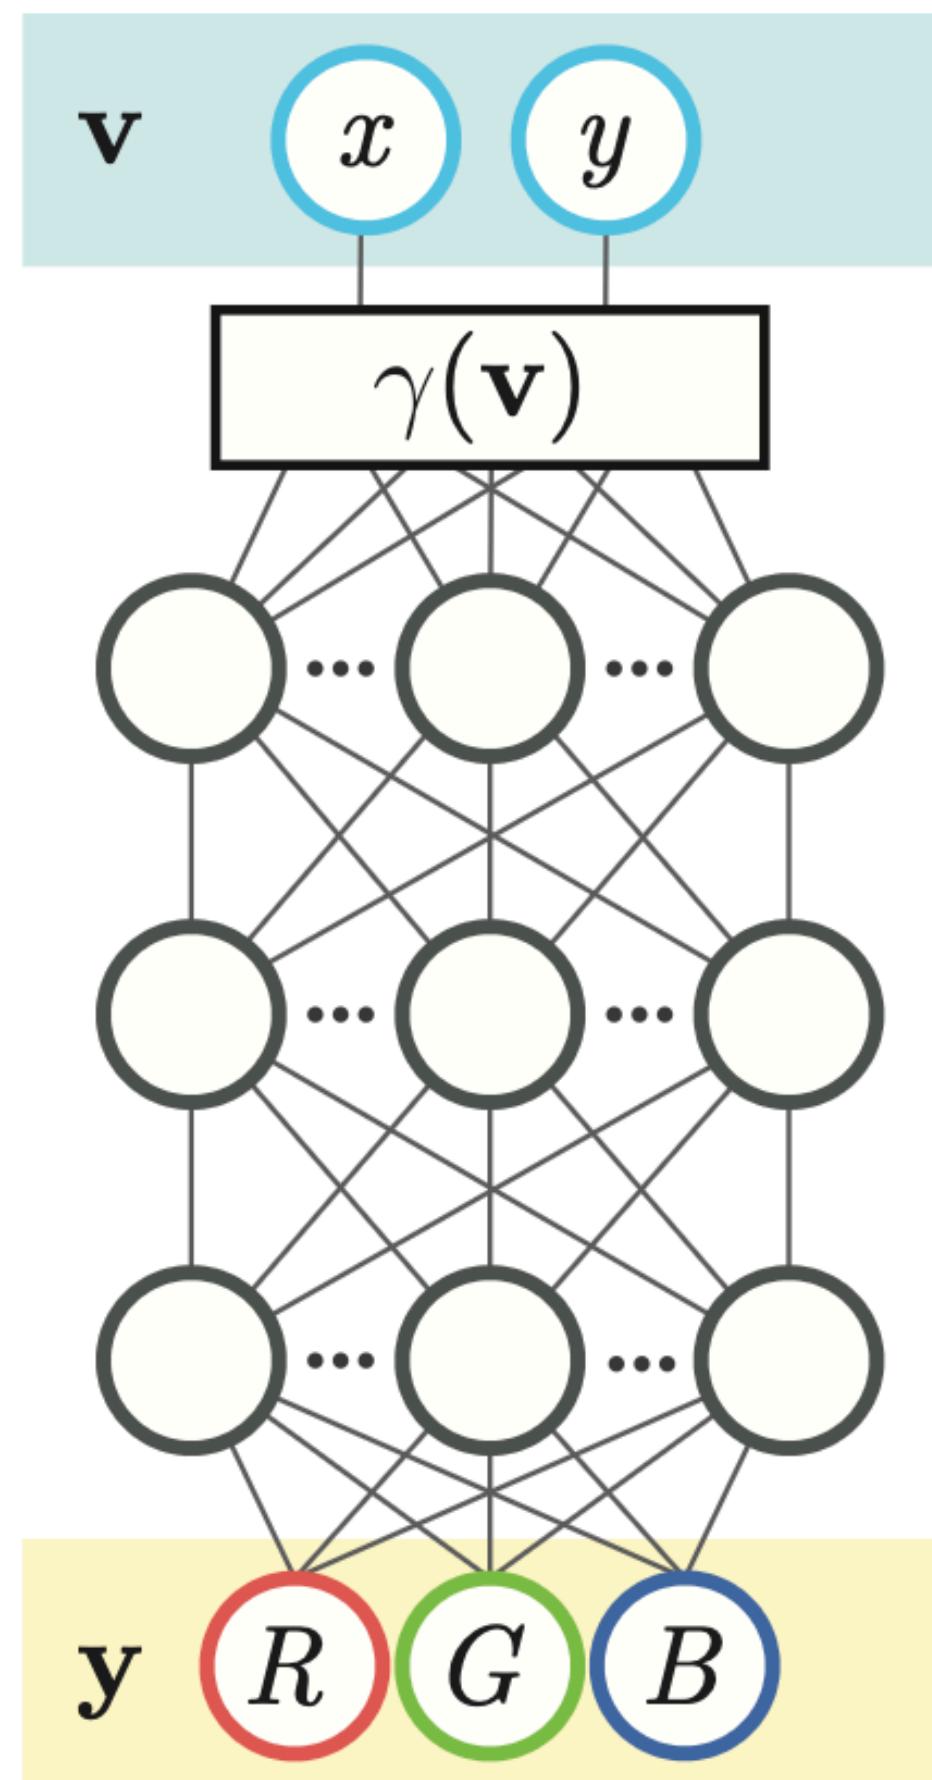


# Feature Embeddings



$$\gamma(\mathbf{x}) = [\gamma_1(\mathbf{x}), \dots, \gamma_n(\mathbf{x})]$$

# Feature Embeddings



$$\gamma(\mathbf{x}) = [\gamma_1(\mathbf{x}), \dots, \gamma_n(\mathbf{x})]$$

$$\gamma_{2i}(\mathbf{x}) = \sin(2^{i-1}\pi x_i)$$

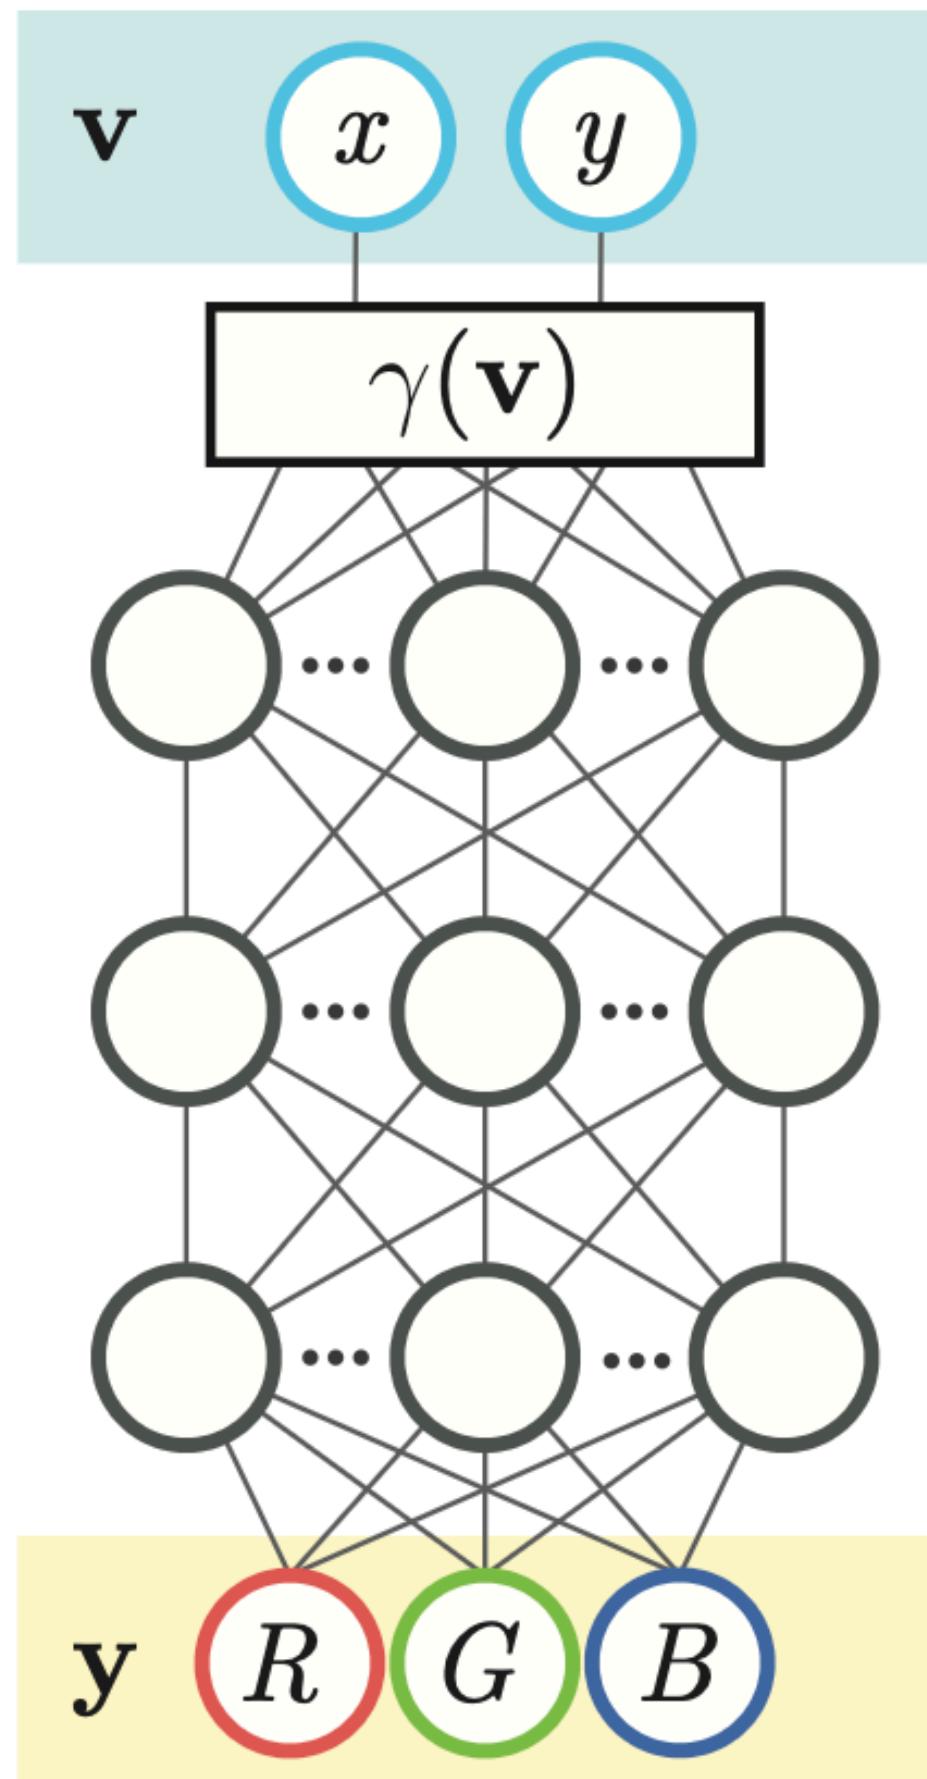
$$\gamma_{(2i+1)}(\mathbf{x}) = \cos(2^{i-1}\pi x_i)$$

**Sinusoidal Embeddings**

Zhong et al. ICLR 2020

Mildenhall et al., ECCV 2020

# Feature Embeddings



$$\gamma(\mathbf{x}) = [\gamma_1(\mathbf{x}), \dots, \gamma_n(\mathbf{x})]$$

$$\gamma_{2i}(\mathbf{x}) = \sin(2^{i-1}\pi x_i)$$

$$\gamma_{(2i+1)}(\mathbf{x}) = \cos(2^{i-1}\pi x_i)$$

**Sinusoidal Embeddings**

Zhong et al. ICLR 2020

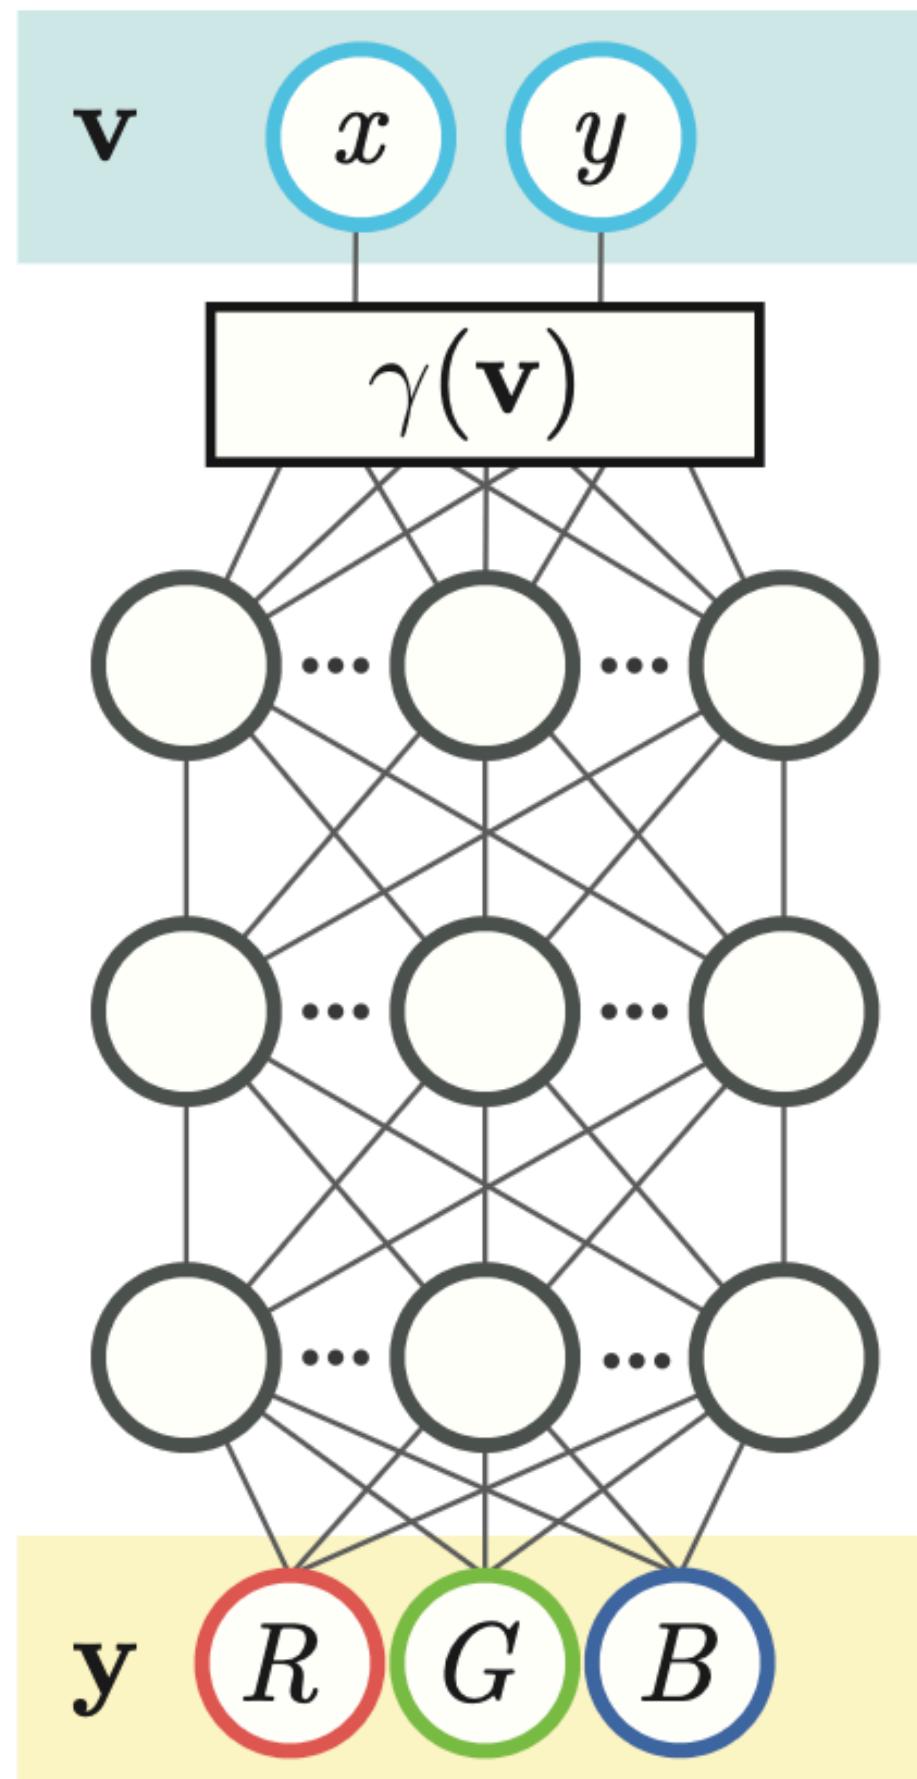
Mildenhall et al., ECCV 2020

$$\gamma(\mathbf{x}) = \exp\left(-\frac{\|t - x\|^2}{2\sigma^2}\right)$$

**Gaussian Embeddings**

Zheng et al., arXiv 2021

# Feature Embeddings



$$\gamma(\mathbf{x}) = [\gamma_1(\mathbf{x}), \dots, \gamma_n(\mathbf{x})]$$

$$\gamma_{2i}(\mathbf{x}) = \sin(2^{i-1}\pi x_i)$$

$$\gamma_{(2i+1)}(\mathbf{x}) = \cos(2^{i-1}\pi x_i)$$

**Sinusoidal Embeddings**

Zhong et al. ICLR 2020

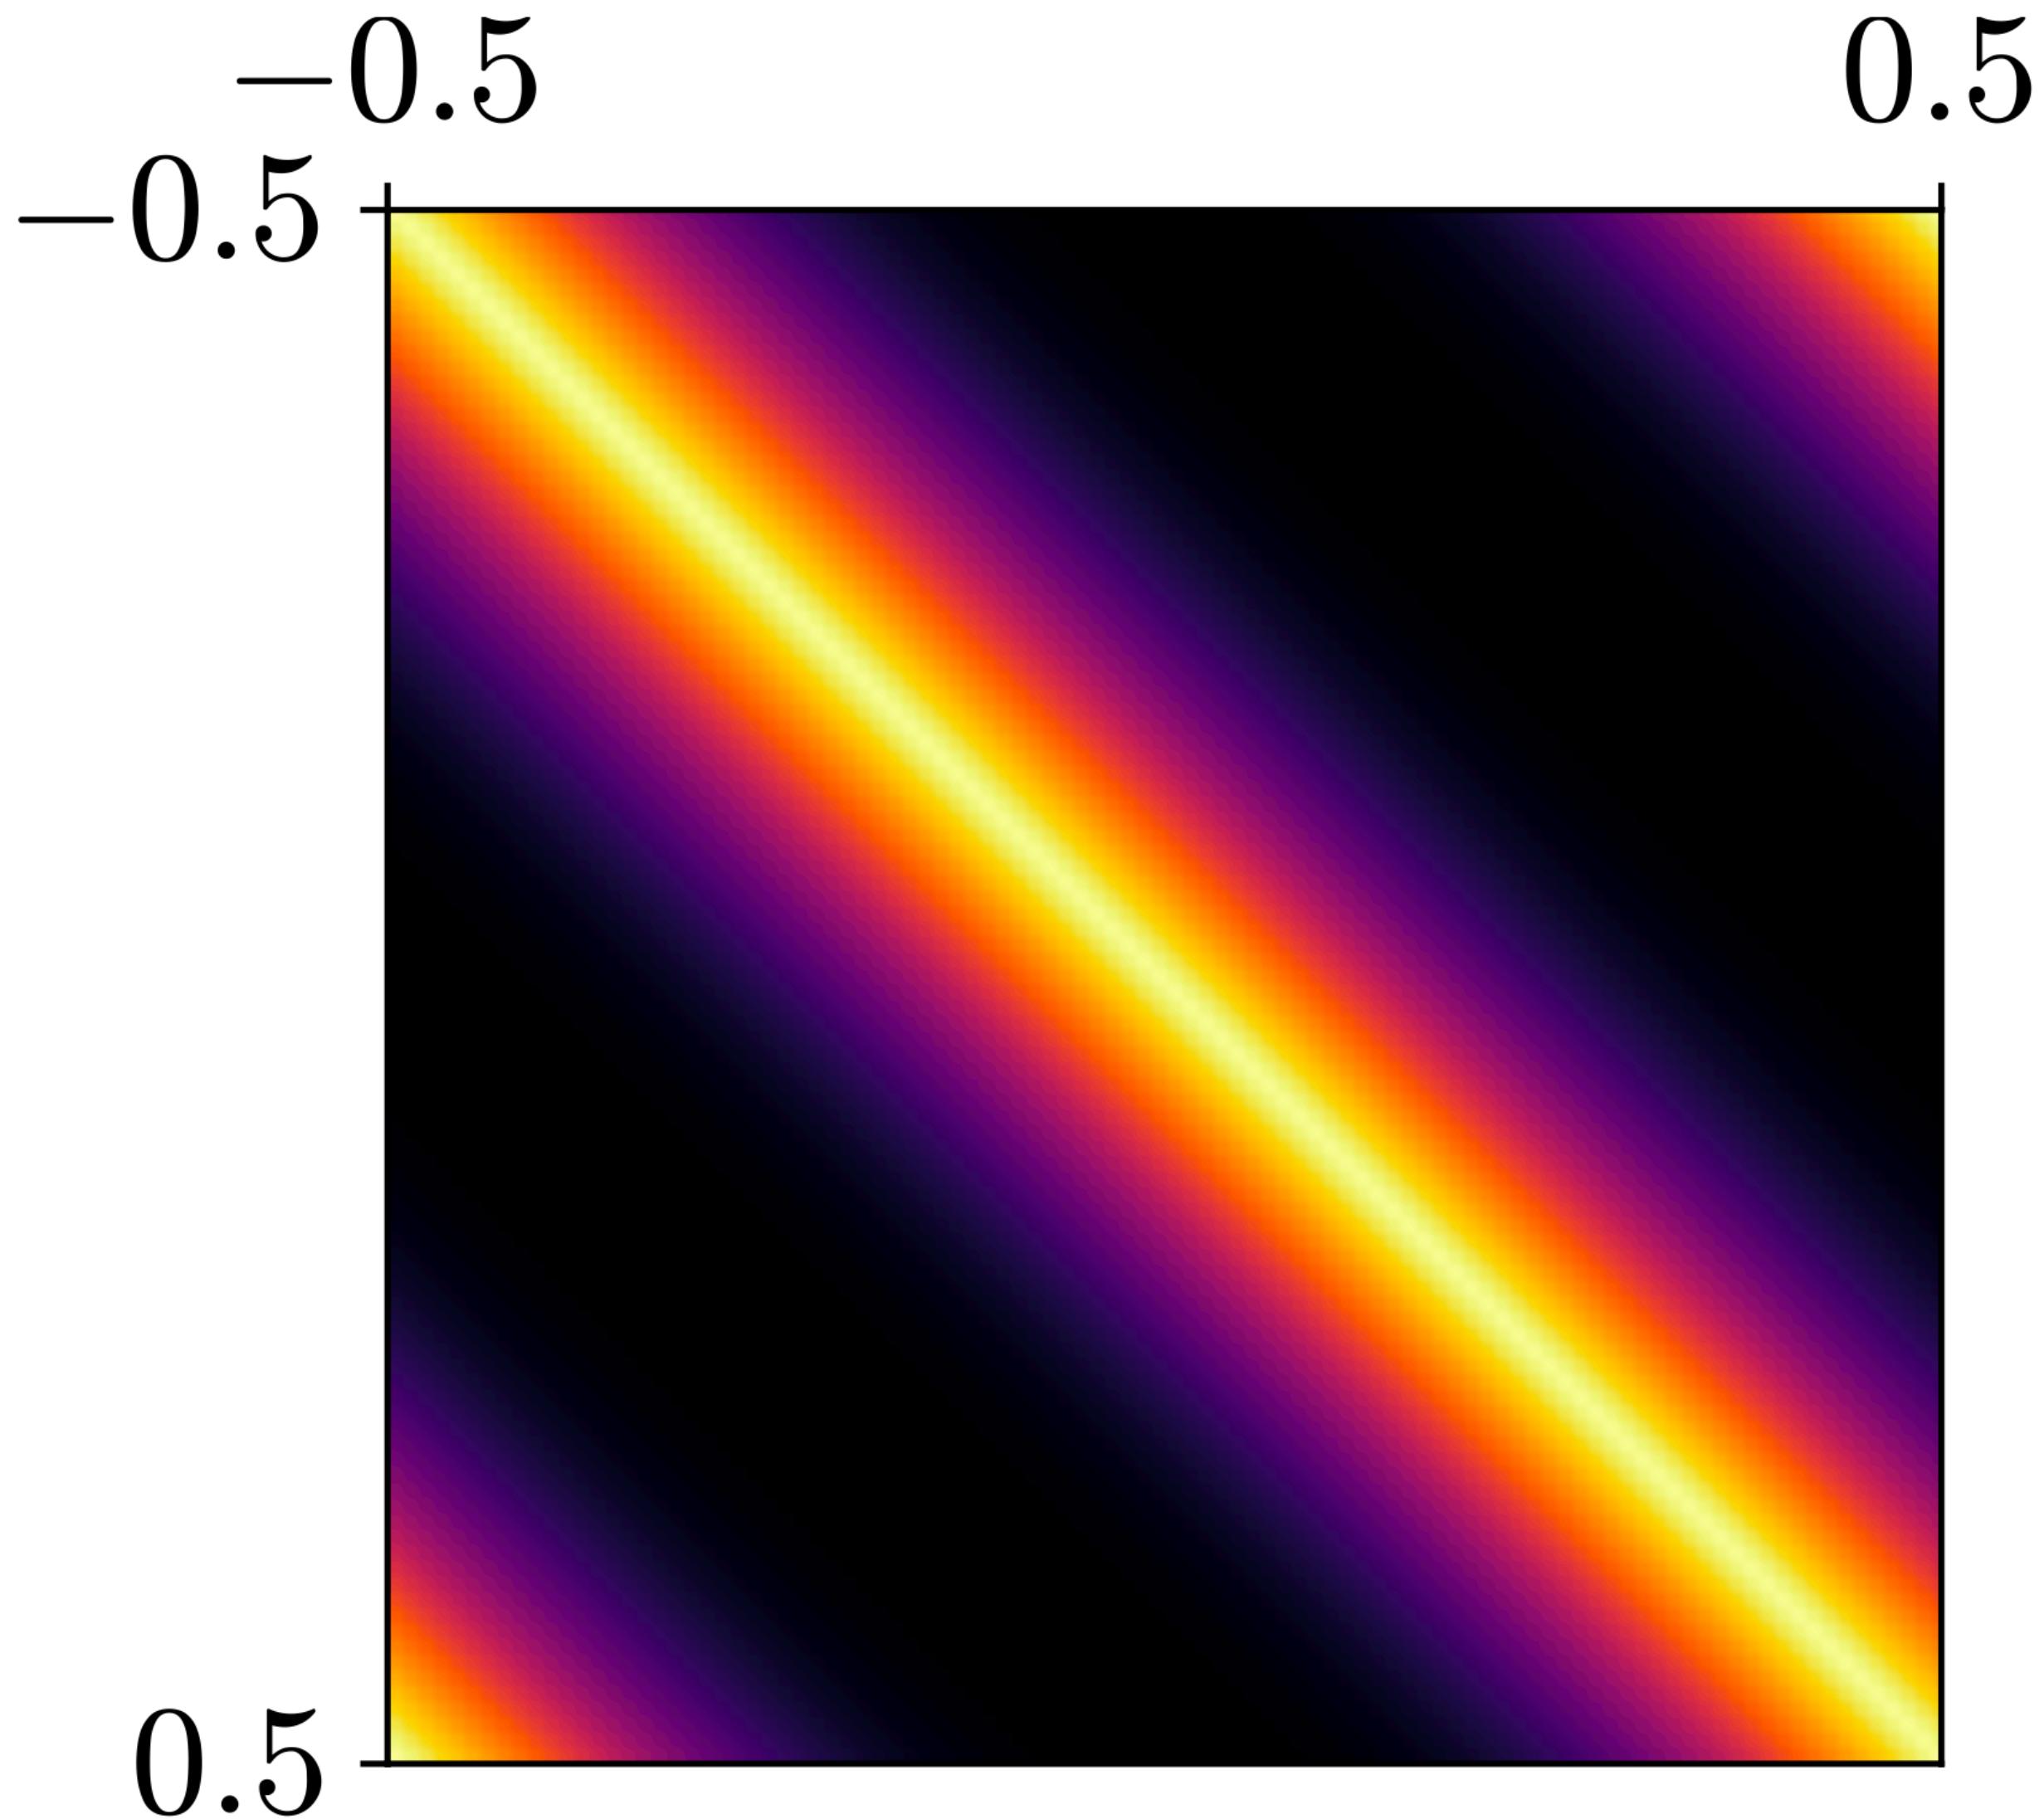
Mildenhall et al., ECCV 2020

$$\gamma(\mathbf{x}) = \exp\left(-\frac{\|t - x\|^2}{2\sigma^2}\right)$$

**Gaussian Embeddings**

Zheng et al., arXiv 2021

• • •



NTK kernel  $k_{NTK}(\gamma(x_i), \gamma(x_j))$  for a 4-layer ReLU MLP with one scalar input and **Sinusoidal Feature Mapping**

$$\gamma(x) = [\cos(2\pi x), \sin(2\pi x)].$$

From “Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains”, Tancik, Srinivasan, Mildenhall et al.

NTK is local (note diagonal).

# SIREN



ReLU

SIREN  
85

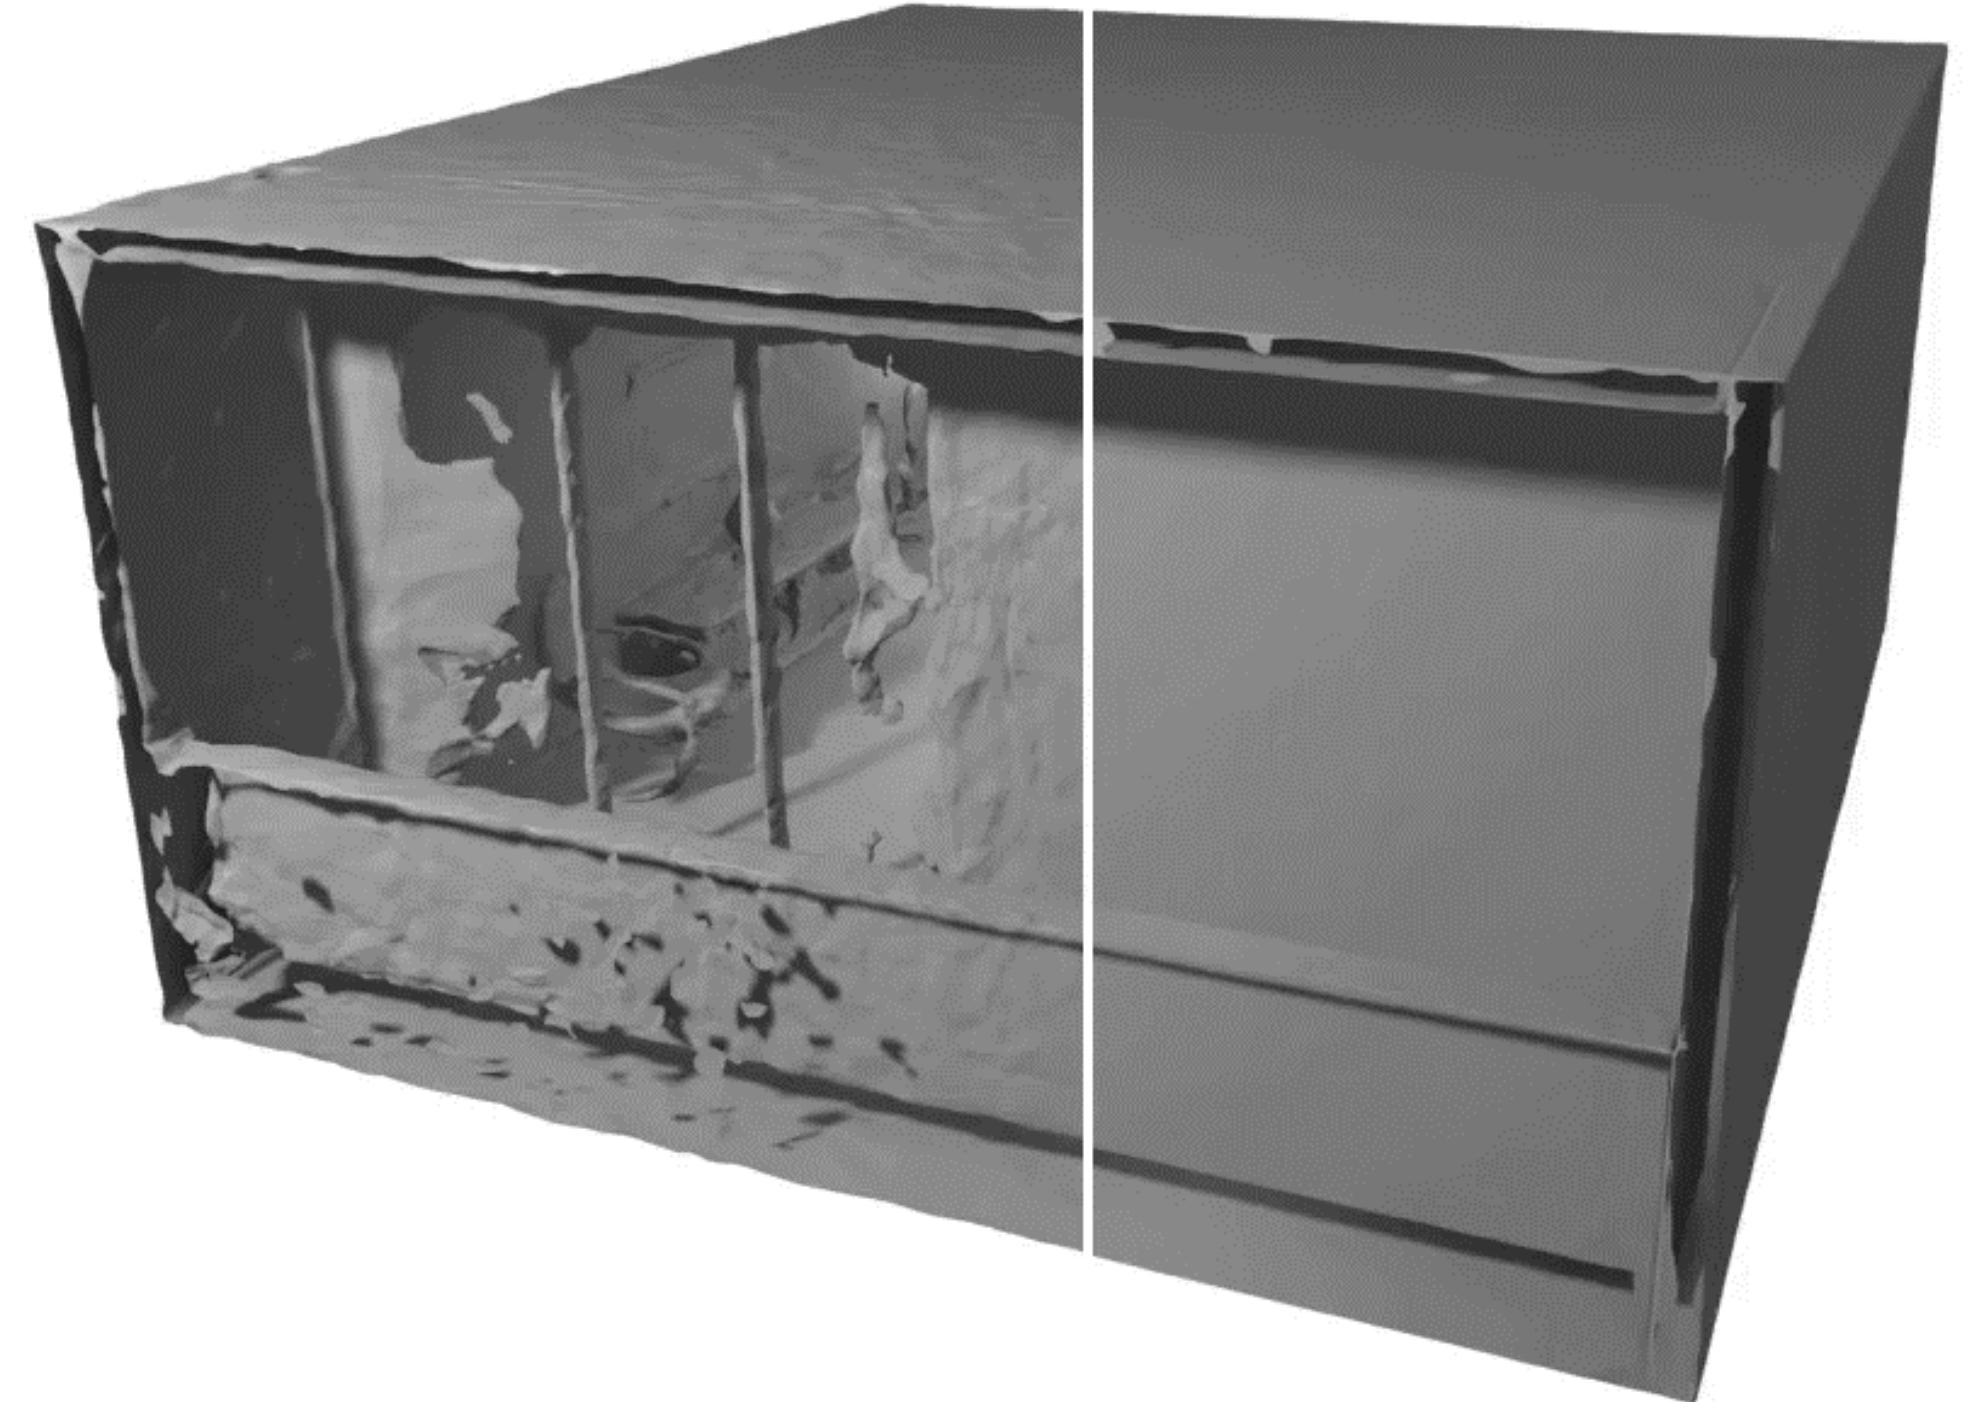
# SIREN



ReLU

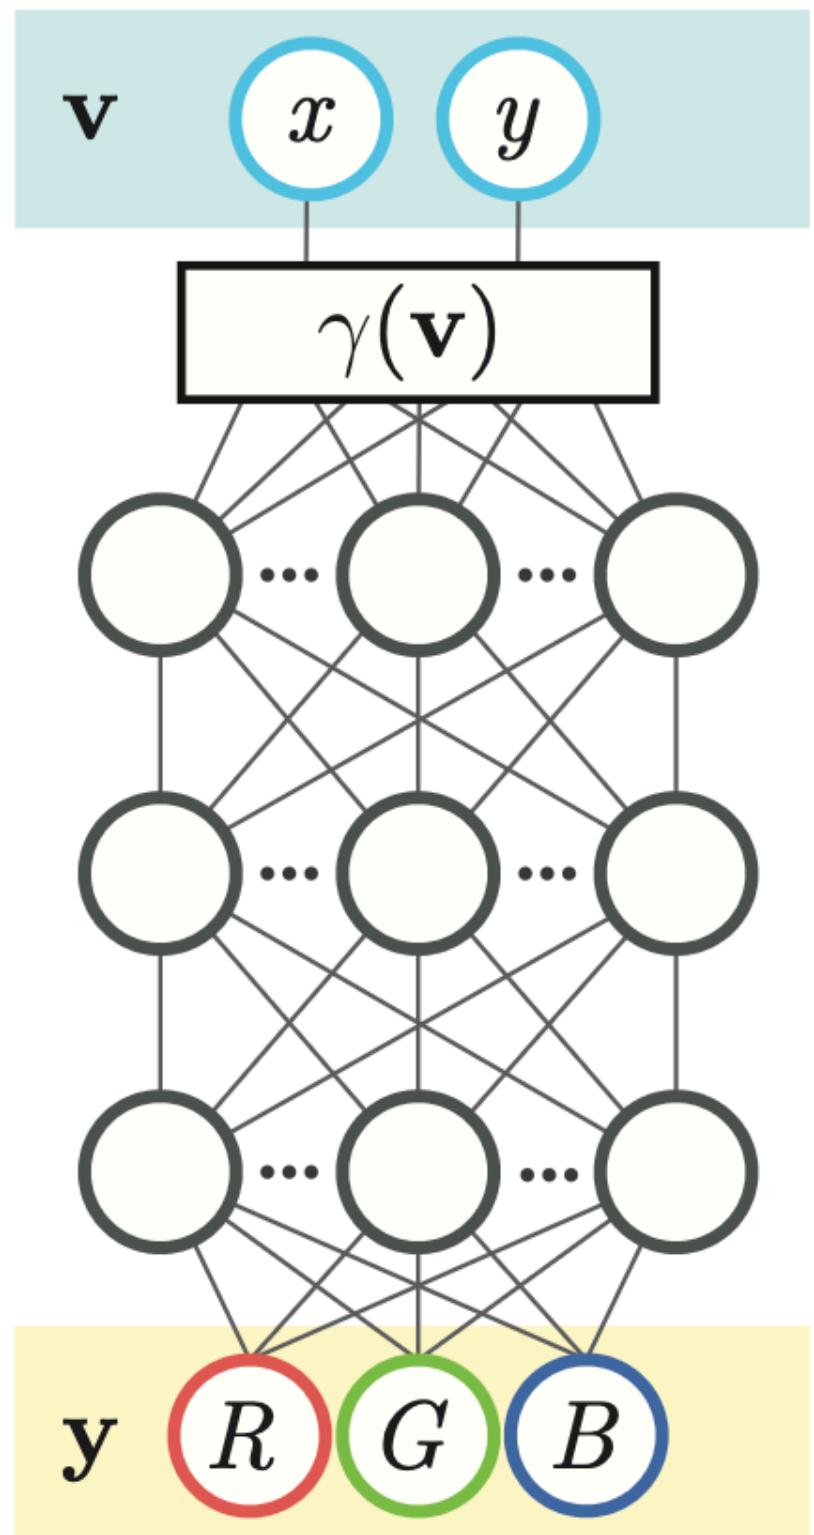
SIREN  
85

# SIREN



~1MB compared to  
110MB of full mesh!

# Fourier Features



(a) Coordinate-based MLP

No Fourier features  
 $\gamma(\mathbf{v}) = \mathbf{v}$

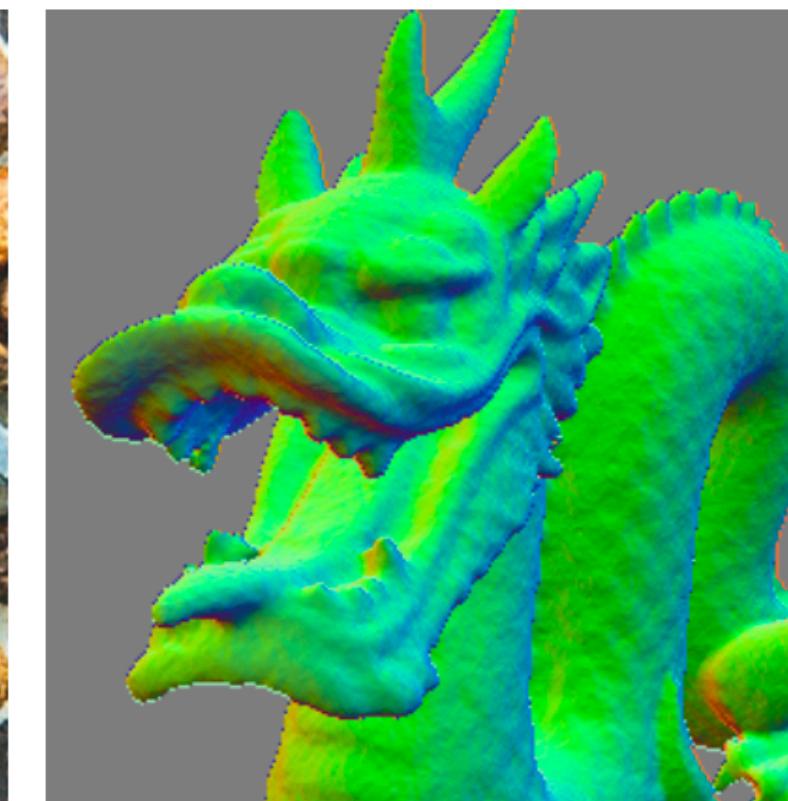
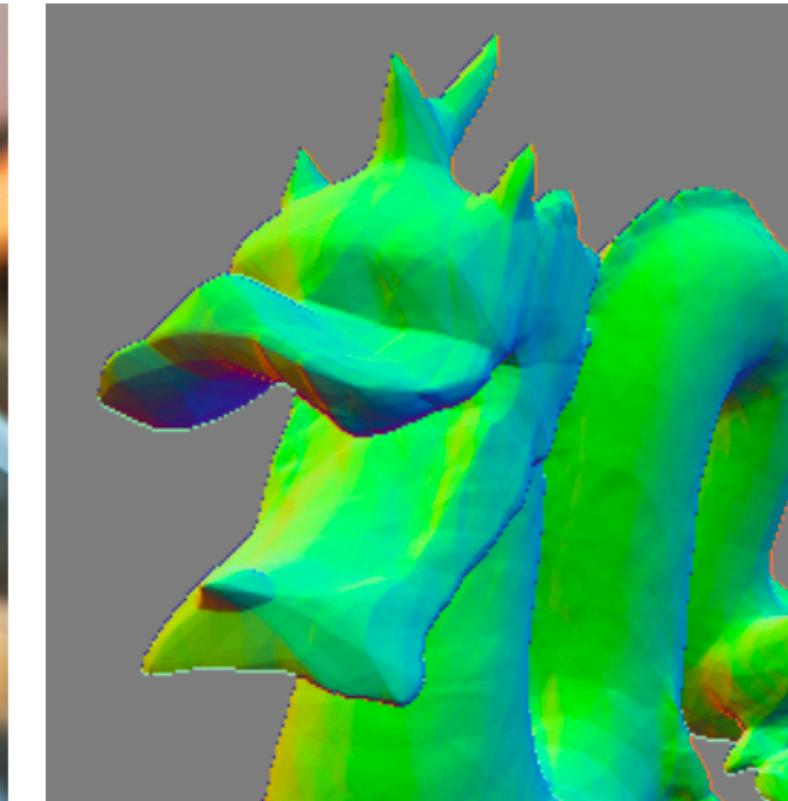


With Fourier features  
 $\gamma(\mathbf{v}) = \text{FF}(\mathbf{v})$



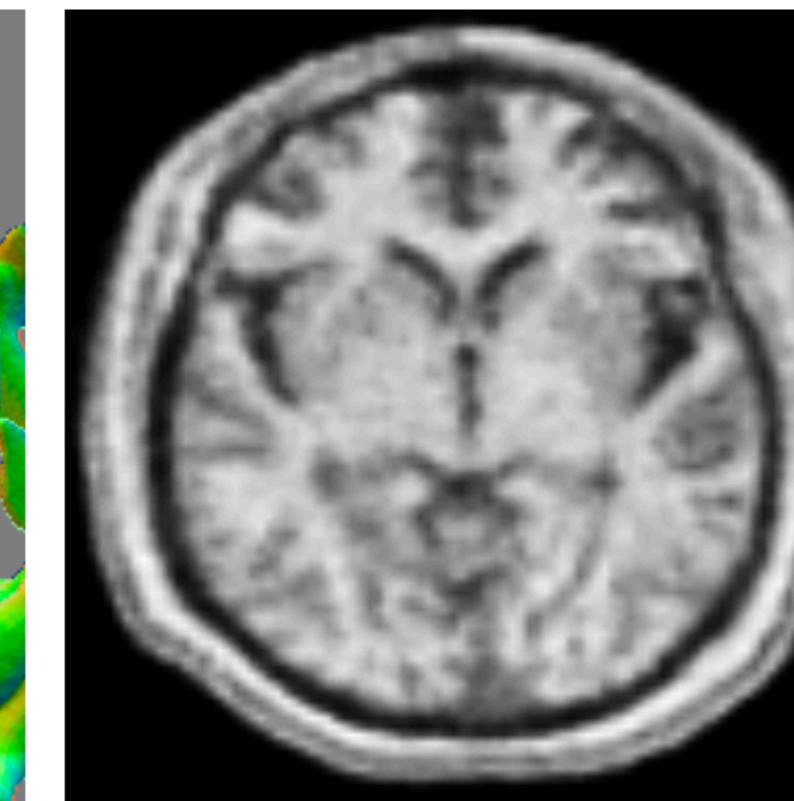
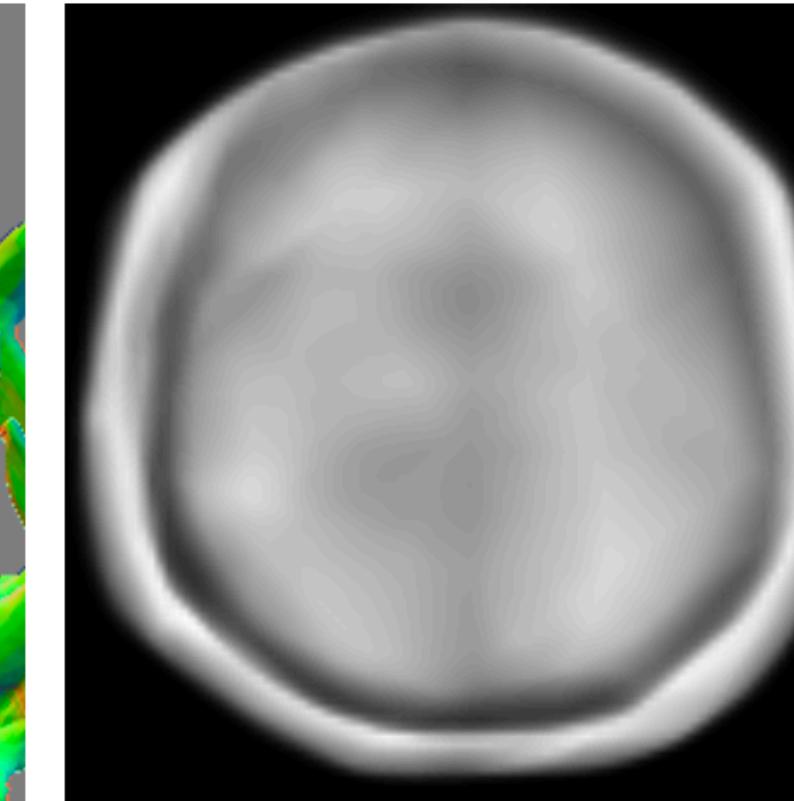
(b) Image regression

$$(x, y) \rightarrow \text{RGB}$$



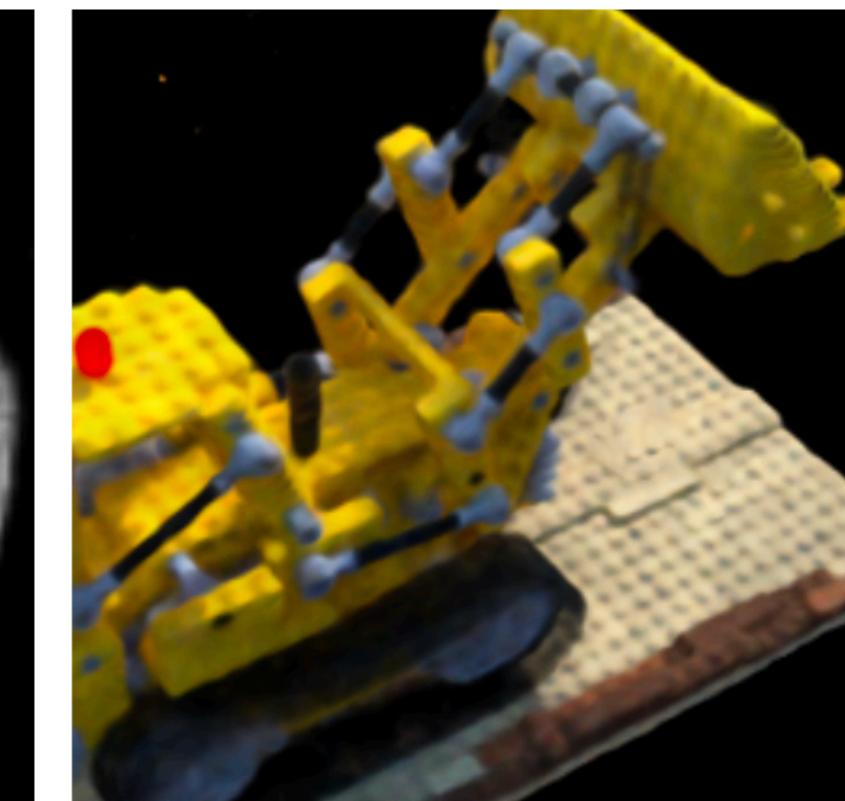
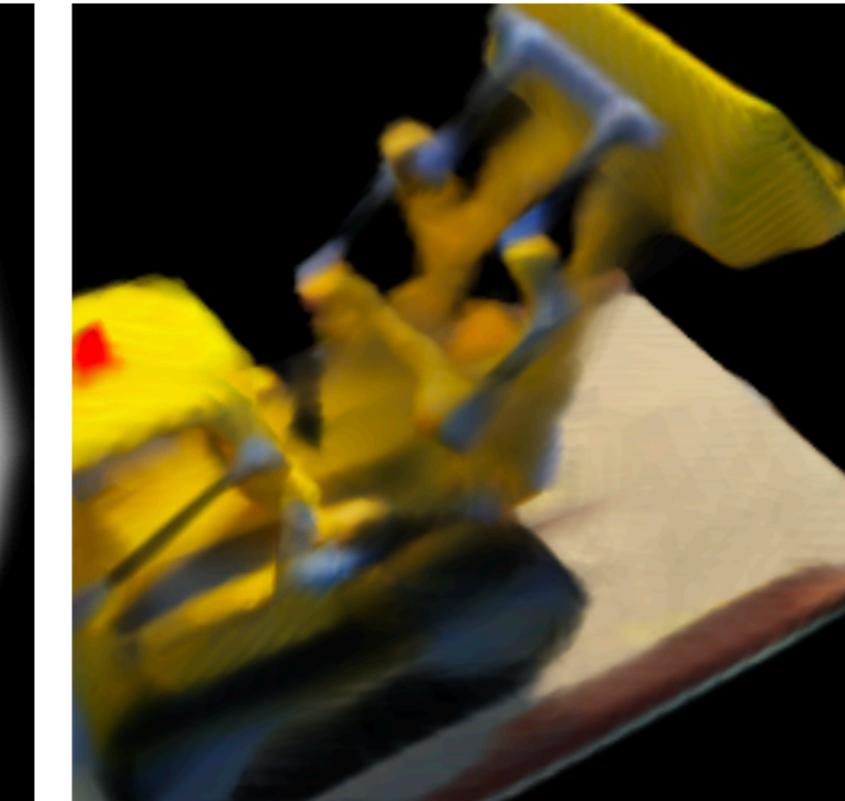
(c) 3D shape regression

$$(x, y, z) \rightarrow \text{occupancy}$$



(d) MRI reconstruction

$$(x, y, z) \rightarrow \text{density}$$



(e) Inverse rendering

$$(x, y, z) \rightarrow \text{RGB, density}$$

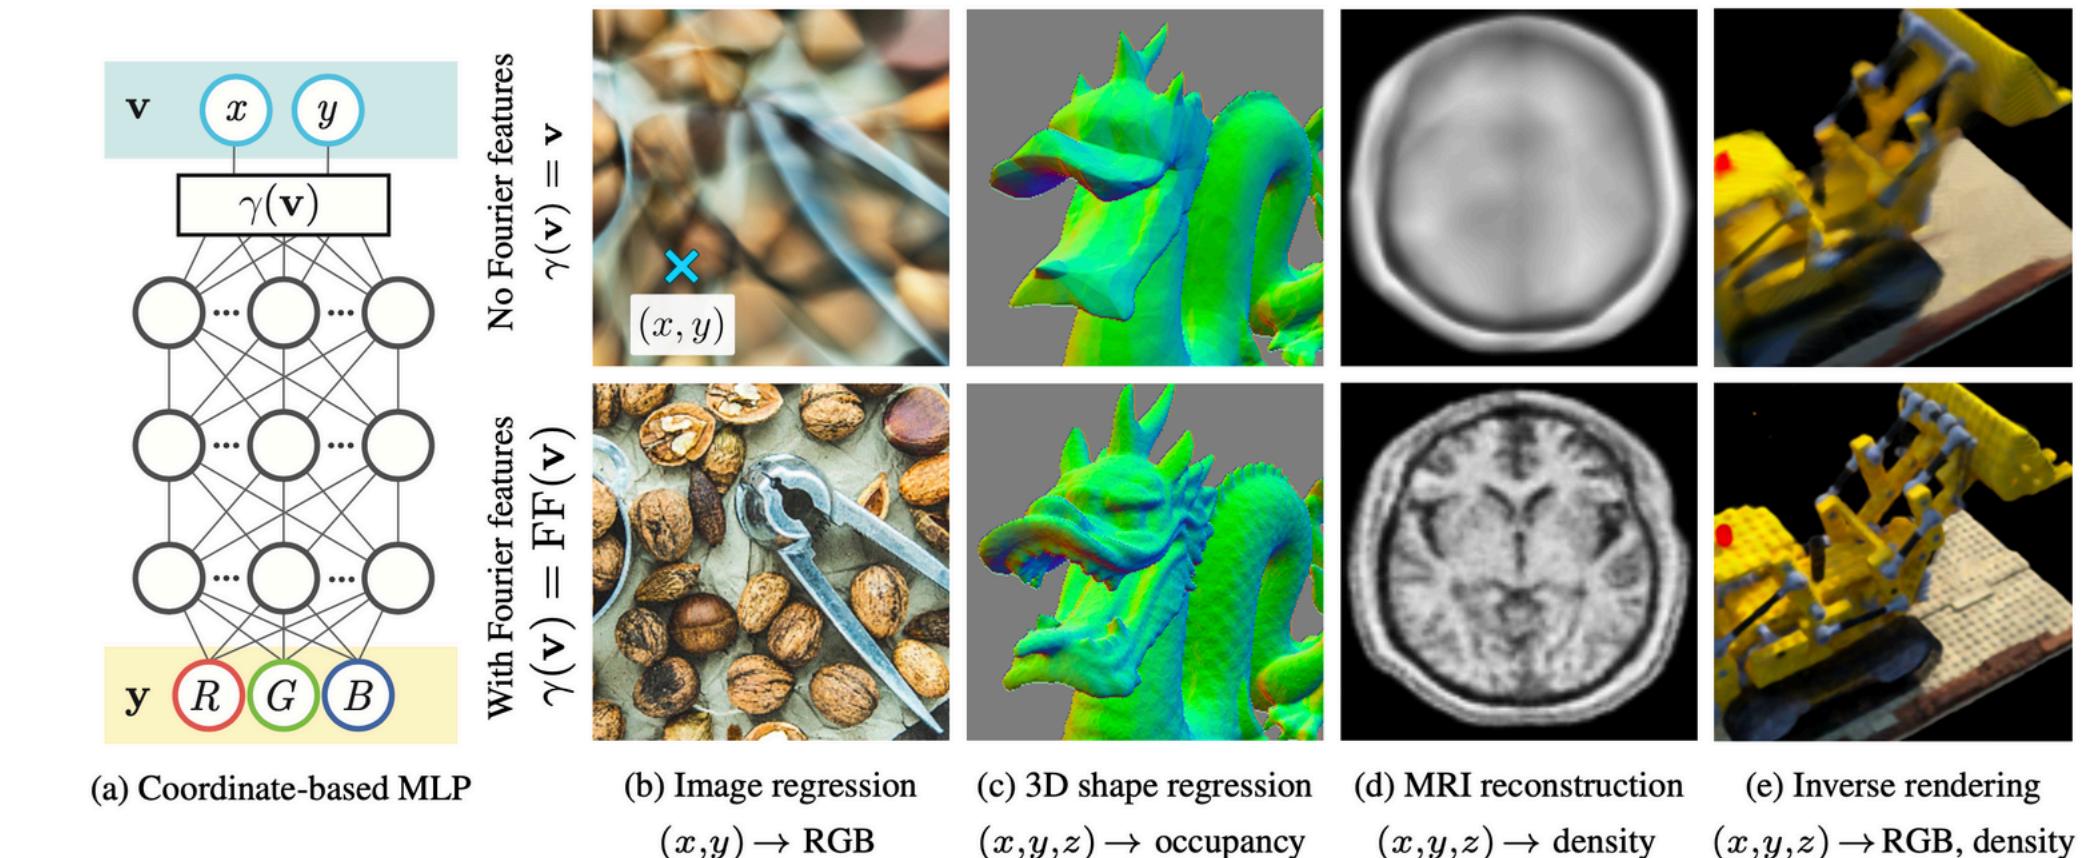
# SIREN: Neural Implicit Representations With Periodic Activation Functions

Sitzmann & Martel et al. NeurIPS 2020



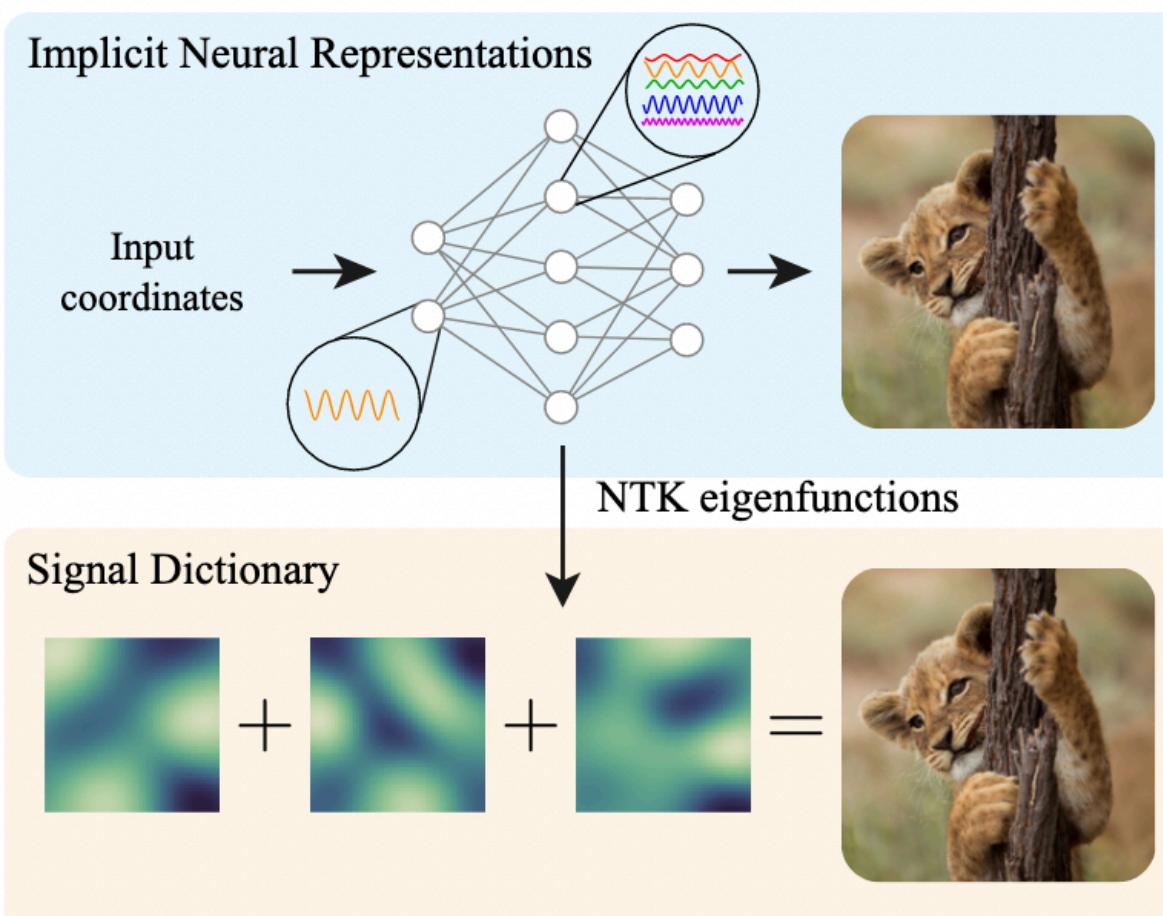
# Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains

Tancik, Srinivasan, Mildenhall NeurIPS 2020



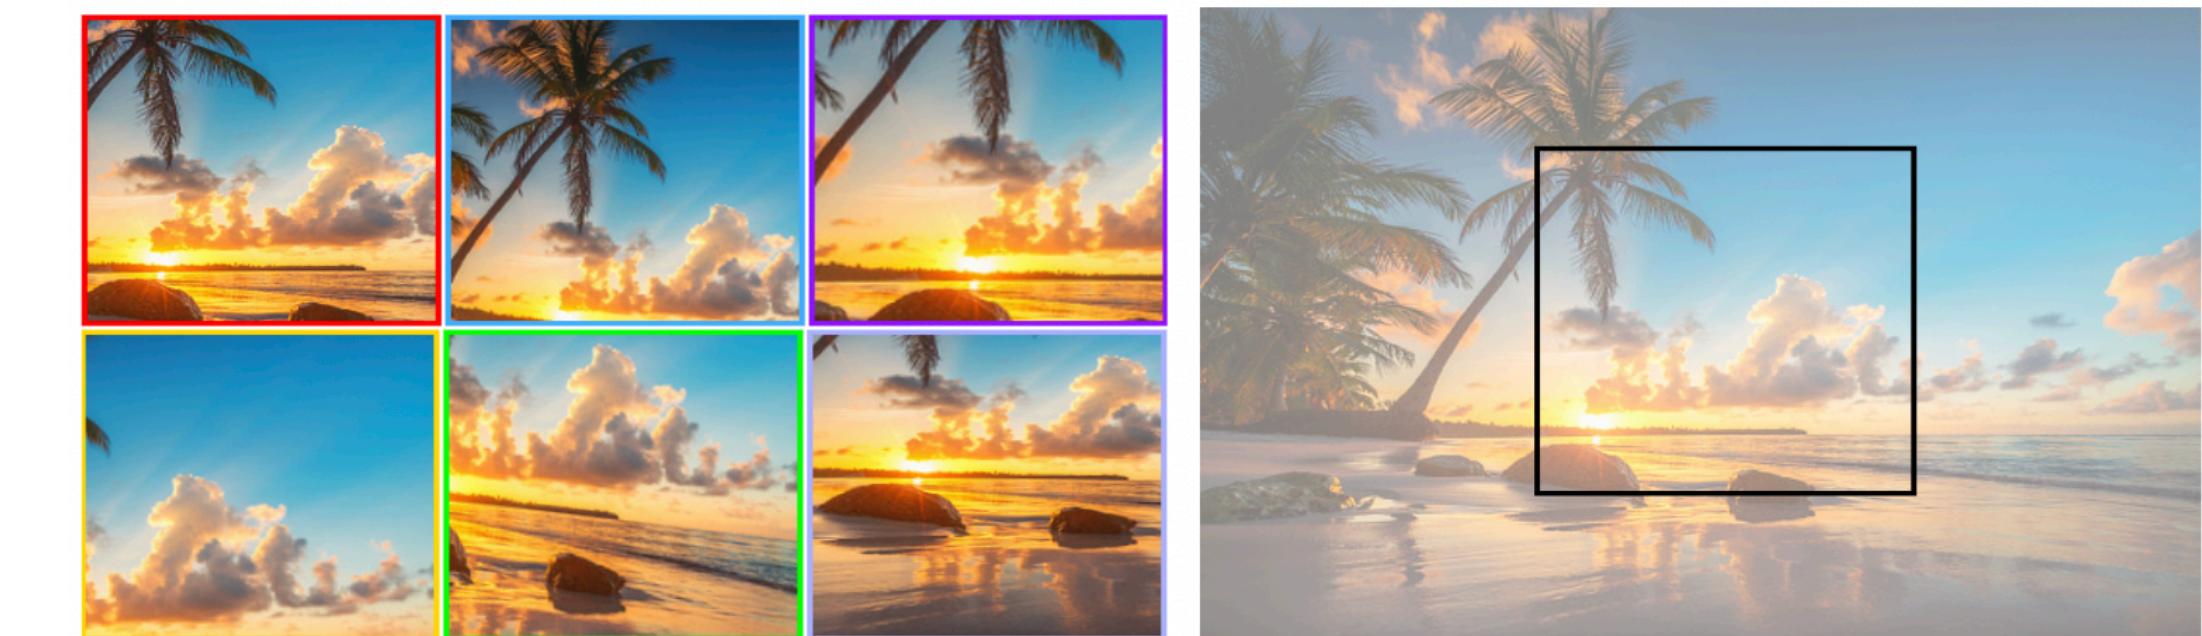
# A Structured Dictionary Perspective on Implicit Neural Representations

Yüce & Ortiz-Jiménez et al. CVPR 2022



# Gaussian Activated Neural Radiance Fields for High Fidelity Reconstruction & Pose Estimation

Chng et al. ECCV 2022



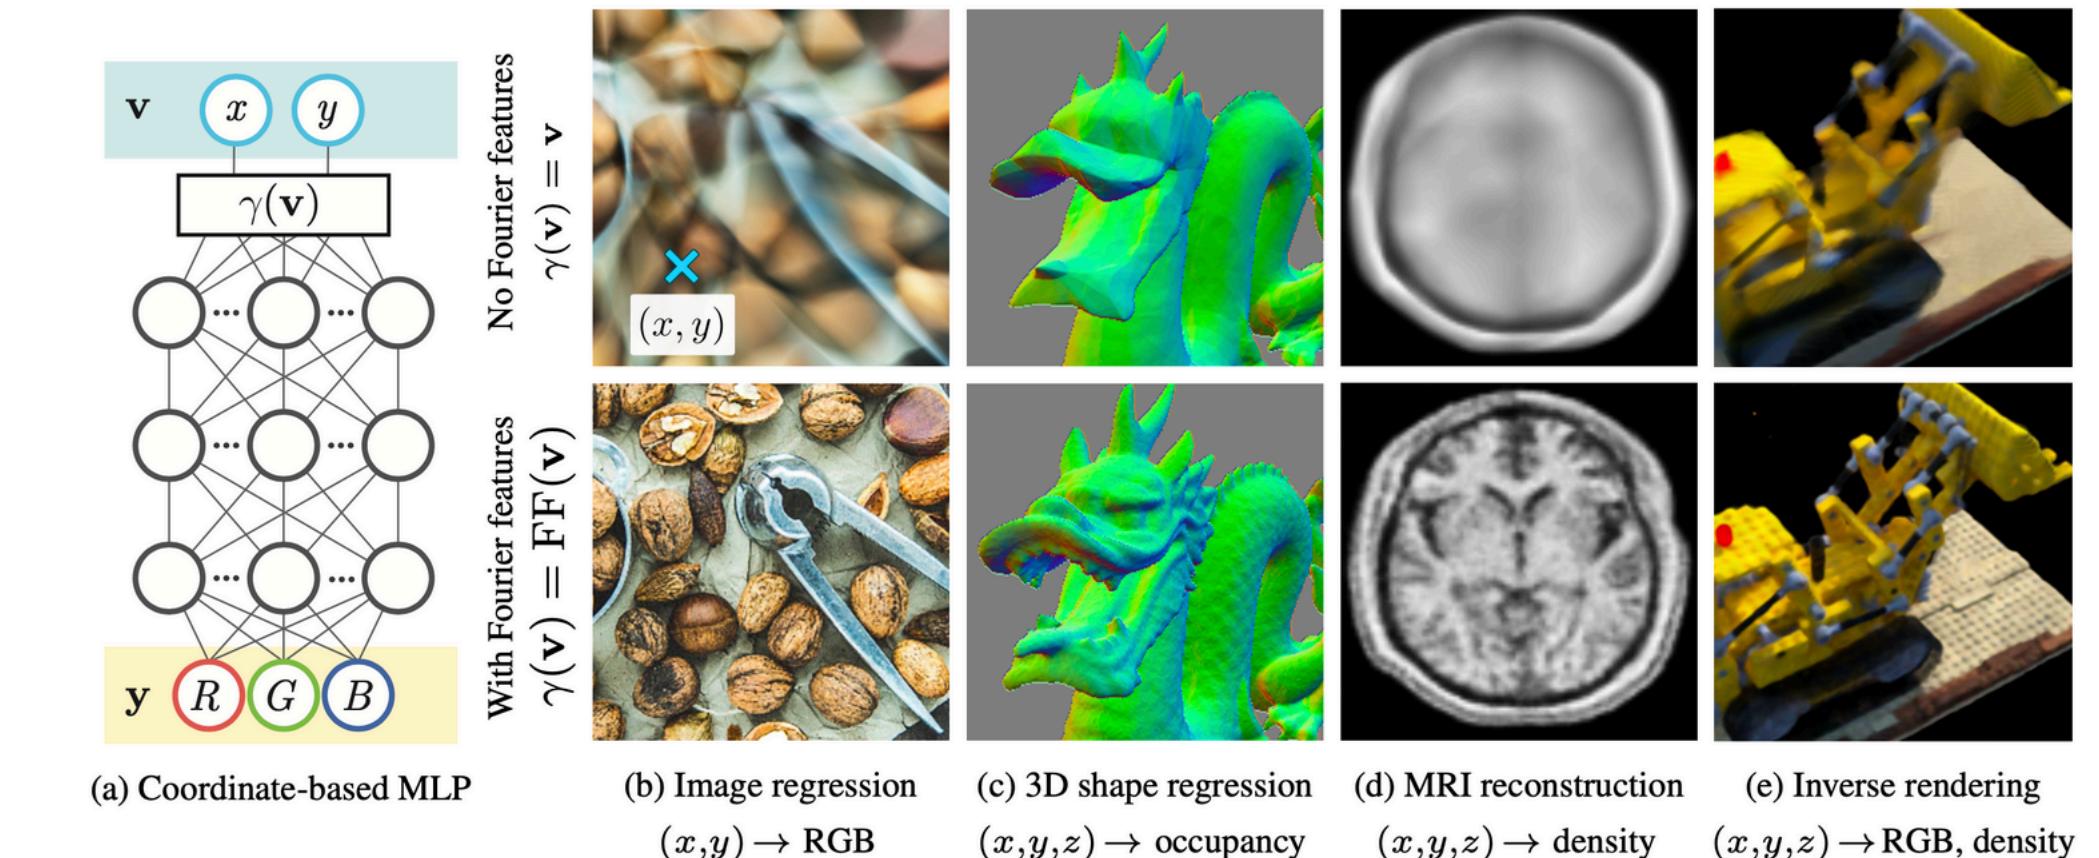
# SIREN: Neural Implicit Representations With Periodic Activation Functions

Sitzmann & Martel et al. NeurIPS 2020



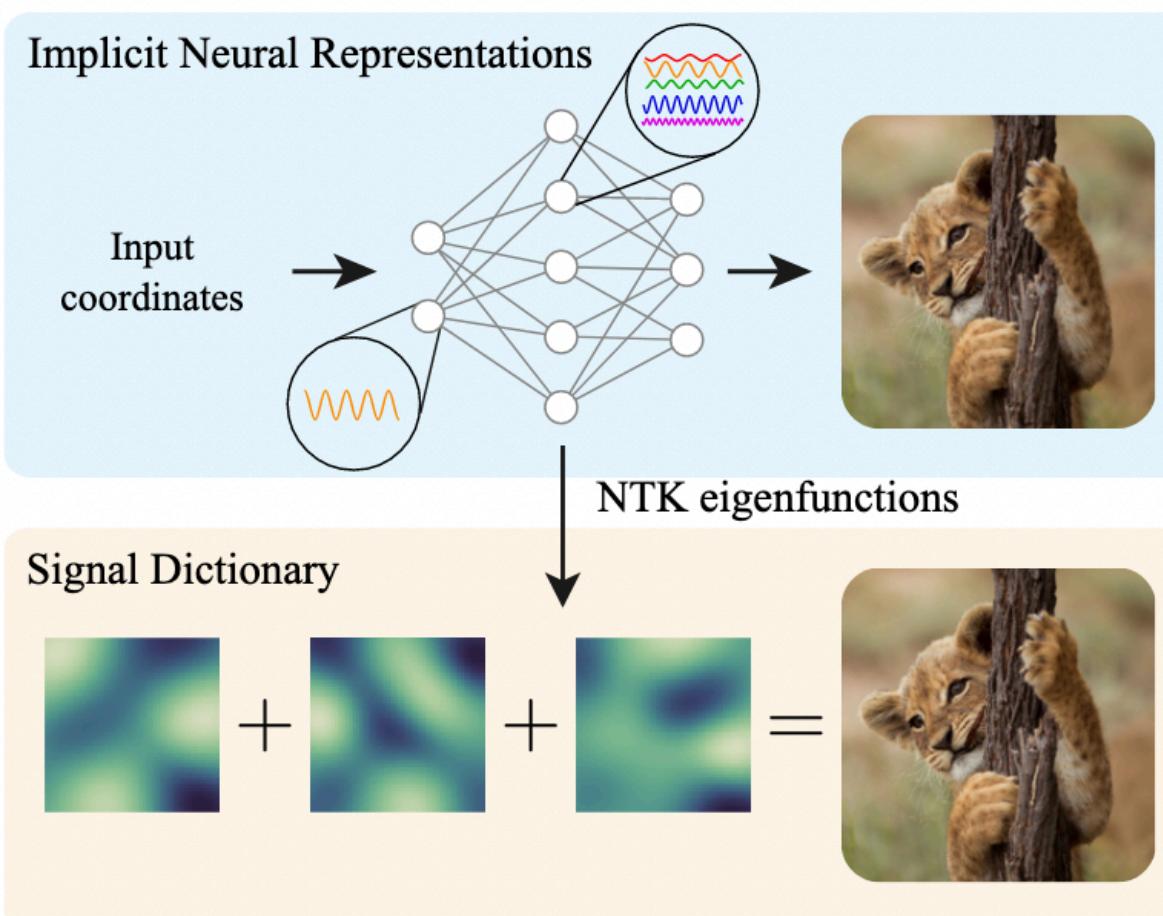
# Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains

Tancik, Srinivasan, Mildenhall NeurIPS 2020



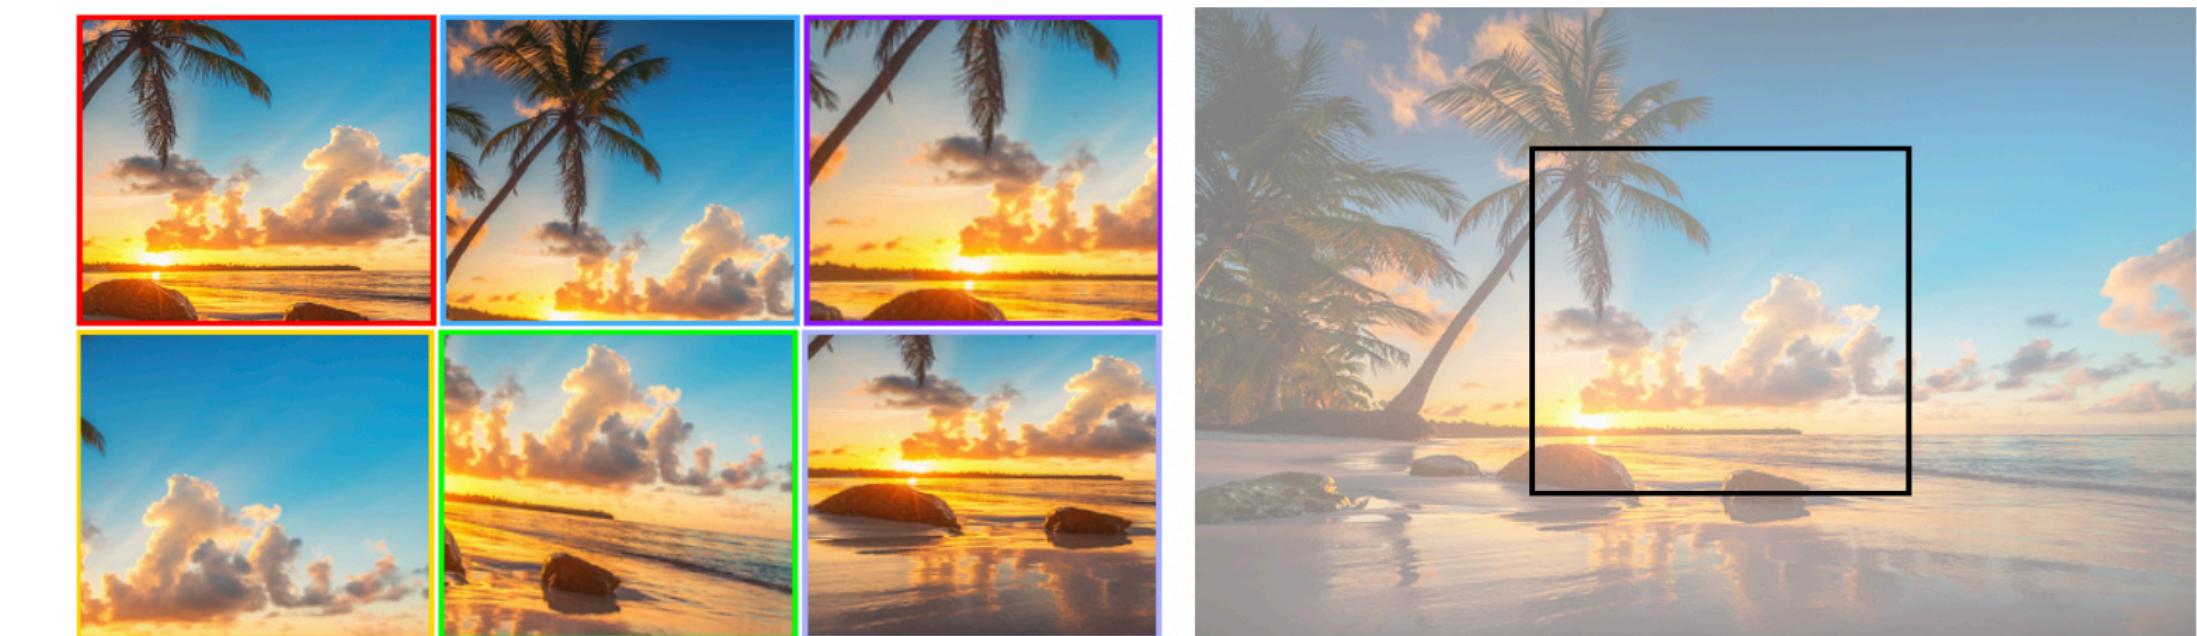
# A Structured Dictionary Perspective on Implicit Neural Representations

Yüce & Ortiz-Jiménez et al. CVPR 2022



# Gaussian Activated Neural Radiance Fields for High Fidelity Reconstruction & Pose Estimation

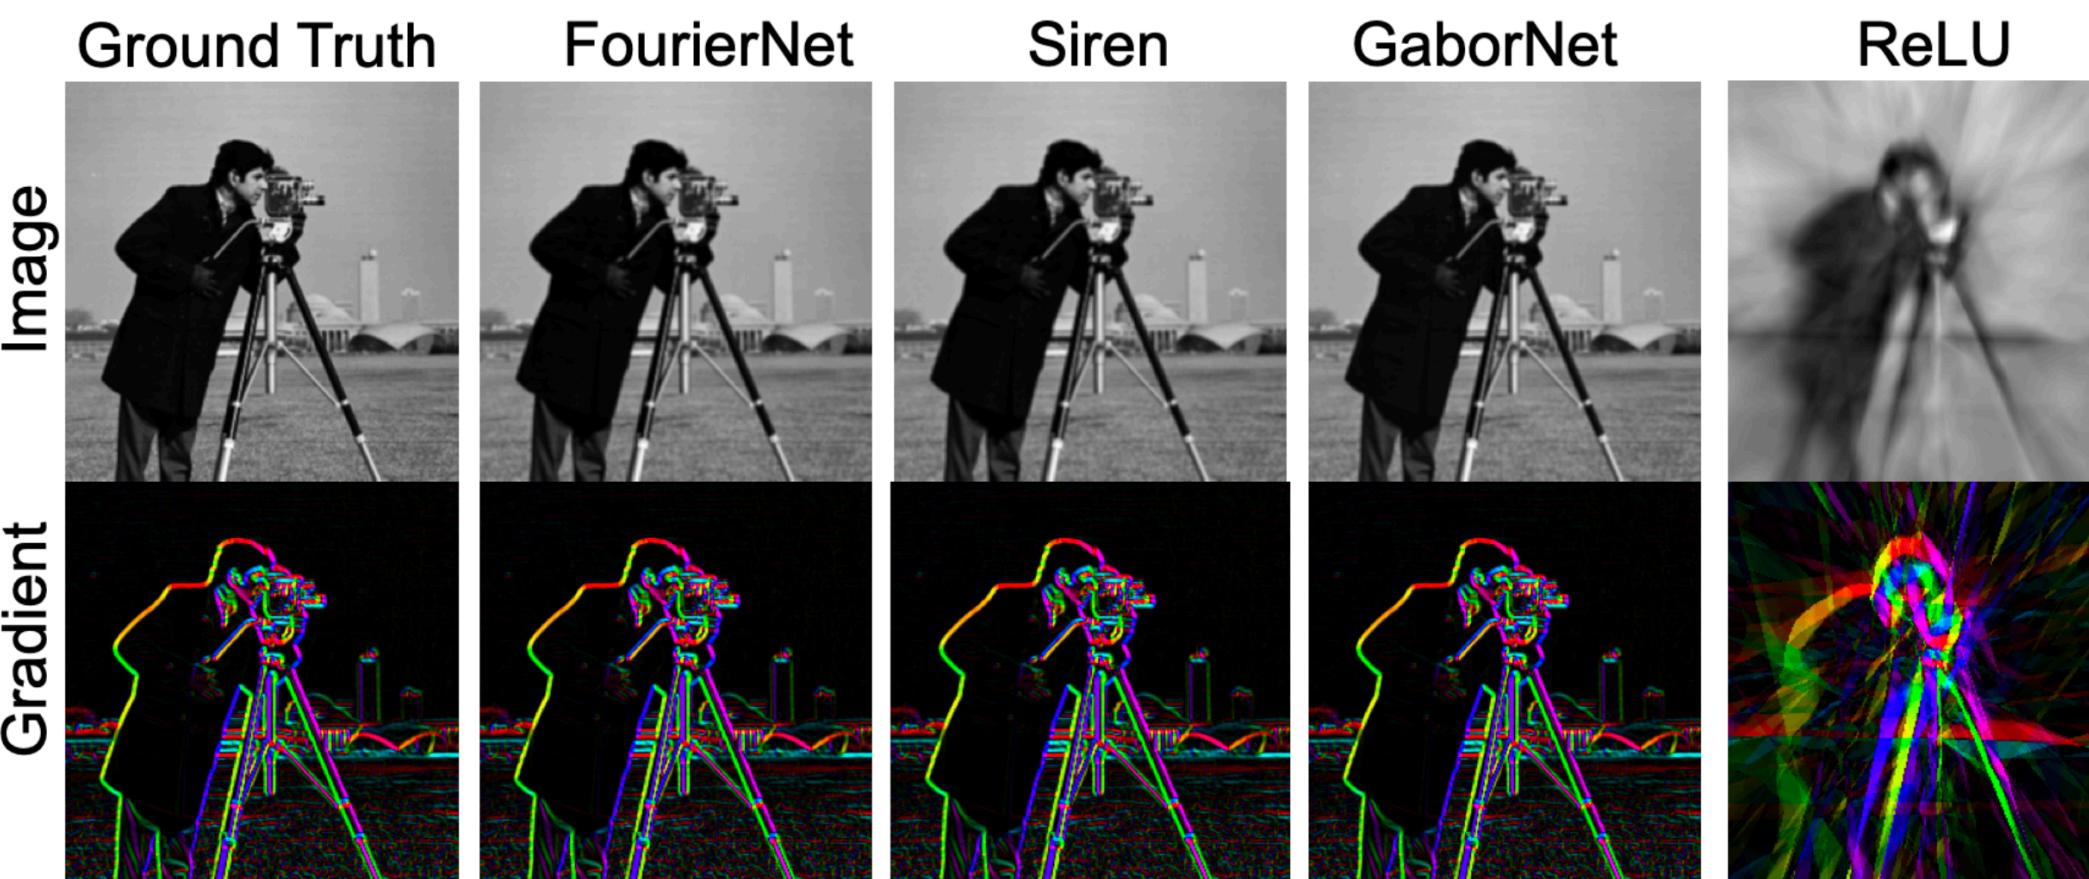
Chng et al. ECCV 2022



# Neural (?) Fields with analytical (but not tractably computable) Fourier spectrum!

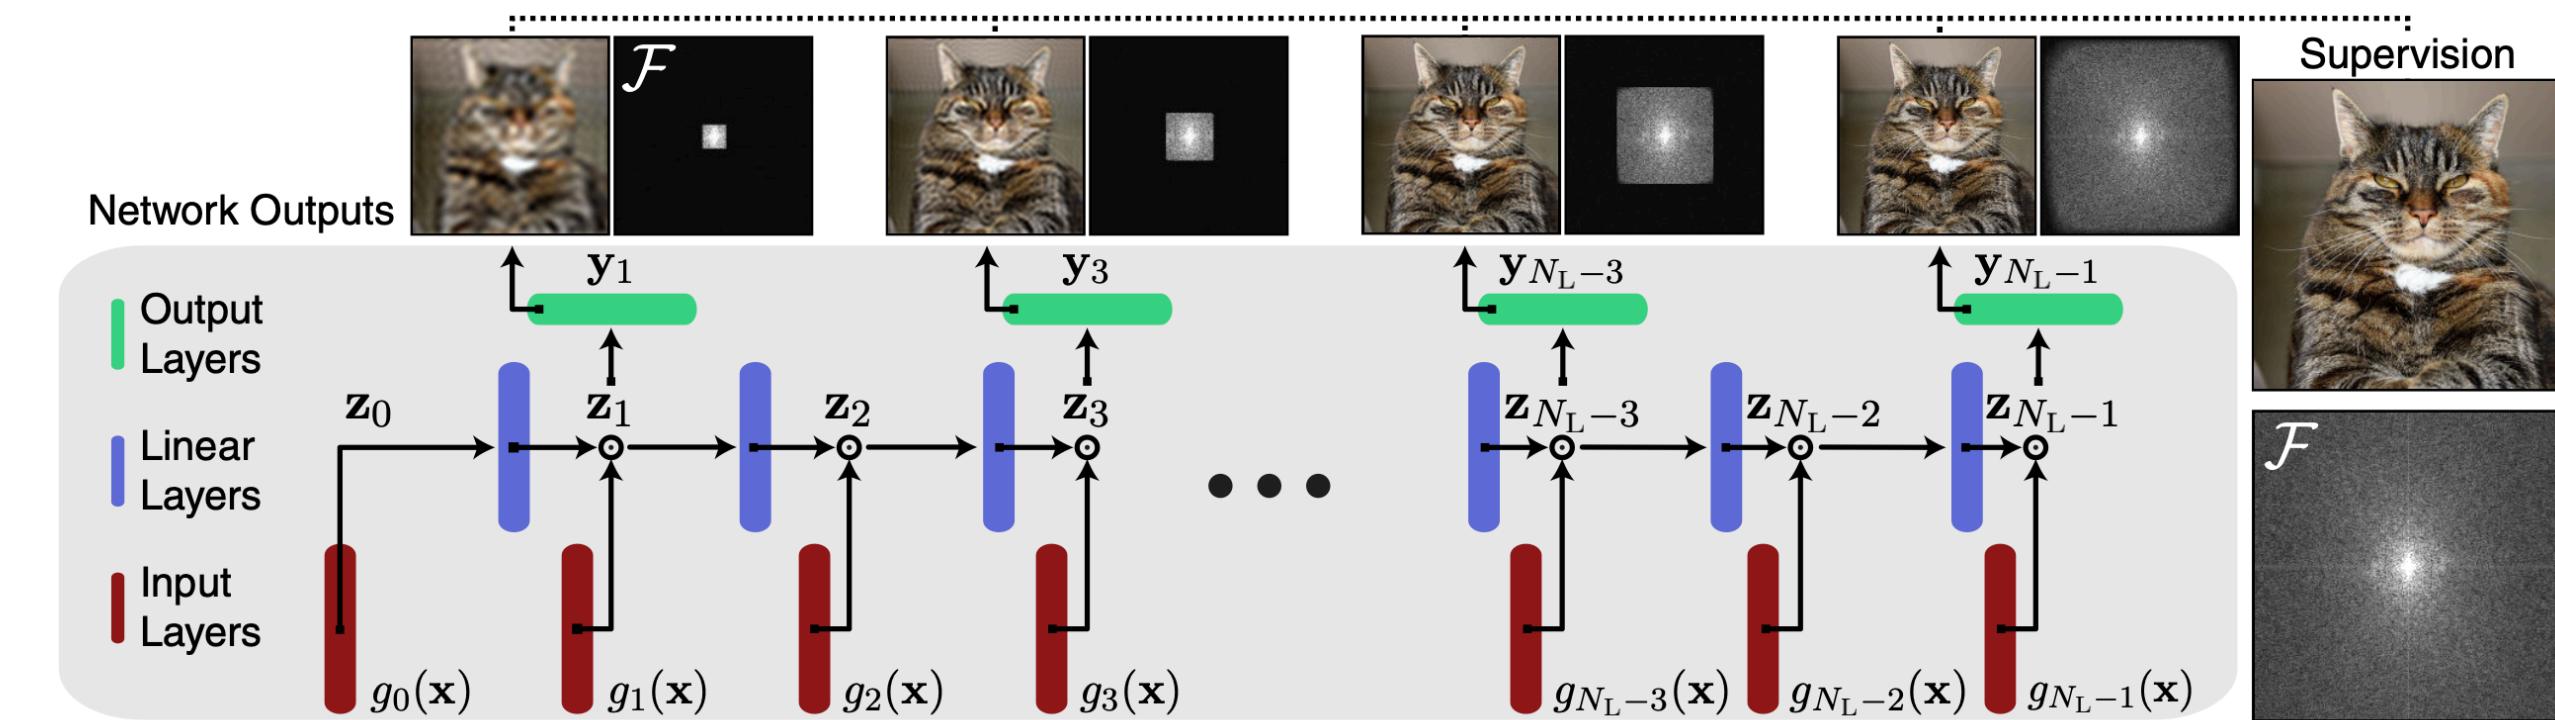
## Multiplicative Filter Networks

Fathony et al. ICLR 2021



## BACON: Band-limited coordinate networks for multiscale scene representation

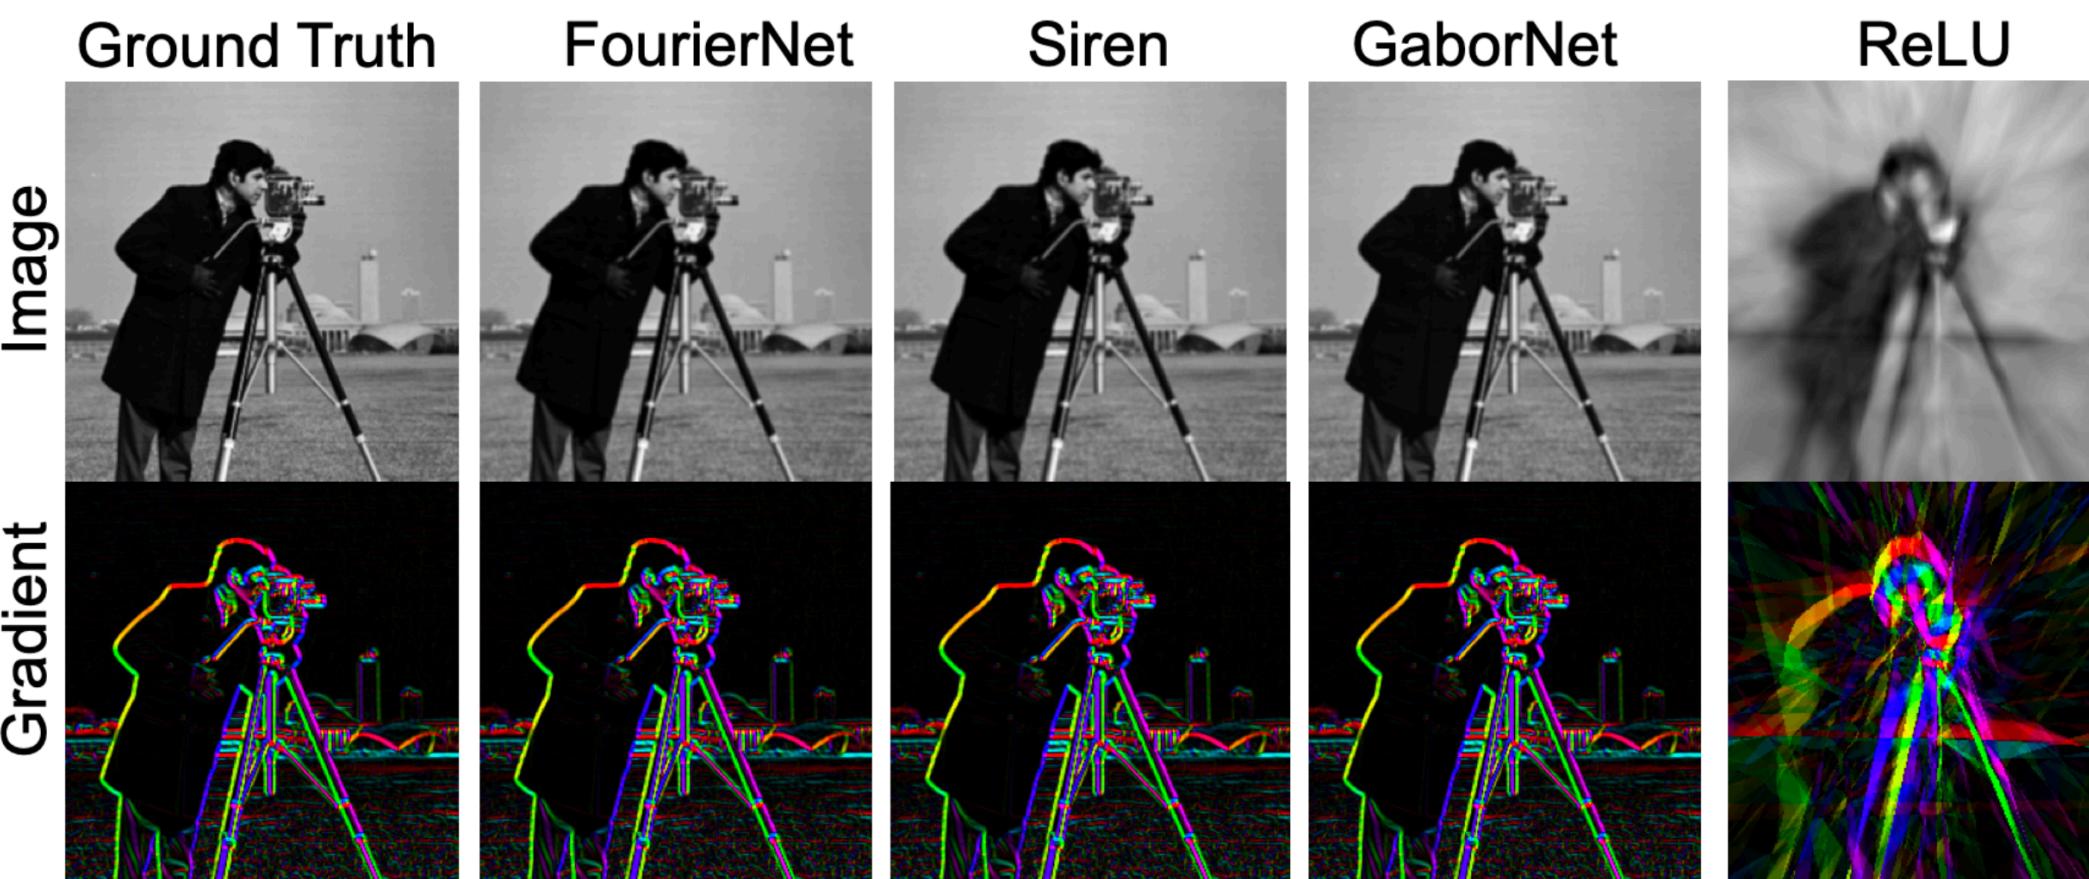
Lindell et al. CVPR 2022



# Neural (?) Fields with analytical (but not tractably computable) Fourier spectrum!

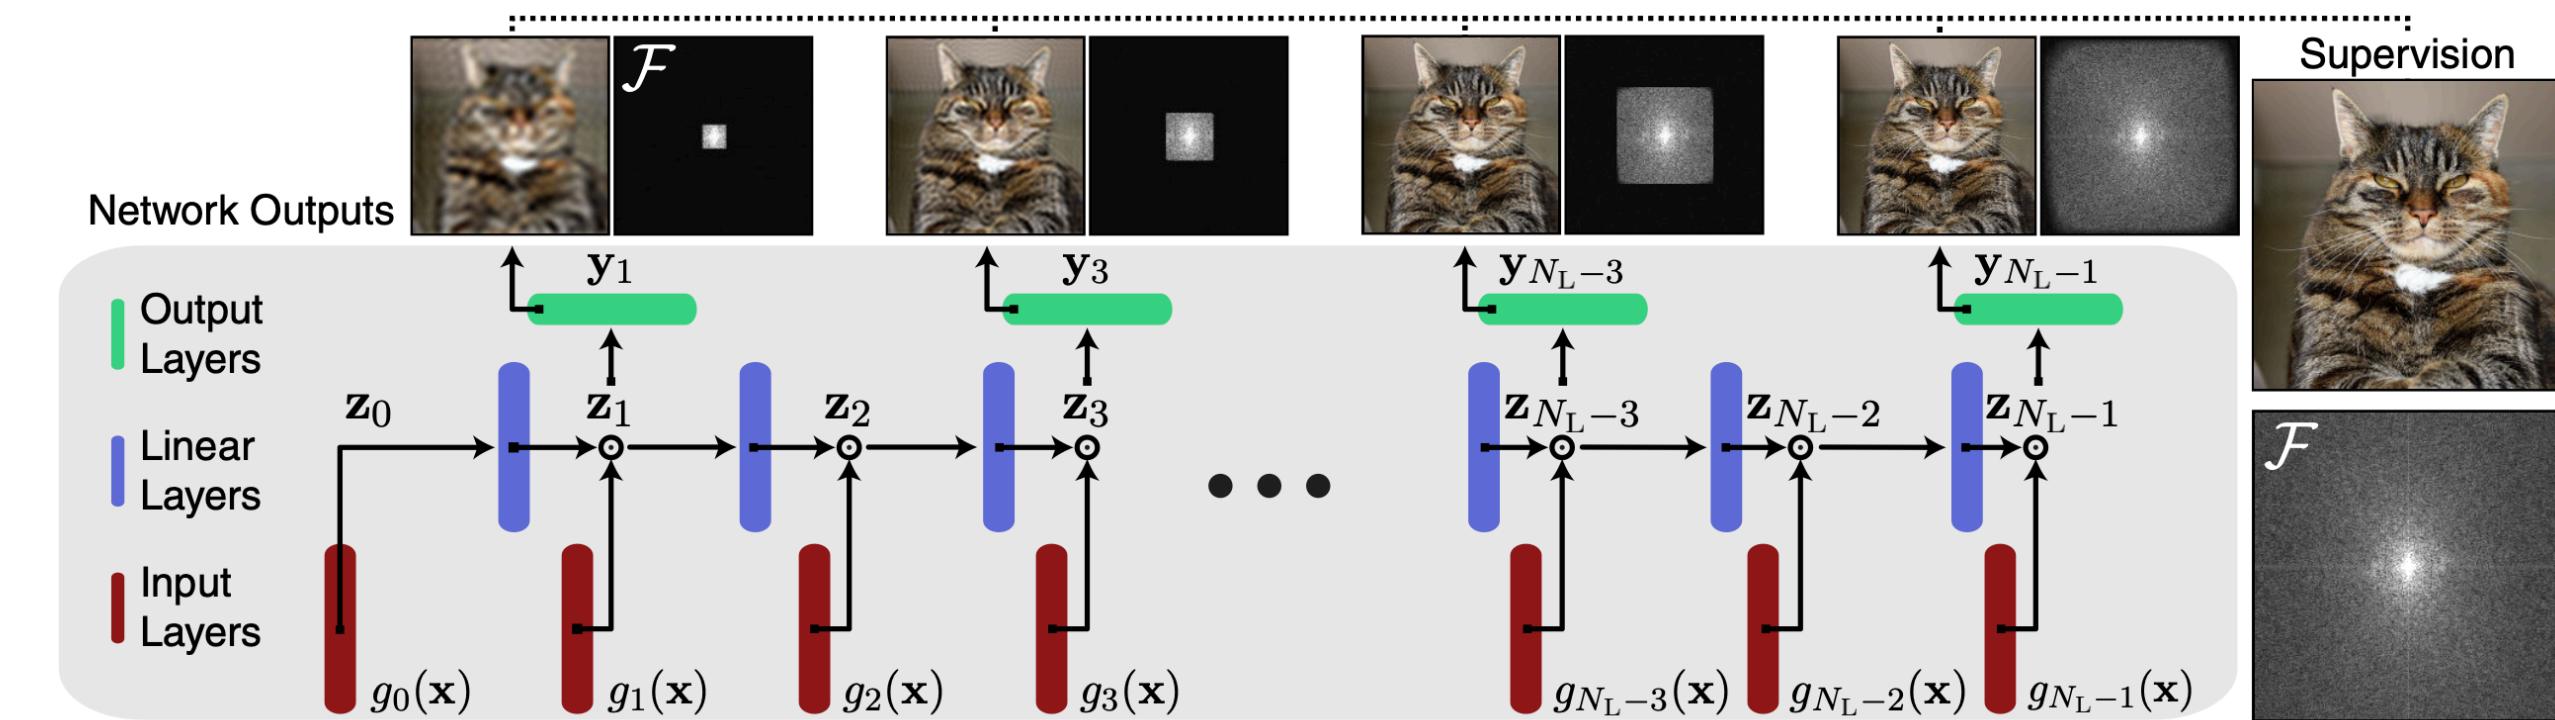
## Multiplicative Filter Networks

Fathony et al. ICLR 2021



## BACON: Band-limited coordinate networks for multiscale scene representation

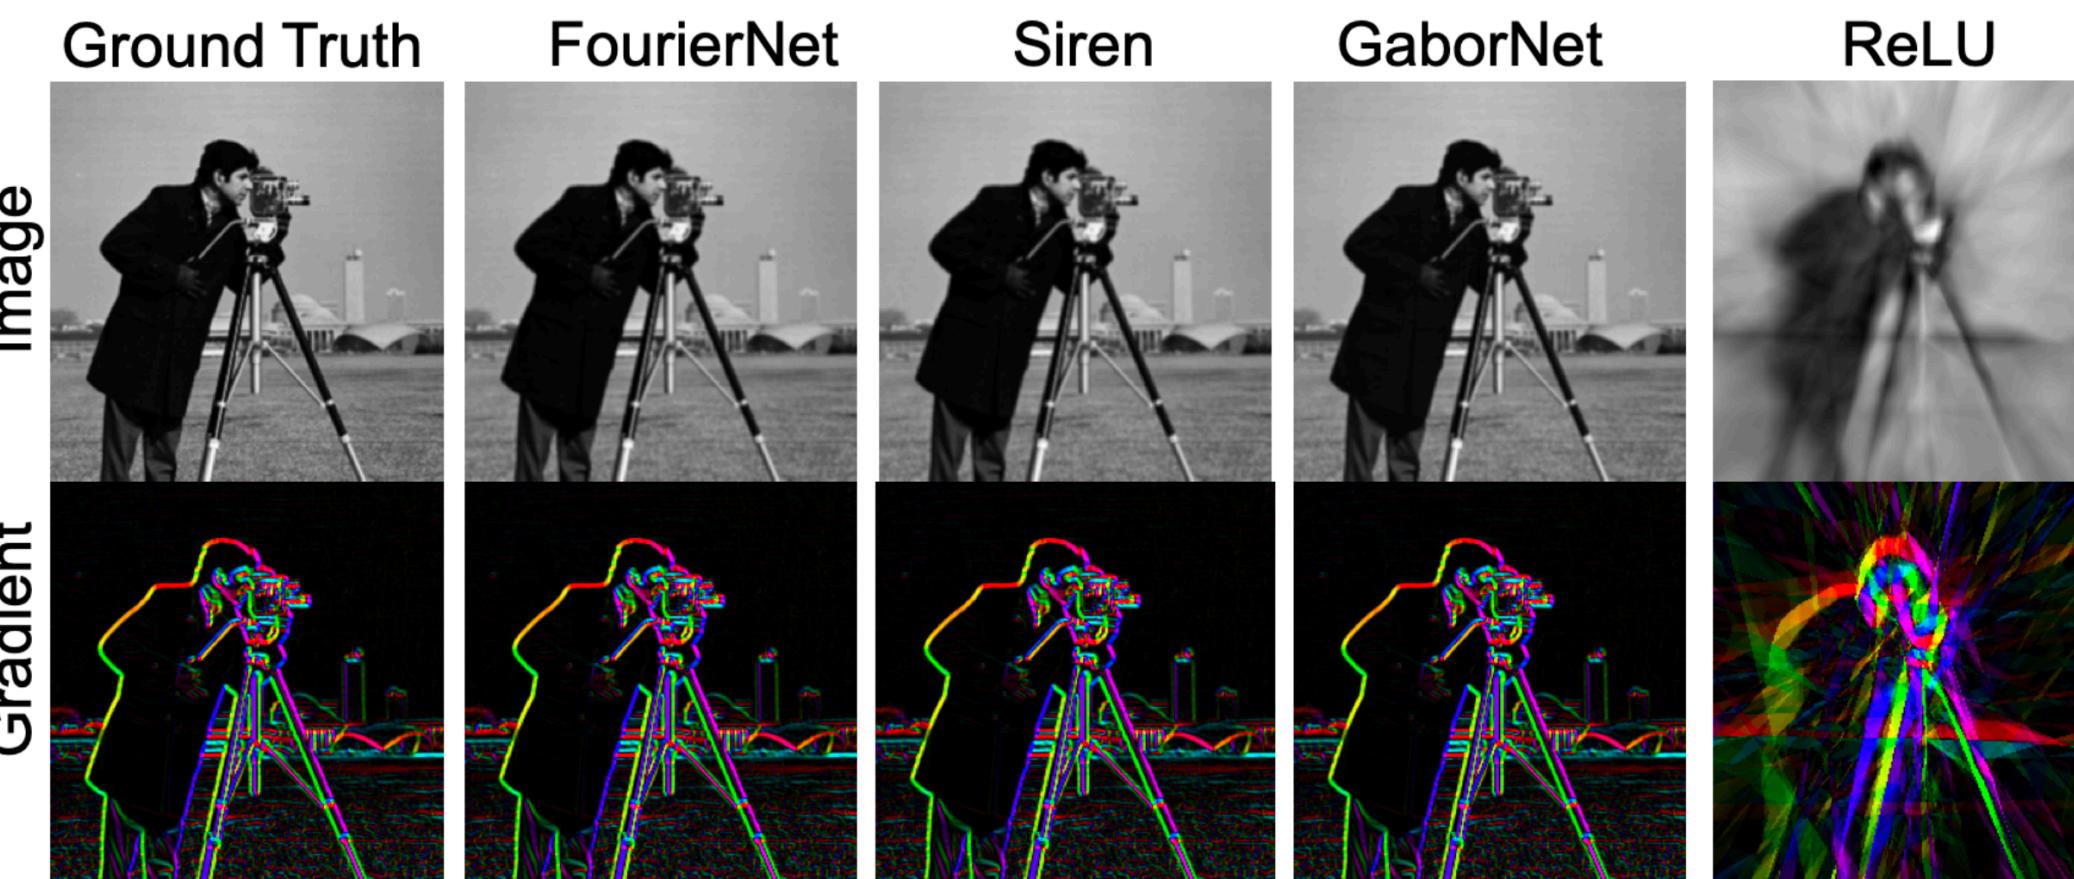
Lindell et al. CVPR 2022



# Neural (?) Fields with analytical (but not tractably computable) spectrum!

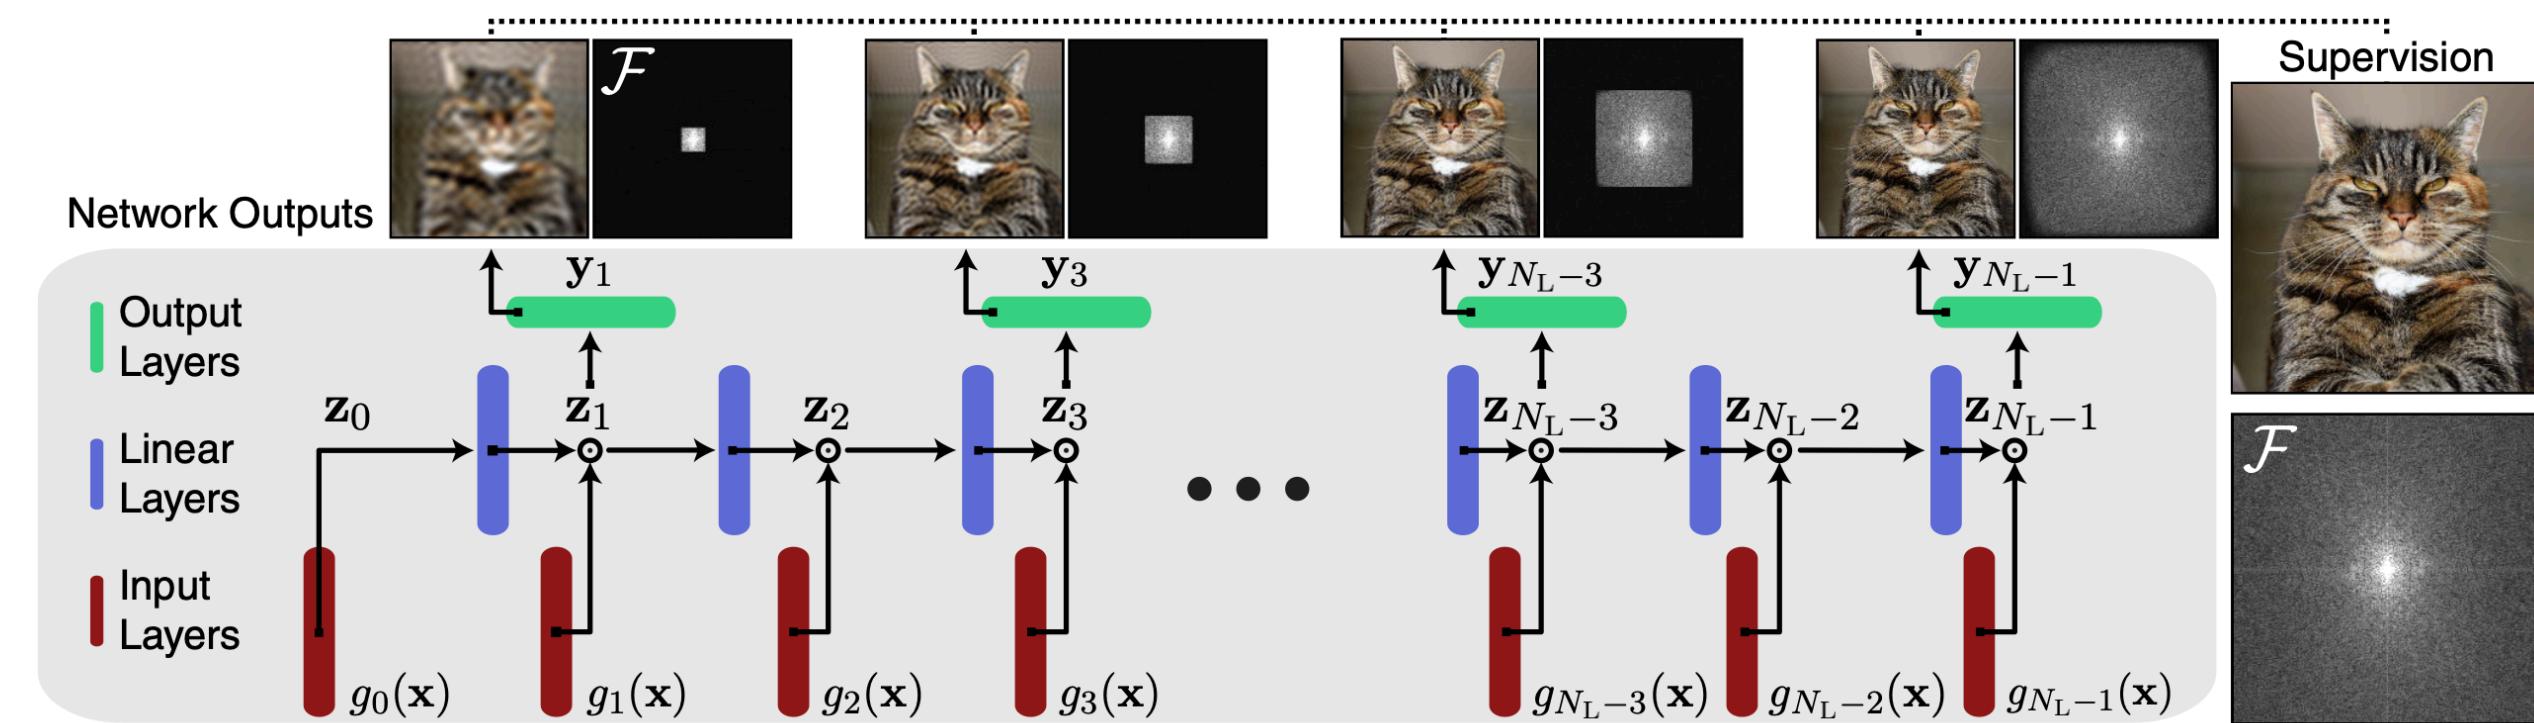
## Multiplicative Filter Networks

Fathony et al. ICLR 2021



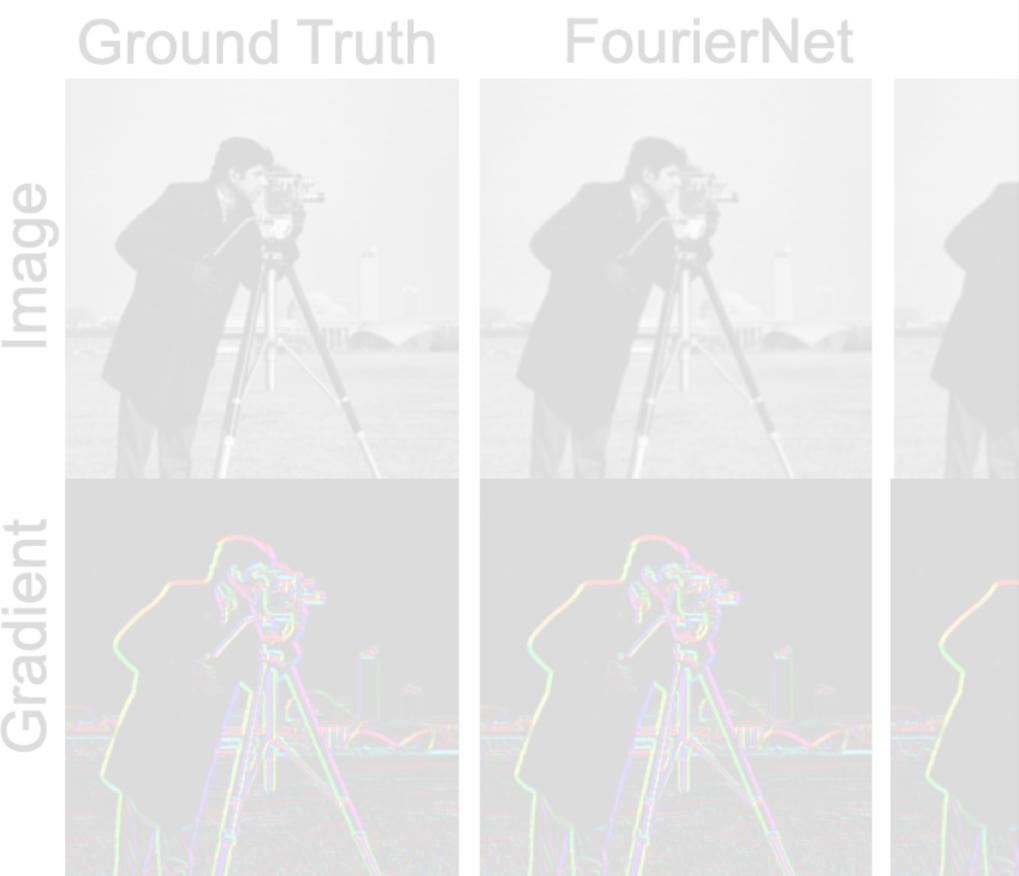
## BACON: Band-limited coordinate networks for multiscale scene representation

Lindell et al. CVPR 2022



(but

## Multiplicative Fathy et al.



arXiv:2111.11426v4 [cs.CV] 5 Apr 2022

EUROGRAPHICS 2022  
D. Meneveaux and G. Patanè  
(Guest Editors)

Volume 41 (2022), Number 2  
STAR – State of The Art Report

rum!

ordinate networks for  
representation  
CVPR 2022

## Neural Fields in Visual Computing and Beyond

Yiheng Xie<sup>1,2</sup> Towaki Takikawa<sup>3,4</sup> Shunsuke Saito<sup>5</sup> Or Litany<sup>4</sup> Shiqin Yan<sup>1</sup> Numair Khan<sup>1</sup> Federico Tombari<sup>6,7</sup>  
James Tompkin<sup>1</sup> Vincent Sitzmann<sup>8†</sup> Srinath Sridhar<sup>1†</sup>

<sup>1</sup>Brown University <sup>2</sup>Unity Technologies <sup>3</sup>University of Toronto <sup>4</sup>NVIDIA <sup>5</sup>Meta Reality Labs Research <sup>6</sup>Google <sup>7</sup>Technical University of Munich  
<sup>8</sup>Massachusetts Institute of Technology <sup>†</sup>Equal advising

<https://neuralfields.cs.brown.edu/>

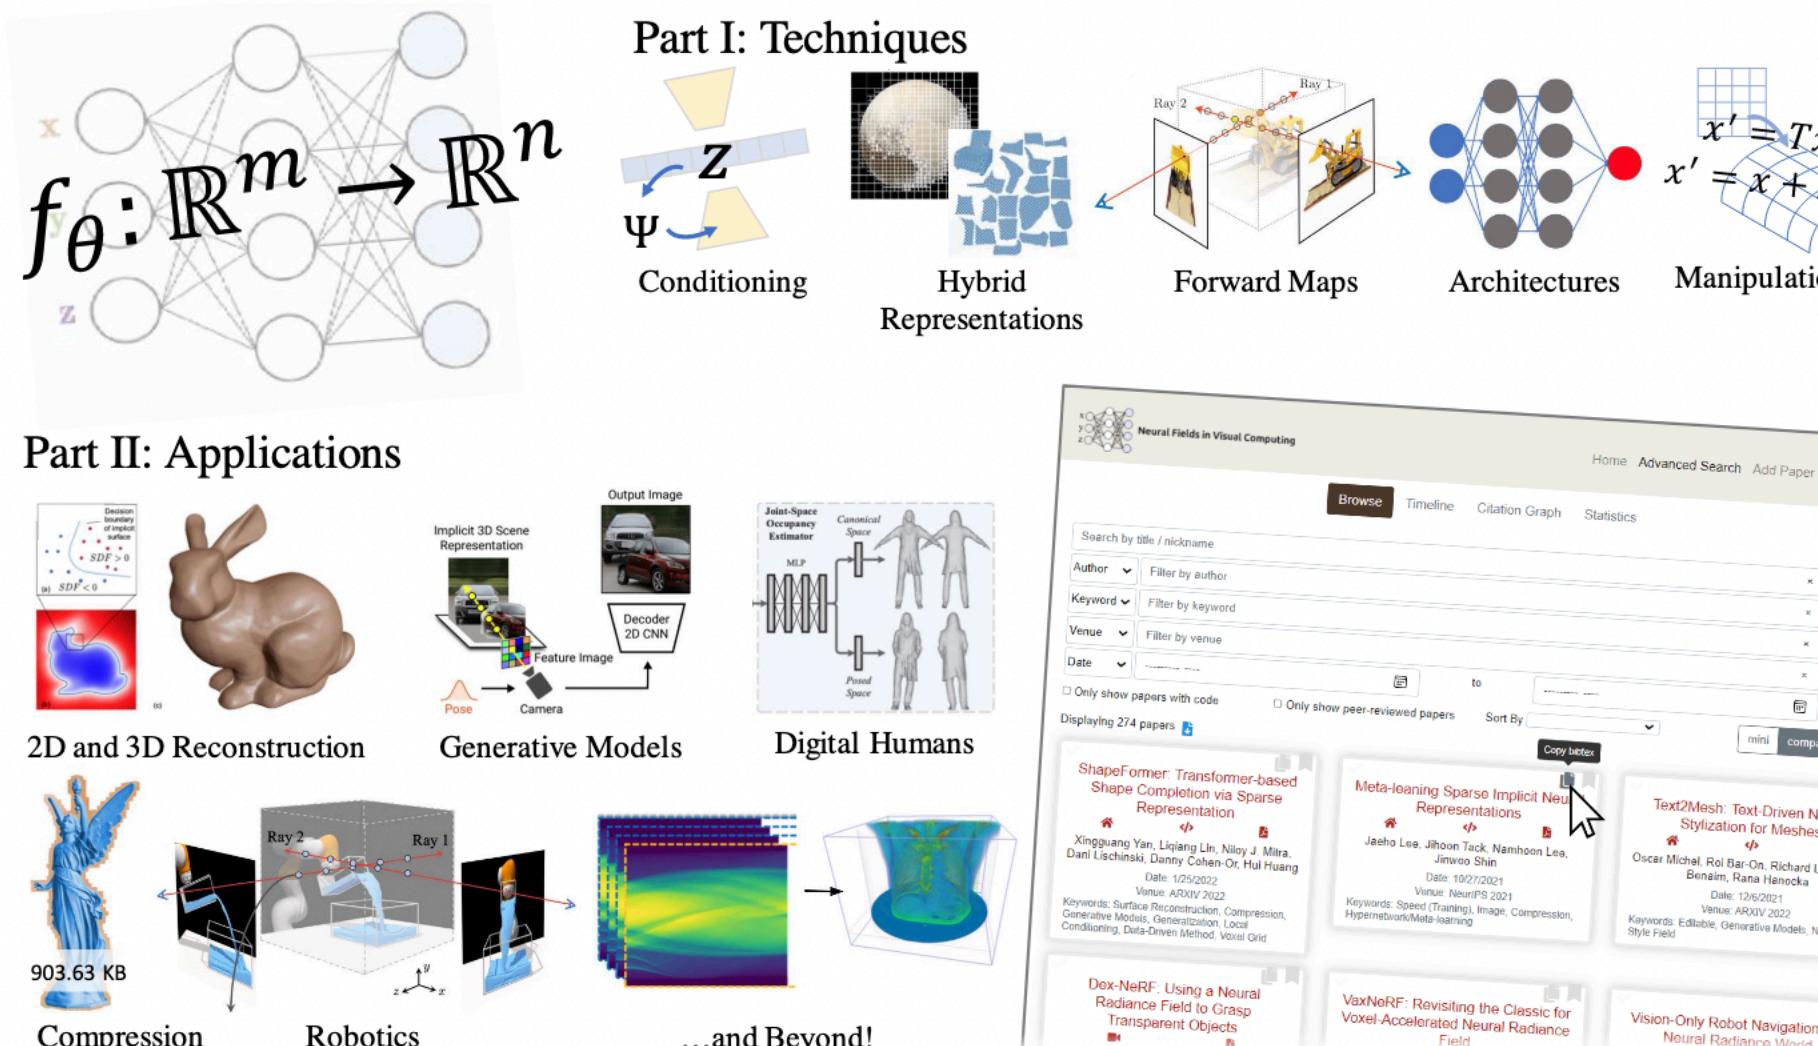
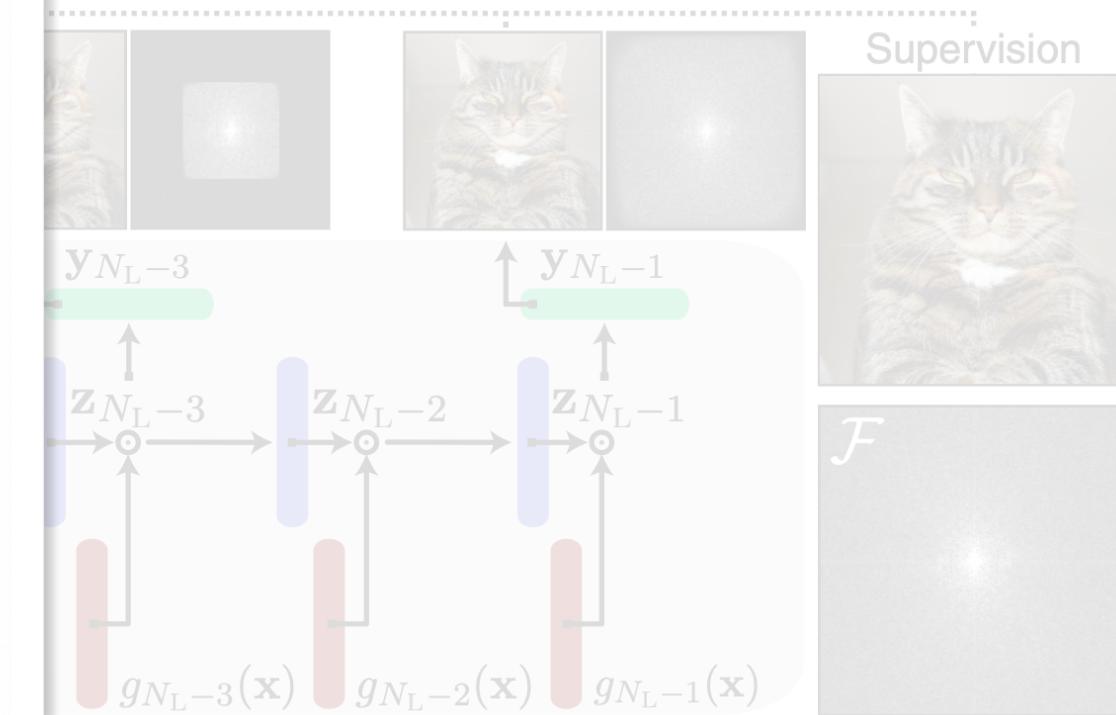


Figure 1: Contribution of this report. Following a survey of over 250 papers, we provide a review of (**Part I**) techniques in neural fields such as prior learning and conditioning, representations, forward maps, architectures, and manipulation, and of (**Part II**) applications in visual computing including 2D image processing, 3D scene reconstruction, generative modeling, digital humans, compression, robotics, and beyond. This report is complemented by a [community-driven website](https://neuralfields.cs.brown.edu/) with search, filtering, bibliographic, and visualization features.

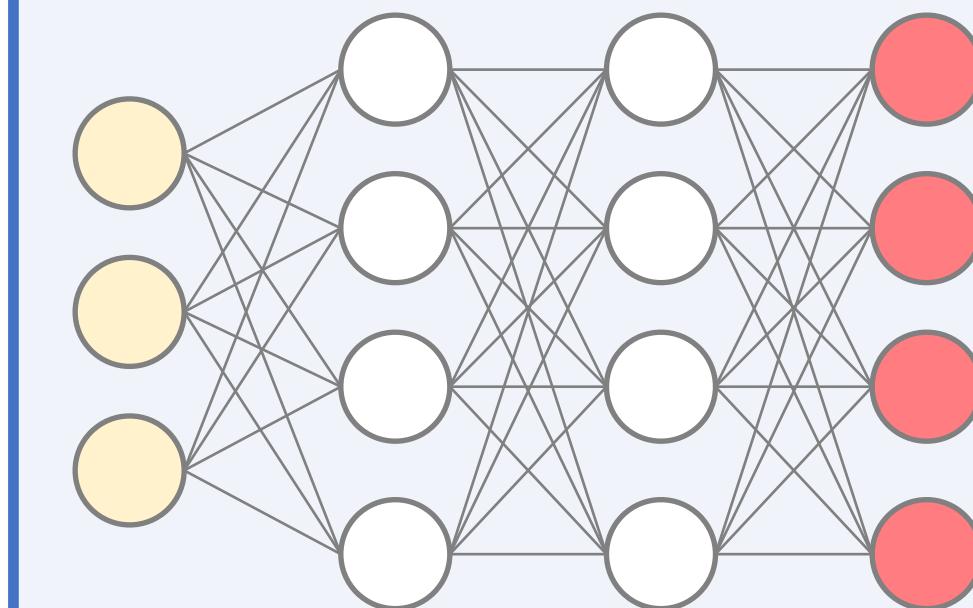


# Neural Fields

- Storage memory **does not grow with spatial resolution or number of spatial dimensions.**
- **Slow** sampling: Each query takes a forward pass through the neural net. Means GPU-memory intensive forward passes.
- Does **not** expose locality: Can't identify set of parameters / direction in parameter space that encodes particular spatial location.
- **Inconvenient processing:** cannot run convolutions.
- But: automatic adaptive resolution. Over optimization, will converge to assign more compute to higher-frequency areas.

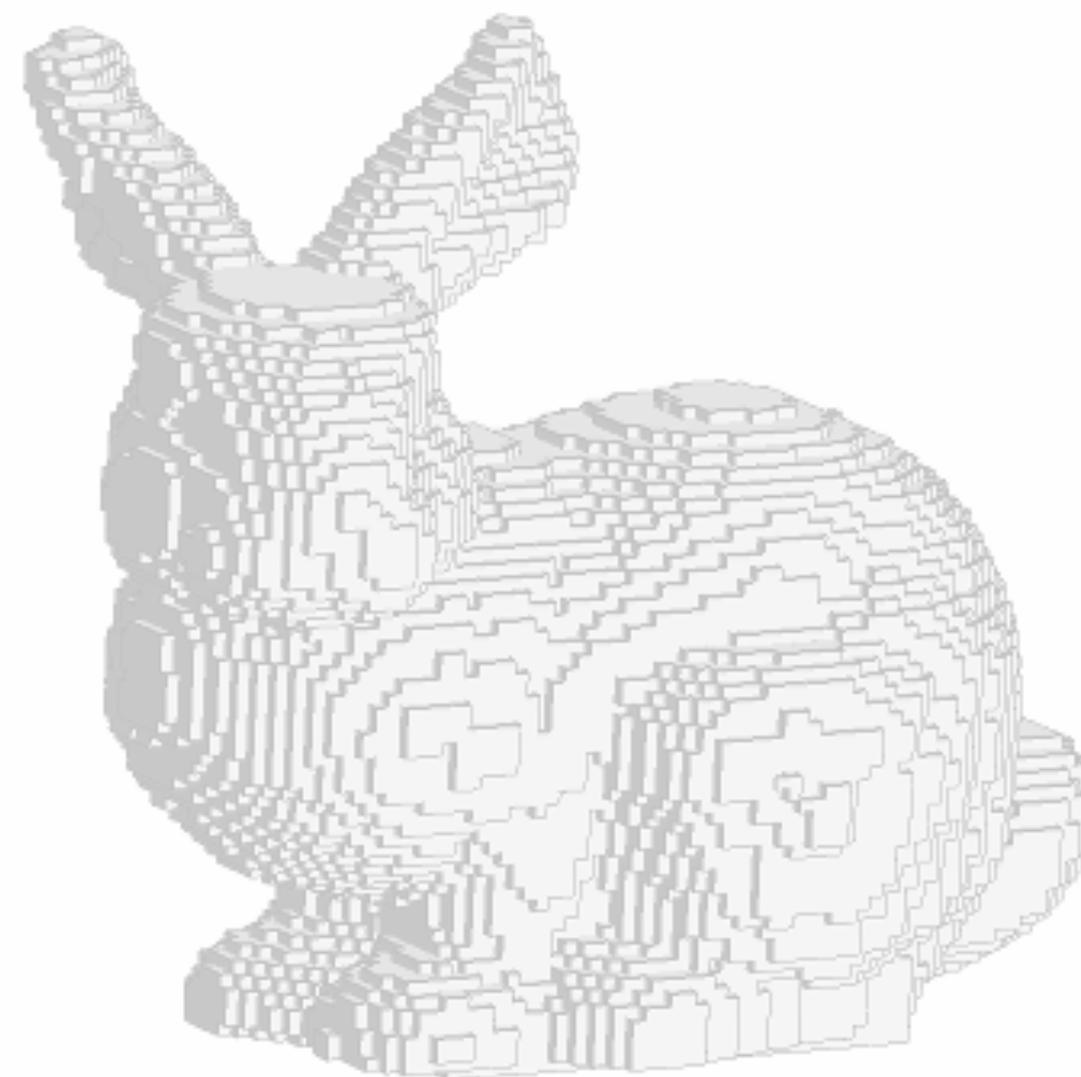
Neural Fields

$$\Phi : \mathbb{R}^3 \rightarrow \mathbb{R}^n$$

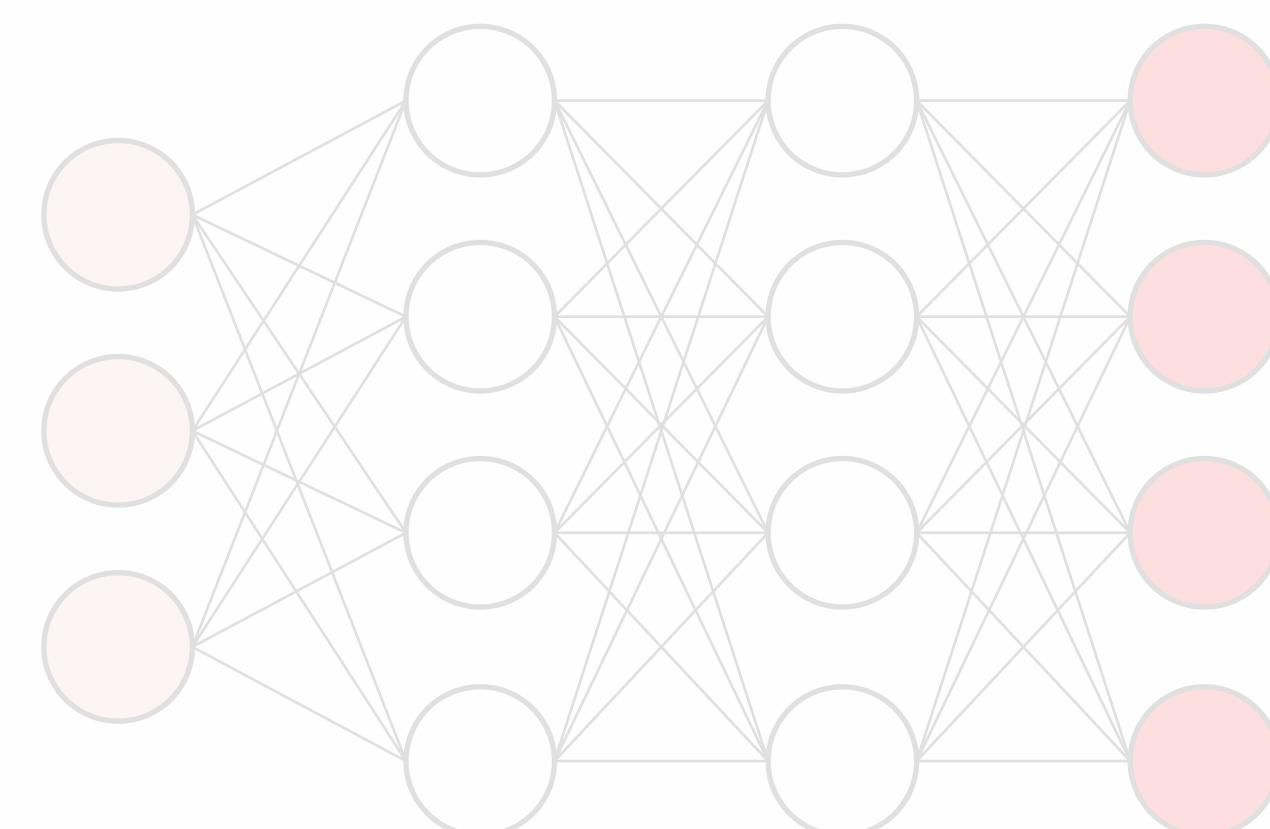


Note: Despite their name, has nothing to do with Machine Learning itself.  
It's as “smart” as a voxelgrid.

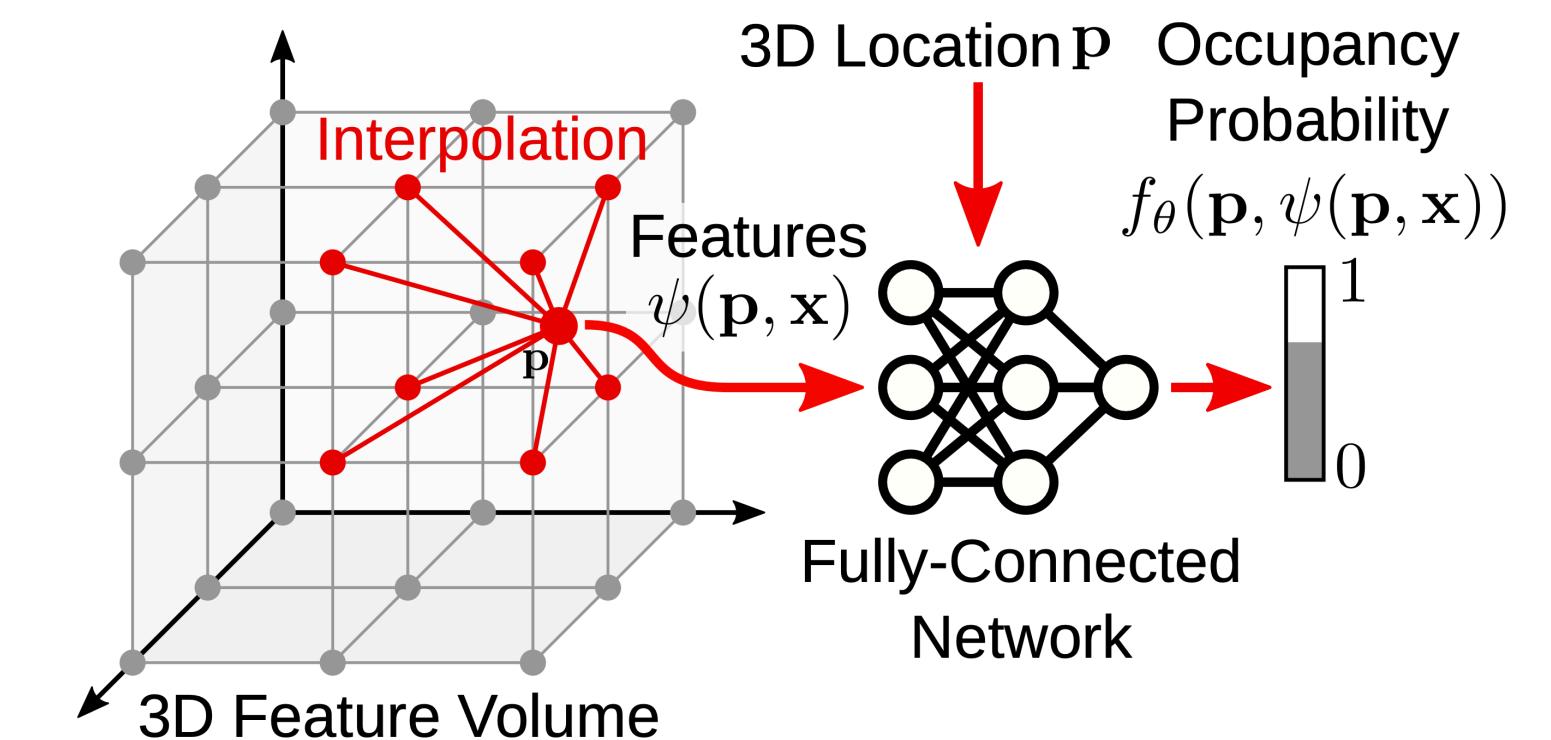
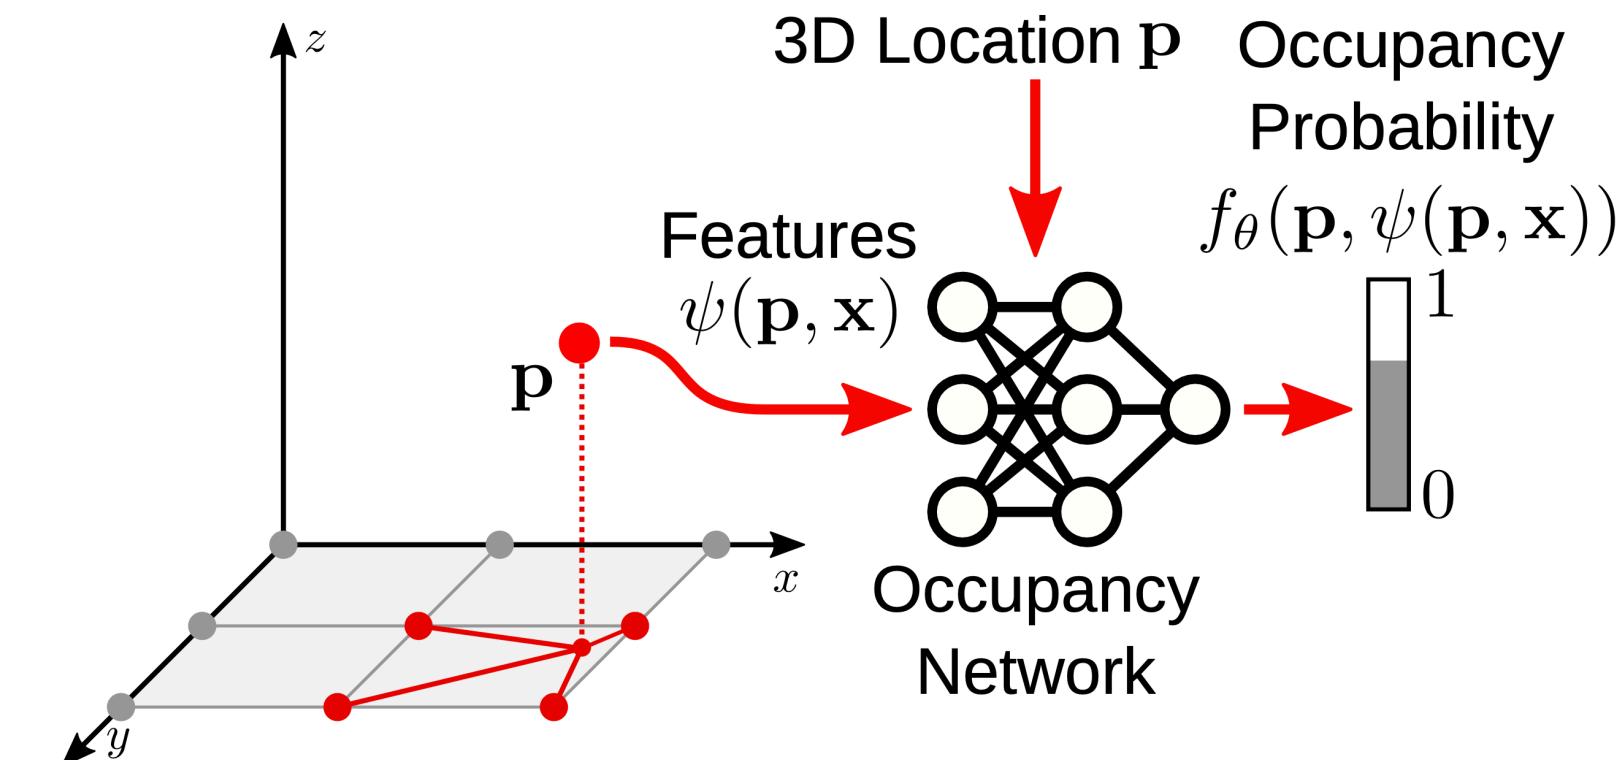
# Hybrid discrete / continuous reps



Discrete Parameterizations

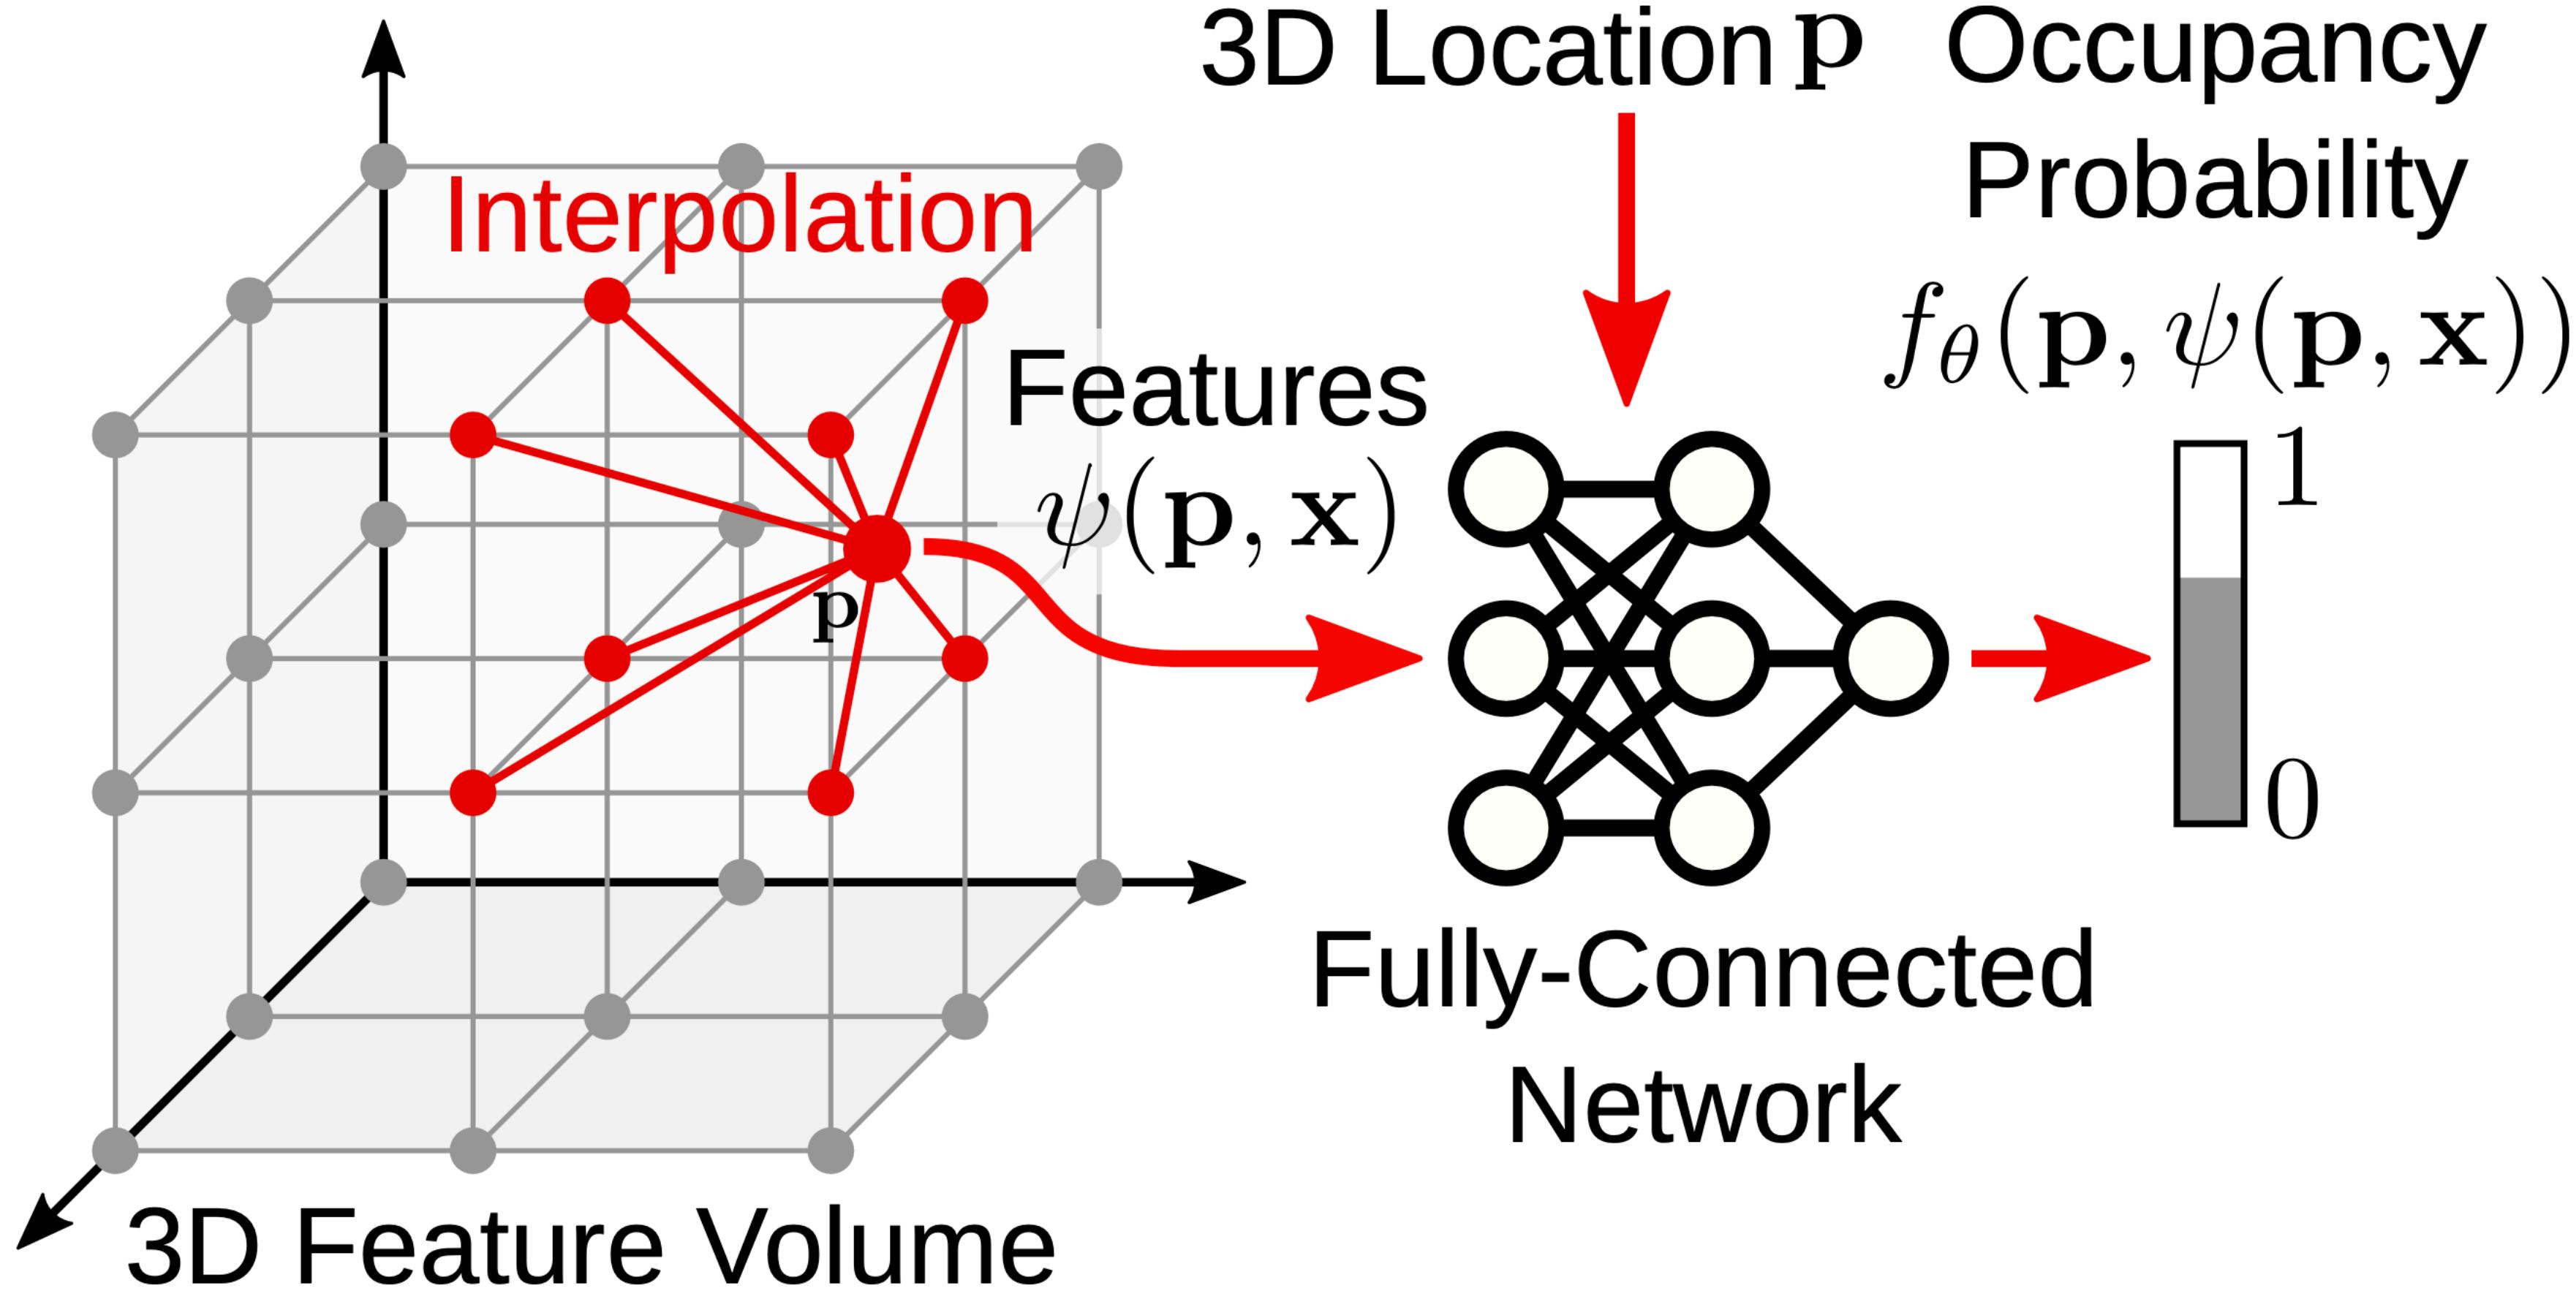


Continuous Parameterizations



Hybrid Parameterizations

# Basic idea: Use Neural Net as Interpolation Kernel



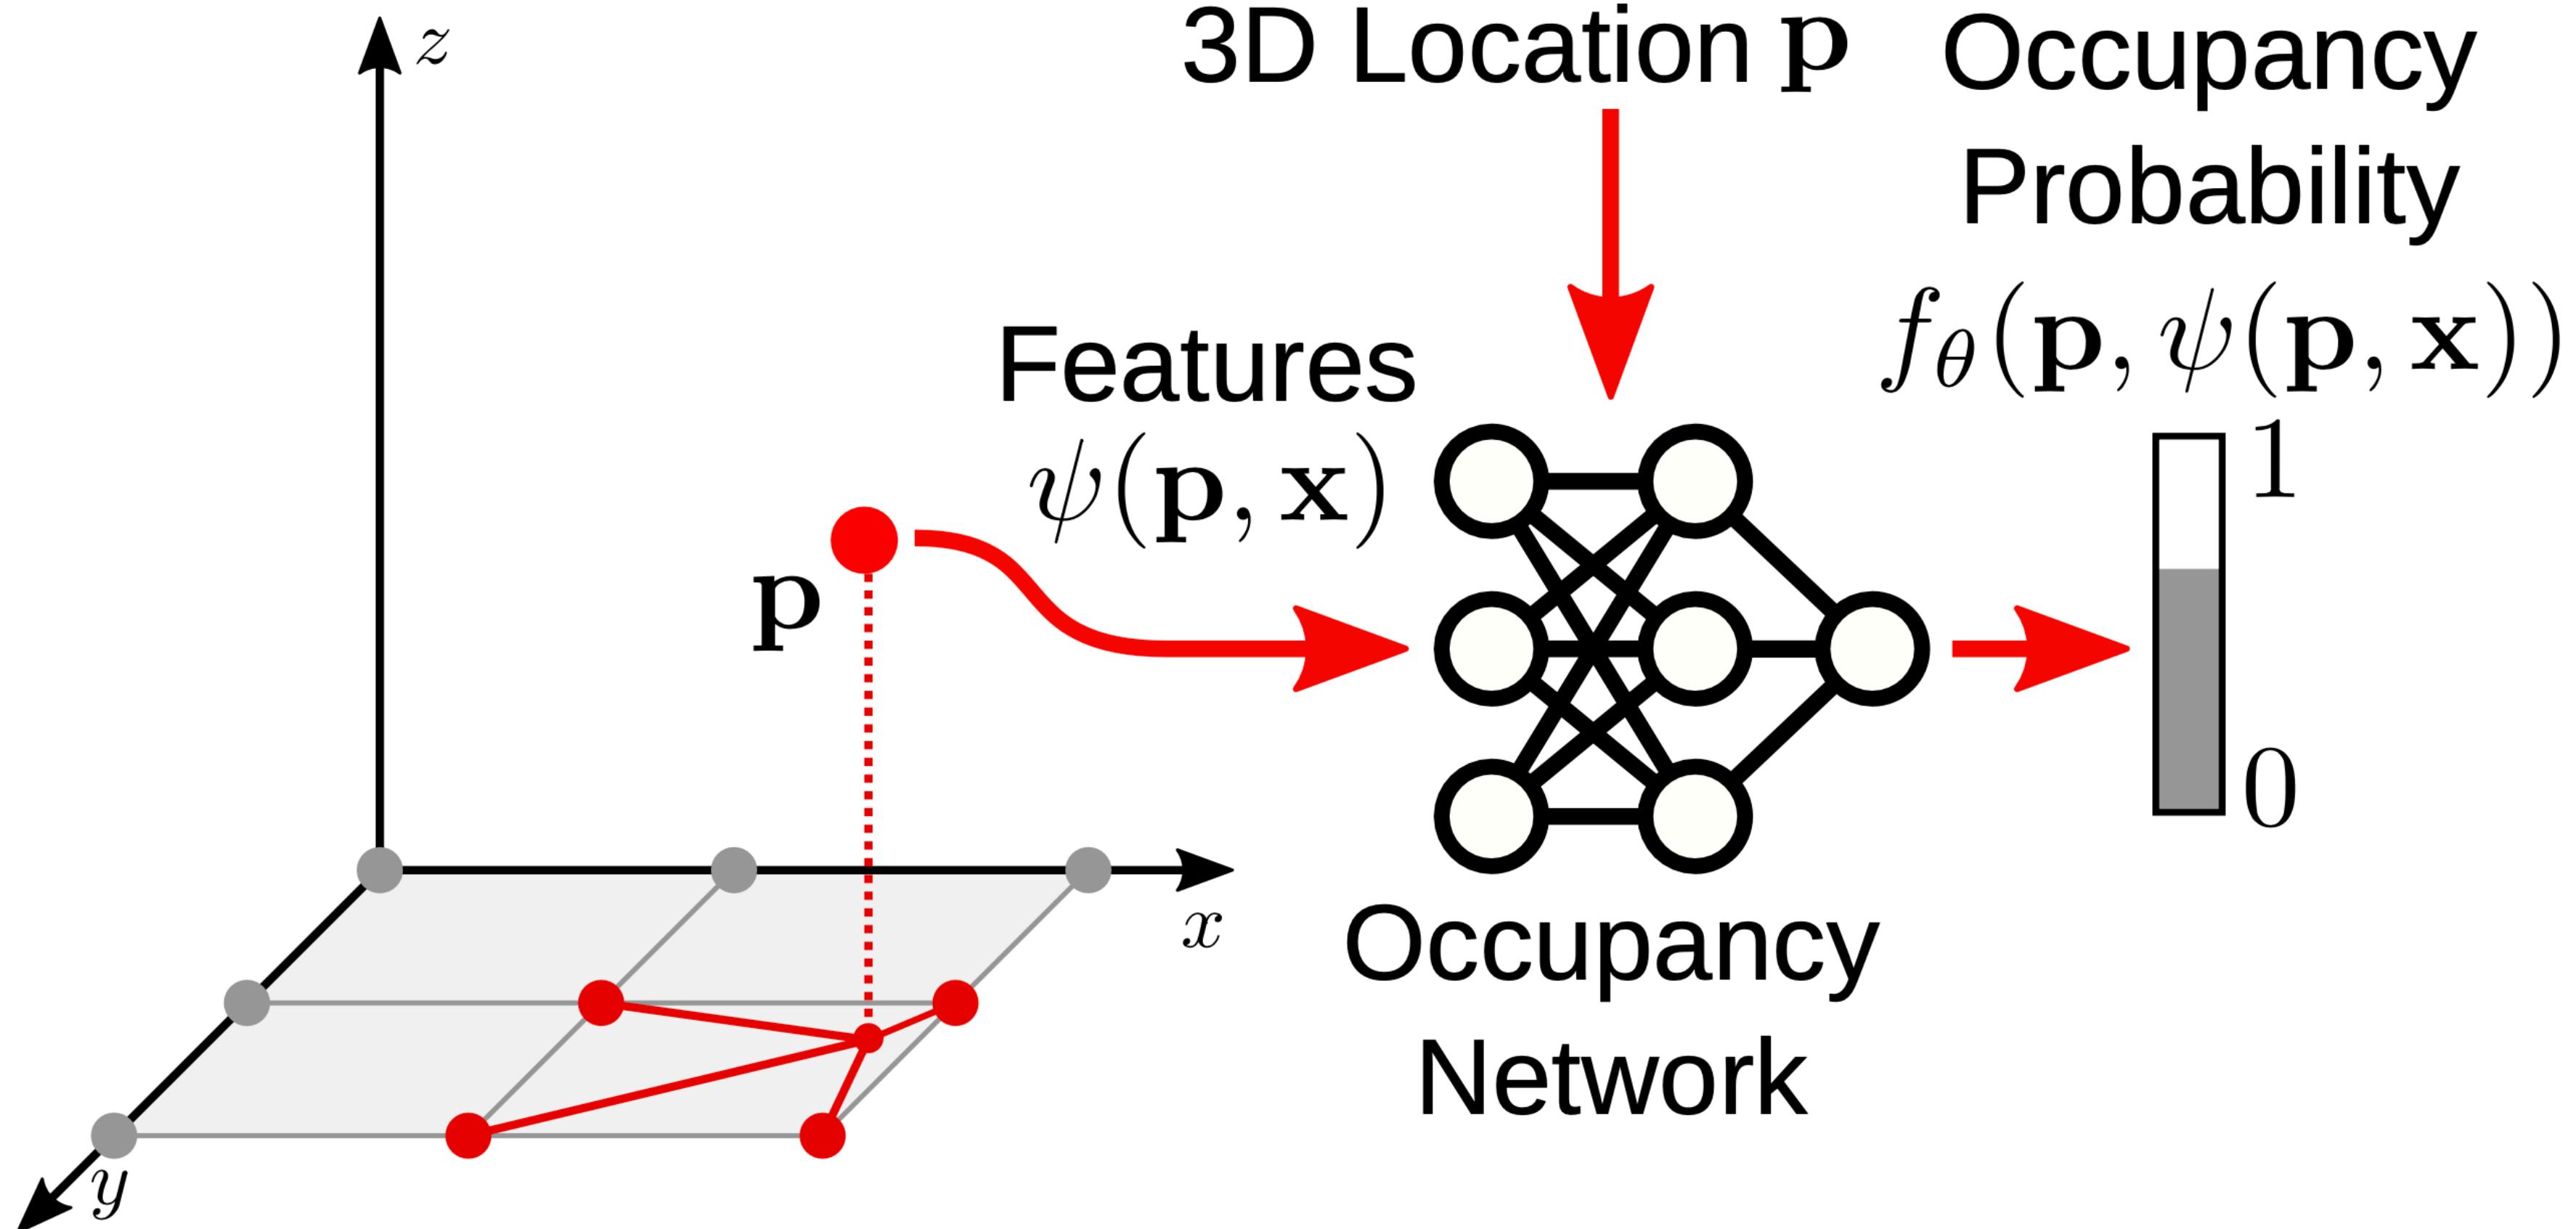
Convolutional Occupancy Networks [Peng et al. 2020]

Local Implicit Grid Representations for 3D Scenes [Jiang et al. 2020]

Implicit Functions in Feature Space for 3D Shape Reconstruction and Completion [Chibane et al. 2020]

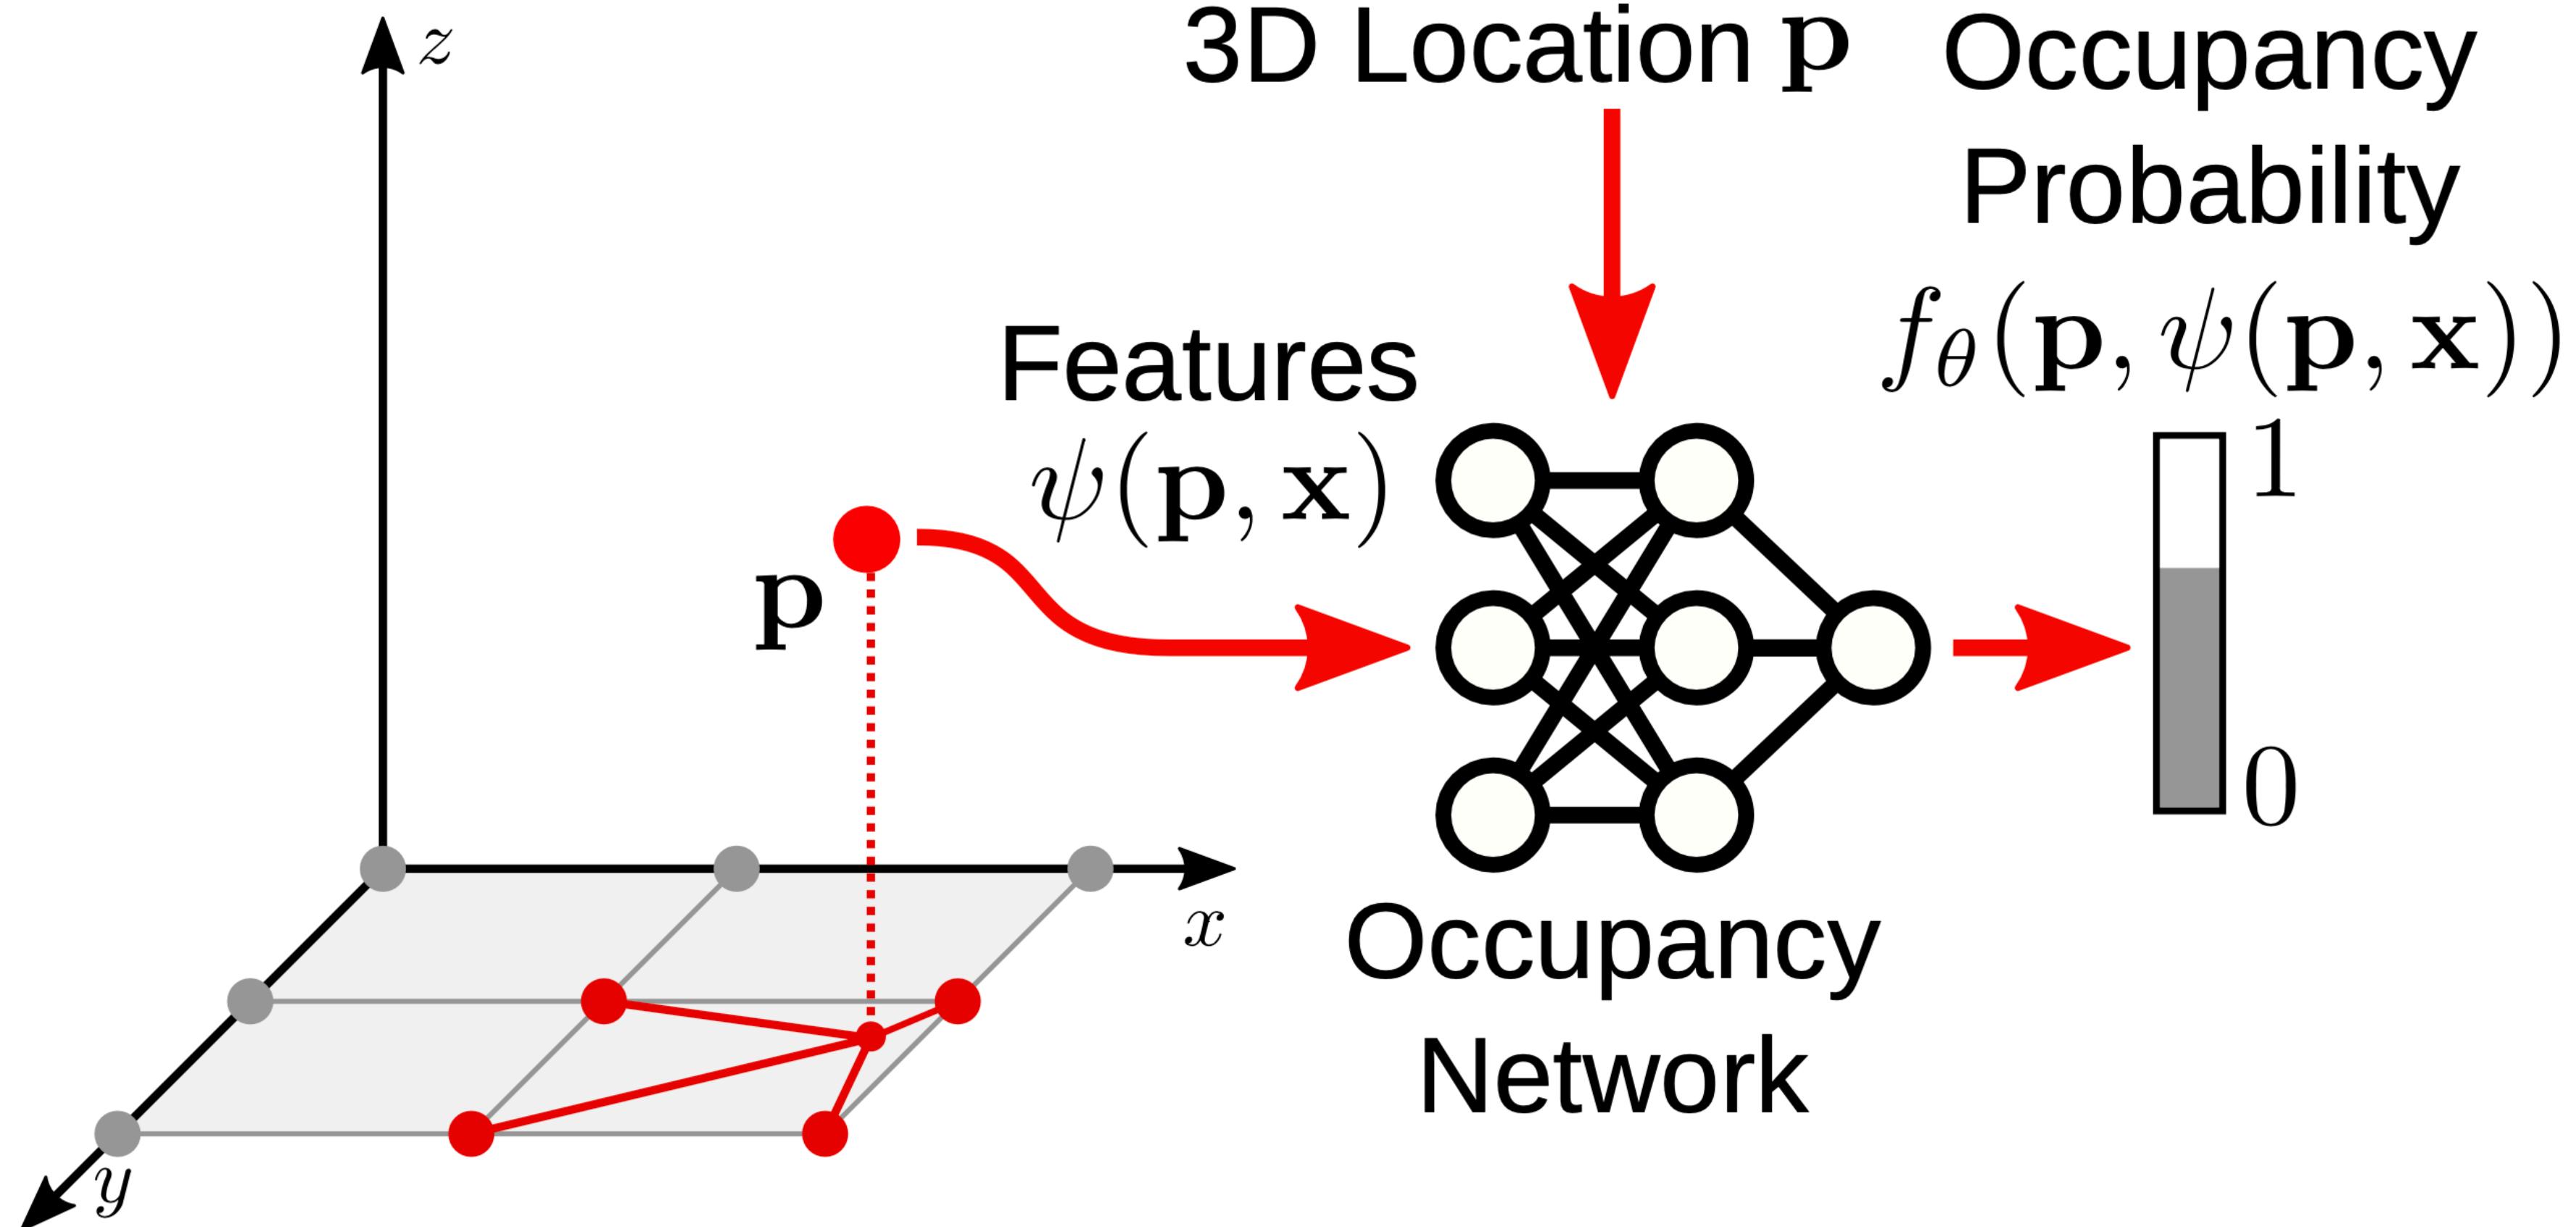
Deep Local Shapes: Learning Local SDF Priors for Detailed 3D Reconstruction [Chabra et al. 2020]

# Ground plan & Orthographic Projection



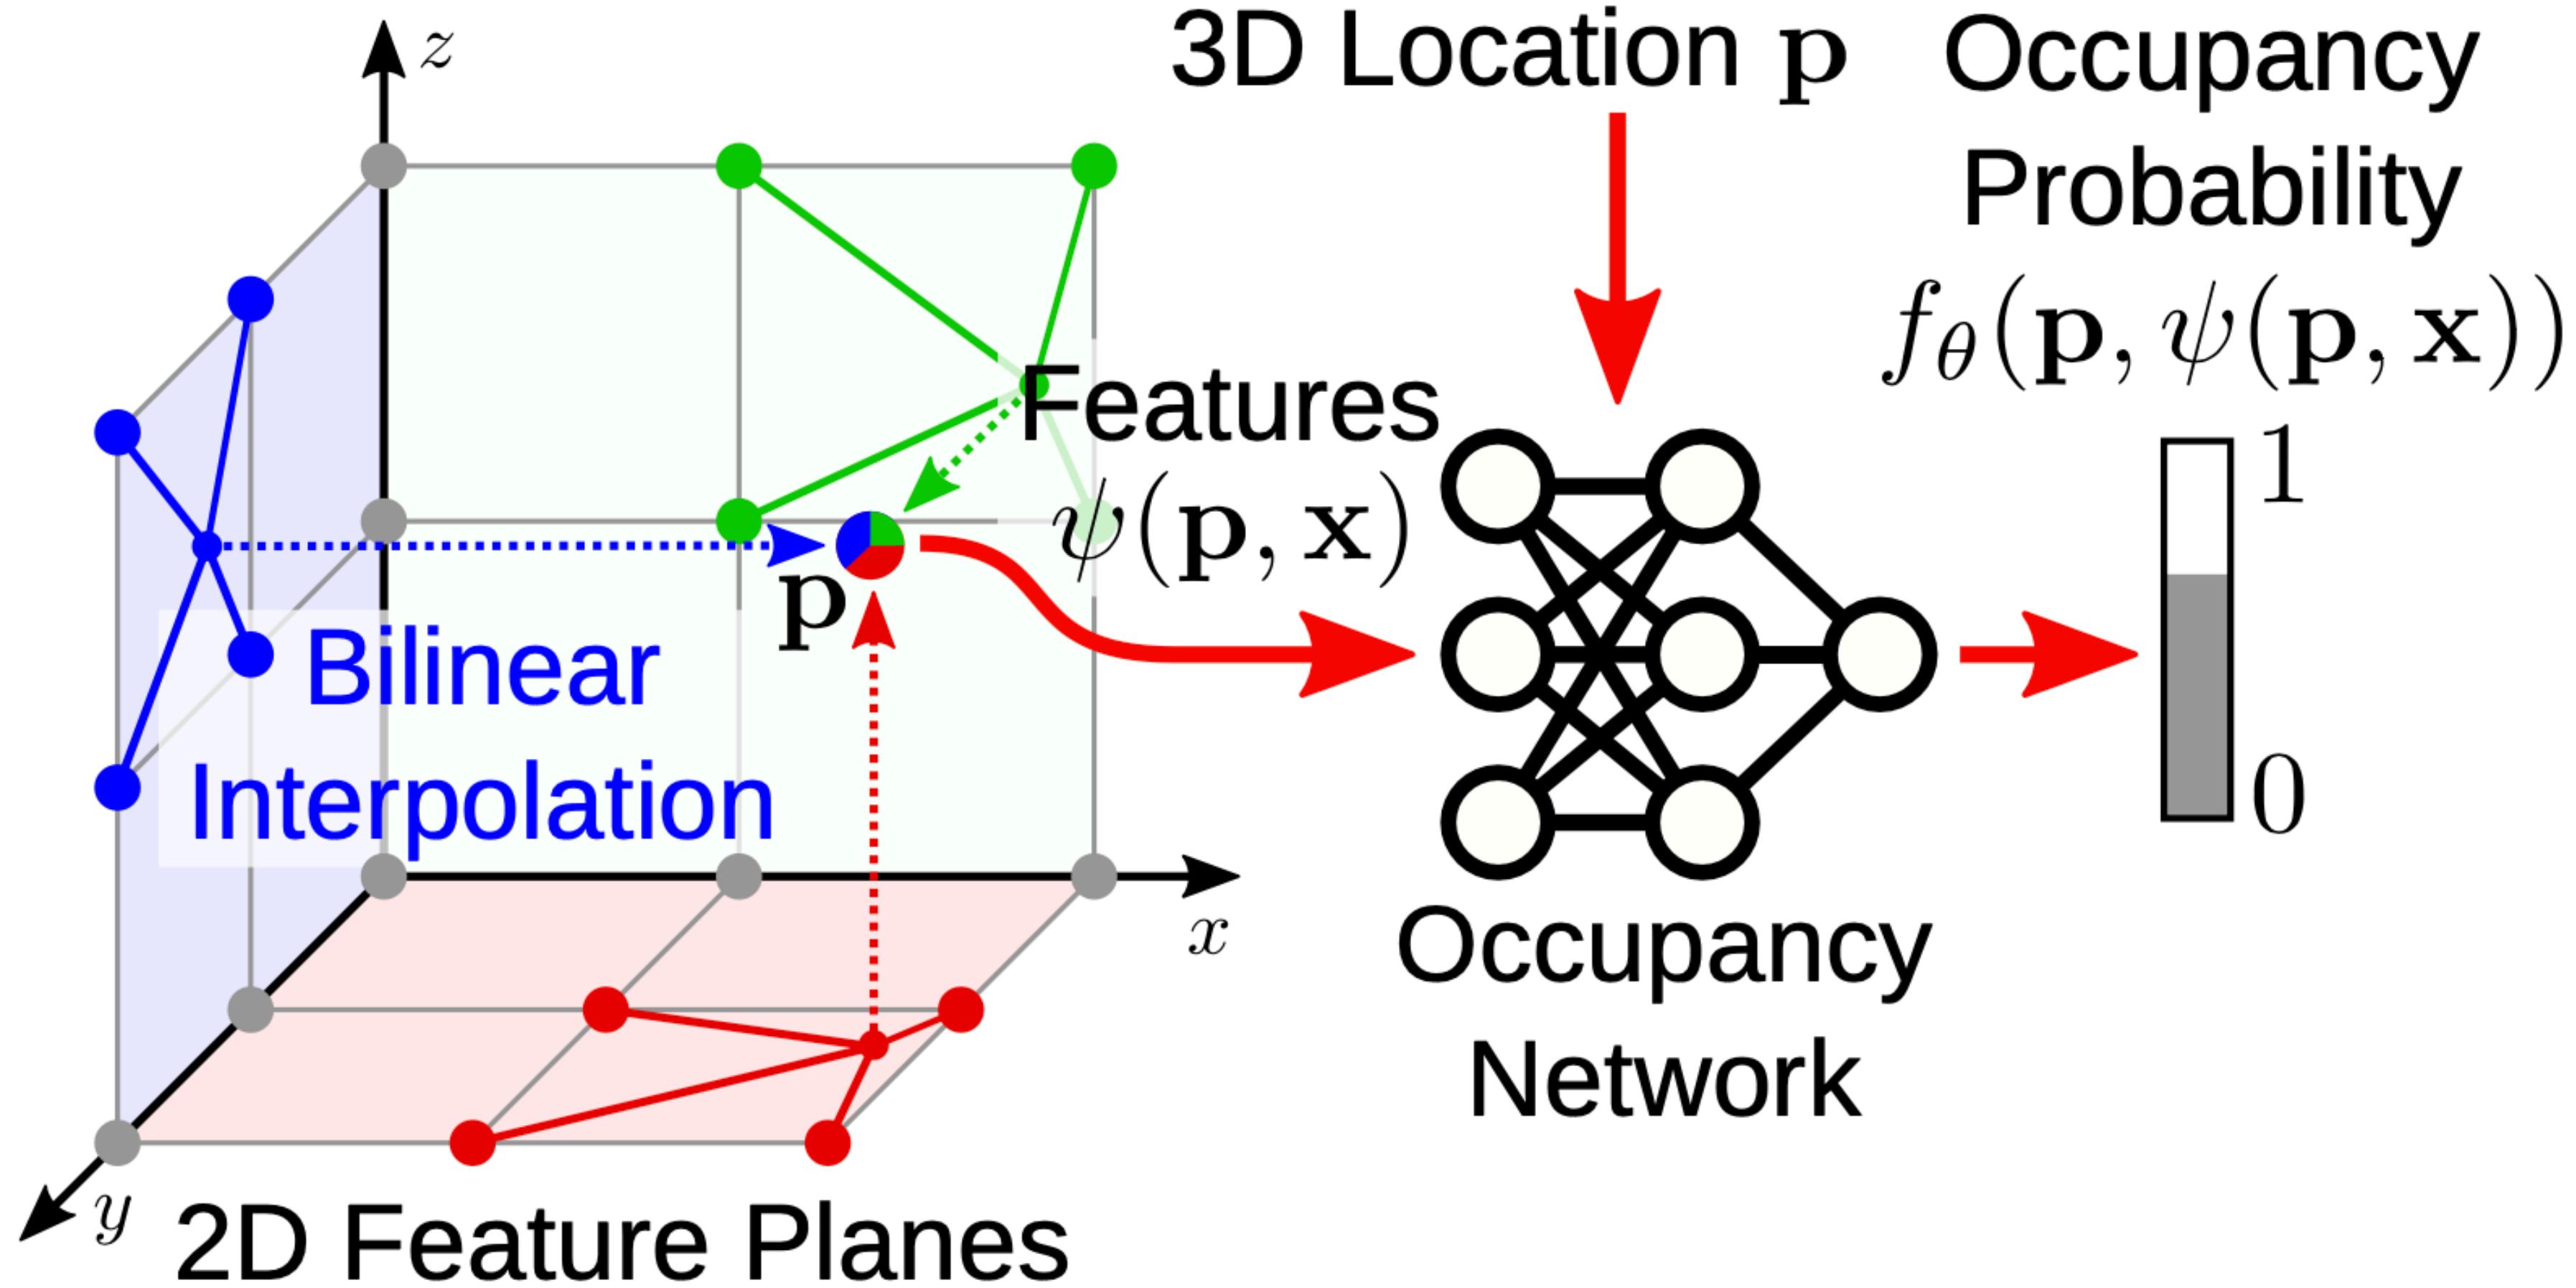
Any assumptions made on scene? Any limitations?

# Ground plan & Orthographic Projection

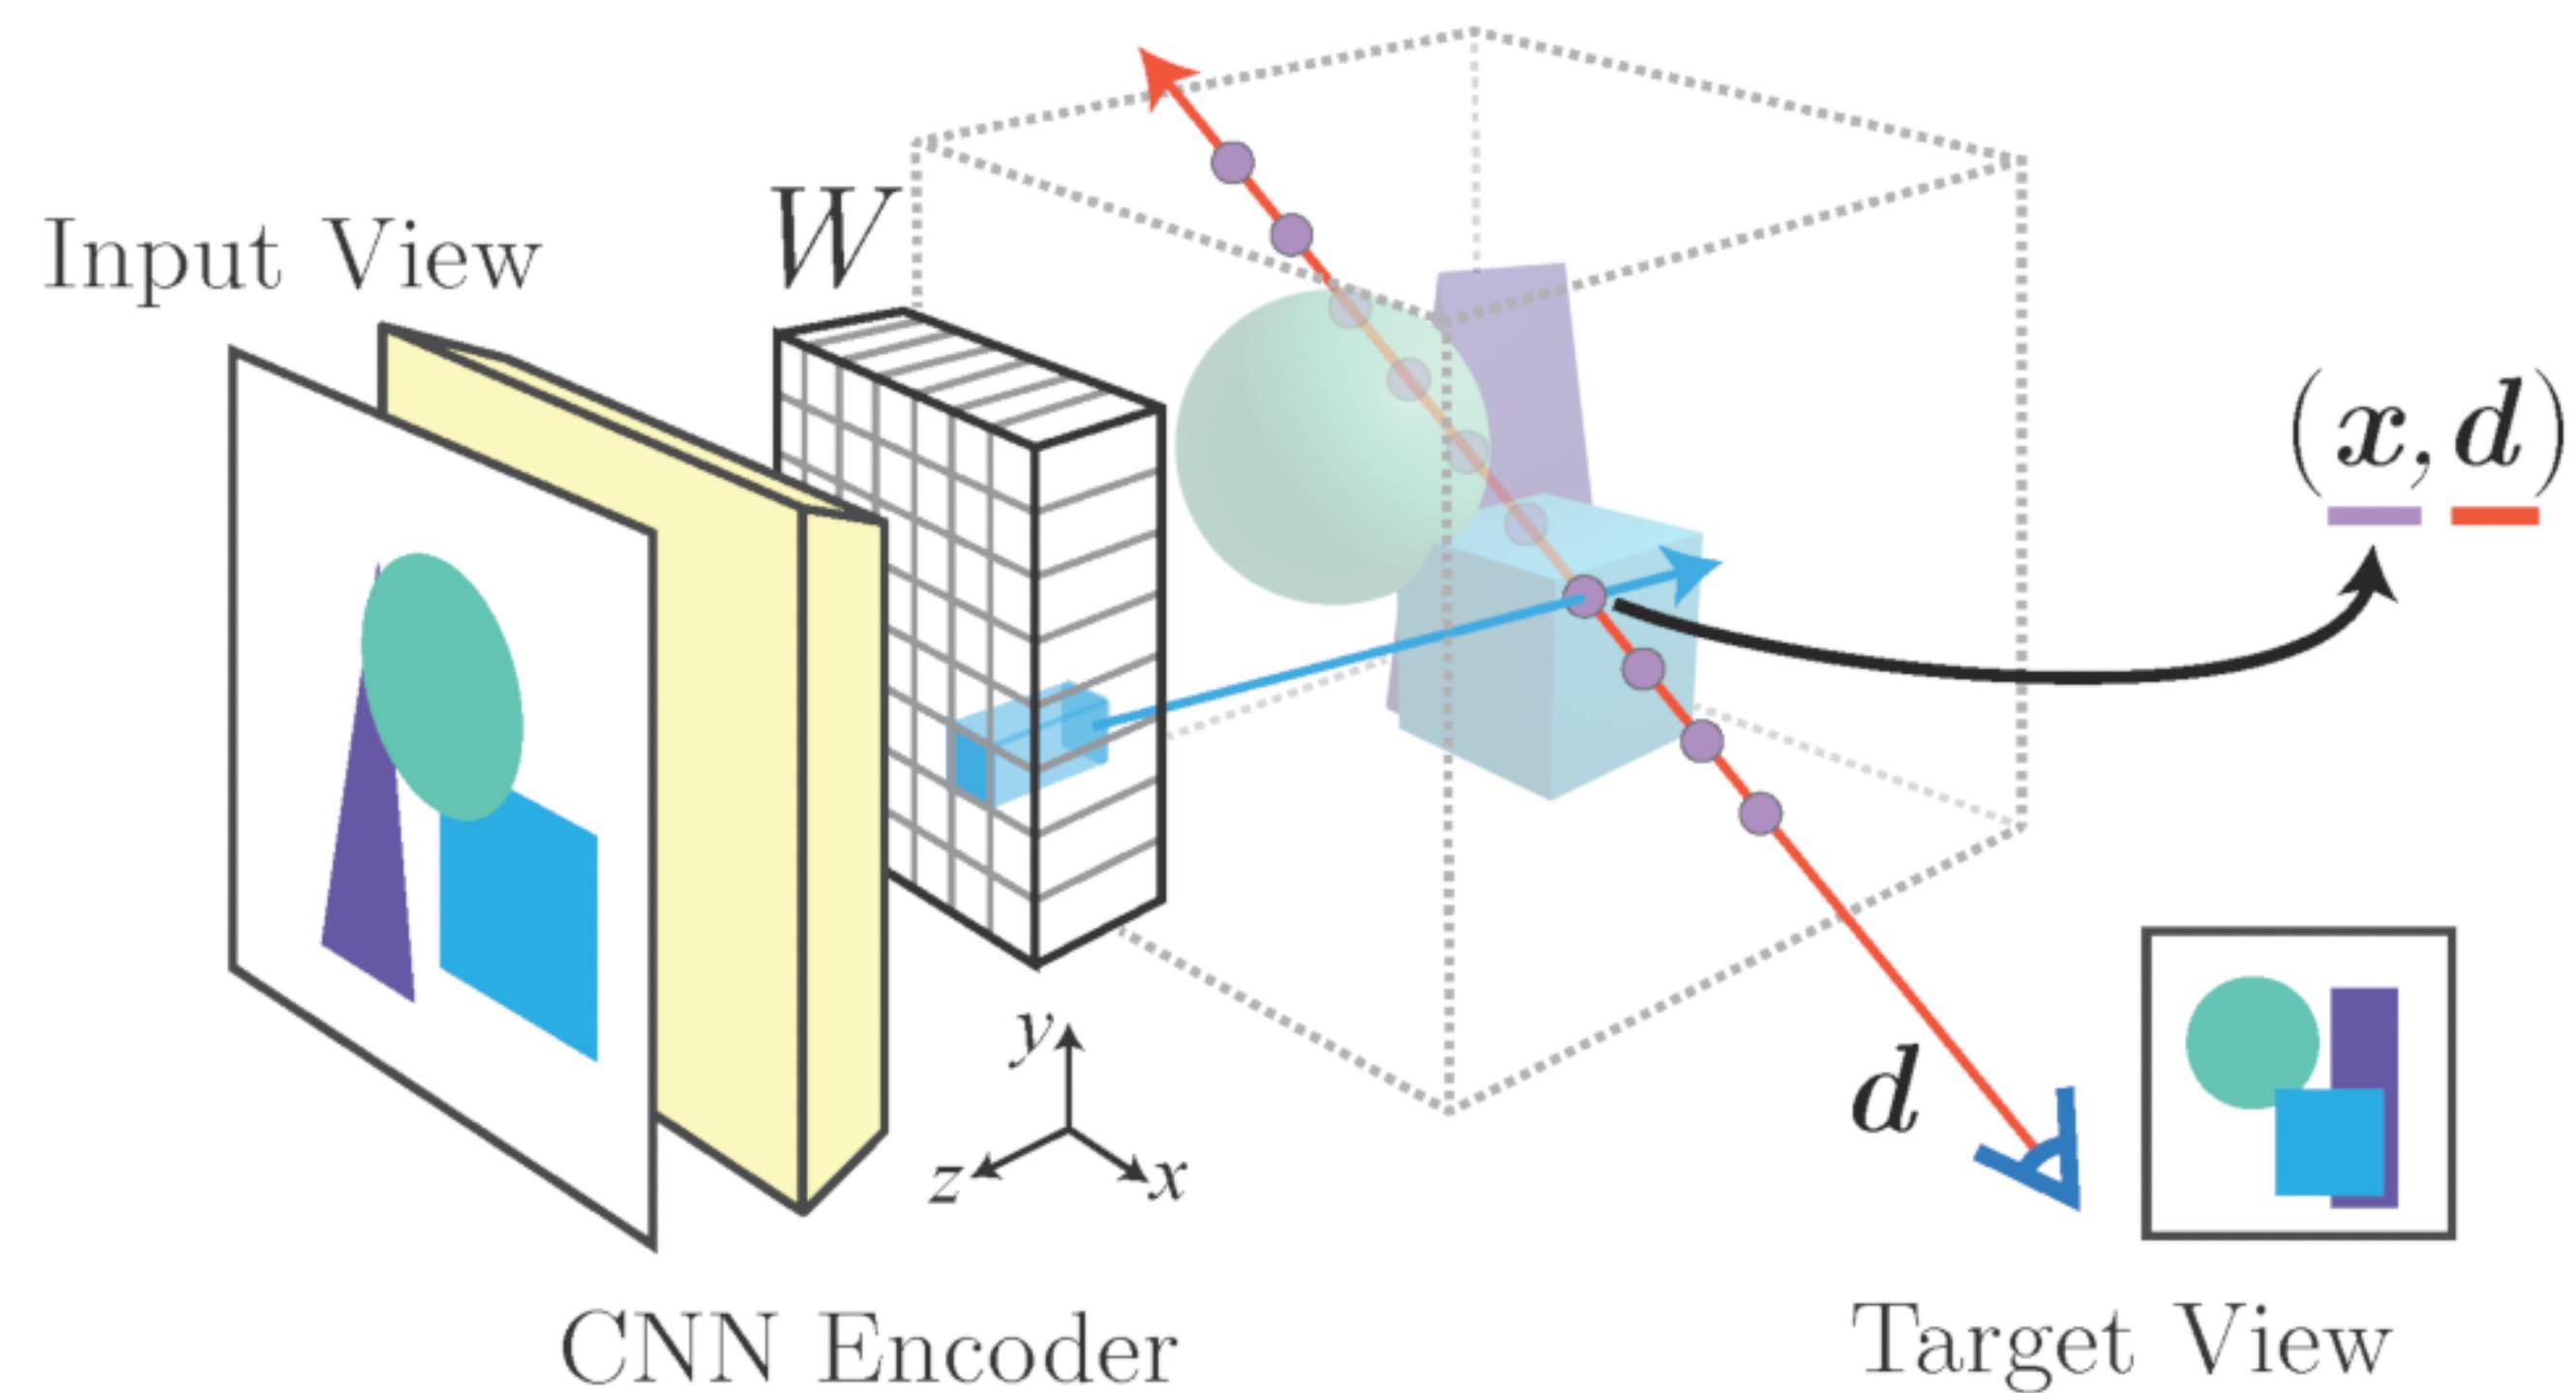


No. You can still represent *any* scene - some info stored in NN, some in grid.

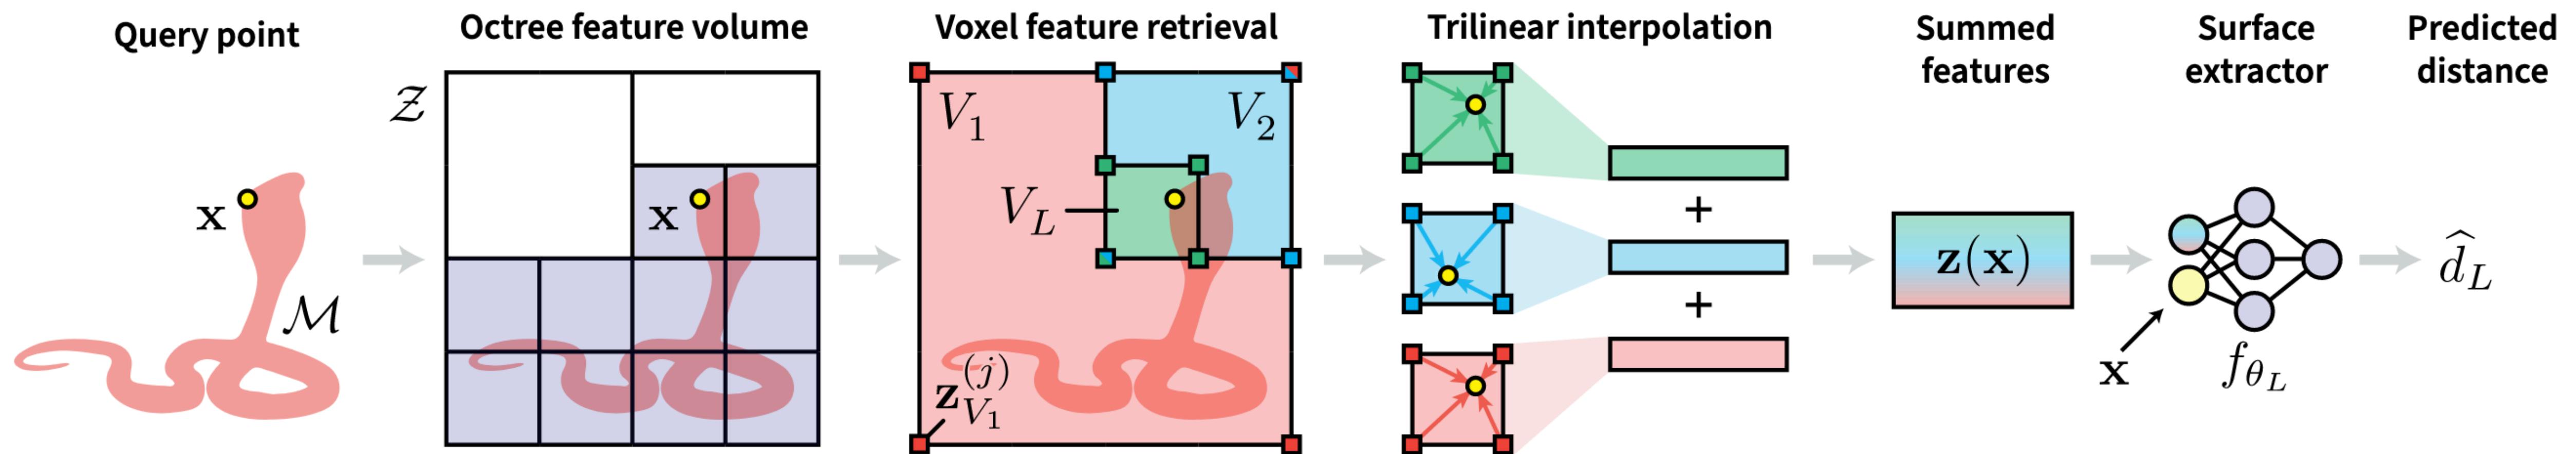
# Tri-Planes and Orthographic Projection



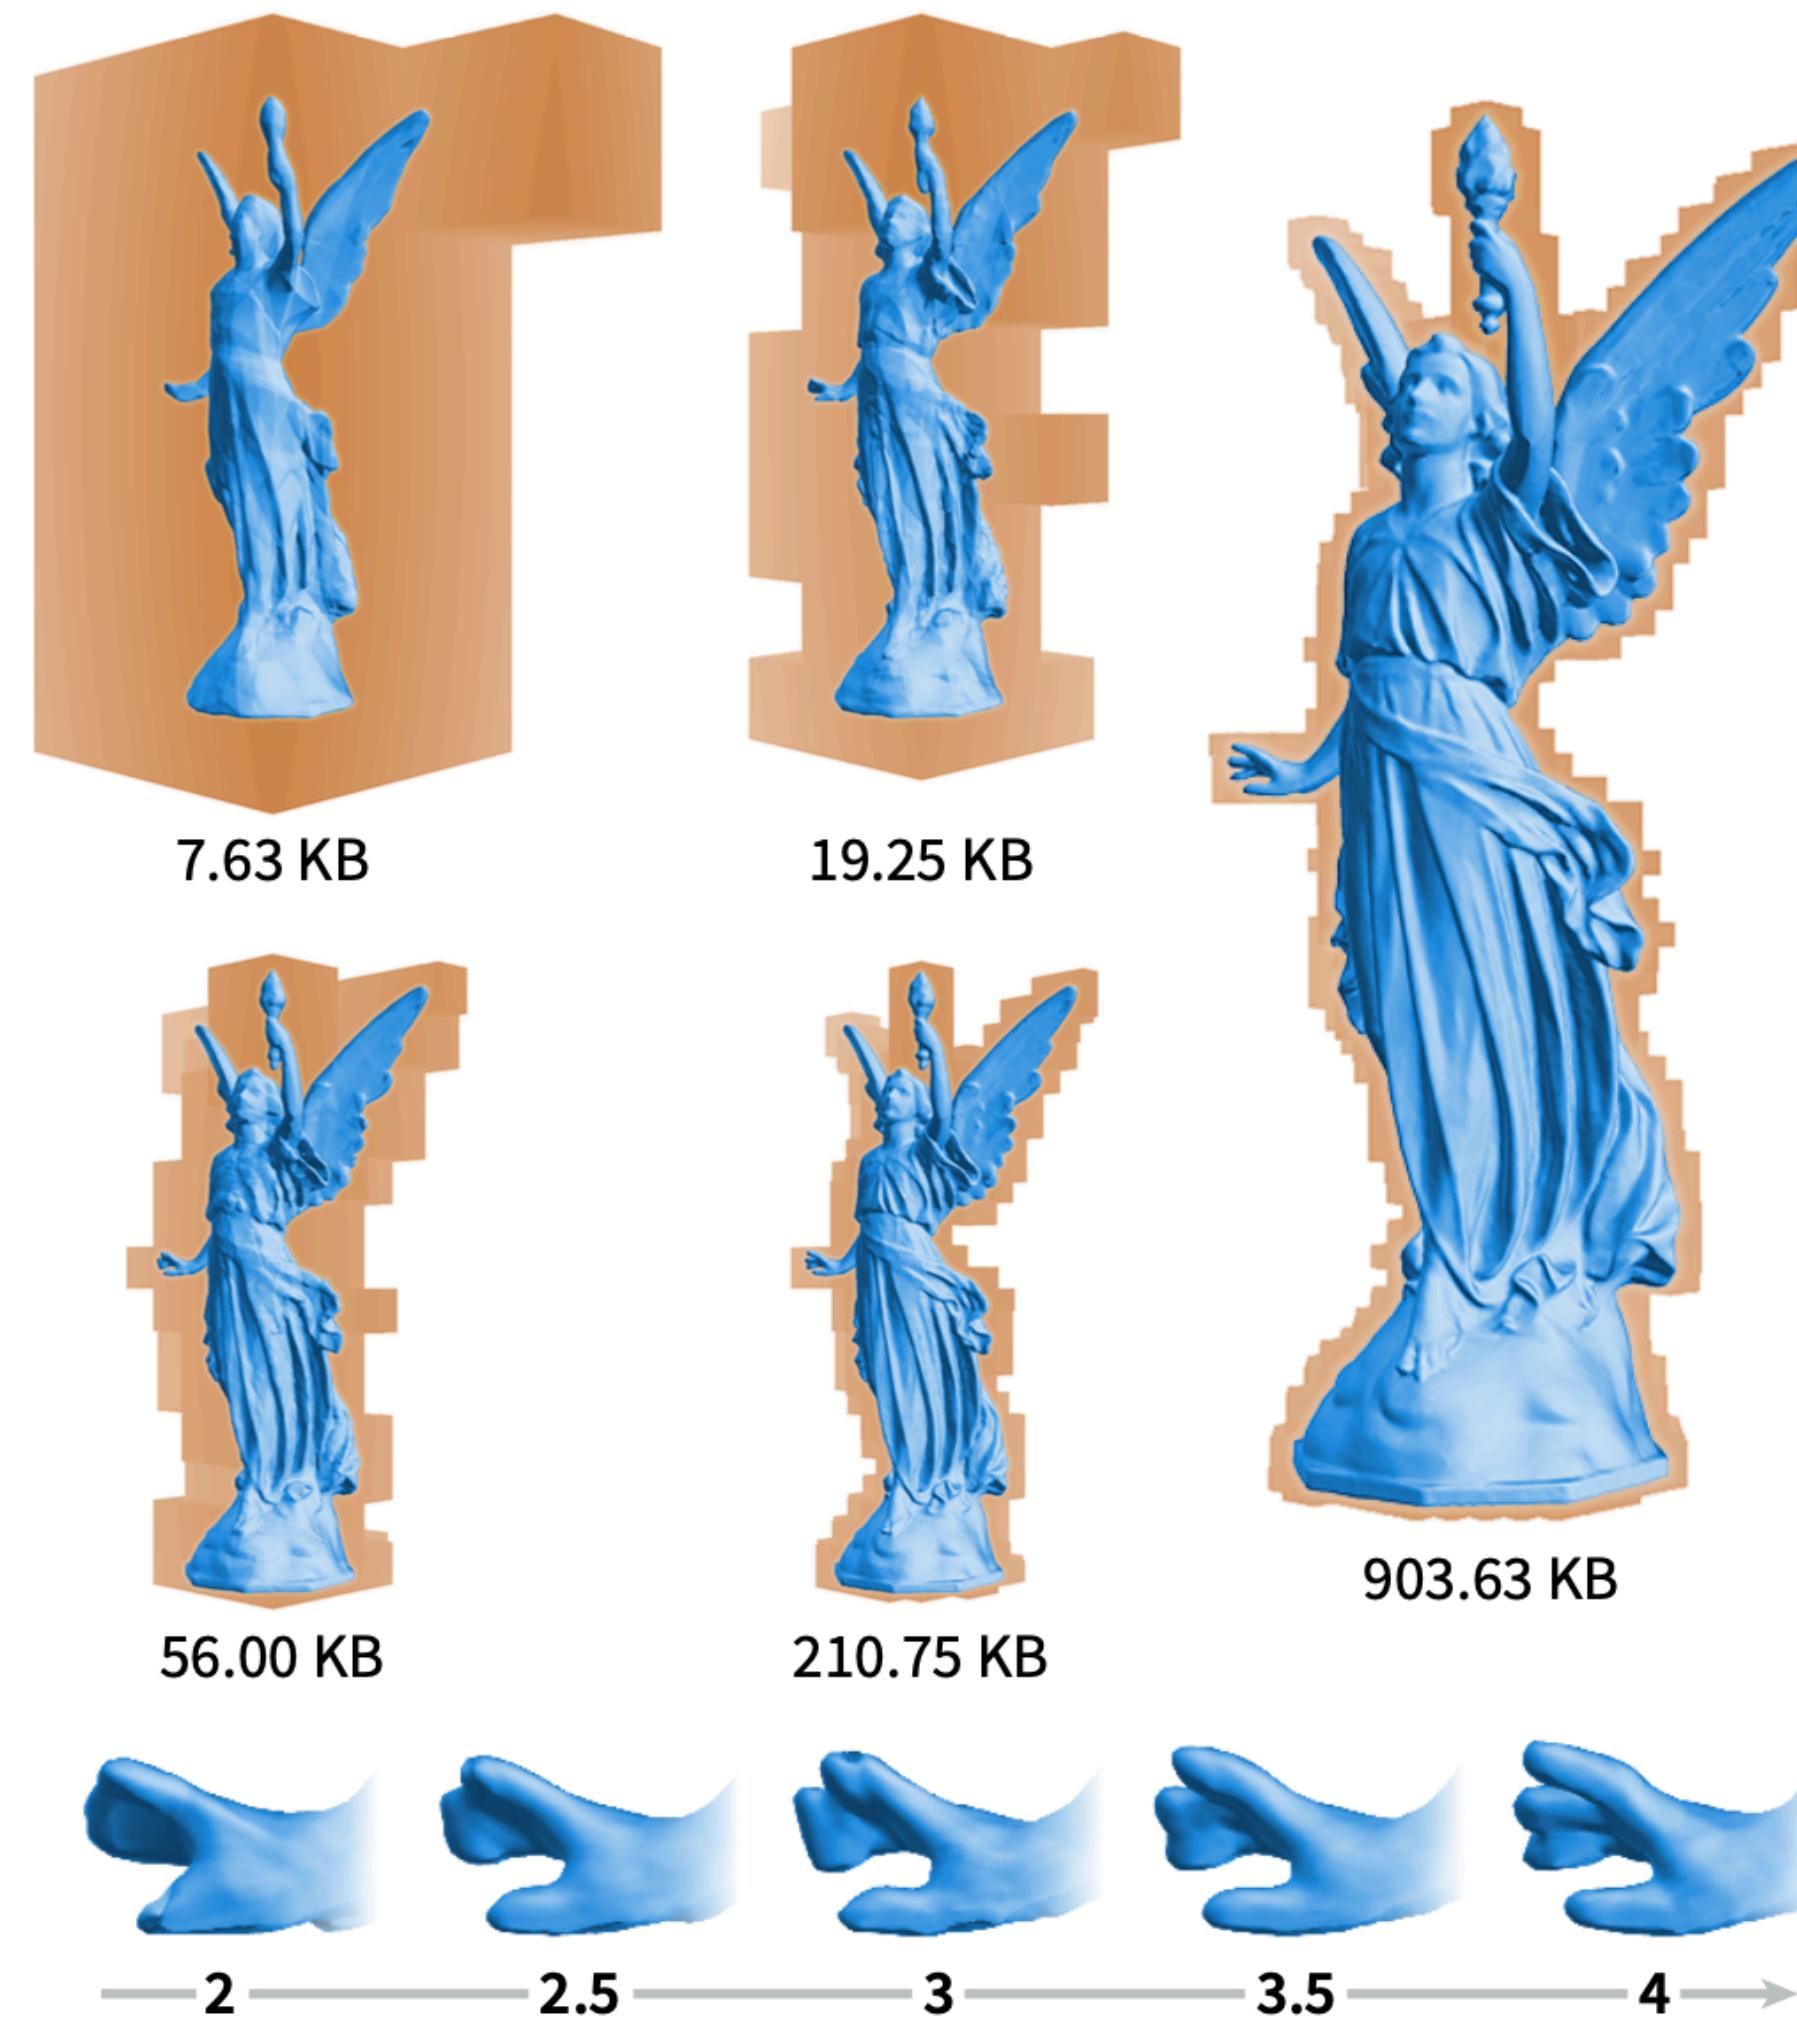
# Image Grids & Perspective Projection



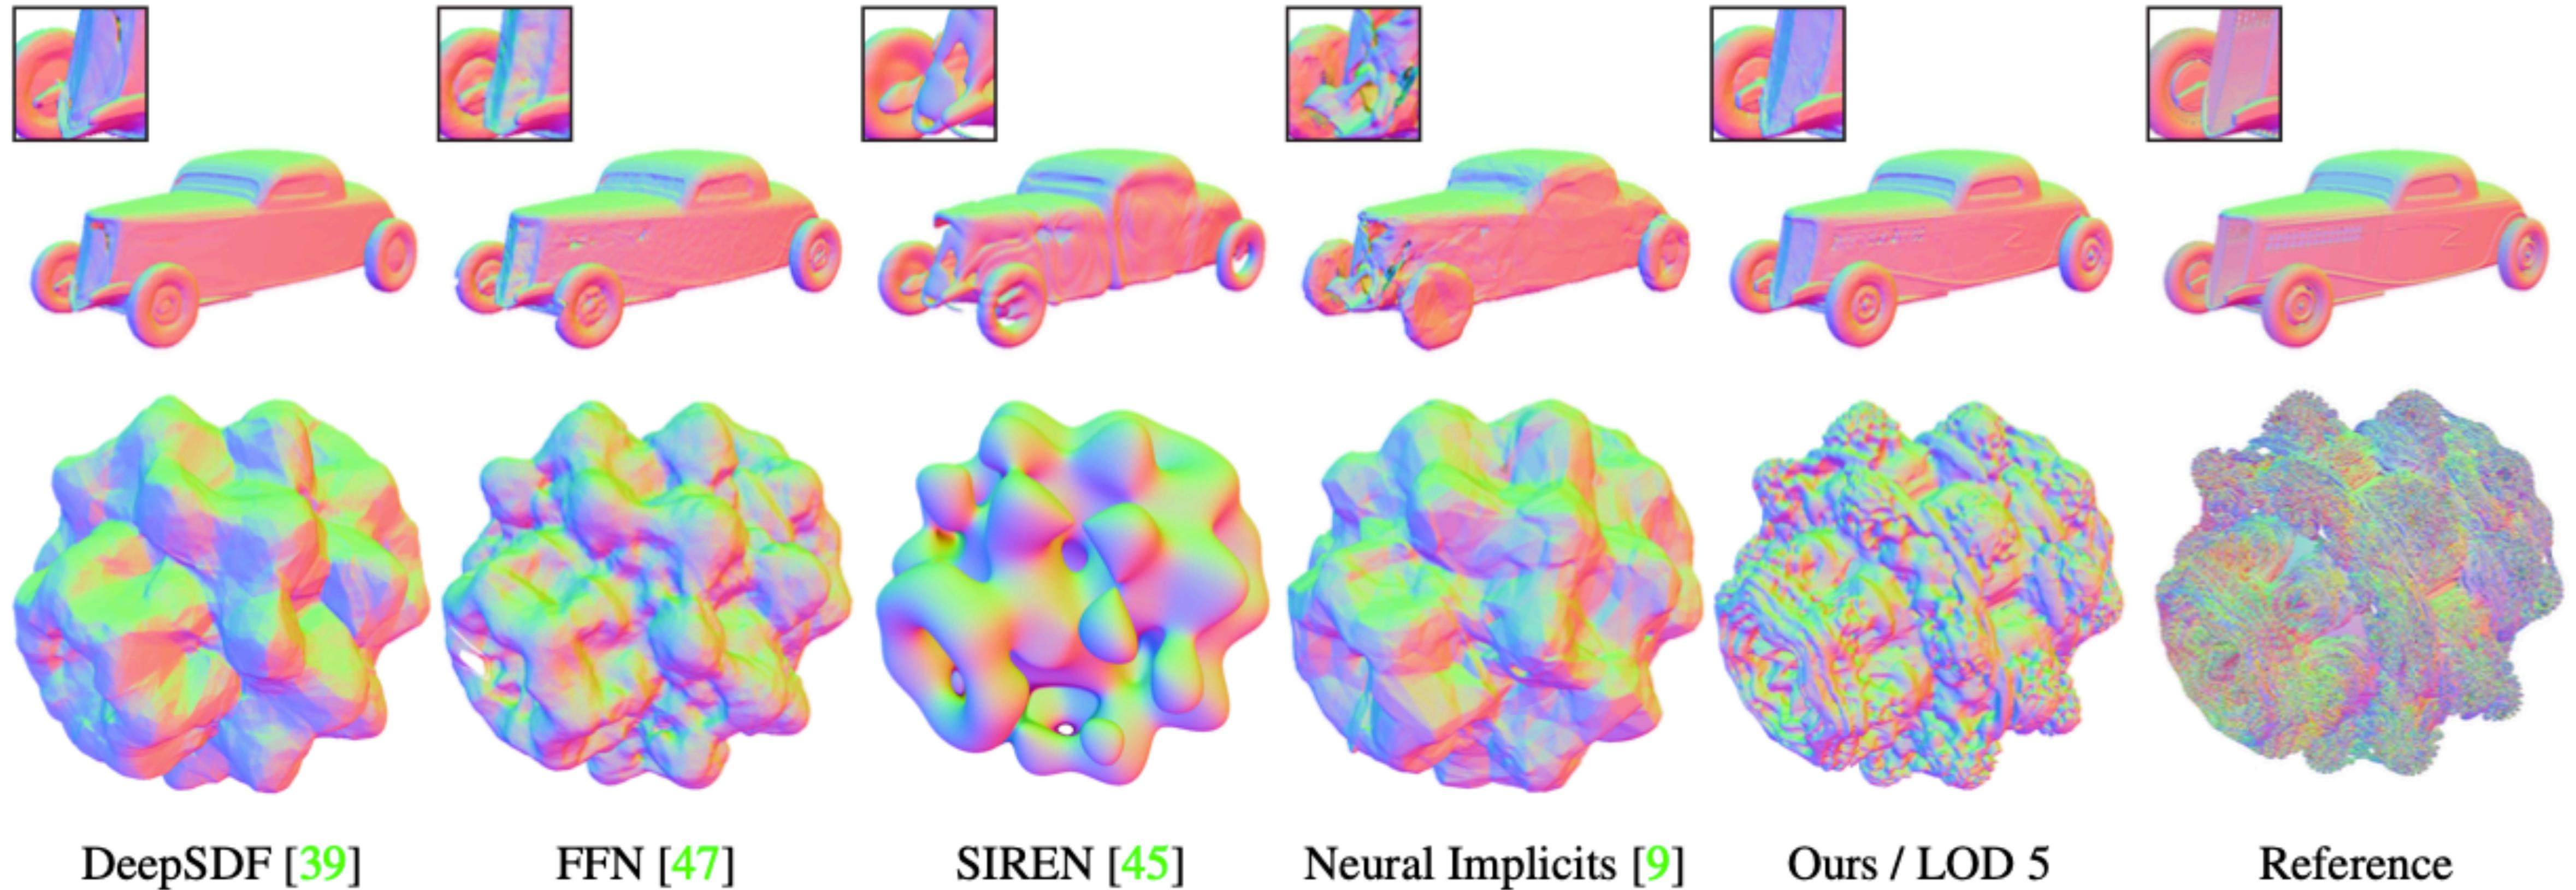
# Multi-scale Voxel Grids



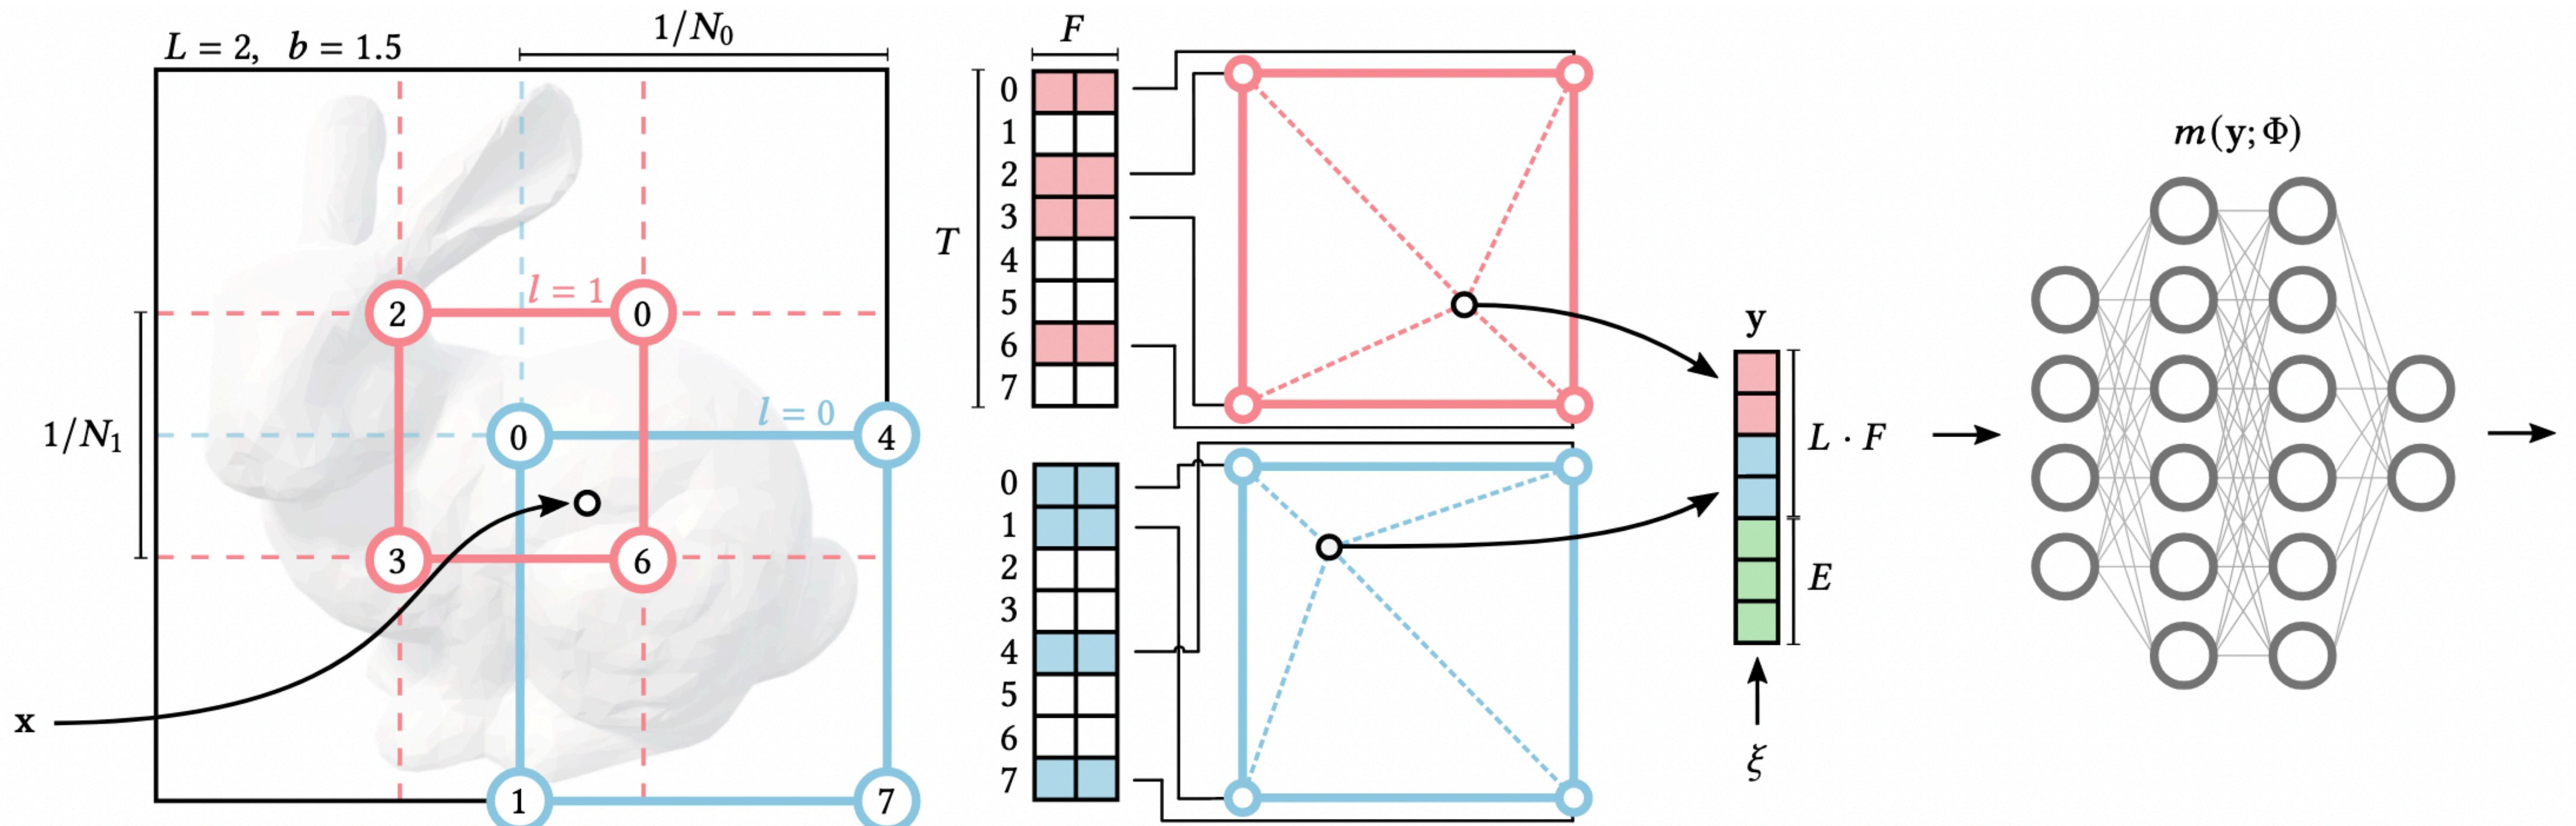
# Multi-scale Voxel Grids



# Multi-scale Voxel Grids



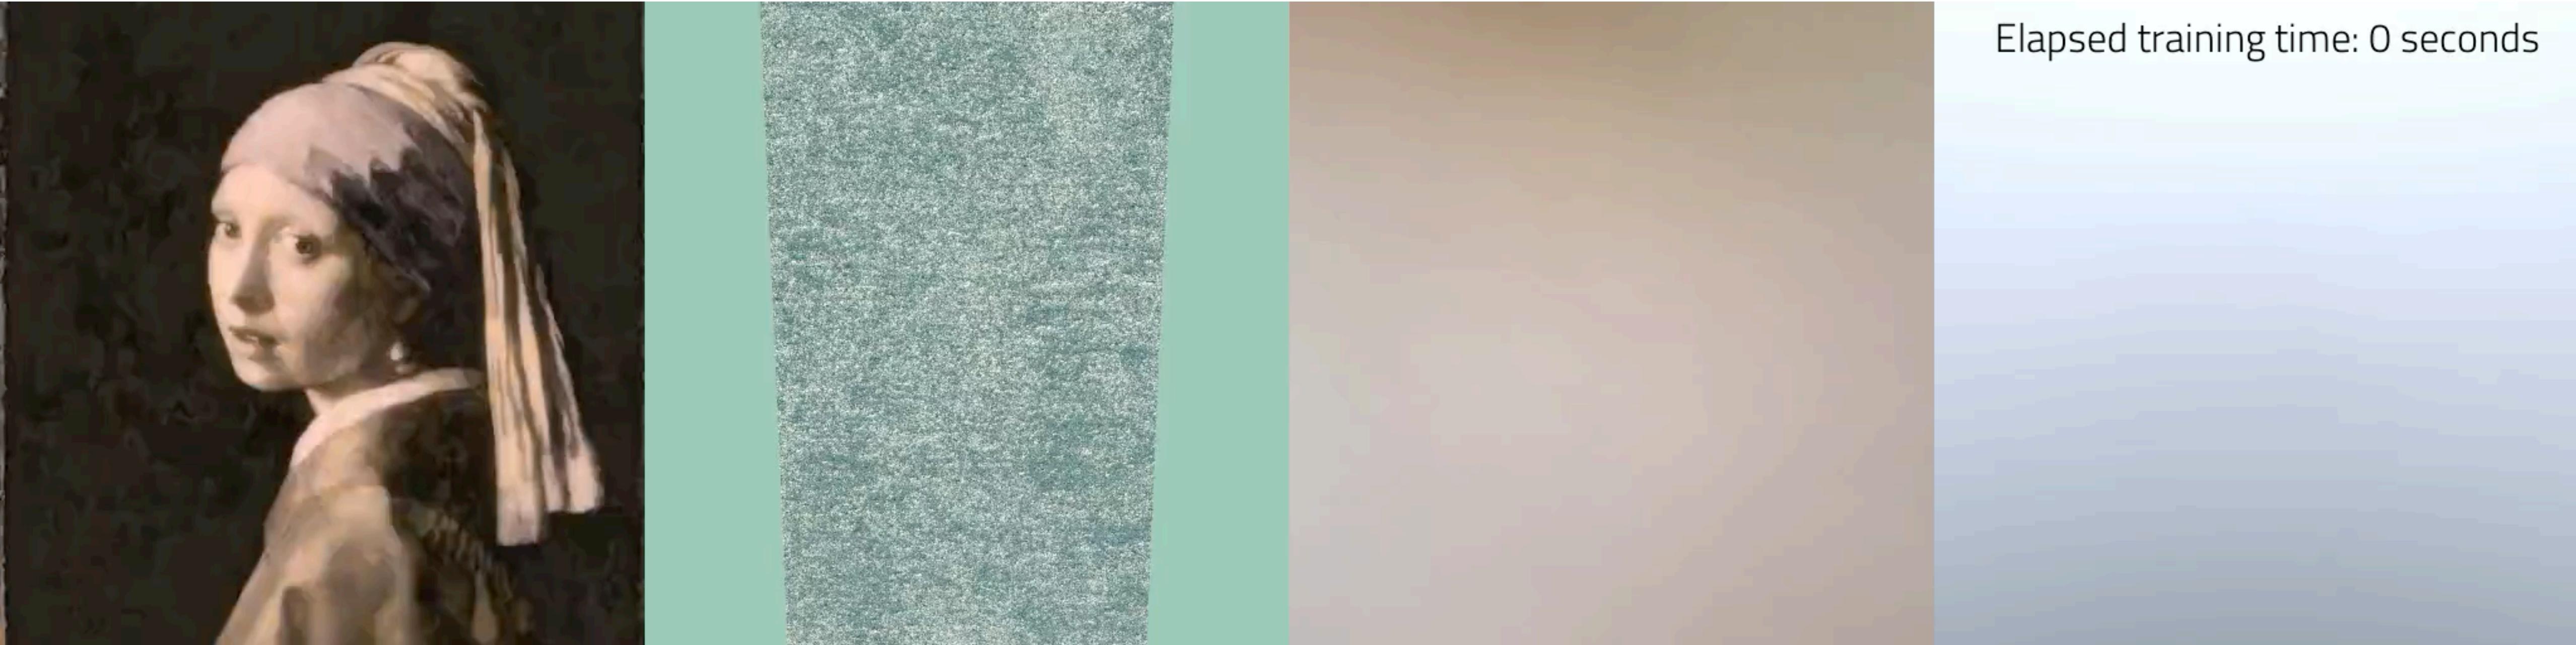
# Multi-scale Voxel Grids + Hash Table



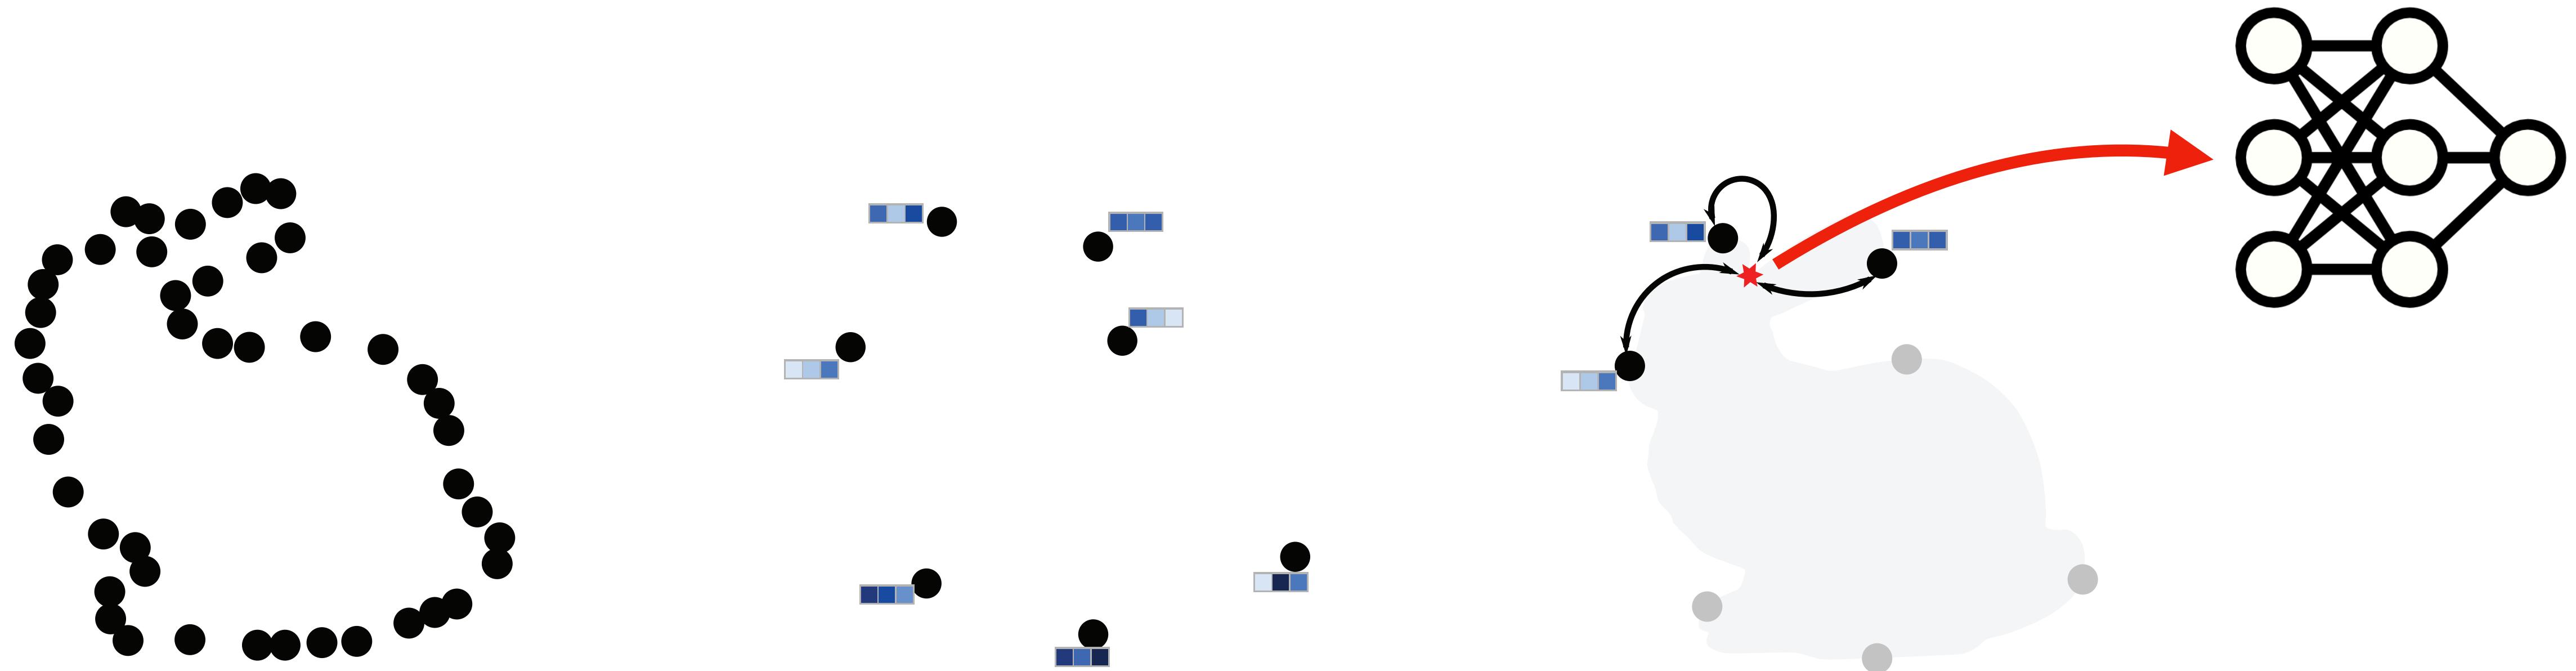
# Multi-scale Voxel Grids + Hash Table



# Multi-scale Voxel Grids + Hash Table

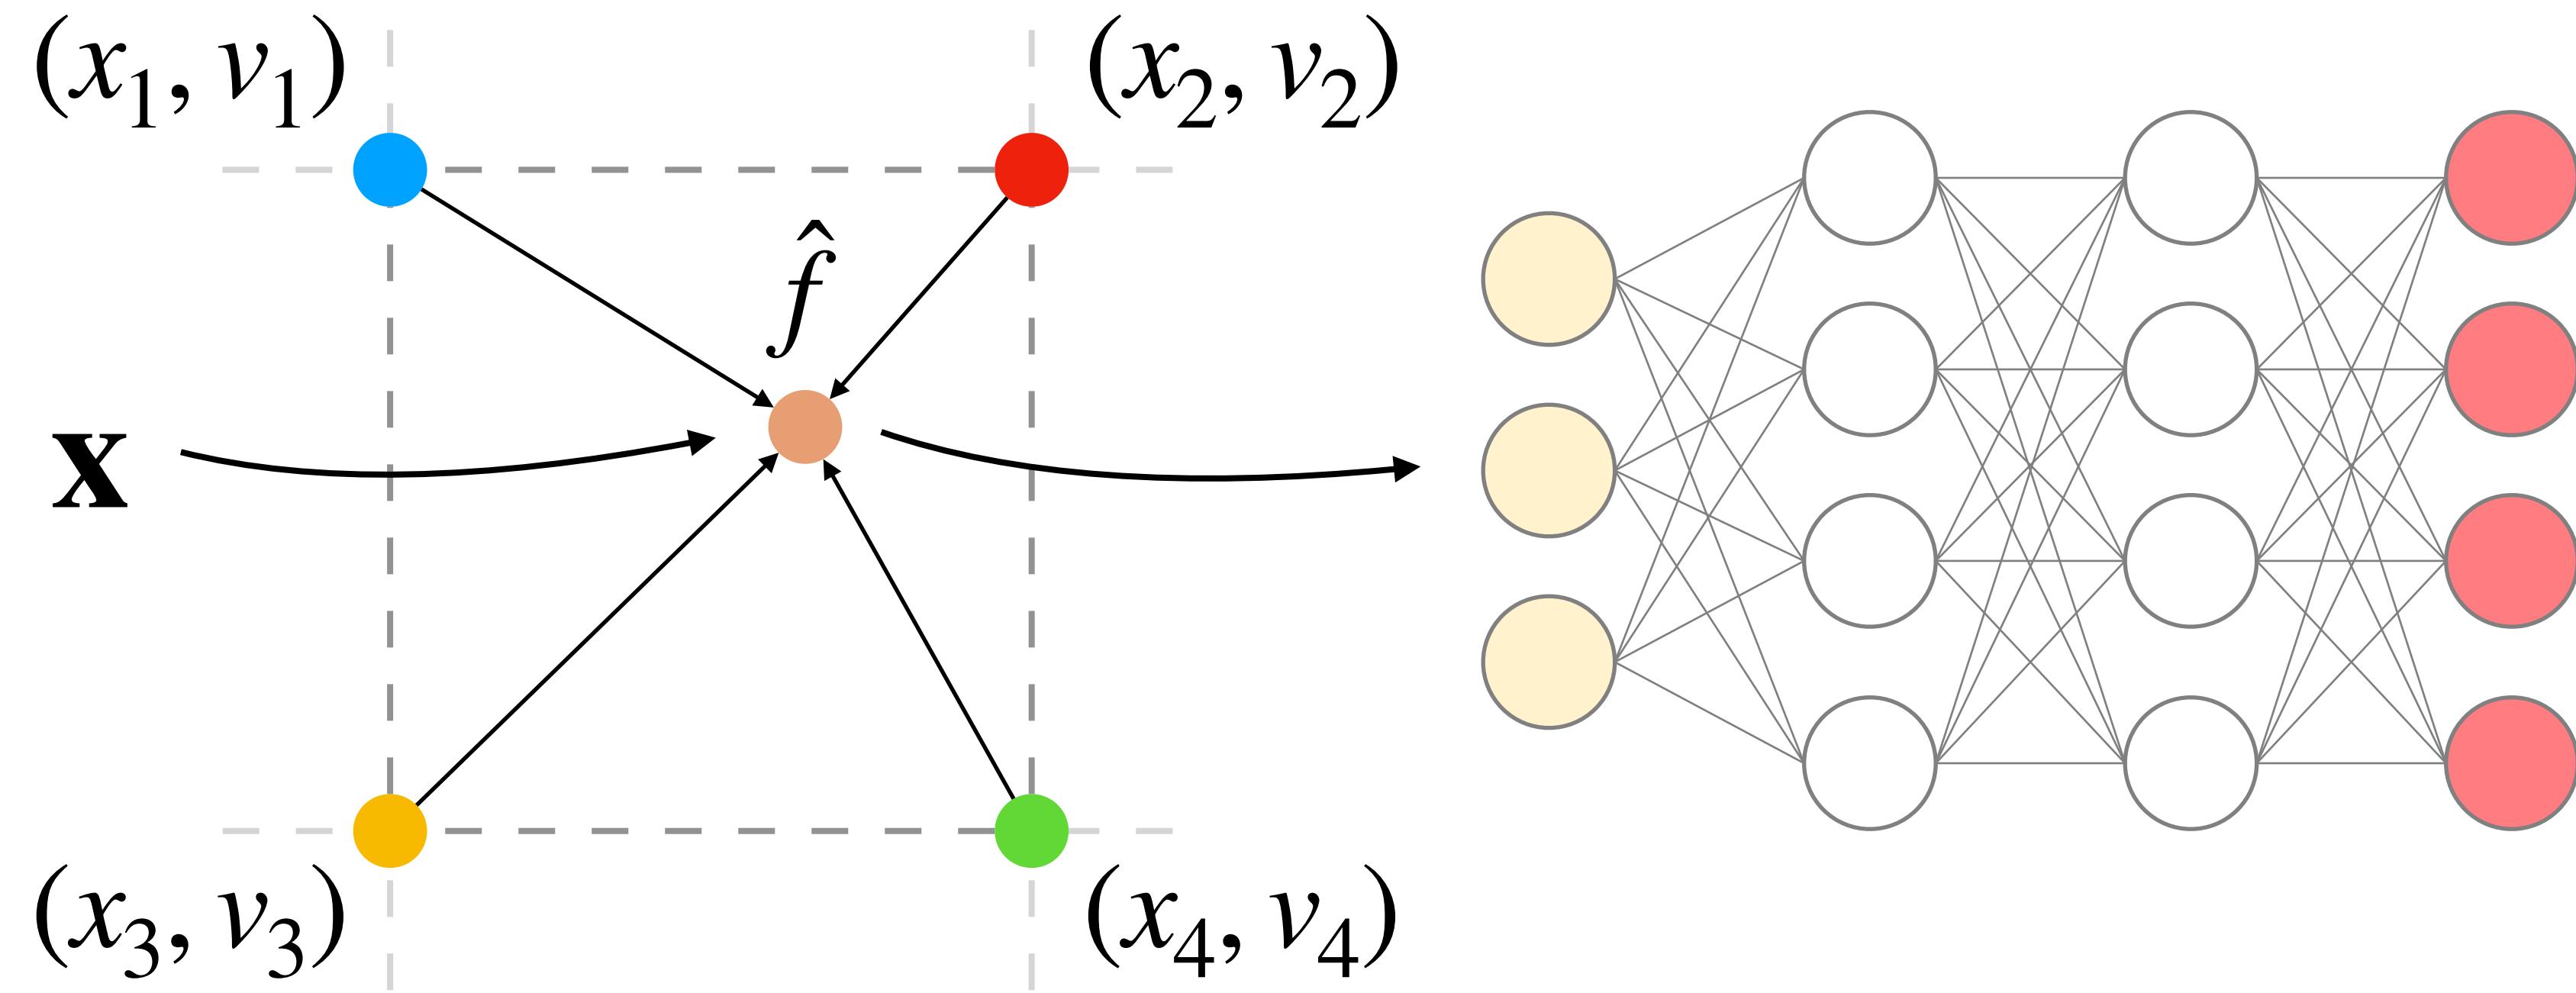


# Point Cloud & Nearest Neighbor

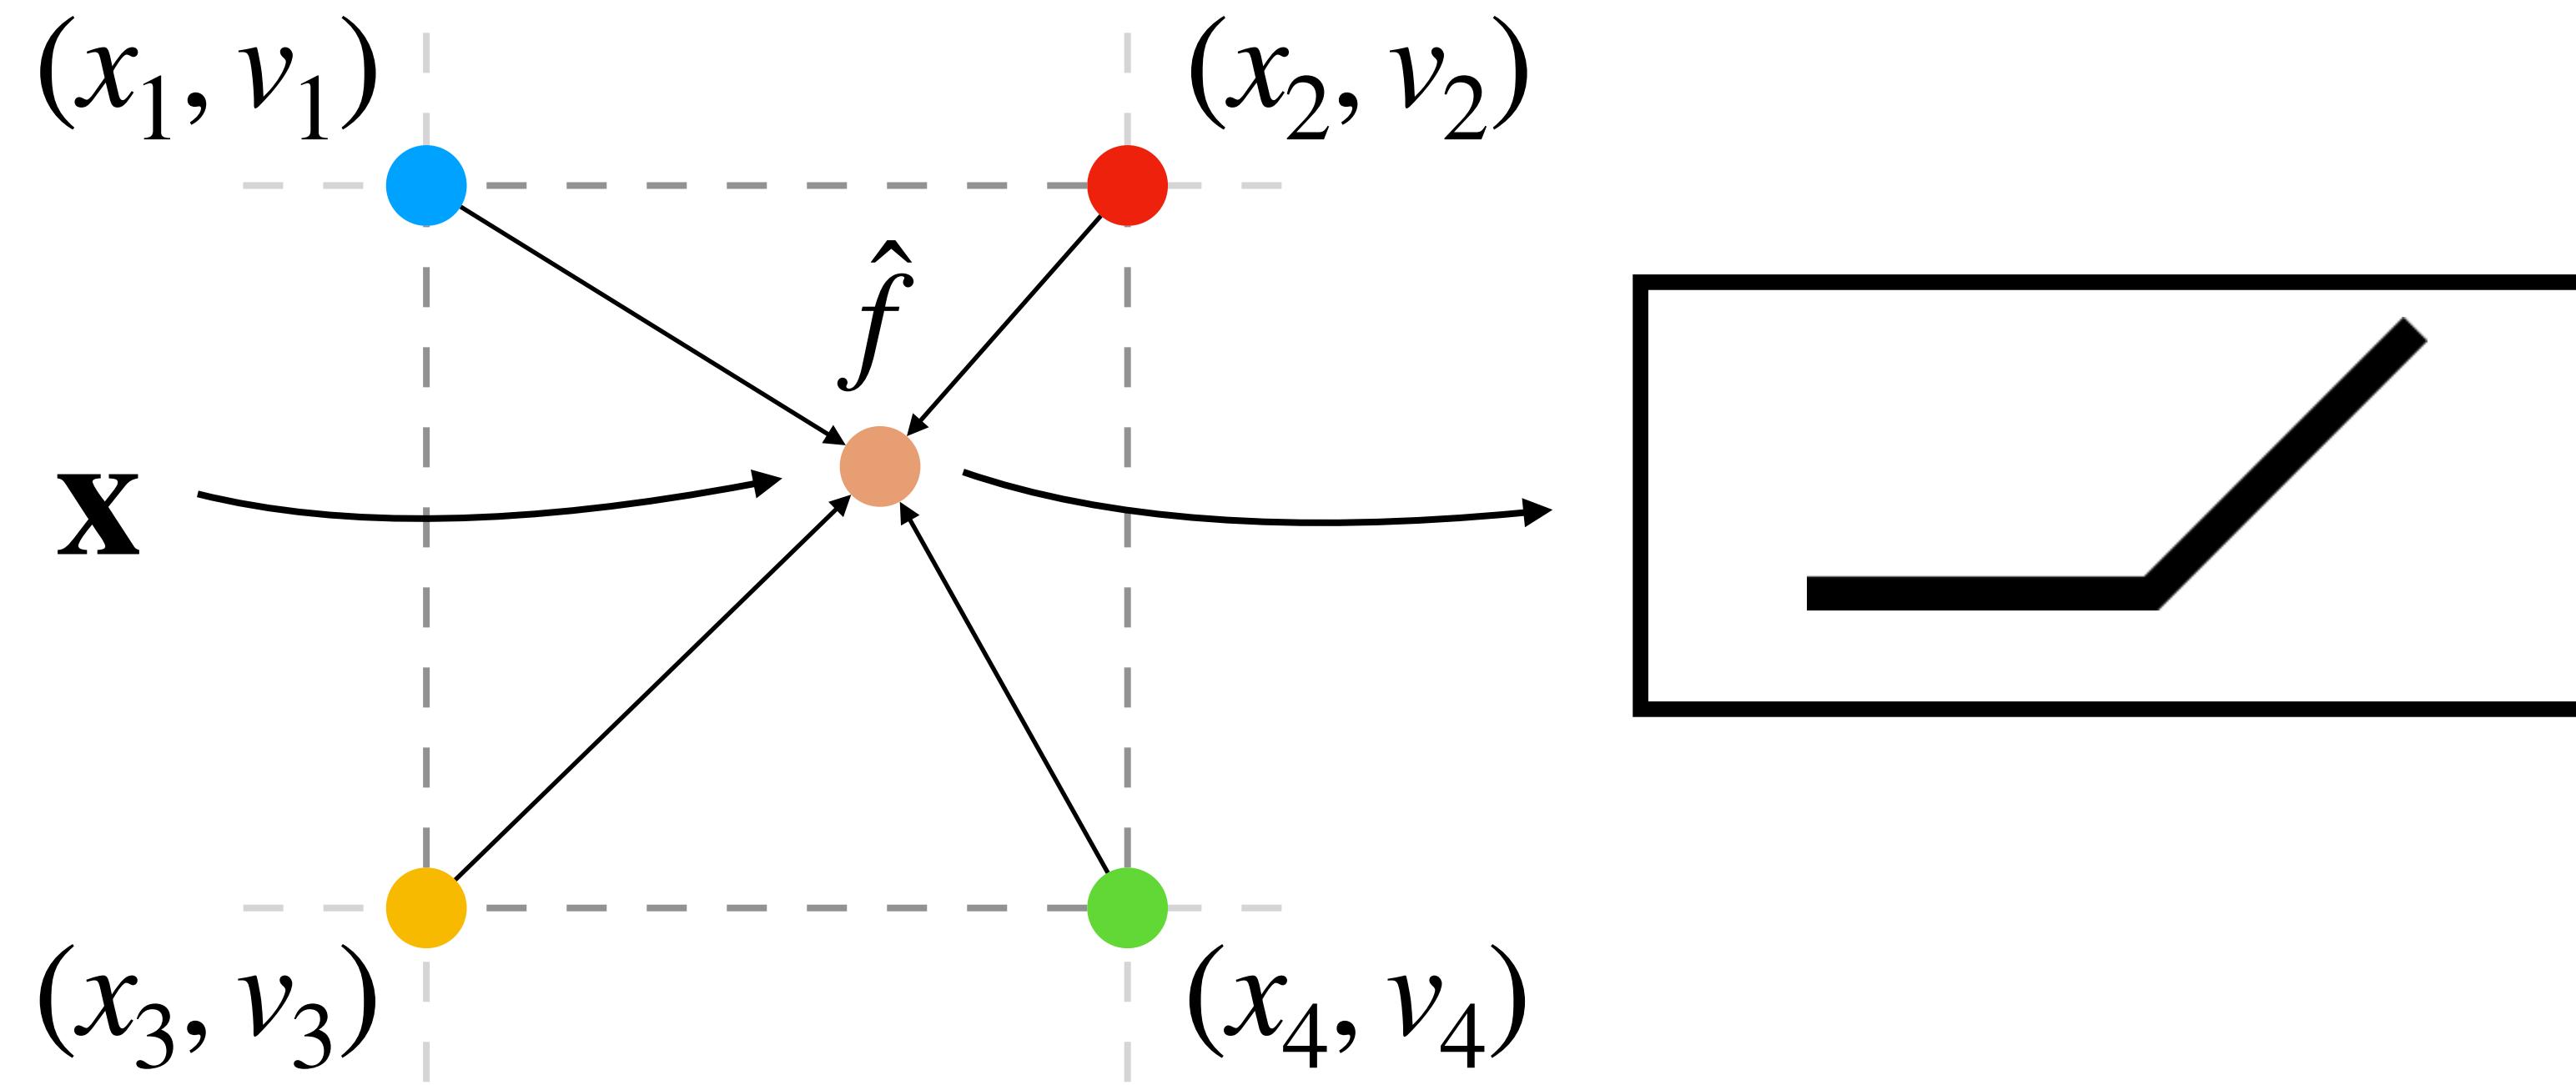


Hybrid representations can be used to decode a surface representation into a volume representation!

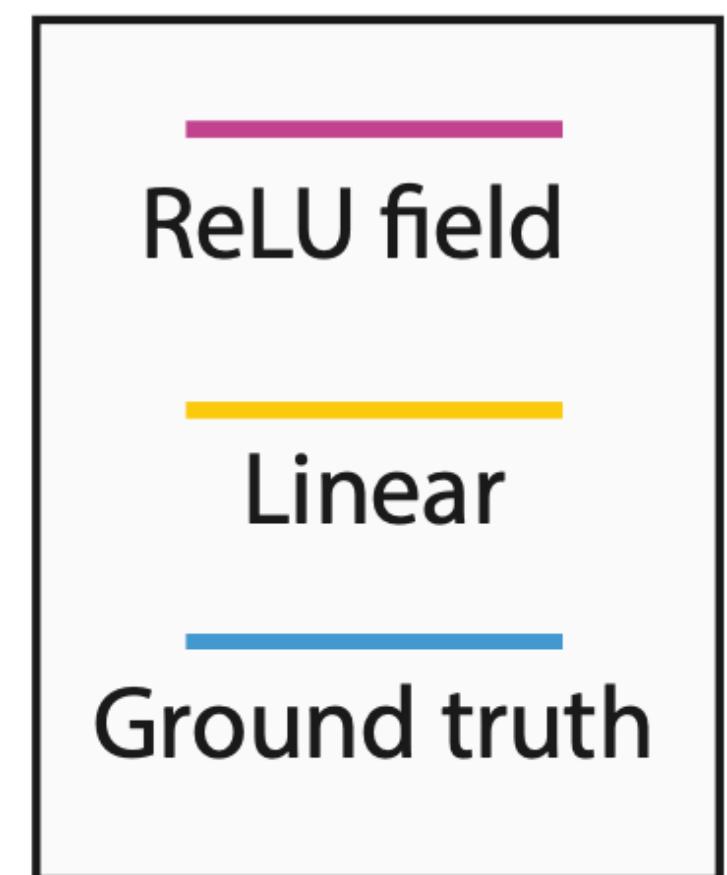
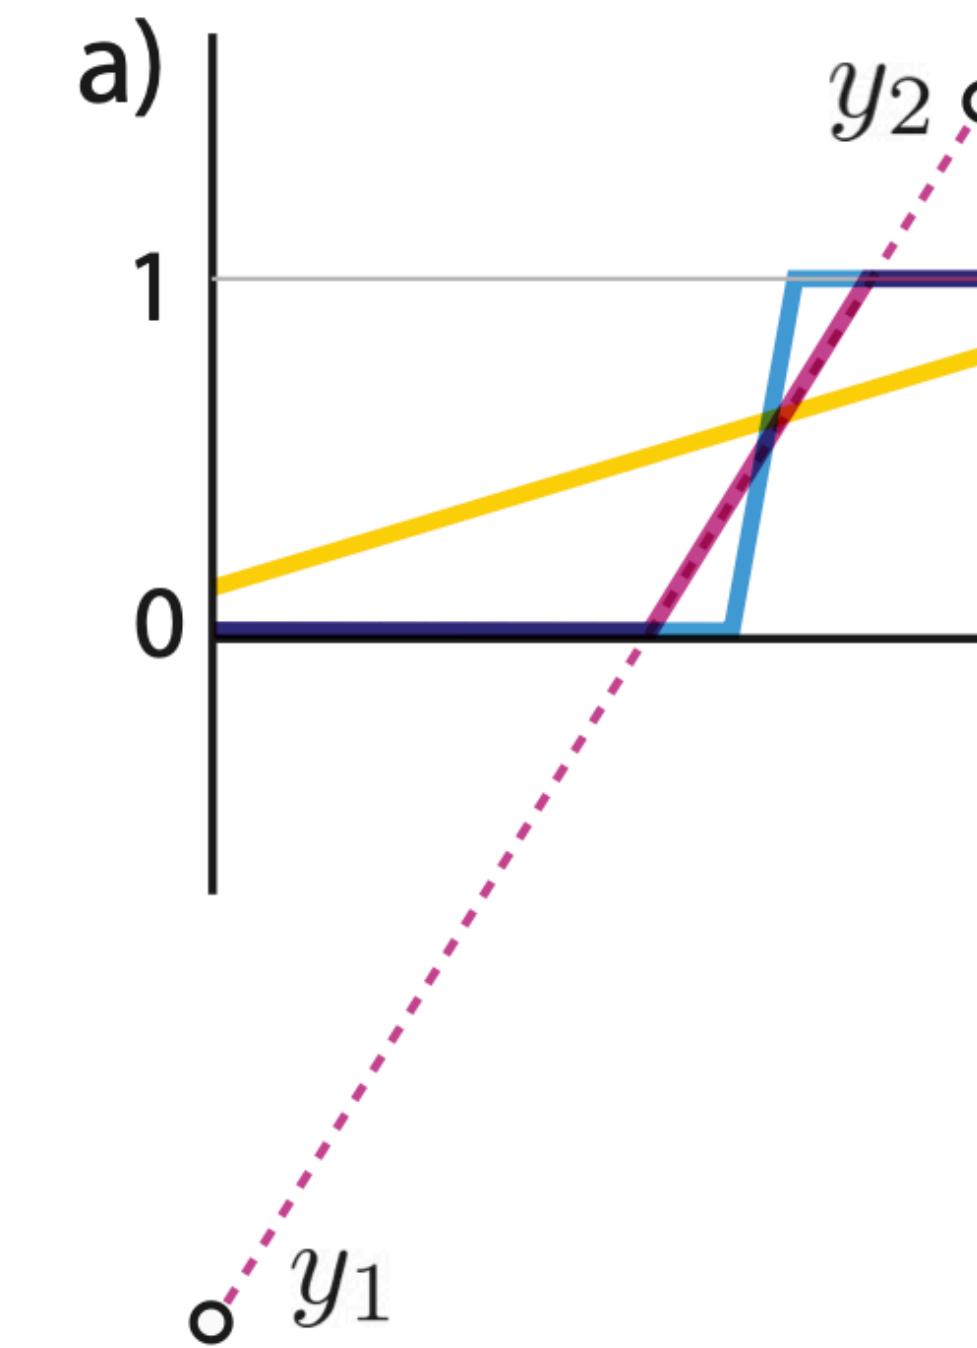
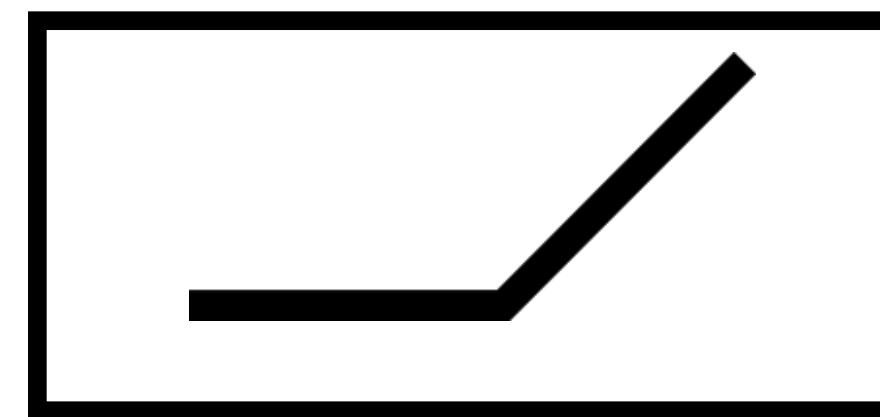
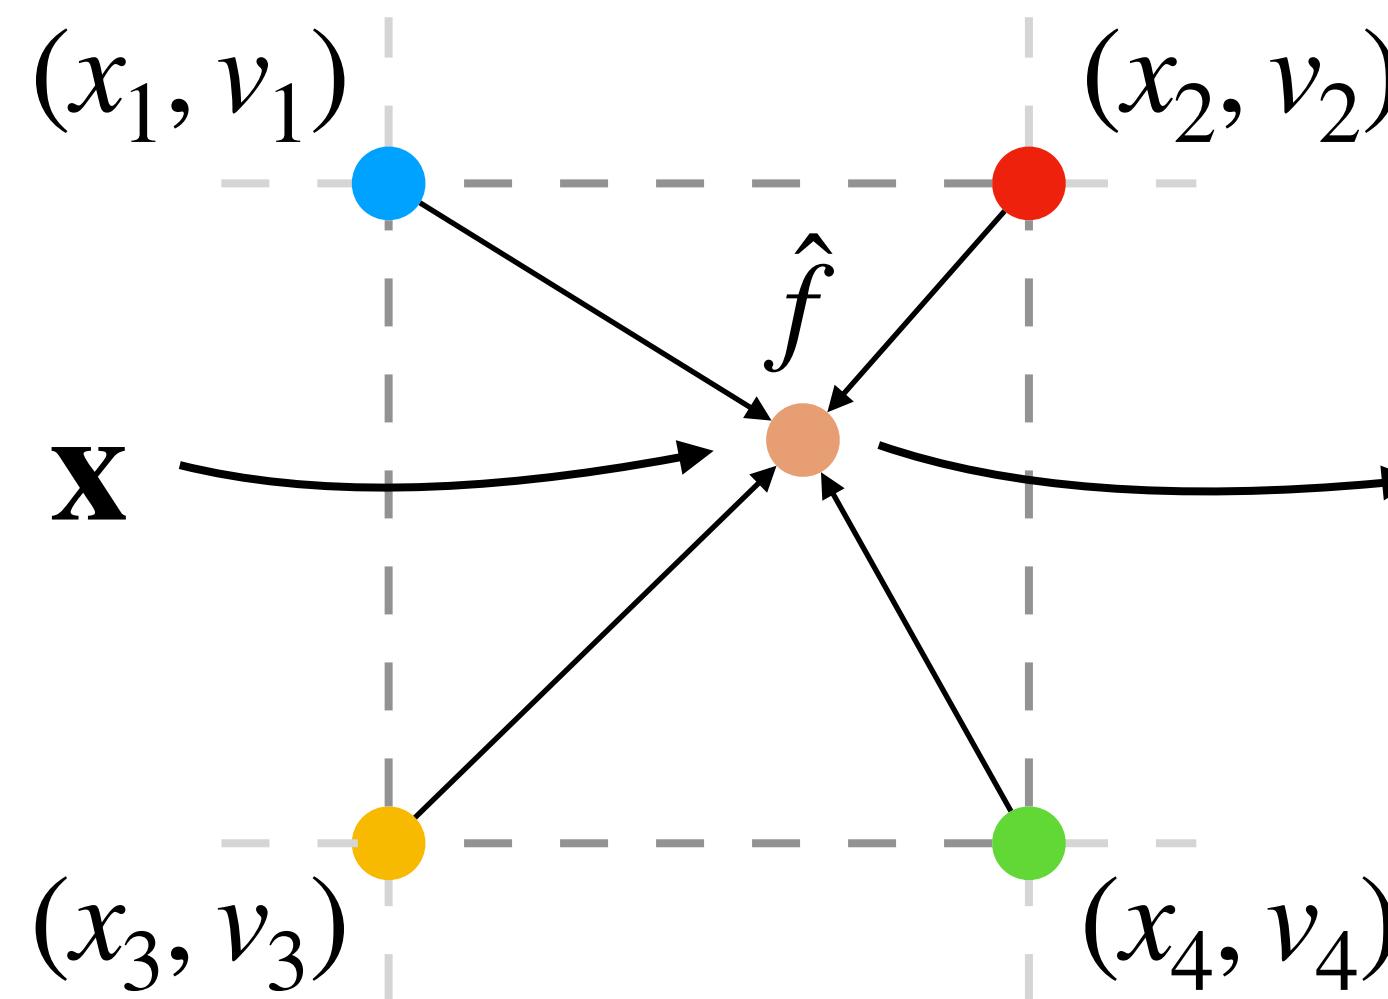
# ReLU Fields: How much MLP do you need?



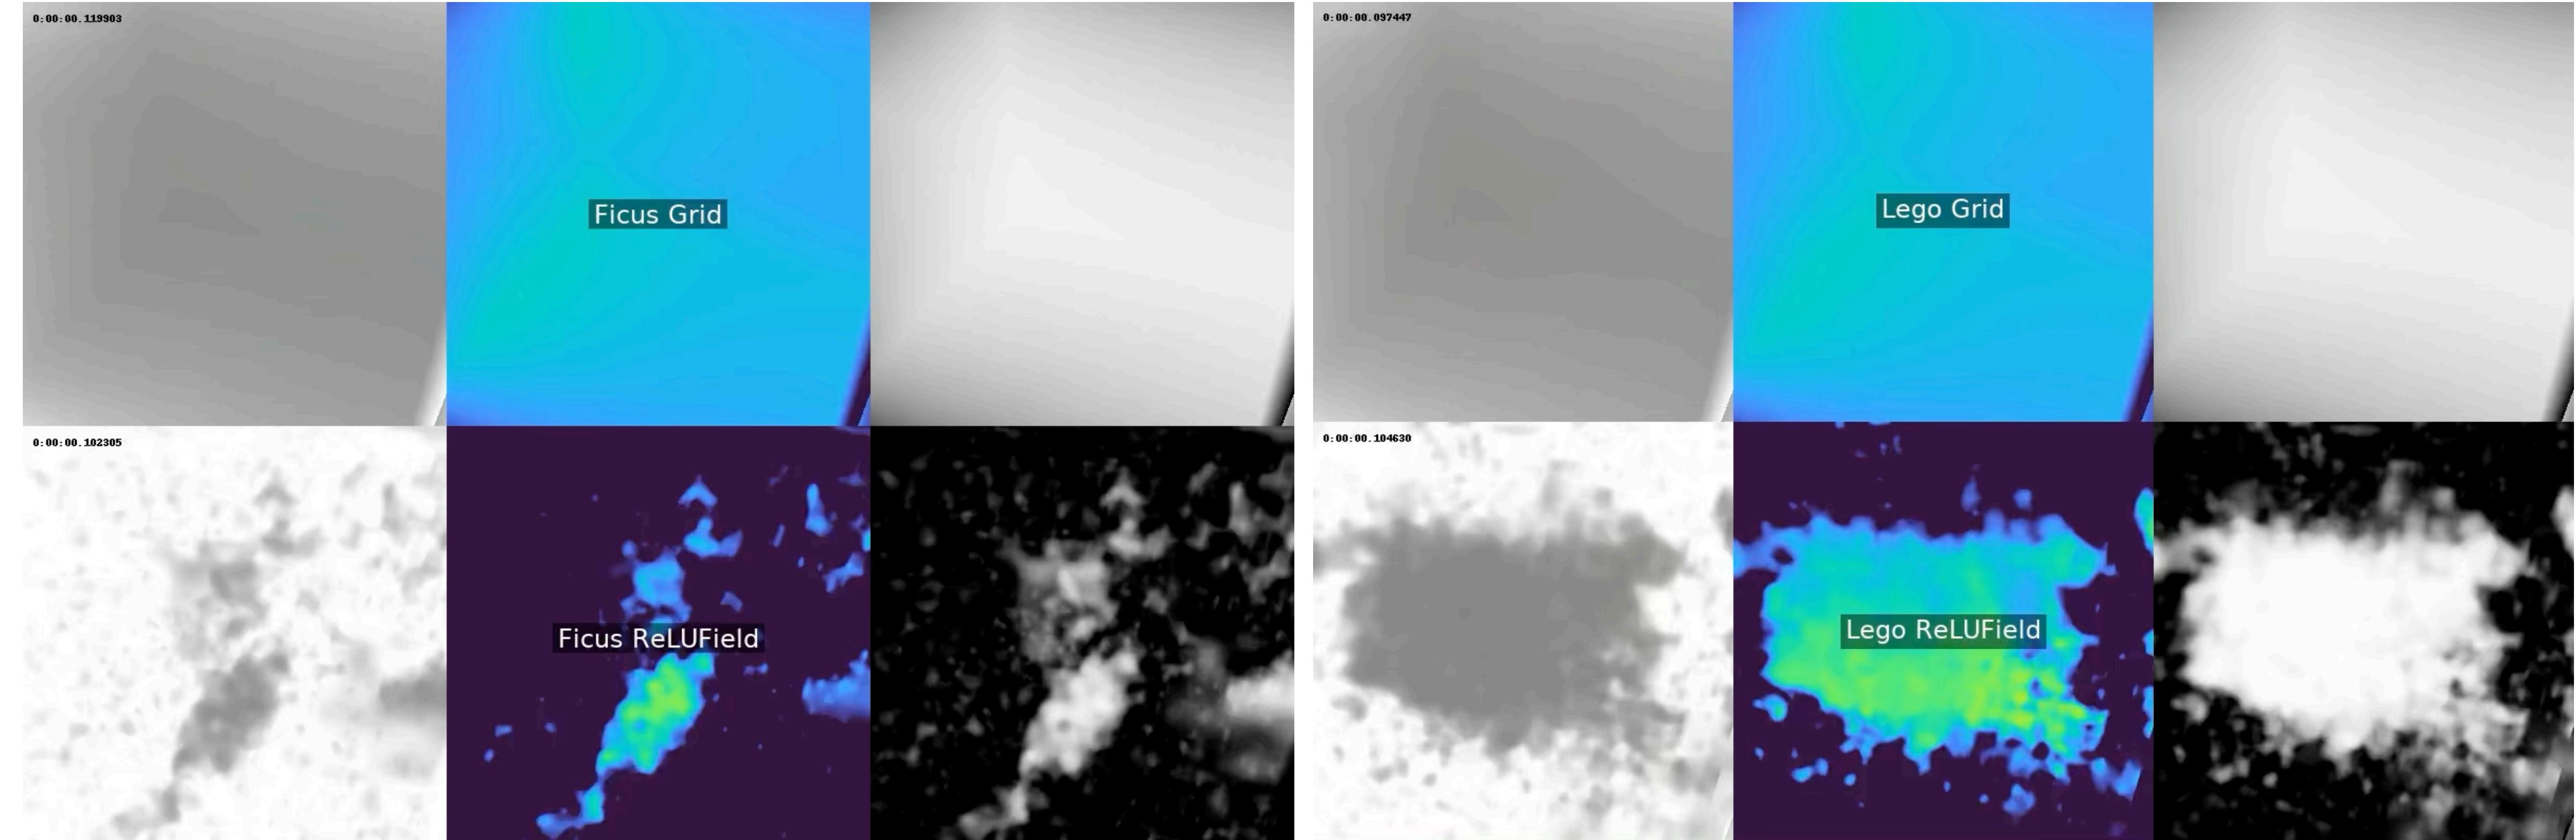
# ReLU Fields: How much MLP do you need?



# ReLU Fields: How much MLP do you need?

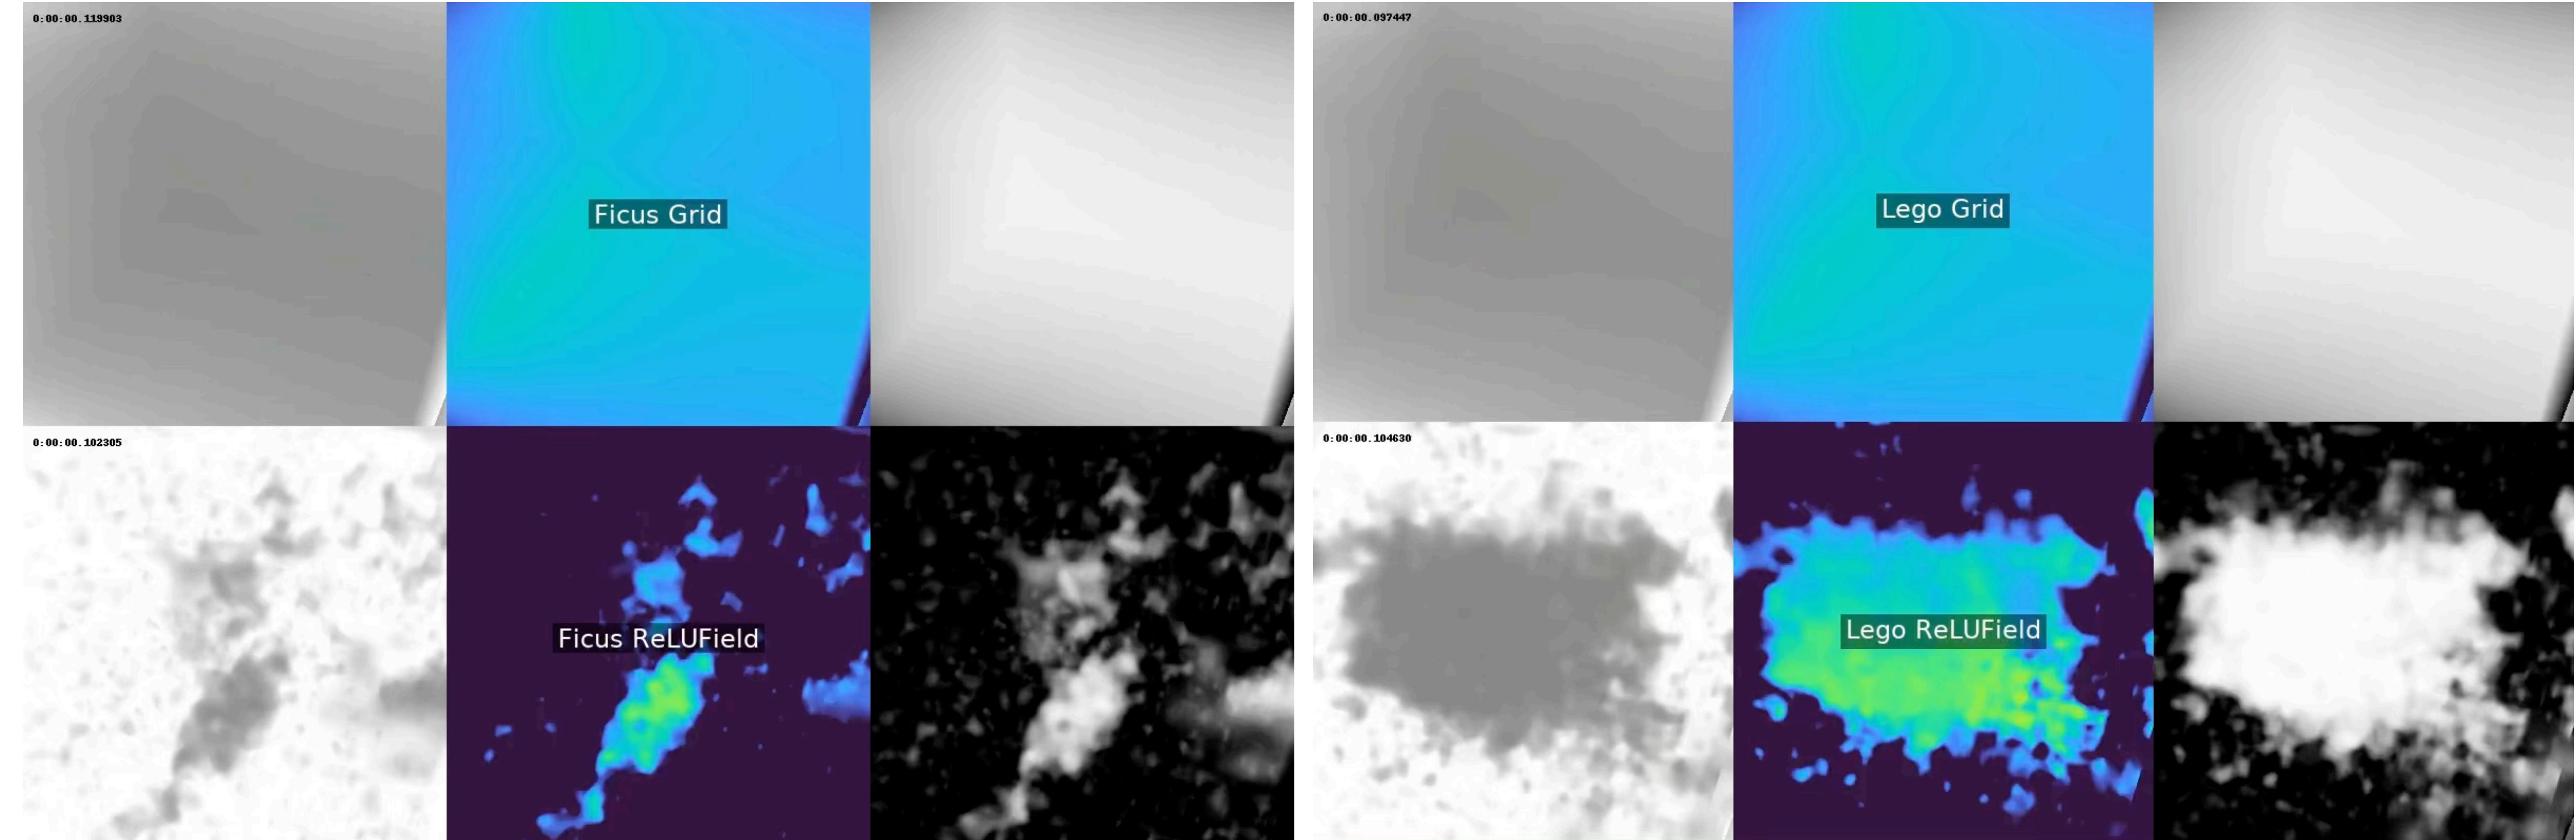


# ReLU Fields: How much MLP do you need?



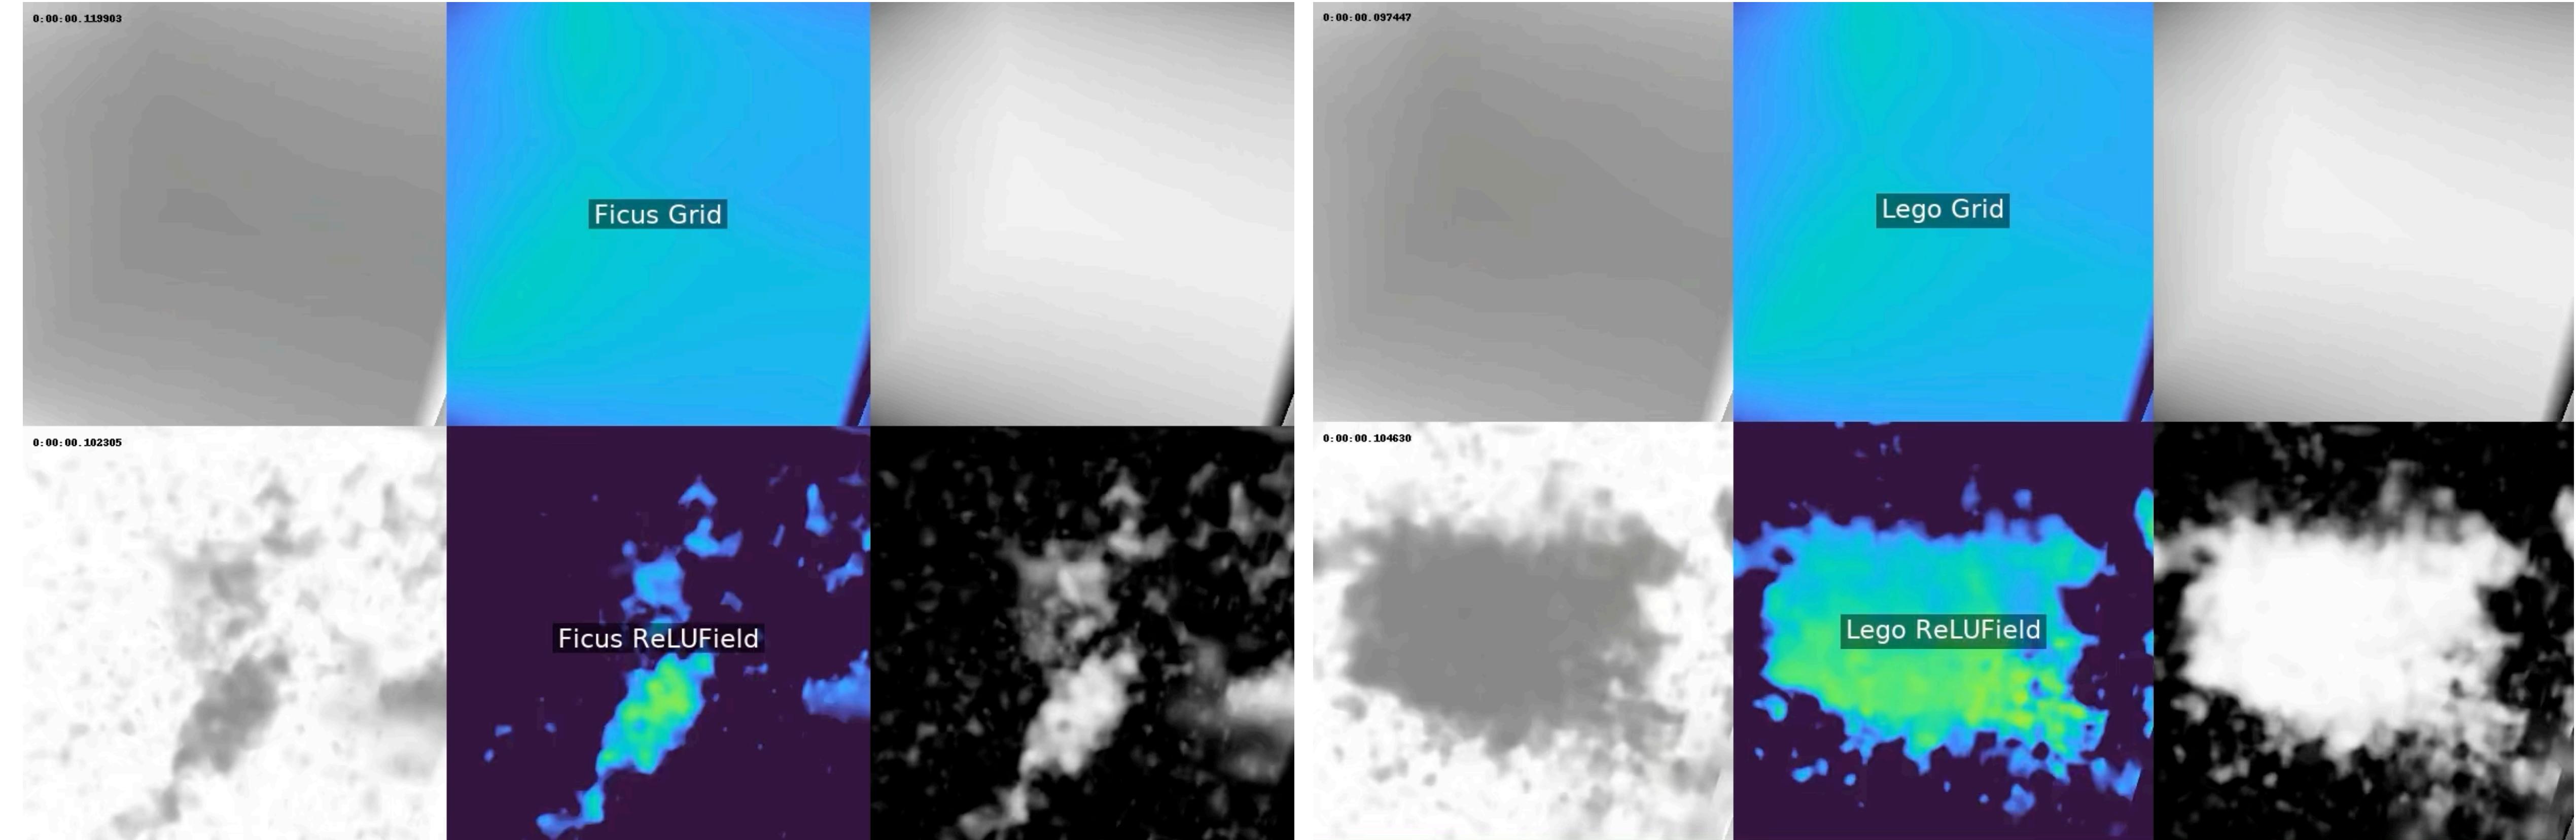
ReLU Fields: The Little Non-linearity That Could (Karnewar et al. 2022)

# ReLU Fields: How much MLP do you need?



ReLU Fields: The Little Non-linearity That Could (Karnewar et al. 2022)

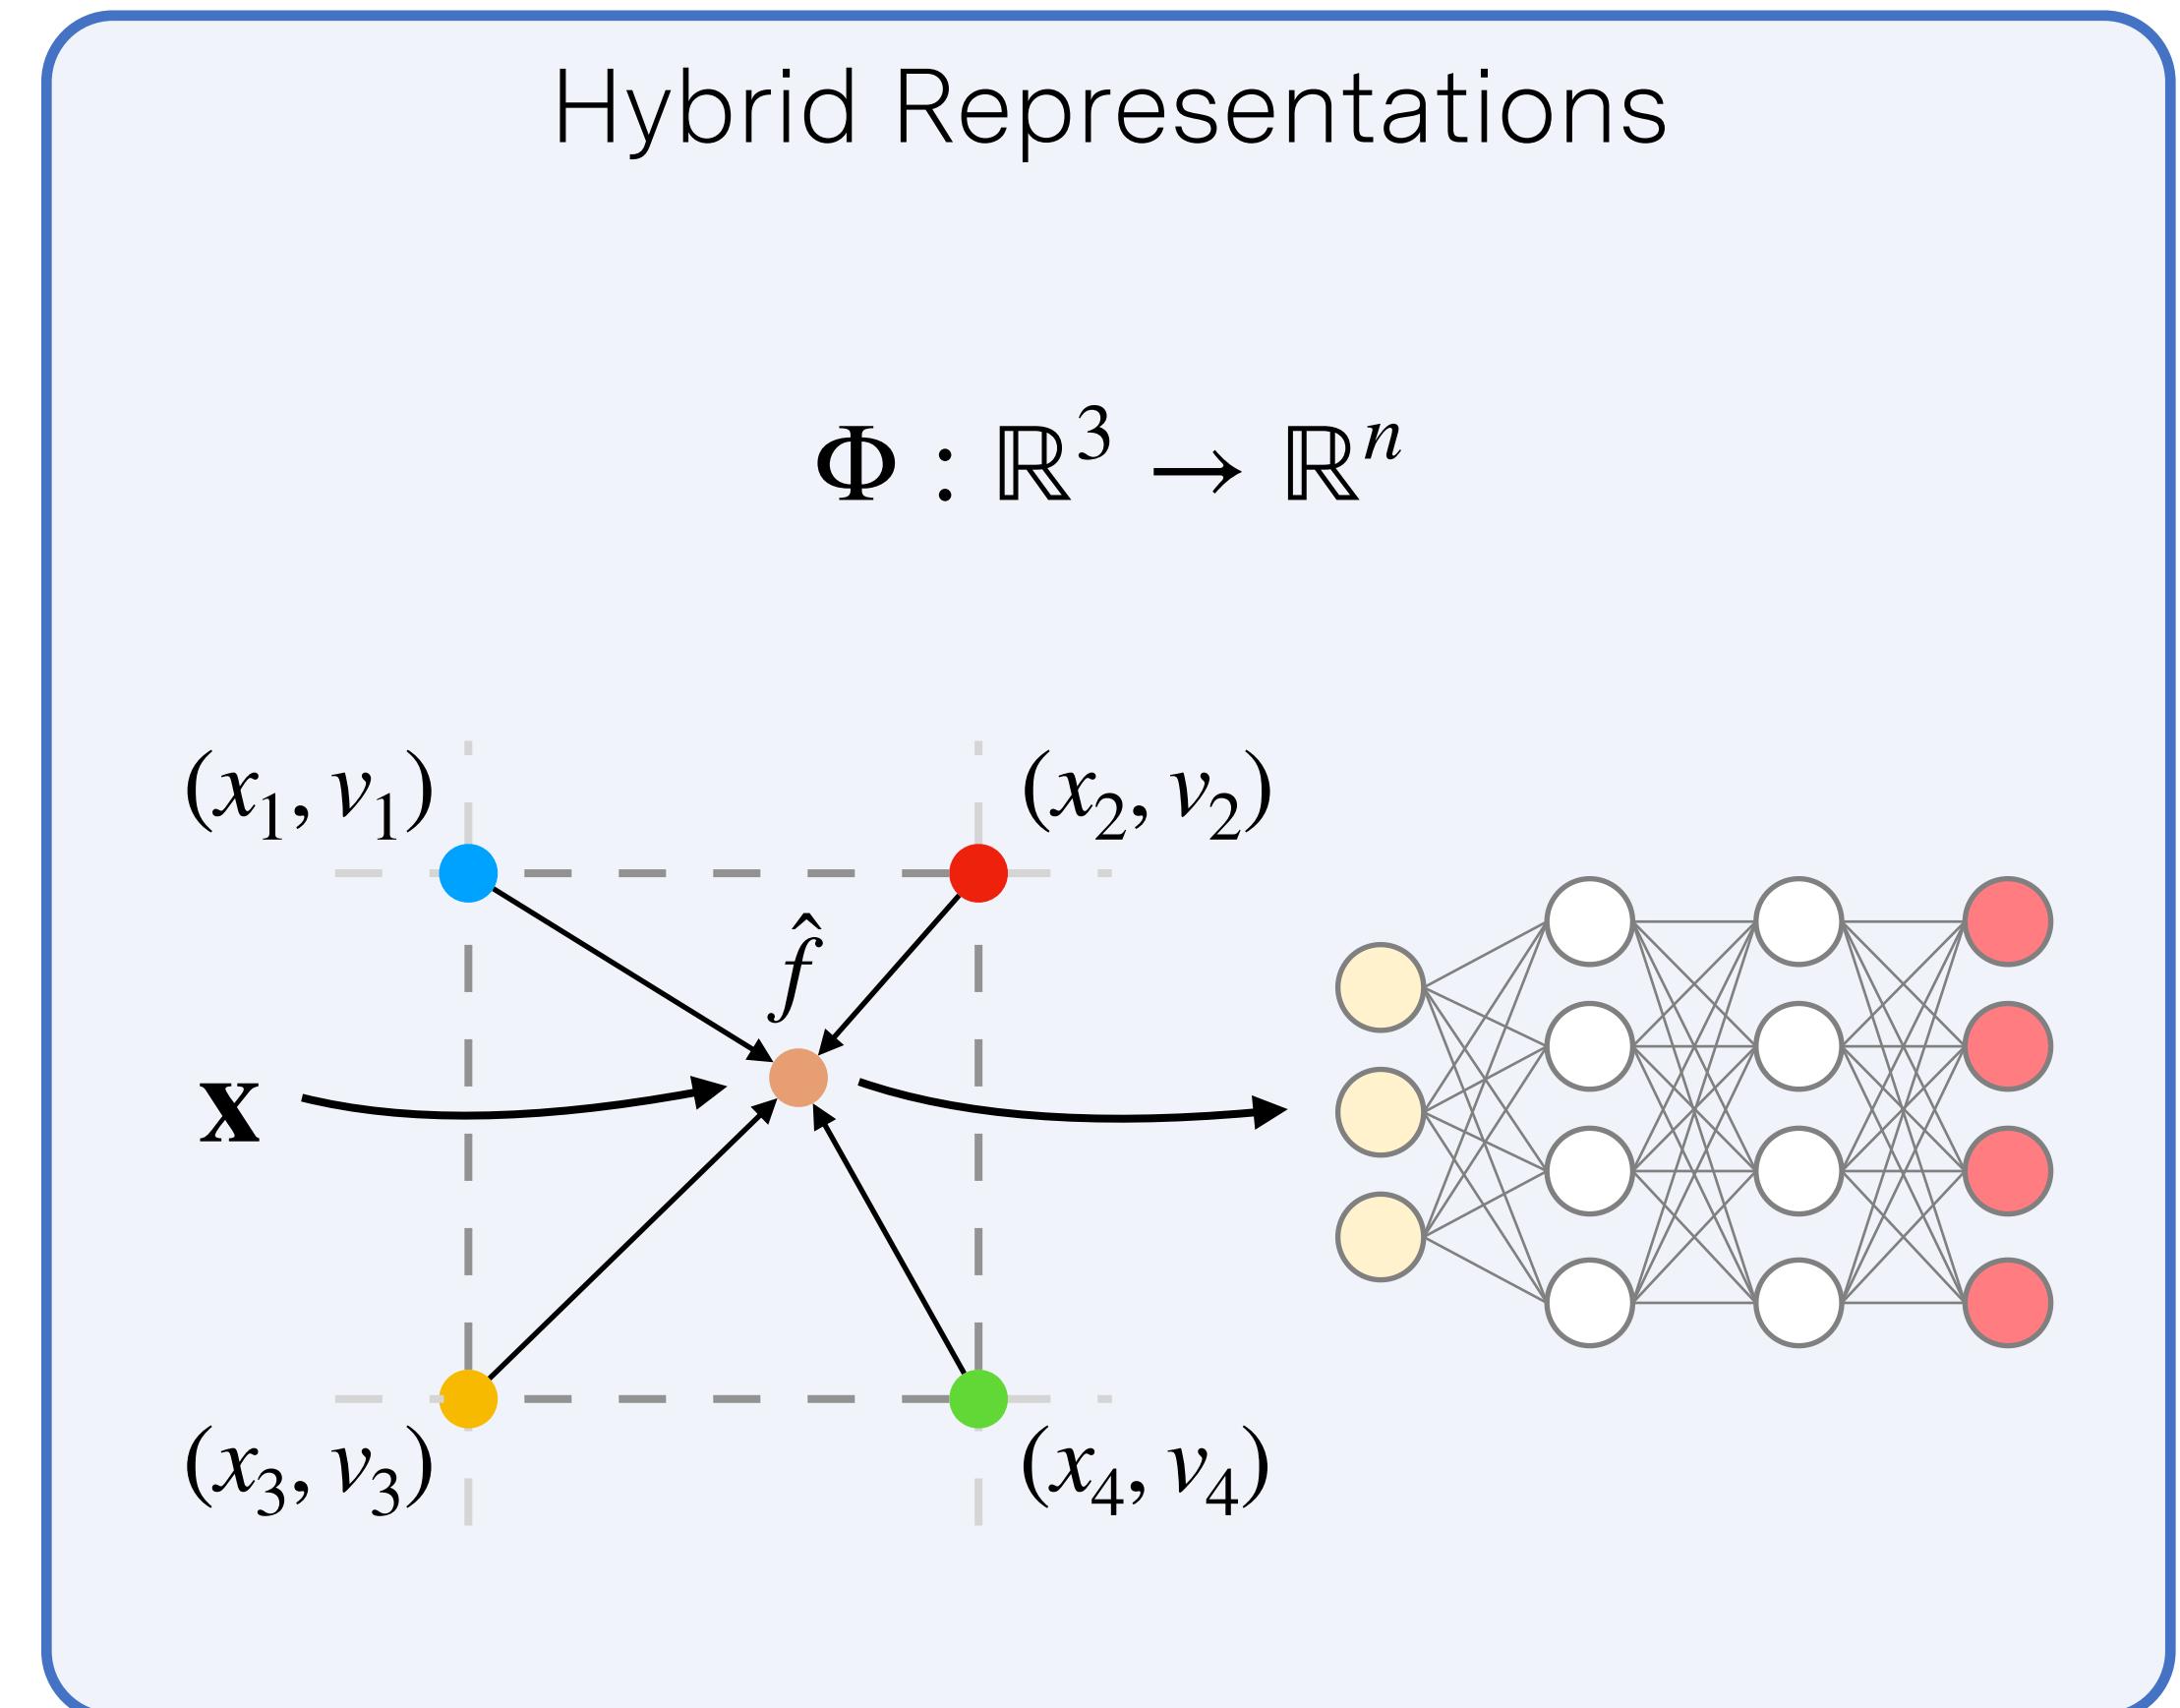
# ReLU Fields: How much MLP do you need?



ReLU Fields: The Little Non-linearity That Could (Karnewar et al. 2022)

# Hybrid Representations

- Very natural to trade off memory with compute.
- Neural network can locally “super resolve” discrete representation: Basically acts as interpolation kernel.
- Can be the best of both worlds: Fast and “continuous”.

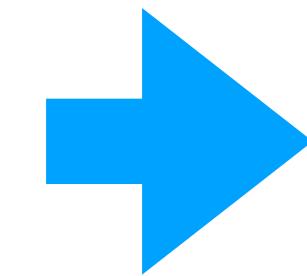
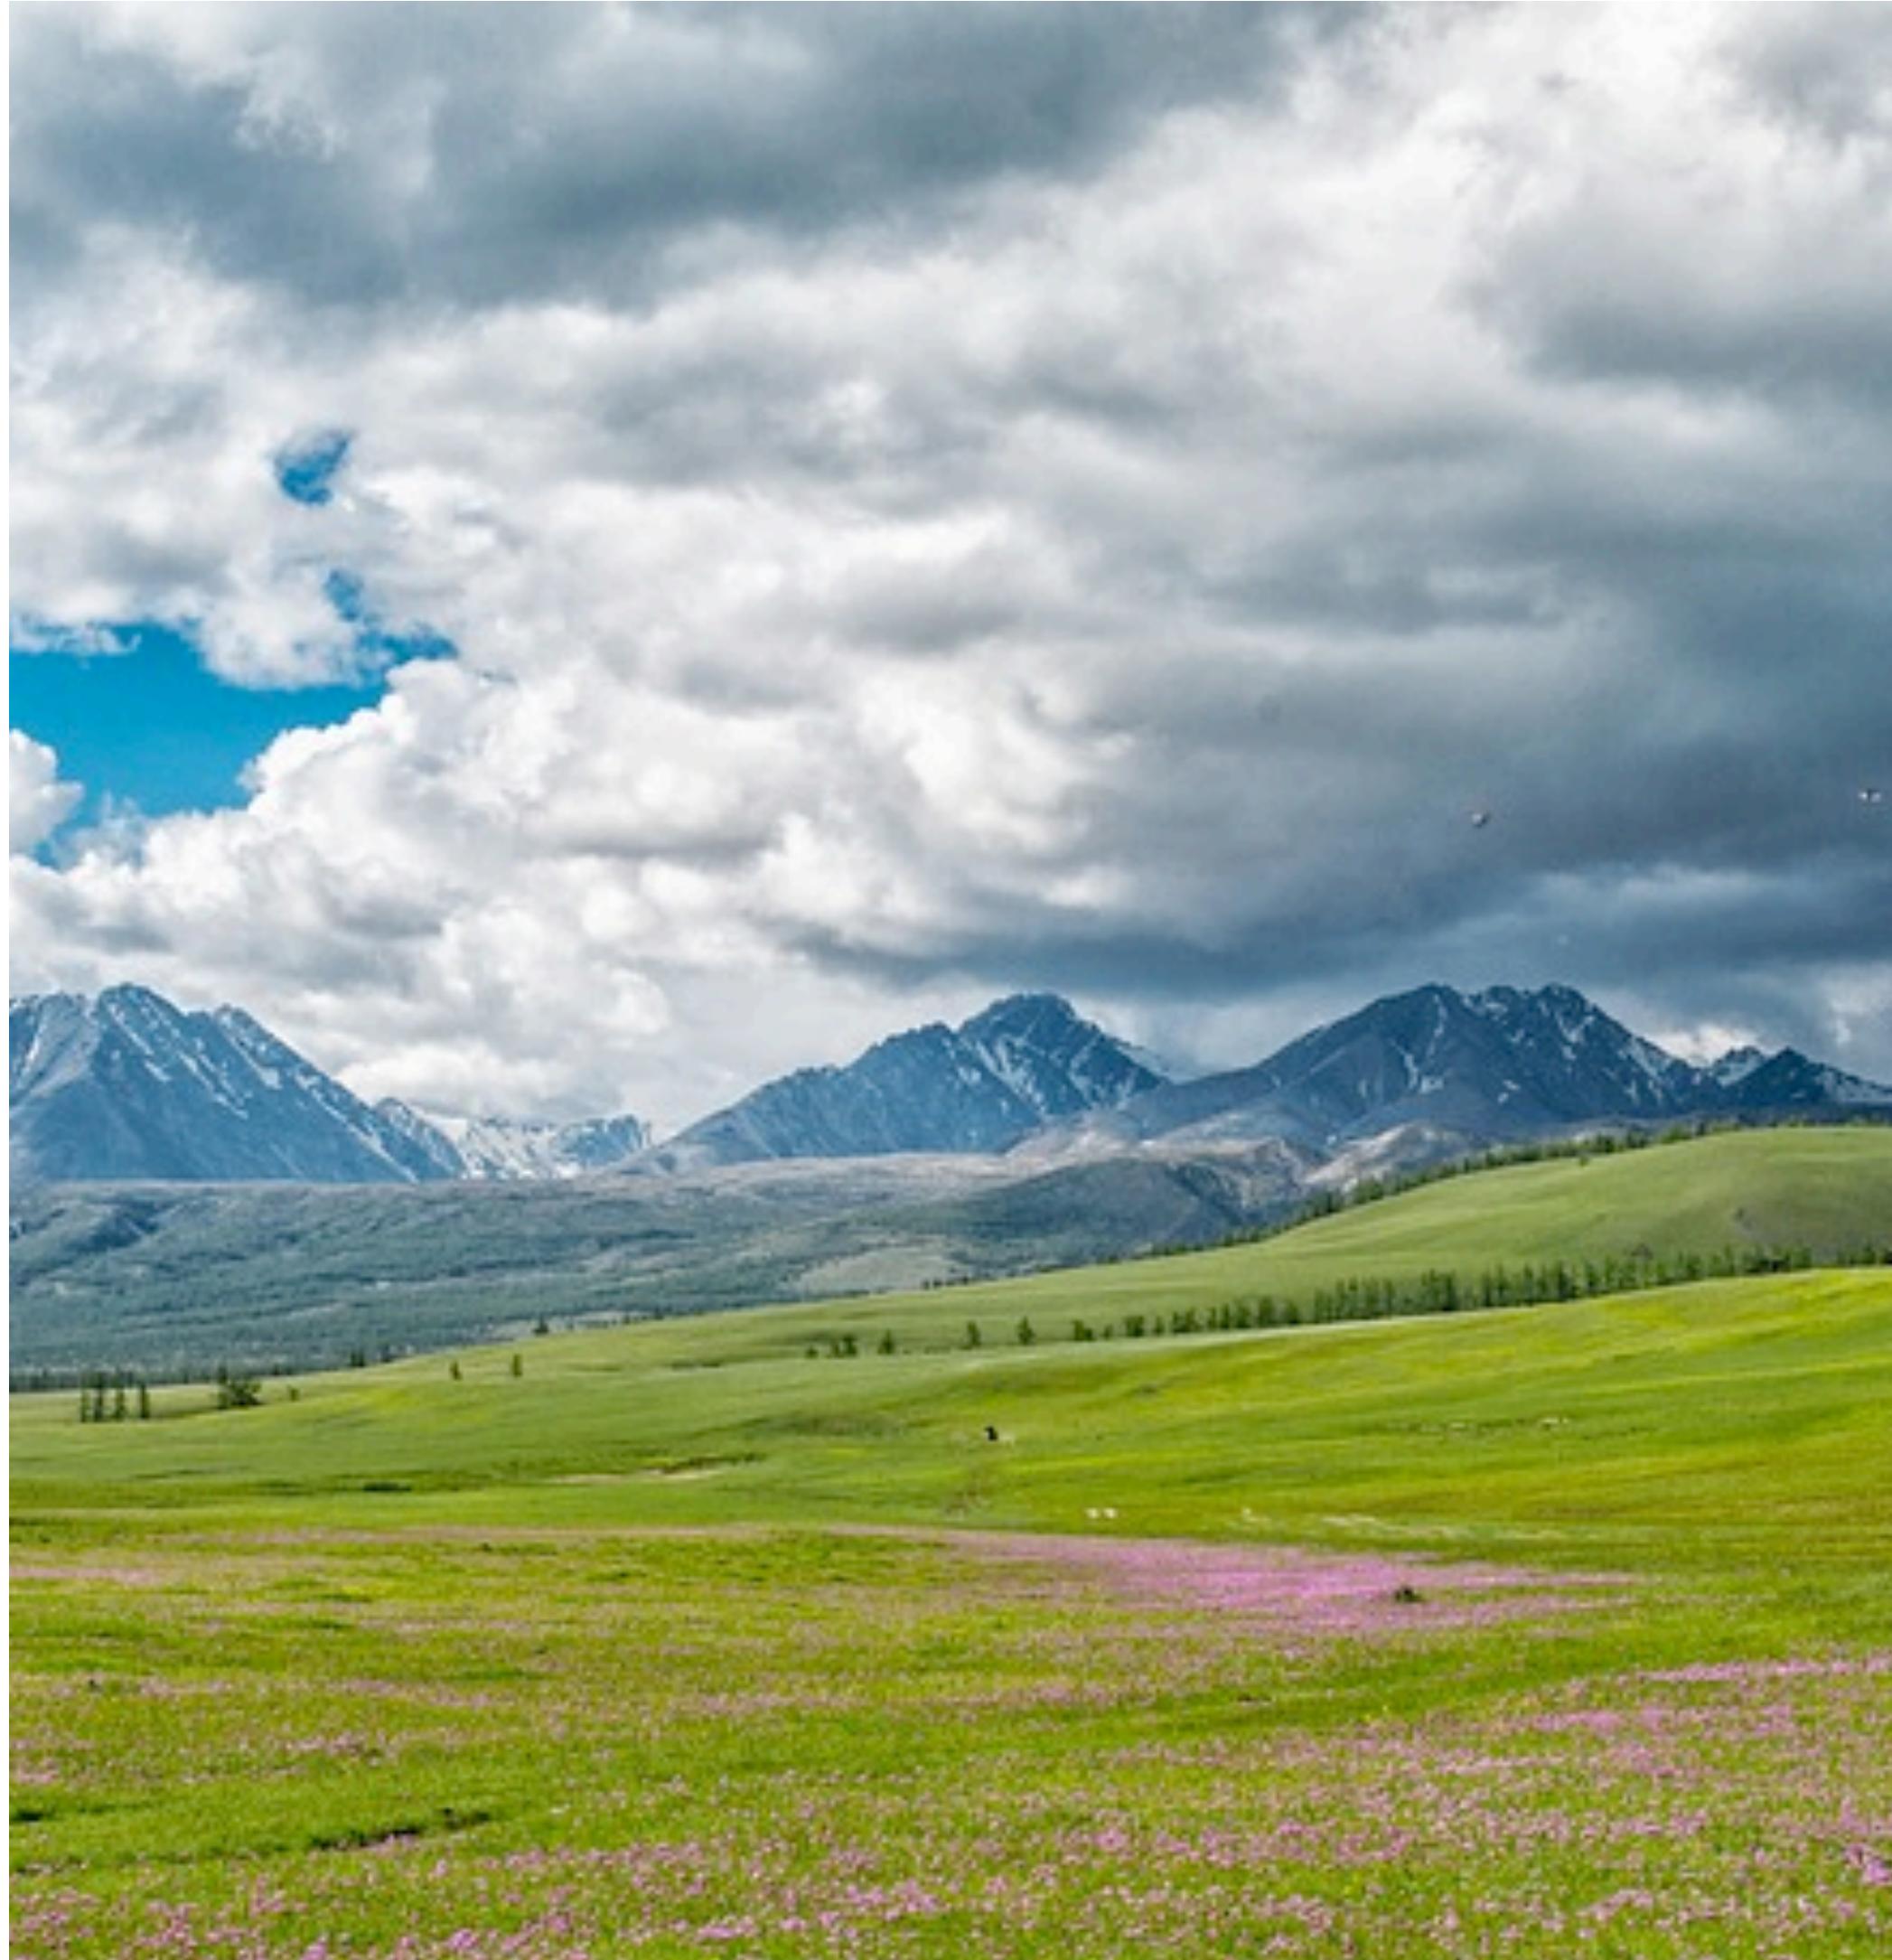


# How can we represent unbounded scenes?



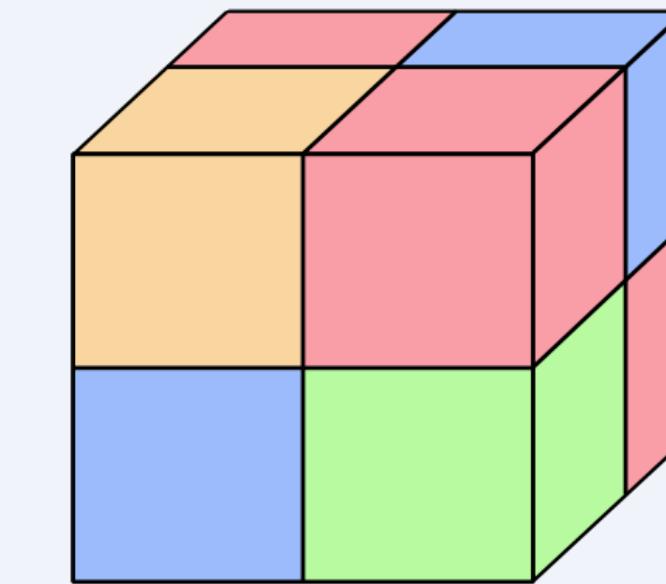
Image Source: Wikipedia

# How can we represent unbounded scenes?

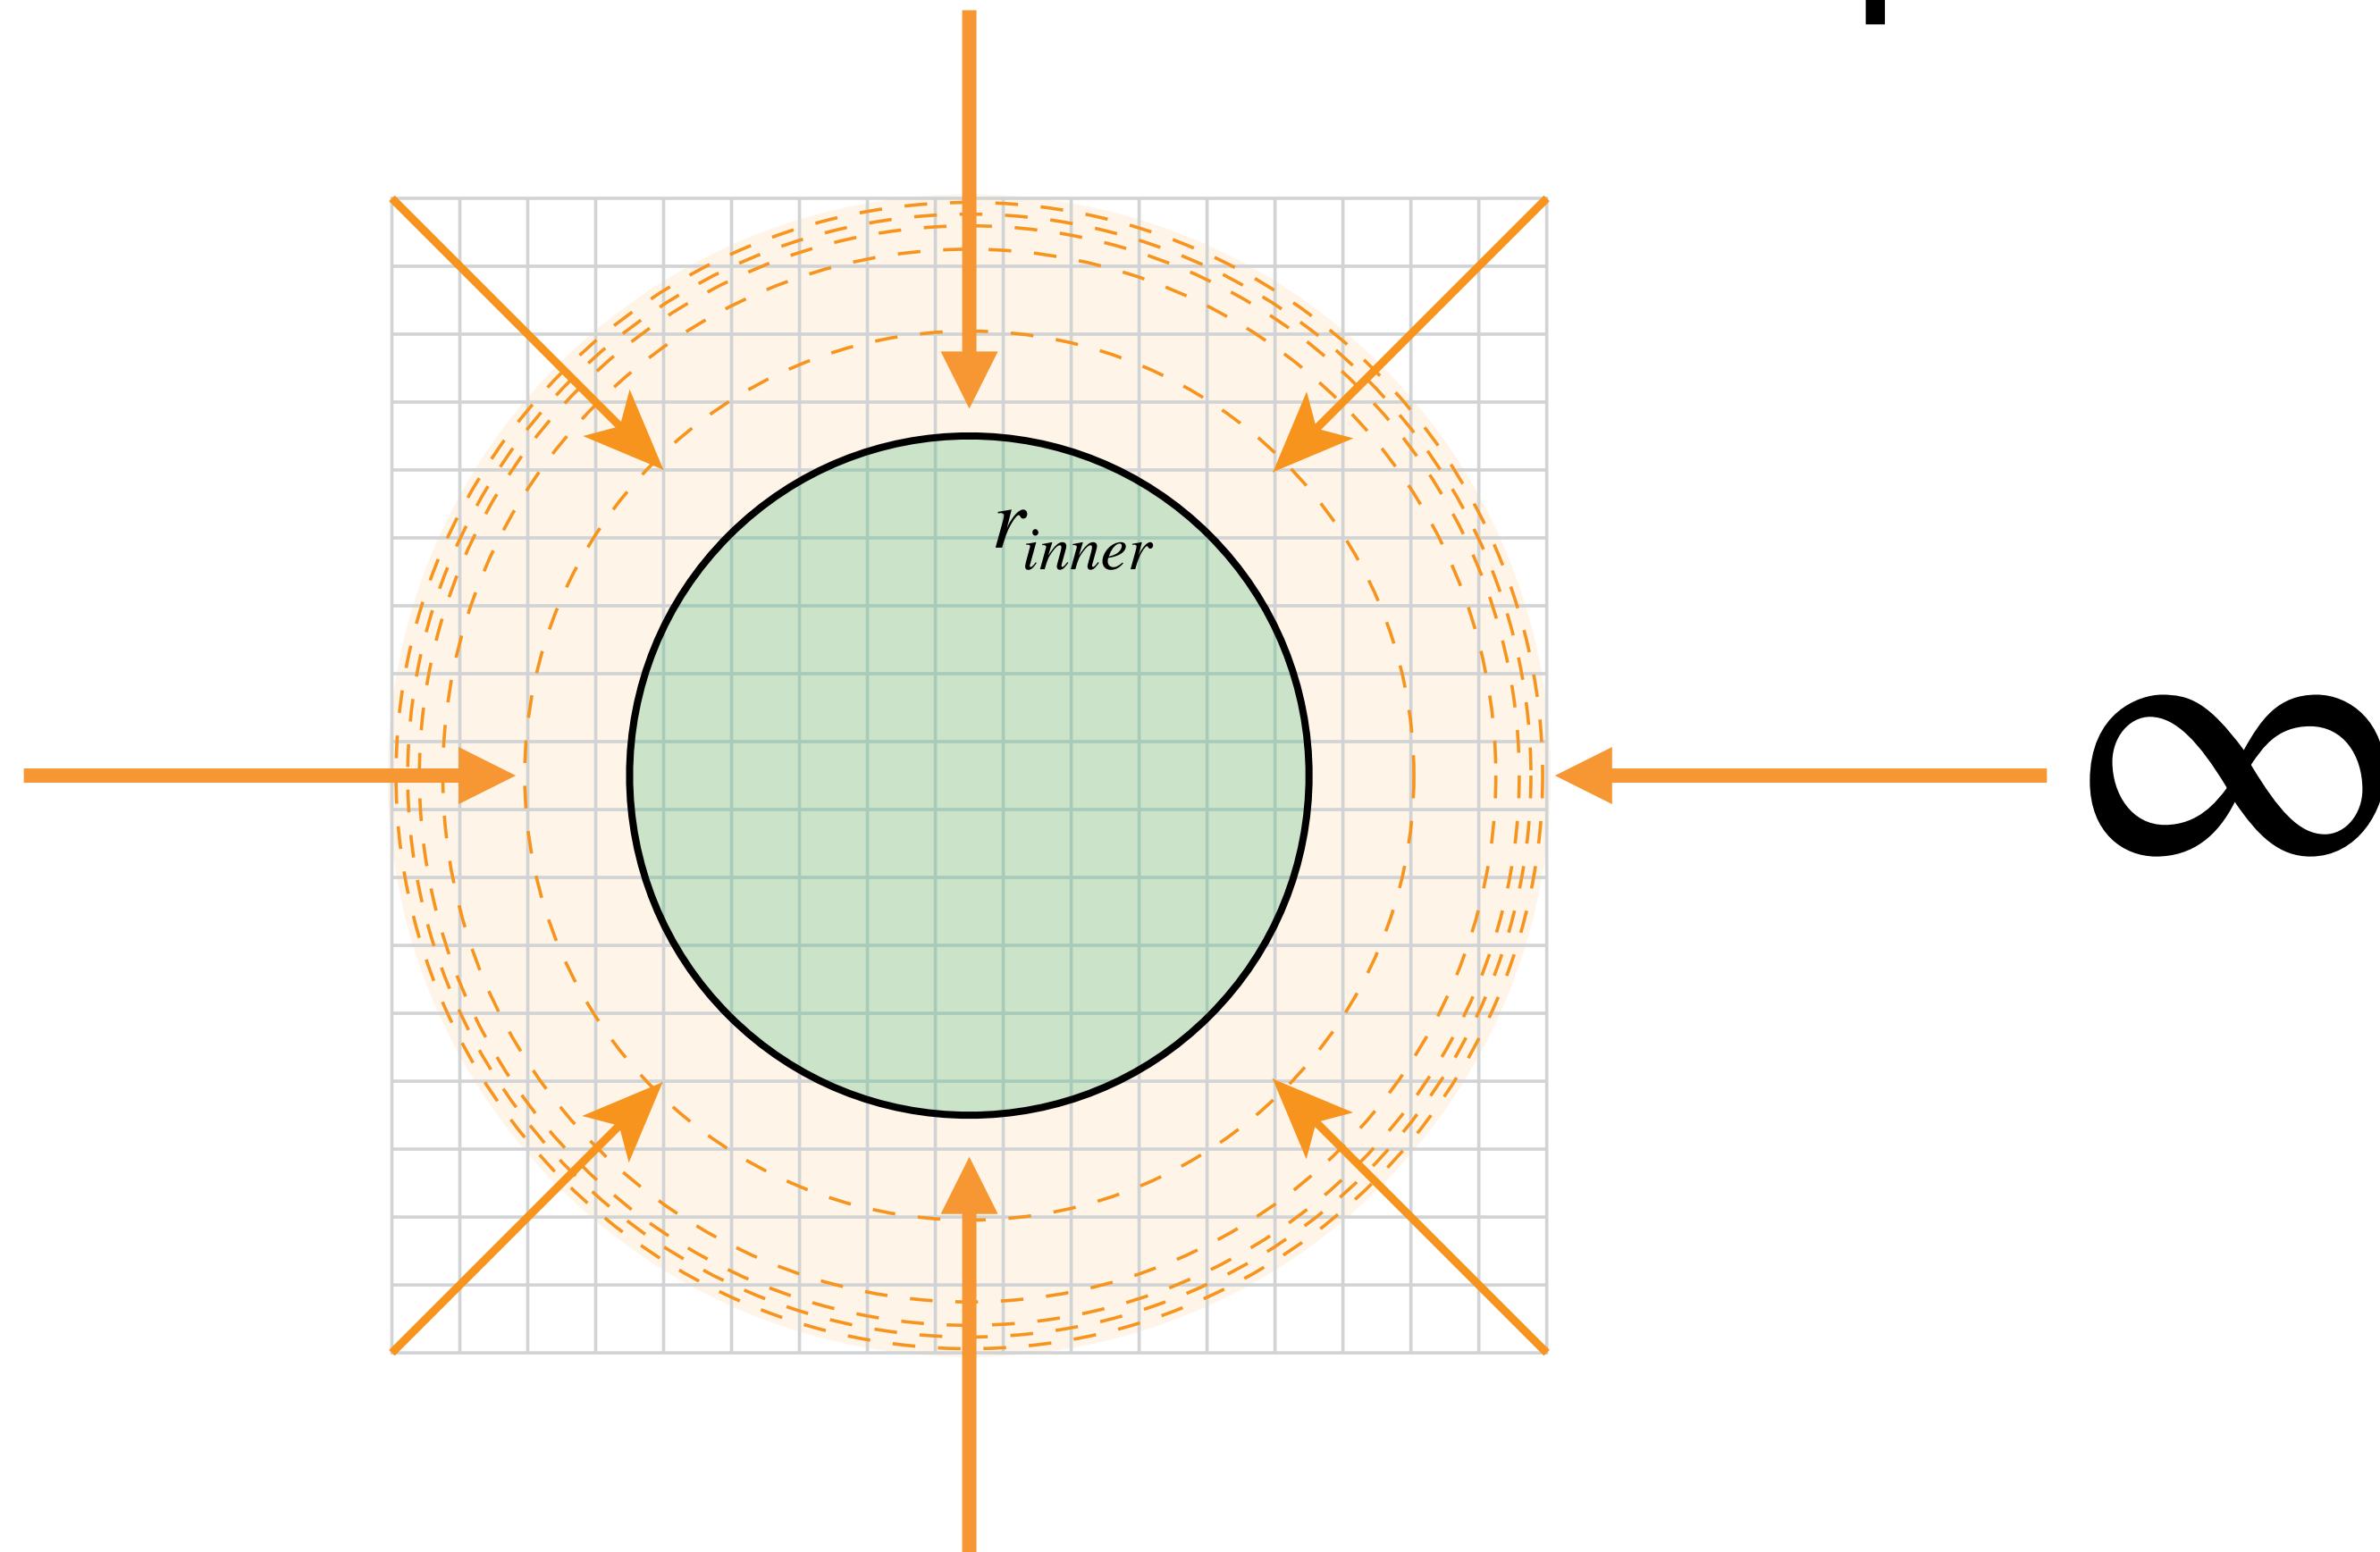


Scene  
Representation

$$\Phi : \mathbb{R}^3 \rightarrow \mathbb{R}^n$$

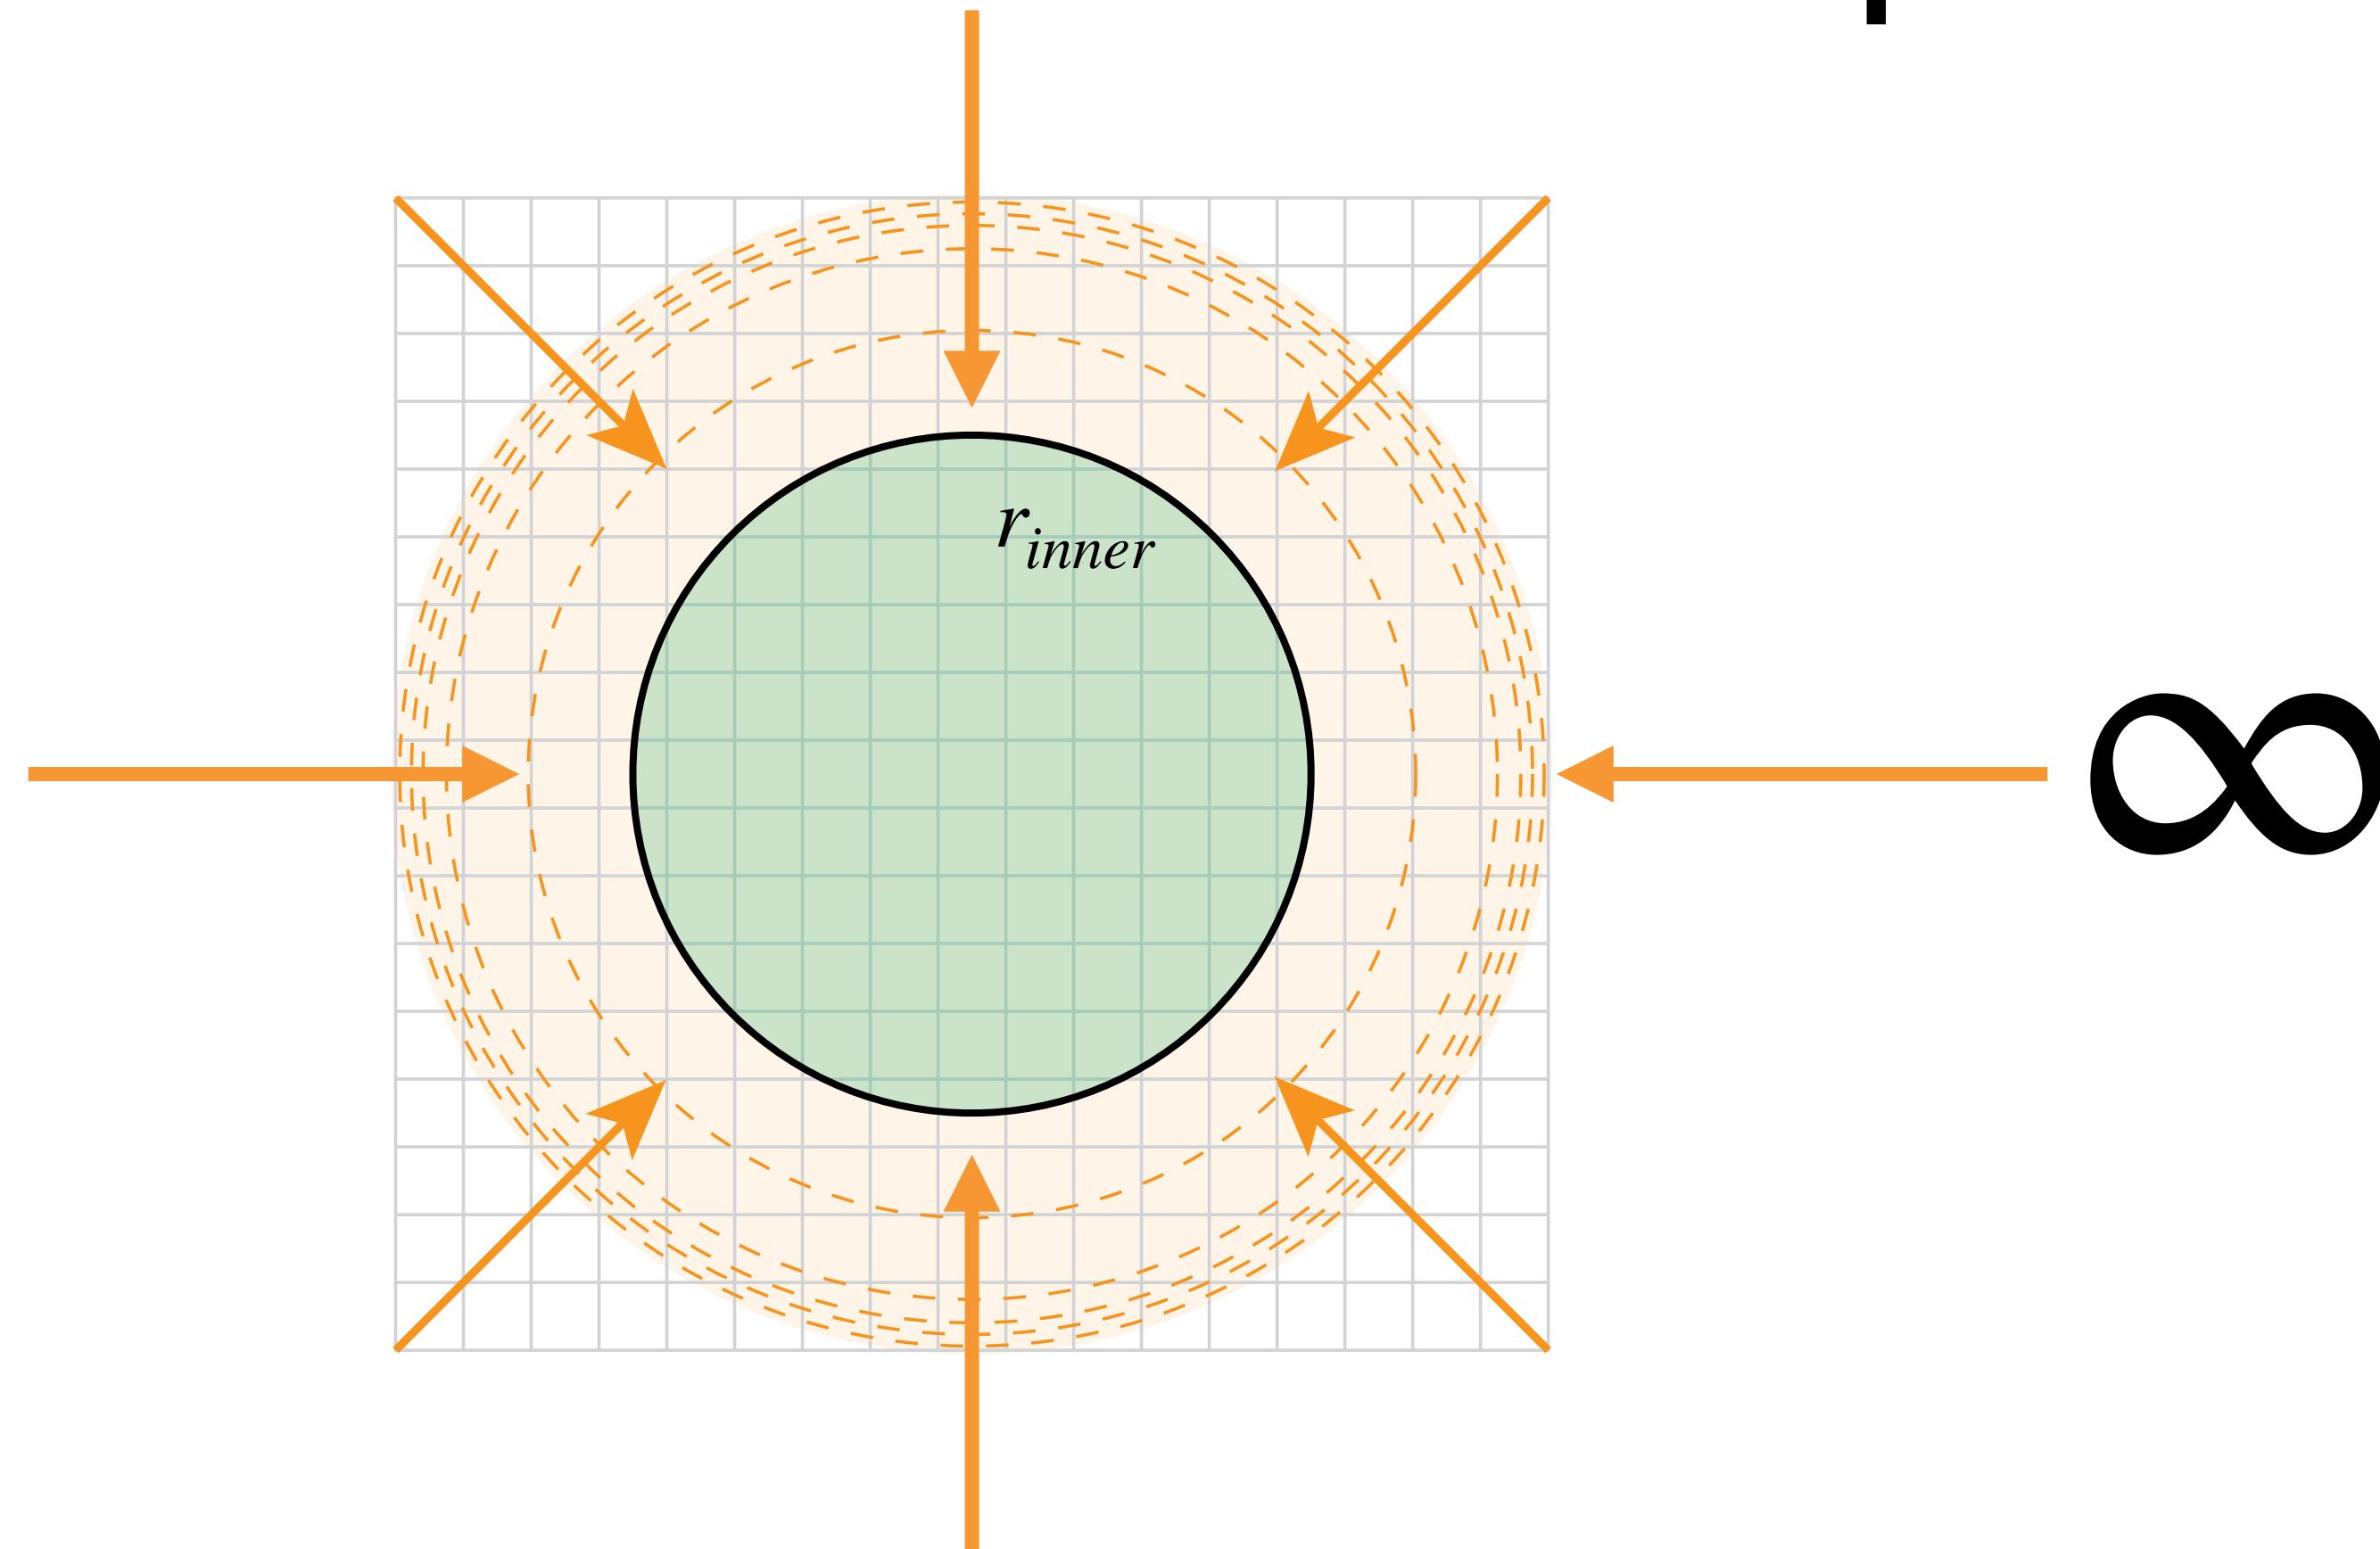


# Parameterizing Unbounded Scenes: Compactification



Come up with nonlinear mapping that “squashes” space, such that infinity is mapped to some **finite** value. Aligns with intuition “further away, less resolution”.

# Parameterizing Unbounded Scenes: Compactification

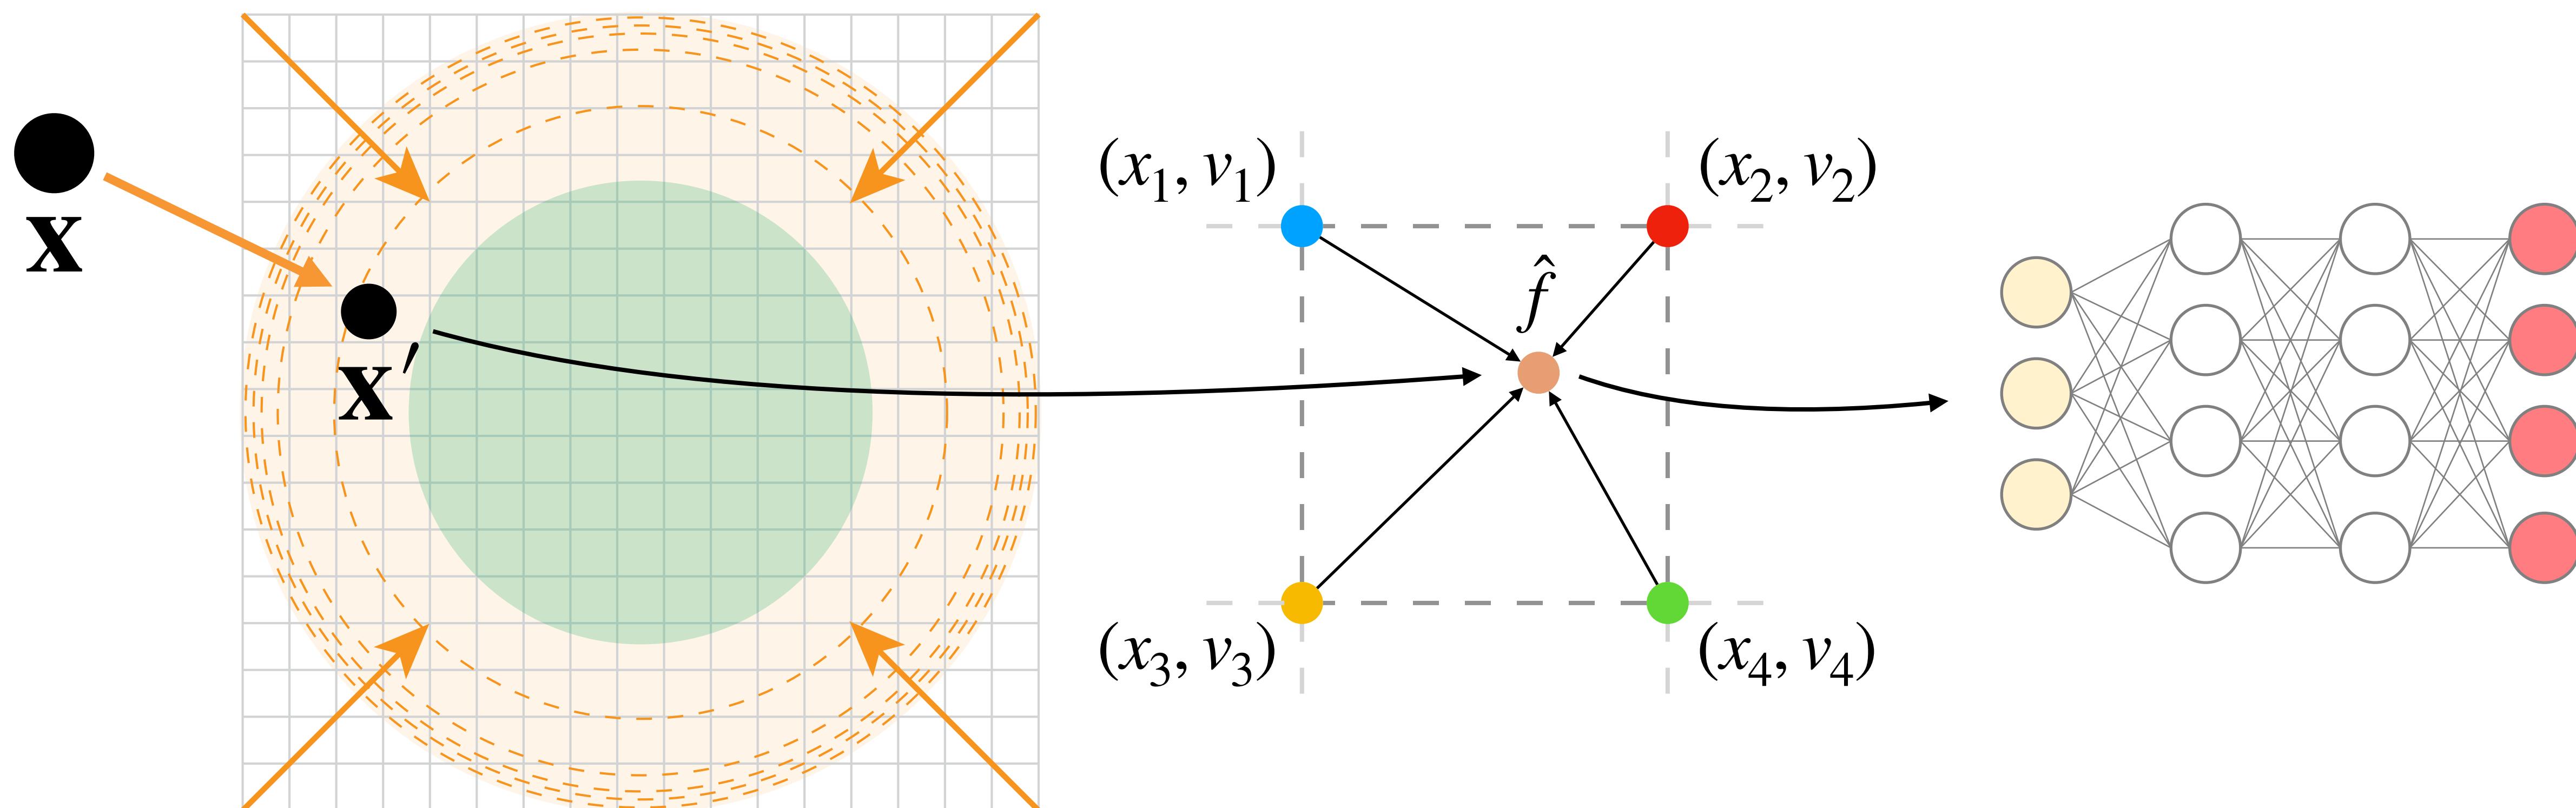


$$\mathbf{x}' = \frac{(1 + k) - k/\|\mathbf{u}\|}{\mathbf{u}/\|\mathbf{u}\|} r_{inner}$$
$$\mathbf{u} = \frac{\mathbf{x}}{r_{inner}}$$

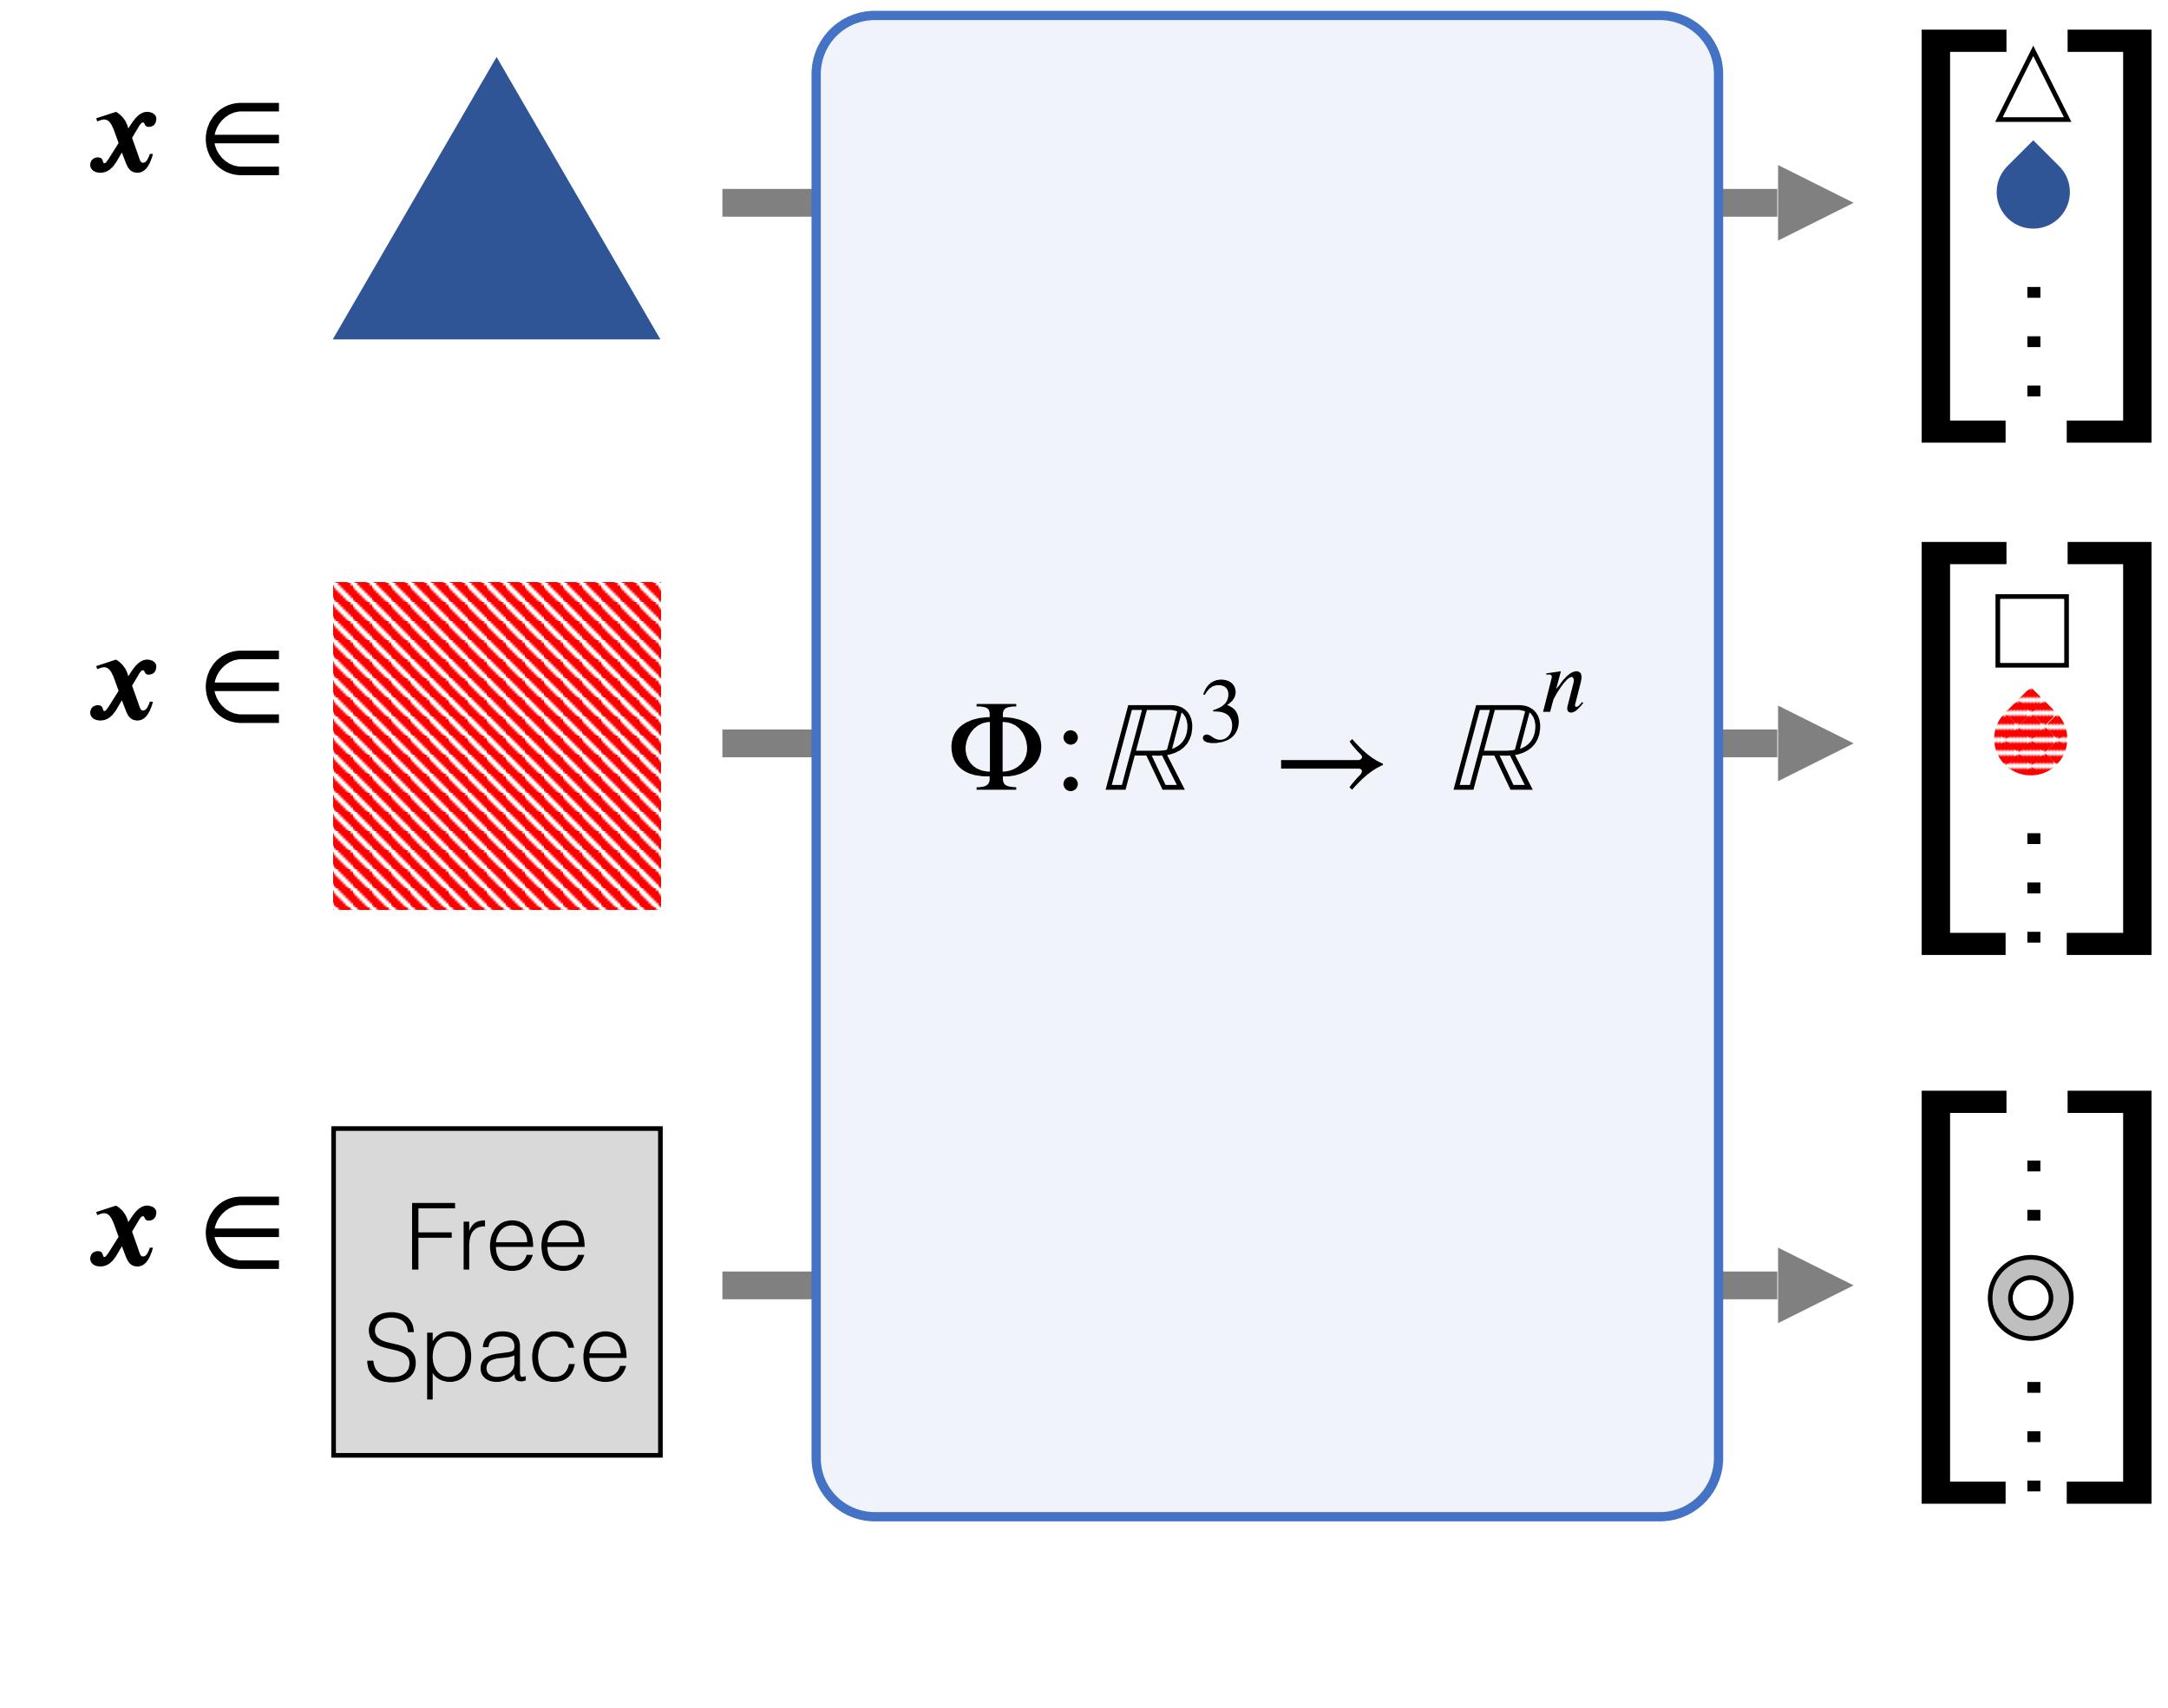
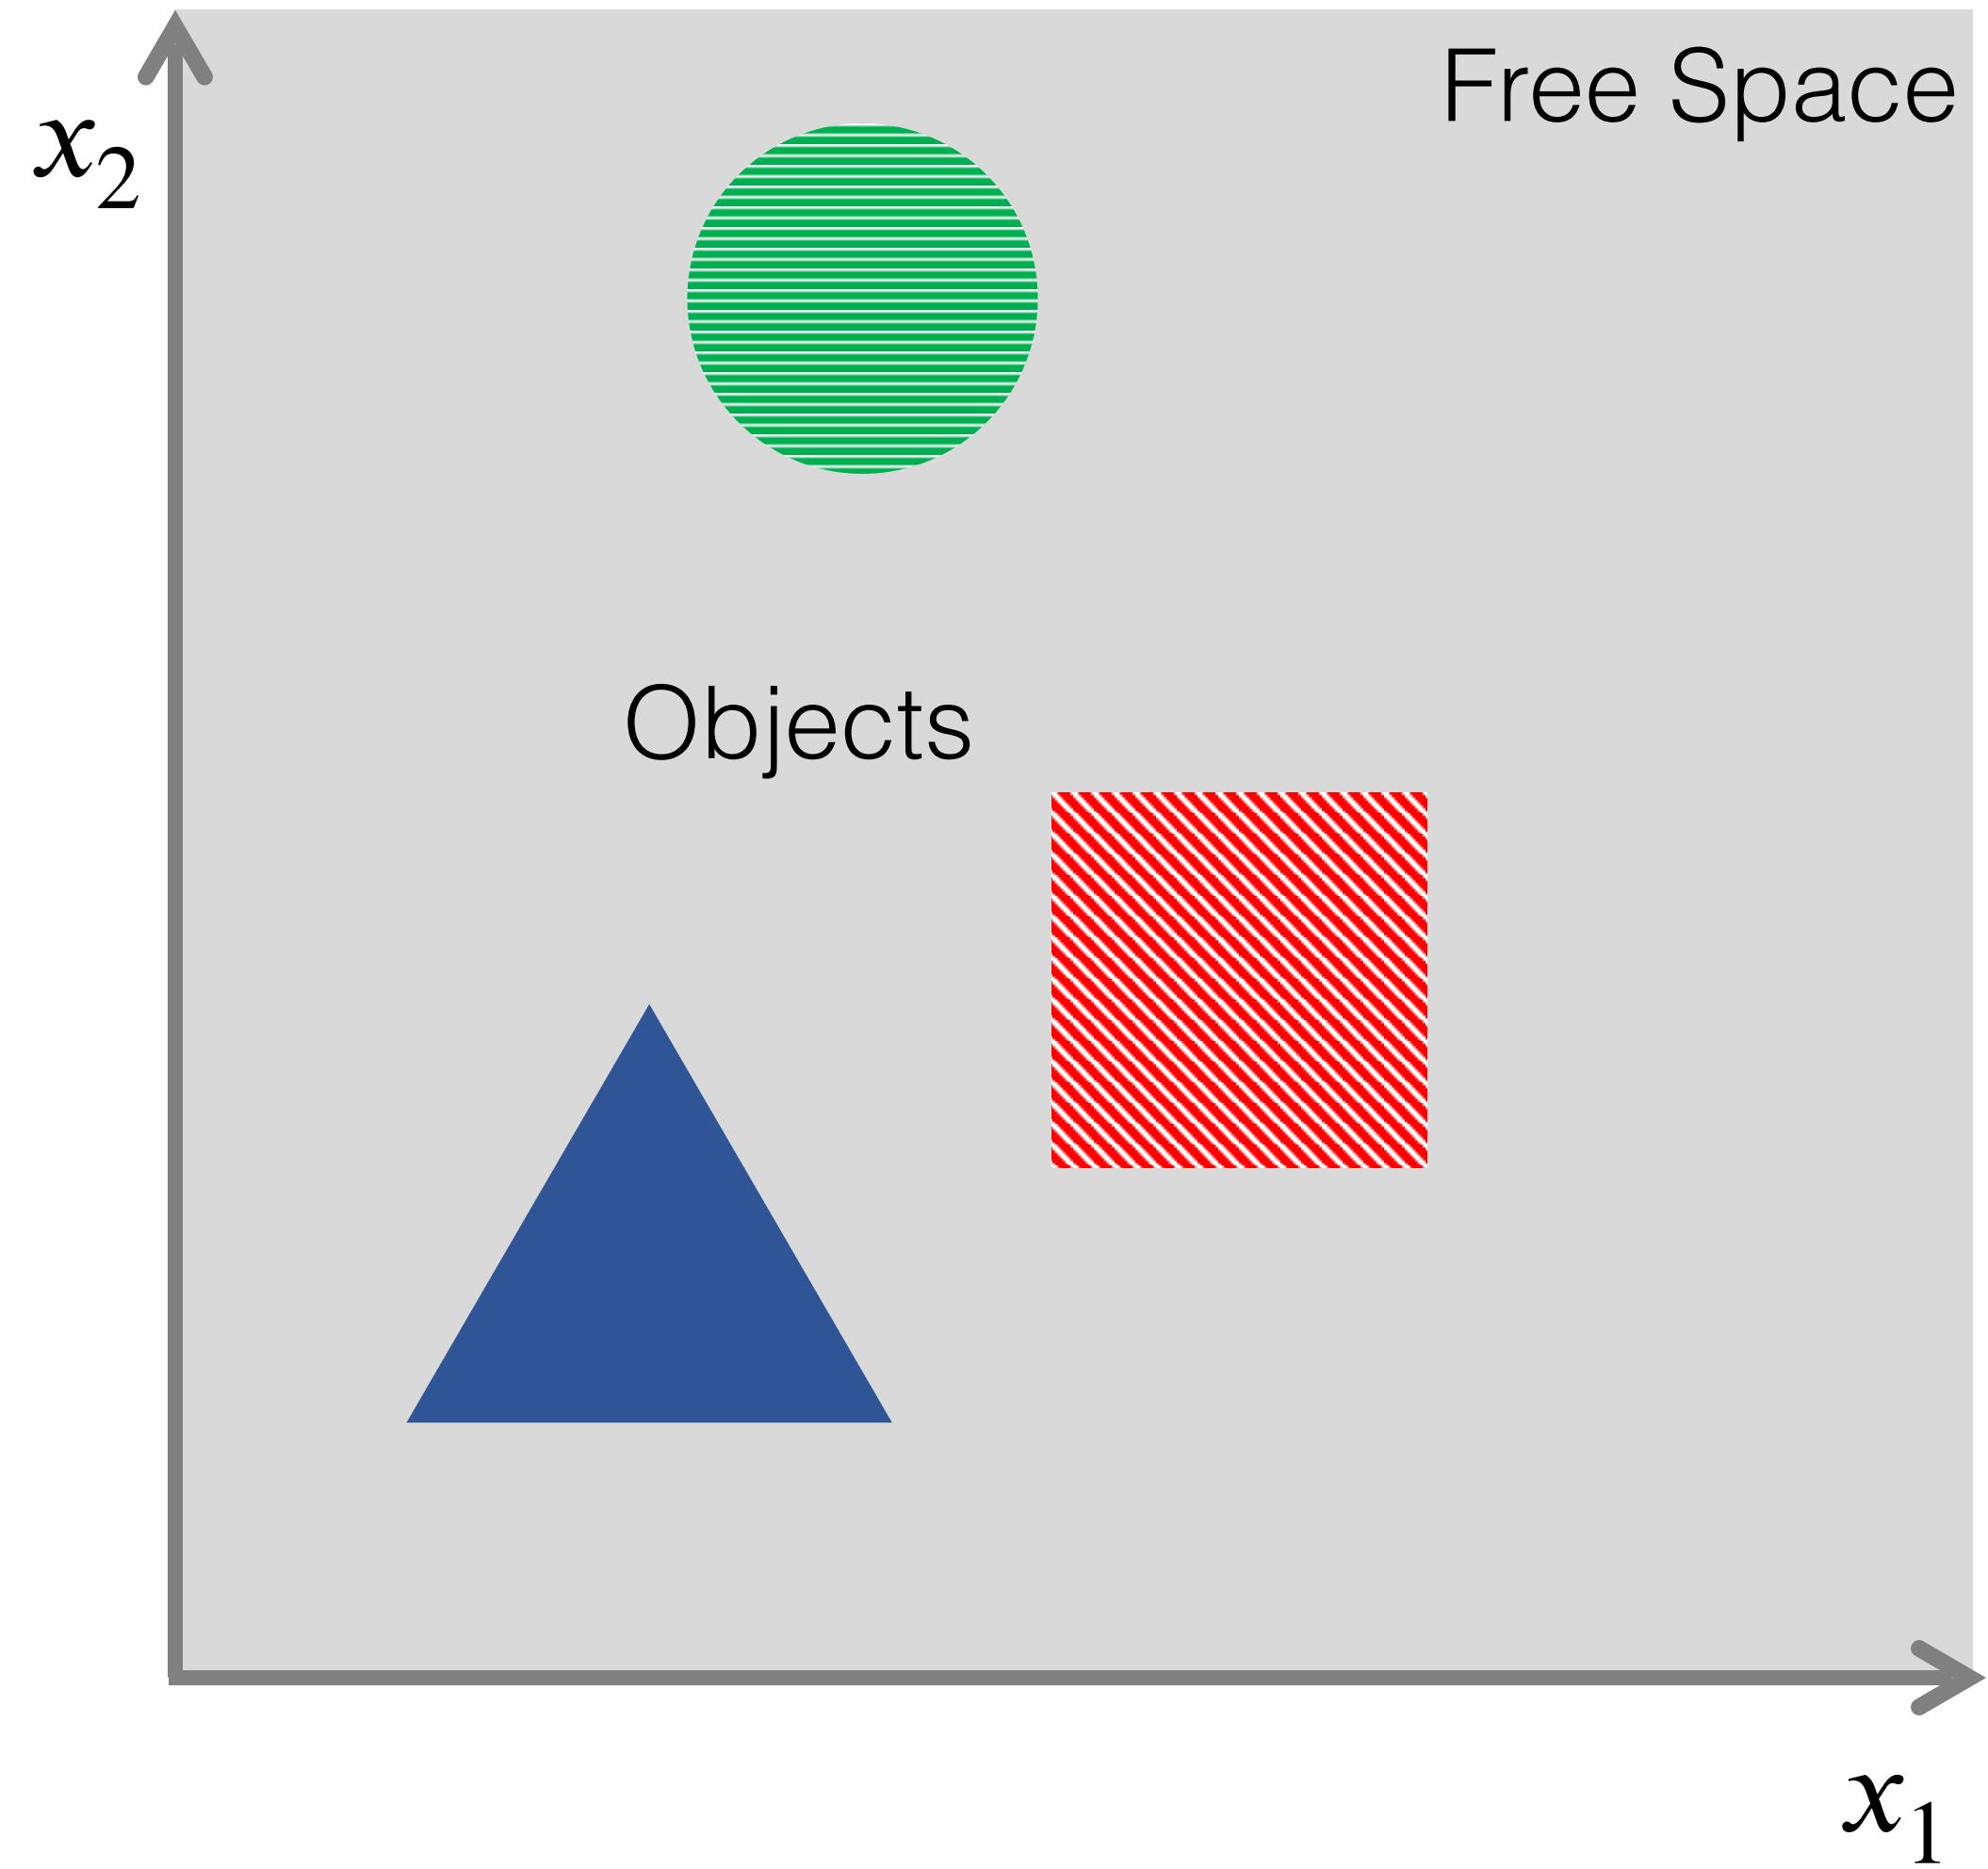
Example Compactification function from  
“Mip-NeRF 360: Unbounded Anti-Aliased Neural  
Radiance Fields”, Barron et al., CVPR 2022

Come up with nonlinear mapping that “squashes” space, such that infinity is mapped to some **finite** value. Aligns with intuition “further away, less resolution”.

# Parameterizing Unbounded Scenes: Compactification



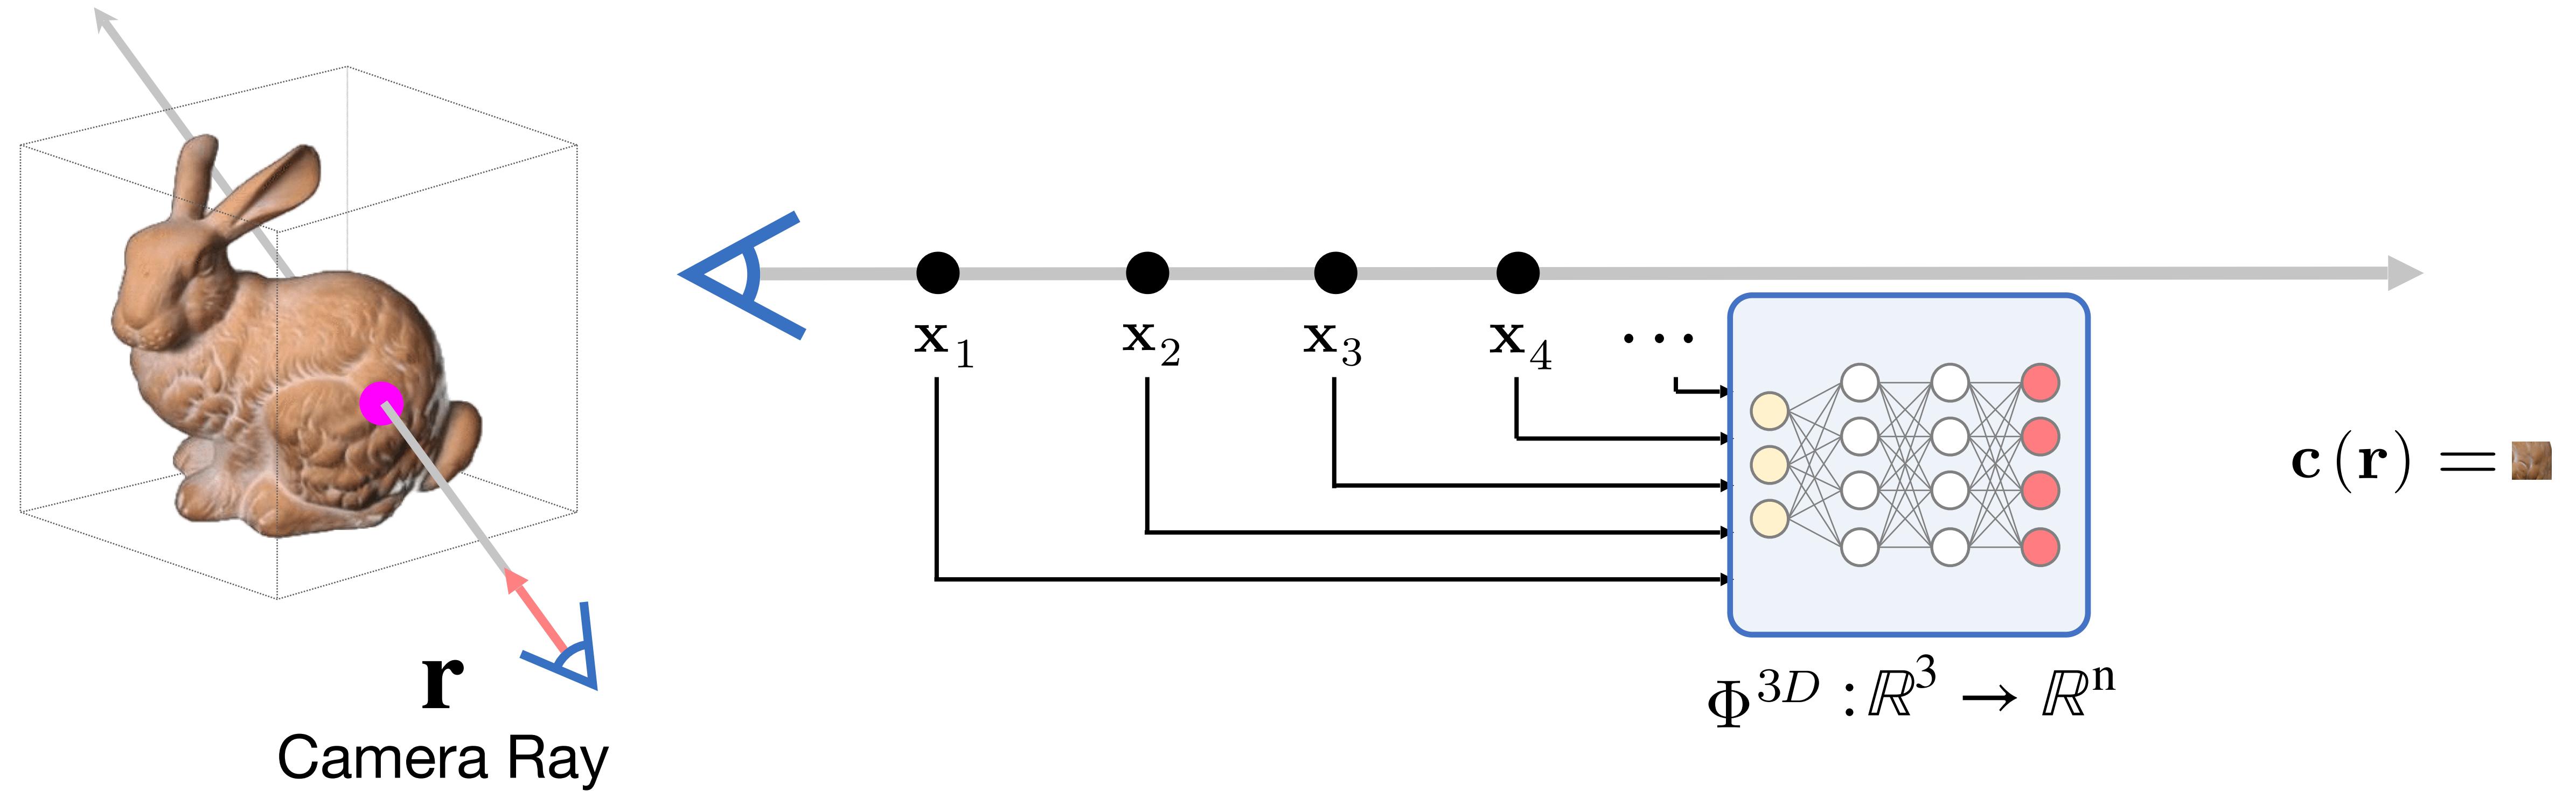
# Do we have to parameterize 3D Scene as 3D Function?



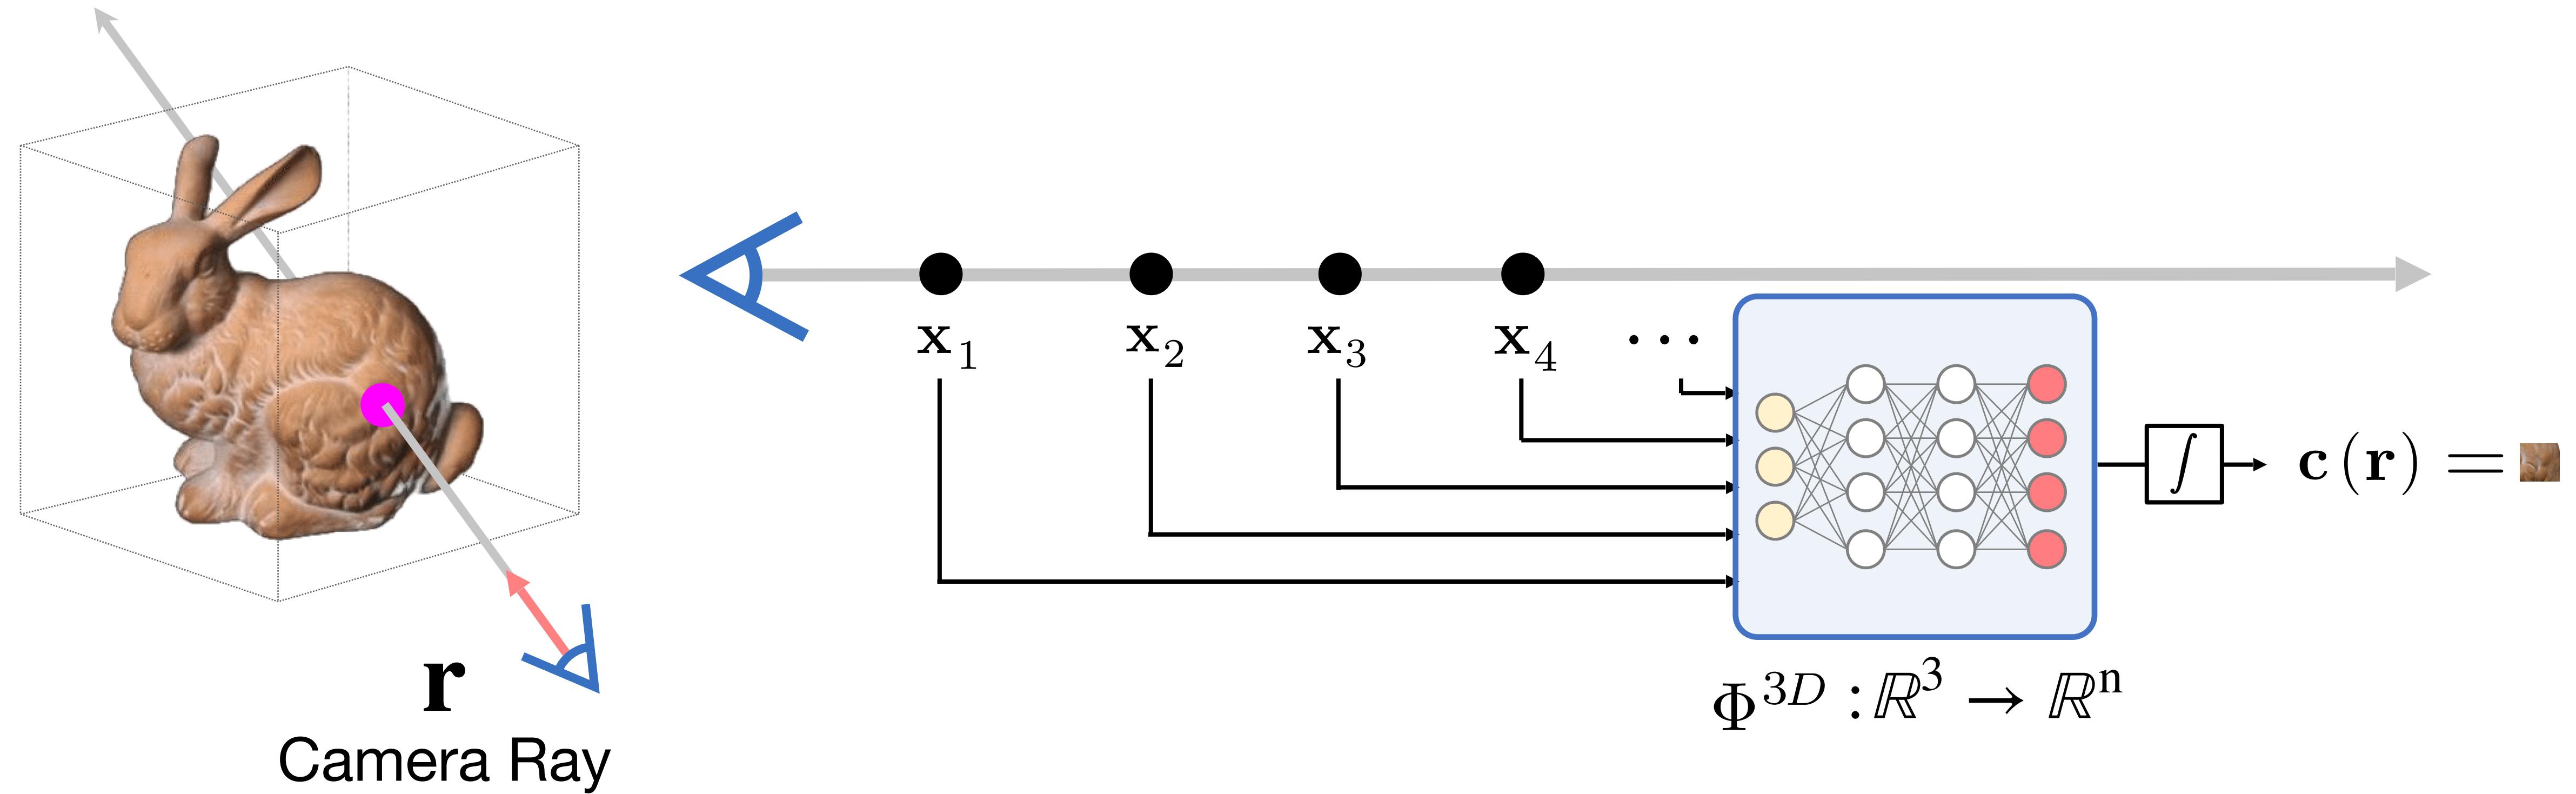
# 3D-structured Representation

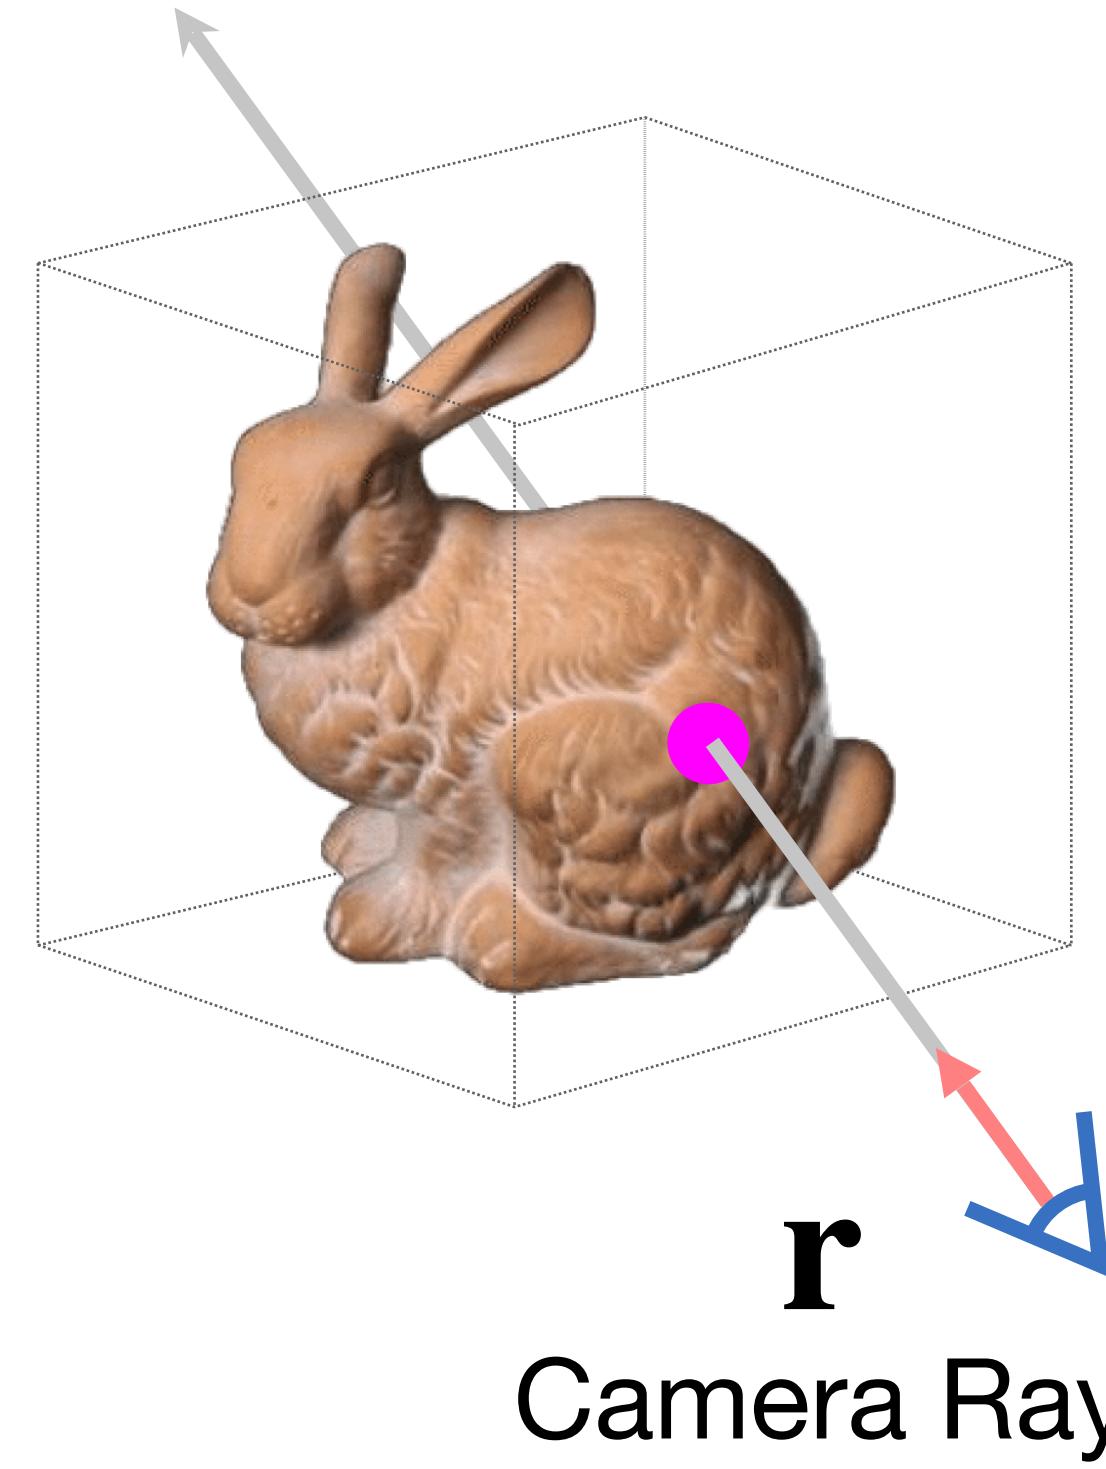


# 3D-structured Representation



# 3D-structured Representation





A

$x_1$

$x_2$

$x_3$

$x_4$

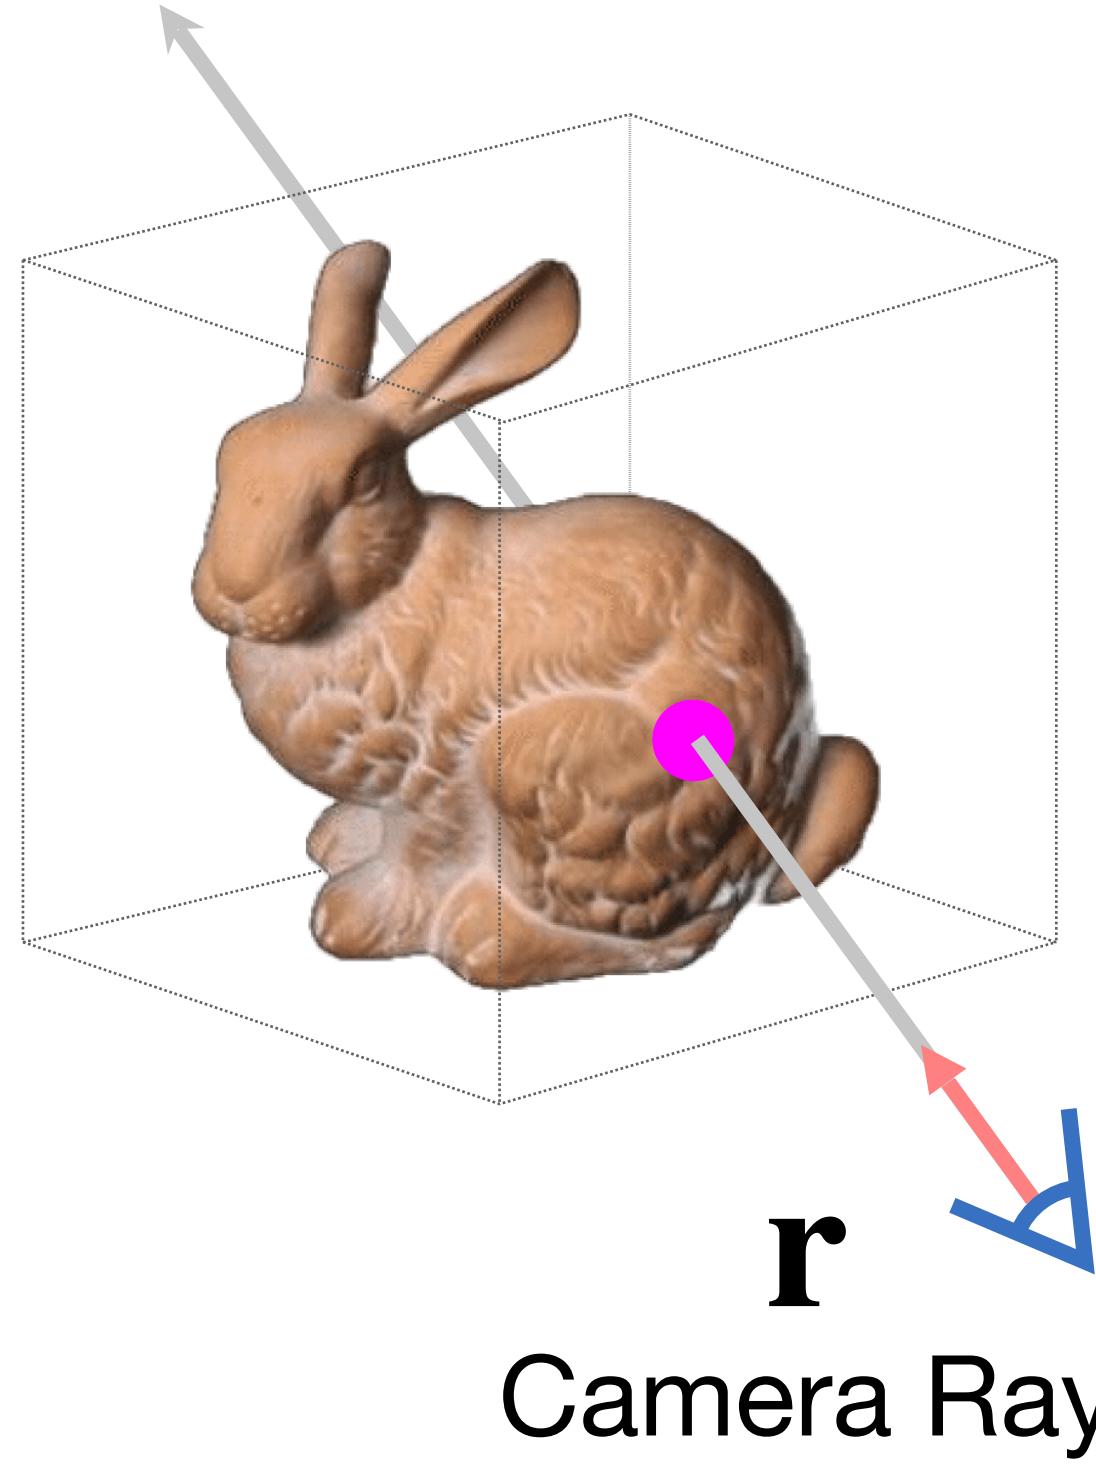
...

## Light Field

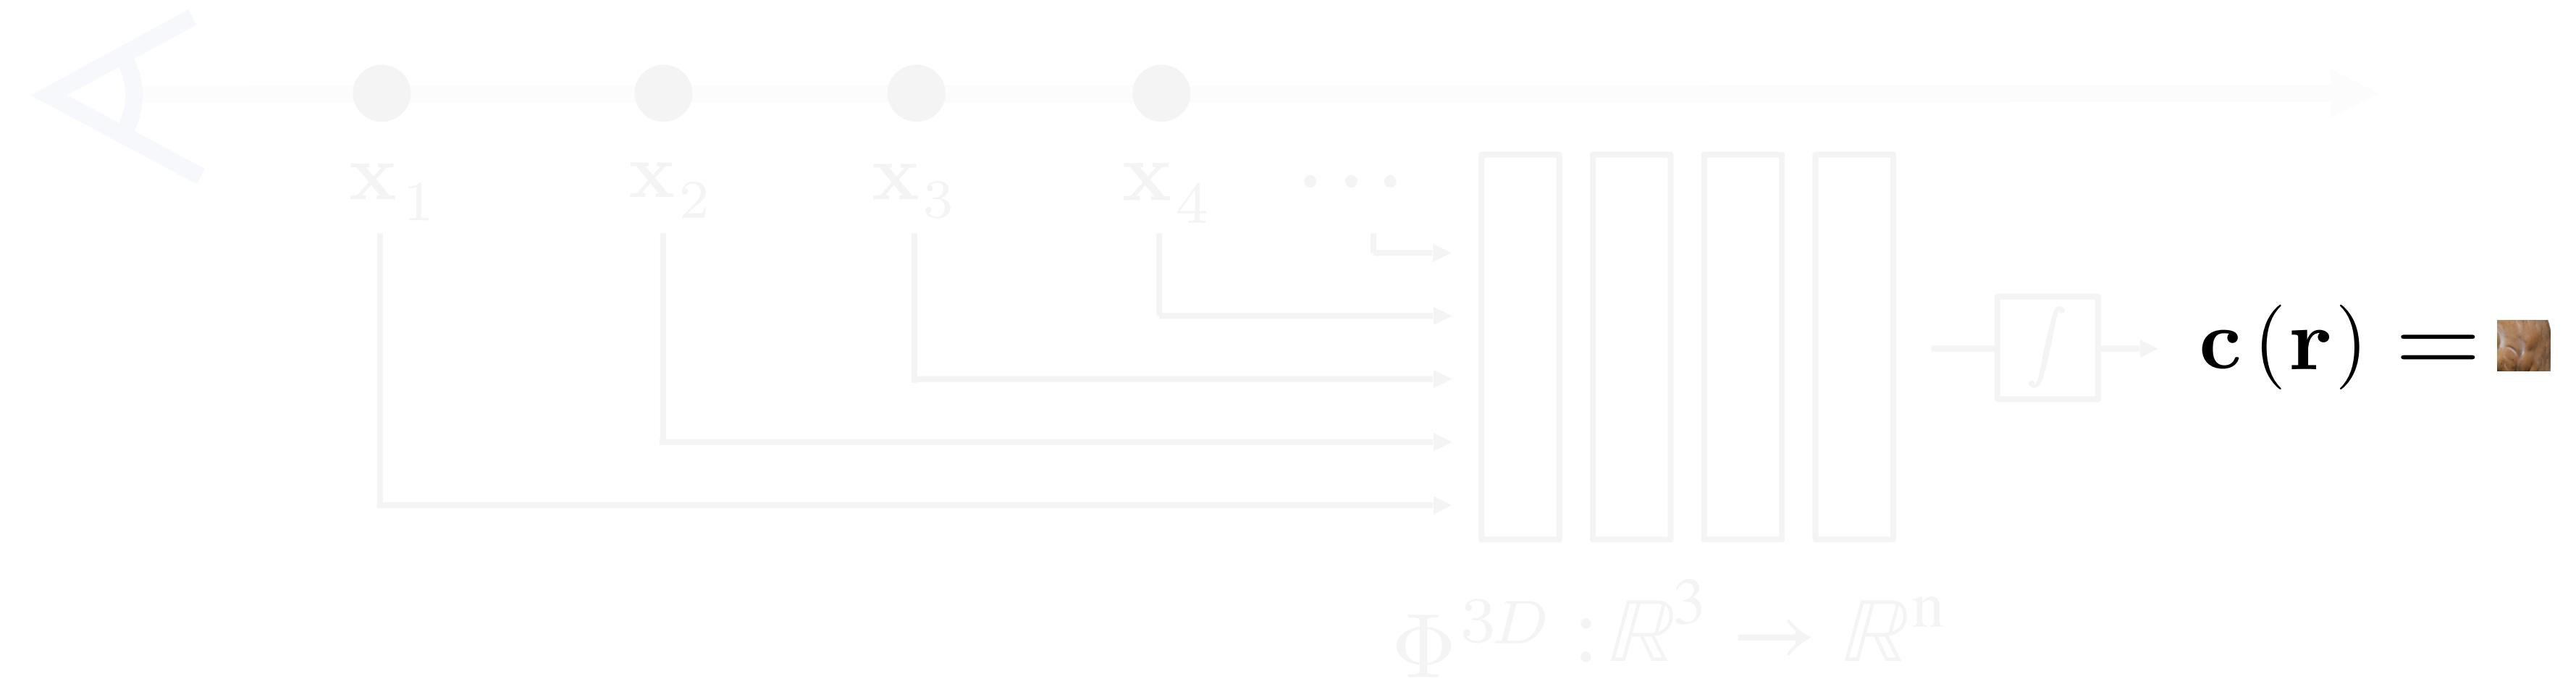
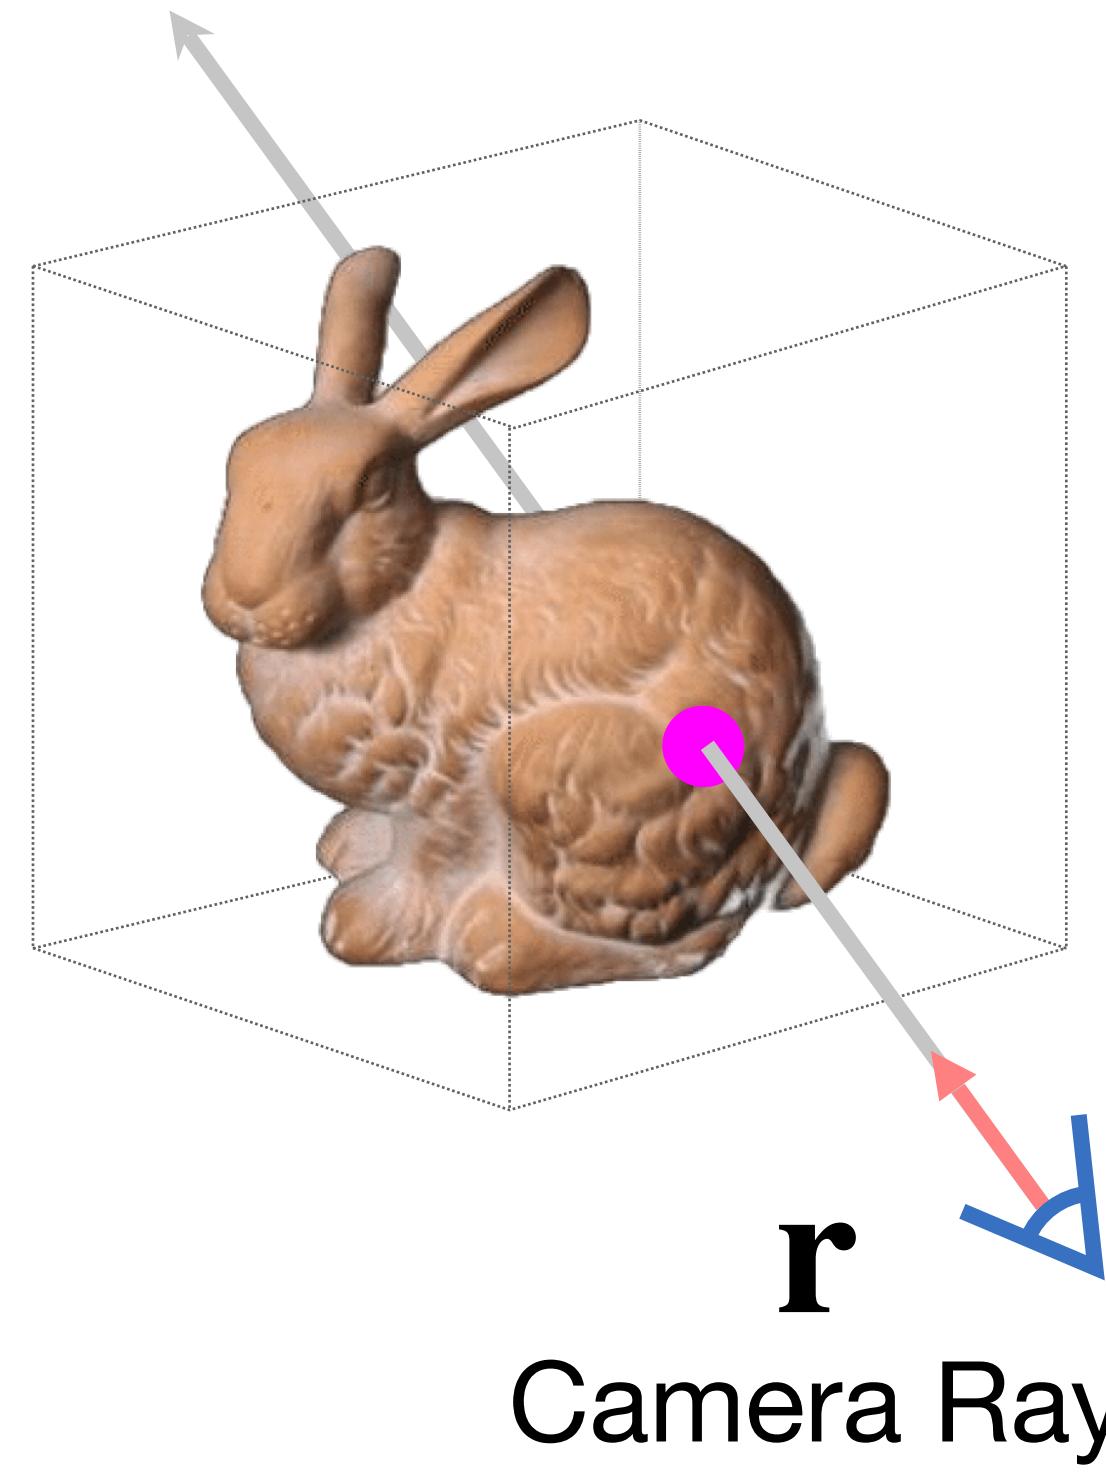
$$\Phi^{3D} : \mathbb{R}^3 \rightarrow \mathbb{R}^n$$



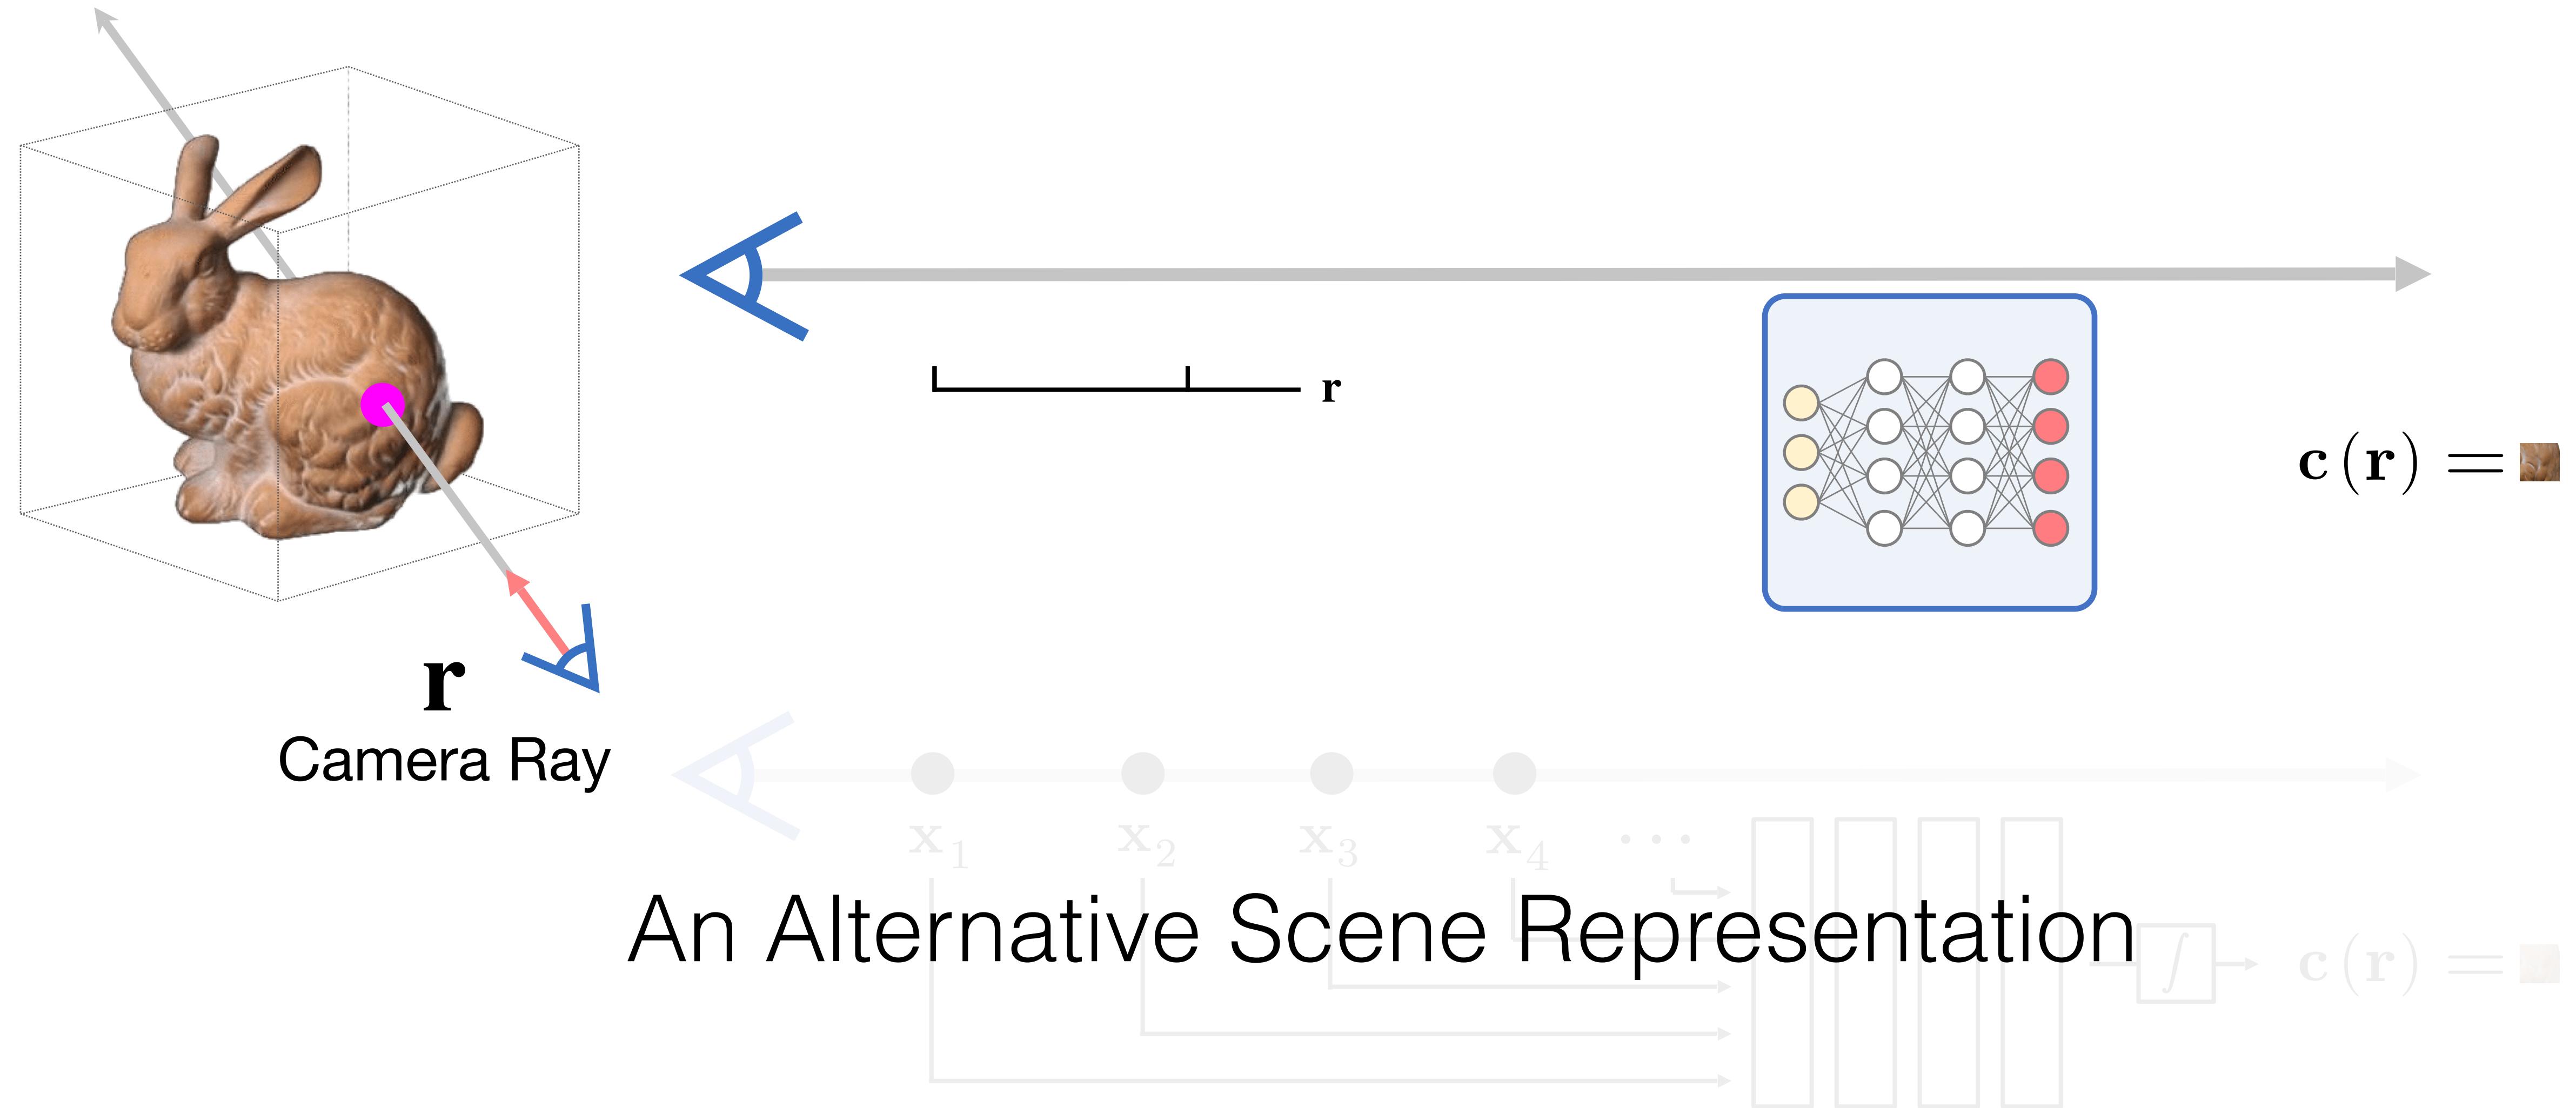
$$c(r) =$$



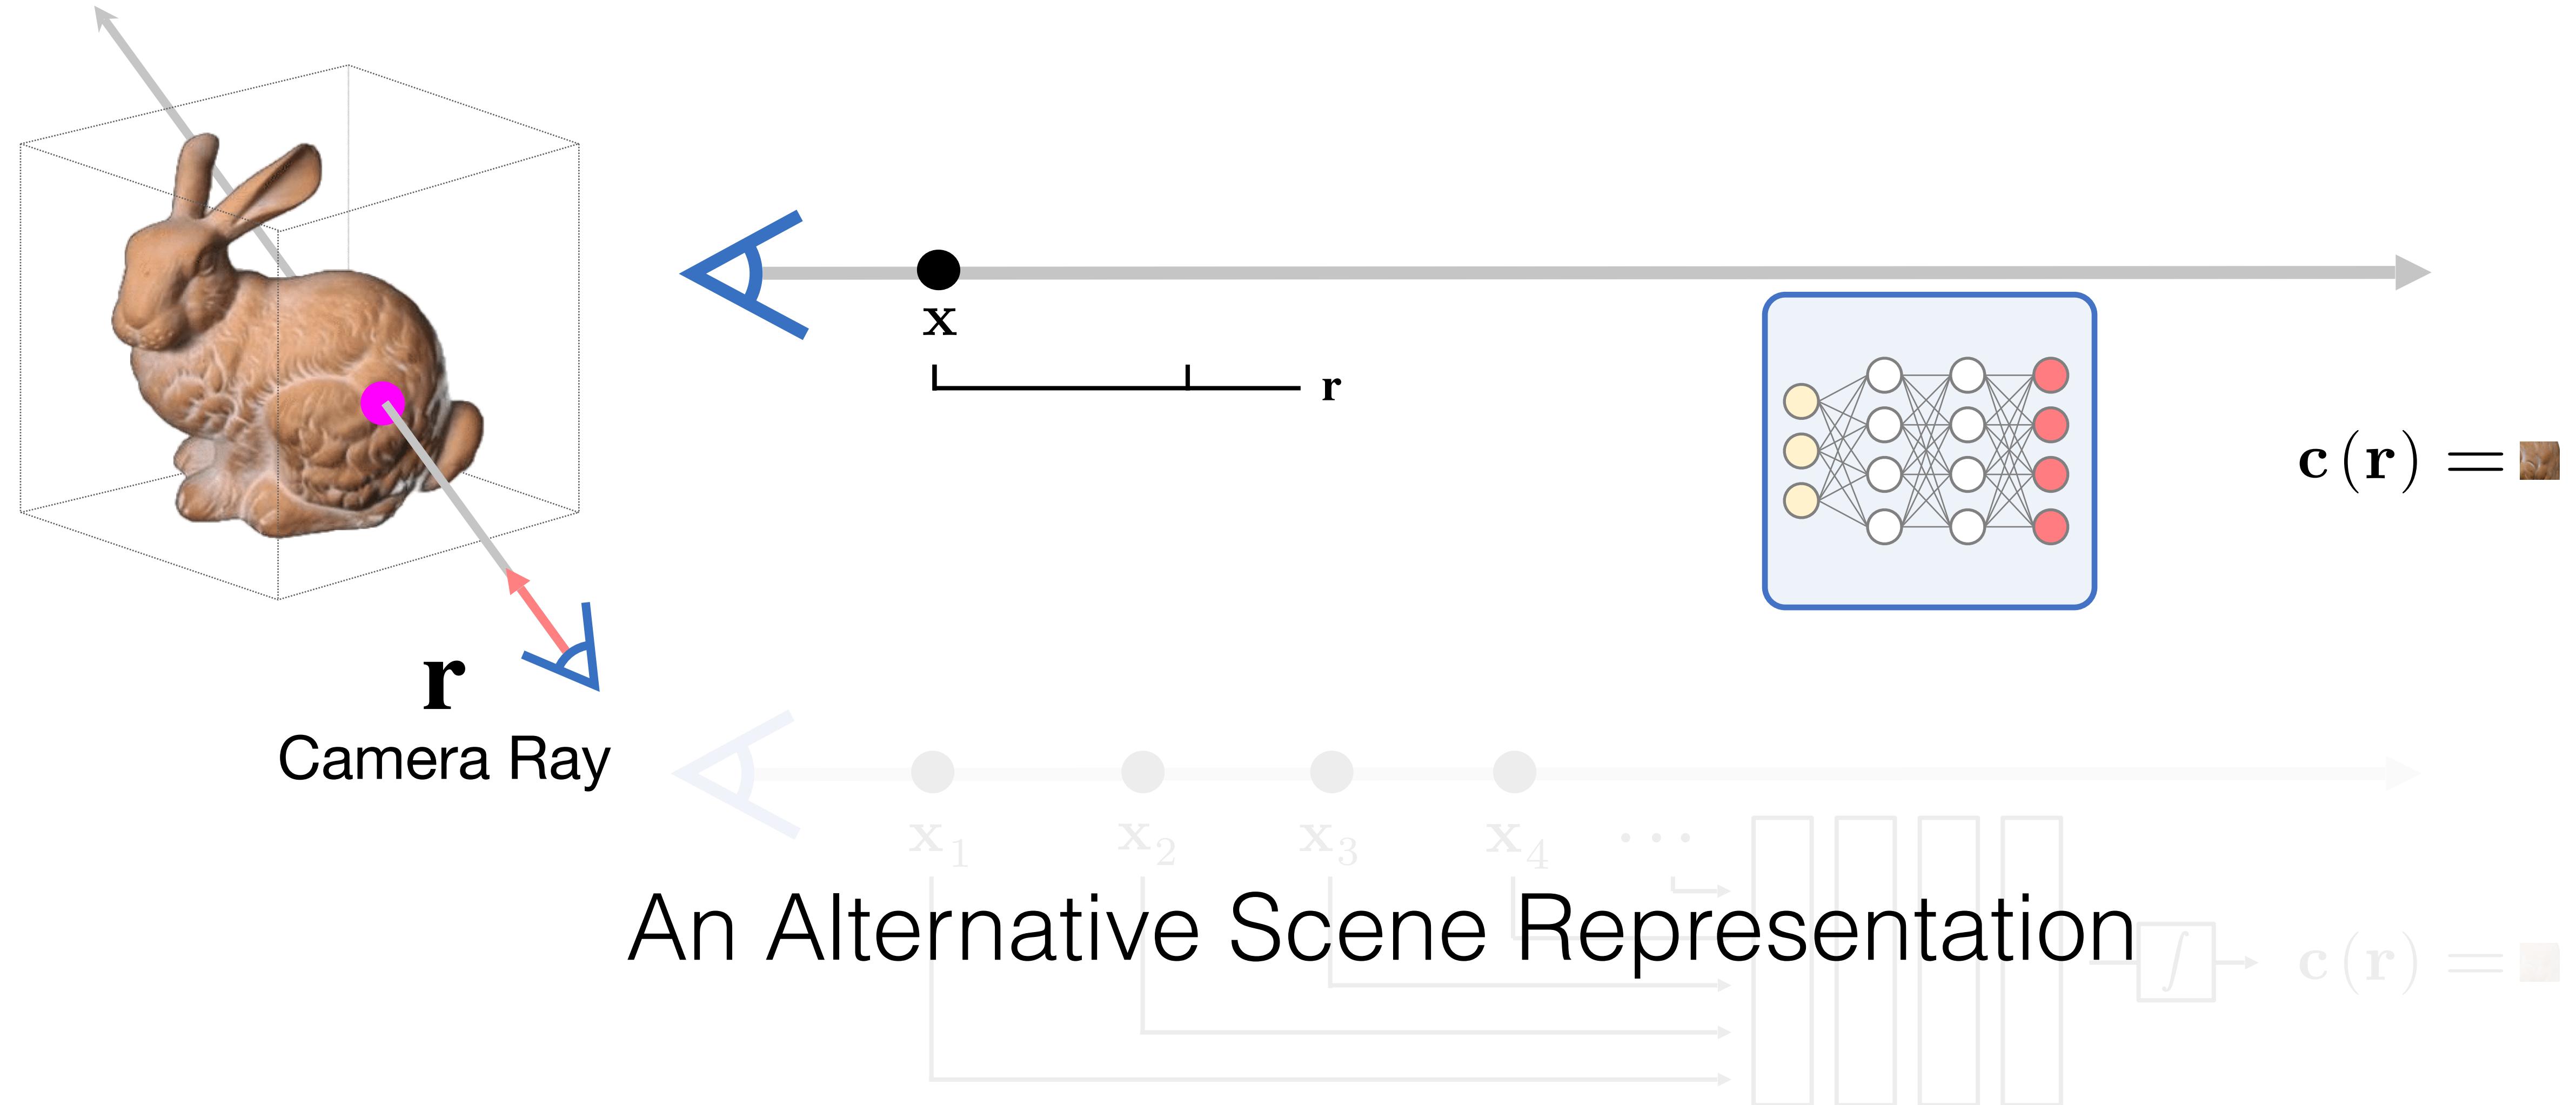
# Light Field Networks



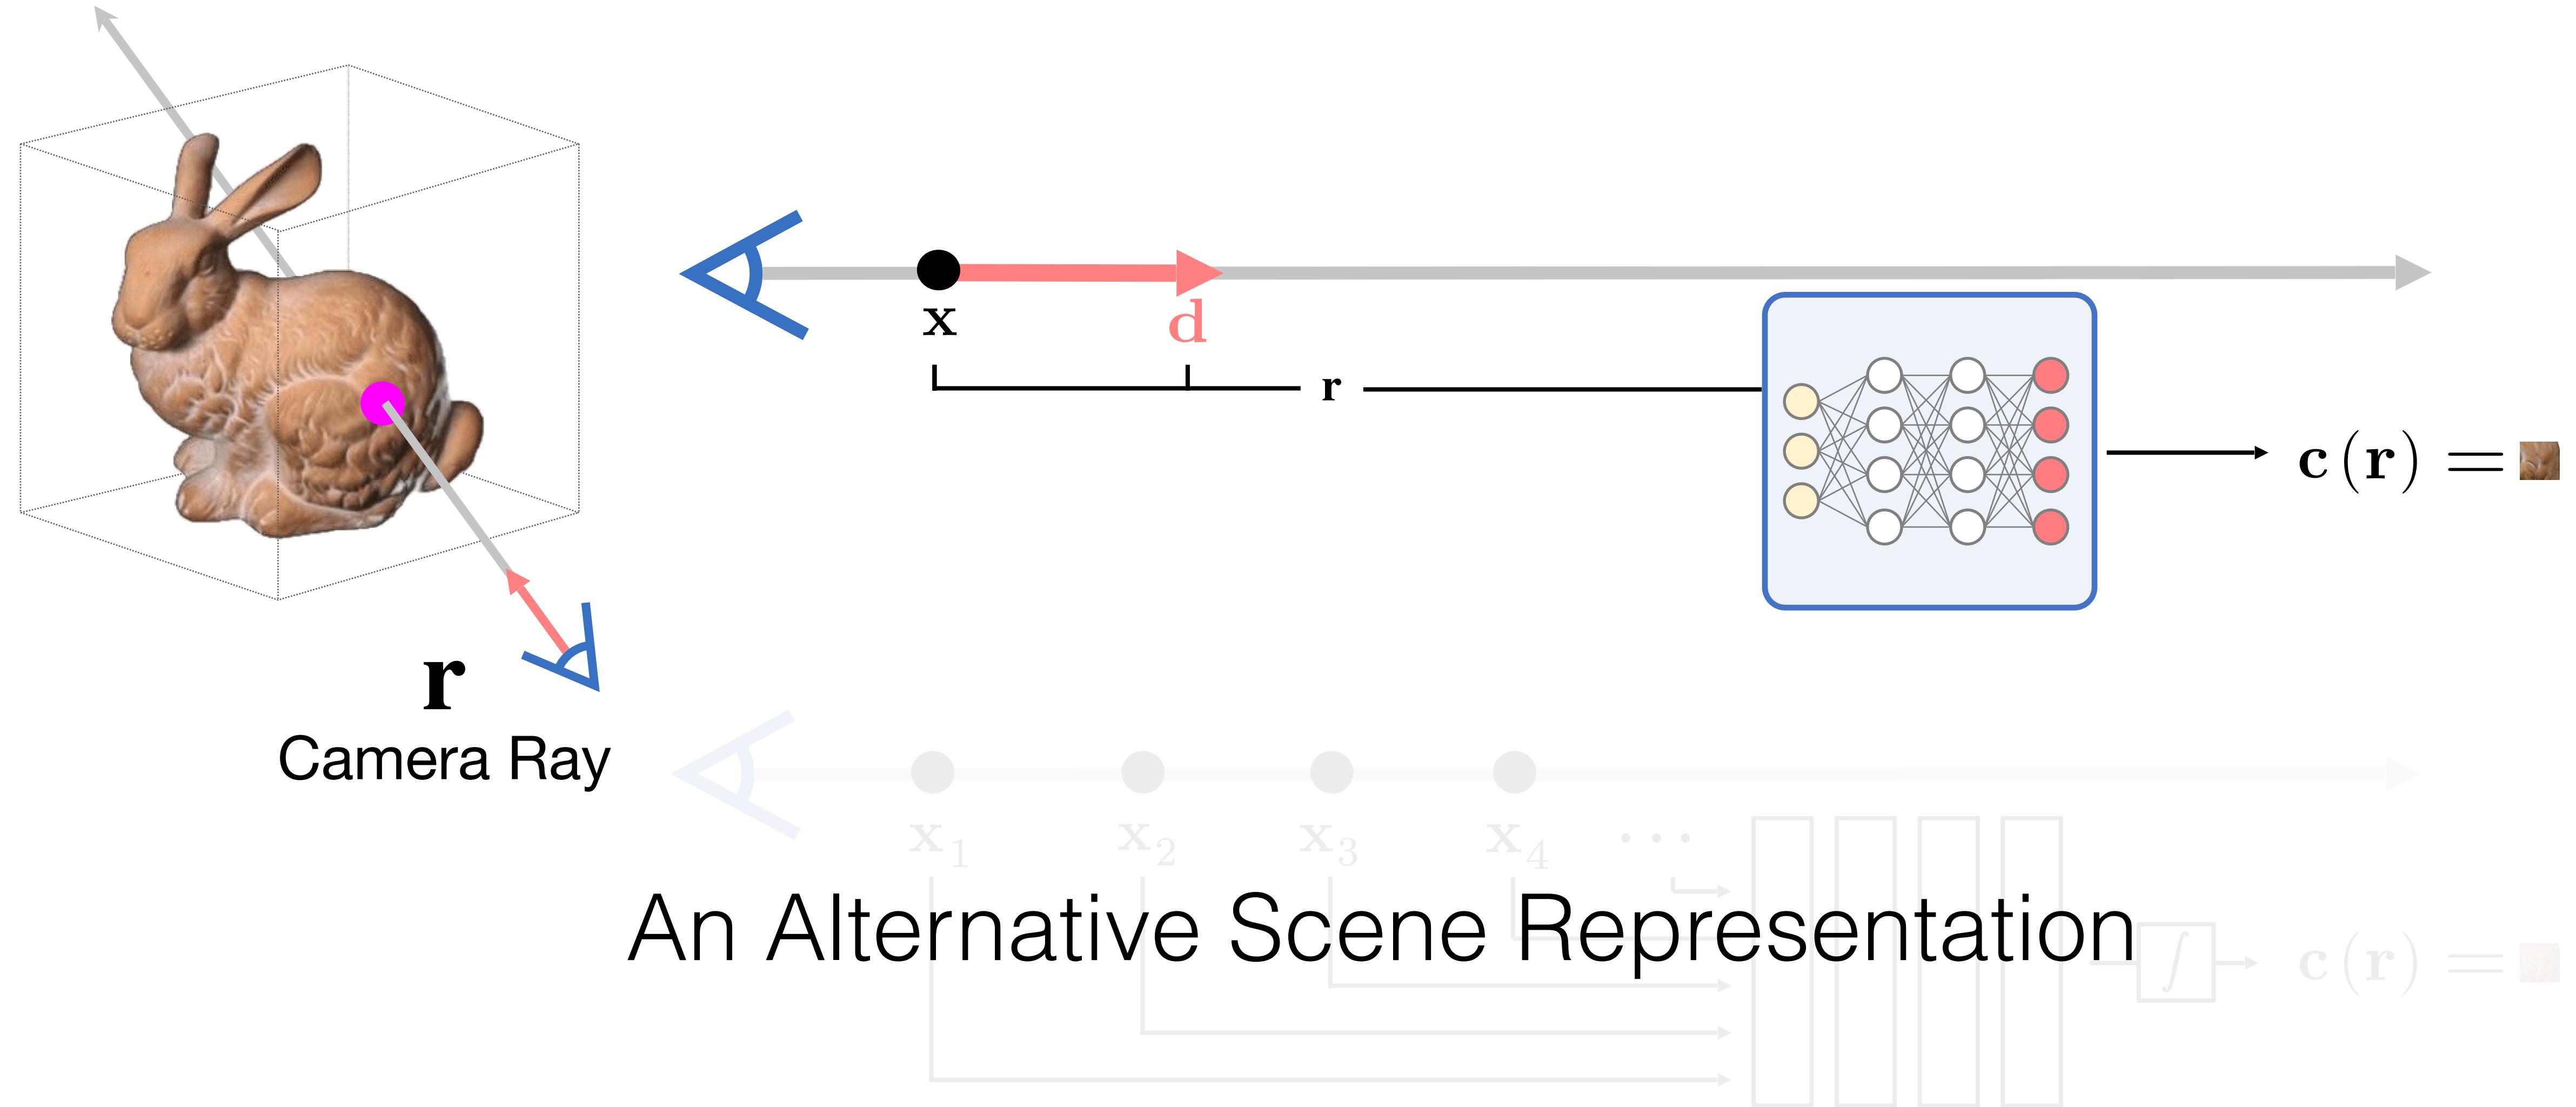
# Light Field Networks



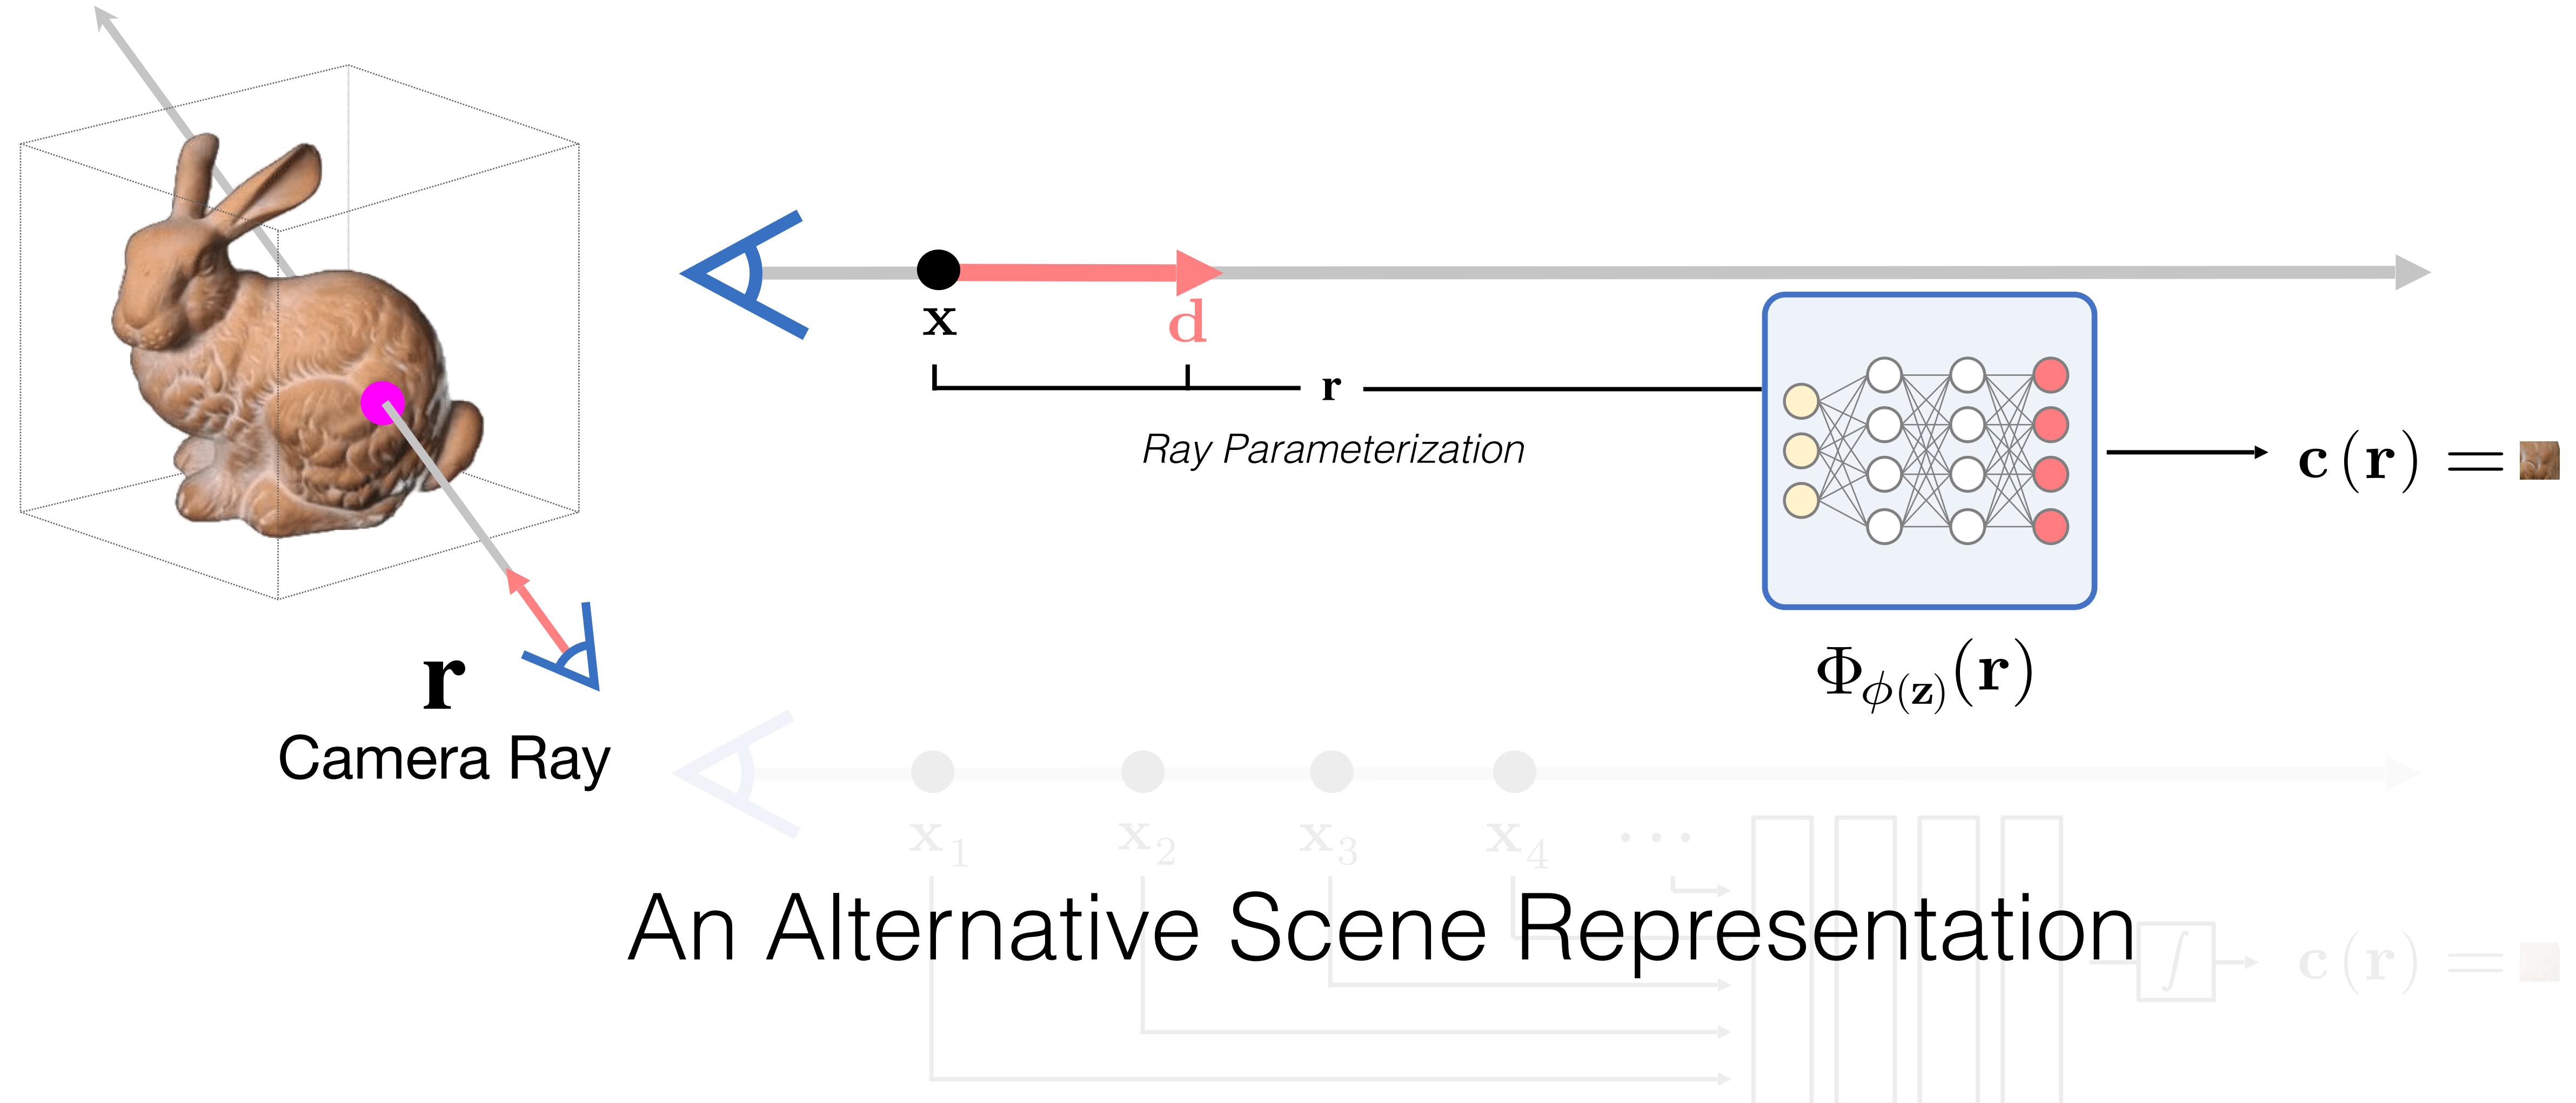
# Light Field Networks



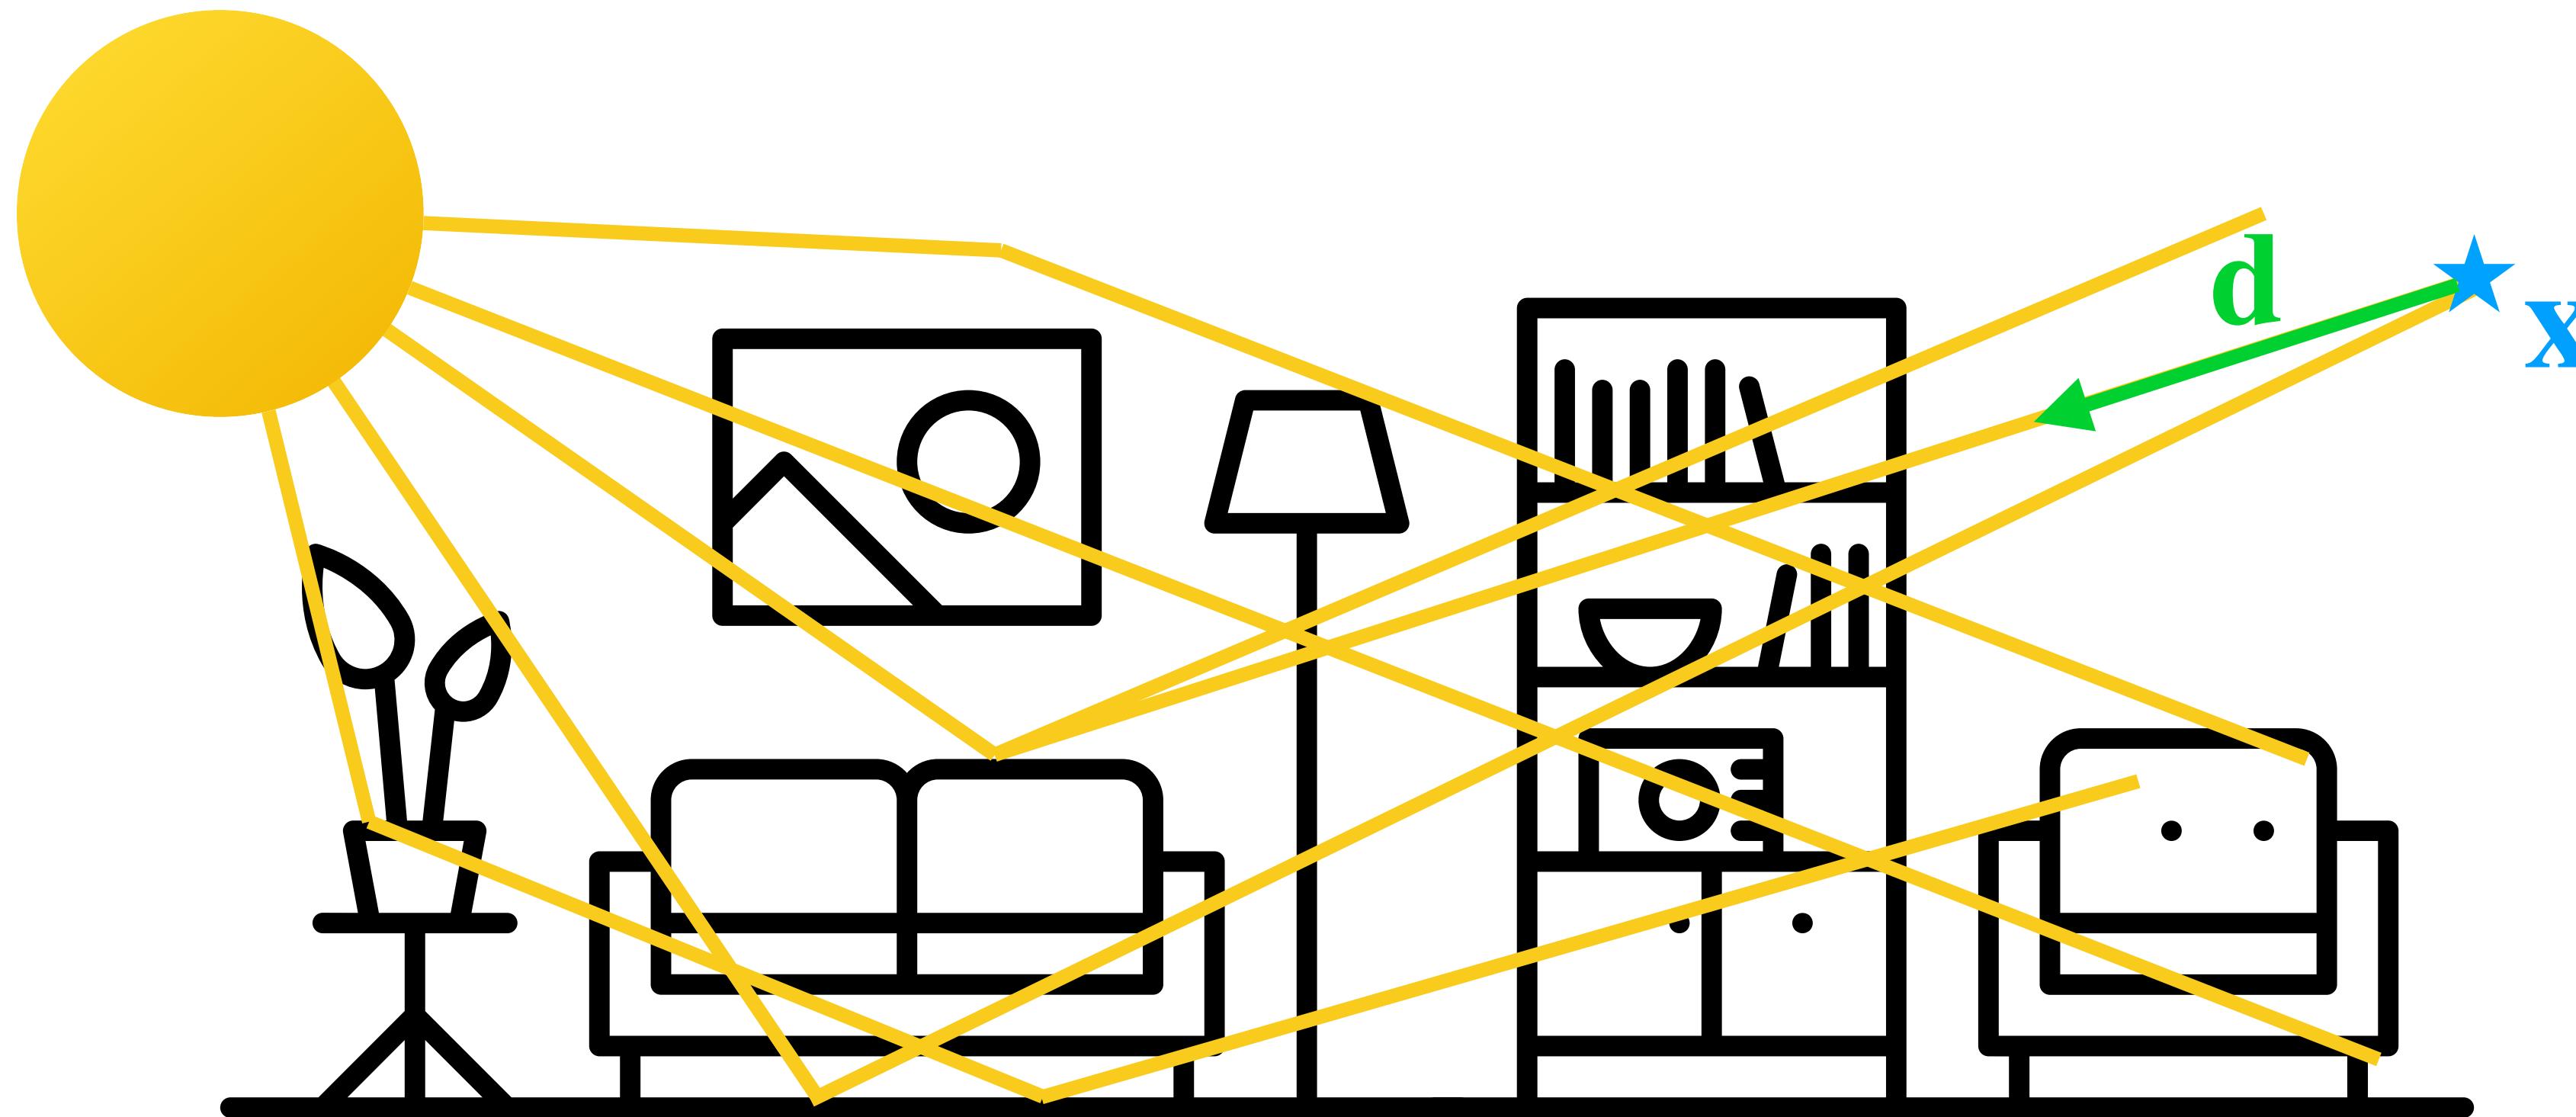
# Light Field Networks



# Light Field Networks



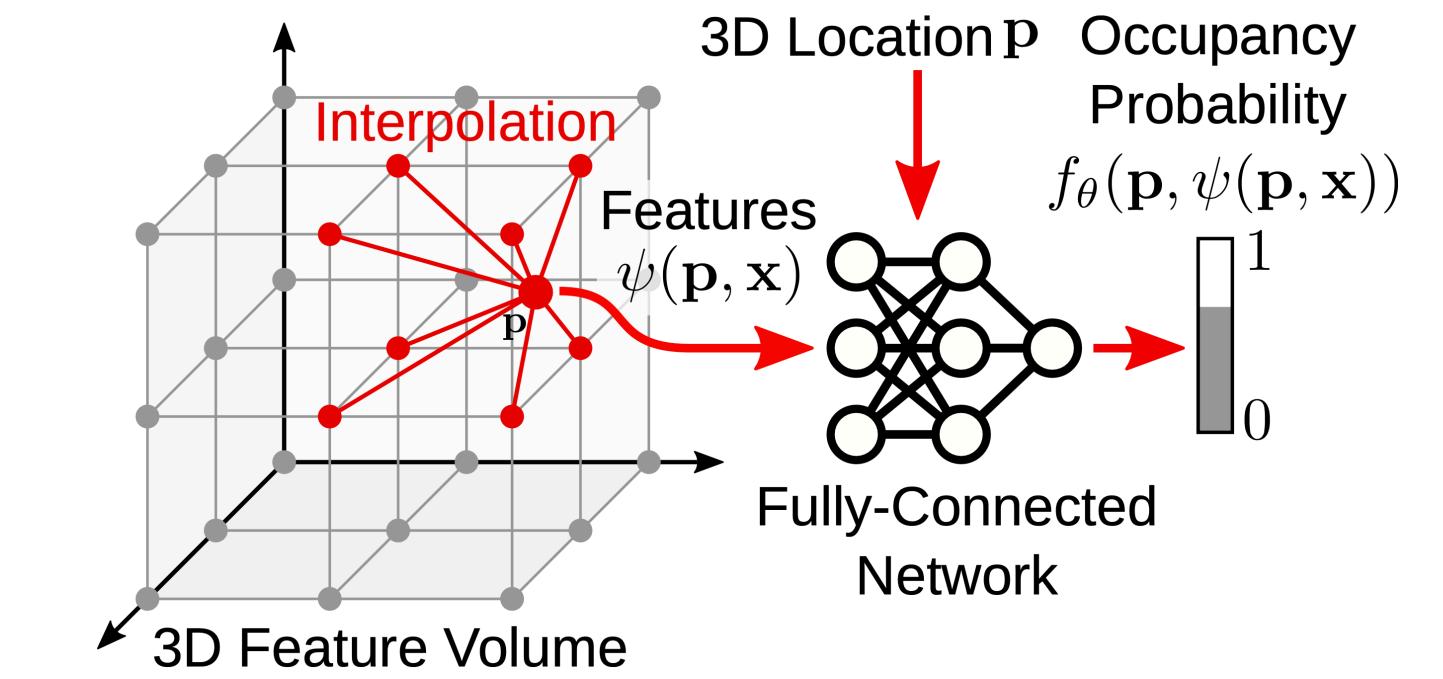
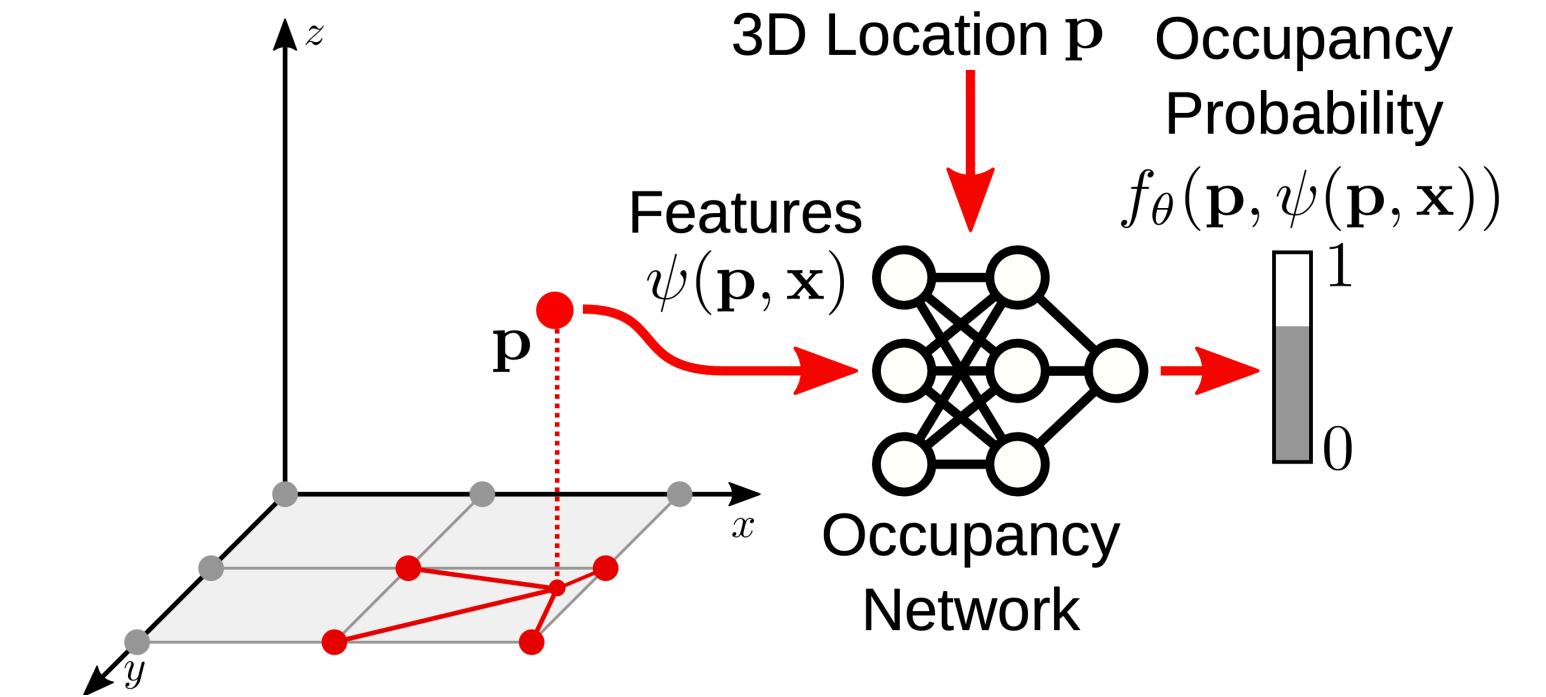
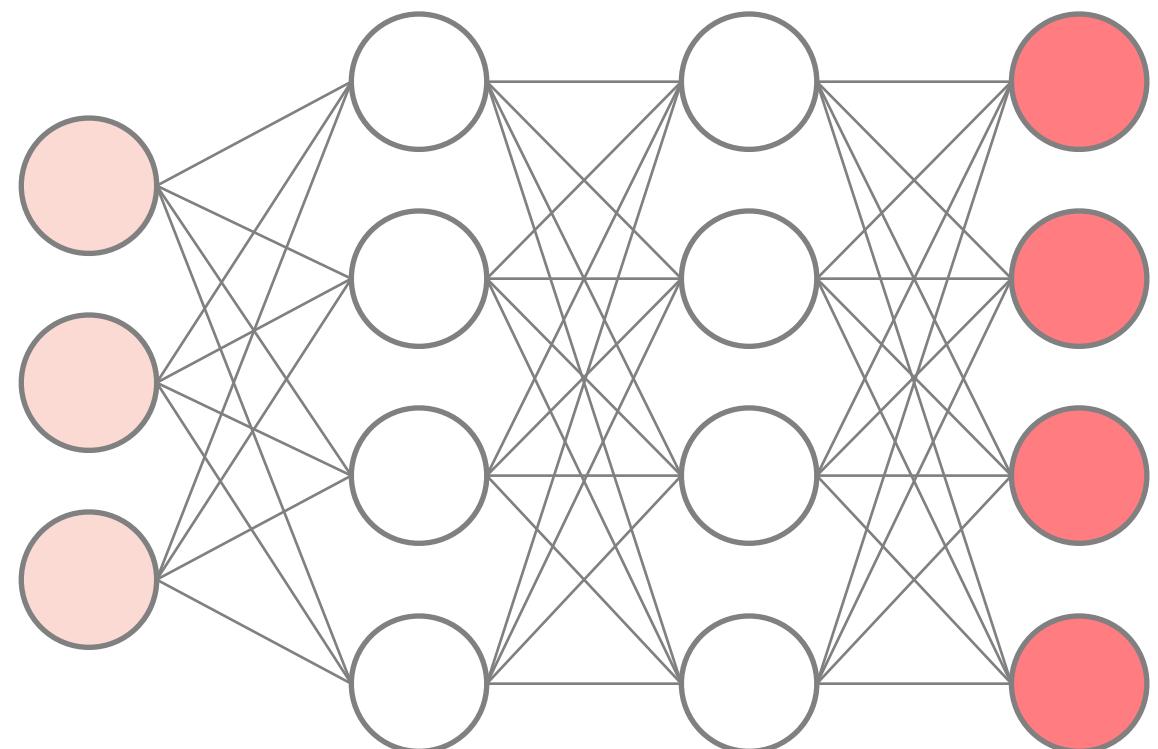
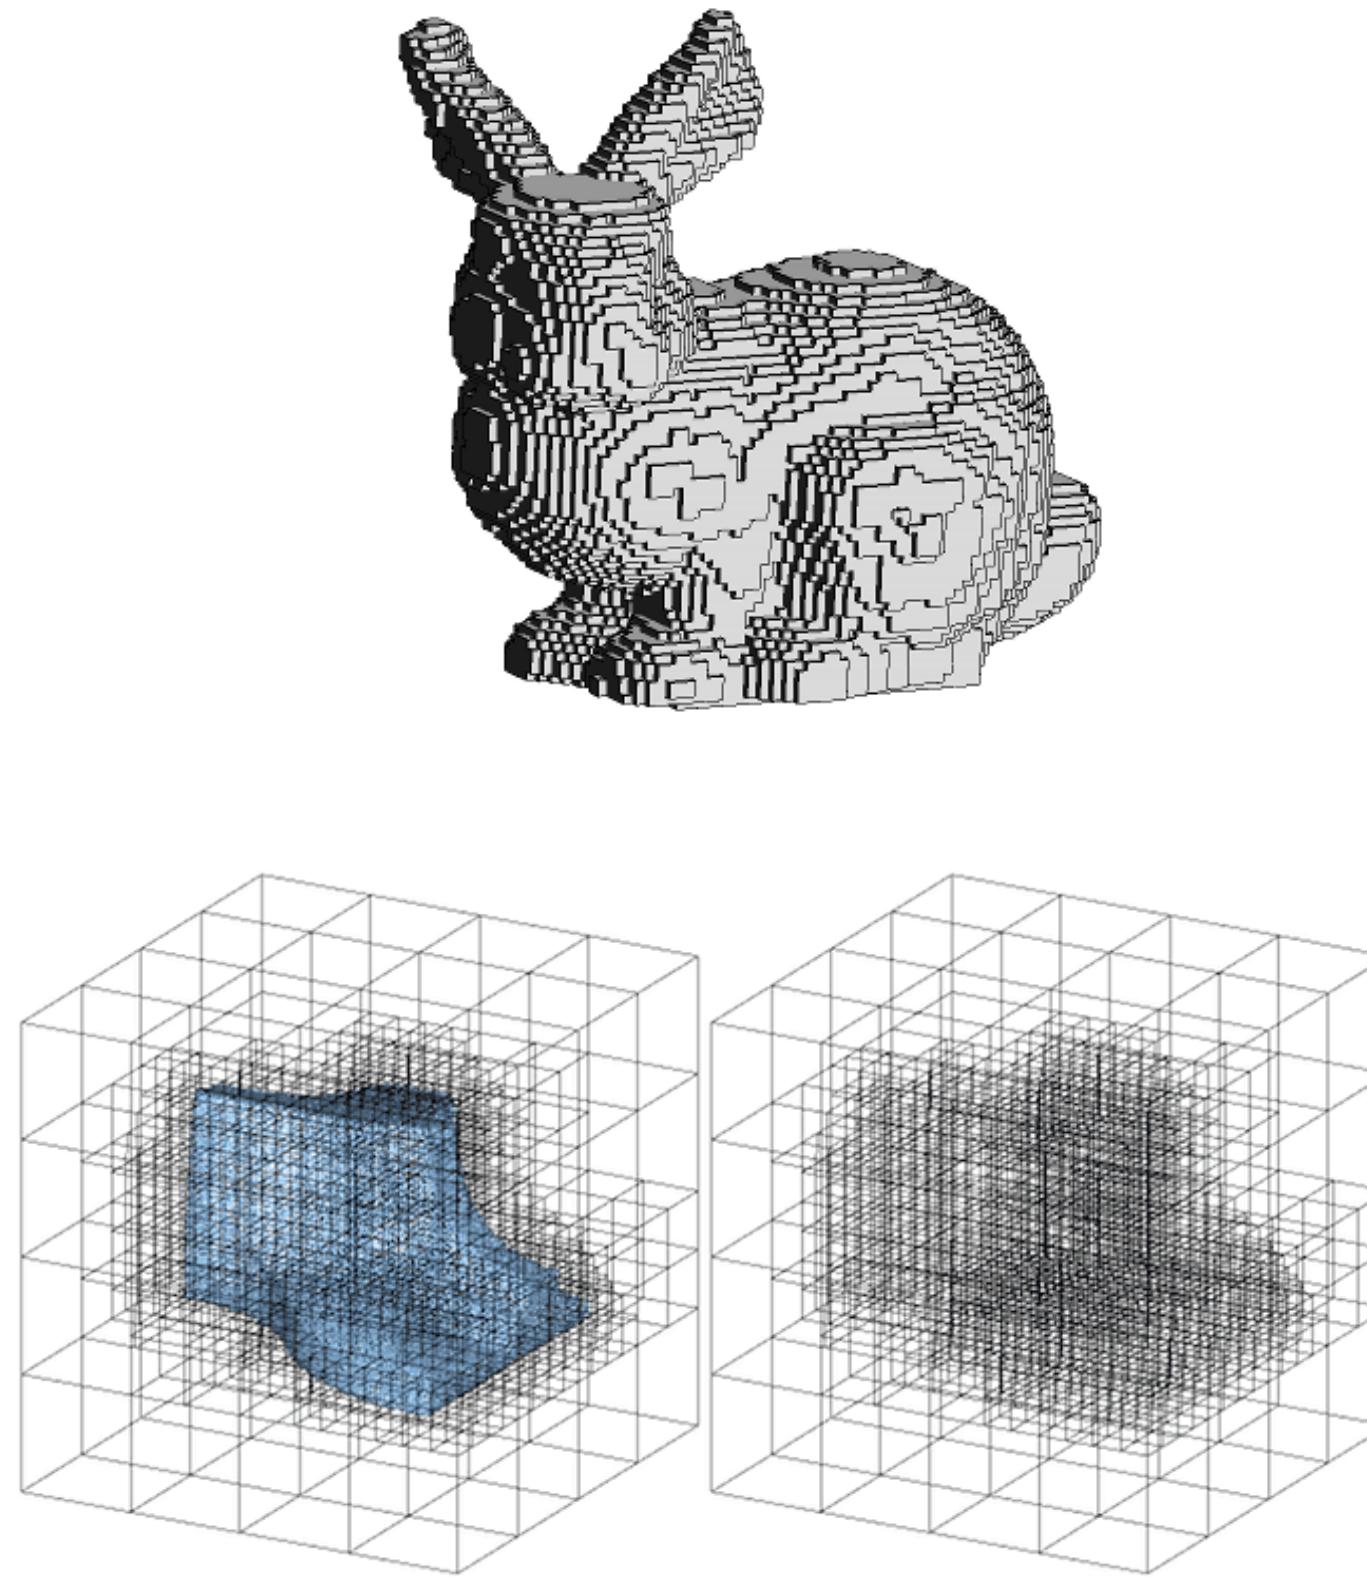
Just because Scene is 3D, don't have to store information in 3D data structure...



$$LF : \mathbb{R}^3 \times \mathbb{S}^2 \rightarrow \mathbb{R}^3, \quad LF(\mathbf{x}, \mathbf{d}) = \mathbf{c}$$

Light Field is 5-dimensional! What parameterizations are possible...?

# Summary



- No “Artificial Intelligence”, “understanding” or “Machine Learning” when fitting a single representation to a single 3D thing, even if there are neural networks in them.
- It becomes *a lot* more interesting when considering *inference*, i.e., we want representation to be output of an encoder! More on that later in “Prior-based reconstruction”.